

# Xenomai Cobalt interface

## 2.99.0

Generated by Doxygen 1.7.1

Thu Jan 5 2012 10:14:34



# Contents

<b>1</b>	<b>Deprecated List</b>	<b>1</b>
<b>2</b>	<b>Module Index</b>	<b>3</b>
2.1	Modules . . . . .	3
<b>3</b>	<b>Data Structure Index</b>	<b>5</b>
3.1	Data Structures . . . . .	5
<b>4</b>	<b>File Index</b>	<b>7</b>
4.1	File List . . . . .	7
<b>5</b>	<b>Module Documentation</b>	<b>11</b>
5.1	Thread state flags. . . . .	11
5.1.1	Detailed Description . . . . .	12
5.1.2	Define Documentation . . . . .	12
5.1.2.1	XNHELD . . . . .	12
5.1.2.2	XNLOCK . . . . .	13
5.1.2.3	XNMIGRATE . . . . .	13
5.1.2.4	XNPEND . . . . .	13
5.1.2.5	XNREADY . . . . .	13
5.1.2.6	XNSUSP . . . . .	13
5.2	Thread information flags. . . . .	13
5.2.1	Detailed Description . . . . .	14
5.3	HAL. . . . .	14
5.3.1	Detailed Description . . . . .	15
5.3.2	Function Documentation . . . . .	15
5.3.2.1	rthal_apc_alloc . . . . .	15
5.3.2.2	rthal_apc_free . . . . .	16
5.4	Thread cancellation. . . . .	16
5.4.1	Detailed Description . . . . .	17

5.4.2	Function Documentation	17
5.4.2.1	pthread_cancel	17
5.4.2.2	pthread_cleanup_pop	18
5.4.2.3	pthread_cleanup_push	18
5.4.2.4	pthread_setcancelstate	19
5.4.2.5	pthread_setcanceltype	19
5.5	Clocks and timers services.	20
5.5.1	Detailed Description	21
5.5.2	Function Documentation	21
5.5.2.1	do_clock_host_realtime	21
5.5.2.2	timer_create	21
5.5.2.3	timer_gettime	22
5.5.2.4	timer_settime	23
5.6	Condition variables services.	23
5.6.1	Detailed Description	24
5.6.2	Function Documentation	24
5.6.2.1	pthread_cond_destroy	24
5.6.2.2	pthread_cond_init	25
5.6.2.3	pthread_condattr_destroy	26
5.6.2.4	pthread_condattr_getclock	26
5.6.2.5	pthread_condattr_getpshared	26
5.6.2.6	pthread_condattr_init	27
5.6.2.7	pthread_condattr_setclock	27
5.6.2.8	pthread_condattr_setpshared	28
5.7	POSIX skin.	28
5.8	Message queues services.	29
5.8.1	Detailed Description	29
5.8.2	Function Documentation	29
5.8.2.1	mq_close	29
5.8.2.2	mq_getattr	30
5.8.2.3	mq_open	31
5.8.2.4	mq_setattr	32
5.8.2.5	mq_unlink	32
5.9	Mutex services.	33
5.9.1	Detailed Description	34
5.9.2	Function Documentation	34

5.9.2.1	<a href="#">pthread_mutexattr_destroy</a> . . . . .	34
5.9.2.2	<a href="#">pthread_mutexattr_getprotocol</a> . . . . .	34
5.9.2.3	<a href="#">pthread_mutexattr_getpshared</a> . . . . .	35
5.9.2.4	<a href="#">pthread_mutexattr_gettype</a> . . . . .	36
5.9.2.5	<a href="#">pthread_mutexattr_init</a> . . . . .	36
5.9.2.6	<a href="#">pthread_mutexattr_setprotocol</a> . . . . .	37
5.9.2.7	<a href="#">pthread_mutexattr_setpshared</a> . . . . .	37
5.9.2.8	<a href="#">pthread_mutexattr_settype</a> . . . . .	38
5.10	Buffer descriptors. . . . .	38
5.10.1	Detailed Description . . . . .	39
5.10.2	Function Documentation . . . . .	41
5.10.2.1	<a href="#">xnbufd_copy_from_kmem</a> . . . . .	41
5.10.2.2	<a href="#">xnbufd_copy_to_kmem</a> . . . . .	42
5.10.2.3	<a href="#">xnbufd_invalidate</a> . . . . .	43
5.10.2.4	<a href="#">xnbufd_map_kread</a> . . . . .	43
5.10.2.5	<a href="#">xnbufd_map_kwrite</a> . . . . .	44
5.10.2.6	<a href="#">xnbufd_map_uread</a> . . . . .	44
5.10.2.7	<a href="#">xnbufd_map_uwrite</a> . . . . .	45
5.10.2.8	<a href="#">xnbufd_reset</a> . . . . .	45
5.10.2.9	<a href="#">xnbufd_unmap_kread</a> . . . . .	46
5.10.2.10	<a href="#">xnbufd_unmap_kwrite</a> . . . . .	46
5.10.2.11	<a href="#">xnbufd_unmap_uread</a> . . . . .	47
5.10.2.12	<a href="#">xnbufd_unmap_uwrite</a> . . . . .	47
5.11	System clock services. . . . .	48
5.11.1	Function Documentation . . . . .	48
5.11.1.1	<a href="#">xnclock_adjust</a> . . . . .	48
5.12	Debugging services. . . . .	49
5.13	Dynamic memory allocation services. . . . .	49
5.13.1	Detailed Description . . . . .	50
5.13.2	Function Documentation . . . . .	50
5.13.2.1	<a href="#">xnheap_alloc</a> . . . . .	50
5.13.2.2	<a href="#">xnheap_destroy</a> . . . . .	51
5.13.2.3	<a href="#">xnheap_extend</a> . . . . .	51
5.13.2.4	<a href="#">xnheap_free</a> . . . . .	52
5.13.2.5	<a href="#">xnheap_init</a> . . . . .	53
5.13.2.6	<a href="#">xnheap_schedule_free</a> . . . . .	54

5.13.2.7	<code>xnheap_set_label</code>	54
5.13.2.8	<code>xnheap_test_and_free</code>	55
5.14	Interrupt management	55
5.14.1	Detailed Description	56
5.14.2	Function Documentation	56
5.14.2.1	<code>xnintr_affinity</code>	56
5.14.2.2	<code>xnintr_attach</code>	56
5.14.2.3	<code>xnintr_destroy</code>	57
5.14.2.4	<code>xnintr_detach</code>	58
5.14.2.5	<code>xnintr_disable</code>	58
5.14.2.6	<code>xnintr_enable</code>	59
5.14.2.7	<code>xnintr_init</code>	59
5.15	Lightweight key-to-object mapping service	61
5.15.1	Detailed Description	61
5.15.2	Function Documentation	62
5.15.2.1	<code>xnmap_create</code>	62
5.15.2.2	<code>xnmap_delete</code>	62
5.15.2.3	<code>xnmap_enter</code>	63
5.15.2.4	<code>xnmap_fetch</code>	63
5.15.2.5	<code>xnmap_fetch_nocheck</code>	64
5.15.2.6	<code>xnmap_remove</code>	65
5.16	Real-time pod services	65
5.16.1	Detailed Description	67
5.16.2	Function Documentation	67
5.16.2.1	<code>__xnpod_reset_thread</code>	67
5.16.2.2	<code>xnpod_abort_thread</code>	68
5.16.2.3	<code>xnpod_add_hook</code>	68
5.16.2.4	<code>xnpod_delete_thread</code>	69
5.16.2.5	<code>xnpod_disable_timesource</code>	70
5.16.2.6	<code>xnpod_dispatch_signals</code>	70
5.16.2.7	<code>xnpod_enable_timesource</code>	70
5.16.2.8	<code>xnpod_handle_exception</code>	71
5.16.2.9	<code>xnpod_init</code>	71
5.16.2.10	<code>xnpod_init_thread</code>	71
5.16.2.11	<code>xnpod_migrate_thread</code>	73
5.16.2.12	<code>xnpod_remove_hook</code>	73

5.16.2.13	xnpod_resume_thread	74
5.16.2.14	xnpod_schedule	75
5.16.2.15	xnpod_set_thread_mode	76
5.16.2.16	xnpod_set_thread_periodic	77
5.16.2.17	xnpod_set_thread_schedparam	78
5.16.2.18	xnpod_set_thread_tslice	79
5.16.2.19	xnpod_shutdown	79
5.16.2.20	xnpod_start_thread	80
5.16.2.21	xnpod_stop_thread	81
5.16.2.22	xnpod_suspend_thread	82
5.16.2.23	xnpod_unblock_thread	83
5.16.2.24	xnpod_wait_thread_period	84
5.16.2.25	xnpod_welcome_thread	84
5.17	Registry services.	85
5.17.1	Detailed Description	85
5.17.2	Function Documentation	86
5.17.2.1	xnregistry_bind	86
5.17.2.2	xnregistry_enter	87
5.17.2.3	xnregistry_fetch	88
5.17.2.4	xnregistry_get	88
5.17.2.5	xnregistry_put	89
5.17.2.6	xnregistry_remove	89
5.17.2.7	xnregistry_remove_safe	90
5.18	File descriptors events multiplexing services.	91
5.18.1	Detailed Description	91
5.18.2	Function Documentation	92
5.18.2.1	xnselect	92
5.18.2.2	xnselect_bind	92
5.18.2.3	xnselect_destroy	93
5.18.2.4	xnselect_init	93
5.18.2.5	xnselect_signal	93
5.18.2.6	xnselector_destroy	94
5.18.2.7	xnselector_init	94
5.19	Real-time shadow services.	94
5.19.1	Detailed Description	95
5.19.2	Function Documentation	95

5.19.2.1	xnshadow_harden	95
5.19.2.2	xnshadow_map	95
5.19.2.3	xnshadow_ppd_get	97
5.19.2.4	xnshadow_relax	97
5.20	Thread synchronization services.	98
5.20.1	Detailed Description	99
5.20.2	Function Documentation	99
5.20.2.1	xnsynch_acquire	99
5.20.2.2	xnsynch_clear_boost	99
5.20.2.3	xnsynch_flush	100
5.20.2.4	xnsynch_forget_sleeper	101
5.20.2.5	xnsynch_init	101
5.20.2.6	xnsynch_peek_pendq	102
5.20.2.7	xnsynch_release	103
5.20.2.8	xnsynch_release_all_ownerships	104
5.20.2.9	xnsynch_requeue_sleeper	104
5.20.2.10	xnsynch_sleep_on	104
5.20.2.11	xnsynch_wakeup_one_sleeper	105
5.20.2.12	xnsynch_wakeup_this_sleeper	106
5.21	Timer services.	107
5.21.1	Detailed Description	107
5.21.2	Function Documentation	108
5.21.2.1	xntimer_destroy	108
5.21.2.2	xntimer_freeze	108
5.21.2.3	xntimer_get_date	109
5.21.2.4	xntimer_get_interval	109
5.21.2.5	xntimer_get_overruns	109
5.21.2.6	xntimer_get_timeout	110
5.21.2.7	xntimer_init	110
5.21.2.8	xntimer_start	111
5.21.2.9	xntimer_stop	111
5.21.2.10	xntimer_tick	112
5.22	Virtual file services	112
5.22.1	Detailed Description	114
5.22.2	Function Documentation	114
5.22.2.1	xnvfile_destroy	114



5.22.2.2	<a href="#">xnvmfile_get_blob</a>	115
5.22.2.3	<a href="#">xnvmfile_get_integer</a>	115
5.22.2.4	<a href="#">xnvmfile_get_string</a>	116
5.22.2.5	<a href="#">xnvmfile_init_dir</a>	116
5.22.2.6	<a href="#">xnvmfile_init_link</a>	116
5.22.2.7	<a href="#">xnvmfile_init_regular</a>	117
5.22.2.8	<a href="#">xnvmfile_init_snapshot</a>	117
5.22.3	<a href="#">Variable Documentation</a>	118
5.22.3.1	<a href="#">nkvfroot</a>	118
5.22.3.2	<a href="#">nkvfroot</a>	118
5.23	<a href="#">Inter-Driver API</a>	119
5.23.1	<a href="#">Function Documentation</a>	121
5.23.1.1	<a href="#">rtdm_accept</a>	121
5.23.1.2	<a href="#">rtdm_bind</a>	121
5.23.1.3	<a href="#">rtdm_close</a>	121
5.23.1.4	<a href="#">rtdm_connect</a>	121
5.23.1.5	<a href="#">rtdm_context_get</a>	121
5.23.1.6	<a href="#">rtdm_context_lock</a>	122
5.23.1.7	<a href="#">rtdm_context_put</a>	123
5.23.1.8	<a href="#">rtdm_context_unlock</a>	123
5.23.1.9	<a href="#">rtdm_getpeername</a>	124
5.23.1.10	<a href="#">rtdm_getsockname</a>	124
5.23.1.11	<a href="#">rtdm_getsockopt</a>	124
5.23.1.12	<a href="#">rtdm_ioctl</a>	124
5.23.1.13	<a href="#">rtdm_listen</a>	124
5.23.1.14	<a href="#">rtdm_open</a>	125
5.23.1.15	<a href="#">rtdm_read</a>	125
5.23.1.16	<a href="#">rtdm_recv</a>	125
5.23.1.17	<a href="#">rtdm_recvfrom</a>	125
5.23.1.18	<a href="#">rtdm_recvmsg</a>	125
5.23.1.19	<a href="#">rtdm_select_bind</a>	126
5.23.1.20	<a href="#">rtdm_send</a>	126
5.23.1.21	<a href="#">rtdm_sendmsg</a>	127
5.23.1.22	<a href="#">rtdm_sendto</a>	127
5.23.1.23	<a href="#">rtdm_setsockopt</a>	127
5.23.1.24	<a href="#">rtdm_shutdown</a>	127

5.23.1.25	<code>rtdm_socket</code>	127
5.23.1.26	<code>rtdm_write</code>	128
5.24	Device Registration Services	128
5.24.1	Define Documentation	131
5.24.1.1	<code>RTDM_CLOSING</code>	131
5.24.1.2	<code>RTDM_CREATED_IN_NRT</code>	131
5.24.1.3	<code>RTDM_DEVICE_TYPE_MASK</code>	131
5.24.1.4	<code>RTDM_EXCLUSIVE</code>	131
5.24.1.5	<code>RTDM_NAMED_DEVICE</code>	131
5.24.1.6	<code>RTDM_PROTOCOL_DEVICE</code>	131
5.24.2	Typedef Documentation	131
5.24.2.1	<code>rtdm_close_handler_t</code>	131
5.24.2.2	<code>rtdm_ioctl_handler_t</code>	132
5.24.2.3	<code>rtdm_open_handler_t</code>	132
5.24.2.4	<code>rtdm_read_handler_t</code>	133
5.24.2.5	<code>rtdm_recvmmsg_handler_t</code>	133
5.24.2.6	<code>rtdm_select_bind_handler_t</code>	134
5.24.2.7	<code>rtdm_sendmsg_handler_t</code>	134
5.24.2.8	<code>rtdm_socket_handler_t</code>	135
5.24.2.9	<code>rtdm_write_handler_t</code>	135
5.24.3	Function Documentation	135
5.24.3.1	<code>rtdm_context_to_private</code>	135
5.24.3.2	<code>rtdm_dev_register</code>	136
5.24.3.3	<code>rtdm_dev_unregister</code>	136
5.24.3.4	<code>rtdm_private_to_context</code>	137
5.25	Driver Development API	138
5.25.1	Detailed Description	139
5.26	Clock Services	139
5.26.1	Function Documentation	139
5.26.1.1	<code>rtdm_clock_read</code>	139
5.26.1.2	<code>rtdm_clock_read_monotonic</code>	140
5.27	Task Services	140
5.27.1	Typedef Documentation	142
5.27.1.1	<code>rtdm_task_proc_t</code>	142
5.27.2	Function Documentation	142
5.27.2.1	<code>rtdm_task_busy_sleep</code>	142

5.27.2.2	<a href="#">rt dm_task_current</a>	143
5.27.2.3	<a href="#">rt dm_task_destroy</a>	143
5.27.2.4	<a href="#">rt dm_task_init</a>	143
5.27.2.5	<a href="#">rt dm_task_join_nrt</a>	144
5.27.2.6	<a href="#">rt dm_task_set_period</a>	145
5.27.2.7	<a href="#">rt dm_task_set_priority</a>	145
5.27.2.8	<a href="#">rt dm_task_sleep</a>	145
5.27.2.9	<a href="#">rt dm_task_sleep_abs</a>	146
5.27.2.10	<a href="#">rt dm_task_sleep_until</a>	147
5.27.2.11	<a href="#">rt dm_task_unblock</a>	147
5.27.2.12	<a href="#">rt dm_task_wait_period</a>	148
5.28	<a href="#">Timer Services</a>	148
5.28.1	<a href="#">Typedef Documentation</a>	149
5.28.1.1	<a href="#">rt dm_timer_handler_t</a>	149
5.28.2	<a href="#">Enumeration Type Documentation</a>	149
5.28.2.1	<a href="#">rt dm_timer_mode</a>	149
5.28.3	<a href="#">Function Documentation</a>	150
5.28.3.1	<a href="#">rt dm_timer_destroy</a>	150
5.28.3.2	<a href="#">rt dm_timer_init</a>	150
5.28.3.3	<a href="#">rt dm_timer_start</a>	151
5.28.3.4	<a href="#">rt dm_timer_start_in_handler</a>	151
5.28.3.5	<a href="#">rt dm_timer_stop</a>	152
5.28.3.6	<a href="#">rt dm_timer_stop_in_handler</a>	152
5.29	<a href="#">Synchronisation Services</a>	153
5.29.1	<a href="#">Define Documentation</a>	156
5.29.1.1	<a href="#">RTDM_EXECUTE_ATOMICALLY</a>	156
5.29.1.2	<a href="#">rt dm_lock_get</a>	156
5.29.1.3	<a href="#">rt dm_lock_get_irqsave</a>	157
5.29.1.4	<a href="#">rt dm_lock_init</a>	157
5.29.1.5	<a href="#">rt dm_lock_irqrestore</a>	158
5.29.1.6	<a href="#">rt dm_lock_irqsave</a>	158
5.29.1.7	<a href="#">rt dm_lock_put</a>	158
5.29.1.8	<a href="#">rt dm_lock_put_irqrestore</a>	159
5.29.2	<a href="#">Enumeration Type Documentation</a>	159
5.29.2.1	<a href="#">rt dm_selecttype</a>	159
5.29.3	<a href="#">Function Documentation</a>	159

5.29.3.1	<code>rtdm_event_clear</code>	159
5.29.3.2	<code>rtdm_event_destroy</code>	160
5.29.3.3	<code>rtdm_event_init</code>	160
5.29.3.4	<code>rtdm_event_pulse</code>	161
5.29.3.5	<code>rtdm_event_select_bind</code>	161
5.29.3.6	<code>rtdm_event_signal</code>	162
5.29.3.7	<code>rtdm_event_timedwait</code>	162
5.29.3.8	<code>rtdm_event_wait</code>	163
5.29.3.9	<code>rtdm_mutex_destroy</code>	164
5.29.3.10	<code>rtdm_mutex_init</code>	164
5.29.3.11	<code>rtdm_mutex_lock</code>	164
5.29.3.12	<code>rtdm_mutex_timedlock</code>	165
5.29.3.13	<code>rtdm_mutex_unlock</code>	166
5.29.3.14	<code>rtdm_select_bind</code>	166
5.29.3.15	<code>rtdm_sem_destroy</code>	167
5.29.3.16	<code>rtdm_sem_down</code>	167
5.29.3.17	<code>rtdm_sem_init</code>	168
5.29.3.18	<code>rtdm_sem_select_bind</code>	168
5.29.3.19	<code>rtdm_sem_timeddown</code>	169
5.29.3.20	<code>rtdm_sem_up</code>	170
5.29.3.21	<code>rtdm_toseq_init</code>	170
5.30	Interrupt Management Services	171
5.30.1	Define Documentation	172
5.30.1.1	<code>rtdm_irq_get_arg</code>	172
5.30.2	Typedef Documentation	173
5.30.2.1	<code>rtdm_irq_handler_t</code>	173
5.30.3	Function Documentation	173
5.30.3.1	<code>rtdm_irq_disable</code>	173
5.30.3.2	<code>rtdm_irq_enable</code>	174
5.30.3.3	<code>rtdm_irq_free</code>	174
5.30.3.4	<code>rtdm_irq_request</code>	175
5.31	Non-Real-Time Signalling Services	176
5.31.1	Detailed Description	176
5.31.2	Typedef Documentation	176
5.31.2.1	<code>rtdm_nrtsig_handler_t</code>	176
5.31.3	Function Documentation	177

5.31.3.1	<a href="#">rt dm_nrtsig_destroy</a>	177
5.31.3.2	<a href="#">rt dm_nrtsig_init</a>	177
5.31.3.3	<a href="#">rt dm_nrtsig_pend</a>	178
5.32	Utility Services	178
5.32.1	Function Documentation	179
5.32.1.1	<a href="#">rt dm_copy_from_user</a>	179
5.32.1.2	<a href="#">rt dm_copy_to_user</a>	180
5.32.1.3	<a href="#">rt dm_free</a>	181
5.32.1.4	<a href="#">rt dm_in_rt_context</a>	181
5.32.1.5	<a href="#">rt dm_iomap_to_user</a>	182
5.32.1.6	<a href="#">rt dm_malloc</a>	183
5.32.1.7	<a href="#">rt dm_mmap_to_user</a>	183
5.32.1.8	<a href="#">rt dm_munmap</a>	184
5.32.1.9	<a href="#">rt dm_printk</a>	185
5.32.1.10	<a href="#">rt dm_read_user_ok</a>	185
5.32.1.11	<a href="#">rt dm_rt_capable</a>	186
5.32.1.12	<a href="#">rt dm_rw_user_ok</a>	186
5.32.1.13	<a href="#">rt dm_safe_copy_from_user</a>	187
5.32.1.14	<a href="#">rt dm_safe_copy_to_user</a>	187
5.32.1.15	<a href="#">rt dm_strncpy_from_user</a>	188
5.33	Device Profiles	189
5.33.1	Detailed Description	190
5.33.2	Define Documentation	190
5.33.2.1	<a href="#">RTIOC_DEVICE_INFO</a>	190
5.33.2.2	<a href="#">RTIOC_PURGE</a>	191
5.34	Semaphores services.	191
5.34.1	Detailed Description	191
5.34.2	Function Documentation	192
5.34.2.1	<a href="#">sem_close</a>	192
5.34.2.2	<a href="#">sem_destroy</a>	192
5.34.2.3	<a href="#">sem_getvalue</a>	193
5.34.2.4	<a href="#">sem_open</a>	193
5.34.2.5	<a href="#">sem_post</a>	194
5.34.2.6	<a href="#">sem_timedwait</a>	194
5.34.2.7	<a href="#">sem_trywait</a>	195
5.34.2.8	<a href="#">sem_unlink</a>	196

5.34.2.9	<code>sem_wait</code>	196
5.35	Threads management services.	197
5.35.1	Detailed Description	198
5.35.2	Function Documentation	198
5.35.2.1	<code>pthread_create</code>	198
5.35.2.2	<code>pthread_getschedparam</code>	199
5.35.2.3	<code>pthread_getschedparam_ex</code>	199
5.35.2.4	<code>pthread_make_periodic_np</code>	200
5.35.2.5	<code>pthread_set_mode_np</code>	201
5.35.2.6	<code>pthread_set_name_np</code>	201
5.35.2.7	<code>pthread_setschedparam</code>	202
5.35.2.8	<code>pthread_setschedparam_ex</code>	203
5.36	CAN Devices	204
5.36.1	Detailed Description	213
5.36.2	Define Documentation	215
5.36.2.1	<code>CAN_CTRLMODE_LISTENONLY</code>	215
5.36.2.2	<code>CAN_CTRLMODE_LOOPBACK</code>	215
5.36.2.3	<code>CAN_ERR_LOSTARB_UNSPEC</code>	216
5.36.2.4	<code>CAN_RAW_ERR_FILTER</code>	216
5.36.2.5	<code>CAN_RAW_FILTER</code>	216
5.36.2.6	<code>CAN_RAW_LOOPBACK</code>	217
5.36.2.7	<code>CAN_RAW_RECV_OWN_MSGS</code>	217
5.36.2.8	<code>RTCAN_RTIOC_RCV_TIMEOUT</code>	217
5.36.2.9	<code>RTCAN_RTIOC_SND_TIMEOUT</code>	218
5.36.2.10	<code>RTCAN_RTIOC_TAKE_TIMESTAMP</code>	219
5.36.2.11	<code>SIOCGCANBAUDRATE</code>	219
5.36.2.12	<code>SIOCGCANCTRLMODE</code>	220
5.36.2.13	<code>SIOCGCANCUSTOMBITTIME</code>	220
5.36.2.14	<code>SIOCGCANSTATE</code>	221
5.36.2.15	<code>SIOCGIFINDEX</code>	221
5.36.2.16	<code>SIOCSCANBAUDRATE</code>	222
5.36.2.17	<code>SIOCSCANCTRLMODE</code>	223
5.36.2.18	<code>SIOCSCANCUSTOMBITTIME</code>	223
5.36.2.19	<code>SIOCSCANMODE</code>	224
5.36.2.20	<code>SOL_CAN_RAW</code>	225
5.36.3	Typedef Documentation	225

5.36.3.1	<code>can_filter_t</code>	225
5.36.3.2	<code>can_frame_t</code>	225
5.36.4	Enumeration Type Documentation	225
5.36.4.1	<code>CAN_BITTIME_TYPE</code>	225
5.36.4.2	<code>CAN_MODE</code>	226
5.36.4.3	<code>CAN_STATE</code>	226
5.37	Real-Time Driver Model	226
5.37.1	Detailed Description	227
5.37.2	Define Documentation	227
5.37.2.1	<code>RTDM_TIMEOUT_INFINITE</code>	227
5.37.2.2	<code>RTDM_TIMEOUT_NONE</code>	228
5.37.3	Typedef Documentation	228
5.37.3.1	<code>nanosecs_abs_t</code>	228
5.37.3.2	<code>nanosecs_rel_t</code>	228
5.38	User API	228
5.38.1	Detailed Description	230
5.38.2	Function Documentation	230
5.38.2.1	<code>rt_dev_accept</code>	230
5.38.2.2	<code>rt_dev_bind</code>	230
5.38.2.3	<code>rt_dev_close</code>	231
5.38.2.4	<code>rt_dev_connect</code>	231
5.38.2.5	<code>rt_dev_getpeername</code>	232
5.38.2.6	<code>rt_dev_getsockname</code>	232
5.38.2.7	<code>rt_dev_getsockopt</code>	233
5.38.2.8	<code>rt_dev_ioctl</code>	233
5.38.2.9	<code>rt_dev_listen</code>	234
5.38.2.10	<code>rt_dev_open</code>	234
5.38.2.11	<code>rt_dev_read</code>	234
5.38.2.12	<code>rt_dev_recv</code>	235
5.38.2.13	<code>rt_dev_recvfrom</code>	235
5.38.2.14	<code>rt_dev_recvmsg</code>	236
5.38.2.15	<code>rt_dev_send</code>	236
5.38.2.16	<code>rt_dev_sendmsg</code>	237
5.38.2.17	<code>rt_dev_sendto</code>	237
5.38.2.18	<code>rt_dev_setsockopt</code>	238
5.38.2.19	<code>rt_dev_shutdown</code>	238

5.38.2.20	rt_dev_socket	239
5.38.2.21	rt_dev_write	239
5.39	Real-time IPC protocols	240
5.39.1	Detailed Description	243
5.39.2	Define Documentation	243
5.39.2.1	BUFP_BUFSZ	243
5.39.2.2	BUFP_LABEL	243
5.39.2.3	IDDP_LABEL	244
5.39.2.4	IDDP_POOLSZ	245
5.39.2.5	SO_RCVTIMEO	245
5.39.2.6	SO_SNDTIMEO	246
5.39.2.7	XDDP_BUFSZ	246
5.39.2.8	XDDP_EVTDOWN	247
5.39.2.9	XDDP_EVTIN	247
5.39.2.10	XDDP_EVTNOBUF	247
5.39.2.11	XDDP_EVTOUT	247
5.39.2.12	XDDP_LABEL	247
5.39.2.13	XDDP_MONITOR	248
5.39.2.14	XDDP_POOLSZ	249
5.39.3	Enumeration Type Documentation	250
5.39.3.1	"@14	250
5.39.4	Function Documentation	250
5.39.4.1	bind__AF_RTIPC	250
5.39.4.2	close__AF_RTIPC	252
5.39.4.3	connect__AF_RTIPC	252
5.39.4.4	getpeername__AF_RTIPC	253
5.39.4.5	getsockname__AF_RTIPC	253
5.39.4.6	getsockopt__AF_RTIPC	253
5.39.4.7	recvmsg__AF_RTIPC	254
5.39.4.8	sendmsg__AF_RTIPC	254
5.39.4.9	setsockopt__AF_RTIPC	255
5.39.4.10	socket__AF_RTIPC	255
5.40	Serial Devices	256
5.40.1	Detailed Description	261
5.40.2	Define Documentation	261
5.40.2.1	RTSER_RTIOC_BREAK_CTL	261



5.40.2.2	RTSER_RTIOC_GET_CONFIG . . . . .	262
5.40.2.3	RTSER_RTIOC_GET_CONTROL . . . . .	262
5.40.2.4	RTSER_RTIOC_GET_STATUS . . . . .	263
5.40.2.5	RTSER_RTIOC_SET_CONFIG . . . . .	263
5.40.2.6	RTSER_RTIOC_SET_CONTROL . . . . .	264
5.40.2.7	RTSER_RTIOC_WAIT_EVENT . . . . .	264
5.41	Testing Devices . . . . .	265
5.41.1	Detailed Description . . . . .	266
5.42	Sched . . . . .	267
5.42.1	Function Documentation . . . . .	268
5.42.1.1	xnsched_rotate . . . . .	268
<b>6</b>	<b>Data Structure Documentation</b>	<b>269</b>
6.1	can_bittime Struct Reference . . . . .	269
6.1.1	Detailed Description . . . . .	270
6.2	can_bittime_btr Struct Reference . . . . .	270
6.2.1	Detailed Description . . . . .	270
6.3	can_bittime_std Struct Reference . . . . .	270
6.3.1	Detailed Description . . . . .	271
6.4	can_filter Struct Reference . . . . .	271
6.4.1	Detailed Description . . . . .	271
6.4.2	Field Documentation . . . . .	272
6.4.2.1	can_id . . . . .	272
6.4.2.2	can_mask . . . . .	272
6.5	can_frame Struct Reference . . . . .	272
6.5.1	Detailed Description . . . . .	272
6.5.2	Field Documentation . . . . .	272
6.5.2.1	can_id . . . . .	272
6.6	rtdm_dev_context Struct Reference . . . . .	273
6.6.1	Detailed Description . . . . .	274
6.7	rtdm_device Struct Reference . . . . .	274
6.7.1	Detailed Description . . . . .	276
6.7.2	Field Documentation . . . . .	276
6.7.2.1	open_rt . . . . .	276
6.7.2.2	socket_rt . . . . .	276
6.8	rtdm_device_info Struct Reference . . . . .	276
6.8.1	Detailed Description . . . . .	277

6.9	rtm_operations Struct Reference . . . . .	277
6.9.1	Detailed Description . . . . .	278
6.9.2	Field Documentation . . . . .	278
6.9.2.1	close_rt . . . . .	278
6.10	rtipc_port_label Struct Reference . . . . .	278
6.10.1	Detailed Description . . . . .	279
6.10.2	Field Documentation . . . . .	279
6.10.2.1	label . . . . .	279
6.11	rtser_config Struct Reference . . . . .	279
6.11.1	Detailed Description . . . . .	280
6.12	rtser_event Struct Reference . . . . .	280
6.12.1	Detailed Description . . . . .	280
6.13	rtser_status Struct Reference . . . . .	281
6.13.1	Detailed Description . . . . .	281
6.14	sockaddr_can Struct Reference . . . . .	281
6.14.1	Detailed Description . . . . .	281
6.14.2	Field Documentation . . . . .	281
6.14.2.1	can_ifindex . . . . .	281
6.15	sockaddr_ipc Struct Reference . . . . .	282
6.15.1	Detailed Description . . . . .	282
6.15.2	Field Documentation . . . . .	282
6.15.2.1	sipc_port . . . . .	282
6.16	xnpod Struct Reference . . . . .	282
6.16.1	Detailed Description . . . . .	283
6.16.2	Field Documentation . . . . .	283
6.16.2.1	refcnt . . . . .	283
6.16.2.2	sched . . . . .	283
6.16.2.3	status . . . . .	283
6.16.2.4	tdeleteq . . . . .	283
6.16.2.5	threadq . . . . .	283
6.16.2.6	timerlck . . . . .	284
6.16.2.7	tstartq . . . . .	284
6.16.2.8	tswitchq . . . . .	284
6.17	xnsched Struct Reference . . . . .	284
6.17.1	Detailed Description . . . . .	284
6.17.2	Field Documentation . . . . .	284

6.17.2.1	curr	284
6.17.2.2	htimer	285
6.17.2.3	inesting	285
6.17.2.4	lflags	285
6.17.2.5	rootcb	285
6.17.2.6	rt	285
6.17.2.7	status	285
6.18	xnthread_info Struct Reference	285
6.18.1	Detailed Description	286
6.18.2	Field Documentation	286
6.18.2.1	affinity	286
6.18.2.2	bprio	286
6.18.2.3	cprio	286
6.18.2.4	cpu	287
6.18.2.5	ctxswitches	287
6.18.2.6	exectime	287
6.18.2.7	modeswitches	287
6.18.2.8	name	287
6.18.2.9	pagefaults	287
6.18.2.10	relpoint	287
6.18.2.11	state	287
6.18.2.12	syscalls	287
6.19	xnfile_lock_ops Struct Reference	287
6.19.1	Detailed Description	288
6.19.2	Field Documentation	288
6.19.2.1	get	288
6.19.2.2	put	288
6.20	xnfile_regular_iterator Struct Reference	288
6.20.1	Detailed Description	289
6.20.2	Field Documentation	289
6.20.2.1	pos	289
6.20.2.2	private	289
6.20.2.3	seq	289
6.20.2.4	vfile	289
6.21	xnfile_regular_ops Struct Reference	289
6.21.1	Detailed Description	290

6.21.2	Field Documentation	290
6.21.2.1	begin	290
6.21.2.2	end	290
6.21.2.3	next	291
6.21.2.4	rewind	291
6.21.2.5	show	292
6.21.2.6	store	292
6.22	xnvfile_rev_tag Struct Reference	293
6.22.1	Detailed Description	293
6.22.2	Field Documentation	293
6.22.2.1	rev	293
6.23	xnvfile_snapshot Struct Reference	293
6.23.1	Detailed Description	294
6.24	xnvfile_snapshot_iterator Struct Reference	294
6.24.1	Detailed Description	295
6.24.2	Field Documentation	295
6.24.2.1	databuf	295
6.24.2.2	endfn	295
6.24.2.3	nrdata	295
6.24.2.4	private	295
6.24.2.5	seq	295
6.24.2.6	vfile	295
6.25	xnvfile_snapshot_ops Struct Reference	296
6.25.1	Detailed Description	296
6.25.2	Field Documentation	296
6.25.2.1	begin	296
6.25.2.2	end	297
6.25.2.3	next	297
6.25.2.4	rewind	297
6.25.2.5	show	298
6.25.2.6	store	299
<b>7</b>	<b>File Documentation</b>	<b>301</b>
7.1	include/cobalt/nucleus/bufd.h File Reference	301
7.1.1	Detailed Description	302
7.2	include/cobalt/nucleus/clock.h File Reference	303
7.2.1	Detailed Description	303

7.3	<a href="#">include/cobalt/nucleus/hostrt.h File Reference</a>	304
7.3.1	<a href="#">Detailed Description</a>	305
7.4	<a href="#">include/cobalt/nucleus/map.h File Reference</a>	306
7.4.1	<a href="#">Detailed Description</a>	307
7.5	<a href="#">include/cobalt/nucleus/pod.h File Reference</a>	307
7.5.1	<a href="#">Detailed Description</a>	309
7.6	<a href="#">include/cobalt/nucleus/registry.h File Reference</a>	310
7.6.1	<a href="#">Detailed Description</a>	311
7.7	<a href="#">include/cobalt/nucleus/sched-idle.h File Reference</a>	311
7.7.1	<a href="#">Detailed Description</a>	312
7.8	<a href="#">include/cobalt/nucleus/sched-rt.h File Reference</a>	312
7.8.1	<a href="#">Detailed Description</a>	312
7.9	<a href="#">include/cobalt/nucleus/sched-sporadic.h File Reference</a>	313
7.9.1	<a href="#">Detailed Description</a>	313
7.10	<a href="#">include/cobalt/nucleus/sched-tp.h File Reference</a>	313
7.10.1	<a href="#">Detailed Description</a>	314
7.11	<a href="#">include/cobalt/nucleus/sched.h File Reference</a>	314
7.11.1	<a href="#">Detailed Description</a>	315
7.12	<a href="#">include/cobalt/nucleus/select.h File Reference</a>	316
7.12.1	<a href="#">Detailed Description</a>	317
7.13	<a href="#">include/cobalt/nucleus/timer.h File Reference</a>	318
7.13.1	<a href="#">Detailed Description</a>	319
7.14	<a href="#">include/cobalt/nucleus/vdso.h File Reference</a>	319
7.14.1	<a href="#">Detailed Description</a>	320
7.15	<a href="#">include/cobalt/nucleus/vfile.h File Reference</a>	320
7.15.1	<a href="#">Detailed Description</a>	322
7.16	<a href="#">include/rtdm/rtdm.h File Reference</a>	323
7.16.1	<a href="#">Detailed Description</a>	331
7.17	<a href="#">include/rtdm/rtdm_driver.h File Reference</a>	331
7.17.1	<a href="#">Detailed Description</a>	333
7.18	<a href="#">include/rtdm/rtdm_driver.h File Reference</a>	334
7.18.1	<a href="#">Detailed Description</a>	340
7.19	<a href="#">include/rtdm/rtdm_driver.h File Reference</a>	341
7.19.1	<a href="#">Detailed Description</a>	344
7.20	<a href="#">include/rtdm/rtdm_driver.h File Reference</a>	344
7.20.1	<a href="#">Detailed Description</a>	348

7.21	<a href="#">include/rtdm/rtesting.h File Reference</a>	348
7.21.1	<a href="#">Detailed Description</a>	350
7.22	<a href="#">kernel/cobalt/arch/arm/hal.c File Reference</a>	350
7.22.1	<a href="#">Detailed Description</a>	351
7.22.2	<a href="#">Function Documentation</a>	351
7.22.2.1	<a href="#">rthal_timer_release</a>	351
7.22.2.2	<a href="#">rthal_timer_request</a>	351
7.23	<a href="#">kernel/cobalt/arch/blackfin/hal.c File Reference</a>	352
7.23.1	<a href="#">Detailed Description</a>	352
7.24	<a href="#">kernel/cobalt/arch/generic/hal.c File Reference</a>	353
7.24.1	<a href="#">Detailed Description</a>	353
7.25	<a href="#">kernel/cobalt/arch/nios2/hal.c File Reference</a>	354
7.25.1	<a href="#">Detailed Description</a>	354
7.26	<a href="#">kernel/cobalt/arch/powerpc/hal.c File Reference</a>	354
7.26.1	<a href="#">Detailed Description</a>	354
7.27	<a href="#">kernel/cobalt/arch/sh/hal.c File Reference</a>	355
7.27.1	<a href="#">Detailed Description</a>	355
7.28	<a href="#">kernel/cobalt/arch/x86/hal.c File Reference</a>	355
7.28.1	<a href="#">Detailed Description</a>	356
7.29	<a href="#">kernel/cobalt/arch/x86/smi.c File Reference</a>	356
7.29.1	<a href="#">Detailed Description</a>	356
7.30	<a href="#">kernel/cobalt/nucleus/clock.c File Reference</a>	357
7.30.1	<a href="#">Detailed Description</a>	357
7.31	<a href="#">kernel/cobalt/rtdm/module.c File Reference</a>	358
7.31.1	<a href="#">Detailed Description</a>	358
7.32	<a href="#">kernel/cobalt/nucleus/bufd.c File Reference</a>	359
7.32.1	<a href="#">Detailed Description</a>	360
7.33	<a href="#">kernel/cobalt/nucleus/debug.c File Reference</a>	360
7.33.1	<a href="#">Detailed Description</a>	360
7.34	<a href="#">kernel/cobalt/nucleus/heap.c File Reference</a>	361
7.34.1	<a href="#">Detailed Description</a>	362
7.35	<a href="#">kernel/cobalt/nucleus/intr.c File Reference</a>	362
7.35.1	<a href="#">Detailed Description</a>	363
7.36	<a href="#">kernel/cobalt/nucleus/map.c File Reference</a>	364
7.36.1	<a href="#">Detailed Description</a>	364
7.37	<a href="#">kernel/cobalt/nucleus/pod.c File Reference</a>	365

7.37.1 Detailed Description . . . . .	367
7.38 kernel/cobalt/nucleus/registry.c File Reference . . . . .	367
7.38.1 Detailed Description . . . . .	368
7.39 kernel/cobalt/nucleus/sched-idle.c File Reference . . . . .	368
7.39.1 Detailed Description . . . . .	369
7.40 kernel/cobalt/nucleus/sched-rt.c File Reference . . . . .	369
7.40.1 Detailed Description . . . . .	370
7.41 kernel/cobalt/nucleus/sched-sporadic.c File Reference . . . . .	370
7.41.1 Detailed Description . . . . .	371
7.42 kernel/cobalt/nucleus/sched-tp.c File Reference . . . . .	371
7.42.1 Detailed Description . . . . .	372
7.43 kernel/cobalt/nucleus/sched.c File Reference . . . . .	373
7.43.1 Detailed Description . . . . .	373
7.44 kernel/cobalt/nucleus/select.c File Reference . . . . .	373
7.44.1 Detailed Description . . . . .	374
7.45 kernel/cobalt/nucleus/shadow.c File Reference . . . . .	375
7.45.1 Detailed Description . . . . .	375
7.46 kernel/cobalt/nucleus/synch.c File Reference . . . . .	376
7.46.1 Detailed Description . . . . .	377
7.47 kernel/cobalt/nucleus/timer.c File Reference . . . . .	378
7.47.1 Detailed Description . . . . .	379
7.48 kernel/cobalt/nucleus/vfile.c File Reference . . . . .	379
7.48.1 Detailed Description . . . . .	380
7.49 kernel/cobalt/rtdm/core.c File Reference . . . . .	381
7.49.1 Detailed Description . . . . .	384
7.50 kernel/cobalt/rtdm/device.c File Reference . . . . .	384
7.50.1 Detailed Description . . . . .	385
7.51 kernel/cobalt/rtdm/drvlib.c File Reference . . . . .	385
7.51.1 Detailed Description . . . . .	390
7.52 kernel/cobalt/syscall.c File Reference . . . . .	390
7.52.1 Detailed Description . . . . .	391
<b>8 Example Documentation . . . . .</b>	<b>393</b>
8.1 bufp-label.c . . . . .	393
8.2 bufp-readwrite.c . . . . .	393
8.3 cross-link.c . . . . .	393
8.4 iddp-label.c . . . . .	393

---

8.5	<a href="#">iddp-sendrecv.c</a>	393
8.6	<a href="#">rtcan_rtt.c</a>	393
8.7	<a href="#">rtcanconfig.c</a>	393
8.8	<a href="#">rtcanrecv.c</a>	394
8.9	<a href="#">rtcansend.c</a>	394
8.10	<a href="#">xddp-echo.c</a>	394
8.11	<a href="#">xddp-label.c</a>	394
8.12	<a href="#">xddp-stream.c</a>	394



# Chapter 1

## Deprecated List

Global `rtdm_device::open_rt` Only use non-real-time open handler in new drivers.

Global `rtdm_device::socket_rt` Only use non-real-time socket creation handler in new drivers.

Global `rtdm_operations::close_rt` Only use non-real-time close handler in new drivers.

Global `rtdm_task_sleep_until(nanosecs_abs_t wakeup_time)` Use `rtdm_task_sleep_abs` instead!



# Chapter 2

## Module Index

### 2.1 Modules

Here is a list of all modules:

Thread state flags. . . . .	11
Thread information flags. . . . .	13
HAL. . . . .	14
Clocks and timers services. . . . .	20
Condition variables services. . . . .	23
POSIX skin. . . . .	28
Message queues services. . . . .	29
Mutex services. . . . .	33
Buffer descriptors. . . . .	38
System clock services. . . . .	48
Debugging services. . . . .	49
Dynamic memory allocation services. . . . .	49
Interrupt management. . . . .	55
Lightweight key-to-object mapping service . . . . .	61
Real-time pod services. . . . .	65
Registry services. . . . .	85
File descriptors events multiplexing services. . . . .	91
Real-time shadow services. . . . .	94
Thread synchronization services. . . . .	98
Timer services. . . . .	107
Virtual file services . . . . .	112
Semaphores services. . . . .	191
Threads management services. . . . .	197
Thread cancellation. . . . .	16
Real-Time Driver Model . . . . .	226
Driver Development API . . . . .	138
Inter-Driver API . . . . .	119
Device Registration Services . . . . .	128
Synchronisation Services . . . . .	153
Clock Services . . . . .	139
Task Services . . . . .	140
Timer Services . . . . .	148
Synchronisation Services . . . . .	153

Interrupt Management Services . . . . .	171
Non-Real-Time Signalling Services . . . . .	176
Utility Services . . . . .	178
Device Profiles . . . . .	189
CAN Devices . . . . .	204
Real-time IPC protocols . . . . .	240
Serial Devices . . . . .	256
Testing Devices . . . . .	265
User API . . . . .	228
Sched . . . . .	267

## Chapter 3

# Data Structure Index

### 3.1 Data Structures

Here are the data structures with brief descriptions:

<a href="#">can_bittime</a> (Custom CAN bit-time definition ) . . . . .	269
<a href="#">can_bittime_btr</a> (Hardware-specific BTR bit-times ) . . . . .	270
<a href="#">can_bittime_std</a> (Standard bit-time parameters according to Bosch ) . . . . .	270
<a href="#">can_filter</a> (Filter for reception of CAN messages ) . . . . .	271
<a href="#">can_frame</a> (Raw CAN frame ) . . . . .	272
<a href="#">rtdm_dev_context</a> (Device context ) . . . . .	273
<a href="#">rtdm_device</a> (RTDM device ) . . . . .	274
<a href="#">rtdm_device_info</a> (Device information ) . . . . .	276
<a href="#">rtdm_operations</a> (Device operations ) . . . . .	277
<a href="#">rtipc_port_label</a> (Port label information structure ) . . . . .	278
<a href="#">rtser_config</a> (Serial device configuration ) . . . . .	279
<a href="#">rtser_event</a> (Additional information about serial device events ) . . . . .	280
<a href="#">rtser_status</a> (Serial device status ) . . . . .	281
<a href="#">sockaddr_can</a> (Socket address structure for the CAN address family ) . . . . .	281
<a href="#">sockaddr_ipc</a> (Socket address structure for the RTIPC address family ) . . . . .	282
<a href="#">xnpod</a> (Real-time pod descriptor ) . . . . .	282
<a href="#">xnsched</a> (Scheduling information structure ) . . . . .	284
<a href="#">xnthread_info</a> (Structure containing thread information ) . . . . .	285
<a href="#">xnfile_lock_ops</a> (Vfile locking operations ) . . . . .	287
<a href="#">xnfile_regular_iterator</a> (Regular vfile iterator ) . . . . .	288
<a href="#">xnfile_regular_ops</a> (Regular vfile operation descriptor ) . . . . .	289
<a href="#">xnfile_rev_tag</a> (Snapshot revision tag ) . . . . .	293
<a href="#">xnfile_snapshot</a> (Snapshot vfile descriptor ) . . . . .	293
<a href="#">xnfile_snapshot_iterator</a> (Snapshot-driven vfile iterator ) . . . . .	294
<a href="#">xnfile_snapshot_ops</a> (Snapshot vfile operation descriptor ) . . . . .	296



# Chapter 4

## File Index

### 4.1 File List

Here is a list of all documented files with brief descriptions:

include/cobalt/ <b>core.h</b> . . . . .	??
include/cobalt/ <b>errno.h</b> . . . . .	??
include/cobalt/ <b>fcntl.h</b> . . . . .	??
include/cobalt/ <b>mqueue.h</b> . . . . .	??
include/cobalt/ <b>posix.h</b> . . . . .	??
include/cobalt/ <b>pthread.h</b> . . . . .	??
include/cobalt/ <b>sched.h</b> . . . . .	??
include/cobalt/ <b>semaphore.h</b> . . . . .	??
include/cobalt/ <b>signal.h</b> . . . . .	??
include/cobalt/ <b>stdio.h</b> . . . . .	??
include/cobalt/ <b>stdlib.h</b> . . . . .	??
include/cobalt/ <b>syscall.h</b> . . . . .	??
include/cobalt/ <b>syslog.h</b> . . . . .	??
include/cobalt/ <b>time.h</b> . . . . .	??
include/cobalt/ <b>unistd.h</b> . . . . .	??
include/cobalt/ <b>wrappers.h</b> . . . . .	??
include/cobalt/nucleus/ <b>assert.h</b> . . . . .	??
include/cobalt/nucleus/ <b>bheap.h</b> . . . . .	??
include/cobalt/nucleus/ <b>bufd.h</b> . . . . .	301
include/cobalt/nucleus/ <b>clock.h</b> . . . . .	303
include/cobalt/nucleus/ <b>compiler.h</b> . . . . .	??
include/cobalt/nucleus/ <b>debug.h</b> . . . . .	??
include/cobalt/nucleus/ <b>heap.h</b> . . . . .	??
include/cobalt/nucleus/ <b>hostrt.h</b> (Definitions for global semaphore heap shared objects ) . . . . .	304
include/cobalt/nucleus/ <b>intr.h</b> . . . . .	??
include/cobalt/nucleus/ <b>map.h</b> . . . . .	306
include/cobalt/nucleus/ <b>pipe.h</b> . . . . .	??
include/cobalt/nucleus/ <b>pod.h</b> (Real-time pod interface header ) . . . . .	307
include/cobalt/nucleus/ <b>ppd.h</b> . . . . .	??
include/cobalt/nucleus/ <b>queue.h</b> . . . . .	??
include/cobalt/nucleus/ <b>registry.h</b> (This file is part of the Xenomai project ) . . . . .	310
include/cobalt/nucleus/ <b>sched-idle.h</b> (Definitions for the IDLE scheduling class ) . . . . .	311
include/cobalt/nucleus/ <b>sched-rt.h</b> (Definitions for the RT scheduling class ) . . . . .	312

include/cobalt/nucleus/ <a href="#">sched-sporadic.h</a> (Definitions for the SSP scheduling class ) . . .	313
include/cobalt/nucleus/ <a href="#">sched-tp.h</a> (Definitions for the TP scheduling class ) . . . . .	313
include/cobalt/nucleus/ <a href="#">sched.h</a> (Scheduler interface header ) . . . . .	314
include/cobalt/nucleus/ <a href="#">schedparam.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">schedqueue.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">select.h</a> (File descriptors events multiplexing header ) . . . . .	316
include/cobalt/nucleus/ <a href="#">seqlock.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">shadow.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">stat.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">synch.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">sys_ppd.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">system.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">thread.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">timer.h</a> . . . . .	318
include/cobalt/nucleus/ <a href="#">trace.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">types.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">vdso.h</a> (Definitions for global semaphore heap shared objects ) .	319
include/cobalt/nucleus/ <a href="#">version.h</a> . . . . .	??
include/cobalt/nucleus/ <a href="#">vfile.h</a> (This file is part of the Xenomai project ) . . . . .	320
include/cobalt/sys/ <a href="#">ioctl.h</a> . . . . .	??
include/cobalt/sys/ <a href="#">mman.h</a> . . . . .	??
include/cobalt/sys/ <a href="#">select.h</a> . . . . .	??
include/cobalt/sys/ <a href="#">socket.h</a> . . . . .	??
include/cobalt/sys/ <a href="#">time.h</a> . . . . .	??
include/rtdm/ <a href="#">rtcan.h</a> (Real-Time Driver Model for RT-Socket-CAN, CAN device profile header ) . . . . .	323
include/rtdm/ <a href="#">rtdm.h</a> (Real-Time Driver Model for Xenomai, user API header ) . . . . .	331
include/rtdm/ <a href="#">rtdm_driver.h</a> (Real-Time Driver Model for Xenomai, driver API header )	334
include/rtdm/ <a href="#">rtipc.h</a> (This file is part of the Xenomai project ) . . . . .	341
include/rtdm/ <a href="#">rtserial.h</a> (Real-Time Driver Model for Xenomai, serial device profile header ) . . . . .	344
include/rtdm/ <a href="#">rttesting.h</a> (Real-Time Driver Model for Xenomai, testing device profile header ) . . . . .	348
include/rtdm/ <a href="#">syscall.h</a> . . . . .	??
kernel/cobalt/ <a href="#">cancel.h</a> . . . . .	??
kernel/cobalt/ <a href="#">clock.h</a> . . . . .	??
kernel/cobalt/ <a href="#">cond.h</a> . . . . .	??
kernel/cobalt/ <a href="#">internal.h</a> . . . . .	??
kernel/cobalt/ <a href="#">monitor.h</a> . . . . .	??
kernel/cobalt/ <a href="#">mq.h</a> . . . . .	??
kernel/cobalt/ <a href="#">mutex.h</a> . . . . .	??
kernel/cobalt/ <a href="#">registry.h</a> . . . . .	??
kernel/cobalt/ <a href="#">sem.h</a> . . . . .	??
kernel/cobalt/ <a href="#">syscall.c</a> (This file is part of the Xenomai project ) . . . . .	390
kernel/cobalt/ <a href="#">thread.h</a> . . . . .	??
kernel/cobalt/ <a href="#">timer.h</a> . . . . .	??
kernel/cobalt/arch/arm/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for ARM ) . . .	350
kernel/cobalt/arch/blackfin/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for the Blackfin architecture ) . . . . .	352
kernel/cobalt/arch/generic/ <a href="#">hal.c</a> (Generic Real-Time HAL ) . . . . .	353
kernel/cobalt/arch/nios2/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture ) . . . . .	354
kernel/cobalt/arch/powerpc/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for PowerPC ) . . . . .	354



kernel/cobalt/arch/sh/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for the SuperH architecture) . . . . .	355
kernel/cobalt/arch/x86/ <a href="#">hal.c</a> (Adeos-based Real-Time Abstraction Layer for x86) . . . . .	355
kernel/cobalt/arch/x86/ <a href="#">smi.c</a> (SMI workaround for x86) . . . . .	356
kernel/cobalt/nucleus/ <a href="#">bufd.c</a> . . . . .	359
kernel/cobalt/nucleus/ <a href="#">clock.c</a> . . . . .	357
kernel/cobalt/nucleus/ <a href="#">debug.c</a> (Debug services) . . . . .	360
kernel/cobalt/nucleus/ <a href="#">heap.c</a> (Dynamic memory allocation services) . . . . .	361
kernel/cobalt/nucleus/ <a href="#">intr.c</a> (Interrupt management) . . . . .	362
kernel/cobalt/nucleus/ <a href="#">map.c</a> . . . . .	364
kernel/cobalt/nucleus/ <a href="#">pod.c</a> (Real-time pod services) . . . . .	365
kernel/cobalt/nucleus/ <a href="#">registry.c</a> (This file is part of the Xenomai project) . . . . .	367
kernel/cobalt/nucleus/ <a href="#">sched-idle.c</a> (Idle scheduling class implementation (i.e. Linux placeholder)) . . . . .	368
kernel/cobalt/nucleus/ <a href="#">sched-rt.c</a> (Common real-time scheduling class implementation (FIFO + RR)) . . . . .	369
kernel/cobalt/nucleus/ <a href="#">sched-sporadic.c</a> (POSIX SCHED_SPORADIC scheduling class) . . . . .	370
kernel/cobalt/nucleus/ <a href="#">sched-tp.c</a> (Temporal partitioning (typical of IMA systems)) . . . . .	371
kernel/cobalt/nucleus/ <a href="#">sched.c</a> . . . . .	373
kernel/cobalt/nucleus/ <a href="#">select.c</a> (File descriptors events multiplexing) . . . . .	373
kernel/cobalt/nucleus/ <a href="#">shadow.c</a> (Real-time shadow services) . . . . .	375
kernel/cobalt/nucleus/ <a href="#">synch.c</a> (Thread synchronization services) . . . . .	376
kernel/cobalt/nucleus/ <a href="#">timer.c</a> . . . . .	378
kernel/cobalt/nucleus/ <a href="#">vfile.c</a> (This file is part of the Xenomai project) . . . . .	379
kernel/cobalt/rtdm/ <a href="#">core.c</a> (Real-Time Driver Model for Xenomai, device operation multiplexing) . . . . .	381
kernel/cobalt/rtdm/ <a href="#">device.c</a> (Real-Time Driver Model for Xenomai, device management) . . . . .	384
kernel/cobalt/rtdm/ <a href="#">drvlib.c</a> (Real-Time Driver Model for Xenomai, driver library) . . . . .	385
kernel/cobalt/rtdm/ <a href="#">internal.h</a> . . . . .	??
kernel/cobalt/rtdm/ <a href="#">module.c</a> (Real-Time Driver Model for Xenomai) . . . . .	358



# Chapter 5

## Module Documentation

### 5.1 Thread state flags.

Bits reporting permanent or transient states of thread.

#### Defines

- #define `XNSUSP` 0x00000001  
*Suspended.*
- #define `XNPEND` 0x00000002  
*Sleep-wait for a resource.*
- #define `XNDELAY` 0x00000004  
*Delayed.*
- #define `XNREADY` 0x00000008  
*Linked to the ready queue.*
- #define `XNDORMANT` 0x00000010  
*Not started yet or killed.*
- #define `XNZOMBIE` 0x00000020  
*Zombie thread in deletion process.*
- #define `XNSTARTED` 0x00000080  
*Thread has been started.*
- #define `XNMAPPED` 0x00000100  
*Mapped to a regular Linux task (shadow only).*
- #define `XNRELAX` 0x00000200  
*Relaxed shadow thread (blocking bit).*
- #define `XNMIGRATE` 0x00000400

*Thread is currently migrating to another CPU.*

- `#define XNHELD 0x00000800`  
*Thread is held to process emergency.*
- `#define XNBOOST 0x00001000`  
*Undergoes a PIP boost.*
- `#define XNDEBUG 0x00002000`  
*Hit a debugger breakpoint (shadow only).*
- `#define XNLOCK 0x00004000`  
*Holds the scheduler lock (i.e.*
- `#define XNRRB 0x00008000`  
*Undergoes a round-robin scheduling.*
- `#define XNASDI 0x00010000`  
*ASR are disabled.*
- `#define XNDEFCAN 0x00020000`  
*Deferred cancelability mode (self-set only).*
- `#define XNTRAPSW 0x00040000`  
*Trap execution mode switches.*
- `#define XNFPU 0x00080000`  
*Thread uses FPU.*
- `#define XNSHADOW 0x00100000`  
*Shadow thread.*
- `#define XNROOT 0x00200000`  
*Root thread (that is, Linux/IDLE).*
- `#define XNOTHER 0x00400000`  
*Non real-time shadow (prio=0).*

### 5.1.1 Detailed Description

Bits reporting permanent or transient states of thread.

### 5.1.2 Define Documentation

#### 5.1.2.1 `#define XNHELD 0x00000800`

Thread is held to process emergency.

Referenced by `xnpod_resume_thread()`, and `xnpod_suspend_thread()`.

**5.1.2.2 #define XNLOCK 0x00004000**

Holds the scheduler lock (i.e.  
not preemptible)

Referenced by `__xnpod_reset_thread()`, `pthread_set_mode_np()`, `xnpod_set_thread_mode()`, and `xnpod_welcome_thread()`.

**5.1.2.3 #define XNMIGRATE 0x00000400**

Thread is currently migrating to another CPU.

Referenced by `xnpod_delete_thread()`.

**5.1.2.4 #define XNPEND 0x00000002**

Sleep-wait for a resource.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, `xnpod_unblock_thread()`, `xnsynch_acquire()`, `xnsynch_flush()`, `xnsynch_forget_sleeper()`, `xnsynch_sleep_on()`, `xnsynch_wakeup_one_sleeper()`, and `xnsynch_wakeup_this_sleeper()`.

**5.1.2.5 #define XNREADY 0x00000008**

Linked to the ready queue.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, and `xnpod_suspend_thread()`.

**5.1.2.6 #define XNSUSP 0x00000001**

Suspended.

Referenced by `__xnpod_reset_thread()`, `xnpod_handle_exception()`, `xnpod_init_thread()`, `xnpod_start_thread()`, and `xnpod_suspend_thread()`.

**5.2 Thread information flags.**

Bits reporting events notified to the thread.

**Defines**

- **#define XNTIMEO 0x00000001**  
*Woken up due to a timeout condition.*
- **#define XNRMID 0x00000002**  
*Pending on a removed resource.*
- **#define XNBREAK 0x00000004**  
*Forcibly awoken from a wait state.*

- `#define XNKICKED 0x00000008`  
*Forced out of primary mode (shadow only).*
- `#define XNWAKEN 0x00000010`  
*Thread waken up upon resource availability.*
- `#define XNROBBED 0x00000020`  
*Robbed from resource ownership.*
- `#define XNAFFSET 0x00000040`  
*CPU affinity changed from primary mode.*
- `#define XNPRISET 0x00000080`  
*Priority changed from primary mode.*
- `#define XNABORT 0x00000100`  
*Thread is being aborted.*
- `#define XNCANPND 0x00000200`  
*Cancellation request is pending.*
- `#define XNSWREP 0x00000400`  
*Mode switch already reported.*

### 5.2.1 Detailed Description

Bits reporting events notified to the thread.

## 5.3 HAL.

Generic Adeos-based hardware abstraction layer.

### Files

- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for ARM.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.*
- file [hal.c](#)  
*Generic Real-Time HAL.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture.*

- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for PowerPC.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the SuperH architecture.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for x86.*
- file [smi.c](#)  
*SMI workaround for x86.*

## Functions

- int [rthal\\_apc\\_alloc](#) (const char \*name, void(\*handler)(void \*cookie), void \*cookie)  
*Allocate an APC slot.*
- void [rthal\\_apc\\_free](#) (int apc)  
*Releases an APC slot.*

### 5.3.1 Detailed Description

Generic Adeos-based hardware abstraction layer.

### 5.3.2 Function Documentation

#### 5.3.2.1 int rthal\_apc\_alloc ( const char \* name, void(\*) (void \*cookie) handler, void \* cookie )

Allocate an APC slot.

APC is the acronym for Asynchronous Procedure Call, a mean by which activities from the Xenomai domain can schedule deferred invocations of handlers to be run into the Linux domain, as soon as possible when the Linux kernel gets back in control. Up to BITS\_PER\_LONG APC slots can be active at any point in time. APC support is built upon Adeos's virtual interrupt support.

The HAL guarantees that any Linux kernel service which would be callable from a regular Linux interrupt handler is also available to APC handlers.

#### Parameters

**name** is a symbolic name identifying the APC which will get reported through the /proc/xenomai/apc interface. Passing NULL to create an anonymous APC is allowed.

**handler** The address of the fault handler to call upon exception condition. The handle will be passed the *cookie* value unmodified.

**cookie** A user-defined opaque cookie the HAL will pass to the APC handler as its sole argument.

**Returns**

an valid APC id. is returned upon success, or a negative error code otherwise:

- -EINVAL is returned if *handler* is invalid.
- -EBUSY is returned if no more APC slots are available.

Environments:

This service can be called from:

- Linux domain context.

**5.3.2.2 int rthal\_apc\_free ( int *apc* )**

Releases an APC slot.

This service deallocates an APC slot obtained by [rthal\\_apc\\_alloc\(\)](#).

**Parameters**

*apc* The APC id. to release, as returned by a successful call to the [rthal\\_apc\\_alloc\(\)](#) service.

Environments:

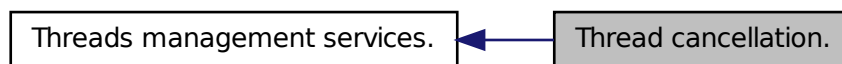
This service can be called from:

- Any domain context.

**5.4 Thread cancellation.**

Thread cancellation.

Collaboration diagram for Thread cancellation.:

**Functions**

- int [pthread\\_cancel](#) (pthread\_t thread)  
*Cancel a thread.*



- void [pthread\\_cleanup\\_push](#) (cleanup\_routine\_t \*routine, void \*arg)  
*Register a cleanup handler to be executed at the time of cancellation.*
- void [pthread\\_cleanup\\_pop](#) (int execute)  
*Unregister the last registered cleanup handler.*
- int [pthread\\_setcanceltype](#) (int type, int \*oldtype\_ptr)  
*Set cancelability type of the current thread.*
- int [pthread\\_setcancelstate](#) (int state, int \*oldstate\_ptr)  
*Set cancelability state of the current thread.*

### 5.4.1 Detailed Description

Thread cancellation. Cancellation is the mechanism by which a thread can terminate the execution of a Xenomai POSIX skin thread (created with [pthread\\_create\(\)](#)). More precisely, a thread can send a cancellation request to a Xenomai POSIX skin thread and depending on its cancelability type (see [pthread\\_setcanceltype\(\)](#)) and state (see [pthread\\_setcancelstate\(\)](#)), the target thread can then either ignore the request, honor it immediately, or defer it till it reaches a cancellation point. When threads are first created by [pthread\\_create\(\)](#), they always defer cancellation requests.

When a thread eventually honors a cancellation request, it behaves as if `pthread_exit(PTHREAD_CANCELLED)` was called. All cleanup handlers are executed in reverse order, finalization functions for thread-specific data are called, and finally the thread stops executing. If the canceled thread was joinable, the return value PTHREAD\_CANCELLED is provided to whichever thread calls `pthread_join()` on it. See `pthread_exit()` for more information.

Cancellation points are the points where the thread checks for pending cancellation requests and performs them. The POSIX threads functions `pthread_join()`, `pthread_cond_wait()`, `pthread_cond_timedwait()`, `pthread_testcancel()`, [sem\\_wait\(\)](#), [sem\\_timedwait\(\)](#), `sigwait()`, `sigwaitinfo()` and `sigtimedwait()` are cancellation points.

See also

[Specification.](#)

### 5.4.2 Function Documentation

#### 5.4.2.1 int pthread\_cancel ( pthread\_t thread )

Cancel a thread.

This service sends a cancellation request to the thread *thread* and returns immediately. Depending on the target thread cancelability state (see [pthread\\_setcancelstate\(\)](#)) and type (see [pthread\\_setcanceltype\(\)](#)), its termination is either immediate, deferred or ignored.

When the cancellation request is handled and before the thread is terminated, the cancellation cleanup handlers (registered with the [pthread\\_cleanup\\_push\(\)](#) service) are called, then the thread-specific data destructor functions (registered with `pthread_key_create()`).

Returns

0 on success;  
an error number if:

- ESRCH, the thread *thread* was not found.

See also

[Specification.](#)

References `xnpod_schedule()`, and `xnpod_unblock_thread()`.

#### 5.4.2.2 void pthread\_cleanup\_pop ( int *execute* )

Unregister the last registered cleanup handler.

If the calling thread is a Xenomai POSIX skin thread (i.e. created with `pthread_create()`), this service unregisters the last routine which was registered with `pthread_cleanup_push()` and call it if *execute* is not null.

If the caller context is invalid (not a Xenomai POSIX skin thread), this service has no effect.

This service may be called at any place, but for maximal portability, should only called in the same lexical scope as the matching call to `pthread_cleanup_push()`.

##### Parameters

*execute* if non zero, the last registered cleanup handler should be executed before it is un-registered.

##### Valid contexts:

- Xenomai POSIX skin kernel-space thread,
- Xenomai POSIX skin user-space thread (switches to primary mode).

See also

[Specification.](#)

#### 5.4.2.3 void pthread\_cleanup\_push ( cleanup\_routine\_t \* *routine*, void \* *arg* )

Register a cleanup handler to be executed at the time of cancellation.

This service registers the given *routine* to be executed at the time of cancellation of the calling thread, if this thread is a Xenomai POSIX skin thread (i.e. created with the `pthread_create()` service). If the caller context is invalid (not a Xenomai POSIX skin thread), this service has no effect.

If allocation from the system heap fails (because the system heap size is too small), this service fails silently.

The routines registered with this service get called in LIFO order when the calling thread calls `pthread_exit()` or is canceled, or when it calls the `pthread_cleanup_pop()` service with a non null argument.

##### Parameters

*routine* the cleanup routine to be registered;

*arg* the argument associated with this routine.

**Valid contexts:**

- Xenomai POSIX skin kernel-space thread,
- Xenomai POSIX skin user-space thread (switches to primary mode).

**See also**

[Specification.](#)

**5.4.2.4 int pthread\_setcancelstate ( int state, int \* oldstate\_ptr )**

Set cancelability state of the current thread.

This service atomically set the cancelability state of the calling thread and returns its previous value at the address *oldstate\_ptr*, if the calling thread is a Xenomai POSIX skin thread (i.e. created with the `pthread_create` service).

The cancelability state of a POSIX thread may be:

- `PTHREAD_CANCEL_ENABLE`, meaning that cancellation requests will be handled if received;
- `PTHREAD_CANCEL_DISABLE`, meaning that cancellation requests will not be handled if received.

**Parameters**

*state* new cancelability state of the calling thread;

*oldstate\_ptr* address where the old cancelability state will be stored on success.

**Returns**

0 on success;

an error number if:

- `EINVAL`, *state* is not a valid cancelability state;
- `EPERM`, the caller context is invalid.

**Valid contexts:**

- Xenomai POSIX skin kernel-space thread,
- Xenomai POSIX skin user-space thread (switches to primary mode).

**See also**

[Specification.](#)

**5.4.2.5 int pthread\_setcanceltype ( int type, int \* oldtype\_ptr )**

Set cancelability type of the current thread.

This service atomically sets the cancelability type of the calling thread, and return its previous value at the address *oldtype\_ptr*, if this thread is a Xenomai POSIX skin thread (i.e. was created with the `pthread_create()` service).

The cancelability type of a POSIX thread may be:

- `PTHREAD_CANCEL_DEFERRED`, meaning that cancellation requests are only handled in services which are cancellation points;
- `PTHREAD_CANCEL_ASYNCHRONOUS`, meaning that cancellation requests are handled as soon as they are sent.

### Parameters

*type* new cancelability type of the calling thread;

*oldtype\_ptr* address where the old cancelability type will be stored on success.

### Returns

0 on success;

an error number if:

- `EINVAL`, *type* is not a valid cancelability type;
- `EPERM`, the caller context is invalid.

### Valid contexts:

- Xenomai POSIX skin kernel-space thread,
- Xenomai POSIX skin user-space thread (switches to primary mode).

### See also

[Specification.](#)

## 5.5 Clocks and timers services.

Clocks and timers services.

### Functions

- static int [do\\_clock\\_host\\_realtime](#) (struct timespec \*tp)  
*Read the host-synchronised realtime clock.*
- static int [timer\\_create](#) (clockid\_t clockid, const struct sigevent \*\_\_restrict\_\_ evp, timer\_t \*\_\_restrict\_\_ timerid)  
*Create a timer object.*
- static int [timer\\_settime](#) (timer\_t timerid, int flags, const struct itimerspec \*\_\_restrict\_\_ value, struct itimerspec \*\_\_restrict\_\_ ovalue)  
*Start or stop a timer.*
- static int [timer\\_gettime](#) (timer\_t timerid, struct itimerspec \*value)  
*Get timer next expiration date and reload value.*

### 5.5.1 Detailed Description

Clocks and timers services. Xenomai POSIX skin supports two clocks:

CLOCK\_REALTIME maps to the nucleus system clock, keeping time as the amount of time since the Epoch, with a resolution of one nanosecond.

CLOCK\_MONOTONIC maps to an architecture-dependent high resolution counter, so is suitable for measuring short time intervals. However, when used for sleeping (with `clock_nanosleep()`), the CLOCK\_MONOTONIC clock has a resolution of one nanosecond, like the CLOCK\_REALTIME clock.

CLOCK\_MONOTONIC\_RAW is Linux-specific, and provides monotonic time values from a hardware timer which is not adjusted by NTP. This is strictly equivalent to CLOCK\_MONOTONIC with Xenomai, which is not NTP adjusted either.

Timer objects may be created with the `timer_create()` service using either of the two clocks. The resolution of these timers is one nanosecond, as is the case for `clock_nanosleep()`.

See also

[Specification.](#)

### 5.5.2 Function Documentation

#### 5.5.2.1 `static int do_clock_host_realtime ( struct timespec * tp ) [static]`

Read the host-synchronised realtime clock.

Obtain the current time with NTP corrections from the Linux domain

##### Parameters

*tp* pointer to a struct timespec

##### Return values

0 on success;

-1 if no suitable NTP-corrected clocksource is available

See also

[Specification.](#)

#### 5.5.2.2 `static int timer_create ( clockid_t clockid, const struct sigevent *__restrict__ evp, timer_t *__restrict__ timerid ) [inline, static]`

Create a timer object.

This service creates a time object using the clock *clockid*.

If *evp* is not `NULL`, it describes the notification mechanism used on timer expiration. Only notification via signal delivery is supported (member *sigev\_notify* of *evp* set to `SIGEV_SIGNAL`). The signal will be sent to the thread starting the timer with the `timer_settime()` service. If *evp* is `NULL`, the `SIGALRM` signal will be used.

Note that signals sent to user-space threads will cause them to switch to secondary mode.

If this service succeeds, an identifier for the created timer is returned at the address *timerid*. The timer is unarmed until started with the [timer\\_settime\(\)](#) service.

### Parameters

- clockid* clock used as a timing base;
- evp* description of the asynchronous notification to occur when the timer expires;
- timerid* address where the identifier of the created timer will be stored on success.

### Return values

- 0 on success;
- 1 with *errno* set if:
  - EINVAL, the clock *clockid* is invalid;
  - EINVAL, the member *sigev\_notify* of the **sigevent** structure at the address *evp* is not SIGEV\_SIGNAL;
  - EINVAL, the member *sigev\_signo* of the **sigevent** structure is an invalid signal number;
  - EAGAIN, the maximum number of timers was exceeded, recompile with a larger value.

### See also

[Specification.](#)

References `sem_getvalue()`, and `xntimer_init()`.

#### 5.5.2.3 static int timer\_gettime ( timer\_t timerid, struct itimerspec \* value ) [inline, static]

Get timer next expiration date and reload value.

This service stores, at the address *value*, the expiration date (member *it\_value*) and reload value (member *it\_interval*) of the timer *timerid*. The values are returned as time intervals, and as multiples of the system clock tick duration (see note in section [Clocks and timers services](#) for details on the duration of the system clock tick). If the timer was not started, the returned members *it\_value* and *it\_interval* of *value* are zero.

### Parameters

- timerid* timer identifier;
- value* address where the timer expiration date and reload value are stored on success.

### Return values

- 0 on success;
- 1 with *errno* set if:
  - EINVAL, *timerid* is invalid;
  - EPERM, the timer *timerid* does not belong to the current process.

### See also

[Specification.](#)

**5.5.2.4** `static int timer_settime ( timer_t timerid, int flags, const struct itimerspec *__restrict__ value, struct itimerspec *__restrict__ ovalue ) [inline, static]`

Start or stop a timer.

This service sets a timer expiration date and reload value of the timer *timerid*. If *ovalue* is not *NULL*, the current expiration date and reload value are stored at the address *ovalue* as with [timer\\_gettime\(\)](#).

If the member *it\_value* of the **itimerspec** structure at *value* is zero, the timer is stopped, otherwise the timer is started. If the member *it\_interval* is not zero, the timer is periodic. The current thread must be a POSIX skin thread (created with [pthread\\_create\(\)](#)) and will be notified via signal of timer expirations. Note that these notifications will cause user-space threads to switch to secondary mode.

When starting the timer, if *flags* is *TIMER\_ABSTIME*, the expiration value is interpreted as an absolute date of the clock passed to the [timer\\_create\(\)](#) service. Otherwise, the expiration value is interpreted as a time interval.

Expiration date and reload value are rounded to an integer count of nanoseconds.

#### Parameters

*timerid* identifier of the timer to be started or stopped;

*flags* one of 0 or *TIMER\_ABSTIME*;

*value* address where the specified timer expiration date and reload value are read;

*ovalue* address where the specified timer previous expiration date and reload value are stored if not *NULL*.

#### Return values

0 on success;

-1 with *errno* set if:

- *EPERM*, the caller context is invalid;
- *EINVAL*, the specified timer identifier, expiration date or reload value is invalid;
- *EPERM*, the timer *timerid* does not belong to the current process.

#### Valid contexts:

- Xenomai kernel-space POSIX skin thread,
- kernel-space thread cancellation cleanup routine,
- Xenomai POSIX skin user-space thread (switches to primary mode),
- user-space thread cancellation cleanup routine.

#### See also

[Specification.](#)

References [xntimer\\_start\(\)](#), and [xntimer\\_stop\(\)](#).

## 5.6 Condition variables services.

Condition variables services.

## Functions

- static int [pthread\\_cond\\_init](#) (struct \_\_shadow\_cond \*cnd, const pthread\_condattr\_t \*attr)  
*Initialize a condition variable.*
- static int [pthread\\_cond\\_destroy](#) (struct \_\_shadow\_cond \*cnd)  
*Destroy a condition variable.*
- static int [pthread\\_condattr\\_init](#) (pthread\_condattr\_t \*attr)  
*Initialize a condition variable attributes object.*
- static int [pthread\\_condattr\\_destroy](#) (pthread\_condattr\_t \*attr)  
*Destroy a condition variable attributes object.*
- static int [pthread\\_condattr\\_getclock](#) (const pthread\_condattr\_t \*attr, clockid\_t \*clk\_id)  
*Get the clock selection attribute from a condition variable attributes object.*
- static int [pthread\\_condattr\\_setclock](#) (pthread\_condattr\_t \*attr, clockid\_t clk\_id)  
*Set the clock selection attribute of a condition variable attributes object.*
- static int [pthread\\_condattr\\_getpshared](#) (const pthread\_condattr\_t \*attr, int \*pshared)  
*Get the process-shared attribute from a condition variable attributes object.*
- static int [pthread\\_condattr\\_setpshared](#) (pthread\_condattr\_t \*attr, int pshared)  
*Set the process-shared attribute of a condition variable attributes object.*

### 5.6.1 Detailed Description

Condition variables services. A condition variable is a synchronization object that allows threads to suspend execution until some predicate on shared data is satisfied. The basic operations on conditions are: signal the condition (when the predicate becomes true), and wait for the condition, suspending the thread execution until another thread signals the condition.

A condition variable must always be associated with a mutex, to avoid the race condition where a thread prepares to wait on a condition variable and another thread signals the condition just before the first thread actually waits on it.

Before it can be used, a condition variable has to be initialized with [pthread\\_cond\\_init\(\)](#). An attribute object, which reference may be passed to this service, allows to select the features of the created condition variable, namely the *clock* used by the [pthread\\_cond\\_timedwait\(\)](#) service (*CLOCK\_REALTIME* is used by default), and whether it may be shared between several processes (it may not be shared by default, see [pthread\\_condattr\\_setpshared\(\)](#)).

Note that only [pthread\\_cond\\_init\(\)](#) may be used to initialize a condition variable, using the static initializer *PTHREAD\_COND\_INITIALIZER* is not supported.

### 5.6.2 Function Documentation

#### 5.6.2.1 static int pthread\_cond\_destroy ( struct \_\_shadow\_cond \* cnd ) [inline, static]

Destroy a condition variable.



This service destroys the condition variable *cnd*, if no thread is currently blocked on it. The condition variable becomes invalid for all condition variable services (they all return the EINVAL error) except [pthread\\_cond\\_init\(\)](#).

#### Parameters

*cnd* the condition variable to be destroyed.

#### Returns

0 on succes,

an error number if:

- EINVAL, the condition variable *cnd* is invalid;
- EPERM, the condition variable is not process-shared and does not belong to the current process;
- EBUSY, some thread is currently using the condition variable.

#### See also

[Specification.](#)

**5.6.2.2** `static int pthread_cond_init ( struct __shadow_cond * cnd, const pthread_condattr_t * attr ) [inline, static]`

Initialize a condition variable.

This service initializes the condition variable *cnd*, using the condition variable attributes object *attr*. If *attr* is *NULL* or this service is called from user-space, default attributes are used (see [pthread\\_condattr\\_init\(\)](#)).

#### Parameters

*cnd* the condition variable to be initialized;

*attr* the condition variable attributes object.

#### Returns

0 on succes,

an error number if:

- EINVAL, the condition variable attributes object *attr* is invalid or uninitialized;
- EBUSY, the condition variable *cnd* was already initialized;
- ENOMEM, insufficient memory exists in the system heap to initialize the condition variable, increase CONFIG\_XENO\_OPT\_SYS\_HEAPSZ.

#### See also

[Specification.](#)

References [xnheap\\_alloc\(\)](#), [xnheap\\_free\(\)](#), and [xnsynch\\_init\(\)](#).

### 5.6.2.3 static int pthread\_condattr\_destroy ( pthread\_condattr\_t \* attr ) [inline, static]

Destroy a condition variable attributes object.

This service destroys the condition variable attributes object *attr*. The object becomes invalid for all condition variable services (they all return EINVAL) except [pthread\\_condattr\\_init\(\)](#).

#### Parameters

*attr* the initialized mutex attributes object to be destroyed.

#### Returns

0 on success;

an error number if:

- EINVAL, the mutex attributes object *attr* is invalid.

#### See also

[Specification.](#)

### 5.6.2.4 static int pthread\_condattr\_getclock ( const pthread\_condattr\_t \* attr, clockid\_t \* clk\_id ) [inline, static]

Get the clock selection attribute from a condition variable attributes object.

This service stores, at the address *clk\_id*, the value of the *clock* attribute in the condition variable attributes object *attr*.

See [pthread\\_cond\\_timedwait\(\)](#) documentation for a description of the effect of this attribute on a condition variable. The clock ID returned is *CLOCK\_REALTIME*, *CLOCK\_MONOTONIC* or *CLOCK\_MONOTONIC\_RAW*.

#### Parameters

*attr* an initialized condition variable attributes object,

*clk\_id* address where the *clock* attribute value will be stored on success.

#### Returns

0 on success,

an error number if:

- EINVAL, the attribute object *attr* is invalid.

#### See also

[Specification.](#)

### 5.6.2.5 static int pthread\_condattr\_getpshared ( const pthread\_condattr\_t \* attr, int \* pshared ) [inline, static]

Get the process-shared attribute from a condition variable attributes object.

This service stores, at the address *pshared*, the value of the *pshared* attribute in the condition variable attributes object *attr*.

The *pshared* attribute may only be one of `PTHREAD_PROCESS_PRIVATE` or `PTHREAD_PROCESS_SHARED`. See [pthread\\_condattr\\_setpshared\(\)](#) for the meaning of these two constants.

#### Parameters

*attr* an initialized condition variable attributes object.

*pshared* address where the value of the *pshared* attribute will be stored on success.

#### Returns

0 on success,

an error number if:

- `EINVAL`, the *pshared* address is invalid;
- `EINVAL`, the condition variable attributes object *attr* is invalid.

#### See also

[Specification.](#)

#### 5.6.2.6 `static int pthread_condattr_init ( pthread_condattr_t * attr ) [inline, static]`

Initialize a condition variable attributes object.

This services initializes the condition variable attributes object *attr* with default values for all attributes. Default value are:

- for the *clock* attribute, `CLOCK_REALTIME`;
- for the *pshared* attribute `PTHREAD_PROCESS_PRIVATE`.

If this service is called specifying a condition variable attributes object that was already initialized, the attributes object is reinitialized.

#### Parameters

*attr* the condition variable attributes object to be initialized.

#### Returns

0 on success;

an error number if:

- `ENOMEM`, the condition variable attribute object pointer *attr* is `NULL`.

#### See also

[Specification.](#)

#### 5.6.2.7 `static int pthread_condattr_setclock ( pthread_condattr_t * attr, clockid_t clk_id ) [inline, static]`

Set the clock selection attribute of a condition variable attributes object.

This service set the *clock* attribute of the condition variable attributes object *attr*.

See `pthread_cond_timedwait()` documentation for a description of the effect of this attribute on a condition variable.

**Parameters**

*attr* an initialized condition variable attributes object,  
*clk\_id* value of the *clock* attribute, may be *CLOCK\_REALTIME*, *CLOCK\_MONOTONIC* or *CLOCK\_MONOTONIC\_RAW*.

**Returns**

0 on success,  
 an error number if:

- EINVAL, the condition variable attributes object *attr* is invalid;
- EINVAL, the value of *clk\_id* is invalid for the *clock* attribute.

**See also**

[Specification.](#)

### 5.6.2.8 static int pthread\_condattr\_setpshared ( pthread\_condattr\_t \* attr, int pshared ) [inline, static]

Set the process-shared attribute of a condition variable attributes object.

This service set the *pshared* attribute of the condition variable attributes object *attr*.

**Parameters**

*attr* an initialized condition variable attributes object.

*pshared* value of the *pshared* attribute, may be one of:

- PTHREAD\_PROCESS\_PRIVATE, meaning that a condition variable created with the attributes object *attr* will only be accessible by threads within the same process as the thread that initialized the condition variable;
- PTHREAD\_PROCESS\_SHARED, meaning that a condition variable created with the attributes object *attr* will be accessible by any thread that has access to the memory where the condition variable is allocated.

**Returns**

0 on success,  
 an error status if:

- EINVAL, the condition variable attributes object *attr* is invalid;
- EINVAL, the value of *pshared* is invalid.

**See also**

[Specification.](#)

## 5.7 POSIX skin.

Xenomai POSIX skin is an implementation of a small subset of the Single Unix specification over Xenomai generic RTOS core.

Xenomai POSIX skin is an implementation of a small subset of the Single Unix specification over Xenomai generic RTOS core. The following table gives equivalence between Alchemy services and Cobalt services.

Alchemy services	Cobalt services
alchemy_alarm	<a href="#">Clocks and timers services.</a>
alchemy_cond	<a href="#">Condition variables services.</a>
alchemy_event	no direct equivalence, see <a href="#">Condition variables services.</a>
alchemy_heap	no direct equivalence
alchemy_mutex	<a href="#">Mutex services.</a>
alchemy_pipe	no direct equivalence, see <a href="#">Message queues services.</a>
alchemy_queue	<a href="#">Message queues services.</a>
alchemy_sem	<a href="#">Semaphores services.</a>
alchemy_task	<a href="#">Threads management services.</a>
alchemy_timer	<a href="#">Clocks and timers services.</a>

## 5.8 Message queues services.

Message queues services.

### Functions

- static mqd\_t [mq\\_open](#) (const char \*name, int oflags,...)  
*Open a message queue.*
- static int [mq\\_close](#) (mqd\_t fd)  
*Close a message queue.*
- static int [mq\\_unlink](#) (const char \*name)  
*Unlink a message queue.*
- static int [mq\\_getattr](#) (mqd\_t fd, struct mq\_attr \*attr)  
*Get the attributes object of a message queue.*
- static int [mq\\_setattr](#) (mqd\_t fd, const struct mq\_attr \*\_\_restrict\_\_ attr, struct mq\_attr \*\_\_restrict\_\_ oattr)  
*Set flags of a message queue.*

### 5.8.1 Detailed Description

Message queues services. A message queue allow exchanging data between real-time threads. For a POSIX message queue, maximum message length and maximum number of messages are fixed when it is created with [mq\\_open\(\)](#).

### 5.8.2 Function Documentation

#### 5.8.2.1 static int mq\_close ( mqd\_t fd ) [inline, static]

Close a message queue.

This service closes the message queue descriptor *fd*. The message queue is destroyed only when all open descriptors are closed, and when unlinked with a call to the [mq\\_unlink\(\)](#) service.

### Parameters

*fd* message queue descriptor.

### Return values

0 on success;

-1 with *errno* set if:

- EBADF, *fd* is an invalid message queue descriptor;
- EPERM, the caller context is invalid.

### Valid contexts:

- kernel module initialization or cleanup routine;
- kernel-space cancellation cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode);
- user-space cancellation cleanup routine.

### See also

[Specification.](#)

#### 5.8.2.2 static int mq\_getattr ( mqd\_t *fd*, struct mq\_attr \* *attr* ) [inline, static]

Get the attributes object of a message queue.

This service stores, at the address *attr*, the attributes of the messages queue descriptor *fd*.

The following attributes are set:

- *mq\_flags*, flags of the message queue descriptor *fd*;
- *mq\_maxmsg*, maximum number of messages in the message queue;
- *mq\_msgsize*, maximum message size;
- *mq\_curmsgs*, number of messages currently in the queue.

### Parameters

*fd* message queue descriptor;

*attr* address where the message queue attributes will be stored on success.

### Return values

0 on success;

-1 with *errno* set if:

- EBADF, *fd* is not a valid descriptor.

### See also

[Specification.](#)

### 5.8.2.3 static mqd\_t mq\_open ( const char \* *name*, int *oflags*, ... ) [static]

Open a message queue.

This service establishes a connection between the message queue named *name* and the calling context (kernel-space as a whole, or user-space process).

One of the following values should be set in *oflags*:

- **O\_RDONLY**, meaning that the returned queue descriptor may only be used for receiving messages;
- **O\_WRONLY**, meaning that the returned queue descriptor may only be used for sending messages;
- **O\_RDWR**, meaning that the returned queue descriptor may be used for both sending and receiving messages.

If no message queue named *name* exists, and *oflags* has the **O\_CREAT** bit set, the message queue is created by this function, taking two more arguments:

- a *mode* argument, of type **mode\_t**, currently ignored;
- an *attr* argument, pointer to an **mq\_attr** structure, specifying the attributes of the new message queue.

If *oflags* has the two bits **O\_CREAT** and **O\_EXCL** set and the message queue already exists, this service fails.

If the **O\_NONBLOCK** bit is set in *oflags*, the **mq\_send()**, **mq\_receive()**, **mq\_timedsend()** and **mq\_timedreceive()** services return -1 with *errno* set to **EAGAIN** instead of blocking their caller.

The following arguments of the **mq\_attr** structure at the address *attr* are used when creating a message queue:

- *mq\_maxmsg* is the maximum number of messages in the queue (128 by default);
- *mq\_msgsize* is the maximum size of each message (128 by default).

*name* may be any arbitrary string, in which slashes have no particular meaning. However, for portability, using a name which starts with a slash and contains no other slash is recommended.

#### Parameters

*name* name of the message queue to open;

*oflags* flags.

#### Returns

a message queue descriptor on success;

-1 with *errno* set if:

- **ENAMETOOLONG**, the length of the *name* argument exceeds 64 characters;
- **EEXIST**, the bits **O\_CREAT** and **O\_EXCL** were set in *oflags* and the message queue already exists;
- **ENOENT**, the bit **O\_CREAT** is not set in *oflags* and the message queue does not exist;

- ENOSPC, allocation of system memory failed, or insufficient memory exists in the system heap to create the queue, try increasing CONFIG\_XENO\_OPT\_SYS\_HEAPSZ;
- EPERM, attempting to create a message queue from an invalid context;
- EINVAL, the *attr* argument is invalid;
- EMFILE, too many descriptors are currently open.

#### Valid contexts:

When creating a message queue, only the following contexts are valid:

- kernel module initialization or cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode).

#### See also

[Specification.](#)

**5.8.2.4** `static int mq_setattr ( mqd_t fd, const struct mq_attr *__restrict attr, struct mq_attr *__restrict oattr ) [inline, static]`

Set flags of a message queue.

This service sets the flags of the *fd* descriptor to the value of the member *mq\_flags* of the **mq\_attr** structure pointed to by *attr*.

The previous value of the message queue attributes are stored at the address *oattr* if it is not *NULL*.

Only setting or clearing the O\_NONBLOCK flag has an effect.

#### Parameters

*fd* message queue descriptor;

*attr* pointer to new attributes (only *mq\_flags* is used);

*oattr* if not *NULL*, address where previous message queue attributes will be stored on success.

#### Return values

*0* on success;

*-1* with *errno* set if:

- EBADF, *fd* is not a valid message queue descriptor.

#### See also

[Specification.](#)

**5.8.2.5** `static int mq_unlink ( const char * name ) [inline, static]`

Unlink a message queue.

This service unlinks the message queue named *name*. The message queue is not destroyed until all queue descriptors obtained with the [mq\\_open\(\)](#) service are closed with the [mq\\_close\(\)](#) service. However, after a call to this service, the unlinked queue may no longer be reached with the [mq\\_open\(\)](#) service.



**Parameters**

*name* name of the message queue to be unlinked.

**Return values**

*0* on success;

*-1* with *errno* set if:

- EPERM, the caller context is invalid;
- ENAMETOOLONG, the length of the *name* argument exceeds 64 characters;
- ENOENT, the message queue does not exist.

**Valid contexts:**

- kernel module initialization or cleanup routine;
- kernel-space cancellation cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode);
- user-space cancellation cleanup routine.

**See also**

[Specification.](#)

## 5.9 Mutex services.

Mutex services.

**Functions**

- static int [pthread\\_mutexattr\\_init](#) (pthread\_mutexattr\_t \*attr)  
*Initialize a mutex attributes object.*
- static int [pthread\\_mutexattr\\_destroy](#) (pthread\_mutexattr\_t \*attr)  
*Destroy a mutex attributes object.*
- static int [pthread\\_mutexattr\\_gettype](#) (const pthread\_mutexattr\_t \*attr, int \*type)  
*Get the mutex type attribute from a mutex attributes object.*
- static int [pthread\\_mutexattr\\_settype](#) (pthread\_mutexattr\_t \*attr, int type)  
*Set the mutex type attribute of a mutex attributes object.*
- static int [pthread\\_mutexattr\\_getprotocol](#) (const pthread\_mutexattr\_t \*attr, int \*proto)  
*Get the protocol attribute from a mutex attributes object.*
- static int [pthread\\_mutexattr\\_setprotocol](#) (pthread\_mutexattr\_t \*attr, int proto)  
*Set the protocol attribute of a mutex attributes object.*
- static int [pthread\\_mutexattr\\_getpshared](#) (const pthread\_mutexattr\_t \*attr, int \*pshared)  
*Get the process-shared attribute of a mutex attributes object.*

- static int [pthread\\_mutexattr\\_setpshared](#) (pthread\_mutexattr\_t \*attr, int pshared)

*Set the process-shared attribute of a mutex attributes object.*

### 5.9.1 Detailed Description

Mutex services. A mutex is a MUTual EXclusion device, and is useful for protecting shared data structures from concurrent modifications, and implementing critical sections and monitors.

A mutex has two possible states: unlocked (not owned by any thread), and locked (owned by one thread). A mutex can never be owned by two different threads simultaneously. A thread attempting to lock a mutex that is already locked by another thread is suspended until the owning thread unlocks the mutex first.

Before it can be used, a mutex has to be initialized with `pthread_mutex_init()`. An attribute object, which reference may be passed to this service, allows to select the features of the created mutex, namely its *type* (see [pthread\\_mutexattr\\_settype\(\)](#)), the priority *protocol* it uses (see [pthread\\_mutexattr\\_setprotocol\(\)](#)) and whether it may be shared between several processes (see [pthread\\_mutexattr\\_setpshared\(\)](#)).

By default, Xenomai POSIX skin mutexes are of the normal type, use no priority protocol and may not be shared between several processes.

Note that only `pthread_mutex_init()` may be used to initialize a mutex, using the static initializer `PTHREAD_MUTEX_INITIALIZER` is not supported.

### 5.9.2 Function Documentation

#### 5.9.2.1 static int pthread\_mutexattr\_destroy ( pthread\_mutexattr\_t \* attr ) [inline, static]

Destroy a mutex attributes object.

This service destroys the mutex attributes object *attr*. The object becomes invalid for all mutex services (they all return EINVAL) except [pthread\\_mutexattr\\_init\(\)](#).

#### Parameters

*attr* the initialized mutex attributes object to be destroyed.

#### Returns

- 0 on success;
- an error number if:
  - EINVAL, the mutex attributes object *attr* is invalid.

#### See also

[Specification.](#)

#### 5.9.2.2 static int pthread\_mutexattr\_getprotocol ( const pthread\_mutexattr\_t \* attr, int \* proto ) [inline, static]

Get the protocol attribute from a mutex attributes object.

This service stores, at the address *proto*, the value of the *protocol* attribute in the mutex attributes object *attr*.

The *protocol* attribute may only be one of *PTHREAD\_PRIO\_NONE* or *PTHREAD\_PRIO\_INHERIT*. See [pthread\\_mutexattr\\_setprotocol\(\)](#) for the meaning of these two constants.

#### Parameters

*attr* an initialized mutex attributes object;

*proto* address where the value of the *protocol* attribute will be stored on success.

#### Returns

0 on success,  
an error number if:

- EINVAL, the *proto* address is invalid;
- EINVAL, the mutex attributes object *attr* is invalid.

#### See also

[Specification.](#)

#### 5.9.2.3 static int pthread\_mutexattr\_getpshared ( const pthread\_mutexattr\_t \* attr, int \* pshared ) [inline, static]

Get the process-shared attribute of a mutex attributes object.

This service stores, at the address *pshared*, the value of the *pshared* attribute in the mutex attributes object *attr*.

The *pshared* attribute may only be one of *PTHREAD\_PROCESS\_PRIVATE* or *PTHREAD\_PROCESS\_SHARED*. See [pthread\\_mutexattr\\_setpshared\(\)](#) for the meaning of these two constants.

#### Parameters

*attr* an initialized mutex attributes object;

*pshared* address where the value of the *pshared* attribute will be stored on success.

#### Returns

0 on success;  
an error number if:

- EINVAL, the *pshared* address is invalid;
- EINVAL, the mutex attributes object *attr* is invalid.

#### See also

[Specification.](#)

#### 5.9.2.4 static int pthread\_mutexattr\_gettype ( const pthread\_mutexattr\_t \* *attr*, int \* *type* ) [inline, static]

Get the mutex type attribute from a mutex attributes object.

This service stores, at the address *type*, the value of the *type* attribute in the mutex attributes object *attr*.

See pthread\_mutex\_lock() and pthread\_mutex\_unlock() documentations for a description of the values of the *type* attribute and their effect on a mutex.

##### Parameters

- attr* an initialized mutex attributes object,
- type* address where the *type* attribute value will be stored on success.

##### Returns

- 0 on success,
- an error number if:
  - EINVAL, the *type* address is invalid;
  - EINVAL, the mutex attributes object *attr* is invalid.

##### See also

[Specification.](#)

#### 5.9.2.5 static int pthread\_mutexattr\_init ( pthread\_mutexattr\_t \* *attr* ) [inline, static]

Initialize a mutex attributes object.

This services initializes the mutex attributes object *attr* with default values for all attributes. Default value are :

- for the *type* attribute, *PTHREAD\_MUTEX\_NORMAL*;
- for the *protocol* attribute, *PTHREAD\_PRIO\_NONE*;
- for the *pshared* attribute, *PTHREAD\_PROCESS\_PRIVATE*.

If this service is called specifying a mutex attributes object that was already initialized, the attributes object is reinitialized.

##### Parameters

- attr* the mutex attributes object to be initialized.

##### Returns

- 0 on success;
- an error number if:
  - ENOMEM, the mutex attributes object pointer *attr* is *NULL*.

##### See also

[Specification.](#)

### 5.9.2.6 `static int pthread_mutexattr_setprotocol ( pthread_mutexattr_t * attr, int proto )` `[inline, static]`

Set the protocol attribute of a mutex attributes object.

This service set the *type* attribute of the mutex attributes object *attr*.

#### Parameters

*attr* an initialized mutex attributes object,

*proto* value of the *protocol* attribute, may be one of:

- PTHREAD\_PRIO\_NONE, meaning that a mutex created with the attributes object *attr* will not follow any priority protocol;
- PTHREAD\_PRIO\_INHERIT, meaning that a mutex created with the attributes object *attr*, will follow the priority inheritance protocol.

The value PTHREAD\_PRIO\_PROTECT (priority ceiling protocol) is unsupported.

#### Returns

0 on success,

an error number if:

- EINVAL, the mutex attributes object *attr* is invalid;
- ENOTSUP, the value of *proto* is unsupported;
- EINVAL, the value of *proto* is invalid.

#### See also

[Specification.](#)

### 5.9.2.7 `static int pthread_mutexattr_setpshared ( pthread_mutexattr_t * attr, int pshared )` `[inline, static]`

Set the process-shared attribute of a mutex attributes object.

This service set the *pshared* attribute of the mutex attributes object *attr*.

#### Parameters

*attr* an initialized mutex attributes object.

*pshared* value of the *pshared* attribute, may be one of:

- PTHREAD\_PROCESS\_PRIVATE, meaning that a mutex created with the attributes object *attr* will only be accessible by threads within the same process as the thread that initialized the mutex;
- PTHREAD\_PROCESS\_SHARED, meaning that a mutex created with the attributes object *attr* will be accessible by any thread that has access to the memory where the mutex is allocated.

#### Returns

0 on success,

an error status if:

- EINVAL, the mutex attributes object *attr* is invalid;
- EINVAL, the value of *pshared* is invalid.

See also

[Specification.](#)

#### 5.9.2.8 static int pthread\_mutexattr\_settype ( pthread\_mutexattr\_t \* attr, int type ) [inline, static]

Set the mutex type attribute of a mutex attributes object.

This service set the *type* attribute of the mutex attributes object *attr*.

See pthread\_mutex\_lock() and pthread\_mutex\_unlock() documentations for a description of the values of the *type* attribute and their effect on a mutex.

The PTHREAD\_MUTEX\_DEFAULT default *type* is the same as PTHREAD\_MUTEX\_NORMAL. Note that using a Xenomai POSIX skin recursive mutex with a Xenomai POSIX skin condition variable is safe (see pthread\_cond\_wait() documentation).

#### Parameters

*attr* an initialized mutex attributes object,  
*type* value of the *type* attribute.

#### Returns

- 0 on success,  
 an error number if:
- EINVAL, the mutex attributes object *attr* is invalid;
  - EINVAL, the value of *type* is invalid for the *type* attribute.

See also

[Specification.](#)

## 5.10 Buffer descriptors.

### Files

- file [bufd.h](#)
- file [bufd.c](#)

### Functions

- static void [xnbufd\\_map\\_uread](#) (struct xnbufd \*bufd, const void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from user memory.*
- static void [xnbufd\\_map\\_uwrite](#) (struct xnbufd \*bufd, void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to user memory.*

- `ssize_t xnbuofd_unmap_uread` (struct xnbuofd \*buofd)  
*Finalize a buffer descriptor obtained from `xnbuofd_map_uread()`.*
- `ssize_t xnbuofd_unmap_uwrite` (struct xnbuofd \*buofd)  
*Finalize a buffer descriptor obtained from `xnbuofd_map_uwrite()`.*
- static void `xnbuofd_map_kread` (struct xnbuofd \*buofd, const void \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from kernel memory.*
- static void `xnbuofd_map_kwrite` (struct xnbuofd \*buofd, void \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to kernel memory.*
- `ssize_t xnbuofd_unmap_kread` (struct xnbuofd \*buofd)  
*Finalize a buffer descriptor obtained from `xnbuofd_map_kread()`.*
- `ssize_t xnbuofd_unmap_kwrite` (struct xnbuofd \*buofd)  
*Finalize a buffer descriptor obtained from `xnbuofd_map_kwrite()`.*
- `ssize_t xnbuofd_copy_to_kmem` (void \*ptr, struct xnbuofd \*buofd, size\_t len)  
*Copy memory covered by a buffer descriptor to kernel memory.*
- `ssize_t xnbuofd_copy_from_kmem` (struct xnbuofd \*buofd, void \*from, size\_t len)  
*Copy kernel memory to the area covered by a buffer descriptor.*
- void `xnbuofd_invalidate` (struct xnbuofd \*buofd)  
*Invalidate a buffer descriptor.*
- static void `xnbuofd_reset` (struct xnbuofd \*buofd)  
*Reset a buffer descriptor.*

### 5.10.1 Detailed Description

A buffer descriptor is a simple abstraction dealing with copy operations to/from memory buffers which may belong to different address spaces.

To this end, the buffer descriptor library provides a small set of copy routines which are aware of address space restrictions when moving data, and a generic container type which can hold a reference to - or cover - a particular memory area, either present in kernel space, or in any of the existing user memory contexts.

The goal of the buffer descriptor abstraction is to hide address space specifics from Xenomai services dealing with memory areas, allowing them to operate on multiple address spaces seamlessly.

The common usage patterns are as follows:

- Implementing a Xenomai syscall returning a bulk of data to the caller, which may have to be copied back to either kernel or user space:

```

[Syscall implementation]
ssize_t rt_bulk_read_inner(struct xnbufd *bufd)
{
    ssize_t ret;
    size_t len;
    void *bulk;

    bulk = get_next_readable_bulk(&len);
    ret = xnbufd_copy_from_kmem(bufd, bulk, min(bufd->b_len, len));
    free_bulk(bulk);

    ret = this_may_fail();
    if (ret)
        xnbufd_invalidate(bufd);

    return ret;
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_read(void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_kwrite(&bufd);

    return ret;
}

[Userland trampoline for user syscalls]
int __rt_bulk_read(struct pt_regs *regs)
{
    struct xnbufd bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbufd_map_uwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_uwrite(&bufd);

    return ret;
}

```

- Implementing a Xenomai syscall receiving a bulk of data from the caller, which may have to be read from either kernel or user space:

```

[Syscall implementation]
ssize_t rt_bulk_write_inner(struct xnbufd *bufd)
{
    void *bulk = get_free_bulk(bufd->b_len);
    return xnbufd_copy_to_kmem(bulk, bufd, bufd->b_len);
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_write(const void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kread(&bufd, ptr, len);

```



```

    ret = rt_bulk_write_inner(&bufd);
    xnbuid_unmap_kread(&bufd);

    return ret;
}

[Userland trampoline for user syscalls]
int __rt_bulk_write(struct pt_regs *regs)
{
    struct xnbuid bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbuid_map_uread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbuid_unmap_uread(&bufd);

    return ret;
}

```

## 5.10.2 Function Documentation

### 5.10.2.1 `ssize_t xnbuid_copy_from_kmem ( struct xnbuid * bufd, void * from, size_t len )`

Copy kernel memory to the area covered by a buffer descriptor.

This routine copies *len* bytes from the kernel memory starting at *from* to the area referred to by the buffer descriptor *bufd*. `xnbuid_copy_from_kmem()` tracks the write offset within the destination memory internally, so that it may be called several times in a loop, until the entire memory area is stored.

The destination address space is dealt with, according to the following rules:

- if *bufd* refers to a writable kernel area (i.e. see `xnbuid_map_kwrite()`), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a writable user area (i.e. see `xnbuid_map_uwrite()`), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while writing to that memory.
- if *bufd* refers to a user area which may not be immediately written to from the current context, the copy is postponed until `xnbuid_unmap_uwrite()` is invoked for *ubufd*, at which point the copy will take place. In such a case, the source memory is transferred to a carry over buffer allocated internally; this operation may lead to request dynamic memory from the nucleus heap if *len* is greater than 64 bytes.

#### Parameters

*bufd* The address of the buffer descriptor covering the user memory to copy data to.

*from* The start address of the kernel memory to copy from.

*len* The length of the kernel memory to copy to *bufd*.

**Returns**

The number of bytes written so far to the memory area covered by *ubufd*. Otherwise,

- -ENOMEM is returned when no memory is available from the nucleus heap to allocate the carry over buffer.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while writing to the user area. For that reason, [xnbufd\\_copy\\_from\\_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

**Note**

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

**5.10.2.2 ssize\_t xnbufd\_copy\_to\_kmem ( void \* to, struct xnbufd \* bufd, size\_t len )**

Copy memory covered by a buffer descriptor to kernel memory.

This routine copies *len* bytes from the area referred to by the buffer descriptor *bufd* to the kernel memory area *to*. [xnbufd\\_copy\\_to\\_kmem\(\)](#) tracks the read offset within the source memory internally, so that it may be called several times in a loop, until the entire memory area is loaded.

The source address space is dealt with, according to the following rules:

- if *bufd* refers to readable kernel area (i.e. see [xnbufd\\_map\\_kread\(\)](#)), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a readable user area (i.e. see [xnbufd\\_map\\_uread\(\)](#)), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while reading from that memory.
- any attempt to read from *bufd* from a non-suitable context is considered as a bug, and will raise a panic assertion when the nucleus is compiled in debug mode.

**Parameters**

*to* The start address of the kernel memory to copy to.

*bufd* The address of the buffer descriptor covering the user memory to copy data from.

*len* The length of the user memory to copy from *bufd*.

**Returns**

The number of bytes read so far from the memory area covered by *ubufd*. Otherwise:

- -EINVAL is returned upon attempt to read from the user area from an invalid context. This error is only returned when the debug mode is disabled; otherwise a panic assertion is raised.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while reading from the user area. For that reason, [xnbufd\\_copy\\_to\\_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

#### Note

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

#### 5.10.2.3 void xnbufd\_invalidate ( struct xnbufd \* *bufd* )

Invalidate a buffer descriptor.

The buffer descriptor is invalidated, making it unusable for further copy operations. If an outstanding carry over buffer was allocated by a previous call to [xnbufd\\_copy\\_from\\_kmem\(\)](#), it is immediately freed so that no data transfer will happen when the descriptor is finalized.

The only action that may subsequently be performed on an invalidated descriptor is calling the relevant unmapping routine for it. For that reason, [xnbufd\\_invalidate\(\)](#) should be invoked on the error path when data may have been transferred to the carry over buffer.

#### Parameters

*bufd* The address of the buffer descriptor to invalidate.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 5.10.2.4 void xnbufd\_map\_kread ( struct xnbufd \* *bufd*, const void \* *ptr*, size\_t *len* ) [inline, static]

Initialize a buffer descriptor for reading from kernel memory.

The new buffer descriptor may be used to copy data from kernel memory. This routine should be used in pair with [xnbufd\\_unmap\\_kread\(\)](#).

**Parameters**

*bufd* The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

*ptr* The start of the kernel buffer to map.

*len* The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

**5.10.2.5 void xnbuid\_map\_kwrite ( struct xnbuid \* *bufd*, void \* *ptr*, size\_t *len* ) [inline, static]**

Initialize a buffer descriptor for writing to kernel memory.

The new buffer descriptor may be used to copy data to kernel memory. This routine should be used in pair with [xnbuid\\_unmap\\_kwrite\(\)](#).

**Parameters**

*bufd* The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

*ptr* The start of the kernel buffer to map.

*len* The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

**5.10.2.6 void xnbuid\_map\_uread ( struct xnbuid \* *bufd*, const void \_\_user \* *ptr*, size\_t *len* ) [inline, static]**

Initialize a buffer descriptor for reading from user memory.

The new buffer descriptor may be used to copy data from user memory. This routine should be used in pair with [xnbuid\\_unmap\\_uread\(\)](#).

**Parameters**

*bufd* The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

*ptr* The start of the user buffer to map.

*len* The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

**5.10.2.7 void xnbuid\_map\_uwrite ( struct xnbuid \* bufd, void \_\_user \* ptr, size\_t len )  
[inline, static]**

Initialize a buffer descriptor for writing to user memory.

The new buffer descriptor may be used to copy data to user memory. This routine should be used in pair with [xnbuid\\_unmap\\_uwrite\(\)](#).

**Parameters**

*bufd* The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

*ptr* The start of the user buffer to map.

*len* The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

**5.10.2.8 void xnbuid\_reset ( struct xnbuid \* bufd ) [inline, static]**

Reset a buffer descriptor.

The buffer descriptor is reset, so that all data already copied is forgotten. Any carry over buffer allocated is kept, though.

**Parameters**

*bufd* The address of the buffer descriptor to reset.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 5.10.2.9 `ssize_t xnbuid_unmap_kread ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_kread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_kread\(\)](#), to read data from a kernel area.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes read so far from the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 5.10.2.10 `ssize_t xnbuid_unmap_kwrite ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_kwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_kwrite\(\)](#), to write data to a kernel area.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes written so far to the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 5.10.2.11 `ssize_t xnbuid_unmap_uread ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_uread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_uread\(\)](#), to read data from a user area.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes read so far from the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

##### Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

#### 5.10.2.12 `ssize_t xnbuid_unmap_uwrite ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_uwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_uwrite\(\)](#), to write data to a user area.

The main action taken is to write the contents of the kernel memory area passed to [xnbuid\\_copy\\_from\\_kmem\(\)](#) whenever the copy operation was postponed at that time; the carry over buffer is eventually released as needed. If [xnbuid\\_copy\\_from\\_kmem\(\)](#) was allowed to copy to the destination user memory at once, then [xnbuid\\_unmap\\_uwrite\(\)](#) leads to a no-op.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes written so far to the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

#### Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

## 5.11 System clock services.

### Files

- file `clock.h`
- file `clock.c`

### Functions

- void `xnclock_adjust` (`xnsticks_t` `delta`)  
*Adjust the clock time for the system.*

#### 5.11.1 Function Documentation

##### 5.11.1.1 void `xnclock_adjust` ( `xnsticks_t` *delta* )

Adjust the clock time for the system.

Xenomai tracks the current time as a monotonously increasing count of ticks since the epoch. The epoch is initially the same as the underlying machine time.

This service changes the epoch for the system by applying the specified tick delta on the wallclock offset.

#### Parameters

*delta* The adjustment of the system time expressed in ticks.

#### Note

This routine must be entered `nklock` locked, interrupts off.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.



## 5.12 Debugging services.

### Files

- file [debug.c](#)  
*Debug services.*

## 5.13 Dynamic memory allocation services.

### Files

- file [heap.c](#)  
*Dynamic memory allocation services.*

### Functions

- int [xnheap\\_init](#) (xnheap\_t \*heap, void \*heapaddr, u\_long heapsize, u\_long pagesize)  
*Initialize a memory heap.*
- void [xnheap\\_set\\_label](#) (xnheap\_t \*heap, const char \*label,...)  
*Set the heap's label string.*
- void [xnheap\\_destroy](#) (xnheap\_t \*heap, void(\*flushfn)(xnheap\_t \*heap, void \*extaddr, u\_long extsize, void \*cookie), void \*cookie)  
*Destroys a memory heap.*
- void \* [xnheap\\_alloc](#) (xnheap\_t \*heap, u\_long size)  
*Allocate a memory block from a memory heap.*
- int [xnheap\\_test\\_and\\_free](#) (xnheap\_t \*heap, void \*block, int(\*ckfn)(void \*block))  
*Test and release a memory block to a memory heap.*
- int [xnheap\\_free](#) (xnheap\_t \*heap, void \*block)  
*Release a memory block to a memory heap.*
- int [xnheap\\_extend](#) (xnheap\_t \*heap, void \*extaddr, u\_long extsize)  
*Extend a memory heap.*
- void [xnheap\\_schedule\\_free](#) (xnheap\_t \*heap, void \*block, xnholder\_t \*link)  
*Schedule a memory block for release.*

### 5.13.1 Detailed Description

Dynamic memory allocation services.

The implementation of the memory allocator follows the algorithm described in a USENIX 1988 paper called "Design of a General Purpose Memory Allocator for the 4.3BSD Unix Kernel" by Marshall K. McKusick and Michael J. Karels. You can find it at various locations on the net, including <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf>. A minor variation allows this implementation to have 'extendable' heaps when needed, with multiple memory extents providing autonomous page address spaces.

The data structures hierarchy is as follows:

```

HEAP {
    block_buckets[]
    extent_queue -----+
}
|
V
    EXTENT #1 {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    } -+
|
|
V
    EXTENT #n {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    }

```

### 5.13.2 Function Documentation

#### 5.13.2.1 void\* xnheap\_alloc ( xnheap\_t \* heap, u\_long size )

Allocate a memory block from a memory heap.

Allocates a contiguous region of memory from an active memory heap. Such allocation is guaranteed to be time-bounded.

#### Parameters

*heap* The descriptor address of the heap to get memory from.

*size* The size in bytes of the requested block. Sizes lower or equal to the page size are rounded either to the minimum allocation size if lower than this value, or to the minimum alignment size if greater or equal to this value. In the current implementation, with MINALLOC = 8 and MINALIGN = 16, a 7 bytes request will be rounded to 8 bytes, and a 17 bytes request will be rounded to 32.

#### Returns

The address of the allocated region upon success, or NULL if no memory is available from the specified heap.

#### Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `pthread_cond_init()`, and `xnshadow_map()`.

**5.13.2.2** `void xnheap_destroy ( xnheap_t * heap, void(*) (xnheap_t *heap, void *extaddr, u_long extsize, void *cookie) flushfn, void * cookie )`

Destroys a memory heap.

Destroys a memory heap.

#### Parameters

*heap* The descriptor address of the destroyed heap.

*flushfn* If non-NULL, the address of a flush routine which will be called for each extent attached to the heap. This routine can be used by the calling code to further release the heap memory.

*cookie* If *flushfn* is non-NULL, *cookie* is an opaque pointer which will be passed unmodified to *flushfn*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`, and `xnpod_shutdown()`.

**5.13.2.3** `int xnheap_extend ( xnheap_t * heap, void * extaddr, u_long extsize )`

Extend a memory heap.

Add a new extent to an existing memory heap.

#### Parameters

*heap* The descriptor address of the heap to add an extent to.

*extaddr* The address of the extent memory.

*extsize* The size of the extent memory (in bytes). In the current implementation, this size must match the one of the initial extent passed to `xnheap_init()`.

**Returns**

0 is returned upon success, or -EINVAL is returned if *extsize* differs from the initial extent's size.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

**5.13.2.4 int xnheap\_free ( xnheap\_t \* heap, void \* block )**

Release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from.

**Parameters**

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

**Returns**

0 is returned upon success, or one of the following error codes:

- -EFAULT is returned whenever the memory address is outside the heap address space.
- -EINVAL is returned whenever the memory address does not represent a valid block.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References xnheap\_test\_and\_free().

Referenced by pthread\_cond\_init().

### 5.13.2.5 `int xnheap_init ( xnheap_t * heap, void * heapaddr, u_long heapsize, u_long pagesize )`

Initialize a memory heap.

Initializes a memory heap suitable for time-bounded allocation requests of dynamic memory.

#### Parameters

*heap* The address of a heap descriptor which will be used to store the allocation data. This descriptor must always be valid while the heap is active therefore it must be allocated in permanent memory.

*heapaddr* The address of the heap storage area. All allocations will be made from the given area in time-bounded mode. Since additional extents can be added to a heap, this parameter is also known as the "initial extent".

*heapsize* The size in bytes of the initial extent pointed at by *heapaddr*. *heapsize* must be a multiple of *pagesize* and lower than 16 Mbytes. *heapsize* must be large enough to contain a dynamically-sized internal header. The following formula gives the size of this header:

$H = \text{heapsize}, P = \text{pagesize}, M = \text{sizeof}(\text{struct pagemap}), E = \text{sizeof}(\text{xnextent\_t})$

$\text{hdrsize} = ((H - E) * M) / (M + 1)$

This value is then aligned on the next 16-byte boundary. The routine `xnheap_overhead()` computes the corrected heap size according to the previous formula.

#### Parameters

*pagesize* The size in bytes of the fundamental memory page which will be used to subdivide the heap internally. Choosing the right page size is important regarding performance and memory fragmentation issues, so it might be a good idea to take a look at <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf> to pick the best one for your needs. In the current implementation, *pagesize* must be a power of two in the range [ 8 .. 32768 ] inclusive.

#### Returns

0 is returned upon success, or one of the following error codes:

- -EINVAL is returned whenever a parameter is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

### 5.13.2.6 void xnheap\_schedule\_free ( xnheap\_t \* heap, void \* block, xnholder\_t \* link )

Schedule a memory block for release.

This routine records a block for later release by `xnheap_finalize_free()`. This service is useful to lazily free blocks of heap memory when immediate release is not an option, e.g. when active references are still pending on the object for a short time after the call. `xnheap_finalize_free()` is expected to be eventually called by the client code at some point in the future when actually freeing the idle objects is deemed safe.

#### Parameters

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

*link* The address of a link member, likely but not necessarily within the released object, which will be used by the heap manager to hold the block in the queue of idle objects.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

### 5.13.2.7 void xnheap\_set\_label ( xnheap\_t \* heap, const char \* label, ... )

Set the heap's label string.

Set the heap label that will be used in statistic outputs.

#### Parameters

*heap* The address of a heap descriptor.

*label* Label string displayed in statistic outputs. This parameter can be a format string, in which case succeeding parameters will be used to resolve the final label.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

### 5.13.2.8 `int xnheap_test_and_free ( xnheap_t * heap, void * block, int(*)(void *block) ckfn )`

Test and release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from. Before the actual release is performed, an optional user-defined can be invoked to check for additional criteria with respect to the request consistency.

#### Parameters

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

*ckfn* The address of a user-supplied verification routine which is to be called after the memory address specified by *block* has been checked for validity. The routine is expected to proceed to further consistency checks, and either return zero upon success, or non-zero upon error. In the latter case, the release process is aborted, and *ckfn*'s return value is passed back to the caller of this service as its error return code. *ckfn* must not trigger the rescheduling procedure either directly or indirectly.

#### Returns

0 is returned upon success, or -EINVAL is returned whenever the block is not a valid region of the specified heap. Additional return codes can also be defined locally by the *ckfn* routine.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnheap_free()`.

## 5.14 Interrupt management.

### Files

- file [intr.c](#)  
*Interrupt management.*

### Functions

- int [xnintr\\_init](#) (xnintr\_t \*intr, const char \*name, unsigned irq, xnintr\_isr\_t isr, xnintr\_iack\_t iack, xnintr\_flags\_t flags)  
*Initialize an interrupt object.*

- `int xnintr_destroy (xnintr_t *intr)`  
*Destroy an interrupt object.*
- `int xnintr_attach (xnintr_t *intr, void *cookie)`  
*Attach an interrupt object.*
- `int xnintr_detach (xnintr_t *intr)`  
*Detach an interrupt object.*
- `void xnintr_enable (xnintr_t *intr)`  
*Enable an interrupt object.*
- `void xnintr_disable (xnintr_t *intr)`  
*Disable an interrupt object.*
- `void xnintr_affinity (xnintr_t *intr, xnarch_cpumask_t cpumask)`  
*Set interrupt's processor affinity.*

### 5.14.1 Detailed Description

Interrupt management.

### 5.14.2 Function Documentation

#### 5.14.2.1 `void xnintr_affinity ( xnintr_t * intr, xnarch_cpumask_t cpumask )`

Set interrupt's processor affinity.

Causes the IRQ associated with the interrupt object *intr* to be received only on processors which bits are set in *cpumask*.

#### Parameters

*intr* The descriptor address of the interrupt object which affinity is to be changed.  
*cpumask* The new processor affinity of the interrupt object.

#### Note

Depending on architectures, setting more than one bit in *cpumask* could be meaningless.

#### 5.14.2.2 `int xnintr_attach ( xnintr_t * intr, void * cookie )`

Attach an interrupt object.

Attach an interrupt object previously initialized by `xnintr_init()`. After this operation is completed, all IRQs received from the corresponding interrupt channel are directed to the object's ISR.

#### Parameters

*intr* The descriptor address of the interrupt object to attach.



*cookie* A user-defined opaque value which is stored into the interrupt object descriptor for further retrieval by the ISR/ISR handlers.

### Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while attaching the interrupt.
- -EBUSY is returned if the interrupt object was already attached.

### Note

The caller **must not** hold `nklock` when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

### Note

Attaching an interrupt resets the tracked number of receipts to zero.

Referenced by `rtdm_irq_request()`.

#### 5.14.2.3 `int xnintr_destroy ( xnintr_t * intr )`

Destroy an interrupt object.

Destroys an interrupt object previously initialized by `xnintr_init()`. The interrupt object is automatically detached by a call to `xnintr_detach()`. No more IRQs will be dispatched by this object after this service has returned.

### Parameters

*intr* The descriptor address of the interrupt object to destroy.

### Returns

0 is returned on success. Otherwise, -EINVAL is returned if an error occurred while detaching the interrupt (see `xnintr_detach()`).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

References `xnintr_detach()`.

#### 5.14.2.4 `int xnintr_detach ( xnintr_t * intr )`

Detach an interrupt object.

Detach an interrupt object previously attached by `xnintr_attach()`. After this operation is completed, no more IRQs are directed to the object's ISR, but the interrupt object itself remains valid. A detached interrupt object can be attached again by a subsequent call to `xnintr_attach()`.

##### Parameters

*intr* The descriptor address of the interrupt object to detach.

##### Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while detaching the interrupt, or if the interrupt object was not attached. In both cases, no action is performed.

##### Note

The caller **must not** hold `nklock` when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by `xnintr_destroy()`.

#### 5.14.2.5 `void xnintr_disable ( xnintr_t * intr )`

Disable an interrupt object.

Disables the hardware interrupt line associated with an interrupt object. This operation invalidates further interrupt requests from the given source until the IRQ line is re-enabled anew.

##### Parameters

*intr* The descriptor address of the interrupt object to disable.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

**5.14.2.6 void xnintr\_enable ( xnintr\_t \* intr )**

Enable an interrupt object.

Enables the hardware interrupt line associated with an interrupt object. Over real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

**Parameters**

*intr* The descriptor address of the interrupt object to enable.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by `rtdm_irq_request()`.

**5.14.2.7 int xnintr\_init ( xnintr\_t \* intr, const char \* name, unsigned irq, xn\_isr\_t isr, xniack\_t iack, xnflags\_t flags )**

Initialize an interrupt object.

Associates an interrupt object with an IRQ line.

When an interrupt occurs on the given *irq* line, the ISR is fired in order to deal with the hardware event. The interrupt service code may call any non-suspensive service from the nucleus.

Upon receipt of an IRQ, the ISR is immediately called on behalf of the interrupted stack context, the rescheduling procedure is locked, and the interrupt source is masked at hardware level. The status value returned by the ISR is then checked for the following values:

- `XN_ISR_HANDLED` indicates that the interrupt request has been fulfilled by the ISR.
- `XN_ISR_NONE` indicates the opposite to `XN_ISR_HANDLED`. The ISR must always return this value when it determines that the interrupt request has not been issued by the dedicated hardware device.

In addition, one of the following bits may be set by the ISR :

NOTE: use these bits with care and only when you do understand their effect on the system. The ISR is not encouraged to use these bits in case it shares the IRQ line with other ISRs in the real-time domain.

- `XN_ISR_NOENABLE` causes the nucleus to ask the real-time control layer `_not_` to re-enable the IRQ line (read the following section). `ipipe_end_irq()` must be called to re-enable the IRQ line later.

- `XN_ISR_PROPAGATE` tells the nucleus to require the real-time control layer to forward the IRQ. For instance, this would cause the Adeos control layer to propagate the interrupt down the interrupt pipeline to other Adeos domains, such as Linux. This is the regular way to share interrupts between the nucleus and the host system. In effect, `XN_ISR_PROPAGATE` implies `XN_ISR_NOENABLE` since it would make no sense to re-enable the interrupt channel before the next domain down the pipeline has had a chance to process the propagated interrupt.

The nucleus re-enables the IRQ line by default. Over some real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

A count of interrupt receipts is tracked into the interrupt descriptor, and reset to zero each time the interrupt object is attached. Since this count could wrap around, it should be used as an indication of interrupt activity only.

### Parameters

*intr* The address of a interrupt object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

*name* An ASCII string standing for the symbolic name of the interrupt object or NULL ("`<unknown>`" will be applied then).

*irq* The hardware interrupt channel associated with the interrupt object. This value is architecture-dependent. An interrupt object must then be attached to the hardware interrupt vector using the `xnintr_attach()` service for the associated IRQs to be directed to this object.

*isr* The address of a valid low-level interrupt service routine if this parameter is non-zero. This handler will be called each time the corresponding IRQ is delivered on behalf of an interrupt context. When called, the ISR is passed the descriptor address of the interrupt object.

*iack* The address of an optional interrupt acknowledge routine, aimed at replacing the default one. Only very specific situations actually require to override the default setting for this parameter, like having to acknowledge non-standard PIC hardware. *iack* should return a non-zero value to indicate that the interrupt has been properly acknowledged. If *iack* is NULL, the default routine will be used instead.

*flags* A set of creation flags affecting the operation. The valid flags are:

- `XN_ISR_SHARED` enables IRQ-sharing with other interrupt objects.
- `XN_ISR_EDGE` is an additional flag need to be set together with `XN_ISR_SHARED` to enable IRQ-sharing of edge-triggered interrupts.

### Returns

0 is returned on success. Otherwise, `-EINVAL` is returned if *irq* is not a valid interrupt number.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

Rescheduling: never.

Referenced by `rtm_irq_request()`, and `xnpod_enable_timesource()`.

## 5.15 Lightweight key-to-object mapping service

### Files

- file [map.h](#)
- file [map.c](#)

### Functions

- `xnmap_t * xnmap_create` (int nkeys, int reserve, int offset)  
*Create a map.*
- `void xnmap_delete` (xnmap\_t \*map)  
*Delete a map.*
- `int xnmap_enter` (xnmap\_t \*map, int key, void \*objaddr)  
*Index an object into a map.*
- `int xnmap_remove` (xnmap\_t \*map, int key)  
*Remove an object reference from a map.*
- `static void * xnmap_fetch_nocheck` (xnmap\_t \*map, int key)  
*Search an object into a map - unchecked form.*
- `static void * xnmap_fetch` (xnmap\_t \*map, int key)  
*Search an object into a map.*

### 5.15.1 Detailed Description

A map is a simple indexing structure which associates unique integer keys with pointers to objects. The current implementation supports reservation, for naming/indexing the real-time objects skins create, either on a fixed, user-provided integer (i.e. a reserved key value), or by drawing the next available key internally if the caller did not specify any fixed key. For instance, in some given map, the key space ranging from 0 to 255 could be reserved for fixed keys, whilst the range from 256 to 511 could be available for drawing free keys dynamically.

A maximum of 1024 unique keys per map is supported on 32bit machines.

(This implementation should not be confused with C++ STL maps, which are dynamically expandable and allow arbitrary key types; Xenomai maps don't).

## 5.15.2 Function Documentation

### 5.15.2.1 `xnmap_t* xnmap_create ( int nkeys, int reserve, int offset )`

Create a map.

Allocates a new map with the specified addressing capabilities. The memory is obtained from the Xenomai system heap.

#### Parameters

*nkeys* The maximum number of unique keys the map will be able to hold. This value cannot exceed the static limit represented by `XNMAP_MAX_KEYS`, and must be a power of two.

*reserve* The number of keys which should be kept for reservation within the index space. Reserving a key means to specify a valid key to the `xnmap_enter()` service, which will then attempt to register this exact key, instead of drawing the next available key from the unreserved index space. When reservation is in effect, the unreserved index space will hold key values greater than *reserve*, keeping the low key values for the reserved space. For instance, passing *reserve* = 32 would cause the index range [ 0 .. 31 ] to be kept for reserved keys. When non-zero, *reserve* is rounded to the next multiple of `BITS_PER_LONG`. If *reserve* is zero no reservation will be available from the map.

*offset* The lowest key value `xnmap_enter()` will return to the caller. Key values will be in the range [ 0 + *offset* .. *nkeys* + *offset* - 1 ]. Negative offsets are valid.

#### Returns

the address of the new map is returned on success; otherwise, NULL is returned if *nkeys* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

### 5.15.2.2 `void xnmap_delete ( xnmap_t* map )`

Delete a map.

Deletes a map, freeing any associated memory back to the Xenomai system heap.

#### Parameters

*map* The address of the map to delete.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

### 5.15.2.3 `int xnmap_enter ( xnmap_t * map, int key, void * objaddr )`

Index an object into a map.

Insert a new object into the given map.

#### Parameters

*map* The address of the map to insert into.

*key* The key to index the object on. If this key is within the valid index range [ 0 - offset .. nkeys - offset - 1 ], then an attempt to reserve this exact key is made. If *key* has an out-of-range value lower or equal to 0 - offset - 1, then an attempt is made to draw a free key from the unreserved index space.

*objaddr* The address of the object to index on the key. This value will be returned by a successful call to [xnmap\\_fetch\(\)](#) with the same key.

#### Returns

a valid key is returned on success, either *key* if reserved, or the next free key. Otherwise:

- -EEXIST is returned upon attempt to reserve a busy key.
- -ENOSPC when no more free key is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

### 5.15.2.4 `void xnmap_fetch ( xnmap_t * map, int key ) [inline, static]`

Search an object into a map.

Retrieve an object reference from the given map by its index key.

#### Parameters

*map* The address of the map to retrieve from.

*key* The key to be searched for in the map index.

### Returns

The indexed object address is returned on success, otherwise NULL is returned when *key* is invalid or no object is currently indexed on it.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 5.15.2.5 void xnmap\_fetch\_nocheck ( xnmap\_t \* map, int key ) [inline, static]

Search an object into a map - unchecked form.

Retrieve an object reference from the given map by its index key, but does not perform any sanity check on the provided key.

### Parameters

*map* The address of the map to retrieve from.

*key* The key to be searched for in the map index.

### Returns

The indexed object address is returned on success, otherwise NULL is returned when no object is currently indexed on *key*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.



#### 5.15.2.6 `int xnmap_remove ( xnmap_t * map, int key )`

Remove an object reference from a map.

Removes an object reference from the given map, releasing the associated key.

##### Parameters

*map* The address of the map to remove from.

*key* The key the object reference to be removed is indexed on.

##### Returns

0 is returned on success. Otherwise:

- -ESRCH is returned if *key* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

## 5.16 Real-time pod services.

### Data Structures

- struct [xnpod](#)  
*Real-time pod descriptor.*

### Files

- file [pod.h](#)  
*Real-time pod interface header.*
- file [pod.c](#)  
*Real-time pod services.*

## Functions

- void [\\_\\_xnpod\\_reset\\_thread](#) (struct xnthread \*thread)  
*Reset the thread.*
- int [xnpod\\_init](#) (void)  
*Initialize the core pod.*
- int [xnpod\\_enable\\_timesource](#) (void)  
*Activate the core time source.*
- void [xnpod\\_disable\\_timesource](#) (void)  
*Stop the core time source.*
- void [xnpod\\_shutdown](#) (int xtype)  
*Shutdown the current pod.*
- int [xnpod\\_init\\_thread](#) (struct xnthread \*thread, const struct xnthread\_init\_attr \*attr, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Initialize a new thread.*
- int [xnpod\\_start\\_thread](#) (xnthread\_t \*thread, const struct xnthread\_start\_attr \*attr)  
*Initial start of a newly created thread.*
- void [xnpod\\_stop\\_thread](#) (xnthread\_t \*thread)  
*Stop a thread.*
- void [xnpod\\_delete\\_thread](#) (xnthread\_t \*thread)  
*Delete a thread.*
- void [xnpod\\_abort\\_thread](#) (xnthread\_t \*thread)  
*Abort a thread.*
- xnflags\_t [xnpod\\_set\\_thread\\_mode](#) (xnthread\_t \*thread, xnflags\_t clrmask, xnflags\_t set-mask)  
*Change a thread's control mode.*
- void [xnpod\\_suspend\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \*wchan)  
*Suspend a thread.*
- void [xnpod\\_resume\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask)  
*Resume a thread.*
- int [xnpod\\_unblock\\_thread](#) (xnthread\_t \*thread)  
*Unblock a thread.*
- int [xnpod\\_set\\_thread\\_schedparam](#) (struct xnthread \*thread, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*

- int [xnpod\\_migrate\\_thread](#) (int cpu)  
*Migrate the current thread.*
- void [xnpod\\_dispatch\\_signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- static void [xnpod\\_schedule](#) (void)  
*Rescheduling procedure entry point.*
- int [xnpod\\_handle\\_exception](#) (struct ipipe\_trap\_data \*d)  
*Exception handler.*
- int [xnpod\\_set\\_thread\\_periodic](#) (xnthread\_t \*thread, xnticks\_t idate, xntmode\_t timeout\_mode, xnticks\_t period)  
*Make a thread periodic.*
- int [xnpod\\_wait\\_thread\\_period](#) (unsigned long \*overruns\_r)  
*Wait for the next periodic release point.*
- int [xnpod\\_set\\_thread\\_tslice](#) (struct xnthread \*thread, xnticks\_t quantum)  
*Set thread time-slicing information.*
- int [xnpod\\_add\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Install a nucleus hook.*
- int [xnpod\\_remove\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Remove a nucleus hook.*
- void [xnpod\\_welcome\\_thread](#) (xnthread\_t \*thread, int imask)  
*Thread prologue.*

### 5.16.1 Detailed Description

Real-time pod services.

### 5.16.2 Function Documentation

#### 5.16.2.1 void \_\_xnpod\_reset\_thread ( struct xnthread \* thread )

Reset the thread.

**For internal use only.**

This internal routine resets the state of a thread so that it can be subsequently stopped or restarted.

References `XNLOCK`, `xnpod_resume_thread()`, `xnpod_unblock_thread()`, `XNSUSP`, and `xnsynch_release_all_ownerships()`.

Referenced by `xnpod_stop_thread()`.

### 5.16.2.2 void xnpod\_abort\_thread ( xntthread\_t \* thread )

Abort a thread.

Unconditionally terminates a thread and releases all the nucleus resources it currently holds, regardless of whether the target thread is currently active in kernel or user-space. [xnpod\\_abort\\_thread\(\)](#) should be reserved for use by skin cleanup routines; [xnpod\\_delete\\_thread\(\)](#) should be preferred as the common method for removing threads from a running system.

#### Parameters

*thread* The descriptor address of the terminated thread.

This service forces a call to [xnpod\\_delete\\_thread\(\)](#) for the target thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References XNABORT, XNDORMANT, [xnpod\\_delete\\_thread\(\)](#), and [xnpod\\_suspend\\_thread\(\)](#).

### 5.16.2.3 int xnpod\_add\_hook ( int type, void(\*) (xntthread\_t \*) routine )

Install a nucleus hook.

The nucleus allows to register user-defined routines which get called whenever a specific scheduling event occurs. Multiple hooks can be chained for a single event type, and get called on a FIFO basis.

The scheduling is locked while a hook is executing.

#### Parameters

*type* Defines the kind of hook to install:

- XNHOOK\_THREAD\_START: The user-defined routine will be called on behalf of the starter thread whenever a new thread starts. The descriptor address of the started thread is passed to the routine.
- XNHOOK\_THREAD\_DELETE: The user-defined routine will be called on behalf of the deleter thread whenever a thread is deleted. The descriptor address of the deleted thread is passed to the routine.
- XNHOOK\_THREAD\_SWITCH: The user-defined routine will be called on behalf of the resuming thread whenever a context switch takes place. The descriptor address of the thread which has been switched out is passed to the routine.

#### Parameters

*routine* The address of the user-supplied routine to call.

**Returns**

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- -EINVAL is returned if type is incorrect.
- -ENOMEM is returned if not enough memory is available from the system heap to add the new hook.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

**5.16.2.4 void xnpod\_delete\_thread ( xnthread\_t \* thread )**

Delete a thread.

Terminates a thread and releases all the nucleus resources it currently holds. A thread exists in the system since [xnpod\\_init\\_thread\(\)](#) has been called to create it, so this service must be called in order to destroy it afterwards.

**Parameters**

*thread* The descriptor address of the terminated thread.

The target thread's resources may not be immediately removed if this is an active shadow thread running in user-space. In such a case, the mated Linux task is sent a termination signal instead, and the actual deletion is deferred until the task exit event is called.

The DELETE hooks are called on behalf of the calling context (if any). The information stored in the thread control block remains valid until all hooks have been called.

Self-terminating a thread is allowed. In such a case, this service does not return to the caller.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References `xnsched::curr`, `xnsched::status`, `XNABORT`, `XNCANPND`, `XNDEFKAN`, `XNDORMANT`, `XNMIGRATE`, `XNPEND`, `xnpod_schedule()`, `xnpod_unblock_thread()`, `XNREADY`, `XNROOT`, `xnselector_destroy()`, `xnsynch_forget_sleeper()`, `xnsynch_release_all_ownerships()`, `xntimer_destroy()`, and `XNZOMBIE`.

Referenced by `rtdm_task_init()`, `xnpod_abort_thread()`, and `xnpod_shutdown()`.

### 5.16.2.5 void xnpod\_disable\_timesource ( void )

Stop the core time source.

Releases the hardware timer, and deactivates the system clock.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

References xntimer\_freeze().

Referenced by xnpod\_shutdown().

### 5.16.2.6 void xnpod\_dispatch\_signals ( void )

Deliver pending asynchronous signals to the running thread.

**For internal use only.**

This internal routine checks for the presence of asynchronous signals directed to the running thread, and attempts to start the asynchronous service routine (ASR) if any. Called with nklock locked, interrupts off.

References XNASDI.

Referenced by xnpod\_welcome\_thread(), and xnshadow\_harden().

### 5.16.2.7 int xnpod\_enable\_timesource ( void )

Activate the core time source.

On every architecture, Xenomai directly manages a hardware timer clocked in one-shot mode, to support any number of software timers internally. Timings are always specified as a count of nanoseconds.

The [xnpod\\_enable\\_timesource\(\)](#) service configures the hardware timer chip. Because Xenomai most often interposes on the system timer used by the Linux kernel, a software timer may be started to relay periodic ticks to the host kernel if needed.

#### Returns

0 is returned on success. Otherwise:

- -ENODEV is returned if a failure occurred while configuring the hardware timer.
- -ENOSYS is returned if no active pod exists.

Environments:

This service can be called from:

- Regular Linux kernel context.

Rescheduling: never.

References `xnsched::htimer`, `xnintr_init()`, and `xntimer_start()`.

Referenced by `xnpod_init()`.

#### 5.16.2.8 `int xnpod_handle_exception ( struct ipipe_trap_data * d )`

Exception handler.

This is the handler which is called whenever an exception/fault is caught over the primary domain.

##### Parameters

- d* A pointer to the trap information block received from the pipeline core.

References `xnpod_suspend_thread()`, `xnshadow_relax()`, and `XNSUSP`.

#### 5.16.2.9 `int xnpod_init ( void )`

Initialize the core pod.

Initializes the core interface pod which can subsequently be used to start real-time activities. Once the core pod is active, real-time skins can be stacked over. There can only be a single core pod active in the host environment.

##### Returns

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if the memory manager fails to initialize.

Environments:

This service can be called from:

- Kernel module initialization code

References `xnpod::refcnt`, `xnsched::rootcb`, `xnpod::sched`, `xnpod::status`, `xnpod::tdeleteq`, `xnpod::threadq`, `xnpod::timerlck`, `xnpod::tstartq`, `xnpod::tswitchq`, `xnheap_destroy()`, `xnheap_init()`, `xnheap_set_label()`, `xnpod_enable_timesource()`, and `xnpod_shutdown()`.

#### 5.16.2.10 `int xnpod_init_thread ( struct xnthread * thread, const struct xnthread_init_attr * attr, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param )`

Initialize a new thread.

Initializes a new thread attached to the active pod. The thread is left in an innocuous state until it is actually started by [xnpod\\_start\\_thread\(\)](#).

## Parameters

*thread* The address of a thread descriptor the nucleus will use to store the thread-specific data. This descriptor must always be valid while the thread is active therefore it must be allocated in permanent memory.

## Warning

Some architectures may require the descriptor to be properly aligned in memory; this is an additional reason for descriptors not to be laid in the program stack where alignment constraints might not always be satisfied.

## Parameters

*attr* A pointer to an attribute block describing the initial properties of the new thread. Members of this structure are defined as follows:

- **name:** An ASCII string standing for the symbolic name of the thread. This name is copied to a safe place into the thread descriptor. This name might be used in various situations by the nucleus for issuing human-readable diagnostic messages, so it is usually a good idea to provide a sensible value here. NULL is fine though and means "anonymous".
- **flags:** A set of creation flags affecting the operation. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the created thread:
- **XNSUSP** creates the thread in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod\\_resume\\_thread\(\)](#) service for its execution to actually begin, additionally to issuing [xnpod\\_start\\_thread\(\)](#) for it. This flag can also be specified when invoking [xnpod\\_start\\_thread\(\)](#) as a starting mode.
- **XNFPU** (enable FPU) tells the nucleus that the new thread will use the floating-point unit. In such a case, the nucleus will handle the FPU context save/restore ops upon thread switches at the expense of a few additional cycles per context switch. By default, a thread is not expected to use the FPU. This flag is simply ignored when the nucleus runs on behalf of a userspace-based real-time control layer since the FPU management is always active if present.
- **stacksize:** The size of the stack (in bytes) for the new thread. If zero is passed, the nucleus will use a reasonable pre-defined size depending on the underlying real-time control layer.
- **ops:** A pointer to a structure defining the class-level operations available for this thread. Fields from this structure must have been set appropriately by the caller.

## Parameters

*sched\_class* The initial scheduling class the new thread should be assigned to.

*sched\_param* The initial scheduling parameters to set for the new thread; *sched\_param* must be valid within the context of *sched\_class*.

## Returns

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:



- -EINVAL is returned if *attr->flags* has invalid bits set.
- -ENOMEM is returned if not enough memory is available from the system heap to create the new thread's stack.

Side-effect: This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

References XNDORMANT, XNFPU, *xnpod\_suspend\_thread()*, XNSHADOW, and XNSUSP.

Referenced by *pthread\_create()*, and *rtdm\_task\_init()*.

#### 5.16.2.11 *int xnpod\_migrate\_thread ( int *cpu* )*

Migrate the current thread.

This call makes the current thread migrate to another CPU if its affinity allows it.

##### Parameters

*cpu* The destination CPU.

##### Return values

- 0 if the thread could migrate ;
- -EPERM if the calling context is asynchronous, or the current thread affinity forbids this migration ;
- -EBUSY if the scheduler is locked.

References *xnpod\_schedule()*.

#### 5.16.2.12 *int xnpod\_remove\_hook ( int *type*, void(\*)(*xnthread\_t* \*) *routine* )*

Remove a nucleus hook.

This service removes a nucleus hook previously registered using [xnpod\\_add\\_hook\(\)](#).

##### Parameters

- *type* Defines the kind of hook to remove among XNHOOK\_THREAD\_START, XNHOOK\_THREAD\_DELETE and XNHOOK\_THREAD\_SWITCH.
- *routine* The address of the user-supplied routine to remove.

## Returns

0 is returned on success. Otherwise, -EINVAL is returned if type is incorrect or if the routine has never been registered before.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

### 5.16.2.13 void xnpod\_resume\_thread ( xnthread\_t \* thread, xnflags\_t mask )

Resume a thread.

Resumes the execution of a thread previously suspended by one or more calls to [xnpod\\_suspend\\_thread\(\)](#). This call removes a suspensive condition affecting the target thread. When all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

## Parameters

*thread* The descriptor address of the resumed thread.

*mask* The suspension mask specifying the suspensive condition to remove from the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag removes the explicit suspension condition. This condition might be additive to the XNPEND condition.
- XNDELAY. This flag removes the counted delay wait condition.
- XNPEND. This flag removes the resource wait condition. If a watchdog is armed, it is automatically disarmed by this call. Unlike the two previous conditions, only the current thread can set this condition for itself, i.e. no thread can force another one to pend on a resource.

When the thread is eventually resumed by one or more calls to [xnpod\\_resume\\_thread\(\)](#), the caller of [xnpod\\_suspend\\_thread\(\)](#) in the awakened thread that suspended itself should check for the following bits in its own information mask to determine what caused its wake up:

- XNRMID means that the caller must assume that the pended synchronization object has been destroyed (see [xnsynch\\_flush\(\)](#)).
- XNTIMEO means that the delay elapsed, or the watchdog went off before the corresponding synchronization object was signaled.
- XNBREAK means that the wait has been forcibly broken by a call to [xnpod\\_unblock\\_thread\(\)](#).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNDELAY, XNHELD, XNPEND, XNREADY, xnsynch\_forget\_sleeper(), and xntimer\_stop().

Referenced by \_\_xnpod\_reset\_thread(), xnpod\_start\_thread(), xnpod\_unblock\_thread(), xnsynch\_flush(), xnsynch\_wakeup\_one\_sleeper(), and xnsynch\_wakeup\_this\_sleeper().

#### 5.16.2.14 void xnpod\_schedule ( void ) [inline, static]

Rescheduling procedure entry point.

This is the central rescheduling routine which should be called to validate and apply changes which have previously been made to the nucleus scheduling state, such as suspending, resuming or changing the priority of threads. This call first determines if a thread switch should take place, and performs it as needed. [xnpod\\_schedule\(\)](#) schedules out the current thread if:

- the current thread is now blocked or deleted.
- a runnable thread from a higher priority scheduling class is waiting for the CPU.
- the current thread does not lead the runnable threads from its own scheduling class (e.g. round-robin in the RT class).

The nucleus implements a lazy rescheduling scheme so that most of the services affecting the threads state MUST be followed by a call to the rescheduling procedure for the new scheduling state to be applied. In other words, multiple changes on the scheduler state can be done in a row, waking threads up, blocking others, without being immediately translated into the corresponding context switches, like it would be necessary would it appear that a higher priority thread than the current one became runnable for instance. When all changes have been applied, the rescheduling procedure is then called to consider those changes, and possibly replace the current thread by another one.

As a notable exception to the previous principle however, every action which ends up suspending or deleting the current thread begets an immediate call to the rescheduling procedure on behalf of the service causing the state transition. For instance, self-suspension, self-destruction, or sleeping on a synchronization object automatically leads to a call to the rescheduling procedure, therefore the caller does not need to explicitly issue [xnpod\\_schedule\(\)](#) after such operations.

The rescheduling procedure always leads to a null-effect if it is called on behalf of an ISR or callout. Any outstanding scheduler lock held by the outgoing thread will be restored when the thread is scheduled back in.

Calling this procedure with no applicable context switch pending is harmless and simply leads to a null-effect.

Side-effects:

- If an asynchronous service routine exists, the pending asynchronous signals are delivered to a resuming thread or on behalf of the caller before it returns from the procedure if no context switch has taken place. This behaviour can be disabled by setting the XNASDI flag in the thread's status mask by calling [xnpod\\_set\\_thread\\_mode\(\)](#).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine, although this leads to a no-op.
- Kernel-based task
- User-space task

#### Note

The switch hooks are called on behalf of the resuming thread.

References `xnsched::lflags`, and `xnsched::status`.

Referenced by `pthread_cancel()`, `pthread_set_mode_np()`, `pthread_setschedparam()`, `pthread_setschedparam_ex()`, `rtdm_event_signal()`, `rtdm_sem_up()`, `xnpod_delete_thread()`, `xnpod_migrate_thread()`, `xnpod_shutdown()`, `xnpod_start_thread()`, `xnpod_stop_thread()`, `xnpod_suspend_thread()`, `xnregistry_enter()`, `xnregistry_put()`, `xnselect_bind()`, and `xnselect_destroy()`.

#### 5.16.2.15 `xnflags_t xnpod_set_thread_mode ( xnthread_t * thread, xnflags_t clrmask, xnflags_t setmask )`

Change a thread's control mode.

Change the control mode of a given thread. The control mode affects the behaviour of the nucleus regarding the specified thread.

#### Parameters

*thread* The descriptor address of the affected thread.

*clrmask* Clears the corresponding bits from the control field before *setmask* is applied. The scheduler lock held by the current thread can be forcibly released by passing the XNLOCK bit in this mask. In this case, the lock nesting count is also reset to zero.

*setmask* The new thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the thread:

- XNLOCK causes the thread to lock the scheduler. The target thread will have to call the `xnpod_unlock_sched()` service to unlock the scheduler or clear the XNLOCK bit forcibly using this service. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- XNASDI disables the asynchronous signal handling for this thread. See [xnpod\\_schedule\(\)](#) for more on this.

Environments:

This service can be called from:

- Kernel-based task
- User-space task in primary mode.

Rescheduling: never, therefore, the caller should reschedule if XNLOCK has been passed into *clrmask*.

References XNLOCK.

Referenced by `pthread_set_mode_np()`.

#### 5.16.2.16 `int xnpod_set_thread_periodic ( xntthread_t * thread, xnticks_t idate, xntmode_t timeout_mode, xnticks_t period )`

Make a thread periodic.

Make a thread periodic by programming its first release point and its period in the processor time line. Subsequent calls to `xnpod_wait_thread_period()` will delay the thread until the next periodic release point in the processor timeline is reached.

##### Parameters

*thread* The descriptor address of the affected thread. This thread is immediately delayed until the first periodic release point is reached.

*idate* The initial (absolute) date of the first release point, expressed in nanoseconds. The affected thread will be delayed by the first call to `xnpod_wait_thread_period()` until this point is reached. If *idate* is equal to XN\_INFINITE, the current system date is used, and no initial delay takes place. In the latter case, *timeout\_mode* is not considered and can have any valid value.

*timeout\_mode* The mode of the *idate* parameter. It can either be set to XN\_ABSOLUTE or XN\_REALTIME with *idate* different from XN\_INFINITE (see also `xntimer_start()`).

*period* The period of the thread, expressed in nanoseconds. As a side-effect, passing XN\_INFINITE attempts to stop the thread's periodic timer; in the latter case, the routine always exits successfully, regardless of the previous state of this timer.

##### Returns

0 is returned upon success. Otherwise:

- -ETIMEDOUT is returned *idate* is different from XN\_INFINITE and represents a date in the past.
- -EINVAL is returned if *period* is different from XN\_INFINITE but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`. -EINVAL is also returned if *timeout\_mode* is not compatible with *idate*, such as XN\_RELATIVE with *idate* different from XN\_INFINITE.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

- User-space task

Rescheduling: none.

References `xntimer_start()`, and `xntimer_stop()`.

Referenced by `pthread_make_periodic_np()`, and `rtdm_task_init()`.

**5.16.2.17** `int xnpod_set_thread_schedparam ( struct xnthread * thread, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param )`

Change the base scheduling parameters of a thread.

Changes the base scheduling policy and parameters of a thread. If the thread is currently blocked, waiting in priority-pending mode (`XNSYNCH_PRIO`) for a synchronization object to be signaled, the nucleus will attempt to reorder the object's wait queue so that it reflects the new sleeper's priority, unless the `XNSYNCH_DREORD` flag has been set for the pended object.

#### Parameters

*thread* The descriptor address of the affected thread.

*sched\_class* The new scheduling class the thread should be assigned to.

*sched\_param* The scheduling parameters to set for the thread; *sched\_param* must be valid within the context of *sched\_class*.

It is absolutely required to use this service to change a thread priority, in order to have all the needed housekeeping chores correctly performed. i.e. Do *not* call `xnsched_set_policy()` directly or worse, change the `thread.cprio` field by hand in any case.

#### Returns

0 is returned on success. Otherwise, a negative error code indicates the cause of a failure that happened in the scheduling class implementation for *sched\_class*. Invalid parameters passed into *sched\_param* are common causes of error.

Side-effects:

- This service does not call the rescheduling procedure but may affect the state of the runnable queue for the previous and new scheduling classes.
- Assigning the same scheduling class and parameters to a running or ready thread moves it to the end of the runnable queue, thus causing a manual round-robin.
- If the thread is a user-space shadow, this call propagates the request to the mated Linux task.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine

- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `pthread_setschedparam()`, and `pthread_setschedparam_ex()`.

#### 5.16.2.18 `int xnpod_set_thread_tslice ( struct xnthread * thread, xnticks_t quantum )`

Set thread time-slicing information.

Update the time-slicing information for a given thread. This service enables or disables round-robin scheduling for the thread, depending on the value of *quantum*. By default, times-slicing is disabled for a new thread initialized by a call to `xnpod_init_thread()`.

##### Parameters

*thread* The descriptor address of the affected thread.

*quantum* The time quantum assigned to the thread expressed in nanoseconds. If *quantum* is different from `XN_INFINITE`, the time-slice for the thread is set to that value and its current time credit is refilled (i.e. the thread is given a full time-slice to run next). Otherwise, if *quantum* equals `XN_INFINITE`, time-slicing is stopped for that thread.

##### Returns

0 is returned upon success. Otherwise:

- `-EINVAL` is returned if *quantum* is not `XN_INFINITE`, and the base scheduling class of the target thread does not support time-slicing.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

References `XNRRB`, `xntimer_start()`, and `xntimer_stop()`.

Referenced by `pthread_create()`, `pthread_setschedparam()`, and `pthread_setschedparam_ex()`.

#### 5.16.2.19 `void xnpod_shutdown ( int xtype )`

Shutdown the current pod.

Forcibly shutdowns the active pod. All existing nucleus threads (but the root one) are terminated, and the system heap is freed.

##### Parameters

*xtype* An exit code passed to the host environment who started the nucleus. Zero is always interpreted as a successful return.

The nucleus never calls this routine directly. Skins should provide their own shutdown handlers which end up calling [xnpod\\_shutdown\(\)](#) after their own housekeeping chores have been carried out.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References [xnheap\\_destroy\(\)](#), [xnpod\\_delete\\_thread\(\)](#), [xnpod\\_disable\\_timesource\(\)](#), [xnpod\\_schedule\(\)](#), and [XNROOT](#).

Referenced by [xnpod\\_init\(\)](#).

#### 5.16.2.20 `int xnpod_start_thread ( xnthread_t * thread, const struct xnthread_start_attr * attr )`

Initial start of a newly created thread.

Starts a (newly) created thread, scheduling it for the first time. This call releases the target thread from the [XNDORMANT](#) state. This service also sets the initial mode and interrupt mask for the new thread.

##### Parameters

*thread* The descriptor address of the affected thread which must have been previously initialized by the [xnpod\\_init\\_thread\(\)](#) service.

*attr* A pointer to an attribute block describing the execution properties of the new thread. Members of this structure are defined as follows:

- **mode**: The initial thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the started thread:
- **XNLOCK** causes the thread to lock the scheduler when it starts. The target thread will have to call the [xnpod\\_unlock\\_sched\(\)](#) service to unlock the scheduler. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- **XNASDI** disables the asynchronous signal handling for this thread. See [xnpod\\_schedule\(\)](#) for more on this.
- **XNSUSP** makes the thread start in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod\\_resume\\_thread\(\)](#) service for its execution to actually begin.
- **imask**: The interrupt mask that should be asserted when the thread starts. The processor interrupt state will be set to the given value when the thread starts running. The interpretation of this value might be different across real-time layers, but a non-zero value should always mark an interrupt masking in effect (e.g. [local\\_irq\\_disable\(\)](#)). Conversely, a zero value should always mark a fully preemptible state regarding interrupts (e.g. [local\\_irq\\_enable\(\)](#)).



- **affinity**: The processor affinity of this thread. Passing `XNPOD_ALL_CPUS` or an empty affinity set means "any cpu".
- **entry**: The address of the thread's body routine. In other words, it is the thread entry point.
- **cookie**: A user-defined opaque cookie the nucleus will pass to the emerging thread as the sole argument of its entry point.

The START hooks are called on behalf of the calling context (if any).

#### Return values

- `0` if *thread* could be started ;
- `-EBUSY` if *thread* was not dormant or stopped ;
- `-EINVAL` if the value of *attr->affinity* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References `XNDORMANT`, `xnpod_resume_thread()`, `xnpod_schedule()`, `XNSHADOW`, `XN-STARTED`, and `XNSUSP`.

Referenced by `rtm_task_init()`, and `xnshadow_map()`.

#### 5.16.2.21 void xnpod\_stop\_thread ( xnthread\_t \* thread )

Stop a thread.

Stop a previously started thread. The thread is put back into the dormant state; however, it is not deleted from the system.

#### Parameters

*thread* The descriptor address of the affected thread which must have been previously started by the [xnpod\\_start\\_thread\(\)](#) service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References `__xnpod_reset_thread()`, `XNDORMANT`, `xnpod_schedule()`, `xnpod_suspend_thread()`, and `XNROOT`.

### 5.16.2.22 void xnpod\_suspend\_thread ( xnthread\_t \* thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \* wchan )

Suspend a thread.

Suspends the execution of a thread according to a given suspensive condition. This thread will not be eligible for scheduling until all the pending suspensive conditions set by this service are removed by one or more calls to [xnpod\\_resume\\_thread\(\)](#).

#### Parameters

*thread* The descriptor address of the suspended thread.

*mask* The suspension mask specifying the suspensive condition to add to the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag forcibly suspends a thread, regardless of any resource to wait for. A reverse call to [xnpod\\_resume\\_thread\(\)](#) specifying the XNSUSP bit must be issued to remove this condition, which is cumulative with other suspension bits. *wchan* should be NULL when using this suspending mode.
- XNDELAY. This flag denotes a counted delay wait (in ticks) which duration is defined by the value of the timeout parameter.
- XNPEND. This flag denotes a wait for a synchronization object to be signaled. The *wchan* argument must point to this object. A timeout value can be passed to bound the wait. This suspending mode should not be used directly by the client interface, but rather through the [xnsynch\\_sleep\\_on\(\)](#) call.

#### Parameters

*timeout* The timeout which may be used to limit the time the thread spends on a resource. This value is a wait time given in nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing XN\_INFINITE and setting *timeout\_mode* to XN\_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer. If the current operation mode of the system timer is oneshot and *timeout* elapses before [xnpod\\_suspend\\_thread\(\)](#) has completed, then the target thread will not be suspended, and this routine leads to a null effect.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

*wchan* The address of a pending resource. This parameter is used internally by the synchronization object implementation code to specify on which object the suspended thread spends. NULL is a legitimate value when this parameter does not apply to the current suspending mode (e.g. XNSUSP).

#### Note

If the target thread is a shadow which has received a Linux-originated signal, then this service immediately exits without suspending the thread, but raises the XNBREAK condition in its information mask.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread suspends itself.

References `xnsched::curr`, `XNBREAK`, `XNDELAY`, `XNDORMANT`, `XNHELD`, `XNKICKED`, `xnpod_schedule()`, `XNREADY`, `XNRELAX`, `XNRMID`, `XNROBBED`, `XNROOT`, `XNSHADOW`, `XNSUSP`, `xnsynch_forget_sleeper()`, `XNTIMEO`, `xntimer_start()`, and `XNWAKEN`.

Referenced by `xnpod_abort_thread()`, `xnpod_handle_exception()`, `xnpod_init_thread()`, `xnpod_stop_thread()`, `xnpod_wait_thread_period()`, `xnshadow_map()`, `xnshadow_relax()`, `xnsynch_acquire()`, and `xnsynch_sleep_on()`.

#### 5.16.2.23 `int xnpod_unblock_thread ( xntthread_t * thread )`

Unblock a thread.

Breaks the thread out of any wait it is currently in. This call removes the `XNDELAY` and `XNPEND` suspensive conditions previously put by `xnpod_suspend_thread()` on the target thread. If all suspensive conditions are gone, the thread is left in a `READY` state at which point it becomes eligible anew for scheduling.

##### Parameters

*thread* The descriptor address of the unblocked thread.

This call neither releases the thread from the `XNSUSP`, `XNRELAX`, `XNDORMANT` or `XNHELD` suspensive conditions.

When the thread resumes execution, the `XNBREAK` bit is set in the unblocked thread's information mask. Unblocking a non-blocked thread is perfectly harmless.

##### Returns

non-zero is returned if the thread was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References `XNBREAK`, `XNDELAY`, `XNPEND`, and `xnpod_resume_thread()`.

Referenced by `__xnpod_reset_thread()`, `pthread_cancel()`, and `xnpod_delete_thread()`.

#### 5.16.2.24 `int xnpod_wait_thread_period ( unsigned long * overruns_r )`

Wait for the next periodic release point.

Make the current thread wait for the next periodic release point in the processor time line.

##### Parameters

*overruns\_r* If non-NULL, *overruns\_r* must be a pointer to a memory location which will be written with the count of pending overruns. This value is copied only when `xnpod_wait_thread_period()` returns -ETIMEDOUT or success; the memory location remains unmodified otherwise. If NULL, this count will never be copied back.

##### Returns

0 is returned upon success; if *overruns\_r* is valid, zero is copied to the pointed memory location. Otherwise:

- -EWOULDBLOCK is returned if `xnpod_set_thread_periodic()` has not previously been called for the calling thread.
- -EINTR is returned if `xnpod_unblock_thread()` has been called for the waiting thread before the next periodic release point has been reached. In this case, the overrun counter is reset too.
- -ETIMEDOUT is returned if the timer has overrun, which indicates that one or more previous release points have been missed by the calling thread. If *overruns\_r* is valid, the count of pending overruns is copied to the pointed memory location.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always, unless the current release point has already been reached. In the latter case, the current thread immediately returns from this service without being delayed.

References XNBREAK, XNDELAY, `xnpod_suspend_thread()`, and `xntimer_get_overruns()`.

#### 5.16.2.25 `void xnpod_welcome_thread ( xnthread_t * thread, int imask )`

Thread prologue.

**For internal use only.**

This internal routine is called on behalf of a (re)starting thread's prologue before the user entry point is invoked. This call is reserved for internal housekeeping chores and cannot be inlined.

Entered with `nklock` locked, `irqs` off.

References XNLOCK, and `xnpod_dispatch_signals()`.

## 5.17 Registry services.

### Files

- file [registry.h](#)

*This file is part of the Xenomai project.*

- file [registry.c](#)

*This file is part of the Xenomai project.*

### Functions

- int [xnregistry\\_enter](#) (const char \*key, void \*objaddr, xnhandle\_t \*phandle, struct xnnode \*pnode)

*Register a real-time object.*

- int [xnregistry\\_bind](#) (const char \*key, xnticks\_t timeout, int timeout\_mode, xnhandle\_t \*phandle)

*Bind to a real-time object.*

- int [xnregistry\\_remove](#) (xnhandle\_t handle)

*Forcibly unregister a real-time object.*

- int [xnregistry\\_remove\\_safe](#) (xnhandle\_t handle, xnticks\_t timeout)

*Unregister an idle real-time object.*

- void \* [xnregistry\\_get](#) (xnhandle\_t handle)

*Find and lock a real-time object into the registry.*

- u\_long [xnregistry\\_put](#) (xnhandle\_t handle)

*Unlock a real-time object from the registry.*

- void \* [xnregistry\\_fetch](#) (xnhandle\_t handle)

*Find a real-time object into the registry.*

### 5.17.1 Detailed Description

The registry provides a mean to index real-time object descriptors created by Xenomai skins on unique alphanumeric keys. When labeled this way, a real-time object is globally exported; it can be searched for, and its descriptor returned to the caller for further use; the latter operation is called a "binding". When no object has been registered under the given name yet, the registry can be asked to set up a rendez-vous, blocking the caller until the object is eventually registered.

## 5.17.2 Function Documentation

**5.17.2.1** `int xnregistry_bind ( const char * key, xnticks_t timeout, int timeout_mode, xnhandle_t * phandle )`

Bind to a real-time object.

This service retrieves the registry handle of a given object identified by its key. Unless otherwise specified, this service will block the caller if the object is not registered yet, waiting for such registration to occur.

### Parameters

*key* A valid NULL-terminated string which identifies the object to bind to.

*timeout* The timeout which may be used to limit the time the thread wait for the object to be registered. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic (XN\_ABSOLUTE), or absolute adjustable (XN\_REALTIME) depending on *timeout\_mode*. Passing XN\_INFINITE **and** setting *timeout\_mode* to XN\_RELATIVE specifies an unbounded wait. Passing XN\_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry. All other values are used as a wait limit.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

*phandle* A pointer to a memory location which will be written upon success with the generic handle defined by the registry for the retrieved object. Contents of this memory is undefined upon failure.

### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* is NULL.
- -EINTR is returned if [xnpod\\_unblock\\_thread\(\)](#) has been called for the waiting thread before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to XN\_NONBLOCK and the searched object is not registered on entry. As a special exception, this error is also returned if this service should block, but was called from a context which cannot sleep (e.g. interrupt, non-realtime or scheduler locked).
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN\_NONBLOCK.
- Kernel-based thread.

Rescheduling: always unless the request is immediately satisfied or *timeout* specifies a non-blocking operation.

References XNBREAK, xnsynch\_sleep\_on(), and XNTIMEO.

#### 5.17.2.2 `int xnregistry_enter ( const char * key, void * objaddr, xnhandle_t * phandle, struct xnpnode * pnode )`

Register a real-time object.

This service allocates a new registry slot for an associated object, and indexes it by an alphanumeric key for later retrieval.

##### Parameters

**key** A valid NULL-terminated string by which the object will be indexed and later retrieved in the registry. Since it is assumed that such key is stored into the registered object, it will *\*not\** be copied but only kept by reference in the registry. Pass an empty string if the object shall only occupy a registry slot for handle-based lookups.

**objaddr** An opaque pointer to the object to index by *key*.

**phandle** A pointer to a generic handle defined by the registry which will uniquely identify the indexed object, until the latter is unregistered using the [xnregistry\\_remove\(\)](#) service.

**pnode** A pointer to an optional /proc node class descriptor. This structure provides the information needed to export all objects from the given class through the /proc filesystem, under the /proc/xenomai/registry entry. Passing NULL indicates that no /proc support is available for the newly registered object.

##### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *objaddr* are NULL, or if *key* contains an invalid '/' character.
- -ENOMEM is returned if the system fails to get enough dynamic memory from the global real-time heap in order to register the object.
- -EEXIST is returned if the *key* is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: possible.

References xnpod\_schedule(), and xnsynch\_init().

### 5.17.2.3 void\* xnregistry\_fetch ( xnhandle\_t handle )

Find a real-time object into the registry.

This service retrieves an object from its handle into the registry and returns the memory address of its descriptor.

#### Parameters

*handle* The generic handle of the object to fetch. If XNOBJECT\_SELF is passed, the object is the calling Xenomai thread.

#### Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT\_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT\_SELF.
- Kernel-based thread

Rescheduling: never.

### 5.17.2.4 void\* xnregistry\_get ( xnhandle\_t handle )

Find and lock a real-time object into the registry.

This service retrieves an object from its handle into the registry and prevents its removal atomically. A locking count is tracked, so that [xnregistry\\_get\(\)](#) and [xnregistry\\_put\(\)](#) must be used in pair.

#### Parameters

*handle* The generic handle of the object to find and lock. If XNOBJECT\_SELF is passed, the object is the calling Xenomai thread.

#### Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT\_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT\_SELF.
- Kernel-based thread.

Rescheduling: never.



#### 5.17.2.5 `u_long xnregistry_put ( xnhandle_t handle )`

Unlock a real-time object from the registry.

This service decrements the lock count of a registered object previously locked by a call to `xnregistry_get()`. The object is actually unlocked from the registry when the locking count falls down to zero, thus waking up any thread currently blocked on `xnregistry_remove()` for unregistering it.

##### Parameters

*handle* The generic handle of the object to unlock. If `XNOBJECT_SELF` is passed, the object is the calling Xenomai thread.

##### Returns

The decremented lock count is returned upon success. Zero is also returned if *handle* does not reference a registered object, or if *handle* is equal to `XNOBJECT_SELF` but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from `XNOBJECT_SELF`.
- Kernel-based thread

Rescheduling: possible if the lock count falls down to zero and some thread is currently waiting for the object to be unlocked.

References `xnpod_schedule()`, and `xnsynch_flush()`.

#### 5.17.2.6 `int xnregistry_remove ( xnhandle_t handle )`

Forcibly unregister a real-time object.

This service forcibly removes an object from the registry. The removal is performed regardless of the current object's locking status.

##### Parameters

*handle* The generic handle of the object to remove.

##### Returns

0 is returned upon success. Otherwise:

- `-ESRCH` is returned if *handle* does not reference a registered object.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: never.

Referenced by `xnregistry_remove_safe()`.

#### 5.17.2.7 `int xnregistry_remove_safe ( xnhandle_t handle, xnticks_t timeout )`

Unregister an idle real-time object.

This service removes an object from the registry. The caller might sleep as a result of waiting for the target object to be unlocked prior to the removal (see [xnregistry\\_put\(\)](#)).

##### Parameters

*handle* The generic handle of the object to remove.

*timeout* If the object is locked on entry, *param* gives the number of nanoseconds to wait for the unlocking to occur. Passing `XN_INFINITE` causes the caller to block indefinitely until the object is unlocked. Passing `XN_NONBLOCK` causes the service to return immediately without waiting if the object is locked on entry.

##### Returns

0 is returned upon success. Otherwise:

- `-ESRCH` is returned if *handle* does not reference a registered object.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `XN_NONBLOCK` and the object is locked on entry.
- `-EBUSY` is returned if *handle* refers to a locked object and the caller could not sleep until it is unlocked.
- `-ETIMEDOUT` is returned if the object cannot be removed within the specified amount of time.
- `-EINTR` is returned if [xnpod\\_unblock\\_thread\(\)](#) has been called for the calling thread waiting for the object to be unlocked.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to `XN_NONBLOCK`.
- Kernel-based thread.

Rescheduling: possible if the object to remove is currently locked and the calling context can sleep.

References `XNBREAK`, `xnregistry_remove()`, `xnsynch_sleep_on()`, and `XNTIMEO`.

## 5.18 File descriptors events multiplexing services.

### Files

- file [select.h](#)  
*file descriptors events multiplexing header.*
- file [select.c](#)  
*file descriptors events multiplexing.*

### Functions

- void [xnselect\\_init](#) (struct xnselect \*select\_block)  
*Initialize a struct xnselect structure.*
- int [xnselect\\_bind](#) (struct xnselect \*select\_block, struct xnselect\_binding \*binding, struct xnselector \*selector, unsigned type, unsigned index, unsigned state)  
*Bind a file descriptor (represented by its xnselect structure) to a selector block.*
- static int [xnselect\\_signal](#) (struct xnselect \*select\_block, unsigned state)  
*Signal a file descriptor state change.*
- void [xnselect\\_destroy](#) (struct xnselect \*select\_block)  
*Destroy the xnselect structure associated with a file descriptor.*
- int [xnselector\\_init](#) (struct xnselector \*selector)  
*Initialize a selector structure.*
- int [xnselect](#) (struct xnselector \*selector, fd\_set \*out\_fds[XNSELECT\_MAX\_TYPES], fd\_set \*in\_fds[XNSELECT\_MAX\_TYPES], int nfds, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.*
- void [xnselector\\_destroy](#) (struct xnselector \*selector)  
*Destroy a selector block.*

#### 5.18.1 Detailed Description

File descriptors events multiplexing services.

This module implements the services needed for implementing the posix "select" service, or any other events multiplexing services.

Following the implementation of the posix select service, this module defines three types of events:

- `XNSELECT_READ` meaning that a file descriptor is ready for reading;
- `XNSELECT_WRITE` meaning that a file descriptor is ready for writing;

- `XNSELECT_EXCEPT` meaning that a file descriptor received an exceptional event.

It works by defining two structures:

- a *struct xnselect* structure, which should be added to every file descriptor for every event type (read, write, or except);
- a *struct xnselector* structure, the selection structure, passed by the thread calling the `xnselect` service, where this service does all its housekeeping.

## 5.18.2 Function Documentation

**5.18.2.1** `int xnselect ( struct xnselector * selector, fd_set * out_fds[XNSELECT_MAX_TYPES], fd_set * in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode )`

Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.

### Parameters

*selector* structure to check for pending events

*out\_fds* The set of descriptors with pending events if a strictly positive number is returned, or the set of descriptors not yet bound if `-ECHRNG` is returned;

*in\_fds* the set of descriptors which events should be checked

*nfds* the highest-numbered descriptor in any of the *in\_fds* sets, plus 1;

*timeout* the timeout, whose meaning depends on *timeout\_mode*, note that `xnselect()` pass *timeout* and *timeout\_mode* unchanged to `xnsynch_sleep_on`, so passing a relative value different from `XN_INFINITE` as a timeout with *timeout\_mode* set to `XN_RELATIVE`, will cause a longer sleep than expected if the sleep is interrupted.

*timeout\_mode* the mode of *timeout*.

### Return values

`-EINVAL` if *nfds* is negative;

`-ECHRNG` if some of the descriptors passed in *in\_fds* have not yet been registered with `xnselect_bind()`, *out\_fds* contains the set of such descriptors;

`-EINTR` if *xnselect* was interrupted while waiting;

`0` in case of timeout.

*the* number of file descriptors having received an event.

References `XNBREAK`, `xnsynch_sleep_on()`, and `XNTIMEO`.

**5.18.2.2** `int xnselect_bind ( struct xnselect * select_block, struct xnselect_binding * binding, struct xnselector * selector, unsigned type, unsigned index, unsigned state )`

Bind a file descriptor (represented by its *xnselect* structure) to a selector block.

### Parameters

*select\_block* pointer to the *struct xnselect* to be bound;

*binding* pointer to a newly allocated (using `xnmalloc`) *struct xnselect\_binding*;  
*selector* pointer to the selector structure;  
*type* type of events (`XNSELECT_READ`, `XNSELECT_WRITE`, or `XNSELECT_EXCEPT`);  
*index* index of the file descriptor (represented by *select\_block*) in the bit fields used by the *selector* structure;  
*state* current state of the file descriptor>.

*select\_block* must have been initialized with `xnselect_init()`, the *xnselector* structure must have been initialized with `xnselector_init()`, *binding* may be uninitialized.

This service must be called with `nklock` locked, `irqs` off. For this reason, the *binding* parameter must have been allocated by the caller outside the locking section.

### Return values

`-EINVAL` if *type* or *index* is invalid;  
`0` otherwise.

References `xnpod_schedule()`.

Referenced by `rtdm_event_select_bind()`, and `rtdm_sem_select_bind()`.

#### 5.18.2.3 void xnselect\_destroy ( struct xnselect \* select\_block )

Destroy the *xnselect* structure associated with a file descriptor.

Any binding with a *xnselector* block is destroyed.

### Parameters

*select\_block* pointer to the *xnselect* structure associated with a file descriptor

References `xnpod_schedule()`.

#### 5.18.2.4 void xnselect\_init ( struct xnselect \* select\_block )

Initialize a *struct xnselect* structure.

This service must be called to initialize a *struct xnselect* structure before it is bound to a selector by the means of `xnselect_bind()`.

### Parameters

*select\_block* pointer to the *xnselect* structure to be initialized

Referenced by `rtdm_event_init()`, and `rtdm_sem_init()`.

#### 5.18.2.5 static int xnselect\_signal ( struct xnselect \* select\_block, unsigned state ) [inline, static]

Signal a file descriptor state change.

**Parameters**

*select\_block* pointer to an *xnselect* structure representing the file descriptor whose state changed;

*state* new value of the state.

**Return values**

*1* if rescheduling is needed;

*0* otherwise.

Referenced by `rtdm_event_clear()`, `rtdm_event_signal()`, `rtdm_event_timedwait()`, `rtdm_sem_timeddown()`, and `rtdm_sem_up()`.

**5.18.2.6 void xnselector\_destroy ( struct xnselector \* selector )**

Destroy a selector block.

All bindings with file descriptor are destroyed.

**Parameters**

*selector* the selector block to be destroyed

Referenced by `xnpod_delete_thread()`.

**5.18.2.7 int xnselector\_init ( struct xnselector \* selector )**

Initialize a selector structure.

**Parameters**

*selector* The selector structure to be initialized.

**Return values**

*0*

References `xnsynch_init()`.

**5.19 Real-time shadow services.****Files**

- file [shadow.c](#)

*Real-time shadow services.*

## Functions

- `int xnshadow_harden (void)`  
*Migrate a Linux task to the Xenomai domain.*
- `void xnshadow_relax (int notify, int reason)`  
*Switch a shadow thread back to the Linux domain.*
- `int xnshadow_map (xnthread_t *thread, xncompletion_t __user *u_completion, unsigned long __user *u_window_offset)`  
*Create a shadow thread context.*
- `xnshadow_ppd_t * xnshadow_ppd_get (unsigned int muxid)`  
*Return the per-process data attached to the calling process.*

### 5.19.1 Detailed Description

Real-time shadow services.

### 5.19.2 Function Documentation

#### 5.19.2.1 `int xnshadow_harden ( void )`

Migrate a Linux task to the Xenomai domain.

**For internal use only.**

This service causes the transition of "current" from the Linux domain to Xenomai. The shadow will resume in the Xenomai domain as returning from `schedule()`.

Environments:

This service can be called from:

- User-space thread operating in secondary (i.e. relaxed) mode.

Rescheduling: always.

References `XNDEBUG`, `xnpod_dispatch_signals()`, `XNRELAX`, and `xnshadow_relax()`.

Referenced by `xnshadow_map()`.

#### 5.19.2.2 `int xnshadow_map ( xnthread_t * thread, xncompletion_t __user * u_completion, unsigned long __user * u_window_offset )`

Create a shadow thread context.

**For internal use only.**

This call maps a nucleus thread to the "current" Linux task. The priority and scheduling class of the underlying Linux task are not affected; it is assumed that the interface library did set them appropriately before issuing the shadow mapping request.

**Parameters**

*thread* The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to [xnpod\\_init\\_thread\(\)](#).

*u\_completion* is the address of an optional completion descriptor aimed at synchronizing our parent thread with us. If non-NULL, the information [xnshadow\\_map\(\)](#) will store into the completion block will be later used to wake up the parent thread when the current shadow has been initialized. In the latter case, the new shadow thread is left in a dormant state (XNDORMANT) after its creation, leading to the suspension of "current" in the Linux domain, only processing signals. Otherwise, the shadow thread is immediately started and "current" immediately resumes in the Xenomai domain from this service.

*u\_window\_offset* will receive the offset of the per-thread "u\_window" structure in the process shared heap associated to *thread*. This structure reflects thread state information visible from userland through a shared memory window.

**Returns**

0 is returned on success. Otherwise:

- -ERESTARTSYS is returned if the current Linux task has received a signal, thus preventing the final migration to the Xenomai domain (i.e. in order to process the signal in the Linux domain). This error should not be considered as fatal.
- -EPERM is returned if the shadow thread has been killed before the current task had a chance to return to the caller. In such a case, the real-time mapping operation has failed globally, and no Xenomai resource remains attached to it.
- -EINVAL is returned if the thread control block does not bear the XNSHADOW bit.
- -EBUSY is returned if either the current Linux task or the associated shadow thread is already involved in a shadow mapping.

Environments:

This service can be called from:

- Regular user-space process.

Rescheduling: always.

References [xnheap\\_alloc\(\)](#), [XNMAPPED](#), [XNOTHER](#), [xnpod\\_start\\_thread\(\)](#), [xnpod\\_suspend\\_thread\(\)](#), [XNPRIOSSET](#), [XNRELAX](#), [XNSHADOW](#), and [xnshadow\\_harden\(\)](#).



**5.19.2.3 `xnshadow_ppd_t* xnshadow_ppd_get ( unsigned int muxid )`**

Return the per-process data attached to the calling process.

This service returns the per-process data attached to the calling process for the skin whose *muxid* is *muxid*. It must be called with `nklock` locked, `irqs` off.

See `xnshadow_register_interface()` documentation for information on the way to attach a per-process data to a process.

**Parameters**

*muxid* the skin *muxid*.

**Returns**

the per-process data if the current context is a user-space process;  
NULL otherwise.

**5.19.2.4 `void xnshadow_relax ( int notify, int reason )`**

Switch a shadow thread back to the Linux domain.

**For internal use only.**

This service yields the control of the running shadow back to Linux. This is obtained by suspending the shadow and scheduling a wake up call for the mated user task inside the Linux domain. The Linux task will resume on return from `xnpod_suspend_thread()` on behalf of the root thread.

**Parameters**

*notify* A boolean flag indicating whether threads monitored from secondary mode switches should be sent a SIGDEBUG signal. For instance, some internal operations like task exit should not trigger such signal.

*reason* The reason to report along with the SIGDEBUG signal.

Environments:

This service can be called from:

- User-space thread operating in primary (i.e. `harden`) mode.

Rescheduling: always.

**Note**

"current" is valid here since the shadow runs with the properties of the Linux task.

References `XNAFFSET`, `xnpod_suspend_thread()`, `XNPRIOSSET`, `XNRELAX`, `XNROOT`, and `XNTRAPSW`.

Referenced by `xnpod_handle_exception()`, and `xnshadow_harden()`.

## 5.20 Thread synchronization services.

### Files

- file [synch.c](#)  
*Thread synchronization services.*

### Functions

- void [xnsynch\\_init](#) (struct xnsynch \*synch, xnflags\_t flags, xnarch\_atomic\_t \*fastlock)  
*Initialize a synchronization object.*
- xnflags\_t [xnsynch\\_sleep\\_on](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Sleep on an ownerless synchronization object.*
- struct xnthread \* [xnsynch\\_wakeup\\_one\\_sleeper](#) (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*
- struct xnpholder \* [xnsynch\\_wakeup\\_this\\_sleeper](#) (struct xnsynch \*synch, struct xnpholder \*holder)  
*Give the resource ownership to a given waiting thread.*
- xnflags\_t [xnsynch\\_acquire](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Acquire the ownership of a synchronization object.*
- static void [xnsynch\\_clear\\_boost](#) (struct xnsynch \*synch, struct xnthread \*owner)  
*Clear the priority boost.*
- void [xnsynch\\_requeue\\_sleeper](#) (struct xnthread \*thread)  
*Change a sleeper's priority.*
- struct xnthread \* [xnsynch\\_peek\\_pendq](#) (struct xnsynch \*synch)  
*Access the thread leading a synch object wait queue.*
- int [xnsynch\\_flush](#) (struct xnsynch \*synch, xnflags\_t reason)  
*Unblock all waiters pending on a resource.*
- void [xnsynch\\_forget\\_sleeper](#) (struct xnthread \*thread)  
*Abort a wait for a resource.*
- void [xnsynch\\_release\\_all\\_ownerships](#) (struct xnthread \*thread)  
*Release all ownerships.*
- static struct xnthread \* [xnsynch\\_release](#) (struct xnsynch \*synch, struct xnthread \*thread)  
*Give the resource ownership to the next waiting thread.*

## 5.20.1 Detailed Description

Thread synchronization services.

## 5.20.2 Function Documentation

### 5.20.2.1 `xnflags_t xnsynch_acquire ( struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode )`

Acquire the ownership of a synchronization object.

This service should be called by upper interfaces wanting the current thread to acquire the ownership of the given resource. If the resource is already assigned to a thread, the caller is suspended.

This service must be used only with synchronization objects that track ownership (XNSYNCH\_OWNER set).

#### Parameters

*synch* The descriptor address of the synchronization object to acquire.

*timeout* The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing XN\_INFINITY and setting *mode* to XN\_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

#### Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References XNBOOST, XNBREAK, XNOTHER, XNPEND, `xnpod_suspend_thread()`, XNRMID, XNROBBED, XNTIMEO, and XNWAKEN.

Referenced by `rtdm_mutex_timedlock()`.

### 5.20.2.2 `void xnsynch_clear_boost ( struct xnsynch * synch, struct xnthread * owner ) [static]`

Clear the priority boost.

**For internal use only.**

This service is called internally whenever a synchronization object is not claimed anymore by sleepers to reset the object owner's priority to its initial level.

**Parameters**

*synch* The descriptor address of the synchronization object.

*owner* The descriptor address of the thread which currently owns the synchronization object.

**Note**

This routine must be entered nklock locked, interrupts off.

References XNBOOST.

Referenced by `xnsynch_flush()`, and `xnsynch_forget_sleeper()`.

**5.20.2.3 int xnsynch\_flush ( struct xnsynch \* *synch*, xnflags\_t *reason* )**

Unblock all waiters pending on a resource.

This service atomically releases all threads which currently sleep on a given resource.

This service should be called by upper interfaces under circumstances requiring that the pending queue of a given resource is cleared, such as before the resource is deleted.

**Parameters**

*synch* The descriptor address of the synchronization object to be flushed.

*reason* Some flags to set in the information mask of every unblocked thread. Zero is an acceptable value. The following bits are pre-defined by the nucleus:

- XNRMID should be set to indicate that the synchronization object is about to be destroyed (see [xnpod\\_resume\\_thread\(\)](#)).
- XNBREAK should be set to indicate that the wait has been forcibly interrupted (see [xnpod\\_unblock\\_thread\(\)](#)).

**Returns**

XNSYNCH\_RESCHED is returned if at least one thread is unblocked, which means the caller should invoke [xnpod\\_schedule\(\)](#) for applying the new scheduling state. Otherwise, XNSYNCH\_DONE is returned.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object is no more owned by any thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, xnpod\_resume\_thread(), and xnsynch\_clear\_boost().

Referenced by rtdm\_event\_signal(), and xnregistry\_put().

#### 5.20.2.4 void xnsynch\_forget\_sleeper ( struct xnthread \* *thread* )

Abort a wait for a resource.

**For internal use only.**

Performs all the necessary housekeeping chores to stop a thread from waiting on a given synchronization object.

##### Parameters

*thread* The descriptor address of the affected thread.

When the trace support is enabled (i.e. MVM), the idle state is posted to the synchronization object's state diagram (if any) whenever no thread remains blocked on it. The real-time interfaces must ensure that such condition (i.e. EMPTY/IDLE) is mapped to state #0.

##### Note

This routine must be entered nklock locked, interrupts off.

References XNPEND, and xnsynch\_clear\_boost().

Referenced by xnpod\_delete\_thread(), xnpod\_resume\_thread(), and xnpod\_suspend\_thread().

#### 5.20.2.5 void xnsynch\_init ( struct xnsynch \* *synch*, xnflags\_t *flags*, xnarch\_atomic\_t \* *fastlock* )

Initialize a synchronization object.

Initializes a new specialized object which can subsequently be used to synchronize real-time activities. The Xenomai nucleus provides a basic synchronization object which can be used to build higher resource objects. Nucleus threads can wait for and signal such objects in order to synchronize their activities.

This object has built-in support for priority inheritance.

##### Parameters

*synch* The address of a synchronization object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

*flags* A set of creation flags affecting the operation. The valid flags are:

- XNSYNCH\_PRIO causes the threads waiting for the resource to pend in priority order. Otherwise, FIFO ordering is used (XNSYNCH\_FIFO).
- XNSYNCH\_OWNER indicates that the synchronization object shall track its owning thread (required if XNSYNCH\_PIP is selected). Note that setting this flag implies the use `xnsynch_acquire` and `xnsynch_release` instead of `xnsynch_sleep_on` and `xnsynch_wakeup_one_sleeper/xnsynch_wakeup_this_sleeper`.
- XNSYNCH\_PIP causes the priority inheritance mechanism to be automatically activated when a priority inversion is detected among threads using this object. Otherwise, no priority inheritance takes place upon priority inversion (XNSYNCH\_NOPIP).
- XNSYNCH\_DREORD (Disable REORDERing) tells the nucleus that the wait queue should not be reordered whenever the priority of a blocked thread it holds is changed. If this flag is not specified, changing the priority of a blocked thread using `xnpod_set_thread_schedparam()` will cause this object's wait queue to be reordered according to the new priority level, provided the synchronization object makes the waiters wait by priority order on the awaited resource (XNSYNCH\_PRIO).

#### Parameters

*fastlock* Address of the fast lock word to be associated with the synchronization object. If NULL is passed or XNSYNCH\_OWNER is not set, fast-lock support is disabled.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `pthread_cond_init()`, `pthread_create()`, `rtdm_event_init()`, `rtdm_mutex_init()`, `rtdm_sem_init()`, `xnregistry_enter()`, and `xnselector_init()`.

#### 5.20.2.6 `struct xnthread* xnsynch_peek_pendq ( struct xnsynch * synch ) [read]`

Access the thread leading a synch object wait queue.

This services returns the descriptor address of to the thread leading a synchronization object wait queue.

#### Parameters

*synch* The descriptor address of the target synchronization object.

#### Returns

The descriptor address of the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

**5.20.2.7** `struct xnthread * xnsynch_release ( struct xnsynch * synch, struct xnthread * owner )` **[static, read]**

Give the resource ownership to the next waiting thread.

This service releases the ownership of the given synchronization object. The thread which is currently leading the object's pending list, if any, is unblocked from its pending state. However, no reschedule is performed.

This service must be used only with synchronization objects that track ownership (XNSYNCH\_OWNER set).

#### Parameters

*synch* The descriptor address of the synchronization object whose ownership is changed.  
*owner* The descriptor address of the current owner.

#### Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNOTHER.

Referenced by `xnsynch_release_all_ownerships()`.

#### 5.20.2.8 void xnsynch\_release\_all\_ownerships ( struct xnthread \* *thread* )

Release all ownerships.

**For internal use only.**

This call is used internally to release all the ownerships obtained by a thread on synchronization objects. This routine must be entered interrupts off.

##### Parameters

*thread* The descriptor address of the affected thread.

##### Note

This routine must be entered nklock locked, interrupts off.

References xnsynch\_release().

Referenced by \_\_xnpod\_reset\_thread(), and xnpod\_delete\_thread().

#### 5.20.2.9 void xnsynch\_requeue\_sleeper ( struct xnthread \* *thread* )

Change a sleeper's priority.

**For internal use only.**

This service is used by the PIP code to update the pending priority of a sleeping thread.

##### Parameters

*thread* The descriptor address of the affected thread.

##### Note

This routine must be entered nklock locked, interrupts off.

References XNBOOST.

#### 5.20.2.10 xnflags\_t xnsynch\_sleep\_on ( struct xnsynch \* *synch*, xnticks\_t *timeout*, xntmode\_t *timeout\_mode* )

Sleep on an ownerless synchronization object.

Makes the calling thread sleep on the specified synchronization object, waiting for it to be signaled.

This service should be called by upper interfaces wanting the current thread to pend on the given resource. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH\_OWNER).

##### Parameters

*synch* The descriptor address of the synchronization object to sleep on.



***timeout*** The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing `XN_INFINITE` **and** setting *mode* to `XN_RELATIVE` specifies an unbounded wait. All other values are used to initialize a watchdog timer.

***timeout\_mode*** The mode of the *timeout* parameter. It can either be set to `XN_RELATIVE`, `XN_ABSOLUTE`, or `XN_REALTIME` (see also [xntimer\\_start\(\)](#)).

### Returns

A bitmask which may include zero or one information bit among `XNRMID`, `XNTIMEO` and `XNBREAK`, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always.

References `XNBREAK`, `XNPEND`, `xnpod_suspend_thread()`, `XNRMID`, and `XNTIMEO`.

Referenced by `rtdm_event_timedwait()`, `rtdm_sem_timeddown()`, `xnregistry_bind()`, `xnregistry_remove_safe()`, and `xnselect()`.

#### 5.20.2.11 `struct xnthread* xnsynch_wakeup_one_sleeper ( struct xnsynch * synch ) [read]`

Give the resource ownership to the next waiting thread.

This service wakes up the thread which is currently leading the synchronization object's pending list. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a single waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (`XNSYNCH_OWNER` not set).

### Parameters

***synch*** The descriptor address of the synchronization object whose ownership is changed.

### Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod\_resume\_thread().

Referenced by rtdm\_sem\_up().

#### 5.20.2.12 `struct xnpholder* xnsynch_wakeup_this_sleeper ( struct xnsynch * synch, struct xnpholder * holder ) [read]`

Give the resource ownership to a given waiting thread.

This service wakes up a specific thread which is currently pending on the given synchronization object. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a specific waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH\_OWNER not set).

#### Parameters

*synch* The descriptor address of the synchronization object whose ownership is changed.

*holder* The link holder address of the thread to unblock (&thread->plink) which MUST be currently linked to the synchronization object's pending queue (i.e. synch->pendq).

#### Returns

The link address of the unblocked thread in the synchronization object's pending queue.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod\_resume\_thread().

## 5.21 Timer services.

### Files

- file [timer.h](#)
- file [timer.c](#)

### Functions

- int [xntimer\\_start](#) (xntimer\_t \*timer, xnticks\_t value, xnticks\_t interval, xntmode\_t mode)  
*Arm a timer.*
- xnticks\_t [xntimer\\_get\\_date](#) (xntimer\_t \*timer)  
*Return the absolute expiration date.*
- xnticks\_t [xntimer\\_get\\_timeout](#) (xntimer\_t \*timer)  
*Return the relative expiration date.*
- xnticks\_t [xntimer\\_get\\_interval](#) (xntimer\_t \*timer)  
*Return the timer interval value.*
- static void [xntimer\\_stop](#) (xntimer\_t \*timer)  
*Disarm a timer.*
- void [xntimer\\_tick](#) (void)  
*Process a timer tick.*
- void [xntimer\\_init](#) (xntimer\_t \*timer, void(\*handler)(xntimer\_t \*timer))  
*Initialize a timer object.*
- void [xntimer\\_destroy](#) (xntimer\_t \*timer)  
*Release a timer object.*
- unsigned long [xntimer\\_get\\_overruns](#) (xntimer\_t \*timer, xnticks\_t now)  
*Get the count of overruns for the last tick.*
- void [xntimer\\_freeze](#) (void)  
*Freeze all timers (from every time bases).*

#### 5.21.1 Detailed Description

The Xenomai timer facility always operate the timer hardware in oneshot mode, regardless of the time base in effect. Periodic timing is obtained through a software emulation, using cascading timers.

The timer object stores time as a count of CPU ticks (e.g. TSC values).

## 5.21.2 Function Documentation

### 5.21.2.1 void xntimer\_destroy ( xntimer\_t \* *timer* )

Release a timer object.

Destroys a timer. After it has been destroyed, all resources associated with the timer have been released. The timer is automatically deactivated before deletion if active on entry.

#### Parameters

*timer* The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References xntimer\_stop().

Referenced by rtdm\_timer\_destroy(), and xnpod\_delete\_thread().

### 5.21.2.2 void xntimer\_freeze ( void )

Freeze all timers (from every time bases).

**For internal use only.**

This routine deactivates all active timers atomically.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

References xnsched::status.

Referenced by xnpod\_disable\_timesource().

### 5.21.2.3 `xnticks_t xntimer_get_date ( xntimer_t * timer )`

Return the absolute expiration date.

Return the next expiration date of a timer as an absolute count of nanoseconds.

#### Parameters

*timer* The address of a valid timer descriptor.

#### Returns

The expiration date in nanoseconds. The special value `XN_INFINITE` is returned if *timer* is currently disabled.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

### 5.21.2.4 `xnticks_t xntimer_get_interval ( xntimer_t * timer )`

Return the timer interval value.

Return the timer interval value in nanoseconds.

#### Parameters

*timer* The address of a valid timer descriptor.

#### Returns

The duration of a period in nanoseconds. The special value `XN_INFINITE` is returned if *timer* is currently disabled or one shot.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

### 5.21.2.5 `unsigned long xntimer_get_overruns ( xntimer_t * timer, xnticks_t now )`

Get the count of overruns for the last tick.

This service returns the count of pending overruns for the last tick of a given timer, as measured by the difference between the expected expiry date of the timer and the date *now* passed as argument.

#### Parameters

*timer* The address of a valid timer descriptor.

*now* current date (in the monotonic time base)

### Returns

the number of overruns of *timer* at date *now*

Referenced by `xnpod_wait_thread_period()`.

#### 5.21.2.6 `xnticks_t xntimer_get_timeout ( xntimer_t * timer )`

Return the relative expiration date.

This call returns the count of nanoseconds remaining until the timer expires.

### Parameters

*timer* The address of a valid timer descriptor.

### Returns

The count of nanoseconds until expiry. The special value `XN_INFINITE` is returned if *timer* is currently disabled. It might happen that the timer expires when this service runs (even if the associated handler has not been fired yet); in such a case, 1 is returned.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

#### 5.21.2.7 `void xntimer_init ( xntimer_t * timer, void(*)(xntimer_t *timer) handler )`

Initialize a timer object.

Creates a timer. When created, a timer is left disarmed; it must be started using [xntimer\\_start\(\)](#) in order to be activated.

### Parameters

*timer* The address of a timer descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

*handler* The routine to call upon expiration of the timer.

There is no limitation on the number of timers which can be created/active concurrently.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

Referenced by `timer_create()`.

### 5.21.2.8 `int xntimer_start ( xntimer_t * timer, xnticks_t value, xnticks_t interval, xntmode_t mode )`

Arm a timer.

Activates a timer so that the associated timeout handler will be fired after each expiration time. A timer can be either periodic or one-shot, depending on the reload value passed to this routine. The given timer must have been previously initialized.

#### Parameters

*timer* The address of a valid timer descriptor.

*value* The date of the initial timer shot, expressed in nanoseconds.

*interval* The reload value of the timer. It is a periodic interval value to be used for reprogramming the next timer shot, expressed in nanoseconds. If *interval* is equal to XN\_INFINITE, the timer will not be reloaded after it has expired.

*mode* The timer mode. It can be XN\_RELATIVE if *value* shall be interpreted as a relative date, XN\_ABSOLUTE for an absolute date based on the monotonic clock of the related time base (as returned by `xnclock_read_monotonic()`), or XN\_REALTIME if the absolute date is based on the adjustable real-time clock (obtained from `xnclock_read()`).

#### Returns

0 is returned upon success, or -ETIMEDOUT if an absolute date in the past has been given.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

#### Note

Must be called with `nklock` held, IRQs off.

Referenced by `rtm_timer_start()`, `timer_settime()`, `xnpod_enable_timesource()`, `xnpod_set_thread_periodic()`, `xnpod_set_thread_tslice()`, and `xnpod_suspend_thread()`.

### 5.21.2.9 `int xntimer_stop ( xntimer_t * timer ) [inline, static]`

Disarm a timer.

This service deactivates a timer previously armed using [xntimer\\_start\(\)](#). Once disarmed, the timer can be subsequently re-armed using the latter service.

#### Parameters

*timer* The address of a valid timer descriptor.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

#### Note

Must be called with nklock held, IRQs off.

Referenced by `rtm_timer_stop()`, `timer_settime()`, `xnpod_resume_thread()`, `xnpod_set_thread_periodic()`, `xnpod_set_thread_tslice()`, and `xntimer_destroy()`.

#### 5.21.2.10 void xntimer\_tick ( void )

Process a timer tick.

**For internal use only.**

This routine informs all active timers that the clock has been updated by processing the outstanding timer list. Elapsed timer actions will be fired.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

References `xnsched::htimer`, `xnsched::lflags`, and `xnsched::status`.

## 5.22 Virtual file services

### Data Structures

- struct [xnvfile\\_lock\\_ops](#)  
*Vfile locking operations.*
- struct [xnvfile\\_regular\\_ops](#)  
*Regular vfile operation descriptor.*
- struct [xnvfile\\_regular\\_iterator](#)  
*Regular vfile iterator.*
- struct [xnvfile\\_snapshot\\_ops](#)  
*Snapshot vfile operation descriptor.*
- struct [xnvfile\\_rev\\_tag](#)  
*Snapshot revision tag.*
- struct [xnvfile\\_snapshot](#)



*Snapshot vfile descriptor.*

- struct [xnvmfile\\_snapshot\\_iterator](#)  
*Snapshot-driven vfile iterator.*

## Files

- file [vfile.h](#)  
*This file is part of the Xenomai project.*

## Functions

- int [xnvmfile\\_init\\_snapshot](#) (const char \*name, struct [xnvmfile\\_snapshot](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a snapshot-driven vfile.*
- int [xnvmfile\\_init\\_regular](#) (const char \*name, struct [xnvmfile\\_regular](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a regular vfile.*
- int [xnvmfile\\_init\\_dir](#) (const char \*name, struct [xnvmfile\\_directory](#) \*vdir, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual directory entry.*
- int [xnvmfile\\_init\\_link](#) (const char \*from, const char \*to, struct [xnvmfile\\_link](#) \*vlink, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual link entry.*
- void [xnvmfile\\_destroy](#) (struct [xnvmfile](#) \*vfile)  
*Removes a virtual file entry.*
- ssize\_t [xnvmfile\\_get\\_blob](#) (struct [xnvmfile\\_input](#) \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_string](#) (struct [xnvmfile\\_input](#) \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_integer](#) (struct [xnvmfile\\_input](#) \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct [xnvmfile\\_directory](#) [nkvfroot](#)  
*Xenomai vfile root directory.*
- struct [xnvmfile\\_directory](#) [nkvfroot](#)  
*Xenomai vfile root directory.*

### 5.22.1 Detailed Description

Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces. This encapsulation is aimed at:

- supporting consistent collection of very large record-based output, without incurring latency peaks for undergoing real-time activities.
- in the future, hiding discrepancies between linux kernel releases, regarding the proper way to export kernel object states to userland, either via the /proc interface or by any other mean.

This virtual file implementation offers record-based read support based on seq\_files, single-buffer write support, directory and link handling, all visible from the /proc namespace.

The vfile support exposes four filesystem object types:

- snapshot-driven file (struct `xnvfile_snapshot`). This is commonly used to export real-time object states via the /proc filesystem. To minimize the latency involved in protecting the vfile routines from changes applied by real-time code on such objects, a snapshot of the data to output is first taken under proper locking, before the collected data is formatted and sent out in a lockless manner.

Because a large number of records may have to be output, the data collection phase is not strictly atomic as a whole, but only protected at record level. The vfile implementation can be notified of updates to the underlying data set, and restart the collection from scratch until the snapshot is fully consistent.

- regular sequential file (struct `xnvfile_regular`). This is basically an encapsulated sequential file object as available from the host kernel (i.e. `seq_file`), with a few additional features to make it more handy in a Xenomai environment, like implicit locking support and shortened declaration for simplest, single-record output.
- virtual link (struct `xnvfile_link`). This is a symbolic link feature integrated with the vfile semantics. The link target is computed dynamically at creation time from a user-given helper routine.
- virtual directory (struct `xnvfile_directory`). A directory object, which can be used to create a hierarchy for ordering a set of vfile objects.

### 5.22.2 Function Documentation

#### 5.22.2.1 void `xnvfile_destroy` ( struct `xnvfile` \* *vfile* )

Removes a virtual file entry.

##### Parameters

*vfile* A pointer to the virtual file descriptor to remove.

References `nkvfroot`.

### 5.22.2.2 `ssize_t xnvfile_get_blob ( struct xnvfile_input * input, void * data, size_t size )`

Read in a data bulk written to the vfile.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_blob()` retrieves this data as an untyped binary blob, and copies it back to the caller's buffer.

#### Parameters

*input* A pointer to the input descriptor passed to the `store()` handler.

*data* The address of the destination buffer to copy the input data to.

*size* The maximum number of bytes to copy to the destination buffer. If *size* is larger than the actual data size, the input is truncated to *size*.

#### Returns

The number of bytes read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- `-EFAULT` indicates an invalid source buffer address.

Referenced by `xnvfile_get_integer()`, and `xnvfile_get_string()`.

### 5.22.2.3 `ssize_t xnvfile_get_integer ( struct xnvfile_input * input, long * valp )`

Evaluate the string written to the vfile as a long integer.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_integer()` retrieves and interprets this data as a long integer, and copies the resulting value back to *valp*.

The long integer can be expressed in decimal, octal or hexadecimal bases depending on the prefix found.

#### Parameters

*input* A pointer to the input descriptor passed to the `store()` handler.

*valp* The address of a long integer variable to receive the value.

#### Returns

The number of characters read while evaluating the input as a long integer upon success. Otherwise, a negative error code is returned:

- `-EINVAL` indicates a parse error on the input stream; the written text cannot be evaluated as a long integer.
- `-EFAULT` indicates an invalid source buffer address.

References `xnvfile_get_blob()`.

#### 5.22.2.4 `ssize_t xnvfile_get_string ( struct xnvfile_input * input, char * s, size_t maxlen )`

Read in a C-string written to the vfile.

When writing to a vfile, the associated store() handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_string()` retrieves this data as a null-terminated character string, and copies it back to the caller's buffer.

##### Parameters

- input* A pointer to the input descriptor passed to the store() handler.
- s* The address of the destination string buffer to copy the input data to.
- maxlen* The maximum number of bytes to copy to the destination buffer, including the ending null character. If *maxlen* is larger than the actual string length, the input is truncated to *maxlen*.

##### Returns

The number of characters read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- -EFAULT indicates an invalid source buffer address.

References `xnvfile_get_blob()`.

#### 5.22.2.5 `int xnvfile_init_dir ( const char * name, struct xnvfile_directory * vdir, struct xnvfile_directory * parent )`

Initialize a virtual directory entry.

##### Parameters

- name* The name which should appear in the pseudo-filesystem, identifying the vdir entry.
- vdir* A pointer to the virtual directory descriptor to initialize.
- parent* A pointer to a virtual directory descriptor standing for the parent directory of the new vdir. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfroot* vdir.

##### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual directory entry cannot be created in the /proc hierarchy.

#### 5.22.2.6 `int xnvfile_init_link ( const char * from, const char * to, struct xnvfile_link * vlink, struct xnvfile_directory * parent )`

Initialize a virtual link entry.

##### Parameters

- from* The name which should appear in the pseudo-filesystem, identifying the vlink entry.

*to* The target file name which should be referred to symbolically by *name*.

*vlink* A pointer to the virtual link descriptor to initialize.

*parent* A pointer to a virtual directory descriptor standing for the parent directory of the new vlink. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfsroot* vdir.

### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual link entry cannot be created in the /proc hierarchy.

#### 5.22.2.7 `int xnvfile_init_regular ( const char * name, struct xnvfile_regular * vfile, struct xnvfile_directory * parent )`

Initialize a regular vfile.

### Parameters

*name* The name which should appear in the pseudo-filesystem, identifying the vfile entry.

*vfile* A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- .privsz is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- entry.lockops is a pointer to a [locking](#) descriptor", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case no locking will be applied.
- .ops is a pointer to an [operation descriptor](#).

### Parameters

*parent* A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfsroot* vdir.

### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

#### 5.22.2.8 `int xnvfile_init_snapshot ( const char * name, struct xnvfile_snapshot * vfile, struct xnvfile_directory * parent )`

Initialize a snapshot-driven vfile.

### Parameters

*name* The name which should appear in the pseudo-filesystem, identifying the vfile entry.

*vfile* A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- `.privsz` is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- `.datasz` is the size (in bytes) of a single record to be collected by the [next\(\) handler](#) from the [operation descriptor](#).
- `.tag` is a pointer to a mandatory vfile revision tag structure (struct [xnvmfile\\_rev\\_tag](#)). This tag will be monitored for changes by the vfile core while collecting data to output, so that any update detected will cause the current snapshot data to be dropped, and the collection to restart from the beginning. To this end, any change to the data which may be part of the collected records, should also invoke `xnvmfile_touch()` on the associated tag.
- `entry.lockops` is a pointer to a [locking descriptor](#)", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case the operations on the nucleus lock (i.e. `nklock`) will be used internally around calls to data collection handlers (see [operation descriptor](#)).
- `.ops` is a pointer to an [operation descriptor](#).

#### Parameters

*parent* A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the `/proc` root directory will be used. `/proc/xenomai` is mapped on the globally available `nkvmfroot` vdir.

#### Returns

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if the virtual file entry cannot be created in the `/proc` hierarchy.

References `xnvmfile_snapshot_ops::store`.

### 5.22.3 Variable Documentation

#### 5.22.3.1 struct xnvmfile\_directory nkvmfroot

Xenomai vfile root directory.

This vdir maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by `xnvmfile_destroy()`.

#### 5.22.3.2 struct xnvmfile\_directory nkvmfroot

Xenomai vfile root directory.

This vdir maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by `xnvmfile_destroy()`.

## 5.23 Inter-Driver API

Collaboration diagram for Inter-Driver API:



### Functions

- struct `rt dm_dev_context` \* `rt dm_context_get` (int fd)  
*Retrieve and lock a device context.*
- int `rt dm_select_bind` (int fd, `rt dm_selector_t` \*selector, enum `rt dm_selecttype` type, unsigned fd\_index)  
*Bind a selector to specified event types of a given file descriptor.*
- void `rt dm_context_lock` (struct `rt dm_dev_context` \*context)  
*Increment context reference counter.*
- void `rt dm_context_unlock` (struct `rt dm_dev_context` \*context)  
*Decrement context reference counter.*
- void `rt dm_context_put` (struct `rt dm_dev_context` \*context)  
*Release a device context obtained via `rt dm_context_get()`.*
- int `rt dm_open` (const char \*path, int oflag,...)  
*Open a device.*
- int `rt dm_socket` (int protocol\_family, int socket\_type, int protocol)  
*Create a socket.*
- int `rt dm_close` (int fd)  
*Close a device or socket.*
- int `rt dm_ioctl` (int fd, int request,...)  
*Issue an IOCTL.*
- ssize\_t `rt dm_read` (int fd, void \*buf, size\_t nbyte)  
*Read from device.*
- ssize\_t `rt dm_write` (int fd, const void \*buf, size\_t nbyte)  
*Write to device.*

- `ssize_t rtdm_recvmmsg` (int fd, struct msghdr \*msg, int flags)  
*Receive message from socket.*
- `ssize_t rtdm_recvfrom` (int fd, void \*buf, size\_t len, int flags, struct sockaddr \*from, socklen\_t \*fromlen)  
*Receive message from socket.*
- `ssize_t rtdm_recv` (int fd, void \*buf, size\_t len, int flags)  
*Receive message from socket.*
- `ssize_t rtdm_sendmsg` (int fd, const struct msghdr \*msg, int flags)  
*Transmit message to socket.*
- `ssize_t rtdm_sendto` (int fd, const void \*buf, size\_t len, int flags, const struct sockaddr \*to, socklen\_t tolen)  
*Transmit message to socket.*
- `ssize_t rtdm_send` (int fd, const void \*buf, size\_t len, int flags)  
*Transmit message to socket.*
- `int rtdm_bind` (int fd, const struct sockaddr \*my\_addr, socklen\_t addrlen)  
*Bind to local address.*
- `int rtdm_connect` (int fd, const struct sockaddr \*serv\_addr, socklen\_t addrlen)  
*Connect to remote address.*
- `int rtdm_listen` (int fd, int backlog)  
*Listen for incoming connection requests.*
- `int rtdm_accept` (int fd, struct sockaddr \*addr, socklen\_t \*addrlen)  
*Accept a connection requests.*
- `int rtdm_shutdown` (int fd, int how)  
*Shut down parts of a connection.*
- `int rtdm_getsockopt` (int fd, int level, int optname, void \*optval, socklen\_t \*optlen)  
*Get socket option.*
- `int rtdm_setsockopt` (int fd, int level, int optname, const void \*optval, socklen\_t optlen)  
*Set socket option.*
- `int rtdm_getsockname` (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get local socket address.*
- `int rtdm_getpeername` (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get socket destination address.*



### 5.23.1 Function Documentation

#### 5.23.1.1 `int rtdm_accept ( int fd, struct sockaddr * addr, socklen_t * addrlen )`

Accept a connection requests.

Refer to [rt\\_dev\\_accept\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### 5.23.1.2 `int rtdm_bind ( int fd, const struct sockaddr * my_addr, socklen_t addrlen )`

Bind to local address.

Refer to [rt\\_dev\\_bind\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### 5.23.1.3 `int rtdm_close ( int fd )`

Close a device or socket.

Refer to [rt\\_dev\\_close\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### 5.23.1.4 `int rtdm_connect ( int fd, const struct sockaddr * serv_addr, socklen_t addrlen )`

Connect to remote address.

Refer to [rt\\_dev\\_connect\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### 5.23.1.5 `struct rtdm_dev_context* rtdm_context_get ( int fd ) [read]`

Retrieve and lock a device context.

##### Parameters

[in] *fd* File descriptor

##### Returns

Pointer to associated device context, or NULL on error

**Note**

The device context has to be unlocked using [rt dm\\_context\\_put\(\)](#) when it is no longer referenced.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `rt dm_dev_context::close_lock_count`.

Referenced by `rt dm_select_bind()`.

**5.23.1.6 void rt dm\_context\_lock ( struct rt dm\_dev\_context \* context )**

Increment context reference counter.

**Parameters**

[in] *context* Device context

**Note**

[rt dm\\_context\\_get\(\)](#) automatically increments the lock counter. You only need to call this function in special scenarios, e.g. when keeping additional references to the context structure that have different lifetimes. Only use [rt dm\\_context\\_lock\(\)](#) on contexts that are currently locked via an earlier [rt dm\\_context\\_get\(\)](#)/`rt dm_context_lock()` or while running a device operation handler.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.23.1.7 void rtdm\_context\_put ( struct rtdm\_dev\_context \* context )

Release a device context obtained via [rtdm\\_context\\_get\(\)](#).

##### Parameters

[in] *context* Device context

##### Note

Every successful call to [rtdm\\_context\\_get\(\)](#) must be matched by a [rtdm\\_context\\_put\(\)](#) invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.23.1.8 void rtdm\_context\_unlock ( struct rtdm\_dev\_context \* context )

Decrement context reference counter.

##### Parameters

[in] *context* Device context

##### Note

Every call to [rtdm\\_context\\_locked\(\)](#) must be matched by a [rtdm\\_context\\_unlock\(\)](#) invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Referenced by [rtdm\\_select\\_bind\(\)](#).

**5.23.1.9 int rtdm\_getpeername ( int *fd*, struct sockaddr \* *name*, socklen\_t \* *namelen* )**

Get socket destination address.

Refer to [rt\\_dev\\_getpeername\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.10 int rtdm\_getsockname ( int *fd*, struct sockaddr \* *name*, socklen\_t \* *namelen* )**

Get local socket address.

Refer to [rt\\_dev\\_getsockname\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.11 int rtdm\_getsockopt ( int *fd*, int *level*, int *optname*, void \* *optval*, socklen\_t \* *optlen* )**

Get socket option.

Refer to [rt\\_dev\\_getsockopt\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.12 int rtdm\_ioctl ( int *fd*, int *request*, ... )**

Issue an IOCTL.

Refer to [rt\\_dev\\_ioctl\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.13 int rtdm\_listen ( int *fd*, int *backlog* )**

Listen for incoming connection requests.

Refer to [rt\\_dev\\_listen\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.14** `int rtdm_open ( const char * path, int oflag, ... )`

Open a device.

Refer to [rt\\_dev\\_open\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.15** `ssize_t rtdm_read ( int fd, void * buf, size_t nbyte )`

Read from device.

Refer to [rt\\_dev\\_read\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.16** `ssize_t rtdm_recv ( int fd, void * buf, size_t len, int flags )`

Receive message from socket.

Refer to [rt\\_dev\\_recv\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.17** `ssize_t rtdm_recvfrom ( int fd, void * buf, size_t len, int flags, struct sockaddr * from, socklen_t * fromlen )`

Receive message from socket.

Refer to [rt\\_dev\\_recvfrom\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.18** `ssize_t rtdm_recvmsg ( int fd, struct msghdr * msg, int flags )`

Receive message from socket.

Refer to [rt\\_dev\\_recvmsg\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### 5.23.1.19 `int rtdm_select_bind ( int fd, rtdm_selector_t * selector, enum rtdm_selecttype type, unsigned fd_index )`

Bind a selector to specified event types of a given file descriptor.

**For internal use only.**

This function is invoked by higher RTOS layers implementing select-like services. It shall not be called directly by RTDM drivers.

#### Parameters

- [in] *fd* File descriptor to bind to
- [in,out] *selector* Selector object that shall be bound to the given event
- [in] *type* Event type the caller is interested in
- [in] *fd\_index* Index in the file descriptor set of the caller

#### Returns

0 on success, otherwise:

- -EBADF is returned if the file descriptor *fd* cannot be resolved.
- -EINVAL is returned if *type* or *fd\_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `rtdm_dev_context::ops`, `rtdm_context_get()`, `rtdm_context_unlock()`, and `rtdm_operations::select_bind`.

### 5.23.1.20 `ssize_t rtdm_send ( int fd, const void * buf, size_t len, int flags )`

Transmit message to socket.

Refer to [rt\\_dev\\_send\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.21** `ssize_t rtdm_sendmsg ( int fd, const struct msghdr * msg, int flags )`

Transmit message to socket.

Refer to [rt\\_dev\\_sendmsg\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.22** `ssize_t rtdm_sendto ( int fd, const void * buf, size_t len, int flags, const struct sockaddr * to, socklen_t to_len )`

Transmit message to socket.

Refer to [rt\\_dev\\_sendto\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.23** `int rtdm_setsockopt ( int fd, int level, int optname, const void * optval, socklen_t optlen )`

Set socket option.

Refer to [rt\\_dev\\_setsockopt\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.24** `int rtdm_shutdown ( int fd, int how )`

Shut down parts of a connection.

Refer to [rt\\_dev\\_shutdown\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**5.23.1.25** `int rtdm_socket ( int protocol_family, int socket_type, int protocol )`

Create a socket.

Refer to [rt\\_dev\\_socket\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### 5.23.1.26 ssize\_t rtdm\_write ( int fd, const void \* buf, size\_t nbyte )

Write to device.

Refer to [rt\\_dev\\_write\(\)](#) for parameters and return values

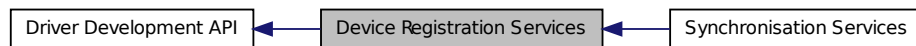
Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

## 5.24 Device Registration Services

Collaboration diagram for Device Registration Services:



### Data Structures

- struct [rtdm\\_operations](#)  
*Device operations.*
- struct [rtdm\\_dev\\_context](#)  
*Device context.*
- struct [rtdm\\_device](#)  
*RTDM device.*

### Modules

- [Synchronisation Services](#)

### Functions

- int [rtdm\\_dev\\_register](#) (struct [rtdm\\_device](#) \*device)  
*Register a RTDM device.*
- int [rtdm\\_dev\\_unregister](#) (struct [rtdm\\_device](#) \*device, unsigned int poll\_delay)  
*Unregisters a RTDM device.*
- static void \* [rtdm\\_context\\_to\\_private](#) (struct [rtdm\\_dev\\_context](#) \*context)  
*Locate the driver private area associated to a device context structure.*



- static struct `rtm_dev_context * rtm_private_to_context` (void \*dev\_private)

*Locate a device context structure from its driver private area.*

## Operation Handler Prototypes

- typedef int(\* `rtm_open_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, int oflag)

*Named device open handler.*

- typedef int(\* `rtm_socket_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, int protocol)

*Socket creation handler for protocol devices.*

- typedef int(\* `rtm_close_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info)

*Close handler.*

- typedef int(\* `rtm_ioctl_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, unsigned int request, void \_\_user \*arg)

*IOCTL handler.*

- typedef int(\* `rtm_select_bind_handler_t` )(struct `rtm_dev_context` \*context, rtm\_selector\_t \*selector, enum `rtm_selecttype` type, unsigned fd\_index)

*Select binding handler.*

- typedef ssize\_t(\* `rtm_read_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, void \*buf, size\_t nbyte)

*Read handler.*

- typedef ssize\_t(\* `rtm_write_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, const void \*buf, size\_t nbyte)

*Write handler.*

- typedef ssize\_t(\* `rtm_recvmmsg_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, struct msghdr \*msg, int flags)

*Receive message handler.*

- typedef ssize\_t(\* `rtm_sendmsg_handler_t` )(struct `rtm_dev_context` \*context, rtm\_user\_info\_t \*user\_info, const struct msghdr \*msg, int flags)

*Transmit message handler.*

## Device Flags

Static flags describing a RTDM device

- #define `RTDM_EXCLUSIVE` 0x0001

*If set, only a single instance of the device can be requested by an application.*

- `#define RTDM_NAMED_DEVICE 0x0010`  
*If set, the device is addressed via a clear-text name.*
- `#define RTDM_PROTOCOL_DEVICE 0x0020`  
*If set, the device is addressed via a combination of protocol ID and socket type.*
- `#define RTDM_DEVICE_TYPE_MASK 0x00F0`  
*Mask selecting the device type.*

## Context Flags

Dynamic flags describing the state of an open RTDM device (bit numbers)

- `#define RTDM_CREATED_IN_NRT 0`  
*Set by RTDM if the device instance was created in non-real-time context.*
- `#define RTDM_CLOSING 1`  
*Set by RTDM when the device is being closed.*
- `#define RTDM_USER_CONTEXT_FLAG 8`  
*Lowest bit number the driver developer can use freely.*

## Driver Versioning

Current revisions of RTDM structures, encoding of driver versions. See [API Versioning](#) for the interface revision.

- `#define RTDM_DEVICE_STRUCT_VER 5`  
*Version of struct `rtdm_device`.*
- `#define RTDM_CONTEXT_STRUCT_VER 3`  
*Version of struct `rtdm_dev_context`.*
- `#define RTDM_SECURE_DEVICE 0x80000000`  
*Flag indicating a secure variant of RTDM (not supported here).*
- `#define RTDM_DRIVER_VER(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))`  
*Version code constructor for driver revisions.*
- `#define RTDM_DRIVER_MAJOR_VER(ver) (((ver) >> 16) & 0xFF)`  
*Get major version number from driver revision code.*
- `#define RTDM_DRIVER_MINOR_VER(ver) (((ver) >> 8) & 0xFF)`  
*Get minor version number from driver revision code.*

- `#define RTDM_DRIVER_PATCH_VER(ver) ((ver) & 0xFF)`

*Get patch version number from driver revision code.*

## 5.24.1 Define Documentation

### 5.24.1.1 `#define RTDM_CLOSING 1`

Set by RTDM when the device is being closed.

### 5.24.1.2 `#define RTDM_CREATED_IN_NRT 0`

Set by RTDM if the device instance was created in non-real-time context.

### 5.24.1.3 `#define RTDM_DEVICE_TYPE_MASK 0x00F0`

Mask selecting the device type.

Referenced by `rtdm_dev_register()`, and `rtdm_dev_unregister()`.

### 5.24.1.4 `#define RTDM_EXCLUSIVE 0x0001`

If set, only a single instance of the device can be requested by an application.

Referenced by `rtdm_dev_register()`.

### 5.24.1.5 `#define RTDM_NAMED_DEVICE 0x0010`

If set, the device is addressed via a clear-text name.

Referenced by `rtdm_dev_register()`, and `rtdm_dev_unregister()`.

### 5.24.1.6 `#define RTDM_PROTOCOL_DEVICE 0x0020`

If set, the device is addressed via a combination of protocol ID and socket type.

Referenced by `rtdm_dev_register()`.

## 5.24.2 Typedef Documentation

### 5.24.2.1 `typedef int(* rtdm_close_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info)`

Close handler.

#### Parameters

[in] *context* Context structure associated with opened device instance

[in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode or deferred user mode call

### Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, -EAGAIN to request a recall after a grace period, or a valid negative error code according to IEEE Std 1003.1.

### Note

Drivers must be prepared for that case that the close handler is invoked more than once per open context (even if the handler already completed an earlier run successfully). The driver has to avoid releasing resources twice as well as returning false errors on successive close invocations.

### See also

close() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.24.2.2** `typedef int(* rtdm_ioctl_handler_t)(struct rtdm_dev_context *context,  
rtdm_user_info_t *user_info, unsigned int request, void __user *arg)`

IOCTL handler.

### Parameters

[in] *context* Context structure associated with opened device instance  
 [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call  
 [in] *request* Request number as passed by the user  
 [in,out] *arg* Request argument as passed by the user

### Returns

A positive value or 0 on success. On failure return either -ENOSYS, to request that the function be called again from the opposite realtime/non-realtime context, or another negative error code.

### See also

ioctl() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.24.2.3** `typedef int(* rtdm_open_handler_t)(struct rtdm_dev_context *context,  
rtdm_user_info_t *user_info, int oflag)`

Named device open handler.

### Parameters

[in] *context* Context structure associated with opened device instance  
 [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *oflag* Open flags as passed by the user

### Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

### See also

open() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.24.2.4 `typedef ssize_t(* rtdm_read_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, void *buf, size_t nbyte)`

Read handler.

### Parameters

- [in] *context* Context structure associated with opened device instance
- [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call
- [out] *buf* Input buffer as passed by the user
- [in] *nbyte* Number of bytes the user requests to read

### Returns

On success, the number of bytes read. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

### See also

read() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.24.2.5 `typedef ssize_t(* rtdm_recvmmsg_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, struct msghdr *msg, int flags)`

Receive message handler.

### Parameters

- [in] *context* Context structure associated with opened device instance
- [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call
- [in,out] *msg* Message descriptor as passed by the user, automatically mirrored to safe kernel memory in case of user mode call
- [in] *flags* Message flags as passed by the user

### Returns

On success, the number of bytes received. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

**See also**

recvmsg() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.24.2.6** `typedef int(* rtdm_select_bind_handler_t)(struct rtdm_dev_context *context, rtdm_selector_t *selector, enum rtdm_selecttype type, unsigned fd_index)`

Select binding handler.

**Parameters**

- [in] *context* Context structure associated with opened device instance
- [in,out] *selector* Object that shall be bound to the given event
- [in] *type* Event type the selector is interested in
- [in] *fd\_index* Opaque value, to be passed to rtdm\_event\_select\_bind or rtdm\_sem\_select\_bind unmodified

**Returns**

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

**5.24.2.7** `typedef ssize_t(* rtdm_sendmsg_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, const struct msghdr *msg, int flags)`

Transmit message handler.

**Parameters**

- [in] *context* Context structure associated with opened device instance
- [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call
- [in] *msg* Message descriptor as passed by the user, automatically mirrored to safe kernel memory in case of user mode call
- [in] *flags* Message flags as passed by the user

**Returns**

On success, the number of bytes transmitted. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

**See also**

sendmsg() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.24.2.8** `typedef int(* rtdm_socket_handler_t)(struct rtdm_dev_context *context,  
rtdm_user_info_t *user_info, int protocol)`

Socket creation handler for protocol devices.

#### Parameters

- [in] *context* Context structure associated with opened device instance
- [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call
- [in] *protocol* Protocol number as passed by the user

#### Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

#### See also

socket() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.24.2.9** `typedef ssize_t(* rtdm_write_handler_t)(struct rtdm_dev_context *context,  
rtdm_user_info_t *user_info, const void *buf, size_t nbyte)`

Write handler.

#### Parameters

- [in] *context* Context structure associated with opened device instance
- [in] *user\_info* Opaque pointer to information about user mode caller, NULL if kernel mode call
- [in] *buf* Output buffer as passed by the user
- [in] *nbyte* Number of bytes the user requests to write

#### Returns

On success, the number of bytes written. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

#### See also

write() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.24.3 Function Documentation

**5.24.3.1** `static void* rtdm_context_to_private ( struct rtdm_dev_context * context )  
[inline, static]`

Locate the driver private area associated to a device context structure.

**Parameters**

[in] *context* Context structure associated with opened device instance

**Returns**

The address of the private driver area associated to *context*.

References `rtm_dev_context::dev_private`.

**5.24.3.2 int rtdm\_dev\_register ( struct rtdm\_device \* device )**

Register a RTDM device.

**Parameters**

[in] *device* Pointer to structure describing the new device.

**Returns**

0 is returned upon success. Otherwise:

- -EINVAL is returned if the device structure contains invalid entries. Check kernel log in this case.
- -ENOMEM is returned if the context for an exclusive device cannot be allocated.
- -EEXIST is returned if the specified device name of protocol ID is already in use.
- -EAGAIN is returned if some /proc entry cannot be created.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References `rtm_operations::close_nrt`, `rtm_operations::close_rt`, `rtm_device::context_size`, `rtm_device::device_class`, `rtm_device::device_flags`, `rtm_device::device_name`, `rtm_device::device_sub_class`, `rtm_device::driver_version`, `rtm_device::open_rt`, `rtm_device::ops`, `rtm_device::proc_name`, `rtm_device::profile_version`, `rtm_device::protocol_family`, `rtm_device::reserved`, `RTDM_DEVICE_STRUCT_VER`, `RTDM_DEVICE_TYPE_MASK`, `RTDM_EXCLUSIVE`, `RTDM_NAMED_DEVICE`, `RTDM_PROTOCOL_DEVICE`, `rtm_operations::select_bind`, `rtm_device::socket_rt`, `rtm_device::socket_type`, and `rtm_device::struct_version`.

**5.24.3.3 int rtdm\_dev\_unregister ( struct rtdm\_device \* device, unsigned int poll\_delay )**

Unregisters a RTDM device.

**Parameters**

[in] *device* Pointer to structure describing the device to be unregistered.



[in] *poll\_delay* Polling delay in milliseconds to check repeatedly for open instances of *device*, or 0 for non-blocking mode.

### Returns

0 is returned upon success. Otherwise:

- -ENODEV is returned if the device was not registered.
- -EAGAIN is returned if the device is busy with open instances and 0 has been passed for *poll\_delay*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References `rtm_device::device_flags`, `rtm_device::device_name`, `rtm_device::protocol_family`, `rtm_device::reserved`, `RTDM_DEVICE_TYPE_MASK`, `RTDM_NAMED_DEVICE`, and `rtm_device::socket_type`.

#### 5.24.3.4 `static struct rtdm_dev_context* rtdm_private_to_context ( void * dev_private )` [static, read]

Locate a device context structure from its driver private area.

### Parameters

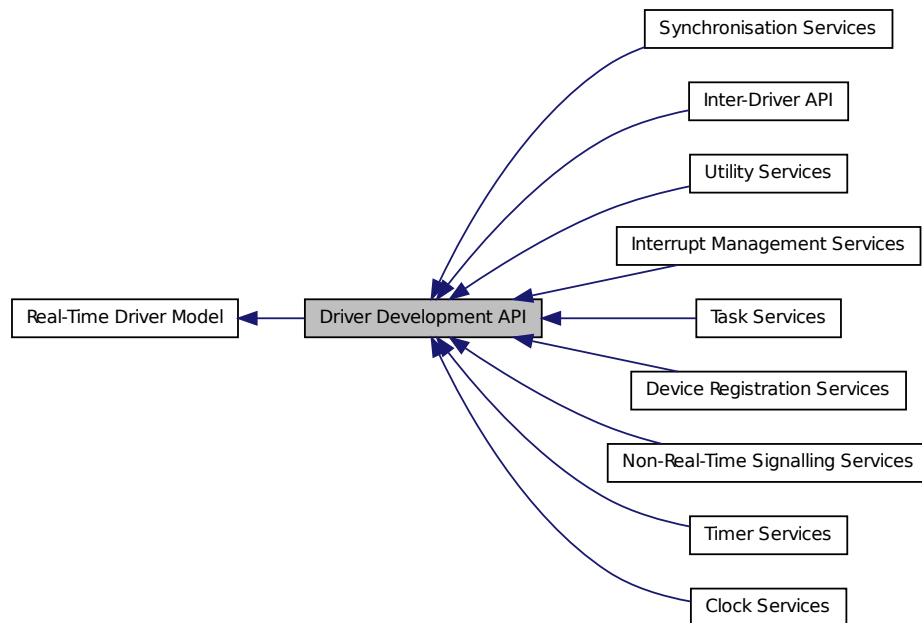
[in] *dev\_private* Address of a private context area

### Returns

The address of the device context structure defining *dev\_private*.

## 5.25 Driver Development API

Collaboration diagram for Driver Development API:



### Modules

- [Inter-Driver API](#)
- [Device Registration Services](#)
- [Clock Services](#)
- [Task Services](#)
- [Timer Services](#)
- [Synchronisation Services](#)
- [Interrupt Management Services](#)
- [Non-Real-Time Signalling Services](#)
- [Utility Services](#)

### Files

- file [rtdm\\_driver.h](#)

*Real-Time Driver Model for Xenomai, driver API header.*

### 5.25.1 Detailed Description

This is the lower interface of RTDM provided to device drivers, currently limited to kernel-space. Real-time drivers should only use functions of this interface in order to remain portable.

## 5.26 Clock Services

Collaboration diagram for Clock Services:



### Functions

- `nanosecs_abs_t rtdm_clock_read (void)`  
*Get system time.*
- `nanosecs_abs_t rtdm_clock_read_monotonic (void)`  
*Get monotonic time.*

### 5.26.1 Function Documentation

#### 5.26.1.1 `nanosecs_abs_t rtdm_clock_read ( void )`

Get system time.

#### Returns

The system time in nanoseconds is returned

#### Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period.

The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.26.1.2 `nanosecs_abs_t rtdm_clock_read_monotonic ( void )`

Get monotonic time.

##### Returns

The monotonic time in nanoseconds is returned

##### Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period.

The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:

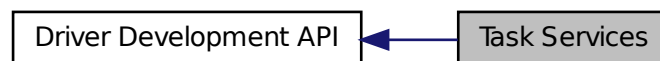
This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

## 5.27 Task Services

Collaboration diagram for Task Services:



## Typedefs

- typedef void(\* [rtdm\\_task\\_proc\\_t](#))(void \*arg)

*Real-time task procedure.*

## Functions

- int [rtdm\\_task\\_init](#) (rtdm\_task\_t \*task, const char \*name, [rtdm\\_task\\_proc\\_t](#) task\_proc, void \*arg, int priority, [nanosecs\\_rel\\_t](#) period)

*Intialise and start a real-time task.*

- void [rtdm\\_task\\_destroy](#) (rtdm\_task\_t \*task)

*Destroy a real-time task.*

- void [rtdm\\_task\\_set\\_priority](#) (rtdm\_task\_t \*task, int priority)

*Adjust real-time task priority.*

- int [rtdm\\_task\\_set\\_period](#) (rtdm\_task\_t \*task, [nanosecs\\_rel\\_t](#) period)

*Adjust real-time task period.*

- int [rtdm\\_task\\_wait\\_period](#) (void)

*Wait on next real-time task period.*

- int [rtdm\\_task\\_unblock](#) (rtdm\_task\_t \*task)

*Activate a blocked real-time task.*

- rtdm\_task\_t \* [rtdm\\_task\\_current](#) (void)

*Get current real-time task.*

- int [rtdm\\_task\\_sleep](#) ([nanosecs\\_rel\\_t](#) delay)

*Sleep a specified amount of time.*

- int [rtdm\\_task\\_sleep\\_until](#) ([nanosecs\\_abs\\_t](#) wakeup\_time)

*Sleep until a specified absolute time.*

- int [rtdm\\_task\\_sleep\\_abs](#) ([nanosecs\\_abs\\_t](#) wakeup\_time, enum [rtdm\\_timer\\_mode](#) mode)

*Sleep until a specified absolute time.*

- void [rtdm\\_task\\_join\\_nrt](#) (rtdm\_task\_t \*task, unsigned int poll\_delay)

*Wait on a real-time task to terminate.*

- void [rtdm\\_task\\_busy\\_sleep](#) ([nanosecs\\_rel\\_t](#) delay)

*Busy-wait a specified amount of time.*

## Task Priority Range

Maximum and minimum task priorities

- `#define RTDM_TASK_LOWEST_PRIORITY XNSCHED_LOW_PRIO`
- `#define RTDM_TASK_HIGHEST_PRIORITY XNSCHED_HIGH_PRIO`

## Task Priority Modification

Raise or lower task priorities by one level

- `#define RTDM_TASK_RAISE_PRIORITY (+1)`
- `#define RTDM_TASK_LOWER_PRIORITY (-1)`

### 5.27.1 Typedef Documentation

#### 5.27.1.1 `typedef void(* rtdm_task_proc_t)(void *arg)`

Real-time task procedure.

##### Parameters

[in,out] *arg* argument as passed to `rtdm_task_init()`

### 5.27.2 Function Documentation

#### 5.27.2.1 `void rtdm_task_busy_sleep ( nanosecs_rel_t delay )`

Busy-wait a specified amount of time.

##### Parameters

[in] *delay* Delay in nanoseconds. Note that a zero delay does **not** have the meaning of `RTDM_TIMEOUT_INFINITE` here.

##### Note

The caller must not be migratable to different CPUs while executing this service. Otherwise, the actual delay will be undefined.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (should be avoided or kept short)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never (except due to external interruptions).

### 5.27.2.2 `rtm_task_t* rtdm_task_current ( void )`

Get current real-time task.

#### Returns

Pointer to task handle

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.27.2.3 `void rtdm_task_destroy ( rtdm_task_t * task )`

Destroy a real-time task.

#### Parameters

[in,out] *task* Task handle as returned by [rtdm\\_task\\_init\(\)](#)

#### Note

Passing the same task handle to RTDM services after the completion of this function is not allowed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.27.2.4 `int rtdm_task_init ( rtdm_task_t * task, const char * name, rtdm_task_proc_t task_proc, void * arg, int priority, nanosecs_rel_t period )`

Initialise and start a real-time task.

After initialising a task, the task handle remains valid and can be passed to RTDM services until either [rtdm\\_task\\_destroy\(\)](#) or [rtdm\\_task\\_join\\_nrt\(\)](#) was invoked.

#### Parameters

[in,out] *task* Task handle

[in] *name* Optional task name

- [in] *task\_proc* Procedure to be executed by the task
- [in] *arg* Custom argument passed to `task_proc()` on entry
- [in] *priority* Priority of the task, see also [Task Priority Range](#)
- [in] *period* Period in nanoseconds of a cyclic task, 0 for non-cyclic mode

### Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References `xnpod_delete_thread()`, `xnpod_init_thread()`, `xnpod_set_thread_periodic()`, and `xnpod_start_thread()`.

#### 5.27.2.5 void `rtdm_task_join_nrt ( rtdm_task_t * task, unsigned int poll_delay )`

Wait on a real-time task to terminate.

### Parameters

- [in,out] *task* Task handle as returned by [rtdm\\_task\\_init\(\)](#)
- [in] *poll\_delay* Delay in milliseconds between periodic tests for the state of the real-time task. This parameter is ignored if the termination is internally realised without polling.

### Note

Passing the same task handle to RTDM services after the completion of this function is not allowed.

This service does not trigger the termination of the targeted task. The user has to take of this, otherwise [rtdm\\_task\\_join\\_nrt\(\)](#) will never return.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

References XNZOMBIE.



#### 5.27.2.6 `int rtdm_task_set_period ( rtdm_task_t * task, nanosecs_rel_t period )`

Adjust real-time task period.

##### Parameters

[in,out] *task* Task handle as returned by [rtdm\\_task\\_init\(\)](#)

[in] *period* New period in nanoseconds of a cyclic task, 0 for non-cyclic mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.27.2.7 `void rtdm_task_set_priority ( rtdm_task_t * task, int priority )`

Adjust real-time task priority.

##### Parameters

[in,out] *task* Task handle as returned by [rtdm\\_task\\_init\(\)](#)

[in] *priority* New priority of the task, see also [Task Priority Range](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.27.2.8 `int rtdm_task_sleep ( nanosecs_rel_t delay )`

Sleep a specified amount of time.

##### Parameters

[in] *delay* Delay in nanoseconds, see [RTDM\\_TIMEOUT\\_xxx](#) for special values.

### Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rt dm\\_task\\_unblock\(\)](#).
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always.

#### 5.27.2.9 `int rtdm_task_sleep_abs ( nanosecs_abs_t wakeup_time, enum rtdm_timer_mode mode )`

Sleep until a specified absolute time.

### Parameters

[in] *wakeup\_time* Absolute timeout in nanoseconds

[in] *mode* Selects the timer mode, see RTDM\_TIMERMODE\_XXX for details

### Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rt dm\\_task\\_unblock\(\)](#).
- -EPERM *may* be returned if an illegal invocation environment is detected.
- -EINVAL is returned if an invalid parameter was passed.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

### 5.27.2.10 `int rtdm_task_sleep_until ( nanosecs_abs_t wakeup_time )`

Sleep until a specified absolute time.

#### Deprecated

Use `rtdm_task_sleep_abs` instead!

#### Parameters

[in] *wakeup\_time* Absolute timeout in nanoseconds

#### Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtdm\\_task\\_unblock\(\)](#).
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

### 5.27.2.11 `int rtdm_task_unblock ( rtdm_task_t * task )`

Activate a blocked real-time task.

#### Returns

Non-zero is returned if the task was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

### 5.27.2.12 `int rtdm_task_wait_period ( void )`

Wait on next real-time task period.

#### Returns

0 on success, otherwise:

- -EINVAL is returned if calling task is not in periodic mode.
- -ETIMEDOUT is returned if a timer overrun occurred, which indicates that a previous release point has been missed by the calling task.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless a timer overrun occurred.

## 5.28 Timer Services

Collaboration diagram for Timer Services:



### Typedefs

- `typedef void(* rtdm\_timer\_handler\_t)(rtdm_timer_t *timer)`  
*Timer handler.*

### Functions

- `int rtdm\_timer\_init (rtdm_timer_t *timer, rtdm\_timer\_handler\_t handler, const char *name)`  
*Initialise a timer.*

- void `rtdm_timer_destroy` (`rtdm_timer_t *timer`)  
*Destroy a timer.*
- int `rtdm_timer_start` (`rtdm_timer_t *timer`, `nanosecs_abs_t` expiry, `nanosecs_rel_t` interval, enum `rtdm_timer_mode` mode)  
*Start a timer.*
- void `rtdm_timer_stop` (`rtdm_timer_t *timer`)  
*Stop a timer.*
- int `rtdm_timer_start_in_handler` (`rtdm_timer_t *timer`, `nanosecs_abs_t` expiry, `nanosecs_rel_t` interval, enum `rtdm_timer_mode` mode)  
*Start a timer from inside a timer handler.*
- void `rtdm_timer_stop_in_handler` (`rtdm_timer_t *timer`)  
*Stop a timer from inside a timer handler.*

## RTDM\_TIMERMODE\_XXX

Timer operation modes

- enum `rtdm_timer_mode` { `RTDM_TIMERMODE_RELATIVE` = `XN_RELATIVE`, `RTDM_TIMERMODE_ABSOLUTE` = `XN_ABSOLUTE`, `RTDM_TIMERMODE_REALTIME` = `XN_REALTIME` }

### 5.28.1 Typedef Documentation

#### 5.28.1.1 typedef void(\* rtdm\_timer\_handler\_t)(rtdm\_timer\_t \*timer)

Timer handler.

#### Parameters

[in] *timer* Timer handle as returned by `rtdm_timer_init()`

### 5.28.2 Enumeration Type Documentation

#### 5.28.2.1 enum rtdm\_timer\_mode

#### Enumerator:

*RTDM\_TIMERMODE\_RELATIVE* Monotonic timer with relative timeout.

*RTDM\_TIMERMODE\_ABSOLUTE* Monotonic timer with absolute timeout.

*RTDM\_TIMERMODE\_REALTIME* Adjustable timer with absolute timeout.

### 5.28.3 Function Documentation

#### 5.28.3.1 void rtdm\_timer\_destroy ( rtdm\_timer\_t \* *timer* )

Destroy a timer.

##### Parameters

[in,out] *timer* Timer handle as returned by [rtdm\\_timer\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xntimer\\_destroy\(\)](#).

#### 5.28.3.2 int rtdm\_timer\_init ( rtdm\_timer\_t \* *timer*, rtdm\_timer\_handler\_t *handler*, const char \* *name* )

Initialise a timer.

##### Parameters

[in,out] *timer* Timer handle

[in] *handler* Handler to be called on timer expiry

[in] *name* Optional timer name

##### Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.28.3.3 `int rtdm_timer_start ( rtdm_timer_t * timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode )`

Start a timer.

#### Parameters

- [in,out] *timer* Timer handle as returned by [rtdm\\_timer\\_init\(\)](#)
- [in] *expiry* Firing time of the timer, mode defines if relative or absolute
- [in] *interval* Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers
- [in] *mode* Defines the operation mode, see [RTDM\\_TIMERMODE\\_xxx](#) for possible values

#### Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if *expiry* describes an absolute date in the past.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xntimer_start()`.

### 5.28.3.4 `int rtdm_timer_start_in_handler ( rtdm_timer_t * timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode )`

Start a timer from inside a timer handler.

#### Parameters

- [in,out] *timer* Timer handle as returned by [rtdm\\_timer\\_init\(\)](#)
- [in] *expiry* Firing time of the timer, mode defines if relative or absolute
- [in] *interval* Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers
- [in] *mode* Defines the operation mode, see [RTDM\\_TIMERMODE\\_xxx](#) for possible values

#### Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if *expiry* describes an absolute date in the past.

Environments:

This service can be called from:

- Timer handler

Rescheduling: never.

#### 5.28.3.5 void rtdm\_timer\_stop ( rtdm\_timer\_t \* *timer* )

Stop a timer.

##### Parameters

[in,out] *timer* Timer handle as returned by [rtdm\\_timer\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xntimer\\_stop\(\)](#).

#### 5.28.3.6 void rtdm\_timer\_stop\_in\_handler ( rtdm\_timer\_t \* *timer* )

Stop a timer from inside a timer handler.

##### Parameters

[in,out] *timer* Timer handle as returned by [rtdm\\_timer\\_init\(\)](#)

Environments:

This service can be called from:

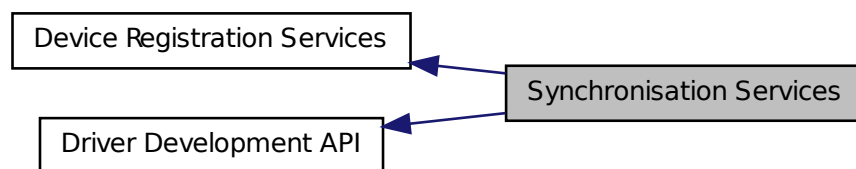
- Timer handler

Rescheduling: never.



## 5.29 Synchronisation Services

Collaboration diagram for Synchronisation Services:



### Functions

- `int rtdm_select_bind(int fd, rtdm_selector_t *selector, enum rtdm_selecttype type, unsigned fd_index)`  
*Bind a selector to specified event types of a given file descriptor.*

### RTDM\_SELECTTYPE\_xxx

Event types select can bind to

- `enum rtdm_selecttype { RTDM_SELECTTYPE_READ = XNSELECT_READ, RTDM_SELECTTYPE_WRITE = XNSELECT_WRITE, RTDM_SELECTTYPE_EXCEPT = XNSELECT_EXCEPT }`

### Spinlock with Preemption Deactivation

- `typedef ipipe_spinlock_t rtdm_lock_t`  
*Lock variable.*
- `typedef unsigned long rtdm_lockctx_t`  
*Variable to save the context while holding a lock.*
- `#define RTDM_LOCK_UNLOCKED IPIPE_SPIN_LOCK_UNLOCKED`  
*Static lock initialisation.*
- `#define rtdm_lock_init(lock) spin_lock_init(lock)`  
*Dynamic lock initialisation.*
- `#define rtdm_lock_get(lock) spin_lock(lock)`  
*Acquire lock from non-preemptible contexts.*

- #define `rt dm_lock_put(lock)` `spin_unlock(lock)`  
*Release lock without preemption restoration.*
- #define `rt dm_lock_get_irqsave(lock, context)` `spin_lock_irqsave(lock, context)`  
*Acquire lock and disable preemption.*
- #define `rt dm_lock_put_irqrestore(lock, context)` `spin_unlock_irqrestore(lock, context)`  
*Release lock and restore preemption state.*
- #define `rt dm_lock_irqsave(context)` `splhigh(context)`  
*Disable preemption locally.*
- #define `rt dm_lock_irqrestore(context)` `splexit(context)`  
*Restore preemption state.*

## Timeout Sequence Management

- void `rt dm_toseq_init (rt dm_toseq_t *timeout_seq, nanosecs_rel_t timeout)`  
*Initialise a timeout sequence.*

## Event Services

- void `rt dm_event_init (rt dm_event_t *event, unsigned long pending)`  
*Initialise an event.*
- void `rt dm_event_destroy (rt dm_event_t *event)`  
*Destroy an event.*
- void `rt dm_event_pulse (rt dm_event_t *event)`  
*Signal an event occurrence to currently listening waiters.*
- void `rt dm_event_signal (rt dm_event_t *event)`  
*Signal an event occurrence.*
- int `rt dm_event_wait (rt dm_event_t *event)`  
*Wait on event occurrence.*
- int `rt dm_event_timedwait (rt dm_event_t *event, nanosecs_rel_t timeout, rt dm_toseq_t *timeout_seq)`  
*Wait on event occurrence with timeout.*
- void `rt dm_event_clear (rt dm_event_t *event)`  
*Clear event state.*
- int `rt dm_event_select_bind (rt dm_event_t *event, rt dm_selector_t *selector, enum rt dm_selecttype type, unsigned fd_index)`  
*Bind a selector to an event.*

## Semaphore Services

- void `rt dm_sem_init` (rt dm\_sem\_t \*sem, unsigned long value)  
*Initialise a semaphore.*
- void `rt dm_sem_destroy` (rt dm\_sem\_t \*sem)  
*Destroy a semaphore.*
- int `rt dm_sem_down` (rt dm\_sem\_t \*sem)  
*Decrement a semaphore.*
- int `rt dm_sem_timeddown` (rt dm\_sem\_t \*sem, `nanosecs_rel_t` timeout, rt dm\_toseq\_t \*timeout\_seq)  
*Decrement a semaphore with timeout.*
- void `rt dm_sem_up` (rt dm\_sem\_t \*sem)  
*Increment a semaphore.*
- int `rt dm_sem_select_bind` (rt dm\_sem\_t \*sem, rt dm\_selector\_t \*selector, enum `rt dm_selecttype` type, unsigned fd\_index)  
*Bind a selector to a semaphore.*

## Mutex Services

- void `rt dm_mutex_init` (rt dm\_mutex\_t \*mutex)  
*Initialise a mutex.*
- void `rt dm_mutex_destroy` (rt dm\_mutex\_t \*mutex)  
*Destroy a mutex.*
- void `rt dm_mutex_unlock` (rt dm\_mutex\_t \*mutex)  
*Release a mutex.*
- int `rt dm_mutex_lock` (rt dm\_mutex\_t \*mutex)  
*Request a mutex.*
- int `rt dm_mutex_timedlock` (rt dm\_mutex\_t \*mutex, `nanosecs_rel_t` timeout, rt dm\_toseq\_t \*timeout\_seq)  
*Request a mutex with timeout.*

## Global Lock across Scheduler Invocation

- #define `RTDM_EXECUTE_ATOMICALY`(code\_block)  
*Execute code block atomically.*

## 5.29.1 Define Documentation

### 5.29.1.1 `#define RTDM_EXECUTE_ATOMICALY( code_block )`

**Value:**

```
{
    <ENTER_ATOMIC_SECTION>
    code_block;
    <LEAVE_ATOMIC_SECTION>
}
```

Execute code block atomically.

Generally, it is illegal to suspend the current task by calling `rtdm_task_sleep()`, `rtdm_event_wait()`, etc. while holding a spinlock. In contrast, this macro allows to combine several operations including a potentially rescheduling call to an atomic code block with respect to other `RTDM_EXECUTE_ATOMICALY()` blocks. The macro is a light-weight alternative for protecting code blocks via mutexes, and it can even be used to synchronise real-time and non-real-time contexts.

#### Parameters

*code\_block* Commands to be executed atomically

#### Note

It is not allowed to leave the code block explicitly by using `break`, `return`, `goto`, etc. This would leave the global lock held during the code block execution in an inconsistent state. Moreover, do not embed complex operations into the code block. Consider that they will be executed under preemption lock with interrupts switched-off. Also note that invocation of rescheduling calls may break the atomicity until the task gains the CPU again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible, depends on functions called within *code\_block*.

### 5.29.1.2 `#define rtdm_lock_get( lock ) spin_lock(lock)`

Acquire lock from non-preemptible contexts.

#### Parameters

*lock* Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.29.1.3 `#define rtdm_lock_get_irqsave( lock, context ) spin_lock_irqsave(lock, context)`

Acquire lock and disable preemption.

##### Parameters

*lock* Address of lock variable

*context* name of local variable to store the context in

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.29.1.4 `#define rtdm_lock_init( lock ) spin_lock_init(lock)`

Dynamic lock initialisation.

##### Parameters

*lock* Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.29.1.5 `#define rtdm_lock_irqrestore( context ) splexit(context)`

Restore preemption state.

##### Parameters

*context* name of local variable which stored the context

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.29.1.6 `#define rtdm_lock_irqsave( context ) splhigh(context)`

Disable preemption locally.

##### Parameters

*context* name of local variable to store the context in

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.29.1.7 `#define rtdm_lock_put( lock ) spin_unlock(lock)`

Release lock without preemption restoration.

##### Parameters

*lock* Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.29.1.8 `#define rtdm_lock_put_irqrestore( lock, context ) spin_unlock_irqrestore(lock, context)`

Release lock and restore preemption state.

##### Parameters

*lock* Address of lock variable

*context* name of local variable which stored the context

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

## 5.29.2 Enumeration Type Documentation

### 5.29.2.1 `enum rtdm_selecttype`

#### Enumerator:

*RTDM\_SELECTTYPE\_READ* Select input data availability events.

*RTDM\_SELECTTYPE\_WRITE* Select output buffer availability events.

*RTDM\_SELECTTYPE\_EXCEPT* Select exceptional events.

## 5.29.3 Function Documentation

### 5.29.3.1 `void rtdm_event_clear ( rtdm_event_t * event )`

Clear event state.

#### Parameters

[in,out] *event* Event handle as returned by [rtdm\\_event\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_signal()`.

### 5.29.3.2 `void rtdm_event_destroy ( rtdm_event_t * event )`

Destroy an event.

#### Parameters

[in,out] *event* Event handle as returned by `rtdm_event_init()`

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

### 5.29.3.3 `void rtdm_event_init ( rtdm_event_t * event, unsigned long pending )`

Initialise an event.

#### Parameters

[in,out] *event* Event handle

[in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_init()`, and `xnsynch_init()`.



#### 5.29.3.4 void rtdm\_event\_pulse ( rtdm\_event\_t \* event )

Signal an event occurrence to currently listening waiters.

This function wakes up all current waiters of the given event, but it does not change the event state. Subsequently callers of [rtdm\\_event\\_wait\(\)](#) or [rtdm\\_event\\_timedwait\(\)](#) will therefore be blocked first.

##### Parameters

[in,out] *event* Event handle as returned by [rtdm\\_event\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.29.3.5 int rtdm\_event\_select\_bind ( rtdm\_event\_t \* event, rtdm\_selector\_t \* selector, enum rtdm\_selecttype type, unsigned fd\_index )

Bind a selector to an event.

This functions binds the given selector to an event so that the former is notified when the event state changes. Typically the select binding handler will invoke this service.

##### Parameters

[in,out] *event* Event handle as returned by [rtdm\\_event\\_init\(\)](#)

[in,out] *selector* Selector as passed to the select binding handler

[in] *type* Type of the bound event as passed to the select binding handler

[in] *fd\_index* File descriptor index as passed to the select binding handler

##### Returns

0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.
- -EINVAL is returned if *type* or *fd\_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_bind()`.

#### 5.29.3.6 `void rtdm_event_signal ( rtdm_event_t * event )`

Signal an event occurrence.

This function sets the given event and wakes up all current waiters. If no waiter is presently registered, the next call to `rtdm_event_wait()` or `rtdm_event_timedwait()` will return immediately.

##### Parameters

[in,out] *event* Event handle as returned by `rtdm_event_init()`

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References `xnpod_schedule()`, `xnselect_signal()`, and `xnsynch_flush()`.

#### 5.29.3.7 `int rtdm_event_timedwait ( rtdm_event_t * event, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq )`

Wait on event occurrence with timeout.

This function waits or tests for the occurrence of the given event, taking the provided timeout into account. On successful return, the event is reset.

##### Parameters

[in,out] *event* Event handle as returned by `rtdm_event_init()`

[in] *timeout* Relative timeout in nanoseconds, see `RTDM_TIMEOUT_xxx` for special values

[in,out] *timeout\_seq* Handle of a timeout sequence as returned by `rtdm_toseq_init()` or NULL

##### Returns

0 on success, otherwise:

- `-ETIMEDOUT` is returned if the request has not been satisfied within the specified amount of time.

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rt dm\\_task\\_unblock\(\)](#).
- -EIDRM is returned if *event* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.
- -EWOULDBLOCK is returned if a negative *timeout* (i.e., non-blocking operation) has been specified.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References XNBREAK, XNRMID, xselect\_signal(), xnsynch\_sleep\_on(), and XNTIMEO.

Referenced by [rt dm\\_event\\_wait\(\)](#).

#### 5.29.3.8 int [rt dm\\_event\\_wait](#) ( [rt dm\\_event\\_t](#) \* *event* )

Wait on event occurrence.

This is the light-weight version of [rt dm\\_event\\_timedwait\(\)](#), implying an infinite timeout.

##### Parameters

[in,out] *event* Event handle as returned by [rt dm\\_event\\_init\(\)](#)

##### Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rt dm\\_task\\_unblock\(\)](#).
- -EIDRM is returned if *event* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [rt dm\\_event\\_timedwait\(\)](#).

### 5.29.3.9 void rtdm\_mutex\_destroy ( rtdm\_mutex\_t \* mutex )

Destroy a mutex.

#### Parameters

[in,out] *mutex* Mutex handle as returned by [rtdm\\_mutex\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

### 5.29.3.10 void rtdm\_mutex\_init ( rtdm\_mutex\_t \* mutex )

Initialise a mutex.

This function initialises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

#### Parameters

[in,out] *mutex* Mutex handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xnsynch\\_init\(\)](#).

### 5.29.3.11 int rtdm\_mutex\_lock ( rtdm\_mutex\_t \* mutex )

Request a mutex.

This is the light-weight version of [rtdm\\_mutex\\_timedlock\(\)](#), implying an infinite timeout.

#### Parameters

[in,out] *mutex* Mutex handle as returned by [rtdm\\_mutex\\_init\(\)](#)

**Returns**

0 on success, otherwise:

- -EIDRM is returned if *mutex* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References `rtdm_mutex_timedlock()`.

**5.29.3.12** `int rtdm_mutex_timedlock ( rtdm_mutex_t * mutex, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq )`

Request a mutex with timeout.

This function tries to acquire the given mutex. If it is not available, the caller is blocked unless non-blocking operation was selected.

**Parameters**

[in,out] *mutex* Mutex handle as returned by `rtdm_mutex_init()`

[in] *timeout* Relative timeout in nanoseconds, see `RTDM_TIMEOUT_XXX` for special values

[in,out] *timeout\_seq* Handle of a timeout sequence as returned by `rtdm_toseq_init()` or NULL

**Returns**

0 on success, otherwise:

- -ETIMEDOUT is returned if the request has not been satisfied within the specified amount of time.
- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.
- -EIDRM is returned if *mutex* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

References XNBREAK, XNRMID, xnsynch\_acquire(), and XNTIMEO.

Referenced by `rtdm_mutex_lock()`.

### 5.29.3.13 void rtdm\_mutex\_unlock ( rtdm\_mutex\_t \* *mutex* )

Release a mutex.

This function releases the given mutex, waking up a potential waiter which was blocked upon `rtdm_mutex_lock()` or `rtdm_mutex_timedlock()`.

#### Parameters

[in,out] *mutex* Mutex handle as returned by `rtdm_mutex_init()`

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

### 5.29.3.14 int rtdm\_select\_bind ( int *fd*, rtdm\_selector\_t \* *selector*, enum rtdm\_selecttype *type*, unsigned *fd\_index* )

Bind a selector to specified event types of a given file descriptor.

#### For internal use only.

This function is invoked by higher RTOS layers implementing select-like services. It shall not be called directly by RTDM drivers.

#### Parameters

[in] *fd* File descriptor to bind to

[in,out] *selector* Selector object that shall be bound to the given event

[in] *type* Event type the caller is interested in

[in] *fd\_index* Index in the file descriptor set of the caller

#### Returns

0 on success, otherwise:

- -EBADF is returned if the file descriptor *fd* cannot be resolved.
- -EINVAL is returned if *type* or *fd\_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `rtm_dev_context::ops`, `rtm_context_get()`, `rtm_context_unlock()`, and `rtm_operations::select_bind`.

#### 5.29.3.15 `void rtm_sem_destroy ( rtm_sem_t * sem )`

Destroy a semaphore.

##### Parameters

[in,out] *sem* Semaphore handle as returned by [rtm\\_sem\\_init\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.29.3.16 `int rtm_sem_down ( rtm_sem_t * sem )`

Decrement a semaphore.

This is the light-weight version of [rtm\\_sem\\_timeddown\(\)](#), implying an infinite timeout.

##### Parameters

[in,out] *sem* Semaphore handle as returned by [rtm\\_sem\\_init\(\)](#)

##### Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtm\\_task\\_unblock\(\)](#).
- -EIDRM is returned if *sem* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References `rtdm_sem_timeddown()`.

#### 5.29.3.17 `void rtdm_sem_init ( rtdm_sem_t * sem, unsigned long value )`

Initialise a semaphore.

##### Parameters

- [in,out] *sem* Semaphore handle
- [in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_init()`, and `xnsynch_init()`.

#### 5.29.3.18 `int rtdm_sem_select_bind ( rtdm_sem_t * sem, rtdm_selector_t * selector, enum rtdm_selecttype type, unsigned fd_index )`

Bind a selector to a semaphore.

This functions binds the given selector to the semaphore so that the former is notified when the semaphore state changes. Typically the select binding handler will invoke this service.

##### Parameters

- [in,out] *sem* Semaphore handle as returned by [rtdm\\_sem\\_init\(\)](#)
- [in,out] *selector* Selector as passed to the select binding handler
- [in] *type* Type of the bound event as passed to the select binding handler
- [in] *fd\_index* File descriptor index as passed to the select binding handler

##### Returns

0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.



- -EINVAL is returned if *type* or *fd\_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_bind()`.

**5.29.3.19** `int rtdm_sem_timeddown ( rtdm_sem_t * sem, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq )`

Decrement a semaphore with timeout.

This function tries to decrement the given semaphore's value if it is positive on entry. If not, the caller is blocked unless non-blocking operation was selected.

#### Parameters

- [in,out] *sem* Semaphore handle as returned by [rtdm\\_sem\\_init\(\)](#)
- [in] *timeout* Relative timeout in nanoseconds, see [RTDM\\_TIMEOUT\\_xxx](#) for special values
- [in,out] *timeout\_seq* Handle of a timeout sequence as returned by [rtdm\\_toseq\\_init\(\)](#) or NULL

#### Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the request has not been satisfied within the specified amount of time.
- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.
- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtdm\\_task\\_unblock\(\)](#).
- -EIDRM is returned if *sem* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

References `sem_timedwait()`, `XNBREAK`, `XNRMID`, `xnselect_signal()`, `xnsynch_sleep_on()`, and `XNTIMEO`.

Referenced by `rtdm_sem_down()`.

### 5.29.3.20 void rtdm\_sem\_up ( rtdm\_sem\_t \* sem )

Increment a semaphore.

This function increments the given semaphore's value, waking up a potential waiter which was blocked upon `rtdm_sem_down()`.

#### Parameters

[in,out] *sem* Semaphore handle as returned by `rtdm_sem_init()`

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References `xnpod_schedule()`, `xnselect_signal()`, and `xnsynch_wakeup_one_sleeper()`.

### 5.29.3.21 void rtdm\_toseq\_init ( rtdm\_toseq\_t \* timeout\_seq, nanosecs\_rel\_t timeout )

Initialise a timeout sequence.

This service initialises a timeout sequence handle according to the given timeout value. Timeout sequences allow to maintain a continuous *timeout* across multiple calls of blocking synchronisation services. A typical application scenario is given below.

#### Parameters

[in,out] *timeout\_seq* Timeout sequence handle

[in] *timeout* Relative timeout in nanoseconds, see `RTDM_TIMEOUT_XXX` for special values

Application Scenario:

```
int device_service_routine(...)
{
    rtdm_toseq_t timeout_seq;
    ...

    rtdm_toseq_init(&timeout_seq, timeout);
}
```

```

...
while (received < requested) {
    ret = rtdm_event_timedwait(&data_available, timeout, &timeout_seq
);
    if (ret < 0) // including -ETIMEDOUT
        break;

    // receive some data
    ...
}
...
}

```

Using a timeout sequence in such a scenario avoids that the user-provided relative `timeout` is restarted on every call to `rtdm_event_timedwait()`, potentially causing an overall delay that is larger than specified by `timeout`. Moreover, all functions supporting timeout sequences also interpret special timeout values (infinite and non-blocking), disburdening the driver developer from handling them separately.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: never.

## 5.30 Interrupt Management Services

Collaboration diagram for Interrupt Management Services:



### Defines

- #define `rtdm_irq_get_arg`(`irq_handle`, `type`) ((`type *`)`irq_handle->cookie`)  
*Retrieve IRQ handler argument.*

### Typedefs

- typedef int(\* `rtdm_irq_handler_t`)(`rtdm_irq_t *``irq_handle`)  
*Interrupt handler.*

## Functions

- int `rt dm_irq_request` (`rt dm_irq_t` \*irq\_handle, unsigned int irq\_no, `rt dm_irq_handler_t` handler, unsigned long flags, const char \*device\_name, void \*arg)  
*Register an interrupt handler.*
- int `rt dm_irq_free` (`rt dm_irq_t` \*irq\_handle)  
*Release an interrupt handler.*
- int `rt dm_irq_enable` (`rt dm_irq_t` \*irq\_handle)  
*Enable interrupt line.*
- int `rt dm_irq_disable` (`rt dm_irq_t` \*irq\_handle)  
*Disable interrupt line.*

## RTDM\_IRQTYPE\_XXX

Interrupt registrations flags

- #define `RTDM_IRQTYPE_SHARED` XN\_ISR\_SHARED  
*Enable IRQ-sharing with other real-time drivers.*
- #define `RTDM_IRQTYPE_EDGE` XN\_ISR\_EDGE  
*Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.*

## RTDM\_IRQ\_XXX

Return flags of interrupt handlers

- #define `RTDM_IRQ_NONE` XN\_ISR\_NONE  
*Unhandled interrupt.*
- #define `RTDM_IRQ_HANDLED` XN\_ISR\_HANDLED  
*Denote handled interrupt.*

### 5.30.1 Define Documentation

#### 5.30.1.1 #define `rt dm_irq_get_arg( irq_handle, type )` ((type \*)irq\_handle->cookie)

Retrieve IRQ handler argument.

#### Parameters

*irq\_handle* IRQ handle

*type* Type of the pointer to return

**Returns**

The argument pointer registered on [rt dm\\_irq\\_request\(\)](#) is returned, type-casted to the specified *type*.

Environments:

This service can be called from:

- Interrupt service routine

Rescheduling: never.

### 5.30.2 Typedef Documentation

#### 5.30.2.1 typedef int(\* rt dm\_irq\_handler\_t)(rt dm\_irq\_t \*irq\_handle)

Interrupt handler.

**Parameters**

[in] *irq\_handle* IRQ handle as returned by [rt dm\\_irq\\_request\(\)](#)

**Returns**

0 or a combination of [RTDM\\_IRQ\\_XXX](#) flags

### 5.30.3 Function Documentation

#### 5.30.3.1 int rt dm\_irq\_disable ( rt dm\_irq\_t \* irq\_handle )

Disable interrupt line.

**Parameters**

[in,out] *irq\_handle* IRQ handle as returned by [rt dm\\_irq\\_request\(\)](#)

**Returns**

0 on success, otherwise negative error code

**Note**

This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.30.3.2 `int rtdm_irq_enable ( rtdm_irq_t * irq_handle )`

Enable interrupt line.

#### Parameters

[in,out] *irq\_handle* IRQ handle as returned by [rtdm\\_irq\\_request\(\)](#)

#### Returns

0 on success, otherwise negative error code

#### Note

This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

### 5.30.3.3 `int rtdm_irq_free ( rtdm_irq_t * irq_handle )`

Release an interrupt handler.

#### Parameters

[in,out] *irq\_handle* IRQ handle as returned by [rtdm\\_irq\\_request\(\)](#)

#### Returns

0 on success, otherwise negative error code

#### Note

The caller is responsible for shutting down the IRQ source at device level before invoking this service. In turn, `rtdm_irq_free` ensures that any pending event on the given IRQ line is fully processed on return from this service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.30.3.4** `int rtdm_irq_request ( rtdm_irq_t * irq_handle, unsigned int irq_no,  
rtdm_irq_handler_t handler, unsigned long flags, const char * device_name, void *  
arg )`

Register an interrupt handler.

This function registers the provided handler with an IRQ line and enables the line.

#### Parameters

- [in,out] *irq\_handle* IRQ handle
- [in] *irq\_no* Line number of the addressed IRQ
- [in] *handler* Interrupt handler
- [in] *flags* Registration flags, see [RTDM\\_IRQTYPE\\_xxx](#) for details
- [in] *device\_name* Device name to show up in real-time IRQ lists
- [in] *arg* Pointer to be passed to the interrupt handler on invocation

#### Returns

0 on success, otherwise:

- -EINVAL is returned if an invalid parameter was passed.
- -EBUSY is returned if the specified IRQ line is already in use.

Environments:

This service can be called from:

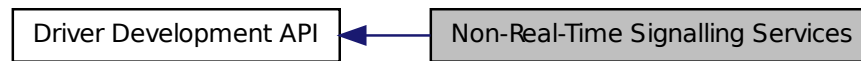
- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnintr_attach()`, `xnintr_enable()`, and `xnintr_init()`.

## 5.31 Non-Real-Time Signalling Services

Collaboration diagram for Non-Real-Time Signalling Services:



### Typedefs

- typedef void(\* [rtdm\\_nrtsig\\_handler\\_t](#))(rtdm\_nrtsig\_t nrt\_sig, void \*arg)  
*Non-real-time signal handler.*

### Functions

- int [rtdm\\_nrtsig\\_init](#) (rtdm\_nrtsig\_t \*nrt\_sig, [rtdm\\_nrtsig\\_handler\\_t](#) handler, void \*arg)  
*Register a non-real-time signal handler.*
- void [rtdm\\_nrtsig\\_destroy](#) (rtdm\_nrtsig\_t \*nrt\_sig)  
*Release a non-realtime signal handler.*
- void [rtdm\\_nrtsig\\_pend](#) (rtdm\_nrtsig\_t \*nrt\_sig)  
*Trigger non-real-time signal.*

#### 5.31.1 Detailed Description

These services provide a mechanism to request the execution of a specified handler in non-real-time context. The triggering can safely be performed in real-time context without suffering from unknown delays. The handler execution will be deferred until the next time the real-time subsystem releases the CPU to the non-real-time part.

#### 5.31.2 Typedef Documentation

##### 5.31.2.1 typedef void(\* rtdm\_nrtsig\_handler\_t)(rtdm\_nrtsig\_t nrt\_sig, void \*arg)

Non-real-time signal handler.

##### Parameters

- [in] *nrt\_sig* Signal handle as returned by [rtdm\\_nrtsig\\_init\(\)](#)
- [in] *arg* Argument as passed to [rtdm\\_nrtsig\\_init\(\)](#)



**Note**

The signal handler will run in soft-IRQ context of the non-real-time subsystem. Note the implications of this context, e.g. no invocation of blocking operations.

**5.31.3 Function Documentation****5.31.3.1 void rtdm\_nrtsig\_destroy ( rtdm\_nrtsig\_t \* *nrt\_sig* )**

Release a non-realtime signal handler.

**Parameters**

[in,out] *nrt\_sig* Signal handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.31.3.2 int rtdm\_nrtsig\_init ( rtdm\_nrtsig\_t \* *nrt\_sig*, rtdm\_nrtsig\_handler\_t *handler*, void \* *arg* )**

Register a non-real-time signal handler.

**Parameters**

[in,out] *nrt\_sig* Signal handle

[in] *handler* Non-real-time signal handler

[in] *arg* Custom argument passed to handler() on each invocation

**Returns**

0 on success, otherwise:

- -EAGAIN is returned if no free signal slot is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.31.3.3 void rtdm\_nrtsig\_pend ( rtdm\_nrtsig\_t \* nrt\_sig )

Trigger non-real-time signal.

#### Parameters

[in,out] *nrt\_sig* Signal handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

## 5.32 Utility Services

Collaboration diagram for Utility Services:



### Functions

- int [rtdm\\_mmap\\_to\\_user](#) (rtdm\_user\_info\_t \*user\_info, void \*src\_addr, size\_t len, int prot, void \*\*pptr, struct vm\_operations\_struct \*vm\_ops, void \*vm\_private\_data)  
*Map a kernel memory range into the address space of the user.*
- int [rtdm\\_iomap\\_to\\_user](#) (rtdm\_user\_info\_t \*user\_info, phys\_addr\_t src\_addr, size\_t len, int prot, void \*\*pptr, struct vm\_operations\_struct \*vm\_ops, void \*vm\_private\_data)  
*Map an I/O memory range into the address space of the user.*
- int [rtdm\\_munmap](#) (rtdm\_user\_info\_t \*user\_info, void \*ptr, size\_t len)  
*Unmap a user memory range.*
- void [rtdm\\_printk](#) (const char \*format,...)  
*Real-time safe message printing on kernel console.*

- void \* [rtm\\_malloc](#) (size\_t size)  
*Allocate memory block in real-time context.*
- void [rtm\\_free](#) (void \*ptr)  
*Release real-time memory block.*
- int [rtm\\_read\\_user\\_ok](#) (rtm\_user\_info\_t \*user\_info, const void \_\_user \*ptr, size\_t size)  
*Check if read access to user-space memory block is safe.*
- int [rtm\\_rw\\_user\\_ok](#) (rtm\_user\_info\_t \*user\_info, const void \_\_user \*ptr, size\_t size)  
*Check if read/write access to user-space memory block is safe.*
- int [rtm\\_copy\\_from\\_user](#) (rtm\_user\_info\_t \*user\_info, void \*dst, const void \_\_user \*src, size\_t size)  
*Copy user-space memory block to specified buffer.*
- int [rtm\\_safe\\_copy\\_from\\_user](#) (rtm\_user\_info\_t \*user\_info, void \*dst, const void \_\_user \*src, size\_t size)  
*Check if read access to user-space memory block and copy it to specified buffer.*
- int [rtm\\_copy\\_to\\_user](#) (rtm\_user\_info\_t \*user\_info, void \_\_user \*dst, const void \*src, size\_t size)  
*Copy specified buffer to user-space memory block.*
- int [rtm\\_safe\\_copy\\_to\\_user](#) (rtm\_user\_info\_t \*user\_info, void \_\_user \*dst, const void \*src, size\_t size)  
*Check if read/write access to user-space memory block is safe and copy specified buffer to it.*
- int [rtm\\_strncpy\\_from\\_user](#) (rtm\_user\_info\_t \*user\_info, char \*dst, const char \_\_user \*src, size\_t count)  
*Copy user-space string to specified buffer.*
- int [rtm\\_in\\_rt\\_context](#) (void)  
*Test if running in a real-time task.*
- int [rtm\\_rt\\_capable](#) (rtm\_user\_info\_t \*user\_info)  
*Test if the caller is capable of running in real-time context.*

### 5.32.1 Function Documentation

#### 5.32.1.1 int rtm\_copy\_from\_user ( rtm\_user\_info\_t \* user\_info, void \* dst, const void \_\_user \* src, size\_t size )

Copy user-space memory block to specified buffer.

##### Parameters

[in] *user\_info* User information pointer as passed to the invoked device operation handler

[in] *dst* Destination buffer address  
 [in] *src* Address of the user-space memory block  
 [in] *size* Size of the memory block

### Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

### Note

Before invoking this service, verify via [rtm\\_read\\_user\\_ok\(\)](#) that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.2** `int rtdm_copy_to_user ( rtdm_user_info_t * user_info, void __user * dst, const void * src, size_t size )`

Copy specified buffer to user-space memory block.

### Parameters

[in] *user\_info* User information pointer as passed to the invoked device operation handler  
 [in] *dst* Address of the user-space memory block  
 [in] *src* Source buffer address  
 [in] *size* Size of the memory block

### Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

### Note

Before invoking this service, verify via [rtm\\_rw\\_user\\_ok\(\)](#) that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.32.1.3 void rtdm\_free ( void \* ptr )

Release real-time memory block.

##### Parameters

[in] *ptr* Pointer to memory block as returned by [rtdm\\_malloc\(\)](#)

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.32.1.4 int rtdm\_in\_rt\_context ( void )

Test if running in a real-time task.

##### Returns

Non-zero is returned if the caller resides in real-time context, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.5** `int rtdm_iomap_to_user ( rtdm_user_info_t * user_info, phys_addr_t src_addr, size_t len, int prot, void ** pptr, struct vm_operations_struct * vm_ops, void * vm_private_data )`

Map an I/O memory range into the address space of the user.

#### Parameters

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *src\_addr* physical I/O address to be mapped
- [in] *len* Length of the memory range
- [in] *prot* Protection flags for the user's memory range, typically either PROT\_READ or PROT\_READ|PROT\_WRITE
- [in,out] *pptr* Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return
- [in] *vm\_ops* vm\_operations to be executed on the vma\_area of the user memory range or NULL
- [in] *vm\_private\_data* Private data to be stored in the vma\_area, primarily useful for vm\_-operation handlers

#### Returns

- 0 on success, otherwise (most common values):
- -EINVAL is returned if an invalid start address, size, or destination address was passed.
- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.
- -EAGAIN is returned if too much memory has been already locked by the user process.
- -EPERM *may* be returned if an illegal invocation environment is detected.

#### Note

RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via `rtdm_munmap()`, either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard `munmap()` or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of `rtdm_iomap_to_user()` can pass a `vm_operations_struct` on invocation, defining a close handler for the `vm_area`. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

**5.32.1.6 void\* rtdm\_malloc ( size\_t size )**

Allocate memory block in real-time context.

**Parameters**

[in] *size* Requested size of the memory block

**Returns**

The pointer to the allocated block is returned on success, NULL otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.7 int rtdm\_mmap\_to\_user ( rtdm\_user\_info\_t \* user\_info, void \* src\_addr, size\_t len, int prot, void \*\* pptr, struct vm\_operations\_struct \* vm\_ops, void \* vm\_private\_data )**

Map a kernel memory range into the address space of the user.

**Parameters**

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *src\_addr* Kernel virtual address to be mapped
- [in] *len* Length of the memory range
- [in] *prot* Protection flags for the user's memory range, typically either PROT\_READ or PROT\_READ|PROT\_WRITE
- [in,out] *pptr* Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return
- [in] *vm\_ops* vm\_operations to be executed on the vma\_area of the user memory range or NULL
- [in] *vm\_private\_data* Private data to be stored in the vma\_area, primarily useful for vm\_-operation handlers

**Returns**

0 on success, otherwise (most common values):

- -EINVAL is returned if an invalid start address, size, or destination address was passed.
- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.

- -EAGAIN is returned if too much memory has been already locked by the user process.
- -EPERM *may* be returned if an illegal invocation environment is detected.

#### Note

This service only works on memory regions allocated via `kmalloc()` or `vmalloc()`. To map physical I/O memory to user-space use `rt dm_iomap_to_user()` instead.

RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via `rt dm_munmap()`, either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard `munmap()` or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of `rt dm_mmap_to_user()` can pass a `vm_operations_struct` on invocation, defining a close handler for the `vm_area`. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

#### 5.32.1.8 `int rt dm_munmap ( rt dm_user_info_t * user_info, void * ptr, size_t len )`

Unmap a user memory range.

#### Parameters

- [in] *user\_info* User information pointer as passed to `rt dm_mmap_to_user()` when requesting to map the memory range
- [in] *ptr* User address or the memory range
- [in] *len* Length of the memory range

#### Returns

0 on success, otherwise:

- -EINVAL is returned if an invalid address or size was passed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.



#### 5.32.1.9 void rtdm\_printk ( const char \* *format*, ... )

Real-time safe message printing on kernel console.

##### Parameters

- [in] *format* Format string (conforming standard printf())
- ... Arguments referred by *format*

##### Returns

On success, this service returns the number of characters printed. Otherwise, a negative error code is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

#### 5.32.1.10 int rtdm\_read\_user\_ok ( rtdm\_user\_info\_t \* *user\_info*, const void \_\_user \* *ptr*, size\_t *size* )

Check if read access to user-space memory block is safe.

##### Parameters

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *ptr* Address of the user-provided memory block
- [in] *size* Size of the memory block

##### Returns

Non-zero is return when it is safe to read from the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.32.1.11 `int rtdm_rt_capable ( rtdm_user_info_t * user_info )`

Test if the caller is capable of running in real-time context.

#### Parameters

[in] *user\_info* User information pointer as passed to the invoked device operation handler

#### Returns

Non-zero is returned if the caller is able to execute in real-time context (independent of its current execution mode), 0 otherwise.

#### Note

This function can be used by drivers that provide different implementations for the same service depending on the execution mode of the caller. If a caller requests such a service in non-real-time context but is capable of running in real-time as well, it might be appropriate for the driver to reject the request via -ENOSYS so that RTDM can switch the caller and restart the request in real-time context.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.32.1.12 `int rtdm_rw_user_ok ( rtdm_user_info_t * user_info, const void __user * ptr, size_t size )`

Check if read/write access to user-space memory block is safe.

#### Parameters

[in] *user\_info* User information pointer as passed to the invoked device operation handler  
[in] *ptr* Address of the user-provided memory block  
[in] *size* Size of the memory block

#### Returns

Non-zero is return when it is safe to read from or write to the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.13** `int rtdm_safe_copy_from_user ( rtdm_user_info_t * user_info, void * dst, const void __user * src, size_t size )`

Check if read access to user-space memory block and copy it to specified buffer.

#### Parameters

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *dst* Destination buffer address
- [in] *src* Address of the user-space memory block
- [in] *size* Size of the memory block

#### Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

#### Note

This service is a combination of `rtdm_read_user_ok` and `rtdm_copy_from_user`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.14** `int rtdm_safe_copy_to_user ( rtdm_user_info_t * user_info, void __user * dst, const void * src, size_t size )`

Check if read/write access to user-space memory block is safe and copy specified buffer to it.

#### Parameters

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *dst* Address of the user-space memory block
- [in] *src* Source buffer address
- [in] *size* Size of the memory block

#### Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

**Note**

This service is a combination of `rtm_rw_user_ok` and `rtm_copy_to_user`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.32.1.15** `int rtdm_strncpy_from_user ( rtdm_user_info_t * user_info, char * dst, const char __user * src, size_t count )`

Copy user-space string to specified buffer.

**Parameters**

- [in] *user\_info* User information pointer as passed to the invoked device operation handler
- [in] *dst* Destination buffer address
- [in] *src* Address of the user-space string
- [in] *count* Maximum number of bytes to copy, including the trailing '0'

**Returns**

Length of the string on success (not including the trailing '0'), otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

**Note**

This services already includes a check of the source address, calling `rtm_read_user_ok()` for *src* explicitly is not required.

Environments:

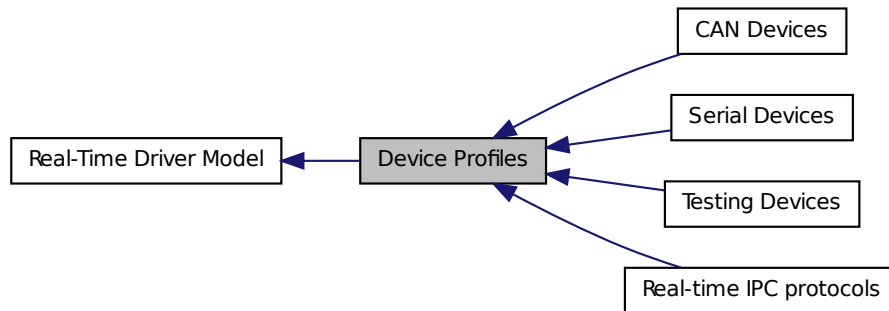
This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

## 5.33 Device Profiles

Collaboration diagram for Device Profiles:



### Data Structures

- struct `rtdm_device_info`  
*Device information.*

### Modules

- [CAN Devices](#)
  - [Real-time IPC protocols](#)
- Profile Revision: 1*

- [Serial Devices](#)
- [Testing Devices](#)

### Typedefs

- typedef struct `rtdm_device_info` `rtdm_device_info_t`  
*Device information.*

### RTDM\_CLASS\_XXX

Device classes

- `#define RTDM_CLASS_PARPORT 1`
- `#define RTDM_CLASS_SERIAL 2`
- `#define RTDM_CLASS_CAN 3`

- `#define RTDM_CLASS_NETWORK 4`
- `#define RTDM_CLASS_RTMAC 5`
- `#define RTDM_CLASS_TESTING 6`
- `#define RTDM_CLASS_RTIPC 7`
- `#define RTDM_CLASS_EXPERIMENTAL 224`
- `#define RTDM_CLASS_MAX 255`

## Device Naming

Maximum length of device names (excluding the final null character)

- `#define RTDM_MAX_DEVNAME_LEN 31`

## RTDM\_PURGE\_xxx\_BUFFER

Flags selecting buffers to be purged

- `#define RTDM_PURGE_RX_BUFFER 0x0001`
- `#define RTDM_PURGE_TX_BUFFER 0x0002`

## Common IOCTLs

The following IOCTLs are common to all device profiles.

- `#define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)`  
*Retrieve information about a device or socket.*
- `#define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)`  
*Purge internal device or socket buffers.*

### 5.33.1 Detailed Description

Device profiles define which operation handlers a driver of a certain class has to implement, which name or protocol it has to register, which IOCTLs it has to provide, and further details. Sub-classes can be defined in order to extend a device profile with more hardware-specific functions.

### 5.33.2 Define Documentation

#### 5.33.2.1 `#define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)`

Retrieve information about a device or socket.

#### Parameters

[out] *arg* Pointer to information buffer (struct `rtdm_device_info`)

### 5.33.2.2 #define RTIOC\_PURGE\_IOW(RTIOC\_TYPE\_COMMON, 0x10, int)

Purge internal device or socket buffers.

#### Parameters

[in] *arg* Purge mask, see [RTDM\\_PURGE\\_xxx\\_BUFFER](#)

## 5.34 Semaphores services.

Semaphores services.

### Functions

- static int [sem\\_destroy](#) (struct \_\_shadow\_sem \*sm)  
*Destroy an unnamed semaphore.*
- sem\_t \* [sem\\_open](#) (const char \*name, int oflags,...)  
*Open a named semaphore.*
- int [sem\\_close](#) (struct \_\_shadow\_sem \*sm)  
*Close a named semaphore.*
- int [sem\\_unlink](#) (const char \*name)  
*Unlink a named semaphore.*
- static int [sem\\_trywait](#) (cobalt\_sem\_t \*sem)  
*Attempt to decrement a semaphore.*
- static int [sem\\_wait](#) (cobalt\_sem\_t \*sem)  
*Decrement a semaphore.*
- static int [sem\\_timedwait](#) (cobalt\_sem\_t \*sem, const struct timespec \*abs\_timeout)  
*Attempt, during a bounded time, to decrement a semaphore.*
- static int [sem\\_post](#) (cobalt\_sem\_t \*sm)  
*Post a semaphore.*
- int [sem\\_getvalue](#) (cobalt\_sem\_t \*sem, int \*value)  
*Get the value of a semaphore.*

### 5.34.1 Detailed Description

Semaphores services. Semaphores are counters for resources shared between threads. The basic operations on semaphores are: increment the counter atomically, and wait until the counter is non-null and decrement it atomically.

Semaphores have a maximum value past which they cannot be incremented. The macro `SEM_VALUE_MAX` is defined to be this maximum value.

## 5.34.2 Function Documentation

### 5.34.2.1 `int sem_close ( struct __shadow_sem * sm )`

Close a named semaphore.

This service closes the semaphore *sm*. The semaphore is destroyed only when unlinked with a call to the [sem\\_unlink\(\)](#) service and when each call to [sem\\_open\(\)](#) matches a call to this service.

When a semaphore is destroyed, the memory it used is returned to the system heap, so that further references to this semaphore are not guaranteed to fail, as is the case for unnamed semaphores.

This service fails if *sm* is an unnamed semaphore.

#### Parameters

*sm* the semaphore to be closed.

#### Return values

0 on success;

-1 with *errno* set if:

- EINVAL, the semaphore *sm* is invalid or is an unnamed semaphore.

#### See also

[Specification.](#)

### 5.34.2.2 `static int sem_destroy ( struct __shadow_sem * sm ) [static]`

Destroy an unnamed semaphore.

This service destroys the semaphore *sm*. Threads currently blocked on *sm* are unblocked and the service they called return -1 with *errno* set to EINVAL. The semaphore is then considered invalid by all semaphore services (they all fail with *errno* set to EINVAL) except [sem\\_init\(\)](#).

This service fails if *sm* is a named semaphore.

#### Parameters

*sm* the semaphore to be destroyed.

#### Return values

*always* 0 on success if SEM\_WARNDEL was not mentioned via [sem\\_init\\_np\(\)](#). If SEM\_WARNDEL was mentioned, then a strictly positive value is returned to warn the caller if threads were pending on the semaphore, or zero otherwise.

-1 with *errno* set if:

- EINVAL, the semaphore *sm* is invalid or a named semaphore;
- EPERM, the semaphore *sm* is not process-shared and does not belong to the current process.

#### See also

[Specification.](#)



**5.34.2.3 int sem\_getvalue ( cobalt\_sem\_t \* sem, int \* value )**

Get the value of a semaphore.

This service stores at the address *value*, the current count of the semaphore *sm*. The state of the semaphore is unchanged.

If the semaphore is currently fully depleted, the value stored is zero, unless SEM\_REPORT was mentioned for a non-standard semaphore (see sem\_init\_np()), in which case the current number of waiters is returned as the semaphore's negative value (e.g. -2 would mean the semaphore is fully depleted AND two threads are currently pending on it).

**Parameters**

*sem* a semaphore;

*value* address where the semaphore count will be stored on success.

**Return values**

0 on success;

-1 with *errno* set if:

- EINVAL, the semaphore is invalid or uninitialized;
- EPERM, the semaphore *sm* is not process-shared and does not belong to the current process.

**See also**

[Specification.](#)

Referenced by timer\_create().

**5.34.2.4 sem\_t\* sem\_open ( const char \* name, int oflags, ... )**

Open a named semaphore.

This service establishes a connection between the semaphore named *name* and the calling context (kernel-space as a whole, or user-space process).

If no semaphore named *name* exists and *oflags* has the O\_CREAT bit set, the semaphore is created by this function, using two more arguments:

- a *mode* argument, of type **mode\_t**, currently ignored;
- a *value* argument, of type **unsigned**, specifying the initial value of the created semaphore.

If *oflags* has the two bits O\_CREAT and O\_EXCL set and the semaphore already exists, this service fails.

*name* may be any arbitrary string, in which slashes have no particular meaning. However, for portability, using a name which starts with a slash and contains no other slash is recommended.

If [sem\\_open\(\)](#) is called from the same context (kernel-space as a whole, or user-space process) several times with the same value of *name*, the same address is returned.

**Parameters**

*name* the name of the semaphore to be created;

*oflags* flags.

### Returns

the address of the named semaphore on success;

SEM\_FAILED with *errno* set if:

- ENAMETOOLONG, the length of the *name* argument exceeds 64 characters;
- EEXIST, the bits *O\_CREAT* and *O\_EXCL* were set in *oflags* and the named semaphore already exists;
- ENOENT, the bit *O\_CREAT* is not set in *oflags* and the named semaphore does not exist;
- ENOSPC, insufficient memory exists in the system heap to create the semaphore, increase CONFIG\_XENO\_OPT\_SYS\_HEAPSZ;
- EINVAL, the *value* argument exceeds SEM\_VALUE\_MAX.

### See also

[Specification.](#)

#### 5.34.2.5 static int sem\_post ( cobalt\_sem\_t \* *sm* ) [static]

Post a semaphore.

This service posts the semaphore *sm*.

If no thread is currently blocked on this semaphore, its count is incremented unless "pulse" mode is enabled for it (see sem\_init\_np(), SEM\_PULSE). If a thread is blocked on the semaphore, the thread heading the wait queue is unblocked.

### Parameters

*sm* the semaphore to be signaled.

### Return values

0 on success;

-1 with *errno* set if:

- EINVAL, the specified semaphore is invalid or uninitialized;
- EPERM, the semaphore *sm* is not process-shared and does not belong to the current process;
- EAGAIN, the semaphore count is SEM\_VALUE\_MAX.

### See also

[Specification.](#)

#### 5.34.2.6 static int sem\_timedwait ( cobalt\_sem\_t \* *sem*, const struct timespec \* *abs\_timeout* ) [static]

Attempt, during a bounded time, to decrement a semaphore.

This service is equivalent to [sem\\_wait\(\)](#), except that the caller is only blocked until the timeout *abs\_timeout* expires.

**Parameters**

*sem* the semaphore to be decremented;  
*abs\_timeout* the timeout, expressed as an absolute value of the relevant clock for the semaphore, either `CLOCK_MONOTONIC` if `SEM_RAWCLOCK` was mentioned via `sem_init_np()`, or `CLOCK_REALTIME` otherwise.

**Return values**

*0* on success;  
*-1* with *errno* set if:

- `EPERM`, the caller context is invalid;
- `EINVAL`, the semaphore is invalid or uninitialized;
- `EINVAL`, the specified timeout is invalid;
- `EPERM`, the semaphore *sm* is not process-shared and does not belong to the current process;
- `EINTR`, the caller was interrupted by a signal while blocked in this service;
- `ETIMEDOUT`, the semaphore could not be decremented and the specified timeout expired.

**Valid contexts:**

- Xenomai kernel-space thread,
- Xenomai user-space thread (switches to primary mode).

**See also**

[Specification.](#)

Referenced by `rtm_sem_timeddown()`.

**5.34.2.7 static int sem\_trywait ( cobalt\_sem\_t \* sem ) [static]**

Attempt to decrement a semaphore.

This service is equivalent to [sem\\_wait\(\)](#), except that it returns immediately if the semaphore *sm* is currently depleted, and that it is not a cancellation point.

**Parameters**

*sem* the semaphore to be decremented.

**Return values**

*0* on success;  
*-1* with *errno* set if:

- `EINVAL`, the specified semaphore is invalid or uninitialized;
- `EPERM`, the semaphore *sm* is not process-shared and does not belong to the current process;
- `EAGAIN`, the specified semaphore is currently fully depleted.

\*

**See also**

[Specification.](#)

### 5.34.2.8 `int sem_unlink ( const char * name )`

Unlink a named semaphore.

This service unlinks the semaphore named *name*. This semaphore is not destroyed until all references obtained with `sem_open()` are closed by calling `sem_close()`. However, the unlinked semaphore may no longer be reached with the `sem_open()` service.

When a semaphore is destroyed, the memory it used is returned to the system heap, so that further references to this semaphore are not guaranteed to fail, as is the case for unnamed semaphores.

#### Parameters

*name* the name of the semaphore to be unlinked.

#### Return values

0 on success;

-1 with *errno* set if:

- ENAMETOOLONG, the length of the *name* argument exceeds 64 characters;
- ENOENT, the named semaphore does not exist.

#### See also

[Specification.](#)

### 5.34.2.9 `static int sem_wait ( cobalt_sem_t * sem ) [static]`

Decrement a semaphore.

This service decrements the semaphore *sm* if it is currently if its value is greater than 0. If the semaphore's value is currently zero, the calling thread is suspended until the semaphore is posted, or a signal is delivered to the calling thread.

This service is a cancellation point for Xenomai POSIX skin threads (created with the `pthread_create()` service). When such a thread is cancelled while blocked in a call to this service, the semaphore state is left unchanged before the cancellation cleanup handlers are called.

#### Parameters

*sem* the semaphore to be decremented.

#### Return values

0 on success;

-1 with *errno* set if:

- EPERM, the caller context is invalid;
- EINVAL, the semaphore is invalid or uninitialized;
- EPERM, the semaphore *sm* is not process-shared and does not belong to the current process;
- EINTR, the caller was interrupted by a signal while blocked in this service.

#### Valid contexts:

- Xenomai kernel-space thread,

- Xenomai user-space thread (switches to primary mode).

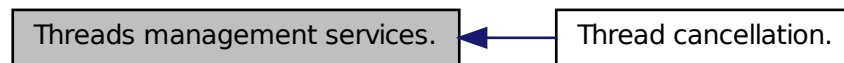
See also

[Specification.](#)

## 5.35 Threads management services.

Threads management services.

Collaboration diagram for Threads management services.:



### Modules

- [Thread cancellation.](#)

*Thread cancellation.*

### Functions

- static int [pthread\\_getschedparam](#) (pthread\_t tid, int \*pol, struct sched\_param \*par)  
*Get the scheduling policy and parameters of the specified thread.*
- static int [pthread\\_getschedparam\\_ex](#) (pthread\_t tid, int \*pol, struct sched\_param\_ex \*par)  
*Get the extended scheduling policy and parameters of the specified thread.*
- static int [pthread\\_create](#) (pthread\_t \*tid, const pthread\_attr\_t \*attr)  
*Create a thread.*
- static int [pthread\\_make\\_periodic\\_np](#) (pthread\_t thread, clockid\_t clock\_id, struct timespec \*starttp, struct timespec \*periodtp)  
*Make a thread periodic.*
- static int [pthread\\_set\\_mode\\_np](#) (int clrmask, int setmask, int \*mode\_r)  
*Set the mode of the current thread.*
- static int [pthread\\_set\\_name\\_np](#) (pthread\_t thread, const char \*name)  
*Set a thread name.*

- static int `pthread_setschedparam` (pthread\_t tid, int pol, const struct sched\_param \*par)  
*Set the scheduling policy and parameters of the specified thread.*
- static int `pthread_setschedparam_ex` (pthread\_t tid, int pol, const struct sched\_param\_ex \*par)  
*Set the extended scheduling policy and parameters of the specified thread.*

### 5.35.1 Detailed Description

Threads management services.

See also

[Specification.](#)

### 5.35.2 Function Documentation

#### 5.35.2.1 static int pthread\_create ( pthread\_t \* tid, const pthread\_attr\_t \* attr ) [inline, static]

Create a thread.

This service creates a Cobalt thread control block. The created thread may use Cobalt API services. The new thread control block can be mapped over a regular Linux thread, forming a Xenomai shadow.

The new thread signal mask is inherited from the current thread, if it was also created with `pthread_create()`, otherwise the new thread signal mask is empty.

Other attributes of the new thread depend on the *attr* argument. If *attr* is null, default values for these attributes are used.

Returning from the *start* routine has the same effect as calling `pthread_exit()` with the return value.

#### Parameters

*tid* address where the identifier of the new thread will be stored on success;

*attr* thread attributes;

#### Returns

0 on success;

an error number if:

- EINVAL, *attr* is invalid;
- EAGAIN, insufficient memory exists in the system heap to create a new thread, increase CONFIG\_XENO\_OPT\_SYS\_HEAPSZ;
- EINVAL, thread attribute *inheritsched* is set to PTHREAD\_INHERIT\_SCHED and the calling thread does not belong to the POSIX skin;

See also

[Specification.](#)

**Note**

When creating or shadowing a Xenomai thread for the first time in user-space, Xenomai installs a handler for the SIGWINCH signal. If you had installed a handler before that, it will be automatically called by Xenomai for SIGWINCH signals that it has not sent.

If, however, you install a signal handler for SIGWINCH after creating or shadowing the first Xenomai thread, you have to explicitly call the function `xeno_sigwinch_handler` at the beginning of your signal handler, using its return to know if the signal was in fact an internal signal of Xenomai (in which case it returns 1), or if you should handle the signal (in which case it returns 0). `xeno_sigwinch_handler` prototype is:

```
int xeno_sigwinch_handler(int sig, siginfo_t *si, void *ctxt);
```

Which means that you should register your handler with `sigaction`, using the `SA_SIGINFO` flag, and pass all the arguments you received to `xeno_sigwinch_handler`.

References `pthread_getschedparam_ex()`, `xnpod_init_thread()`, `xnpod_set_thread_tslice()`, and `xnsynch_init()`.

### 5.35.2.2 `static int pthread_getschedparam ( pthread_t tid, int * pol, struct sched_param * par ) [inline, static]`

Get the scheduling policy and parameters of the specified thread.

This service returns, at the addresses `pol` and `par`, the current scheduling policy and scheduling parameters (i.e. priority) of the Xenomai POSIX skin thread `tid`. If this service is called from user-space and `tid` is not the identifier of a Xenomai POSIX skin thread, this service fallback to Linux regular `pthread_getschedparam` service.

**Parameters**

*tid* target thread;

*pol* address where the scheduling policy of *tid* is stored on success;

*par* address where the scheduling parameters of *tid* is stored on success.

**Returns**

0 on success;

an error number if:

- `ESRCH`, *tid* is invalid.

**See also**

[Specification.](#)

### 5.35.2.3 `static int pthread_getschedparam_ex ( pthread_t tid, int * pol, struct sched_param_ex * par ) [inline, static]`

Get the extended scheduling policy and parameters of the specified thread.

This service is an extended version of [pthread\\_getschedparam\(\)](#), that also supports Xenomai-specific or additional POSIX scheduling policies, which are not available with the host Linux environment.

Typically, SCHED\_SPORADIC parameters can be retrieved from this call.

### Parameters

- tid* target thread;
- pol* address where the scheduling policy of *tid* is stored on success;
- par* address where the scheduling parameters of *tid* is stored on success.

### Returns

- 0 on success;
- an error number if:
  - ESRCH, *tid* is invalid.

### See also

[Specification.](#)

References XNRRB.

Referenced by pthread\_create().

**5.35.2.4 static int pthread\_make\_periodic\_np ( pthread\_t thread, clockid\_t clock\_id, struct timespec \* starttp, struct timespec \* periodtp ) [inline, static]**

Make a thread periodic.

This service make the POSIX skin thread *thread* periodic.

This service is a non-portable extension of the POSIX interface.

### Parameters

- thread* thread identifier. This thread is immediately delayed until the first periodic release point is reached.
- clock\_id* clock identifier, either CLOCK\_REALTIME, CLOCK\_MONOTONIC or CLOCK\_MONOTONIC\_RAW.
- starttp* start time, expressed as an absolute value of the clock *clock\_id*. The affected thread will be delayed until this point is reached.
- periodtp* period, expressed as a time interval.

### Returns

- 0 on success;
- an error number if:
  - ESRCH, *thread* is invalid;
  - ETIMEDOUT, the start time has already passed.
  - ENOTSUP, the specified clock is unsupported;

Rescheduling: always, until the *starttp* start time has been reached.

References xnpod\_set\_thread\_periodic().



### 5.35.2.5 `static int pthread_set_mode_np ( int clrmask, int setmask, int * mode_r ) [inline, static]`

Set the mode of the current thread.

This service sets the mode of the calling thread. *clrmask* and *setmask* are two bit masks which are respectively cleared and set in the calling thread status. They are a bitwise OR of the following values:

- `PTHREAD_LOCK_SCHED`, when set, locks the scheduler, which prevents the current thread from being switched out until the scheduler is unlocked;
- `PTHREAD_WARN_SW`, when set, causes the signal `SIGXCPU` to be sent to the current thread, whenever it involontary switches to secondary mode;
- `PTHREAD_CONFORMING` can be passed in *setmask* to switch the current user-space task to its preferred runtime mode. The only meaningful use of this switch is to force a real-time shadow back to primary mode. Any other use either cause to a nop, or an error.

`PTHREAD_LOCK_SCHED` is valid for any Xenomai thread, the other bits are only valid for Xenomai user-space threads.

This service is a non-portable extension of the POSIX interface.

#### Parameters

*clrmask* set of bits to be cleared;

*setmask* set of bits to be set.

*mode\_r* If non-NULL, *mode\_r* must be a pointer to a memory location which will be written upon success with the previous set of active mode bits. If NULL, the previous set of active mode bits will not be returned.

#### Returns

0 on success;

an error number if:

- `EINVAL`, some bit in *clrmask* or *setmask* is invalid.

References `XNLOCK`, `xnpod_schedule()`, `xnpod_set_thread_mode()`, and `XNSHADOW`.

### 5.35.2.6 `static int pthread_set_name_np ( pthread_t thread, const char * name ) [inline, static]`

Set a thread name.

This service set to *name*, the name of *thread*. This name is used for displaying information in `/proc/xenomai/sched`.

This service is a non-portable extension of the POSIX interface.

#### Parameters

*thread* target thread;

*name* name of the thread.

**Returns**

- 0 on success;
- an error number if:
  - ESRCH, *thread* is invalid.

**5.35.2.7** `static int pthread_setschedparam ( pthread_t tid, int pol, const struct sched_param * par ) [inline, static]`

Set the scheduling policy and parameters of the specified thread.

This service set the scheduling policy of the Xenomai POSIX skin thread *tid* to the value *pol*, and its scheduling parameters (i.e. its priority) to the value pointed to by *par*.

When used in user-space, passing the current thread ID as *tid* argument, this service turns the current thread into a Xenomai POSIX skin thread. If *tid* is neither the identifier of the current thread nor the identifier of a Xenomai POSIX skin thread this service falls back to the regular [pthread\\_setschedparam\(\)](#) service, hereby causing the current thread to switch to secondary mode if it is Xenomai thread.

**Parameters**

- tid* target thread;
- pol* scheduling policy, one of SCHED\_FIFO, SCHED\_COBALT, SCHED\_RR or SCHED\_OTHER;
- par* scheduling parameters address.

**Returns**

- 0 on success;
- an error number if:
  - ESRCH, *tid* is invalid;
  - EINVAL, *pol* or *par->sched\_priority* is invalid;
  - EAGAIN, in user-space, insufficient memory exists in the system heap, increase CONFIG\_XENO\_OPT\_SYS\_HEAPSZ;
  - EFAULT, in user-space, *par* is an invalid address;
  - EPERM, in user-space, the calling process does not have superuser permissions.

**See also**

[Specification.](#)

**Note**

When creating or shadowing a Xenomai thread for the first time in user-space, Xenomai installs a handler for the SIGWINCH signal. If you had installed a handler before that, it will be automatically called by Xenomai for SIGWINCH signals that it has not sent.

If, however, you install a signal handler for SIGWINCH after creating or shadowing the first Xenomai thread, you have to explicitly call the function `xeno_sigwinch_handler` at the beginning of your signal handler, using its return to know if the signal was in fact an internal signal of

Xenomai (in which case it returns 1), or if you should handle the signal (in which case it returns 0). `xeno_sigwinch_handler` prototype is:

```
int xeno_sigwinch_handler(int sig, siginfo_t *si, void *ctxt);
```

Which means that you should register your handler with `sigaction`, using the `SA_SIGINFO` flag, and pass all the arguments you received to `xeno_sigwinch_handler`.

References `xnpod_schedule()`, `xnpod_set_thread_schedparam()`, and `xnpod_set_thread_tslice()`.

Referenced by `pthread_setschedparam_ex()`.

**5.35.2.8** `static int pthread_setschedparam_ex ( pthread_t tid, int pol, const struct sched_param_ex * par ) [inline, static]`

Set the extended scheduling policy and parameters of the specified thread.

This service is an extended version of [pthread\\_setschedparam\(\)](#), that supports Xenomai-specific or additional scheduling policies, which are not available with the host Linux environment.

Typically, a Xenomai thread policy can be set to `SCHED_SPORADIC` using this call.

#### Parameters

*tid* target thread;

*pol* address where the scheduling policy of *tid* is stored on success;

*par* address where the scheduling parameters of *tid* is stored on success.

#### Returns

0 on success;  
an error number if:

- `ESRCH`, *tid* is invalid.
- `EINVAL`, *par* contains invalid parameters.
- `ENOMEM`, lack of memory to perform the operation.

#### See also

[Specification.](#)

References `pthread_setschedparam()`, `xnpod_schedule()`, `xnpod_set_thread_schedparam()`, and `xnpod_set_thread_tslice()`.

## 5.36 CAN Devices

Collaboration diagram for CAN Devices:



### Data Structures

- struct [can\\_bittime\\_std](#)  
*Standard bit-time parameters according to Bosch.*
- struct [can\\_bittime\\_btr](#)  
*Hardware-specific BTR bit-times.*
- struct [can\\_bittime](#)  
*Custom CAN bit-time definition.*
- struct [can\\_filter](#)  
*Filter for reception of CAN messages.*
- struct [sockaddr\\_can](#)  
*Socket address structure for the CAN address family.*
- struct [can\\_frame](#)  
*Raw CAN frame.*

### Files

- file [rtcan.h](#)  
*Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.*

### Defines

- #define [AF\\_CAN](#) 29  
*CAN address family.*
- #define [PF\\_CAN](#) AF\_CAN  
*CAN protocol family.*

- `#define SOL_CAN_RAW 103`  
*CAN socket levels.*

## Typedefs

- `typedef uint32_t can_id_t`  
*Type of CAN id (see `CAN_xxx_MASK` and `CAN_xxx_FLAG`).*
- `typedef can_id_t can_err_mask_t`  
*Type of CAN error mask.*
- `typedef uint32_t can_baudrate_t`  
*Baudrate definition in bits per second.*
- `typedef enum CAN_BITTIME_TYPE can_bittime_type_t`  
*See `CAN_BITTIME_TYPE`.*
- `typedef enum CAN_MODE can_mode_t`  
*See `CAN_MODE`.*
- `typedef int can_ctrlmode_t`  
*See `CAN_CTRLMODE`.*
- `typedef enum CAN_STATE can_state_t`  
*See `CAN_STATE`.*
- `typedef struct can_filter can_filter_t`  
*Filter for reception of CAN messages.*
- `typedef struct can_frame can_frame_t`  
*Raw CAN frame.*

## Enumerations

- `enum CAN_BITTIME_TYPE { CAN_BITTIME_STD, CAN_BITTIME_BTR }`  
*Supported CAN bit-time types.*

## CAN operation modes

Modes into which CAN controllers can be set

- `enum CAN_MODE { CAN_MODE_STOP = 0, CAN_MODE_START, CAN_MODE_SLEEP }`

## CAN controller states

States a CAN controller can be in.

- enum `CAN_STATE` {  
    `CAN_STATE_ACTIVE` = 0, `CAN_STATE_BUS_WARNING`, `CAN_STATE_BUS_PASSIVE`,  
    `CAN_STATE_BUS_OFF`,  
    `CAN_STATE_SCANNING_BAUDRATE`,     `CAN_STATE_STOPPED`,     `CAN_STATE_SLEEPING` }

## CAN ID masks

Bit masks for masking CAN IDs

- #define `CAN_EFF_MASK` 0x1FFFFFFF  
    *Bit mask for extended CAN IDs.*
- #define `CAN_SFF_MASK` 0x000007FF  
    *Bit mask for standard CAN IDs.*

## CAN ID flags

Flags within a CAN ID indicating special CAN frame attributes

- #define `CAN_EFF_FLAG` 0x80000000  
    *Extended frame.*
- #define `CAN_RTR_FLAG` 0x40000000  
    *Remote transmission frame.*
- #define `CAN_ERR_FLAG` 0x20000000  
    *Error frame (see [Errors](#)), not valid in struct `can_filter`.*
- #define `CAN_INV_FILTER` `CAN_ERR_FLAG`  
    *Invert CAN filter definition, only valid in struct `can_filter`.*

## Particular CAN protocols

Possible protocols for the PF\_CAN protocol family

Currently only the RAW protocol is supported.

- #define `CAN_RAW` 1  
    *Raw protocol of PF\_CAN, applicable to socket type `SOCK_RAW`.*

## CAN controller modes

Special CAN controllers modes, which can be or'ed together.

### Note

These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.

- #define `CAN_CTRLMODE_LISTENONLY` 0x1
- #define `CAN_CTRLMODE_LOOPBACK` 0x2

## Timestamp switches

Arguments to pass to `RTCAN_RTIOC_TAKE_TIMESTAMP`

- #define `RTCAN_TAKE_NO_TIMESTAMPS` 0  
*Switch off taking timestamps.*
- #define `RTCAN_TAKE_TIMESTAMPS` 1  
*Do take timestamps.*

## RAW socket options

Setting and getting CAN RAW socket options.

- #define `CAN_RAW_FILTER` 0x1  
*CAN filter definition.*
- #define `CAN_RAW_ERR_FILTER` 0x2  
*CAN error mask.*
- #define `CAN_RAW_LOOPBACK` 0x3  
*CAN TX loopback.*
- #define `CAN_RAW_RECV_OWN_MSGS` 0x4  
*CAN receive own messages.*

## IOCTLs

CAN device IOCTLs

- #define `SIOCGIFINDEX` `defined_by_kernel_header_file`  
*Get CAN interface index by name.*
- #define `SIOCSCANBAUDRATE` `_IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)`  
*Set baud rate.*

- #define [SIOCGCANBAUDRATE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x02, struct ifreq)  
*Get baud rate.*
- #define [SIOCSCANCUSTOMBITTIME](#) \_IOW(RTIOC\_TYPE\_CAN, 0x03, struct ifreq)  
*Set custom bit time parameter.*
- #define [SIOCGCANCUSTOMBITTIME](#) \_IOW(RTIOC\_TYPE\_CAN, 0x04, struct ifreq)  
*Get custom bit-time parameters.*
- #define [SIOCSCANMODE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x05, struct ifreq)  
*Set operation mode of CAN controller.*
- #define [SIOCGCANSTATE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x06, struct ifreq)  
*Get current state of CAN controller.*
- #define [SIOCSCANCTRLMODE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x07, struct ifreq)  
*Set special controller modes.*
- #define [SIOCGCANCTRLMODE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x08, struct ifreq)  
*Get special controller modes.*
- #define [RTCAN\\_RTIOC\\_TAKE\\_TIMESTAMP](#) \_IOW(RTIOC\_TYPE\_CAN, 0x09, int)  
*Enable or disable storing a high precision timestamp upon reception of a CAN frame.*
- #define [RTCAN\\_RTIOC\\_RCV\\_TIMEOUT](#) \_IOW(RTIOC\_TYPE\_CAN, 0x0A, nanosecs\_rel\_t)  
*Specify a reception timeout for a socket.*
- #define [RTCAN\\_RTIOC\\_SND\\_TIMEOUT](#) \_IOW(RTIOC\_TYPE\_CAN, 0x0B, nanosecs\_rel\_t)  
*Specify a transmission timeout for a socket.*

## Error mask

Error class (mask) in `can_id` field of struct [can\\_frame](#) to be used with [CAN\\_RAW\\_ERR\\_FILTER](#).

**Note:** Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.

**Note:** In case of a bus-off error condition ([CAN\\_ERR\\_BUSOFF](#)), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling [CAN\\_MODE\\_START](#).

**Note:** Bus error interrupts ([CAN\\_ERR\\_BUSERROR](#)) are enabled when an application is calling a [Recv](#) function on a socket listening on bus errors (using [CAN\\_RAW\\_ERR\\_FILTER](#)). After one bus error has occurred, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.

- #define [CAN\\_ERR\\_TX\\_TIMEOUT](#) 0x00000001U  
*TX timeout (netdevice driver).*



- #define [CAN\\_ERR\\_LOSTARB](#) 0x00000002U  
*Lost arbitration (see [data\[0\]](#)).*
- #define [CAN\\_ERR\\_CRTL](#) 0x00000004U  
*Controller problems (see [data\[1\]](#)).*
- #define [CAN\\_ERR\\_PROT](#) 0x00000008U  
*Protocol violations (see [data\[2\]](#), [data\[3\]](#)).*
- #define [CAN\\_ERR\\_TRX](#) 0x00000010U  
*Transceiver status (see [data\[4\]](#)).*
- #define [CAN\\_ERR\\_ACK](#) 0x00000020U  
*Received no ACK on transmission.*
- #define [CAN\\_ERR\\_BUSOFF](#) 0x00000040U  
*Bus off.*
- #define [CAN\\_ERR\\_BUSERROR](#) 0x00000080U  
*Bus error (may flood!).*
- #define [CAN\\_ERR\\_RESTARTED](#) 0x00000100U  
*Controller restarted.*
- #define [CAN\\_ERR\\_MASK](#) 0x1FFFFFFFU  
*Omit EFF, RTR, ERR flags.*

### Arbitration lost error

Error in the [data\[0\]](#) field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_LOSTARB\\_UNSPEC](#) 0x00  
*unspecified*

### Controller problems

Error in the [data\[1\]](#) field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_CRTL\\_UNSPEC](#) 0x00  
*unspecified*
- #define [CAN\\_ERR\\_CRTL\\_RX\\_OVERFLOW](#) 0x01  
*RX buffer overflow.*
- #define [CAN\\_ERR\\_CRTL\\_TX\\_OVERFLOW](#) 0x02  
*TX buffer overflow.*

- #define [CAN\\_ERR\\_CTRL\\_RX\\_WARNING](#) 0x04  
*reached warning level for RX errors*
- #define [CAN\\_ERR\\_CTRL\\_TX\\_WARNING](#) 0x08  
*reached warning level for TX errors*
- #define [CAN\\_ERR\\_CTRL\\_RX\\_PASSIVE](#) 0x10  
*reached passive level for RX errors*
- #define [CAN\\_ERR\\_CTRL\\_TX\\_PASSIVE](#) 0x20  
*reached passive level for TX errors*

## Protocol error type

Error in the data[2] field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_PROT\\_UNSPEC](#) 0x00  
*unspecified*
- #define [CAN\\_ERR\\_PROT\\_BIT](#) 0x01  
*single bit error*
- #define [CAN\\_ERR\\_PROT\\_FORM](#) 0x02  
*frame format error*
- #define [CAN\\_ERR\\_PROT\\_STUFF](#) 0x04  
*bit stuffing error*
- #define [CAN\\_ERR\\_PROT\\_BIT0](#) 0x08  
*unable to send dominant bit*
- #define [CAN\\_ERR\\_PROT\\_BIT1](#) 0x10  
*unable to send recessive bit*
- #define [CAN\\_ERR\\_PROT\\_OVERLOAD](#) 0x20  
*bus overload*
- #define [CAN\\_ERR\\_PROT\\_ACTIVE](#) 0x40  
*active error announcement*
- #define [CAN\\_ERR\\_PROT\\_TX](#) 0x80  
*error occurred on transmission*

## Protocol error location

Error in the data[4] field of struct `can_frame`.

- #define `CAN_ERR_PROT_LOC_UNSPEC` 0x00  
*unspecified*
- #define `CAN_ERR_PROT_LOC_SOF` 0x03  
*start of frame*
- #define `CAN_ERR_PROT_LOC_ID28_21` 0x02  
*ID bits 28 - 21 (SFF: 10 - 3).*
- #define `CAN_ERR_PROT_LOC_ID20_18` 0x06  
*ID bits 20 - 18 (SFF: 2 - 0 ).*
- #define `CAN_ERR_PROT_LOC_SRTR` 0x04  
*substitute RTR (SFF: RTR)*
- #define `CAN_ERR_PROT_LOC_IDE` 0x05  
*identifier extension*
- #define `CAN_ERR_PROT_LOC_ID17_13` 0x07  
*ID bits 17-13.*
- #define `CAN_ERR_PROT_LOC_ID12_05` 0x0F  
*ID bits 12-5.*
- #define `CAN_ERR_PROT_LOC_ID04_00` 0x0E  
*ID bits 4-0.*
- #define `CAN_ERR_PROT_LOC_RTR` 0x0C  
*RTR.*
- #define `CAN_ERR_PROT_LOC_RES1` 0x0D  
*reserved bit 1*
- #define `CAN_ERR_PROT_LOC_RES0` 0x09  
*reserved bit 0*
- #define `CAN_ERR_PROT_LOC_DLC` 0x0B  
*data length code*
- #define `CAN_ERR_PROT_LOC_DATA` 0x0A  
*data section*
- #define `CAN_ERR_PROT_LOC_CRC_SEQ` 0x08  
*CRC sequence.*
- #define `CAN_ERR_PROT_LOC_CRC_DEL` 0x18

*CRC delimiter.*

- #define `CAN_ERR_PROT_LOC_ACK` 0x19  
*ACK slot.*
- #define `CAN_ERR_PROT_LOC_ACK_DEL` 0x1B  
*ACK delimiter.*
- #define `CAN_ERR_PROT_LOC_EOF` 0x1A  
*end of frame*
- #define `CAN_ERR_PROT_LOC_INTERM` 0x12  
*intermission*
- #define `CAN_ERR_TRX_UNSPEC` 0x00  
*0000 0000*
- #define `CAN_ERR_TRX_CANH_NO_WIRE` 0x04  
*0000 0100*
- #define `CAN_ERR_TRX_CANH_SHORT_TO_BAT` 0x05  
*0000 0101*
- #define `CAN_ERR_TRX_CANH_SHORT_TO_VCC` 0x06  
*0000 0110*
- #define `CAN_ERR_TRX_CANH_SHORT_TO_GND` 0x07  
*0000 0111*
- #define `CAN_ERR_TRX_CANL_NO_WIRE` 0x40  
*0100 0000*
- #define `CAN_ERR_TRX_CANL_SHORT_TO_BAT` 0x50  
*0101 0000*
- #define `CAN_ERR_TRX_CANL_SHORT_TO_VCC` 0x60  
*0110 0000*
- #define `CAN_ERR_TRX_CANL_SHORT_TO_GND` 0x70  
*0111 0000*
- #define `CAN_ERR_TRX_CANL_SHORT_TO_CANH` 0x80  
*1000 0000*

### 5.36.1 Detailed Description

This is the common interface a RTDM-compliant CAN device has to provide. Feel free to report bugs and comments on this profile to the "Socketcan" mailing list ([Socketcan-core@lists.berlios.de](mailto:Socketcan-core@lists.berlios.de)) or directly to the authors ([wg@grandegger.com](mailto:wg@grandegger.com) or [Sebastian.Smolorz@stud.uni-hannover.de](mailto:Sebastian.Smolorz@stud.uni-hannover.de)).

**Profile Revision:** 2

#### Device Characteristics

**Device Flags:** RTDM\_PROTOCOL\_DEVICE

**Protocol Family:** PF\_CAN

**Socket Type:** SOCK\_RAW

**Device Class:** RTDM\_CLASS\_CAN

#### Supported Operations

##### Socket

Environments: non-RT (RT optional, deprecated)

Specific return values:

- -EPROTONOSUPPORT (Protocol is not supported by the driver. See [CAN protocols](#) for possible protocols.)

##### Close

Blocking calls to any of the [Send](#) or [Receive](#) functions will be unblocked when the socket is closed and return with an error.

Environments: non-RT (RT optional, deprecated)

Specific return values: none

##### IOCTL

Mandatory Environments: see [below](#)

Specific return values: see [below](#)

##### Bind

Binds a socket to one or all CAN devices (see struct [sockaddr\\_can](#)). If a filter list has been defined with setsockopt (see [Sockopts](#)), it will be used upon reception of CAN frames to decide whether the bound socket will receive a frame. If no filter has been defined, the socket will receive **all** CAN frames on the specified interface(s).

Binding to special interface index 0 will make the socket receive CAN frames from all CAN interfaces.

Binding to an interface index is also relevant for the [Send](#) functions because they will transmit a message over the interface the socket is bound to when no socket address is given to them.

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -ENOMEM (Not enough memory to fulfill the operation)
- -EINVAL (Invalid address family, or invalid length of address structure)
- -ENODEV (Invalid CAN interface index)
- -ENOSPC (No enough space for filter list)
- -EBADF (Socket is about to be closed)
- -EAGAIN (Too many receivers. Old binding (if any) is still active. Close some sockets and try again.)

##### Setsockopt, Getsockopt

These functions allow to set and get various socket options. Currently, only CAN raw sockets are supported.

Supported Levels and Options:

- Level **SOL\_CAN\_RAW** : CAN RAW protocol (see [CAN\\_RAW](#))
  - Option [CAN\\_RAW\\_FILTER](#) : CAN filter list
  - Option [CAN\\_RAW\\_ERR\\_FILTER](#) : CAN error mask
  - Option [CAN\\_RAW\\_LOOPBACK](#) : CAN TX loopback to local sockets

Environments: non-RT (RT optional)

Specific return values: see links to options above.

### **Recv, Recvfrom, Recvmsg**

These functions receive CAN messages from a socket. Only one message per call can be received, so only one buffer with the correct length must be passed. For **SOCK\_RAW**, this is the size of struct [can\\_frame](#).

Unlike a call to one of the [Send](#) functions, a Recv function will not return with an error if an interface is down (due to bus-off or setting of stop mode) or in sleep mode. Moreover, in such a case there may still be some CAN messages in the socket buffer which could be read out successfully.

It is possible to receive a high precision timestamp with every CAN message. The condition is a former instruction to the socket via [RTCAN\\_RTIOC\\_TAKE\\_TIMESTAMP](#). The timestamp will be copied to the `msg_control` buffer of struct `msg_hdr` if it points to a valid memory location with size of [nanosecs\\_abs\\_t](#). If this is a NULL pointer the timestamp will be discarded silently.

**Note:** A `msg_controllen` of 0 upon completion of the function call indicates that no timestamp is available for that message.

Supported Flags [in]:

- **MSG\_DONTWAIT** (By setting this flag the operation will only succeed if it would not block, i.e. if there is a message in the socket buffer. This flag takes precedence over a timeout specified by [RTCAN\\_RTIOC\\_RCV\\_TIMEOUT](#).)
- **MSG\_PEEK** (Receive a message but leave it in the socket buffer. The next receive operation will get that message again.)

Supported Flags [out]: none

Environments: RT (non-RT optional)

Specific return values:

- Non-negative value (Indicating the successful reception of a CAN message. For **SOCK\_RAW**, this is the size of struct [can\\_frame](#) regardless of the actual size of the payload.)
- **-EFAULT** (It was not possible to access user space memory area at one of the specified addresses.)
- **-EINVAL** (Unsupported flag detected, or invalid length of socket address buffer, or invalid length of message control buffer)
- **-EMSGSIZE** (Zero or more than one iovec buffer passed, or buffer too small)
- **-EAGAIN** (No data available in non-blocking mode)
- **-EBADF** (Socket was closed.)
- **-EINTR** (Operation was interrupted explicitly or by signal.)
- **-ETIMEDOUT** (Timeout)

### **Send, Sendto, Sendmsg**

These functions send out CAN messages. Only one message per call can be transmitted, so only one buffer with the correct length must be passed. For **SOCK\_RAW**, this is the size of struct [can\\_frame](#).

The following only applies to `SOCK_RAW`: If a socket address of struct `sockaddr_can` is given, only `can_ifindex` is used. It is also possible to omit the socket address. Then the interface the socket is bound to will be used for sending messages.

If an interface goes down (due to bus-off or setting of stop mode) all senders that were blocked on this interface will be woken up.

Supported Flags:

- `MSG_DONTWAIT` (By setting this flag the transmit operation will only succeed if it would not block. This flag takes precedence over a timeout specified by `RTCAN_RTIOTC_SND_TIMEOUT`.)

Environments: RT (non-RT optional)

Specific return values:

- Non-negative value equal to given buffer size (Indicating the successful completion of the function call. See also note.)
- `-EOPNOTSUPP` (`MSG_OOB` flag is not supported.)
- `-EINVAL` (Unsupported flag detected *or*: Invalid length of socket address *or*: Invalid address family *or*: Data length code of CAN frame not between 0 and 15 *or*: CAN standard frame has got an ID not between 0 and 2031)
- `-EMSGSIZE` (Zero or more than one buffer passed or invalid size of buffer)
- `-EFAULT` (It was not possible to access user space memory area at one of the specified addresses.)
- `-ENXIO` (Invalid CAN interface index - 0 is not allowed here - or socket not bound or rather bound to all interfaces.)
- `-ENETDOWN` (Controller is bus-off or in stopped state.)
- `-ECOMM` (Controller is sleeping)
- `-EAGAIN` (Cannot transmit without blocking but a non-blocking call was requested.)
- `-EINTR` (Operation was interrupted explicitly or by signal)
- `-EBADF` (Socket was closed.)
- `-ETIMEDOUT` (Timeout)

**Note:** A successful completion of the function call does not implicate a successful transmission of the message.

## 5.36.2 Define Documentation

### 5.36.2.1 `#define CAN_CTRLMODE_LISTENONLY 0x1`

Listen-Only mode

In this mode the CAN controller would give no acknowledge to the CAN-bus, even if a message is received successfully and messages would not be transmitted. This mode might be useful for bus-monitoring, hot-plugging or throughput analysis.

### 5.36.2.2 `#define CAN_CTRLMODE_LOOPBACK 0x2`

Loopback mode

In this mode the CAN controller does an internal loop-back, a message is transmitted and simultaneously received. That mode can be used for self test operation.

### 5.36.2.3 #define CAN\_ERR\_LOSTARB\_UNSPEC 0x00

unspecified

else bit number in bitstream

### 5.36.2.4 #define CAN\_RAW\_ERR\_FILTER 0x2

CAN error mask.

A CAN error mask (see [Errors](#)) can be set with `setsockopt`. This mask is then used to decide if error frames are delivered to this socket in case of error conditions. The error frames are marked with the `CAN_ERR_FLAG` or `CAN_XXX_FLAG` and must be handled by the application properly. A detailed description of the errors can be found in the `can_id` and the data fields of struct `can_frame` (see [Errors](#) for further details).

#### Parameters

- [in] *level* `SOL_CAN_RAW`
- [in] *optname* `CAN_RAW_ERR_FILTER`
- [in] *optval* Pointer to error mask of type `can_err_mask_t`.
- [in] *optlen* Size of error mask: `sizeof(can_err_mask_t)`.

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -EINVAL (Invalid length "optlen")

### 5.36.2.5 #define CAN\_RAW\_FILTER 0x1

CAN filter definition.

A CAN raw filter list with elements of struct `can_filter` can be installed with `setsockopt`. This list is used upon reception of CAN frames to decide whether the bound socket will receive a frame. An empty filter list can also be defined using `optlen = 0`, which is recommended for write-only sockets.

If the socket was already bound with [Bind](#), the old filter list gets replaced with the new one. Be aware that already received, but not read out CAN frames may stay in the socket buffer.

#### Parameters

- [in] *level* `SOL_CAN_RAW`
- [in] *optname* `CAN_RAW_FILTER`
- [in] *optval* Pointer to array of struct `can_filter`.
- [in] *optlen* Size of filter list: `count * sizeof( struct can_filter)`.

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)



- -ENOMEM (Not enough memory to fulfill the operation)
- -EINVAL (Invalid length "optlen")
- -ENOSPC (No space to store filter list, check RT-Socket-CAN kernel parameters)

#### 5.36.2.6 #define CAN\_RAW\_LOOPBACK 0x3

CAN TX loopback.

The TX loopback to other local sockets can be selected with this `setsockopt`.

##### Note

The TX loopback feature must be enabled in the kernel and then the loopback to other local TX sockets is enabled by default.

##### Parameters

- [in] *level* SOL\_CAN\_RAW
- [in] *optname* CAN\_RAW\_LOOPBACK
- [in] *optval* Pointer to integer value.
- [in] *optlen* Size of int: sizeof(int).

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -EINVAL (Invalid length "optlen")
- -EOPNOTSUPP (not supported, check RT-Socket-CAN kernel parameters).

#### 5.36.2.7 #define CAN\_RAW\_RECV\_OWN\_MSGS 0x4

CAN receive own messages.

Not supported by RT-Socket-CAN, but defined for compatibility with Socket-CAN.

#### 5.36.2.8 #define RTCAN\_RTIOC\_RCV\_TIMEOUT \_IOW(RTIOC\_TYPE\_CAN, 0x0A, nanosecs\_rel\_t)

Specify a reception timeout for a socket.

Defines a timeout for all receive operations via a socket which will take effect when one of the [receive functions](#) is called without the MSG\_DONTWAIT flag set.

The default value for a newly created socket is an infinite timeout.

##### Note

The setting of the timeout value is not done atomically to avoid locks. Please set the value before receiving messages from the socket.

### Parameters

[in] *arg* Pointer to [nanosecs\\_rel\\_t](#) variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See [Timeouts](#) for special timeouts.

### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.36.2.9 `#define RTCAN_RTIOC_SND_TIMEOUT_IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_t)`

Specify a transmission timeout for a socket.

Defines a timeout for all send operations via a socket which will take effect when one of the [send functions](#) is called without the MSG\_DONTWAIT flag set.

The default value for a newly created socket is an infinite timeout.

### Note

The setting of the timeout value is not done atomically to avoid locks. Please set the value before sending messages to the socket.

### Parameters

[in] *arg* Pointer to [nanosecs\\_rel\\_t](#) variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See [Timeouts](#) for special timeouts.

### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.36.2.10 `#define RTCAN_RTIOC_TAKE_TIMESTAMP _IOW(RTIOC_TYPE_CAN, 0x09, int)`

Enable or disable storing a high precision timestamp upon reception of a CAN frame.

A newly created socket takes no timestamps by default.

##### Parameters

[in] *arg* int variable, see [Timestamp switches](#)

##### Returns

0 on success.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

##### Note

Activating taking timestamps only has an effect on newly received CAN messages from the bus. Frames that already are in the socket buffer do not have timestamps if it was deactivated before. See [Receive](#) for more details.

Rescheduling: never.

#### 5.36.2.11 `#define SIOCGCANBAUDRATE _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)`

Get baud rate.

##### Parameters

[in,out] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). *ifr\_name* must hold a valid CAN interface name, *ifr\_ifru* will be filled with an instance of [can\\_baudrate\\_t](#).

##### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.36.2.12 `#define SIOCGCANCTRLMODE _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)`

Get special controller modes.

##### Parameters

[in] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru must be filled with an instance of [can\\_ctrlmode\\_t](#).

##### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.36.2.13 `#define SIOCGCANCUSTOMBITTIME _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)`

Get custom bit-time parameters.

##### Parameters

[in,out] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru will be filled with an instance of struct [can\\_bittime](#).

##### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.36.2.14 `#define SIOCGCANSTATE_IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)`

Get current state of CAN controller.

States are divided into main states and additional error indicators. A CAN controller is always in exactly one main state. CAN bus errors are registered by the CAN hardware and collected by the driver. There is one error indicator (bit) per error type. If this IOCTL is triggered the error types which occurred since the last call of this IOCTL are reported and thereafter the error indicators are cleared. See also [CAN controller states](#).

##### Parameters

[in,out] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h).  
ifr\_name must hold a valid CAN interface name, ifr\_ifru will be filled with an instance of [can\\_mode\\_t](#).

##### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

#### 5.36.2.15 `#define SIOCGIFINDEX defined_by_kernel_header_file`

Get CAN interface index by name.

##### Parameters

[in,out] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h).  
If ifr\_name holds a valid CAN interface name ifr\_ifindex will be filled with the corresponding interface index.

## Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

### 5.36.2.16 #define SIOCSCANBAUDRATE\_IOW(RTIOC\_TYPE\_CAN, 0x01, struct ifreq)

Set baud rate.

The baudrate must be specified in bits per second. The driver will try to calculate resonable CAN bit-timing parameters. You can use [SIOCSCANCUSTOMBITTIME](#) to set custom bit-timing.

## Parameters

[in] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru must be filled with an instance of [can\\_baudrate\\_t](#).

## Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can\\_baudrate\\_t](#).
- -EDOM : Baud rate not possible.
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

## Note

Setting the baud rate is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

### 5.36.2.17 #define SIOCSCANCTRLMODE \_IOW(RTIOC\_TYPE\_CAN, 0x07, struct ifreq)

Set special controller modes.

Various special controller modes could be or'ed together (see [CAN\\_CTRLMODE](#) for further information).

#### Parameters

[in] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru must be filled with an instance of [can\\_ctrlmode\\_t](#).

#### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can\\_baudrate\\_t](#).
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

#### Note

Setting special controller modes is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

### 5.36.2.18 #define SIOCSCANCUSTOMBITTIME \_IOW(RTIOC\_TYPE\_CAN, 0x03, struct ifreq)

Set custom bit time parameter.

Custem-bit time could be defined in various formats (see struct [can\\_bittime](#)).

#### Parameters

[in] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru must be filled with an instance of struct [can\\_bittime](#).

#### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can\\_baudrate\\_t](#).
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

#### Note

Setting the bit-time is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

#### 5.36.2.19 #define SIOCSCANMODE\_IOW(RTIOC\_TYPE\_CAN, 0x05, struct ifreq)

Set operation mode of CAN controller.

See [CAN controller modes](#) for available modes.

#### Parameters

[in] *arg* Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr\_name must hold a valid CAN interface name, ifr\_ifru must be filled with an instance of [can\\_mode\\_t](#).

#### Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EAGAIN: ([CAN\\_MODE\\_START](#), [CAN\\_MODE\\_STOP](#)) Could not successfully set mode, hardware is busy. Try again.
- -EINVAL: ([CAN\\_MODE\\_START](#)) Cannot start controller, set baud rate first.
- -ENETDOWN: ([CAN\\_MODE\\_SLEEP](#)) Cannot go into sleep mode because controller is stopped or bus off.
- -EOPNOTSUPP: unknown mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)



**Note**

Setting a CAN controller into normal operation after a bus-off can take some time (128 occurrences of 11 consecutive recessive bits). In such a case, although this IOCTL will return immediately with success and [SIOCGCANSTATE](#) will report [CAN\\_STATE\\_ACTIVE](#), bus-off recovery may still be in progress.

If a controller is bus-off, setting it into stop mode will return no error but the controller remains bus-off.

Rescheduling: possible.

**5.36.2.20 #define SOL\_CAN\_RAW 103**

CAN socket levels.

Used for [Sockopts](#) for the particular protocols.

**5.36.3 Typedef Documentation****5.36.3.1 typedef struct can\_filter can\_filter\_t**

Filter for reception of CAN messages.

This filter works as follows: A received CAN ID is AND'ed bitwise with `can_mask` and then compared to `can_id`. This also includes the [CAN\\_EFF\\_FLAG](#) and [CAN\\_RTR\\_FLAG](#) of [CAN\\_XXX\\_FLAG](#). If this comparison is true, the message will be received by the socket. The logic can be inverted with the `can_id` flag [CAN\\_INV\\_FILTER](#):

```
if (can_id & CAN_INV_FILTER) {
    if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))
        accept-message;
} else {
    if ((received_can_id & can_mask) == can_id)
        accept-message;
}
```

Multiple filters can be arranged in a filter list and set with [Sockopts](#). If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

**5.36.3.2 typedef struct can\_frame can\_frame\_t**

Raw CAN frame.

Central structure for receiving and sending CAN frames.

**5.36.4 Enumeration Type Documentation****5.36.4.1 enum CAN\_BITTIME\_TYPE**

Supported CAN bit-time types.

**Enumerator:**

[CAN\\_BITTIME\\_STD](#) Standard bit-time definition according to Bosch.

*CAN\_BITTIME\_BTR* Hardware-specific BTR bit-time definition.

#### 5.36.4.2 enum CAN\_MODE

Enumerator:

*CAN\_MODE\_STOP* Set controller in Stop mode (no reception / transmission possible)

*CAN\_MODE\_START* Set controller into normal operation.

Coming from stopped mode or bus off, the controller begins with no errors in [CAN\\_STATE\\_ACTIVE](#).

*CAN\_MODE\_SLEEP* Set controller into Sleep mode.

This is only possible if the controller is not stopped or bus-off.

Notice that sleep mode will only be entered when there is no bus activity. If the controller detects bus activity while "sleeping" it will go into operating mode again.

To actively leave sleep mode again trigger *CAN\_MODE\_START*.

#### 5.36.4.3 enum CAN\_STATE

Enumerator:

*CAN\_STATE\_ACTIVE* CAN controller is error active.

*CAN\_STATE\_BUS\_WARNING* CAN controller is error active, warning level is reached.

*CAN\_STATE\_BUS\_PASSIVE* CAN controller is error passive.

*CAN\_STATE\_BUS\_OFF* CAN controller went into Bus Off.

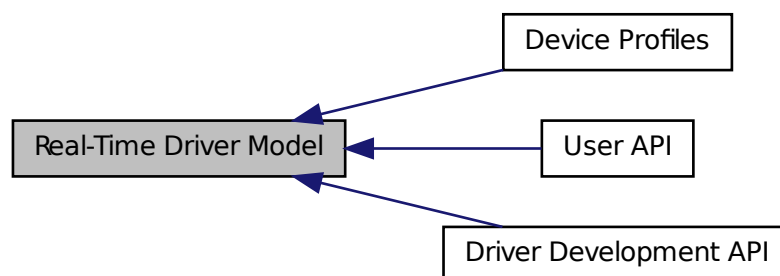
*CAN\_STATE\_SCANNING\_BAUDRATE* CAN controller is scanning to get the baudrate.

*CAN\_STATE\_STOPPED* CAN controller is in stopped mode.

*CAN\_STATE\_SLEEPING* CAN controller is in Sleep mode.

## 5.37 Real-Time Driver Model

Collaboration diagram for Real-Time Driver Model:



## Modules

- [Driver Development API](#)
- [Device Profiles](#)
- [User API](#)

## Typedefs

- typedef uint64\_t [nanosecs\\_abs\\_t](#)  
*RTDM type for representing absolute dates.*
- typedef int64\_t [nanosecs\\_rel\\_t](#)  
*RTDM type for representing relative intervals.*

## API Versioning

- #define [RTDM\\_API\\_VER](#) 8  
*Common user and driver API version.*
- #define [RTDM\\_API\\_MIN\\_COMPAT\\_VER](#) 6  
*Minimum API revision compatible with the current release.*

## RTDM\_TIMEOUT\_XXX

Special timeout values

- #define [RTDM\\_TIMEOUT\\_INFINITE](#) 0  
*Block forever.*
- #define [RTDM\\_TIMEOUT\\_NONE](#) (-1)  
*Any negative timeout means non-blocking.*

### 5.37.1 Detailed Description

The Real-Time Driver Model (RTDM) provides a unified interface to both users and developers of real-time device drivers. Specifically, it addresses the constraints of mixed RT/non-RT systems like Xenomai. RTDM conforms to POSIX semantics (IEEE Std 1003.1) where available and applicable.

**API Revision:** 8

### 5.37.2 Define Documentation

#### 5.37.2.1 #define RTDM\_TIMEOUT\_INFINITE 0

Block forever.

### 5.37.2.2 `#define RTDM_TIMEOUT_NONE (-1)`

Any negative timeout means non-blocking.

## 5.37.3 Typedef Documentation

### 5.37.3.1 `typedef uint64_t nanosecs_abs_t`

RTDM type for representing absolute dates.

Its base type is a 64 bit unsigned integer. The unit is 1 nanosecond.

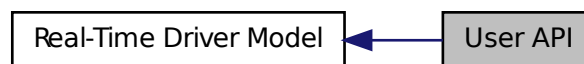
### 5.37.3.2 `typedef int64_t nanosecs_rel_t`

RTDM type for representing relative intervals.

Its base type is a 64 bit signed integer. The unit is 1 nanosecond. Relative intervals can also encode the special timeouts "infinite" and "non-blocking", see [RTDM\\_TIMEOUT\\_xxx](#).

## 5.38 User API

Collaboration diagram for User API:



## Files

- file [rtdm.h](#)  
*Real-Time Driver Model for Xenomai, user API header.*

## Functions

- int [rt\\_dev\\_open](#) (const char \*path, int oflag,...)  
*Open a device.*
- int [rt\\_dev\\_socket](#) (int protocol\_family, int socket\_type, int protocol)  
*Create a socket.*
- int [rt\\_dev\\_close](#) (int fd)

*Close a device or socket.*

- `int rt_dev_ioctl (int fd, int request,...)`  
*Issue an IOCTL.*
- `ssize_t rt_dev_read (int fd, void *buf, size_t nbyte)`  
*Read from device.*
- `ssize_t rt_dev_write (int fd, const void *buf, size_t nbyte)`  
*Write to device.*
- `ssize_t rt_dev_recvmsg (int fd, struct msghdr *msg, int flags)`  
*Receive message from socket.*
- `ssize_t rt_dev_recvfrom (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)`  
*Receive message from socket.*
- `ssize_t rt_dev_recv (int fd, void *buf, size_t len, int flags)`  
*Receive message from socket.*
- `ssize_t rt_dev_sendmsg (int fd, const struct msghdr *msg, int flags)`  
*Transmit message to socket.*
- `ssize_t rt_dev_sendto (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)`  
*Transmit message to socket.*
- `ssize_t rt_dev_send (int fd, const void *buf, size_t len, int flags)`  
*Transmit message to socket.*
- `int rt_dev_bind (int fd, const struct sockaddr *my_addr, socklen_t addrlen)`  
*Bind to local address.*
- `int rt_dev_connect (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)`  
*Connect to remote address.*
- `int rt_dev_listen (int fd, int backlog)`  
*Listen for incoming connection requests.*
- `int rt_dev_accept (int fd, struct sockaddr *addr, socklen_t *addrlen)`  
*Accept a connection requests.*
- `int rt_dev_shutdown (int fd, int how)`  
*Shut down parts of a connection.*
- `int rt_dev_getsockopt (int fd, int level, int optname, void *optval, socklen_t *optlen)`  
*Get socket option.*
- `int rt_dev_setsockopt (int fd, int level, int optname, const void *optval, socklen_t optlen)`

*Set socket option.*

- int `rt_dev_getsockname` (int *fd*, struct sockaddr \**name*, socklen\_t \**namelen*)  
*Get local socket address.*
- int `rt_dev_getpeername` (int *fd*, struct sockaddr \**name*, socklen\_t \**namelen*)  
*Get socket destination address.*

### 5.38.1 Detailed Description

This is the upper interface of RTDM provided to application programs both in kernel and user space. Note that certain functions may not be implemented by every device. Refer to the [Device Profiles](#) for precise information.

### 5.38.2 Function Documentation

#### 5.38.2.1 int `rt_dev_accept` ( int *fd*, struct sockaddr \* *addr*, socklen\_t \* *addrlen* )

Accept a connection requests.

##### Parameters

- [in] *fd* File descriptor as returned by `rt_dev_socket()`
- [out] *addr* Buffer for remote address
- [in,out] *addrlen* Address buffer size

##### Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

##### See also

`accept()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.38.2.2 int `rt_dev_bind` ( int *fd*, const struct sockaddr \* *my\_addr*, socklen\_t *addrlen* )

Bind to local address.

##### Parameters

- [in] *fd* File descriptor as returned by `rt_dev_socket()`
- [in] *my\_addr* Address buffer
- [in] *addrlen* Address buffer size

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

bind() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.3 int rt\_dev\_close ( int *fd* )**

Close a device or socket.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_open\(\)](#) or [rt\\_dev\\_socket\(\)](#)

**Returns**

0 on success, otherwise a negative error code.

**Note**

If the matching [rt\\_dev\\_open\(\)](#) or [rt\\_dev\\_socket\(\)](#) call took place in non-real-time context, [rt\\_dev\\_close\(\)](#) must be issued within non-real-time as well. Otherwise, the call will fail.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

close() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.4 int rt\_dev\_connect ( int *fd*, const struct sockaddr \* *serv\_addr*, socklen\_t *addrlen* )**

Connect to remote address.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[in] *serv\_addr* Address buffer

[in] *addrlen* Address buffer size

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

connect() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.5 int rt\_dev\_getpeername ( int *fd*, struct sockaddr \* *name*, socklen\_t \* *namelen* )

Get socket destination address.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[out] *name* Address buffer

[in,out] *namelen* Address buffer size

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

getpeername() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.6 int rt\_dev\_getsockname ( int *fd*, struct sockaddr \* *name*, socklen\_t \* *namelen* )

Get local socket address.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[out] *name* Address buffer

[in,out] *namelen* Address buffer size

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

getsockname() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>



### 5.38.2.7 `int rt_dev_getsockopt ( int fd, int level, int optname, void * optval, socklen_t * optlen )`

Get socket option.

#### Parameters

- [in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)
- [in] *level* Addressed stack level
- [in] *optname* Option name ID
- [out] *optval* Value buffer
- [in,out] *optlen* Value buffer size

#### Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### See also

`getsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.8 `int rt_dev_ioctl ( int fd, int request, ... )`

Issue an IOCTL.

#### Parameters

- [in] *fd* File descriptor as returned by [rt\\_dev\\_open\(\)](#) or [rt\\_dev\\_socket\(\)](#)
- [in] *request* IOCTL code
- ... Optional third argument, depending on IOCTL function (void \* or unsigned long)

#### Returns

Positiv value on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### See also

`ioctl()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.9 `int rt_dev_listen ( int fd, int backlog )`

Listen for incoming connection requests.

#### Parameters

- [in] *fd* File descriptor as returned by `rt_dev_socket()`
- [in] *backlog* Maximum queue length

#### Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### See also

`lsiten()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.10 `int rt_dev_open ( const char * path, int oflag, ... )`

Open a device.

#### Parameters

- [in] *path* Device name
- [in] *oflag* Open flags
- ... Further parameters will be ignored.

#### Returns

Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

#### See also

`open()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.11 `ssize_t rt_dev_read ( int fd, void * buf, size_t nbyte )`

Read from device.

#### Parameters

- [in] *fd* File descriptor as returned by `rt_dev_open()`

[out] *buf* Input buffer  
[in] *nbyte* Number of bytes to read

### Returns

Number of bytes read, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### See also

`read()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.38.2.12 `ssize_t rt_dev_recv ( int fd, void * buf, size_t len, int flags )`

Receive message from socket.

### Parameters

[in] *fd* File descriptor as returned by `rt_dev_socket()`  
[out] *buf* Message buffer  
[in] *len* Message buffer size  
[in] *flags* Message flags

### Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### See also

`recv()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.38.2.13 `ssize_t rt_dev_recvfrom ( int fd, void * buf, size_t len, int flags, struct sockaddr * from, socklen_t * fromlen )`

Receive message from socket.

### Parameters

[in] *fd* File descriptor as returned by `rt_dev_socket()`  
[out] *buf* Message buffer  
[in] *len* Message buffer size

[in] *flags* Message flags  
[out] *from* Buffer for message sender address  
[in,out] *fromlen* Address buffer size

### Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### See also

recvfrom() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.38.2.14 ssize\_t rt\_dev\_recvmsg ( int *fd*, struct msghdr \* *msg*, int *flags* )

Receive message from socket.

### Parameters

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)  
[in,out] *msg* Message descriptor  
[in] *flags* Message flags

### Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

### See also

recvmsg() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

#### 5.38.2.15 ssize\_t rt\_dev\_send ( int *fd*, const void \* *buf*, size\_t *len*, int *flags* )

Transmit message to socket.

### Parameters

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)  
[in] *buf* Message buffer  
[in] *len* Message buffer size  
[in] *flags* Message flags

**Returns**

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

send() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.16 ssize\_t rt\_dev\_sendmsg ( int *fd*, const struct msghdr \* *msg*, int *flags* )**

Transmit message to socket.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[in] *msg* Message descriptor

[in] *flags* Message flags

**Returns**

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

sendmsg() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.17 ssize\_t rt\_dev\_sendto ( int *fd*, const void \* *buf*, size\_t *len*, int *flags*, const struct sockaddr \* *to*, socklen\_t *tolen* )**

Transmit message to socket.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[in] *buf* Message buffer

[in] *len* Message buffer size

[in] *flags* Message flags

[in] *to* Buffer for message destination address

[in] *tolen* Address buffer size

**Returns**

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

`sendto()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.18** `int rt_dev_setsockopt ( int fd, int level, int optname, const void * optval,  
socklen_t optlen )`

Set socket option.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[in] *level* Addressed stack level

[in] *optname* Option name ID

[in] *optval* Value buffer

[in] *optlen* Value buffer size

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

`setsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

**5.38.2.19** `int rt_dev_shutdown ( int fd, int how )`

Shut down parts of a connection.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_socket\(\)](#)

[in] *how* Specifies the part to be shut down (SHUT\_XXX)

**Returns**

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

shutdown() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.20 int rt\_dev\_socket ( int *protocol\_family*, int *socket\_type*, int *protocol* )

Create a socket.

**Parameters**

[in] *protocol\_family* Protocol family (PF\_XXX)

[in] *socket\_type* Socket type (SOCK\_XXX)

[in] *protocol* Protocol ID, 0 for default

**Returns**

Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

socket() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

### 5.38.2.21 ssize\_t rt\_dev\_write ( int *fd*, const void \* *buf*, size\_t *nbyte* )

Write to device.

**Parameters**

[in] *fd* File descriptor as returned by [rt\\_dev\\_open\(\)](#)

[in] *buf* Output buffer

[in] *nbyte* Number of bytes to write

**Returns**

Number of bytes written, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

**See also**

write() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

## 5.39 Real-time IPC protocols

**Profile Revision:** 1

Collaboration diagram for Real-time IPC protocols:



### Data Structures

- struct [rtipc\\_port\\_label](#)  
*Port label information structure.*
- struct [sockaddr\\_ipc](#)  
*Socket address structure for the RTIPC address family.*

### Files

- file [rtipc.h](#)  
*This file is part of the Xenomai project.*

### Typedefs

- typedef int16\_t [rtipc\\_port\\_t](#)  
*Port number type for the RTIPC address family.*

### RTIPC protocol list

protocols for the PF\_RTIPC protocol family

- enum { [IPCPROTO\\_IPC](#) = 0, [IPCPROTO\\_XDDP](#) = 1, [IPCPROTO\\_IDDP](#) = 2, [IPCPROTO\\_-BUFP](#) = 3 }

### Supported operations

Standard socket operations supported by the RTIPC protocols.



- int `socket__AF_RTIPC` (int domain=AF\_RTIPC, int type=SOCK\_DGRAM, int protocol)  
*Create an endpoint for communication in the AF\_RTIPC domain.*
- int `close__AF_RTIPC` (int sockfd)  
*Close a RTIPC socket descriptor.*
- int `bind__AF_RTIPC` (int sockfd, const struct `sockaddr_ipc` \*addr, socklen\_t addrlen)  
*Bind a RTIPC socket to a port.*
- int `connect__AF_RTIPC` (int sockfd, const struct `sockaddr_ipc` \*addr, socklen\_t addrlen)  
*Initiate a connection on a RTIPC socket.*
- int `setsockopt__AF_RTIPC` (int sockfd, int level, int optname, const void \*optval, socklen\_t optlen)  
*Set options on RTIPC sockets.*
- int `getsockopt__AF_RTIPC` (int sockfd, int level, int optname, void \*optval, socklen\_t \*optlen)  
*Get options on RTIPC sockets.*
- ssize\_t `sendmsg__AF_RTIPC` (int sockfd, const struct msghdr \*msg, int flags)  
*Send a message on a RTIPC socket.*
- ssize\_t `recvmsg__AF_RTIPC` (int sockfd, struct msghdr \*msg, int flags)  
*Receive a message from a RTIPC socket.*
- int `getsockname__AF_RTIPC` (int sockfd, struct `sockaddr_ipc` \*addr, socklen\_t \*addrlen)  
*Get socket name.*
- int `getpeername__AF_RTIPC` (int sockfd, struct `sockaddr_ipc` \*addr, socklen\_t \*addrlen)  
*Get socket peer.*

## XDDP socket options

Setting and getting XDDP socket options.

- #define `XDDP_LABEL` 1  
*XDDP label assignment.*
- #define `XDDP_POOLSZ` 2  
*XDDP local pool size configuration.*
- #define `XDDP_BUFSZ` 3  
*XDDP streaming buffer size configuration.*
- #define `XDDP_MONITOR` 4  
*XDDP monitoring callback.*

## XDDP events

Specific events occurring on XDDP channels, which can be monitored via the [XDDP\\_MONITOR](#) socket option.

- #define [XDDP\\_EVTIN](#) 1  
*Monitor writes to the non real-time endpoint.*
- #define [XDDP\\_EVTOUT](#) 2  
*Monitor reads from the non real-time endpoint.*
- #define [XDDP\\_EVTDOWN](#) 3  
*Monitor close from the non real-time endpoint.*
- #define [XDDP\\_EVTNOBUF](#) 4  
*Monitor memory shortage for non real-time datagrams.*

## IDDP socket options

Setting and getting IDDP socket options.

- #define [IDDP\\_LABEL](#) 1  
*IDDP label assignment.*
- #define [IDDP\\_POOLSZ](#) 2  
*IDDP local pool size configuration.*

## BUFP socket options

Setting and getting BUFP socket options.

- #define [BUFP\\_LABEL](#) 1  
*BUFP label assignment.*
- #define [BUFP\\_BUFSZ](#) 2  
*BUFP buffer size configuration.*

## Socket level options

Setting and getting supported standard socket level options.

- #define [SO\\_SNDTIMEO](#) defined\_by\_kernel\_header\_file  
*IPCPROTO\_IDDP and IPCPROTO\_BUFP protocols support the standard SO\_SNDTIMEO socket option, from the SOL\_SOCKET level.*
- #define [SO\\_RCVTIMEO](#) defined\_by\_kernel\_header\_file  
*All RTIPC protocols support the standard SO\_RCVTIMEO socket option, from the SOL\_SOCKET level.*

### 5.39.1 Detailed Description

Profile Revision: 1

#### Device Characteristics

Device Flags: RTDM\_PROTOCOL\_DEVICE

Protocol Family: PF\_RTIPC

Socket Type: SOCK\_DGRAM

Device Class: RTDM\_CLASS\_RTIPC

### 5.39.2 Define Documentation

#### 5.39.2.1 #define BUFP\_BUFSZ 2

BUFP buffer size configuration.

All messages written to a BUFP socket are buffered in a single per-socket memory area. Configuring the size of such buffer prior to binding the socket to a destination port is mandatory.

It is not allowed to configure a buffer size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

#### Note

: the buffer memory is obtained from the host allocator by the [bind call](#).

#### Parameters

[in] *level* [SOL\\_BUFP](#)

[in] *optname* BUFP\_BUFSZ

[in] *optval* Pointer to a variable of type `size_t`, containing the required size of the buffer to reserve at binding time

[in] *optlen* `sizeof(size_t)`

#### Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid or *\*optval* is zero)

#### Calling context:

RT/non-RT

#### 5.39.2.2 #define BUFP\_LABEL 1

BUFP label assignment.

ASCII label strings can be attached to BUFP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see [BUFP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

#### Parameters

- [in] *level* [SOL\\_BUFP](#)
- [in] *optname* [BUFP\\_LABEL](#)
- [in] *optval* Pointer to struct [rtipc\\_port\\_label](#)
- [in] *optlen* sizeof(struct rtipc\_port\_label)

#### Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid)

#### Calling context:

RT/non-RT

#### 5.39.2.3 #define IDDP\_LABEL 1

IDDP label assignment.

ASCII label strings can be attached to IDDP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see [IDDP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

#### Parameters

- [in] *level* [SOL\\_IDDP](#)
- [in] *optname* [IDDP\\_LABEL](#)
- [in] *optval* Pointer to struct [rtipc\\_port\\_label](#)
- [in] *optlen* sizeof(struct rtipc\_port\_label)

#### Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid)

**Calling context:**

RT/non-RT

**5.39.2.4 #define IDDP\_POOLSZ 2**

IDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

**Note**

: the pool memory is obtained from the host allocator by the [bind call](#).

**Parameters**

[in] *level* [SOL\\_IDDP](#)

[in] *optname* IDDP\_POOLSZ

[in] *optval* Pointer to a variable of type `size_t`, containing the required size of the local pool to reserve at binding time

[in] *optlen* `sizeof(size_t)`

**Returns**

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid or *\*optval* is zero)

**Calling context:**

RT/non-RT

**5.39.2.5 #define SO\_RCVTIMEO defined\_by\_kernel\_header\_file**

All RTIPC protocols support the standard SO\_RCVTIMEO socket option, from the SOL\_SOCKET level.

**See also**

`setsockopt()`, `getsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399/>

### 5.39.2.6 #define SO\_SNDTIMEO defined\_by\_kernel\_header\_file

[IPPROTO\\_IPDDP](#) and [IPPROTO\\_BUFDP](#) protocols support the standard SO\_SNDTIMEO socket option, from the SOL\_SOCKET level.

#### See also

setsockopt(), getsockopt() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399/>

### 5.39.2.7 #define XDDP\_BUFSZ 3

XDDP streaming buffer size configuration.

In addition to sending datagrams, real-time threads may stream data in a byte-oriented mode through the port as well. This increases the bandwidth and reduces the overhead, when the overall data to send to the Linux domain is collected by bits, and keeping the message boundaries is not required.

This feature is enabled when a non-zero buffer size is set for the socket. In that case, the real-time data accumulates into the streaming buffer when MSG\_MORE is passed to any of the [send functions](#), until:

- the receiver from the Linux domain wakes up and consumes it,
- a different source port attempts to send data to the same destination port,
- MSG\_MORE is absent from the send flags,
- the buffer is full,

whichever comes first.

Setting *\*optval* to zero disables the streaming buffer, in which case all sendings are conveyed in separate datagrams, regardless of MSG\_MORE.

#### Note

only a single streaming buffer exists per socket. When this buffer is full, the real-time data stops accumulating and sending operations resume in mere datagram mode. Accumulation may happen again after some or all data in the streaming buffer is consumed from the Linux domain endpoint.

The streaming buffer size may be adjusted multiple times during the socket lifetime; the latest configuration change will take effect when the accumulation resumes after the previous buffer was flushed.

#### Parameters

[in] *level* [SOL\\_XDDP](#)

[in] *optname* XDDP\_BUFSZ

[in] *optval* Pointer to a variable of type `size_t`, containing the required size of the streaming buffer

[in] *optlen* `sizeof(size_t)`

**Returns**

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -ENOMEM (Not enough memory)
- -EINVAL (*optlen* is invalid)

**Calling context:**

RT/non-RT

**5.39.2.8 #define XDDP\_EVTDOWN 3**

**Monitor** close from the non real-time endpoint.

XDDP\_EVTDOWN is sent when the non real-time endpoint is closed. The argument is always 0.

**5.39.2.9 #define XDDP\_EVTIN 1**

**Monitor** writes to the non real-time endpoint.

XDDP\_EVTIN is sent when data is written to the non real-time endpoint the socket is bound to (i.e. via /dev/rtpN), which means that some input is pending for the real-time endpoint. The argument is the size of the incoming message.

**5.39.2.10 #define XDDP\_EVTNOBUF 4**

**Monitor** memory shortage for non real-time datagrams.

XDDP\_EVTNOBUF is sent when no memory is available from the pool to hold the message currently sent from the non real-time endpoint. The argument is the size of the failed allocation. Upon return from the callback, the caller will block and retry until enough space is available from the pool; during that process, the callback might be invoked multiple times, each time a new attempt to get the required memory fails.

**5.39.2.11 #define XDDP\_EVTOUT 2**

**Monitor** reads from the non real-time endpoint.

XDDP\_EVTOUT is sent when the non real-time endpoint successfully reads a complete message (i.e. via /dev/rtpN). The argument is the size of the outgoing message.

**5.39.2.12 #define XDDP\_LABEL 1**

XDDP label assignment.

ASCII label strings can be attached to XDDP ports, so that opening the non-RT endpoint can be done by specifying this symbolic device name rather than referring to a raw pseudo-device entry (i.e. /dev/rtpN).

When available, this label will be registered when binding, in addition to the port number (see [XDDP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

#### Parameters

- [in] *level* [SOL\\_XDDP](#)
- [in] *optname* [XDDP\\_LABEL](#)
- [in] *optval* Pointer to struct [rtipc\\_port\\_label](#)
- [in] *optlen* sizeof(struct rtipc\_port\_label)

#### Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* invalid)

#### Calling context:

RT/non-RT

#### 5.39.2.13 #define XDDP\_MONITOR 4

XDDP monitoring callback.

Other RTDM drivers may install a user-defined callback via the [rtdm\\_setsockopt](#) call from the inter-driver API, in order to collect particular events occurring on the channel.

This notification mechanism is particularly useful to monitor a channel asynchronously while performing other tasks.

The user-provided routine will be passed the RTDM file descriptor of the socket receiving the event, the event code, and an optional argument. Four events are currently defined, see [XDDP\\_EVENTS](#).

The XDDP\_EVTIN and XDDP\_EVTOUT events are fired on behalf of a fully atomic context; therefore, care must be taken to keep their overhead low. In those cases, the Xenomai services that may be called from the callback are restricted to the set allowed to a real-time interrupt handler.

#### Parameters

- [in] *level* [SOL\\_XDDP](#)
- [in] *optname* [XDDP\\_MONITOR](#)
- [in] *optval* Pointer to a pointer to function of type `int (*)(int fd, int event, long arg)`, containing the address of the user-defined callback. Passing a NULL callback pointer in *optval* disables monitoring.
- [in] *optlen* sizeof(int (\*)(int fd, int event, long arg))



**Returns**

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EPERM (Operation not allowed from user-space)
- -EINVAL (*optlen* is invalid)

**Calling context:**

RT/non-RT, kernel space only

**5.39.2.14 #define XDDP\_POOLSZ 2**

XDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

**Note**

: the pool memory is obtained from the host allocator by the [bind call](#).

**Parameters**

[in] *level* [SOL\\_XDDP](#)

[in] *optname* XDDP\_POOLSZ

[in] *optval* Pointer to a variable of type `size_t`, containing the required size of the local pool to reserve at binding time

[in] *optlen* `sizeof(size_t)`

**Returns**

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* invalid or *\*optval* is zero)

**Calling context:**

RT/non-RT

### 5.39.3 Enumeration Type Documentation

#### 5.39.3.1 anonymous enum

Enumerator:

**IPCPROTO\_IPC** Default protocol (IDDP).

**IPCPROTO\_XDDP** Cross-domain datagram protocol (RT <-> non-RT). Real-time Xenomai threads and regular Linux threads may want to exchange data in a way that does not require the former to leave the real-time domain (i.e. secondary mode). The RTDM-based XDDP protocol is available for this purpose.

On the Linux domain side, pseudo-device files named `/dev/rtp<minor>` give regular POSIX threads access to non real-time communication endpoints, via the standard character-based I/O interface. On the Xenomai domain side, sockets may be bound to XDDP ports, which act as proxies to send and receive data to/from the associated pseudo-device files. Ports and pseudo-device minor numbers are paired, meaning that e.g. port 7 will proxy the traffic to/from `/dev/rtp7`.

All data sent through a bound/connected XDDP socket via `sendto(2)` or `write(2)` will be passed to the peer endpoint in the Linux domain, and made available for reading via the standard `read(2)` system call. Conversely, all data sent using `write(2)` through the non real-time endpoint will be conveyed to the real-time socket endpoint, and made available to the `recvfrom(2)` or `read(2)` system calls.

**IPCPROTO\_IDDP** Intra-domain datagram protocol (RT <-> RT). The RTDM-based IDDP protocol enables real-time threads to exchange datagrams within the Xenomai domain, via socket endpoints.

**IPCPROTO\_BUFP** Buffer protocol (RT <-> RT, byte-oriented). The RTDM-based BUFP protocol implements a lightweight, byte-oriented, one-way Producer-Consumer data path. All messages written are buffered into a single memory area in strict FIFO order, until read by the consumer.

This protocol always prevents short writes, and only allows short reads when a potential deadlock situation arises (i.e. readers and writers waiting for each other indefinitely).

### 5.39.4 Function Documentation

#### 5.39.4.1 `int bind__AF_RTIPC ( int sockfd, const struct sockaddr_ipc * addr, socklen_t addrlen )`

Bind a RTIPC socket to a port.

Bind the socket to a destination port.

Parameters

[in] *sockfd* The RTDM file descriptor obtained from the socket creation call.

[in] *addr* The address to bind the socket to (see struct [sockaddr\\_ipc](#)). The meaning of such address depends on the RTIPC protocol in use for the socket:

- **IPCPROTO\_XDDP**

This action creates an endpoint for channelling traffic between the Xenomai and Linux domains.

*sipc\_family* must be `AF_RTIPC`, *sipc\_port* is either `-1`, or a valid free port number between 0 and `CONFIG_XENO_OPT_PIPE_NRDEV-1`.

If *sipc\_port* is -1, a free port will be assigned automatically.

Upon success, the pseudo-device /dev/rtpN will be reserved for this communication channel, where *N* is the assigned port number. The non real-time side shall open this device to exchange data over the bound socket.

If a label was assigned (see [XDDP\\_LABEL](#)) prior to binding the socket to a port, a registry link referring to the created pseudo-device will be automatically set up as /proc/xenomai/registry/rtipc/xddp/*label*, where *label* is the label string passed to setsockopt() for the [XDDP\\_LABEL](#) option.

- IPCPROTO\_IDDP

This action creates an endpoint for exchanging datagrams within the Xenomai domain.

*sipc\_family* must be AF\_RTIPC, *sipc\_port* is either -1, or a valid free port number between 0 and CONFIG\_XENO\_OPT\_IDDP\_NRPORT-1.

If *sipc\_port* is -1, a free port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see [IDDP\\_LABEL](#)) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as /proc/xenomai/registry/rtipc/iddp/*label*, where *label* is the label string passed to setsockopt() for the [IDDP\\_LABEL](#) option.

- IPCPROTO\_BUF

This action creates an endpoint for a one-way byte stream within the Xenomai domain.

*sipc\_family* must be AF\_RTIPC, *sipc\_port* is either -1, or a valid free port number between 0 and CONFIG\_XENO\_OPT\_BUF\_NRPORT-1.

If *sipc\_port* is -1, an available port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see [BUFP\\_LABEL](#)) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as /proc/xenomai/registry/rtipc/bufp/*label*, where *label* is the label string passed to setsockopt() for the [BUFP\\_LABEL](#) option.

### Parameters

[in] *addrlen* The size in bytes of the structure pointed to by *addr*.

### Returns

In addition to the standard error codes for bind(2), the following specific error code may be returned:

- -EFAULT (Invalid data address given)
- -ENOMEM (Not enough memory)
- -EINVAL (Invalid parameter)
- -EADDRINUSE (Socket already bound to a port, or no port available)

### Calling context:

non-RT

#### 5.39.4.2 `int close__AF_RTIPC ( int sockfd )`

Close a RTIPC socket descriptor.

Blocking calls to any of the [sendmsg](#) or [recvmsg](#) functions will be unblocked when the socket is closed and return with an error.

##### Returns

In addition to the standard error codes for `close(2)`, the following specific error code may be returned: none

##### Calling context:

non-RT

#### 5.39.4.3 `int connect__AF_RTIPC ( int sockfd, const struct sockaddr_ipc * addr, socklen_t addrlen )`

Initiate a connection on a RTIPC socket.

##### Parameters

[in] *sockfd* The RTDM file descriptor obtained from the socket creation call.

[in] *addr* The address to connect the socket to (see struct [sockaddr\\_ipc](#)).

- If `sipc_port` is a valid port for the protocol, it is used verbatim and the connection succeeds immediately, regardless of whether the destination is bound at the time of the call.
- If `sipc_port` is -1 and a label was assigned to the socket, `connect()` blocks for the requested amount of time (see [SO\\_RCVTIMEO](#)) until a socket is bound to the same label via `bind(2)` (see [XDDP\\_LABEL](#), [IDDP\\_LABEL](#), [BUFP\\_LABEL](#)), in which case a connection is established between both endpoints.
- If `sipc_port` is -1 and no label was assigned to the socket, the default destination address is cleared, meaning that any subsequent write to the socket will return `-EDESTADDRREQ`, until a valid destination address is set via `connect(2)` or `bind(2)`.

##### Parameters

[in] *addrlen* The size in bytes of the structure pointed to by *addr*.

##### Returns

In addition to the standard error codes for `connect(2)`, the following specific error code may be returned: none.

##### Calling context:

RT/non-RT

#### 5.39.4.4 `int getpeername__AF_RTIPC ( int sockfd, struct sockaddr_ipc * addr, socklen_t * addrlen )`

Get socket peer.

The name of the remote endpoint for the socket is copied back (see struct [sockaddr\\_ipc](#)). This is the default destination address for messages sent on the socket. It can be set either explicitly via `connect(2)`, or implicitly via `bind(2)` if no `connect(2)` was called prior to binding the socket to a port, in which case both the local and remote names are equal.

##### Returns

In addition to the standard error codes for `getpeername(2)`, the following specific error code may be returned: none.

##### Calling context:

RT/non-RT

#### 5.39.4.5 `int getsockname__AF_RTIPC ( int sockfd, struct sockaddr_ipc * addr, socklen_t * addrlen )`

Get socket name.

The name of the local endpoint for the socket is copied back (see struct [sockaddr\\_ipc](#)).

##### Returns

In addition to the standard error codes for `getsockname(2)`, the following specific error code may be returned: none.

##### Calling context:

RT/non-RT

#### 5.39.4.6 `int getsockopt__AF_RTIPC ( int sockfd, int level, int optname, void * optval, socklen_t * optlen )`

Get options on RTIPC sockets.

These functions allow to get various socket options. Supported Levels and Options:

- Level [SOL\\_SOCKET](#)
- Level [SOL\\_XDDP](#)
- Level [SOL\\_IDDP](#)
- Level [SOL\\_BUFP](#)

##### Returns

In addition to the standard error codes for `getsockopt(2)`, the following specific error code may be returned: follow the option links above.

##### Calling context:

RT/non-RT

#### 5.39.4.7 `ssize_t recvmsg__AF_RTIPC ( int sockfd, struct msghdr * msg, int flags )`

Receive a message from a RTIPC socket.

##### Parameters

- [in] *sockfd* The RTDM file descriptor obtained from the socket creation call.
- [out] *msg* The address the message header will be copied at.
- [in] *flags* Operation flags:
  - MSG\_DONTWAIT Non-blocking I/O operation. The caller will not be blocked whenever no message is immediately available for receipt at the time of the call, but will rather return with -EWOULDBLOCK.

##### Note

[IPCPROTO\\_BUF](#) does not allow for short reads and always returns the requested amount of bytes, except in one situation: whenever some writer is waiting for sending data upon a buffer full condition, while the caller would have to wait for receiving a complete message. This is usually the sign of a pathological use of the BUFP socket, like defining an incorrect buffer size via [BUFP\\_BUFSZ](#). In that case, a short read is allowed to prevent a deadlock.

##### Returns

In addition to the standard error codes for `recvmsg(2)`, the following specific error code may be returned: none.

##### Calling context:

RT

#### 5.39.4.8 `ssize_t sendmsg__AF_RTIPC ( int sockfd, const struct msghdr * msg, int flags )`

Send a message on a RTIPC socket.

##### Parameters

- [in] *sockfd* The RTDM file descriptor obtained from the socket creation call.
- [in] *msg* The address of the message header conveying the datagram.
- [in] *flags* Operation flags:
  - MSG\_OOB Send out-of-band message. For all RTIPC protocols except [IPCPROTO\\_BUF](#), sending out-of-band data actually means pushing them to the head of the receiving queue, so that the reader will always receive them before normal messages. [IPCPROTO\\_BUF](#) does not support out-of-band sending.
  - MSG\_DONTWAIT Non-blocking I/O operation. The caller will not be blocked whenever the message cannot be sent immediately at the time of the call (e.g. memory shortage), but will rather return with -EWOULDBLOCK. Unlike other RTIPC protocols, [IPCPROTO\\_XDDP](#) accepts but never considers MSG\_DONTWAIT since writing to a real-time XDDP endpoint is inherently a non-blocking operation.

- `MSG_MORE` Accumulate data before sending. This flag is accepted by the [IPCPROTO\\_XDDP](#) protocol only, and tells the send service to accumulate the outgoing data into an internal streaming buffer, instead of issuing a datagram immediately for it. See [XDDP\\_BUFSZ](#) for more.

#### Note

No RTIPC protocol allows for short writes, and only complete messages are sent to the peer.

#### Returns

In addition to the standard error codes for `sendmsg(2)`, the following specific error code may be returned: none.

#### Calling context:

RT

**5.39.4.9** `int setsockopt__AF_RTIPC ( int sockfd, int level, int optname, const void * optval, socklen_t optlen )`

Set options on RTIPC sockets.

These functions allow to set various socket options. Supported Levels and Options:

- Level [SOL\\_SOCKET](#)
- Level [SOL\\_XDDP](#)
- Level [SOL\\_IDDP](#)
- Level [SOL\\_BUF](#)

#### Returns

In addition to the standard error codes for `setsockopt(2)`, the following specific error code may be returned: follow the option links above.

#### Calling context:

non-RT

**5.39.4.10** `int socket__AF_RTIPC ( int domain = AF_RTIPC, int type = SOCK_DGRAM, int protocol )`

Create an endpoint for communication in the `AF_RTIPC` domain.

#### Parameters

- [in] *domain* The communication domain. Must be `AF_RTIPC`.
- [in] *type* The socket type. Must be `SOCK_DGRAM`.
- [in] *protocol* Any of [IPCPROTO\\_XDDP](#), [IPCPROTO\\_IDDP](#), or [IPCPROTO\\_BUF](#). [IPCPROTO\\_IPC](#) is also valid, and refers to the default RTIPC protocol, namely [IPCPROTO\\_IDDP](#).

### Returns

In addition to the standard error codes for `socket(2)`, the following specific error code may be returned:

- `-ENOPROTOOPT` (Protocol is known, but not compiled in the RTIPC driver). See [RTIPC protocols](#) for available protocols.

### Calling context:

non-RT

## 5.40 Serial Devices

Collaboration diagram for Serial Devices:



### Data Structures

- struct [rtser\\_config](#)  
*Serial device configuration.*
- struct [rtser\\_status](#)  
*Serial device status.*
- struct [rtser\\_event](#)  
*Additional information about serial device events.*

### Files

- file [rtserial.h](#)  
*Real-Time Driver Model for Xenomai, serial device profile header.*

### Defines

- `#define` [RTSER\\_RTIOC\\_BREAK\\_CTL](#) `_IOR(RTIOC_TYPE_SERIAL, 0x06, int)`  
*Set or clear break on UART output line.*



## RTSER\_BREAK\_xxx

Break control

- typedef struct [rtser\\_config](#) [rtser\\_config\\_t](#)  
*Serial device configuration.*
- typedef struct [rtser\\_status](#) [rtser\\_status\\_t](#)  
*Serial device status.*
- typedef struct [rtser\\_event](#) [rtser\\_event\\_t](#)  
*Additional information about serial device events.*
- #define [RTSER\\_BREAK\\_CLR](#) 0x00  
*Serial device configuration.*
- #define [RTSER\\_BREAK\\_SET](#) 0x01  
*Serial device configuration.*
- #define [RTIOC\\_TYPE\\_SERIAL](#) RTDM\_CLASS\_SERIAL  
*Serial device configuration.*

## RTSER\_DEF\_BAUD

Default baud rate

- #define [RTSER\\_DEF\\_BAUD](#) 9600

## RTSER\_xxx\_PARITY

Number of parity bits

- #define [RTSER\\_NO\\_PARITY](#) 0x00
- #define [RTSER\\_ODD\\_PARITY](#) 0x01
- #define [RTSER\\_EVEN\\_PARITY](#) 0x03
- #define [RTSER\\_DEF\\_PARITY](#) [RTSER\\_NO\\_PARITY](#)

## RTSER\_xxx\_BITS

Number of data bits

- #define [RTSER\\_5\\_BITS](#) 0x00
- #define [RTSER\\_6\\_BITS](#) 0x01
- #define [RTSER\\_7\\_BITS](#) 0x02
- #define [RTSER\\_8\\_BITS](#) 0x03
- #define [RTSER\\_DEF\\_BITS](#) [RTSER\\_8\\_BITS](#)

## RTSER\_xxx\_STOPB

Number of stop bits

- #define [RTSER\\_1\\_STOPB](#) 0x00  
*valid only in combination with 5 data bits*
- #define [RTSER\\_1\\_5\\_STOPB](#) 0x01  
*valid only in combination with 5 data bits*
- #define [RTSER\\_2\\_STOPB](#) 0x01  
*valid only in combination with 5 data bits*
- #define [RTSER\\_DEF\\_STOPB](#) RTSER\_1\_STOPB  
*valid only in combination with 5 data bits*

## RTSER\_xxx\_HAND

Handshake mechanisms

- #define [RTSER\\_NO\\_HAND](#) 0x00
- #define [RTSER\\_RTSCTS\\_HAND](#) 0x01
- #define [RTSER\\_DEF\\_HAND](#) RTSER\_NO\_HAND

## RTSER\_FIFO\_xxx

Reception FIFO interrupt threshold

- #define [RTSER\\_FIFO\\_DEPTH\\_1](#) 0x00
- #define [RTSER\\_FIFO\\_DEPTH\\_4](#) 0x40
- #define [RTSER\\_FIFO\\_DEPTH\\_8](#) 0x80
- #define [RTSER\\_FIFO\\_DEPTH\\_14](#) 0xC0
- #define [RTSER\\_DEF\\_FIFO\\_DEPTH](#) RTSER\_FIFO\_DEPTH\_1

## RTSER\_TIMEOUT\_xxx

Special timeout values, see also [RTDM\\_TIMEOUT\\_xxx](#)

- #define [RTSER\\_TIMEOUT\\_INFINITE](#) RTDM\_TIMEOUT\_INFINITE
- #define [RTSER\\_TIMEOUT\\_NONE](#) RTDM\_TIMEOUT\_NONE
- #define [RTSER\\_DEF\\_TIMEOUT](#) RTDM\_TIMEOUT\_INFINITE

## RTSER\_xxx\_TIMESTAMP\_HISTORY

Timestamp history control

- #define [RTSER\\_RX\\_TIMESTAMP\\_HISTORY](#) 0x01
- #define [RTSER\\_DEF\\_TIMESTAMP\\_HISTORY](#) 0x00

## RTSER\_EVENT\_xxx

Events bits

- #define RTSER\_EVENT\_RXPEND 0x01
- #define RTSER\_EVENT\_ERRPEND 0x02
- #define RTSER\_EVENT\_MODEMHI 0x04
- #define RTSER\_EVENT\_MODEMLO 0x08
- #define RTSER\_DEF\_EVENT\_MASK 0x00

## RTSER\_SET\_xxx

Configuration mask bits

- #define RTSER\_SET\_BAUD 0x0001
- #define RTSER\_SET\_PARITY 0x0002
- #define RTSER\_SET\_DATA\_BITS 0x0004
- #define RTSER\_SET\_STOP\_BITS 0x0008
- #define RTSER\_SET\_HANDSHAKE 0x0010
- #define RTSER\_SET\_FIFO\_DEPTH 0x0020
- #define RTSER\_SET\_TIMEOUT\_RX 0x0100
- #define RTSER\_SET\_TIMEOUT\_TX 0x0200
- #define RTSER\_SET\_TIMEOUT\_EVENT 0x0400
- #define RTSER\_SET\_TIMESTAMP\_HISTORY 0x0800
- #define RTSER\_SET\_EVENT\_MASK 0x1000

## RTSER\_LSR\_xxx

Line status bits

- #define RTSER\_LSR\_DATA 0x01
- #define RTSER\_LSR\_OVERRUN\_ERR 0x02
- #define RTSER\_LSR\_PARITY\_ERR 0x04
- #define RTSER\_LSR\_FRAMING\_ERR 0x08
- #define RTSER\_LSR\_BREAK\_IND 0x10
- #define RTSER\_LSR\_THR\_EMPTY 0x20
- #define RTSER\_LSR\_TRANSM\_EMPTY 0x40
- #define RTSER\_LSR\_FIFO\_ERR 0x80
- #define RTSER\_SOFT\_OVERRUN\_ERR 0x0100

## RTSER\_MSR\_xxx

Modem status bits

- #define RTSER\_MSR\_DCTS 0x01
- #define RTSER\_MSR\_DDSD 0x02
- #define RTSER\_MSR\_TERI 0x04
- #define RTSER\_MSR\_DDCD 0x08

- `#define RTSER_MSR_CTS 0x10`
- `#define RTSER_MSR_DSR 0x20`
- `#define RTSER_MSR_RI 0x40`
- `#define RTSER_MSR_DCD 0x80`

## RTSER\_MCR\_xxx

Modem control bits

- `#define RTSER_MCR_DTR 0x01`
- `#define RTSER_MCR_RTS 0x02`
- `#define RTSER_MCR_OUT1 0x04`
- `#define RTSER_MCR_OUT2 0x08`
- `#define RTSER_MCR_LOOP 0x10`

## Sub-Classes of RTDM\_CLASS\_SERIAL

- `#define RTDM_SUBCLASS_16550A 0`

## IOCTLs

Serial device IOCTLs

- `#define RTSER_RTIOC_GET_CONFIG _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_config)`  
*Get serial device configuration.*
- `#define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_config)`  
*Set serial device configuration.*
- `#define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_status)`  
*Get serial device status.*
- `#define RTSER_RTIOC_GET_CONTROL _IOR(RTIOC_TYPE_SERIAL, 0x03, int)`  
*Get serial device's modem control register.*
- `#define RTSER_RTIOC_SET_CONTROL _IOW(RTIOC_TYPE_SERIAL, 0x04, int)`  
*Set serial device's modem control register.*
- `#define RTSER_RTIOC_WAIT_EVENT _IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_event)`  
*Wait on serial device events according to previously set mask.*

### 5.40.1 Detailed Description

This is the common interface a RTDM-compliant serial device has to provide. Feel free to comment on this profile via the Xenomai mailing list ([Xenomai-core@gna.org](mailto:Xenomai-core@gna.org)) or directly to the author ([jan.kiszka@web.de](mailto:jan.kiszka@web.de)).

**Profile Revision:** 3

#### Device Characteristics

**Device Flags:** RTDM\_NAMED\_DEVICE, RTDM\_EXCLUSIVE

**Device Name:** "rtser<N>", N >= 0

**Device Class:** RTDM\_CLASS\_SERIAL

#### Supported Operations

##### Open

Environments: non-RT (RT optional, deprecated)

Specific return values: none

##### Close

Environments: non-RT (RT optional, deprecated)

Specific return values: none

##### IOCTL

Mandatory Environments: see [below](#)

Specific return values: see [below](#)

##### Read

Environments: RT (non-RT optional)

Specific return values:

- -ETIMEDOUT
- -EINTR (interrupted explicitly or by signal)
- -EAGAIN (no data available in non-blocking mode)
- -EBADF (device has been closed while reading)
- -EIO (hardware error or broken bit stream)

##### Write

Environments: RT (non-RT optional)

Specific return values:

- -ETIMEDOUT
- -EINTR (interrupted explicitly or by signal)
- -EAGAIN (no data written in non-blocking mode)
- -EBADF (device has been closed while writing)

### 5.40.2 Define Documentation

#### 5.40.2.1 #define RTSER\_RTIOC\_BREAK\_CTL\_IOR(RTIOC\_TYPE\_SERIAL, 0x06, int)

Set or clear break on UART output line.

#### Parameters

[in] *arg* RTSER\_BREAK\_SET or RTSER\_BREAK\_CLR (int)

**Returns**

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

**Note**

A set break condition may also be cleared on UART line reconfiguration.

Rescheduling: never.

**5.40.2.2 #define RTSER\_RTIOC\_GET\_CONFIG \_IOR(RTIOC\_TYPE\_SERIAL, 0x00, struct rtser\_config)**

Get serial device configuration.

**Parameters**

[out] *arg* Pointer to configuration buffer (struct [rtser\\_config](#))

**Returns**

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

**5.40.2.3 #define RTSER\_RTIOC\_GET\_CONTROL \_IOR(RTIOC\_TYPE\_SERIAL, 0x03, int)**

Get serial device's modem control register.

**Parameters**

[out] *arg* Pointer to variable receiving the content (int, see [RTSER\\_MCR\\_xxx](#))

**Returns**

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.40.2.4 `#define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_status)`

Get serial device status.

##### Parameters

[out] *arg* Pointer to status buffer (struct [rtser\\_status](#))

##### Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

##### Note

The error states `RTSER_LSR_OVERRUN_ERR`, `RTSER_LSR_PARITY_ERR`, `RTSER_LSR_FRAMING_ERR`, and `RTSER_SOFT_OVERRUN_ERR` that may have occurred during previous read accesses to the device will be saved for being reported via this IOCTL. Upon return from `RTSER_RTIOC_GET_STATUS`, the saved state will be cleared.

Rescheduling: never.

#### 5.40.2.5 `#define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_config)`

Set serial device configuration.

##### Parameters

[in] *arg* Pointer to configuration buffer (struct [rtser\\_config](#))

##### Returns

0 on success, otherwise:

- -EPERM is returned if the caller's context is invalid, see note below.
- -ENOMEM is returned if a new history buffer for timestamps cannot be allocated.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

#### Note

If [rtser\\_config](#) contains a valid timestamp\_history and the addressed device has been opened in non-real-time context, this IOCTL must be issued in non-real-time context as well. Otherwise, this command will fail.

Rescheduling: never.

#### 5.40.2.6 #define RTSER\_RTIOC\_SET\_CONTROL\_IOW(RTIOC\_TYPE\_SERIAL, 0x04, int)

Set serial device's modem control register.

#### Parameters

[in] *arg* New control register content (int, see [RTSER\\_MCR\\_xxx](#))

#### Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

#### 5.40.2.7 #define RTSER\_RTIOC\_WAIT\_EVENT\_IOR(RTIOC\_TYPE\_SERIAL, 0x05, struct rtser\_event)

Wait on serial device events according to previously set mask.

#### Parameters

[out] *arg* Pointer to event information buffer (struct [rtser\\_event](#))



**Returns**

0 on success, otherwise:

- -EBUSY is returned if another task is already waiting on events of this device.
- -EBADF is returned if the file descriptor is invalid or the device has just been closed.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

## 5.41 Testing Devices

Collaboration diagram for Testing Devices:

**Files**

- file [rttesting.h](#)  
*Real-Time Driver Model for Xenomai, testing device profile header.*

**Sub-Classes of RTDM\_CLASS\_TESTING**

- `#define RTDM_SUBCLASS_TIMERBENCH 0`  
*subclass name: "timerbench"*
- `#define RTDM_SUBCLASS_IRQBENCH 1`  
*subclass name: "irqbench"*
- `#define RTDM_SUBCLASS_SWITCHTEST 2`  
*subclass name: "switchtest"*
- `#define RTDM_SUBCLASS_RTDMTTEST 3`  
*subclass name: "rtdm"*

## IOCTLs

Testing device IOCTLs

- `#define RTTST_RTIOC_INTERM_BENCH_RES _IOWR(RTIOC_TYPE_TESTING, 0x00, struct rttst_interm_bench_res)`
- `#define RTTST_RTIOC_TMBENCH_START _IOW(RTIOC_TYPE_TESTING, 0x10, struct rttst_tmbench_config)`
- `#define RTTST_RTIOC_TMBENCH_STOP _IOWR(RTIOC_TYPE_TESTING, 0x11, struct rttst_overall_bench_res)`
- `#define RTTST_RTIOC_IRQBENCH_START _IOW(RTIOC_TYPE_TESTING, 0x20, struct rttst_irqbench_config)`
- `#define RTTST_RTIOC_IRQBENCH_STOP _IO(RTIOC_TYPE_TESTING, 0x21)`
- `#define RTTST_RTIOC_IRQBENCH_GET_STATS _IOR(RTIOC_TYPE_TESTING, 0x22, struct rttst_irqbench_stats)`
- `#define RTTST_RTIOC_IRQBENCH_WAIT_IRQ _IO(RTIOC_TYPE_TESTING, 0x23)`
- `#define RTTST_RTIOC_IRQBENCH_REPLY_IRQ _IO(RTIOC_TYPE_TESTING, 0x24)`
- `#define RTTST_RTIOC_SWTEST_SET_TASKS_COUNT _IOW(RTIOC_TYPE_TESTING, 0x30, unsigned long)`
- `#define RTTST_RTIOC_SWTEST_SET_CPU _IOW(RTIOC_TYPE_TESTING, 0x31, unsigned long)`
- `#define RTTST_RTIOC_SWTEST_REGISTER_UTASK _IOW(RTIOC_TYPE_TESTING, 0x32, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_CREATE_KTASK _IOWR(RTIOC_TYPE_TESTING, 0x33, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_PEND _IOR(RTIOC_TYPE_TESTING, 0x34, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_SWITCH_TO _IOR(RTIOC_TYPE_TESTING, 0x35, struct rttst_swtest_dir)`
- `#define RTTST_RTIOC_SWTEST_GET_SWITCHES_COUNT _IOR(RTIOC_TYPE_TESTING, 0x36, unsigned long)`
- `#define RTTST_RTIOC_SWTEST_GET_LAST_ERROR _IOR(RTIOC_TYPE_TESTING, 0x37, struct rttst_swtest_error)`
- `#define RTTST_RTIOC_SWTEST_SET_PAUSE _IOW(RTIOC_TYPE_TESTING, 0x38, unsigned long)`
- `#define RTTST_RTIOC_RTDM_DEFER_CLOSE _IOW(RTIOC_TYPE_TESTING, 0x40, unsigned long)`

### 5.41.1 Detailed Description

This group of devices is intended to provide in-kernel testing results. Feel free to comment on this profile via the Xenomai mailing list ([xenomai-core@gna.org](mailto:xenomai-core@gna.org)) or directly to the author ([jan.kiszka@web.de](mailto:jan.kiszka@web.de)).

**Profile Revision:** 2

#### Device Characteristics

**Device Flags:** RTDM\_NAMED\_DEVICE

**Device Name:** "rttest[-<subclass>]<N>", N >= 0, optional subclass name to simplify device discovery

**Device Class:** RTDM\_CLASS\_TESTING

## Supported Operations

### Open

Environments: non-RT (RT optional, deprecated)

Specific return values: none

### Close

Environments: non-RT (RT optional, deprecated)

Specific return values: none

### IOCTL

Mandatory Environments: see [TSTIOCTLs](#) below

Specific return values: see [TSTIOCTLs](#) below

## 5.42 Sched

## Data Structures

- struct [xnsched](#)  
*Scheduling information structure.*

## Files

- file [sched.h](#)  
*Scheduler interface header.*
- file [sched-idle.c](#)  
*Idle scheduling class implementation (i.e. Linux placeholder).*
- file [sched-rt.c](#)  
*Common real-time scheduling class implementation (FIFO + RR).*
- file [sched-sporadic.c](#)  
*POSIX SCHED\_SPORADIC scheduling class.*
- file [sched-tp.c](#)  
*Temporal partitioning (typical of IMA systems).*
- file [sched.c](#)

## Typedefs

- typedef struct [xnsched](#) [xnsched\\_t](#)  
*Scheduling information structure.*

## Functions

- static void [xnsched\\_rotate](#) (struct [xnsched](#) \*sched, struct [xnsched\\_class](#) \*sched\_class, const union [xnsched\\_policy\\_param](#) \*sched\_param)

*Rotate a scheduler runqueue.*

### 5.42.1 Function Documentation

**5.42.1.1** `void xnsched_rotate ( struct xnsched * sched, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param ) [inline, static]`

Rotate a scheduler runqueue.

The specified scheduling class is requested to rotate its runqueue for the given scheduler. Rotation is performed according to the scheduling parameter specified by *sched\_param*.

#### Note

The nucleus supports round-robin scheduling for the members of the RT class.

#### Parameters

*sched* The per-CPU scheduler hosting the target scheduling class.

*sched\_class* The scheduling class which should rotate its runqueue.

*sched\_param* The scheduling parameter providing rotation information to the specified scheduling class.

Environments:

This service should be called from:

- Kernel-based task
- Interrupt service routine
- User-space task (primary mode only)

Rescheduling: never.

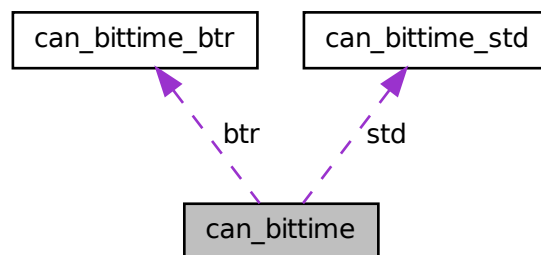
## Chapter 6

# Data Structure Documentation

### 6.1 `can_bittime` Struct Reference

Custom CAN bit-time definition.

Collaboration diagram for `can_bittime`:



#### Data Fields

- [`can\_bittime\_type\_t`](#) type  
*Type of bit-time definition.*
- struct [`can\_bittime\_std`](#) `std`  
*Standard bit-time.*
- struct [`can\_bittime\_btr`](#) `btr`  
*Hardware-specific BTR bit-time.*

### 6.1.1 Detailed Description

Custom CAN bit-time definition.

The documentation for this struct was generated from the following file:

- `include/rtdm/rtcan.h`

## 6.2 `can_bittime_btr` Struct Reference

Hardware-specific BTR bit-times.

### Data Fields

- `uint8_t btr0`  
*Bus timing register 0.*
- `uint8_t btr1`  
*Bus timing register 1.*

### 6.2.1 Detailed Description

Hardware-specific BTR bit-times.

The documentation for this struct was generated from the following file:

- `include/rtdm/rtcan.h`

## 6.3 `can_bittime_std` Struct Reference

Standard bit-time parameters according to Bosch.

### Data Fields

- `uint32_t brp`  
*Baud rate prescaler.*
- `uint8_t prop\_seg`  
*from 1 to 8*
- `uint8_t phase\_seg1`  
*from 1 to 8*
- `uint8_t phase\_seg2`  
*from 1 to 8*

- uint8\_t [sjw](#):7  
*from 1 to 4*
- uint8\_t [sam](#):1  
*1 - enable triple sampling*

### 6.3.1 Detailed Description

Standard bit-time parameters according to Bosch.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtdmcan.h](#)

## 6.4 can\_filter Struct Reference

Filter for reception of CAN messages.

### Data Fields

- uint32\_t [can\\_id](#)  
*CAN ID which must match with incoming IDs after passing the mask.*
- uint32\_t [can\\_mask](#)  
*Mask which is applied to incoming IDs.*

### 6.4.1 Detailed Description

Filter for reception of CAN messages. This filter works as follows: A received CAN ID is AND'ed bitwise with [can\\_mask](#) and then compared to [can\\_id](#). This also includes the [CAN\\_EFF\\_FLAG](#) and [CAN\\_RTR\\_FLAG](#) of [CAN\\_xxx\\_FLAG](#). If this comparison is true, the message will be received by the socket. The logic can be inverted with the [can\\_id](#) flag [CAN\\_INV\\_FILTER](#) :

```
if (can_id & CAN_INV_FILTER) {  
    if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))  
        accept-message;  
} else {  
    if ((received_can_id & can_mask) == can_id)  
        accept-message;  
}
```

Multiple filters can be arranged in a filter list and set with [Sockopts](#). If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

## 6.4.2 Field Documentation

### 6.4.2.1 `uint32_t can_filter::can_id`

CAN ID which must match with incoming IDs after passing the mask.

The filter logic can be inverted with the flag [CAN\\_INV\\_FILTER](#).

### 6.4.2.2 `uint32_t can_filter::can_mask`

Mask which is applied to incoming IDs.

See [CAN ID masks](#) if exactly one CAN ID should come through.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtcan.h](#)

## 6.5 `can_frame` Struct Reference

Raw CAN frame.

### Public Member Functions

- `uint8_t data[8] \_\_attribute\_\_\(\(aligned\(8\)\)\)`  
*Payload data bytes.*

### Data Fields

- `can\_id\_t can\_id`  
*CAN ID of the frame.*
- `uint8_t can\_dlc`  
*Size of the payload in bytes.*

### 6.5.1 Detailed Description

Raw CAN frame. Central structure for receiving and sending CAN frames.

### 6.5.2 Field Documentation

#### 6.5.2.1 `can_id_t can_frame::can_id`

CAN ID of the frame.

See [CAN ID flags](#) for special bits.

The documentation for this struct was generated from the following file:

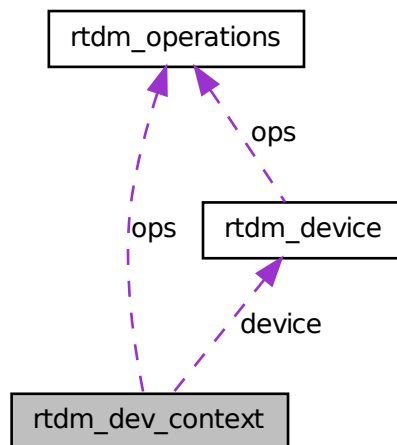


- `include/rtdm/rtdm.h`

## 6.6 rtdm\_dev\_context Struct Reference

Device context.

Collaboration diagram for `rtdm_dev_context`:



### Data Fields

- unsigned long `context_flags`  
*Context flags, see [Context Flags](#) for details.*
- int `fd`  
*Associated file descriptor.*
- atomic\_t `close_lock_count`  
*Lock counter of context, held while structure is referenced by an operation handler.*
- struct `rtdm_operations` \* `ops`  
*Set of active device operation handlers.*
- struct `rtdm_device` \* `device`  
*Reference to owning device.*
- struct `rtdm_devctx_reserved` `reserved`  
*Data stored by RTDM inside a device context (internal use only).*

- char [dev\\_private](#) [0]

*Begin of driver defined context data structure.*

### 6.6.1 Detailed Description

Device context. A device context structure is associated with every open device instance. RTDM takes care of its creation and destruction and passes it to the operation handlers when being invoked.

Drivers can attach arbitrary data immediately after the official structure. The size of this data is provided via [rtdm\\_device.context\\_size](#) during device registration.

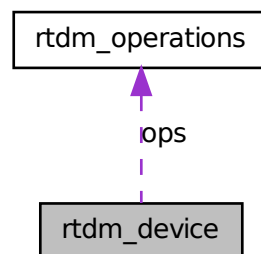
The documentation for this struct was generated from the following file:

- [include/rtdm/rtdm\\_driver.h](#)

## 6.7 rtdm\_device Struct Reference

RTDM device.

Collaboration diagram for rtdm\_device:



### Data Fields

- int [struct\\_version](#)

*Revision number of this structure, see [Driver Versioning](#) defines.*

- int [device\\_flags](#)

*Device flags, see [Device Flags](#) for details.*

- size\_t [context\\_size](#)

*Size of driver defined appendix to struct [rtdm\\_dev\\_context](#).*

- char [device\\_name](#) [RTDM\_MAX\_DEVNAME\_LEN+1]  
*Named device identification (orthogonal to Linux device name space).*
- int [protocol\\_family](#)  
*Protocol device identification: protocol family (PF\_XXX).*
- int [socket\\_type](#)  
*Protocol device identification: socket type (SOCK\_XXX).*
- [rtdm\\_open\\_handler\\_t](#) [open\\_rt](#)  
*Named device instance creation for real-time contexts, optional (but deprecated) if [open\\_nrt](#) is non-NULL, ignored for protocol devices.*
- [rtdm\\_open\\_handler\\_t](#) [open\\_nrt](#)  
*Named device instance creation for non-real-time contexts, optional if [open\\_rt](#) is non-NULL, ignored for protocol devices.*
- [rtdm\\_socket\\_handler\\_t](#) [socket\\_rt](#)  
*Protocol socket creation for real-time contexts, optional (but deprecated) if [socket\\_nrt](#) is non-NULL, ignored for named devices.*
- [rtdm\\_socket\\_handler\\_t](#) [socket\\_nrt](#)  
*Protocol socket creation for non-real-time contexts, optional if [socket\\_rt](#) is non-NULL, ignored for named devices.*
- struct [rtdm\\_operations](#) [ops](#)  
*Default operations on newly opened device instance.*
- int [device\\_class](#)  
*Device class ID, see [RTDM\\_CLASS\\_XXX](#).*
- int [device\\_sub\\_class](#)  
*Device sub-class, see [RTDM\\_SUBCLASS\\_XXX](#) definition in the [Device Profiles](#).*
- int [profile\\_version](#)  
*Supported device profile version.*
- const char \* [driver\\_name](#)  
*Informational driver name (reported via /proc).*
- int [driver\\_version](#)  
*Driver version, see [Driver Versioning](#) defines.*
- const char \* [peripheral\\_name](#)  
*Informational peripheral name the device is attached to (reported via /proc).*
- const char \* [provider\\_name](#)  
*Informational driver provider name (reported via /proc).*
- const char \* [proc\\_name](#)

*Name of /proc entry for the device, must not be NULL.*

- int [device\\_id](#)  
*Driver definable device ID.*
- void \* [device\\_data](#)  
*Driver definable device data.*
- struct rtdm\_dev\_reserved [reserved](#)  
*Data stored by RTDM inside a registered device (internal use only).*

### 6.7.1 Detailed Description

RTDM device. This structure specifies a RTDM device. As some fields, especially the reserved area, will be modified by RTDM during runtime, the structure must not reside in write-protected memory.

### 6.7.2 Field Documentation

#### 6.7.2.1 [rtdm\\_open\\_handler\\_t](#) [rtdm\\_device::open\\_rt](#)

Named device instance creation for real-time contexts, optional (but deprecated) if [open\\_nrt](#) is non-NULL, ignored for protocol devices.

#### Deprecated

Only use non-real-time open handler in new drivers.

Referenced by [rtdm\\_dev\\_register\(\)](#).

#### 6.7.2.2 [rtdm\\_socket\\_handler\\_t](#) [rtdm\\_device::socket\\_rt](#)

Protocol socket creation for real-time contexts, optional (but deprecated) if [socket\\_nrt](#) is non-NULL, ignored for named devices.

#### Deprecated

Only use non-real-time socket creation handler in new drivers.

Referenced by [rtdm\\_dev\\_register\(\)](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/rtdm\\_driver.h](#)

## 6.8 [rtdm\\_device\\_info](#) Struct Reference

Device information.

## Data Fields

- int [device\\_flags](#)  
*Device flags, see [Device Flags](#) for details.*
- int [device\\_class](#)  
*Device class ID, see [RTDM\\_CLASS\\_xxx](#).*
- int [device\\_sub\\_class](#)  
*Device sub-class, either [RTDM\\_SUBCLASS\\_GENERIC](#) or a [RTDM\\_SUBCLASS\\_xxx](#) definition of the related [Device Profile](#).*
- int [profile\\_version](#)  
*Supported device profile version.*

### 6.8.1 Detailed Description

Device information.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtdm.h](#)

## 6.9 rtdm\_operations Struct Reference

Device operations.

## Data Fields

### Common Operations

- [rtdm\\_close\\_handler\\_t close\\_rt](#)  
*Close handler for real-time contexts (optional, deprecated).*
- [rtdm\\_close\\_handler\\_t close\\_nrt](#)  
*Close handler for non-real-time contexts (required).*
- [rtdm\\_ioctl\\_handler\\_t ioctl\\_rt](#)  
*IOCTL from real-time context (optional).*
- [rtdm\\_ioctl\\_handler\\_t ioctl\\_nrt](#)  
*IOCTL from non-real-time context (optional).*
- [rtdm\\_select\\_bind\\_handler\\_t select\\_bind](#)  
*Select binding handler for any context (optional).*

### Stream-Oriented Device Operations

- [rt dm\\_read\\_handler\\_t read\\_rt](#)  
*Read handler for real-time context (optional).*
- [rt dm\\_read\\_handler\\_t read\\_nrt](#)  
*Read handler for non-real-time context (optional).*
- [rt dm\\_write\\_handler\\_t write\\_rt](#)  
*Write handler for real-time context (optional).*
- [rt dm\\_write\\_handler\\_t write\\_nrt](#)  
*Write handler for non-real-time context (optional).*

### Message-Oriented Device Operations

- [rt dm\\_recvmsg\\_handler\\_t recvmsg\\_rt](#)  
*Receive message handler for real-time context (optional).*
- [rt dm\\_recvmsg\\_handler\\_t recvmsg\\_nrt](#)  
*Receive message handler for non-real-time context (optional).*
- [rt dm\\_sendmsg\\_handler\\_t sendmsg\\_rt](#)  
*Transmit message handler for real-time context (optional).*
- [rt dm\\_sendmsg\\_handler\\_t sendmsg\\_nrt](#)  
*Transmit message handler for non-real-time context (optional).*

## 6.9.1 Detailed Description

Device operations.

## 6.9.2 Field Documentation

### 6.9.2.1 [rt dm\\_close\\_handler\\_t](#) [rt dm\\_operations::close\\_rt](#)

Close handler for real-time contexts (optional, deprecated).

#### Deprecated

Only use non-real-time close handler in new drivers.

Referenced by [rt dm\\_dev\\_register\(\)](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/rtdm\\_driver.h](#)

## 6.10 [rt ipc\\_port\\_label](#) Struct Reference

Port label information structure.

## Data Fields

- char [label](#) [XNOBJECT\_NAME\_LEN]  
*Port label string, null-terminated.*

### 6.10.1 Detailed Description

Port label information structure.

### 6.10.2 Field Documentation

#### 6.10.2.1 char rtipc\_port\_label::label[XNOBJECT\_NAME\_LEN]

Port label string, null-terminated.

The documentation for this struct was generated from the following file:

- include/rtdm/[rtipc.h](#)

## 6.11 rtser\_config Struct Reference

Serial device configuration.

## Data Fields

- int [config\\_mask](#)  
*mask specifying valid fields, see [RTSER\\_SET\\_xxx](#)*
- int [baud\\_rate](#)  
*baud rate, default [RTSER\\_DEF\\_BAUD](#)*
- int [parity](#)  
*number of parity bits, see [RTSER\\_xxx\\_PARITY](#)*
- int [data\\_bits](#)  
*number of data bits, see [RTSER\\_xxx\\_BITS](#)*
- int [stop\\_bits](#)  
*number of stop bits, see [RTSER\\_xxx\\_STOPB](#)*
- int [handshake](#)  
*handshake mechanisms, see [RTSER\\_xxx\\_HAND](#)*
- int [fifo\\_depth](#)  
*reception FIFO interrupt threshold, see [RTSER\\_FIFO\\_xxx](#)*
- [nanosecs\\_rel\\_t rx\\_timeout](#)

*reception timeout, see [RTSER\\_TIMEOUT\\_xxx](#) for special values*

- [nanosecs\\_rel\\_t tx\\_timeout](#)  
*transmission timeout, see [RTSER\\_TIMEOUT\\_xxx](#) for special values*
- [nanosecs\\_rel\\_t event\\_timeout](#)  
*event timeout, see [RTSER\\_TIMEOUT\\_xxx](#) for special values*
- [int timestamp\\_history](#)  
*enable timestamp history, see [RTSER\\_xxx\\_TIMESTAMP\\_HISTORY](#)*
- [int event\\_mask](#)  
*event mask to be used with [RTSER\\_RTIOC\\_WAIT\\_EVENT](#), see [RTSER\\_EVENT\\_xxx](#)*

### 6.11.1 Detailed Description

Serial device configuration.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtserial.h](#)

## 6.12 rtser\_event Struct Reference

Additional information about serial device events.

### Data Fields

- [int events](#)  
*signalled events, see [RTSER\\_EVENT\\_xxx](#)*
- [int rx\\_pending](#)  
*number of pending input characters*
- [nanosecs\\_abs\\_t last\\_timestamp](#)  
*last interrupt timestamp*
- [nanosecs\\_abs\\_t rxpend\\_timestamp](#)  
*reception timestamp of oldest character in input queue*

### 6.12.1 Detailed Description

Additional information about serial device events.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtserial.h](#)



## 6.13 rtser\_status Struct Reference

Serial device status.

### Data Fields

- int [line\\_status](#)  
*line status register, see [RTSER\\_LSR\\_xxx](#)*
- int [modem\\_status](#)  
*modem status register, see [RTSER\\_MSR\\_xxx](#)*

### 6.13.1 Detailed Description

Serial device status.

The documentation for this struct was generated from the following file:

- [include/rtdm/rtserial.h](#)

## 6.14 sockaddr\_can Struct Reference

Socket address structure for the CAN address family.

### Data Fields

- sa\_family\_t [can\\_family](#)  
*CAN address family, must be [AF\\_CAN](#).*
- int [can\\_ifindex](#)  
*Interface index of CAN controller.*

### 6.14.1 Detailed Description

Socket address structure for the CAN address family.

### 6.14.2 Field Documentation

#### 6.14.2.1 int sockaddr\_can::can\_ifindex

Interface index of CAN controller.

See [SIOCGIFINDEX](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/rtcan.h](#)

## 6.15 sockaddr\_ipc Struct Reference

Socket address structure for the RTIPC address family.

### Data Fields

- `sa_family_t sipc_family`  
*RTIPC address family, must be AF\_RTIPC.*
- `rtipc_port_t sipc_port`  
*Port number.*

### 6.15.1 Detailed Description

Socket address structure for the RTIPC address family.

### 6.15.2 Field Documentation

#### 6.15.2.1 `rtipc_port_t sockaddr_ipc::sipc_port`

Port number.

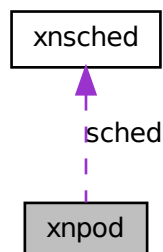
The documentation for this struct was generated from the following file:

- `include/rtdm/rtipc.h`

## 6.16 xnpod Struct Reference

Real-time pod descriptor.

Collaboration diagram for xnpod:



## Data Fields

- `xnflags_t` [status](#)
- `xnsched_t` [sched](#) [XNARCH\_NR\_CPUS]
- `xnqueue_t` [threadq](#)
- `xnqueue_t` [tstartq](#)
- `xnqueue_t` [tswitchq](#)
- `xnqueue_t` [tdeleteq](#)
- `atomic_counter_t` [timerlck](#)
- `int` [refcnt](#)

### 6.16.1 Detailed Description

Real-time pod descriptor. The source of all Xenomai magic.

### 6.16.2 Field Documentation

#### 6.16.2.1 `int xnpod::refcnt`

Reference count.

Referenced by `xnpod_init()`.

#### 6.16.2.2 `xnsched_t xnpod::sched[XNARCH_NR_CPUS]`

Per-cpu scheduler slots.

Referenced by `xnpod_init()`.

#### 6.16.2.3 `xnflags_t xnpod::status`

Status bitmask.

Referenced by `xnpod_init()`.

#### 6.16.2.4 `xnqueue_t xnpod::tdeleteq`

Thread delete hook queue.

Referenced by `xnpod_init()`.

#### 6.16.2.5 `xnqueue_t xnpod::threadq`

All existing threads.

Referenced by `xnpod_init()`.

#### 6.16.2.6 `atomic_counter_t xnpod::timerlck`

Timer lock depth.

Referenced by `xnpod_init()`.

#### 6.16.2.7 `xnqueue_t xnpod::tstartq`

Thread start hook queue.

Referenced by `xnpod_init()`.

#### 6.16.2.8 `xnqueue_t xnpod::tswitchq`

Thread switch hook queue.

Referenced by `xnpod_init()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/pod.h`

## 6.17 `xnsched` Struct Reference

Scheduling information structure.

### Data Fields

- `xnflags_t status`
- `xnflags_t lflags`
- `struct xnthread * curr`
- `struct xnsched_rt rt`
- `volatile unsigned inesting`
- `struct xntimer htimer`
- `struct xnthread rootcb`

#### 6.17.1 Detailed Description

Scheduling information structure.

#### 6.17.2 Field Documentation

##### 6.17.2.1 `struct xnthread* xnsched::curr`

Current thread.

Referenced by `xnpod_delete_thread()`, and `xnpod_suspend_thread()`.

#### 6.17.2.2 struct xntimer xnsched::htimer

Host timer.

Referenced by xnpod\_enable\_timesource(), and xntimer\_tick().

#### 6.17.2.3 volatile unsigned xnsched::inesting

Interrupt nesting level.

#### 6.17.2.4 xnflags\_t xnsched::lflags

Scheduler specific local flags bitmask.

Referenced by xnpod\_schedule(), and xntimer\_tick().

#### 6.17.2.5 struct xntthread xnsched::rootcb

Root thread control block.

Referenced by xnpod\_init().

#### 6.17.2.6 struct xnsched\_rt xnsched::rt

Context of built-in real-time class.

#### 6.17.2.7 xnflags\_t xnsched::status

Scheduler specific status bitmask.

Referenced by xnpod\_delete\_thread(), xnpod\_schedule(), xntimer\_freeze(), and xntimer\_tick().

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/[sched.h](#)

## 6.18 xntthread\_info Struct Reference

Structure containing thread information.

### Data Fields

- unsigned long [state](#)  
*Thread state.*
- int [bprio](#)  
*Base priority.*
- int [cprio](#)

*Current priority.*

- int `cpu`  
*CPU the thread currently runs on.*
- unsigned long `affinity`  
*Thread's CPU affinity.*
- unsigned long long `relpoint`  
*Time of next release.*
- unsigned long long `exectime`  
*Execution time in primary mode in nanoseconds.*
- unsigned long `modeswitches`  
*Number of primary->secondary mode switches.*
- unsigned long `ctxswitches`  
*Number of context switches.*
- unsigned long `pagefaults`  
*Number of triggered page faults.*
- unsigned long `syscalls`  
*Number of Xenomai syscalls.*
- char `name` [XNOBJECT\_NAME\_LEN]  
*Symbolic name assigned at creation.*

### 6.18.1 Detailed Description

Structure containing thread information.

### 6.18.2 Field Documentation

#### 6.18.2.1 unsigned long `xnthread_info::affinity`

Thread's CPU affinity.

#### 6.18.2.2 int `xnthread_info::bprio`

Base priority.

#### 6.18.2.3 int `xnthread_info::cprio`

Current priority.

May change through Priority Inheritance.

**6.18.2.4 int xnthread\_info::cpu**

CPU the thread currently runs on.

**6.18.2.5 unsigned long xnthread\_info::ctxswitches**

Number of context switches.

**6.18.2.6 unsigned long long xnthread\_info::exectime**

Execution time in primary mode in nanoseconds.

**6.18.2.7 unsigned long xnthread\_info::modeswitches**

Number of primary->secondary mode switches.

**6.18.2.8 char xnthread\_info::name[XNOBJECT\_NAME\_LEN]**

Symbolic name assigned at creation.

**6.18.2.9 unsigned long xnthread\_info::pagefaults**

Number of triggered page faults.

**6.18.2.10 unsigned long long xnthread\_info::relpoint**

Time of next release.

**6.18.2.11 unsigned long xnthread\_info::state**

Thread state,.

**See also**

[Thread state flags.](#)

**6.18.2.12 unsigned long xnthread\_info::syscalls**

Number of Xenomai syscalls.

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/thread.h

## 6.19 xnvfile\_lock\_ops Struct Reference

Vfile locking operations.

## Data Fields

- `int(* get)(struct xnvfile *vfile)`
- `void(* put)(struct xnvfile *vfile)`

### 6.19.1 Detailed Description

Vfile locking operations.

This structure describes the operations to be provided for implementing locking support on vfiles. They apply to both snapshot-driven and regular vfiles.

### 6.19.2 Field Documentation

#### 6.19.2.1 `int(* xnvfile_lock_ops::get)(struct xnvfile *vfile)`

This handler should grab the desired lock.

##### Parameters

*vfile* A pointer to the virtual file which needs locking.

##### Returns

zero should be returned if the call succeeds. Otherwise, a negative error code can be returned; upon error, the current vfile operation is aborted, and the user-space caller is passed back the error value.

#### 6.19.2.2 `void(* xnvfile_lock_ops::put)(struct xnvfile *vfile)`

This handler should release the lock previously grabbed by the [get\(\)](#) handler.

##### Parameters

*vfile* A pointer to the virtual file which currently holds the lock to release.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

## 6.20 `xnvfile_regular_iterator` Struct Reference

Regular vfile iterator.

## Data Fields

- `loff_t pos`  
*Current record position while iterating.*



- struct seq\_file \* [seq](#)  
*Backlink to the host sequential file supporting the vfile.*
- struct xnvfile\_regular \* [vfile](#)  
*Backlink to the vfile being read.*
- char [private](#) [0]  
*Start of private area.*

### 6.20.1 Detailed Description

Regular vfile iterator.

This structure defines an iterator over a regular vfile.

### 6.20.2 Field Documentation

#### 6.20.2.1 loff\_t xnvfile\_regular\_iterator::pos

Current record position while iterating.

#### 6.20.2.2 char xnvfile\_regular\_iterator::private[0]

Start of private area.

Use xnvfile\_iterator\_priv() to address it.

#### 6.20.2.3 struct seq\_file\* xnvfile\_regular\_iterator::seq

Backlink to the host sequential file supporting the vfile.

#### 6.20.2.4 struct xnvfile\_regular\* xnvfile\_regular\_iterator::vfile

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/[vfile.h](#)

## 6.21 xnvfile\_regular\_ops Struct Reference

Regular vfile operation descriptor.

### Data Fields

- int(\* [rewind](#) )(struct [xnvfile\\_regular\\_iterator](#) \*it)
- void \*(\* [begin](#) )(struct [xnvfile\\_regular\\_iterator](#) \*it)

- void *(\* next)* (struct [xnvmfile\\_regular\\_iterator](#) \*it)
- void *(\* end)* (struct [xnvmfile\\_regular\\_iterator](#) \*it)
- int *(\* show)* (struct [xnvmfile\\_regular\\_iterator](#) \*it, void \*data)
- ssize\_t *(\* store)* (struct [xnvmfile\\_input](#) \*input)

### 6.21.1 Detailed Description

Regular vfile operation descriptor.

This structure describes the operations available with a regular vfile. It defines handlers for sending back formatted kernel data upon a user-space read request, and for obtaining user data upon a user-space write request.

### 6.21.2 Field Documentation

#### 6.21.2.1 void *(\* xnvmfile\_regular\_ops::begin)* (struct [xnvmfile\\_regular\\_iterator](#) \*it)

This handler should prepare for iterating over the records upon a read request, starting from the specified position.

##### Parameters

- it* A pointer to the current vfile iterator. On entry, *it->pos* is set to the (0-based) position of the first record to output. This handler may be called multiple times with different position requests.

##### Returns

A pointer to the first record to format and output, to be passed to the [show\(\)](#) handler as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- VFILE\_SEQ\_START, a special value indicating that [the show\(\)](#) handler should receive a NULL data pointer first, in order to output a header.
- ERR\_PTR(errno), where errno is a negative error code; upon error, the current operation will be aborted immediately.

##### Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the [show\(\)](#) handler will be called only once for a read operation, with a NULL *data* parameter. This particular setting is convenient for simple regular vfiles having a single, fixed record to output.

#### 6.21.2.2 void *(\* xnvmfile\_regular\_ops::end)* (struct [xnvmfile\\_regular\\_iterator](#) \*it)

This handler is called after all records have been output.

**Parameters**

*it* A pointer to the current vfile iterator.

**Note**

This handler is optional and the pointer may be NULL.

**6.21.2.3 void\*(\* xnvfile\_regular\_ops::next)(struct xnvfile\_regular\_iterator \*it)**

This handler should return the address of the next record to format and output by the [show\(\)](#) handler".

**Parameters**

*it* A pointer to the current vfile iterator. On entry, *it->pos* is set to the (0-based) position of the next record to output.

**Returns**

A pointer to the next record to format and output, to be passed to the [show\(\)](#) handler as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- ERR\_PTR(errno), where *errno* is a negative error code; upon error, the current operation will be aborted immediately.

**Note**

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the read operation will stop after the first invocation of the [show\(\)](#) handler.

**6.21.2.4 int(\* xnvfile\_regular\_ops::rewind)(struct xnvfile\_regular\_iterator \*it)**

This handler is called only once, when the virtual file is opened, before the [begin\(\)](#) handler is invoked.

**Parameters**

*it* A pointer to the vfile iterator which will be used to read the file contents.

**Returns**

Zero should be returned upon success. Otherwise, a negative error code aborts the operation, and is passed back to the reader.

**Note**

This handler is optional. It should not be used to allocate resources but rather to perform consistency checks, since no closure call is issued in case the open sequence eventually fails.

### 6.21.2.5 `int(* xnvfile_regular_ops::show)(struct xnvfile_iterator *it, void *data)`

This handler should format and output a record.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

#### Parameters

*it* A pointer to the current vfile iterator.

*data* A pointer to the record to format then output. The first call to the handler may receive a NULL *data* pointer, depending on the presence and/or return of a [handler](#); the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

#### Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

### 6.21.2.6 `ssize_t(* xnvfile_regular_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating some kernel setting, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, `xnvfile_get_string()` will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, `xnvfile_get_integer()` will attempt to return a long integer from the input data.

#### Parameters

*input* A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

#### Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the `store()` handler.

#### Note

This handler is optional, and may be omitted for read-only vfiles.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

## 6.22 xnvfile\_rev\_tag Struct Reference

Snapshot revision tag.

### Data Fields

- int [rev](#)  
*Current revision number.*

### 6.22.1 Detailed Description

Snapshot revision tag.

This structure defines a revision tag to be used with [snapshot-driven vfiles](#).

### 6.22.2 Field Documentation

#### 6.22.2.1 int xnvfile\_rev\_tag::rev

Current revision number.

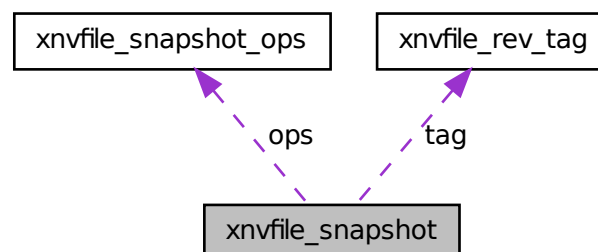
The documentation for this struct was generated from the following file:

- [include/cobalt/nucleus/vfile.h](#)

## 6.23 xnvfile\_snapshot Struct Reference

Snapshot vfile descriptor.

Collaboration diagram for xnvfile\_snapshot:



### 6.23.1 Detailed Description

Snapshot vfile descriptor.

This structure describes a snapshot-driven vfile. Reading from such a vfile involves a preliminary data collection phase under lock protection, and a subsequent formatting and output phase of the collected data records. Locking is done in a way that does not increase worst-case latency, regardless of the number of records to be collected for output.

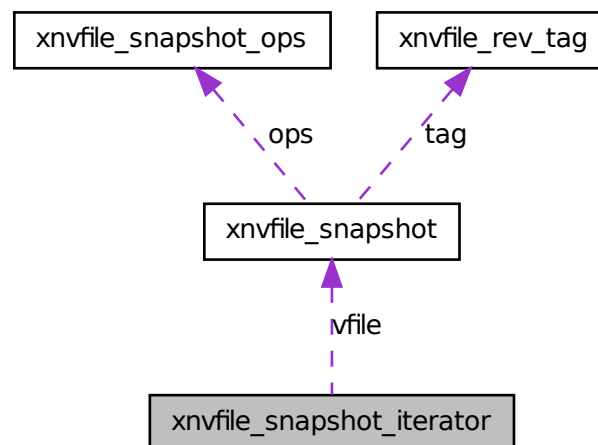
The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

## 6.24 xnvfile\_snapshot\_iterator Struct Reference

Snapshot-driven vfile iterator.

Collaboration diagram for xnvfile\_snapshot\_iterator:



### Data Fields

- `int nrdata`  
*Number of collected records.*
- `caddr_t databuf`  
*Address of record buffer.*
- `struct seq_file * seq`  
*Backlink to the host sequential file supporting the vfile.*

- struct `xnvfile_snapshot * vfile`  
*Backlink to the vfile being read.*
- void(\* `endfn`)(struct `xnvfile_snapshot_iterator *it`, void \*buf)  
*Buffer release handler.*
- char `private` [0]  
*Start of private area.*

### 6.24.1 Detailed Description

Snapshot-driven vfile iterator.

This structure defines an iterator over a snapshot-driven vfile.

### 6.24.2 Field Documentation

#### 6.24.2.1 `caddr_t xnvfile_snapshot_iterator::databuf`

Address of record buffer.

#### 6.24.2.2 `void(* xnvfile_snapshot_iterator::endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`

Buffer release handler.

#### 6.24.2.3 `int xnvfile_snapshot_iterator::nrdata`

Number of collected records.

#### 6.24.2.4 `char xnvfile_snapshot_iterator::private[0]`

Start of private area.

Use `xnvfile_iterator_priv()` to address it.

#### 6.24.2.5 `struct seq_file* xnvfile_snapshot_iterator::seq`

Backlink to the host sequential file supporting the vfile.

#### 6.24.2.6 `struct xnvfile_snapshot* xnvfile_snapshot_iterator::vfile`

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

## 6.25 xnvfile\_snapshot\_ops Struct Reference

Snapshot vfile operation descriptor.

### Data Fields

- `int(* rewind )(struct xnvfile\_snapshot\_iterator *it)`
- `void *(* begin )(struct xnvfile\_snapshot\_iterator *it)`
- `void(* end )(struct xnvfile\_snapshot\_iterator *it, void *buf)`
- `int(* next )(struct xnvfile\_snapshot\_iterator *it, void *data)`
- `int(* show )(struct xnvfile\_snapshot\_iterator *it, void *data)`
- `ssize_t(* store )(struct xnvfile\_input *input)`

### 6.25.1 Detailed Description

Snapshot vfile operation descriptor.

This structure describes the operations available with a snapshot-driven vfile. It defines handlers for returning a printable snapshot of some Xenomai object contents upon a user-space read request, and for updating this object upon a user-space write request.

### 6.25.2 Field Documentation

#### 6.25.2.1 `void *(* xnvfile\_snapshot\_ops::begin )(struct xnvfile\_snapshot\_iterator *it)`

This handler should allocate the snapshot buffer to hold records during the data collection phase. When specified, all records collected via the [next\(\)](#) handler" will be written to a cell from the memory area returned by [begin\(\)](#).

#### Parameters

*it* A pointer to the current snapshot iterator.

#### Returns

A pointer to the record buffer, if the call succeeds. Otherwise:

- NULL in case of allocation error. This will abort the data collection, and return -ENOMEM to the reader.
- VFILE\_SEQ\_EMPTY, a special value indicating that no record will be output. In such a case, the [next\(\) handler](#) will not be called, and the data collection will stop immediately. However, the [show\(\) handler](#) will still be called once, with a NULL data pointer (i.e. header display request).

#### Note

This handler is optional; if none is given, an internal allocation depending on the value returned by the [rewind\(\) handler](#) can be obtained.



### 6.25.2.2 void(\* xnvfile\_snapshot\_ops::end)(struct xnvfile\_snapshot\_iterator \*it, void \*buf)

This handler releases the memory buffer previously obtained from [begin\(\)](#). It is usually called after the snapshot data has been output by [show\(\)](#), but it may also be called before rewinding the vfile after a revision change, to release the dropped buffer.

#### Parameters

- it* A pointer to the current snapshot iterator.
- buf* A pointer to the buffer to release.

#### Note

This routine is optional and the pointer may be NULL. It is not needed upon internal buffer allocation; see the description of the [rewind\(\)](#) handler".

### 6.25.2.3 int(\* xnvfile\_snapshot\_ops::next)(struct xnvfile\_snapshot\_iterator \*it, void \*data)

This handler fetches the next record, as part of the snapshot data to be sent back to the reader via the [show\(\)](#).

#### Parameters

- it* A pointer to the current snapshot iterator.
- data* A pointer to the record to fill in.

#### Returns

a strictly positive value, if the call succeeds and leaves a valid record into *data*, which should be passed to the [show\(\) handler\(\)](#) during the formatting and output phase. Otherwise:

- A negative error code. This will abort the data collection, and return this status to the reader.
- VFILE\_SEQ\_SKIP, a special value indicating that the current record should be skipped. In such a case, the *data* pointer is not advanced to the next position before the [next\(\) handler](#) is called anew.

#### Note

This handler is called with the vfile lock held. Before each invocation of this handler, the vfile core checks whether the revision tag has been touched, in which case the data collection is restarted from scratch. A data collection phase succeeds whenever all records can be fetched via the [next\(\) handler](#), while the revision tag remains unchanged, which indicates that a consistent snapshot of the object state was taken.

### 6.25.2.4 int(\* xnvfile\_snapshot\_ops::rewind)(struct xnvfile\_snapshot\_iterator \*it)

This handler (re-)initializes the data collection, moving the seek pointer at the first record. When the file revision tag is touched while collecting data, the current reading is aborted, all collected data dropped, and the vfile is eventually rewound.

### Parameters

*it* A pointer to the current snapshot iterator. Two useful information can be retrieved from this iterator in this context:

- *it->vfile* is a pointer to the descriptor of the virtual file being rewound.
- `xnvfile_iterator_priv(it)` returns a pointer to the private data area, available from the descriptor, which size is *vfile->privsz*. If the latter size is zero, the returned pointer is meaningless and should not be used.

### Returns

A negative error code aborts the data collection, and is passed back to the reader. Otherwise:

- a strictly positive value is interpreted as the total number of records which will be returned by the [next\(\) handler](#) during the data collection phase. If no [begin\(\) handler](#) is provided in the [operation descriptor](#), this value is used to allocate the snapshot buffer internally. The size of this buffer would then be *vfile->datasz* \* value.
- zero leaves the allocation to the [begin\(\) handler](#) if present, or indicates that no record is to be output in case such handler is not given.

### Note

This handler is optional; a NULL value indicates that nothing needs to be done for rewinding the *vfile*. It is called with the *vfile* lock held.

#### 6.25.2.5 `int(* xnvfile_snapshot_ops::show)(struct xnvfile_snapshot_iterator *it, void *data)`

This handler should format and output a record from the collected data.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

### Parameters

*it* A pointer to the current snapshot iterator.

*data* A pointer to the record to format then output. The first call to the handler is always passed a NULL *data* pointer; the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

### Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

#### 6.25.2.6 ssize\_t(\* xnvfile\_snapshot\_ops::store)(struct xnvfile\_input \*input)

This handler receives data written to the vfile, likely for updating the associated Xenomai object's state, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile\\_get\\_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile\\_get\\_integer\(\)](#) will attempt to return a long integer from the input data.

##### Parameters

*input* A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

##### Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

##### Note

This handler is optional, and may be omitted for read-only vfiles.

Referenced by [xnvfile\\_init\\_snapshot\(\)](#).

The documentation for this struct was generated from the following file:

- [include/cobalt/nucleus/vfile.h](#)

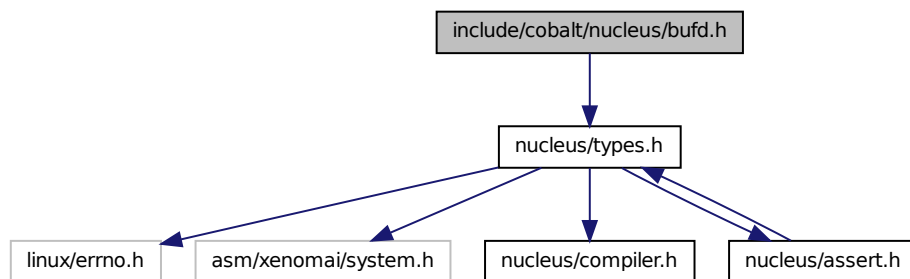


## Chapter 7

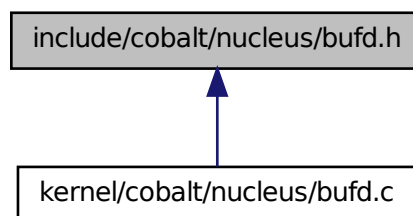
# File Documentation

### 7.1 include/cobalt/nucleus/bufd.h File Reference

Include dependency graph for bufd.h:



This graph shows which files directly or indirectly include this file:



## Functions

- static void [xnbufd\\_map\\_uread](#) (struct xnbufd \*bufd, const void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from user memory.*
- static void [xnbufd\\_map\\_uwrite](#) (struct xnbufd \*bufd, void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to user memory.*
- ssize\_t [xnbufd\\_unmap\\_uread](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_uread\(\)](#).*
- ssize\_t [xnbufd\\_unmap\\_uwrite](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_uwrite\(\)](#).*
- static void [xnbufd\\_map\\_kread](#) (struct xnbufd \*bufd, const void \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from kernel memory.*
- static void [xnbufd\\_map\\_kwrite](#) (struct xnbufd \*bufd, void \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to kernel memory.*
- ssize\_t [xnbufd\\_unmap\\_kread](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_kread\(\)](#).*
- ssize\_t [xnbufd\\_unmap\\_kwrite](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_kwrite\(\)](#).*
- ssize\_t [xnbufd\\_copy\\_to\\_kmem](#) (void \*ptr, struct xnbufd \*bufd, size\_t len)  
*Copy memory covered by a buffer descriptor to kernel memory.*
- ssize\_t [xnbufd\\_copy\\_from\\_kmem](#) (struct xnbufd \*bufd, void \*from, size\_t len)  
*Copy kernel memory to the area covered by a buffer descriptor.*
- void [xnbufd\\_invalidate](#) (struct xnbufd \*bufd)  
*Invalidate a buffer descriptor.*
- static void [xnbufd\\_reset](#) (struct xnbufd \*bufd)  
*Reset a buffer descriptor.*

### 7.1.1 Detailed Description

#### Note

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

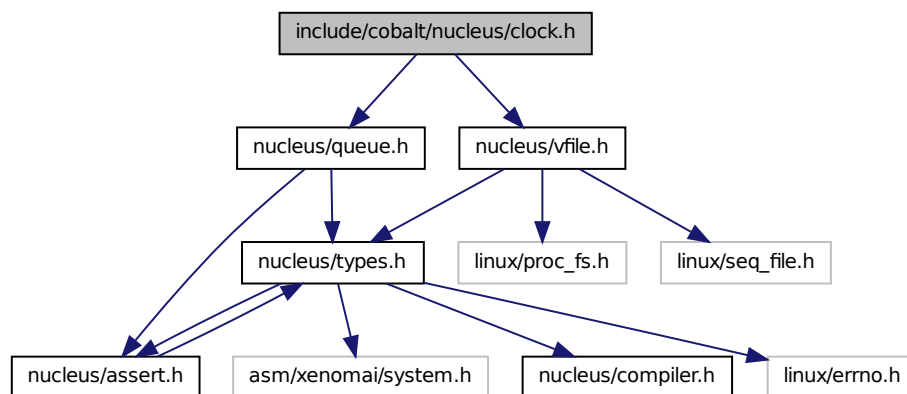
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

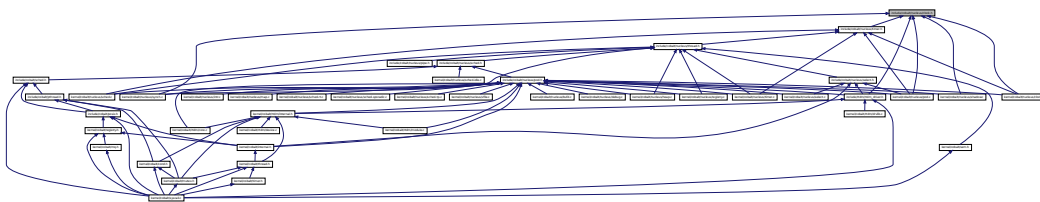
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.2 include/cobalt/nucleus/clock.h File Reference

Include dependency graph for clock.h:



This graph shows which files directly or indirectly include this file:



### Functions

- void [xnclock\\_adjust](#) (xnsticks\_t delta)

*Adjust the clock time for the system.*

### 7.2.1 Detailed Description

#### Note

Copyright (C) 2006,2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

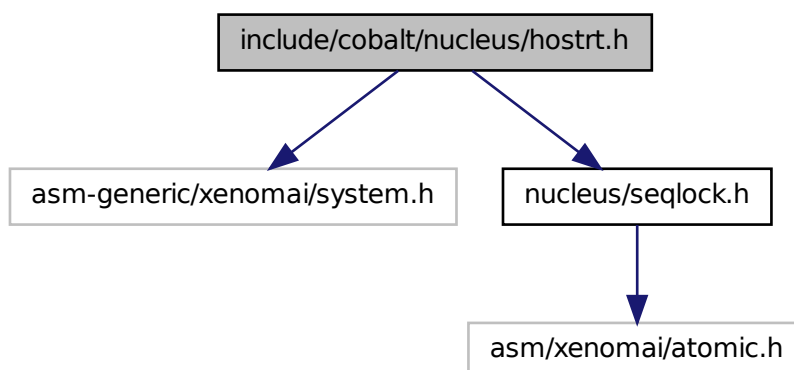
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

### 7.3 include/cobalt/nucleus/hostrt.h File Reference

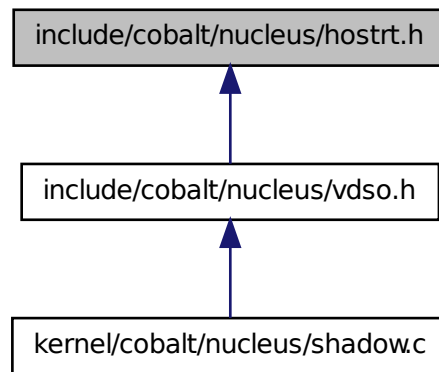
Definitions for global semaphore heap shared objects.

Include dependency graph for hostrt.h:





This graph shows which files directly or indirectly include this file:



### 7.3.1 Detailed Description

Definitions for global semaphore heap shared objects.

#### Author

Wolfgang Mauerer

Copyright (C) 2010 Wolfgang Mauerer <[wolfgang.mauerer@siemens.com](mailto:wolfgang.mauerer@siemens.com)>.

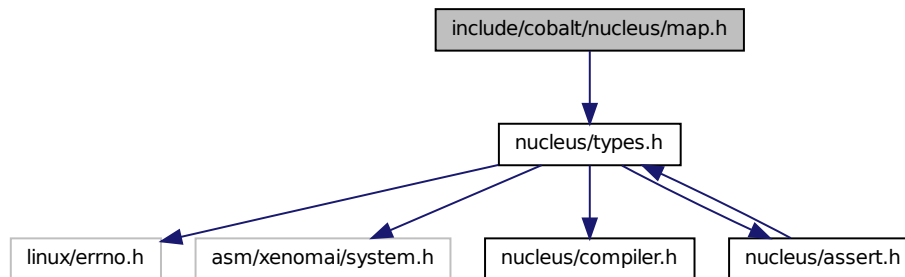
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

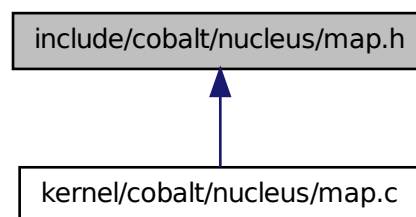
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.4 include/cobalt/nucleus/map.h File Reference

Include dependency graph for map.h:



This graph shows which files directly or indirectly include this file:



### Functions

- `xnmap_t * xnmap\_create` (`int nkeys`, `int reserve`, `int offset`)  
*Create a map.*
- `void xnmap\_delete` (`xnmap_t *map`)  
*Delete a map.*
- `int xnmap\_enter` (`xnmap_t *map`, `int key`, `void *objaddr`)  
*Index an object into a map.*
- `int xnmap\_remove` (`xnmap_t *map`, `int key`)  
*Remove an object reference from a map.*

- static void \* [xnmap\\_fetch\\_nocheck](#) (xnmap\_t \*map, int key)

*Search an object into a map - unchecked form.*

- static void \* [xnmap\\_fetch](#) (xnmap\_t \*map, int key)

*Search an object into a map.*

### 7.4.1 Detailed Description

#### Note

Copyright (C) 2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

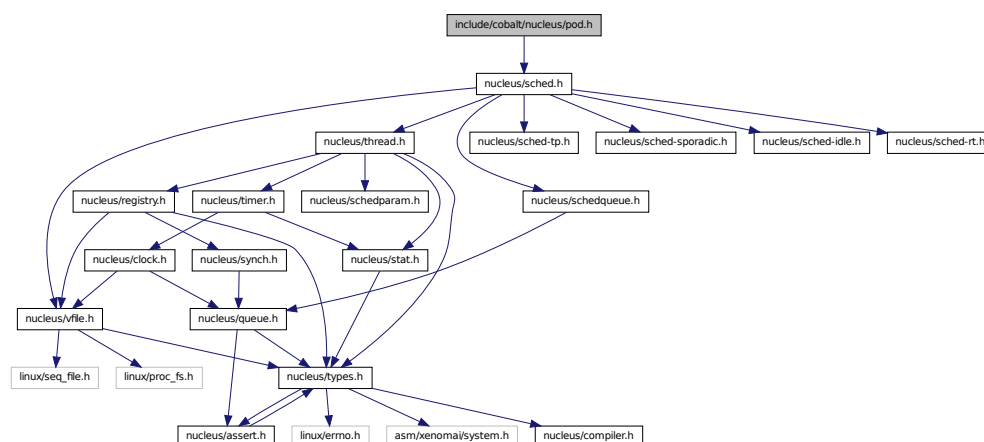
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

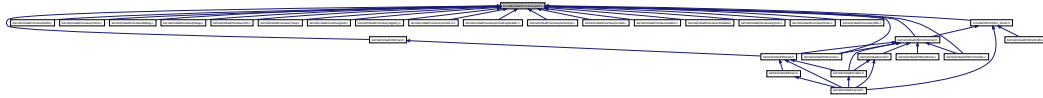
## 7.5 include/cobalt/nucleus/pod.h File Reference

Real-time pod interface header.

Include dependency graph for pod.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [xnpod](#)  
*Real-time pod descriptor.*

## Functions

- void [\\_\\_xnpod\\_reset\\_thread](#) (struct xnthread \*thread)  
*Reset the thread.*
- int [xnpod\\_init](#) (void)  
*Initialize the core pod.*
- int [xnpod\\_enable\\_timesource](#) (void)  
*Activate the core time source.*
- void [xnpod\\_disable\\_timesource](#) (void)  
*Stop the core time source.*
- void [xnpod\\_shutdown](#) (int xtype)  
*Shutdown the current pod.*
- int [xnpod\\_init\\_thread](#) (struct xnthread \*thread, const struct xnthread\_init\_attr \*attr, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Initialize a new thread.*
- int [xnpod\\_start\\_thread](#) (xnthread\_t \*thread, const struct xnthread\_start\_attr \*attr)  
*Initial start of a newly created thread.*
- void [xnpod\\_stop\\_thread](#) (xnthread\_t \*thread)  
*Stop a thread.*
- void [xnpod\\_delete\\_thread](#) (xnthread\_t \*thread)  
*Delete a thread.*
- void [xnpod\\_abort\\_thread](#) (xnthread\_t \*thread)  
*Abort a thread.*
- xnflags\_t [xnpod\\_set\\_thread\\_mode](#) (xnthread\_t \*thread, xnflags\_t clrmask, xnflags\_t set-mask)  
*Change a thread's control mode.*

- void [xn timer suspend thread](#) (xn timer thread\_t \*thread, xn timer flags\_t mask, xn timer ticks\_t timeout, xn timer mode\_t timeout\_mode, struct xn timer synch \*wchan)  
*Suspend a thread.*
- void [xn timer resume thread](#) (xn timer thread\_t \*thread, xn timer flags\_t mask)  
*Resume a thread.*
- int [xn timer unblock thread](#) (xn timer thread\_t \*thread)  
*Unblock a thread.*
- int [xn timer set thread sched param](#) (struct xn timer thread \*thread, struct xn timer sched\_class \*sched\_class, const union xn timer sched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*
- int [xn timer migrate thread](#) (int cpu)  
*Migrate the current thread.*
- void [xn timer dispatch signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- static void [xn timer schedule](#) (void)  
*Rescheduling procedure entry point.*
- int [xn timer handle exception](#) (struct ipipe\_trap\_data \*d)  
*Exception handler.*
- int [xn timer set thread periodic](#) (xn timer thread\_t \*thread, xn timer ticks\_t idate, xn timer mode\_t timeout\_mode, xn timer ticks\_t period)  
*Make a thread periodic.*
- int [xn timer wait thread period](#) (unsigned long \*overruns\_r)  
*Wait for the next periodic release point.*
- int [xn timer set thread tslice](#) (struct xn timer thread \*thread, xn timer ticks\_t quantum)  
*Set thread time-slicing information.*
- int [xn timer add hook](#) (int type, void(\*routine)(xn timer thread\_t \*))  
*Install a nucleus hook.*
- int [xn timer remove hook](#) (int type, void(\*routine)(xn timer thread\_t \*))  
*Remove a nucleus hook.*

### 7.5.1 Detailed Description

Real-time pod interface header.

#### Author

Philippe Gerum

Copyright (C) 2001-2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2004 The Xenomai project <<http://www.xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

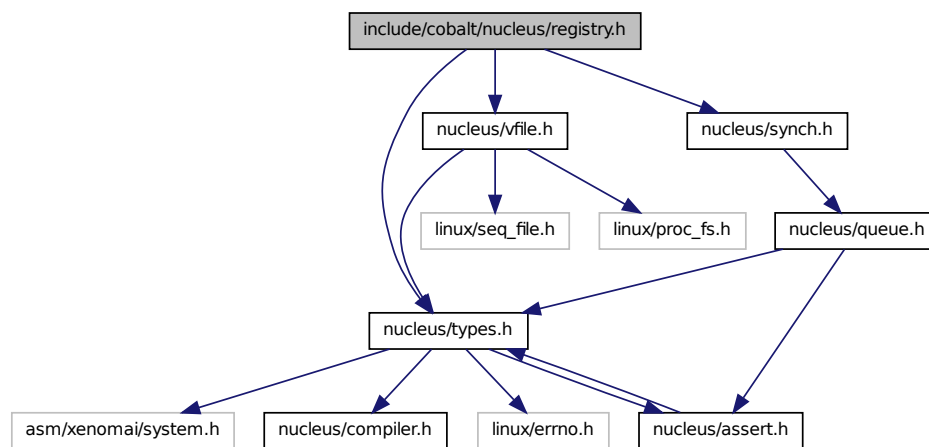
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

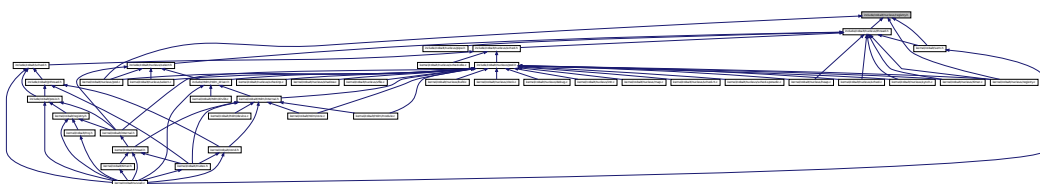
## 7.6 include/cobalt/nucleus/registry.h File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.h:



This graph shows which files directly or indirectly include this file:



## Functions

- int [xnregistry\\_enter](#) (const char \*key, void \*objaddr, xnhandle\_t \*phandle, struct xninode \*pnode)  
*Register a real-time object.*
- int [xnregistry\\_bind](#) (const char \*key, xnticks\_t timeout, int timeout\_mode, xnhandle\_t \*phandle)  
*Bind to a real-time object.*
- int [xnregistry\\_remove](#) (xnhandle\_t handle)  
*Forcibly unregister a real-time object.*
- int [xnregistry\\_remove\\_safe](#) (xnhandle\_t handle, xnticks\_t timeout)  
*Unregister an idle real-time object.*
- void \* [xnregistry\\_get](#) (xnhandle\_t handle)  
*Find and lock a real-time object into the registry.*
- void \* [xnregistry\\_fetch](#) (xnhandle\_t handle)  
*Find a real-time object into the registry.*
- u\_long [xnregistry\\_put](#) (xnhandle\_t handle)  
*Unlock a real-time object from the registry.*

### 7.6.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2004 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

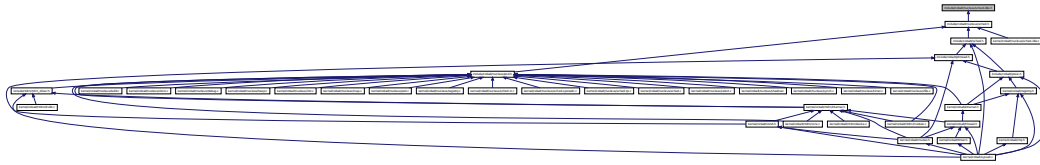
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.7 include/cobalt/nucleus/sched-idle.h File Reference

Definitions for the IDLE scheduling class.

This graph shows which files directly or indirectly include this file:



### 7.7.1 Detailed Description

Definitions for the IDLE scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

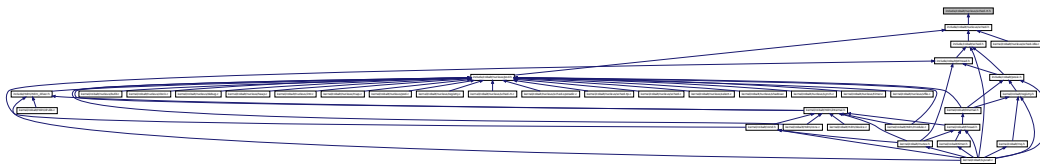
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.8 include/cobalt/nucleus/sched-rt.h File Reference

Definitions for the RT scheduling class.

This graph shows which files directly or indirectly include this file:



### 7.8.1 Detailed Description

Definitions for the RT scheduling class.

#### Author

Philippe Gerum



Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

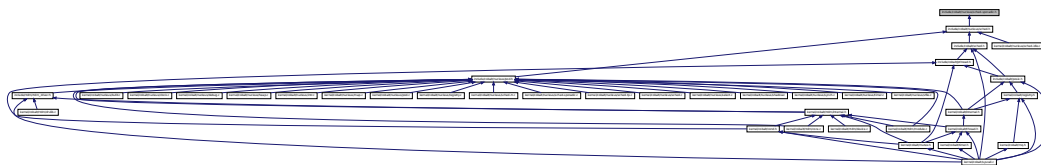
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.9 include/cobalt/nucleus/sched-sporadic.h File Reference

Definitions for the SSP scheduling class.

This graph shows which files directly or indirectly include this file:



### 7.9.1 Detailed Description

Definitions for the SSP scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

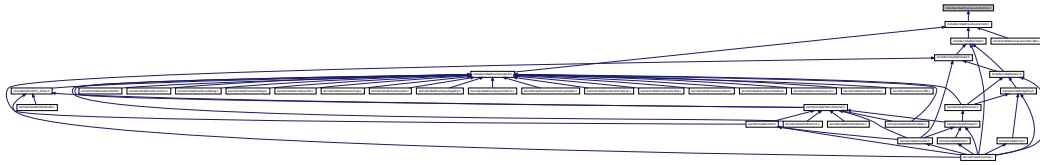
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.10 include/cobalt/nucleus/sched-tp.h File Reference

Definitions for the TP scheduling class.

This graph shows which files directly or indirectly include this file:



### 7.10.1 Detailed Description

Definitions for the TP scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

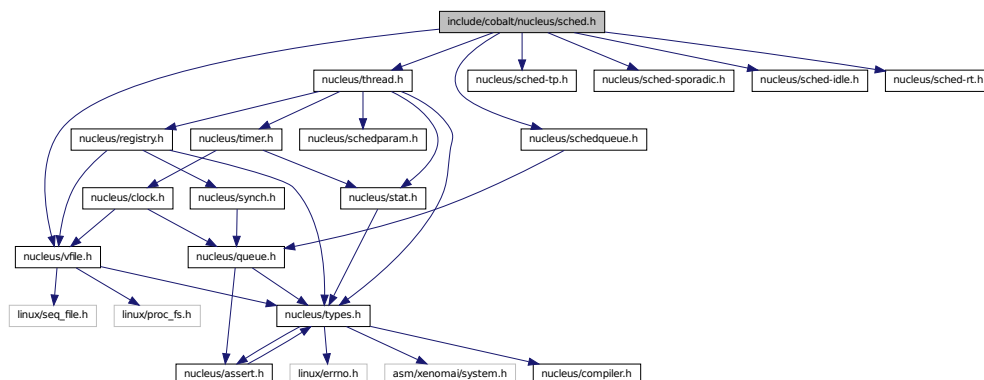
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

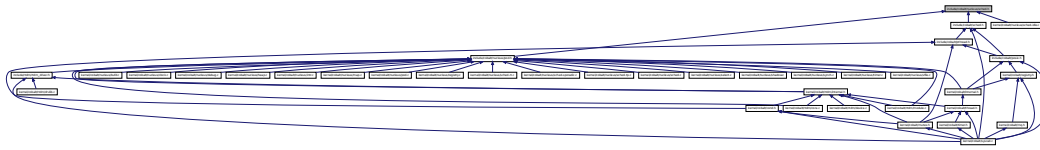
## 7.11 include/cobalt/nucleus/sched.h File Reference

Scheduler interface header.

Include dependency graph for sched.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [xnsched](#)  
*Scheduling information structure.*

## Typedefs

- typedef struct [xnsched](#) [xnsched\\_t](#)  
*Scheduling information structure.*

## Functions

- static void [xnsched\\_rotate](#) (struct [xnsched](#) \*sched, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Rotate a scheduler runqueue.*

### 7.11.1 Detailed Description

Scheduler interface header.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

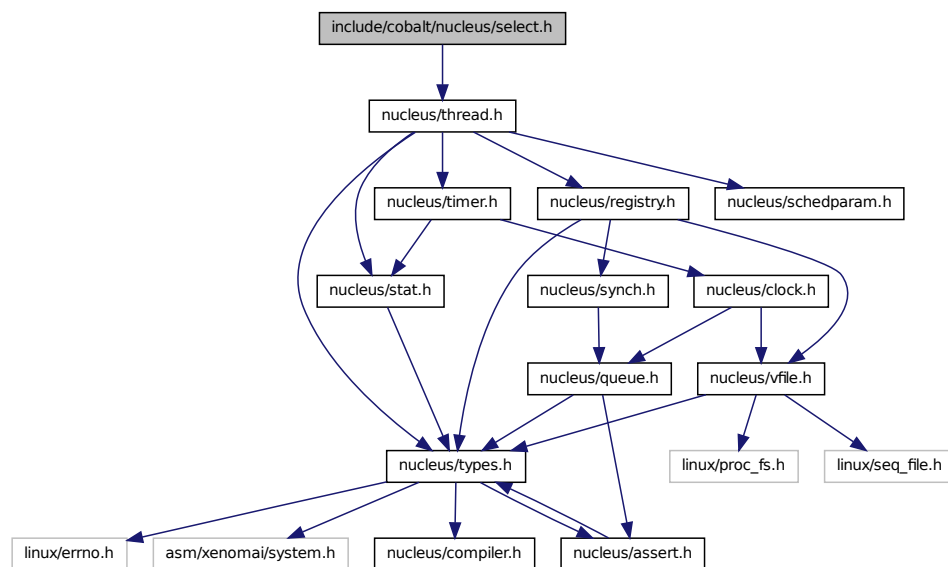
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

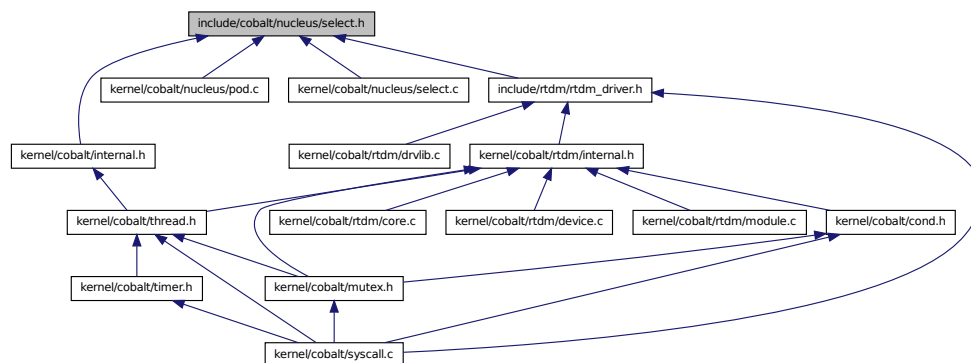
## 7.12 include/cobalt/nucleus/select.h File Reference

file descriptors events multiplexing header.

Include dependency graph for select.h:



This graph shows which files directly or indirectly include this file:



## Functions

- void [xnselect\\_init](#) (struct xnselect \*select\_block)  
*Initialize a struct xnselect structure.*

- int [xnselect\\_bind](#) (struct xnselect \*select\_block, struct xnselect\_binding \*binding, struct xnselector \*selector, unsigned type, unsigned index, unsigned state)

*Bind a file descriptor (represented by its xnselect structure) to a selector block.*

- static int [xnselect\\_signal](#) (struct xnselect \*select\_block, unsigned state)

*Signal a file descriptor state change.*

- void [xnselect\\_destroy](#) (struct xnselect \*select\_block)

*Destroy the xnselect structure associated with a file descriptor.*

- int [xnselector\\_init](#) (struct xnselector \*selector)

*Initialize a selector structure.*

- int [xnselect](#) (struct xnselector \*selector, fd\_set \*out\_fds[XNSELECT\_MAX\_TYPES], fd\_set \*in\_fds[XNSELECT\_MAX\_TYPES], int nfds, xnticks\_t timeout, xntmode\_t timeout\_mode)

*Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.*

- void [xnselector\\_destroy](#) (struct xnselector \*selector)

*Destroy a selector block.*

### 7.12.1 Detailed Description

file descriptors events multiplexing header.

#### Author

Gilles Chanteperdrix

Copyright (C) 2008 Efixo <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>

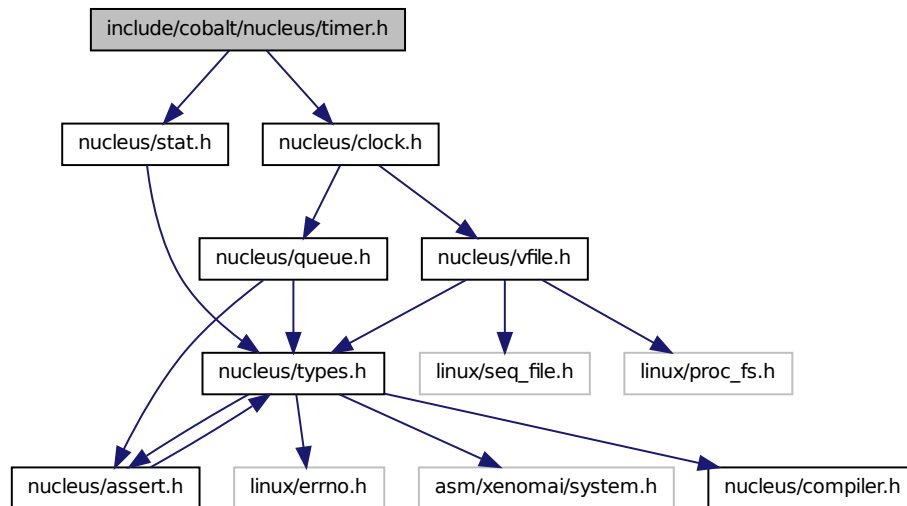
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

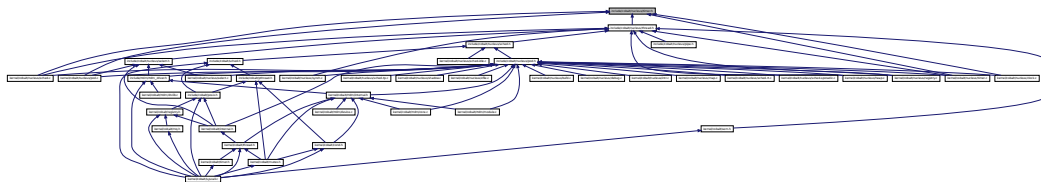
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.13 include/cobalt/nucleus/timer.h File Reference

Include dependency graph for timer.h:



This graph shows which files directly or indirectly include this file:



## Functions

- void [`xntimer\_destroy`](#) (`xntimer_t *timer`)  
*Release a timer object.*
- int [`xntimer\_start`](#) (`xntimer_t *timer`, `xnticks_t value`, `xnticks_t interval`, `xntmode_t mode`)  
*Arm a timer.*
- `xnticks_t` [`xntimer\_get\_date`](#) (`xntimer_t *timer`)  
*Return the absolute expiration date.*
- `xnticks_t` [`xntimer\_get\_timeout`](#) (`xntimer_t *timer`)  
*Return the relative expiration date.*

- `xnticks_t xntimer_get_interval` (`xntimer_t *timer`)  
*Return the timer interval value.*
- `static void xntimer_stop` (`xntimer_t *timer`)  
*Disarm a timer.*
- `unsigned long xntimer_get_overruns` (`xntimer_t *timer`, `xnticks_t now`)  
*Get the count of overruns for the last tick.*
- `void xntimer_freeze` (`void`)  
*Freeze all timers (from every time bases).*
- `void xntimer_tick` (`void`)  
*Process a timer tick.*

### 7.13.1 Detailed Description

#### Note

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

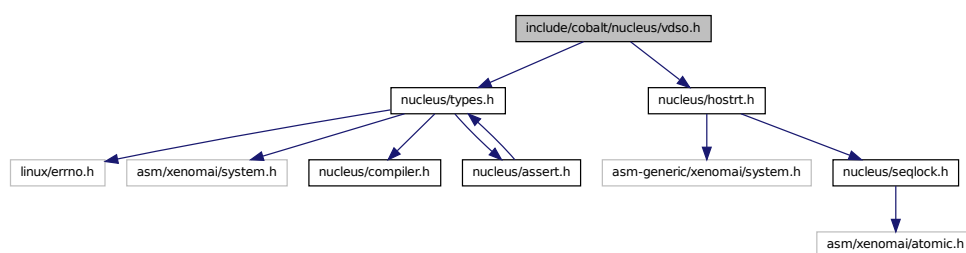
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

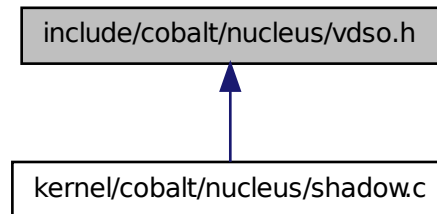
## 7.14 include/cobalt/nucleus/vdso.h File Reference

Definitions for global semaphore heap shared objects.

Include dependency graph for `vdso.h`:



This graph shows which files directly or indirectly include this file:



### 7.14.1 Detailed Description

Definitions for global semaphore heap shared objects.

#### Author

Wolfgang Mauerer

Copyright (C) 2009 Wolfgang Mauerer <[wolfgang.mauerer@siemens.com](mailto:wolfgang.mauerer@siemens.com)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

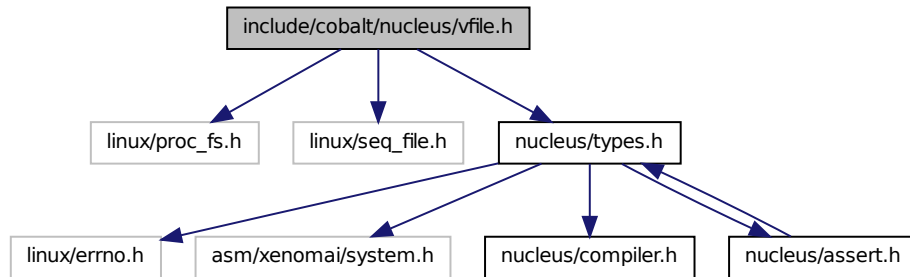
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.15 include/cobalt/nucleus/vfile.h File Reference

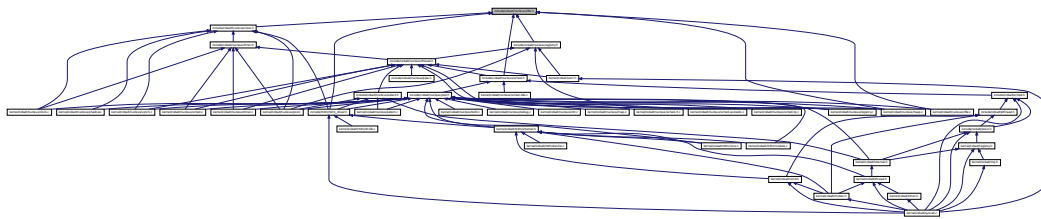
This file is part of the Xenomai project.



Include dependency graph for vfile.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [xnvfile\\_lock\\_ops](#)  
*Vfile locking operations.*
- struct [xnvfile\\_regular\\_ops](#)  
*Regular vfile operation descriptor.*
- struct [xnvfile\\_regular\\_iterator](#)  
*Regular vfile iterator.*
- struct [xnvfile\\_snapshot\\_ops](#)  
*Snapshot vfile operation descriptor.*
- struct [xnvfile\\_rev\\_tag](#)  
*Snapshot revision tag.*
- struct [xnvfile\\_snapshot](#)  
*Snapshot vfile descriptor.*
- struct [xnvfile\\_snapshot\\_iterator](#)  
*Snapshot-driven vfile iterator.*

## Functions

- int [xnvmfile\\_init\\_snapshot](#) (const char \*name, struct [xnvmfile\\_snapshot](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a snapshot-driven vfile.*
- int [xnvmfile\\_init\\_regular](#) (const char \*name, struct [xnvmfile\\_regular](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a regular vfile.*
- int [xnvmfile\\_init\\_dir](#) (const char \*name, struct [xnvmfile\\_directory](#) \*vdir, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual directory entry.*
- int [xnvmfile\\_init\\_link](#) (const char \*from, const char \*to, struct [xnvmfile\\_link](#) \*vlink, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual link entry.*
- void [xnvmfile\\_destroy](#) (struct [xnvmfile](#) \*vfile)  
*Removes a virtual file entry.*
- ssize\_t [xnvmfile\\_get\\_blob](#) (struct [xnvmfile\\_input](#) \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_string](#) (struct [xnvmfile\\_input](#) \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_integer](#) (struct [xnvmfile\\_input](#) \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct [xnvmfile\\_directory](#) [nkvfroot](#)  
*Xenomai vfile root directory.*

### 7.15.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2010 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

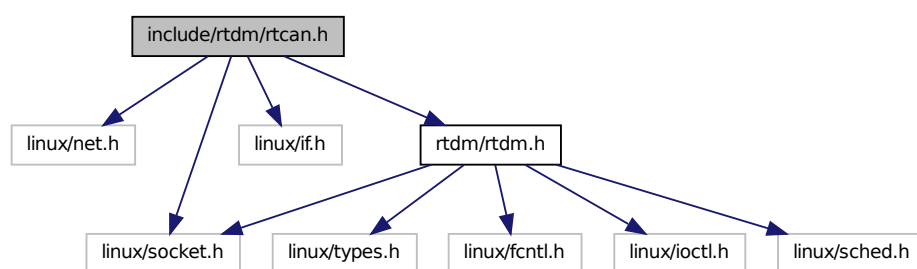
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.16 include/rtdm/rtdm.h File Reference

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

Include dependency graph for rtdm.h:



### Data Structures

- struct [can\\_bittime\\_std](#)  
*Standard bit-time parameters according to Bosch.*
- struct [can\\_bittime\\_btr](#)  
*Hardware-specific BTR bit-times.*
- struct [can\\_bittime](#)  
*Custom CAN bit-time definition.*
- struct [can\\_filter](#)  
*Filter for reception of CAN messages.*
- struct [sockaddr\\_can](#)  
*Socket address structure for the CAN address family.*
- struct [can\\_frame](#)  
*Raw CAN frame.*

### Defines

- #define [AF\\_CAN](#) 29

*CAN address family.*

- `#define PF_CAN AF_CAN`  
*CAN protocol family.*
- `#define SOL_CAN_RAW 103`  
*CAN socket levels.*

### CAN ID masks

*Bit masks for masking CAN IDs*

- `#define CAN_EFF_MASK 0x1FFFFFFF`  
*Bit mask for extended CAN IDs.*
- `#define CAN_SFF_MASK 0x000007FF`  
*Bit mask for standard CAN IDs.*

### CAN ID flags

*Flags within a CAN ID indicating special CAN frame attributes*

- `#define CAN_EFF_FLAG 0x80000000`  
*Extended frame.*
- `#define CAN_RTR_FLAG 0x40000000`  
*Remote transmission frame.*
- `#define CAN_ERR_FLAG 0x20000000`  
*Error frame (see [Errors](#)), not valid in struct `can_filter`.*
- `#define CAN_INV_FILTER CAN_ERR_FLAG`  
*Invert CAN filter definition, only valid in struct `can_filter`.*

### Particular CAN protocols

*Possible protocols for the PF\_CAN protocol family*

*Currently only the RAW protocol is supported.*

- `#define CAN_RAW 1`  
*Raw protocol of PF\_CAN, applicable to socket type `SOCK_RAW`.*

### CAN controller modes

*Special CAN controllers modes, which can be or'ed together.*

#### Note

*These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.*

- `#define CAN_CTRLMODE_LISTENONLY 0x1`
- `#define CAN_CTRLMODE_LOOPBACK 0x2`

### Timestamp switches

Arguments to pass to [RTCAN\\_RTIOC\\_TAKE\\_TIMESTAMP](#)

- #define [RTCAN\\_TAKE\\_NO\\_TIMESTAMPS](#) 0  
*Switch off taking timestamps.*
- #define [RTCAN\\_TAKE\\_TIMESTAMPS](#) 1  
*Do take timestamps.*

### RAW socket options

Setting and getting CAN RAW socket options.

- #define [CAN\\_RAW\\_FILTER](#) 0x1  
*CAN filter definition.*
- #define [CAN\\_RAW\\_ERR\\_FILTER](#) 0x2  
*CAN error mask.*
- #define [CAN\\_RAW\\_LOOPBACK](#) 0x3  
*CAN TX loopback.*
- #define [CAN\\_RAW\\_RECV\\_OWN\\_MSGS](#) 0x4  
*CAN receive own messages.*

### IOCTLs

CAN device IOCTLs

- #define [SIOCGIFINDEX](#) defined\_by\_kernel\_header\_file  
*Get CAN interface index by name.*
- #define [SIOCSCANBAUDRATE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x01, struct ifreq)  
*Set baud rate.*
- #define [SIOCGCANBAUDRATE](#) \_IOWR(RTIOC\_TYPE\_CAN, 0x02, struct ifreq)  
*Get baud rate.*
- #define [SIOCSCANCUSTOMBITTIME](#) \_IOW(RTIOC\_TYPE\_CAN, 0x03, struct ifreq)  
*Set custom bit time parameter.*
- #define [SIOCGCANCUSTOMBITTIME](#) \_IOWR(RTIOC\_TYPE\_CAN, 0x04, struct ifreq)  
*Get custom bit-time parameters.*
- #define [SIOCSCANMODE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x05, struct ifreq)  
*Set operation mode of CAN controller.*
- #define [SIOCGCANSTATE](#) \_IOWR(RTIOC\_TYPE\_CAN, 0x06, struct ifreq)  
*Get current state of CAN controller.*
- #define [SIOCSCANCTRLMODE](#) \_IOW(RTIOC\_TYPE\_CAN, 0x07, struct ifreq)  
*Set special controller modes.*
- #define [SIOCGCANCTRLMODE](#) \_IOWR(RTIOC\_TYPE\_CAN, 0x08, struct ifreq)

*Get special controller modes.*

- #define [RTCAN\\_RTIOC\\_TAKE\\_TIMESTAMP](#) \_IOW(RTIOC\_TYPE\_CAN, 0x09, int)  
*Enable or disable storing a high precision timestamp upon reception of a CAN frame.*
- #define [RTCAN\\_RTIOC\\_RCV\\_TIMEOUT](#) \_IOW(RTIOC\_TYPE\_CAN, 0x0A, nanosecs\_rel\_t)  
*Specify a reception timeout for a socket.*
- #define [RTCAN\\_RTIOC\\_SND\\_TIMEOUT](#) \_IOW(RTIOC\_TYPE\_CAN, 0x0B, nanosecs\_rel\_t)  
*Specify a transmission timeout for a socket.*

### Error mask

Error class (mask) in `can_id` field of struct `can_frame` to be used with [CAN\\_RAW\\_ERR\\_FILTER](#).

**Note:** Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.

**Note:** In case of a bus-off error condition ([CAN\\_ERR\\_BUSOFF](#)), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling [CAN\\_MODE\\_START](#).

**Note:** Bus error interrupts ([CAN\\_ERR\\_BUSERROR](#)) are enabled when an application is calling a [Recv](#) function on a socket listening on bus errors (using [CAN\\_RAW\\_ERR\\_FILTER](#)). After one bus error has occurred, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.

- #define [CAN\\_ERR\\_TX\\_TIMEOUT](#) 0x00000001U  
*TX timeout (netdevice driver).*
- #define [CAN\\_ERR\\_LOSTARB](#) 0x00000002U  
*Lost arbitration (see [data\[0\]](#)).*
- #define [CAN\\_ERR\\_CRTL](#) 0x00000004U  
*Controller problems (see [data\[1\]](#)).*
- #define [CAN\\_ERR\\_PROT](#) 0x00000008U  
*Protocol violations (see [data\[2\]](#), [data\[3\]](#)).*
- #define [CAN\\_ERR\\_TRX](#) 0x00000010U  
*Transceiver status (see [data\[4\]](#)).*
- #define [CAN\\_ERR\\_ACK](#) 0x00000020U  
*Received no ACK on transmission.*
- #define [CAN\\_ERR\\_BUSOFF](#) 0x00000040U  
*Bus off.*
- #define [CAN\\_ERR\\_BUSERROR](#) 0x00000080U  
*Bus error (may flood!).*
- #define [CAN\\_ERR\\_RESTARTED](#) 0x00000100U  
*Controller restarted.*

- #define [CAN\\_ERR\\_MASK](#) 0x1FFFFFFU  
*Omit EFF, RTR, ERR flags.*

#### Arbitration lost error

Error in the data[0] field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_LOSTARB\\_UNSPEC](#) 0x00  
*unspecified*

#### Controller problems

Error in the data[1] field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_CRTL\\_UNSPEC](#) 0x00  
*unspecified*
- #define [CAN\\_ERR\\_CRTL\\_RX\\_OVERFLOW](#) 0x01  
*RX buffer overflow.*
- #define [CAN\\_ERR\\_CRTL\\_TX\\_OVERFLOW](#) 0x02  
*TX buffer overflow.*
- #define [CAN\\_ERR\\_CRTL\\_RX\\_WARNING](#) 0x04  
*reached warning level for RX errors*
- #define [CAN\\_ERR\\_CRTL\\_TX\\_WARNING](#) 0x08  
*reached warning level for TX errors*
- #define [CAN\\_ERR\\_CRTL\\_RX\\_PASSIVE](#) 0x10  
*reached passive level for RX errors*
- #define [CAN\\_ERR\\_CRTL\\_TX\\_PASSIVE](#) 0x20  
*reached passive level for TX errors*

#### Protocol error type

Error in the data[2] field of struct [can\\_frame](#).

- #define [CAN\\_ERR\\_PROT\\_UNSPEC](#) 0x00  
*unspecified*
- #define [CAN\\_ERR\\_PROT\\_BIT](#) 0x01  
*single bit error*
- #define [CAN\\_ERR\\_PROT\\_FORM](#) 0x02  
*frame format error*
- #define [CAN\\_ERR\\_PROT\\_STUFF](#) 0x04  
*bit stuffing error*
- #define [CAN\\_ERR\\_PROT\\_BIT0](#) 0x08  
*unable to send dominant bit*

- #define `CAN_ERR_PROT_BIT1` 0x10  
*unable to send recessive bit*
- #define `CAN_ERR_PROT_OVERLOAD` 0x20  
*bus overload*
- #define `CAN_ERR_PROT_ACTIVE` 0x40  
*active error announcement*
- #define `CAN_ERR_PROT_TX` 0x80  
*error occurred on transmission*

### Protocol error location

Error in the `data[4]` field of struct `can_frame`.

- #define `CAN_ERR_PROT_LOC_UNSPEC` 0x00  
*unspecified*
- #define `CAN_ERR_PROT_LOC_SOF` 0x03  
*start of frame*
- #define `CAN_ERR_PROT_LOC_ID28_21` 0x02  
*ID bits 28 - 21 (SFF: 10 - 3).*
- #define `CAN_ERR_PROT_LOC_ID20_18` 0x06  
*ID bits 20 - 18 (SFF: 2 - 0).*
- #define `CAN_ERR_PROT_LOC_SRTR` 0x04  
*substitute RTR (SFF: RTR)*
- #define `CAN_ERR_PROT_LOC_IDE` 0x05  
*identifier extension*
- #define `CAN_ERR_PROT_LOC_ID17_13` 0x07  
*ID bits 17-13.*
- #define `CAN_ERR_PROT_LOC_ID12_05` 0x0F  
*ID bits 12-5.*
- #define `CAN_ERR_PROT_LOC_ID04_00` 0x0E  
*ID bits 4-0.*
- #define `CAN_ERR_PROT_LOC_RTR` 0x0C  
*RTR.*
- #define `CAN_ERR_PROT_LOC_RES1` 0x0D  
*reserved bit 1*
- #define `CAN_ERR_PROT_LOC_RES0` 0x09  
*reserved bit 0*



- #define [CAN\\_ERR\\_PROT\\_LOC\\_DLC](#) 0x0B  
*data length code*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_DATA](#) 0x0A  
*data section*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_CRC\\_SEQ](#) 0x08  
*CRC sequence.*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_CRC\\_DEL](#) 0x18  
*CRC delimiter.*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_ACK](#) 0x19  
*ACK slot.*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_ACK\\_DEL](#) 0x1B  
*ACK delimiter.*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_EOF](#) 0x1A  
*end of frame*
- #define [CAN\\_ERR\\_PROT\\_LOC\\_INTERM](#) 0x12  
*intermission*
- #define [CAN\\_ERR\\_TRX\\_UNSPEC](#) 0x00  
*0000 0000*
- #define [CAN\\_ERR\\_TRX\\_CANH\\_NO\\_WIRE](#) 0x04  
*0000 0100*
- #define [CAN\\_ERR\\_TRX\\_CANH\\_SHORT\\_TO\\_BAT](#) 0x05  
*0000 0101*
- #define [CAN\\_ERR\\_TRX\\_CANH\\_SHORT\\_TO\\_VCC](#) 0x06  
*0000 0110*
- #define [CAN\\_ERR\\_TRX\\_CANH\\_SHORT\\_TO\\_GND](#) 0x07  
*0000 0111*
- #define [CAN\\_ERR\\_TRX\\_CANL\\_NO\\_WIRE](#) 0x40  
*0100 0000*
- #define [CAN\\_ERR\\_TRX\\_CANL\\_SHORT\\_TO\\_BAT](#) 0x50  
*0101 0000*
- #define [CAN\\_ERR\\_TRX\\_CANL\\_SHORT\\_TO\\_VCC](#) 0x60  
*0110 0000*
- #define [CAN\\_ERR\\_TRX\\_CANL\\_SHORT\\_TO\\_GND](#) 0x70  
*0111 0000*
- #define [CAN\\_ERR\\_TRX\\_CANL\\_SHORT\\_TO\\_CANH](#) 0x80  
*1000 0000*

## Typedefs

- typedef uint32\_t [can\\_id\\_t](#)  
*Type of CAN id (see [CAN\\_xxx\\_MASK](#) and [CAN\\_xxx\\_FLAG](#)).*
- typedef [can\\_id\\_t](#) [can\\_err\\_mask\\_t](#)  
*Type of CAN error mask.*
- typedef uint32\_t [can\\_baudrate\\_t](#)  
*Baudrate definition in bits per second.*
- typedef enum [CAN\\_BITTIME\\_TYPE](#) [can\\_bittime\\_type\\_t](#)  
*See [CAN\\_BITTIME\\_TYPE](#).*
- typedef enum [CAN\\_MODE](#) [can\\_mode\\_t](#)  
*See [CAN\\_MODE](#).*
- typedef int [can\\_ctrlmode\\_t](#)  
*See [CAN\\_CTRLMODE](#).*
- typedef enum [CAN\\_STATE](#) [can\\_state\\_t](#)  
*See [CAN\\_STATE](#).*
- typedef struct [can\\_filter](#) [can\\_filter\\_t](#)  
*Filter for reception of CAN messages.*
- typedef struct [can\\_frame](#) [can\\_frame\\_t](#)  
*Raw CAN frame.*

## Enumerations

- enum [CAN\\_BITTIME\\_TYPE](#) { [CAN\\_BITTIME\\_STD](#), [CAN\\_BITTIME\\_BTR](#) }  
*Supported CAN bit-time types.*

### CAN operation modes

*Modes into which CAN controllers can be set*

- enum [CAN\\_MODE](#) { [CAN\\_MODE\\_STOP](#) = 0, [CAN\\_MODE\\_START](#), [CAN\\_MODE\\_SLEEP](#) }

### CAN controller states

*States a CAN controller can be in.*

- enum [CAN\\_STATE](#) {  
    [CAN\\_STATE\\_ACTIVE](#) = 0, [CAN\\_STATE\\_BUS\\_WARNING](#), [CAN\\_STATE\\_BUS\\_PASSIVE](#), [CAN\\_STATE\\_BUS\\_OFF](#),  
    [CAN\\_STATE\\_SCANNING\\_BAUDRATE](#), [CAN\\_STATE\\_STOPPED](#), [CAN\\_STATE\\_SLEEPING](#) }

### 7.16.1 Detailed Description

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

#### Note

Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>

Copyright (C) 2005, 2006 Sebastian Smolorz <Sebastian.Smolorz@stud.uni-hannover.de>

This RTDM CAN device profile header is based on:

include/linux/can.h, include/linux/socket.h, net/can/pf\_can.h in linux-can.patch, a CAN socket framework for Linux

Copyright (C) 2004, 2005, Robert Schwebel, Benedikt Spranger, Marc Kleine-Budde, Pengutronix

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU

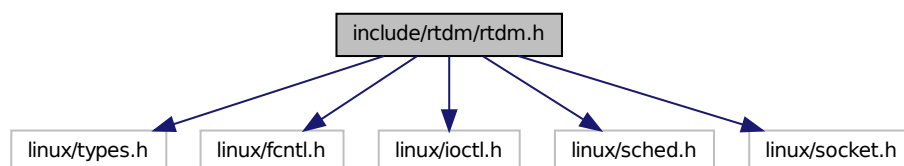
General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

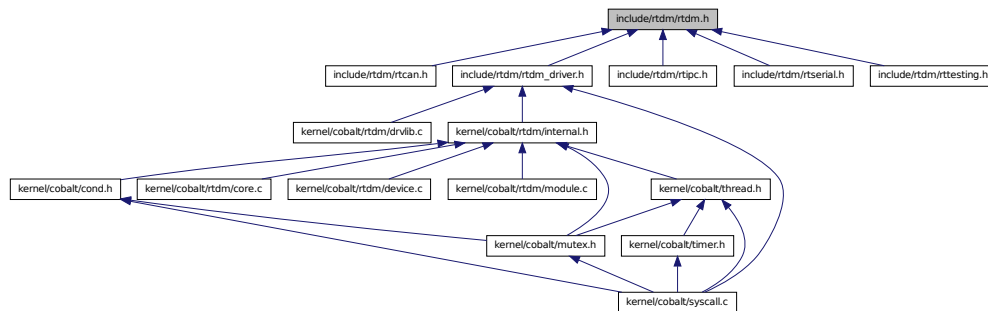
## 7.17 include/rtdm/rtdm.h File Reference

Real-Time Driver Model for Xenomai, user API header.

Include dependency graph for rtdm.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [rtdm\\_device\\_info](#)  
*Device information.*

## Defines

### API Versioning

- #define [RTDM\\_API\\_VER](#) 8  
*Common user and driver API version.*
- #define [RTDM\\_API\\_MIN\\_COMPAT\\_VER](#) 6  
*Minimum API revision compatible with the current release.*

### RTDM\_TIMEOUT\_XXX

*Special timeout values*

- #define [RTDM\\_TIMEOUT\\_INFINITE](#) 0  
*Block forever.*
- #define [RTDM\\_TIMEOUT\\_NONE](#) (-1)  
*Any negative timeout means non-blocking.*

### RTDM\_CLASS\_XXX

*Device classes*

- #define [RTDM\\_CLASS\\_PARPORT](#) 1
- #define [RTDM\\_CLASS\\_SERIAL](#) 2
- #define [RTDM\\_CLASS\\_CAN](#) 3
- #define [RTDM\\_CLASS\\_NETWORK](#) 4
- #define [RTDM\\_CLASS\\_RTMAC](#) 5
- #define [RTDM\\_CLASS\\_TESTING](#) 6

- #define **RTDM\_CLASS\_RTIPC** 7
- #define **RTDM\_CLASS\_EXPERIMENTAL** 224
- #define **RTDM\_CLASS\_MAX** 255

### Device Naming

*Maximum length of device names (excluding the final null character)*

- #define **RTDM\_MAX\_DEVNAME\_LEN** 31

### RTDM\_PURGE\_XXX\_BUFFER

*Flags selecting buffers to be purged*

- #define **RTDM\_PURGE\_RX\_BUFFER** 0x0001
- #define **RTDM\_PURGE\_TX\_BUFFER** 0x0002

### Common IOCTLs

*The following IOCTLs are common to all device profiles.*

- #define **RTIOC\_DEVICE\_INFO** \_IOR(RTIOC\_TYPE\_COMMON, 0x00, struct rtdm\_device\_info)  
*Retrieve information about a device or socket.*
- #define **RTIOC\_PURGE** \_IOW(RTIOC\_TYPE\_COMMON, 0x10, int)  
*Purge internal device or socket buffers.*

## Typedefs

- typedef uint64\_t **nanosecs\_abs\_t**  
*RTDM type for representing absolute dates.*
- typedef int64\_t **nanosecs\_rel\_t**  
*RTDM type for representing relative intervals.*
- typedef struct **rtdm\_device\_info** **rtdm\_device\_info\_t**  
*Device information.*

### 7.17.1 Detailed Description

Real-Time Driver Model for Xenomai, user API header.

#### Note

Copyright (C) 2005, 2006 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>  
Copyright (C) 2005 Joerg Langenberg <[joerg.langenberg@gmx.net](mailto:joerg.langenberg@gmx.net)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.



- struct [rtdm\\_device](#)

*RTDM device.*

## Defines

- #define [rtdm\\_irq\\_get\\_arg](#)(irq\_handle, type) ((type \*)irq\_handle->cookie)

*Retrieve IRQ handler argument.*

## Device Flags

*Static flags describing a RTDM device*

- #define [RTDM\\_EXCLUSIVE](#) 0x0001  
*If set, only a single instance of the device can be requested by an application.*
- #define [RTDM\\_NAMED\\_DEVICE](#) 0x0010  
*If set, the device is addressed via a clear-text name.*
- #define [RTDM\\_PROTOCOL\\_DEVICE](#) 0x0020  
*If set, the device is addressed via a combination of protocol ID and socket type.*
- #define [RTDM\\_DEVICE\\_TYPE\\_MASK](#) 0x00F0  
*Mask selecting the device type.*

## Context Flags

*Dynamic flags describing the state of an open RTDM device (bit numbers)*

- #define [RTDM\\_CREATED\\_IN\\_NRT](#) 0  
*Set by RTDM if the device instance was created in non-real-time context.*
- #define [RTDM\\_CLOSING](#) 1  
*Set by RTDM when the device is being closed.*
- #define [RTDM\\_USER\\_CONTEXT\\_FLAG](#) 8  
*Lowest bit number the driver developer can use freely.*

## Driver Versioning

*Current revisions of RTDM structures, encoding of driver versions. See [API Versioning](#) for the interface revision.*

- #define [RTDM\\_DEVICE\\_STRUCT\\_VER](#) 5  
*Version of struct [rtdm\\_device](#).*
- #define [RTDM\\_CONTEXT\\_STRUCT\\_VER](#) 3  
*Version of struct [rtdm\\_dev\\_context](#).*
- #define [RTDM\\_SECURE\\_DEVICE](#) 0x80000000  
*Flag indicating a secure variant of RTDM (not supported here).*

- #define **RTDM\_DRIVER\_VER**(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))  
*Version code constructor for driver revisions.*
- #define **RTDM\_DRIVER\_MAJOR\_VER**(ver) (((ver) >> 16) & 0xFF)  
*Get major version number from driver revision code.*
- #define **RTDM\_DRIVER\_MINOR\_VER**(ver) (((ver) >> 8) & 0xFF)  
*Get minor version number from driver revision code.*
- #define **RTDM\_DRIVER\_PATCH\_VER**(ver) ((ver) & 0xFF)  
*Get patch version number from driver revision code.*

### Global Lock across Scheduler Invocation

- #define **RTDM\_EXECUTE\_ATOMICALLY**(code\_block)  
*Execute code block atomically.*

### RTDM\_IRQTYPE\_xxx

*Interrupt registrations flags*

- #define **RTDM\_IRQTYPE\_SHARED** XN\_ISR\_SHARED  
*Enable IRQ-sharing with other real-time drivers.*
- #define **RTDM\_IRQTYPE\_EDGE** XN\_ISR\_EDGE  
*Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.*

### RTDM\_IRQ\_xxx

*Return flags of interrupt handlers*

- #define **RTDM\_IRQ\_NONE** XN\_ISR\_NONE  
*Unhandled interrupt.*
- #define **RTDM\_IRQ\_HANDLED** XN\_ISR\_HANDLED  
*Denote handled interrupt.*

### Task Priority Range

*Maximum and minimum task priorities*

- #define **RTDM\_TASK\_LOWEST\_PRIORITY** XNSCHED\_LOW\_PRIO
- #define **RTDM\_TASK\_HIGHEST\_PRIORITY** XNSCHED\_HIGH\_PRIO

### Task Priority Modification

*Raise or lower task priorities by one level*

- #define **RTDM\_TASK\_RAISE\_PRIORITY** (+1)
- #define **RTDM\_TASK\_LOWER\_PRIORITY** (-1)



## Typedefs

- typedef int(\* [rtdm\\_irq\\_handler\\_t](#))(rtdm\_irq\_t \*irq\_handle)  
*Interrupt handler.*
- typedef void(\* [rtdm\\_nrtsig\\_handler\\_t](#))(rtdm\_nrtsig\_t nrt\_sig, void \*arg)  
*Non-real-time signal handler.*
- typedef void(\* [rtdm\\_timer\\_handler\\_t](#))(rtdm\_timer\_t \*timer)  
*Timer handler.*
- typedef void(\* [rtdm\\_task\\_proc\\_t](#))(void \*arg)  
*Real-time task procedure.*

## Operation Handler Prototypes

- typedef int(\* [rtdm\\_open\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, int oflag)  
*Named device open handler.*
- typedef int(\* [rtdm\\_socket\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, int protocol)  
*Socket creation handler for protocol devices.*
- typedef int(\* [rtdm\\_close\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info)  
*Close handler.*
- typedef int(\* [rtdm\\_ioctl\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, unsigned int request, void \_\_user \*arg)  
*IOCTL handler.*
- typedef int(\* [rtdm\\_select\\_bind\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_selector\_t \*selector, enum [rtdm\\_selecttype](#) type, unsigned fd\_index)  
*Select binding handler.*
- typedef ssize\_t(\* [rtdm\\_read\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, void \*buf, size\_t nbyte)  
*Read handler.*
- typedef ssize\_t(\* [rtdm\\_write\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, const void \*buf, size\_t nbyte)  
*Write handler.*
- typedef ssize\_t(\* [rtdm\\_recvmmsg\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, struct msghdr \*msg, int flags)  
*Receive message handler.*
- typedef ssize\_t(\* [rtdm\\_sendmsg\\_handler\\_t](#))(struct [rtdm\\_dev\\_context](#) \*context, rtdm\_user\_info\_t \*user\_info, const struct msghdr \*msg, int flags)  
*Transmit message handler.*

## Enumerations

### RTDM\_SELECTTYPE\_XXX

*Event types select can bind to*

- enum `rtdm_selecttype` { `RTDM_SELECTTYPE_READ` = `XNSELECT_READ`, `RTDM_SELECTTYPE_WRITE` = `XNSELECT_WRITE`, `RTDM_SELECTTYPE_EXCEPT` = `XNSELECT_EXCEPT` }

### RTDM\_TIMERMODE\_XXX

*Timer operation modes*

- enum `rtdm_timer_mode` { `RTDM_TIMERMODE_RELATIVE` = `XN_RELATIVE`, `RTDM_TIMERMODE_ABSOLUTE` = `XN_ABSOLUTE`, `RTDM_TIMERMODE_REALTIME` = `XN_REALTIME` }

## Functions

- static void \* `rtdm_context_to_private` (struct `rtdm_dev_context` \*context)  
*Locate the driver private area associated to a device context structure.*
- static struct `rtdm_dev_context` \* `rtdm_private_to_context` (void \*dev\_private)  
*Locate a device context structure from its driver private area.*
- int `rtdm_dev_register` (struct `rtdm_device` \*device)  
*Register a RTDM device.*
- int `rtdm_dev_unregister` (struct `rtdm_device` \*device, unsigned int poll\_delay)  
*Unregisters a RTDM device.*
- struct `rtdm_dev_context` \* `rtdm_context_get` (int fd)  
*Retrieve and lock a device context.*
- int `rtdm_select_bind` (int fd, rtdm\_selector\_t \*selector, enum `rtdm_selecttype` type, unsigned fd\_index)  
*Bind a selector to specified event types of a given file descriptor.*
- int `rtdm_irq_request` (rtdm\_irq\_t \*irq\_handle, unsigned int irq\_no, `rtdm_irq_handler_t` handler, unsigned long flags, const char \*device\_name, void \*arg)  
*Register an interrupt handler.*
- void `rtdm_timer_destroy` (rtdm\_timer\_t \*timer)  
*Destroy a timer.*
- int `rtdm_timer_start` (rtdm\_timer\_t \*timer, `nanosecs_abs_t` expiry, `nanosecs_rel_t` interval, enum `rtdm_timer_mode` mode)  
*Start a timer.*
- void `rtdm_timer_stop` (rtdm\_timer\_t \*timer)  
*Stop a timer.*

- int [rtdm\\_task\\_init](#) (rtdm\_task\_t \*task, const char \*name, [rtdm\\_task\\_proc\\_t](#) task\_proc, void \*arg, int priority, [nanosecs\\_rel\\_t](#) period)  
*Initialise and start a real-time task.*
- void [rtdm\\_task\\_busy\\_sleep](#) ([nanosecs\\_rel\\_t](#) delay)  
*Busy-wait a specified amount of time.*
- void [rtdm\\_toseq\\_init](#) (rtdm\_toseq\_t \*timeout\_seq, [nanosecs\\_rel\\_t](#) timeout)  
*Initialise a timeout sequence.*
- void [rtdm\\_event\\_init](#) (rtdm\_event\_t \*event, unsigned long pending)  
*Initialise an event.*
- int [rtdm\\_event\\_select\\_bind](#) (rtdm\_event\_t \*event, rtdm\_selector\_t \*selector, enum [rtdm\\_selecttype](#) type, unsigned fd\_index)  
*Bind a selector to an event.*
- int [rtdm\\_event\\_wait](#) (rtdm\_event\_t \*event)  
*Wait on event occurrence.*
- int [rtdm\\_event\\_timedwait](#) (rtdm\_event\_t \*event, [nanosecs\\_rel\\_t](#) timeout, rtdm\_toseq\_t \*timeout\_seq)  
*Wait on event occurrence with timeout.*
- void [rtdm\\_event\\_signal](#) (rtdm\_event\_t \*event)  
*Signal an event occurrence.*
- void [rtdm\\_event\\_clear](#) (rtdm\_event\_t \*event)  
*Clear event state.*
- void [rtdm\\_sem\\_init](#) (rtdm\_sem\_t \*sem, unsigned long value)  
*Initialise a semaphore.*
- int [rtdm\\_sem\\_select\\_bind](#) (rtdm\_sem\_t \*sem, rtdm\_selector\_t \*selector, enum [rtdm\\_selecttype](#) type, unsigned fd\_index)  
*Bind a selector to a semaphore.*
- int [rtdm\\_sem\\_down](#) (rtdm\_sem\_t \*sem)  
*Decrement a semaphore.*
- int [rtdm\\_sem\\_timeddown](#) (rtdm\_sem\_t \*sem, [nanosecs\\_rel\\_t](#) timeout, rtdm\_toseq\_t \*timeout\_seq)  
*Decrement a semaphore with timeout.*
- void [rtdm\\_sem\\_up](#) (rtdm\_sem\_t \*sem)  
*Increment a semaphore.*
- void [rtdm\\_mutex\\_init](#) (rtdm\_mutex\_t \*mutex)  
*Initialise a mutex.*

- int [rtdm\\_mutex\\_lock](#) (rtdm\_mutex\_t \*mutex)  
*Request a mutex.*
- int [rtdm\\_mutex\\_timedlock](#) (rtdm\_mutex\_t \*mutex, [nanosecs\\_rel\\_t](#) timeout, rtdm\_toseq\_t \*timeout\_seq)  
*Request a mutex with timeout.*

## Spinlock with Preemption Deactivation

- #define [RTDM\\_LOCK\\_UNLOCKED](#) IPIPE\_SPIN\_LOCK\_UNLOCKED  
*Static lock initialisation.*
- #define [rtdm\\_lock\\_init](#)(lock) spin\_lock\_init(lock)  
*Dynamic lock initialisation.*
- #define [rtdm\\_lock\\_get](#)(lock) spin\_lock(lock)  
*Acquire lock from non-preemptible contexts.*
- #define [rtdm\\_lock\\_put](#)(lock) spin\_unlock(lock)  
*Release lock without preemption restoration.*
- #define [rtdm\\_lock\\_get\\_irqsave](#)(lock, context) spin\_lock\_irqsave(lock, context)  
*Acquire lock and disable preemption.*
- #define [rtdm\\_lock\\_put\\_irqrestore](#)(lock, context) spin\_unlock\_irqrestore(lock, context)  
*Release lock and restore preemption state.*
- #define [rtdm\\_lock\\_irqsave](#)(context) splhigh(context)  
*Disable preemption locally.*
- #define [rtdm\\_lock\\_irqrestore](#)(context) splexit(context)  
*Restore preemption state.*
- typedef ipipe\_spinlock\_t [rtdm\\_lock\\_t](#)  
*Lock variable.*
- typedef unsigned long [rtdm\\_lockctx\\_t](#)  
*Variable to save the context while holding a lock.*

### 7.18.1 Detailed Description

Real-Time Driver Model for Xenomai, driver API header.

#### Note

Copyright (C) 2005-2007 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>  
 Copyright (C) 2005 Joerg Langenberg <[joerg.langenberg@gmx.net](mailto:joerg.langenberg@gmx.net)>  
 Copyright (C) 2008 Gilles Chanteperdrix <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

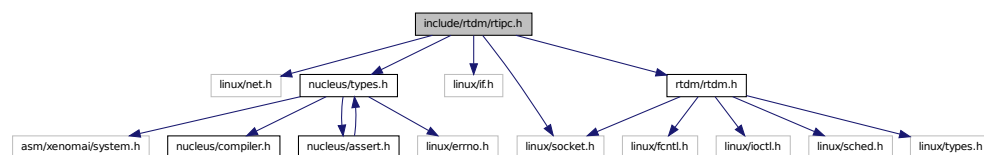
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.19 include/rtdm/rtipc.h File Reference

This file is part of the Xenomai project.

Include dependency graph for rtipc.h:



### Data Structures

- struct [rtipc\\_port\\_label](#)  
*Port label information structure.*
- struct [sockaddr\\_ipc](#)  
*Socket address structure for the RTIPC address family.*

### Defines

#### XDDP socket options

*Setting and getting XDDP socket options.*

- #define [XDDP\\_LABEL](#) 1  
*XDDP label assignment.*
- #define [XDDP\\_POOLSZ](#) 2  
*XDDP local pool size configuration.*
- #define [XDDP\\_BUFSZ](#) 3  
*XDDP streaming buffer size configuration.*
- #define [XDDP\\_MONITOR](#) 4

*XDDP monitoring callback.*

### XDDP events

*Specific events occurring on XDDP channels, which can be monitored via the [XDDP\\_MONITOR](#) socket option.*

- `#define XDDP_EVTIN 1`  
*Monitor writes to the non real-time endpoint.*
- `#define XDDP_EVTOUT 2`  
*Monitor reads from the non real-time endpoint.*
- `#define XDDP_EVTDOWN 3`  
*Monitor close from the non real-time endpoint.*
- `#define XDDP_EVTNOBUF 4`  
*Monitor memory shortage for non real-time datagrams.*

### IDDP socket options

*Setting and getting IDDP socket options.*

- `#define IDDP_LABEL 1`  
*IDDP label assignment.*
- `#define IDDP_POOLSZ 2`  
*IDDP local pool size configuration.*

### BUFP socket options

*Setting and getting BUFP socket options.*

- `#define BUFP_LABEL 1`  
*BUFP label assignment.*
- `#define BUFP_BUFSZ 2`  
*BUFP buffer size configuration.*

### Socket level options

*Setting and getting supported standard socket level options.*

- `#define SO_SNDTIMEO defined_by_kernel_header_file`  
*IPPROTO\_IDDP and IPPROTO\_BUFP protocols support the standard SO\_SNDTIMEO socket option, from the SOL\_SOCKET level.*
- `#define SO_RCVTIMEO defined_by_kernel_header_file`  
*All RTIPC protocols support the standard SO\_RCVTIMEO socket option, from the SOL\_SOCKET level.*

## Typedefs

- typedef int16\_t [rtipc\\_port\\_t](#)  
*Port number type for the RTIPC address family.*

## Enumerations

### RTIPC protocol list

*protocols for the PF\_RTIPC protocol family*

- enum { [IPCPROTO\\_IPC](#) = 0, [IPCPROTO\\_XDDP](#) = 1, [IPCPROTO\\_IDDP](#) = 2, [IPCPROTO\\_BUFP](#) = 3 }

## Functions

### Supported operations

*Standard socket operations supported by the RTIPC protocols.*

- int [socket\\_\\_AF\\_RTIPC](#) (int domain=AF\_RTIPC, int type=SOCK\_DGRAM, int protocol)  
*Create an endpoint for communication in the AF\_RTIPC domain.*
- int [close\\_\\_AF\\_RTIPC](#) (int sockfd)  
*Close a RTIPC socket descriptor.*
- int [bind\\_\\_AF\\_RTIPC](#) (int sockfd, const struct [sockaddr\\_ipc](#) \*addr, socklen\_t addrlen)  
*Bind a RTIPC socket to a port.*
- int [connect\\_\\_AF\\_RTIPC](#) (int sockfd, const struct [sockaddr\\_ipc](#) \*addr, socklen\_t addrlen)  
*Initiate a connection on a RTIPC socket.*
- int [setsockopt\\_\\_AF\\_RTIPC](#) (int sockfd, int level, int optname, const void \*optval, socklen\_t optlen)  
*Set options on RTIPC sockets.*
- int [getsockopt\\_\\_AF\\_RTIPC](#) (int sockfd, int level, int optname, void \*optval, socklen\_t \*optlen)  
*Get options on RTIPC sockets.*
- ssize\_t [sendmsg\\_\\_AF\\_RTIPC](#) (int sockfd, const struct msghdr \*msg, int flags)  
*Send a message on a RTIPC socket.*
- ssize\_t [recvmsg\\_\\_AF\\_RTIPC](#) (int sockfd, struct msghdr \*msg, int flags)  
*Receive a message from a RTIPC socket.*
- int [getsockname\\_\\_AF\\_RTIPC](#) (int sockfd, struct [sockaddr\\_ipc](#) \*addr, socklen\_t \*addrlen)  
*Get socket name.*
- int [getpeername\\_\\_AF\\_RTIPC](#) (int sockfd, struct [sockaddr\\_ipc](#) \*addr, socklen\_t \*addrlen)  
*Get socket peer.*

### 7.19.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

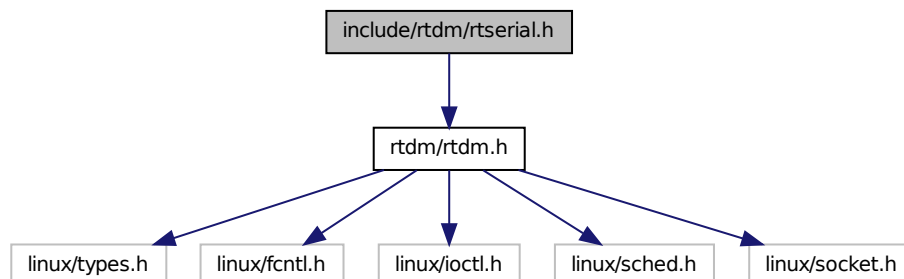
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.20 include/rtdm/rtserial.h File Reference

Real-Time Driver Model for Xenomai, serial device profile header.

Include dependency graph for rtserial.h:



### Data Structures

- struct [rtser\\_config](#)  
*Serial device configuration.*
- struct [rtser\\_status](#)  
*Serial device status.*
- struct [rtser\\_event](#)  
*Additional information about serial device events.*



## Defines

- #define [RTSER\\_RTIOC\\_BREAK\\_CTL](#) \_IOR(RTIOC\_TYPE\_SERIAL, 0x06, int)  
*Set or clear break on UART output line.*

### RTSER\_DEF\_BAUD

*Default baud rate*

- #define [RTSER\\_DEF\\_BAUD](#) 9600

### RTSER\_xxx\_PARITY

*Number of parity bits*

- #define [RTSER\\_NO\\_PARITY](#) 0x00
- #define [RTSER\\_ODD\\_PARITY](#) 0x01
- #define [RTSER\\_EVEN\\_PARITY](#) 0x03
- #define [RTSER\\_DEF\\_PARITY](#) [RTSER\\_NO\\_PARITY](#)

### RTSER\_xxx\_BITS

*Number of data bits*

- #define [RTSER\\_5\\_BITS](#) 0x00
- #define [RTSER\\_6\\_BITS](#) 0x01
- #define [RTSER\\_7\\_BITS](#) 0x02
- #define [RTSER\\_8\\_BITS](#) 0x03
- #define [RTSER\\_DEF\\_BITS](#) [RTSER\\_8\\_BITS](#)

### RTSER\_xxx\_STOPB

*Number of stop bits*

- #define [RTSER\\_1\\_STOPB](#) 0x00  
*valid only in combination with 5 data bits*
- #define [RTSER\\_1\\_5\\_STOPB](#) 0x01  
*valid only in combination with 5 data bits*
- #define [RTSER\\_2\\_STOPB](#) 0x01  
*valid only in combination with 5 data bits*
- #define [RTSER\\_DEF\\_STOPB](#) [RTSER\\_1\\_STOPB](#)  
*valid only in combination with 5 data bits*

### RTSER\_xxx\_HAND

*Handshake mechanisms*

- #define [RTSER\\_NO\\_HAND](#) 0x00
- #define [RTSER\\_RTSCTS\\_HAND](#) 0x01
- #define [RTSER\\_DEF\\_HAND](#) [RTSER\\_NO\\_HAND](#)

### RTSER\_FIFO\_xxx

*Reception FIFO interrupt threshold*

- #define **RTSER\_FIFO\_DEPTH\_1** 0x00
- #define **RTSER\_FIFO\_DEPTH\_4** 0x40
- #define **RTSER\_FIFO\_DEPTH\_8** 0x80
- #define **RTSER\_FIFO\_DEPTH\_14** 0xC0
- #define **RTSER\_DEF\_FIFO\_DEPTH** RTSER\_FIFO\_DEPTH\_1

#### **RTSER\_TIMEOUT\_xxx**

*Special timeout values, see also [RTDM\\_TIMEOUT\\_xxx](#)*

- #define **RTSER\_TIMEOUT\_INFINITE** RTDM\_TIMEOUT\_INFINITE
- #define **RTSER\_TIMEOUT\_NONE** RTDM\_TIMEOUT\_NONE
- #define **RTSER\_DEF\_TIMEOUT** RTDM\_TIMEOUT\_INFINITE

#### **RTSER\_xxx\_TIMESTAMP\_HISTORY**

*Timestamp history control*

- #define **RTSER\_RX\_TIMESTAMP\_HISTORY** 0x01
- #define **RTSER\_DEF\_TIMESTAMP\_HISTORY** 0x00

#### **RTSER\_EVENT\_xxx**

*Events bits*

- #define **RTSER\_EVENT\_RXPEND** 0x01
- #define **RTSER\_EVENT\_ERRPEND** 0x02
- #define **RTSER\_EVENT\_MODEMHI** 0x04
- #define **RTSER\_EVENT\_MODEMLO** 0x08
- #define **RTSER\_DEF\_EVENT\_MASK** 0x00

#### **RTSER\_SET\_xxx**

*Configuration mask bits*

- #define **RTSER\_SET\_BAUD** 0x0001
- #define **RTSER\_SET\_PARITY** 0x0002
- #define **RTSER\_SET\_DATA\_BITS** 0x0004
- #define **RTSER\_SET\_STOP\_BITS** 0x0008
- #define **RTSER\_SET\_HANDSHAKE** 0x0010
- #define **RTSER\_SET\_FIFO\_DEPTH** 0x0020
- #define **RTSER\_SET\_TIMEOUT\_RX** 0x0100
- #define **RTSER\_SET\_TIMEOUT\_TX** 0x0200
- #define **RTSER\_SET\_TIMEOUT\_EVENT** 0x0400
- #define **RTSER\_SET\_TIMESTAMP\_HISTORY** 0x0800
- #define **RTSER\_SET\_EVENT\_MASK** 0x1000

#### **RTSER\_LSR\_xxx**

*Line status bits*

- #define **RTSER\_LSR\_DATA** 0x01
- #define **RTSER\_LSR\_OVERRUN\_ERR** 0x02
- #define **RTSER\_LSR\_PARITY\_ERR** 0x04
- #define **RTSER\_LSR\_FRAMING\_ERR** 0x08
- #define **RTSER\_LSR\_BREAK\_IND** 0x10
- #define **RTSER\_LSR\_THR\_EMPTY** 0x20
- #define **RTSER\_LSR\_TRANSM\_EMPTY** 0x40
- #define **RTSER\_LSR\_FIFO\_ERR** 0x80
- #define **RTSER\_SOFT\_OVERRUN\_ERR** 0x0100

**RTSER\_MSR\_xxx***Modem status bits*

- #define **RTSER\_MSR\_DCTS** 0x01
- #define **RTSER\_MSR\_DDSR** 0x02
- #define **RTSER\_MSR\_TERI** 0x04
- #define **RTSER\_MSR\_DDCD** 0x08
- #define **RTSER\_MSR\_CTS** 0x10
- #define **RTSER\_MSR\_DSR** 0x20
- #define **RTSER\_MSR\_RI** 0x40
- #define **RTSER\_MSR\_DCD** 0x80

**RTSER\_MCR\_xxx***Modem control bits*

- #define **RTSER\_MCR\_DTR** 0x01
- #define **RTSER\_MCR\_RTS** 0x02
- #define **RTSER\_MCR\_OUT1** 0x04
- #define **RTSER\_MCR\_OUT2** 0x08
- #define **RTSER\_MCR\_LOOP** 0x10

**Sub-Classes of RTDM\_CLASS\_SERIAL**

- #define **RTDM\_SUBCLASS\_16550A** 0

**IOCTLs***Serial device IOCTLs*

- #define **RTSER\_RTIOC\_GET\_CONFIG** \_IOR(RTIOC\_TYPE\_SERIAL, 0x00, struct rtser\_config)  
*Get serial device configuration.*
- #define **RTSER\_RTIOC\_SET\_CONFIG** \_IOW(RTIOC\_TYPE\_SERIAL, 0x01, struct rtser\_config)  
*Set serial device configuration.*
- #define **RTSER\_RTIOC\_GET\_STATUS** \_IOR(RTIOC\_TYPE\_SERIAL, 0x02, struct rtser\_status)  
*Get serial device status.*
- #define **RTSER\_RTIOC\_GET\_CONTROL** \_IOR(RTIOC\_TYPE\_SERIAL, 0x03, int)  
*Get serial device's modem control register.*
- #define **RTSER\_RTIOC\_SET\_CONTROL** \_IOW(RTIOC\_TYPE\_SERIAL, 0x04, int)  
*Set serial device's modem control register.*
- #define **RTSER\_RTIOC\_WAIT\_EVENT** \_IOR(RTIOC\_TYPE\_SERIAL, 0x05, struct rtser\_event)  
*Wait on serial device events according to previously set mask.*

## RTSER\_BREAK\_xxx

Break control

- `#define RTSER_BREAK_CLR 0x00`  
*Serial device configuration.*
- `#define RTSER_BREAK_SET 0x01`  
*Serial device configuration.*
- `#define RTIOC_TYPE_SERIAL RTDM_CLASS_SERIAL`  
*Serial device configuration.*
- `typedef struct rtser_config rtser_config_t`  
*Serial device configuration.*
- `typedef struct rtser_status rtser_status_t`  
*Serial device status.*
- `typedef struct rtser_event rtser_event_t`  
*Additional information about serial device events.*

### 7.20.1 Detailed Description

Real-Time Driver Model for Xenomai, serial device profile header.

#### Note

Copyright (C) 2005-2007 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

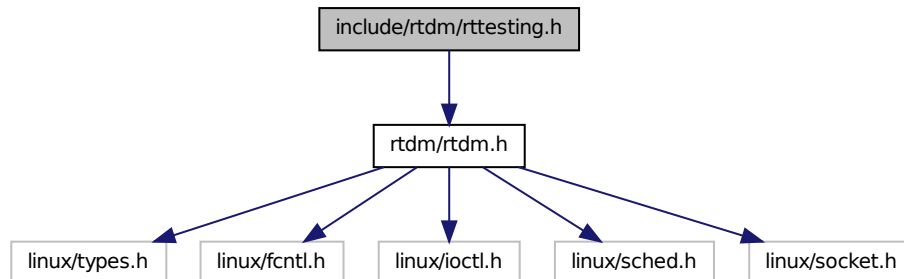
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.21 include/rtdm/rtesting.h File Reference

Real-Time Driver Model for Xenomai, testing device profile header.

Include dependency graph for rttesting.h:



## Defines

### Sub-Classes of RTDM\_CLASS\_TESTING

- #define [RTDM\\_SUBCLASS\\_TIMERBENCH](#) 0  
    *subclass name: "timerbench"*
- #define [RTDM\\_SUBCLASS\\_IRQBENCH](#) 1  
    *subclass name: "irqbench"*
- #define [RTDM\\_SUBCLASS\\_SWITCHTEST](#) 2  
    *subclass name: "switchtest"*
- #define [RTDM\\_SUBCLASS\\_RTDMTTEST](#) 3  
    *subclass name: "rtdm"*

## IOCTLs

### Testing device IOCTLs

- #define [RTTST\\_RTIOC\\_INTERM\\_BENCH\\_RES](#) \_IOWR(RTIOC\_TYPE\_TESTING, 0x00, struct rttst\_interm\_bench\_res)
- #define [RTTST\\_RTIOC\\_TMBENCH\\_START](#) \_IOW(RTIOC\_TYPE\_TESTING, 0x10, struct rttst\_tmbench\_config)
- #define [RTTST\\_RTIOC\\_TMBENCH\\_STOP](#) \_IOWR(RTIOC\_TYPE\_TESTING, 0x11, struct rttst\_overall\_bench\_res)
- #define [RTTST\\_RTIOC\\_IRQBENCH\\_START](#) \_IOW(RTIOC\_TYPE\_TESTING, 0x20, struct rttst\_irqbench\_config)
- #define [RTTST\\_RTIOC\\_IRQBENCH\\_STOP](#) \_IO(RTIOC\_TYPE\_TESTING, 0x21)
- #define [RTTST\\_RTIOC\\_IRQBENCH\\_GET\\_STATS](#) \_IOR(RTIOC\_TYPE\_TESTING, 0x22, struct rttst\_irqbench\_stats)
- #define [RTTST\\_RTIOC\\_IRQBENCH\\_WAIT\\_IRQ](#) \_IO(RTIOC\_TYPE\_TESTING, 0x23)
- #define [RTTST\\_RTIOC\\_IRQBENCH\\_REPLY\\_IRQ](#) \_IO(RTIOC\_TYPE\_TESTING, 0x24)
- #define [RTTST\\_RTIOC\\_SWTEST\\_SET\\_TASKS\\_COUNT](#) \_IOW(RTIOC\_TYPE\_TESTING, 0x30, unsigned long)

- `#define RTTST_RTIOC_SWTEST_SET_CPU _IOW(RTIOC_TYPE_TESTING, 0x31, unsigned long)`
- `#define RTTST_RTIOC_SWTEST_REGISTER_UTASK _IOW(RTIOC_TYPE_TESTING, 0x32, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_CREATE_KTASK _IOWR(RTIOC_TYPE_TESTING, 0x33, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_PEND _IOR(RTIOC_TYPE_TESTING, 0x34, struct rttst_swtest_task)`
- `#define RTTST_RTIOC_SWTEST_SWITCH_TO _IOR(RTIOC_TYPE_TESTING, 0x35, struct rttst_swtest_dir)`
- `#define RTTST_RTIOC_SWTEST_GET_SWITCHES_COUNT _IOR(RTIOC_TYPE_TESTING, 0x36, unsigned long)`
- `#define RTTST_RTIOC_SWTEST_GET_LAST_ERROR _IOR(RTIOC_TYPE_TESTING, 0x37, struct rttst_swtest_error)`
- `#define RTTST_RTIOC_SWTEST_SET_PAUSE _IOW(RTIOC_TYPE_TESTING, 0x38, unsigned long)`
- `#define RTTST_RTIOC_RTDM_DEFER_CLOSE _IOW(RTIOC_TYPE_TESTING, 0x40, unsigned long)`

### 7.21.1 Detailed Description

Real-Time Driver Model for Xenomai, testing device profile header.

#### Note

Copyright (C) 2005 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

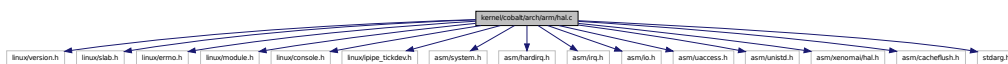
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.22 kernel/cobalt/arch/arm/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for ARM.

Include dependency graph for hal.c:



### Functions

- `int rthal_timer_request(void(*tick_handler)(void), void(*mode_emul)(enum clock_event_mode mode, struct clock_event_device *cdev), int(*tick_emul)(unsigned long delay, struct`

clock\_event\_device \*cdev), int cpu)

*Grab the hardware timer.*

- void [rthal\\_timer\\_release](#) (int cpu)

*Release the hardware timer.*

### 7.22.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for ARM. ARM port Copyright (C) 2005 Stelian Pop

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

ARM-specific HAL services.

### 7.22.2 Function Documentation

#### 7.22.2.1 void [rthal\\_timer\\_release](#) ( int *cpu* )

Release the hardware timer.

Releases the hardware timer, thus reverting the effect of a previous call to [rthal\\_timer\\_request\(\)](#). In case the timer hardware is shared with Linux, a periodic setup suitable for the Linux kernel will be reset.

#### Parameters

*cpu* The CPU number the timer was grabbed from.

Environments:

This service can be called from:

- Linux domain context.

#### 7.22.2.2 int [rthal\\_timer\\_request](#) ( void(\*) (void) *tick\_handler*, void(\*) (enum clock\_event\_mode mode, struct clock\_event\_device \*cdev) *mode\_emul*, int(\*) (unsigned long delay, struct clock\_event\_device \*cdev) *tick\_emul*, int *cpu* )

Grab the hardware timer.

[rthal\\_timer\\_request\(\)](#) grabs and tunes the hardware timer in oneshot mode in order to clock the master time base. GENERIC\_CLOCKEVENTS is required from the host kernel.

A user-defined routine is registered as the clock tick handler. This handler will always be invoked on behalf of the Xenomai domain for each incoming tick.

Host tick emulation is a way to share the clockchip hardware between Linux and Xenomai, when the former provides support for oneshot timing (i.e. high resolution timers and no-HZ scheduler ticking).

### Parameters

*tick\_handler* The address of the Xenomai tick handler which will process each incoming tick.

*mode\_emul* The optional address of a callback to be invoked upon mode switch of the host tick device, notified by the Linux kernel.

*tick\_emul* The optional address of a callback to be invoked upon setup of the next shot date for the host tick device, notified by the Linux kernel.

*cpu* The CPU number to grab the timer from.

### Returns

a positive value is returned on success, representing the duration of a Linux periodic tick expressed as a count of nanoseconds; zero should be returned when the Linux kernel does not undergo periodic timing on the given CPU (e.g. oneshot mode). Otherwise:

- -EBUSY is returned if the hardware timer has already been grabbed. [rthal\\_timer\\_request\(\)](#) must be issued before [rthal\\_timer\\_request\(\)](#) is called again.
- -ENODEV is returned if the hardware timer cannot be used. This situation may occur after the kernel disabled the timer due to invalid calibration results; in such a case, such hardware is unusable for any timing duties.

Environments:

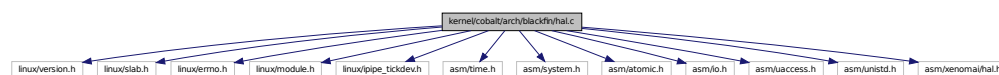
This service can be called from:

- Linux domain context.

## 7.23 kernel/cobalt/arch/blackfin/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.

Include dependency graph for hal.c:



### 7.23.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the Blackfin architecture. Copyright (C) 2005-2006 Philippe Gerum.



Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

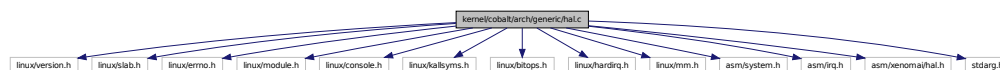
You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Blackfin-specific HAL services.

## 7.24 kernel/cobalt/arch/generic/hal.c File Reference

Generic Real-Time HAL.

Include dependency graph for hal.c:



### Functions

- `int rthal_apc_alloc` (const char \*name, void(\*handler)(void \*cookie), void \*cookie)  
*Allocate an APC slot.*
- `void rthal_apc_free` (int apc)  
*Releases an APC slot.*

#### 7.24.1 Detailed Description

Generic Real-Time HAL. Copyright ©2005 Philippe Gerum.

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

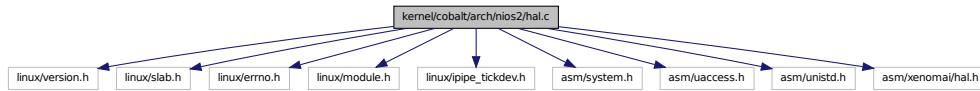
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.25 kernel/cobalt/arch/nios2/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture.

Include dependency graph for hal.c:



### 7.25.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture. Copyright (C) 2009 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

NIOS2-specific HAL services.

## 7.26 kernel/cobalt/arch/powerpc/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for PowerPC.

Include dependency graph for hal.c:



### 7.26.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for PowerPC. Copyright (C) 2004-2006 Philippe Gerum.

64-bit PowerPC adoption copyright (C) 2005 Taneli Vähäkangas and Heikki Lindholm

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

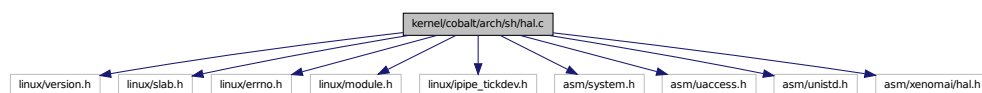
You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

PowerPC-specific HAL services.

## 7.27 kernel/cobalt/arch/sh/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the SuperH architecture.

Include dependency graph for hal.c:



### 7.27.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the SuperH architecture. Copyright (C) 2011 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

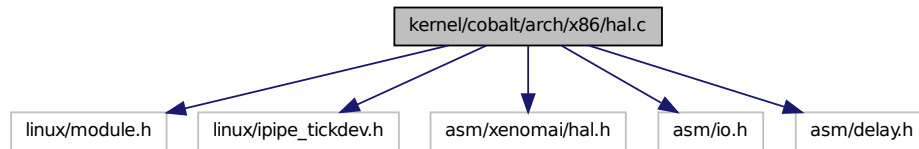
You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

SuperH-specific HAL services.

## 7.28 kernel/cobalt/arch/x86/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for x86.

Include dependency graph for hal.c:



### 7.28.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86. Common code of i386 and x86\_64.

Copyright (C) 2007-2012 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

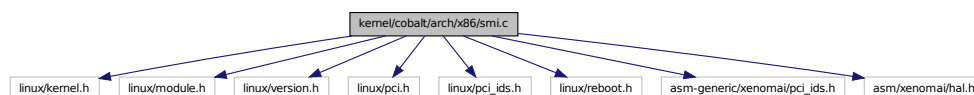
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.29 kernel/cobalt/arch/x86/smi.c File Reference

SMI workaround for x86.

Include dependency graph for smi.c:



### 7.29.1 Detailed Description

SMI workaround for x86. Cut/Pasted from Vitor Angelo "smi" module. Adapted by Gilles Chanteperdrix <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>.

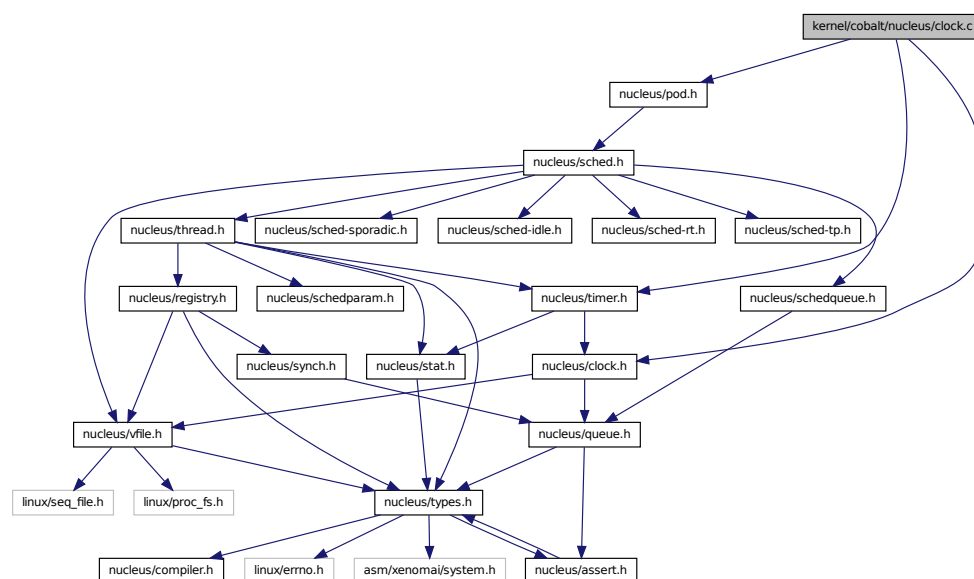
This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.30 kernel/cobalt/nucleus/clock.c File Reference

Include dependency graph for clock.c:



## Functions

- void [xnclock\\_adjust](#) (xnsticks\_t delta)

*Adjust the clock time for the system.*

### 7.30.1 Detailed Description

#### Note

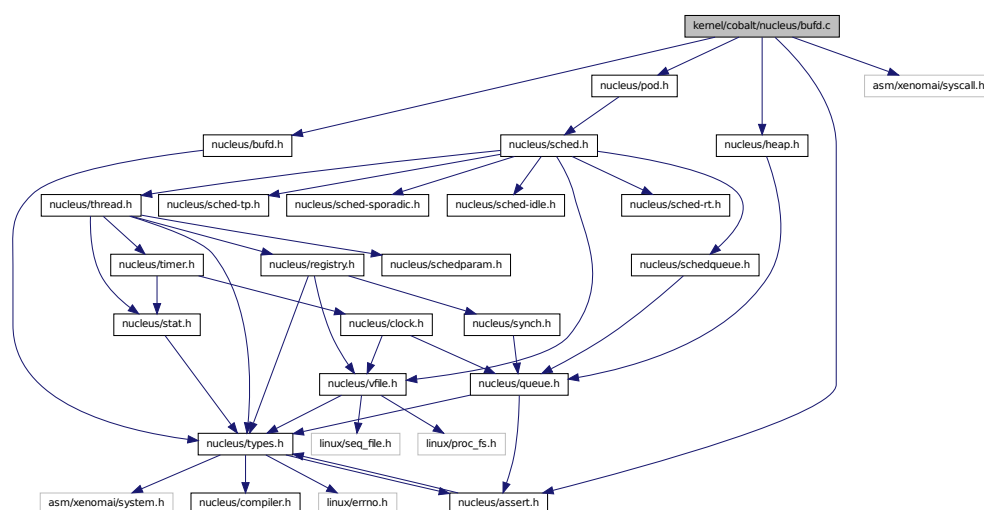
Copyright (C) 2006-2011 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.



## 7.32 kernel/cobalt/nucleus/bufd.c File Reference

Include dependency graph for bufd.c:



## Functions

- `ssize_t xnbufd_copy_to_kmem (void *ptr, struct xnbufd *bufd, size_t len)`  
Copy memory covered by a buffer descriptor to kernel memory.
- `ssize_t xnbufd_copy_from_kmem (struct xnbufd *bufd, void *from, size_t len)`  
Copy kernel memory to the area covered by a buffer descriptor.
- `ssize_t xnbufd_unmap_uread (struct xnbufd *bufd)`  
Finalize a buffer descriptor obtained from `xnbufd_map_uread()`.
- `ssize_t xnbufd_unmap_uwrite (struct xnbufd *bufd)`  
Finalize a buffer descriptor obtained from `xnbufd_map_uwrite()`.
- `void xnbufd_invalidate (struct xnbufd *bufd)`  
Invalidate a buffer descriptor.
- `ssize_t xnbufd_unmap_kread (struct xnbufd *bufd)`  
Finalize a buffer descriptor obtained from `xnbufd_map_kread()`.
- `ssize_t xnbufd_unmap_kwrite (struct xnbufd *bufd)`  
Finalize a buffer descriptor obtained from `xnbufd_map_kwrite()`.

### 7.32.1 Detailed Description

#### Note

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

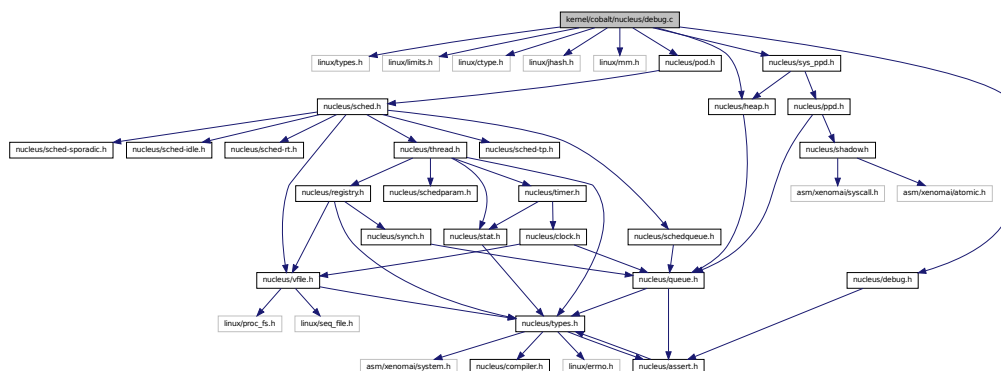
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.33 kernel/cobalt/nucleus/debug.c File Reference

Debug services.

Include dependency graph for debug.c:



### 7.33.1 Detailed Description

Debug services.

#### Author

Philippe Gerum

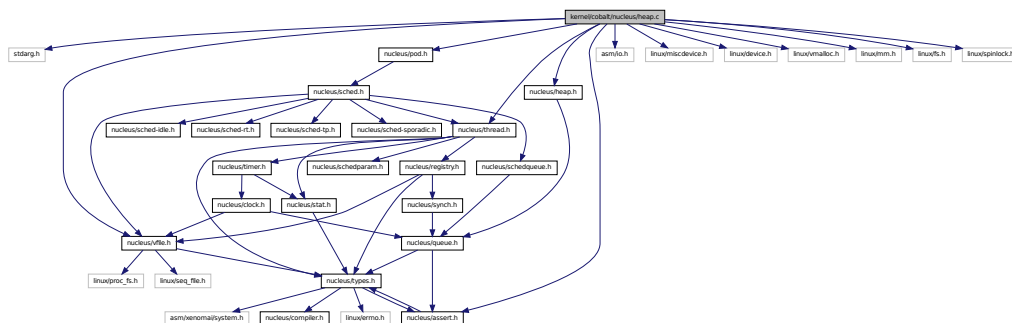
Copyright (C) 2010 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.



You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Dynamic memory allocation services.  
Include dependency graph for heap.c:



- int `xnheap_init` (xnheap\_t \*heap, void \*heapaddr, u\_long heapsize, u\_long pagesize)  
*Initialize a memory heap.*
- void `xnheap_set_label` (xnheap\_t \*heap, const char \*label,...)  
*Set the heap's label string.*
- void `xnheap_destroy` (xnheap\_t \*heap, void(\*flushfn)(xnheap\_t \*heap, void \*extaddr, u\_long extsize, void \*cookie), void \*cookie)  
*Destroys a memory heap.*
- void \* `xnheap_alloc` (xnheap\_t \*heap, u\_long size)  
*Allocate a memory block from a memory heap.*
- int `xnheap_test_and_free` (xnheap\_t \*heap, void \*block, int(\*ckfn)(void \*block))  
*Test and release a memory block to a memory heap.*
- int `xnheap_free` (xnheap\_t \*heap, void \*block)  
*Release a memory block to a memory heap.*
- int `xnheap_extend` (xnheap\_t \*heap, void \*extaddr, u\_long extsize)  
*Extend a memory heap.*

- void [xnheap\\_schedule\\_free](#) (xnheap\_t \*heap, void \*block, xnholder\_t \*link)

*Schedule a memory block for release.*

### 7.34.1 Detailed Description

Dynamic memory allocation services.

#### Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

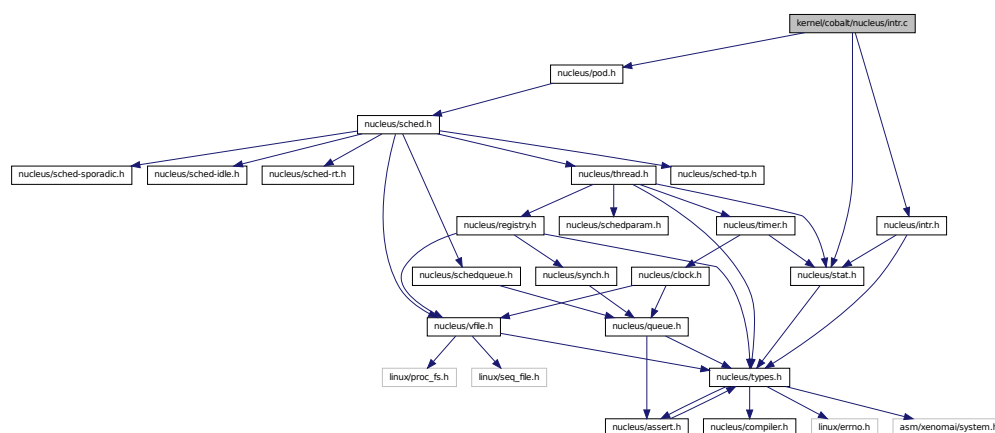
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.35 kernel/cobalt/nucleus/intr.c File Reference

Interrupt management.

Include dependency graph for intr.c:



## Functions

- int [xnintr\\_init](#) (xnintr\_t \*intr, const char \*name, unsigned irq, xnintr\_t isr, xniack\_t iack, xnflags\_t flags)  
*Initialize an interrupt object.*
- int [xnintr\\_destroy](#) (xnintr\_t \*intr)  
*Destroy an interrupt object.*
- int [xnintr\\_attach](#) (xnintr\_t \*intr, void \*cookie)  
*Attach an interrupt object.*
- int [xnintr\\_detach](#) (xnintr\_t \*intr)  
*Detach an interrupt object.*
- void [xnintr\\_enable](#) (xnintr\_t \*intr)  
*Enable an interrupt object.*
- void [xnintr\\_disable](#) (xnintr\_t \*intr)  
*Disable an interrupt object.*
- void [xnintr\\_affinity](#) (xnintr\_t \*intr, xnarch\_cpumask\_t cpumask)  
*Set interrupt's processor affinity.*

### 7.35.1 Detailed Description

Interrupt management.

#### Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2005,2006 Dmitry Adamushko <[dmitry.adamushko@gmail.com](mailto:dmitry.adamushko@gmail.com)>. Copyright (C) 2007 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>.

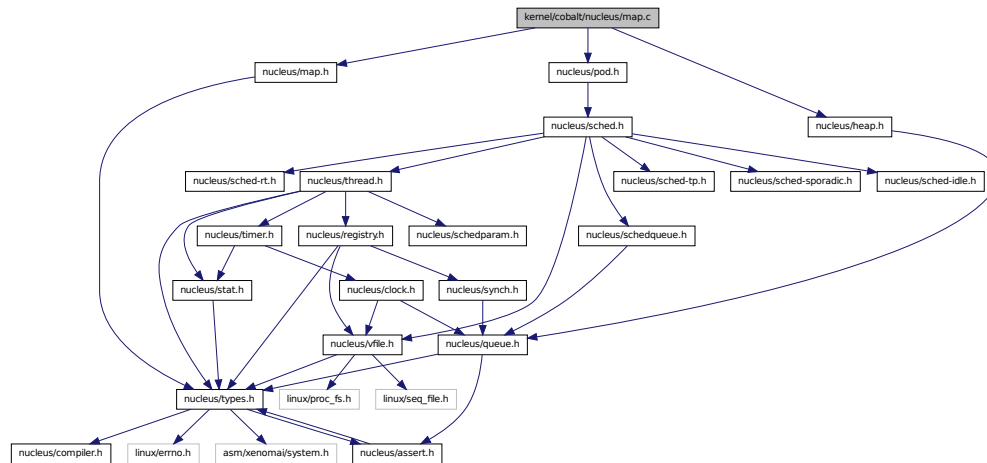
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.36 kernel/cobalt/nucleus/map.c File Reference

Include dependency graph for map.c:



### Functions

- `xnmap_t * xnmap\_create (int nkeys, int reserve, int offset)`  
*Create a map.*
- `void xnmap\_delete (xnmap_t *map)`  
*Delete a map.*
- `int xnmap\_enter (xnmap_t *map, int key, void *objaddr)`  
*Index an object into a map.*
- `int xnmap\_remove (xnmap_t *map, int key)`  
*Remove an object reference from a map.*

### 7.36.1 Detailed Description

#### Note

Copyright (C) 2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

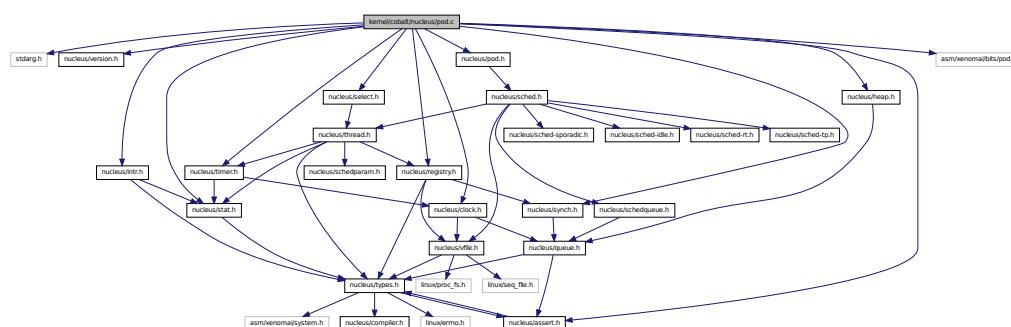
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.37 kernel/cobalt/nucleus/pod.c File Reference

Real-time pod services.

Include dependency graph for pod.c:



## Functions

- `int xnpod_init (void)`  
*Initialize the core pod.*
- `void xnpod_shutdown (int xtype)`  
*Shutdown the current pod.*
- `int xnpod_init_thread (struct xnthread *thread, const struct xnthread_init_attr *attr, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)`  
*Initialize a new thread.*
- `int xnpod_start_thread (xnthread_t *thread, const struct xnthread_start_attr *attr)`  
*Initial start of a newly created thread.*
- `void __xnpod_reset_thread (struct xnthread *thread)`  
*Reset the thread.*
- `void xnpod_stop_thread (xnthread_t *thread)`  
*Stop a thread.*
- `xnflags_t xnpod_set_thread_mode (xnthread_t *thread, xnflags_t clrmask, xnflags_t set-mask)`  
*Change a thread's control mode.*
- `void xnpod_delete_thread (xnthread_t *thread)`

*Delete a thread.*

- void [xn timer abort thread](#) (xn timer\_t \*thread)  
*Abort a thread.*
- void [xn timer suspend thread](#) (xn timer\_t \*thread, xn timer\_flags\_t mask, xn timer\_ticks\_t timeout, xn timer\_mode\_t timeout\_mode, struct xn timer\_synch \*wchan)  
*Suspend a thread.*
- void [xn timer resume thread](#) (xn timer\_t \*thread, xn timer\_flags\_t mask)  
*Resume a thread.*
- int [xn timer unblock thread](#) (xn timer\_t \*thread)  
*Unblock a thread.*
- int [xn timer set thread sched param](#) (struct xn timer \*thread, struct xn timer\_sched\_class \*sched\_class, const union xn timer\_sched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*
- int [xn timer migrate thread](#) (int cpu)  
*Migrate the current thread.*
- void [xn timer dispatch signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- void [xn timer welcome thread](#) (xn timer\_t \*thread, int imask)  
*Thread prologue.*
- int [xn timer add hook](#) (int type, void(\*routine)(xn timer\_t \*))  
*Install a nucleus hook.*
- int [xn timer remove hook](#) (int type, void(\*routine)(xn timer\_t \*))  
*Remove a nucleus hook.*
- int [xn timer handle exception](#) (struct ipipe\_trap\_data \*d)  
*Exception handler.*
- int [xn timer enable timesource](#) (void)  
*Activate the core time source.*
- void [xn timer disable timesource](#) (void)  
*Stop the core time source.*
- int [xn timer set thread periodic](#) (xn timer\_t \*thread, xn timer\_ticks\_t idate, xn timer\_mode\_t timeout\_mode, xn timer\_ticks\_t period)  
*Make a thread periodic.*
- int [xn timer wait thread period](#) (unsigned long \*overruns\_r)  
*Wait for the next periodic release point.*

- int [xnpod\\_set\\_thread\\_tslice](#) (struct xnthread \*thread, xnticks\_t quantum)

*Set thread time-slicing information.*

### 7.37.1 Detailed Description

Real-time pod services.

#### Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.Xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

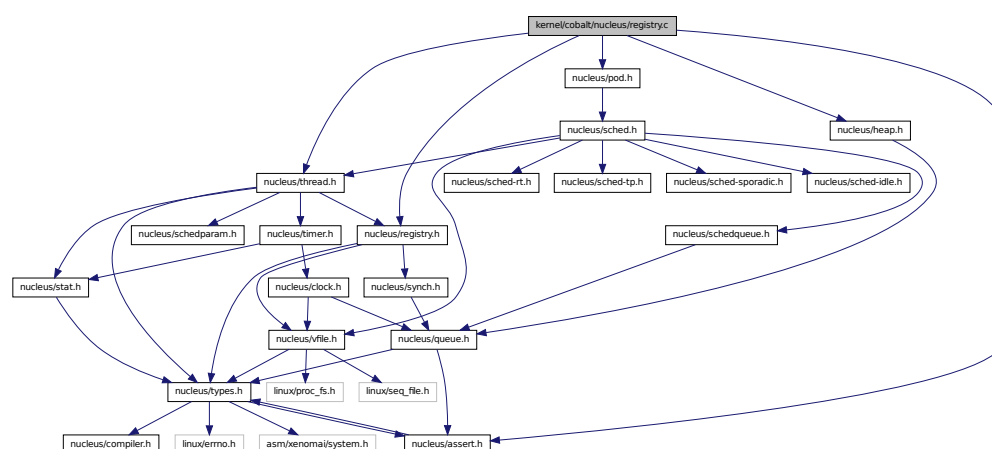
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.38 kernel/cobalt/nucleus/registry.c File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.c:



## Functions

- int [xnregistry\\_enter](#) (const char \*key, void \*objaddr, xnhandle\_t \*phandle, struct xninode \*pnode)  
*Register a real-time object.*
- int [xnregistry\\_bind](#) (const char \*key, xnticks\_t timeout, int timeout\_mode, xnhandle\_t \*phandle)  
*Bind to a real-time object.*
- int [xnregistry\\_remove](#) (xnhandle\_t handle)  
*Forcibly unregister a real-time object.*
- int [xnregistry\\_remove\\_safe](#) (xnhandle\_t handle, xnticks\_t timeout)  
*Unregister an idle real-time object.*
- void \* [xnregistry\\_get](#) (xnhandle\_t handle)  
*Find and lock a real-time object into the registry.*
- u\_long [xnregistry\\_put](#) (xnhandle\_t handle)  
*Unlock a real-time object from the registry.*
- void \* [xnregistry\\_fetch](#) (xnhandle\_t handle)  
*Find a real-time object into the registry.*

### 7.38.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2004 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

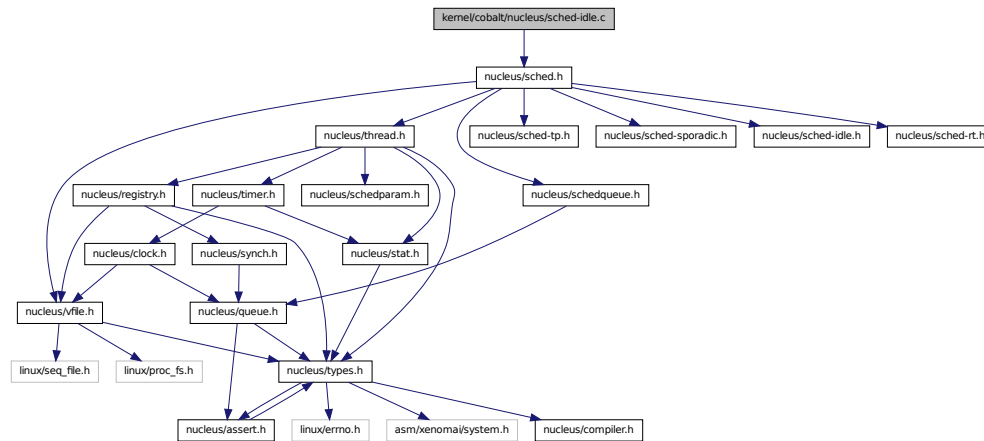
You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.39 kernel/cobalt/nucleus/sched-idle.c File Reference

Idle scheduling class implementation (i.e. Linux placeholder).



Include dependency graph for sched-idle.c:



### 7.39.1 Detailed Description

Idle scheduling class implementation (i.e. Linux placeholder).

#### Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

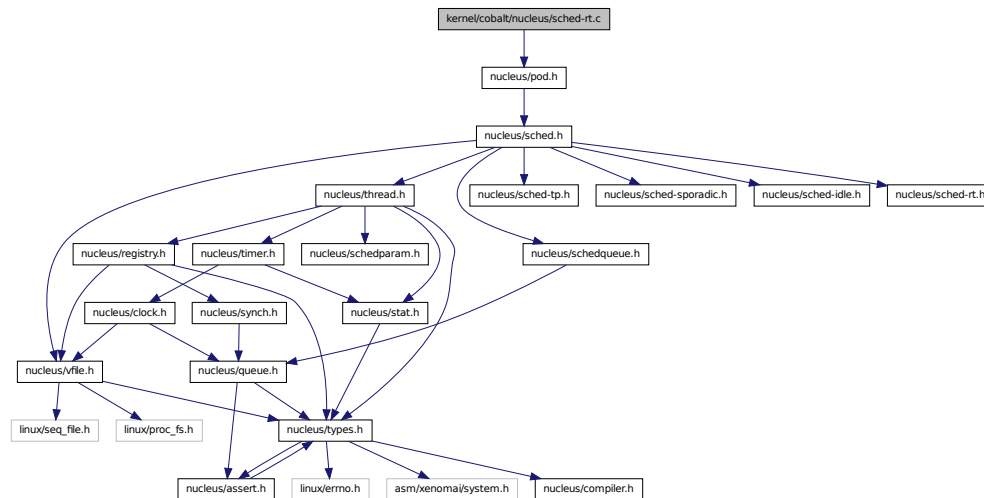
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.40 kernel/cobalt/nucleus/sched-rt.c File Reference

Common real-time scheduling class implementation (FIFO + RR).

Include dependency graph for sched-rt.c:



### 7.40.1 Detailed Description

Common real-time scheduling class implementation (FIFO + RR).

#### Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

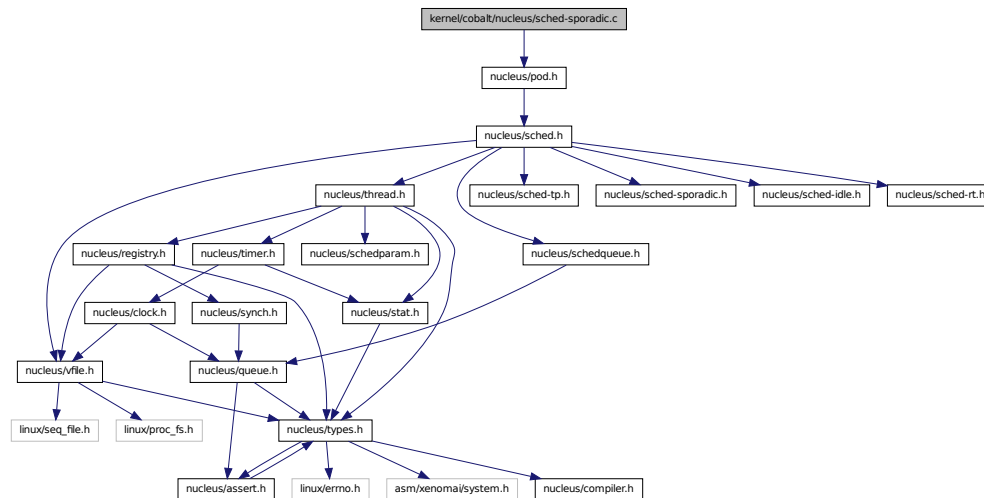
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.41 kernel/cobalt/nucleus/sched-sporadic.c File Reference

POSIX SCHED\_SPORADIC scheduling class.

Include dependency graph for sched-sporadic.c:



### 7.41.1 Detailed Description

POSIX SCHED\_SPORADIC scheduling class.

#### Author

Philippe Gerum Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

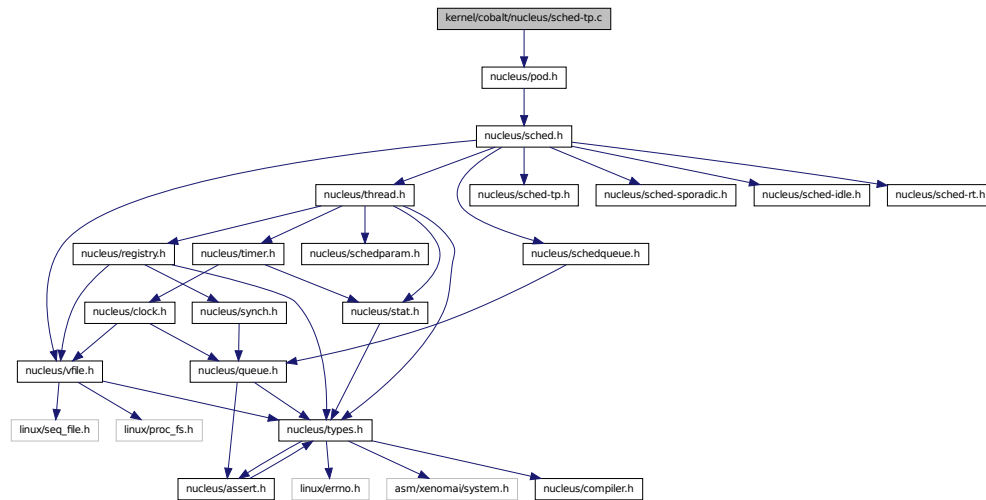
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.42 kernel/cobalt/nucleus/sched-tp.c File Reference

Temporal partitioning (typical of IMA systems).

Include dependency graph for sched-tp.c:



### 7.42.1 Detailed Description

Temporal partitioning (typical of IMA systems).

#### Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

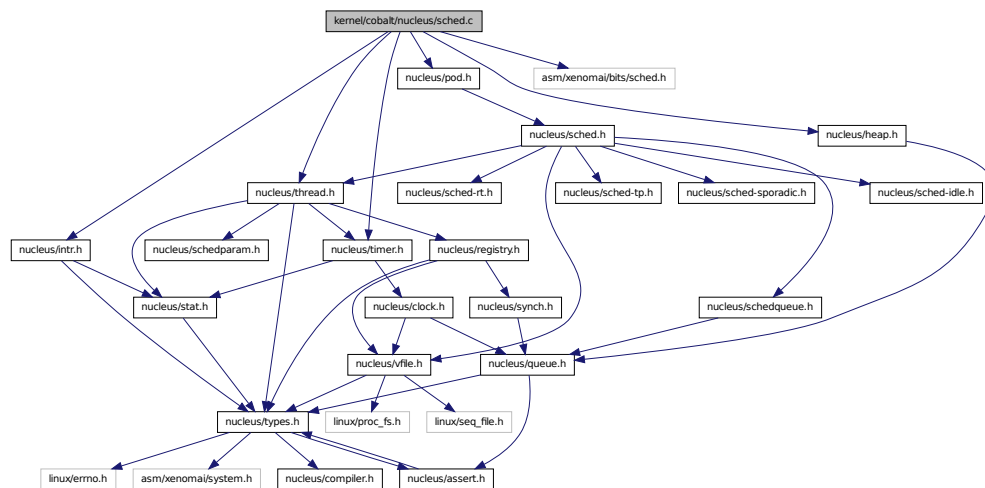
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.43 kernel/cobalt/nucleus/sched.c File Reference

Include dependency graph for sched.c:



### 7.43.1 Detailed Description

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

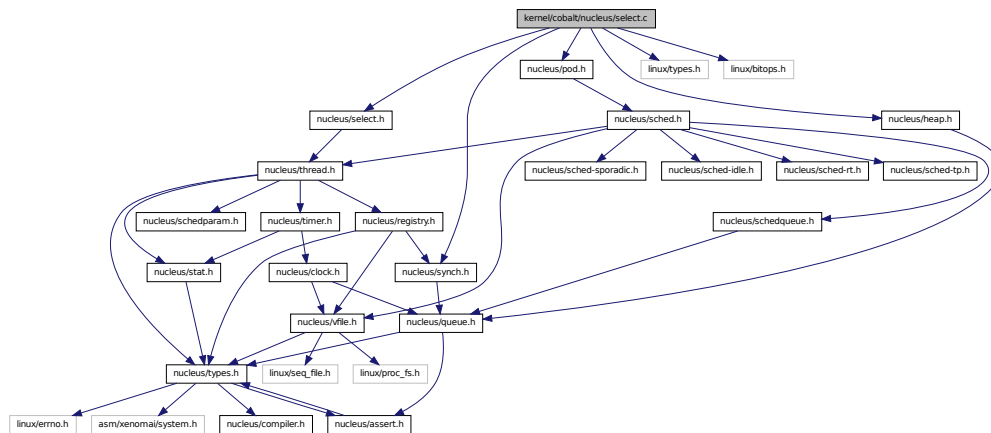
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.44 kernel/cobalt/nucleus/select.c File Reference

file descriptors events multiplexing.

Include dependency graph for select.c:



## Functions

- void [xnselect\\_init](#) (struct xnselect \*select\_block)  
*Initialize a struct xnselect structure.*
- int [xnselect\\_bind](#) (struct xnselect \*select\_block, struct xnselect\_binding \*binding, struct xnselector \*selector, unsigned type, unsigned index, unsigned state)  
*Bind a file descriptor (represented by its xnselect structure) to a selector block.*
- void [xnselect\\_destroy](#) (struct xnselect \*select\_block)  
*Destroy the xnselect structure associated with a file descriptor.*
- int [xnselector\\_init](#) (struct xnselector \*selector)  
*Initialize a selector structure.*
- int [xnselect](#) (struct xnselector \*selector, fd\_set \*out\_fds[XNSELECT\_MAX\_TYPES], fd\_set \*in\_fds[XNSELECT\_MAX\_TYPES], int nfds, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.*
- void [xnselector\\_destroy](#) (struct xnselector \*selector)  
*Destroy a selector block.*

### 7.44.1 Detailed Description

file descriptors events multiplexing.

#### Author

Gilles Chantepedrix

Copyright (C) 2008 Efixo <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

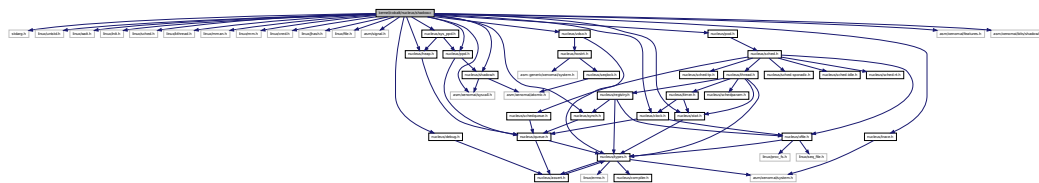
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.45 kernel/cobalt/nucleus/shadow.c File Reference

Real-time shadow services.

Include dependency graph for shadow.c:



### Functions

- `int xnshadow_harden` (void)  
*Migrate a Linux task to the Xenomai domain.*
- `void xnshadow_relax` (int notify, int reason)  
*Switch a shadow thread back to the Linux domain.*
- `int xnshadow_map` (xnthread\_t \*thread, xncompletion\_t \_\_user \*u\_completion, unsigned long \_\_user \*u\_window\_offset)  
*Create a shadow thread context.*
- `xnshadow_ppd_t * xnshadow_ppd_get` (unsigned int muxid)  
*Return the per-process data attached to the calling process.*

### 7.45.1 Detailed Description

Real-time shadow services.

#### Author

Philippe Gerum

Copyright (C) 2001-2012 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.xenomai.org>> Copyright (C) 2006 Gilles Chantepedrix <[gilles.chantepedrix@xenomai.org](mailto:gilles.chantepedrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

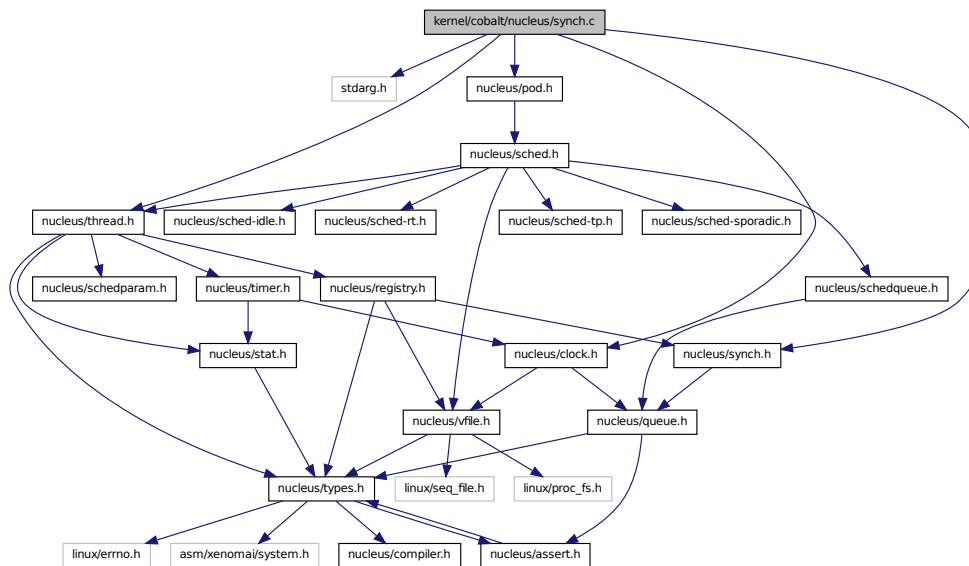
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.46 kernel/cobalt/nucleus/synch.c File Reference

Thread synchronization services.

Include dependency graph for synch.c:



### Functions

- void [xnsynch\\_init](#) (struct xnsynch \*synch, xnflags\_t flags, xnarch\_atomic\_t \*fastlock)  
*Initialize a synchronization object.*
- xnflags\_t [xnsynch\\_sleep\\_on](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)



*Sleep on an ownerless synchronization object.*

- struct xnthread \* [xnsynch\\_wakeup\\_one\\_sleeper](#) (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*
- struct xnpholder \* [xnsynch\\_wakeup\\_this\\_sleeper](#) (struct xnsynch \*synch, struct xnpholder \*holder)  
*Give the resource ownership to a given waiting thread.*
- xnflags\_t [xnsynch\\_acquire](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Acquire the ownership of a synchronization object.*
- static void [xnsynch\\_clear\\_boost](#) (struct xnsynch \*synch, struct xnthread \*owner)  
*Clear the priority boost.*
- void [xnsynch\\_requeue\\_sleeper](#) (struct xnthread \*thread)  
*Change a sleeper's priority.*
- struct xnthread \* [xnsynch\\_peek\\_pendq](#) (struct xnsynch \*synch)  
*Access the thread leading a synch object wait queue.*
- int [xnsynch\\_flush](#) (struct xnsynch \*synch, xnflags\_t reason)  
*Unblock all waiters pending on a resource.*
- void [xnsynch\\_forget\\_sleeper](#) (struct xnthread \*thread)  
*Abort a wait for a resource.*
- void [xnsynch\\_release\\_all\\_ownerships](#) (struct xnthread \*thread)  
*Release all ownerships.*

## 7.46.1 Detailed Description

Thread synchronization services.

### Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

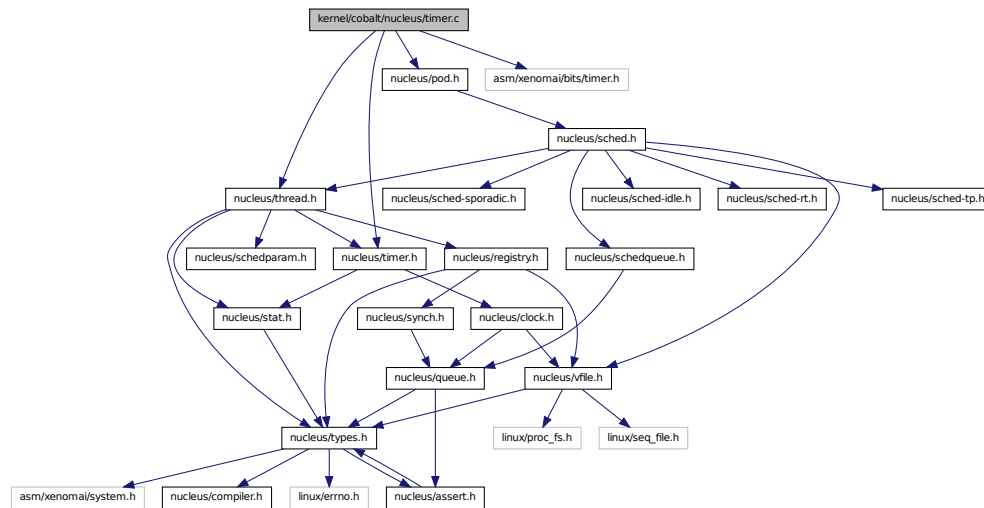
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.47 kernel/cobalt/nucleus/timer.c File Reference

Include dependency graph for timer.c:



### Functions

- `int xntimer_start(xntimer_t *timer, xnticks_t value, xnticks_t interval, xntmode_t mode)`  
*Arm a timer.*
- `xnticks_t xntimer_get_date(xntimer_t *timer)`  
*Return the absolute expiration date.*
- `xnticks_t xntimer_get_timeout(xntimer_t *timer)`  
*Return the relative expiration date.*
- `xnticks_t xntimer_get_interval(xntimer_t *timer)`  
*Return the timer interval value.*
- `void xntimer_tick(void)`  
*Process a timer tick.*
- `void xntimer_init(xntimer_t *timer, void(*handler)(xntimer_t *timer))`  
*Initialize a timer object.*
- `void xntimer_destroy(xntimer_t *timer)`  
*Release a timer object.*
- `unsigned long xntimer_get_overruns(xntimer_t *timer, xnticks_t now)`  
*Get the count of overruns for the last tick.*

- void `xntimer_freeze` (void)

*Freeze all timers (from every time bases).*

### 7.47.1 Detailed Description

#### Note

Copyright (C) 2001,2002,2003,2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 Gilles Chantepredrix <[gilles.chantepredrix@xenomai.org](mailto:gilles.chantepredrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

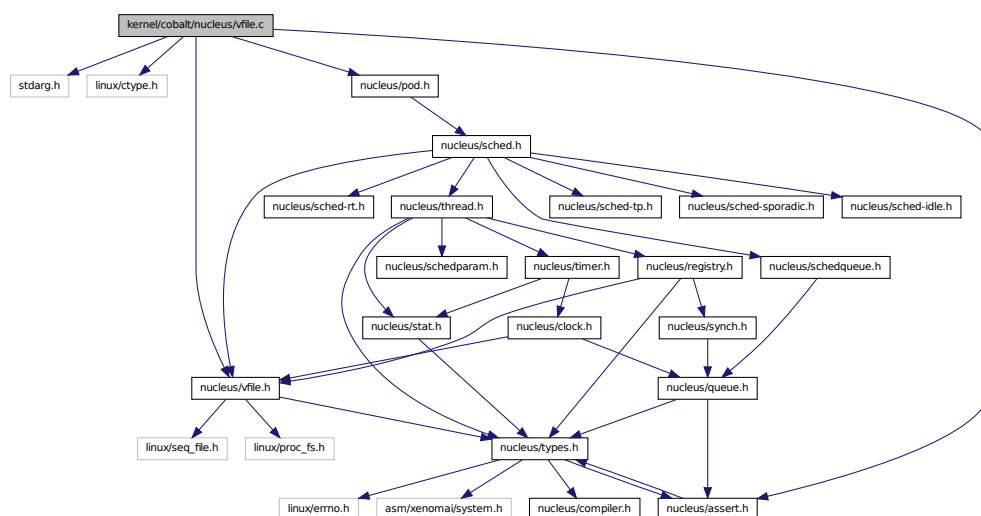
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.48 kernel/cobalt/nucleus/vfile.c File Reference

This file is part of the Xenomai project.

Include dependency graph for `vfile.c`:



## Functions

- int [xnvmfile\\_init\\_snapshot](#) (const char \*name, struct [xnvmfile\\_snapshot](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a snapshot-driven vfile.*
- int [xnvmfile\\_init\\_regular](#) (const char \*name, struct [xnvmfile\\_regular](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a regular vfile.*
- int [xnvmfile\\_init\\_dir](#) (const char \*name, struct [xnvmfile\\_directory](#) \*vdir, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual directory entry.*
- int [xnvmfile\\_init\\_link](#) (const char \*from, const char \*to, struct [xnvmfile\\_link](#) \*vlink, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual link entry.*
- void [xnvmfile\\_destroy](#) (struct [xnvmfile](#) \*vfile)  
*Removes a virtual file entry.*
- ssize\_t [xnvmfile\\_get\\_blob](#) (struct [xnvmfile\\_input](#) \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_string](#) (struct [xnvmfile\\_input](#) \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_integer](#) (struct [xnvmfile\\_input](#) \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct [xnvmfile\\_directory](#) [nkvfroot](#)  
*Xenomai vfile root directory.*

### 7.48.1 Detailed Description

This file is part of the Xenomai project.

#### Note

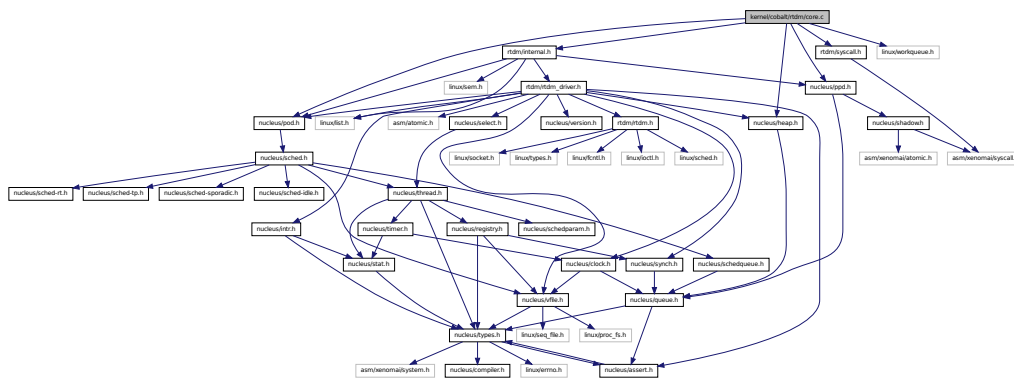
Copyright (C) 2010 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

## 7.49 kernel/cobalt/rtdm/core.c File Reference

Include dependency graph for core.c:



- struct `rtdm_dev_context * rtdm_context_get` (int fd)  
*Retrieve and lock a device context.*
- int `rtdm_select_bind` (int fd, rtdm\_selector\_t \*selector, enum `rtdm_selecttype` type, unsigned fd\_index)  
*Bind a selector to specified event types of a given file descriptor.*
- void `rtdm_context_lock` (struct `rtdm_dev_context *context`)  
*Increment context reference counter.*
- void `rtdm_context_unlock` (struct `rtdm_dev_context *context`)  
*Decrement context reference counter.*
- void `rtdm_context_put` (struct `rtdm_dev_context *context`)  
*Release a device context obtained via `rtdm_context_get()`.*
- int `rtdm_open` (const char \*path, int oflag,...)  
*Open a device.*
- int `rtdm_socket` (int protocol\_family, int socket\_type, int protocol)  
*Create a socket.*
- int `rtdm_close` (int fd)

*Close a device or socket.*

- `int rtdm_ioctl` (int fd, int request,...)  
*Issue an IOCTL.*
- `ssize_t rtdm_read` (int fd, void \*buf, size\_t nbyte)  
*Read from device.*
- `ssize_t rtdm_write` (int fd, const void \*buf, size\_t nbyte)  
*Write to device.*
- `ssize_t rtdm_recvmmsg` (int fd, struct msghdr \*msg, int flags)  
*Receive message from socket.*
- `ssize_t rtdm_recvfrom` (int fd, void \*buf, size\_t len, int flags, struct sockaddr \*from, socklen\_t \*fromlen)  
*Receive message from socket.*
- `ssize_t rtdm_recv` (int fd, void \*buf, size\_t len, int flags)  
*Receive message from socket.*
- `ssize_t rtdm_sendmsg` (int fd, const struct msghdr \*msg, int flags)  
*Transmit message to socket.*
- `ssize_t rtdm_sendto` (int fd, const void \*buf, size\_t len, int flags, const struct sockaddr \*to, socklen\_t tolen)  
*Transmit message to socket.*
- `ssize_t rtdm_send` (int fd, const void \*buf, size\_t len, int flags)  
*Transmit message to socket.*
- `int rtdm_bind` (int fd, const struct sockaddr \*my\_addr, socklen\_t addrlen)  
*Bind to local address.*
- `int rtdm_connect` (int fd, const struct sockaddr \*serv\_addr, socklen\_t addrlen)  
*Connect to remote address.*
- `int rtdm_listen` (int fd, int backlog)  
*Listen for incoming connection requests.*
- `int rtdm_accept` (int fd, struct sockaddr \*addr, socklen\_t \*addrlen)  
*Accept a connection requests.*
- `int rtdm_shutdown` (int fd, int how)  
*Shut down parts of a connection.*
- `int rtdm_getsockopt` (int fd, int level, int optname, void \*optval, socklen\_t \*optlen)  
*Get socket option.*
- `int rtdm_setsockopt` (int fd, int level, int optname, const void \*optval, socklen\_t optlen)

*Set socket option.*

- int [rtdm\\_getsockname](#) (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get local socket address.*
- int [rtdm\\_getpeername](#) (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get socket destination address.*
- int [rt\\_dev\\_open](#) (const char \*path, int oflag,...)  
*Open a device.*
- int [rt\\_dev\\_socket](#) (int protocol\_family, int socket\_type, int protocol)  
*Create a socket.*
- int [rt\\_dev\\_close](#) (int fd)  
*Close a device or socket.*
- int [rt\\_dev\\_ioctl](#) (int fd, int request,...)  
*Issue an IOCTL.*
- ssize\_t [rt\\_dev\\_read](#) (int fd, void \*buf, size\_t nbyte)  
*Read from device.*
- ssize\_t [rt\\_dev\\_write](#) (int fd, const void \*buf, size\_t nbyte)  
*Write to device.*
- ssize\_t [rt\\_dev\\_recvmsg](#) (int fd, struct msghdr \*msg, int flags)  
*Receive message from socket.*
- ssize\_t [rt\\_dev\\_recvfrom](#) (int fd, void \*buf, size\_t len, int flags, struct sockaddr \*from, socklen\_t \*fromlen)  
*Receive message from socket.*
- ssize\_t [rt\\_dev\\_recv](#) (int fd, void \*buf, size\_t len, int flags)  
*Receive message from socket.*
- ssize\_t [rt\\_dev\\_sendmsg](#) (int fd, const struct msghdr \*msg, int flags)  
*Transmit message to socket.*
- ssize\_t [rt\\_dev\\_sendto](#) (int fd, const void \*buf, size\_t len, int flags, const struct sockaddr \*to, socklen\_t tolen)  
*Transmit message to socket.*
- ssize\_t [rt\\_dev\\_send](#) (int fd, const void \*buf, size\_t len, int flags)  
*Transmit message to socket.*
- int [rt\\_dev\\_bind](#) (int fd, const struct sockaddr \*my\_addr, socklen\_t addrlen)  
*Bind to local address.*
- int [rt\\_dev\\_connect](#) (int fd, const struct sockaddr \*serv\_addr, socklen\_t addrlen)

*Connect to remote address.*

- `int rt_dev_listen` (int fd, int backlog)  
*Listen for incoming connection requests.*
- `int rt_dev_accept` (int fd, struct sockaddr \*addr, socklen\_t \*addrlen)  
*Accept a connection requests.*
- `int rt_dev_shutdown` (int fd, int how)  
*Shut down parts of a connection.*
- `int rt_dev_getsockopt` (int fd, int level, int optname, void \*optval, socklen\_t \*optlen)  
*Get socket option.*
- `int rt_dev_setsockopt` (int fd, int level, int optname, const void \*optval, socklen\_t optlen)  
*Set socket option.*
- `int rt_dev_getsockname` (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get local socket address.*
- `int rt_dev_getpeername` (int fd, struct sockaddr \*name, socklen\_t \*namelen)  
*Get socket destination address.*

### 7.49.1 Detailed Description

Real-Time Driver Model for Xenomai, device operation multiplexing.

#### Note

Copyright (C) 2005 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>

Copyright (C) 2005 Joerg Langenberg <[joerg.langenberg@gmx.net](mailto:joerg.langenberg@gmx.net)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.50 kernel/cobalt/rtdm/device.c File Reference

Real-Time Driver Model for Xenomai, device management.



[illegible]

- `int rtdm_dev_register` (struct `rtdm_device` \*device)  
*Register a RTDM device.*
- `int rtdm_dev_unregister` (struct `rtdm_device` \*device, unsigned int poll\_delay)  
*Unregisters a RTDM device.*

## Real-Time Driver Model for Xenomai, device management.

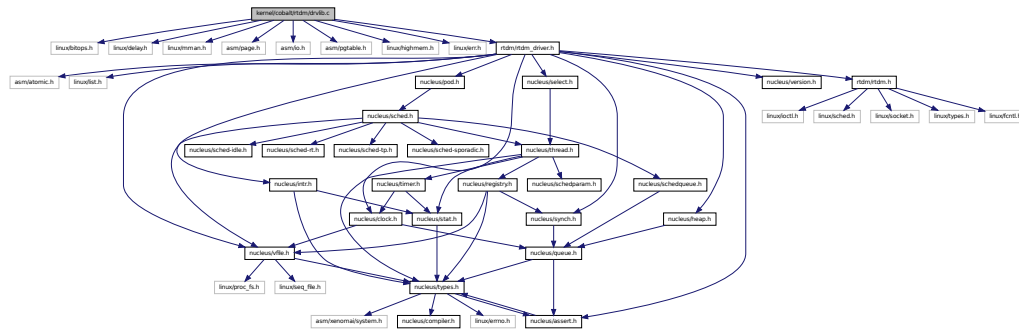
Copyright (C) 2005 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>  
Copyright (C) 2005 Joerg Langenberg <[joerg.langenberg@gmx.net](mailto:joerg.langenberg@gmx.net)>

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## Real-Time Driver Model for Xenomai, driver library.

Include dependency graph for drvlib.c:



## Functions

- [nanosecs\\_abs\\_t rtdm\\_clock\\_read](#) (void)  
*Get system time.*
- [nanosecs\\_abs\\_t rtdm\\_clock\\_read\\_monotonic](#) (void)  
*Get monotonic time.*
- [int rtdm\\_task\\_init](#) (rtdm\_task\_t \*task, const char \*name, [rtdm\\_task\\_proc\\_t](#) task\_proc, void \*arg, int priority, [nanosecs\\_rel\\_t](#) period)  
*Initialise and start a real-time task.*
- [void rtdm\\_task\\_destroy](#) (rtdm\_task\_t \*task)  
*Destroy a real-time task.*
- [void rtdm\\_task\\_set\\_priority](#) (rtdm\_task\_t \*task, int priority)  
*Adjust real-time task priority.*
- [int rtdm\\_task\\_set\\_period](#) (rtdm\_task\_t \*task, [nanosecs\\_rel\\_t](#) period)  
*Adjust real-time task period.*
- [int rtdm\\_task\\_wait\\_period](#) (void)  
*Wait on next real-time task period.*
- [int rtdm\\_task\\_unblock](#) (rtdm\_task\_t \*task)  
*Activate a blocked real-time task.*
- [rtdm\\_task\\_t \\* rtdm\\_task\\_current](#) (void)  
*Get current real-time task.*
- [int rtdm\\_task\\_sleep](#) ([nanosecs\\_rel\\_t](#) delay)  
*Sleep a specified amount of time.*
- [int rtdm\\_task\\_sleep\\_until](#) ([nanosecs\\_abs\\_t](#) wakeup\_time)  
*Sleep until a specified absolute time.*

- int [rtdm\\_task\\_sleep\\_abs](#) ([nanosecs\\_abs\\_t](#) wakeup\_time, enum [rtdm\\_timer\\_mode](#) mode)  
*Sleep until a specified absolute time.*
- void [rtdm\\_task\\_join\\_nrt](#) ([rtdm\\_task\\_t](#) \*task, unsigned int poll\_delay)  
*Wait on a real-time task to terminate.*
- void [rtdm\\_task\\_busy\\_sleep](#) ([nanosecs\\_rel\\_t](#) delay)  
*Busy-wait a specified amount of time.*
- int [rtdm\\_timer\\_init](#) ([rtdm\\_timer\\_t](#) \*timer, [rtdm\\_timer\\_handler\\_t](#) handler, const char \*name)  
*Initialise a timer.*
- void [rtdm\\_timer\\_destroy](#) ([rtdm\\_timer\\_t](#) \*timer)  
*Destroy a timer.*
- int [rtdm\\_timer\\_start](#) ([rtdm\\_timer\\_t](#) \*timer, [nanosecs\\_abs\\_t](#) expiry, [nanosecs\\_rel\\_t](#) interval, enum [rtdm\\_timer\\_mode](#) mode)  
*Start a timer.*
- void [rtdm\\_timer\\_stop](#) ([rtdm\\_timer\\_t](#) \*timer)  
*Stop a timer.*
- int [rtdm\\_timer\\_start\\_in\\_handler](#) ([rtdm\\_timer\\_t](#) \*timer, [nanosecs\\_abs\\_t](#) expiry, [nanosecs\\_rel\\_t](#) interval, enum [rtdm\\_timer\\_mode](#) mode)  
*Start a timer from inside a timer handler.*
- void [rtdm\\_timer\\_stop\\_in\\_handler](#) ([rtdm\\_timer\\_t](#) \*timer)  
*Stop a timer from inside a timer handler.*
- int [rtdm\\_irq\\_request](#) ([rtdm\\_irq\\_t](#) \*irq\_handle, unsigned int irq\_no, [rtdm\\_irq\\_handler\\_t](#) handler, unsigned long flags, const char \*device\_name, void \*arg)  
*Register an interrupt handler.*
- int [rtdm\\_irq\\_free](#) ([rtdm\\_irq\\_t](#) \*irq\_handle)  
*Release an interrupt handler.*
- int [rtdm\\_irq\\_enable](#) ([rtdm\\_irq\\_t](#) \*irq\_handle)  
*Enable interrupt line.*
- int [rtdm\\_irq\\_disable](#) ([rtdm\\_irq\\_t](#) \*irq\_handle)  
*Disable interrupt line.*
- int [rtdm\\_nrtsig\\_init](#) ([rtdm\\_nrtsig\\_t](#) \*nrt\_sig, [rtdm\\_nrtsig\\_handler\\_t](#) handler, void \*arg)  
*Register a non-real-time signal handler.*
- void [rtdm\\_nrtsig\\_destroy](#) ([rtdm\\_nrtsig\\_t](#) \*nrt\_sig)  
*Release a non-realtime signal handler.*

- void [rt dm\\_nrtsig\\_pend](#) (rt dm\_nrtsig\_t \*nrt\_sig)  
*Trigger non-real-time signal.*
- int [rt dm\\_mmap\\_to\\_user](#) (rt dm\_user\_info\_t \*user\_info, void \*src\_addr, size\_t len, int prot, void \*\*pptr, struct vm\_operations\_struct \*vm\_ops, void \*vm\_private\_data)  
*Map a kernel memory range into the address space of the user.*
- int [rt dm\\_iomap\\_to\\_user](#) (rt dm\_user\_info\_t \*user\_info, phys\_addr\_t src\_addr, size\_t len, int prot, void \*\*pptr, struct vm\_operations\_struct \*vm\_ops, void \*vm\_private\_data)  
*Map an I/O memory range into the address space of the user.*
- int [rt dm\\_munmap](#) (rt dm\_user\_info\_t \*user\_info, void \*ptr, size\_t len)  
*Unmap a user memory range.*
- void [rt dm\\_printk](#) (const char \*format,...)  
*Real-time safe message printing on kernel console.*
- void \* [rt dm\\_malloc](#) (size\_t size)  
*Allocate memory block in real-time context.*
- void [rt dm\\_free](#) (void \*ptr)  
*Release real-time memory block.*
- int [rt dm\\_read\\_user\\_ok](#) (rt dm\_user\_info\_t \*user\_info, const void \_\_user \*ptr, size\_t size)  
*Check if read access to user-space memory block is safe.*
- int [rt dm\\_rw\\_user\\_ok](#) (rt dm\_user\_info\_t \*user\_info, const void \_\_user \*ptr, size\_t size)  
*Check if read/write access to user-space memory block is safe.*
- int [rt dm\\_copy\\_from\\_user](#) (rt dm\_user\_info\_t \*user\_info, void \*dst, const void \_\_user \*src, size\_t size)  
*Copy user-space memory block to specified buffer.*
- int [rt dm\\_safe\\_copy\\_from\\_user](#) (rt dm\_user\_info\_t \*user\_info, void \*dst, const void \_\_user \*src, size\_t size)  
*Check if read access to user-space memory block and copy it to specified buffer.*
- int [rt dm\\_copy\\_to\\_user](#) (rt dm\_user\_info\_t \*user\_info, void \_\_user \*dst, const void \*src, size\_t size)  
*Copy specified buffer to user-space memory block.*
- int [rt dm\\_safe\\_copy\\_to\\_user](#) (rt dm\_user\_info\_t \*user\_info, void \_\_user \*dst, const void \*src, size\_t size)  
*Check if read/write access to user-space memory block is safe and copy specified buffer to it.*
- int [rt dm\\_strncpy\\_from\\_user](#) (rt dm\_user\_info\_t \*user\_info, char \*dst, const char \_\_user \*src, size\_t count)  
*Copy user-space string to specified buffer.*
- int [rt dm\\_in\\_rt\\_context](#) (void)

*Test if running in a real-time task.*

- int [rtdm\\_rt\\_capable](#) (rtdm\_user\_info\_t \*user\_info)

*Test if the caller is capable of running in real-time context.*

## Timeout Sequence Management

- void [rtdm\\_toseq\\_init](#) (rtdm\_toseq\_t \*timeout\_seq, [nanosecs\\_rel\\_t](#) timeout)

*Initialise a timeout sequence.*

## Event Services

- void [rtdm\\_event\\_init](#) (rtdm\_event\_t \*event, unsigned long pending)

*Initialise an event.*

- void [rtdm\\_event\\_destroy](#) (rtdm\_event\_t \*event)

*Destroy an event.*

- void [rtdm\\_event\\_pulse](#) (rtdm\_event\_t \*event)

*Signal an event occurrence to currently listening waiters.*

- void [rtdm\\_event\\_signal](#) (rtdm\_event\_t \*event)

*Signal an event occurrence.*

- int [rtdm\\_event\\_wait](#) (rtdm\_event\_t \*event)

*Wait on event occurrence.*

- int [rtdm\\_event\\_timedwait](#) (rtdm\_event\_t \*event, [nanosecs\\_rel\\_t](#) timeout, rtdm\_toseq\_t \*timeout\_seq)

*Wait on event occurrence with timeout.*

- void [rtdm\\_event\\_clear](#) (rtdm\_event\_t \*event)

*Clear event state.*

- int [rtdm\\_event\\_select\\_bind](#) (rtdm\_event\_t \*event, rtdm\_selector\_t \*selector, enum [rtdm\\_selecttype](#) type, unsigned fd\_index)

*Bind a selector to an event.*

## Semaphore Services

- void [rtdm\\_sem\\_init](#) (rtdm\_sem\_t \*sem, unsigned long value)

*Initialise a semaphore.*

- void [rtdm\\_sem\\_destroy](#) (rtdm\_sem\_t \*sem)

*Destroy a semaphore.*

- int [rtdm\\_sem\\_down](#) (rtdm\_sem\_t \*sem)

*Decrement a semaphore.*

- int [rtdm\\_sem\\_timeddown](#) (rtdm\_sem\_t \*sem, [nanosecs\\_rel\\_t](#) timeout, rtdm\_toseq\_t \*timeout\_seq)

*Decrement a semaphore with timeout.*

- void `rtdm_sem_up` (rtdm\_sem\_t \*sem)  
*Increment a semaphore.*
- int `rtdm_sem_select_bind` (rtdm\_sem\_t \*sem, rtdm\_selector\_t \*selector, enum `rtdm_selecttype` type, unsigned fd\_index)  
*Bind a selector to a semaphore.*

## Mutex Services

- void `rtdm_mutex_init` (rtdm\_mutex\_t \*mutex)  
*Initialise a mutex.*
- void `rtdm_mutex_destroy` (rtdm\_mutex\_t \*mutex)  
*Destroy a mutex.*
- void `rtdm_mutex_unlock` (rtdm\_mutex\_t \*mutex)  
*Release a mutex.*
- int `rtdm_mutex_lock` (rtdm\_mutex\_t \*mutex)  
*Request a mutex.*
- int `rtdm_mutex_timedlock` (rtdm\_mutex\_t \*mutex, `nanosecs_rel_t` timeout, rtdm\_toseq\_t \*timeout\_seq)  
*Request a mutex with timeout.*

### 7.51.1 Detailed Description

Real-Time Driver Model for Xenomai, driver library.

#### Note

Copyright (C) 2005-2007 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>  
 Copyright (C) 2005 Joerg Langenberg <[joerg.langenberg@gmx.net](mailto:joerg.langenberg@gmx.net)>  
 Copyright (C) 2008 Gilles Chanteperdrix <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 7.52 kernel/cobalt/syscall.c File Reference

This file is part of the Xenomai project.







## Chapter 8

# Example Documentation

8.1 `bufp-label.c`

8.2 `bufp-readwrite.c`

8.3 `cross-link.c`

8.4 `iddp-label.c`

8.5 `iddp-sendrecv.c`

8.6 `rtcan_rtt.c`

8.7 `rtcanconfig.c`

## 8.8 `rtcanrecv.c`

## 8.9 `rtcansend.c`

## 8.10 `xddp-echo.c`

## 8.11 `xddp-label.c`

## 8.12 `xddp-stream.c`

# Index

- [\\_\\_xnpod\\_reset\\_thread](#)
  - [pod](#), [67](#)
- [affinity](#)
  - [xnthread\\_info](#), [286](#)
- [arm/hal.c](#)
  - [rthal\\_timer\\_release](#), [351](#)
  - [rthal\\_timer\\_request](#), [351](#)
- [begin](#)
  - [xnvmfile\\_regular\\_ops](#), [290](#)
  - [xnvmfile\\_snapshot\\_ops](#), [296](#)
- [bind\\_\\_AF\\_RTIPC](#)
  - [rtipc](#), [250](#)
- [bprio](#)
  - [xnthread\\_info](#), [286](#)
- [bufd](#)
  - [xnbufd\\_copy\\_from\\_kmem](#), [41](#)
  - [xnbufd\\_copy\\_to\\_kmem](#), [42](#)
  - [xnbufd\\_invalidate](#), [43](#)
  - [xnbufd\\_map\\_kread](#), [43](#)
  - [xnbufd\\_map\\_kwrite](#), [44](#)
  - [xnbufd\\_map\\_uread](#), [44](#)
  - [xnbufd\\_map\\_uwrite](#), [45](#)
  - [xnbufd\\_reset](#), [45](#)
  - [xnbufd\\_unmap\\_kread](#), [46](#)
  - [xnbufd\\_unmap\\_kwrite](#), [46](#)
  - [xnbufd\\_unmap\\_uread](#), [47](#)
  - [xnbufd\\_unmap\\_uwrite](#), [47](#)
- [Buffer descriptors.](#), [38](#)
- [BUFP\\_BUFSZ](#)
  - [rtipc](#), [243](#)
- [BUFP\\_LABEL](#)
  - [rtipc](#), [243](#)
- [CAN Devices](#), [204](#)
- [CAN\\_BITTIME\\_BTR](#)
  - [rtcan](#), [225](#)
- [CAN\\_BITTIME\\_STD](#)
  - [rtcan](#), [225](#)
- [CAN\\_MODE\\_SLEEP](#)
  - [rtcan](#), [226](#)
- [CAN\\_MODE\\_START](#)
  - [rtcan](#), [226](#)
- [CAN\\_MODE\\_STOP](#)
  - [rtcan](#), [226](#)
- [rtcan](#), [226](#)
- [CAN\\_STATE\\_ACTIVE](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_BUS\\_OFF](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_BUS\\_PASSIVE](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_BUS\\_WARNING](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_SCANNING\\_BAUDRATE](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_SLEEPING](#)
  - [rtcan](#), [226](#)
- [CAN\\_STATE\\_STOPPED](#)
  - [rtcan](#), [226](#)
- [can\\_bittime](#), [269](#)
- [can\\_bittime\\_btr](#), [270](#)
- [can\\_bittime\\_std](#), [270](#)
- [CAN\\_BITTIME\\_TYPE](#)
  - [rtcan](#), [225](#)
- [CAN\\_CTRLMODE\\_LISTENONLY](#)
  - [rtcan](#), [215](#)
- [CAN\\_CTRLMODE\\_LOOPBACK](#)
  - [rtcan](#), [215](#)
- [CAN\\_ERR\\_LOSTARB\\_UNSPEC](#)
  - [rtcan](#), [215](#)
- [can\\_filter](#), [271](#)
  - [can\\_id](#), [272](#)
  - [can\\_mask](#), [272](#)
- [can\\_filter\\_t](#)
  - [rtcan](#), [225](#)
- [can\\_frame](#), [272](#)
  - [can\\_id](#), [272](#)
- [can\\_frame\\_t](#)
  - [rtcan](#), [225](#)
- [can\\_id](#)
  - [can\\_filter](#), [272](#)
  - [can\\_frame](#), [272](#)
- [can\\_ifindex](#)
  - [sockaddr\\_can](#), [281](#)
- [can\\_mask](#)
  - [can\\_filter](#), [272](#)
- [CAN\\_MODE](#)
  - [rtcan](#), [226](#)

- CAN\_RAW\_ERR\_FILTER
  - rtcan, 216
- CAN\_RAW\_FILTER
  - rtcan, 216
- CAN\_RAW\_LOOPBACK
  - rtcan, 217
- CAN\_RAW\_RECV\_OWN\_MSGS
  - rtcan, 217
- CAN\_STATE
  - rtcan, 226
- clock
  - xnclock\_adjust, 48
- Clock Services, 139
- Clocks and timers services., 20
- close\_\_AF\_RTIPC
  - rtipc, 251
- close\_rt
  - rtdm\_operations, 278
- cobalt\_cancel
  - pthread\_cancel, 17
  - pthread\_cleanup\_pop, 18
  - pthread\_cleanup\_push, 18
  - pthread\_setcancelstate, 19
  - pthread\_setcanceltype, 19
- cobalt\_cond
  - pthread\_cond\_destroy, 24
  - pthread\_cond\_init, 25
  - pthread\_condattr\_destroy, 25
  - pthread\_condattr\_getclock, 26
  - pthread\_condattr\_getpshared, 26
  - pthread\_condattr\_init, 27
  - pthread\_condattr\_setclock, 27
  - pthread\_condattr\_setpshared, 28
- cobalt\_mq
  - mq\_close, 29
  - mq\_getattr, 30
  - mq\_open, 30
  - mq\_setattr, 32
  - mq\_unlink, 32
- cobalt\_mutex
  - pthread\_mutexattr\_destroy, 34
  - pthread\_mutexattr\_getprotocol, 34
  - pthread\_mutexattr\_getpshared, 35
  - pthread\_mutexattr\_gettype, 35
  - pthread\_mutexattr\_init, 36
  - pthread\_mutexattr\_setprotocol, 36
  - pthread\_mutexattr\_setpshared, 37
  - pthread\_mutexattr\_settype, 38
- cobalt\_sem
  - sem\_close, 192
  - sem\_destroy, 192
  - sem\_getvalue, 192
  - sem\_open, 193
  - sem\_post, 194
  - sem\_timedwait, 194
  - sem\_trywait, 195
  - sem\_unlink, 195
  - sem\_wait, 196
- cobalt\_thread
  - pthread\_create, 198
  - pthread\_getschedparam, 199
  - pthread\_getschedparam\_ex, 199
  - pthread\_make\_periodic\_np, 200
  - pthread\_set\_mode\_np, 200
  - pthread\_set\_name\_np, 201
  - pthread\_setschedparam, 202
  - pthread\_setschedparam\_ex, 203
- cobalt\_time
  - do\_clock\_host\_realtime, 21
  - timer\_create, 21
  - timer\_gettime, 22
  - timer\_settime, 22
- Condition variables services., 23
- connect\_\_AF\_RTIPC
  - rtipc, 252
- cprio
  - xnthread\_info, 286
- cpu
  - xnthread\_info, 286
- ctxswitches
  - xnthread\_info, 287
- curr
  - xnsched, 284
- databuf
  - xnvfile\_snapshot\_iterator, 295
- Debugging services., 49
- Device Profiles, 189
- Device Registration Services, 128
- devregister
  - rtdm\_close\_handler\_t, 131
  - RTDM\_CLOSING, 131
  - rtdm\_context\_to\_private, 135
  - RTDM\_CREATED\_IN\_NRT, 131
  - rtdm\_dev\_register, 136
  - rtdm\_dev\_unregister, 136
  - RTDM\_DEVICE\_TYPE\_MASK, 131
  - RTDM\_EXCLUSIVE, 131
  - rtdm\_ioctl\_handler\_t, 132
  - RTDM\_NAMED\_DEVICE, 131
  - rtdm\_open\_handler\_t, 132
  - rtdm\_private\_to\_context, 137
  - RTDM\_PROTOCOL\_DEVICE, 131
  - rtdm\_read\_handler\_t, 133
  - rtdm\_recvmmsg\_handler\_t, 133
  - rtdm\_select\_bind\_handler\_t, 134
  - rtdm\_sendmsg\_handler\_t, 134
  - rtdm\_socket\_handler\_t, 134

- rtdm\_write\_handler\_t, 135
- do\_clock\_host\_realtime
  - cobalt\_time, 21
- Driver Development API, 138
- Dynamic memory allocation services., 49
- end
  - xnvfile\_regular\_ops, 290
  - xnvfile\_snapshot\_ops, 296
- endfn
  - xnvfile\_snapshot\_iterator, 295
- exectime
  - xnthread\_info, 287
- File descriptors events multiplexing services., 91
- get
  - xnvfile\_lock\_ops, 288
- getpeername\_\_AF\_RTIPC
  - rtpc, 252
- getsockname\_\_AF\_RTIPC
  - rtpc, 253
- getsockopt\_\_AF\_RTIPC
  - rtpc, 253
- hal
  - rthal\_apc\_alloc, 15
  - rthal\_apc\_free, 16
- HAL., 14
- heap
  - xnheap\_alloc, 50
  - xnheap\_destroy, 51
  - xnheap\_extend, 51
  - xnheap\_free, 52
  - xnheap\_init, 52
  - xnheap\_schedule\_free, 53
  - xnheap\_set\_label, 54
  - xnheap\_test\_and\_free, 54
- htimer
  - xnsched, 284
- IDDP\_LABEL
  - rtpc, 244
- IDDP\_POOLSZ
  - rtpc, 245
- include/cobalt/nucleus/bufd.h, 301
- include/cobalt/nucleus/clock.h, 303
- include/cobalt/nucleus/hostrt.h, 304
- include/cobalt/nucleus/map.h, 306
- include/cobalt/nucleus/pod.h, 307
- include/cobalt/nucleus/registry.h, 310
- include/cobalt/nucleus/sched-idle.h, 311
- include/cobalt/nucleus/sched-rt.h, 312
- include/cobalt/nucleus/sched-sporadic.h, 313
- include/cobalt/nucleus/sched-tp.h, 313
- include/cobalt/nucleus/sched.h, 314
- include/cobalt/nucleus/select.h, 316
- include/cobalt/nucleus/timer.h, 318
- include/cobalt/nucleus/vdso.h, 319
- include/cobalt/nucleus/vfile.h, 320
- include/rtdm/rtdm.h, 331
- include/rtdm/rtdm\_driver.h, 334
- include/rtdm/rtpc.h, 341
- include/rtdm/rtserial.h, 344
- include/rtdm/rttesting.h, 348
- inesting
  - xnsched, 285
- Inter-Driver API, 119
- interdrv
  - rtdm\_accept, 121
  - rtdm\_bind, 121
  - rtdm\_close, 121
  - rtdm\_connect, 121
  - rtdm\_context\_get, 121
  - rtdm\_context\_lock, 122
  - rtdm\_context\_put, 122
  - rtdm\_context\_unlock, 123
  - rtdm\_getpeername, 123
  - rtdm\_getsockname, 124
  - rtdm\_getsockopt, 124
  - rtdm\_ioctl, 124
  - rtdm\_listen, 124
  - rtdm\_open, 124
  - rtdm\_read, 125
  - rtdm\_recv, 125
  - rtdm\_recvfrom, 125
  - rtdm\_recvmsg, 125
  - rtdm\_select\_bind, 125
  - rtdm\_send, 126
  - rtdm\_sendmsg, 126
  - rtdm\_sendto, 127
  - rtdm\_setsockopt, 127
  - rtdm\_shutdown, 127
  - rtdm\_socket, 127
  - rtdm\_write, 127
- Interrupt Management Services, 171
- Interrupt management., 55
- intr
  - xnintr\_affinity, 56
  - xnintr\_attach, 56
  - xnintr\_destroy, 57
  - xnintr\_detach, 57
  - xnintr\_disable, 58
  - xnintr\_enable, 58
  - xnintr\_init, 59
- IPCPROTO\_BUF
  - rtpc, 250

- IPCPROTO\_IDDP
  - rtipc, 250
- IPCPROTO\_IPC
  - rtipc, 250
- IPCPROTO\_XDDP
  - rtipc, 250
- kernel/cobalt/arch/arm/hal.c, 350
- kernel/cobalt/arch/blackfin/hal.c, 352
- kernel/cobalt/arch/generic/hal.c, 353
- kernel/cobalt/arch/nios2/hal.c, 354
- kernel/cobalt/arch/powerpc/hal.c, 354
- kernel/cobalt/arch/sh/hal.c, 355
- kernel/cobalt/arch/x86/hal.c, 355
- kernel/cobalt/arch/x86/smi.c, 356
- kernel/cobalt/nucleus/bufd.c, 359
- kernel/cobalt/nucleus/clock.c, 357
- kernel/cobalt/nucleus/debug.c, 360
- kernel/cobalt/nucleus/heap.c, 361
- kernel/cobalt/nucleus/intr.c, 362
- kernel/cobalt/nucleus/map.c, 364
- kernel/cobalt/nucleus/pod.c, 365
- kernel/cobalt/nucleus/registry.c, 367
- kernel/cobalt/nucleus/sched-idle.c, 368
- kernel/cobalt/nucleus/sched-rt.c, 369
- kernel/cobalt/nucleus/sched-sporadic.c, 370
- kernel/cobalt/nucleus/sched-tp.c, 371
- kernel/cobalt/nucleus/sched.c, 373
- kernel/cobalt/nucleus/select.c, 373
- kernel/cobalt/nucleus/shadow.c, 375
- kernel/cobalt/nucleus/synch.c, 376
- kernel/cobalt/nucleus/timer.c, 378
- kernel/cobalt/nucleus/vfile.c, 379
- kernel/cobalt/rtdm/core.c, 381
- kernel/cobalt/rtdm/device.c, 384
- kernel/cobalt/rtdm/drvlib.c, 385
- kernel/cobalt/rtdm/module.c, 358
- kernel/cobalt/syscall.c, 390
- label
  - rtipc\_port\_label, 279
- lflags
  - xnsched, 285
- Lightweight key-to-object mapping service, 61
- map
  - xnmap\_create, 62
  - xnmap\_delete, 62
  - xnmap\_enter, 63
  - xnmap\_fetch, 63
  - xnmap\_fetch\_nocheck, 64
  - xnmap\_remove, 64
- Message queues services., 29
- modeswitches
  - xnthread\_info, 287
- mq\_close
  - cobalt\_mq, 29
- mq\_getattr
  - cobalt\_mq, 30
- mq\_open
  - cobalt\_mq, 30
- mq\_setattr
  - cobalt\_mq, 32
- mq\_unlink
  - cobalt\_mq, 32
- Mutex services., 33
- name
  - xnthread\_info, 287
- nanosecs\_abs\_t
  - rtdm, 228
- nanosecs\_rel\_t
  - rtdm, 228
- next
  - xnvfile\_regular\_ops, 291
  - xnvfile\_snapshot\_ops, 297
- nkvfroot
  - vfile, 118
- Non-Real-Time Signalling Services, 176
- nrdata
  - xnvfile\_snapshot\_iterator, 295
- nrtsignal
  - rtdm\_nrtsig\_destroy, 177
  - rtdm\_nrtsig\_handler\_t, 176
  - rtdm\_nrtsig\_init, 177
  - rtdm\_nrtsig\_pend, 177
- nucleus\_state\_flags
  - XNHELD, 12
  - XNLOCK, 12
  - XNMIGRATE, 13
  - XNPEND, 13
  - XNREADY, 13
  - XNSUSP, 13
- open\_rt
  - rtdm\_device, 276
- pagefaults
  - xnthread\_info, 287
- pod
  - \_\_xnpod\_reset\_thread, 67
  - xnpod\_abort\_thread, 67
  - xnpod\_add\_hook, 68
  - xnpod\_delete\_thread, 69
  - xnpod\_disable\_timesource, 69
  - xnpod\_dispatch\_signals, 70
  - xnpod\_enable\_timesource, 70
  - xnpod\_handle\_exception, 71

- xnpod\_init, [71](#)
- xnpod\_init\_thread, [71](#)
- xnpod\_migrate\_thread, [73](#)
- xnpod\_remove\_hook, [73](#)
- xnpod\_resume\_thread, [74](#)
- xnpod\_schedule, [75](#)
- xnpod\_set\_thread\_mode, [76](#)
- xnpod\_set\_thread\_periodic, [77](#)
- xnpod\_set\_thread\_schedparam, [78](#)
- xnpod\_set\_thread\_tslice, [79](#)
- xnpod\_shutdown, [79](#)
- xnpod\_start\_thread, [80](#)
- xnpod\_stop\_thread, [81](#)
- xnpod\_suspend\_thread, [81](#)
- xnpod\_unblock\_thread, [83](#)
- xnpod\_wait\_thread\_period, [83](#)
- xnpod\_welcome\_thread, [84](#)
- pos
  - xnvfile\_regular\_iterator, [289](#)
- POSIX skin., [28](#)
- private
  - xnvfile\_regular\_iterator, [289](#)
  - xnvfile\_snapshot\_iterator, [295](#)
- profiles
  - RTIOC\_DEVICE\_INFO, [190](#)
  - RTIOC\_PURGE, [190](#)
- pthread\_cancel
  - cobalt\_cancel, [17](#)
- pthread\_cleanup\_pop
  - cobalt\_cancel, [18](#)
- pthread\_cleanup\_push
  - cobalt\_cancel, [18](#)
- pthread\_cond\_destroy
  - cobalt\_cond, [24](#)
- pthread\_cond\_init
  - cobalt\_cond, [25](#)
- pthread\_condattr\_destroy
  - cobalt\_cond, [25](#)
- pthread\_condattr\_getclock
  - cobalt\_cond, [26](#)
- pthread\_condattr\_getpshared
  - cobalt\_cond, [26](#)
- pthread\_condattr\_init
  - cobalt\_cond, [27](#)
- pthread\_condattr\_setclock
  - cobalt\_cond, [27](#)
- pthread\_condattr\_setpshared
  - cobalt\_cond, [28](#)
- pthread\_create
  - cobalt\_thread, [198](#)
- pthread\_getschedparam
  - cobalt\_thread, [199](#)
- pthread\_getschedparam\_ex
  - cobalt\_thread, [199](#)
- pthread\_make\_periodic\_np
  - cobalt\_thread, [200](#)
- pthread\_mutexattr\_destroy
  - cobalt\_mutex, [34](#)
- pthread\_mutexattr\_getprotocol
  - cobalt\_mutex, [34](#)
- pthread\_mutexattr\_getpshared
  - cobalt\_mutex, [35](#)
- pthread\_mutexattr\_gettype
  - cobalt\_mutex, [35](#)
- pthread\_mutexattr\_init
  - cobalt\_mutex, [36](#)
- pthread\_mutexattr\_setprotocol
  - cobalt\_mutex, [36](#)
- pthread\_mutexattr\_setpshared
  - cobalt\_mutex, [37](#)
- pthread\_mutexattr\_settype
  - cobalt\_mutex, [38](#)
- pthread\_set\_mode\_np
  - cobalt\_thread, [200](#)
- pthread\_set\_name\_np
  - cobalt\_thread, [201](#)
- pthread\_setcancelstate
  - cobalt\_cancel, [19](#)
- pthread\_setcanceltype
  - cobalt\_cancel, [19](#)
- pthread\_setschedparam
  - cobalt\_thread, [202](#)
- pthread\_setschedparam\_ex
  - cobalt\_thread, [203](#)
- put
  - xnvfile\_lock\_ops, [288](#)
- Real-Time Driver Model, [226](#)
- Real-time IPC protocols, [240](#)
- Real-time pod services., [65](#)
- Real-time shadow services., [94](#)
- recvmsg\_\_AF\_RTIPC
  - rtipc, [253](#)
- refcnt
  - xnpod, [283](#)
- registry
  - xnregistry\_bind, [86](#)
  - xnregistry\_enter, [87](#)
  - xnregistry\_fetch, [87](#)
  - xnregistry\_get, [88](#)
  - xnregistry\_put, [88](#)
  - xnregistry\_remove, [89](#)
  - xnregistry\_remove\_safe, [90](#)
- Registry services., [85](#)
- relpoint
  - xnthread\_info, [287](#)
- rev
  - xnvfile\_rev\_tag, [293](#)

- rewind
  - xnvfile\_regular\_ops, 291
  - xnvfile\_snapshot\_ops, 297
- rootcb
  - xnsched, 285
- rt
  - xnsched, 285
- rt\_dev\_accept
  - userapi, 230
- rt\_dev\_bind
  - userapi, 230
- rt\_dev\_close
  - userapi, 231
- rt\_dev\_connect
  - userapi, 231
- rt\_dev\_getpeername
  - userapi, 232
- rt\_dev\_getsockname
  - userapi, 232
- rt\_dev\_getsockopt
  - userapi, 232
- rt\_dev\_ioctl
  - userapi, 233
- rt\_dev\_listen
  - userapi, 233
- rt\_dev\_open
  - userapi, 234
- rt\_dev\_read
  - userapi, 234
- rt\_dev\_recv
  - userapi, 235
- rt\_dev\_recvfrom
  - userapi, 235
- rt\_dev\_recvmsg
  - userapi, 236
- rt\_dev\_send
  - userapi, 236
- rt\_dev\_sendmsg
  - userapi, 237
- rt\_dev\_sendto
  - userapi, 237
- rt\_dev\_setsockopt
  - userapi, 238
- rt\_dev\_shutdown
  - userapi, 238
- rt\_dev\_socket
  - userapi, 239
- rt\_dev\_write
  - userapi, 239
- rtcan
  - CAN\_BITTIME\_BTR, 225
  - CAN\_BITTIME\_STD, 225
  - CAN\_MODE\_SLEEP, 226
  - CAN\_MODE\_START, 226
  - CAN\_MODE\_STOP, 226
  - CAN\_STATE\_ACTIVE, 226
  - CAN\_STATE\_BUS\_OFF, 226
  - CAN\_STATE\_BUS\_PASSIVE, 226
  - CAN\_STATE\_BUS\_WARNING, 226
  - CAN\_STATE\_SCANNING\_BAUDRATE, 226
  - CAN\_STATE\_SLEEPING, 226
  - CAN\_STATE\_STOPPED, 226
  - CAN\_BITTIME\_TYPE, 225
  - CAN\_CTRLMODE\_LISTENONLY, 215
  - CAN\_CTRLMODE\_LOOPBACK, 215
  - CAN\_ERR\_LOSTARB\_UNSPEC, 215
  - can\_filter\_t, 225
  - can\_frame\_t, 225
  - CAN\_MODE, 226
  - CAN\_RAW\_ERR\_FILTER, 216
  - CAN\_RAW\_FILTER, 216
  - CAN\_RAW\_LOOPBACK, 217
  - CAN\_RAW\_RECV\_OWN\_MSGS, 217
  - CAN\_STATE, 226
  - RTCAN\_RTIOC\_RCV\_TIMEOUT, 217
  - RTCAN\_RTIOC\_SND\_TIMEOUT, 218
  - RTCAN\_RTIOC\_TAKE\_TIMESTAMP, 218
  - SIOCGCANBAUDRATE, 219
  - SIOCGCANCTRLMODE, 220
  - SIOCGCANCUSTOMBITTIME, 220
  - SIOCGCANSTATE, 221
  - SIOCGIFINDEX, 221
  - SIOCSCANBAUDRATE, 222
  - SIOCSCANCTRLMODE, 222
  - SIOCSCANCUSTOMBITTIME, 223
  - SIOCSCANMODE, 224
  - SOL\_CAN\_RAW, 225
  - RTCAN\_RTIOC\_RCV\_TIMEOUT
    - rtcan, 217
  - RTCAN\_RTIOC\_SND\_TIMEOUT
    - rtcan, 218
  - RTCAN\_RTIOC\_TAKE\_TIMESTAMP
    - rtcan, 218
- rtdm
  - nanosecs\_abs\_t, 228
  - nanosecs\_rel\_t, 228
  - RTDM\_TIMEOUT\_INFINITE, 227
  - RTDM\_TIMEOUT\_NONE, 227
  - RTDM\_SELECTTYPE\_EXCEPT
    - rtdmsync, 159
  - RTDM\_SELECTTYPE\_READ
    - rtdmsync, 159
  - RTDM\_SELECTTYPE\_WRITE
    - rtdmsync, 159
  - RTDM\_TIMERMODE\_ABSOLUTE
    - rtdmtimer, 149
  - RTDM\_TIMERMODE\_REALTIME



- rtdmtimer, [149](#)
- RTDM\_TIMERMODE\_RELATIVE
  - rtdmtimer, [149](#)
- rtdm\_accept
  - interdrv, [121](#)
- rtdm\_bind
  - interdrv, [121](#)
- rtdm\_clock\_read
  - rtdmclock, [139](#)
- rtdm\_clock\_read\_monotonic
  - rtdmclock, [140](#)
- rtdm\_close
  - interdrv, [121](#)
- rtdm\_close\_handler\_t
  - devregister, [131](#)
- RTDM\_CLOSING
  - devregister, [131](#)
- rtdm\_connect
  - interdrv, [121](#)
- rtdm\_context\_get
  - interdrv, [121](#)
- rtdm\_context\_lock
  - interdrv, [122](#)
- rtdm\_context\_put
  - interdrv, [122](#)
- rtdm\_context\_to\_private
  - devregister, [135](#)
- rtdm\_context\_unlock
  - interdrv, [123](#)
- rtdm\_copy\_from\_user
  - util, [179](#)
- rtdm\_copy\_to\_user
  - util, [180](#)
- RTDM\_CREATED\_IN\_NRT
  - devregister, [131](#)
- rtdm\_dev\_context, [273](#)
- rtdm\_dev\_register
  - devregister, [136](#)
- rtdm\_dev\_unregister
  - devregister, [136](#)
- rtdm\_device, [274](#)
  - open\_rt, [276](#)
  - socket\_rt, [276](#)
- rtdm\_device\_info, [276](#)
- RTDM\_DEVICE\_TYPE\_MASK
  - devregister, [131](#)
- rtdm\_event\_clear
  - rtdmsync, [159](#)
- rtdm\_event\_destroy
  - rtdmsync, [160](#)
- rtdm\_event\_init
  - rtdmsync, [160](#)
- rtdm\_event\_pulse
  - rtdmsync, [160](#)
- rtdm\_event\_select\_bind
  - rtdmsync, [161](#)
- rtdm\_event\_signal
  - rtdmsync, [162](#)
- rtdm\_event\_timedwait
  - rtdmsync, [162](#)
- rtdm\_event\_wait
  - rtdmsync, [163](#)
- RTDM\_EXCLUSIVE
  - devregister, [131](#)
- RTDM\_EXECUTE\_ATOMICALY
  - rtdmsync, [156](#)
- rtdm\_free
  - util, [181](#)
- rtdm\_getpeername
  - interdrv, [123](#)
- rtdm\_getsockname
  - interdrv, [124](#)
- rtdm\_getsockopt
  - interdrv, [124](#)
- rtdm\_in\_rt\_context
  - util, [181](#)
- rtdm\_ioctl
  - interdrv, [124](#)
- rtdm\_ioctl\_handler\_t
  - devregister, [132](#)
- rtdm\_iomap\_to\_user
  - util, [181](#)
- rtdm\_irq\_disable
  - rtdmirq, [173](#)
- rtdm\_irq\_enable
  - rtdmirq, [174](#)
- rtdm\_irq\_free
  - rtdmirq, [174](#)
- rtdm\_irq\_get\_arg
  - rtdmirq, [172](#)
- rtdm\_irq\_handler\_t
  - rtdmirq, [173](#)
- rtdm\_irq\_request
  - rtdmirq, [175](#)
- rtdm\_listen
  - interdrv, [124](#)
- rtdm\_lock\_get
  - rtdmsync, [156](#)
- rtdm\_lock\_get\_irqsave
  - rtdmsync, [157](#)
- rtdm\_lock\_init
  - rtdmsync, [157](#)
- rtdm\_lock\_irqrestore
  - rtdmsync, [157](#)
- rtdm\_lock\_irqsave
  - rtdmsync, [158](#)
- rtdm\_lock\_put
  - rtdmsync, [158](#)

- rt dm\_lock\_put\_irqrestore
  - rt dmsync, 159
- rt dm\_malloc
  - util, 182
- rt dm\_mmap\_to\_user
  - util, 183
- rt dm\_munmap
  - util, 184
- rt dm\_mutex\_destroy
  - rt dmsync, 163
- rt dm\_mutex\_init
  - rt dmsync, 164
- rt dm\_mutex\_lock
  - rt dmsync, 164
- rt dm\_mutex\_timedlock
  - rt dmsync, 165
- rt dm\_mutex\_unlock
  - rt dmsync, 166
- RTDM\_NAMED\_DEVICE
  - devregister, 131
- rt dm\_nrt sig\_destroy
  - nrt signal, 177
- rt dm\_nrt sig\_handler\_t
  - nrt signal, 176
- rt dm\_nrt sig\_init
  - nrt signal, 177
- rt dm\_nrt sig\_pend
  - nrt signal, 177
- rt dm\_open
  - interdrv, 124
- rt dm\_open\_handler\_t
  - devregister, 132
- rt dm\_operations, 277
  - close\_rt, 278
- rt dm\_printk
  - util, 184
- rt dm\_private\_to\_context
  - devregister, 137
- RTDM\_PROTOCOL\_DEVICE
  - devregister, 131
- rt dm\_read
  - interdrv, 125
- rt dm\_read\_handler\_t
  - devregister, 133
- rt dm\_read\_user\_ok
  - util, 185
- rt dm\_rcv
  - interdrv, 125
- rt dm\_rcvfrom
  - interdrv, 125
- rt dm\_rcvmsg
  - interdrv, 125
- rt dm\_rcvmsg\_handler\_t
  - devregister, 133
- rt dm\_rt\_capable
  - util, 185
- rt dm\_rw\_user\_ok
  - util, 186
- rt dm\_safe\_copy\_from\_user
  - util, 186
- rt dm\_safe\_copy\_to\_user
  - util, 187
- rt dm\_select\_bind
  - interdrv, 125
  - rt dmsync, 166
- rt dm\_select\_bind\_handler\_t
  - devregister, 134
- rt dm\_selecttype
  - rt dmsync, 159
- rt dm\_sem\_destroy
  - rt dmsync, 167
- rt dm\_sem\_down
  - rt dmsync, 167
- rt dm\_sem\_init
  - rt dmsync, 168
- rt dm\_sem\_select\_bind
  - rt dmsync, 168
- rt dm\_sem\_timeddown
  - rt dmsync, 169
- rt dm\_sem\_up
  - rt dmsync, 170
- rt dm\_send
  - interdrv, 126
- rt dm\_sendmsg
  - interdrv, 126
- rt dm\_sendmsg\_handler\_t
  - devregister, 134
- rt dm\_sendto
  - interdrv, 127
- rt dm\_setsockopt
  - interdrv, 127
- rt dm\_shutdown
  - interdrv, 127
- rt dm\_socket
  - interdrv, 127
- rt dm\_socket\_handler\_t
  - devregister, 134
- rt dm\_strncpy\_from\_user
  - util, 188
- rt dm\_task\_busy\_sleep
  - rt dmtask, 142
- rt dm\_task\_current
  - rt dmtask, 142
- rt dm\_task\_destroy
  - rt dmtask, 143
- rt dm\_task\_init
  - rt dmtask, 143
- rt dm\_task\_join\_nrt

- rtmdtask, [144](#)
- rtmd\_task\_proc\_t
  - rtmdtask, [142](#)
- rtmd\_task\_set\_period
  - rtmdtask, [144](#)
- rtmd\_task\_set\_priority
  - rtmdtask, [145](#)
- rtmd\_task\_sleep
  - rtmdtask, [145](#)
- rtmd\_task\_sleep\_abs
  - rtmdtask, [146](#)
- rtmd\_task\_sleep\_until
  - rtmdtask, [146](#)
- rtmd\_task\_unblock
  - rtmdtask, [147](#)
- rtmd\_task\_wait\_period
  - rtmdtask, [147](#)
- RTDM\_TIMEOUT\_INFINITE
  - rtmd, [227](#)
- RTDM\_TIMEOUT\_NONE
  - rtmd, [227](#)
- rtmd\_timer\_destroy
  - rtmdtimer, [150](#)
- rtmd\_timer\_handler\_t
  - rtmdtimer, [149](#)
- rtmd\_timer\_init
  - rtmdtimer, [150](#)
- rtmd\_timer\_mode
  - rtmdtimer, [149](#)
- rtmd\_timer\_start
  - rtmdtimer, [150](#)
- rtmd\_timer\_start\_in\_handler
  - rtmdtimer, [151](#)
- rtmd\_timer\_stop
  - rtmdtimer, [152](#)
- rtmd\_timer\_stop\_in\_handler
  - rtmdtimer, [152](#)
- rtmd\_toseq\_init
  - rtmdsync, [170](#)
- rtmd\_write
  - interdrv, [127](#)
- rtmd\_write\_handler\_t
  - devregister, [135](#)
- rtmdclock
  - rtmd\_clock\_read, [139](#)
  - rtmd\_clock\_read\_monotonic, [140](#)
- rtdmirq
  - rtmd\_irq\_disable, [173](#)
  - rtmd\_irq\_enable, [174](#)
  - rtmd\_irq\_free, [174](#)
  - rtmd\_irq\_get\_arg, [172](#)
  - rtmd\_irq\_handler\_t, [173](#)
  - rtmd\_irq\_request, [175](#)
- rtmdsync
  - RTDM\_SELECTTYPE\_EXCEPT, [159](#)
  - RTDM\_SELECTTYPE\_READ, [159](#)
  - RTDM\_SELECTTYPE\_WRITE, [159](#)
  - rtmd\_event\_clear, [159](#)
  - rtmd\_event\_destroy, [160](#)
  - rtmd\_event\_init, [160](#)
  - rtmd\_event\_pulse, [160](#)
  - rtmd\_event\_select\_bind, [161](#)
  - rtmd\_event\_signal, [162](#)
  - rtmd\_event\_timedwait, [162](#)
  - rtmd\_event\_wait, [163](#)
  - RTDM\_EXECUTE\_ATOMICALY, [156](#)
  - rtmd\_lock\_get, [156](#)
  - rtmd\_lock\_get\_irqsave, [157](#)
  - rtmd\_lock\_init, [157](#)
  - rtmd\_lock\_irqrestore, [157](#)
  - rtmd\_lock\_irqsave, [158](#)
  - rtmd\_lock\_put, [158](#)
  - rtmd\_lock\_put\_irqrestore, [159](#)
  - rtmd\_mutex\_destroy, [163](#)
  - rtmd\_mutex\_init, [164](#)
  - rtmd\_mutex\_lock, [164](#)
  - rtmd\_mutex\_timedlock, [165](#)
  - rtmd\_mutex\_unlock, [166](#)
  - rtmd\_select\_bind, [166](#)
  - rtmd\_selecttype, [159](#)
  - rtmd\_sem\_destroy, [167](#)
  - rtmd\_sem\_down, [167](#)
  - rtmd\_sem\_init, [168](#)
  - rtmd\_sem\_select\_bind, [168](#)
  - rtmd\_sem\_timeddown, [169](#)
  - rtmd\_sem\_up, [170](#)
  - rtmd\_toseq\_init, [170](#)
- rtmdtask
  - rtmd\_task\_busy\_sleep, [142](#)
  - rtmd\_task\_current, [142](#)
  - rtmd\_task\_destroy, [143](#)
  - rtmd\_task\_init, [143](#)
  - rtmd\_task\_join\_nrt, [144](#)
  - rtmd\_task\_proc\_t, [142](#)
  - rtmd\_task\_set\_period, [144](#)
  - rtmd\_task\_set\_priority, [145](#)
  - rtmd\_task\_sleep, [145](#)
  - rtmd\_task\_sleep\_abs, [146](#)
  - rtmd\_task\_sleep\_until, [146](#)
  - rtmd\_task\_unblock, [147](#)
  - rtmd\_task\_wait\_period, [147](#)
- rtmdtimer
  - RTDM\_TIMERMODE\_ABSOLUTE, [149](#)
  - RTDM\_TIMERMODE\_REALTIME, [149](#)
  - RTDM\_TIMERMODE\_RELATIVE, [149](#)
  - rtmd\_timer\_destroy, [150](#)
  - rtmd\_timer\_handler\_t, [149](#)
  - rtmd\_timer\_init, [150](#)

- rtdm\_timer\_mode, [149](#)
  - rtdm\_timer\_start, [150](#)
  - rtdm\_timer\_start\_in\_handler, [151](#)
  - rtdm\_timer\_stop, [152](#)
  - rtdm\_timer\_stop\_in\_handler, [152](#)
- rthal\_apc\_alloc
  - hal, [15](#)
- rthal\_apc\_free
  - hal, [16](#)
- rthal\_timer\_release
  - arm/hal.c, [351](#)
- rthal\_timer\_request
  - arm/hal.c, [351](#)
- RTIOC\_DEVICE\_INFO
  - profiles, [190](#)
- RTIOC\_PURGE
  - profiles, [190](#)
- rtipc
  - bind\_\_AF\_RTIPC, [250](#)
  - BUFP\_BUFSZ, [243](#)
  - BUFP\_LABEL, [243](#)
  - close\_\_AF\_RTIPC, [251](#)
  - connect\_\_AF\_RTIPC, [252](#)
  - getpeername\_\_AF\_RTIPC, [252](#)
  - getsockname\_\_AF\_RTIPC, [253](#)
  - getsockopt\_\_AF\_RTIPC, [253](#)
  - IDDP\_LABEL, [244](#)
  - IDDP\_POOLSZ, [245](#)
  - IPCPROTO\_BUFP, [250](#)
  - IPCPROTO\_IDDP, [250](#)
  - IPCPROTO\_IPC, [250](#)
  - IPCPROTO\_XDDP, [250](#)
  - recvmsg\_\_AF\_RTIPC, [253](#)
  - sendmsg\_\_AF\_RTIPC, [254](#)
  - setsockopt\_\_AF\_RTIPC, [255](#)
  - SO\_RCVTIMEO, [245](#)
  - SO\_SNDTIMEO, [245](#)
  - socket\_\_AF\_RTIPC, [255](#)
  - XDDP\_BUFSZ, [246](#)
  - XDDP\_EVTDOWN, [247](#)
  - XDDP\_EVTIN, [247](#)
  - XDDP\_EVTNOBUF, [247](#)
  - XDDP\_EVTOUT, [247](#)
  - XDDP\_LABEL, [247](#)
  - XDDP\_MONITOR, [248](#)
  - XDDP\_POOLSZ, [249](#)
- rtipc\_port\_label, [278](#)
  - label, [279](#)
- rtser\_config, [279](#)
- rtser\_event, [280](#)
- RTSER\_RTIOC\_BREAK\_CTL
  - rtserial, [261](#)
- RTSER\_RTIOC\_GET\_CONFIG
  - rtserial, [262](#)
- RTSER\_RTIOC\_GET\_CONTROL
  - rtserial, [262](#)
- RTSER\_RTIOC\_GET\_STATUS
  - rtserial, [263](#)
- RTSER\_RTIOC\_SET\_CONFIG
  - rtserial, [263](#)
- RTSER\_RTIOC\_SET\_CONTROL
  - rtserial, [264](#)
- RTSER\_RTIOC\_WAIT\_EVENT
  - rtserial, [264](#)
- rtser\_status, [281](#)
- rtserial
  - RTSER\_RTIOC\_BREAK\_CTL, [261](#)
  - RTSER\_RTIOC\_GET\_CONFIG, [262](#)
  - RTSER\_RTIOC\_GET\_CONTROL, [262](#)
  - RTSER\_RTIOC\_GET\_STATUS, [263](#)
  - RTSER\_RTIOC\_SET\_CONFIG, [263](#)
  - RTSER\_RTIOC\_SET\_CONTROL, [264](#)
  - RTSER\_RTIOC\_WAIT\_EVENT, [264](#)
- Sched, [267](#)
- sched
  - xnpod, [283](#)
  - xnsched\_rotate, [268](#)
- select
  - xnselect, [92](#)
  - xnselect\_bind, [92](#)
  - xnselect\_destroy, [93](#)
  - xnselect\_init, [93](#)
  - xnselect\_signal, [93](#)
  - xnselector\_destroy, [94](#)
  - xnselector\_init, [94](#)
- sem\_close
  - cobalt\_sem, [192](#)
- sem\_destroy
  - cobalt\_sem, [192](#)
- sem\_getvalue
  - cobalt\_sem, [192](#)
- sem\_open
  - cobalt\_sem, [193](#)
- sem\_post
  - cobalt\_sem, [194](#)
- sem\_timedwait
  - cobalt\_sem, [194](#)
- sem\_trywait
  - cobalt\_sem, [195](#)
- sem\_unlink
  - cobalt\_sem, [195](#)
- sem\_wait
  - cobalt\_sem, [196](#)
- Semaphores services., [191](#)
- sendmsg\_\_AF\_RTIPC
  - rtipc, [254](#)
- seq

- xnvfile\_regular\_iterator, [289](#)
- xnvfile\_snapshot\_iterator, [295](#)
- Serial Devices, [256](#)
- setsockopt\_\_AF\_RTIPC
  - rtipc, [255](#)
- shadow
  - xnshadow\_harden, [95](#)
  - xnshadow\_map, [95](#)
  - xnshadow\_ppd\_get, [96](#)
  - xnshadow\_relax, [97](#)
- show
  - xnvfile\_regular\_ops, [291](#)
  - xnvfile\_snapshot\_ops, [298](#)
- SIOCGCANBAUDRATE
  - rtcan, [219](#)
- SIOCGCANCTRLMODE
  - rtcan, [220](#)
- SIOCGCANCUSTOMBITTIME
  - rtcan, [220](#)
- SIOCGCANSTATE
  - rtcan, [221](#)
- SIOCGIFINDEX
  - rtcan, [221](#)
- SIOGSCANBAUDRATE
  - rtcan, [222](#)
- SIOGSCANCTRLMODE
  - rtcan, [222](#)
- SIOGSCANCUSTOMBITTIME
  - rtcan, [223](#)
- SIOGSCANMODE
  - rtcan, [224](#)
- sipc\_port
  - sockaddr\_ipc, [282](#)
- SO\_RCVTIMEO
  - rtipc, [245](#)
- SO\_SNDTIMEO
  - rtipc, [245](#)
- sockaddr\_can, [281](#)
  - can\_ifindex, [281](#)
- sockaddr\_ipc, [282](#)
  - sipc\_port, [282](#)
- socket\_\_AF\_RTIPC
  - rtipc, [255](#)
- socket\_rt
  - rtdm\_device, [276](#)
- SOL\_CAN\_RAW
  - rtcan, [225](#)
- state
  - xnthread\_info, [287](#)
- status
  - xnpod, [283](#)
  - xnsched, [285](#)
- store
  - xnvfile\_regular\_ops, [292](#)
  - xnvfile\_snapshot\_ops, [298](#)
- synch
  - xnsynch\_acquire, [99](#)
  - xnsynch\_clear\_boost, [99](#)
  - xnsynch\_flush, [100](#)
  - xnsynch\_forget\_sleeper, [101](#)
  - xnsynch\_init, [101](#)
  - xnsynch\_peek\_pendq, [102](#)
  - xnsynch\_release, [103](#)
  - xnsynch\_release\_all\_ownerships, [103](#)
  - xnsynch\_requeue\_sleeper, [104](#)
  - xnsynch\_sleep\_on, [104](#)
  - xnsynch\_wakeup\_one\_sleeper, [105](#)
  - xnsynch\_wakeup\_this\_sleeper, [106](#)
- Synchronisation Services, [153](#)
- syscalls
  - xnthread\_info, [287](#)
- System clock services., [48](#)
- Task Services, [140](#)
- tdeleteq
  - xnpod, [283](#)
- Testing Devices, [265](#)
- Thread cancellation., [16](#)
- Thread information flags., [13](#)
- Thread state flags., [11](#)
- Thread synchronization services., [98](#)
- threadq
  - xnpod, [283](#)
- Threads management services., [197](#)
- timer
  - xntimer\_destroy, [108](#)
  - xntimer\_freeze, [108](#)
  - xntimer\_get\_date, [108](#)
  - xntimer\_get\_interval, [109](#)
  - xntimer\_get\_overruns, [109](#)
  - xntimer\_get\_timeout, [110](#)
  - xntimer\_init, [110](#)
  - xntimer\_start, [110](#)
  - xntimer\_stop, [111](#)
  - xntimer\_tick, [112](#)
- Timer Services, [148](#)
- Timer services., [107](#)
- timer\_create
  - cobalt\_time, [21](#)
- timer\_gettime
  - cobalt\_time, [22](#)
- timer\_settime
  - cobalt\_time, [22](#)
- timerlck
  - xnpod, [283](#)
- tstartq
  - xnpod, [284](#)
- tswitchq

- xnpod, [284](#)
- User API, [228](#)
- userapi
  - rt\_dev\_accept, [230](#)
  - rt\_dev\_bind, [230](#)
  - rt\_dev\_close, [231](#)
  - rt\_dev\_connect, [231](#)
  - rt\_dev\_getpeername, [232](#)
  - rt\_dev\_getsockname, [232](#)
  - rt\_dev\_getsockopt, [232](#)
  - rt\_dev\_ioctl, [233](#)
  - rt\_dev\_listen, [233](#)
  - rt\_dev\_open, [234](#)
  - rt\_dev\_read, [234](#)
  - rt\_dev\_recv, [235](#)
  - rt\_dev\_recvfrom, [235](#)
  - rt\_dev\_recvmsg, [236](#)
  - rt\_dev\_send, [236](#)
  - rt\_dev\_sendmsg, [237](#)
  - rt\_dev\_sendto, [237](#)
  - rt\_dev\_setsockopt, [238](#)
  - rt\_dev\_shutdown, [238](#)
  - rt\_dev\_socket, [239](#)
  - rt\_dev\_write, [239](#)
- util
  - rtdm\_copy\_from\_user, [179](#)
  - rtdm\_copy\_to\_user, [180](#)
  - rtdm\_free, [181](#)
  - rtdm\_in\_rt\_context, [181](#)
  - rtdm\_iomap\_to\_user, [181](#)
  - rtdm\_malloc, [182](#)
  - rtdm\_mmap\_to\_user, [183](#)
  - rtdm\_munmap, [184](#)
  - rtdm\_printk, [184](#)
  - rtdm\_read\_user\_ok, [185](#)
  - rtdm\_rt\_capable, [185](#)
  - rtdm\_rw\_user\_ok, [186](#)
  - rtdm\_safe\_copy\_from\_user, [186](#)
  - rtdm\_safe\_copy\_to\_user, [187](#)
  - rtdm\_strncpy\_from\_user, [188](#)
- Utility Services, [178](#)
- vfile
  - nkvfroot, [118](#)
  - xnvfile\_destroy, [114](#)
  - xnvfile\_get\_blob, [114](#)
  - xnvfile\_get\_integer, [115](#)
  - xnvfile\_get\_string, [115](#)
  - xnvfile\_init\_dir, [116](#)
  - xnvfile\_init\_link, [116](#)
  - xnvfile\_init\_regular, [117](#)
  - xnvfile\_init\_snapshot, [117](#)
  - xnvfile\_regular\_iterator, [289](#)
  - xnvfile\_snapshot\_iterator, [295](#)
- Virtual file services, [112](#)
- XDDP\_BUFSZ
  - rtipc, [246](#)
- XDDP\_EVTDOWN
  - rtipc, [247](#)
- XDDP\_EVTIN
  - rtipc, [247](#)
- XDDP\_EVTNOBUF
  - rtipc, [247](#)
- XDDP\_EVTOUT
  - rtipc, [247](#)
- XDDP\_LABEL
  - rtipc, [247](#)
- XDDP\_MONITOR
  - rtipc, [248](#)
- XDDP\_POOLSZ
  - rtipc, [249](#)
- xnbufd\_copy\_from\_kmem
  - bufd, [41](#)
- xnbufd\_copy\_to\_kmem
  - bufd, [42](#)
- xnbufd\_invalidate
  - bufd, [43](#)
- xnbufd\_map\_kread
  - bufd, [43](#)
- xnbufd\_map\_kwrite
  - bufd, [44](#)
- xnbufd\_map\_uread
  - bufd, [44](#)
- xnbufd\_map\_uwrite
  - bufd, [45](#)
- xnbufd\_reset
  - bufd, [45](#)
- xnbufd\_unmap\_kread
  - bufd, [46](#)
- xnbufd\_unmap\_kwrite
  - bufd, [46](#)
- xnbufd\_unmap\_uread
  - bufd, [47](#)
- xnbufd\_unmap\_uwrite
  - bufd, [47](#)
- xnclock\_adjust
  - clock, [48](#)
- xnheap\_alloc
  - heap, [50](#)
- xnheap\_destroy
  - heap, [51](#)
- xnheap\_extend
  - heap, [51](#)
- xnheap\_free
  - heap, [52](#)
- xnheap\_init

- heap, 52
- xnheap\_schedule\_free
  - heap, 53
- xnheap\_set\_label
  - heap, 54
- xnheap\_test\_and\_free
  - heap, 54
- XNHeld
  - nucleus\_state\_flags, 12
- xnintr\_affinity
  - intr, 56
- xnintr\_attach
  - intr, 56
- xnintr\_destroy
  - intr, 57
- xnintr\_detach
  - intr, 57
- xnintr\_disable
  - intr, 58
- xnintr\_enable
  - intr, 58
- xnintr\_init
  - intr, 59
- XNLOCK
  - nucleus\_state\_flags, 12
- xnmap\_create
  - map, 62
- xnmap\_delete
  - map, 62
- xnmap\_enter
  - map, 63
- xnmap\_fetch
  - map, 63
- xnmap\_fetch\_noccheck
  - map, 64
- xnmap\_remove
  - map, 64
- XNMIGRATE
  - nucleus\_state\_flags, 13
- XNPEND
  - nucleus\_state\_flags, 13
- xnpod, 282
  - refcnt, 283
  - sched, 283
  - status, 283
  - tdeleteq, 283
  - threadq, 283
  - timerlck, 283
  - tstartq, 284
  - tswitchq, 284
- xnpod\_abort\_thread
  - pod, 67
- xnpod\_add\_hook
  - pod, 68
- xnpod\_delete\_thread
  - pod, 69
- xnpod\_disable\_timesource
  - pod, 69
- xnpod\_dispatch\_signals
  - pod, 70
- xnpod\_enable\_timesource
  - pod, 70
- xnpod\_handle\_exception
  - pod, 71
- xnpod\_init
  - pod, 71
- xnpod\_init\_thread
  - pod, 71
- xnpod\_migrate\_thread
  - pod, 73
- xnpod\_remove\_hook
  - pod, 73
- xnpod\_resume\_thread
  - pod, 74
- xnpod\_schedule
  - pod, 75
- xnpod\_set\_thread\_mode
  - pod, 76
- xnpod\_set\_thread\_periodic
  - pod, 77
- xnpod\_set\_thread\_schedparam
  - pod, 78
- xnpod\_set\_thread\_tslice
  - pod, 79
- xnpod\_shutdown
  - pod, 79
- xnpod\_start\_thread
  - pod, 80
- xnpod\_stop\_thread
  - pod, 81
- xnpod\_suspend\_thread
  - pod, 81
- xnpod\_unblock\_thread
  - pod, 83
- xnpod\_wait\_thread\_period
  - pod, 83
- xnpod\_welcome\_thread
  - pod, 84
- XNREADY
  - nucleus\_state\_flags, 13
- xnregistry\_bind
  - registry, 86
- xnregistry\_enter
  - registry, 87
- xnregistry\_fetch
  - registry, 87
- xnregistry\_get
  - registry, 88

- xnregistry\_put
  - registry, 88
- xnregistry\_remove
  - registry, 89
- xnregistry\_remove\_safe
  - registry, 90
- xnsched, 284
  - curr, 284
  - htimer, 284
  - inesting, 285
  - lflags, 285
  - rootcb, 285
  - rt, 285
  - status, 285
- xnsched\_rotate
  - sched, 268
- xnselect
  - select, 92
- xnselect\_bind
  - select, 92
- xnselect\_destroy
  - select, 93
- xnselect\_init
  - select, 93
- xnselect\_signal
  - select, 93
- xnselector\_destroy
  - select, 94
- xnselector\_init
  - select, 94
- xnshadow\_harden
  - shadow, 95
- xnshadow\_map
  - shadow, 95
- xnshadow\_ppd\_get
  - shadow, 96
- xnshadow\_relax
  - shadow, 97
- XNSUSP
  - nucleus\_state\_flags, 13
- xnsynch\_acquire
  - synch, 99
- xnsynch\_clear\_boost
  - synch, 99
- xnsynch\_flush
  - synch, 100
- xnsynch\_forget\_sleeper
  - synch, 101
- xnsynch\_init
  - synch, 101
- xnsynch\_peek\_pendq
  - synch, 102
- xnsynch\_release
  - synch, 103
- xnsynch\_release\_all\_ownerships
  - synch, 103
- xnsynch\_requeue\_sleeper
  - synch, 104
- xnsynch\_sleep\_on
  - synch, 104
- xnsynch\_wakeup\_one\_sleeper
  - synch, 105
- xnsynch\_wakeup\_this\_sleeper
  - synch, 106
- xnthread\_info, 285
  - affinity, 286
  - bprio, 286
  - cprio, 286
  - cpu, 286
  - ctxswitches, 287
  - exectime, 287
  - modeswitches, 287
  - name, 287
  - pagefaults, 287
  - relpoint, 287
  - state, 287
  - syscalls, 287
- xntimer\_destroy
  - timer, 108
- xntimer\_freeze
  - timer, 108
- xntimer\_get\_date
  - timer, 108
- xntimer\_get\_interval
  - timer, 109
- xntimer\_get\_overruns
  - timer, 109
- xntimer\_get\_timeout
  - timer, 110
- xntimer\_init
  - timer, 110
- xntimer\_start
  - timer, 110
- xntimer\_stop
  - timer, 111
- xntimer\_tick
  - timer, 112
- xnvfile\_destroy
  - vfile, 114
- xnvfile\_get\_blob
  - vfile, 114
- xnvfile\_get\_integer
  - vfile, 115
- xnvfile\_get\_string
  - vfile, 115
- xnvfile\_init\_dir
  - vfile, 116
- xnvfile\_init\_link



- vfile, [116](#)
- xnvfile\_init\_regular
  - vfile, [117](#)
- xnvfile\_init\_snapshot
  - vfile, [117](#)
- xnvfile\_lock\_ops, [287](#)
  - get, [288](#)
  - put, [288](#)
- xnvfile\_regular\_iterator, [288](#)
  - pos, [289](#)
  - private, [289](#)
  - seq, [289](#)
  - vfile, [289](#)
- xnvfile\_regular\_ops, [289](#)
  - begin, [290](#)
  - end, [290](#)
  - next, [291](#)
  - rewind, [291](#)
  - show, [291](#)
  - store, [292](#)
- xnvfile\_rev\_tag, [293](#)
  - rev, [293](#)
- xnvfile\_snapshot, [293](#)
- xnvfile\_snapshot\_iterator, [294](#)
  - databuf, [295](#)
  - endfn, [295](#)
  - nrdata, [295](#)
  - private, [295](#)
  - seq, [295](#)
  - vfile, [295](#)
- xnvfile\_snapshot\_ops, [296](#)
  - begin, [296](#)
  - end, [296](#)
  - next, [297](#)
  - rewind, [297](#)
  - show, [298](#)
  - store, [298](#)