

# Xenomai nanokernel API

## 2.5.93

Generated by Doxygen 1.7.1

Wed Sep 28 2011 00:02:26



# Contents

<b>1</b>	<b>Module Index</b>	<b>1</b>
1.1	Modules . . . . .	1
<b>2</b>	<b>Data Structure Index</b>	<b>3</b>
2.1	Data Structures . . . . .	3
<b>3</b>	<b>File Index</b>	<b>5</b>
3.1	File List . . . . .	5
<b>4</b>	<b>Module Documentation</b>	<b>7</b>
4.1	Thread state flags. . . . .	7
4.1.1	Detailed Description . . . . .	9
4.1.2	Define Documentation . . . . .	9
4.1.2.1	XNHELD . . . . .	9
4.1.2.2	XNLOCK . . . . .	9
4.1.2.3	XNMIGRATE . . . . .	9
4.1.2.4	XNPEND . . . . .	9
4.1.2.5	XNREADY . . . . .	9
4.1.2.6	XNSUSP . . . . .	10
4.2	Thread information flags. . . . .	10
4.2.1	Detailed Description . . . . .	11
4.3	Buffer descriptors. . . . .	11
4.3.1	Detailed Description . . . . .	12
4.3.2	Function Documentation . . . . .	14
4.3.2.1	xnbufd_copy_from_kmem . . . . .	14
4.3.2.2	xnbufd_copy_to_kmem . . . . .	15
4.3.2.3	xnbufd_invalidate . . . . .	16
4.3.2.4	xnbufd_map_kread . . . . .	16
4.3.2.5	xnbufd_map_kwrite . . . . .	17

4.3.2.6	<a href="#">xnbufd_map_uread</a>	17
4.3.2.7	<a href="#">xnbufd_map_uwrite</a>	18
4.3.2.8	<a href="#">xnbufd_reset</a>	18
4.3.2.9	<a href="#">xnbufd_unmap_kread</a>	19
4.3.2.10	<a href="#">xnbufd_unmap_kwrite</a>	19
4.3.2.11	<a href="#">xnbufd_unmap_uread</a>	20
4.3.2.12	<a href="#">xnbufd_unmap_uwrite</a>	20
4.4	Dynamic memory allocation services.	21
4.4.1	Detailed Description	22
4.4.2	Function Documentation	22
4.4.2.1	<a href="#">xnheap_alloc</a>	22
4.4.2.2	<a href="#">xnheap_destroy</a>	23
4.4.2.3	<a href="#">xnheap_extend</a>	23
4.4.2.4	<a href="#">xnheap_free</a>	24
4.4.2.5	<a href="#">xnheap_init</a>	25
4.4.2.6	<a href="#">xnheap_schedule_free</a>	26
4.4.2.7	<a href="#">xnheap_set_label</a>	26
4.4.2.8	<a href="#">xnheap_test_and_free</a>	27
4.5	Interrupt management.	27
4.5.1	Detailed Description	28
4.5.2	Function Documentation	28
4.5.2.1	<a href="#">xnintr_affinity</a>	28
4.5.2.2	<a href="#">xnintr_attach</a>	29
4.5.2.3	<a href="#">xnintr_destroy</a>	29
4.5.2.4	<a href="#">xnintr_detach</a>	30
4.5.2.5	<a href="#">xnintr_disable</a>	31
4.5.2.6	<a href="#">xnintr_enable</a>	31
4.5.2.7	<a href="#">xnintr_init</a>	32
4.6	Lightweight key-to-object mapping service	33
4.6.1	Detailed Description	34
4.6.2	Function Documentation	34
4.6.2.1	<a href="#">xnmap_create</a>	34
4.6.2.2	<a href="#">xnmap_delete</a>	35
4.6.2.3	<a href="#">xnmap_enter</a>	35
4.6.2.4	<a href="#">xnmap_fetch</a>	36
4.6.2.5	<a href="#">xnmap_fetch_nocheck</a>	37

4.6.2.6	<a href="#">xnmap_remove</a>	37
4.7	<a href="#">Xenomai nucleus.</a>	38
4.7.1	<a href="#">Detailed Description</a>	39
4.8	<a href="#">Real-time pod services.</a>	39
4.8.1	<a href="#">Detailed Description</a>	41
4.8.2	<a href="#">Function Documentation</a>	42
4.8.2.1	<a href="#">xnpod_abort_thread</a>	42
4.8.2.2	<a href="#">xnpod_add_hook</a>	42
4.8.2.3	<a href="#">xnpod_delete_thread</a>	43
4.8.2.4	<a href="#">xnpod_disable_timesource</a>	44
4.8.2.5	<a href="#">xnpod_dispatch_signals</a>	44
4.8.2.6	<a href="#">xnpod_enable_timesource</a>	44
4.8.2.7	<a href="#">xnpod_init</a>	45
4.8.2.8	<a href="#">xnpod_init_thread</a>	46
4.8.2.9	<a href="#">xnpod_migrate_thread</a>	47
4.8.2.10	<a href="#">xnpod_remove_hook</a>	48
4.8.2.11	<a href="#">xnpod_restart_thread</a>	48
4.8.2.12	<a href="#">xnpod_resume_thread</a>	49
4.8.2.13	<a href="#">xnpod_schedule</a>	50
4.8.2.14	<a href="#">xnpod_set_thread_mode</a>	51
4.8.2.15	<a href="#">xnpod_set_thread_periodic</a>	52
4.8.2.16	<a href="#">xnpod_set_thread_schedparam</a>	53
4.8.2.17	<a href="#">xnpod_set_thread_tslice</a>	54
4.8.2.18	<a href="#">xnpod_shutdown</a>	54
4.8.2.19	<a href="#">xnpod_start_thread</a>	55
4.8.2.20	<a href="#">xnpod_stop_thread</a>	56
4.8.2.21	<a href="#">xnpod_suspend_thread</a>	57
4.8.2.22	<a href="#">xnpod_trap_fault</a>	58
4.8.2.23	<a href="#">xnpod_unblock_thread</a>	58
4.8.2.24	<a href="#">xnpod_wait_thread_period</a>	59
4.8.2.25	<a href="#">xnpod_welcome_thread</a>	60
4.9	<a href="#">Registry services.</a>	60
4.9.1	<a href="#">Detailed Description</a>	61
4.9.2	<a href="#">Function Documentation</a>	61
4.9.2.1	<a href="#">xnregistry_bind</a>	61
4.9.2.2	<a href="#">xnregistry_enter</a>	62

4.9.2.3	xnregistry_fetch . . . . .	63
4.9.2.4	xnregistry_get . . . . .	64
4.9.2.5	xnregistry_put . . . . .	64
4.9.2.6	xnregistry_remove . . . . .	65
4.9.2.7	xnregistry_remove_safe . . . . .	65
4.10	File descriptors events multiplexing services. . . . .	67
4.10.1	Detailed Description . . . . .	68
4.10.2	Function Documentation . . . . .	68
4.10.2.1	xnselect . . . . .	68
4.10.2.2	xnselect_bind . . . . .	69
4.10.2.3	xnselect_destroy . . . . .	69
4.10.2.4	xnselect_init . . . . .	69
4.10.2.5	xnselector_destroy . . . . .	70
4.10.2.6	xnselector_init . . . . .	70
4.11	Real-time shadow services. . . . .	70
4.11.1	Detailed Description . . . . .	71
4.11.2	Function Documentation . . . . .	71
4.11.2.1	xnshadow_harden . . . . .	71
4.11.2.2	xnshadow_map . . . . .	71
4.11.2.3	xnshadow_ppd_get . . . . .	73
4.11.2.4	xnshadow_relax . . . . .	73
4.12	Thread synchronization services. . . . .	74
4.12.1	Detailed Description . . . . .	75
4.12.2	Function Documentation . . . . .	75
4.12.2.1	xnsynch_acquire . . . . .	75
4.12.2.2	xnsynch_clear_boost . . . . .	76
4.12.2.3	xnsynch_flush . . . . .	76
4.12.2.4	xnsynch_forget_sleeper . . . . .	77
4.12.2.5	xnsynch_init . . . . .	78
4.12.2.6	xnsynch_peek_pendq . . . . .	79
4.12.2.7	xnsynch_release . . . . .	79
4.12.2.8	xnsynch_release_all_ownerships . . . . .	80
4.12.2.9	xnsynch_requeue_sleeper . . . . .	80
4.12.2.10	xnsynch_sleep_on . . . . .	81
4.12.2.11	xnsynch_wakeup_one_sleeper . . . . .	81
4.12.2.12	xnsynch_wakeup_this_sleeper . . . . .	82

4.13 Time base services. . . . .	83
4.13.1 Detailed Description . . . . .	84
4.13.2 Function Documentation . . . . .	84
4.13.2.1 xntbase_adjust_time . . . . .	84
4.13.2.2 xntbase_alloc . . . . .	85
4.13.2.3 xntbase_convert . . . . .	86
4.13.2.4 xntbase_free . . . . .	87
4.13.2.5 xntbase_get_time . . . . .	87
4.13.2.6 xntbase_start . . . . .	88
4.13.2.7 xntbase_stop . . . . .	88
4.13.2.8 xntbase_switch . . . . .	89
4.13.2.9 xntbase_tick . . . . .	89
4.13.2.10 xntbase_update . . . . .	90
4.14 Timer services. . . . .	91
4.14.1 Detailed Description . . . . .	92
4.14.2 Function Documentation . . . . .	92
4.14.2.1 xntimer_destroy . . . . .	92
4.14.2.2 xntimer_freeze . . . . .	92
4.14.2.3 xntimer_get_date . . . . .	93
4.14.2.4 xntimer_get_interval . . . . .	94
4.14.2.5 xntimer_get_overruns . . . . .	94
4.14.2.6 xntimer_get_timeout . . . . .	95
4.14.2.7 xntimer_init . . . . .	95
4.14.2.8 xntimer_start . . . . .	96
4.14.2.9 xntimer_stop . . . . .	97
4.14.2.10 xntimer_tick_aperiodic . . . . .	97
4.14.2.11 xntimer_tick_periodic . . . . .	98
4.15 Virtual file services . . . . .	98
4.15.1 Detailed Description . . . . .	100
4.15.2 Function Documentation . . . . .	101
4.15.2.1 xnvfile_destroy . . . . .	101
4.15.2.2 xnvfile_get_blob . . . . .	101
4.15.2.3 xnvfile_get_integer . . . . .	101
4.15.2.4 xnvfile_get_string . . . . .	102
4.15.2.5 xnvfile_init_dir . . . . .	102
4.15.2.6 xnvfile_init_link . . . . .	103

4.15.2.7	xnvfile_init_regular . . . . .	103
4.15.2.8	xnvfile_init_snapshot . . . . .	104
4.15.3	Variable Documentation . . . . .	105
4.15.3.1	nkvfroot . . . . .	105
4.15.3.2	nkvfroot . . . . .	105
4.16	HAL. . . . .	105
4.16.1	Detailed Description . . . . .	107
4.16.2	Function Documentation . . . . .	107
4.16.2.1	rthal_apc_alloc . . . . .	107
4.16.2.2	rthal_apc_free . . . . .	108
4.16.2.3	rthal_irq_disable . . . . .	108
4.16.2.4	rthal_irq_enable . . . . .	108
4.16.2.5	rthal_irq_host_release . . . . .	109
4.16.2.6	rthal_irq_host_request . . . . .	109
4.16.2.7	rthal_irq_release . . . . .	110
4.16.2.8	rthal_irq_request . . . . .	111
4.16.2.9	rthal_timer_release . . . . .	111
4.16.2.10	rthal_timer_request . . . . .	112
4.16.2.11	rthal_trap_catch . . . . .	113
4.17	Sched . . . . .	113
4.17.1	Function Documentation . . . . .	114
4.17.1.1	xnsched_rotate . . . . .	114
<b>5</b>	<b>Data Structure Documentation</b>	<b>115</b>
5.1	xnpod Struct Reference . . . . .	115
5.1.1	Detailed Description . . . . .	116
5.1.2	Field Documentation . . . . .	116
5.1.2.1	refcnt . . . . .	116
5.1.2.2	sched . . . . .	116
5.1.2.3	status . . . . .	116
5.1.2.4	tdeleteq . . . . .	116
5.1.2.5	threadq . . . . .	116
5.1.2.6	timerlck . . . . .	116
5.1.2.7	tsliced . . . . .	116
5.1.2.8	tslicer . . . . .	117
5.1.2.9	tstartq . . . . .	117
5.1.2.10	tswitchq . . . . .	117



5.2	xnsched Struct Reference	117
5.2.1	Detailed Description	117
5.2.2	Field Documentation	117
5.2.2.1	curr	117
5.2.2.2	htimer	118
5.2.2.3	inesting	118
5.2.2.4	lflags	118
5.2.2.5	rootcb	118
5.2.2.6	rt	118
5.2.2.7	status	118
5.3	xnthread_info Struct Reference	118
5.3.1	Detailed Description	119
5.3.2	Field Documentation	119
5.3.2.1	affinity	119
5.3.2.2	bprio	119
5.3.2.3	cprio	119
5.3.2.4	cpu	120
5.3.2.5	ctxswitches	120
5.3.2.6	exectime	120
5.3.2.7	modeswitches	120
5.3.2.8	name	120
5.3.2.9	pagefaults	120
5.3.2.10	relpoint	120
5.3.2.11	state	120
5.4	xnvfile_lock_ops Struct Reference	120
5.4.1	Detailed Description	121
5.4.2	Field Documentation	121
5.4.2.1	get	121
5.4.2.2	put	121
5.5	xnvfile_regular_iterator Struct Reference	121
5.5.1	Detailed Description	122
5.5.2	Field Documentation	122
5.5.2.1	pos	122
5.5.2.2	private	122
5.5.2.3	seq	122
5.5.2.4	vfile	122

5.6	xnvfile_regular_ops Struct Reference	122
5.6.1	Detailed Description	123
5.6.2	Field Documentation	123
5.6.2.1	begin	123
5.6.2.2	end	123
5.6.2.3	next	124
5.6.2.4	rewind	124
5.6.2.5	show	125
5.6.2.6	store	125
5.7	xnvfile_rev_tag Struct Reference	126
5.7.1	Detailed Description	126
5.7.2	Field Documentation	126
5.7.2.1	rev	126
5.8	xnvfile_snapshot Struct Reference	126
5.8.1	Detailed Description	127
5.9	xnvfile_snapshot_iterator Struct Reference	127
5.9.1	Detailed Description	128
5.9.2	Field Documentation	128
5.9.2.1	databuf	128
5.9.2.2	endfn	128
5.9.2.3	nrdata	128
5.9.2.4	private	128
5.9.2.5	seq	128
5.9.2.6	vfile	128
5.10	xnvfile_snapshot_ops Struct Reference	129
5.10.1	Detailed Description	129
5.10.2	Field Documentation	129
5.10.2.1	begin	129
5.10.2.2	end	130
5.10.2.3	next	130
5.10.2.4	rewind	130
5.10.2.5	show	131
5.10.2.6	store	132
<b>6</b>	<b>File Documentation</b>	<b>133</b>
6.1	include/nucleus/bufd.h File Reference	133
6.1.1	Detailed Description	134

6.2	<a href="#">include/nucleus/hostrt.h File Reference</a>	135
6.2.1	<a href="#">Detailed Description</a>	136
6.3	<a href="#">include/nucleus/map.h File Reference</a>	136
6.3.1	<a href="#">Detailed Description</a>	137
6.4	<a href="#">include/nucleus/pod.h File Reference</a>	138
6.4.1	<a href="#">Detailed Description</a>	140
6.5	<a href="#">include/nucleus/registry.h File Reference</a>	140
6.5.1	<a href="#">Detailed Description</a>	142
6.6	<a href="#">include/nucleus/sched-idle.h File Reference</a>	142
6.6.1	<a href="#">Detailed Description</a>	142
6.7	<a href="#">include/nucleus/sched-rt.h File Reference</a>	143
6.7.1	<a href="#">Detailed Description</a>	143
6.8	<a href="#">include/nucleus/sched-sporadic.h File Reference</a>	143
6.8.1	<a href="#">Detailed Description</a>	144
6.9	<a href="#">include/nucleus/sched-tp.h File Reference</a>	144
6.9.1	<a href="#">Detailed Description</a>	144
6.10	<a href="#">include/nucleus/sched.h File Reference</a>	145
6.10.1	<a href="#">Detailed Description</a>	146
6.11	<a href="#">include/nucleus/select.h File Reference</a>	146
6.11.1	<a href="#">Detailed Description</a>	147
6.12	<a href="#">include/nucleus/timebase.h File Reference</a>	148
6.12.1	<a href="#">Detailed Description</a>	149
6.13	<a href="#">include/nucleus/timer.h File Reference</a>	150
6.13.1	<a href="#">Detailed Description</a>	151
6.14	<a href="#">include/nucleus/vdso.h File Reference</a>	151
6.14.1	<a href="#">Detailed Description</a>	152
6.15	<a href="#">include/nucleus/vfile.h File Reference</a>	152
6.15.1	<a href="#">Detailed Description</a>	154
6.16	<a href="#">ksrc/arch/arm/hal.c File Reference</a>	155
6.16.1	<a href="#">Detailed Description</a>	155
6.17	<a href="#">ksrc/arch/blackfin/hal.c File Reference</a>	156
6.17.1	<a href="#">Detailed Description</a>	156
6.18	<a href="#">ksrc/arch/generic/hal.c File Reference</a>	157
6.18.1	<a href="#">Detailed Description</a>	157
6.19	<a href="#">ksrc/arch/nios2/hal.c File Reference</a>	157
6.19.1	<a href="#">Detailed Description</a>	158

6.20	<a href="#">ksrc/arch/powerpc/hal.c File Reference</a>	158
6.20.1	<a href="#">Detailed Description</a>	159
6.21	<a href="#">ksrc/arch/sh/hal.c File Reference</a>	160
6.21.1	<a href="#">Detailed Description</a>	160
6.22	<a href="#">ksrc/arch/x86/hal-common.c File Reference</a>	161
6.22.1	<a href="#">Detailed Description</a>	161
6.23	<a href="#">ksrc/arch/x86/hal_32.c File Reference</a>	162
6.23.1	<a href="#">Detailed Description</a>	162
6.24	<a href="#">ksrc/arch/x86/hal_64.c File Reference</a>	163
6.24.1	<a href="#">Detailed Description</a>	163
6.25	<a href="#">ksrc/arch/x86/smi.c File Reference</a>	163
6.25.1	<a href="#">Detailed Description</a>	163
6.26	<a href="#">ksrc/nucleus/bufd.c File Reference</a>	164
6.26.1	<a href="#">Detailed Description</a>	165
6.27	<a href="#">ksrc/nucleus/heap.c File Reference</a>	165
6.27.1	<a href="#">Detailed Description</a>	166
6.28	<a href="#">ksrc/nucleus/intr.c File Reference</a>	166
6.28.1	<a href="#">Detailed Description</a>	167
6.29	<a href="#">ksrc/nucleus/map.c File Reference</a>	168
6.29.1	<a href="#">Detailed Description</a>	169
6.30	<a href="#">ksrc/nucleus/pod.c File Reference</a>	169
6.30.1	<a href="#">Detailed Description</a>	171
6.31	<a href="#">ksrc/nucleus/registry.c File Reference</a>	171
6.31.1	<a href="#">Detailed Description</a>	172
6.32	<a href="#">ksrc/nucleus/sched-idle.c File Reference</a>	173
6.32.1	<a href="#">Detailed Description</a>	173
6.33	<a href="#">ksrc/nucleus/sched-rt.c File Reference</a>	174
6.33.1	<a href="#">Detailed Description</a>	174
6.34	<a href="#">ksrc/nucleus/sched-sporadic.c File Reference</a>	175
6.34.1	<a href="#">Detailed Description</a>	175
6.35	<a href="#">ksrc/nucleus/sched-tp.c File Reference</a>	175
6.35.1	<a href="#">Detailed Description</a>	176
6.36	<a href="#">ksrc/nucleus/sched.c File Reference</a>	177
6.36.1	<a href="#">Detailed Description</a>	177
6.37	<a href="#">ksrc/nucleus/select.c File Reference</a>	177
6.37.1	<a href="#">Detailed Description</a>	178

---

6.38	<a href="#">ksrc/nucleus/shadow.c File Reference</a>	179
6.38.1	<a href="#">Detailed Description</a>	179
6.39	<a href="#">ksrc/nucleus/synch.c File Reference</a>	180
6.39.1	<a href="#">Detailed Description</a>	181
6.40	<a href="#">ksrc/nucleus/timebase.c File Reference</a>	182
6.40.1	<a href="#">Detailed Description</a>	183
6.41	<a href="#">ksrc/nucleus/timer.c File Reference</a>	183
6.41.1	<a href="#">Detailed Description</a>	184
6.42	<a href="#">ksrc/nucleus/vfile.c File Reference</a>	184
6.42.1	<a href="#">Detailed Description</a>	186



# Chapter 1

## Module Index

### 1.1 Modules

Here is a list of all modules:

Xenomai nucleus. . . . .	38
Thread state flags. . . . .	7
Thread information flags. . . . .	10
Buffer descriptors. . . . .	11
Dynamic memory allocation services. . . . .	21
Interrupt management. . . . .	27
Lightweight key-to-object mapping service . . . . .	33
Real-time pod services. . . . .	39
Registry services. . . . .	60
File descriptors events multiplexing services. . . . .	67
Real-time shadow services. . . . .	70
Thread synchronization services. . . . .	74
Time base services. . . . .	83
Timer services. . . . .	91
Virtual file services . . . . .	98
HAL. . . . .	105
Sched . . . . .	113





## Chapter 2

# Data Structure Index

### 2.1 Data Structures

Here are the data structures with brief descriptions:

<a href="#">xnpod</a> (Real-time pod descriptor ) . . . . .	115
<a href="#">xnsched</a> (Scheduling information structure ) . . . . .	117
<a href="#">xnthread_info</a> (Structure containing thread information ) . . . . .	118
<a href="#">xnvfile_lock_ops</a> (Vfile locking operations ) . . . . .	120
<a href="#">xnvfile_regular_iterator</a> (Regular vfile iterator ) . . . . .	121
<a href="#">xnvfile_regular_ops</a> (Regular vfile operation descriptor ) . . . . .	122
<a href="#">xnvfile_rev_tag</a> (Snapshot revision tag ) . . . . .	126
<a href="#">xnvfile_snapshot</a> (Snapshot vfile descriptor ) . . . . .	126
<a href="#">xnvfile_snapshot_iterator</a> (Snapshot-driven vfile iterator ) . . . . .	127
<a href="#">xnvfile_snapshot_ops</a> (Snapshot vfile operation descriptor ) . . . . .	129



## Chapter 3

# File Index

### 3.1 File List

Here is a list of all documented files with brief descriptions:

include/nucleus/assert.h . . . . .	??
include/nucleus/bheap.h . . . . .	??
include/nucleus/bufd.h . . . . .	133
include/nucleus/compiler.h . . . . .	??
include/nucleus/heap.h . . . . .	??
include/nucleus/hostrt.h (Definitions for global semaphore heap shared objects ) . . . .	135
include/nucleus/intr.h . . . . .	??
include/nucleus/jhash.h . . . . .	??
include/nucleus/map.h . . . . .	136
include/nucleus/module.h . . . . .	??
include/nucleus/pipe.h . . . . .	??
include/nucleus/pod.h (Real-time pod interface header ) . . . . .	138
include/nucleus/ppd.h . . . . .	??
include/nucleus/queue.h . . . . .	??
include/nucleus/registry.h (This file is part of the Xenomai project ) . . . . .	140
include/nucleus/sched-idle.h (Definitions for the IDLE scheduling class ) . . . . .	142
include/nucleus/sched-rt.h (Definitions for the RT scheduling class ) . . . . .	143
include/nucleus/sched-sporadic.h (Definitions for the SSP scheduling class ) . . . . .	143
include/nucleus/sched-tp.h (Definitions for the TP scheduling class ) . . . . .	144
include/nucleus/sched.h (Scheduler interface header ) . . . . .	145
include/nucleus/schedparam.h . . . . .	??
include/nucleus/schedqueue.h . . . . .	??
include/nucleus/select.h (File descriptors events multiplexing header ) . . . . .	146
include/nucleus/seqlock.h . . . . .	??
include/nucleus/shadow.h . . . . .	??
include/nucleus/stat.h . . . . .	??
include/nucleus/synch.h . . . . .	??
include/nucleus/sys_ppd.h . . . . .	??
include/nucleus/system.h . . . . .	??
include/nucleus/thread.h . . . . .	??
include/nucleus/timebase.h . . . . .	148
include/nucleus/timer.h . . . . .	150
include/nucleus/trace.h . . . . .	??

include/nucleus/ <b>types.h</b> . . . . .	??
include/nucleus/ <b>vdso.h</b> (Definitions for global semaphore heap shared objects) . . . . .	151
include/nucleus/ <b>version.h</b> . . . . .	??
include/nucleus/ <b>vfile.h</b> (This file is part of the Xenomai project) . . . . .	152
include/nucleus/ <b>xenomai.h</b> . . . . .	??
ksrc/arch/arm/ <b>hal.c</b> (Adeos-based Real-Time Abstraction Layer for ARM) . . . . .	155
ksrc/arch/blackfin/ <b>hal.c</b> (Adeos-based Real-Time Abstraction Layer for the Blackfin architecture) . . . . .	156
ksrc/arch/generic/ <b>hal.c</b> (Generic Real-Time HAL) . . . . .	157
ksrc/arch/nios2/ <b>hal.c</b> (Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture) . . . . .	157
ksrc/arch/powerpc/ <b>hal.c</b> (Adeos-based Real-Time Abstraction Layer for PowerPC) . . . . .	158
ksrc/arch/sh/ <b>hal.c</b> (Adeos-based Real-Time Abstraction Layer for the SuperH architecture) . . . . .	160
ksrc/arch/x86/ <b>hal-common.c</b> (Adeos-based Real-Time Abstraction Layer for x86) . . . . .	161
ksrc/arch/x86/ <b>hal_32.c</b> (Adeos-based Real-Time Abstraction Layer for x86) . . . . .	162
ksrc/arch/x86/ <b>hal_64.c</b> (Adeos-based Real-Time Abstraction Layer for x86_64) . . . . .	163
ksrc/arch/x86/ <b>smi.c</b> (SMI workaround for x86) . . . . .	163
ksrc/nucleus/ <b>bufd.c</b> . . . . .	164
ksrc/nucleus/ <b>heap.c</b> (Dynamic memory allocation services) . . . . .	165
ksrc/nucleus/ <b>intr.c</b> (Interrupt management) . . . . .	166
ksrc/nucleus/ <b>map.c</b> . . . . .	168
ksrc/nucleus/ <b>pod.c</b> (Real-time pod services) . . . . .	169
ksrc/nucleus/ <b>registry.c</b> (This file is part of the Xenomai project) . . . . .	171
ksrc/nucleus/ <b>sched-idle.c</b> (Idle scheduling class implementation (i.e. Linux placeholder)) . . . . .	173
ksrc/nucleus/ <b>sched-rt.c</b> (Common real-time scheduling class implementation (FIFO + RR)) . . . . .	174
ksrc/nucleus/ <b>sched-sporadic.c</b> (POSIX SCHED_SPORADIC scheduling class) . . . . .	175
ksrc/nucleus/ <b>sched-tp.c</b> (Temporal partitioning (typical of IMA systems)) . . . . .	175
ksrc/nucleus/ <b>sched.c</b> . . . . .	177
ksrc/nucleus/ <b>select.c</b> (File descriptors events multiplexing) . . . . .	177
ksrc/nucleus/ <b>shadow.c</b> (Real-time shadow services) . . . . .	179
ksrc/nucleus/ <b>synch.c</b> (Thread synchronization services) . . . . .	180
ksrc/nucleus/ <b>timebase.c</b> . . . . .	182
ksrc/nucleus/ <b>timer.c</b> . . . . .	183
ksrc/nucleus/ <b>vfile.c</b> (This file is part of the Xenomai project) . . . . .	184

# Chapter 4

## Module Documentation

### 4.1 Thread state flags.

Bits reporting permanent or transient states of thread.

Collaboration diagram for Thread state flags.:



#### Defines

- #define `XNSUSP` 0x00000001  
*Suspended.*
- #define `XNPEND` 0x00000002  
*Sleep-wait for a resource.*
- #define `XNDELAY` 0x00000004  
*Delayed.*
- #define `XNREADY` 0x00000008  
*Linked to the ready queue.*
- #define `XNDORMANT` 0x00000010  
*Not started yet or killed.*
- #define `XNZOMBIE` 0x00000020

*Zombie thread in deletion process.*

- #define `XNRESTART` 0x00000040  
*Restarting thread.*
- #define `XNSTARTED` 0x00000080  
*Thread has been started.*
- #define `XNMAPPED` 0x00000100  
*Mapped to a regular Linux task (shadow only).*
- #define `XNRELAX` 0x00000200  
*Relaxed shadow thread (blocking bit).*
- #define `XNMIGRATE` 0x00000400  
*Thread is currently migrating to another CPU.*
- #define `XNHELD` 0x00000800  
*Thread is held to process emergency.*
- #define `XNBOOST` 0x00001000  
*Undergoes a PIP boost.*
- #define `XNDEBUG` 0x00002000  
*Hit a debugger breakpoint (shadow only).*
- #define `XNLOCK` 0x00004000  
*Holds the scheduler lock (i.e.*
- #define `XNRRB` 0x00008000  
*Undergoes a round-robin scheduling.*
- #define `XNASDI` 0x00010000  
*ASR are disabled.*
- #define `XNDEFCAN` 0x00020000  
*Deferred cancelability mode (self-set only).*
- #define `XNTRAPSW` 0x00040000  
*Trap execution mode switches.*
- #define `XNRPIOFF` 0x00080000  
*Stop priority coupling (shadow only).*
- #define `XNFPU` 0x00100000  
*Thread uses FPU.*
- #define `XNSHADOW` 0x00200000  
*Shadow thread.*

- `#define XNROOT 0x00400000`  
*Root thread (that is, Linux/IDLE).*
- `#define XNOTHER 0x00800000`  
*Non real-time shadow (prio=0).*

### 4.1.1 Detailed Description

Bits reporting permanent or transient states of thread.

### 4.1.2 Define Documentation

#### 4.1.2.1 `#define XNHELD 0x00000800`

Thread is held to process emergency.

Referenced by `xnpod_resume_thread()`, and `xnpod_suspend_thread()`.

#### 4.1.2.2 `#define XNLOCK 0x00004000`

Holds the scheduler lock (i.e.  
not preemptible)

Referenced by `xnpod_set_thread_mode()`, and `xnpod_welcome_thread()`.

#### 4.1.2.3 `#define XNMIGRATE 0x00000400`

Thread is currently migrating to another CPU.

Referenced by `xnpod_delete_thread()`.

#### 4.1.2.4 `#define XNPEND 0x00000002`

Sleep-wait for a resource.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, `xnpod_unblock_thread()`, `xnsynch_acquire()`, `xnsynch_flush()`, `xnsynch_forget_sleeper()`, `xnsynch_sleep_on()`, `xnsynch_wakeup_one_sleeper()`, and `xnsynch_wakeup_this_sleeper()`.

#### 4.1.2.5 `#define XNREADY 0x00000008`

Linked to the ready queue.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, `xnpod_start_thread()`, and `xnpod_suspend_thread()`.

#### 4.1.2.6 #define XNSUSP 0x00000001


Suspended.

Referenced by `xnpod_init_thread()`, `xnpod_start_thread()`, `xnpod_suspend_thread()`, and `xnpod_trap_fault()`.

## 4.2 Thread information flags.

Bits reporting events notified to the thread.

Collaboration diagram for Thread information flags.:



## Defines

- #define [XNTIMEO](#) 0x00000001  
*Woken up due to a timeout condition.*
- #define [XNRMID](#) 0x00000002  
*Pending on a removed resource.*
- #define [XNBREAK](#) 0x00000004  
*Forcibly awoken from a wait state.*
- #define [XNKICKED](#) 0x00000008  
*Kicked upon Linux signal (shadow only).*
- #define [XNWAKEN](#) 0x00000010  
*Thread waken up upon resource availability.*
- #define [XNROBBED](#) 0x00000020  
*Robbed from resource ownership.*
- #define [XNATOMIC](#) 0x00000040  
*In atomic switch from secondary to primary mode.*
- #define [XNAFFSET](#) 0x00000080  
*CPU affinity changed from primary mode.*
- #define [XNPRIOSET](#) 0x00000100



*Priority changed from primary mode.*

- #define [XNABORT](#) 0x00000200  
*Thread is being aborted.*
- #define [XNCANPND](#) 0x00000400  
*Cancellation request is pending.*
- #define [XNAMOK](#) 0x00000800  
*Runaway, watchdog signal pending (shadow only).*
- #define [XNSWREP](#) 0x00001000  
*Mode switch already reported.*

### 4.2.1 Detailed Description

Bits reporting events notified to the thread.

## 4.3 Buffer descriptors.

Collaboration diagram for Buffer descriptors.:



### Files

- file [bufd.h](#)
- file [bufd.c](#)

### Functions

- static void [xnbufd\\_map\\_uread](#) (struct xnbufd \*bufd, const void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from user memory.*
- static void [xnbufd\\_map\\_uwrite](#) (struct xnbufd \*bufd, void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to user memory.*
- ssize\_t [xnbufd\\_unmap\\_uread](#) (struct xnbufd \*bufd)

Finalize a buffer descriptor obtained from `xnbufd_map_uread()`.

- `ssize_t xnbufd_unmap_uwrite` (struct xnbufd \*bufd)  
Finalize a buffer descriptor obtained from `xnbufd_map_uwrite()`.
- static void `xnbufd_map_kread` (struct xnbufd \*bufd, const void \*ptr, size\_t len)  
Initialize a buffer descriptor for reading from kernel memory.
- static void `xnbufd_map_kwrite` (struct xnbufd \*bufd, void \*ptr, size\_t len)  
Initialize a buffer descriptor for writing to kernel memory.
- `ssize_t xnbufd_unmap_kread` (struct xnbufd \*bufd)  
Finalize a buffer descriptor obtained from `xnbufd_map_kread()`.
- `ssize_t xnbufd_unmap_kwrite` (struct xnbufd \*bufd)  
Finalize a buffer descriptor obtained from `xnbufd_map_kwrite()`.
- `ssize_t xnbufd_copy_to_kmem` (void \*ptr, struct xnbufd \*bufd, size\_t len)  
Copy memory covered by a buffer descriptor to kernel memory.
- `ssize_t xnbufd_copy_from_kmem` (struct xnbufd \*bufd, void \*from, size\_t len)  
Copy kernel memory to the area covered by a buffer descriptor.
- void `xnbufd_invalidate` (struct xnbufd \*bufd)  
Invalidate a buffer descriptor.
- static void `xnbufd_reset` (struct xnbufd \*bufd)  
Reset a buffer descriptor.

### 4.3.1 Detailed Description

A buffer descriptor is a simple abstraction dealing with copy operations to/from memory buffers which may belong to different address spaces.

To this end, the buffer descriptor library provides a small set of copy routines which are aware of address space restrictions when moving data, and a generic container type which can hold a reference to - or cover - a particular memory area, either present in kernel space, or in any of the existing user memory contexts.

The goal of the buffer descriptor abstraction is to hide address space specifics from Xenomai services dealing with memory areas, allowing them to operate on multiple address spaces seamlessly.

The common usage patterns are as follows:

- Implementing a Xenomai syscall returning a bulk of data to the caller, which may have to be copied back to either kernel or user space:

```
[Syscall implementation]
ssize_t rt_bulk_read_inner(struct xnbufd *bufd)
{
    ssize_t ret;
```

```

    size_t len;
    void *bulk;

    bulk = get_next_readable_bulk(&len);
    ret = xnbufd_copy_from_kmem(bufd, bulk, min(bufd->b_len, len));
    free_bulk(bulk);

    ret = this_may_fail();
    if (ret)
        xnbufd_invalidate(bufd);

    return ret;
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_read(void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_kwrite(&bufd);

    return ret;
}

[Userland trampoline for user syscalls]
int __rt_bulk_read(struct pt_regs *regs)
{
    struct xnbufd bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbufd_map_uwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_uwrite(&bufd);

    return ret;
}

```

- Implementing a Xenomai syscall receiving a bulk of data from the caller, which may have to be read from either kernel or user space:

```

[Syscall implementation]
ssize_t rt_bulk_write_inner(struct xnbufd *bufd)
{
    void *bulk = get_free_bulk(bufd->b_len);
    return xnbufd_copy_to_kmem(bulk, bufd, bufd->b_len);
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_write(const void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbufd_unmap_kread(&bufd);

    return ret;
}

```

```

}

[Userland trampoline for user syscalls]
int __rt_bulk_write(struct pt_regs *regs)
{
    struct xnbuid bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbuid_map_uread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbuid_unmap_uread(&bufd);

    return ret;
}

```

### 4.3.2 Function Documentation

#### 4.3.2.1 `ssize_t xnbuid_copy_from_kmem ( struct xnbuid * bufd, void * from, size_t len )`

Copy kernel memory to the area covered by a buffer descriptor.

This routine copies *len* bytes from the kernel memory starting at *from* to the area referred to by the buffer descriptor *bufd*. `xnbuid_copy_from_kmem()` tracks the write offset within the destination memory internally, so that it may be called several times in a loop, until the entire memory area is stored.

The destination address space is dealt with, according to the following rules:

- if *bufd* refers to a writable kernel area (i.e. see `xnbuid_map_kwrite()`), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a writable user area (i.e. see `xnbuid_map_uwrite()`), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while writing to that memory.
- if *bufd* refers to a user area which may not be immediately written to from the current context, the copy is postponed until `xnbuid_unmap_uwrite()` is invoked for *ubufd*, at which point the copy will take place. In such a case, the source memory is transferred to a carry over buffer allocated internally; this operation may lead to request dynamic memory from the nucleus heap if *len* is greater than 64 bytes.

#### Parameters

*bufd* The address of the buffer descriptor covering the user memory to copy data to.

*from* The start address of the kernel memory to copy from.

*len* The length of the kernel memory to copy to *bufd*.

#### Returns

The number of bytes written so far to the memory area covered by *ubufd*. Otherwise,

- -ENOMEM is returned when no memory is available from the nucleus heap to allocate the carry over buffer.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while writing to the user area. For that reason, [xnbufd\\_copy\\_from\\_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

#### Note

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

#### 4.3.2.2 `ssize_t xnbufd_copy_to_kmem ( void * to, struct xnbufd * bufd, size_t len )`

Copy memory covered by a buffer descriptor to kernel memory.

This routine copies *len* bytes from the area referred to by the buffer descriptor *bufd* to the kernel memory area *to*. [xnbufd\\_copy\\_to\\_kmem\(\)](#) tracks the read offset within the source memory internally, so that it may be called several times in a loop, until the entire memory area is loaded.

The source address space is dealt with, according to the following rules:

- if *bufd* refers to readable kernel area (i.e. see [xnbufd\\_map\\_kread\(\)](#)), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a readable user area (i.e. see [xnbufd\\_map\\_uread\(\)](#)), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while reading from that memory.
- any attempt to read from *bufd* from a non-suitable context is considered as a bug, and will raise a panic assertion when the nucleus is compiled in debug mode.

#### Parameters

*to* The start address of the kernel memory to copy to.

*bufd* The address of the buffer descriptor covering the user memory to copy data from.

*len* The length of the user memory to copy from *bufd*.

#### Returns

The number of bytes read so far from the memory area covered by *ubufd*. Otherwise:

- -EINVAL is returned upon attempt to read from the user area from an invalid context. This error is only returned when the debug mode is disabled; otherwise a panic assertion is raised.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while reading from the user area. For that reason, [xnbufd\\_copy\\_to\\_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

#### Note

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

#### 4.3.2.3 void xnbufd\_invalidate ( struct xnbufd \* *bufd* )

Invalidate a buffer descriptor.

The buffer descriptor is invalidated, making it unusable for further copy operations. If an outstanding carry over buffer was allocated by a previous call to [xnbufd\\_copy\\_from\\_kmem\(\)](#), it is immediately freed so that no data transfer will happen when the descriptor is finalized.

The only action that may subsequently be performed on an invalidated descriptor is calling the relevant unmapping routine for it. For that reason, [xnbufd\\_invalidate\(\)](#) should be invoked on the error path when data may have been transferred to the carry over buffer.

#### Parameters

*bufd* The address of the buffer descriptor to invalidate.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 4.3.2.4 void xnbufd\_map\_kread ( struct xnbufd \* *bufd*, const void \* *ptr*, size\_t *len* ) [inline, static]

Initialize a buffer descriptor for reading from kernel memory.

The new buffer descriptor may be used to copy data from kernel memory. This routine should be used in pair with [xnbufd\\_unmap\\_kread\(\)](#).

#### Parameters

*bufd* The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

*ptr* The start of the kernel buffer to map.

*len* The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

**4.3.2.5 void xnbuid\_map\_kwrite ( struct xnbuid \* *bufd*, void \* *ptr*, size\_t *len* ) [inline, static]**

Initialize a buffer descriptor for writing to kernel memory.

The new buffer descriptor may be used to copy data to kernel memory. This routine should be used in pair with [xnbuid\\_unmap\\_kwrite\(\)](#).

#### Parameters

*bufd* The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

*ptr* The start of the kernel buffer to map.

*len* The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

**4.3.2.6 void xnbuid\_map\_uread ( struct xnbuid \* *bufd*, const void \_\_user \* *ptr*, size\_t *len* ) [inline, static]**

Initialize a buffer descriptor for reading from user memory.

The new buffer descriptor may be used to copy data from user memory. This routine should be used in pair with [xnbuid\\_unmap\\_uread\(\)](#).

#### Parameters

*bufd* The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

*ptr* The start of the user buffer to map.

*len* The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

**4.3.2.7 void xnbuid\_map\_uwrite ( struct xnbuid \* bufd, void \_\_user \* ptr, size\_t len )  
[inline, static]**

Initialize a buffer descriptor for writing to user memory.

The new buffer descriptor may be used to copy data to user memory. This routine should be used in pair with [xnbuid\\_unmap\\_uwrite\(\)](#).

#### Parameters

*bufd* The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

*ptr* The start of the user buffer to map.

*len* The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

**4.3.2.8 void xnbuid\_reset ( struct xnbuid \* bufd ) [inline, static]**

Reset a buffer descriptor.

The buffer descriptor is reset, so that all data already copied is forgotten. Any carry over buffer allocated is kept, though.

#### Parameters

*bufd* The address of the buffer descriptor to reset.

Environments:

This service can be called from:

- Kernel code (including from primary mode)



- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 4.3.2.9 `ssize_t xnbuid_unmap_kread ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_kread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_kread\(\)](#), to read data from a kernel area.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes read so far from the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 4.3.2.10 `ssize_t xnbuid_unmap_kwrite ( struct xnbuid * bufd )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_kwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_kwrite\(\)](#), to write data to a kernel area.

##### Parameters

*bufd* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes written so far to the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

#### 4.3.2.11 `ssize_t xnbuid_unmap_uread ( struct xnbuid * buid )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_uread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_uread\(\)](#), to read data from a user area.

##### Parameters

*buid* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes read so far from the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

##### Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

#### 4.3.2.12 `ssize_t xnbuid_unmap_uwrite ( struct xnbuid * buid )`

Finalize a buffer descriptor obtained from [xnbuid\\_map\\_uwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid\\_map\\_uwrite\(\)](#), to write data to a user area.

The main action taken is to write the contents of the kernel memory area passed to [xnbuid\\_copy\\_from\\_kmem\(\)](#) whenever the copy operation was postponed at that time; the carry over buffer is eventually released as needed. If [xnbuid\\_copy\\_from\\_kmem\(\)](#) was allowed to copy to the destination user memory at once, then [xnbuid\\_unmap\\_uwrite\(\)](#) leads to a no-op.

##### Parameters

*buid* The address of the buffer descriptor to finalize.

##### Returns

The number of bytes written so far to the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.


Rescheduling: never.

#### Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

## 4.4 Dynamic memory allocation services.

Collaboration diagram for Dynamic memory allocation services.:



### Files

- file [heap.c](#)  
*Dynamic memory allocation services.*

### Functions

- `int xnheap\_init (xnheap_t *heap, void *heapaddr, u_long heapsize, u_long pagesize)`  
*Initialize a memory heap.*
- `void xnheap\_set\_label (xnheap_t *heap, const char *label,...)`  
*Set the heap's label string.*
- `void xnheap\_destroy (xnheap_t *heap, void(*flushfn)(xnheap_t *heap, void *extaddr, u_long extsize, void *cookie), void *cookie)`  
*Destroys a memory heap.*
- `void * xnheap\_alloc (xnheap_t *heap, u_long size)`  
*Allocate a memory block from a memory heap.*
- `int xnheap\_test\_and\_free (xnheap_t *heap, void *block, int(*ckfn)(void *block))`  
*Test and release a memory block to a memory heap.*
- `int xnheap\_free (xnheap_t *heap, void *block)`  
*Release a memory block to a memory heap.*
- `int xnheap\_extend (xnheap_t *heap, void *extaddr, u_long extsize)`

*Extend a memory heap.*

- void [xnheap\\_schedule\\_free](#) (xnheap\_t \*heap, void \*block, xnholder\_t \*link)

*Schedule a memory block for release.*

### 4.4.1 Detailed Description

Dynamic memory allocation services.

The implementation of the memory allocator follows the algorithm described in a USENIX 1988 paper called "Design of a General Purpose Memory Allocator for the 4.3BSD Unix Kernel" by Marshall K. McKusick and Michael J. Karels. You can find it at various locations on the net, including <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf>. A minor variation allows this implementation to have 'extendable' heaps when needed, with multiple memory extents providing autonomous page address spaces.

The data structures hierarchy is as follows:

```

HEAP {
    block_buckets[]
    extent_queue -----+
}
|
V
    EXTENT #1 {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    } --+
|
|
V
    EXTENT #n {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    }

```

### 4.4.2 Function Documentation

#### 4.4.2.1 void\* xnheap\_alloc ( xnheap\_t \* heap, u\_long size )

Allocate a memory block from a memory heap.

Allocates a contiguous region of memory from an active memory heap. Such allocation is guaranteed to be time-bounded.

#### Parameters

*heap* The descriptor address of the heap to get memory from.

*size* The size in bytes of the requested block. Sizes lower or equal to the page size are rounded either to the minimum allocation size if lower than this value, or to the minimum alignment size if greater or equal to this value. In the current implementation, with MINALLOC = 8 and MINALIGN = 16, a 7 bytes request will be rounded to 8 bytes, and a 17 bytes request will be rounded to 32.

**Returns**

The address of the allocated region upon success, or NULL if no memory is available from the specified heap.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnshadow_map()`.

**4.4.2.2** `void xnheap_destroy ( xnheap_t * heap, void(*)(xnheap_t *heap, void *extaddr, u_long extsize, void *cookie) flushfn, void * cookie )`

Destroys a memory heap.

Destroys a memory heap.

**Parameters**

*heap* The descriptor address of the destroyed heap.

*flushfn* If non-NULL, the address of a flush routine which will be called for each extent attached to the heap. This routine can be used by the calling code to further release the heap memory.

*cookie* If *flushfn* is non-NULL, *cookie* is an opaque pointer which will be passed unmodified to *flushfn*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`, and `xnpod_shutdown()`.

**4.4.2.3** `int xnheap_extend ( xnheap_t * heap, void * extaddr, u_long extsize )`

Extend a memory heap.

Add a new extent to an existing memory heap.

### Parameters

*heap* The descriptor address of the heap to add an extent to.

*extaddr* The address of the extent memory.

*extsize* The size of the extent memory (in bytes). In the current implementation, this size must match the one of the initial extent passed to [xnheap\\_init\(\)](#).

### Returns

0 is returned upon success, or -EINVAL is returned if *extsize* differs from the initial extent's size.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.4.2.4 `int xnheap_free ( xnheap_t* heap, void* block )`

Release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from.

### Parameters

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

### Returns

0 is returned upon success, or one of the following error codes:

- -EFAULT is returned whenever the memory address is outside the heap address space.
- -EINVAL is returned whenever the memory address does not represent a valid block.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References [xnheap\\_test\\_and\\_free\(\)](#).

#### 4.4.2.5 `int xnheap_init ( xnheap_t * heap, void * heapaddr, u_long heapsize, u_long pagesize )`

Initialize a memory heap.

Initializes a memory heap suitable for time-bounded allocation requests of dynamic memory.

##### Parameters

*heap* The address of a heap descriptor which will be used to store the allocation data. This descriptor must always be valid while the heap is active therefore it must be allocated in permanent memory.

*heapaddr* The address of the heap storage area. All allocations will be made from the given area in time-bounded mode. Since additional extents can be added to a heap, this parameter is also known as the "initial extent".

*heapsize* The size in bytes of the initial extent pointed at by *heapaddr*. *heapsize* must be a multiple of *pagesize* and lower than 16 Mbytes. *heapsize* must be large enough to contain a dynamically-sized internal header. The following formula gives the size of this header:

$H = \text{heapsize}, P = \text{pagesize}, M = \text{sizeof}(\text{struct pagemap}), E = \text{sizeof}(\text{xnextent\_t})$

$\text{hdrsize} = ((H - E) * M) / (M + 1)$

This value is then aligned on the next 16-byte boundary. The routine `xnheap_overhead()` computes the corrected heap size according to the previous formula.

##### Parameters

*pagesize* The size in bytes of the fundamental memory page which will be used to subdivide the heap internally. Choosing the right page size is important regarding performance and memory fragmentation issues, so it might be a good idea to take a look at <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf> to pick the best one for your needs. In the current implementation, *pagesize* must be a power of two in the range [ 8 .. 32768 ] inclusive.

##### Returns

0 is returned upon success, or one of the following error codes:

- -EINVAL is returned whenever a parameter is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

#### 4.4.2.6 void xnheap\_schedule\_free ( xnheap\_t \* heap, void \* block, xnholder\_t \* link )

Schedule a memory block for release.

This routine records a block for later release by `xnheap_finalize_free()`. This service is useful to lazily free blocks of heap memory when immediate release is not an option, e.g. when active references are still pending on the object for a short time after the call. `xnheap_finalize_free()` is expected to be eventually called by the client code at some point in the future when actually freeing the idle objects is deemed safe.

##### Parameters

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

*link* The address of a link member, likely but not necessarily within the released object, which will be used by the heap manager to hold the block in the queue of idle objects.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.4.2.7 void xnheap\_set\_label ( xnheap\_t \* heap, const char \* label, ... )

Set the heap's label string.

Set the heap label that will be used in statistic outputs.

##### Parameters

*heap* The address of a heap descriptor.

*label* Label string displayed in statistic outputs. This parameter can be a format string, in which case succeeding parameters will be used to resolve the final label.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.



#### 4.4.2.8 int xnheap\_test\_and\_free ( xnheap\_t \* heap, void \* block, int(\*)(void \*block) ckfn )

Test and release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from. Before the actual release is performed, an optional user-defined can be invoked to check for additional criteria with respect to the request consistency.

##### Parameters

*heap* The descriptor address of the heap to release memory to.

*block* The address of the region to be returned to the heap.

*ckfn* The address of a user-supplied verification routine which is to be called after the memory address specified by *block* has been checked for validity. The routine is expected to proceed to further consistency checks, and either return zero upon success, or non-zero upon error. In the latter case, the release process is aborted, and *ckfn's* return value is passed back to the caller of this service as its error return code. *ckfn* must not trigger the rescheduling procedure either directly or indirectly.

##### Returns

0 is returned upon success, or -EINVAL is returned whenever the block is not a valid region of the specified heap. Additional return codes can also be defined locally by the *ckfn* routine.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by xnheap\_free().

## 4.5 Interrupt management.

Collaboration diagram for Interrupt management.:



## Files

- file [intr.c](#)  
*Interrupt management.*

## Functions

- int [xnintr\\_init](#) (xnintr\_t \*intr, const char \*name, unsigned irq, xnintr\_t isr, xniack\_t iack, xnflags\_t flags)  
*Initialize an interrupt object.*
- int [xnintr\\_destroy](#) (xnintr\_t \*intr)  
*Destroy an interrupt object.*
- int [xnintr\\_attach](#) (xnintr\_t \*intr, void \*cookie)  
*Attach an interrupt object.*
- int [xnintr\\_detach](#) (xnintr\_t \*intr)  
*Detach an interrupt object.*
- int [xnintr\\_enable](#) (xnintr\_t \*intr)  
*Enable an interrupt object.*
- int [xnintr\\_disable](#) (xnintr\_t \*intr)  
*Disable an interrupt object.*
- void [xnintr\\_affinity](#) (xnintr\_t \*intr, xnarch\_cpumask\_t cpumask)  
*Set interrupt's processor affinity.*

### 4.5.1 Detailed Description

Interrupt management.

### 4.5.2 Function Documentation

#### 4.5.2.1 void xnintr\_affinity ( xnintr\_t\* intr, xnarch\_cpumask\_t cpumask )

Set interrupt's processor affinity.

Causes the IRQ associated with the interrupt object *intr* to be received only on processors which bits are set in *cpumask*.

#### Parameters

- intr* The descriptor address of the interrupt object which affinity is to be changed.
- cpumask* The new processor affinity of the interrupt object.

**Returns**

the previous cpumask on success, or an empty mask on failure.

**Note**

Depending on architectures, setting more than one bit in *cpumask* could be meaningless.

**4.5.2.2 int xnintr\_attach ( xnintr\_t \* intr, void \* cookie )**

Attach an interrupt object.

Attach an interrupt object previously initialized by [xnintr\\_init\(\)](#). After this operation is completed, all IRQs received from the corresponding interrupt channel are directed to the object's ISR.

**Parameters**

*intr* The descriptor address of the interrupt object to attach.

*cookie* A user-defined opaque value which is stored into the interrupt object descriptor for further retrieval by the ISR/ISR handlers.

**Returns**

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while attaching the interrupt.
- -EBUSY is returned if the interrupt object was already attached.

**Note**

The caller **must not** hold nklock when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

**Note**

Attaching an interrupt resets the tracked number of receipts to zero.

**4.5.2.3 int xnintr\_destroy ( xnintr\_t \* intr )**

Destroy an interrupt object.

Destroys an interrupt object previously initialized by [xnintr\\_init\(\)](#). The interrupt object is automatically detached by a call to [xnintr\\_detach\(\)](#). No more IRQs will be dispatched by this object after this service has returned.

### Parameters

*intr* The descriptor address of the interrupt object to destroy.

### Returns

0 is returned on success. Otherwise, -EINVAL is returned if an error occurred while detaching the interrupt (see [xnintr\\_detach\(\)](#)).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

References [xnintr\\_detach\(\)](#).

#### 4.5.2.4 int xnintr\_detach ( xnintr\_t \* intr )

Detach an interrupt object.

Detach an interrupt object previously attached by [xnintr\\_attach\(\)](#). After this operation is completed, no more IRQs are directed to the object's ISR, but the interrupt object itself remains valid. A detached interrupt object can be attached again by a subsequent call to [xnintr\\_attach\(\)](#).

### Parameters

*intr* The descriptor address of the interrupt object to detach.

### Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while detaching the interrupt, or if the interrupt object was not attached. In both cases, no action is performed.

### Note

The caller **must not** hold nklock when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by [xnintr\\_destroy\(\)](#).

#### 4.5.2.5 `int xnintr_disable ( xnintr_t * intr )`

Disable an interrupt object.

Disables the hardware interrupt line associated with an interrupt object. This operation invalidates further interrupt requests from the given source until the IRQ line is re-enabled anew.

##### Parameters

*intr* The descriptor address of the interrupt object to disable.

##### Returns

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while disabling the interrupt.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

#### 4.5.2.6 `int xnintr_enable ( xnintr_t * intr )`

Enable an interrupt object.

Enables the hardware interrupt line associated with an interrupt object. Over real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

##### Parameters

*intr* The descriptor address of the interrupt object to enable.

##### Returns

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while enabling the interrupt.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

**4.5.2.7** `int xnintr_init ( xnintr_t * intr, const char * name, unsigned irq, xnintr_t isr, xniack_t iack, xnflags_t flags )`

Initialize an interrupt object.

Associates an interrupt object with an IRQ line.

When an interrupt occurs on the given *irq* line, the ISR is fired in order to deal with the hardware event. The interrupt service code may call any non-suspensive service from the nucleus.

Upon receipt of an IRQ, the ISR is immediately called on behalf of the interrupted stack context, the rescheduling procedure is locked, and the interrupt source is masked at hardware level. The status value returned by the ISR is then checked for the following values:

- `XN_ISR_HANDLED` indicates that the interrupt request has been fulfilled by the ISR.
- `XN_ISR_NONE` indicates the opposite to `XN_ISR_HANDLED`. The ISR must always return this value when it determines that the interrupt request has not been issued by the dedicated hardware device.

In addition, one of the following bits may be set by the ISR :

NOTE: use these bits with care and only when you do understand their effect on the system. The ISR is not encouraged to use these bits in case it shares the IRQ line with other ISRs in the real-time domain.

- `XN_ISR_NOENABLE` causes the nucleus to ask the real-time control layer `_not_` to re-enable the IRQ line (read the following section). `xnarch_end_irq()` must be called to re-enable the IRQ line later.
- `XN_ISR_PROPAGATE` tells the nucleus to require the real-time control layer to forward the IRQ. For instance, this would cause the Adeos control layer to propagate the interrupt down the interrupt pipeline to other Adeos domains, such as Linux. This is the regular way to share interrupts between the nucleus and the host system. In effect, `XN_ISR_PROPAGATE` implies `XN_ISR_NOENABLE` since it would make no sense to re-enable the interrupt channel before the next domain down the pipeline has had a chance to process the propagated interrupt.

The nucleus re-enables the IRQ line by default. Over some real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

A count of interrupt receipts is tracked into the interrupt descriptor, and reset to zero each time the interrupt object is attached. Since this count could wrap around, it should be used as an indication of interrupt activity only.

#### Parameters

- intr* The address of a interrupt object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.
- name* An ASCII string standing for the symbolic name of the interrupt object or NULL ("`<unknown>`" will be applied then).
- irq* The hardware interrupt channel associated with the interrupt object. This value is architecture-dependent. An interrupt object must then be attached to the hardware

interrupt vector using the [xnintr\\_attach\(\)](#) service for the associated IRQs to be directed to this object.

*isr* The address of a valid low-level interrupt service routine if this parameter is non-zero. This handler will be called each time the corresponding IRQ is delivered on behalf of an interrupt context. When called, the ISR is passed the descriptor address of the interrupt object.

*iack* The address of an optional interrupt acknowledge routine, aimed at replacing the default one. Only very specific situations actually require to override the default setting for this parameter, like having to acknowledge non-standard PIC hardware. *iack* should return a non-zero value to indicate that the interrupt has been properly acknowledged. If *iack* is NULL, the default routine will be used instead.

*flags* A set of creation flags affecting the operation. The valid flags are:

- XN\_ISR\_SHARED enables IRQ-sharing with other interrupt objects.
- XN\_ISR\_EDGE is an additional flag need to be set together with XN\_ISR\_SHARED to enable IRQ-sharing of edge-triggered interrupts.

### Returns

0 is returned on success. Otherwise, -EINVAL is returned if *irq* is not a valid interrupt number.

Environments:

This service can be called from:


- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by `xnpod_enable_timesource()`.

## 4.6 Lightweight key-to-object mapping service

Collaboration diagram for Lightweight key-to-object mapping service:



### Files

- file [map.h](#)
- file [map.c](#)

## Functions

- `xnmap_t * xnmap_create (int nkeys, int reserve, int offset)`  
*Create a map.*
- `void xnmap_delete (xnmap_t *map)`  
*Delete a map.*
- `int xnmap_enter (xnmap_t *map, int key, void *objaddr)`  
*Index an object into a map.*
- `int xnmap_remove (xnmap_t *map, int key)`  
*Remove an object reference from a map.*
- `static void * xnmap_fetch_nocheck (xnmap_t *map, int key)`  
*Search an object into a map - unchecked form.*
- `static void * xnmap_fetch (xnmap_t *map, int key)`  
*Search an object into a map.*

### 4.6.1 Detailed Description

A map is a simple indexing structure which associates unique integer keys with pointers to objects. The current implementation supports reservation, for naming/indexing the real-time objects skins create, either on a fixed, user-provided integer (i.e. a reserved key value), or by drawing the next available key internally if the caller did not specify any fixed key. For instance, in some given map, the key space ranging from 0 to 255 could be reserved for fixed keys, whilst the range from 256 to 511 could be available for drawing free keys dynamically.

A maximum of 1024 unique keys per map is supported on 32bit machines.

(This implementation should not be confused with C++ STL maps, which are dynamically expandable and allow arbitrary key types; Xenomai maps don't).

### 4.6.2 Function Documentation

#### 4.6.2.1 `xnmap_t * xnmap_create ( int nkeys, int reserve, int offset )`

Create a map.

Allocates a new map with the specified addressing capabilities. The memory is obtained from the Xenomai system heap.

#### Parameters

***nkeys*** The maximum number of unique keys the map will be able to hold. This value cannot exceed the static limit represented by `XNMAP_MAX_KEYS`, and must be a power of two.

***reserve*** The number of keys which should be kept for reservation within the index space. Reserving a key means to specify a valid key to the `xnmap_enter()` service, which will then attempt to register this exact key, instead of drawing the next available key from



the unreserved index space. When reservation is in effect, the unreserved index space will hold key values greater than *reserve*, keeping the low key values for the reserved space. For instance, passing *reserve* = 32 would cause the index range [ 0 .. 31 ] to be kept for reserved keys. When non-zero, *reserve* is rounded to the next multiple of BITS\_PER\_LONG. If *reserve* is zero no reservation will be available from the map.

*offset* The lowest key value `xnmap_enter()` will return to the caller. Key values will be in the range [ 0 + offset .. *nkeys* + offset - 1 ]. Negative offsets are valid.

### Returns

the address of the new map is returned on success; otherwise, NULL is returned if *nkeys* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.6.2.2 void xnmap\_delete ( xnmap\_t \* map )

Delete a map.

Deletes a map, freeing any associated memory back to the Xenomai system heap.

### Parameters

*map* The address of the map to delete.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.6.2.3 int xnmap\_enter ( xnmap\_t \* map, int key, void \* objaddr )

Index an object into a map.

Insert a new object into the given map.

### Parameters

*map* The address of the map to insert into.

*key* The key to index the object on. If this key is within the valid index range [ 0 - offset .. nkeys - offset - 1 ], then an attempt to reserve this exact key is made. If *key* has an out-of-range value lower or equal to 0 - offset - 1, then an attempt is made to draw a free key from the unreserved index space.

*objaddr* The address of the object to index on the key. This value will be returned by a successful call to [xnmap\\_fetch\(\)](#) with the same key.

### Returns

a valid key is returned on success, either *key* if reserved, or the next free key. Otherwise:

- -EEXIST is returned upon attempt to reserve a busy key.
- -ENOSPC when no more free key is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.6.2.4 void xnmap\_fetch ( xnmap\_t \* map, int key ) [inline, static]

Search an object into a map.

Retrieve an object reference from the given map by its index key.

### Parameters

*map* The address of the map to retrieve from.

*key* The key to be searched for in the map index.

### Returns

The indexed object address is returned on success, otherwise NULL is returned when *key* is invalid or no object is currently indexed on it.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.6.2.5 void xnmap\_fetch\_nocheck ( xnmap\_t \* map, int key ) [inline, static]

Search an object into a map - unchecked form.

Retrieve an object reference from the given map by its index key, but does not perform any sanity check on the provided key.

##### Parameters

*map* The address of the map to retrieve from.

*key* The key to be searched for in the map index.

##### Returns

The indexed object address is returned on success, otherwise NULL is returned when no object is currently indexed on *key*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.6.2.6 int xnmap\_remove ( xnmap\_t \* map, int key )

Remove an object reference from a map.

Removes an object reference from the given map, releasing the associated key.

##### Parameters

*map* The address of the map to remove from.

*key* The key the object reference to be removed is indexed on.

##### Returns

0 is returned on success. Otherwise:

- -ESRCH is returned if *key* is invalid.

Environments:

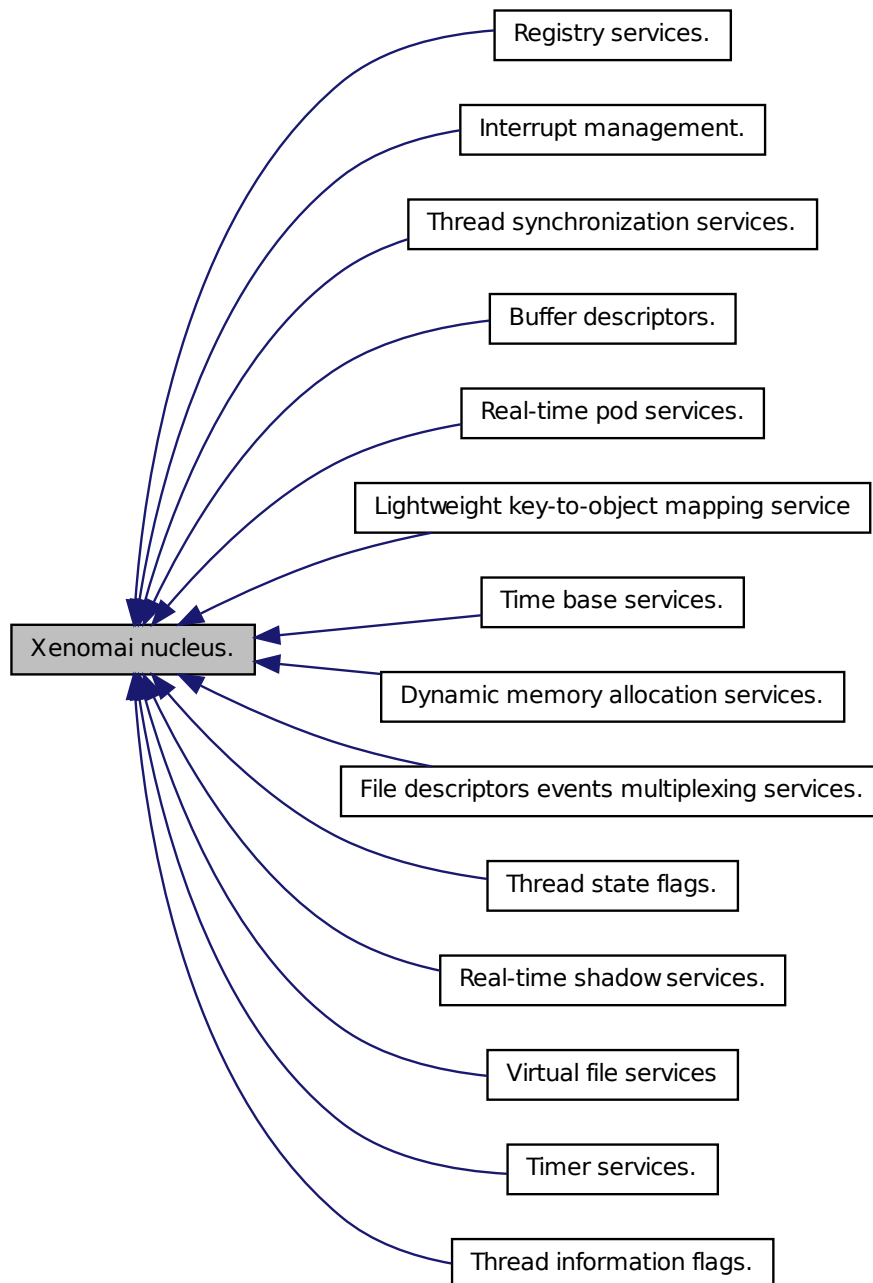
This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

## 4.7 Xenomai nucleus.

Collaboration diagram for Xenomai nucleus.:



## Modules

- [Thread state flags.](#)  
*Bits reporting permanent or transient states of thread.*
- [Thread information flags.](#)  
*Bits reporting events notified to the thread.*
- [Buffer descriptors.](#)
- [Dynamic memory allocation services.](#)
- [Interrupt management.](#)
- [Lightweight key-to-object mapping service](#)
- [Real-time pod services.](#)
- [Registry services.](#)
- [File descriptors events multiplexing services.](#)
- [Real-time shadow services.](#)
- [Thread synchronization services.](#)
- [Time base services.](#)
- [Timer services.](#)
- [Virtual file services](#)

### 4.7.1 Detailed Description

An abstract RTOS core.

## 4.8 Real-time pod services.

Collaboration diagram for Real-time pod services.:



## Data Structures

- struct [xnpod](#)  
*Real-time pod descriptor.*

## Files

- file [pod.h](#)  
*Real-time pod interface header.*
- file [pod.c](#)  
*Real-time pod services.*

## Functions

- int [xnpod\\_init](#) (void)  
*Initialize the core pod.*
- int [xnpod\\_enable\\_timesource](#) (void)  
*Activate the core time source.*
- void [xnpod\\_disable\\_timesource](#) (void)  
*Stop the core time source.*
- void [xnpod\\_shutdown](#) (int xtype)  
*Shutdown the current pod.*
- int [xnpod\\_init\\_thread](#) (struct xnthread \*thread, const struct xnthread\_init\_attr \*attr, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Initialize a new thread.*
- int [xnpod\\_start\\_thread](#) (xnthread\_t \*thread, const struct xnthread\_start\_attr \*attr)  
*Initial start of a newly created thread.*
- void [xnpod\\_stop\\_thread](#) (xnthread\_t \*thread)  
*Stop a thread.*
- void [xnpod\\_restart\\_thread](#) (xnthread\_t \*thread)  
*Restart a thread.*
- void [xnpod\\_delete\\_thread](#) (xnthread\_t \*thread)  
*Delete a thread.*
- void [xnpod\\_abort\\_thread](#) (xnthread\_t \*thread)  
*Abort a thread.*
- xnflags\_t [xnpod\\_set\\_thread\\_mode](#) (xnthread\_t \*thread, xnflags\_t clrmask, xnflags\_t set-mask)  
*Change a thread's control mode.*
- void [xnpod\\_suspend\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \*wchan)  
*Suspend a thread.*

- void [xnpod\\_resume\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask)  
*Resume a thread.*
- int [xnpod\\_unblock\\_thread](#) (xnthread\_t \*thread)  
*Unblock a thread.*
- int [xnpod\\_set\\_thread\\_schedparam](#) (struct xnthread \*thread, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*
- int [xnpod\\_migrate\\_thread](#) (int cpu)  
*Migrate the current thread.*
- void [xnpod\\_dispatch\\_signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- static void [xnpod\\_schedule](#) (void)  
*Rescheduling procedure entry point.*
- int [xnpod\\_set\\_thread\\_periodic](#) (xnthread\_t \*thread, xnticks\_t idate, xnticks\_t period)  
*Make a thread periodic.*
- int [xnpod\\_wait\\_thread\\_period](#) (unsigned long \*overruns\_r)  
*Wait for the next periodic release point.*
- int [xnpod\\_set\\_thread\\_tslice](#) (struct xnthread \*thread, xnticks\_t quantum)  
*Set thread time-slicing information.*
- int [xnpod\\_add\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Install a nucleus hook.*
- int [xnpod\\_remove\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Remove a nucleus hook.*
- void [xnpod\\_welcome\\_thread](#) (xnthread\_t \*thread, int imask)  
*Thread prologue.*
- int [xnpod\\_trap\\_fault](#) (xnarch\_fltinfo\_t \*fltinfo)  
*Default fault handler.*

### 4.8.1 Detailed Description

Real-time pod services.

## 4.8.2 Function Documentation

### 4.8.2.1 void xnpod\_abort\_thread ( xntthread\_t \* thread )

Abort a thread.

Unconditionally terminates a thread and releases all the nucleus resources it currently holds, regardless of whether the target thread is currently active in kernel or user-space. [xnpod\\_abort\\_thread\(\)](#) should be reserved for use by skin cleanup routines; [xnpod\\_delete\\_thread\(\)](#) should be preferred as the common method for removing threads from a running system.

#### Parameters

*thread* The descriptor address of the terminated thread.

This service forces a call to [xnpod\\_delete\\_thread\(\)](#) for the target thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References XNABORT, XNDORMANT, [xnpod\\_delete\\_thread\(\)](#), and [xnpod\\_suspend\\_thread\(\)](#).

### 4.8.2.2 int xnpod\_add\_hook ( int type, void(\*)(xntthread\_t \*) routine )

Install a nucleus hook.

The nucleus allows to register user-defined routines which get called whenever a specific scheduling event occurs. Multiple hooks can be chained for a single event type, and get called on a FIFO basis.

The scheduling is locked while a hook is executing.

#### Parameters

*type* Defines the kind of hook to install:

- XNHOOK\_THREAD\_START: The user-defined routine will be called on behalf of the starter thread whenever a new thread starts. The descriptor address of the started thread is passed to the routine.
- XNHOOK\_THREAD\_DELETE: The user-defined routine will be called on behalf of the deleter thread whenever a thread is deleted. The descriptor address of the deleted thread is passed to the routine.
- XNHOOK\_THREAD\_SWITCH: The user-defined routine will be called on behalf of the resuming thread whenever a context switch takes place. The descriptor address of the thread which has been switched out is passed to the routine.



### Parameters

*routine* The address of the user-supplied routine to call.

### Returns

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- -EINVAL is returned if type is incorrect.
- -ENOMEM is returned if not enough memory is available from the system heap to add the new hook.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.8.2.3 void `xnpod_delete_thread` ( `xnthread_t * thread` )

Delete a thread.

Terminates a thread and releases all the nucleus resources it currently holds. A thread exists in the system since `xnpod_init_thread()` has been called to create it, so this service must be called in order to destroy it afterwards.

### Parameters

*thread* The descriptor address of the terminated thread.

The target thread's resources may not be immediately removed if this is an active shadow thread running in user-space. In such a case, the mated Linux task is sent a termination signal instead, and the actual deletion is deferred until the task exit event is called.

The DELETE hooks are called on behalf of the calling context (if any). The information stored in the thread control block remains valid until all hooks have been called.

Self-terminating a thread is allowed. In such a case, this service does not return to the caller.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References `xnsched::curr`, `xnsched::status`, `XNABORT`, `XNCANPND`, `XNDEFCAN`, `XNDORMANT`, `XNMIGRATE`, `XNPEND`, `xnpod_schedule()`, `xnpod_unblock_thread()`, `XNREADY`, `XNROOT`, `xnselector_destroy()`, `xnsynch_forget_sleeper()`, `xnsynch_release_all_ownerships()`, `xntimer_destroy()`, and `XNZOMBIE`.

Referenced by `xnpod_abort_thread()`, and `xnpod_shutdown()`.

#### 4.8.2.4 `void xnpod_disable_timesource ( void )`

Stop the core time source.

Releases the hardware timer, and deactivates the master time base.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

References `xntimer_freeze()`.

Referenced by `xnpod_shutdown()`.

#### 4.8.2.5 `void xnpod_dispatch_signals ( void )`

Deliver pending asynchronous signals to the running thread.

**For internal use only.**

This internal routine checks for the presence of asynchronous signals directed to the running thread, and attempts to start the asynchronous service routine (ASR) if any. Called with `nklock` locked, interrupts off.

References `XNASDI`.

Referenced by `xnpod_welcome_thread()`, and `xnshadow_harden()`.

#### 4.8.2.6 `int xnpod_enable_timesource ( void )`

Activate the core time source.

Xenomai implements the notion of time base, by which software timers that belong to different skins may be clocked separately according to distinct frequencies, or aperiodically. In the periodic case, delays and timeouts are given in counts of ticks; the duration of a tick is specified by the time base. In the aperiodic case, timings are directly specified in nanoseconds.

Only a single aperiodic (i.e. tick-less) time base may exist in the system, and the nucleus provides for it through the `nktbase` object. All skins depending on aperiodic timings should bind to the latter, also known as the master time base. Skins depending on periodic timings may create and bind to their own time base. Such a periodic time base is managed as a slave object of the

master one. A cascading software timer, which is fired by the master time base according to the appropriate frequency, triggers in turn the update process of the associated slave time base, which eventually fires the elapsed software timers controlled by the latter.

Xenomai always controls the underlying hardware timer in a tick-less fashion, also known as the oneshot mode. The [xnpod\\_enable\\_timesource\(\)](#) service configures the timer chip as needed, and activates the master time base.

### Returns

0 is returned on success. Otherwise:

- -ENODEV is returned if a failure occurred while configuring the hardware timer.
- -ENOSYS is returned if no active pod exists.

Side-effect: A host timing service is started in order to relay the canonical periodical tick to the underlying architecture, regardless of the frequency used for Xenomai's system tick. This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

### Note

Built-in support for periodic timing depends on `CONFIG_XENO_OPT_TIMING_PERIODIC`.

References `xnsched::htimer`, `xnintr_init()`, and `xntimer_start()`.

Referenced by `xnpod_init()`.

#### 4.8.2.7 `int xnpod_init ( void )`

Initialize the core pod.

Initializes the core interface pod which can subsequently be used to start real-time activities. Once the core pod is active, real-time skins can be stacked over. There can only be a single core pod active in the host environment. Such environment can be confined to a process (e.g. simulator), or expand machine-wide (e.g. I-pipe).

### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the memory manager fails to initialize.

Environments:

This service can be called from:

- Kernel module initialization code

References `xnpod::refcnt`, `xnsched::rootcb`, `xnpod::sched`, `xnpod::status`, `xnpod::tdeleteq`, `xnpod::threadq`, `xnpod::timerlck`, `xnpod::tsliced`, `xnpod::tslicer`, `xnpod::tstartq`, `xnpod::tswitchq`, `xnheap_destroy()`, `xnheap_init()`, `xnheap_set_label()`, `xnpod_enable_timesource()`, `xnpod_shutdown()`, and `xntimer_init()`.

**4.8.2.8** `int xnpod_init_thread ( struct xnthread * thread, const struct xnthread_init_attr * attr, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param )`

Initialize a new thread.

Initializes a new thread attached to the active pod. The thread is left in an innocuous state until it is actually started by [xnpod\\_start\\_thread\(\)](#).

#### Parameters

*thread* The address of a thread descriptor the nucleus will use to store the thread-specific data. This descriptor must always be valid while the thread is active therefore it must be allocated in permanent memory.

#### Warning

Some architectures may require the descriptor to be properly aligned in memory; this is an additional reason for descriptors not to be laid in the program stack where alignment constraints might not always be satisfied.

#### Parameters

*attr* A pointer to an attribute block describing the initial properties of the new thread. Members of this structure are defined as follows:

- **name**: An ASCII string standing for the symbolic name of the thread. This name is copied to a safe place into the thread descriptor. This name might be used in various situations by the nucleus for issuing human-readable diagnostic messages, so it is usually a good idea to provide a sensible value here. The simulator even uses this name intensively to identify threads in the debugging GUI it provides. However, passing NULL here is always legal and means "anonymous".
- **tbase**: The time base descriptor to refer to for all timed operations issued by the new thread. See [xntbase\\_alloc\(\)](#) for detailed explanations about time bases.
- **flags**: A set of creation flags affecting the operation. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the created thread:
  - **XNSUSP** creates the thread in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod\\_resume\\_thread\(\)](#) service for its execution to actually begin, additionally to issuing [xnpod\\_start\\_thread\(\)](#) for it. This flag can also be specified when invoking [xnpod\\_start\\_thread\(\)](#) as a starting mode.
  - **XNFPU** (enable FPU) tells the nucleus that the new thread will use the floating-point unit. In such a case, the nucleus will handle the FPU context save/restore ops upon thread switches

at the expense of a few additional cycles per context switch. By default, a thread is not expected to use the FPU. This flag is simply ignored when the nucleus runs on behalf of a userspace-based real-time control layer since the FPU management is always active if present.

- **stacksize:** The size of the stack (in bytes) for the new thread. If zero is passed, the nucleus will use a reasonable pre-defined size depending on the underlying real-time control layer.
- **ops:** A pointer to a structure defining the class-level operations available for this thread. Fields from this structure must have been set appropriately by the caller.

#### Parameters

*sched\_class* The initial scheduling class the new thread should be assigned to.

*sched\_param* The initial scheduling parameters to set for the new thread; *sched\_param* must be valid within the context of *sched\_class*.

#### Returns

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- -EINVAL is returned if *attr->flags* has invalid bits set.
- -ENOMEM is returned if not enough memory is available from the system heap to create the new thread's stack.

Side-effect: This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

References XNDORMANT, XNFPU, *xnpod\_suspend\_thread()*, XNSHADOW, and XNSUSP.

#### 4.8.2.9 `int xnpod_migrate_thread ( int cpu )`

Migrate the current thread.

This call makes the current thread migrate to another CPU if its affinity allows it.

#### Parameters

*cpu* The destination CPU.

#### Return values

0 if the thread could migrate ;

- EPERM* if the calling context is asynchronous, or the current thread affinity forbids this migration ;
- EBUSY* if the scheduler is locked.

References `xnpod_schedule()`.

#### 4.8.2.10 `int xnpod_remove_hook ( int type, void(*) (xnthread_t *) routine )`

Remove a nucleus hook.

This service removes a nucleus hook previously registered using [xnpod\\_add\\_hook\(\)](#).

##### Parameters

- type* Defines the kind of hook to remove among `XNHOOK_THREAD_START`, `XNHOOK_THREAD_DELETE` and `XNHOOK_THREAD_SWITCH`.
- routine* The address of the user-supplied routine to remove.

##### Returns

0 is returned on success. Otherwise, `-EINVAL` is returned if *type* is incorrect or if the routine has never been registered before.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.8.2.11 `void xnpod_restart_thread ( xnthread_t * thread )`

Restart a thread.

Restarts a previously started thread. The thread is first terminated then respawned using the same information that prevailed when it was first started, including the mode bits and interrupt mask initially passed to the [xnpod\\_start\\_thread\(\)](#) service. As a consequence of this call, the thread entry point is rerun.

##### Parameters

- thread* The descriptor address of the affected thread which must have been previously started by the [xnpod\\_start\\_thread\(\)](#) service.

Self-restarting a thread is allowed. However, restarting the root thread is not. Restarting a thread which was never started once leads to a null-effect.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: possible.

References `xnpod_schedule()`, `XNRESTART`, `XNROOT`, `XNSHADOW`, and `XNSTARTED`.

#### 4.8.2.12 `void xnpod_resume_thread ( xntthread_t * thread, xnflags_t mask )`

Resume a thread.

Resumes the execution of a thread previously suspended by one or more calls to `xnpod_suspend_thread()`. This call removes a suspensive condition affecting the target thread. When all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

##### Parameters

*thread* The descriptor address of the resumed thread.

*mask* The suspension mask specifying the suspensive condition to remove from the thread's wait mask. Possible values usable by the caller are:

- `XNSUSP`. This flag removes the explicit suspension condition. This condition might be additive to the `XNPEND` condition.
- `XNDELAY`. This flag removes the counted delay wait condition.
- `XNPEND`. This flag removes the resource wait condition. If a watchdog is armed, it is automatically disarmed by this call. Unlike the two previous conditions, only the current thread can set this condition for itself, i.e. no thread can force another one to pend on a resource.

When the thread is eventually resumed by one or more calls to `xnpod_resume_thread()`, the caller of `xnpod_suspend_thread()` in the awakened thread that suspended itself should check for the following bits in its own information mask to determine what caused its wake up:

- `XNRMID` means that the caller must assume that the pended synchronization object has been destroyed (see `xnsynch_flush()`).
- `XNTIMEO` means that the delay elapsed, or the watchdog went off before the corresponding synchronization object was signaled.
- `XNBREAK` means that the wait has been forcibly broken by a call to `xnpod_unblock_thread()`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine

- Kernel-based task
- User-space task

Rescheduling: never.

References `xnsched::curr`, `XNDELAY`, `XNHELD`, `XNPEND`, `XNREADY`, `xnsynch_forget_sleeper()`, and `xntimer_stop()`.

Referenced by `xnpod_start_thread()`, `xnpod_unblock_thread()`, `xnsynch_flush()`, `xnsynch_wakeup_one_sleeper()`, and `xnsynch_wakeup_this_sleeper()`.

#### 4.8.2.13 `void xnpod_schedule ( void ) [inline, static]`

Rescheduling procedure entry point.

This is the central rescheduling routine which should be called to validate and apply changes which have previously been made to the nucleus scheduling state, such as suspending, resuming or changing the priority of threads. This call first determines if a thread switch should take place, and performs it as needed. `xnpod_schedule()` schedules out the current thread if:

- the current thread is now blocked or deleted.
- a runnable thread from a higher priority scheduling class is waiting for the CPU.
- the current thread does not lead the runnable threads from its own scheduling class (e.g. round-robin in the RT class).

The nucleus implements a lazy rescheduling scheme so that most of the services affecting the threads state MUST be followed by a call to the rescheduling procedure for the new scheduling state to be applied. In other words, multiple changes on the scheduler state can be done in a row, waking threads up, blocking others, without being immediately translated into the corresponding context switches, like it would be necessary would it appear that a higher priority thread than the current one became runnable for instance. When all changes have been applied, the rescheduling procedure is then called to consider those changes, and possibly replace the current thread by another one.

As a notable exception to the previous principle however, every action which ends up suspending or deleting the current thread begets an immediate call to the rescheduling procedure on behalf of the service causing the state transition. For instance, self-suspension, self-destruction, or sleeping on a synchronization object automatically leads to a call to the rescheduling procedure, therefore the caller does not need to explicitly issue `xnpod_schedule()` after such operations.

The rescheduling procedure always leads to a null-effect if it is called on behalf of an ISR or callout. Any outstanding scheduler lock held by the outgoing thread will be restored when the thread is scheduled back in.

Calling this procedure with no applicable context switch pending is harmless and simply leads to a null-effect.

Side-effects:

- If an asynchronous service routine exists, the pending asynchronous signals are delivered to a resuming thread or on behalf of the caller before it returns from the procedure if no context switch has taken place. This behaviour can be disabled by setting the `XNASDI` flag in the thread's status mask by calling `xnpod_set_thread_mode()`.



Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine, although this leads to a no-op.
- Kernel-based task
- User-space task

#### Note

The switch hooks are called on behalf of the resuming thread.

References `xnsched::lflags`, and `xnsched::status`.

Referenced by `xnpod_delete_thread()`, `xnpod_migrate_thread()`, `xnpod_restart_thread()`, `xnpod_shutdown()`, `xnpod_start_thread()`, `xnpod_stop_thread()`, `xnpod_suspend_thread()`, `xnregistry_enter()`, `xnregistry_put()`, `xnselect_bind()`, and `xnselect_destroy()`.

#### 4.8.2.14 `xnflags_t xnpod_set_thread_mode ( xnthread_t * thread, xnflags_t clrmask, xnflags_t setmask )`

Change a thread's control mode.

Change the control mode of a given thread. The control mode affects the behaviour of the nucleus regarding the specified thread.

#### Parameters

*thread* The descriptor address of the affected thread.

*clrmask* Clears the corresponding bits from the control field before *setmask* is applied. The scheduler lock held by the current thread can be forcibly released by passing the XNLOCK bit in this mask. In this case, the lock nesting count is also reset to zero.

*setmask* The new thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the thread:

- XNLOCK causes the thread to lock the scheduler. The target thread will have to call the `xnpod_unlock_sched()` service to unlock the scheduler or clear the XNLOCK bit forcibly using this service. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- XNASDI disables the asynchronous signal handling for this thread. See [xnpod\\_schedule\(\)](#) for more on this.
- XNRPIOFF disables thread priority coupling between Xenomai and Linux schedulers. This bit prevents the root Linux thread from inheriting the priority of the running shadow Xenomai thread. Use `CONFIG_XENO_OPT_RPIOFF` to globally disable priority coupling.

Environments:

This service can be called from:

- Kernel-based task
- User-space task in primary mode.

Rescheduling: never, therefore, the caller should reschedule if XNLOCK has been passed into *clrmask*.

References XNLOCK.

#### 4.8.2.15 `int xnpod_set_thread_periodic ( xnthread_t * thread, xnticks_t idate, xnticks_t period )`

Make a thread periodic.

Make a thread periodic by programming its first release point and its period in the processor time line. Subsequent calls to `xnpod_wait_thread_period()` will delay the thread until the next periodic release point in the processor timeline is reached.

##### Parameters

*thread* The descriptor address of the affected thread. This thread is immediately delayed until the first periodic release point is reached.

*idate* The initial (absolute) date of the first release point, expressed in clock ticks (see note). The affected thread will be delayed until this point is reached. If *idate* is equal to XN\_INFINITE, the current system date is used, and no initial delay takes place.

*period* The period of the thread, expressed in clock ticks (see note). As a side-effect, passing XN\_INFINITE attempts to stop the thread's periodic timer; in the latter case, the routine always exits successfully, regardless of the previous state of this timer.

##### Returns

0 is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *idate* is different from XN\_INFINITE and represents a date in the past.
- -EWOULDBLOCK is returned if the relevant time base has not been initialized by a call to `xnpod_init_timebase()`.
- -EINVAL is returned if *period* is different from XN\_INFINITE but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the operation affects the current thread and *idate* has not elapsed yet.

**Note**

The *idate* and *period* values will be interpreted as jiffies if *thread* is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

References `XNDELAY`, `xnpod_suspend_thread()`, `xntimer_start()`, and `xntimer_stop()`.

#### 4.8.2.16 `int xnpod_set_thread_schedparam ( struct xnthread * thread, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param )`

Change the base scheduling parameters of a thread.

Changes the base scheduling policy and paramaters of a thread. If the thread is currently blocked, waiting in priority-pending mode (`XNSYNCH_PRIO`) for a synchronization object to be signaled, the nucleus will attempt to reorder the object's wait queue so that it reflects the new sleeper's priority, unless the `XNSYNCH_DREORD` flag has been set for the pended object.

**Parameters**

*thread* The descriptor address of the affected thread.

*sched\_class* The new scheduling class the thread should be assigned to.

*sched\_param* The scheduling parameters to set for the thread; *sched\_param* must be valid within the context of *sched\_class*.

It is absolutely required to use this service to change a thread priority, in order to have all the needed housekeeping chores correctly performed. i.e. Do *not* call `xnsched_set_policy()` directly or worse, change the `thread.cprio` field by hand in any case.

**Returns**

0 is returned on success. Otherwise, a negative error code indicates the cause of a failure that happened in the scheduling class implementation for *sched\_class*. Invalid parameters passed into *sched\_param* are common causes of error.

Side-effects:

- This service does not call the rescheduling procedure but may affect the state of the runnable queue for the previous and new scheduling classes.
- Assigning the same scheduling class and parameters to a running or ready thread moves it to the end of the runnable queue, thus causing a manual round-robin.
- If the thread is a user-space shadow, this call propagates the request to the mated Linux task.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.8.2.17 `int xnpod_set_thread_tslice ( struct xnthread * thread, xnticks_t quantum )`

Set thread time-slicing information.

Update the time-slicing information for a given thread. This service enables or disables round-robin scheduling for the thread, depending on the value of *quantum*. By default, times-slicing is disabled for a new thread initialized by a call to `xnpod_init_thread()`.

##### Parameters

*thread* The descriptor address of the affected thread.

*quantum* The time quantum assigned to the thread expressed in time-slicing ticks (see note). If *quantum* is different from `XN_INFINITE`, the time-slice for the thread is set to that value and its current time credit is refilled (i.e. the thread is given a full time-slice to run next). Otherwise, if *quantum* equals `XN_INFINITE`, time-slicing is stopped for that thread.

##### Returns

0 is returned upon success. Otherwise:

- `-EINVAL` is returned if the base scheduling class of the target thread does not support time-slicing.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

##### Note

If *thread* is bound to a periodic timebase, *quantum* represents the number of periodic ticks in that timebase. Otherwise, if *thread* is bound to the master time base, a full time-slice will last: *quantum* \* `CONFIG_XENO_OPT_TIMING_VIRTICK`.

References `XNRRB`, `xntimer_start()`, and `xntimer_stop()`.

#### 4.8.2.18 `void xnpod_shutdown ( int xtype )`

Shutdown the current pod.

Forcibly shutdowns the active pod. All existing nucleus threads (but the root one) are terminated, and the system heap is freed.

##### Parameters

*xtype* An exit code passed to the host environment who started the nucleus. Zero is always interpreted as a successful return.

The nucleus never calls this routine directly. Skins should provide their own shutdown handlers which end up calling [xnpod\\_shutdown\(\)](#) after their own housekeeping chores have been carried out.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References [xnheap\\_destroy\(\)](#), [xnpod\\_delete\\_thread\(\)](#), [xnpod\\_disable\\_timesource\(\)](#), [xnpod\\_schedule\(\)](#), [XNROOT](#), and [xntimer\\_destroy\(\)](#).

Referenced by [xnpod\\_init\(\)](#).

#### 4.8.2.19 `int xnpod_start_thread ( xnthread_t * thread, const struct xnthread_start_attr * attr )`

Initial start of a newly created thread.

Starts a (newly) created thread, scheduling it for the first time. This call releases the target thread from the XNDORMANT state. This service also sets the initial mode and interrupt mask for the new thread.

##### Parameters

*thread* The descriptor address of the affected thread which must have been previously initialized by the [xnpod\\_init\\_thread\(\)](#) service.

*attr* A pointer to an attribute block describing the execution properties of the new thread. Members of this structure are defined as follows:

- **mode**: The initial thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the started thread:
- **XNLOCK** causes the thread to lock the scheduler when it starts. The target thread will have to call the [xnpod\\_unlock\\_sched\(\)](#) service to unlock the scheduler. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- **XNASDI** disables the asynchronous signal handling for this thread. See [xnpod\\_schedule\(\)](#) for more on this.
- **XNSUSP** makes the thread start in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod\\_resume\\_thread\(\)](#) service for its execution to actually begin.
- **imask**: The interrupt mask that should be asserted when the thread starts. The processor interrupt state will be set to the given value when the thread starts running. The interpretation of this value might be different across real-time layers, but a non-zero value should always mark an interrupt masking in effect (e.g. [local\\_irq\\_disable\(\)](#)). Conversely, a zero value should always mark a fully preemptible state regarding interrupts (e.g. [local\\_irq\\_enable\(\)](#)).

- *affinity*: The processor affinity of this thread. Passing `XNPOD_ALL_CPUS` or an empty affinity set means "any cpu".
- *entry*: The address of the thread's body routine. In other words, it is the thread entry point.
- *cookie*: A user-defined opaque cookie the nucleus will pass to the emerging thread as the sole argument of its entry point.

The START hooks are called on behalf of the calling context (if any).

#### Return values

- `0` if *thread* could be started ;
- `-EBUSY` if *thread* was not dormant or stopped ;
- `-EINVAL` if the value of *attr->affinity* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References `XNDORMANT`, `xnpod_resume_thread()`, `xnpod_schedule()`, `XNREADY`, `XN-SHADOW`, `XNSTARTED`, and `XNSUSP`.

Referenced by `xnshadow_map()`.

#### 4.8.2.20 void xnpod\_stop\_thread ( xnthread\_t \* thread )

Stop a thread.

Stop a previously started thread. The thread is put back into the dormant state; however, it is not deleted from the system.

#### Parameters

- *thread* The descriptor address of the affected thread which must have been previously started by the `xnpod_start_thread()` service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References `XNDORMANT`, `xnpod_schedule()`, `xnpod_suspend_thread()`, and `XNROOT`.

#### 4.8.2.21 void xnpod\_suspend\_thread ( xntthread\_t \* thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \* wchan )

Suspend a thread.

Suspends the execution of a thread according to a given suspensive condition. This thread will not be eligible for scheduling until all the pending suspensive conditions set by this service are removed by one or more calls to [xnpod\\_resume\\_thread\(\)](#).

##### Parameters

*thread* The descriptor address of the suspended thread.

*mask* The suspension mask specifying the suspensive condition to add to the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag forcibly suspends a thread, regardless of any resource to wait for. A reverse call to [xnpod\\_resume\\_thread\(\)](#) specifying the XNSUSP bit must be issued to remove this condition, which is cumulative with other suspension bits. *wchan* should be NULL when using this suspending mode.
- XNDELAY. This flag denotes a counted delay wait (in ticks) which duration is defined by the value of the timeout parameter.
- XNPEND. This flag denotes a wait for a synchronization object to be signaled. The *wchan* argument must point to this object. A timeout value can be passed to bound the wait. This suspending mode should not be used directly by the client interface, but rather through the [xnsynch\\_sleep\\_on\(\)](#) call.

##### Parameters

*timeout* The timeout which may be used to limit the time the thread pends on a resource. This value is a wait time given in ticks (see note). It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing XN\_INFINITE and setting *timeout\_mode* to XN\_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer. If the current operation mode of the system timer is oneshot and *timeout* elapses before [xnpod\\_suspend\\_thread\(\)](#) has completed, then the target thread will not be suspended, and this routine leads to a null effect.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

*wchan* The address of a pended resource. This parameter is used internally by the synchronization object implementation code to specify on which object the suspended thread pends. NULL is a legitimate value when this parameter does not apply to the current suspending mode (e.g. XNSUSP).

##### Note

If the target thread is a shadow which has received a Linux-originated signal, then this service immediately exits without suspending the thread, but raises the XNBREAK condition in its information mask.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread suspends itself.

#### Note

The *timeout* value will be interpreted as jiffies if *thread* is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

References `xnsched::curr`, `XNBREAK`, `XNDELAY`, `XNDORMANT`, `XNHELD`, `XNKICKED`, `xnpod_schedule()`, `XNREADY`, `XNRELAX`, `XNRMID`, `XNROBBED`, `XNROOT`, `XNSHADOW`, `XNSUSP`, `xnsynch_forget_sleeper()`, `XNTIMEO`, `xntimer_start()`, and `XNWAKEN`.

Referenced by `xnpod_abort_thread()`, `xnpod_init_thread()`, `xnpod_set_thread_periodic()`, `xnpod_stop_thread()`, `xnpod_trap_fault()`, `xnpod_wait_thread_period()`, `xnshadow_map()`, `xnshadow_relax()`, `xnsynch_acquire()`, and `xnsynch_sleep_on()`.

#### 4.8.2.22 void xnpod\_trap\_fault ( xnarch\_fltinfo\_t \* fltinfo )

Default fault handler.

This is the default handler which is called whenever an uncontrolled exception or fault is caught. If the fault is caught on behalf of a real-time thread, the fault is not propagated to the host system. Otherwise, the fault is unhandled by the nucleus and simply propagated.

#### Parameters

*fltinfo* An opaque pointer to the arch-specific buffer describing the fault. The actual layout is defined by the `xnarch_fltinfo_t` type in each arch-dependent layer file.

References `xnpod_suspend_thread()`, `xnshadow_relax()`, and `XNSUSP`.

#### 4.8.2.23 int xnpod\_unblock\_thread ( xnthread\_t \* thread )

Unblock a thread.

Breaks the thread out of any wait it is currently in. This call removes the `XNDELAY` and `XNPEND` suspensive conditions previously put by `xnpod_suspend_thread()` on the target thread. If all suspensive conditions are gone, the thread is left in a `READY` state at which point it becomes eligible anew for scheduling.

#### Parameters

*thread* The descriptor address of the unblocked thread.

This call neither releases the thread from the `XNSUSP`, `XNRELAX`, `XNDORMANT` or `XNHELD` suspensive conditions.

When the thread resumes execution, the `XNBREAK` bit is set in the unblocked thread's information mask. Unblocking a non-blocked thread is perfectly harmless.



**Returns**

non-zero is returned if the thread was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNBREAK, XNDELAY, XNPEND, and `xnpod_resume_thread()`.

Referenced by `xnpod_delete_thread()`.

**4.8.2.24 int xnpod\_wait\_thread\_period ( unsigned long \* overruns\_r )**

Wait for the next periodic release point.

Make the current thread wait for the next periodic release point in the processor time line.

**Parameters**

*overruns\_r* If non-NULL, *overruns\_r* must be a pointer to a memory location which will be written with the count of pending overruns. This value is copied only when `xnpod_wait_thread_period()` returns -ETIMEDOUT or success; the memory location remains unmodified otherwise. If NULL, this count will never be copied back.

**Returns**

0 is returned upon success; if *overruns\_r* is valid, zero is copied to the pointed memory location. Otherwise:

- -EWOULDBLOCK is returned if `xnpod_set_thread_periodic()` has not previously been called for the calling thread.
- -EINTR is returned if `xnpod_unblock_thread()` has been called for the waiting thread before the next periodic release point has been reached. In this case, the overrun counter is reset too.
- -ETIMEDOUT is returned if the timer has overrun, which indicates that one or more previous release points have been missed by the calling thread. If *overruns\_r* is valid, the count of pending overruns is copied to the pointed memory location.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task

Rescheduling: always, unless the current release point has already been reached. In the latter case, the current thread immediately returns from this service without being delayed.

References XNBREAK, XNDELAY, `xnpod_suspend_thread()`, and `xntimer_get_overruns()`.

#### 4.8.2.25 void `xnpod_welcome_thread ( xnthread_t * thread, int imask )`

Thread prologue.

**For internal use only.**

This internal routine is called on behalf of a (re)starting thread's prologue before the user entry point is invoked. This call is reserved for internal housekeeping chores and cannot be inlined.

Entered with `nklock` locked, `irqs` off.

References XNLOCK, `xnpod_dispatch_signals()`, and XNRESTART.

## 4.9 Registry services.

Collaboration diagram for Registry services.:



### Files

- file [registry.h](#)  
*This file is part of the Xenomai project.*
- file [registry.c](#)  
*This file is part of the Xenomai project.*

### Functions

- int [xnregistry\\_enter](#) (const char \*key, void \*objaddr, xnhandle\_t \*phandle, struct xnpnode \*pnode)

*Register a real-time object.*

- `int xnregistry_bind (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)`

*Bind to a real-time object.*

- `int xnregistry_remove (xnhandle_t handle)`

*Forcibly unregister a real-time object.*

- `int xnregistry_remove_safe (xnhandle_t handle, xnticks_t timeout)`

*Unregister an idle real-time object.*

- `void * xnregistry_get (xnhandle_t handle)`

*Find and lock a real-time object into the registry.*

- `u_long xnregistry_put (xnhandle_t handle)`

*Unlock a real-time object from the registry.*

- `void * xnregistry_fetch (xnhandle_t handle)`

*Find a real-time object into the registry.*

### 4.9.1 Detailed Description

The registry provides a mean to index real-time object descriptors created by Xenomai skins on unique alphanumeric keys. When labeled this way, a real-time object is globally exported; it can be searched for, and its descriptor returned to the caller for further use; the latter operation is called a "binding". When no object has been registered under the given name yet, the registry can be asked to set up a rendez-vous, blocking the caller until the object is eventually registered.

### 4.9.2 Function Documentation

#### 4.9.2.1 `int xnregistry_bind ( const char * key, xnticks_t timeout, int timeout_mode, xnhandle_t * phandle )`

Bind to a real-time object.

This service retrieves the registry handle of a given object identified by its key. Unless otherwise specified, this service will block the caller if the object is not registered yet, waiting for such registration to occur.

#### Parameters

**key** A valid NULL-terminated string which identifies the object to bind to.

**timeout** The timeout which may be used to limit the time the thread wait for the object to be registered. This value is a wait time given in ticks (see note). It can either be relative, absolute monotonic (XN\_ABSOLUTE), or absolute adjustable (XN\_REALTIME) depending on *timeout\_mode*. Passing XN\_INFINITE and setting *timeout\_mode* to XN\_RELATIVE specifies an unbounded wait. Passing XN\_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry. All other values are used as a wait limit.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

*phandle* A pointer to a memory location which will be written upon success with the generic handle defined by the registry for the retrieved object. Contents of this memory is undefined upon failure.

### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* is NULL.
- -EINTR is returned if [xnpod\\_unblock\\_thread\(\)](#) has been called for the waiting thread before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to XN\_NONBLOCK and the searched object is not registered on entry. As a special exception, this error is also returned if this service should block, but was called from a context which cannot sleep (e.g. interrupt, non-realtime or scheduler locked).
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN\_NONBLOCK.
- Kernel-based thread.

Rescheduling: always unless the request is immediately satisfied or *timeout* specifies a non-blocking operation.

### Note

The *timeout* value will be interpreted as jiffies if *thread* is bound to a periodic time base (see [xnpod\\_init\\_thread](#)), or nanoseconds otherwise.

References XNBREAK, [xnsynch\\_sleep\\_on\(\)](#), [xntbase\\_get\\_time\(\)](#), and XNTIMEO.

**4.9.2.2** `int xnregistry_enter ( const char * key, void * objaddr, xnhandle_t * phandle, struct xnpnode * pnode )`

Register a real-time object.

This service allocates a new registry slot for an associated object, and indexes it by an alphanumeric key for later retrieval.

**Parameters**

*key* A valid NULL-terminated string by which the object will be indexed and later retrieved in the registry. Since it is assumed that such key is stored into the registered object, it will *not* be copied but only kept by reference in the registry. Pass an empty string if the object shall only occupy a registry slot for handle-based lookups.

*objaddr* An opaque pointer to the object to index by *key*.

*phandle* A pointer to a generic handle defined by the registry which will uniquely identify the indexed object, until the latter is unregistered using the [xnregistry\\_remove\(\)](#) service.

*pnode* A pointer to an optional /proc node class descriptor. This structure provides the information needed to export all objects from the given class through the /proc filesystem, under the /proc/xenomai/registry entry. Passing NULL indicates that no /proc support is available for the newly registered object.

**Returns**

0 is returned upon success. Otherwise:

- -EINVAL is returned if *objaddr* are NULL, or if *key* contains an invalid '/' character.
- -ENOMEM is returned if the system fails to get enough dynamic memory from the global real-time heap in order to register the object.
- -EEXIST is returned if the *key* is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: possible.

References [xnpod\\_schedule\(\)](#), and [xnsynch\\_init\(\)](#).

**4.9.2.3 void\* xnregistry\_fetch ( xnhandle\_t handle )**

Find a real-time object into the registry.

This service retrieves an object from its handle into the registry and returns the memory address of its descriptor.

**Parameters**

*handle* The generic handle of the object to fetch. If XNOBJECT\_SELF is passed, the object is the calling Xenomai thread.

**Returns**

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT\_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT\_SELF.
- Kernel-based thread

Rescheduling: never.

#### 4.9.2.4 void\* xnregistry\_get ( xnhandle\_t handle )

Find and lock a real-time object into the registry.

This service retrieves an object from its handle into the registry and prevents its removal atomically. A locking count is tracked, so that [xnregistry\\_get\(\)](#) and [xnregistry\\_put\(\)](#) must be used in pair.

##### Parameters

*handle* The generic handle of the object to find and lock. If XNOBJECT\_SELF is passed, the object is the calling Xenomai thread.

##### Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT\_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT\_SELF.
- Kernel-based thread.

Rescheduling: never.

#### 4.9.2.5 u\_long xnregistry\_put ( xnhandle\_t handle )

Unlock a real-time object from the registry.

This service decrements the lock count of a registered object previously locked by a call to [xnregistry\\_get\(\)](#). The object is actually unlocked from the registry when the locking count falls down to zero, thus waking up any thread currently blocked on [xnregistry\\_remove\(\)](#) for unregistering it.

##### Parameters

*handle* The generic handle of the object to unlock. If XNOBJECT\_SELF is passed, the object is the calling Xenomai thread.

**Returns**

The decremented lock count is returned upon success. Zero is also returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT\_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT\_SELF.
- Kernel-based thread

Rescheduling: possible if the lock count falls down to zero and some thread is currently waiting for the object to be unlocked.

References `xnpod_schedule()`, and `xnsynch_flush()`.

**4.9.2.6 int xnregistry\_remove ( xnhandle\_t handle )**

Forcibly unregister a real-time object.

This service forcibly removes an object from the registry. The removal is performed regardless of the current object's locking status.

**Parameters**

*handle* The generic handle of the object to remove.

**Returns**

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: never.

Referenced by `xnregistry_remove_safe()`.

**4.9.2.7 int xnregistry\_remove\_safe ( xnhandle\_t handle, xnticks\_t timeout )**

Unregister an idle real-time object.

This service removes an object from the registry. The caller might sleep as a result of waiting for the target object to be unlocked prior to the removal (see [xnregistry\\_put\(\)](#)).

## Parameters

*handle* The generic handle of the object to remove.

*timeout* If the object is locked on entry, *param* gives the number of clock ticks to wait for the unlocking to occur (see note). Passing XN\_INFINITE causes the caller to block indefinitely until the object is unlocked. Passing XN\_NONBLOCK causes the service to return immediately without waiting if the object is locked on entry.

## Returns

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.
- -EWOULDBLOCK is returned if *timeout* is equal to XN\_NONBLOCK and the object is locked on entry.
- -EBUSY is returned if *handle* refers to a locked object and the caller could not sleep until it is unlocked.
- -ETIMEDOUT is returned if the object cannot be removed within the specified amount of time.
- -EINTR is returned if [xnpod\\_unblock\\_thread\(\)](#) has been called for the calling thread waiting for the object to be unlocked.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN\_NONBLOCK.
- Kernel-based thread.

Rescheduling: possible if the object to remove is currently locked and the calling context can sleep.

## Note


The *timeout* value will be interpreted as jiffies if the current thread is bound to a periodic time base (see [xnpod\\_init\\_thread\(\)](#)), or nanoseconds otherwise.

References XNBREAK, [xnregistry\\_remove\(\)](#), [xnsynch\\_sleep\\_on\(\)](#), and XNTIMEO.



## 4.10 File descriptors events multiplexing services.

Collaboration diagram for File descriptors events multiplexing services.:



### Files

- file [select.h](#)  
*file descriptors events multiplexing header.*
- file [select.c](#)  
*file descriptors events multiplexing.*

### Functions

- void [xnselect\\_init](#) (struct xnselect \*select\_block)  
*Initialize a struct xnselect structure.*
- int [xnselect\\_bind](#) (struct xnselect \*select\_block, struct xnselect\_binding \*binding, struct xnselector \*selector, unsigned type, unsigned index, unsigned state)  
*Bind a file descriptor (represented by its xnselect structure) to a selector block.*
- void [xnselect\\_destroy](#) (struct xnselect \*select\_block)  
*Destroy the xnselect structure associated with a file descriptor.*
- int [xnselector\\_init](#) (struct xnselector \*selector)  
*Initialize a selector structure.*
- int [xnselect](#) (struct xnselector \*selector, fd\_set \*out\_fds[XNSELECT\_MAX\_TYPES], fd\_set \*in\_fds[XNSELECT\_MAX\_TYPES], int nfds, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.*
- void [xnselector\\_destroy](#) (struct xnselector \*selector)  
*Destroy a selector block.*

### 4.10.1 Detailed Description

File descriptors events multiplexing services.

This module implements the services needed for implementing the posix "select" service, or any other events multiplexing services.

Following the implementation of the posix select service, this module defines three types of events:

- *XNSELECT\_READ* meaning that a file descriptor is ready for reading;
- *XNSELECT\_WRITE* meaning that a file descriptor is ready for writing;
- *XNSELECT\_EXCEPT* meaning that a file descriptor received an exceptional event.

It works by defining two structures:

- a *struct xnselect* structure, which should be added to every file descriptor for every event type (read, write, or except);
- a *struct xnselector* structure, the selection structure, passed by the thread calling the xnselect service, where this service does all its housekeeping.

### 4.10.2 Function Documentation

**4.10.2.1** `int xnselect ( struct xnselector * selector, fd_set * out_fds[XNSELECT_MAX_TYPES], fd_set * in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode )`

Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.

#### Parameters

*selector* structure to check for pending events

*out\_fds* The set of descriptors with pending events if a strictly positive number is returned, or the set of descriptors not yet bound if -ECHRNG is returned;

*in\_fds* the set of descriptors which events should be checked

*nfds* the highest-numbered descriptor in any of the *in\_fds* sets, plus 1;

*timeout* the timeout, whose meaning depends on *timeout\_mode*, note that `xnselect()` pass *timeout* and *timeout\_mode* unchanged to `xnsynch_sleep_on`, so passing a relative value different from `XN_INFINITE` as a timeout with *timeout\_mode* set to `XN_RELATIVE`, will cause a longer sleep than expected if the sleep is interrupted.

*timeout\_mode* the mode of *timeout*.

#### Return values

-*EINVAL* if *nfds* is negative;

-*ECHRNG* if some of the descriptors passed in *in\_fds* have not yet been registered with `xnselect_bind()`, *out\_fds* contains the set of such descriptors;

-*EINTR* if *xnselect* was interrupted while waiting;

0 in case of timeout.

*the* number of file descriptors having received an event.

References `XNBREAK`, `xnsynch_sleep_on()`, and `XNTIMEO`.

#### 4.10.2.2 `int xnselect_bind ( struct xnselect * select_block, struct xnselect_binding * binding, struct xnselector * selector, unsigned type, unsigned index, unsigned state )`

Bind a file descriptor (represented by its *xnselect* structure) to a selector block.

##### Parameters

*select\_block* pointer to the *struct xnselect* to be bound;  
*binding* pointer to a newly allocated (using `xnmalloc`) *struct xnselect\_binding*;  
*selector* pointer to the selector structure;  
*type* type of events (`XNSELECT_READ`, `XNSELECT_WRITE`, or `XNSELECT_EXCEPT`);  
*index* index of the file descriptor (represented by *select\_block*) in the bit fields used by the *selector* structure;  
*state* current state of the file descriptor>.

*select\_block* must have been initialized with `xnselect_init()`, the *xnselector* structure must have been initialized with `xnselector_init()`, *binding* may be uninitialized.

This service must be called with `nklock` locked, `irqs` off. For this reason, the *binding* parameter must have been allocated by the caller outside the locking section.

##### Return values

`-EINVAL` if *type* or *index* is invalid;  
`0` otherwise.

References `xnpod_schedule()`.

#### 4.10.2.3 `void xnselect_destroy ( struct xnselect * select_block )`

Destroy the *xnselect* structure associated with a file descriptor.

Any binding with a *xnselector* block is destroyed.

##### Parameters

*select\_block* pointer to the *xnselect* structure associated with a file descriptor

References `xnpod_schedule()`.

#### 4.10.2.4 `void xnselect_init ( struct xnselect * select_block )`

Initialize a *struct xnselect* structure.

This service must be called to initialize a *struct xnselect* structure before it is bound to a selector by the means of `xnselect_bind()`.

##### Parameters

*select\_block* pointer to the *xnselect* structure to be initialized

#### 4.10.2.5 void xnselector\_destroy ( struct xnselector \* selector )

Destroy a selector block.

All bindings with file descriptor are destroyed.

##### Parameters

*selector* the selector block to be destroyed

Referenced by xnpod\_delete\_thread().

#### 4.10.2.6 int xnselector\_init ( struct xnselector \* selector )

Initialize a selector structure.

##### Parameters

*selector* The selector structure to be initialized.


##### Return values

0

References xnsynch\_init().

## 4.11 Real-time shadow services.

Collaboration diagram for Real-time shadow services.:



## Files

- file [shadow.c](#)  
*Real-time shadow services.*

## Functions

- int [xnshadow\\_harden](#) (void)  
*Migrate a Linux task to the Xenomai domain.*

- void `xnshadow_relax` (int notify, int reason)  
*Switch a shadow thread back to the Linux domain.*
- int `xnshadow_map` (xnthread\_t \*thread, xncompletion\_t \_\_user \*u\_completion, unsigned long \_\_user \*u\_mode\_offset)  
*Create a shadow thread context.*
- xnshadow\_ppd\_t \* `xnshadow_ppd_get` (unsigned muxid)  
*Return the per-process data attached to the calling process.*

### 4.11.1 Detailed Description

Real-time shadow services.

### 4.11.2 Function Documentation

#### 4.11.2.1 int xnshadow\_harden ( void )

Migrate a Linux task to the Xenomai domain.

**For internal use only.**

This service causes the transition of "current" from the Linux domain to Xenomai. This is obtained by asking the gatekeeper to resume the shadow mated with "current" then triggering the rescheduling procedure in the Xenomai domain. The shadow will resume in the Xenomai domain as returning from `schedule()`.

Environments:

This service can be called from:

- User-space thread operating in secondary (i.e. relaxed) mode.

Rescheduling: always.

References `XNATOMIC`, `XNDEBUG`, `xnpod_dispatch_signals()`, `XNRELAX`, and `xnshadow_relax()`.

Referenced by `xnshadow_map()`.

#### 4.11.2.2 int xnshadow\_map ( xnthread\_t \* thread, xncompletion\_t \_\_user \* u\_completion, unsigned long \_\_user \* u\_mode\_offset )

Create a shadow thread context.

**For internal use only.**

This call maps a nucleus thread to the "current" Linux task. The priority and scheduling class of the underlying Linux task are not affected; it is assumed that the interface library did set them appropriately before issuing the shadow mapping request.

**Parameters**

*thread* The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to [xnpod\\_init\\_thread\(\)](#).

*u\_completion* is the address of an optional completion descriptor aimed at synchronizing our parent thread with us. If non-NULL, the information [xnshadow\\_map\(\)](#) will store into the completion block will be later used to wake up the parent thread when the current shadow has been initialized. In the latter case, the new shadow thread is left in a dormant state (XNDORMANT) after its creation, leading to the suspension of "current" in the Linux domain, only processing signals. Otherwise, the shadow thread is immediately started and "current" immediately resumes in the Xenomai domain from this service.

*u\_mode\_offset* is the address of a user space address where we will store the offset of the "u\_mode" thread variable in the process shared heap. This thread variable reflects the current thread mode (primary or secondary). The nucleus will try to update the variable before switching to secondary or after switching from primary mode.

**Returns**

0 is returned on success. Otherwise:

- -ERESTARTSYS is returned if the current Linux task has received a signal, thus preventing the final migration to the Xenomai domain (i.e. in order to process the signal in the Linux domain). This error should not be considered as fatal.
- -EPERM is returned if the shadow thread has been killed before the current task had a chance to return to the caller. In such a case, the real-time mapping operation has failed globally, and no Xenomai resource remains attached to it.
- -EINVAL is returned if the thread control block does not bear the XNSHADOW bit.
- -EBUSY is returned if either the current Linux task or the associated shadow thread is already involved in a shadow mapping.

**Environments:**

This service can be called from:

- Regular user-space process.

Rescheduling: always.

References [xnheap\\_alloc\(\)](#), [XNMAPPED](#), [XNOTHER](#), [xnpod\\_start\\_thread\(\)](#), [xnpod\\_suspend\\_thread\(\)](#), [XNPRIOSSET](#), [XNRELAX](#), [XNSHADOW](#), and [xnshadow\\_harden\(\)](#).

**4.11.2.3 `xnshadow_ppd_t* xnshadow_ppd_get ( unsigned muxid )`**

Return the per-process data attached to the calling process.

This service returns the per-process data attached to the calling process for the skin whose *muxid* is *muxid*. It must be called with `nklock` locked, `irqs` off.

See `xnshadow_register_interface()` documentation for information on the way to attach a per-process data to a process.

**Parameters**

*muxid* the skin *muxid*.

**Returns**

the per-process data if the current context is a user-space process;  
NULL otherwise.

**4.11.2.4 `void xnshadow_relax ( int notify, int reason )`**

Switch a shadow thread back to the Linux domain.

**For internal use only.**

This service yields the control of the running shadow back to Linux. This is obtained by suspending the shadow and scheduling a wake up call for the mated user task inside the Linux domain. The Linux task will resume on return from `xnpod_suspend_thread()` on behalf of the root thread.

**Parameters**

*notify* A boolean flag indicating whether threads monitored from secondary mode switches should be sent a SIGDEBUG signal. For instance, some internal operations like task exit should not trigger such signal.

*reason* The reason to report along with the SIGDEBUG signal.

Environments:

This service can be called from:

- User-space thread operating in primary (i.e. `harden`) mode.

Rescheduling: always.

**Note**


"current" is valid here since the shadow runs with the properties of the Linux task.

References `XNAFFSET`, `xnpod_suspend_thread()`, `XNPRIOSSET`, `XNRELAX`, `XNROOT`, and `XN-TRAPSW`.

Referenced by `xnpod_trap_fault()`, and `xnshadow_harden()`.

## 4.12 Thread synchronization services.

Collaboration diagram for Thread synchronization services.:



### Files

- file [synch.c](#)  
*Thread synchronization services.*

### Functions

- void [xnsynch\\_init](#) (struct xnsynch \*synch, xnflags\_t flags, xnarch\_atomic\_t \*fastlock)  
*Initialize a synchronization object.*
- xnflags\_t [xnsynch\\_sleep\\_on](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Sleep on an ownerless synchronization object.*
- struct xnthread \* [xnsynch\\_wakeup\\_one\\_sleeper](#) (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*
- struct xnpholder \* [xnsynch\\_wakeup\\_this\\_sleeper](#) (struct xnsynch \*synch, struct xnpholder \*holder)  
*Give the resource ownership to a given waiting thread.*
- xnflags\_t [xnsynch\\_acquire](#) (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Acquire the ownership of a synchronization object.*
- static void [xnsynch\\_clear\\_boost](#) (struct xnsynch \*synch, struct xnthread \*owner)  
*Clear the priority boost.*
- void [xnsynch\\_requeue\\_sleeper](#) (struct xnthread \*thread)  
*Change a sleeper's priority.*
- struct xnthread \* [xnsynch\\_release](#) (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*



- struct xnthread \* [xnsynch\\_peek\\_pending](#) (struct xnsynch \*synch)  
*Access the thread leading a synch object wait queue.*
- int [xnsynch\\_flush](#) (struct xnsynch \*synch, xnflags\_t reason)  
*Unblock all waiters pending on a resource.*
- void [xnsynch\\_forget\\_sleeper](#) (struct xnthread \*thread)  
*Abort a wait for a resource.*
- void [xnsynch\\_release\\_all\\_ownerships](#) (struct xnthread \*thread)  
*Release all ownerships.*

### 4.12.1 Detailed Description

Thread synchronization services.

### 4.12.2 Function Documentation

#### 4.12.2.1 xnflags\_t xnsynch\_acquire ( struct xnsynch \* synch, xnticks\_t timeout, xntmode\_t timeout\_mode )

Acquire the ownership of a synchronization object.

This service should be called by upper interfaces wanting the current thread to acquire the ownership of the given resource. If the resource is already assigned to a thread, the caller is suspended.

This service must be used only with synchronization objects that track ownership (XNSYNCH\_OWNER set).

#### Parameters

**synch** The descriptor address of the synchronization object to acquire.

**timeout** The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given in ticks (see note). It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing XN\_INFINITE and setting *mode* to XN\_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

**timeout\_mode** The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

#### Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task

Rescheduling: possible.

#### Note

The *timeout* value will be interpreted as jiffies if the current thread is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

References XNBOOST, XNBREAK, XNOTHER, XNPEND, `xnpod_suspend_thread()`, XNRMID, XNROBBED, XNTIMEO, and XNWAKEN.

#### 4.12.2.2 void xnsynch\_clear\_boost ( struct xnsynch \* *synch*, struct xnthread \* *owner* ) [static]

Clear the priority boost.

**For internal use only.**

This service is called internally whenever a synchronization object is not claimed anymore by sleepers to reset the object owner's priority to its initial level.

#### Parameters

*synch* The descriptor address of the synchronization object.

*owner* The descriptor address of the thread which currently owns the synchronization object.

#### Note

This routine must be entered `nklock` locked, interrupts off.

References XNBOOST.

Referenced by `xnsynch_flush()`, and `xnsynch_forget_sleeper()`.

#### 4.12.2.3 int xnsynch\_flush ( struct xnsynch \* *synch*, xnflags\_t *reason* )

Unblock all waiters pending on a resource.

This service atomically releases all threads which currently sleep on a given resource.

This service should be called by upper interfaces under circumstances requiring that the pending queue of a given resource is cleared, such as before the resource is deleted.

#### Parameters

*synch* The descriptor address of the synchronization object to be flushed.

*reason* Some flags to set in the information mask of every unblocked thread. Zero is an acceptable value. The following bits are pre-defined by the nucleus:

- XNRMID should be set to indicate that the synchronization object is about to be destroyed (see `xnpod_resume_thread()`).

- XNBREAK should be set to indicate that the wait has been forcibly interrupted (see [xnpod\\_unblock\\_thread\(\)](#)).

### Returns

XNSYNCH\_RESCHEDED is returned if at least one thread is unblocked, which means the caller should invoke [xnpod\\_schedule\(\)](#) for applying the new scheduling state. Otherwise, XNSYNCH\_DONE is returned.

### Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object is no more owned by any thread.

### Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, [xnpod\\_resume\\_thread\(\)](#), and [xnsynch\\_clear\\_boost\(\)](#).

Referenced by [xnregistry\\_put\(\)](#).

#### 4.12.2.4 void xnsynch\_forget\_sleeper ( struct xnthread \* *thread* )

Abort a wait for a resource.

### For internal use only.

Performs all the necessary housekeeping chores to stop a thread from waiting on a given synchronization object.

### Parameters

*thread* The descriptor address of the affected thread.

When the trace support is enabled (i.e. MVM), the idle state is posted to the synchronization object's state diagram (if any) whenever no thread remains blocked on it. The real-time interfaces must ensure that such condition (i.e. EMPTY/IDLE) is mapped to state #0.

### Note

This routine must be entered `nklock` locked, interrupts off.

References XNPEND, and [xnsynch\\_clear\\_boost\(\)](#).

Referenced by [xnpod\\_delete\\_thread\(\)](#), [xnpod\\_resume\\_thread\(\)](#), and [xnpod\\_suspend\\_thread\(\)](#).

#### 4.12.2.5 void xnsynch\_init ( struct xnsynch \* *synch*, xnflags\_t *flags*, xnarch\_atomic\_t \* *fastlock* )

Initialize a synchronization object.

Initializes a new specialized object which can subsequently be used to synchronize real-time activities. The Xenomai nucleus provides a basic synchronization object which can be used to build higher resource objects. Nucleus threads can wait for and signal such objects in order to synchronize their activities.

This object has built-in support for priority inheritance.

##### Parameters

*synch* The address of a synchronization object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

*flags* A set of creation flags affecting the operation. The valid flags are:

- XNSYNCH\_PRIO causes the threads waiting for the resource to pend in priority order. Otherwise, FIFO ordering is used (XNSYNCH\_FIFO).
- XNSYNCH\_OWNER indicates that the synchronization object shall track its owning thread (required if XNSYNCH\_PIP is selected). Note that setting this flag implies the use `xnsynch_acquire` and `xnsynch_release` instead of `xnsynch_sleep_on` and `xnsynch_wakeup_one_sleeper/xnsynch_wakeup_this_sleeper`.
- XNSYNCH\_PIP causes the priority inheritance mechanism to be automatically activated when a priority inversion is detected among threads using this object. Otherwise, no priority inheritance takes place upon priority inversion (XNSYNCH\_NOPIP).
- XNSYNCH\_DREORD (Disable REORdering) tells the nucleus that the wait queue should not be reordered whenever the priority of a blocked thread it holds is changed. If this flag is not specified, changing the priority of a blocked thread using `xnpod_set_thread_schedparam()` will cause this object's wait queue to be reordered according to the new priority level, provided the synchronization object makes the waiters wait by priority order on the awaited resource (XNSYNCH\_PRIO).

##### Parameters

*fastlock* Address of the fast lock word to be associated with the synchronization object. If NULL is passed or XNSYNCH\_OWNER is not set, fast-lock support is disabled.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnregistry_enter()`, and `xnselector_init()`.

#### 4.12.2.6 `struct xnthread* xnsynch_peek_pendq ( struct xnsynch * synch ) [read]`

Access the thread leading a synch object wait queue.

This services returns the descriptor address of to the thread leading a synchronization object wait queue.

##### Parameters

*synch* The descriptor address of the target synchronization object.

##### Returns

The descriptor address of the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.12.2.7 `struct xnthread* xnsynch_release ( struct xnsynch * synch ) [read]`

Give the resource ownership to the next waiting thread.

This service releases the ownership of the given synchronization object. The thread which is currently leading the object's pending list, if any, is unblocked from its pending state. However, no reschedule is performed.

This service must be used only with synchronization objects that track ownership (XNSYNCH\_OWNER set).

##### Parameters

*synch* The descriptor address of the synchronization object whose ownership is changed.

##### Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transfered to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.12.2.8 void xnsynch\_release\_all\_ownerships ( struct xnthread \* *thread* )

Release all ownerships.

**For internal use only.**

This call is used internally to release all the ownerships obtained by a thread on synchronization objects. This routine must be entered interrupts off.

##### Parameters

*thread* The descriptor address of the affected thread.

##### Note

This routine must be entered nklock locked, interrupts off.

Referenced by xnpod\_delete\_thread().

#### 4.12.2.9 void xnsynch\_requeue\_sleeper ( struct xnthread \* *thread* )

Change a sleeper's priority.

**For internal use only.**

This service is used by the PIP code to update the pending priority of a sleeping thread.

##### Parameters

*thread* The descriptor address of the affected thread.

##### Note

This routine must be entered nklock locked, interrupts off.

References XNBOOST.

#### 4.12.2.10 `xnflags_t xnsynch_sleep_on ( struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode )`

Sleep on an ownerless synchronization object.

Makes the calling thread sleep on the specified synchronization object, waiting for it to be signaled.

This service should be called by upper interfaces wanting the current thread to pend on the given resource. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH\_OWNER).

##### Parameters

*synch* The descriptor address of the synchronization object to sleep on.

*timeout* The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given in ticks (see note). It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout\_mode*. Passing XN\_INFINITE and setting *mode* to XN\_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

*timeout\_mode* The mode of the *timeout* parameter. It can either be set to XN\_RELATIVE, XN\_ABSOLUTE, or XN\_REALTIME (see also [xntimer\\_start\(\)](#)).

##### Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always.

##### Note

The *timeout* value will be interpreted as jiffies if the current thread is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

References XNBREAK, XNPEND, `xnpod_suspend_thread()`, XNRMID, and XNTIMEO.

Referenced by `xnregistry_bind()`, `xnregistry_remove_safe()`, and `xnselect()`.

#### 4.12.2.11 `struct xnthread* xnsynch_wakeup_one_sleeper ( struct xnsynch * synch ) [read]`

Give the resource ownership to the next waiting thread.

This service wakes up the thread which is currently leading the synchronization object's pending list. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a single waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH\_OWNER not set).

**Parameters**

*synch* The descriptor address of the synchronization object whose ownership is changed.

**Returns**

The descriptor address of the unblocked thread.

**Side-effects:**

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

**Environments:**

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod\_resume\_thread().

#### 4.12.2.12 **struct xnpholder\* xnsynch\_wakeup\_this\_sleeper ( struct xnsynch \* *synch*, struct xnpholder \* *holder* ) [read]**

Give the resource ownership to a given waiting thread.

This service wakes up a specific thread which is currently pending on the given synchronization object. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a specific waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH\_OWNER not set).

**Parameters**

*synch* The descriptor address of the synchronization object whose ownership is changed.

*holder* The link holder address of the thread to unblock (&thread->plink) which MUST be currently linked to the synchronization object's pending queue (i.e. synch->pendq).

**Returns**

The link address of the unblocked thread in the synchronization object's pending queue.

**Side-effects:**

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.



- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod\_resume\_thread().

## 4.13 Time base services.

Collaboration diagram for Time base services.:



### Files

- file [timebase.h](#)
- file [timebase.c](#)

### Functions

- int [xntbase\\_alloc](#) (const char \*name, u\_long period, u\_long flags, xntbase\_t \*\*basep)  
*Allocate a time base.*
- void [xntbase\\_free](#) (xntbase\_t \*base)  
*Free a time base.*
- int [xntbase\\_update](#) (xntbase\_t \*base, u\_long period)  
*Change the period of a time base.*
- int [xntbase\\_switch](#) (const char \*name, u\_long period, xntbase\_t \*\*basep)  
*Replace a time base.*

- void [xntbase\\_start](#) (xntbase\_t \*base)  
*Start a time base.*
- void [xntbase\\_stop](#) (xntbase\_t \*base)  
*Stop a time base.*
- void [xntbase\\_tick](#) (xntbase\_t \*base)  
*Announce a clock tick to a time base.*
- xnticks\_t [xntbase\\_convert](#) (xntbase\_t \*srcbase, xnticks\_t ticks, xntbase\_t \*dstbase)  
*Convert a clock value into another time base.*
- static xnticks\_t [xntbase\\_get\\_time](#) (xntbase\_t \*base)  
*Get the clock time for a given time base.*
- void [xntbase\\_adjust\\_time](#) (xntbase\_t \*base, xnticks\_t delta)  
*Adjust the clock time for the system.*

### 4.13.1 Detailed Description

Xenomai implements the notion of time base, by which software timers that belong to different skins may be clocked separately according to distinct frequencies, or aperiodically. In the periodic case, delays and timeouts are given in counts of ticks; the duration of a tick is specified by the time base. In the aperiodic case, timings are directly specified in nanoseconds.

Only a single aperiodic (i.e. tick-less) time base may exist in the system, and the nucleus provides for it through the `nktbase` object. All skins depending on aperiodic timings should bind to the latter (see [xntbase\\_alloc\(\)](#)), also known as the master time base.

Skins depending on periodic timings may create and bind to their own time base. Such a periodic time base is managed as a timed slave object of the master time base. A cascading software timer fired by the master time base according to the appropriate frequency, triggers in turn the update process of the associated timed slave, which eventually fires the elapsed software timers controlled by the periodic time base. In other words, Xenomai emulates periodic timing over an aperiodic policy.

Xenomai always controls the underlying timer hardware in a tick-less fashion, also known as the oneshot mode.

### 4.13.2 Function Documentation

#### 4.13.2.1 void xntbase\_adjust\_time ( xntbase\_t \* base, xnticks\_t delta )

Adjust the clock time for the system.

Xenomai tracks the current time as a monotonously increasing count of ticks since the epoch. The epoch is initially the same as the underlying machine time, and it is always synchronised across all active time bases.

This service changes the epoch for the system by applying the specified tick delta on the master's wallclock offset and resynchronizing all other time bases.

**Parameters**

*base* The address of the initiating time base.

*delta* The adjustment of the system time expressed in ticks of the specified time base.

**Note**

This routine must be entered nlock locked, interrupts off.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.13.2.2 `int xntbase_alloc ( const char * name, u_long period, u_long flags, xntbase_t ** basep )`

Allocate a time base.

A time base is an abstraction used to provide private clocking information to real-time skins, by which they may operate either in aperiodic or periodic mode, possibly according to distinct clock frequencies in the latter case. This abstraction is required in order to support several RTOS emulators running concurrently, which may exhibit different clocking policies and/or period.

Once allocated, a time base may be attached to all software timers created directly or indirectly by a given skin, and influences all timed services accordingly.

The `xntbase_alloc()` service allocates a new time base to the caller, and returns the address of its descriptor. The new time base is left in a disabled state (unless *period* equals `XN_APERIODIC_TICK`), calling `xntbase_start()` is needed to enable it.

**Parameters**

*name* The symbolic name of the new time base. This information is used to report status information when reading from `/proc/xenomai/timebases`; it has currently no other usage.

*period* The duration of the clock tick for the new time base, given as a count of nanoseconds. The special `XN_APERIODIC_TICK` value may be used to retrieve the master - aperiodic - time base, which is always up and running when a real-time skin has called the `xnpod_init()` service. All other values are meant to define the clock rate of a periodic time base. For instance, passing 1000000 (ns) in the *period* parameter will create a periodic time base clocked at a frequency of 1Khz.

*flags* A bitmask composed as follows:

- `XNTBISO` causes the target timebase to be isolated from global wallclock offset updates as performed by `xntbase_adjust_time()`.

**Parameters**

*basep* A pointer to a memory location which will be written upon success with the address of the allocated time base. If *period* equals `XN_APERIODIC_TICK`, the address of the built-in master time base descriptor will be copied back to this location.

**Returns**

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if no system memory is available to allocate a new time base descriptor.

Environments:

This service can be called from:

- Kernel module initialization code
- User-space task in secondary mode

Rescheduling: never.

**Note**

Any periodic time base allocated by a real-time skin must be released by a call to `xntbase_free()` before the kernel module implementing the skin may be unloaded.

Referenced by `xntbase_switch()`.

#### 4.13.2.3 `xnticks_t xntbase_convert ( xntbase_t * srcbase, xnticks_t ticks, xntbase_t * dstbase )`

Convert a clock value into another time base.

**Parameters**

*srcbase* The descriptor address of the source time base.

*ticks* The clock value expressed in the source time base to convert to the destination time base.

*dstbase* The descriptor address of the destination time base.

**Returns**

The converted count of ticks in the destination time base is returned.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Rescheduling: never.

#### 4.13.2.4 void xntbase\_free ( xntbase\_t \* *base* )

Free a time base.

This service disarms all outstanding timers from the affected periodic time base, destroys the aperiodic cascading timer, then releases the time base descriptor.

##### Parameters

*base* The address of the time base descriptor to release.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

##### Note

Requests to free the master time base are silently caught and discarded; in such a case, outstanding aperiodic timers are left untouched.

Referenced by xntbase\_switch().

#### 4.13.2.5 xnticks\_t xntbase\_get\_time ( xntbase\_t \* *base* ) [inline, static]

Get the clock time for a given time base.

This service returns the (external) clock time as maintained by the specified time base. This value is adjusted with the wallclock offset as defined by [xntbase\\_adjust\\_time\(\)](#).

##### Parameters

*base* The address of the time base to query.

##### Returns

The current time (in jiffies) if the specified time base runs in periodic mode, or the machine time (converted to nanoseconds) as maintained by the hardware if *base* refers to the master time base.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by xnregistry\_bind().

#### 4.13.2.6 void xntbase\_start ( xntbase\_t \* base )

Start a time base.

This service enables a time base, using a cascading timer running in the master time base as the source of periodic clock ticks. The time base is synchronised on the Xenomai system clock. Timers attached to the started time base are immediated armed.

##### Parameters

*base* The address of the time base descriptor to start.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Rescheduling: never.

##### Note

Requests to enable the master time base are silently caught and discarded; only the internal service [xnpod\\_enable\\_timesource\(\)](#) is allowed to start the latter. The master time base remains enabled until no real-time skin remains attached to the nucleus.

Referenced by [xntbase\\_switch\(\)](#).

#### 4.13.2.7 void xntbase\_stop ( xntbase\_t \* base )

Stop a time base.

This service disables a time base, stopping the cascading timer running in the master time base which is used to clock it. Outstanding timers attached to the stopped time base are immediated disarmed.

Stopping a time base also invalidates its clock setting.

##### Parameters

*base* The address of the time base descriptor to stop.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

##### Note

Requests to disable the master time base are silently caught and discarded; only the internal service [xnpod\\_disable\\_timesource\(\)](#) is allowed to stop the latter. The master time base remains enabled until no real-time skin remains attached to the nucleus.

#### 4.13.2.8 `int xntbase_switch ( const char * name, u_long period, xntbase_t ** basep )`

Replace a time base.

This service is useful for switching the current time base of a real-time skin between aperiodic and periodic modes, by providing a new time base descriptor as needed. The original time base descriptor is freed as a result of this operation (unless it refers to the master time base). The new time base is automatically started by a call to [xntbase\\_start\(\)](#) if the original time base was enabled at the time of the call, or left in a disabled state otherwise.

This call handles all mode transitions and configuration changes carefully, i.e. periodic <-> periodic, aperiodic <-> aperiodic, periodic <-> aperiodic.

##### Parameters

*name* The symbolic name of the new time base. This information is used to report status information when reading from `/proc/xenomai/timebases`; it has currently no other usage.

*period* The duration of the clock tick for the time base, given as a count of nanoseconds. This value is meant to define the new clock rate of the new periodic time base (i.e.  $1e9 / \text{period}$ ).

*basep* A pointer to a memory location which will be first read to pick the address of the original time base to be replaced, then written back upon success with the address of the new time base. A null pointer is allowed on input in *basep*, in which case the new time base will be created as if [xntbase\\_alloc\(\)](#) had been called directly.

##### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if no system memory is available to allocate a new time base descriptor.

Environments:

This service can be called from:

- Kernel module initialization code
- User-space task in secondary mode

Rescheduling: never.

References [xntbase\\_alloc\(\)](#), [xntbase\\_free\(\)](#), [xntbase\\_start\(\)](#), and [xntbase\\_update\(\)](#).

#### 4.13.2.9 `void xntbase_tick ( xntbase_t * base )`

Announce a clock tick to a time base.

This service announces a new clock tick to a time base. Normally, only specialized nucleus code would announce clock ticks. However, under certain circumstances, it may be useful to allow client code to send such notifications on their own.

Notifying a clock tick to a time base causes the timer management code to check for outstanding timers, which may in turn fire off elapsed timeout handlers. Additionally, periodic time bases (i.e. all but the master time base) would also update their count of elapsed jiffies, in case the current processor has been defined as the internal time keeper (i.e. `CPU# == XNTIMER_KEEPER_ID`).

### Parameters

*base* The address of the time base descriptor to announce a tick to.

Environments:

This service can be called from:

- Interrupt context only.

Rescheduling: never.

References `xntimer_tick_aperiodic()`.

#### 4.13.2.10 `int xntbase_update ( xntbase_t * base, u_long period )`

Change the period of a time base.

### Parameters

*base* The address of the time base descriptor to update.

*period* The duration of the clock tick for the time base, given as a count of nanoseconds. This value is meant to define the new clock rate of the affected periodic time base (i.e.  $1e9 / \text{period}$ ).

### Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if an attempt is made to set a null period.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Rescheduling: never.

### Note

Requests to update the master time base are silently caught and discarded. The master time base has a fixed aperiodic policy which may not be changed.

Referenced by `xntbase_switch()`.



## 4.14 Timer services.

Collaboration diagram for Timer services.:



### Files

- file [timer.h](#)
- file [timer.c](#)

### Functions

- static int [xntimer\\_start](#) (xntimer\_t \*timer, xnticks\_t value, xnticks\_t interval, xntmode\_t mode)  
*Arm a timer.*
- static void [xntimer\\_stop](#) (xntimer\_t \*timer)  
*Disarm a timer.*
- static xnticks\_t [xntimer\\_get\\_date](#) (xntimer\_t \*timer)  
*Return the absolute expiration date.*
- static xnticks\_t [xntimer\\_get\\_timeout](#) (xntimer\_t \*timer)  
*Return the relative expiration date.*
- static xnticks\_t [xntimer\\_get\\_interval](#) (xntimer\_t \*timer)  
*Return the timer interval value.*
- void [xntimer\\_tick\\_aperiodic](#) (void)  
*Process a timer tick for the aperiodic master time base.*
- void [xntimer\\_tick\\_periodic](#) (xntimer\_t \*mtimer)  
*Process a timer tick for a slave periodic time base.*
- void [xntimer\\_init](#) (xntimer\_t \*timer, xntbase\_t \*base, void(\*handler)(xntimer\_t \*timer))  
*Initialize a timer object.*
- void [xntimer\\_destroy](#) (xntimer\_t \*timer)  
*Release a timer object.*

- unsigned long `xntimer_get_overruns` (`xntimer_t *timer`, `xnticks_t now`)

*Get the count of overruns for the last tick.*

- void `xntimer_freeze` (void)

*Freeze all timers (from every time bases).*

### 4.14.1 Detailed Description

The Xenomai timer facility always operate the timer hardware in oneshot mode, regardless of the time base in effect. Periodic timing is obtained through a software emulation, using cascading timers.

Depending on the time base used, the timer object stores time values either as count of jiffies (periodic), or as count of CPU ticks (aperiodic).

### 4.14.2 Function Documentation

#### 4.14.2.1 void `xntimer_destroy` ( `xntimer_t * timer` )

Release a timer object.

Destroys a timer. After it has been destroyed, all resources associated with the timer have been released. The timer is automatically deactivated before deletion if active on entry.

#### Parameters

*timer* The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References `xntimer_stop`().

Referenced by `xnpod_delete_thread`(), and `xnpod_shutdown`().

#### 4.14.2.2 void `xntimer_freeze` ( void )

Freeze all timers (from every time bases).

**For internal use only.**

This routine deactivates all active timers atomically.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_disable_timesource()`.

**4.14.2.3 `xnticks_t xntimer_get_date ( xntimer_t * timer ) [inline, static]`**

Return the absolute expiration date.

Return the next expiration date of a timer in absolute clock ticks (see note).

**Parameters**

*timer* The address of a valid timer descriptor.

**Returns**

The expiration date converted to the current time unit. The special value `XN_INFINITE` is returned if *timer* is currently inactive.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

**Note**

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

#### 4.14.2.4 `xnticks_t xntimer_get_interval ( xntimer_t * timer ) [inline, static]`

Return the timer interval value.

Return the timer interval value in clock ticks (see note).

##### Parameters

*timer* The address of a valid timer descriptor.

##### Returns

The expiration date converted to the current time unit. The special value `XN_INFINITE` is returned if *timer* is currently inactive or aperiodic.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

##### Note

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

#### 4.14.2.5 `unsigned long xntimer_get_overruns ( xntimer_t * timer, xnticks_t now )`

Get the count of overruns for the last tick.

This service returns the count of pending overruns for the last tick of a given timer, as measured by the difference between the expected expiry date of the timer and the date *now* passed as argument.

##### Parameters

*timer* The address of a valid timer descriptor.

*now* current date (in the monotonic time base)

##### Returns

the number of overruns of *timer* at date *now*

Referenced by `xnpod_wait_thread_period()`.

**4.14.2.6 xnticks\_t xntimer\_get\_timeout ( xntimer\_t \* timer ) [inline, static]**

Return the relative expiration date.

Return the next expiration date of a timer in relative clock ticks (see note).

**Parameters**

*timer* The address of a valid timer descriptor.

**Returns**

The expiration date converted to the current time unit. The special value XN\_INFINITE is returned if *timer* is currently inactive. In oneshot mode, it might happen that the timer has already expired when this service is run (even if the associated handler has not been fired yet); in such a case, 1 is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

**Note**

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

**4.14.2.7 void xntimer\_init ( xntimer\_t \* timer, xntbase\_t \* base, void(\*) (xntimer\_t \*timer) handler )**

Initialize a timer object.

Creates a timer. When created, a timer is left disarmed; it must be started using [xntimer\\_start\(\)](#) in order to be activated.

**Parameters**

*timer* The address of a timer descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

*base* The descriptor address of the time base the new timer depends on. See [xntbase\\_alloc\(\)](#) for detailed explanations about time bases.

*handler* The routine to call upon expiration of the timer.

There is no limitation on the number of timers which can be created/active concurrently.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

**4.14.2.8** `void xntimer_start ( xntimer_t * timer, xnticks_t value, xnticks_t interval, xntmode_t mode ) [inline, static]`

Arm a timer.

Activates a timer so that the associated timeout handler will be fired after each expiration time. A timer can be either periodic or single-shot, depending on the reload value passed to this routine. The given timer must have been previously initialized, and will be clocked according to the policy defined by the time base specified in `xntimer_init()`.

#### Parameters

*timer* The address of a valid timer descriptor.

*value* The date of the initial timer shot, expressed in clock ticks (see note).

*interval* The reload value of the timer. It is a periodic interval value to be used for reprogramming the next timer shot, expressed in clock ticks (see note). If *interval* is equal to `XN_INFINITE`, the timer will not be reloaded after it has expired.

*mode* The timer mode. It can be `XN_RELATIVE` if *value* shall be interpreted as a relative date, `XN_ABSOLUTE` for an absolute date based on the monotonic clock of the related time base (as returned by `xntbase_get_jiffies()`), or `XN_REALTIME` if the absolute date is based on the adjustable real-time clock of the time base (as returned by `xntbase_get_time()`).

#### Returns

0 is returned upon success, or `-ETIMEDOUT` if an absolute date in the past has been given.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### Note

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.  
Must be called with `nklock` held, IRQs off.

Referenced by `xnpod_enable_timesource()`, `xnpod_set_thread_periodic()`, `xnpod_set_thread_tslice()`, and `xnpod_suspend_thread()`.

#### 4.14.2.9 `int xntimer_stop ( xntimer_t * timer ) [inline, static]`

Disarm a timer.

This service deactivates a timer previously armed using [xntimer\\_start\(\)](#). Once disarmed, the timer can be subsequently re-armed using the latter service.

#### Parameters

*timer* The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

#### Note

Must be called with `nklock` held, IRQs off.

Referenced by `xnpod_resume_thread()`, `xnpod_set_thread_periodic()`, `xnpod_set_thread_tslice()`, and `xntimer_destroy()`.

#### 4.14.2.10 `void xntimer_tick_aperiodic ( void )`

Process a timer tick for the aperiodic master time base.

#### For internal use only.

This routine informs all active timers that the clock has been updated by processing the outstanding timer list. Elapsed timer actions will be fired.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

References xnsched::htimer, xnsched::lflags, and xnsched::status.

Referenced by xntbase\_tick().

#### 4.14.2.11 void xntimer\_tick\_periodic ( xntimer\_t \* mtimer )

Process a timer tick for a slave periodic time base.

**For internal use only.**

The periodic timer tick is cascaded from a software timer managed from the master aperiodic time base; in other words, periodic timing is emulated by software timers running in aperiodic timing mode. There may be several concurrent periodic time bases (albeit a single aperiodic time base - i.e. the master one called "nktbase" - may exist at any point in time).

This routine informs all active timers that the clock has been updated by processing the timer wheel. Elapsed timer actions will be fired.

#### Parameters

*mtimer* The address of the cascading timer running in the master time base which announced the tick.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

#### Note

Only active timers are inserted into the timer wheel.

## 4.15 Virtual file services

Collaboration diagram for Virtual file services:





## Data Structures

- struct [xnvmfile\\_lock\\_ops](#)  
*Vfile locking operations.*
- struct [xnvmfile\\_regular\\_ops](#)  
*Regular vfile operation descriptor.*
- struct [xnvmfile\\_regular\\_iterator](#)  
*Regular vfile iterator.*
- struct [xnvmfile\\_snapshot\\_ops](#)  
*Snapshot vfile operation descriptor.*
- struct [xnvmfile\\_rev\\_tag](#)  
*Snapshot revision tag.*
- struct [xnvmfile\\_snapshot](#)  
*Snapshot vfile descriptor.*
- struct [xnvmfile\\_snapshot\\_iterator](#)  
*Snapshot-driven vfile iterator.*

## Files

- file [vfile.h](#)  
*This file is part of the Xenomai project.*

## Functions

- int [xnvmfile\\_init\\_snapshot](#) (const char \*name, struct [xnvmfile\\_snapshot](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a snapshot-driven vfile.*
- int [xnvmfile\\_init\\_regular](#) (const char \*name, struct [xnvmfile\\_regular](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a regular vfile.*
- int [xnvmfile\\_init\\_dir](#) (const char \*name, struct [xnvmfile\\_directory](#) \*vdir, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual directory entry.*
- int [xnvmfile\\_init\\_link](#) (const char \*from, const char \*to, struct [xnvmfile\\_link](#) \*vlink, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual link entry.*
- void [xnvmfile\\_destroy](#) (struct [xnvmfile](#) \*vfile)

*Removes a virtual file entry.*

- ssize\_t [xnvmfile\\_get\\_blob](#) (struct xnvmfile\_input \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_string](#) (struct xnvmfile\_input \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_integer](#) (struct xnvmfile\_input \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct xnvmfile\_directory [nkvfroot](#)  
*Xenomai vfile root directory.*
- struct xnvmfile\_directory [nkvfroot](#)  
*Xenomai vfile root directory.*

### 4.15.1 Detailed Description

Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces. This encapsulation is aimed at:

- supporting consistent collection of very large record-based output, without incurring latency peaks for undergoing real-time activities.
- in the future, hiding discrepancies between linux kernel releases, regarding the proper way to export kernel object states to userland, either via the /proc interface or by any other mean.

This virtual file implementation offers record-based read support based on seq\_files, single-buffer write support, directory and link handling, all visible from the /proc namespace.

The vfile support exposes four filesystem object types:

- snapshot-driven file (struct [xnvmfile\\_snapshot](#)). This is commonly used to export real-time object states via the /proc filesystem. To minimize the latency involved in protecting the vfile routines from changes applied by real-time code on such objects, a snapshot of the data to output is first taken under proper locking, before the collected data is formatted and sent out in a lockless manner.

Because a large number of records may have to be output, the data collection phase is not strictly atomic as a whole, but only protected at record level. The vfile implementation can be notified of updates to the underlying data set, and restart the collection from scratch until the snapshot is fully consistent.

- regular sequential file (struct [xnvmfile\\_regular](#)). This is basically an encapsulated sequential file object as available from the host kernel (i.e. seq\_file), with a few additional features to make it more handy in a Xenomai environment, like implicit locking support and shortened declaration for simplest, single-record output.

- virtual link (struct `xnvfile_link`). This is a symbolic link feature integrated with the vfile semantics. The link target is computed dynamically at creation time from a user-given helper routine.
- virtual directory (struct `xnvfile_directory`). A directory object, which can be used to create a hierarchy for ordering a set of vfile objects.

## 4.15.2 Function Documentation

### 4.15.2.1 void `xnvfile_destroy` ( struct `xnvfile` \* *vfile* )

Removes a virtual file entry.

#### Parameters

*vfile* A pointer to the virtual file descriptor to remove.

References `nkvfroot`.

### 4.15.2.2 ssize\_t `xnvfile_get_blob` ( struct `xnvfile_input` \* *input*, void \* *data*, size\_t *size* )

Read in a data bulk written to the vfile.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_blob()` retrieves this data as an untyped binary blob, and copies it back to the caller's buffer.

#### Parameters

*input* A pointer to the input descriptor passed to the `store()` handler.

*data* The address of the destination buffer to copy the input data to.

*size* The maximum number of bytes to copy to the destination buffer. If *size* is larger than the actual data size, the input is truncated to *size*.

#### Returns

The number of bytes read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- `-EFAULT` indicates an invalid source buffer address.

Referenced by `xnvfile_get_integer()`, and `xnvfile_get_string()`.

### 4.15.2.3 ssize\_t `xnvfile_get_integer` ( struct `xnvfile_input` \* *input*, long \* *valp* )

Evaluate the string written to the vfile as a long integer.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_integer()` retrieves and interprets this data as a long integer, and copies the resulting value back to *valp*.

The long integer can be expressed in decimal, octal or hexadecimal bases depending on the prefix found.

**Parameters**

*input* A pointer to the input descriptor passed to the store() handler.

*valp* The address of a long integer variable to receive the value.

**Returns**

The number of characters read while evaluating the input as a long integer upon success. Otherwise, a negative error code is returned:

- -EINVAL indicates a parse error on the input stream; the written text cannot be evaluated as a long integer.
- -EFAULT indicates an invalid source buffer address.

References xnvfile\_get\_blob().

#### 4.15.2.4 ssize\_t xnvfile\_get\_string ( struct xnvfile\_input \* *input*, char \* *s*, size\_t *maxlen* )

Read in a C-string written to the vfile.

When writing to a vfile, the associated store() handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. [xnvfile\\_get\\_string\(\)](#) retrieves this data as a null-terminated character string, and copies it back to the caller's buffer.

**Parameters**

*input* A pointer to the input descriptor passed to the store() handler.

*s* The address of the destination string buffer to copy the input data to.

*maxlen* The maximum number of bytes to copy to the destination buffer, including the ending null character. If *maxlen* is larger than the actual string length, the input is truncated to *maxlen*.

**Returns**

The number of characters read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- -EFAULT indicates an invalid source buffer address.

References xnvfile\_get\_blob().

#### 4.15.2.5 int xnvfile\_init\_dir ( const char \* *name*, struct xnvfile\_directory \* *vdir*, struct xnvfile\_directory \* *parent* )

Initialize a virtual directory entry.

**Parameters**

*name* The name which should appear in the pseudo-filesystem, identifying the vdir entry.

*vdir* A pointer to the virtual directory descriptor to initialize.

*parent* A pointer to a virtual directory descriptor standing for the parent directory of the new vdir. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfsroot* vdir.

### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual directory entry cannot be created in the /proc hierarchy.

#### 4.15.2.6 `int xnvmfile_init_link ( const char * from, const char * to, struct xnvmfile_link * vlink, struct xnvmfile_directory * parent )`

Initialize a virtual link entry.

### Parameters

*from* The name which should appear in the pseudo-filesystem, identifying the vlink entry.

*to* The target file name which should be referred to symbolically by *name*.

*vlink* A pointer to the virtual link descriptor to initialize.

*parent* A pointer to a virtual directory descriptor standing for the parent directory of the new vlink. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfsroot* vdir.

### Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual link entry cannot be created in the /proc hierarchy.

#### 4.15.2.7 `int xnvmfile_init_regular ( const char * name, struct xnvmfile_regular * vfile, struct xnvmfile_directory * parent )`

Initialize a regular vfile.

### Parameters

*name* The name which should appear in the pseudo-filesystem, identifying the vfile entry.

*vfile* A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- .privsz is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- entry.lockops is a pointer to a [locking](#) descriptor", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case no locking will be applied.
- .ops is a pointer to an [operation descriptor](#).

**Parameters**

*parent* A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvrroot* vdir.

**Returns**

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

#### 4.15.2.8 `int xnvfile_init_snapshot ( const char * name, struct xnvfile_snapshot * vfile, struct xnvfile_directory * parent )`

Initialize a snapshot-driven vfile.

**Parameters**

*name* The name which should appear in the pseudo-filesystem, identifying the vfile entry.

*vfile* A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- .privsz is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- .datasz is the size (in bytes) of a single record to be collected by the [next\(\) handler](#) from the [operation descriptor](#).
- .tag is a pointer to a mandatory vfile revision tag structure (struct [xnvfile\\_rev\\_tag](#)). This tag will be monitored for changes by the vfile core while collecting data to output, so that any update detected will cause the current snapshot data to be dropped, and the collection to restart from the beginning. To this end, any change to the data which may be part of the collected records, should also invoke `xnvfile_touch()` on the associated tag.
- entry.lockops is a pointer to a [locking descriptor](#)", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case the operations on the nucleus lock (i.e. `nklock`) will be used internally around calls to data collection handlers (see [operation descriptor](#)).
- .ops is a pointer to an [operation descriptor](#).

**Parameters**

*parent* A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvrroot* vdir.

**Returns**

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

References `xnvfile_snapshot_ops::store`.

### 4.15.3 Variable Documentation

#### 4.15.3.1 struct xnvfile\_directory nkvfroot

Xenomai vfile root directory.

This vdir maps the /proc/xenomai directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by xnvfile\_destroy().

#### 4.15.3.2 struct xnvfile\_directory nkvfroot

Xenomai vfile root directory.

This vdir maps the /proc/xenomai directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by xnvfile\_destroy().

## 4.16 HAL.

Generic Adeos-based hardware abstraction layer.

### Files

- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for ARM.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.*
- file [hal.c](#)  
*Generic Real-Time HAL.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for PowerPC.*
- file [hal.c](#)  
*Adeos-based Real-Time Abstraction Layer for the SuperH architecture.*
- file [hal-common.c](#)  
*Adeos-based Real-Time Abstraction Layer for x86.*
- file [hal\\_32.c](#)  
*Adeos-based Real-Time Abstraction Layer for x86.*

- file [hal\\_64.c](#)  
*Adeos-based Real-Time Abstraction Layer for x86\_64.*
- file [smi.c](#)  
*SMI workaround for x86.*

## Functions

- int [rthal\\_timer\\_request](#) (void(\*tick\_handler)(void), void(\*mode\_emul)(enum clock\_event\_mode mode, struct clock\_event\_device \*cdev), int(\*tick\_emul)(unsigned long delay, struct clock\_event\_device \*cdev), int cpu)  
*Grab the hardware timer.*
- void [rthal\\_timer\\_release](#) (int cpu)  
*Release the hardware timer.*
- int [rthal\\_irq\\_host\\_request](#) (unsigned irq, rthal\_irq\_host\_handler\_t handler, char \*name, void \*dev\_id)  
*Install a shared Linux interrupt handler.*
- int [rthal\\_irq\\_host\\_release](#) (unsigned irq, void \*dev\_id)  
*Uninstall a shared Linux interrupt handler.*
- int [rthal\\_irq\\_enable](#) (unsigned irq)  
*Enable an interrupt source.*
- int [rthal\\_irq\\_disable](#) (unsigned irq)  
*Disable an interrupt source.*
- int [rthal\\_irq\\_request](#) (unsigned irq, rthal\_irq\_handler\_t handler, rthal\_irq\_ackfn\_t ackfn, void \*cookie)  
*Install a real-time interrupt handler.*
- int [rthal\\_irq\\_release](#) (unsigned irq)  
*Uninstall a real-time interrupt handler.*
- rthal\_trap\_handler\_t [rthal\\_trap\\_catch](#) (rthal\_trap\_handler\_t handler)  
*Installs a fault handler.*
- int [rthal\\_apc\\_alloc](#) (const char \*name, void(\*handler)(void \*cookie), void \*cookie)  
*Allocate an APC slot.*
- void [rthal\\_apc\\_free](#) (int apc)  
*Releases an APC slot.*



### 4.16.1 Detailed Description

Generic Adeos-based hardware abstraction layer. x86\_64-specific HAL services.

i386-specific HAL services.

SuperH-specific HAL services.

PowerPC-specific HAL services.

NIOS2-specific HAL services.

Blackfin-specific HAL services.

ARM-specific HAL services.

### 4.16.2 Function Documentation

**4.16.2.1** `int rthal_apc_alloc ( const char * name, void(*)(void *cookie) handler, void * cookie )`

Allocate an APC slot.

APC is the acronym for Asynchronous Procedure Call, a mean by which activities from the Xenomai domain can schedule deferred invocations of handlers to be run into the Linux domain, as soon as possible when the Linux kernel gets back in control. Up to BITS\_PER\_LONG APC slots can be active at any point in time. APC support is built upon Adeos's virtual interrupt support.

The HAL guarantees that any Linux kernel service which would be callable from a regular Linux interrupt handler is also available to APC handlers.

#### Parameters

*name* is a symbolic name identifying the APC which will get reported through the /proc/xenomai/apc interface. Passing NULL to create an anonymous APC is allowed.

*handler* The address of the fault handler to call upon exception condition. The handle will be passed the *cookie* value unmodified.

*cookie* A user-defined opaque cookie the HAL will pass to the APC handler as its sole argument.

#### Returns

an valid APC id. is returned upon success, or a negative error code otherwise:

- -EINVAL is returned if *handler* is invalid.
- -EBUSY is returned if no more APC slots are available.

Environments:

This service can be called from:

- Linux domain context.

#### 4.16.2.2 `int rthal_apc_free ( int apc )`

Releases an APC slot.

This service deallocates an APC slot obtained by [rthal\\_apc\\_alloc\(\)](#).

##### Parameters

*apc* The APC id. to release, as returned by a successful call to the [rthal\\_apc\\_alloc\(\)](#) service.

Environments:

This service can be called from:

- Any domain context.

#### 4.16.2.3 `int rthal_irq_disable ( unsigned irq )`

Disable an interrupt source.

Disables an interrupt source at PIC level. After this call has returned, no more IRQs from the given source will be allowed, until the latter is enabled again using [rthal\\_irq\\_enable\(\)](#).

##### Parameters

*irq* The interrupt source to disable. This value is architecture-dependent.

##### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

#### 4.16.2.4 `int rthal_irq_enable ( unsigned irq )`

Enable an interrupt source.

Enables an interrupt source at PIC level. Since Adeos masks and acknowledges the associated interrupt source upon IRQ receipt, this action is usually needed whenever the HAL handler does not propagate the IRQ event to the Linux domain, thus preventing the regular Linux interrupt handling code from re-enabling said source. After this call has returned, IRQs from the given source will be enabled again.

**Parameters**

*irq* The interrupt source to enable. This value is architecture-dependent.

**Returns**

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

**4.16.2.5 int rthal\_irq\_host\_release ( unsigned irq, void \* dev\_id )**

Uninstall a shared Linux interrupt handler.

Uninstalls a shared interrupt handler from the Linux domain for the given interrupt source. The handler is removed from the existing list of Linux handlers for this interrupt source.

**Parameters**

*irq* The interrupt source to detach the shared handler from. This value is architecture-dependent.

*dev\_id* is a valid device id, identical in essence to the one requested by the `free_irq()` service provided by the Linux kernel. This value will be used to locate the handler to remove from the chain of existing Linux handlers for the given interrupt source. This parameter must match the device id. passed to [rthal\\_irq\\_host\\_request\(\)](#) for the same handler instance.

**Returns**

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.

Environments:

This service can be called from:

- Linux domain context.

**4.16.2.6 int rthal\_irq\_host\_request ( unsigned irq, rthal\_irq\_host\_handler\_t handler, char \* name, void \* dev\_id )**

Install a shared Linux interrupt handler.

Installs a shared interrupt handler in the Linux domain for the given interrupt source. The handler is appended to the existing list of Linux handlers for this interrupt source.

### Parameters

- irq* The interrupt source to attach the shared handler to. This value is architecture-dependent.
- handler* The address of a valid interrupt service routine. This handler will be called each time the corresponding IRQ is delivered, as part of the chain of existing regular Linux handlers for this interrupt source. The handler prototype is the same as the one required by the `request_irq()` service provided by the Linux kernel.
- name* is a symbolic name identifying the handler which will get reported through the `/proc/interrupts` interface.
- dev\_id* is a unique device id, identical in essence to the one requested by the `request_irq()` service.

### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid or *handler* is NULL.

Environments:

This service can be called from:

- Linux domain context.

#### 4.16.2.7 `int rthal_irq_release ( unsigned irq )`

Uninstall a real-time interrupt handler.

Uninstalls an interrupt handler previously attached using the [rthal\\_irq\\_request\(\)](#) service.

### Parameters

- irq* The hardware interrupt channel to uninstall a handler from. This value is architecture-dependent.

### Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

Referenced by `rthal_timer_release()`.

#### 4.16.2.8 `int rthal_irq_request ( unsigned irq, rthal_irq_handler_t handler, rthal_irq_ackfn_t ackfn, void * cookie )`

Install a real-time interrupt handler.

Installs an interrupt handler for the specified IRQ line by requesting the appropriate Adeos virtualization service. The handler is invoked by Adeos on behalf of the Xenomai domain context. Once installed, the HAL interrupt handler will be called prior to the regular Linux handler for the same interrupt source.

##### Parameters

- irq* The hardware interrupt channel to install a handler on. This value is architecture-dependent.
- handler* The address of a valid interrupt service routine. This handler will be called each time the corresponding IRQ is delivered, and will be passed the *cookie* value unmodified.
- ackfn* The address of an optional interrupt acknowledge routine, aimed at replacing the one provided by Adeos. Only very specific situations actually require to override the default Adeos setting for this parameter, like having to acknowledge non-standard PIC hardware. If *ackfn* is NULL, the default Adeos routine will be used instead.
- cookie* A user-defined opaque cookie the HAL will pass to the interrupt handler as its sole argument.

##### Returns

0 is returned upon success. Otherwise:

- -EBUSY is returned if an interrupt handler is already installed. [rthal\\_irq\\_release\(\)](#) must be issued first before a handler is installed anew.
- -EINVAL is returned if *irq* is invalid or *handler* is NULL.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

Referenced by `rthal_timer_request()`.

#### 4.16.2.9 `void rthal_timer_release ( int cpu )`

Release the hardware timer.

Releases the hardware timer, thus reverting the effect of a previous call to [rthal\\_timer\\_request\(\)](#). In case the timer hardware is shared with Linux, a periodic setup suitable for the Linux kernel will be reset.

##### Parameters

- cpu* The CPU number the timer was grabbed from.

Environments:

This service can be called from:

- Linux domain context.

References `rthal_irq_release()`.

**4.16.2.10** `int rthal_timer_request ( void(*) (void) tick_handler, void(*) (enum clock_event_mode mode, struct clock_event_device *cdev) mode_emul, int(*) (unsigned long delay, struct clock_event_device *cdev) tick_emul, int cpu )`

Grab the hardware timer.

`rthal_timer_request()` grabs and tunes the hardware timer in oneshot mode in order to clock the master time base.

A user-defined routine is registered as the clock tick handler. This handler will always be invoked on behalf of the Xenomai domain for each incoming tick.

Hooks for emulating oneshot mode for the tick device are accepted when `CONFIG_GENERIC_CLOCKEVENTS` is defined for the host kernel. Host tick emulation is a way to share the clockchip hardware between Linux and Xenomai, when the former provides support for oneshot timing (i.e. high resolution timers and no-HZ scheduler ticking).

#### Parameters

- tick\_handler* The address of the Xenomai tick handler which will process each incoming tick.
- mode\_emul* The optional address of a callback to be invoked upon mode switch of the host tick device, notified by the Linux kernel. This parameter is only considered whenever `CONFIG_GENERIC_CLOCKEVENTS` is defined.
- tick\_emul* The optional address of a callback to be invoked upon setup of the next shot date for the host tick device, notified by the Linux kernel. This parameter is only considered whenever `CONFIG_GENERIC_CLOCKEVENTS` is defined.
- cpu* The CPU number to grab the timer from.

#### Returns

a positive value is returned on success, representing the duration of a Linux periodic tick expressed as a count of nanoseconds; zero should be returned when the Linux kernel does not undergo periodic timing on the given CPU (e.g. oneshot mode). Otherwise:

- `-EBUSY` is returned if the hardware timer has already been grabbed. `rthal_timer_request()` must be issued before `rthal_timer_request()` is called again.
- `-ENODEV` is returned if the hardware timer cannot be used. This situation may occur after the kernel disabled the timer due to invalid calibration results; in such a case, such hardware is unusable for any timing duties.

Environments:

This service can be called from:

- Linux domain context.

References `rthal_irq_request()`.

#### 4.16.2.11 `int rthal_trap_catch ( rthal_trap_handler_t handler )`

Installs a fault handler.

The HAL attempts to invoke a fault handler whenever an uncontrolled exception or fault is caught at machine level. This service allows to install a user-defined handler for such events.

##### Parameters

*handler* The address of the fault handler to call upon exception condition. The handler is passed the address of the low-level information block describing the fault as passed by Adeos. Its layout is implementation-dependent.

##### Returns

The address of the fault handler previously installed.

Environments:

This service can be called from:

- Any domain context.

## 4.17 Sched

### Data Structures

- struct [xnsched](#)  
*Scheduling information structure.*

### Files

- file [sched.h](#)  
*Scheduler interface header.*
- file [sched-idle.c](#)  
*Idle scheduling class implementation (i.e. Linux placeholder).*
- file [sched-rt.c](#)  
*Common real-time scheduling class implementation (FIFO + RR).*
- file [sched-sporadic.c](#)  
*POSIX SCHED\_SPORADIC scheduling class.*
- file [sched-tp.c](#)  
*Temporal partitioning (typical of IMA systems).*
- file [sched.c](#)

## Typedefs

- typedef struct [xnsched](#) [xnsched\\_t](#)  
*Scheduling information structure.*

## Functions

- static void [xnsched\\_rotate](#) (struct [xnsched](#) \**sched*, struct [xnsched\\_class](#) \**sched\_class*, const union [xnsched\\_policy\\_param](#) \**sched\_param*)  
*Rotate a scheduler runqueue.*

### 4.17.1 Function Documentation

#### 4.17.1.1 void [xnsched\\_rotate](#) ( struct [xnsched](#) \* *sched*, struct [xnsched\\_class](#) \* *sched\_class*, const union [xnsched\\_policy\\_param](#) \* *param* ) [inline, static]

Rotate a scheduler runqueue.

The specified scheduling class is requested to rotate its runqueue for the given scheduler. Rotation is performed according to the scheduling parameter specified by *sched\_param*.

#### Note

The nucleus supports round-robin scheduling for the members of the RT class.

#### Parameters

*sched* The per-CPU scheduler hosting the target scheduling class.

*sched\_class* The scheduling class which should rotate its runqueue.

*sched\_param* The scheduling parameter providing rotation information to the specified scheduling class.

Environments:

This service should be called from:

- Kernel-based task
- Interrupt service routine
- User-space task (primary mode only)

Rescheduling: never.




## Chapter 5

# Data Structure Documentation

### 5.1 xnpod Struct Reference

Real-time pod descriptor.

Collaboration diagram for xnpod:



#### Data Fields

- `xnflags_t` [status](#)
- `xnsched_t` [sched](#) [XNARCH\_NR\_CPUS]
- `xnqueue_t` [threadq](#)
- `xnqueue_t` [tstartq](#)
- `xnqueue_t` [tswitchq](#)
- `xnqueue_t` [tdeleteq](#)
- `atomic_counter_t` [timerlck](#)
- `xntimer_t` [tslicer](#)
- `int` [tsliced](#)
- `int` [refcnt](#)

### 5.1.1 Detailed Description

Real-time pod descriptor. The source of all Xenomai magic.

### 5.1.2 Field Documentation

#### 5.1.2.1 `int xnpod::refcnt`

Reference count.

Referenced by `xnpod_init()`.

#### 5.1.2.2 `xnsched_t xnpod::sched[XNARCH_NR_CPUS]`

Per-cpu scheduler slots.

Referenced by `xnpod_init()`.

#### 5.1.2.3 `xnflags_t xnpod::status`

Status bitmask.

Referenced by `xnpod_init()`.

#### 5.1.2.4 `xnqueue_t xnpod::tdeleteq`

Thread delete hook queue.

Referenced by `xnpod_init()`.

#### 5.1.2.5 `xnqueue_t xnpod::threadq`

All existing threads.

Referenced by `xnpod_init()`.

#### 5.1.2.6 `atomic_counter_t xnpod::timerlck`

Timer lock depth.

Referenced by `xnpod_init()`.

#### 5.1.2.7 `int xnpod::tsliced`

Number of threads using the slicer

Referenced by `xnpod_init()`.

#### 5.1.2.8 xntimer\_t xnpod::tslicer

Time-slicing timer for aperiodic mode

Referenced by xnpod\_init().

#### 5.1.2.9 xnqueue\_t xnpod::tstartq

Thread start hook queue.

Referenced by xnpod\_init().

#### 5.1.2.10 xnqueue\_t xnpod::tswitchq

Thread switch hook queue.

Referenced by xnpod\_init().

The documentation for this struct was generated from the following file:

- include/nucleus/[pod.h](#)

## 5.2 xnsched Struct Reference

Scheduling information structure.

### Data Fields

- xnflags\_t [status](#)
- xnflags\_t [lflags](#)
- struct xnthread \* [curr](#)
- struct xnsched\_rt [rt](#)
- volatile unsigned [inesting](#)
- struct xntimer [htimer](#)
- struct xnthread [rootcb](#)

### 5.2.1 Detailed Description

Scheduling information structure.

### 5.2.2 Field Documentation

#### 5.2.2.1 struct xnthread\* xnsched::curr

Current thread.

Referenced by xnpod\_delete\_thread(), xnpod\_resume\_thread(), and xnpod\_suspend\_thread().

### 5.2.2.2 struct xntimer xnsched::htimer

Host timer.

Referenced by xnpod\_enable\_timesource(), and xntimer\_tick\_aperiodic().

### 5.2.2.3 volatile unsigned xnsched::inesting

Interrupt nesting level.

### 5.2.2.4 xnflags\_t xnsched::lflags

Scheduler specific local flags bitmask.

Referenced by xnpod\_schedule(), and xntimer\_tick\_aperiodic().

### 5.2.2.5 struct xnthread xnsched::rootcb

Root thread control block.

Referenced by xnpod\_init().

### 5.2.2.6 struct xnsched\_rt xnsched::rt

Context of built-in real-time class.

### 5.2.2.7 xnflags\_t xnsched::status

Scheduler specific status bitmask.

Referenced by xnpod\_delete\_thread(), xnpod\_schedule(), and xntimer\_tick\_aperiodic().

The documentation for this struct was generated from the following file:

- include/nucleus/[sched.h](#)

## 5.3 xnthread\_info Struct Reference

Structure containing thread information.

### Data Fields

- unsigned long [state](#)  
*Thread state.*
- int [bprio](#)  
*Base priority.*
- int [cprio](#)

*Current priority.*

- int `cpu`  
*CPU the thread currently runs on.*
- unsigned long `affinity`  
*Thread's CPU affinity.*
- unsigned long long `relpoint`  
*Time of next release.*
- unsigned long long `exectime`  
*Execution time in primary mode in nanoseconds.*
- unsigned long `modeswitches`  
*Number of primary->secondary mode switches.*
- unsigned long `ctxswitches`  
*Number of context switches.*
- unsigned long `pagefaults`  
*Number of triggered page faults.*
- char `name` [XNOBJECT\_NAME\_LEN]  
*Symbolic name assigned at creation.*

### 5.3.1 Detailed Description

Structure containing thread information.

### 5.3.2 Field Documentation

#### 5.3.2.1 unsigned long xntthread\_info::affinity

Thread's CPU affinity.

#### 5.3.2.2 int xntthread\_info::bprio

Base priority.

#### 5.3.2.3 int xntthread\_info::cprio

Current priority.

May change through Priority Inheritance.

**5.3.2.4 int xnthread\_info::cpu**

CPU the thread currently runs on.

**5.3.2.5 unsigned long xnthread\_info::ctxswitches**

Number of context switches.

**5.3.2.6 unsigned long long xnthread\_info::exectime**

Execution time in primary mode in nanoseconds.

**5.3.2.7 unsigned long xnthread\_info::modeswitches**

Number of primary->secondary mode switches.

**5.3.2.8 char xnthread\_info::name[XNOBJECT\_NAME\_LEN]**

Symbolic name assigned at creation.

**5.3.2.9 unsigned long xnthread\_info::pagefaults**

Number of triggered page faults.

**5.3.2.10 unsigned long long xnthread\_info::relpoint**

Time of next release.

**5.3.2.11 unsigned long xnthread\_info::state**

Thread state,.

**See also**

[Thread state flags.](#)

The documentation for this struct was generated from the following file:

- include/nucleus/thread.h

## 5.4 xnvfile\_lock\_ops Struct Reference

Vfile locking operations.

## Data Fields

- `int(* get)(struct xnvfile *vfile)`
- `void(* put)(struct xnvfile *vfile)`

### 5.4.1 Detailed Description

Vfile locking operations.

This structure describes the operations to be provided for implementing locking support on vfiles. They apply to both snapshot-driven and regular vfiles.

### 5.4.2 Field Documentation

#### 5.4.2.1 `int(* xnvfile_lock_ops::get)(struct xnvfile *vfile)`

This handler should grab the desired lock.

##### Parameters

*vfile* A pointer to the virtual file which needs locking.

##### Returns

zero should be returned if the call succeeds. Otherwise, a negative error code can be returned; upon error, the current vfile operation is aborted, and the user-space caller is passed back the error value.

#### 5.4.2.2 `void(* xnvfile_lock_ops::put)(struct xnvfile *vfile)`

This handler should release the lock previously grabbed by the [get\(\)](#) handler.

##### Parameters

*vfile* A pointer to the virtual file which currently holds the lock to release.

The documentation for this struct was generated from the following file:

- `include/nucleus/vfile.h`

## 5.5 xnvfile\_regular\_iterator Struct Reference

Regular vfile iterator.

## Data Fields

- `loff_t pos`  
*Current record position while iterating.*

- struct seq\_file \* [seq](#)  
*Backlink to the host sequential file supporting the vfile.*
- struct xnvfile\_regular \* [vfile](#)  
*Backlink to the vfile being read.*
- char [private](#) [0]  
*Start of private area.*

### 5.5.1 Detailed Description

Regular vfile iterator.

This structure defines an iterator over a regular vfile.

### 5.5.2 Field Documentation

#### 5.5.2.1 loff\_t xnvfile\_regular\_iterator::pos

Current record position while iterating.

#### 5.5.2.2 char xnvfile\_regular\_iterator::private[0]

Start of private area.

Use xnvfile\_iterator\_priv() to address it.

#### 5.5.2.3 struct seq\_file\* xnvfile\_regular\_iterator::seq

Backlink to the host sequential file supporting the vfile.

#### 5.5.2.4 struct xnvfile\_regular\* xnvfile\_regular\_iterator::vfile

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- include/nucleus/[vfile.h](#)

## 5.6 xnvfile\_regular\_ops Struct Reference

Regular vfile operation descriptor.

### Data Fields

- int(\* [rewind](#) )(struct [xnvfile\\_regular\\_iterator](#) \*it)
- void \*([begin](#) )(struct [xnvfile\\_regular\\_iterator](#) \*it)



- void *next* (struct xnvfile\_regular\_iterator \*it)
- void *end* (struct xnvfile\_regular\_iterator \*it)
- int *show* (struct xnvfile\_regular\_iterator \*it, void \*data)
- ssize\_t *store* (struct xnvfile\_input \*input)

### 5.6.1 Detailed Description

Regular vfile operation descriptor.

This structure describes the operations available with a regular vfile. It defines handlers for sending back formatted kernel data upon a user-space read request, and for obtaining user data upon a user-space write request.

### 5.6.2 Field Documentation

#### 5.6.2.1 void (\* xnvfile\_regular\_ops::begin)(struct xnvfile\_regular\_iterator \*it)

This handler should prepare for iterating over the records upon a read request, starting from the specified position.

##### Parameters

- it* A pointer to the current vfile iterator. On entry, *it->pos* is set to the (0-based) position of the first record to output. This handler may be called multiple times with different position requests.

##### Returns

A pointer to the first record to format and output, to be passed to the [show\(\) handler](#) as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- VFILE\_SEQ\_START, a special value indicating that [the show\(\) handler](#) should receive a NULL data pointer first, in order to output a header.
- ERR\_PTR(errno), where errno is a negative error code; upon error, the current operation will be aborted immediately.

##### Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the [show\(\) handler](#) will be called only once for a read operation, with a NULL *data* parameter. This particular setting is convenient for simple regular vfiles having a single, fixed record to output.

#### 5.6.2.2 void (\* xnvfile\_regular\_ops::end)(struct xnvfile\_regular\_iterator \*it)

This handler is called after all records have been output.

**Parameters**

*it* A pointer to the current vfile iterator.

**Note**

This handler is optional and the pointer may be NULL.

**5.6.2.3 void>(\* xnvfile\_regular\_ops::next)(struct xnvfile\_regular\_iterator \*it)**

This handler should return the address of the next record to format and output by the [show\(\)](#) handler".

**Parameters**

*it* A pointer to the current vfile iterator. On entry, *it->pos* is set to the (0-based) position of the next record to output.

**Returns**

A pointer to the next record to format and output, to be passed to the [show\(\)](#) handler as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- ERR\_PTR(errno), where *errno* is a negative error code; upon error, the current operation will be aborted immediately.

**Note**

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the read operation will stop after the first invocation of the [show\(\)](#) handler.

**5.6.2.4 int(\* xnvfile\_regular\_ops::rewind)(struct xnvfile\_regular\_iterator \*it)**

This handler is called only once, when the virtual file is opened, before the [begin\(\)](#) handler is invoked.

**Parameters**

*it* A pointer to the vfile iterator which will be used to read the file contents.

**Returns**

Zero should be returned upon success. Otherwise, a negative error code aborts the operation, and is passed back to the reader.

**Note**

This handler is optional. It should not be used to allocate resources but rather to perform consistency checks, since no closure call is issued in case the open sequence eventually fails.

### 5.6.2.5 `int(* xnvfile_regular_ops::show)(struct xnvfile_regular_iterator *it, void *data)`

This handler should format and output a record.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

#### Parameters

*it* A pointer to the current vfile iterator.

*data* A pointer to the record to format then output. The first call to the handler may receive a NULL *data* pointer, depending on the presence and/or return of a [handler](#); the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

#### Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

### 5.6.2.6 `ssize_t(* xnvfile_regular_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating some kernel setting, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile\\_get\\_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile\\_get\\_integer\(\)](#) will attempt to return a long integer from the input data.

#### Parameters

*input* A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

#### Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the store() handler.

#### Note

This handler is optional, and may be omitted for read-only vfiles.

The documentation for this struct was generated from the following file:

- `include/nucleus/vfile.h`

## 5.7 xnvfile\_rev\_tag Struct Reference

Snapshot revision tag.

### Data Fields

- int [rev](#)  
*Current revision number.*

### 5.7.1 Detailed Description

Snapshot revision tag.

This structure defines a revision tag to be used with [snapshot-driven vfiles](#).

### 5.7.2 Field Documentation

#### 5.7.2.1 int xnvfile\_rev\_tag::rev

Current revision number.


The documentation for this struct was generated from the following file:

- [include/nucleus/vfile.h](#)

## 5.8 xnvfile\_snapshot Struct Reference

Snapshot vfile descriptor.

Collaboration diagram for xnvfile\_snapshot:



### 5.8.1 Detailed Description

Snapshot vfile descriptor.

This structure describes a snapshot-driven vfile. Reading from such a vfile involves a preliminary data collection phase under lock protection, and a subsequent formatting and output phase of the collected data records. Locking is done in a way that does not increase worst-case latency, regardless of the number of records to be collected for output.


The documentation for this struct was generated from the following file:

- `include/nucleus/vfile.h`

## 5.9 xnvfile\_snapshot\_iterator Struct Reference

Snapshot-driven vfile iterator.

Collaboration diagram for xnvfile\_snapshot\_iterator:



### Data Fields

- `int nrdata`  
*Number of collected records.*
- `caddr_t databuf`  
*Address of record buffer.*
- `struct seq_file * seq`  
*Backlink to the host sequential file supporting the vfile.*

- struct [xnvfile\\_snapshot](#) \* [vfile](#)  
*Backlink to the vfile being read.*
- void(\* [endfn](#))(struct [xnvfile\\_snapshot\\_iterator](#) \*it, void \*buf)  
*Buffer release handler.*
- char [private](#) [0]  
*Start of private area.*

### 5.9.1 Detailed Description

Snapshot-driven vfile iterator.

This structure defines an iterator over a snapshot-driven vfile.

### 5.9.2 Field Documentation

#### 5.9.2.1 `caddr_t xnvfile_snapshot_iterator::databuf`

Address of record buffer.

#### 5.9.2.2 `void(* xnvfile_snapshot_iterator::endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`

Buffer release handler.

#### 5.9.2.3 `int xnvfile_snapshot_iterator::nrdata`

Number of collected records.

#### 5.9.2.4 `char xnvfile_snapshot_iterator::private[0]`

Start of private area.

Use `xnvfile_iterator_priv()` to address it.

#### 5.9.2.5 `struct seq_file* xnvfile_snapshot_iterator::seq`

Backlink to the host sequential file supporting the vfile.

#### 5.9.2.6 `struct xnvfile_snapshot* xnvfile_snapshot_iterator::vfile`

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- `include/nucleus/vfile.h`

## 5.10 xnvfile\_snapshot\_ops Struct Reference

Snapshot vfile operation descriptor.

### Data Fields

- `int(* rewind )(struct xnvfile\_snapshot\_iterator *it)`
- `void *(* begin )(struct xnvfile\_snapshot\_iterator *it)`
- `void(* end )(struct xnvfile\_snapshot\_iterator *it, void *buf)`
- `int(* next )(struct xnvfile\_snapshot\_iterator *it, void *data)`
- `int(* show )(struct xnvfile\_snapshot\_iterator *it, void *data)`
- `ssize_t(* store )(struct xnvfile\_input *input)`

### 5.10.1 Detailed Description

Snapshot vfile operation descriptor.

This structure describes the operations available with a snapshot-driven vfile. It defines handlers for returning a printable snapshot of some Xenomai object contents upon a user-space read request, and for updating this object upon a user-space write request.

### 5.10.2 Field Documentation

#### 5.10.2.1 `void *(* xnvfile\_snapshot\_ops::begin )(struct xnvfile\_snapshot\_iterator *it)`

This handler should allocate the snapshot buffer to hold records during the data collection phase. When specified, all records collected via the [next\(\)](#) handler" will be written to a cell from the memory area returned by [begin\(\)](#).

#### Parameters

*it* A pointer to the current snapshot iterator.

#### Returns

A pointer to the record buffer, if the call succeeds. Otherwise:

- NULL in case of allocation error. This will abort the data collection, and return -ENOMEM to the reader.
- VFILE\_SEQ\_EMPTY, a special value indicating that no record will be output. In such a case, the [next\(\) handler](#) will not be called, and the data collection will stop immediately. However, the [show\(\) handler](#) will still be called once, with a NULL data pointer (i.e. header display request).

#### Note

This handler is optional; if none is given, an internal allocation depending on the value returned by the [rewind\(\) handler](#) can be obtained.

### 5.10.2.2 void(\* xnvfile\_snapshot\_ops::end)(struct xnvfile\_snapshot\_iterator \*it, void \*buf)

This handler releases the memory buffer previously obtained from [begin\(\)](#). It is usually called after the snapshot data has been output by [show\(\)](#), but it may also be called before rewinding the vfile after a revision change, to release the dropped buffer.

#### Parameters

- it* A pointer to the current snapshot iterator.
- buf* A pointer to the buffer to release.

#### Note

This routine is optional and the pointer may be NULL. It is not needed upon internal buffer allocation; see the description of the [rewind\(\)](#) handler".

### 5.10.2.3 int(\* xnvfile\_snapshot\_ops::next)(struct xnvfile\_snapshot\_iterator \*it, void \*data)

This handler fetches the next record, as part of the snapshot data to be sent back to the reader via the [show\(\)](#).

#### Parameters

- it* A pointer to the current snapshot iterator.
- data* A pointer to the record to fill in.

#### Returns

a strictly positive value, if the call succeeds and leaves a valid record into *data*, which should be passed to the [show\(\) handler\(\)](#) during the formatting and output phase. Otherwise:

- A negative error code. This will abort the data collection, and return this status to the reader.
- VFILE\_SEQ\_SKIP, a special value indicating that the current record should be skipped. In such a case, the *data* pointer is not advanced to the next position before the [next\(\) handler](#) is called anew.

#### Note

This handler is called with the vfile lock held. Before each invocation of this handler, the vfile core checks whether the revision tag has been touched, in which case the data collection is restarted from scratch. A data collection phase succeeds whenever all records can be fetched via the [next\(\) handler](#), while the revision tag remains unchanged, which indicates that a consistent snapshot of the object state was taken.

### 5.10.2.4 int(\* xnvfile\_snapshot\_ops::rewind)(struct xnvfile\_snapshot\_iterator \*it)

This handler (re-)initializes the data collection, moving the seek pointer at the first record. When the file revision tag is touched while collecting data, the current reading is aborted, all collected data dropped, and the vfile is eventually rewound.



**Parameters**

*it* A pointer to the current snapshot iterator. Two useful information can be retrieved from this iterator in this context:

- *it->vfile* is a pointer to the descriptor of the virtual file being rewound.
- `xnvfile_iterator_priv(it)` returns a pointer to the private data area, available from the descriptor, which size is *vfile->privsz*. If the latter size is zero, the returned pointer is meaningless and should not be used.

**Returns**

A negative error code aborts the data collection, and is passed back to the reader. Otherwise:

- a strictly positive value is interpreted as the total number of records which will be returned by the [next\(\) handler](#) during the data collection phase. If no [begin\(\) handler](#) is provided in the [operation descriptor](#), this value is used to allocate the snapshot buffer internally. The size of this buffer would then be *vfile->datasz* \* value.
- zero leaves the allocation to the [begin\(\) handler](#) if present, or indicates that no record is to be output in case such handler is not given.

**Note**

This handler is optional; a NULL value indicates that nothing needs to be done for rewinding the *vfile*. It is called with the *vfile* lock held.

**5.10.2.5 `int(* xnvfile_snapshot_ops::show)(struct xnvfile_snapshot_iterator *it, void *data)`**

This handler should format and output a record from the collected data.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

**Parameters**

*it* A pointer to the current snapshot iterator.

*data* A pointer to the record to format then output. The first call to the handler is always passed a NULL *data* pointer; the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

**Returns**

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

#### 5.10.2.6 `ssize_t(* xnvfile_snapshot_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating the associated Xenomai object's state, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile\\_get\\_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile\\_get\\_integer\(\)](#) will attempt to return a long integer from the input data.

##### Parameters

*input* A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

##### Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

##### Note

This handler is optional, and may be omitted for read-only vfiles.

Referenced by `xnvfile_init_snapshot()`.

The documentation for this struct was generated from the following file:


- `include/nucleus/vfile.h`

## Chapter 6


# File Documentation

### 6.1 include/nucleus/bufd.h File Reference

Include dependency graph for bufd.h:



This graph shows which files directly or indirectly include this file:



## Functions

- static void [xnbufd\\_map\\_uread](#) (struct xnbufd \*bufd, const void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from user memory.*
- static void [xnbufd\\_map\\_uwrite](#) (struct xnbufd \*bufd, void \_\_user \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to user memory.*
- ssize\_t [xnbufd\\_unmap\\_uread](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_uread\(\)](#).*
- ssize\_t [xnbufd\\_unmap\\_uwrite](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_uwrite\(\)](#).*
- static void [xnbufd\\_map\\_kread](#) (struct xnbufd \*bufd, const void \*ptr, size\_t len)  
*Initialize a buffer descriptor for reading from kernel memory.*
- static void [xnbufd\\_map\\_kwrite](#) (struct xnbufd \*bufd, void \*ptr, size\_t len)  
*Initialize a buffer descriptor for writing to kernel memory.*
- ssize\_t [xnbufd\\_unmap\\_kread](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_kread\(\)](#).*
- ssize\_t [xnbufd\\_unmap\\_kwrite](#) (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from [xnbufd\\_map\\_kwrite\(\)](#).*
- ssize\_t [xnbufd\\_copy\\_to\\_kmem](#) (void \*ptr, struct xnbufd \*bufd, size\_t len)  
*Copy memory covered by a buffer descriptor to kernel memory.*
- ssize\_t [xnbufd\\_copy\\_from\\_kmem](#) (struct xnbufd \*bufd, void \*from, size\_t len)  
*Copy kernel memory to the area covered by a buffer descriptor.*
- void [xnbufd\\_invalidate](#) (struct xnbufd \*bufd)  
*Invalidate a buffer descriptor.*
- static void [xnbufd\\_reset](#) (struct xnbufd \*bufd)  
*Reset a buffer descriptor.*

### 6.1.1 Detailed Description

#### Note

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.2 include/nucleus/hostrt.h File Reference

Definitions for global semaphore heap shared objects.

Include dependency graph for hostrt.h:



This graph shows which files directly or indirectly include this file:



### 6.2.1 Detailed Description

Definitions for global semaphore heap shared objects.

#### Author

Wolfgang Mauerer

Copyright (C) 2010 Wolfgang Mauerer <[wolfgang.mauerer@siemens.com](mailto:wolfgang.mauerer@siemens.com)>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.3 include/nucleus/map.h File Reference

Include dependency graph for map.h:



This graph shows which files directly or indirectly include this file:



## Functions

- `xnmap_t * xnmap\_create` (int nkeys, int reserve, int offset)  
*Create a map.*
- `void xnmap\_delete` (xnmap\_t \*map)  
*Delete a map.*
- `int xnmap\_enter` (xnmap\_t \*map, int key, void \*objaddr)  
*Index an object into a map.*
- `int xnmap\_remove` (xnmap\_t \*map, int key)  
*Remove an object reference from a map.*
- `static void * xnmap\_fetch\_nocheck` (xnmap\_t \*map, int key)  
*Search an object into a map - unchecked form.*
- `static void * xnmap\_fetch` (xnmap\_t \*map, int key)  
*Search an object into a map.*

### 6.3.1 Detailed Description

#### Note

Copyright (C) 2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.4 include/nucleus/pod.h File Reference

Real-time pod interface header.

Include dependency graph for pod.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [xn timer](#)  
*Real-time pod descriptor.*

## Functions

- int [xn timer\\_init](#) (void)  
*Initialize the core pod.*
- int [xn timer\\_enable\\_timesource](#) (void)  
*Activate the core time source.*
- void [xn timer\\_disable\\_timesource](#) (void)  
*Stop the core time source.*



- void [xnpod\\_shutdown](#) (int xtype)  
*Shutdown the current pod.*
- int [xnpod\\_init\\_thread](#) (struct xnthread \*thread, const struct xnthread\_init\_attr \*attr, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Initialize a new thread.*
- int [xnpod\\_start\\_thread](#) (xnthread\_t \*thread, const struct xnthread\_start\_attr \*attr)  
*Initial start of a newly created thread.*
- void [xnpod\\_stop\\_thread](#) (xnthread\_t \*thread)  
*Stop a thread.*
- void [xnpod\\_restart\\_thread](#) (xnthread\_t \*thread)  
*Restart a thread.*
- void [xnpod\\_delete\\_thread](#) (xnthread\_t \*thread)  
*Delete a thread.*
- void [xnpod\\_abort\\_thread](#) (xnthread\_t \*thread)  
*Abort a thread.*
- xnflags\_t [xnpod\\_set\\_thread\\_mode](#) (xnthread\_t \*thread, xnflags\_t clrmask, xnflags\_t setmask)  
*Change a thread's control mode.*
- void [xnpod\\_suspend\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \*wchan)  
*Suspend a thread.*
- void [xnpod\\_resume\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask)  
*Resume a thread.*
- int [xnpod\\_unblock\\_thread](#) (xnthread\_t \*thread)  
*Unblock a thread.*
- int [xnpod\\_set\\_thread\\_schedparam](#) (struct xnthread \*thread, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*
- int [xnpod\\_migrate\\_thread](#) (int cpu)  
*Migrate the current thread.*
- void [xnpod\\_dispatch\\_signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- static void [xnpod\\_schedule](#) (void)  
*Rescheduling procedure entry point.*

- int [xnpod\\_set\\_thread\\_periodic](#) (xnthread\_t \*thread, xnticks\_t idate, xnticks\_t period)

*Make a thread periodic.*

- int [xnpod\\_wait\\_thread\\_period](#) (unsigned long \*overruns\_r)

*Wait for the next periodic release point.*

- int [xnpod\\_set\\_thread\\_tslice](#) (struct xnthread \*thread, xnticks\_t quantum)

*Set thread time-slicing information.*

- int [xnpod\\_add\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))

*Install a nucleus hook.*

- int [xnpod\\_remove\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))

*Remove a nucleus hook.*

### 6.4.1 Detailed Description

Real-time pod interface header.

#### Author

Philippe Gerum

Copyright (C) 2001-2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2004 The Xenomai project <<http://www.xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.5 include/nucleus/registry.h File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.h:



This graph shows which files directly or indirectly include this file:



## Functions

- `int xnregistry_enter (const char *key, void *objaddr, xnhandle_t *phandle, struct xnnode *pnode)`  
*Register a real-time object.*
- `int xnregistry_bind (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)`  
*Bind to a real-time object.*
- `int xnregistry_remove (xnhandle_t handle)`  
*Forcibly unregister a real-time object.*
- `int xnregistry_remove_safe (xnhandle_t handle, xnticks_t timeout)`  
*Unregister an idle real-time object.*
- `void * xnregistry_get (xnhandle_t handle)`  
*Find and lock a real-time object into the registry.*
- `void * xnregistry_fetch (xnhandle_t handle)`

*Find a real-time object into the registry.*

- u\_long [xnregistry\\_put](#) (xnhandle\_t handle)

*Unlock a real-time object from the registry.*

### 6.5.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2004 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.6 include/nucleus/sched-idle.h File Reference

Definitions for the IDLE scheduling class.

This graph shows which files directly or indirectly include this file:



### 6.6.1 Detailed Description

Definitions for the IDLE scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.7 include/nucleus/sched-rt.h File Reference

### Definitions for the RT scheduling class.

This graph shows which files directly or indirectly include this file:



### 6.7.1 Detailed Description

### Definitions for the RT scheduling class.

**Author**

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.8 include/nucleus/sched-sporadic.h File Reference

### Definitions for the SSP scheduling class.

This graph shows which files directly or indirectly include this file:



### 6.8.1 Detailed Description

Definitions for the SSP scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.9 include/nucleus/sched-tp.h File Reference

Definitions for the TP scheduling class.

This graph shows which files directly or indirectly include this file:



### 6.9.1 Detailed Description

Definitions for the TP scheduling class.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

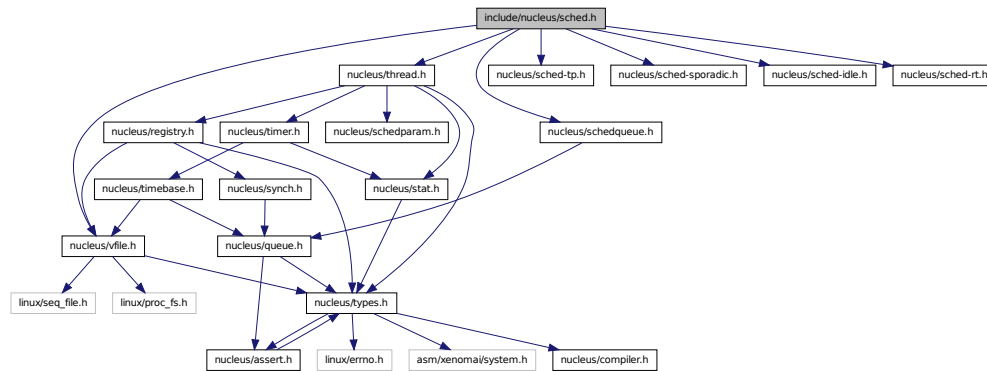
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.10 include/nucleus/sched.h File Reference

Scheduler interface header.

Include dependency graph for sched.h:



This graph shows which files directly or indirectly include this file:



### Data Structures

- struct [xnsched](#)  
*Scheduling information structure.*

### Typedefs

- typedef struct [xnsched](#) [xnsched\\_t](#)  
*Scheduling information structure.*

### Functions

- static void [xnsched\\_rotate](#) (struct [xnsched](#) \*sched, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Rotate a scheduler runqueue.*

### 6.10.1 Detailed Description

Scheduler interface header.

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

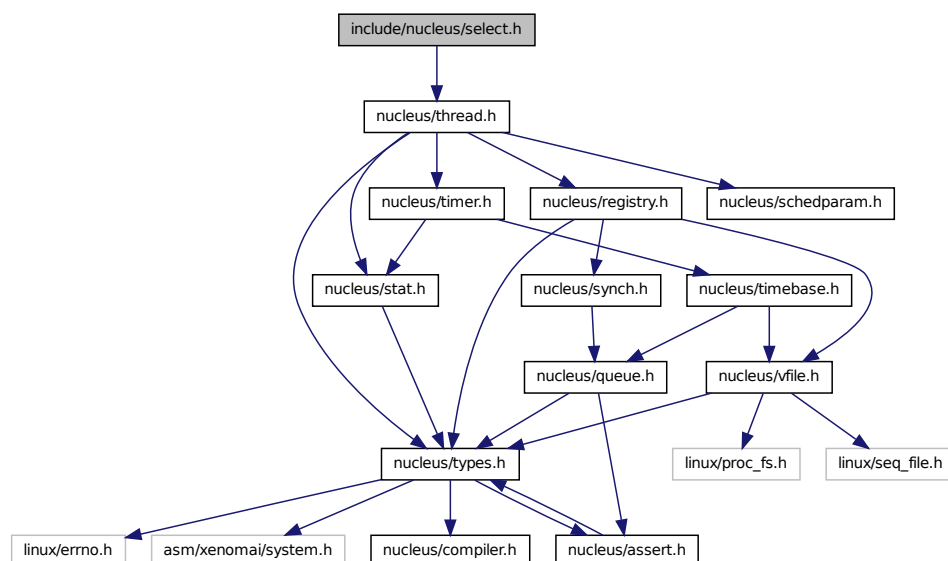
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.11 include/nucleus/select.h File Reference


file descriptors events multiplexing header.

Include dependency graph for select.h:





This graph shows which files directly or indirectly include this file:



### 6.11.1 Detailed Description

file descriptors events multiplexing header.

#### Author

Gilles Chanteperdrix

Copyright (C) 2008 Efixo <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>

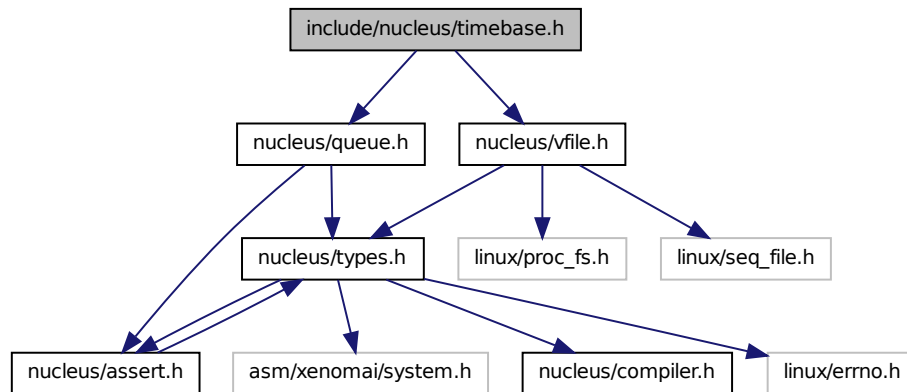
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.12 include/nucleus/timebase.h File Reference

Include dependency graph for timebase.h:



This graph shows which files directly or indirectly include this file:



## Functions

- `int xntbase_alloc` (`const char *name`, `u_long period`, `u_long flags`, `xntbase_t **basep`)  
*Allocate a time base.*
- `void xntbase_free` (`xntbase_t *base`)  
*Free a time base.*
- `int xntbase_update` (`xntbase_t *base`, `u_long period`)  
*Change the period of a time base.*
- `int xntbase_switch` (`const char *name`, `u_long period`, `xntbase_t **basep`)  
*Replace a time base.*
- `void xntbase_start` (`xntbase_t *base`)  
*Start a time base.*
- `void xntbase_stop` (`xntbase_t *base`)

*Stop a time base.*

- void [xntbase\\_tick](#) (xntbase\_t \*base)

*Announce a clock tick to a time base.*

- xnticks\_t [xntbase\\_convert](#) (xntbase\_t \*srcbase, xnticks\_t ticks, xntbase\_t \*dstbase)

*Convert a clock value into another time base.*

- static xnticks\_t [xntbase\\_get\\_time](#) (xntbase\_t \*base)

*Get the clock time for a given time base.*

- void [xntbase\\_adjust\\_time](#) (xntbase\_t \*base, xnticks\_t delta)

*Adjust the clock time for the system.*

### 6.12.1 Detailed Description

#### Note

Copyright (C) 2006,2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.13 include/nucleus/timer.h File Reference

Include dependency graph for timer.h:



This graph shows which files directly or indirectly include this file:



## Functions

- void [xntimer\\_destroy](#) (xntimer\_t \*timer)  
*Release a timer object.*
- static int [xntimer\\_start](#) (xntimer\_t \*timer, xnticks\_t value, xnticks\_t interval, xntmode\_t mode)  
*Arm a timer.*
- static void [xntimer\\_stop](#) (xntimer\_t \*timer)  
*Disarm a timer.*
- static xnticks\_t [xntimer\\_get\\_date](#) (xntimer\_t \*timer)  
*Return the absolute expiration date.*
- static xnticks\_t [xntimer\\_get\\_timeout](#) (xntimer\_t \*timer)

*Return the relative expiration date.*

- static xnticks\_t [xntimer\\_get\\_interval](#) (xntimer\_t \*timer)  
*Return the timer interval value.*
- unsigned long [xntimer\\_get\\_overruns](#) (xntimer\_t \*timer, xnticks\_t now)  
*Get the count of overruns for the last tick.*
- void [xntimer\\_freeze](#) (void)  
*Freeze all timers (from every time bases).*
- void [xntimer\\_tick\\_aperiodic](#) (void)  
*Process a timer tick for the aperiodic master time base.*
- void [xntimer\\_tick\\_periodic](#) (xntimer\_t \*timer)  
*Process a timer tick for a slave periodic time base.*

### 6.13.1 Detailed Description

#### Note

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.14 include/nucleus/vdso.h File Reference

Definitions for global semaphore heap shared objects.

Include dependency graph for vdso.h:



This graph shows which files directly or indirectly include this file:



### 6.14.1 Detailed Description

Definitions for global semaphore heap shared objects.

#### Author

Wolfgang Maurerer

Copyright (C) 2009 Wolfgang Maurerer <[wolfgang.maurerer@siemens.com](mailto:wolfgang.maurerer@siemens.com)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


## 6.15 include/nucleus/vfile.h File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.h:



This graph shows which files directly or indirectly include this file:



## Data Structures

- struct [xnfile\\_lock\\_ops](#)  
*Vfile locking operations.*
- struct [xnfile\\_regular\\_ops](#)  
*Regular vfile operation descriptor.*
- struct [xnfile\\_regular\\_iterator](#)  
*Regular vfile iterator.*
- struct [xnfile\\_snapshot\\_ops](#)  
*Snapshot vfile operation descriptor.*
- struct [xnfile\\_rev\\_tag](#)  
*Snapshot revision tag.*
- struct [xnfile\\_snapshot](#)  
*Snapshot vfile descriptor.*
- struct [xnfile\\_snapshot\\_iterator](#)  
*Snapshot-driven vfile iterator.*

## Functions

- int [xnvmfile\\_init\\_snapshot](#) (const char \*name, struct [xnvmfile\\_snapshot](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a snapshot-driven vfile.*
- int [xnvmfile\\_init\\_regular](#) (const char \*name, struct [xnvmfile\\_regular](#) \*vfile, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a regular vfile.*
- int [xnvmfile\\_init\\_dir](#) (const char \*name, struct [xnvmfile\\_directory](#) \*vdir, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual directory entry.*
- int [xnvmfile\\_init\\_link](#) (const char \*from, const char \*to, struct [xnvmfile\\_link](#) \*vlink, struct [xnvmfile\\_directory](#) \*parent)  
*Initialize a virtual link entry.*
- void [xnvmfile\\_destroy](#) (struct [xnvmfile](#) \*vfile)  
*Removes a virtual file entry.*
- ssize\_t [xnvmfile\\_get\\_blob](#) (struct [xnvmfile\\_input](#) \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_string](#) (struct [xnvmfile\\_input](#) \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- ssize\_t [xnvmfile\\_get\\_integer](#) (struct [xnvmfile\\_input](#) \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct [xnvmfile\\_directory](#) [nkvfroot](#)  
*Xenomai vfile root directory.*

### 6.15.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2010 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.



You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.16 ksrc/arch/arm/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for ARM.

Include dependency graph for hal.c:



### Functions

- `int rthal_timer_request` (void(\*tick\_handler)(void), void(\*mode\_emul)(enum clock\_event\_mode mode, struct clock\_event\_device \*cdev), int(\*tick\_emul)(unsigned long delay, struct clock\_event\_device \*cdev), int cpu)  
*Grab the hardware timer.*
- `void rthal_timer_release` (int cpu)  
*Release the hardware timer.*
- `int rthal_irq_host_request` (unsigned irq, rthal\_irq\_host\_handler\_t handler, char \*name, void \*dev\_id)  
*Install a shared Linux interrupt handler.*
- `int rthal_irq_host_release` (unsigned irq, void \*dev\_id)  
*Uninstall a shared Linux interrupt handler.*

### 6.16.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for ARM. ARM port Copyright (C) 2005 Stelian Pop

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.17 ksrc/arch/blackfin/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.

Include dependency graph for hal.c:



### Functions

- `int rthal_timer_request` (void(\*tick\_handler)(void), void(\*mode\_emul)(enum clock\_event\_mode mode, struct clock\_event\_device \*cdev), int(\*tick\_emul)(unsigned long delay, struct clock\_event\_device \*cdev), int cpu)  
*Grab the hardware timer.*
- `void rthal_timer_release` (int cpu)  
*Release the hardware timer.*
- `int rthal_irq_enable` (unsigned irq)  
*Enable an interrupt source.*
- `int rthal_irq_disable` (unsigned irq)  
*Disable an interrupt source.*
- `int rthal_irq_host_request` (unsigned irq, rthal\_irq\_host\_handler\_t handler, char \*name, void \*dev\_id)  
*Install a shared Linux interrupt handler.*
- `int rthal_irq_host_release` (unsigned irq, void \*dev\_id)  
*Uninstall a shared Linux interrupt handler.*

### 6.17.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the Blackfin architecture. Copyright (C) 2005-2006 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.18 ksrc/arch/generic/hal.c File Reference

Generic Real-Time HAL.

Include dependency graph for hal.c:



### Functions

- int [rthal\\_irq\\_request](#) (unsigned irq, rthal\_irq\_handler\_t handler, rthal\_irq\_ackfn\_t ackfn, void \*cookie)  
*Install a real-time interrupt handler.*
- int [rthal\\_irq\\_release](#) (unsigned irq)  
*Uninstall a real-time interrupt handler.*
- rthal\_trap\_handler\_t [rthal\\_trap\\_catch](#) (rthal\_trap\_handler\_t handler)  
*Installs a fault handler.*
- int [rthal\\_apc\\_alloc](#) (const char \*name, void(\*handler)(void \*cookie), void \*cookie)  
*Allocate an APC slot.*
- void [rthal\\_apc\\_free](#) (int apc)  
*Releases an APC slot.*

### 6.18.1 Detailed Description

Generic Real-Time HAL. Copyright ©2005 Philippe Gerum.

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.19 ksrc/arch/nios2/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture.

Include dependency graph for hal.c:



## Functions

- void [rthal\\_timer\\_release](#) (int cpu)  
*Release the hardware timer.*
- int [rthal\\_irq\\_enable](#) (unsigned irq)  
*Enable an interrupt source.*
- int [rthal\\_irq\\_disable](#) (unsigned irq)  
*Disable an interrupt source.*
- int [rthal\\_irq\\_host\\_request](#) (unsigned irq, rthal\_irq\_host\_handler\_t handler, char \*name, void \*dev\_id)  
*Install a shared Linux interrupt handler.*
- int [rthal\\_irq\\_host\\_release](#) (unsigned irq, void \*dev\_id)  
*Uninstall a shared Linux interrupt handler.*

### 6.19.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the NIOS2 architecture. Copyright (C) 2009 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.20 ksrc/arch/powerpc/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for PowerPC.

Include dependency graph for hal.c:



## Functions

- `int rthal_timer_request` (void(\*tick\_handler)(void), void(\*mode\_emul)(enum clock\_event\_mode mode, struct clock\_event\_device \*cdev), int(\*tick\_emul)(unsigned long delay, struct clock\_event\_device \*cdev), int cpu)

*Grab the hardware timer.*

- `void rthal_timer_release` (int cpu)

*Release the hardware timer.*

- `int rthal_irq_host_request` (unsigned irq, rthal\_irq\_host\_handler\_t handler, char \*name, void \*dev\_id)

*Install a shared Linux interrupt handler.*

- `int rthal_irq_host_release` (unsigned irq, void \*dev\_id)

*Uninstall a shared Linux interrupt handler.*

- `int rthal_irq_enable` (unsigned irq)

*Enable an interrupt source.*

- `int rthal_irq_disable` (unsigned irq)

*Disable an interrupt source.*

### 6.20.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for PowerPC. Copyright (C) 2004-2006 Philippe Gerum.

64-bit PowerPC adoption copyright (C) 2005 Taneli Vähäkangas and Heikki Lindholm

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.21 ksrc/arch/sh/hal.c File Reference

Adeos-based Real-Time Abstraction Layer for the SuperH architecture.

Include dependency graph for hal.c:



### Functions

- `int rthal_timer_request (void(*tick_handler)(void), void(*mode_emul)(enum clock_event_ - mode mode, struct clock_event_device *cdev), int(*tick_emul)(unsigned long delay, struct clock_event_device *cdev), int cpu)`  
*Grab the hardware timer.*
- `void rthal_timer_release (int cpu)`  
*Release the hardware timer.*
- `int rthal_irq_enable (unsigned irq)`  
*Enable an interrupt source.*
- `int rthal_irq_disable (unsigned irq)`  
*Disable an interrupt source.*
- `int rthal_irq_host_request (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)`  
*Install a shared Linux interrupt handler.*
- `int rthal_irq_host_release (unsigned irq, void *dev_id)`  
*Uninstall a shared Linux interrupt handler.*

### 6.21.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the SuperH architecture. Copyright (C) 2011 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.22 ksrc/arch/x86/hal-common.c File Reference

Adeos-based Real-Time Abstraction Layer for x86.

Include dependency graph for hal-common.c:



### Functions

- `int rthal_irq_host_request` (unsigned irq, `rthal_irq_host_handler_t` handler, char \*name, void \*dev\_id)  
*Install a shared Linux interrupt handler.*
- `int rthal_irq_host_release` (unsigned irq, void \*dev\_id)  
*Uninstall a shared Linux interrupt handler.*
- `int rthal_irq_enable` (unsigned irq)  
*Enable an interrupt source.*
- `int rthal_irq_disable` (unsigned irq)  
*Disable an interrupt source.*

### 6.22.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86. Common code of i386 and x86\_64.

Copyright (C) 2007 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.23 ksrc/arch/x86/hal\_32.c File Reference

Adeos-based Real-Time Abstraction Layer for x86.

Include dependency graph for hal\_32.c:



### Functions

- `int rthal_timer_request` (void(\*tick\_handler)(void), void(\*mode\_emul)(enum clock\_event\_mode mode, struct clock\_event\_device \*cdev), int(\*tick\_emul)(unsigned long delay, struct clock\_event\_device \*cdev), int cpu)  
*Grab the hardware timer.*
- `void rthal_timer_release` (int cpu)  
*Release the hardware timer.*

### 6.23.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86. Inspired from original RTAI/x86 HAL interface:

Copyright ©2000 Paolo Mantegazza,

Copyright ©2000 Steve Papacharalambous,

Copyright ©2000 Stuart Hughes,

RTAI/x86 rewrite over Adeos:

Copyright ©2002-2007 Philippe Gerum. SMI workaround:

Copyright ©2004 Gilles Chanteperdrix.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.



You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.24 ksrc/arch/x86/hal\_64.c File Reference

Adeos-based Real-Time Abstraction Layer for x86\_64.

Include dependency graph for hal\_64.c:



### 6.24.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86\_64. Derived from the Xenomai/i386 HAL.

Copyright (C) 2007 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.25 ksrc/arch/x86/smi.c File Reference

SMI workaround for x86.

Include dependency graph for smi.c:



### 6.25.1 Detailed Description

SMI workaround for x86. Cut/Pasted from Vitor Angelo "smi" module. Adapted by Gilles Chanteperdrix <[gilles.chanteperdrix@xenomai.org](mailto:gilles.chanteperdrix@xenomai.org)>.


This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.26 ksrc/nucleus/bufd.c File Reference

Include dependency graph for bufd.c:



### Functions

- `ssize_t xnbufd_copy_to_kmem` (void \*ptr, struct xnbufd \*bufd, size\_t len)  
*Copy memory covered by a buffer descriptor to kernel memory.*
- `ssize_t xnbufd_copy_from_kmem` (struct xnbufd \*bufd, void \*from, size\_t len)  
*Copy kernel memory to the area covered by a buffer descriptor.*
- `ssize_t xnbufd_unmap_uread` (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from `xnbufd_map_uread()`.*
- `ssize_t xnbufd_unmap_uwrite` (struct xnbufd \*bufd)  
*Finalize a buffer descriptor obtained from `xnbufd_map_uwrite()`.*
- `void xnbufd_invalidate` (struct xnbufd \*bufd)



*Set the heap's label string.*

- void [xnheap\\_destroy](#) (xnheap\_t \*heap, void(\*flushfn)(xnheap\_t \*heap, void \*extaddr, u\_long extsize, void \*cookie), void \*cookie)

*Destroys a memory heap.*

- void \* [xnheap\\_alloc](#) (xnheap\_t \*heap, u\_long size)

*Allocate a memory block from a memory heap.*

- int [xnheap\\_test\\_and\\_free](#) (xnheap\_t \*heap, void \*block, int(\*ckfn)(void \*block))

*Test and release a memory block to a memory heap.*

- int [xnheap\\_free](#) (xnheap\_t \*heap, void \*block)

*Release a memory block to a memory heap.*

- int [xnheap\\_extend](#) (xnheap\_t \*heap, void \*extaddr, u\_long extsize)

*Extend a memory heap.*

- void [xnheap\\_schedule\\_free](#) (xnheap\_t \*heap, void \*block, xnholder\_t \*link)

*Schedule a memory block for release.*

## 6.27.1 Detailed Description

Dynamic memory allocation services.

### Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

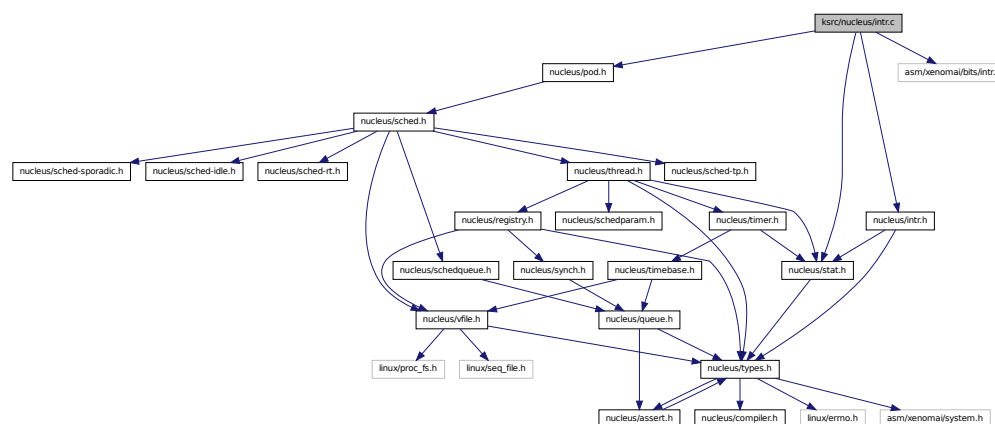
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.28 ksrc/nucleus/intr.c File Reference

Interrupt management.

Include dependency graph for intr.c:



## Functions

- int [xnintr\\_init](#) (xnintr\_t \*intr, const char \*name, unsigned irq, xnintr\_t isr, xniack\_t iack, xnflags\_t flags)  
*Initialize an interrupt object.*
- int [xnintr\\_destroy](#) (xnintr\_t \*intr)  
*Destroy an interrupt object.*
- int [xnintr\\_attach](#) (xnintr\_t \*intr, void \*cookie)  
*Attach an interrupt object.*
- int [xnintr\\_detach](#) (xnintr\_t \*intr)  
*Detach an interrupt object.*
- int [xnintr\\_enable](#) (xnintr\_t \*intr)  
*Enable an interrupt object.*
- int [xnintr\\_disable](#) (xnintr\_t \*intr)  
*Disable an interrupt object.*
- void [xnintr\\_affinity](#) (xnintr\_t \*intr, xnarch\_cpumask\_t cpumask)  
*Set interrupt's processor affinity.*

### 6.28.1 Detailed Description

Interrupt management.

#### Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2005,2006 Dmitry Adamushko <[dmitry.adamushko@gmail.com](mailto:dmitry.adamushko@gmail.com)>. Copyright (C) 2007 Jan Kiszka <[jan.kiszka@web.de](mailto:jan.kiszka@web.de)>.

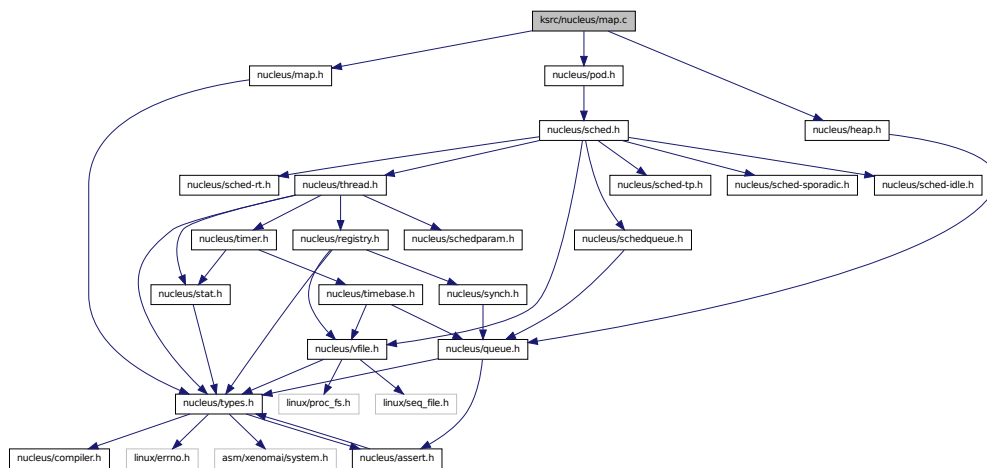
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.29 ksrc/nucleus/map.c File Reference

Include dependency graph for map.c:



## Functions

- `xnmap_t * xnmap\_create` (int nkeys, int reserve, int offset)  
*Create a map.*
- `void xnmap\_delete` (xnmap\_t \*map)  
*Delete a map.*
- `int xnmap\_enter` (xnmap\_t \*map, int key, void \*objaddr)  
*Index an object into a map.*
- `int xnmap\_remove` (xnmap\_t \*map, int key)  
*Remove an object reference from a map.*

## 6.29.1 Detailed Description

### Note

Copyright (C) 2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

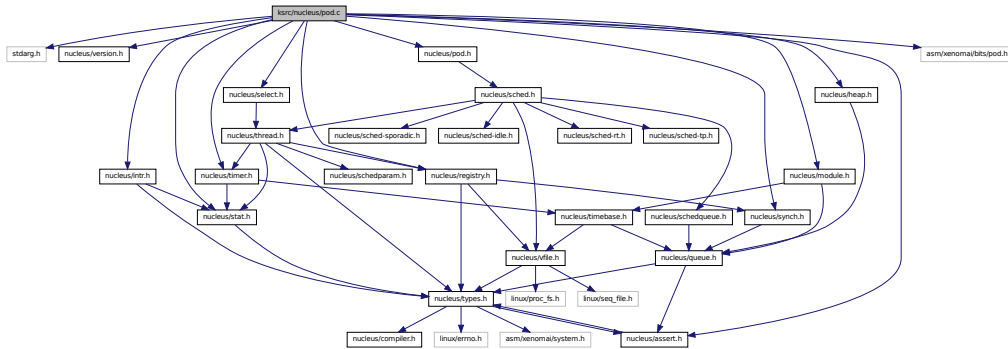
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.30 ksrc/nucleus/pod.c File Reference

Real-time pod services.

Include dependency graph for pod.c:



## Functions

- `int xnpod_init (void)`  
*Initialize the core pod.*
- `void xnpod_shutdown (int xtype)`  
*Shutdown the current pod.*
- `int xnpod_init_thread (struct xnthread *thread, const struct xnthread_init_attr *attr, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)`  
*Initialize a new thread.*
- `int xnpod_start_thread (xnthread_t *thread, const struct xnthread_start_attr *attr)`  
*Initial start of a newly created thread.*

- void [xnpod\\_stop\\_thread](#) (xnthread\_t \*thread)  
*Stop a thread.*
- void [xnpod\\_restart\\_thread](#) (xnthread\_t \*thread)  
*Restart a thread.*
- xnflags\_t [xnpod\\_set\\_thread\\_mode](#) (xnthread\_t \*thread, xnflags\_t clrmask, xnflags\_t set-mask)  
*Change a thread's control mode.*
- void [xnpod\\_delete\\_thread](#) (xnthread\_t \*thread)  
*Delete a thread.*
- void [xnpod\\_abort\\_thread](#) (xnthread\_t \*thread)  
*Abort a thread.*
- void [xnpod\\_suspend\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask, xnticks\_t timeout, xntmode\_t timeout\_mode, struct xnsynch \*wchan)  
*Suspend a thread.*
- void [xnpod\\_resume\\_thread](#) (xnthread\_t \*thread, xnflags\_t mask)  
*Resume a thread.*
- int [xnpod\\_unblock\\_thread](#) (xnthread\_t \*thread)  
*Unblock a thread.*
- int [xnpod\\_set\\_thread\\_schedparam](#) (struct xnthread \*thread, struct xnsched\_class \*sched\_class, const union xnsched\_policy\_param \*sched\_param)  
*Change the base scheduling parameters of a thread.*
- int [xnpod\\_migrate\\_thread](#) (int cpu)  
*Migrate the current thread.*
- void [xnpod\\_dispatch\\_signals](#) (void)  
*Deliver pending asynchronous signals to the running thread.*
- void [xnpod\\_welcome\\_thread](#) (xnthread\_t \*thread, int imask)  
*Thread prologue.*
- int [xnpod\\_add\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Install a nucleus hook.*
- int [xnpod\\_remove\\_hook](#) (int type, void(\*routine)(xnthread\_t \*))  
*Remove a nucleus hook.*
- int [xnpod\\_trap\\_fault](#) (xnarch\_fltinfo\_t \*fltinfo)  
*Default fault handler.*
- int [xnpod\\_enable\\_timesource](#) (void)



*Activate the core time source.*

- void [xnpod\\_disable\\_timesource](#) (void)

*Stop the core time source.*

- int [xnpod\\_set\\_thread\\_periodic](#) (xnthread\_t \*thread, xnticks\_t idate, xnticks\_t period)

*Make a thread periodic.*

- int [xnpod\\_wait\\_thread\\_period](#) (unsigned long \*overruns\_r)

*Wait for the next periodic release point.*

- int [xnpod\\_set\\_thread\\_tslice](#) (struct xnthread \*thread, xnticks\_t quantum)

*Set thread time-slicing information.*

### 6.30.1 Detailed Description

Real-time pod services.

#### Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.Xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.31 ksrc/nucleus/registry.c File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.c:



## Functions

- `int xnregistry_enter` (const char \*key, void \*objaddr, xnhandle\_t \*phandle, struct xninode \*pnode)  
*Register a real-time object.*
- `int xnregistry_bind` (const char \*key, xnticks\_t timeout, int timeout\_mode, xnhandle\_t \*phandle)  
*Bind to a real-time object.*
- `int xnregistry_remove` (xnhandle\_t handle)  
*Forcibly unregister a real-time object.*
- `int xnregistry_remove_safe` (xnhandle\_t handle, xnticks\_t timeout)  
*Unregister an idle real-time object.*
- `void * xnregistry_get` (xnhandle\_t handle)  
*Find and lock a real-time object into the registry.*
- `u_long xnregistry_put` (xnhandle\_t handle)  
*Unlock a real-time object from the registry.*
- `void * xnregistry_fetch` (xnhandle\_t handle)  
*Find a real-time object into the registry.*

### 6.31.1 Detailed Description

This file is part of the Xenomai project.

**Note**

Copyright (C) 2004 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.32 ksrc/nucleus/sched-idle.c File Reference

Idle scheduling class implementation (i.e. Linux placeholder).

Include dependency graph for sched-idle.c:



### 6.32.1 Detailed Description

Idle scheduling class implementation (i.e. Linux placeholder).

**Author**

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.33 ksrc/nucleus/sched-rt.c File Reference

Common real-time scheduling class implementation (FIFO + RR).

Include dependency graph for sched-rt.c:



### 6.33.1 Detailed Description

Common real-time scheduling class implementation (FIFO + RR).

#### Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

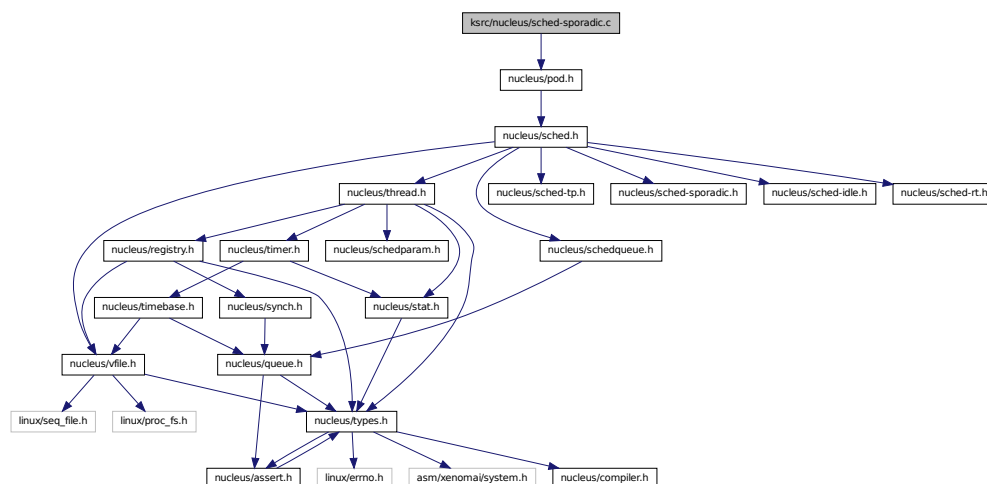
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.34 ksrc/nucleus/sched-sporadic.c File Reference

POSIX SCHED\_SPORADIC scheduling class.

Include dependency graph for sched-sporadic.c:



### 6.34.1 Detailed Description

POSIX SCHED\_SPORADIC scheduling class.

#### Author

Philippe Gerum Copyright (C) 2009 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.35 ksrc/nucleus/sched-tp.c File Reference

Temporal partitioning (typical of IMA systems).

Include dependency graph for sched-tp.c:



### 6.35.1 Detailed Description

Temporal partitioning (typical of IMA systems).

#### Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.36 ksrc/nucleus/sched.c File Reference

Include dependency graph for sched.c:



### 6.36.1 Detailed Description

#### Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

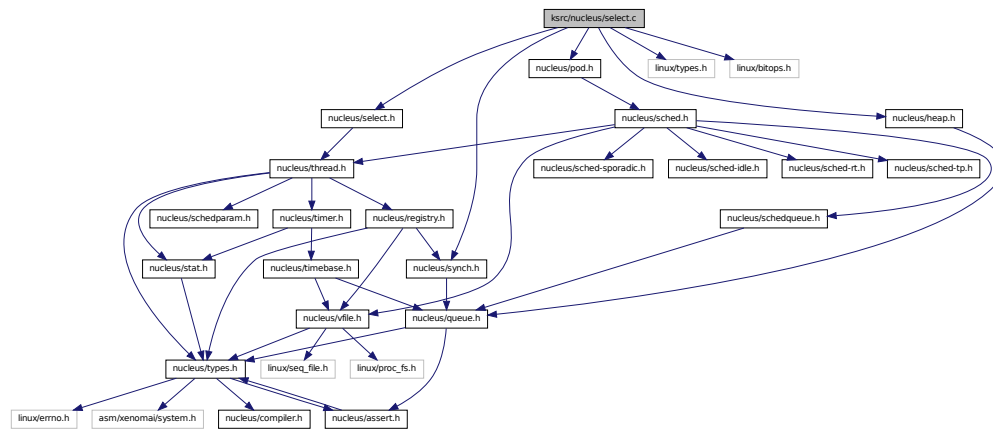
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.37 ksrc/nucleus/select.c File Reference

file descriptors events multiplexing.

Include dependency graph for select.c:



## Functions

- void [xnselect\\_init](#) (struct xnselect \*select\_block)  
*Initialize a struct xnselect structure.*
- int [xnselect\\_bind](#) (struct xnselect \*select\_block, struct xnselect\_binding \*binding, struct xnselector \*selector, unsigned type, unsigned index, unsigned state)  
*Bind a file descriptor (represented by its xnselect structure) to a selector block.*
- void [xnselect\\_destroy](#) (struct xnselect \*select\_block)  
*Destroy the xnselect structure associated with a file descriptor.*
- int [xnselector\\_init](#) (struct xnselector \*selector)  
*Initialize a selector structure.*
- int [xnselect](#) (struct xnselector \*selector, fd\_set \*out\_fds[XNSELECT\_MAX\_TYPES], fd\_set \*in\_fds[XNSELECT\_MAX\_TYPES], int nfds, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.*
- void [xnselector\\_destroy](#) (struct xnselector \*selector)  
*Destroy a selector block.*

### 6.37.1 Detailed Description

file descriptors events multiplexing.

#### Author

Gilles Chanteperdrix



Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Real-time shadow services.

- `int xnshadow_harden` (void)  
*Migrate a Linux task to the Xenomai domain.*
- `void xnshadow_relax` (int notify, int reason)  
*Switch a shadow thread back to the Linux domain.*
- `int xnshadow_map` (xnthread\_t \*thread, xncompletion\_t \_\_user \*u\_completion, unsigned long \_\_user \*u\_mode\_offset)  
*Create a shadow thread context.*
- `xnshadow_ppd_t * xnshadow_ppd_get` (unsigned muxid)  
*Return the per-process data attached to the calling process.*

Real-time shadow services.



- `xnflags_t xnsynch_sleep_on` (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Sleep on an ownerless synchronization object.*
- `struct xnthread * xnsynch_wakeup_one_sleeper` (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*
- `struct xnpholder * xnsynch_wakeup_this_sleeper` (struct xnsynch \*synch, struct xnpholder \*holder)  
*Give the resource ownership to a given waiting thread.*
- `xnflags_t xnsynch_acquire` (struct xnsynch \*synch, xnticks\_t timeout, xntmode\_t timeout\_mode)  
*Acquire the ownership of a synchronization object.*
- `static void xnsynch_clear_boost` (struct xnsynch \*synch, struct xnthread \*owner)  
*Clear the priority boost.*
- `void xnsynch_requeue_sleeper` (struct xnthread \*thread)  
*Change a sleeper's priority.*
- `struct xnthread * xnsynch_release` (struct xnsynch \*synch)  
*Give the resource ownership to the next waiting thread.*
- `struct xnthread * xnsynch_peek_pendq` (struct xnsynch \*synch)  
*Access the thread leading a synch object wait queue.*
- `int xnsynch_flush` (struct xnsynch \*synch, xnflags\_t reason)  
*Unblock all waiters pending on a resource.*
- `void xnsynch_forget_sleeper` (struct xnthread \*thread)  
*Abort a wait for a resource.*
- `void xnsynch_release_all_ownerships` (struct xnthread \*thread)  
*Release all ownerships.*

### 6.39.1 Detailed Description

Thread synchronization services.

#### Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.40 ksrc/nucleus/timebase.c File Reference

Include dependency graph for timebase.c:



### Functions

- `int xntbase_alloc` (const char \*name, u\_long period, u\_long flags, xntbase\_t \*\*basep)  
*Allocate a time base.*
- `void xntbase_free` (xntbase\_t \*base)  
*Free a time base.*
- `int xntbase_update` (xntbase\_t \*base, u\_long period)  
*Change the period of a time base.*
- `int xntbase_switch` (const char \*name, u\_long period, xntbase\_t \*\*basep)  
*Replace a time base.*
- `void xntbase_start` (xntbase\_t \*base)  
*Start a time base.*
- `void xntbase_stop` (xntbase\_t \*base)

*Stop a time base.*

- void [xntbase\\_tick](#) (xntbase\_t \*base)  
*Announce a clock tick to a time base.*
- xnticks\_t [xntbase\\_convert](#) (xntbase\_t \*srcbase, xnticks\_t ticks, xntbase\_t \*dstbase)  
*Convert a clock value into another time base.*
- void [xntbase\\_adjust\\_time](#) (xntbase\_t \*base, xnticks\_t delta)  
*Adjust the clock time for the system.*

### 6.40.1 Detailed Description

#### Note

Copyright (C) 2006,2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.41 ksrc/nucleus/timer.c File Reference

Include dependency graph for timer.c:



## Functions

- void [xntimer\\_tick\\_aperiodic](#) (void)  
*Process a timer tick for the aperiodic master time base.*
- void [xntimer\\_tick\\_periodic](#) (xntimer\_t \*mtimer)  
*Process a timer tick for a slave periodic time base.*
- void [xntimer\\_init](#) (xntimer\_t \*timer, xntbase\_t \*base, void(\*handler)(xntimer\_t \*timer))  
*Initialize a timer object.*
- void [xntimer\\_destroy](#) (xntimer\_t \*timer)  
*Release a timer object.*
- unsigned long [xntimer\\_get\\_overruns](#) (xntimer\_t \*timer, xnticks\_t now)  
*Get the count of overruns for the last tick.*
- void [xntimer\\_freeze](#) (void)  
*Freeze all timers (from every time bases).*

### 6.41.1 Detailed Description

#### Note

Copyright (C) 2001,2002,2003,2007 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>. Copyright (C) 2004 Gilles Chantepredrix <[gilles.chantepredrix@xenomai.org](mailto:gilles.chantepredrix@xenomai.org)>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

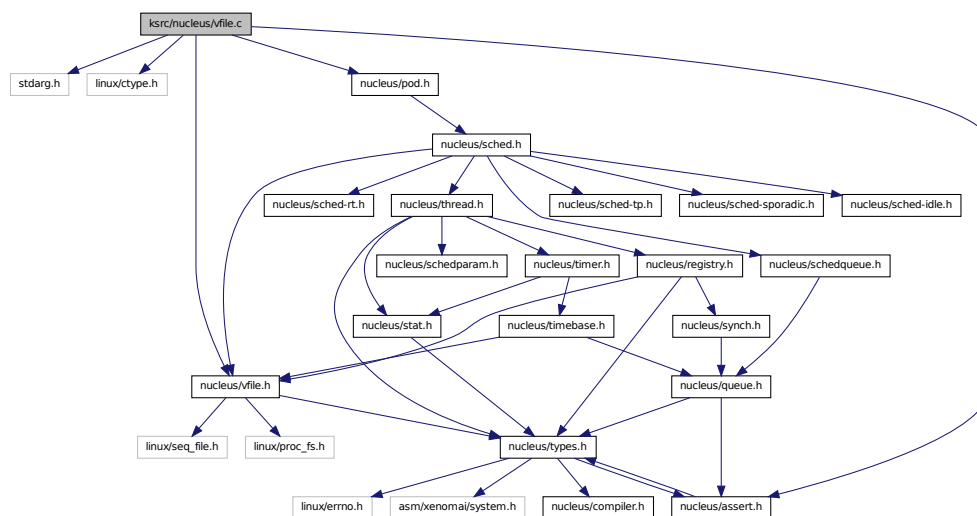
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

## 6.42 ksrc/nucleus/vfile.c File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.c:



## Functions

- `int xnvfile_init_snapshot` (const char \*name, struct `xnvfile_snapshot` \*vfile, struct `xnvfile_directory` \*parent)  
*Initialize a snapshot-driven vfile.*
- `int xnvfile_init_regular` (const char \*name, struct `xnvfile_regular` \*vfile, struct `xnvfile_directory` \*parent)  
*Initialize a regular vfile.*
- `int xnvfile_init_dir` (const char \*name, struct `xnvfile_directory` \*vdir, struct `xnvfile_directory` \*parent)  
*Initialize a virtual directory entry.*
- `int xnvfile_init_link` (const char \*from, const char \*to, struct `xnvfile_link` \*vlink, struct `xnvfile_directory` \*parent)  
*Initialize a virtual link entry.*
- `void xnvfile_destroy` (struct `xnvfile` \*vfile)  
*Removes a virtual file entry.*
- `ssize_t xnvfile_get_blob` (struct `xnvfile_input` \*input, void \*data, size\_t size)  
*Read in a data bulk written to the vfile.*
- `ssize_t xnvfile_get_string` (struct `xnvfile_input` \*input, char \*s, size\_t maxlen)  
*Read in a C-string written to the vfile.*
- `ssize_t xnvfile_get_integer` (struct `xnvfile_input` \*input, long \*valp)  
*Evaluate the string written to the vfile as a long integer.*

## Variables

- struct `xnvfile_directory` [nkvfroot](#)  
*Xenomai vfile root directory.*

### 6.42.1 Detailed Description

This file is part of the Xenomai project.

#### Note

Copyright (C) 2010 Philippe Gerum <[rpm@xenomai.org](mailto:rpm@xenomai.org)>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.



# Index

- affinity
  - xnthread\_info, [119](#)
- begin
  - xnvfile\_regular\_ops, [123](#)
  - xnvfile\_snapshot\_ops, [129](#)
- bprio
  - xnthread\_info, [119](#)
- bufd
  - xnbufd\_copy\_from\_kmem, [14](#)
  - xnbufd\_copy\_to\_kmem, [15](#)
  - xnbufd\_invalidate, [16](#)
  - xnbufd\_map\_kread, [16](#)
  - xnbufd\_map\_kwrite, [17](#)
  - xnbufd\_map\_uread, [17](#)
  - xnbufd\_map\_uwrite, [18](#)
  - xnbufd\_reset, [18](#)
  - xnbufd\_unmap\_kread, [19](#)
  - xnbufd\_unmap\_kwrite, [19](#)
  - xnbufd\_unmap\_uread, [19](#)
  - xnbufd\_unmap\_uwrite, [20](#)
- Buffer descriptors., [11](#)
- cprio
  - xnthread\_info, [119](#)
- cpu
  - xnthread\_info, [119](#)
- ctxswitches
  - xnthread\_info, [120](#)
- curr
  - xnsched, [117](#)
- databuf
  - xnvfile\_snapshot\_iterator, [128](#)
- Dynamic memory allocation services., [21](#)
- end
  - xnvfile\_regular\_ops, [123](#)
  - xnvfile\_snapshot\_ops, [129](#)
- endfn
  - xnvfile\_snapshot\_iterator, [128](#)
- exectime
  - xnthread\_info, [120](#)
- File descriptors events multiplexing services., [67](#)
- get
  - xnvfile\_lock\_ops, [121](#)
- hal
  - rthal\_apc\_alloc, [107](#)
  - rthal\_apc\_free, [107](#)
  - rthal\_irq\_disable, [108](#)
  - rthal\_irq\_enable, [108](#)
  - rthal\_irq\_host\_release, [109](#)
  - rthal\_irq\_host\_request, [109](#)
  - rthal\_irq\_release, [110](#)
  - rthal\_irq\_request, [110](#)
  - rthal\_timer\_release, [111](#)
  - rthal\_timer\_request, [112](#)
  - rthal\_trap\_catch, [112](#)
- HAL., [105](#)
- heap
  - xnheap\_alloc, [22](#)
  - xnheap\_destroy, [23](#)
  - xnheap\_extend, [23](#)
  - xnheap\_free, [24](#)
  - xnheap\_init, [24](#)
  - xnheap\_schedule\_free, [25](#)
  - xnheap\_set\_label, [26](#)
  - xnheap\_test\_and\_free, [26](#)
- htimer
  - xnsched, [117](#)
- include/nucleus/bufd.h, [133](#)
- include/nucleus/hostrt.h, [135](#)
- include/nucleus/map.h, [136](#)
- include/nucleus/pod.h, [138](#)
- include/nucleus/registry.h, [140](#)
- include/nucleus/sched-idle.h, [142](#)
- include/nucleus/sched-rt.h, [143](#)
- include/nucleus/sched-sporadic.h, [143](#)
- include/nucleus/sched-tp.h, [144](#)
- include/nucleus/sched.h, [145](#)
- include/nucleus/select.h, [146](#)
- include/nucleus/timebase.h, [148](#)
- include/nucleus/timer.h, [150](#)
- include/nucleus/vdso.h, [151](#)
- include/nucleus/vfile.h, [152](#)
- inesting
  - xnsched, [118](#)

Interrupt management., 27

intr

- xnintr\_affinity, 28
- xnintr\_attach, 29
- xnintr\_destroy, 29
- xnintr\_detach, 30
- xnintr\_disable, 30
- xnintr\_enable, 31
- xnintr\_init, 31

ksrc/arch/arm/hal.c, 155

ksrc/arch/blackfin/hal.c, 156

ksrc/arch/generic/hal.c, 157

ksrc/arch/nios2/hal.c, 157

ksrc/arch/powerpc/hal.c, 158

ksrc/arch/sh/hal.c, 160

ksrc/arch/x86/hal-common.c, 161

ksrc/arch/x86/hal\_32.c, 162

ksrc/arch/x86/hal\_64.c, 163

ksrc/arch/x86/smi.c, 163

ksrc/nucleus/bufd.c, 164

ksrc/nucleus/heap.c, 165

ksrc/nucleus/intr.c, 166

ksrc/nucleus/map.c, 168

ksrc/nucleus/pod.c, 169

ksrc/nucleus/registry.c, 171

ksrc/nucleus/sched-idle.c, 173

ksrc/nucleus/sched-rt.c, 174

ksrc/nucleus/sched-sporadic.c, 175

ksrc/nucleus/sched-tp.c, 175

ksrc/nucleus/sched.c, 177

ksrc/nucleus/select.c, 177

ksrc/nucleus/shadow.c, 179

ksrc/nucleus/synch.c, 180

ksrc/nucleus/timebase.c, 182

ksrc/nucleus/timer.c, 183

ksrc/nucleus/vfile.c, 184

lflags

- xnsched, 118

Lightweight key-to-object mapping service, 33

map

- xnmap\_create, 34
- xnmap\_delete, 35
- xnmap\_enter, 35
- xnmap\_fetch, 36
- xnmap\_fetch\_nocheck, 36
- xnmap\_remove, 37

modeswitches

- xnthread\_info, 120

name

- xnthread\_info, 120

next

- xnvfile\_regular\_ops, 124
- xnvfile\_snapshot\_ops, 130

nkvfroot

- vfile, 105

nrdata

- xnvfile\_snapshot\_iterator, 128

nucleus\_state\_flags

- XNHELD, 9
- XNLOCK, 9
- XNMIGRATE, 9
- XNPEND, 9
- XNREADY, 9
- XNSUSP, 9

pagefaults

- xnthread\_info, 120

pod

- xnpod\_abort\_thread, 42
- xnpod\_add\_hook, 42
- xnpod\_delete\_thread, 43
- xnpod\_disable\_timesource, 44
- xnpod\_dispatch\_signals, 44
- xnpod\_enable\_timesource, 44
- xnpod\_init, 45
- xnpod\_init\_thread, 46
- xnpod\_migrate\_thread, 47
- xnpod\_remove\_hook, 48
- xnpod\_restart\_thread, 48
- xnpod\_resume\_thread, 49
- xnpod\_schedule, 50
- xnpod\_set\_thread\_mode, 51
- xnpod\_set\_thread\_periodic, 52
- xnpod\_set\_thread\_schedparam, 53
- xnpod\_set\_thread\_tslice, 53
- xnpod\_shutdown, 54
- xnpod\_start\_thread, 55
- xnpod\_stop\_thread, 56
- xnpod\_suspend\_thread, 56
- xnpod\_trap\_fault, 58
- xnpod\_unblock\_thread, 58
- xnpod\_wait\_thread\_period, 59
- xnpod\_welcome\_thread, 60

pos

- xnvfile\_regular\_iterator, 122

private

- xnvfile\_regular\_iterator, 122
- xnvfile\_snapshot\_iterator, 128

put

- xnvfile\_lock\_ops, 121

Real-time pod services., 39

Real-time shadow services., 70

refcnt

- xnpod, 116
- registry
  - xnregistry\_bind, 61
  - xnregistry\_enter, 62
  - xnregistry\_fetch, 63
  - xnregistry\_get, 64
  - xnregistry\_put, 64
  - xnregistry\_remove, 65
  - xnregistry\_remove\_safe, 65
- Registry services., 60
- relpoint
  - xnthread\_info, 120
- rev
  - xnvfile\_rev\_tag, 126
- rewind
  - xnvfile\_regular\_ops, 124
  - xnvfile\_snapshot\_ops, 130
- rootcb
  - xnsched, 118
- rt
  - xnsched, 118
- rthal\_apc\_alloc
  - hal, 107
- rthal\_apc\_free
  - hal, 107
- rthal\_irq\_disable
  - hal, 108
- rthal\_irq\_enable
  - hal, 108
- rthal\_irq\_host\_release
  - hal, 109
- rthal\_irq\_host\_request
  - hal, 109
- rthal\_irq\_release
  - hal, 110
- rthal\_irq\_request
  - hal, 110
- rthal\_timer\_release
  - hal, 111
- rthal\_timer\_request
  - hal, 112
- rthal\_trap\_catch
  - hal, 112
- Sched, 113
- sched
  - xnpod, 116
  - xnsched\_rotate, 114
- select
  - xnselect, 68
  - xnselect\_bind, 68
  - xnselect\_destroy, 69
  - xnselect\_init, 69
  - xnselector\_destroy, 69
  - xnselector\_init, 70
- seq
  - xnvfile\_regular\_iterator, 122
  - xnvfile\_snapshot\_iterator, 128
- shadow
  - xnshadow\_harden, 71
  - xnshadow\_map, 71
  - xnshadow\_ppd\_get, 72
  - xnshadow\_relax, 73
- show
  - xnvfile\_regular\_ops, 124
  - xnvfile\_snapshot\_ops, 131
- state
  - xnthread\_info, 120
- status
  - xnpod, 116
  - xnsched, 118
- store
  - xnvfile\_regular\_ops, 125
  - xnvfile\_snapshot\_ops, 131
- synch
  - xnsynch\_acquire, 75
  - xnsynch\_clear\_boost, 76
  - xnsynch\_flush, 76
  - xnsynch\_forget\_sleeper, 77
  - xnsynch\_init, 77
  - xnsynch\_peek\_pendq, 78
  - xnsynch\_release, 79
  - xnsynch\_release\_all\_ownerships, 80
  - xnsynch\_requeue\_sleeper, 80
  - xnsynch\_sleep\_on, 80
  - xnsynch\_wakeup\_one\_sleeper, 81
  - xnsynch\_wakeup\_this\_sleeper, 82
- tdeleteq
  - xnpod, 116
- Thread information flags., 10
- Thread state flags., 7
- Thread synchronization services., 74
- threadq
  - xnpod, 116
- Time base services., 83
- timebase
  - xntbase\_adjust\_time, 84
  - xntbase\_alloc, 85
  - xntbase\_convert, 86
  - xntbase\_free, 86
  - xntbase\_get\_time, 87
  - xntbase\_start, 87
  - xntbase\_stop, 88
  - xntbase\_switch, 88
  - xntbase\_tick, 89
  - xntbase\_update, 90
- timer

- xntimer\_destroy, [92](#)
- xntimer\_freeze, [92](#)
- xntimer\_get\_date, [93](#)
- xntimer\_get\_interval, [93](#)
- xntimer\_get\_overruns, [94](#)
- xntimer\_get\_timeout, [94](#)
- xntimer\_init, [95](#)
- xntimer\_start, [96](#)
- xntimer\_stop, [97](#)
- xntimer\_tick\_aperiodic, [97](#)
- xntimer\_tick\_periodic, [98](#)
- Timer services., [91](#)
- timerlck
  - xnpod, [116](#)
- tsliced
  - xnpod, [116](#)
- tslicer
  - xnpod, [116](#)
- tstartq
  - xnpod, [117](#)
- tswitchq
  - xnpod, [117](#)
- vfile
  - nkvfroot, [105](#)
  - xnvfile\_destroy, [101](#)
  - xnvfile\_get\_blob, [101](#)
  - xnvfile\_get\_integer, [101](#)
  - xnvfile\_get\_string, [102](#)
  - xnvfile\_init\_dir, [102](#)
  - xnvfile\_init\_link, [103](#)
  - xnvfile\_init\_regular, [103](#)
  - xnvfile\_init\_snapshot, [104](#)
  - xnvfile\_regular\_iterator, [122](#)
  - xnvfile\_snapshot\_iterator, [128](#)
- Virtual file services, [98](#)
- Xenomai nucleus., [38](#)
- xnbufd\_copy\_from\_kmem
  - bufd, [14](#)
- xnbufd\_copy\_to\_kmem
  - bufd, [15](#)
- xnbufd\_invalidate
  - bufd, [16](#)
- xnbufd\_map\_kread
  - bufd, [16](#)
- xnbufd\_map\_kwrite
  - bufd, [17](#)
- xnbufd\_map\_uread
  - bufd, [17](#)
- xnbufd\_map\_uwrite
  - bufd, [18](#)
- xnbufd\_reset
  - bufd, [18](#)
- xnbufd\_unmap\_kread
  - bufd, [19](#)
- xnbufd\_unmap\_kwrite
  - bufd, [19](#)
- xnbufd\_unmap\_uread
  - bufd, [19](#)
- xnbufd\_unmap\_uwrite
  - bufd, [20](#)
- xnheap\_alloc
  - heap, [22](#)
- xnheap\_destroy
  - heap, [23](#)
- xnheap\_extend
  - heap, [23](#)
- xnheap\_free
  - heap, [24](#)
- xnheap\_init
  - heap, [24](#)
- xnheap\_schedule\_free
  - heap, [25](#)
- xnheap\_set\_label
  - heap, [26](#)
- xnheap\_test\_and\_free
  - heap, [26](#)
- XNHeld
  - nucleus\_state\_flags, [9](#)
- xnintr\_affinity
  - intr, [28](#)
- xnintr\_attach
  - intr, [29](#)
- xnintr\_destroy
  - intr, [29](#)
- xnintr\_detach
  - intr, [30](#)
- xnintr\_disable
  - intr, [30](#)
- xnintr\_enable
  - intr, [31](#)
- xnintr\_init
  - intr, [31](#)
- XNLOCK
  - nucleus\_state\_flags, [9](#)
- xnmap\_create
  - map, [34](#)
- xnmap\_delete
  - map, [35](#)
- xnmap\_enter
  - map, [35](#)
- xnmap\_fetch
  - map, [36](#)
- xnmap\_fetch\_nocheck
  - map, [36](#)
- xnmap\_remove
  - map, [37](#)

- XNMIGRATE
  - nucleus\_state\_flags, 9
- XNPEND
  - nucleus\_state\_flags, 9
- xnpod, 115
  - refcnt, 116
  - sched, 116
  - status, 116
  - tdeleteq, 116
  - threadq, 116
  - timerlck, 116
  - tsliced, 116
  - tslicer, 116
  - tstartq, 117
  - tswitchq, 117
- xnpod\_abort\_thread
  - pod, 42
- xnpod\_add\_hook
  - pod, 42
- xnpod\_delete\_thread
  - pod, 43
- xnpod\_disable\_timesource
  - pod, 44
- xnpod\_dispatch\_signals
  - pod, 44
- xnpod\_enable\_timesource
  - pod, 44
- xnpod\_init
  - pod, 45
- xnpod\_init\_thread
  - pod, 46
- xnpod\_migrate\_thread
  - pod, 47
- xnpod\_remove\_hook
  - pod, 48
- xnpod\_restart\_thread
  - pod, 48
- xnpod\_resume\_thread
  - pod, 49
- xnpod\_schedule
  - pod, 50
- xnpod\_set\_thread\_mode
  - pod, 51
- xnpod\_set\_thread\_periodic
  - pod, 52
- xnpod\_set\_thread\_schedparam
  - pod, 53
- xnpod\_set\_thread\_tslice
  - pod, 53
- xnpod\_shutdown
  - pod, 54
- xnpod\_start\_thread
  - pod, 55
- xnpod\_stop\_thread
  - pod, 56
- xnpod\_suspend\_thread
  - pod, 56
- xnpod\_trap\_fault
  - pod, 58
- xnpod\_unblock\_thread
  - pod, 58
- xnpod\_wait\_thread\_period
  - pod, 59
- xnpod\_welcome\_thread
  - pod, 60
- XNREADY
  - nucleus\_state\_flags, 9
- xnregistry\_bind
  - registry, 61
- xnregistry\_enter
  - registry, 62
- xnregistry\_fetch
  - registry, 63
- xnregistry\_get
  - registry, 64
- xnregistry\_put
  - registry, 64
- xnregistry\_remove
  - registry, 65
- xnregistry\_remove\_safe
  - registry, 65
- xnsched, 117
  - curr, 117
  - htimer, 117
  - inesting, 118
  - lflags, 118
  - rootcb, 118
  - rt, 118
  - status, 118
- xnsched\_rotate
  - sched, 114
- xnselect
  - select, 68
- xnselect\_bind
  - select, 68
- xnselect\_destroy
  - select, 69
- xnselect\_init
  - select, 69
- xnselector\_destroy
  - select, 69
- xnselector\_init
  - select, 70
- xnshadow\_harden
  - shadow, 71
- xnshadow\_map
  - shadow, 71
- xnshadow\_ppd\_get

- shadow, [72](#)
- xnshadow\_relax
  - shadow, [73](#)
- XNSUSP
  - nucleus\_state\_flags, [9](#)
- xnsynch\_acquire
  - synch, [75](#)
- xnsynch\_clear\_boost
  - synch, [76](#)
- xnsynch\_flush
  - synch, [76](#)
- xnsynch\_forget\_sleeper
  - synch, [77](#)
- xnsynch\_init
  - synch, [77](#)
- xnsynch\_peek\_pendq
  - synch, [78](#)
- xnsynch\_release
  - synch, [79](#)
- xnsynch\_release\_all\_ownerships
  - synch, [80](#)
- xnsynch\_requeue\_sleeper
  - synch, [80](#)
- xnsynch\_sleep\_on
  - synch, [80](#)
- xnsynch\_wakeup\_one\_sleeper
  - synch, [81](#)
- xnsynch\_wakeup\_this\_sleeper
  - synch, [82](#)
- xntbase\_adjust\_time
  - timebase, [84](#)
- xntbase\_alloc
  - timebase, [85](#)
- xntbase\_convert
  - timebase, [86](#)
- xntbase\_free
  - timebase, [86](#)
- xntbase\_get\_time
  - timebase, [87](#)
- xntbase\_start
  - timebase, [87](#)
- xntbase\_stop
  - timebase, [88](#)
- xntbase\_switch
  - timebase, [88](#)
- xntbase\_tick
  - timebase, [89](#)
- xntbase\_update
  - timebase, [90](#)
- xnthread\_info, [118](#)
  - affinity, [119](#)
  - bprio, [119](#)
  - cprio, [119](#)
  - cpu, [119](#)
  - ctxswitches, [120](#)
  - exectime, [120](#)
  - modeswitches, [120](#)
  - name, [120](#)
  - pagefaults, [120](#)
  - relpoint, [120](#)
  - state, [120](#)
- xntimer\_destroy
  - timer, [92](#)
- xntimer\_freeze
  - timer, [92](#)
- xntimer\_get\_date
  - timer, [93](#)
- xntimer\_get\_interval
  - timer, [93](#)
- xntimer\_get\_overruns
  - timer, [94](#)
- xntimer\_get\_timeout
  - timer, [94](#)
- xntimer\_init
  - timer, [95](#)
- xntimer\_start
  - timer, [96](#)
- xntimer\_stop
  - timer, [97](#)
- xntimer\_tick\_aperiodic
  - timer, [97](#)
- xntimer\_tick\_periodic
  - timer, [98](#)
- xnvfile\_destroy
  - vfile, [101](#)
- xnvfile\_get\_blob
  - vfile, [101](#)
- xnvfile\_get\_integer
  - vfile, [101](#)
- xnvfile\_get\_string
  - vfile, [102](#)
- xnvfile\_init\_dir
  - vfile, [102](#)
- xnvfile\_init\_link
  - vfile, [103](#)
- xnvfile\_init\_regular
  - vfile, [103](#)
- xnvfile\_init\_snapshot
  - vfile, [104](#)
- xnvfile\_lock\_ops, [120](#)
  - get, [121](#)
  - put, [121](#)
- xnvfile\_regular\_iterator, [121](#)
  - pos, [122](#)
  - private, [122](#)
  - seq, [122](#)
  - vfile, [122](#)
- xnvfile\_regular\_ops, [122](#)

- [begin, 123](#)
  - [end, 123](#)
  - [next, 124](#)
  - [rewind, 124](#)
  - [show, 124](#)
  - [store, 125](#)
- [xnvfile\\_rev\\_tag, 126](#)
  - [rev, 126](#)
- [xnvfile\\_snapshot, 126](#)
- [xnvfile\\_snapshot\\_iterator, 127](#)
  - [databuf, 128](#)
  - [endfn, 128](#)
  - [nrdata, 128](#)
  - [private, 128](#)
  - [seq, 128](#)
  - [vfile, 128](#)
- [xnvfile\\_snapshot\\_ops, 129](#)
  - [begin, 129](#)
  - [end, 129](#)
  - [next, 130](#)
  - [rewind, 130](#)
  - [show, 131](#)
  - [store, 131](#)