

Xenomai nanokernel API Reference Manual

2.3.50

Generated by Doxygen 1.4.6

Mon Jan 1 20:33:25 2007

Contents

1	Xenomai nanokernel API Module Index	1
1.1	Xenomai nanokernel API Modules	1
2	Xenomai nanokernel API Hierarchical Index	3
2.1	Xenomai nanokernel API Class Hierarchy	3
3	Xenomai nanokernel API Data Structure Index	5
3.1	Xenomai nanokernel API Data Structures	5
4	Xenomai nanokernel API File Index	7
4.1	Xenomai nanokernel API File List	7
5	Xenomai nanokernel API Module Documentation	9
5.1	Thread state flags.	9
5.2	Thread information flags.	11
5.3	Dynamic memory allocation services.	12
5.4	Interrupt management.	18
5.5	Xenomai nucleus.	24
5.6	Real-time pod services.	25
5.7	Registry services.	48
5.8	Real-time shadow services.	54
5.9	Thread synchronization services.	57
5.10	Time base services.	63
5.11	Timer services.	68
5.12	HAL.	71
6	Xenomai nanokernel API Data Structure Documentation	81
6.1	xnpod Struct Reference	81
6.2	xnsched Struct Reference	84

7 Xenomai nanokernel API File Documentation	85
7.1 include/nucleus/pod.h File Reference	85
7.2 include/nucleus/registry.h File Reference	88
7.3 include/nucleus/timer.h File Reference	89
7.4 ksrc/arch/arm/hal.c File Reference	94
7.5 ksrc/arch/blackfin/hal.c File Reference	96
7.6 ksrc/arch/generic/hal.c File Reference	98
7.7 ksrc/arch/i386/hal.c File Reference	100
7.8 ksrc/arch/ia64/hal.c File Reference	102
7.9 ksrc/arch/powerpc/hal.c File Reference	104
7.10 ksrc/arch/blackfin/nmi.c File Reference	106
7.11 ksrc/arch/generic/nmi.c File Reference	107
7.12 ksrc/arch/i386/nmi.c File Reference	108
7.13 ksrc/arch/i386/smi.c File Reference	109
7.14 ksrc/nucleus/heap.c File Reference	110
7.15 ksrc/nucleus/intr.c File Reference	112
7.16 ksrc/nucleus/pod.c File Reference	114
7.17 ksrc/nucleus/registry.c File Reference	117
7.18 ksrc/nucleus/shadow.c File Reference	118
7.19 ksrc/nucleus/synch.c File Reference	120
7.20 ksrc/nucleus/timebase.c File Reference	122
7.21 ksrc/nucleus/timer.c File Reference	123

Chapter 1

Xenomai nanokernel API Module Index

1.1 Xenomai nanokernel API Modules

Here is a list of all modules:

Xenomai nucleus.	24
Thread state flags.	9
Thread information flags.	11
Dynamic memory allocation services.	12
Interrupt management.	18
Real-time pod services.	25
Registry services.	48
Real-time shadow services.	54
Thread synchronization services.	57
Time base services.	63
Timer services.	68
HAL.	71

Chapter 2

Xenomai nanokernel API Hierarchical Index

2.1 Xenomai nanokernel API Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

xnpod	81
xnsched	84

Chapter 3

Xenomai nanokernel API Data Structure Index

3.1 Xenomai nanokernel API Data Structures

Here are the data structures with brief descriptions:

xnpod (Real-time pod descriptor)	81
xnsched (Scheduling information structure)	84

Chapter 4

Xenomai nanokernel API File Index

4.1 Xenomai nanokernel API File List

Here is a list of all documented files with brief descriptions:

include/nucleus/assert.h	??
include/nucleus/bheap.h	??
include/nucleus/compiler.h	??
include/nucleus/core.h	??
include/nucleus/heap.h	??
include/nucleus/intr.h	??
include/nucleus/jhash.h	??
include/nucleus/ltt.h	??
include/nucleus/module.h	??
include/nucleus/pipe.h	??
include/nucleus/pod.h (Real-time pod interface header)	85
include/nucleus/ppd.h	??
include/nucleus/queue.h	??
include/nucleus/registry.h (This file is part of the Xenomai project)	88
include/nucleus/shadow.h	??
include/nucleus/stat.h	??
include/nucleus/synch.h	??
include/nucleus/system.h	??
include/nucleus/thread.h	??
include/nucleus/timebase.h	??
include/nucleus/timer.h	89
include/nucleus/trace.h	??
include/nucleus/types.h	??
include/nucleus/version.h	??
include/nucleus/xenomai.h	??
ksrc/arch/arm/hal.c (Adeos-based Real-Time Abstraction Layer for PowerPC)	94
ksrc/arch/blackfin/hal.c (Adeos-based Real-Time Abstraction Layer for the Blackfin architecture)	96
ksrc/arch/blackfin/nmi.c (NMI watchdog support)	106
ksrc/arch/generic/hal.c (Generic Real-Time HAL)	98
ksrc/arch/generic/nmi.c (Adeos-based Real-Time Abstraction Layer for x86)	107
ksrc/arch/i386/hal.c (Adeos-based Real-Time Abstraction Layer for x86)	100
ksrc/arch/i386/nmi.c (NMI watchdog for x86, from linux/arch/i386/kernel/nmi.c)	108

ksrc/arch/i386/smi.c (SMI workaround for x86)	109
ksrc/arch/ia64/hal.c (Adeos-based Real-Time Abstraction Layer for ia64)	102
ksrc/arch/powerpc/hal.c (Adeos-based Real-Time Abstraction Layer for PowerPC)	104
ksrc/nucleus/heap.c (Dynamic memory allocation services)	110
ksrc/nucleus/intr.c (Interrupt management)	112
ksrc/nucleus/pod.c (Real-time pod services)	114
ksrc/nucleus/registry.c (This file is part of the Xenomai project)	117
ksrc/nucleus/shadow.c (Real-time shadow services)	118
ksrc/nucleus/synch.c (Thread synchronization services)	120
ksrc/nucleus/timebase.c	122
ksrc/nucleus/timer.c	123

Chapter 5

Xenomai nanokernel API Module Documentation

5.1 Thread state flags.

5.1.1 Detailed Description

Bits reporting permanent or transient states of thread.

Defines

- #define `XNSUSP` 0x00000001
Suspended.
- #define `XNPEND` 0x00000002
Sleep-wait for a resource.
- #define `XNDELAY` 0x00000004
Delayed.
- #define `XNREADY` 0x00000008
Linked to the ready queue.
- #define `XNDORMANT` 0x00000010
Not started yet or killed.
- #define `XNZOMBIE` 0x00000020
Zombie thread in deletion process.
- #define `XNRESTART` 0x00000040
Restarting thread.
- #define `XNSTARTED` 0x00000080
Thread has been started.

- #define [XNMAPPED](#) 0x00000100
Mapped to regular Linux task (shadow only).
- #define [XNRELAX](#) 0x00000200
Relaxed shadow thread (blocking bit).
- #define [XNHELD](#) 0x00000400
Held thread from suspended partition.
- #define [XNBOOST](#) 0x00000800
Undergoes regular PIP boost.
- #define [XNDEBUG](#) 0x00001000
Hit debugger breakpoint (shadow only).
- #define [XNLOCK](#) 0x00002000
Not preemptible.
- #define [XNRRB](#) 0x00004000
Undergoes a round-robin scheduling.
- #define [XNASDI](#) 0x00008000
ASR are disabled.
- #define [XNSHIELD](#) 0x00010000
IRQ shield is enabled (shadow only).
- #define [XNTRAPSW](#) 0x00020000
Trap execution mode switches.
- #define [XNRPIOFF](#) 0x00040000
Stop priority coupling (shadow only).
- #define [XNFPU](#) 0x00100000
Thread uses FPU.
- #define [XNSHADOW](#) 0x00200000
Shadow thread.
- #define [XNROOT](#) 0x00400000
Root thread (that is, Linux/IDLE).
- #define [XNINVPS](#) 0x00800000
Using inverted priority scale.

5.2 Thread information flags.

5.2.1 Detailed Description

Bits reporting events notified to thread.

Defines

- #define `XNTIMEO` 0x00000001
Woken up due to a timeout condition.
- #define `XNRMID` 0x00000002
Pending on a removed resource.
- #define `XNBREAK` 0x00000004
Forcibly awoken from a wait state.
- #define `XNKICKED` 0x00000008
Kicked upon Linux signal (shadow only).
- #define `XNWAKEN` 0x00000010
Thread waken up upon resource availability.
- #define `XNROBBED` 0x00000020
Robbed from resource ownership.

5.3 Dynamic memory allocation services.

5.3.1 Detailed Description

Dynamic memory allocation services.

The implementation of the memory allocator follows the algorithm described in a USENIX 1988 paper called "Design of a General Purpose Memory Allocator for the 4.3BSD Unix Kernel" by Marshall K. McKusick and Michael J. Karels. You can find it at various locations on the net, including <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf>. A minor variation allows this implementation to have 'extendable' heaps when needed, with multiple memory extents providing autonomous page address spaces.

The data structures hierarchy is as follows:

```

HEAP {
    block_buckets[]
    extent_queue -----+
}
                        |
                        V
EXTENT #1 {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
} --+
    |
    |
    V
EXTENT #n {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
}

```

Files

- file [heap.c](#)
Dynamic memory allocation services.

Functions

- int [xnheap_init](#) (xnheap_t *heap, void *heapaddr, u_long heapsize, u_long pagesize)
Initialize a memory heap.
- int [xnheap_destroy](#) (xnheap_t *heap, void(*flushfn)(xnheap_t *heap, void *extaddr, u_long extsize, void *cookie), void *cookie)
Destroys a memory heap.
- void * [xnheap_alloc](#) (xnheap_t *heap, u_long size)
Allocate a memory block from a memory heap.
- int [xnheap_test_and_free](#) (xnheap_t *heap, void *block, int(*ckfn)(void *block))
Test and release a memory block to a memory heap.

- `int xnheap_free (xnheap_t *heap, void *block)`
Release a memory block to a memory heap.
- `int xnheap_extend (xnheap_t *heap, void *extaddr, u_long extsize)`
Extend a memory heap.
- `void xnheap_schedule_free (xnheap_t *heap, void *block, xnholder_t *link)`
Schedule a memory block for release.

5.3.2 Function Documentation

5.3.2.1 `void* xnheap_alloc (xnheap_t * heap, u_long size)`

Allocate a memory block from a memory heap.

Allocates a contiguous region of memory from an active memory heap. Such allocation is guaranteed to be time-bounded.

Parameters:

heap The descriptor address of the heap to get memory from.

size The size in bytes of the requested block. Sizes lower or equal to the page size are rounded either to the minimum allocation size if lower than this value, or to the minimum alignment size if greater or equal to this value. In the current implementation, with MINALLOC = 8 and MINALIGN = 16, a 7 bytes request will be rounded to 8 bytes, and a 17 bytes request will be rounded to 32.

Returns:

The address of the allocated region upon success, or NULL if no memory is available from the specified heap.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.2 `int xnheap_destroy (xnheap_t * heap, void(*) (xnheap_t *heap, void *extaddr, u_long extsize, void *cookie) flushfn, void * cookie)`

Destroys a memory heap.

Destroys a memory heap.

Parameters:

heap The descriptor address of the destroyed heap.

flushfn If non-NULL, the address of a flush routine which will be called for each extent attached to the heap. This routine can be used by the calling code to further release the heap memory.

cookie If *flushfn* is non-NULL, *cookie* is an opaque pointer which will be passed unmodified to *flushfn*.

Returns:

0 is returned on success, or -EBUSY if external mappings are still pending on the heap memory.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.3 int xnheap_extend(xnheap_t * heap, void * extaddr, u_long extsize)

Extend a memory heap.

Add a new extent to an existing memory heap.

Parameters:

heap The descriptor address of the heap to add an extent to.

extaddr The address of the extent memory.

extsize The size of the extent memory (in bytes). In the current implementation, this size must match the one of the initial extent passed to [xnheap_init\(\)](#).

Returns:

0 is returned upon success, or -EINVAL is returned if *extsize* differs from the initial extent's size.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.4 int xnheap_free (xnheap_t * heap, void * block)

Release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from.

Parameters:

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

Returns:

0 is returned upon success, or one of the following error codes:

- -EFAULT is returned whenever the memory address is outside the heap address space.
- -EINVAL is returned whenever the memory address does not represent a valid block.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.5 int xnheap_init (xnheap_t * heap, void * heapaddr, u_long heapsize, u_long pagesize)

Initialize a memory heap.

Initializes a memory heap suitable for time-bounded allocation requests of dynamic memory.

Parameters:

heap The address of a heap descriptor which will be used to store the allocation data. This descriptor must always be valid while the heap is active therefore it must be allocated in permanent memory.

heapaddr The address of the heap storage area. All allocations will be made from the given area in time-bounded mode. Since additional extents can be added to a heap, this parameter is also known as the "initial extent".

heapsize The size in bytes of the initial extent pointed at by *heapaddr*. *heapsize* must be a multiple of *pagesize* and lower than 16 Mbytes. *heapsize* must be large enough to contain an internal header. The following formula gives the size of this header:

$$\text{hdrsize} = (\text{sizeof}(\text{xnextent_t}) + ((\text{heapsize} - \text{sizeof}(\text{xnextent_t})) / (\text{pagesize} + 1) + 15) \& \sim 15.$$

pagesize The size in bytes of the fundamental memory page which will be used to subdivide the heap internally. Choosing the right page size is important regarding performance and memory fragmentation issues, so it might be a good idea to take a look at <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf> to pick the best one for your needs. In the current implementation, *pagesize* must be a power of two in the range [8 .. 32768] inclusive.

Returns:

0 is returned upon success, or one of the following error codes:

- -EINVAL is returned whenever a parameter is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.6 void xnheap_schedule_free (xnheap_t * heap, void * block, xnholder_t * link)

Schedule a memory block for release.

This routine records a block for later release by `xnheap_finalize_free()`. This service is useful to lazily free blocks of heap memory when immediate release is not an option, e.g. when active references are still pending on the object for a short time after the call. `xnheap_finalize_free()` is expected to be eventually called by the client code at some point in the future when actually freeing the idle objects is deemed safe.

Parameters:

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

link The address of a link member, likely but not necessarily within the released object, which will be used by the heap manager to hold the block in the queue of idle objects.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.3.2.7 int xnheap_test_and_free (xnheap_t * heap, void * block, int(*)(void *block) ckfn)

Test and release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from. Before the actual release is performed, an optional user-defined can be invoked to check for additional criteria with respect to the request consistency.

Parameters:

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

ckfn The address of a user-supplied verification routine which is to be called after the memory address specified by *block* has been checked for validity. The routine is expected to proceed to further consistency checks, and either return zero upon success, or non-zero upon error. In the latter case, the release process is aborted, and *ckfn*'s return value is passed back to the caller of this service as its error return code. *ckfn* must not trigger the rescheduling procedure either directly or indirectly.

Returns:

0 is returned upon success, or -EINVAL is returned whenever the block is not a valid region of the specified heap. Additional return codes can also be defined locally by the *ckfn* routine.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.4 Interrupt management.

5.4.1 Detailed Description

Interrupt management.

Files

- file [intr.c](#)
Interrupt management.

Functions

- int [xnintr_init](#) (xnintr_t *intr, const char *name, unsigned irq, xnintr_t isr, xniack_t iack, xnflags_t flags)
Initialize an interrupt object.
- int [xnintr_destroy](#) (xnintr_t *intr)
Destroy an interrupt object.
- int [xnintr_attach](#) (xnintr_t *intr, void *cookie)
Attach an interrupt object.
- int [xnintr_detach](#) (xnintr_t *intr)
Detach an interrupt object.
- int [xnintr_enable](#) (xnintr_t *intr)
Enable an interrupt object.
- int [xnintr_disable](#) (xnintr_t *intr)
Disable an interrupt object.
- xnarch_cpumask_t [xnintr_affinity](#) (xnintr_t *intr, xnarch_cpumask_t cpumask)
Set interrupt's processor affinity.

5.4.2 Function Documentation

5.4.2.1 xnarch_cpumask_t xnintr_affinity (xnintr_t * intr, xnarch_cpumask_t cpumask)

Set interrupt's processor affinity.

Causes the IRQ associated with the interrupt object *intr* to be received only on processors which bits are set in *cpumask*.

Parameters:

- intr* The descriptor address of the interrupt object which affinity is to be changed.
- cpumask* The new processor affinity of the interrupt object.

Returns:

the previous *cpumask* on success, or an empty mask on failure.

Note:

Depending on architectures, setting more than one bit in *cpumask* could be meaningless.

5.4.2.2 int xnintr_attach (xnintr_t * *intr*, void * *cookie*)

Attach an interrupt object.

Attach an interrupt object previously initialized by [xnintr_init\(\)](#). After this operation is completed, all IRQs received from the corresponding interrupt channel are directed to the object's ISR.

Parameters:

intr The descriptor address of the interrupt object to attach.

cookie A user-defined opaque value which is stored into the interrupt object descriptor for further retrieval by the ISR/ISR handlers.

Returns:

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while attaching the interrupt. -EBUSY is specifically returned if the interrupt object was already attached.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Note:

Attaching an interrupt resets the tracked number of receipts to zero.

5.4.2.3 int xnintr_destroy (xnintr_t * *intr*)

Destroy an interrupt object.

Destroys an interrupt object previously initialized by [xnintr_init\(\)](#). The interrupt object is automatically detached by a call to [xnintr_detach\(\)](#). No more IRQs will be dispatched by this object after this service has returned.

Parameters:

intr The descriptor address of the interrupt object to destroy.

Returns:

0 is returned on success. Otherwise, -EBUSY is returned if an error occurred while detaching the interrupt (see [xnintr_detach\(\)](#)).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.4.2.4 `int xnintr_detach (xnintr_t * intr)`

Detach an interrupt object.

Detach an interrupt object previously attached by [xnintr_attach\(\)](#). After this operation is completed, no more IRQs are directed to the object's ISR, but the interrupt object itself remains valid. A detached interrupt object can be attached again by a subsequent call to [xnintr_attach\(\)](#).

Parameters:

intr The descriptor address of the interrupt object to detach.

Returns:

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while detaching the interrupt. Detaching a non-attached interrupt object leads to a null-effect and returns 0.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.4.2.5 `int xnintr_disable (xnintr_t * intr)`

Disable an interrupt object.

Disables the hardware interrupt line associated with an interrupt object. This operation invalidates further interrupt requests from the given source until the IRQ line is re-enabled anew.

Parameters:

intr The descriptor address of the interrupt object to disable.

Returns:

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while disabling the interrupt.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.4.2.6 `int xnintr_enable (xnintr_t * intr)`

Enable an interrupt object.

Enables the hardware interrupt line associated with an interrupt object. Over real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

Parameters:

intr The descriptor address of the interrupt object to enable.

Returns:

0 is returned on success. Otherwise, -EINVAL is returned if a low-level error occurred while enabling the interrupt.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.4.2.7 `int xnintr_init (xnintr_t * intr, const char * name, unsigned irq, xn_isr_t isr, xn_iack_t iack, xnflags_t flags)`

Initialize an interrupt object.

Associates an interrupt object with an IRQ line.

When an interrupt occurs on the given *irq* line, the ISR is fired in order to deal with the hardware event. The interrupt service code may call any non-suspensive service from the nucleus.

Upon receipt of an IRQ, the ISR is immediately called on behalf of the interrupted stack context, the rescheduling procedure is locked, and the interrupt source is masked at hardware level. The status value returned by the ISR is then checked for the following values:

- XN_ISR_HANDLED indicates that the interrupt request has been fulfilled by the ISR.

- `XN_ISR_NONE` indicates the opposite to `XN_ISR_HANDLED`. The ISR must always return this value when it determines that the interrupt request has not been issued by the dedicated hardware device.

In addition, one of the following bits may be set by the ISR :

NOTE: use these bits with care and only when you do understand their effect on the system. The ISR is not encouraged to use these bits in case it shares the IRQ line with other ISRs in the real-time domain.

- `XN_ISR_PROPAGATE` tells the nucleus to require the real-time control layer to forward the IRQ. For instance, this would cause the Adeos control layer to propagate the interrupt down the interrupt pipeline to other Adeos domains, such as Linux. This is the regular way to share interrupts between the nucleus and the host system.
- `XN_ISR_NOENABLE` causes the nucleus to ask the real-time control layer `_not_` to re-enable the IRQ line (read the following section). `xnarch_end_irq()` must be called to re-enable the IRQ line later.

The nucleus re-enables the IRQ line by default. Over some real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

A count of interrupt receipts is tracked into the interrupt descriptor, and reset to zero each time the interrupt object is attached. Since this count could wrap around, it should be used as an indication of interrupt activity only.

Parameters:

intr The address of a interrupt object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

name An ASCII string standing for the symbolic name of the interrupt object.

irq The hardware interrupt channel associated with the interrupt object. This value is architecture-dependent. An interrupt object must then be attached to the hardware interrupt vector using the `xnintr_attach()` service for the associated IRQs to be directed to this object.

isr The address of a valid low-level interrupt service routine if this parameter is non-zero. This handler will be called each time the corresponding IRQ is delivered on behalf of an interrupt context. When called, the ISR is passed the descriptor address of the interrupt object.

iack The address of an optional interrupt acknowledge routine, aimed at replacing the default one. Only very specific situations actually require to override the default setting for this parameter, like having to acknowledge non-standard PIC hardware. *iack* should return a non-zero value to indicate that the interrupt has been properly acknowledged. If *iack* is NULL, the default routine will be used instead.

flags A set of creation flags affecting the operation. The valid flags are:

- `XN_ISR_SHARED` enables IRQ-sharing with other interrupt objects.
- `XN_ISR_EDGE` is an additional flag need to be set together with `XN_ISR_SHARED` to enable IRQ-sharing of edge-triggered interrupts.

Returns:

No error condition being defined, 0 is always returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.5 Xenomai nucleus.

5.5.1 Detailed Description

An abstract RTOS core.

Modules

- [Thread state flags.](#)
Bits reporting permanent or transient states of thread.
- [Thread information flags.](#)
Bits reporting events notified to thread.
- [Dynamic memory allocation services.](#)
- [Interrupt management.](#)
- [Real-time pod services.](#)
- [Registry services.](#)
- [Real-time shadow services.](#)
- [Thread synchronization services.](#)
- [Time base services.](#)
- [Timer services.](#)

5.6 Real-time pod services.

5.6.1 Detailed Description

Real-time pod services.

Files

- file [pod.h](#)
Real-time pod interface header.
- file [pod.c](#)
Real-time pod services.

Data Structures

- struct [xnsched](#)
Scheduling information structure.
- struct [xnpod](#)
Real-time pod descriptor.

Typedefs

- typedef [xnsched](#) [xnsched_t](#)
Scheduling information structure.

Functions

- void [xnpod_schedule_runnable](#) (xnthread_t *thread, int flags)
Hidden rescheduling procedure.
- int [xnpod_init](#) (xnpod_t *pod, int minpri, int maxpri, xnflags_t flags)
Initialize a new pod.
- int [xnpod_enable_timesource](#) (void)
Activate the core time source.
- void [xnpod_disable_timesource](#) (void)
Stop the core time source.
- void [xnpod_shutdown](#) (int xtype)
Shutdown the current pod.

- int [xnpod_init_thread](#) (xnthread_t *thread, xntbase_t *tbase, const char *name, int prio, xnflags_t flags, unsigned stacksize)
Initialize a new thread.
- int [xnpod_start_thread](#) (xnthread_t *thread, xnflags_t mode, int imask, xnarch_cpumask_t affinity, void(*entry)(void *cookie), void *cookie)
Initial start of a newly created thread.
- void [xnpod_restart_thread](#) (xnthread_t *thread)
Restart a thread.
- void [xnpod_delete_thread](#) (xnthread_t *thread)
Delete a thread.
- xnflags_t [xnpod_set_thread_mode](#) (xnthread_t *thread, xnflags_t clrmask, xnflags_t setmask)
Change a thread's control mode.
- void [xnpod_suspend_thread](#) (xnthread_t *thread, xnflags_t mask, xnticks_t timeout, int mode, struct xnsynch *wchan)
Suspend a thread.
- void [xnpod_resume_thread](#) (xnthread_t *thread, xnflags_t mask)
Resume a thread.
- int [xnpod_unblock_thread](#) (xnthread_t *thread)
Unblock a thread.
- void [xnpod_renice_thread](#) (xnthread_t *thread, int prio)
Change the base priority of a thread.
- int [xnpod_migrate_thread](#) (int cpu)
Migrate the current thread.
- void [xnpod_rotate_readyq](#) (int prio)
Rotate a priority level in the ready queue.
- void [xnpod_do_rr](#) (void)
Handle the round-robin scheduling policy.
- void [xnpod_schedule](#) (void)
Rescheduling procedure entry point.
- void [xnpod_dispatch_signals](#) (void)
Deliver pending asynchronous signals to the running thread.
- int [xnpod_announce_tick](#) (struct xnintr *intr)
Announce a new clock tick.
- void [xnpod_activate_rr](#) (xnticks_t quantum)

Globally activate the round-robin scheduling.

- void [xnpod_deactivate_rr](#) (void)
Globally deactivate the round-robin scheduling.
- void [xnpod_set_time](#) (xntbase_t *tbase, xnticks_t newtime)
Set the nucleus idea of time for a given time base.
- int [xnpod_set_thread_periodic](#) (xnthread_t *thread, xnticks_t idate, xnticks_t period)
Make a thread periodic.
- int [xnpod_wait_thread_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.
- xnticks_t [xnpod_get_time](#) (xntbase_t *tbase)
Get the nucleus idea of time for a given time base.
- int [xnpod_add_hook](#) (int type, void(*routine)(xnthread_t *))
Install a nucleus hook.
- int [xnpod_remove_hook](#) (int type, void(*routine)(xnthread_t *))
Remove a nucleus hook.
- void [xnpod_welcome_thread](#) (xnthread_t *thread, int imask)
Thread prologue.
- static void [xnpod_preempt_current_thread](#) (xnsched_t *sched)
Preempts the current thread.
- int [xnpod_trap_fault](#) (void *fltinfo)
Default fault handler.

5.6.2 Function Documentation

5.6.2.1 void xnpod_activate_rr (xnticks_t *quantum*)

Globally activate the round-robin scheduling.

This service activates the round-robin scheduling for all threads which have the XNRRB flag set in their status mask (see [xnpod_set_thread_mode\(\)](#)). Each of them will run for the given time quantum, then preempted and moved to the end of its priority group in the ready queue. This process is repeated until the round-robin scheduling is disabled for those threads.

Parameters:

quantum The time credit which will be given to each rr-enabled thread (in ticks).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.2 `int xnpod_add_hook(int type, void(*)(xnthread_t *) routine)`

Install a nucleus hook.

The nucleus allows to register user-defined routines which get called whenever a specific scheduling event occurs. Multiple hooks can be chained for a single event type, and get called on a FIFO basis.

The scheduling is locked while a hook is executing.

Parameters:

type Defines the kind of hook to install:

- `XNHOOK_THREAD_START`: The user-defined routine will be called on behalf of the starter thread whenever a new thread starts. The descriptor address of the started thread is passed to the routine.
- `XNHOOK_THREAD_DELETE`: The user-defined routine will be called on behalf of the deleter thread whenever a thread is deleted. The descriptor address of the deleted thread is passed to the routine.
- `XNHOOK_THREAD_SWITCH`: The user-defined routine will be called on behalf of the resuming thread whenever a context switch takes place. The descriptor address of the thread which has been switched out is passed to the routine.

Parameters:

routine The address of the user-supplied routine to call.

Returns:

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- `-EINVAL` is returned if type is incorrect.
- `-ENOMEM` is returned if not enough memory is available from the system heap to add the new hook.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.3 `int xnpod_announce_tick (struct xnintr * intr)`

Announce a new clock tick.

This is the default service routine for clock ticks which performs the necessary housekeeping chores for time-related services managed by the nucleus. In a way or another, this routine must be called to announce each incoming clock tick to the nucleus.

Parameters:

intr The descriptor address of the interrupt object associated to the timer interrupt.

Side-effect: Since this routine manages the round-robin scheduling, the running thread (which has been preempted by the timer interrupt) can be switched out as a result of its time credit being exhausted. The nucleus always calls the rescheduling procedure after the outer interrupt has been processed.

Returns:

XN_ISR_HANDLED|XN_ISR_NOENABLE is always returned.

Environments:

This service can be called from:

- Interrupt service routine, must be called with interrupts off.

Rescheduling: possible.

5.6.2.4 `void xnpod_deactivate_rr (void)`

Globally deactivate the round-robin scheduling.

This service deactivates the round-robin scheduling for all threads which have the XNRRB flag set in their status mask (see [xnpod_set_thread_mode\(\)](#)).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.5 `void xnpod_delete_thread (xnthread_t * thread)`

Delete a thread.

Terminates a thread and releases all the nucleus resources it currently holds. A thread exists in the system since [xnpod_init_thread\(\)](#) has been called to create it, so this service must be called in order to destroy it afterwards.

Parameters:

thread The descriptor address of the terminated thread.

The DELETE hooks are called on behalf of the calling context (if any). The information stored in the thread control block remains valid until all hooks have been called.

Self-terminating a thread is allowed. In such a case, this service does not return to the caller.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

5.6.2.6 void xnpod_disable_timesource (void)

Stop the core time source.

Releases the hardware timer, and deactivates the master time base.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

5.6.2.7 void xnpod_dispatch_signals (void)

Deliver pending asynchronous signals to the running thread.

For internal use only.

This internal routine checks for the presence of asynchronous signals directed to the running thread, and attempts to start the asynchronous service routine (ASR) if any. Called with nklock locked, interrupts off.

5.6.2.8 void xnpod_do_rr (void)

Handle the round-robin scheduling policy.

For internal use only.

This routine is called from the slave time base tick handler to enforce the round-robin scheduling policy.

This service can be called from:

- Interrupt service routine, must be called with interrupts off, nklock locked.

Rescheduling: never.

5.6.2.9 int xnpod_enable_timesource (void)

Activate the core time source.

Xenomai implements the notion of time base, by which software timers that belong to different skins may be clocked separately according to distinct frequencies, or aperiodically. In the periodic case, delays and timeouts are given in counts of ticks; the duration of a tick is specified by the time base. In the aperiodic case, timings are directly specified in nanoseconds.

Only a single aperiodic (i.e. tick-less) time base may exist in the system, and the nucleus provides for it through the `nktbase` object. All skins depending on aperiodic timings should bind to the latter, also known as the master time base. Skins depending on periodic timings may create and bind to their own time base. Such a periodic time base is managed as a slave object of the master one. A cascading software timer, which is fired by the master time base according to the appropriate frequency, triggers in turn the update process of the associated slave time base, which eventually fires the elapsed software timers controlled by the latter.

Xenomai always controls the underlying hardware timer in a tick-less fashion, also known as the oneshot mode. The `xnpod_enable_timesource()` service configures the timer chip as needed, and activates the master time base.

Returns:

0 is returned on success. Otherwise:

- -ENODEV is returned if a failure occurred while configuring the hardware timer.
- -ENOSYS is returned if no active pod exists.

Side-effect: A host timing service is started in order to relay the canonical periodical tick to the underlying architecture, regardless of the frequency used for Xenomai's system tick. This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

Note:

Built-in support for periodic timing depends on `CONFIG_XENO_OPT_TIMING_PERIODIC`.

5.6.2.10 xnticks_t xnpod_get_time (xntbase_t * tbase)

Get the nucleus idea of time for a given time base.

This service gets the nucleus (external) clock time.

Parameters:

tbase The address of the affected time base.

Returns:

The current time (in jiffies) if the specified time base runs in periodic mode, or the system time (converted to nanoseconds) as maintained by the CPU if *tbase* refers to the master time base.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.11 `int xnpod_init(xnpod_t * pod, int minpri, int maxpri, xnflags_t flags)`

Initialize a new pod.

Initializes a new pod which can subsequently be used to start real-time activities. Once a pod is active, real-time APIs can be stacked over. There can only be a single pod active in the host environment. Such environment can be confined to a process (e.g. simulator or UVM), or expand machine-wide (e.g. Adeos).

Parameters:

- pod* The address of a pod descriptor the nucleus will use to store the pod-specific data. This descriptor must always be valid while the pod is active therefore it must be allocated in permanent memory.
- minpri* The value of the lowest priority level which is valid for threads created on behalf of this pod.
- maxpri* The value of the highest priority level which is valid for threads created on behalf of this pod.
- flags* A set of creation flags affecting the operation. The only defined flag is XNREUSE, which tells the nucleus that a pre-existing pod exhibiting the same properties as the one which is being registered may be reused. In such a case, the call returns successfully, keeping the active pod unmodified.

minpri might be numerically higher than *maxpri* if the upper real-time interface exhibits a reverse priority scheme. For instance, some APIs may define a range like *minpri*=255, *maxpri*=0 specifying that thread priorities increase as the priority level decreases numerically.

Returns:

0 is returned on success. Otherwise:

- -EBUSY is returned if a pod already exists. As a special exception, if the Xenomai pod is currently loaded with no active attachment onto it, it is forcibly unloaded and replaced by the new pod.
- -ENOMEM is returned if the memory manager fails to initialize.

Environments:

This service can be called from:

- Kernel module initialization code

Note:

No initialization code called by this routine may refer to the global "nkpod" pointer.

5.6.2.12 `int xnpod_init_thread (xnthread_t * thread, xntbase_t * tbase, const char * name, int prio, xnflags_t flags, unsigned stacksize)`

Initialize a new thread.

Initializes a new thread attached to the active pod. The thread is left in an innocuous state until it is actually started by `xnpod_start_thread()`.

Parameters:

thread The address of a thread descriptor the nucleus will use to store the thread-specific data. This descriptor must always be valid while the thread is active therefore it must be allocated in permanent memory.

Warning:

Some architectures may require the descriptor to be properly aligned in memory; this is an additional reason for descriptors not to be laid in the program stack where alignment constraints might not always be satisfied.

Parameters:

name An ASCII string standing for the symbolic name of the thread. This name is copied to a safe place into the thread descriptor. This name might be used in various situations by the nucleus for issuing human-readable diagnostic messages, so it is usually a good idea to provide a sensible value here. The simulator even uses this name intensively to identify threads in the debugging GUI it provides. However, passing NULL here is always legal and means "anonymous".

tbase The time base descriptor to refer to for all timed operations issued by the new thread. See `xntbase_alloc()` for detailed explanations about time bases.

prio The base priority of the new thread. This value must range from [minpri .. maxpri] (inclusive) as specified when calling the `xnpod_init()` service.

flags A set of creation flags affecting the operation. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the created thread:

- XNSUSP creates the thread in a suspended state. In such a case, the thread will have to be explicitly resumed using the `xnpod_resume_thread()` service for its execution to actually begin, additionally to issuing `xnpod_start_thread()` for it. This flag can also be specified when invoking `xnpod_start_thread()` as a starting mode.
- XNFPU (enable FPU) tells the nucleus that the new thread will use the floating-point unit. In such a case, the nucleus will handle the FPU context save/restore ops upon thread switches at the expense of a few additional cycles per context switch. By default, a thread is not expected to use the FPU. This flag is simply ignored when the nucleus runs on behalf of a userspace-based real-time control layer since the FPU management is always active if present.

- **XNINVPS** tells the nucleus that the new thread will use an inverted priority scale with respect to the one enforced by the current pod. This means that the calling skin will still have to normalize the priority levels passed to the nucleus routines so that they conform to the pod's priority scale, but the nucleus will automatically rescale those values when displaying the priority information (e.g. `/proc/xenomai/sched` output). This bit must not be confused with the **XNRPRIO** bit, which is internally set by the nucleus during pod initialization when the low priority level is found to be numerically higher than the high priority bound. Having the **XNINVPS** bit set for a thread running on a pod with **XNRPRIO** unset means that the skin emulates a decreasing priority scale using the pod's increasing priority scale. This is typically the case for skins running over the core pod (see `include/nucleus/core.h`).

Parameters:

stacksize The size of the stack (in bytes) for the new thread. If zero is passed, the nucleus will use a reasonable pre-defined size depending on the underlying real-time control layer.

After creation, the new thread can be set a magic cookie by skins using `xnthread_set_magic()` to unambiguously identify threads created in their realm. This value will be copied as-is to the *magic* field of the thread struct. 0 is a conventional value for "no magic".

Returns:

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- **-EINVAL** is returned if *flags* has invalid bits set.
- **-ENOMEM** is returned if not enough memory is available from the system heap to create the new thread's stack.

Side-effect: This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.13 `int xnpod_migrate_thread (int cpu)`

Migrate the current thread.

This call makes the current thread migrate to another CPU if its affinity allows it.

Parameters:

cpu The destination CPU.

Return values:

0 if the thread could migrate ;

- EPERM* if the calling context is asynchronous, or the current thread affinity forbids this migration ;
- EBUSY* if the scheduler is locked.

5.6.2.14 void xnpod_preempt_current_thread (xnsched_t * sched) [inline, static]

Preempts the current thread.

For internal use only.

Preempts the running thread (because a higher priority thread has just been readied). The thread is re-inserted to the front of its priority group in the ready thread queue. Must be called with nklock locked, interrupts off.

5.6.2.15 int xnpod_remove_hook (int type, void(*)(xnthread_t *) routine)

Remove a nucleus hook.

This service removes a nucleus hook previously registered using [xnpod_add_hook\(\)](#).

Parameters:

- type* Defines the kind of hook to remove among XNHOOK_THREAD_START, XNHOOK_THREAD_DELETE and XNHOOK_THREAD_SWITCH.
- routine* The address of the user-supplied routine to remove.

Returns:

0 is returned on success. Otherwise, -EINVAL is returned if type is incorrect or if the routine has never been registered before.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.16 void xnpod_renice_thread (xnthread_t * thread, int prio)

Change the base priority of a thread.

Changes the base priority of a thread. If the reniced thread is currently blocked, waiting in priority-pending mode (XNSYNCH_PRIO) for a synchronization object to be signaled, the nucleus will attempt to reorder the object's wait queue so that it reflects the new sleeper's priority, unless the XNSYNCH_DREORD flag has been set for the pended object.

Parameters:

- thread* The descriptor address of the affected thread.

prio The new thread priority.

It is absolutely required to use this service to change a thread priority, in order to have all the needed housekeeping chores correctly performed. i.e. Do *not* change the `thread.cprio` field by hand, unless the thread is known to be in an innocuous state (e.g. dormant).

Side-effects:

- This service does not call the rescheduling procedure but may affect the ready queue.
- Assigning the same priority to a running or ready thread moves it to the end of the ready queue, thus causing a manual round-robin.
- If the reniced thread is a user-space shadow, propagate the request to the mated Linux task.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.17 void `xnpod_restart_thread` (`xnthread_t * thread`)

Restart a thread.

Restarts a previously started thread. The thread is first terminated then respawned using the same information that prevailed when it was first started, including the mode bits and interrupt mask initially passed to the `xnpod_start_thread()` service. As a consequence of this call, the thread entry point is rerun.

Parameters:

thread The descriptor address of the affected thread which must have been previously started by the `xnpod_start_thread()` service.

Self-restarting a thread is allowed. However, restarting the root thread is not.

Environments:

This service can be called from:

- Kernel-based task
- User-space task

Rescheduling: possible.

5.6.2.18 void xnpod_resume_thread (xnthread_t * thread, xnflags_t mask)

Resume a thread.

Resumes the execution of a thread previously suspended by one or more calls to [xnpod_suspend_thread\(\)](#). This call removes a suspensive condition affecting the target thread. When all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters:

thread The descriptor address of the resumed thread.

mask The suspension mask specifying the suspensive condition to remove from the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag removes the explicit suspension condition. This condition might be additive to the XNPEND condition.
- XNDELAY. This flag removes the counted delay wait condition.
- XNPEND. This flag removes the resource wait condition. If a watchdog is armed, it is automatically disarmed by this call. Unlike the two previous conditions, only the current thread can set this condition for itself, i.e. no thread can force another one to pend on a resource.

When the thread is eventually resumed by one or more calls to [xnpod_resume_thread\(\)](#), the caller of [xnpod_suspend_thread\(\)](#) in the awakened thread that suspended itself should check for the following bits in its own information mask to determine what caused its wake up:

- XNRMID means that the caller must assume that the pended synchronization object has been destroyed (see [xnsynch_flush\(\)](#)).
- XNTIMEO means that the delay elapsed, or the watchdog went off before the corresponding synchronization object was signaled.
- XNBREAK means that the wait has been forcibly broken by a call to [xnpod_unblock_thread\(\)](#).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.19 void xnpod_rotate_readyq (int prio)

Rotate a priority level in the ready queue.

The thread at the head of the ready queue of the given priority level is moved to the end of this queue. Therefore, the execution of threads having the same priority is switched. Round-robin scheduling policies may be implemented by periodically issuing this call in a given period of time. It should be noted that the nucleus already provides a built-in round-robin mode though (see [xnpod_activate_rr\(\)](#)).

Parameters:

prio The priority level to rotate. if XNPOD_RUNPRIO is given, the running thread priority is used to rotate the queue.

The priority level which is considered is always the base priority of a thread, not the possibly PIP-boosted current priority value. Specifying a priority level with no thread on it is harmless, and will simply lead to a null-effect.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.20 void xnpod_schedule (void)

Rescheduling procedure entry point.

This is the central rescheduling routine which should be called to validate and apply changes which have previously been made to the nucleus scheduling state, such as suspending, resuming or changing the priority of threads. This call first determines if a thread switch should take place, and performs it as needed. [xnpod_schedule\(\)](#) actually switches threads if:

- the running thread has been blocked or deleted.
- or, the running thread has a lower priority than the first ready to run thread.
- or, the running thread does not lead no more the ready threads (round-robin).

The nucleus implements a lazy rescheduling scheme so that most of the services affecting the threads state MUST be followed by a call to the rescheduling procedure for the new scheduling state to be applied. In other words, multiple changes on the scheduler state can be done in a row, waking threads up, blocking others, without being immediately translated into the corresponding context switches, like it would be necessary would it appear that a higher priority thread than the current one became runnable for instance. When all changes have been applied, the rescheduling procedure is then called to consider those changes, and possibly replace the current thread by another one.

As a notable exception to the previous principle however, every action which ends up suspending or deleting the current thread begets an immediate call to the rescheduling procedure on behalf of the service causing the state transition. For instance, self-suspension, self-destruction, or sleeping on a synchronization object automatically leads to a call to the rescheduling procedure, therefore the caller does not need to explicitly issue `xnpod_schedule()` after such operations.

The rescheduling procedure always leads to a null-effect if the scheduler is locked (XNLOCK bit set in the status mask of the running thread), or if it is called on behalf of an ISR or callout.

Calling this procedure with no applicable context switch pending is harmless and simply leads to a null-effect.

Side-effects:

- If an asynchronous service routine exists, the pending asynchronous signals are delivered to a resuming thread or on behalf of the caller before it returns from the procedure if no context switch has taken place. This behaviour can be disabled by setting the XNASDI flag in the thread's status mask by calling `xnpod_set_thread_mode()`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine, although this leads to a no-op.
- Kernel-based task
- User-space task

Note:

The switch hooks are called on behalf of the resuming thread.

5.6.2.21 void xnpod_schedule_runnable (xnthread_t * thread, int flags)

Hidden rescheduling procedure.

For internal use only.

This internal routine should NEVER be used directly by the upper interfaces. It reinserts the given thread into the ready queue then switches to the highest priority runnable thread. It must be called with nklock locked, interrupts off.

Parameters:

thread The descriptor address of the thread to reinsert into the ready queue.

flags A bitmask composed as follows:

- XNPOD_SCHDLIFO causes the target thread to be inserted at front of its priority group in the ready queue. Otherwise, the FIFO ordering is applied.
- XNPOD_NOSWITCH reorders the ready queue without switching contexts. This feature is used to preserve the atomicity of some operations.

5.6.2.22 `xnflags_t xnpod_set_thread_mode(xnthread_t * thread, xnflags_t clrmask, xnflags_t setmask)`

Change a thread's control mode.

Change the control mode of a given thread. The control mode affects the behaviour of the nucleus regarding the specified thread.

Parameters:

thread The descriptor address of the affected thread.

clrmask Clears the corresponding bits from the control field before *setmask* is applied. The scheduler lock held by the current thread can be forcibly released by passing the XNLOCK bit in this mask. In this case, the lock nesting count is also reset to zero.

setmask The new thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the thread:

- XNLOCK causes the thread to lock the scheduler. The target thread will have to call the `xnpod_unlock_sched()` service to unlock the scheduler or clear the XNLOCK bit forcibly using this service.
- XNRRB causes the thread to be marked as undergoing the round-robin scheduling policy. The contents of the `thread.rrperiod` field determines the time quantum (in ticks) allowed for its next slice. If the thread is already undergoing the round-robin scheduling policy at the time this service is called, the time quantum remains unchanged.
- XNASDI disables the asynchronous signal handling for this thread. See [xnpod_schedule\(\)](#) for more on this.
- XNSHIELD enables the interrupt shield for the current user-space task. When engaged, the interrupt shield protects the shadow task running in secondary mode from any preemption by the regular Linux interrupt handlers, without delaying in any way Xenomai's interrupt handling. The shield is operated on a per-task basis at each context switch, depending on the setting of this flag. This feature is only available if the `CONFIG_XENO_OPT_ISHIELD` option has been enabled at configuration time; otherwise, this flag is simply ignored.
- XNRPIOFF disables thread priority coupling between Xenomai and Linux schedulers. This bit prevents the root Linux thread from inheriting the priority of the running shadow Xenomai thread. Use `CONFIG_XENO_OPT_RPIOFF` to globally disable priority coupling.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never, therefore, the caller should reschedule if XNLOCK has been passed into *clrmask*.

5.6.2.23 int xnpod_set_thread_periodic (xnthread_t * thread, xnticks_t idate, xnticks_t period)

Make a thread periodic.

Make a thread periodic by programming its first release point and its period in the processor time line. Subsequent calls to `xnpod_wait_thread_period()` will delay the thread until the next periodic release point in the processor timeline is reached.

Parameters:

thread The descriptor address of the affected thread. This thread is immediately delayed until the first periodic release point is reached.

idate The initial (absolute) date of the first release point, expressed in clock ticks (see note). The affected thread will be delayed until this point is reached. If *idate* is equal to `XN_INFINITE`, the current system date is used, and no initial delay takes place.

period The period of the thread, expressed in clock ticks (see note). As a side-effect, passing `XN_INFINITE` attempts to stop the thread's periodic timer; in the latter case, the routine always exits successfully, regardless of the previous state of this timer.

Returns:

0 is returned upon success. Otherwise:

- `-ETIMEDOUT` is returned if *idate* is different from `XN_INFINITE` and represents a date in the past.
- `-EWOULDBLOCK` is returned if the relevant time base has not been initialized by a call to `xnpod_init_timebase()`.
- `-EINVAL` is returned if *period* is different from `XN_INFINITE` but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the operation affects the current thread and *idate* has not elapsed yet.

Note:

The *idate* and *period* values will be interpreted as jiffies if *thread* is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

5.6.2.24 void xnpod_set_time (xntbase_t * tbase, xnticks_t newtime)

Set the nucleus idea of time for a given time base.

The nucleus tracks the current time as a monotonously increasing count of ticks announced by the timer source since the epoch. The epoch is initially the same as the underlying architecture system time. This service changes the epoch for the specified time base.

Parameters:

- tbase* The address of the affected time base.
- newtime* The new time to set for the specified time base.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.25 void xnpod_shutdown (int xtype)

Shutdown the current pod.

Forcibly shutdowns the active pod. All existing nucleus threads (but the root one) are terminated, and the system heap is freed.

Parameters:

- xtype* An exit code passed to the host environment who started the nucleus. Zero is always interpreted as a successful return.

The nucleus never calls this routine directly. Skins should provide their own shutdown handlers which end up calling [xnpod_shutdown\(\)](#) after their own housekeeping chores have been carried out.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

5.6.2.26 int xnpod_start_thread (xnthread_t * thread, xnflags_t mode, int imask, xnarch_cpumask_t affinity, void(*) (void *cookie) entry, void * cookie)

Initial start of a newly created thread.

Starts a (newly) created thread, scheduling it for the first time. This call releases the target thread from the XNDORMANT state. This service also sets the initial mode and interrupt mask for the new thread.

Parameters:

- thread* The descriptor address of the affected thread which must have been previously initialized by the [xnpod_init_thread\(\)](#) service.

mode The initial thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the started thread:

- **XNLOCK** causes the thread to lock the scheduler when it starts. The target thread will have to call the `xnpod_unlock_sched()` service to unlock the scheduler.
- **XNRRB** causes the thread to be marked as undergoing the round-robin scheduling policy at startup. The contents of the `thread.rrperiod` field determines the time quantum (in ticks) allowed for its next slice.
- **XNASDI** disables the asynchronous signal handling for this thread. See `xnpod_schedule()` for more on this.
- **XNSUSP** makes the thread start in a suspended state. In such a case, the thread will have to be explicitly resumed using the `xnpod_resume_thread()` service for its execution to actually begin.

Parameters:

imask The interrupt mask that should be asserted when the thread starts. The processor interrupt state will be set to the given value when the thread starts running. The interpretation of this value might be different across real-time layers, but a non-zero value should always mark an interrupt masking in effect (e.g. `cli()`). Conversely, a zero value should always mark a fully preemptible state regarding interrupts (i.e. `sti()`).

affinity The processor affinity of this thread. Passing `XNPOD_ALL_CPUS` or an empty affinity set means "any cpu".

entry The address of the thread's body routine. In other words, it is the thread entry point.

cookie A user-defined opaque cookie the nucleus will pass to the emerging thread as the sole argument of its entry point.

The START hooks are called on behalf of the calling context (if any).

Return values:

- **0** if *thread* could be started ;
- **-EBUSY** if *thread* was already started ;
- **-EINVAL** if the value of *affinity* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

5.6.2.27 void `xnpod_suspend_thread` (`xnthread_t * thread`, `xnflags_t mask`, `xnticks_t timeout`, `int mode`, `struct xnsynch * wchan`)

Suspend a thread.

Suspends the execution of a thread according to a given suspensive condition. This thread will not be eligible for scheduling until all the pending suspensive conditions set by this service are removed by one or more calls to [xnpod_resume_thread\(\)](#).

Parameters:

thread The descriptor address of the suspended thread.

mask The suspension mask specifying the suspensive condition to add to the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag forcibly suspends a thread, regardless of any resource to wait for. A reverse call to [xnpod_resume_thread\(\)](#) specifying the XNSUSP bit must be issued to remove this condition, which is cumulative with other suspension bits. *wchan* should be NULL when using this suspending mode.
- XNDELAY. This flag denotes a counted delay wait (in ticks) which duration is defined by the value of the *timeout* parameter.
- XNPEND. This flag denotes a wait for a synchronization object to be signaled. The *wchan* argument must point to this object. A timeout value can be passed to bound the wait. This suspending mode should not be used directly by the upper interface, but rather through the [xnsynch_sleep_on\(\)](#) call.

Parameters:

timeout The timeout which may be used to limit the time the thread pends on a resource. This value is a wait time given in ticks (see note). It can either be relative or absolute depending on *mode*. Passing XN_INFINITE **and** setting *mode* to XNTIMER_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer. If the current operation mode of the system timer is oneshot and *timeout* elapses before [xnpod_suspend_thread\(\)](#) has completed, then the target thread will not be suspended, and this routine leads to a null effect.

mode The mode of the *timeout* parameter. It can either be set to XNTIMER_RELATIVE or XNTIMER_ABSOLUTE.

wchan The address of a pended resource. This parameter is used internally by the synchronization object implementation code to specify on which object the suspended thread pends. NULL is a legitimate value when this parameter does not apply to the current suspending mode (e.g. XNSUSP).

Note:

If the target thread is a shadow which has received a Linux-originated signal, then this service immediately exits without suspending the thread, but raises the XNBREAK condition in its information mask.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread suspends itself.

Note:

The *timeout* value will be interpreted as jiffies if *thread* is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

5.6.2.28 void xnpod_trap_fault (void * *fltinfo*)

Default fault handler.

This is the default handler which is called whenever an uncontrolled exception or fault is caught. If the fault is caught on behalf of a real-time thread, the fault handler stored into the service table (`svctable.faulthandler`) is invoked and the fault is not propagated to the host system. Otherwise, the fault is unhandled by the nucleus and simply propagated.

Parameters:

fltinfo An opaque pointer to the arch-specific buffer describing the fault. The actual layout is defined by the `xnarch_fltinfo_t` type in each arch-dependent layer file.

5.6.2.29 int xnpod_unblock_thread (xnthread_t * *thread*)

Unblock a thread.

Breaks the thread out of any wait it is currently in. This call removes the XNDELAY and XNPEND suspensive conditions previously put by `xnpod_suspend_thread()` on the target thread. If all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters:

thread The descriptor address of the unblocked thread.

This call neither releases the thread from the XNSUSP, XNRELAX nor the XNDORMANT suspensive conditions.

When the thread resumes execution, the XNBREAK bit is set in the unblocked thread's information mask. Unblocking a non-blocked thread is perfectly harmless.

Returns:

non-zero is returned if the thread was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.6.2.30 `int xnpod_wait_thread_period (unsigned long * overruns_r)`

Wait for the next periodic release point.

Make the current thread wait for the next periodic release point in the processor time line.

Parameters:

overruns_r If non-NULL, *overruns_r* must be a pointer to a memory location which will be written with the count of pending overruns. This value is copied only when `xnpod_wait_thread_period()` returns `-ETIMEDOUT` or success; the memory location remains unmodified otherwise. If NULL, this count will never be copied back.

Returns:

0 is returned upon success; if *overruns_r* is valid, zero is copied to the pointed memory location. Otherwise:

- `-EWOULDBLOCK` is returned if `xnpod_set_thread_periodic()` has not previously been called for the calling thread.
- `-EINTR` is returned if `xnpod_unblock_thread()` has been called for the waiting thread before the next periodic release point has been reached. In this case, the overrun counter is reset too.
- `-ETIMEDOUT` is returned if the timer has overrun, which indicates that one or more previous release points have been missed by the calling thread. If *overruns_r* is valid, the count of pending overruns is copied to the pointed memory location.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always, unless the current release point has already been reached. In the latter case, the current thread immediately returns from this service without being delayed.

5.6.2.31 `void xnpod_welcome_thread (xnthread_t * thread, int imask)`

Thread prologue.

For internal use only.

This internal routine is called on behalf of a (re)starting thread's prologue before the user entry point is invoked. This call is reserved for internal housekeeping chores and cannot be inlined.

Entered with nklock locked, irqs off.

5.7 Registry services.

5.7.1 Detailed Description

The registry provides a mean to index real-time object descriptors created by Xenomai skins on unique alphanumeric keys. When labeled this way, a real-time object is globally exported; it can be searched for, and its descriptor returned to the caller for further use; the latter operation is called a "binding". When no object has been registered under the given name yet, the registry can be asked to set up a rendez-vous, blocking the caller until the object is eventually registered.

Files

- file [registry.h](#)
This file is part of the Xenomai project.
- file [registry.c](#)
This file is part of the Xenomai project.

Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, xnpnode_t *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_remove_safe](#) (xnhandle_t handle, xnticks_t timeout)
Unregister an idle real-time object.
- void * [xnregistry_get](#) (xnhandle_t handle)
Find and lock a real-time object into the registry.
- u_long [xnregistry_put](#) (xnhandle_t handle)
Unlock a real-time object from the registry.
- void * [xnregistry_fetch](#) (xnhandle_t handle)
Find a real-time object into the registry.

5.7.2 Function Documentation

5.7.2.1 int xnregistry_bind (const char * key, xnticks_t timeout, xnhandle_t * phandle)

Bind to a real-time object.

This service retrieves the registry handle of a given object identified by its key. Unless otherwise specified, this service will block the caller if the object is not registered yet, waiting for such registration to occur.

Parameters:

key A valid NULL-terminated string which identifies the object to bind to.

timeout The number of clock ticks to wait for the registration to occur (see note). Passing XN_INFINITE causes the caller to block indefinitely until the object is registered. Passing XN_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry.

phandle A pointer to a memory location which will be written upon success with the generic handle defined by the registry for the retrieved object. Contents of this memory is undefined upon failure.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* is NULL.
- -EINTR is returned if [xnpod_unblock_thread\(\)](#) has been called for the waiting thread before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to XN_NONBLOCK and the searched object is not registered on entry. As a special exception, this error is also returned if this service should block, but was called from a context which cannot sleep (e.g. interrupt, non-realtime or scheduler locked).
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN_NONBLOCK.
- Kernel-based thread.

Rescheduling: always unless the request is immediately satisfied or *timeout* specifies a non-blocking operation.

Note:

The *timeout* value will be interpreted as jiffies if *thread* is bound to a periodic time base (see [xnpod_init_thread](#)), or nanoseconds otherwise.

5.7.2.2 `int xnregistry_enter (const char * key, void * objaddr, xnhandle_t * phandle, xnnode_t * pnode)`

Register a real-time object.

This service allocates a new registry slot for an associated object, and indexes it by an alphanumeric key for later retrieval.

Parameters:

key A valid NULL-terminated string by which the object will be indexed and later retrieved in the registry. Since it is assumed that such key is stored into the registered object, it will *not* be copied but only kept by reference in the registry.

objaddr An opaque pointer to the object to index by *key*.

phandle A pointer to a generic handle defined by the registry which will uniquely identify the indexed object, until the latter is unregistered using the [xnregistry_remove\(\)](#) service.

pnode A pointer to an optional /proc node class descriptor. This structure provides the information needed to export all objects from the given class through the /proc filesystem, under the /proc/xenomai/registry entry. Passing NULL indicates that no /proc support is available for the newly registered object.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* or *objaddr* are NULL, or if *key* contains an invalid '/' character.
- -ENOMEM is returned if the system fails to get enough dynamic memory from the global real-time heap in order to register the object.
- -EEXIST is returned if the *key* is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: possible.

5.7.2.3 `void* xnregistry_fetch (xnhandle_t handle)`

Find a real-time object into the registry.

This service retrieves an object from its handle into the registry and returns the memory address of its descriptor.

Parameters:

handle The generic handle of the object to fetch. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns:

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread

Rescheduling: never.

5.7.2.4 void* xnregistry_get (xnhandle_t handle)

Find and lock a real-time object into the registry.

This service retrieves an object from its handle into the registry and prevents its removal atomically. A locking count is tracked, so that [xnregistry_get\(\)](#) and [xnregistry_put\(\)](#) must be used in pair.

Parameters:

handle The generic handle of the object to find and lock. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns:

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread.

Rescheduling: never.

5.7.2.5 u_long xnregistry_put (xnhandle_t handle)

Unlock a real-time object from the registry.

This service decrements the lock count of a registered object previously locked by a call to [xnregistry_get\(\)](#). The object is actually unlocked from the registry when the locking count falls down to zero, thus waking up any thread currently blocked on [xnregistry_remove\(\)](#) for unregistering it.

Parameters:

handle The generic handle of the object to unlock. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns:

The decremented lock count is returned upon success. Zero is also returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread

Rescheduling: possible if the lock count falls down to zero and some thread is currently waiting for the object to be unlocked.

5.7.2.6 int xnregistry_remove (xnhandle_t handle)

Forcibly unregister a real-time object.

This service forcibly removes an object from the registry. The removal is performed regardless of the current object's locking status.

Parameters:

handle The generic handle of the object to remove.

Returns:

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: never.

5.7.2.7 int xnregistry_remove_safe (xnhandle_t handle, xnticks_t timeout)

Unregister an idle real-time object.

This service removes an object from the registry. The caller might sleep as a result of waiting for the target object to be unlocked prior to the removal (see [xnregistry_put\(\)](#)).

Parameters:

handle The generic handle of the object to remove.

timeout If the object is locked on entry, *param* gives the number of clock ticks to wait for the unlocking to occur (see note). Passing XN_INFINITE causes the caller to block indefinitely until the object is unlocked. Passing XN_NONBLOCK causes the service to return immediately without waiting if the object is locked on entry.

Returns:

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.
- -EWOULDBLOCK is returned if *timeout* is equal to XN_NONBLOCK and the object is locked on entry.
- -EBUSY is returned if *handle* refers to a locked object and the caller could not sleep until it is unlocked.
- -ETIMEDOUT is returned if the object cannot be removed within the specified amount of time.
- -EINTR is returned if [xnpod_unblock_thread\(\)](#) has been called for the calling thread waiting for the object to be unlocked.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN_NONBLOCK.
- Kernel-based thread.

Rescheduling: possible if the object to remove is currently locked and the calling context can sleep.

Note:

The *timeout* value will be interpreted as jiffies if the current thread is bound to a periodic time base (see [xnpod_init_thread\(\)](#)), or nanoseconds otherwise.

5.8 Real-time shadow services.

5.8.1 Detailed Description

Real-time shadow services.

Files

- file [shadow.c](#)
Real-time shadow services.

Functions

- int [xnshadow_harden](#) (void)
Migrate a Linux task to the Xenomai domain.
- void [xnshadow_relax](#) (int notify)
Switch a shadow thread back to the Linux domain.
- int [xnshadow_map](#) (xnthread_t *thread, xncompletion_t __user *u_completion)
Create a shadow thread context.
- xnshadow_ppd_t * [xnshadow_ppd_get](#) (unsigned muxid)
Return the per-process data attached to the calling process.

5.8.2 Function Documentation

5.8.2.1 int xnshadow_harden (void)

Migrate a Linux task to the Xenomai domain.

For internal use only.

This service causes the transition of "current" from the Linux domain to Xenomai. This is obtained by asking the gatekeeper to resume the shadow mated with "current" then triggering the rescheduling procedure in the Xenomai domain. The shadow will resume in the Xenomai domain as returning from schedule().

Environments:

This service can be called from:

- User-space thread operating in secondary (i.e. relaxed) mode.

Rescheduling: always.

5.8.2.2 `int xnshadow_map (xnthread_t * thread, xncompletion_t __user * u_completion)`

Create a shadow thread context.

For internal use only.

This call maps a nucleus thread to the "current" Linux task. The priority of the Linux task is set to the priority of the shadow thread bounded to the [0..MAX_RT_PRIO-1] range, and its scheduling policy is set to either SCHED_FIFO for non-zero priority levels, or SCHED_NORMAL otherwise.

Parameters:

thread The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to `xnpod_init_thread()`.

Warning:

The thread must have been set the same magic number (i.e. `xnthread_set_magic()`) than the one used to register the skin it belongs to. Failing to do so might lead to unexpected results.

Parameters:

u_completion is the address of an optional completion descriptor aimed at synchronizing our parent thread with us. If non-NULL, the information `xnshadow_map()` will store into the completion block will be later used to wake up the parent thread when the current shadow has been initialized. In the latter case, the new shadow thread is left in a dormant state (XNDORMANT) after its creation, leading to the suspension of "current" in the Linux domain, only processing signals. Otherwise, the shadow thread is immediately started and "current" immediately resumes in the Xenomai domain from this service.

Returns:

0 is returned on success. Otherwise:

- -ERESTARTSYS is returned if the current Linux task has received a signal, thus preventing the final migration to the Xenomai domain (i.e. in order to process the signal in the Linux domain). This error should not be considered as fatal.
- -EPERM is returned if the shadow thread has been killed before the current task had a chance to return to the caller. In such a case, the real-time mapping operation has failed globally, and no Xenomai resource remains attached to it.

Environments:

This service can be called from:

- Regular user-space process.

Rescheduling: always.

5.8.2.3 `xnshadow_ppd_t* xnshadow_ppd_get (unsigned muxid)`

Return the per-process data attached to the calling process.

This service returns the per-process data attached to the calling process for the skin whose muxid is *muxid*. It must be called with `nklock` locked, `irqs` off.

See `xnshadow_register_interface()` documentation for information on the way to attach a per-process data to a process.

Parameters:

muxid the skin muxid.

Returns:

the per-process data if the current context is a user-space process;
NULL otherwise.

5.8.2.4 void xnshadow_relax (int *notify*)

Switch a shadow thread back to the Linux domain.

For internal use only.

This service yields the control of the running shadow back to Linux. This is obtained by suspending the shadow and scheduling a wake up call for the mated user task inside the Linux domain. The Linux task will resume on return from [xnpod_suspend_thread\(\)](#) on behalf of the root thread.

Parameters:

notify A boolean flag indicating whether threads monitored from secondary mode switches should be sent a SIGXCPU signal. For instance, some internal operations like task exit should not trigger such signal.

Environments:

This service can be called from:

- User-space thread operating in primary (i.e. harden) mode.

Rescheduling: always.

Note:

"current" is valid here since the shadow runs with the properties of the Linux task.

5.9 Thread synchronization services.

5.9.1 Detailed Description

Thread synchronization services.

Files

- file [synch.c](#)
Thread synchronization services.

Functions

- void [xnsynch_init](#) (xnsynch_t *synch, xnflags_t flags)
Initialize a synchronization object.
- void [xnsynch_sleep_on](#) (xnsynch_t *synch, xnticks_t timeout, int mode)
Sleep on a synchronization object.
- static void [xnsynch_clear_boost](#) (xnsynch_t *synch, xnthread_t *lastowner)
Clear the priority boost.
- void [xnsynch_renice_sleeper](#) (xnthread_t *thread)
Change a sleeper's priority.
- xnthread_t * [xnsynch_wakeup_one_sleeper](#) (xnsynch_t *synch)
Give the resource ownership to the next waiting thread.
- xnpholder_t * [xnsynch_wakeup_this_sleeper](#) (xnsynch_t *synch, xnpholder_t *holder)
Give the resource ownership to a given waiting thread.
- int [xnsynch_flush](#) (xnsynch_t *synch, xnflags_t reason)
Unblock all waiters pending on a resource.
- void [xnsynch_forget_sleeper](#) (xnthread_t *thread)
Abort a wait for a resource.
- void [xnsynch_release_all_ownerships](#) (xnthread_t *thread)
Release all ownerships.

5.9.2 Function Documentation

5.9.2.1 void xnsynch_clear_boost (xnsynch_t * *synch*, xnthread_t * *owner*) [static]

Clear the priority boost.

For internal use only.

This service is called internally whenever a synchronization object is not claimed anymore by sleepers to reset the object owner's priority to its initial level.

Parameters:

synch The descriptor address of the synchronization object.

owner The descriptor address of the thread which currently owns the synchronization object.

Note:

This routine must be entered nklock locked, interrupts off.

5.9.2.2 int xnsynch_flush (xnsynch_t * synch, xnflags_t reason)

Unblock all waiters pending on a resource.

This service atomically releases all threads which currently sleep on a given resource.

This service should be called by upper interfaces under circumstances requiring that the pending queue of a given resource is cleared, such as before the resource is deleted.

Parameters:

synch The descriptor address of the synchronization object to be flushed.

reason Some flags to set in the information mask of every unblocked thread. Zero is an acceptable value. The following bits are pre-defined by the nucleus:

- XNRMID should be set to indicate that the synchronization object is about to be destroyed (see [xnpod_resume_thread\(\)](#)).
- XNBREAK should be set to indicate that the wait has been forcibly interrupted (see [xnpod_unblock_thread\(\)](#)).

Returns:

XNSYNCH_RESCHED is returned if at least one thread is unblocked, which means the caller should invoke [xnpod_schedule\(\)](#) for applying the new scheduling state. Otherwise, XNSYNCH_DONE is returned.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object is no more owned by any thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.9.2.3 void xnsynch_forget_sleeper (xnthread_t * thread)

Abort a wait for a resource.

For internal use only.

Performs all the necessary housekeeping chores to stop a thread from waiting on a given synchronization object.

Parameters:

thread The descriptor address of the affected thread.

When the trace support is enabled (i.e. MVM), the idle state is posted to the synchronization object's state diagram (if any) whenever no thread remains blocked on it. The real-time interfaces must ensure that such condition (i.e. EMPTY/IDLE) is mapped to state #0.

Note:

This routine must be entered nklock locked, interrupts off.

5.9.2.4 void xnsynch_init (xnsynch_t * synch, xnflags_t flags)

Initialize a synchronization object.

Initializes a new specialized object which can subsequently be used to synchronize real-time activities. The Xenomai nucleus provides a basic synchronization object which can be used to build higher resource objects. Nucleus threads can wait for and signal such objects in order to synchronize their activities.

This object has built-in support for priority inheritance.

Parameters:

synch The address of a synchronization object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

flags A set of creation flags affecting the operation. The valid flags are:

- XNSYNCH_PRIO causes the threads waiting for the resource to pend in priority order. Otherwise, FIFO ordering is used (XNSYNCH_FIFO).
- XNSYNCH_PIP causes the priority inheritance mechanism to be automatically activated when a priority inversion is detected among threads using this object. Otherwise, no priority inheritance takes place upon priority inversion (XNSYNCH_NOPIP).
- XNSYNCH_DREORD (Disable REORDering) tells the nucleus that the wait queue should not be reordered whenever the priority of a blocked thread it holds is changed. If this flag is not specified, changing the priority of a blocked thread using [xnpod_renice_thread\(\)](#) will cause this object's wait queue to be reordered according to the new priority level, provided the synchronization object makes the waiters wait by priority order on the awaited resource (XNSYNCH_PRIO).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task

Rescheduling: never.

5.9.2.5 void xnsynch_release_all_ownerships (xnthread_t * thread)

Release all ownerships.

For internal use only.

This call is used internally to release all the ownerships obtained by a thread on synchronization objects. This routine must be entered interrupts off.

Parameters:

thread The descriptor address of the affected thread.

Note:

This routine must be entered nklock locked, interrupts off.

5.9.2.6 void xnsynch_renice_sleeper (xnthread_t * thread)

Change a sleeper's priority.

For internal use only.

This service is used by the PIP code to update the pending priority of a sleeping thread.

Parameters:

thread The descriptor address of the affected thread.

Note:

This routine must be entered nklock locked, interrupts off.

5.9.2.7 void xnsynch_sleep_on (xnsynch_t * synch, xnticks_t timeout, int mode)

Sleep on a synchronization object.

Makes the calling thread sleep on the specified synchronization object, waiting for it to be signaled.

This service should be called by upper interfaces wanting the current thread to pend on the given resource.

Parameters:

synch The descriptor address of the synchronization object to sleep on.

timeout The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given in ticks (see note). It can either be relative or absolute depending on *mode*. Passing XN_INFINITE **and** setting *mode* to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

mode The mode of the *timeout* parameter. It can either be set to XN_RELATIVE or XN_ABSOLUTE.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always.

Note:

The *timeout* value will be interpreted as jiffies if the current thread is bound to a periodic time base (see `xnpod_init_thread`), or nanoseconds otherwise.

5.9.2.8 `xnthread_t* xnsynch_wakeup_one_sleeper (xnsynch_t * synch)`

Give the resource ownership to the next waiting thread.

This service gives the ownership of a synchronization object to the thread which is currently leading the object's pending list. The sleeping thread is unblocked, but no action is taken regarding the previous owner of the resource.

This service should be called by upper interfaces wanting to signal the given resource so that a single waiter is resumed.

Parameters:

synch The descriptor address of the synchronization object whose ownership is changed.

Returns:

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.9.2.9 xnpholder_t* xnsynch_wakeup_this_sleeper (xnsynch_t * *synch*, xnpholder_t * *holder*)

Give the resource ownership to a given waiting thread.

This service gives the ownership of a given synchronization object to a specific thread which is currently pending on it. The sleeping thread is unblocked from its pending state. No action is taken regarding the previous resource owner.

This service should be called by upper interfaces wanting to signal the given resource so that a specific waiter is resumed.

Parameters:

synch The descriptor address of the synchronization object whose ownership is changed.

holder The link holder address of the thread to unblock (&thread->plink) which MUST be currently linked to the synchronization object's pending queue (i.e. synch->pendq).

Returns:

The link address of the unblocked thread in the synchronization object's pending queue.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.10 Time base services.

5.10.1 Detailed Description

Xenomai implements the notion of time base, by which software timers that belong to different skins may be clocked separately according to distinct frequencies, or aperiodically. In the periodic case, delays and timeouts are given in counts of ticks; the duration of a tick is specified by the time base. In the aperiodic case, timings are directly specified in nanoseconds.

Only a single aperiodic (i.e. tick-less) time base may exist in the system, and the nucleus provides for it through the `nktbase` object. All skins depending on aperiodic timings should bind to the latter (see `xntbase_alloc()`), also known as the master time base.

Skins depending on periodic timings may create and bind to their own time base. Such a periodic time base is managed as a timed slave object of the master time base. A cascading software timer fired by the master time base according to the appropriate frequency, triggers in turn the update process of the associated timed slave, which eventually fires the elapsed software timers controlled by the periodic time base. In other words, Xenomai emulates periodic timing over an aperiodic policy.

Xenomai always controls the underlying timer hardware in a tick-less fashion, also known as the oneshot mode.

Files

- file `timebase.c`

Functions

- `int xntbase_alloc` (const char *name, u_long period, xntbase_t **basep)
Allocate a time base.
- `void xntbase_free` (xntbase_t *base)
Free a time base.
- `int xntbase_update` (xntbase_t *base, u_long period)
Change the period of a time base.
- `int xntbase_switch` (const char *name, u_long period, xntbase_t **basep)
Replace a time base.
- `void xntbase_start` (xntbase_t *base)
Start a time base.
- `void xntbase_stop` (xntbase_t *base)
Stop a time base.
- `void xntbase_tick` (xntbase_t *base)
Announce a clock tick to a time base.

5.10.2 Function Documentation

5.10.2.1 `int xntbase_alloc (const char * name, u_long period, xntbase_t ** basep)`

Allocate a time base.

A time base is an abstraction used to provide private clocking information to real-time skins, by which they may operate either in aperiodic or periodic mode, possibly according to distinct clock frequencies in the latter case. This abstraction is required in order to support several RTOS emulator running concurrently, which exhibit different clocking policies and/or period.

Once allocated, a time base may be attached to all software timers created directly or indirectly by a given skin, and influences all timed services accordingly.

The `xntbase_alloc()` service allocates a new time base to the caller, and returns the address of its descriptor. The new time base is left in a disabled state (unless *period* equals `XN_APERIODIC_TICK`), calling `xntbase_start()` is needed to enable it.

Parameters:

name The symbolic name of the new time base. This information is used to report status information when reading from `/proc/xenomai/timebases`; it has currently no other usage.

period The duration of clock tick for the new time base, given as a count of nanoseconds. The special `XN_APERIODIC_TICK` value may be used to retrieve the master - aperiodic - time base, which is always up and running when a real-time skin has called the `xnpod_init()` service. All other values are meant to define the clock rate of a periodic time base. For instance, passing 1000000 (ns) in the *period* parameter will create a periodic time base clocked at a frequency of 1Khz.

basep A pointer to a memory location which will be written upon success with the address of the allocated time base. If *period* equals `XN_APERIODIC_TICK`, the address of the built-in master time base descriptor will be copied back to this location.

Returns:

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if no system memory is available to allocate a new time base descriptor.

Environments:

This service can be called from:

- Kernel module initialization code

Note:

Any periodic time base allocated by a real-time skin must be released by a call to `xntbase_free()` before the kernel module implementing the skin may be unloaded.

5.10.2.2 `void xntbase_free (xntbase_t * base)`

Free a time base.

This service disarms all outstanding timers from the affected periodic time base, destroys the aperiodic cascading timer, then releases the time base descriptor.

Parameters:

base The address of the time base descriptor to release.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Note:

Requests to free the master time base are silently caught and discarded; in such a case, outstanding aperiodic timers are left untouched.

5.10.2.3 void xntbase_start (xntbase_t * *base*)

Start a time base.

This service enables a time base, using a cascading timer running in the master time base as the source of periodic clock ticks. Timers attached to the started time base are immediated armed.

Parameters:

base The address of the time base descriptor to start.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Note:

Requests to enable the master time base are silently caught and discarded; only the internal service [xn timer source enable](#) is allowed to start the latter. The master time base remains enabled until no real-time skin remains attached to the nucleus.

5.10.2.4 void xntbase_stop (xntbase_t * *base*)

Stop a time base.

This service disables a time base, stopping the cascading timer running in the master time base which is used to clock it. Outstanding timers attached to the stopped time base are immediated disarmed.

Parameters:

base The address of the time base descriptor to stop.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Note:

Requests to disable the master time base are silently caught and discarded; only the internal service `xnpod_disable_timesource()` is allowed to stop the latter. The master time base remains enabled until no real-time skin remains attached to the nucleus.

5.10.2.5 `int xntbase_switch(const char * name, u_long period, xntbase_t ** basep)`

Replace a time base.

This service is useful for switching the current time base of a real-time skin between aperiodic and periodic modes, by providing a new time base descriptor as needed. The original time base descriptor is freed as a result of this operation (unless it refers to the master time base). The new time base is automatically started by a call to `xntbase_start()` if the original time base was enabled at the time of the call, or left in a disabled state otherwise.

This call handles all mode transitions and configuration changes carefully, i.e. periodic <-> periodic, aperiodic <-> aperiodic, periodic <-> aperiodic.

Parameters:

name The symbolic name of the new time base. This information is used to report status information when reading from `/proc/xenomai/timebases`; it has currently no other usage.

period The duration of clock tick for the time base, given as a count of nanoseconds. This value is meant to define the new clock rate of the affected periodic time base (i.e. $1e9 / \text{period}$).

basep A pointer to a memory location which will be first read to pick the address of the original time base to be replaced, then written back upon success with the address of the new time base. A null pointer is allowed on input in *basep*, in which case the new time base will be created as if `xntbase_alloc()` had been called directly.

Returns:

0 is returned on success. Otherwise:

- -ENOMEM is returned if no system memory is available to allocate a new time base descriptor.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

5.10.2.6 void xntbase_tick (xntbase_t * base)

Announce a clock tick to a time base.

This service announces a new clock tick to a time base. Normally, only specialized nucleus code would announce clock ticks. However, under certain circumstances, it may be useful to allow client code to send such notifications on their own.

Notifying a clock tick to a time base causes the timer management code to check for outstanding timers, which may in turn fire off elapsed timeout handlers. Additionally, periodic time bases (i.e. all but the master time base) would also update their count of elapsed jiffies, in case the current processor has been defined as the internal time keeper (i.e. CPU# == XNTIMER_KEEPER_ID).

Parameters:

base The address of the time base descriptor to announce a tick to.

Environments:

This service can be called from:

- Interrupt context only.

5.10.2.7 int xntbase_update (xntbase_t * base, u_long period)

Change the period of a time base.

Parameters:

base The address of the time base descriptor to update.

period The duration of clock tick for the time base, given as a count of nanoseconds. This value is meant to define the new clock rate of the affected periodic time base (i.e. 1e9 / period).

Returns:

0 is returned on success. Otherwise:

- -EINVAL is returned if an attempt is made to set a null period.

Environments:

This service can be called from:

- Kernel module initialization code
- Kernel-based task
- User-space task

Note:

Requests to update the master time base are silently caught and discarded. The master time base has a fixed aperiodic policy which may not be changed.

5.11 Timer services.

5.11.1 Detailed Description

The Xenomai timer facility always operate the timer hardware in oneshot mode, regardless of the time base in effect. Periodic timing is obtained through a software emulation, using cascading timers.

Depending on the time base used, the timer object stores time values either as count of jiffies (periodic), or as count of CPU ticks (aperiodic).

Files

- file [timer.h](#)
- file [timer.c](#)

Functions

- void [xntimer_tick_aperiodic](#) (void)
Process a timer tick for the aperiodic master time base.
- void [xntimer_tick_periodic](#) (xntimer_t *mtimer)
Process a timer tick for a slave periodic time base.
- void [xntimer_init](#) (xntimer_t *timer, xntbase_t *base, void(*handler)(xntimer_t *timer))
Initialize a timer object.
- void [xntimer_destroy](#) (xntimer_t *timer)
Release a timer object.
- void [xntimer_freeze](#) (void)
Freeze all timers (from every time bases).

5.11.2 Function Documentation

5.11.2.1 void xntimer_destroy (xntimer_t * timer)

Release a timer object.

Destroys a timer. After it has been destroyed, all resources associated with the timer have been released. The timer is automatically deactivated before deletion if active on entry.

Parameters:

timer The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.11.2.2 void xntimer_freeze (void)

Freeze all timers (from every time bases).

For internal use only.

This routine deactivates all active timers atomically.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

5.11.2.3 void xntimer_init (xntimer_t * timer, xntbase_t * base, void(*) (xntimer_t *timer) handler)

Initialize a timer object.

Creates a timer. When created, a timer is left disarmed; it must be started using [xntimer_start\(\)](#) in order to be activated.

Parameters:

timer The address of a timer descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

base The descriptor address of the time base the new timer depends on. See [xntbase_alloc\(\)](#) for detailed explanations about time bases.

handler The routine to call upon expiration of the timer.

There is no limitation on the number of timers which can be created/active concurrently.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

5.11.2.4 void xntimer_tick_aperiodic (void)

Process a timer tick for the aperiodic master time base.

For internal use only.

This routine informs all active timers that the clock has been updated by processing the outstanding timer list. Elapsed timer actions will be fired.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

5.11.2.5 void xntimer_tick_periodic (xntimer_t * mtimer)

Process a timer tick for a slave periodic time base.

For internal use only.

The periodic timer tick is cascaded from a software timer managed from the master aperiodic time base; in other words, periodic timing is emulated by software timers running in aperiodic timing mode. There may be several concurrent periodic time bases (albeit a single aperiodic time base - i.e. the master one called "nktbase" - may exist at any point in time).

This routine informs all active timers that the clock has been updated by processing the timer wheel. Elapsed timer actions will be fired.

Parameters:

mtimer The address of the cascading timer running in the master time base which announced the tick.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

Note:

Only active timers are inserted into the timer wheel.

5.12 HAL.

5.12.1 Detailed Description

Generic Adeos-based hardware abstraction layer.

Files

- file [hal.c](#)
Adeos-based Real-Time Abstraction Layer for PowerPC.
- file [hal.c](#)
Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.
- file [nmi.c](#)
NMI watchdog support.
- file [hal.c](#)
Generic Real-Time HAL.
- file [nmi.c](#)
Adeos-based Real-Time Abstraction Layer for x86.
- file [hal.c](#)
Adeos-based Real-Time Abstraction Layer for x86.
- file [nmi.c](#)
NMI watchdog for x86, from linux/arch/i386/kernel/nmi.c.
- file [smi.c](#)
SMI workaround for x86.
- file [hal.c](#)
Adeos-based Real-Time Abstraction Layer for ia64.
- file [hal.c](#)
Adeos-based Real-Time Abstraction Layer for PowerPC.

Functions

- int [rthal_timer_request](#) (void(*handler)(void))
Grab the hardware timer.
- void [rthal_timer_release](#) (void)
Release the hardware timer.
- int [rthal_irq_host_request](#) (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)

Install a shared Linux interrupt handler.

- int [rthal_irq_host_release](#) (unsigned irq, void *dev_id)
Uninstall a shared Linux interrupt handler.
- int [rthal_irq_enable](#) (unsigned irq)
Enable an interrupt source.
- int [rthal_irq_disable](#) (unsigned irq)
Disable an interrupt source.
- int [rthal_irq_request](#) (unsigned irq, rthal_irq_handler_t handler, rthal_irq_ackfn_t ackfn, void *cookie)
Install a real-time interrupt handler.
- int [rthal_irq_release](#) (unsigned irq)
Uninstall a real-time interrupt handler.
- int [rthal_irq_host_pend](#) (unsigned irq)
Propagate an IRQ event to Linux.
- int [rthal_irq_affinity](#) (unsigned irq, cpumask_t cpumask, cpumask_t *oldmask)
Set/Get processor affinity for external interrupt.
- rthal_trap_handler_t [rthal_trap_catch](#) (rthal_trap_handler_t handler)
Installs a fault handler.
- int [rthal_apc_alloc](#) (const char *name, void(*handler)(void *cookie), void *cookie)
Allocate an APC slot.
- int [rthal_apc_free](#) (int apc)
Releases an APC slot.
- int [rthal_apc_schedule](#) (int apc)
Schedule an APC invocation.

5.12.2 Function Documentation

5.12.2.1 int [rthal_apc_alloc](#) (const char * name, void(*) (void *cookie) handler, void * cookie)

Allocate an APC slot.

APC is the acronym for Asynchronous Procedure Call, a mean by which activities from the Xenomai domain can schedule deferred invocations of handlers to be run into the Linux domain, as soon as possible when the Linux kernel gets back in control. Up to BITS_PER_LONG APC slots can be active at any point in time. APC support is built upon Adeos's virtual interrupt support.

The HAL guarantees that any Linux kernel service which would be callable from a regular Linux interrupt handler is also available to APC handlers, including over PREEMPT_RT kernels exhibiting a threaded IRQ model.

Parameters:

name is a symbolic name identifying the APC which will get reported through the /proc/xenomai/apc interface. Passing NULL to create an anonymous APC is allowed.

handler The address of the fault handler to call upon exception condition. The handle will be passed the *cookie* value unmodified.

cookie A user-defined opaque cookie the HAL will pass to the APC handler as its sole argument.

Returns:

an valid APC id. is returned upon success, or a negative error code otherwise:

- -EINVAL is returned if *handler* is invalid.
- -EBUSY is returned if no more APC slots are available.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.2 int rthal_apc_free (int *apc*)

Releases an APC slot.

This service deallocates an APC slot obtained by [rthal_apc_alloc\(\)](#).

Parameters:

apc The APC id. to release, as returned by a successful call to the [rthal_apc_alloc\(\)](#) service.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *apc* is invalid.

Environments:

This service can be called from:

- Any domain context.

5.12.2.3 int rthal_apc_schedule (int *apc*)

Schedule an APC invocation.

This service marks the APC as pending for the Linux domain, so that its handler will be called as soon as possible, when the Linux domain gets back in control.

When posted from the Linux domain, the APC handler is fired as soon as the interrupt mask is explicitly cleared by some kernel code. When posted from the Xenomai domain, the APC handler is fired as soon as the Linux domain is resumed, i.e. after Xenomai has completed all its pending duties.

Parameters:

apc The APC id. to schedule.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *apc* is invalid.

Environments:

This service can be called from:

- Any domain context, albeit the usual calling place is from the Xenomai domain.

5.12.2.4 int rthal_irq_affinity (unsigned irq, cpumask_t cpumask, cpumask_t * oldmask)

Set/Get processor affinity for external interrupt.

On SMP systems, this service ensures that the given interrupt is preferably dispatched to the specified set of processors. The previous affinity mask is returned by this service.

Parameters:

irq The interrupt source whose processor affinity is affected by the operation. Only external interrupts can have their affinity changed/queried, thus virtual interrupt numbers allocated by `rthal_alloc_virq()` are invalid values for this parameter.

cpumask A list of CPU identifiers passed as a bitmask representing the new affinity for this interrupt. A zero value cause this service to return the current affinity mask without changing it.

oldmask If non-NULL, a pointer to a memory area which will be overwritten by the previous affinity mask used for this interrupt source, or a zeroed mask if an error occurred. This service always returns a zeroed mask on uniprocessor systems.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.5 int rthal_irq_disable (unsigned irq)

Disable an interrupt source.

Disables an interrupt source at PIC level. After this call has returned, no more IRQs from the given source will be allowed, until the latter is enabled again using [rthal_irq_enable\(\)](#).

Parameters:

irq The interrupt source to disable. This value is architecture-dependent.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

5.12.2.6 int rthal_irq_enable (unsigned irq)

Enable an interrupt source.

Enables an interrupt source at PIC level. Since Adeos masks and acknowledges the associated interrupt source upon IRQ receipt, this action is usually needed whenever the HAL handler does not propagate the IRQ event to the Linux domain, thus preventing the regular Linux interrupt handling code from re-enabling said source. After this call has returned, IRQs from the given source will be enabled again.

Parameters:

irq The interrupt source to enable. This value is architecture-dependent.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

5.12.2.7 int rthal_irq_host_pend (unsigned irq)

Propagate an IRQ event to Linux.

Causes the given IRQ to be propagated down to the Adeos pipeline to the Linux kernel. This operation is typically used after the given IRQ has been processed into the Xenomai domain by a real-time interrupt handler (see [rthal_irq_request\(\)](#)), in case such interrupt must also be handled by the Linux kernel.

Parameters:

irq The interrupt source to detach the shared handler from. This value is architecture-dependent.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.

Environments:

This service can be called from:

- Xenomai domain context.

5.12.2.8 int rthal_irq_host_release (unsigned irq, void * dev_id)

Uninstall a shared Linux interrupt handler.

Uninstalls a shared interrupt handler from the Linux domain for the given interrupt source. The handler is removed from the existing list of Linux handlers for this interrupt source.

Parameters:

irq The interrupt source to detach the shared handler from. This value is architecture-dependent.

dev_id is a valid device id, identical in essence to the one requested by the `free_irq()` service provided by the Linux kernel. This value will be used to locate the handler to remove from the chain of existing Linux handlers for the given interrupt source. This parameter must match the device id. passed to `rthal_irq_host_request()` for the same handler instance.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.9 int rthal_irq_host_request (unsigned irq, rthal_irq_host_handler_t handler, char * name, void * dev_id)

Install a shared Linux interrupt handler.

Installs a shared interrupt handler in the Linux domain for the given interrupt source. The handler is appended to the existing list of Linux handlers for this interrupt source.

Parameters:

irq The interrupt source to attach the shared handler to. This value is architecture-dependent.

handler The address of a valid interrupt service routine. This handler will be called each time the corresponding IRQ is delivered, as part of the chain of existing regular Linux handlers for this interrupt source. The handler prototype is the same as the one required by the `request_irq()` service provided by the Linux kernel.

name is a symbolic name identifying the handler which will get reported through the `/proc/interrupts` interface.

dev_id is a unique device id, identical in essence to the one requested by the `request_irq()` service.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid or *handler* is NULL.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.10 `int rthal_irq_release (unsigned irq)`

Uninstall a real-time interrupt handler.

Uninstalls an interrupt handler previously attached using the `rthal_irq_request()` service.

Parameters:

irq The hardware interrupt channel to uninstall a handler from. This value is architecture-dependent.

Returns:

0 is returned upon success. Otherwise:

- -EINVAL is returned if *irq* is invalid.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

5.12.2.11 `int rthal_irq_request (unsigned irq, rthal_irq_handler_t handler, rthal_irq_ackfn_t ackfn, void * cookie)`

Install a real-time interrupt handler.

Installs an interrupt handler for the specified IRQ line by requesting the appropriate Adeos virtualization service. The handler is invoked by Adeos on behalf of the Xenomai domain context. Once installed, the HAL interrupt handler will be called prior to the regular Linux handler for the same interrupt source.

Parameters:

- irq* The hardware interrupt channel to install a handler on. This value is architecture-dependent.
- handler* The address of a valid interrupt service routine. This handler will be called each time the corresponding IRQ is delivered, and will be passed the *cookie* value unmodified.
- ackfn* The address of an optional interrupt acknowledge routine, aimed at replacing the one provided by Adeos. Only very specific situations actually require to override the default Adeos setting for this parameter, like having to acknowledge non-standard PIC hardware. *ackfn* should return a non-zero value to indicate that the interrupt has been properly acknowledged. If *ackfn* is NULL, the default Adeos routine will be used instead.
- cookie* A user-defined opaque cookie the HAL will pass to the interrupt handler as its sole argument.

Returns:

0 is returned upon success. Otherwise:

- -EBUSY is returned if an interrupt handler is already installed. [rthal_irq_release\(\)](#) must be issued first before a handler is installed anew.
- -EINVAL is returned if *irq* is invalid or *handler* is NULL.
- Other error codes might be returned in case some internal error happens at the Adeos level. Such error might be caused by conflicting Adeos requests made by third-party code.

Environments:

This service can be called from:

- Any domain context.

5.12.2.12 void rthal_timer_release (void)

Release the hardware timer.

Releases the hardware timer, thus reverting the effect of a previous call to [rthal_timer_request\(\)](#). In case the timer hardware is shared with Linux, a periodic setup suitable for the Linux kernel will be reset.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.13 int rthal_timer_request (void*)(void) handler)

Grab the hardware timer.

[rthal_timer_request\(\)](#) grabs and tunes the hardware timer in oneshot mode in order to clock the master time base.

A user-defined routine is registered as the clock tick handler. This handler will always be invoked on behalf of the Xenomai domain for each incoming tick.

Parameters:

handler The address of the tick handler which will process each incoming tick.

Returns:

0 is returned on success. Otherwise:

- -EBUSY is returned if the hardware timer has already been grabbed. [rthal_timer_request\(\)](#) must be issued before [rthal_timer_request\(\)](#) is called again.

Environments:

This service can be called from:

- Linux domain context.

5.12.2.14 int rthal_trap_catch (rthal_trap_handler_t handler)

Installs a fault handler.

The HAL attempts to invoke a fault handler whenever an uncontrolled exception or fault is caught at machine level. This service allows to install a user-defined handler for such events.

Parameters:

handler The address of the fault handler to call upon exception condition. The handler is passed the address of the low-level information block describing the fault as passed by Adeos. Its layout is implementation-dependent.

Returns:

The address of the fault handler previously installed.

Environments:

This service can be called from:

- Any domain context.

Chapter 6

Xenomai nanokernel API Data Structure Documentation

6.1 xnpod Struct Reference

```
#include <pod.h>
```

6.1.1 Detailed Description

Real-time pod descriptor.

The source of all Xenomai magic.

Data Fields

- `xnflags_t` [status](#)
- `xntimer_t` [htimer](#)
- `xnsched_t` [sched](#) [XNARCH_NR_CPUS]
- `xnqueue_t` [threadq](#)
- `int` [threadq_rev](#)
- `volatile u_long` [schedlck](#)
- `xnqueue_t` [tstartq](#)
- `xnqueue_t` [tswitchq](#)
- `xnqueue_t` [tdeleteq](#)
- `int` [minpri](#)
- `int` [maxpri](#)
- `int` [root_prio_base](#)
- `int` [refcnt](#)
- `atomic_counter_t` [timerlck](#)
- `struct` {
 - `void(*` [settime](#) `)(xntbase_t *tbase, xnticks_t newtime)`
 - `int(*` [faulthandler](#) `)(xnarch_flinfo_t *flinfo)`
 - `int(*` [unload](#) `)(void)`
- `svctable`

6.1.2 Field Documentation

6.1.2.1 `int(* xnpod::faulthandler)(xnarch_fltinfo_t *fltinfo)`

Trap/exception handler.

6.1.2.2 `xntimer_t xnpod::htimer`

Host timer.

6.1.2.3 `int xnpod::maxpri`

Maximum priority value.

6.1.2.4 `int xnpod::minpri`

Minimum priority value.

6.1.2.5 `int xnpod::refcnt`

Reference count.

6.1.2.6 `int xnpod::root_prio_base`

Base priority of ROOT thread.

6.1.2.7 `xnsched_t xnpod::sched[XNARCH_NR_CPUS]`

Per-cpu scheduler slots.

6.1.2.8 `volatile u_long xnpod::schedlck`

Scheduler lock count.

6.1.2.9 `void(* xnpod::settime)(xntbase_t *tbase, xnticks_t newtime)`

Clock setting hook.

6.1.2.10 `xnflags_t xnpod::status`

Status bitmask.

6.1.2.11 `struct { ... } xnpod::svctable`

Table of overridable service entry points.

6.1.2.12 `xnqueue_t` [xnpod::tdeleteq](#)

Thread delete hook queue.

6.1.2.13 `xnqueue_t` [xnpod::threadq](#)

All existing threads.

6.1.2.14 `int` [xnpod::threadq_rev](#)

Modification counter of threadq.

6.1.2.15 `atomic_counter_t` [xnpod::timerlck](#)

Timer lock depth.

6.1.2.16 `xnqueue_t` [xnpod::tstartq](#)

Thread start hook queue.

6.1.2.17 `xnqueue_t` [xnpod::tswitchq](#)

Thread switch hook queue.

6.1.2.18 `int(*` [xnpod::unload](#)`)(void)`

Unloading hook.

The documentation for this struct was generated from the following file:

- `include/nucleus/pod.h`

6.2 xnsched Struct Reference

```
#include <pod.h>
```

6.2.1 Detailed Description

Scheduling information structure.

Data Fields

- `xnflags_t` [status](#)
- `xnthread_t *` [runthread](#)
- `xnarch_cpumask_t` [resched](#)
- `xnsched_queue_t` [readyq](#)
- volatile unsigned [inesting](#)
- `xnthread_t` [rootcb](#)

6.2.2 Field Documentation

6.2.2.1 volatile unsigned [xnsched::inesting](#)

Interrupt nesting level.

6.2.2.2 `xnsched_queue_t` [xnsched::readyq](#)

Ready-to-run threads (prioritized).

6.2.2.3 `xnarch_cpumask_t` [xnsched::resched](#)

Mask of CPUs needing rescheduling.

6.2.2.4 `xnthread_t` [xnsched::rootcb](#)

Root thread control block.

6.2.2.5 `xnthread_t*` [xnsched::runthread](#)

Current thread (service or user).

6.2.2.6 `xnflags_t` [xnsched::status](#)

Scheduler specific status bitmask

The documentation for this struct was generated from the following file:

- `include/nucleus/pod.h`

Chapter 7

Xenomai nanokernel API File Documentation

7.1 include/nucleus/pod.h File Reference

7.1.1 Detailed Description

Real-time pod interface header.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003,2004,2005 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2004 The Xenomai project <<http://www.xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/thread.h>
```

```
#include <nucleus/intr.h>
```

Data Structures

- struct [xnsched](#)

Scheduling information structure.

- struct [xnpod](#)
Real-time pod descriptor.

Typedefs

- typedef [xnsched](#) [xnsched_t](#)
Scheduling information structure.

Functions

- void [xnpod_schedule_runnable](#) (xnthread_t *thread, int flags)
Hidden rescheduling procedure.
- int [xnpod_init](#) (xnpod_t *pod, int minpri, int maxpri, xnflags_t flags)
Initialize a new pod.
- int [xnpod_enable_timesource](#) (void)
Activate the core time source.
- void [xnpod_disable_timesource](#) (void)
Stop the core time source.
- void [xnpod_shutdown](#) (int xtype)
Shutdown the current pod.
- int [xnpod_init_thread](#) (xnthread_t *thread, xntbase_t *tbase, const char *name, int prio, xnflags_t flags, unsigned stacksize)
Initialize a new thread.
- int [xnpod_start_thread](#) (xnthread_t *thread, xnflags_t mode, int imask, xnarch_cpumask_t affinity, void(*entry)(void *cookie), void *cookie)
Initial start of a newly created thread.
- void [xnpod_restart_thread](#) (xnthread_t *thread)
Restart a thread.
- void [xnpod_delete_thread](#) (xnthread_t *thread)
Delete a thread.
- xnflags_t [xnpod_set_thread_mode](#) (xnthread_t *thread, xnflags_t clrmask, xnflags_t set-mask)
Change a thread's control mode.
- void [xnpod_suspend_thread](#) (xnthread_t *thread, xnflags_t mask, xnticks_t timeout, int mode, struct xnsynch *wchan)
Suspend a thread.

- void [xnpod_resume_thread](#) (xnthread_t *thread, xnflags_t mask)
Resume a thread.
- int [xnpod_unblock_thread](#) (xnthread_t *thread)
Unblock a thread.
- void [xnpod_renice_thread](#) (xnthread_t *thread, int prio)
Change the base priority of a thread.
- int [xnpod_migrate_thread](#) (int cpu)
Migrate the current thread.
- void [xnpod_rotate_readyq](#) (int prio)
Rotate a priority level in the ready queue.
- void [xnpod_do_rr](#) (void)
Handle the round-robin scheduling policy.
- void [xnpod_schedule](#) (void)
Rescheduling procedure entry point.
- void [xnpod_dispatch_signals](#) (void)
Deliver pending asynchronous signals to the running thread.
- int [xnpod_announce_tick](#) (struct xnintr *intr)
Announce a new clock tick.
- void [xnpod_activate_rr](#) (xnticks_t quantum)
Globally activate the round-robin scheduling.
- void [xnpod_deactivate_rr](#) (void)
Globally deactivate the round-robin scheduling.
- void [xnpod_set_time](#) (xntbase_t *tbase, xnticks_t newtime)
Set the nucleus idea of time for a given time base.
- int [xnpod_set_thread_periodic](#) (xnthread_t *thread, xnticks_t idate, xnticks_t period)
Make a thread periodic.
- int [xnpod_wait_thread_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.
- xnticks_t [xnpod_get_time](#) (xntbase_t *tbase)
Get the nucleus idea of time for a given time base.
- int [xnpod_add_hook](#) (int type, void(*routine)(xnthread_t *))
Install a nucleus hook.
- int [xnpod_remove_hook](#) (int type, void(*routine)(xnthread_t *))
Remove a nucleus hook.

7.2 include/nucleus/registry.h File Reference

7.2.1 Detailed Description

This file is part of the Xenomai project.

Note:

Copyright (C) 2004 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/types.h>
#include <nucleus/synch.h>
#include <nucleus/thread.h>
```

Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, xnnode_t *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_remove_safe](#) (xnhandle_t handle, xnticks_t timeout)
Unregister an idle real-time object.
- void * [xnregistry_get](#) (xnhandle_t handle)
Find and lock a real-time object into the registry.
- void * [xnregistry_fetch](#) (xnhandle_t handle)
Find a real-time object into the registry.
- u_long [xnregistry_put](#) (xnhandle_t handle)
Unlock a real-time object from the registry.

7.3 include/nucleus/timer.h File Reference

7.3.1 Detailed Description

Note:

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/timebase.h>
```

Functions

- void [xntimer_init](#) (xntimer_t *timer, xntbase_t *base, void(*handler)(xntimer_t *timer))
Initialize a timer object.
- void [xntimer_destroy](#) (xntimer_t *timer)
Release a timer object.
- static int [xntimer_start](#) (xntimer_t *timer, xnticks_t value, xnticks_t interval, int mode)
Arm a timer.
- static void [xntimer_stop](#) (xntimer_t *timer)
Disarm a timer.
- static xnticks_t [xntimer_get_date](#) (xntimer_t *timer)
Return the absolute expiration date.
- static xnticks_t [xntimer_get_timeout](#) (xntimer_t *timer)
Return the relative expiration date.
- static xnticks_t [xntimer_get_interval](#) (xntimer_t *timer)
Return the timer interval value.
- void [xntimer_freeze](#) (void)
Freeze all timers (from every time bases).
- void [xntimer_tick_aperiodic](#) (void)
Process a timer tick for the aperiodic master time base.
- void [xntimer_tick_periodic](#) (xntimer_t *timer)
Process a timer tick for a slave periodic time base.

7.3.2 Function Documentation

7.3.2.1 `xnticks_t xntimer_get_date (xntimer_t * timer)` [inline, static]

Return the absolute expiration date.

Return the next expiration date of a timer in absolute clock ticks (see note).

Parameters:

timer The address of a valid timer descriptor.

Returns:

The expiration date converted to the current time unit. The special value XN_INFINITE is returned if *timer* is currently inactive.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Note:

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

7.3.2.2 `xnticks_t xntimer_get_interval (xntimer_t * timer)` [inline, static]

Return the timer interval value.

Return the timer interval value in clock ticks (see note).

Parameters:

timer The address of a valid timer descriptor.

Returns:

The expiration date converted to the current time unit. The special value XN_INFINITE is returned if *timer* is currently inactive or aperiodic.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine

- Kernel-based task
- User-space task

Rescheduling: never.

Note:

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

7.3.2.3 `xnticks_t xntimer_get_timeout(xntimer_t * timer)` [inline, static]

Return the relative expiration date.

Return the next expiration date of a timer in relative clock ticks (see note).

Parameters:

timer The address of a valid timer descriptor.

Returns:

The expiration date converted to the current time unit. The special value `XN_INFINITE` is returned if *timer* is currently inactive. In oneshot mode, it might happen that the timer has already expired when this service is run (even if the associated handler has not been fired yet); in such a case, 1 is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Note:

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds.

7.3.2.4 `void xntimer_start(xntimer_t * timer, xnticks_t value, xnticks_t interval, int mode)` [inline, static]

Arm a timer.

Activates a timer so that the associated timeout handler will be fired after each expiration time. A timer can be either periodic or single-shot, depending on the reload value passed to this routine. The given timer must have been previously initialized by a call to `xntimer_init()`.

Parameters:

- timer* The address of a valid timer descriptor.
- value* The date of the initial timer shot, expressed in clock ticks (see note).
- interval* The reload value of the timer. It is a periodic interval value to be used for reprogramming the next timer shot, expressed in clock ticks (see note). If *interval* is equal to `XN_INFINITE`, the timer will not be reloaded after it has expired.
- mode* The timer mode. It can be either `XN_RELATIVE` or `XN_ABSOLUTE` to define if *value* shall be interpreted as a relative or absolute date. Absolute dates are based on the nucleus time returned by `xnpod_get_time()`.

Returns:

0 is returned upon success, or `-ETIMEDOUT` if an absolute date in the past has been given.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Note:

This service is sensitive to the current operation mode of the associated time base, as defined by the `xnpod_init_timebase()` service. In periodic mode, clock ticks are interpreted as periodic jiffies. In oneshot mode, clock ticks are interpreted as nanoseconds. Must be called with `nklock` held, IRQs off.

7.3.2.5 `int xntimer_stop (xntimer_t * timer)` [inline, static]

Disarm a timer.

This service deactivates a timer previously armed using `xntimer_start()`. Once disarmed, the timer can be subsequently re-armed using the latter service.

Parameters:

- timer* The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Note:

Must be called with nklock held, IRQs off.

7.4 ksrc/arch/arm/hal.c File Reference

7.4.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for PowerPC.

ARM port Copyright (C) 2005 Stelian Pop

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <linux/console.h>
#include <asm/system.h>
#include <asm/hardirq.h>
#include <asm/irq.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
#include <stdarg.h>
```

Functions

- `int rthal_timer_request (void(*handler)(void))`
Grab the hardware timer.
- `void rthal_timer_release (void)`
Release the hardware timer.
- `int rthal_irq_host_request (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)`
Install a shared Linux interrupt handler.
- `int rthal_irq_host_release (unsigned irq, void *dev_id)`
Uninstall a shared Linux interrupt handler.

- int [rthal_irq_enable](#) (unsigned irq)
Enable an interrupt source.
- int [rthal_irq_disable](#) (unsigned irq)
Disable an interrupt source.

7.5 ksrc/arch/blackfin/hal.c File Reference

7.5.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for the Blackfin architecture.

Copyright (C) 2005-2006 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <asm/system.h>
#include <asm/atomic.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
```

Functions

- int [rthal_timer_request](#) (void(*handler)(void))
Grab the hardware timer.
- void [rthal_timer_release](#) (void)
Release the hardware timer.
- int [rthal_irq_enable](#) (unsigned irq)
Enable an interrupt source.
- int [rthal_irq_disable](#) (unsigned irq)
Disable an interrupt source.
- int [rthal_irq_host_request](#) (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)
Install a shared Linux interrupt handler.

- int [rthal_irq_host_release](#) (unsigned irq, void *dev_id)
Uninstall a shared Linux interrupt handler.

7.6 ksrc/arch/generic/hal.c File Reference

7.6.1 Detailed Description

Generic Real-Time HAL.

Copyright ©2005 Philippe Gerum.

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <linux/console.h>
#include <linux/kallsyms.h>
#include <asm/system.h>
#include <asm/hardirq.h>
#include <asm/irq.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
#include <stdarg.h>
```

Functions

- int [rthal_irq_request](#) (unsigned irq, rthal_irq_handler_t handler, rthal_irq_ackfn_t ackfn, void *cookie)
Install a real-time interrupt handler.
- int [rthal_irq_release](#) (unsigned irq)
Uninstall a real-time interrupt handler.
- int [rthal_irq_host_pend](#) (unsigned irq)
Propagate an IRQ event to Linux.
- int [rthal_irq_affinity](#) (unsigned irq, cpumask_t cpumask, cpumask_t *oldmask)
Set/Get processor affinity for external interrupt.

- `rthal_trap_handler_t rthal_trap_catch` (`rthal_trap_handler_t handler`)
Installs a fault handler.
- `int rthal_apc_alloc` (`const char *name, void(*handler)(void *cookie), void *cookie`)
Allocate an APC slot.
- `int rthal_apc_free` (`int apc`)
Releases an APC slot.
- `int rthal_apc_schedule` (`int apc`)
Schedule an APC invocation.

7.7 ksrc/arch/i386/hal.c File Reference

7.7.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86.

Inspired from original RTAI/x86 HAL interface:

Copyright ©2000 Paolo Mantegazza,

Copyright ©2000 Steve Papacharalambous,

Copyright ©2000 Stuart Hughes,

RTAI/x86 rewrite over Adeos:

Copyright ©2002 Philippe Gerum. NMI watchdog, SMI workaround:

Copyright ©2004 Gilles Chantepredrix.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <linux/console.h>
#include <asm/system.h>
#include <asm/hardirq.h>
#include <asm/desc.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
#include <stdarg.h>
```

Functions

- int [rthal_timer_request](#) (void(*handler)(void))
Grab the hardware timer.
- void [rthal_timer_release](#) (void)

Release the hardware timer.

- int [rthal_irq_host_request](#) (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)

Install a shared Linux interrupt handler.

- int [rthal_irq_host_release](#) (unsigned irq, void *dev_id)

Uninstall a shared Linux interrupt handler.

- int [rthal_irq_enable](#) (unsigned irq)

Enable an interrupt source.

- int [rthal_irq_disable](#) (unsigned irq)

Disable an interrupt source.

7.8 ksrc/arch/ia64/hal.c File Reference

7.8.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for ia64.

Copyright ©2002-2004 Philippe Gerum Copyright ©2004 The HYADES project
<<http://www.hyades-itea.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <linux/console.h>
#include <asm/system.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
#include <stdarg.h>
```

Functions

- int [rthal_timer_request](#) (void(*handler)(void))
Grab the hardware timer.
- void [rthal_timer_release](#) (void)
Release the hardware timer.
- int [rthal_irq_host_request](#) (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)
Install a shared Linux interrupt handler.
- int [rthal_irq_host_release](#) (unsigned irq, void *dev_id)
Uninstall a shared Linux interrupt handler.
- int [rthal_irq_enable](#) (unsigned irq)

Enable an interrupt source.

- int [rthal_irq_disable](#) (unsigned irq)

Disable an interrupt source.

7.9 ksrc/arch/powerpc/hal.c File Reference

7.9.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for PowerPC.

64-bit PowerPC adoption copyright (C) 2005 Taneli VÃd'hÃd'kangas and Heikki Lindholm

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <linux/irq.h>
#include <linux/console.h>
#include <asm/system.h>
#include <asm/hardirq.h>
#include <asm/hw_irq.h>
#include <asm/irq.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
#include <stdarg.h>
```

Functions

- int [rthal_timer_request](#) (void(*handler)(void))
Grab the hardware timer.
- void [rthal_timer_release](#) (void)
Release the hardware timer.
- int [rthal_irq_host_request](#) (unsigned irq, rthal_irq_host_handler_t handler, char *name, void *dev_id)
Install a shared Linux interrupt handler.

- int [rthal_irq_host_release](#) (unsigned irq, void *dev_id)
Uninstall a shared Linux interrupt handler.
- int [rthal_irq_enable](#) (unsigned irq)
Enable an interrupt source.
- int [rthal_irq_disable](#) (unsigned irq)
Disable an interrupt source.

7.10 ksrc/arch/blackfin/nmi.c File Reference

7.10.1 Detailed Description

NMI watchdog support.

Copyright (C) 2005 Philippe Gerum.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/module.h>
#include <asm/system.h>
#include <asm/atomic.h>
#include <asm/irqchip.h>
#include <asm/io.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/xenomai/hal.h>
```

7.11 ksrc/arch/generic/nmi.c File Reference

7.11.1 Detailed Description

Adeos-based Real-Time Abstraction Layer for x86.

Copyright ©2005 Gilles Chantepredrix.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/version.h>
#include <linux/module.h>
#include <asm/system.h>
#include <asm/unistd.h>
#include <asm/uaccess.h>
#include <asm/xenomai/hal.h>
```

7.12 ksrc/arch/i386/nmi.c File Reference

7.12.1 Detailed Description

NMI watchdog for x86, from linux/arch/i386/kernel/nmi.c.

Original authors: Ingo Molnar, Mikael Pettersson, Pavel Machek.

Adaptation to Xenomai by Gilles Chanteperdrix

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/module.h>
#include <linux/version.h>
#include <linux/nmi.h>
#include <asm/nmi.h>
#include <asm/msr.h>
#include <asm/xenomai/hal.h>
```


7.13 ksrc/arch/i386/smi.c File Reference

7.13.1 Detailed Description

SMI workaround for x86.

Cut/Pasted from Vitor Angelo "smi" module. Adapted by Gilles Chanteperdrix <gilles.chanteperdrix@laposte.net>.

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/version.h>
#include <linux/pci.h>
#include <linux/pci_ids.h>
#include <linux/reboot.h>
#include <asm/xenomai/hal.h>
```

7.14 ksrc/nucleus/heap.c File Reference

7.14.1 Detailed Description

Dynamic memory allocation services.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/pod.h>
#include <nucleus/thread.h>
#include <nucleus/heap.h>
#include <asm/xenomai/bits/heap.h>
#include <asm/io.h>
#include <linux/miscdevice.h>
#include <linux/device.h>
#include <linux/vmalloc.h>
#include <linux/mm.h>
```

Functions

- int [xnheap_init](#) (xnheap_t *heap, void *heapaddr, u_long heapsize, u_long pagesize)
Initialize a memory heap.
- int [xnheap_destroy](#) (xnheap_t *heap, void(*flushfn)(xnheap_t *heap, void *extaddr, u_long extsize, void *cookie), void *cookie)
Destroys a memory heap.
- void * [xnheap_alloc](#) (xnheap_t *heap, u_long size)
Allocate a memory block from a memory heap.
- int [xnheap_test_and_free](#) (xnheap_t *heap, void *block, int(*ckfn)(void *block))
Test and release a memory block to a memory heap.
- int [xnheap_free](#) (xnheap_t *heap, void *block)

Release a memory block to a memory heap.

- int [xnheap_extend](#) (xnheap_t *heap, void *extaddr, u_long extsize)
Extend a memory heap.
- void [xnheap_schedule_free](#) (xnheap_t *heap, void *block, xnholder_t *link)
Schedule a memory block for release.

7.15 ksrc/nucleus/intr.c File Reference

7.15.1 Detailed Description

Interrupt management.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2005,2006 Dmitry Adamushko <dmitry.adamushko@gmail.com>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/pod.h>
#include <nucleus/intr.h>
#include <nucleus/ltt.h>
#include <nucleus/stat.h>
#include <asm/xenomai/bits/intr.h>
```

Functions

- int [xnintr_init](#) (xnintr_t *intr, const char *name, unsigned irq, xnintr_t isr, xniack_t iack, xnflags_t flags)
Initialize an interrupt object.
- int [xnintr_destroy](#) (xnintr_t *intr)
Destroy an interrupt object.
- int [xnintr_attach](#) (xnintr_t *intr, void *cookie)
Attach an interrupt object.
- int [xnintr_detach](#) (xnintr_t *intr)
Detach an interrupt object.
- int [xnintr_enable](#) (xnintr_t *intr)
Enable an interrupt object.
- int [xnintr_disable](#) (xnintr_t *intr)
Disable an interrupt object.

- `xnarch_cpumask_t xnintr_affinity` (`xnintr_t *intr`, `xnarch_cpumask_t cpumask`)
Set interrupt's processor affinity.

7.16 ksrc/nucleus/pod.c File Reference

7.16.1 Detailed Description

Real-time pod services.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003,2004,2005 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.Xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <stdarg.h>
#include <nucleus/pod.h>
#include <nucleus/timer.h>
#include <nucleus/synch.h>
#include <nucleus/heap.h>
#include <nucleus/intr.h>
#include <nucleus/registry.h>
#include <nucleus/module.h>
#include <nucleus/ltt.h>
#include <nucleus/stat.h>
#include <asm/xenomai/bits/pod.h>
#include <nucleus/assert.h>
```

Functions

- int [xnpod_init](#) (xnpod_t *pod, int minpri, int maxpri, xnflags_t flags)
Initialize a new pod.
- void [xnpod_shutdown](#) (int xtype)
Shutdown the current pod.
- int [xnpod_init_thread](#) (xnthread_t *thread, xnbase_t *tbase, const char *name, int prio, xnflags_t flags, unsigned stacksize)

Initialize a new thread.

- int [xnpod_start_thread](#) (xnthread_t *thread, xnflags_t mode, int imask, xnarch_cpumask_t affinity, void(*entry)(void *cookie), void *cookie)

Initial start of a newly created thread.

- void [xnpod_restart_thread](#) (xnthread_t *thread)

Restart a thread.

- xnflags_t [xnpod_set_thread_mode](#) (xnthread_t *thread, xnflags_t clrmask, xnflags_t set-mask)

Change a thread's control mode.

- void [xnpod_delete_thread](#) (xnthread_t *thread)

Delete a thread.

- void [xnpod_suspend_thread](#) (xnthread_t *thread, xnflags_t mask, xnticks_t timeout, int mode, struct xnsynch *wchan)

Suspend a thread.

- void [xnpod_resume_thread](#) (xnthread_t *thread, xnflags_t mask)

Resume a thread.

- int [xnpod_unblock_thread](#) (xnthread_t *thread)

Unblock a thread.

- void [xnpod_renice_thread](#) (xnthread_t *thread, int prio)

Change the base priority of a thread.

- int [xnpod_migrate_thread](#) (int cpu)

Migrate the current thread.

- void [xnpod_rotate_readyq](#) (int prio)

Rotate a priority level in the ready queue.

- void [xnpod_activate_rr](#) (xnticks_t quantum)

Globally activate the round-robin scheduling.

- void [xnpod_deactivate_rr](#) (void)

Globally deactivate the round-robin scheduling.

- void [xnpod_dispatch_signals](#) (void)

Deliver pending asynchronous signals to the running thread.

- void [xnpod_welcome_thread](#) (xnthread_t *thread, int imask)

Thread prologue.

- void [xnpod_do_rr](#) (void)

Handle the round-robin scheduling policy.

- static void [xnpod_preempt_current_thread](#) (xnsched_t *sched)
Preempts the current thread.
- void [xnpod_schedule](#) (void)
Rescheduling procedure entry point.
- void [xnpod_schedule_runnable](#) (xnthread_t *thread, int flags)
Hidden rescheduling procedure.
- void [xnpod_set_time](#) (xntbase_t *tbase, xnticks_t newtime)
Set the nucleus idea of time for a given time base.
- xnticks_t [xnpod_get_time](#) (xntbase_t *tbase)
Get the nucleus idea of time for a given time base.
- int [xnpod_add_hook](#) (int type, void(*routine)(xnthread_t *))
Install a nucleus hook.
- int [xnpod_remove_hook](#) (int type, void(*routine)(xnthread_t *))
Remove a nucleus hook.
- int [xnpod_trap_fault](#) (void *fltinfo)
Default fault handler.
- int [xnpod_enable_timesource](#) (void)
Activate the core time source.
- void [xnpod_disable_timesource](#) (void)
Stop the core time source.
- int [xnpod_announce_tick](#) (struct xnintr *intr)
Announce a new clock tick.
- int [xnpod_set_thread_periodic](#) (xnthread_t *thread, xnticks_t idate, xnticks_t period)
Make a thread periodic.
- int [xnpod_wait_thread_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.

7.17 ksrc/nucleus/registry.c File Reference

7.17.1 Detailed Description

This file is part of the Xenomai project.

Note:

Copyright (C) 2004 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/pod.h>
#include <nucleus/heap.h>
#include <nucleus/registry.h>
#include <nucleus/thread.h>
```

Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, xnnode_t *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_remove_safe](#) (xnhandle_t handle, xnticks_t timeout)
Unregister an idle real-time object.
- void * [xnregistry_get](#) (xnhandle_t handle)
Find and lock a real-time object into the registry.
- u_long [xnregistry_put](#) (xnhandle_t handle)
Unlock a real-time object from the registry.
- void * [xnregistry_fetch](#) (xnhandle_t handle)
Find a real-time object into the registry.

7.18 ksrc/nucleus/shadow.c File Reference

7.18.1 Detailed Description

Real-time shadow services.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003,2004,2005,2006 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.xenomai.org>> Copyright (C) 2006 Gilles Chanteperdrix <gilles.chanteperdrix@laposte.net>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <stdarg.h>
#include <linux/unistd.h>
#include <linux/wait.h>
#include <linux/init.h>
#include <linux/kthread.h>
#include <asm/signal.h>
#include <nucleus/pod.h>
#include <nucleus/heap.h>
#include <nucleus/synch.h>
#include <nucleus/module.h>
#include <nucleus/shadow.h>
#include <nucleus/core.h>
#include <nucleus/ltt.h>
#include <nucleus/jhash.h>
#include <nucleus/ppd.h>
#include <nucleus/trace.h>
#include <nucleus/stat.h>
#include <asm/xenomai/features.h>
#include <asm/xenomai/syscall.h>
```

```
#include <asm/xenomai/bits/shadow.h>
```

```
#include <nucleus/assert.h>
```

Functions

- int [xnshadow_harden](#) (void)
Migrate a Linux task to the Xenomai domain.
- void [xnshadow_relax](#) (int notify)
Switch a shadow thread back to the Linux domain.
- int [xnshadow_map](#) (xnthread_t *thread, xncompletion_t __user *u_completion)
Create a shadow thread context.
- xnshadow_ppd_t * [xnshadow_ppd_get](#) (unsigned muxid)
Return the per-process data attached to the calling process.

7.19 ksrc/nucleus/synch.c File Reference

7.19.1 Detailed Description

Thread synchronization services.

Author:

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <stdarg.h>
#include <nucleus/pod.h>
#include <nucleus/synch.h>
#include <nucleus/thread.h>
#include <nucleus/module.h>
#include <nucleus/ltt.h>
```

Functions

- void [xnsynch_init](#) (xnsynch_t *synch, xnflags_t flags)
Initialize a synchronization object.
- void [xnsynch_sleep_on](#) (xnsynch_t *synch, xnticks_t timeout, int mode)
Sleep on a synchronization object.
- static void [xnsynch_clear_boost](#) (xnsynch_t *synch, xnthread_t *lastowner)
Clear the priority boost.
- void [xnsynch_renice_sleeper](#) (xnthread_t *thread)
Change a sleeper's priority.
- xnthread_t * [xnsynch_wakeup_one_sleeper](#) (xnsynch_t *synch)
Give the resource ownership to the next waiting thread.
- xnpholder_t * [xnsynch_wakeup_this_sleeper](#) (xnsynch_t *synch, xnpholder_t *holder)
Give the resource ownership to a given waiting thread.

- int [xnsynch_flush](#) (xnsynch_t *synch, xnflags_t reason)
Unblock all waiters pending on a resource.
- void [xnsynch_forget_sleeper](#) (xnthread_t *thread)
Abort a wait for a resource.
- void [xnsynch_release_all_ownerships](#) (xnthread_t *thread)
Release all ownerships.

7.20 ksrc/nucleus/timebase.c File Reference

7.20.1 Detailed Description

Note:

Copyright (C) 2006,2007 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/pod.h>
```

```
#include <nucleus/timer.h>
```

Functions

- int [xntbase_alloc](#) (const char *name, u_long period, xntbase_t **basep)
Allocate a time base.
- void [xntbase_free](#) (xntbase_t *base)
Free a time base.
- int [xntbase_update](#) (xntbase_t *base, u_long period)
Change the period of a time base.
- int [xntbase_switch](#) (const char *name, u_long period, xntbase_t **basep)
Replace a time base.
- void [xntbase_start](#) (xntbase_t *base)
Start a time base.
- void [xntbase_stop](#) (xntbase_t *base)
Stop a time base.
- void [xntbase_tick](#) (xntbase_t *base)
Announce a clock tick to a time base.

7.21 ksrc/nucleus/timer.c File Reference

7.21.1 Detailed Description

Note:

Copyright (C) 2001,2002,2003,2007 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 Gilles Chantepredrix <gilles.chantepredrix@laposte.net>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```
#include <nucleus/pod.h>
#include <nucleus/thread.h>
#include <nucleus/timer.h>
#include <asm/xenomai/bits/timer.h>
```

Functions

- void [xntimer_tick_aperiodic](#) (void)
Process a timer tick for the aperiodic master time base.
- void [xntimer_tick_periodic](#) (xntimer_t *mtimer)
Process a timer tick for a slave periodic time base.
- void [xntimer_init](#) (xntimer_t *timer, xntbase_t *base, void(*handler)(xntimer_t *timer))
Initialize a timer object.
- void [xntimer_destroy](#) (xntimer_t *timer)
Release a timer object.
- void [xntimer_freeze](#) (void)
Freeze all timers (from every time bases).

Index

Dynamic memory allocation services., [12](#)

faulthandler
 xnpod, [82](#)

hal
 rthal_apc_alloc, [72](#)
 rthal_apc_free, [73](#)
 rthal_apc_schedule, [73](#)
 rthal_irq_affinity, [74](#)
 rthal_irq_disable, [74](#)
 rthal_irq_enable, [75](#)
 rthal_irq_host_pend, [75](#)
 rthal_irq_host_release, [76](#)
 rthal_irq_host_request, [76](#)
 rthal_irq_release, [77](#)
 rthal_irq_request, [77](#)
 rthal_timer_release, [78](#)
 rthal_timer_request, [78](#)
 rthal_trap_catch, [79](#)

HAL., [71](#)

heap
 xnheap_alloc, [13](#)
 xnheap_destroy, [13](#)
 xnheap_extend, [14](#)
 xnheap_free, [14](#)
 xnheap_init, [15](#)
 xnheap_schedule_free, [16](#)
 xnheap_test_and_free, [16](#)

htimer
 xnpod, [82](#)

include/nucleus/pod.h, [85](#)
include/nucleus/registry.h, [88](#)
include/nucleus/timer.h, [89](#)

inesting
 xnsched, [84](#)

Interrupt management., [18](#)

intr
 xnintr_affinity, [18](#)
 xnintr_attach, [19](#)
 xnintr_destroy, [19](#)
 xnintr_detach, [20](#)
 xnintr_disable, [20](#)
 xnintr_enable, [21](#)

 xnintr_init, [21](#)

ksrc/arch/arm/hal.c, [94](#)
ksrc/arch/blackfin/hal.c, [96](#)
ksrc/arch/blackfin/nmi.c, [106](#)
ksrc/arch/generic/hal.c, [98](#)
ksrc/arch/generic/nmi.c, [107](#)
ksrc/arch/i386/hal.c, [100](#)
ksrc/arch/i386/nmi.c, [108](#)
ksrc/arch/i386/smi.c, [109](#)
ksrc/arch/ia64/hal.c, [102](#)
ksrc/arch/powerpc/hal.c, [104](#)
ksrc/nucleus/heap.c, [110](#)
ksrc/nucleus/intr.c, [112](#)
ksrc/nucleus/pod.c, [114](#)
ksrc/nucleus/registry.c, [117](#)
ksrc/nucleus/shadow.c, [118](#)
ksrc/nucleus/synch.c, [120](#)
ksrc/nucleus/timebase.c, [122](#)
ksrc/nucleus/timer.c, [123](#)

maxpri
 xnpod, [82](#)

minpri
 xnpod, [82](#)

pod
 xnpod_activate_rr, [27](#)
 xnpod_add_hook, [28](#)
 xnpod_announce_tick, [28](#)
 xnpod_deactivate_rr, [29](#)
 xnpod_delete_thread, [29](#)
 xnpod_disable_timesource, [30](#)
 xnpod_dispatch_signals, [30](#)
 xnpod_do_rr, [30](#)
 xnpod_enable_timesource, [30](#)
 xnpod_get_time, [31](#)
 xnpod_init, [32](#)
 xnpod_init_thread, [33](#)
 xnpod_migrate_thread, [34](#)
 xnpod_preempt_current_thread, [35](#)
 xnpod_remove_hook, [35](#)
 xnpod_renice_thread, [35](#)
 xnpod_restart_thread, [36](#)
 xnpod_resume_thread, [36](#)

- xnpod_rotate_readyq, 37
 - xnpod_schedule, 38
 - xnpod_schedule_runnable, 39
 - xnpod_set_thread_mode, 39
 - xnpod_set_thread_periodic, 40
 - xnpod_set_time, 41
 - xnpod_shutdown, 42
 - xnpod_start_thread, 42
 - xnpod_suspend_thread, 43
 - xnpod_trap_fault, 45
 - xnpod_unblock_thread, 45
 - xnpod_wait_thread_period, 46
 - xnpod_welcome_thread, 46
- readyq
 - xnsched, 84
- Real-time pod services., 25
- Real-time shadow services., 54
- refcnt
 - xnpod, 82
- registry
 - xnregistry_bind, 48
 - xnregistry_enter, 49
 - xnregistry_fetch, 50
 - xnregistry_get, 51
 - xnregistry_put, 51
 - xnregistry_remove, 52
 - xnregistry_remove_safe, 52
- Registry services., 48
- resched
 - xnsched, 84
- root_prio_base
 - xnpod, 82
- rootcb
 - xnsched, 84
- rthal_apc_alloc
 - hal, 72
- rthal_apc_free
 - hal, 73
- rthal_apc_schedule
 - hal, 73
- rthal_irq_affinity
 - hal, 74
- rthal_irq_disable
 - hal, 74
- rthal_irq_enable
 - hal, 75
- rthal_irq_host_pend
 - hal, 75
- rthal_irq_host_release
 - hal, 76
- rthal_irq_host_request
 - hal, 76
- rthal_irq_release
 - hal, 77
- rthal_irq_request
 - hal, 77
- rthal_timer_release
 - hal, 78
- rthal_timer_request
 - hal, 78
- rthal_trap_catch
 - hal, 79
- runthread
 - xnsched, 84
- sched
 - xnpod, 82
- schedlck
 - xnpod, 82
- settime
 - xnpod, 82
- shadow
 - xnshadow_harden, 54
 - xnshadow_map, 54
 - xnshadow_ppd_get, 55
 - xnshadow_relax, 56
- status
 - xnpod, 82
 - xnsched, 84
- svctable
 - xnpod, 82
- synch
 - xnsynch_clear_boost, 57
 - xnsynch_flush, 58
 - xnsynch_forget_sleeper, 58
 - xnsynch_init, 59
 - xnsynch_release_all_ownerships, 60
 - xnsynch_renice_sleeper, 60
 - xnsynch_sleep_on, 60
 - xnsynch_wakeup_one_sleeper, 61
 - xnsynch_wakeup_this_sleeper, 61
- tdeleteq
 - xnpod, 82
- Thread information flags., 11
- Thread state flags., 9
- Thread synchronization services., 57
- threadq
 - xnpod, 83
- threadq_rev
 - xnpod, 83
- Time base services., 63
- timebase
 - xntbase_alloc, 64
 - xntbase_free, 64
 - xntbase_start, 65
 - xntbase_stop, 65

- xntbase_switch, [66](#)
- xntbase_tick, [66](#)
- xntbase_update, [67](#)
- timer
 - xntimer_destroy, [68](#)
 - xntimer_freeze, [69](#)
 - xntimer_init, [69](#)
 - xntimer_tick_aperiodic, [69](#)
 - xntimer_tick_periodic, [70](#)
- Timer services., [68](#)
- timer.h
 - xntimer_get_date, [90](#)
 - xntimer_get_interval, [90](#)
 - xntimer_get_timeout, [91](#)
 - xntimer_start, [91](#)
 - xntimer_stop, [92](#)
- timerlck
 - xnpod, [83](#)
- tstartq
 - xnpod, [83](#)
- tswitchq
 - xnpod, [83](#)
- unload
 - xnpod, [83](#)
- Xenomai nucleus., [24](#)
- xnheap_alloc
 - heap, [13](#)
- xnheap_destroy
 - heap, [13](#)
- xnheap_extend
 - heap, [14](#)
- xnheap_free
 - heap, [14](#)
- xnheap_init
 - heap, [15](#)
- xnheap_schedule_free
 - heap, [16](#)
- xnheap_test_and_free
 - heap, [16](#)
- xnintr_affinity
 - intr, [18](#)
- xnintr_attach
 - intr, [19](#)
- xnintr_destroy
 - intr, [19](#)
- xnintr_detach
 - intr, [20](#)
- xnintr_disable
 - intr, [20](#)
- xnintr_enable
 - intr, [21](#)
- xnintr_init
 - intr, [21](#)
- xnpod, [81](#)
 - faulthandler, [82](#)
 - htimer, [82](#)
 - maxpri, [82](#)
 - minpri, [82](#)
 - refcnt, [82](#)
 - root_prio_base, [82](#)
 - sched, [82](#)
 - schedlck, [82](#)
 - settime, [82](#)
 - status, [82](#)
 - svctable, [82](#)
 - tdeleteq, [82](#)
 - threadq, [83](#)
 - threadq_rev, [83](#)
 - timerlck, [83](#)
 - tstartq, [83](#)
 - tswitchq, [83](#)
 - unload, [83](#)
- xnpod_activate_rr
 - pod, [27](#)
- xnpod_add_hook
 - pod, [28](#)
- xnpod_announce_tick
 - pod, [28](#)
- xnpod_deactivate_rr
 - pod, [29](#)
- xnpod_delete_thread
 - pod, [29](#)
- xnpod_disable_timesource
 - pod, [30](#)
- xnpod_dispatch_signals
 - pod, [30](#)
- xnpod_do_rr
 - pod, [30](#)
- xnpod_enable_timesource
 - pod, [30](#)
- xnpod_get_time
 - pod, [31](#)
- xnpod_init
 - pod, [32](#)
- xnpod_init_thread
 - pod, [33](#)
- xnpod_migrate_thread
 - pod, [34](#)
- xnpod_preempt_current_thread
 - pod, [35](#)
- xnpod_remove_hook
 - pod, [35](#)
- xnpod_renice_thread
 - pod, [35](#)
- xnpod_restart_thread
 - pod, [36](#)

xnpod_resume_thread
 pod, 36
xnpod_rotate_readyq
 pod, 37
xnpod_schedule
 pod, 38
xnpod_schedule_runnable
 pod, 39
xnpod_set_thread_mode
 pod, 39
xnpod_set_thread_periodic
 pod, 40
xnpod_set_time
 pod, 41
xnpod_shutdown
 pod, 42
xnpod_start_thread
 pod, 42
xnpod_suspend_thread
 pod, 43
xnpod_trap_fault
 pod, 45
xnpod_unblock_thread
 pod, 45
xnpod_wait_thread_period
 pod, 46
xnpod_welcome_thread
 pod, 46
xnregistry_bind
 registry, 48
xnregistry_enter
 registry, 49
xnregistry_fetch
 registry, 50
xnregistry_get
 registry, 51
xnregistry_put
 registry, 51
xnregistry_remove
 registry, 52
xnregistry_remove_safe
 registry, 52
xnsched, 84
 inesting, 84
 readyq, 84
 resched, 84
 rootcb, 84
 runthread, 84
 status, 84
xnshadow_harden
 shadow, 54
xnshadow_map
 shadow, 54
xnshadow_ppd_get
 shadow, 55
xnshadow_relax
 shadow, 56
xnsynch_clear_boost
 synch, 57
xnsynch_flush
 synch, 58
xnsynch_forget_sleeper
 synch, 58
xnsynch_init
 synch, 59
xnsynch_release_all_ownerships
 synch, 60
xnsynch_renice_sleeper
 synch, 60
xnsynch_sleep_on
 synch, 60
xnsynch_wakeup_one_sleeper
 synch, 61
xnsynch_wakeup_this_sleeper
 synch, 61
xntbase_alloc
 timebase, 64
xntbase_free
 timebase, 64
xntbase_start
 timebase, 65
xntbase_stop
 timebase, 65
xntbase_switch
 timebase, 66
xntbase_tick
 timebase, 66
xntbase_update
 timebase, 67
xntimer_destroy
 timer, 68
xntimer_freeze
 timer, 69
xntimer_get_date
 timer.h, 90
xntimer_get_interval
 timer.h, 90
xntimer_get_timeout
 timer.h, 91
xntimer_init
 timer, 69
xntimer_start
 timer.h, 91
xntimer_stop
 timer.h, 92
xntimer_tick_aperiodic
 timer, 69
xntimer_tick_periodic

timer, [70](#)