

Xenomai core
2.99.0

Generated by Doxygen 1.7.1

Thu Jan 5 2012 10:14:32

Contents

1	Module Index	1
1.1	Modules	1
2	Data Structure Index	3
2.1	Data Structures	3
3	File Index	5
3.1	File List	5
4	Module Documentation	7
4.1	Thread state flags.	7
4.1.1	Detailed Description	8
4.1.2	Define Documentation	8
4.1.2.1	XNHELD	8
4.1.2.2	XNLOCK	9
4.1.2.3	XNMIGRATE	9
4.1.2.4	XNPEND	9
4.1.2.5	XNREADY	9
4.1.2.6	XNSUSP	9
4.2	Thread information flags.	9
4.2.1	Detailed Description	10
4.3	Buffer descriptors.	10
4.3.1	Detailed Description	11
4.3.2	Function Documentation	13
4.3.2.1	xnbufd_copy_from_kmem	13
4.3.2.2	xnbufd_copy_to_kmem	14
4.3.2.3	xnbufd_invalidate	15
4.3.2.4	xnbufd_map_kread	15
4.3.2.5	xnbufd_map_kwrite	16

4.3.2.6	xnbufd_map_uread	16
4.3.2.7	xnbufd_map_uwrite	17
4.3.2.8	xnbufd_reset	17
4.3.2.9	xnbufd_unmap_kread	18
4.3.2.10	xnbufd_unmap_kwrite	18
4.3.2.11	xnbufd_unmap_uread	19
4.3.2.12	xnbufd_unmap_uwrite	19
4.4	System clock services.	20
4.4.1	Function Documentation	20
4.4.1.1	xnclock_adjust	20
4.5	Debugging services.	21
4.6	Dynamic memory allocation services.	21
4.6.1	Detailed Description	22
4.6.2	Function Documentation	22
4.6.2.1	xnheap_alloc	22
4.6.2.2	xnheap_destroy	23
4.6.2.3	xnheap_extend	23
4.6.2.4	xnheap_free	24
4.6.2.5	xnheap_init	25
4.6.2.6	xnheap_schedule_free	26
4.6.2.7	xnheap_set_label	26
4.6.2.8	xnheap_test_and_free	27
4.7	Interrupt management.	27
4.7.1	Detailed Description	28
4.7.2	Function Documentation	28
4.7.2.1	xnintr_affinity	28
4.7.2.2	xnintr_attach	28
4.7.2.3	xnintr_destroy	29
4.7.2.4	xnintr_detach	30
4.7.2.5	xnintr_disable	30
4.7.2.6	xnintr_enable	31
4.7.2.7	xnintr_init	31
4.8	Lightweight key-to-object mapping service	33
4.8.1	Detailed Description	33
4.8.2	Function Documentation	33
4.8.2.1	xnmap_create	33

4.8.2.2	xnmap_delete	34
4.8.2.3	xnmap_enter	35
4.8.2.4	xnmap_fetch	35
4.8.2.5	xnmap_fetch_nocheck	36
4.8.2.6	xnmap_remove	36
4.9	Real-time pod services.	37
4.9.1	Detailed Description	39
4.9.2	Function Documentation	39
4.9.2.1	__xnpod_reset_thread	39
4.9.2.2	xnpod_abort_thread	39
4.9.2.3	xnpod_add_hook	40
4.9.2.4	xnpod_delete_thread	41
4.9.2.5	xnpod_disable_timesource	41
4.9.2.6	xnpod_dispatch_signals	42
4.9.2.7	xnpod_enable_timesource	42
4.9.2.8	xnpod_handle_exception	43
4.9.2.9	xnpod_init	43
4.9.2.10	xnpod_init_thread	43
4.9.2.11	xnpod_migrate_thread	45
4.9.2.12	xnpod_remove_hook	45
4.9.2.13	xnpod_resume_thread	46
4.9.2.14	xnpod_schedule	47
4.9.2.15	xnpod_set_thread_mode	48
4.9.2.16	xnpod_set_thread_periodic	48
4.9.2.17	xnpod_set_thread_schedparam	49
4.9.2.18	xnpod_set_thread_tslice	50
4.9.2.19	xnpod_shutdown	51
4.9.2.20	xnpod_start_thread	51
4.9.2.21	xnpod_stop_thread	53
4.9.2.22	xnpod_suspend_thread	53
4.9.2.23	xnpod_unblock_thread	54
4.9.2.24	xnpod_wait_thread_period	55
4.9.2.25	xnpod_welcome_thread	56
4.10	Registry services.	56
4.10.1	Detailed Description	57
4.10.2	Function Documentation	57

4.10.2.1	xnregistry_bind	57
4.10.2.2	xnregistry_enter	58
4.10.2.3	xnregistry_fetch	59
4.10.2.4	xnregistry_get	60
4.10.2.5	xnregistry_put	60
4.10.2.6	xnregistry_remove	61
4.10.2.7	xnregistry_remove_safe	61
4.11	File descriptors events multiplexing services.	62
4.11.1	Detailed Description	63
4.11.2	Function Documentation	63
4.11.2.1	xnselect	63
4.11.2.2	xnselect_bind	64
4.11.2.3	xnselect_destroy	65
4.11.2.4	xnselect_init	65
4.11.2.5	xnselect_signal	65
4.11.2.6	xnselector_destroy	65
4.11.2.7	xnselector_init	66
4.12	Real-time shadow services.	66
4.12.1	Detailed Description	66
4.12.2	Function Documentation	66
4.12.2.1	xnshadow_harden	66
4.12.2.2	xnshadow_map	67
4.12.2.3	xnshadow_ppd_get	68
4.12.2.4	xnshadow_relax	68
4.13	Thread synchronization services.	69
4.13.1	Detailed Description	70
4.13.2	Function Documentation	70
4.13.2.1	xnsynch_acquire	70
4.13.2.2	xnsynch_clear_boost	71
4.13.2.3	xnsynch_flush	71
4.13.2.4	xnsynch_forget_sleeper	72
4.13.2.5	xnsynch_init	73
4.13.2.6	xnsynch_peek_pendq	74
4.13.2.7	xnsynch_release	74
4.13.2.8	xnsynch_release_all_ownerships	75
4.13.2.9	xnsynch_requeue_sleeper	75

4.13.2.10	xnsynch_sleep_on	76
4.13.2.11	xnsynch_wakeup_one_sleeper	76
4.13.2.12	xnsynch_wakeup_this_sleeper	77
4.14	Timer services	78
4.14.1	Detailed Description	79
4.14.2	Function Documentation	79
4.14.2.1	xntimer_destroy	79
4.14.2.2	xntimer_freeze	79
4.14.2.3	xntimer_get_date	80
4.14.2.4	xntimer_get_interval	80
4.14.2.5	xntimer_get_overruns	81
4.14.2.6	xntimer_get_timeout	81
4.14.2.7	xntimer_init	81
4.14.2.8	xntimer_start	82
4.14.2.9	xntimer_stop	83
4.14.2.10	xntimer_tick	83
4.15	Virtual file services	83
4.15.1	Detailed Description	85
4.15.2	Function Documentation	86
4.15.2.1	xnvfile_destroy	86
4.15.2.2	xnvfile_get_blob	86
4.15.2.3	xnvfile_get_integer	86
4.15.2.4	xnvfile_get_string	87
4.15.2.5	xnvfile_init_dir	87
4.15.2.6	xnvfile_init_link	88
4.15.2.7	xnvfile_init_regular	88
4.15.2.8	xnvfile_init_snapshot	89
4.15.3	Variable Documentation	90
4.15.3.1	nkvfroot	90
4.15.3.2	nkvfroot	90
4.16	Sched	90
4.16.1	Function Documentation	91
4.16.1.1	xnsched_rotate	91
5	Data Structure Documentation	93
5.1	xnpod Struct Reference	93
5.1.1	Detailed Description	94

5.1.2	Field Documentation	94
5.1.2.1	refcnt	94
5.1.2.2	sched	94
5.1.2.3	status	94
5.1.2.4	tdeleteq	94
5.1.2.5	threadq	94
5.1.2.6	timerlck	94
5.1.2.7	tstartq	94
5.1.2.8	tswitchq	95
5.2	xnsched Struct Reference	95
5.2.1	Detailed Description	95
5.2.2	Field Documentation	95
5.2.2.1	curr	95
5.2.2.2	htimer	95
5.2.2.3	inesting	95
5.2.2.4	lflags	96
5.2.2.5	rootcb	96
5.2.2.6	rt	96
5.2.2.7	status	96
5.3	xnthread_info Struct Reference	96
5.3.1	Detailed Description	97
5.3.2	Field Documentation	97
5.3.2.1	affinity	97
5.3.2.2	bprio	97
5.3.2.3	cprio	97
5.3.2.4	cpu	97
5.3.2.5	ctxswitches	97
5.3.2.6	exectime	98
5.3.2.7	modeswitches	98
5.3.2.8	name	98
5.3.2.9	pagefaults	98
5.3.2.10	relpoint	98
5.3.2.11	state	98
5.3.2.12	syscalls	98
5.4	xnvfile_lock_ops Struct Reference	98
5.4.1	Detailed Description	99

5.4.2	Field Documentation	99
5.4.2.1	get	99
5.4.2.2	put	99
5.5	xnvfile_regular_iterator Struct Reference	99
5.5.1	Detailed Description	100
5.5.2	Field Documentation	100
5.5.2.1	pos	100
5.5.2.2	private	100
5.5.2.3	seq	100
5.5.2.4	vfile	100
5.6	xnvfile_regular_ops Struct Reference	100
5.6.1	Detailed Description	101
5.6.2	Field Documentation	101
5.6.2.1	begin	101
5.6.2.2	end	101
5.6.2.3	next	102
5.6.2.4	rewind	102
5.6.2.5	show	102
5.6.2.6	store	103
5.7	xnvfile_rev_tag Struct Reference	103
5.7.1	Detailed Description	104
5.7.2	Field Documentation	104
5.7.2.1	rev	104
5.8	xnvfile_snapshot Struct Reference	104
5.8.1	Detailed Description	104
5.9	xnvfile_snapshot_iterator Struct Reference	105
5.9.1	Detailed Description	106
5.9.2	Field Documentation	106
5.9.2.1	databuf	106
5.9.2.2	endfn	106
5.9.2.3	nrdata	106
5.9.2.4	private	106
5.9.2.5	seq	106
5.9.2.6	vfile	106
5.10	xnvfile_snapshot_ops Struct Reference	106
5.10.1	Detailed Description	107

5.10.2	Field Documentation	107
5.10.2.1	begin	107
5.10.2.2	end	107
5.10.2.3	next	108
5.10.2.4	rewind	108
5.10.2.5	show	109
5.10.2.6	store	109
6	File Documentation	111
6.1	include/cobalt/nucleus/bufd.h File Reference	111
6.1.1	Detailed Description	112
6.2	include/cobalt/nucleus/clock.h File Reference	113
6.2.1	Detailed Description	113
6.3	include/cobalt/nucleus/hostrt.h File Reference	114
6.3.1	Detailed Description	115
6.4	include/cobalt/nucleus/map.h File Reference	116
6.4.1	Detailed Description	117
6.5	include/cobalt/nucleus/pod.h File Reference	117
6.5.1	Detailed Description	119
6.6	include/cobalt/nucleus/registry.h File Reference	120
6.6.1	Detailed Description	121
6.7	include/cobalt/nucleus/sched-idle.h File Reference	121
6.7.1	Detailed Description	122
6.8	include/cobalt/nucleus/sched-rt.h File Reference	122
6.8.1	Detailed Description	122
6.9	include/cobalt/nucleus/sched-sporadic.h File Reference	123
6.9.1	Detailed Description	123
6.10	include/cobalt/nucleus/sched-tp.h File Reference	123
6.10.1	Detailed Description	124
6.11	include/cobalt/nucleus/sched.h File Reference	124
6.11.1	Detailed Description	125
6.12	include/cobalt/nucleus/select.h File Reference	125
6.12.1	Detailed Description	127
6.13	include/cobalt/nucleus/timer.h File Reference	128
6.13.1	Detailed Description	129
6.14	include/cobalt/nucleus/vdso.h File Reference	129
6.14.1	Detailed Description	130

6.15	include/cobalt/nucleus/vfile.h File Reference	130
6.15.1	Detailed Description	132
6.16	kernel/cobalt/nucleus/bufd.c File Reference	133
6.16.1	Detailed Description	134
6.17	kernel/cobalt/nucleus/clock.c File Reference	134
6.17.1	Detailed Description	135
6.18	kernel/cobalt/nucleus/debug.c File Reference	135
6.18.1	Detailed Description	135
6.19	kernel/cobalt/nucleus/heap.c File Reference	136
6.19.1	Detailed Description	137
6.20	kernel/cobalt/nucleus/intr.c File Reference	137
6.20.1	Detailed Description	138
6.21	kernel/cobalt/nucleus/map.c File Reference	139
6.21.1	Detailed Description	139
6.22	kernel/cobalt/nucleus/pod.c File Reference	140
6.22.1	Detailed Description	142
6.23	kernel/cobalt/nucleus/registry.c File Reference	142
6.23.1	Detailed Description	143
6.24	kernel/cobalt/nucleus/sched-idle.c File Reference	143
6.24.1	Detailed Description	144
6.25	kernel/cobalt/nucleus/sched-rt.c File Reference	144
6.25.1	Detailed Description	145
6.26	kernel/cobalt/nucleus/sched-sporadic.c File Reference	145
6.26.1	Detailed Description	146
6.27	kernel/cobalt/nucleus/sched-tp.c File Reference	146
6.27.1	Detailed Description	147
6.28	kernel/cobalt/nucleus/sched.c File Reference	148
6.28.1	Detailed Description	148
6.29	kernel/cobalt/nucleus/select.c File Reference	148
6.29.1	Detailed Description	149
6.30	kernel/cobalt/nucleus/shadow.c File Reference	150
6.30.1	Detailed Description	150
6.31	kernel/cobalt/nucleus/synch.c File Reference	151
6.31.1	Detailed Description	152
6.32	kernel/cobalt/nucleus/timer.c File Reference	153
6.32.1	Detailed Description	154

6.33	kernel/cobalt/nucleus/vfile.c File Reference	154
6.33.1	Detailed Description	155

Chapter 1

Module Index

1.1 Modules

Here is a list of all modules:

Thread state flags.	7
Thread information flags.	9
Buffer descriptors.	10
System clock services.	20
Debugging services.	21
Dynamic memory allocation services.	21
Interrupt management.	27
Lightweight key-to-object mapping service	33
Real-time pod services.	37
Registry services.	56
File descriptors events multiplexing services.	62
Real-time shadow services.	66
Thread synchronization services.	69
Timer services.	78
Virtual file services	83
Sched	90

Chapter 2

Data Structure Index

2.1 Data Structures

Here are the data structures with brief descriptions:

xnpod (Real-time pod descriptor)	93
xnsched (Scheduling information structure)	95
xnthread_info (Structure containing thread information)	96
xnvfile_lock_ops (Vfile locking operations)	98
xnvfile_regular_iterator (Regular vfile iterator)	99
xnvfile_regular_ops (Regular vfile operation descriptor)	100
xnvfile_rev_tag (Snapshot revision tag)	103
xnvfile_snapshot (Snapshot vfile descriptor)	104
xnvfile_snapshot_iterator (Snapshot-driven vfile iterator)	105
xnvfile_snapshot_ops (Snapshot vfile operation descriptor)	106

Chapter 3

File Index

3.1 File List

Here is a list of all documented files with brief descriptions:

include/cobalt/nucleus/ assert.h	??
include/cobalt/nucleus/ bheap.h	??
include/cobalt/nucleus/ bufd.h	111
include/cobalt/nucleus/ clock.h	113
include/cobalt/nucleus/ compiler.h	??
include/cobalt/nucleus/ debug.h	??
include/cobalt/nucleus/ heap.h	??
include/cobalt/nucleus/ hostrt.h (Definitions for global semaphore heap shared objects)	114
include/cobalt/nucleus/ intr.h	??
include/cobalt/nucleus/ map.h	116
include/cobalt/nucleus/ pipe.h	??
include/cobalt/nucleus/ pod.h (Real-time pod interface header)	117
include/cobalt/nucleus/ ppd.h	??
include/cobalt/nucleus/ queue.h	??
include/cobalt/nucleus/ registry.h (This file is part of the Xenomai project)	120
include/cobalt/nucleus/ sched-idle.h (Definitions for the IDLE scheduling class)	121
include/cobalt/nucleus/ sched-rt.h (Definitions for the RT scheduling class)	122
include/cobalt/nucleus/ sched-sporadic.h (Definitions for the SSP scheduling class)	123
include/cobalt/nucleus/ sched-tp.h (Definitions for the TP scheduling class)	123
include/cobalt/nucleus/ sched.h (Scheduler interface header)	124
include/cobalt/nucleus/ schedparam.h	??
include/cobalt/nucleus/ schedqueue.h	??
include/cobalt/nucleus/ select.h (File descriptors events multiplexing header)	125
include/cobalt/nucleus/ seqlock.h	??
include/cobalt/nucleus/ shadow.h	??
include/cobalt/nucleus/ stat.h	??
include/cobalt/nucleus/ synch.h	??
include/cobalt/nucleus/ sys_ppd.h	??
include/cobalt/nucleus/ system.h	??
include/cobalt/nucleus/ thread.h	??
include/cobalt/nucleus/ timer.h	128
include/cobalt/nucleus/ trace.h	??
include/cobalt/nucleus/ types.h	??

include/cobalt/nucleus/ vdso.h (Definitions for global semaphore heap shared objects) .	129
include/cobalt/nucleus/ version.h	??
include/cobalt/nucleus/ vfile.h (This file is part of the Xenomai project)	130
kernel/cobalt/nucleus/ bufd.c	133
kernel/cobalt/nucleus/ clock.c	134
kernel/cobalt/nucleus/ debug.c (Debug services)	135
kernel/cobalt/nucleus/ heap.c (Dynamic memory allocation services)	136
kernel/cobalt/nucleus/ intr.c (Interrupt management)	137
kernel/cobalt/nucleus/ map.c	139
kernel/cobalt/nucleus/ pod.c (Real-time pod services)	140
kernel/cobalt/nucleus/ registry.c (This file is part of the Xenomai project)	142
kernel/cobalt/nucleus/ sched-idle.c (Idle scheduling class implementation (i.e. Linux placeholder))	143
kernel/cobalt/nucleus/ sched-rt.c (Common real-time scheduling class implementation (FIFO + RR))	144
kernel/cobalt/nucleus/ sched-sporadic.c (POSIX SCHED_SPORADIC scheduling class) .	145
kernel/cobalt/nucleus/ sched-tp.c (Temporal partitioning (typical of IMA systems)) . . .	146
kernel/cobalt/nucleus/ sched.c	148
kernel/cobalt/nucleus/ select.c (File descriptors events multiplexing)	148
kernel/cobalt/nucleus/ shadow.c (Real-time shadow services)	150
kernel/cobalt/nucleus/ synch.c (Thread synchronization services)	151
kernel/cobalt/nucleus/ timer.c	153
kernel/cobalt/nucleus/ vfile.c (This file is part of the Xenomai project)	154

Chapter 4

Module Documentation

4.1 Thread state flags.

Bits reporting permanent or transient states of thread.

Defines

- #define `XNSUSP` 0x00000001
Suspended.
- #define `XNPEND` 0x00000002
Sleep-wait for a resource.
- #define `XNDELAY` 0x00000004
Delayed.
- #define `XNREADY` 0x00000008
Linked to the ready queue.
- #define `XNDORMANT` 0x00000010
Not started yet or killed.
- #define `XNZOMBIE` 0x00000020
Zombie thread in deletion process.
- #define `XNSTARTED` 0x00000080
Thread has been started.
- #define `XNMAPPED` 0x00000100
Mapped to a regular Linux task (shadow only).
- #define `XNRELAX` 0x00000200
Relaxed shadow thread (blocking bit).
- #define `XNMIGRATE` 0x00000400

Thread is currently migrating to another CPU.

- `#define XNHELD 0x00000800`
Thread is held to process emergency.
- `#define XNBOOST 0x00001000`
Undergoes a PIP boost.
- `#define XNDEBUG 0x00002000`
Hit a debugger breakpoint (shadow only).
- `#define XNLOCK 0x00004000`
Holds the scheduler lock (i.e.
- `#define XNRRB 0x00008000`
Undergoes a round-robin scheduling.
- `#define XNASDI 0x00010000`
ASR are disabled.
- `#define XNDEFCAN 0x00020000`
Deferred cancelability mode (self-set only).
- `#define XNTRAPSW 0x00040000`
Trap execution mode switches.
- `#define XNFPU 0x00080000`
Thread uses FPU.
- `#define XNSHADOW 0x00100000`
Shadow thread.
- `#define XNROOT 0x00200000`
Root thread (that is, Linux/IDLE).
- `#define XNOTHER 0x00400000`
Non real-time shadow (prio=0).

4.1.1 Detailed Description

Bits reporting permanent or transient states of thread.

4.1.2 Define Documentation

4.1.2.1 `#define XNHELD 0x00000800`

Thread is held to process emergency.

Referenced by `xnpod_resume_thread()`, and `xnpod_suspend_thread()`.

4.1.2.2 #define XNLOCK 0x00004000

Holds the scheduler lock (i.e.
not preemptible)

Referenced by `__xnpod_reset_thread()`, `xnpod_set_thread_mode()`, and `xnpod_welcome_thread()`.

4.1.2.3 #define XNMIGRATE 0x00000400

Thread is currently migrating to another CPU.

Referenced by `xnpod_delete_thread()`.

4.1.2.4 #define XNPEND 0x00000002

Sleep-wait for a resource.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, `xnpod_unblock_thread()`, `xnsynch_acquire()`, `xnsynch_flush()`, `xnsynch_forget_sleeper()`, `xnsynch_sleep_on()`, `xnsynch_wakeup_one_sleeper()`, and `xnsynch_wakeup_this_sleeper()`.

4.1.2.5 #define XNREADY 0x00000008

Linked to the ready queue.

Referenced by `xnpod_delete_thread()`, `xnpod_resume_thread()`, and `xnpod_suspend_thread()`.

4.1.2.6 #define XNSUSP 0x00000001

Suspended.

Referenced by `__xnpod_reset_thread()`, `xnpod_handle_exception()`, `xnpod_init_thread()`, `xnpod_start_thread()`, and `xnpod_suspend_thread()`.

4.2 Thread information flags.

Bits reporting events notified to the thread.

Defines

- #define [XNTIMEO](#) 0x00000001
Woken up due to a timeout condition.
- #define [XNRMID](#) 0x00000002
Pending on a removed resource.
- #define [XNBREAK](#) 0x00000004
Forcibly awoken from a wait state.

- #define [XNKICKED](#) 0x00000008
Forced out of primary mode (shadow only).
- #define [XNWAKEN](#) 0x00000010
Thread waken up upon resource availability.
- #define [XNROBBED](#) 0x00000020
Robbed from resource ownership.
- #define [XNAFFSET](#) 0x00000040
CPU affinity changed from primary mode.
- #define [XNPRIOSET](#) 0x00000080
Priority changed from primary mode.
- #define [XNABORT](#) 0x00000100
Thread is being aborted.
- #define [XNCANPND](#) 0x00000200
Cancellation request is pending.
- #define [XNSWREP](#) 0x00000400
Mode switch already reported.

4.2.1 Detailed Description

Bits reporting events notified to the thread.

4.3 Buffer descriptors.

Files

- file [bufd.h](#)
- file [bufd.c](#)

Functions

- static void [xnbufd_map_uread](#) (struct xnbufd *bufd, const void __user *ptr, size_t len)
Initialize a buffer descriptor for reading from user memory.
- static void [xnbufd_map_uwrite](#) (struct xnbufd *bufd, void __user *ptr, size_t len)
Initialize a buffer descriptor for writing to user memory.
- ssize_t [xnbufd_unmap_uread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uread\(\)](#).

- `ssize_t xnbuofd_unmap_uwrite` (struct xnbuofd *bufd)
Finalize a buffer descriptor obtained from `xnbuofd_map_uwrite()`.
- static void `xnbuofd_map_kread` (struct xnbuofd *bufd, const void *ptr, size_t len)
Initialize a buffer descriptor for reading from kernel memory.
- static void `xnbuofd_map_kwrite` (struct xnbuofd *bufd, void *ptr, size_t len)
Initialize a buffer descriptor for writing to kernel memory.
- `ssize_t xnbuofd_unmap_kread` (struct xnbuofd *bufd)
Finalize a buffer descriptor obtained from `xnbuofd_map_kread()`.
- `ssize_t xnbuofd_unmap_kwrite` (struct xnbuofd *bufd)
Finalize a buffer descriptor obtained from `xnbuofd_map_kwrite()`.
- `ssize_t xnbuofd_copy_to_kmem` (void *ptr, struct xnbuofd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- `ssize_t xnbuofd_copy_from_kmem` (struct xnbuofd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- void `xnbuofd_invalidate` (struct xnbuofd *bufd)
Invalidate a buffer descriptor.
- static void `xnbuofd_reset` (struct xnbuofd *bufd)
Reset a buffer descriptor.

4.3.1 Detailed Description

A buffer descriptor is a simple abstraction dealing with copy operations to/from memory buffers which may belong to different address spaces.

To this end, the buffer descriptor library provides a small set of copy routines which are aware of address space restrictions when moving data, and a generic container type which can hold a reference to - or cover - a particular memory area, either present in kernel space, or in any of the existing user memory contexts.

The goal of the buffer descriptor abstraction is to hide address space specifics from Xenomai services dealing with memory areas, allowing them to operate on multiple address spaces seamlessly.

The common usage patterns are as follows:

- Implementing a Xenomai syscall returning a bulk of data to the caller, which may have to be copied back to either kernel or user space:

```
[Syscall implementation]
ssize_t rt_bulk_read_inner(struct xnbuofd *bufd)
{
    ssize_t ret;
    size_t len;
    void *bulk;
```

```

    bulk = get_next_readable_bulk(&len);
    ret = xnbufd_copy_from_kmem(bufd, bulk, min(bufd->b_len, len));
    free_bulk(bulk);

    ret = this_may_fail();
    if (ret)
        xnbufd_invalidate(bufd);

    return ret;
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_read(void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_kwrite(&bufd);

    return ret;
}

[Userland trampoline for user syscalls]
int __rt_bulk_read(struct pt_regs *regs)
{
    struct xnbufd bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbufd_map_uwrite(&bufd, ptr, len);
    ret = rt_bulk_read_inner(&bufd);
    xnbufd_unmap_uwrite(&bufd);

    return ret;
}

```

- Implementing a Xenomai syscall receiving a bulk of data from the caller, which may have to be read from either kernel or user space:

```

[Syscall implementation]
ssize_t rt_bulk_write_inner(struct xnbufd *bufd)
{
    void *bulk = get_free_bulk(bufd->b_len);
    return xnbufd_copy_to_kmem(bulk, bufd, bufd->b_len);
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_write(const void *ptr, size_t len)
{
    struct xnbufd bufd;
    ssize_t ret;

    xnbufd_map_kread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbufd_unmap_kread(&bufd);

    return ret;
}

```



```
[Userland trampoline for user syscalls]
int __rt_bulk_write(struct pt_regs *regs)
{
    struct xnbuid bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbuid_map_uread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbuid_unmap_uread(&bufd);

    return ret;
}
```

4.3.2 Function Documentation

4.3.2.1 `ssize_t xnbuid_copy_from_kmem (struct xnbuid * bufd, void * from, size_t len)`

Copy kernel memory to the area covered by a buffer descriptor.

This routine copies *len* bytes from the kernel memory starting at *from* to the area referred to by the buffer descriptor *bufd*. `xnbuid_copy_from_kmem()` tracks the write offset within the destination memory internally, so that it may be called several times in a loop, until the entire memory area is stored.

The destination address space is dealt with, according to the following rules:

- if *bufd* refers to a writable kernel area (i.e. see `xnbuid_map_kwrite()`), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a writable user area (i.e. see `xnbuid_map_uwrite()`), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while writing to that memory.
- if *bufd* refers to a user area which may not be immediately written to from the current context, the copy is postponed until `xnbuid_unmap_uwrite()` is invoked for *ubufd*, at which point the copy will take place. In such a case, the source memory is transferred to a carry over buffer allocated internally; this operation may lead to request dynamic memory from the nucleus heap if *len* is greater than 64 bytes.

Parameters

bufd The address of the buffer descriptor covering the user memory to copy data to.

from The start address of the kernel memory to copy from.

len The length of the kernel memory to copy to *bufd*.

Returns

The number of bytes written so far to the memory area covered by *ubufd*. Otherwise,

- -ENOMEM is returned when no memory is available from the nucleus heap to allocate the carry over buffer.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while writing to the user area. For that reason, [xnbufd_copy_from_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

Note

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

4.3.2.2 ssize_t xnbufd_copy_to_kmem (void * to, struct xnbufd * bufd, size_t len)

Copy memory covered by a buffer descriptor to kernel memory.

This routine copies *len* bytes from the area referred to by the buffer descriptor *bufd* to the kernel memory area *to*. [xnbufd_copy_to_kmem\(\)](#) tracks the read offset within the source memory internally, so that it may be called several times in a loop, until the entire memory area is loaded.

The source address space is dealt with, according to the following rules:

- if *bufd* refers to readable kernel area (i.e. see [xnbufd_map_kread\(\)](#)), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a readable user area (i.e. see [xnbufd_map_uread\(\)](#)), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while reading from that memory.
- any attempt to read from *bufd* from a non-suitable context is considered as a bug, and will raise a panic assertion when the nucleus is compiled in debug mode.

Parameters

to The start address of the kernel memory to copy to.

bufd The address of the buffer descriptor covering the user memory to copy data from.

len The length of the user memory to copy from *bufd*.

Returns

The number of bytes read so far from the memory area covered by *ubufd*. Otherwise:

- -EINVAL is returned upon attempt to read from the user area from an invalid context. This error is only returned when the debug mode is disabled; otherwise a panic assertion is raised.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: may switch the caller to secondary mode if a page fault occurs while reading from the user area. For that reason, [xnbufd_copy_to_kmem\(\)](#) may only be called from a preemptible section (Linux-wise).

Note

Holding the nklock or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

4.3.2.3 void xnbufd_invalidate (struct xnbufd * *bufd*)

Invalidate a buffer descriptor.

The buffer descriptor is invalidated, making it unusable for further copy operations. If an outstanding carry over buffer was allocated by a previous call to [xnbufd_copy_from_kmem\(\)](#), it is immediately freed so that no data transfer will happen when the descriptor is finalized.

The only action that may subsequently be performed on an invalidated descriptor is calling the relevant unmapping routine for it. For that reason, [xnbufd_invalidate\(\)](#) should be invoked on the error path when data may have been transferred to the carry over buffer.

Parameters

bufd The address of the buffer descriptor to invalidate.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.4 void xnbufd_map_kread (struct xnbufd * *bufd*, const void * *ptr*, size_t *len*) [inline, static]

Initialize a buffer descriptor for reading from kernel memory.

The new buffer descriptor may be used to copy data from kernel memory. This routine should be used in pair with [xnbufd_unmap_kread\(\)](#).

Parameters

bufd The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

ptr The start of the kernel buffer to map.

len The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.5 `void xnbufd_map_kwrite (struct xnbufd * bufd, void * ptr, size_t len) [inline, static]`

Initialize a buffer descriptor for writing to kernel memory.

The new buffer descriptor may be used to copy data to kernel memory. This routine should be used in pair with `xnbufd_unmap_kwrite()`.

Parameters

bufd The address of the buffer descriptor which will map a *len* bytes kernel memory area, starting from *ptr*.

ptr The start of the kernel buffer to map.

len The length of the kernel buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.6 `void xnbufd_map_uread (struct xnbufd * bufd, const void __user * ptr, size_t len) [inline, static]`

Initialize a buffer descriptor for reading from user memory.

The new buffer descriptor may be used to copy data from user memory. This routine should be used in pair with `xnbufd_unmap_uread()`.

Parameters

bufd The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

ptr The start of the user buffer to map.

len The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

4.3.2.7 void `xnbufd_map_uwrite` (struct `xnbufd` * *bufd*, void __user * *ptr*, size_t *len*)
[inline, static]

Initialize a buffer descriptor for writing to user memory.

The new buffer descriptor may be used to copy data to user memory. This routine should be used in pair with [xnbufd_unmap_uwrite\(\)](#).

Parameters

bufd The address of the buffer descriptor which will map a *len* bytes user memory area, starting from *ptr*. *ptr* is never dereferenced directly, since it may refer to a buffer that lives in another address space.

ptr The start of the user buffer to map.

len The length of the user buffer starting at *ptr*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

4.3.2.8 void `xnbufd_reset` (struct `xnbufd` * *bufd*) **[inline, static]**

Reset a buffer descriptor.

The buffer descriptor is reset, so that all data already copied is forgotten. Any carry over buffer allocated is kept, though.

Parameters

bufd The address of the buffer descriptor to reset.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.9 `ssize_t xnbuid_unmap_kread (struct xnbuid * buid)`

Finalize a buffer descriptor obtained from [xnbuid_map_kread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_kread\(\)](#), to read data from a kernel area.

Parameters

buid The address of the buffer descriptor to finalize.

Returns

The number of bytes read so far from the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.10 `ssize_t xnbuid_unmap_kwrite (struct xnbuid * buid)`

Finalize a buffer descriptor obtained from [xnbuid_map_kwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_kwrite\(\)](#), to write data to a kernel area.

Parameters

buid The address of the buffer descriptor to finalize.

Returns

The number of bytes written so far to the memory area covered by *ubuid*.

Environments:

This service can be called from:

- Kernel code (including from primary mode)
- Kernel-based task
- Interrupt service routine

Rescheduling: never.

4.3.2.11 `ssize_t xnbufl_unmap_uread (struct xnbufl * bufl)`

Finalize a buffer descriptor obtained from `xnbufl_map_uread()`.

This routine finalizes a buffer descriptor previously initialized by a call to `xnbufl_map_uread()`, to read data from a user area.

Parameters

bufl The address of the buffer descriptor to finalize.

Returns

The number of bytes read so far from the memory area covered by *bufl*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

4.3.2.12 `ssize_t xnbufl_unmap_uwrite (struct xnbufl * bufl)`

Finalize a buffer descriptor obtained from `xnbufl_map_uwrite()`.

This routine finalizes a buffer descriptor previously initialized by a call to `xnbufl_map_uwrite()`, to write data to a user area.

The main action taken is to write the contents of the kernel memory area passed to `xnbufl_copy_from_kmem()` whenever the copy operation was postponed at that time; the carry over buffer is eventually released as needed. If `xnbufl_copy_from_kmem()` was allowed to copy to the destination user memory at once, then `xnbufl_unmap_uwrite()` leads to a no-op.

Parameters

bufl The address of the buffer descriptor to finalize.

Returns

The number of bytes written so far to the memory area covered by *bufl*.

Environments:

This service can be called from:

- Kernel code (including from primary mode) except Xenomai kernel-based task and interrupt service routines.

Rescheduling: never.

Note

Holding the `nklock` or running real-time interrupts disabled is invalid when calling this routine, and doing so would trigger a debug assertion.

4.4 System clock services.

Files

- file [clock.h](#)
- file [clock.c](#)

Functions

- void [xnclock_adjust](#) (xnsticks_t delta)
Adjust the clock time for the system.

4.4.1 Function Documentation

4.4.1.1 void xnclock_adjust (xnsticks_t delta)

Adjust the clock time for the system.

Xenomai tracks the current time as a monotonously increasing count of ticks since the epoch. The epoch is initially the same as the underlying machine time.

This service changes the epoch for the system by applying the specified tick delta on the wallclock offset.

Parameters

delta The adjustment of the system time expressed in ticks.

Note

This routine must be entered `nklock` locked, interrupts off.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

4.5 Debugging services.

Files

- file [debug.c](#)
Debug services.

4.6 Dynamic memory allocation services.

Files

- file [heap.c](#)
Dynamic memory allocation services.

Functions

- int [xnheap_init](#) (xnheap_t *heap, void *heapaddr, u_long heapsize, u_long pagesize)
Initialize a memory heap.
- void [xnheap_set_label](#) (xnheap_t *heap, const char *label,...)
Set the heap's label string.
- void [xnheap_destroy](#) (xnheap_t *heap, void(*flushfn)(xnheap_t *heap, void *extaddr, u_long extsize, void *cookie), void *cookie)
Destroys a memory heap.
- void * [xnheap_alloc](#) (xnheap_t *heap, u_long size)
Allocate a memory block from a memory heap.
- int [xnheap_test_and_free](#) (xnheap_t *heap, void *block, int(*ckfn)(void *block))
Test and release a memory block to a memory heap.
- int [xnheap_free](#) (xnheap_t *heap, void *block)
Release a memory block to a memory heap.
- int [xnheap_extend](#) (xnheap_t *heap, void *extaddr, u_long extsize)
Extend a memory heap.
- void [xnheap_schedule_free](#) (xnheap_t *heap, void *block, xnholder_t *link)
Schedule a memory block for release.

4.6.1 Detailed Description

Dynamic memory allocation services.

The implementation of the memory allocator follows the algorithm described in a USENIX 1988 paper called "Design of a General Purpose Memory Allocator for the 4.3BSD Unix Kernel" by Marshall K. McKusick and Michael J. Karels. You can find it at various locations on the net, including <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf>. A minor variation allows this implementation to have 'extendable' heaps when needed, with multiple memory extents providing autonomous page address spaces.

The data structures hierarchy is as follows:

```

HEAP {
    block_buckets[]
    extent_queue -----+
}
|
V
    EXTENT #1 {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    } -+
|
|
V
    EXTENT #n {
    {static header}
    page_map[npages]
    page_array[npages][pagesize]
    }

```

4.6.2 Function Documentation

4.6.2.1 void* xnheap_alloc (xnheap_t * heap, u_long size)

Allocate a memory block from a memory heap.

Allocates a contiguous region of memory from an active memory heap. Such allocation is guaranteed to be time-bounded.

Parameters

heap The descriptor address of the heap to get memory from.

size The size in bytes of the requested block. Sizes lower or equal to the page size are rounded either to the minimum allocation size if lower than this value, or to the minimum alignment size if greater or equal to this value. In the current implementation, with MINALLOC = 8 and MINALIGN = 16, a 7 bytes request will be rounded to 8 bytes, and a 17 bytes request will be rounded to 32.

Returns

The address of the allocated region upon success, or NULL if no memory is available from the specified heap.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnshadow_map()`.

4.6.2.2 `void xnheap_destroy (xnheap_t * heap, void(*) (xnheap_t *heap, void *extaddr, u_long extsize, void *cookie) flushfn, void * cookie)`

Destroys a memory heap.

Destroys a memory heap.

Parameters

heap The descriptor address of the destroyed heap.

flushfn If non-NULL, the address of a flush routine which will be called for each extent attached to the heap. This routine can be used by the calling code to further release the heap memory.

cookie If *flushfn* is non-NULL, *cookie* is an opaque pointer which will be passed unmodified to *flushfn*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`, and `xnpod_shutdown()`.

4.6.2.3 `int xnheap_extend (xnheap_t * heap, void * extaddr, u_long extsize)`

Extend a memory heap.

Add a new extent to an existing memory heap.

Parameters

heap The descriptor address of the heap to add an extent to.

extaddr The address of the extent memory.

extsize The size of the extent memory (in bytes). In the current implementation, this size must match the one of the initial extent passed to `xnheap_init()`.

Returns

0 is returned upon success, or -EINVAL is returned if *extsize* differs from the initial extent's size.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.6.2.4 int xnheap_free (xnheap_t * heap, void * block)

Release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from.

Parameters

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

Returns

0 is returned upon success, or one of the following error codes:

- -EFAULT is returned whenever the memory address is outside the heap address space.
- -EINVAL is returned whenever the memory address does not represent a valid block.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References xnheap_test_and_free().

4.6.2.5 `int xnheap_init (xnheap_t * heap, void * heapaddr, u_long heapsize, u_long pagesize)`

Initialize a memory heap.

Initializes a memory heap suitable for time-bounded allocation requests of dynamic memory.

Parameters

heap The address of a heap descriptor which will be used to store the allocation data. This descriptor must always be valid while the heap is active therefore it must be allocated in permanent memory.

heapaddr The address of the heap storage area. All allocations will be made from the given area in time-bounded mode. Since additional extents can be added to a heap, this parameter is also known as the "initial extent".

heapsize The size in bytes of the initial extent pointed at by *heapaddr*. *heapsize* must be a multiple of *pagesize* and lower than 16 Mbytes. *heapsize* must be large enough to contain a dynamically-sized internal header. The following formula gives the size of this header:

$H = \text{heapsize}, P = \text{pagesize}, M = \text{sizeof}(\text{struct pagemap}), E = \text{sizeof}(\text{xnextent_t})$

$\text{hdrsize} = ((H - E) * M) / (M + 1)$

This value is then aligned on the next 16-byte boundary. The routine `xnheap_overhead()` computes the corrected heap size according to the previous formula.

Parameters

pagesize The size in bytes of the fundamental memory page which will be used to subdivide the heap internally. Choosing the right page size is important regarding performance and memory fragmentation issues, so it might be a good idea to take a look at <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf> to pick the best one for your needs. In the current implementation, *pagesize* must be a power of two in the range [8 .. 32768] inclusive.

Returns

0 is returned upon success, or one of the following error codes:

- -EINVAL is returned whenever a parameter is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

4.6.2.6 void xnheap_schedule_free (xnheap_t * heap, void * block, xnholder_t * link)

Schedule a memory block for release.

This routine records a block for later release by `xnheap_finalize_free()`. This service is useful to lazily free blocks of heap memory when immediate release is not an option, e.g. when active references are still pending on the object for a short time after the call. `xnheap_finalize_free()` is expected to be eventually called by the client code at some point in the future when actually freeing the idle objects is deemed safe.

Parameters

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

link The address of a link member, likely but not necessarily within the released object, which will be used by the heap manager to hold the block in the queue of idle objects.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.6.2.7 void xnheap_set_label (xnheap_t * heap, const char * label, ...)

Set the heap's label string.

Set the heap label that will be used in statistic outputs.

Parameters

heap The address of a heap descriptor.

label Label string displayed in statistic outputs. This parameter can be a format string, in which case succeeding parameters will be used to resolve the final label.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnpod_init()`.

4.6.2.8 `int xnheap_test_and_free (xnheap_t * heap, void * block, int(*) (void *block) ckfn)`

Test and release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from. Before the actual release is performed, an optional user-defined can be invoked to check for additional criteria with respect to the request consistency.

Parameters

heap The descriptor address of the heap to release memory to.

block The address of the region to be returned to the heap.

ckfn The address of a user-supplied verification routine which is to be called after the memory address specified by *block* has been checked for validity. The routine is expected to proceed to further consistency checks, and either return zero upon success, or non-zero upon error. In the latter case, the release process is aborted, and *ckfn*'s return value is passed back to the caller of this service as its error return code. *ckfn* must not trigger the rescheduling procedure either directly or indirectly.

Returns

0 is returned upon success, or -EINVAL is returned whenever the block is not a valid region of the specified heap. Additional return codes can also be defined locally by the *ckfn* routine.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnheap_free()`.

4.7 Interrupt management.

Files

- file [intr.c](#)
Interrupt management.

Functions

- int [xnintr_init](#) (xnintr_t *intr, const char *name, unsigned irq, xnintr_isr_t isr, xnintr_iack_t iack, xnintr_flags_t flags)
Initialize an interrupt object.

- int [xnintr_destroy](#) (xnintr_t *intr)
Destroy an interrupt object.
- int [xnintr_attach](#) (xnintr_t *intr, void *cookie)
Attach an interrupt object.
- int [xnintr_detach](#) (xnintr_t *intr)
Detach an interrupt object.
- void [xnintr_enable](#) (xnintr_t *intr)
Enable an interrupt object.
- void [xnintr_disable](#) (xnintr_t *intr)
Disable an interrupt object.
- void [xnintr_affinity](#) (xnintr_t *intr, xnarch_cpumask_t cpumask)
Set interrupt's processor affinity.

4.7.1 Detailed Description

Interrupt management.

4.7.2 Function Documentation

4.7.2.1 void [xnintr_affinity](#) (xnintr_t * *intr*, xnarch_cpumask_t *cpumask*)

Set interrupt's processor affinity.

Causes the IRQ associated with the interrupt object *intr* to be received only on processors which bits are set in *cpumask*.

Parameters

intr The descriptor address of the interrupt object which affinity is to be changed.
cpumask The new processor affinity of the interrupt object.

Note

Depending on architectures, setting more than one bit in *cpumask* could be meaningless.

4.7.2.2 int [xnintr_attach](#) (xnintr_t * *intr*, void * *cookie*)

Attach an interrupt object.

Attach an interrupt object previously initialized by [xnintr_init\(\)](#). After this operation is completed, all IRQs received from the corresponding interrupt channel are directed to the object's ISR.

Parameters

intr The descriptor address of the interrupt object to attach.

cookie A user-defined opaque value which is stored into the interrupt object descriptor for further retrieval by the ISR/ISR handlers.

Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while attaching the interrupt.
- -EBUSY is returned if the interrupt object was already attached.

Note

The caller **must not** hold `nklock` when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Note

Attaching an interrupt resets the tracked number of receipts to zero.

4.7.2.3 `int xnintr_destroy (xnintr_t * intr)`

Destroy an interrupt object.

Destroys an interrupt object previously initialized by `xnintr_init()`. The interrupt object is automatically detached by a call to `xnintr_detach()`. No more IRQs will be dispatched by this object after this service has returned.

Parameters

intr The descriptor address of the interrupt object to destroy.

Returns

0 is returned on success. Otherwise, -EINVAL is returned if an error occurred while detaching the interrupt (see `xnintr_detach()`).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

References `xnintr_detach()`.

4.7.2.4 int xnintr_detach (xnintr_t * intr)

Detach an interrupt object.

Detach an interrupt object previously attached by [xnintr_attach\(\)](#). After this operation is completed, no more IRQs are directed to the object's ISR, but the interrupt object itself remains valid. A detached interrupt object can be attached again by a subsequent call to [xnintr_attach\(\)](#).

Parameters

intr The descriptor address of the interrupt object to detach.

Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while detaching the interrupt, or if the interrupt object was not attached. In both cases, no action is performed.

Note

The caller **must not** hold nklock when invoking this service, this would cause deadlocks.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by [xnintr_destroy\(\)](#).

4.7.2.5 void xnintr_disable (xnintr_t * intr)

Disable an interrupt object.

Disables the hardware interrupt line associated with an interrupt object. This operation invalidates further interrupt requests from the given source until the IRQ line is re-enabled anew.

Parameters

intr The descriptor address of the interrupt object to disable.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

4.7.2.6 void xnintr_enable (xnintr_t * intr)

Enable an interrupt object.

Enables the hardware interrupt line associated with an interrupt object. Over real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

Parameters

intr The descriptor address of the interrupt object to enable.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

4.7.2.7 int xnintr_init (xnintr_t * intr, const char * name, unsigned irq, xnintr_t isr, xniack_t iack, xnflags_t flags)

Initialize an interrupt object.

Associates an interrupt object with an IRQ line.

When an interrupt occurs on the given *irq* line, the ISR is fired in order to deal with the hardware event. The interrupt service code may call any non-suspensive service from the nucleus.

Upon receipt of an IRQ, the ISR is immediately called on behalf of the interrupted stack context, the rescheduling procedure is locked, and the interrupt source is masked at hardware level. The status value returned by the ISR is then checked for the following values:

- XN_ISR_HANDLED indicates that the interrupt request has been fulfilled by the ISR.
- XN_ISR_NONE indicates the opposite to XN_ISR_HANDLED. The ISR must always return this value when it determines that the interrupt request has not been issued by the dedicated hardware device.

In addition, one of the following bits may be set by the ISR :

NOTE: use these bits with care and only when you do understand their effect on the system. The ISR is not encouraged to use these bits in case it shares the IRQ line with other ISRs in the real-time domain.

- XN_ISR_NOENABLE causes the nucleus to ask the real-time control layer *_not_* to re-enable the IRQ line (read the following section). *ipipe_end_irq()* must be called to re-enable the IRQ line later.
- XN_ISR_PROPAGATE tells the nucleus to require the real-time control layer to forward the IRQ. For instance, this would cause the Adeos control layer to propagate the interrupt down the interrupt pipeline to other Adeos domains, such as Linux. This is the regular

way to share interrupts between the nucleus and the host system. In effect, `XN_ISR_PROPAGATE` implies `XN_ISR_NOENABLE` since it would make no sense to re-enable the interrupt channel before the next domain down the pipeline has had a chance to process the propagated interrupt.

The nucleus re-enables the IRQ line by default. Over some real-time control layers which mask and acknowledge IRQs, this operation is necessary to revalidate the interrupt channel so that more interrupts can be notified.

A count of interrupt receipts is tracked into the interrupt descriptor, and reset to zero each time the interrupt object is attached. Since this count could wrap around, it should be used as an indication of interrupt activity only.

Parameters

- intr* The address of a interrupt object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.
- name* An ASCII string standing for the symbolic name of the interrupt object or NULL ("`<unknown>`" will be applied then).
- irq* The hardware interrupt channel associated with the interrupt object. This value is architecture-dependent. An interrupt object must then be attached to the hardware interrupt vector using the `xnintr_attach()` service for the associated IRQs to be directed to this object.
- isr* The address of a valid low-level interrupt service routine if this parameter is non-zero. This handler will be called each time the corresponding IRQ is delivered on behalf of an interrupt context. When called, the ISR is passed the descriptor address of the interrupt object.
- iack* The address of an optional interrupt acknowledge routine, aimed at replacing the default one. Only very specific situations actually require to override the default setting for this parameter, like having to acknowledge non-standard PIC hardware. *iack* should return a non-zero value to indicate that the interrupt has been properly acknowledged. If *iack* is NULL, the default routine will be used instead.
- flags* A set of creation flags affecting the operation. The valid flags are:
 - `XN_ISR_SHARED` enables IRQ-sharing with other interrupt objects.
 - `XN_ISR_EDGE` is an additional flag need to be set together with `XN_ISR_SHARED` to enable IRQ-sharing of edge-triggered interrupts.

Returns

0 is returned on success. Otherwise, `-EINVAL` is returned if *irq* is not a valid interrupt number.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: never.

Referenced by `xnpod_enable_timesource()`.

4.8 Lightweight key-to-object mapping service

Files

- file [map.h](#)
- file [map.c](#)

Functions

- `xnmap_t * xnmap_create (int nkeys, int reserve, int offset)`
Create a map.
- `void xnmap_delete (xnmap_t *map)`
Delete a map.
- `int xnmap_enter (xnmap_t *map, int key, void *objaddr)`
Index an object into a map.
- `int xnmap_remove (xnmap_t *map, int key)`
Remove an object reference from a map.
- `static void * xnmap_fetch_nocheck (xnmap_t *map, int key)`
Search an object into a map - unchecked form.
- `static void * xnmap_fetch (xnmap_t *map, int key)`
Search an object into a map.

4.8.1 Detailed Description

A map is a simple indexing structure which associates unique integer keys with pointers to objects. The current implementation supports reservation, for naming/indexing the real-time objects skins create, either on a fixed, user-provided integer (i.e. a reserved key value), or by drawing the next available key internally if the caller did not specify any fixed key. For instance, in some given map, the key space ranging from 0 to 255 could be reserved for fixed keys, whilst the range from 256 to 511 could be available for drawing free keys dynamically.

A maximum of 1024 unique keys per map is supported on 32bit machines.

(This implementation should not be confused with C++ STL maps, which are dynamically expandable and allow arbitrary key types; Xenomai maps don't).

4.8.2 Function Documentation

4.8.2.1 `xnmap_t * xnmap_create (int nkeys, int reserve, int offset)`

Create a map.

Allocates a new map with the specified addressing capabilities. The memory is obtained from the Xenomai system heap.

Parameters

nkeys The maximum number of unique keys the map will be able to hold. This value cannot exceed the static limit represented by `XNMAP_MAX_KEYS`, and must be a power of two.

reserve The number of keys which should be kept for reservation within the index space. Reserving a key means to specify a valid key to the `xnmap_enter()` service, which will then attempt to register this exact key, instead of drawing the next available key from the unreserved index space. When reservation is in effect, the unreserved index space will hold key values greater than *reserve*, keeping the low key values for the reserved space. For instance, passing *reserve* = 32 would cause the index range [0 .. 31] to be kept for reserved keys. When non-zero, *reserve* is rounded to the next multiple of `BITS_PER_LONG`. If *reserve* is zero no reservation will be available from the map.

offset The lowest key value `xnmap_enter()` will return to the caller. Key values will be in the range [0 + offset .. *nkeys* + offset - 1]. Negative offsets are valid.

Returns

the address of the new map is returned on success; otherwise, NULL is returned if *nkeys* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

4.8.2.2 void xnmap_delete (xnmap_t * map)

Delete a map.

Deletes a map, freeing any associated memory back to the Xenomai system heap.

Parameters

map The address of the map to delete.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

4.8.2.3 int xnmap_enter (xnmap_t* map, int key, void* objaddr)

Index an object into a map.

Insert a new object into the given map.

Parameters

map The address of the map to insert into.

key The key to index the object on. If this key is within the valid index range [0 - offset .. nkeys - offset - 1], then an attempt to reserve this exact key is made. If *key* has an out-of-range value lower or equal to 0 - offset - 1, then an attempt is made to draw a free key from the unreserved index space.

objaddr The address of the object to index on the key. This value will be returned by a successful call to [xnmap_fetch\(\)](#) with the same key.

Returns

a valid key is returned on success, either *key* if reserved, or the next free key. Otherwise:

- -EEXIST is returned upon attempt to reserve a busy key.
- -ENOSPC when no more free key is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.8.2.4 void xnmap_fetch (xnmap_t* map, int key) [inline, static]

Search an object into a map.

Retrieve an object reference from the given map by its index key.

Parameters

map The address of the map to retrieve from.

key The key to be searched for in the map index.

Returns

The indexed object address is returned on success, otherwise NULL is returned when *key* is invalid or no object is currently indexed on it.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.8.2.5 void `xnmap_fetch_nocheck (xnmap_t* map, int key)` [inline, static]

Search an object into a map - unchecked form.

Retrieve an object reference from the given map by its index key, but does not perform any sanity check on the provided key.

Parameters

map The address of the map to retrieve from.

key The key to be searched for in the map index.

Returns

The indexed object address is returned on success, otherwise NULL is returned when no object is currently indexed on *key*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.8.2.6 int `xnmap_remove (xnmap_t* map, int key)`

Remove an object reference from a map.

Removes an object reference from the given map, releasing the associated key.

Parameters

map The address of the map to remove from.

key The key the object reference to be removed is indexed on.

Returns

0 is returned on success. Otherwise:

- -ESRCH is returned if *key* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.9 Real-time pod services.

Data Structures

- struct [xnpod](#)
Real-time pod descriptor.

Files

- file [pod.h](#)
Real-time pod interface header.
- file [pod.c](#)
Real-time pod services.

Functions

- void [__xnpod_reset_thread](#) (struct xntthread *thread)
Reset the thread.
- int [xnpod_init](#) (void)
Initialize the core pod.
- int [xnpod_enable_timesource](#) (void)
Activate the core time source.
- void [xnpod_disable_timesource](#) (void)
Stop the core time source.

- void [xn timer shutdown](#) (int xtype)
Shutdown the current pod.
- int [xn timer init thread](#) (struct xn timer *thread, const struct xn timer_init_attr *attr, struct xn timer_sched_class *sched_class, const union xn timer_sched_policy_param *sched_param)
Initialize a new thread.
- int [xn timer start thread](#) (xn timer_t *thread, const struct xn timer_start_attr *attr)
Initial start of a newly created thread.
- void [xn timer stop thread](#) (xn timer_t *thread)
Stop a thread.
- void [xn timer delete thread](#) (xn timer_t *thread)
Delete a thread.
- void [xn timer abort thread](#) (xn timer_t *thread)
Abort a thread.
- xn timer_flags_t [xn timer set thread mode](#) (xn timer_t *thread, xn timer_flags_t clrmask, xn timer_flags_t set-mask)
Change a thread's control mode.
- void [xn timer suspend thread](#) (xn timer_t *thread, xn timer_flags_t mask, xn timer_ticks_t timeout, xn timer_mode_t timeout_mode, struct xn timer_synch *wchan)
Suspend a thread.
- void [xn timer resume thread](#) (xn timer_t *thread, xn timer_flags_t mask)
Resume a thread.
- int [xn timer unblock thread](#) (xn timer_t *thread)
Unblock a thread.
- int [xn timer set thread sched param](#) (struct xn timer *thread, struct xn timer_sched_class *sched_class, const union xn timer_sched_policy_param *sched_param)
Change the base scheduling parameters of a thread.
- int [xn timer migrate thread](#) (int cpu)
Migrate the current thread.
- void [xn timer dispatch signals](#) (void)
Deliver pending asynchronous signals to the running thread.
- static void [xn timer schedule](#) (void)
Rescheduling procedure entry point.
- int [xn timer handle exception](#) (struct ipipe_trap_data *d)
Exception handler.

- int [xnpod_set_thread_periodic](#) (xnthread_t *thread, xnticks_t idate, xntmode_t timeout_mode, xnticks_t period)
Make a thread periodic.
- int [xnpod_wait_thread_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.
- int [xnpod_set_thread_tslice](#) (struct xnthread *thread, xnticks_t quantum)
Set thread time-slicing information.
- int [xnpod_add_hook](#) (int type, void(*routine)(xnthread_t *))
Install a nucleus hook.
- int [xnpod_remove_hook](#) (int type, void(*routine)(xnthread_t *))
Remove a nucleus hook.
- void [xnpod_welcome_thread](#) (xnthread_t *thread, int imask)
Thread prologue.

4.9.1 Detailed Description

Real-time pod services.

4.9.2 Function Documentation

4.9.2.1 void __xnpod_reset_thread (struct xnthread * thread)

Reset the thread.

For internal use only.

This internal routine resets the state of a thread so that it can be subsequently stopped or restarted.

References `XNLOCK`, `xnpod_resume_thread()`, `xnpod_unblock_thread()`, `XNSUSP`, and `xnsynch_release_all_ownerships()`.

Referenced by `xnpod_stop_thread()`.

4.9.2.2 void xnpod_abort_thread (xnthread_t * thread)

Abort a thread.

Unconditionally terminates a thread and releases all the nucleus resources it currently holds, regardless of whether the target thread is currently active in kernel or user-space. [xnpod_abort_thread\(\)](#) should be reserved for use by skin cleanup routines; [xnpod_delete_thread\(\)](#) should be preferred as the common method for removing threads from a running system.

Parameters

thread The descriptor address of the terminated thread.

This service forces a call to `xnpod_delete_thread()` for the target thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References XNABORT, XNDORMANT, `xnpod_delete_thread()`, and `xnpod_suspend_thread()`.

4.9.2.3 `int xnpod_add_hook (int type, void(*)(xnthread_t *) routine)`

Install a nucleus hook.

The nucleus allows to register user-defined routines which get called whenever a specific scheduling event occurs. Multiple hooks can be chained for a single event type, and get called on a FIFO basis.

The scheduling is locked while a hook is executing.

Parameters

type Defines the kind of hook to install:

- XNHOOK_THREAD_START: The user-defined routine will be called on behalf of the starter thread whenever a new thread starts. The descriptor address of the started thread is passed to the routine.
- XNHOOK_THREAD_DELETE: The user-defined routine will be called on behalf of the deleter thread whenever a thread is deleted. The descriptor address of the deleted thread is passed to the routine.
- XNHOOK_THREAD_SWITCH: The user-defined routine will be called on behalf of the resuming thread whenever a context switch takes place. The descriptor address of the thread which has been switched out is passed to the routine.

Parameters

routine The address of the user-supplied routine to call.

Returns

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- -EINVAL is returned if type is incorrect.
- -ENOMEM is returned if not enough memory is available from the system heap to add the new hook.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

4.9.2.4 void xnpod_delete_thread (xntthread_t * thread)

Delete a thread.

Terminates a thread and releases all the nucleus resources it currently holds. A thread exists in the system since [xnpod_init_thread\(\)](#) has been called to create it, so this service must be called in order to destroy it afterwards.

Parameters

thread The descriptor address of the terminated thread.

The target thread's resources may not be immediately removed if this is an active shadow thread running in user-space. In such a case, the mated Linux task is sent a termination signal instead, and the actual deletion is deferred until the task exit event is called.

The DELETE hooks are called on behalf of the calling context (if any). The information stored in the thread control block remains valid until all hooks have been called.

Self-terminating a thread is allowed. In such a case, this service does not return to the caller.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread self-deletes.

References [xnsched::curr](#), [xnsched::status](#), [XNABORT](#), [XNCANPND](#), [XNDEFKAN](#), [XNDORMANT](#), [XNMIGRATE](#), [XNPEND](#), [xnpod_schedule\(\)](#), [xnpod_unblock_thread\(\)](#), [XNREADY](#), [XNROOT](#), [xnselector_destroy\(\)](#), [xnsynch_forget_sleeper\(\)](#), [xnsynch_release_all_ownerships\(\)](#), [xntimer_destroy\(\)](#), and [XNZOMBIE](#).

Referenced by [xnpod_abort_thread\(\)](#), and [xnpod_shutdown\(\)](#).

4.9.2.5 void xnpod_disable_timesource (void)

Stop the core time source.

Releases the hardware timer, and deactivates the system clock.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task in secondary mode

Rescheduling: never.

References `xntimer_freeze()`.

Referenced by `xnpod_shutdown()`.

4.9.2.6 `void xnpod_dispatch_signals (void)`

Deliver pending asynchronous signals to the running thread.

For internal use only.

This internal routine checks for the presence of asynchronous signals directed to the running thread, and attempts to start the asynchronous service routine (ASR) if any. Called with `nklock` locked, interrupts off.

References `XNASDI`.

Referenced by `xnpod_welcome_thread()`, and `xnshadow_harden()`.

4.9.2.7 `int xnpod_enable_timesource (void)`

Activate the core time source.

On every architecture, Xenomai directly manages a hardware timer clocked in one-shot mode, to support any number of software timers internally. Timings are always specified as a count of nanoseconds.

The `xnpod_enable_timesource()` service configures the hardware timer chip. Because Xenomai most often interposes on the system timer used by the Linux kernel, a software timer may be started to relay periodic ticks to the host kernel if needed.

Returns

0 is returned on success. Otherwise:

- `-ENODEV` is returned if a failure occurred while configuring the hardware timer.
- `-ENOSYS` is returned if no active pod exists.

Environments:

This service can be called from:

- Regular Linux kernel context.

Rescheduling: never.

References `xnsched::htimer`, `xnintr_init()`, and `xntimer_start()`.

Referenced by `xnpod_init()`.

4.9.2.8 int xnpod_handle_exception (struct ipipe_trap_data * d)

Exception handler.

This is the handler which is called whenever an exception/fault is caught over the primary domain.

Parameters

d A pointer to the trap information block received from the pipeline core.

References xnpod_suspend_thread(), xnshadow_relax(), and XNSUSP.

4.9.2.9 int xnpod_init (void)

Initialize the core pod.

Initializes the core interface pod which can subsequently be used to start real-time activities. Once the core pod is active, real-time skins can be stacked over. There can only be a single core pod active in the host environment.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the memory manager fails to initialize.

Environments:

This service can be called from:

- Kernel module initialization code

References xnpod::refcnt, xnsched::rootcb, xnpod::sched, xnpod::status, xnpod::tdeleteq, xnpod::threadq, xnpod::timerlck, xnpod::tstartq, xnpod::tswitchq, xnheap_destroy(), xnheap_init(), xnheap_set_label(), xnpod_enable_timesource(), and xnpod_shutdown().

4.9.2.10 int xnpod_init_thread (struct xnthread * thread, const struct xnthread_init_attr * attr, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param)

Initialize a new thread.

Initializes a new thread attached to the active pod. The thread is left in an innocuous state until it is actually started by [xnpod_start_thread\(\)](#).

Parameters

thread The address of a thread descriptor the nucleus will use to store the thread-specific data. This descriptor must always be valid while the thread is active therefore it must be allocated in permanent memory.

Warning

Some architectures may require the descriptor to be properly aligned in memory; this is an additional reason for descriptors not to be laid in the program stack where alignment constraints might not always be satisfied.

Parameters

attr A pointer to an attribute block describing the initial properties of the new thread. Members of this structure are defined as follows:

- **name**: An ASCII string standing for the symbolic name of the thread. This name is copied to a safe place into the thread descriptor. This name might be used in various situations by the nucleus for issuing human-readable diagnostic messages, so it is usually a good idea to provide a sensible value here. NULL is fine though and means "anonymous".
- **flags**: A set of creation flags affecting the operation. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the created thread:
- **XNSUSP** creates the thread in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod_resume_thread\(\)](#) service for its execution to actually begin, additionally to issuing [xnpod_start_thread\(\)](#) for it. This flag can also be specified when invoking [xnpod_start_thread\(\)](#) as a starting mode.
- **XNFPU** (enable FPU) tells the nucleus that the new thread will use the floating-point unit. In such a case, the nucleus will handle the FPU context save/restore ops upon thread switches at the expense of a few additional cycles per context switch. By default, a thread is not expected to use the FPU. This flag is simply ignored when the nucleus runs on behalf of a userspace-based real-time control layer since the FPU management is always active if present.
- **stacksize**: The size of the stack (in bytes) for the new thread. If zero is passed, the nucleus will use a reasonable pre-defined size depending on the underlying real-time control layer.
- **ops**: A pointer to a structure defining the class-level operations available for this thread. Fields from this structure must have been set appropriately by the caller.

Parameters

sched_class The initial scheduling class the new thread should be assigned to.

sched_param The initial scheduling parameters to set for the new thread; *sched_param* must be valid within the context of *sched_class*.

Returns

0 is returned on success. Otherwise, one of the following error codes indicates the cause of the failure:

- **-EINVAL** is returned if *attr->flags* has invalid bits set.
- **-ENOMEM** is returned if not enough memory is available from the system heap to create the new thread's stack.

Side-effect: This routine does not call the rescheduling procedure.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task

Rescheduling: never.

References XNDORMANT, XNFPU, xnpod_suspend_thread(), XNSHADOW, and XNSUSP.

4.9.2.11 int xnpod_migrate_thread (int *cpu*)

Migrate the current thread.

This call makes the current thread migrate to another CPU if its affinity allows it.

Parameters

cpu The destination CPU.

Return values

- 0 if the thread could migrate ;
- EPERM if the calling context is asynchronous, or the current thread affinity forbids this migration ;
- EBUSY if the scheduler is locked.

References xnpod_schedule().

4.9.2.12 int xnpod_remove_hook (int *type*, void(*)(*xnthread_t* *) *routine*)

Remove a nucleus hook.

This service removes a nucleus hook previously registered using [xnpod_add_hook\(\)](#).

Parameters

- type* Defines the kind of hook to remove among XNHOOK_THREAD_START, XNHOOK_THREAD_DELETE and XNHOOK_THREAD_SWITCH.
- routine* The address of the user-supplied routine to remove.

Returns

0 is returned on success. Otherwise, -EINVAL is returned if type is incorrect or if the routine has never been registered before.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

4.9.2.13 void xnpod_resume_thread (xntthread_t * thread, xnflags_t mask)

Resume a thread.

Resumes the execution of a thread previously suspended by one or more calls to [xnpod_suspend_thread\(\)](#). This call removes a suspensive condition affecting the target thread. When all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters

thread The descriptor address of the resumed thread.

mask The suspension mask specifying the suspensive condition to remove from the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag removes the explicit suspension condition. This condition might be additive to the XNPEND condition.
- XNDELAY. This flag removes the counted delay wait condition.
- XNPEND. This flag removes the resource wait condition. If a watchdog is armed, it is automatically disarmed by this call. Unlike the two previous conditions, only the current thread can set this condition for itself, i.e. no thread can force another one to pend on a resource.

When the thread is eventually resumed by one or more calls to [xnpod_resume_thread\(\)](#), the caller of [xnpod_suspend_thread\(\)](#) in the awakened thread that suspended itself should check for the following bits in its own information mask to determine what caused its wake up:

- XNRMID means that the caller must assume that the pended synchronization object has been destroyed (see [xnsynch_flush\(\)](#)).
- XNTIMEO means that the delay elapsed, or the watchdog went off before the corresponding synchronization object was signaled.
- XNBREAK means that the wait has been forcibly broken by a call to [xnpod_unblock_thread\(\)](#).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNDELAY, XNHELD, XNPEND, XNREADY, [xnsynch_forget_sleeper\(\)](#), and [xntimer_stop\(\)](#).

Referenced by [__xnpod_reset_thread\(\)](#), [xnpod_start_thread\(\)](#), [xnpod_unblock_thread\(\)](#), [xnsynch_flush\(\)](#), [xnsynch_wakeup_one_sleeper\(\)](#), and [xnsynch_wakeup_this_sleeper\(\)](#).

4.9.2.14 void xnpod_schedule (void) [inline, static]

Rescheduling procedure entry point.

This is the central rescheduling routine which should be called to validate and apply changes which have previously been made to the nucleus scheduling state, such as suspending, resuming or changing the priority of threads. This call first determines if a thread switch should take place, and performs it as needed. [xnpod_schedule\(\)](#) schedules out the current thread if:

- the current thread is now blocked or deleted.
- a runnable thread from a higher priority scheduling class is waiting for the CPU.
- the current thread does not lead the runnable threads from its own scheduling class (e.g. round-robin in the RT class).

The nucleus implements a lazy rescheduling scheme so that most of the services affecting the threads state MUST be followed by a call to the rescheduling procedure for the new scheduling state to be applied. In other words, multiple changes on the scheduler state can be done in a row, waking threads up, blocking others, without being immediately translated into the corresponding context switches, like it would be necessary would it appear that a higher priority thread than the current one became runnable for instance. When all changes have been applied, the rescheduling procedure is then called to consider those changes, and possibly replace the current thread by another one.

As a notable exception to the previous principle however, every action which ends up suspending or deleting the current thread begets an immediate call to the rescheduling procedure on behalf of the service causing the state transition. For instance, self-suspension, self-destruction, or sleeping on a synchronization object automatically leads to a call to the rescheduling procedure, therefore the caller does not need to explicitly issue [xnpod_schedule\(\)](#) after such operations.

The rescheduling procedure always leads to a null-effect if it is called on behalf of an ISR or callout. Any outstanding scheduler lock held by the outgoing thread will be restored when the thread is scheduled back in.

Calling this procedure with no applicable context switch pending is harmless and simply leads to a null-effect.

Side-effects:

- If an asynchronous service routine exists, the pending asynchronous signals are delivered to a resuming thread or on behalf of the caller before it returns from the procedure if no context switch has taken place. This behaviour can be disabled by setting the XNASDI flag in the thread's status mask by calling [xnpod_set_thread_mode\(\)](#).

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine, although this leads to a no-op.
- Kernel-based task
- User-space task

Note

The switch hooks are called on behalf of the resuming thread.

References `xnsched::lflags`, and `xnsched::status`.

Referenced by `xnpod_delete_thread()`, `xnpod_migrate_thread()`, `xnpod_shutdown()`, `xnpod_start_thread()`, `xnpod_stop_thread()`, `xnpod_suspend_thread()`, `xnregistry_enter()`, `xnregistry_put()`, `xnselect_bind()`, and `xnselect_destroy()`.

4.9.2.15 `xnflags_t xnpod_set_thread_mode (xnthread_t * thread, xnflags_t clrmask, xnflags_t setmask)`

Change a thread's control mode.

Change the control mode of a given thread. The control mode affects the behaviour of the nucleus regarding the specified thread.

Parameters

thread The descriptor address of the affected thread.

clrmask Clears the corresponding bits from the control field before *setmask* is applied. The scheduler lock held by the current thread can be forcibly released by passing the XNLOCK bit in this mask. In this case, the lock nesting count is also reset to zero.

setmask The new thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the thread:

- XNLOCK causes the thread to lock the scheduler. The target thread will have to call the `xnpod_unlock_sched()` service to unlock the scheduler or clear the XNLOCK bit forcibly using this service. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- XNASDI disables the asynchronous signal handling for this thread. See [xnpod_schedule\(\)](#) for more on this.

Environments:

This service can be called from:

- Kernel-based task
- User-space task in primary mode.

Rescheduling: never, therefore, the caller should reschedule if XNLOCK has been passed into *clrmask*.

References XNLOCK.

4.9.2.16 `int xnpod_set_thread_periodic (xnthread_t * thread, xnticks_t idate, xntmode_t timeout_mode, xnticks_t period)`

Make a thread periodic.

Make a thread periodic by programming its first release point and its period in the processor time line. Subsequent calls to [xnpod_wait_thread_period\(\)](#) will delay the thread until the next periodic release point in the processor timeline is reached.

Parameters

thread The descriptor address of the affected thread. This thread is immediately delayed until the first periodic release point is reached.

idate The initial (absolute) date of the first release point, expressed in nanoseconds. The affected thread will be delayed by the first call to `xnpod_wait_thread_period()` until this point is reached. If *idate* is equal to `XN_INFINITE`, the current system date is used, and no initial delay takes place. In the latter case, *timeout_mode* is not considered and can have any valid value.

timeout_mode The mode of the *idate* parameter. It can either be set to `XN_ABSOLUTE` or `XN_REALTIME` with *idate* different from `XN_INFINITE` (see also `xntimer_start()`).

period The period of the thread, expressed in nanoseconds. As a side-effect, passing `XN_INFINITE` attempts to stop the thread's periodic timer; in the latter case, the routine always exits successfully, regardless of the previous state of this timer.

Returns

0 is returned upon success. Otherwise:

- `-ETIMEDOUT` is returned if *idate* is different from `XN_INFINITE` and represents a date in the past.
- `-EINVAL` is returned if *period* is different from `XN_INFINITE` but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`. `-EINVAL` is also returned if *timeout_mode* is not compatible with *idate*, such as `XN_RELATIVE` with *idate* different from `XN_INFINITE`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: none.

References `xntimer_start()`, and `xntimer_stop()`.

4.9.2.17 `int xnpod_set_thread_schedparam (struct xnthread * thread, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param)`

Change the base scheduling parameters of a thread.

Changes the base scheduling policy and parameters of a thread. If the thread is currently blocked, waiting in priority-pending mode (`XNSYNCH_PRIO`) for a synchronization object to be signaled, the nucleus will attempt to reorder the object's wait queue so that it reflects the new sleeper's priority, unless the `XNSYNCH_DREORD` flag has been set for the pending object.

Parameters

thread The descriptor address of the affected thread.

sched_class The new scheduling class the thread should be assigned to.

sched_param The scheduling parameters to set for the thread; *sched_param* must be valid within the context of *sched_class*.

It is absolutely required to use this service to change a thread priority, in order to have all the needed housekeeping chores correctly performed. i.e. Do *not* call `xnsched_set_policy()` directly or worse, change the `thread.cprio` field by hand in any case.

Returns

0 is returned on success. Otherwise, a negative error code indicates the cause of a failure that happened in the scheduling class implementation for *sched_class*. Invalid parameters passed into *sched_param* are common causes of error.

Side-effects:

- This service does not call the rescheduling procedure but may affect the state of the runnable queue for the previous and new scheduling classes.
- Assigning the same scheduling class and parameters to a running or ready thread moves it to the end of the runnable queue, thus causing a manual round-robin.
- If the thread is a user-space shadow, this call propagates the request to the mated Linux task.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.9.2.18 `int xnpod_set_thread_tslice (struct xnthread * thread, xnticks_t quantum)`

Set thread time-slicing information.

Update the time-slicing information for a given thread. This service enables or disables round-robin scheduling for the thread, depending on the value of *quantum*. By default, times-slicing is disabled for a new thread initialized by a call to [xnpod_init_thread\(\)](#).

Parameters

thread The descriptor address of the affected thread.

quantum The time quantum assigned to the thread expressed in nanoseconds. If *quantum* is different from `XN_INFINITE`, the time-slice for the thread is set to that value and its current time credit is refilled (i.e. the thread is given a full time-slice to run next). Otherwise, if *quantum* equals `XN_INFINITE`, time-slicing is stopped for that thread.

Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *quantum* is not XN_INFINITE, and the base scheduling class of the target thread does not support time-slicing.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

References XNRRB, xntimer_start(), and xntimer_stop().

4.9.2.19 void xnpod_shutdown (int xtype)

Shutdown the current pod.

Forcibly shutdowns the active pod. All existing nucleus threads (but the root one) are terminated, and the system heap is freed.

Parameters

xtype An exit code passed to the host environment who started the nucleus. Zero is always interpreted as a successful return.

The nucleus never calls this routine directly. Skins should provide their own shutdown handlers which end up calling [xnpod_shutdown\(\)](#) after their own housekeeping chores have been carried out.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References xnheap_destroy(), xnpod_delete_thread(), xnpod_disable_timesource(), xnpod_schedule(), and XNROOT.

Referenced by xnpod_init().

4.9.2.20 int xnpod_start_thread (xnthread_t * thread, const struct xnthread_start_attr * attr)

Initial start of a newly created thread.

Starts a (newly) created thread, scheduling it for the first time. This call releases the target thread from the XNDORMANT state. This service also sets the initial mode and interrupt mask for the new thread.

Parameters

thread The descriptor address of the affected thread which must have been previously initialized by the [xnpod_init_thread\(\)](#) service.

attr A pointer to an attribute block describing the execution properties of the new thread. Members of this structure are defined as follows:

- *mode*: The initial thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the started thread:
- **XNLOCK** causes the thread to lock the scheduler when it starts. The target thread will have to call the [xnpod_unlock_sched\(\)](#) service to unlock the scheduler. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
- **XNASDI** disables the asynchronous signal handling for this thread. See [xnpod_schedule\(\)](#) for more on this.
- **XNSUSP** makes the thread start in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnpod_resume_thread\(\)](#) service for its execution to actually begin.
- *imask*: The interrupt mask that should be asserted when the thread starts. The processor interrupt state will be set to the given value when the thread starts running. The interpretation of this value might be different across real-time layers, but a non-zero value should always mark an interrupt masking in effect (e.g. [local_irq_disable\(\)](#)). Conversely, a zero value should always mark a fully preemptible state regarding interrupts (e.g. [local_irq_enable\(\)](#)).
- *affinity*: The processor affinity of this thread. Passing **XNPOD_ALL_CPUS** or an empty affinity set means "any cpu".
- *entry*: The address of the thread's body routine. In other words, it is the thread entry point.
- *cookie*: A user-defined opaque cookie the nucleus will pass to the emerging thread as the sole argument of its entry point.

The START hooks are called on behalf of the calling context (if any).

Return values

- **0** if *thread* could be started ;
- **-EBUSY** if *thread* was not dormant or stopped ;
- **-EINVAL** if the value of *attr->affinity* is invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

- User-space task

Rescheduling: possible.

References XNDORMANT, `xnpod_resume_thread()`, `xnpod_schedule()`, XNSHADOW, XN-STARTED, and XNSUSP.

Referenced by `xnshadow_map()`.

4.9.2.21 `void xnpod_stop_thread (xntthread_t * thread)`

Stop a thread.

Stop a previously started thread. The thread is put back into the dormant state; however, it is not deleted from the system.

Parameters

thread The descriptor address of the affected thread which must have been previously started by the `xnpod_start_thread()` service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References `__xnpod_reset_thread()`, XNDORMANT, `xnpod_schedule()`, `xnpod_suspend_thread()`, and XNROOT.

4.9.2.22 `void xnpod_suspend_thread (xntthread_t * thread, xnflags_t mask, xnticks_t timeout, xntmode_t timeout_mode, struct xnsynch * wchan)`

Suspend a thread.

Suspends the execution of a thread according to a given suspensive condition. This thread will not be eligible for scheduling until all the pending suspensive conditions set by this service are removed by one or more calls to `xnpod_resume_thread()`.

Parameters

thread The descriptor address of the suspended thread.

mask The suspension mask specifying the suspensive condition to add to the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag forcibly suspends a thread, regardless of any resource to wait for. A reverse call to `xnpod_resume_thread()` specifying the XNSUSP bit must be issued to remove this condition, which is cumulative with other suspension bits. *wchan* should be NULL when using this suspending mode.

- **XNDELAY.** This flag denotes a counted delay wait (in ticks) which duration is defined by the value of the timeout parameter.
- **XNPEND.** This flag denotes a wait for a synchronization object to be signaled. The *wchan* argument must point to this object. A timeout value can be passed to bound the wait. This suspending mode should not be used directly by the client interface, but rather through the [xnsynch_sleep_on\(\)](#) call.

Parameters

timeout The timeout which may be used to limit the time the thread pends on a resource. This value is a wait time given in nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout_mode*. Passing **XN_INFINITE** and setting *timeout_mode* to **XN_RELATIVE** specifies an unbounded wait. All other values are used to initialize a watchdog timer. If the current operation mode of the system timer is oneshot and *timeout* elapses before [xnpod_suspend_thread\(\)](#) has completed, then the target thread will not be suspended, and this routine leads to a null effect.

timeout_mode The mode of the *timeout* parameter. It can either be set to **XN_RELATIVE**, **XN_ABSOLUTE**, or **XN_REALTIME** (see also [xntimer_start\(\)](#)).

wchan The address of a pended resource. This parameter is used internally by the synchronization object implementation code to specify on which object the suspended thread pends. **NULL** is a legitimate value when this parameter does not apply to the current suspending mode (e.g. **XNSUSP**).

Note

If the target thread is a shadow which has received a Linux-originated signal, then this service immediately exits without suspending the thread, but raises the **XNBREAK** condition in its information mask.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: possible if the current thread suspends itself.

References [xnsched::curr](#), **XNBREAK**, **XNDELAY**, **XNDORMANT**, **XNHELD**, **XNKICKED**, [xnpod_schedule\(\)](#), **XNREADY**, **XNRELAX**, **XNRMID**, **XNROBBED**, **XNROOT**, **XNSHADOW**, **XNSUSP**, [xnsynch_forget_sleeper\(\)](#), **XNTIMEO**, [xntimer_start\(\)](#), and **XNWAKEN**.

Referenced by [xnpod_abort_thread\(\)](#), [xnpod_handle_exception\(\)](#), [xnpod_init_thread\(\)](#), [xnpod_stop_thread\(\)](#), [xnpod_wait_thread_period\(\)](#), [xnshadow_map\(\)](#), [xnshadow_relax\(\)](#), [xnsynch_acquire\(\)](#), and [xnsynch_sleep_on\(\)](#).

4.9.2.23 int xnpod_unblock_thread (xnthread_t * thread)

Unblock a thread.

Breaks the thread out of any wait it is currently in. This call removes the XNDELAY and XNPEND suspensive conditions previously put by [xnpod_suspend_thread\(\)](#) on the target thread. If all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters

thread The descriptor address of the unblocked thread.

This call neither releases the thread from the XNSUSP, XNRELAX, XNDORMANT or XNHELD suspensive conditions.

When the thread resumes execution, the XNBREAK bit is set in the unblocked thread's information mask. Unblocking a non-blocked thread is perfectly harmless.

Returns

non-zero is returned if the thread was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNBREAK, XNDELAY, XNPEND, and [xnpod_resume_thread\(\)](#).

Referenced by [__xnpod_reset_thread\(\)](#), and [xnpod_delete_thread\(\)](#).

4.9.2.24 `int xnpod_wait_thread_period (unsigned long * overruns_r)`

Wait for the next periodic release point.

Make the current thread wait for the next periodic release point in the processor time line.

Parameters

overruns_r If non-NULL, *overruns_r* must be a pointer to a memory location which will be written with the count of pending overruns. This value is copied only when [xnpod_wait_thread_period\(\)](#) returns -ETIMEDOUT or success; the memory location remains unmodified otherwise. If NULL, this count will never be copied back.

Returns

0 is returned upon success; if *overruns_r* is valid, zero is copied to the pointed memory location. Otherwise:

- -EWOULDBLOCK is returned if [xnpod_set_thread_periodic\(\)](#) has not previously been called for the calling thread.

- -EINTR is returned if [xnpod_unblock_thread\(\)](#) has been called for the waiting thread before the next periodic release point has been reached. In this case, the overrun counter is reset too.
- -ETIMEDOUT is returned if the timer has overrun, which indicates that one or more previous release points have been missed by the calling thread. If *overruns_r* is valid, the count of pending overruns is copied to the pointed memory location.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always, unless the current release point has already been reached. In the latter case, the current thread immediately returns from this service without being delayed.

References XNBREAK, XNDELAY, [xnpod_suspend_thread\(\)](#), and [xntimer_get_overruns\(\)](#).

4.9.2.25 void [xnpod_welcome_thread](#) ([xnthread_t](#) * *thread*, int *imask*)

Thread prologue.

For internal use only.

This internal routine is called on behalf of a (re)starting thread's prologue before the user entry point is invoked. This call is reserved for internal housekeeping chores and cannot be inlined.

Entered with *nklock* locked, *irqs* off.

References XNLOCK, and [xnpod_dispatch_signals\(\)](#).

4.10 Registry services.

Files

- file [registry.h](#)
This file is part of the Xenomai project.
- file [registry.c](#)
This file is part of the Xenomai project.

Functions

- int [xnregistry_enter](#) (const char **key*, void **objaddr*, [xnhandle_t](#) **phandle*, struct [xnnode](#) **pnode*)

Register a real-time object.

- `int xnregistry_bind (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)`

Bind to a real-time object.

- `int xnregistry_remove (xnhandle_t handle)`

Forcibly unregister a real-time object.

- `int xnregistry_remove_safe (xnhandle_t handle, xnticks_t timeout)`

Unregister an idle real-time object.

- `void * xnregistry_get (xnhandle_t handle)`

Find and lock a real-time object into the registry.

- `u_long xnregistry_put (xnhandle_t handle)`

Unlock a real-time object from the registry.

- `void * xnregistry_fetch (xnhandle_t handle)`

Find a real-time object into the registry.

4.10.1 Detailed Description

The registry provides a mean to index real-time object descriptors created by Xenomai skins on unique alphanumeric keys. When labeled this way, a real-time object is globally exported; it can be searched for, and its descriptor returned to the caller for further use; the latter operation is called a "binding". When no object has been registered under the given name yet, the registry can be asked to set up a rendez-vous, blocking the caller until the object is eventually registered.

4.10.2 Function Documentation

4.10.2.1 `int xnregistry_bind (const char * key, xnticks_t timeout, int timeout_mode, xnhandle_t * phandle)`

Bind to a real-time object.

This service retrieves the registry handle of a given object identified by its key. Unless otherwise specified, this service will block the caller if the object is not registered yet, waiting for such registration to occur.

Parameters

key A valid NULL-terminated string which identifies the object to bind to.

timeout The timeout which may be used to limit the time the thread wait for the object to be registered. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic (XN_ABSOLUTE), or absolute adjustable (XN_REALTIME) depending on *timeout_mode*. Passing XN_INFINITE and setting *timeout_mode* to XN_RELATIVE specifies an unbounded wait. Passing XN_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry. All other values are used as a wait limit.

timeout_mode The mode of the *timeout* parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also [xntimer_start\(\)](#)).

phandle A pointer to a memory location which will be written upon success with the generic handle defined by the registry for the retrieved object. Contents of this memory is undefined upon failure.

Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* is NULL.
- -EINTR is returned if [xnpod_unblock_thread\(\)](#) has been called for the waiting thread before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to XN_NONBLOCK and the searched object is not registered on entry. As a special exception, this error is also returned if this service should block, but was called from a context which cannot sleep (e.g. interrupt, non-realtime or scheduler locked).
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN_NONBLOCK.
- Kernel-based thread.

Rescheduling: always unless the request is immediately satisfied or *timeout* specifies a non-blocking operation.

References XNBREAK, xnsynch_sleep_on(), and XNTIMEO.

4.10.2.2 `int xnregistry_enter (const char * key, void * objaddr, xnhandle_t * phandle, struct xnpnode * pnode)`

Register a real-time object.

This service allocates a new registry slot for an associated object, and indexes it by an alphanumeric key for later retrieval.

Parameters

key A valid NULL-terminated string by which the object will be indexed and later retrieved in the registry. Since it is assumed that such key is stored into the registered object, it will **not** be copied but only kept by reference in the registry. Pass an empty string if the object shall only occupy a registry slot for handle-based lookups.

objaddr An opaque pointer to the object to index by *key*.

phandle A pointer to a generic handle defined by the registry which will uniquely identify the indexed object, until the latter is unregistered using the [xnregistry_remove\(\)](#) service.

pnode A pointer to an optional /proc node class descriptor. This structure provides the information needed to export all objects from the given class through the /proc filesystem, under the /proc/xenomai/registry entry. Passing NULL indicates that no /proc support is available for the newly registered object.

Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *objaddr* are NULL, or if *key* contains an invalid '/' character.
- -ENOMEM is returned if the system fails to get enough dynamic memory from the global real-time heap in order to register the object.
- -EEXIST is returned if the *key* is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: possible.

References [xnpod_schedule\(\)](#), and [xnsynch_init\(\)](#).

4.10.2.3 void* xnregistry_fetch (xnhandle_t handle)

Find a real-time object into the registry.

This service retrieves an object from its handle into the registry and returns the memory address of its descriptor.

Parameters

handle The generic handle of the object to fetch. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread

Rescheduling: never.

4.10.2.4 void* xnregistry_get (xnhandle_t handle)

Find and lock a real-time object into the registry.

This service retrieves an object from its handle into the registry and prevents its removal atomically. A locking count is tracked, so that [xnregistry_get\(\)](#) and [xnregistry_put\(\)](#) must be used in pair.

Parameters

handle The generic handle of the object to find and lock. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread.

Rescheduling: never.

4.10.2.5 u_long xnregistry_put (xnhandle_t handle)

Unlock a real-time object from the registry.

This service decrements the lock count of a registered object previously locked by a call to [xnregistry_get\(\)](#). The object is actually unlocked from the registry when the locking count falls down to zero, thus waking up any thread currently blocked on [xnregistry_remove\(\)](#) for unregistering it.

Parameters

handle The generic handle of the object to unlock. If XNOBJECT_SELF is passed, the object is the calling Xenomai thread.

Returns

The decremented lock count is returned upon success. Zero is also returned if *handle* does not reference a registered object, or if *handle* is equal to XNOBJECT_SELF but the current context is not a real-time thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine only if *handle* is different from XNOBJECT_SELF.
- Kernel-based thread

Rescheduling: possible if the lock count falls down to zero and some thread is currently waiting for the object to be unlocked.

References `xnpod_schedule()`, and `xnsynch_flush()`.

4.10.2.6 `int xnregistry_remove (xnhandle_t handle)`

Forcibly unregister a real-time object.

This service forcibly removes an object from the registry. The removal is performed regardless of the current object's locking status.

Parameters

handle The generic handle of the object to remove.

Returns

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based thread

Rescheduling: never.

Referenced by `xnregistry_remove_safe()`.

4.10.2.7 `int xnregistry_remove_safe (xnhandle_t handle, xnticks_t timeout)`

Unregister an idle real-time object.

This service removes an object from the registry. The caller might sleep as a result of waiting for the target object to be unlocked prior to the removal (see [xnregistry_put\(\)](#)).

Parameters

handle The generic handle of the object to remove.

timeout If the object is locked on entry, *param* gives the number of nanoseconds to wait for the unlocking to occur. Passing XN_INFINITE causes the caller to block indefinitely until the object is unlocked. Passing XN_NONBLOCK causes the service to return immediately without waiting if the object is locked on entry.

Returns

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.
- -EWOULDBLOCK is returned if *timeout* is equal to XN_NONBLOCK and the object is locked on entry.
- -EBUSY is returned if *handle* refers to a locked object and the caller could not sleep until it is unlocked.
- -ETIMEDOUT is returned if the object cannot be removed within the specified amount of time.
- -EINTR is returned if [xnpod_unblock_thread\(\)](#) has been called for the calling thread waiting for the object to be unlocked.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine only if *timeout* is equal to XN_NONBLOCK.
- Kernel-based thread.

Rescheduling: possible if the object to remove is currently locked and the calling context can sleep.

References XNBREAK, [xnregistry_remove\(\)](#), [xnsynch_sleep_on\(\)](#), and XNTIMEO.

4.11 File descriptors events multiplexing services.

Files

- file [select.h](#)
file descriptors events multiplexing header.
- file [select.c](#)
file descriptors events multiplexing.

Functions

- void [xnselect_init](#) (struct xnselect *select_block)
Initialize a struct xnselect structure.
- int [xnselect_bind](#) (struct xnselect *select_block, struct xnselect_binding *binding, struct xnselector *selector, unsigned type, unsigned index, unsigned state)
Bind a file descriptor (represented by its xnselect structure) to a selector block.
- static int [xnselect_signal](#) (struct xnselect *select_block, unsigned state)

Signal a file descriptor state change.

- void `xnselect_destroy` (struct xnselect *select_block)
Destroy the xnselect structure associated with a file descriptor.
- int `xnselector_init` (struct xnselector *selector)
Initialize a selector structure.
- int `xnselect` (struct xnselector *selector, fd_set *out_fds[XNSELECT_MAX_TYPES], fd_set *in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)
Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.
- void `xnselector_destroy` (struct xnselector *selector)
Destroy a selector block.

4.11.1 Detailed Description

File descriptors events multiplexing services.

This module implements the services needed for implementing the posix "select" service, or any other events multiplexing services.

Following the implementation of the posix select service, this module defines three types of events:

- `XNSELECT_READ` meaning that a file descriptor is ready for reading;
- `XNSELECT_WRITE` meaning that a file descriptor is ready for writing;
- `XNSELECT_EXCEPT` meaning that a file descriptor received an exceptional event.

It works by defining two structures:

- a *struct xnselect* structure, which should be added to every file descriptor for every event type (read, write, or except);
- a *struct xnselector* structure, the selection structure, passed by the thread calling the xnselect service, where this service does all its housekeeping.

4.11.2 Function Documentation

4.11.2.1 `int xnselect (struct xnselector * selector, fd_set * out_fds[XNSELECT_MAX_TYPES], fd_set * in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)`

Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.

Parameters

selector structure to check for pending events

out_fds The set of descriptors with pending events if a strictly positive number is returned, or the set of descriptors not yet bound if -ECHRNG is returned;

in_fds the set of descriptors which events should be checked

nfds the highest-numbered descriptor in any of the *in_fds* sets, plus 1;

timeout the timeout, whose meaning depends on *timeout_mode*, note that `xnselect()` pass *timeout* and *timeout_mode* unchanged to `xnsynch_sleep_on`, so passing a relative value different from `XN_INFINITE` as a timeout with *timeout_mode* set to `XN_RELATIVE`, will cause a longer sleep than expected if the sleep is interrupted.

timeout_mode the mode of *timeout*.

Return values

`-EINVAL` if *nfds* is negative;

`-ECHRNG` if some of the descriptors passed in *in_fds* have not yet been registered with `xnselect_bind()`, *out_fds* contains the set of such descriptors;

`-EINTR` if `xnselect` was interrupted while waiting;

0 in case of timeout.

the number of file descriptors having received an event.

References `XNBREAK`, `xnsynch_sleep_on()`, and `XNTIMEO`.

4.11.2.2 `int xnselect_bind (struct xnselect * select_block, struct xnselect_binding * binding, struct xnselector * selector, unsigned type, unsigned index, unsigned state)`

Bind a file descriptor (represented by its `xnselect` structure) to a selector block.

Parameters

select_block pointer to the `struct xnselect` to be bound;

binding pointer to a newly allocated (using `xnmalloc`) `struct xnselect_binding`;

selector pointer to the selector structure;

type type of events (`XNSELECT_READ`, `XNSELECT_WRITE`, or `XNSELECT_EXCEPT`);

index index of the file descriptor (represented by *select_block*) in the bit fields used by the *selector* structure;

state current state of the file descriptor>.

select_block must have been initialized with `xnselect_init()`, the `xnselector` structure must have been initialized with `xnselector_init()`, *binding* may be uninitialized.

This service must be called with `nklock` locked, `irqs` off. For this reason, the *binding* parameter must have been allocated by the caller outside the locking section.

Return values

`-EINVAL` if *type* or *index* is invalid;

0 otherwise.

References `xnpod_schedule()`.

4.11.2.3 void xnselect_destroy (struct xnselect * *select_block*)

Destroy the *xnselect* structure associated with a file descriptor.

Any binding with a *xnselector* block is destroyed.

Parameters

select_block pointer to the *xnselect* structure associated with a file descriptor

References `xnpod_schedule()`.

4.11.2.4 void xnselect_init (struct xnselect * *select_block*)

Initialize a *struct xnselect* structure.

This service must be called to initialize a *struct xnselect* structure before it is bound to a selector by the means of `xnselect_bind()`.

Parameters

select_block pointer to the *xnselect* structure to be initialized

**4.11.2.5 static int xnselect_signal (struct xnselect * *select_block*, unsigned *state*)
[inline, static]**

Signal a file descriptor state change.

Parameters

select_block pointer to an *xnselect* structure representing the file descriptor whose state changed;

state new value of the state.

Return values

1 if rescheduling is needed;

0 otherwise.

4.11.2.6 void xnselector_destroy (struct xnselector * *selector*)

Destroy a selector block.

All bindings with file descriptor are destroyed.

Parameters

selector the selector block to be destroyed

Referenced by `xnpod_delete_thread()`.

4.11.2.7 int xnselector_init (struct xnselector * selector)

Initialize a selector structure.

Parameters

selector The selector structure to be initialized.

Return values

0

References xnsynch_init().

4.12 Real-time shadow services.

Files

- file [shadow.c](#)
Real-time shadow services.

Functions

- int [xnshadow_harden](#) (void)
Migrate a Linux task to the Xenomai domain.
- void [xnshadow_relax](#) (int notify, int reason)
Switch a shadow thread back to the Linux domain.
- int [xnshadow_map](#) (xnthread_t *thread, xncompletion_t __user *u_completion, unsigned long __user *u_window_offset)
Create a shadow thread context.
- xnshadow_ppd_t * [xnshadow_ppd_get](#) (unsigned int muxid)
Return the per-process data attached to the calling process.

4.12.1 Detailed Description

Real-time shadow services.

4.12.2 Function Documentation

4.12.2.1 int xnshadow_harden (void)

Migrate a Linux task to the Xenomai domain.

For internal use only.

This service causes the transition of "current" from the Linux domain to Xenomai. The shadow will resume in the Xenomai domain as returning from `schedule()`.

Environments:

This service can be called from:

- User-space thread operating in secondary (i.e. relaxed) mode.

Rescheduling: always.

References `XNDEBUG`, `xnpod_dispatch_signals()`, `XNRELAX`, and `xnshadow_relax()`.

Referenced by `xnshadow_map()`.

4.12.2.2 `int xnshadow_map (xntthread_t * thread, xncompletion_t __user * u_completion, unsigned long __user * u_window_offset)`

Create a shadow thread context.

For internal use only.

This call maps a nucleus thread to the "current" Linux task. The priority and scheduling class of the underlying Linux task are not affected; it is assumed that the interface library did set them appropriately before issuing the shadow mapping request.

Parameters

thread The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to `xnpod_init_thread()`.

u_completion is the address of an optional completion descriptor aimed at synchronizing our parent thread with us. If non-NULL, the information `xnshadow_map()` will store into the completion block will be later used to wake up the parent thread when the current shadow has been initialized. In the latter case, the new shadow thread is left in a dormant state (XNDORMANT) after its creation, leading to the suspension of "current" in the Linux domain, only processing signals. Otherwise, the shadow thread is immediately started and "current" immediately resumes in the Xenomai domain from this service.

u_window_offset will receive the offset of the per-thread "u_window" structure in the process shared heap associated to *thread*. This structure reflects thread state information visible from userland through a shared memory window.

Returns

0 is returned on success. Otherwise:

- `-ERESTARTSYS` is returned if the current Linux task has received a signal, thus preventing the final migration to the Xenomai domain (i.e. in order to process the signal in the Linux domain). This error should not be considered as fatal.
- `-EPERM` is returned if the shadow thread has been killed before the current task had a chance to return to the caller. In such a case, the real-time mapping operation has failed globally, and no Xenomai resource remains attached to it.

- -EINVAL is returned if the thread control block does not bear the XNSHADOW bit.
- -EBUSY is returned if either the current Linux task or the associated shadow thread is already involved in a shadow mapping.

Environments:

This service can be called from:

- Regular user-space process.

Rescheduling: always.

References `xnheap_alloc()`, `XNMAPPED`, `XNOTHER`, `xnpod_start_thread()`, `xnpod_suspend_thread()`, `XNPRIOSSET`, `XNRELAX`, `XNSHADOW`, and `xnshadow_harden()`.

4.12.2.3 `xnshadow_ppd_t* xnshadow_ppd_get (unsigned int muxid)`

Return the per-process data attached to the calling process.

This service returns the per-process data attached to the calling process for the skin whose *muxid* is *muxid*. It must be called with `nklock` locked, `irqs` off.

See `xnshadow_register_interface()` documentation for information on the way to attach a per-process data to a process.

Parameters

muxid the skin *muxid*.

Returns

the per-process data if the current context is a user-space process;
NULL otherwise.

4.12.2.4 `void xnshadow_relax (int notify, int reason)`

Switch a shadow thread back to the Linux domain.

For internal use only.

This service yields the control of the running shadow back to Linux. This is obtained by suspending the shadow and scheduling a wake up call for the mated user task inside the Linux domain. The Linux task will resume on return from `xnpod_suspend_thread()` on behalf of the root thread.

Parameters

notify A boolean flag indicating whether threads monitored from secondary mode switches should be sent a SIGDEBUG signal. For instance, some internal operations like task exit should not trigger such signal.

reason The reason to report along with the SIGDEBUG signal.

Environments:

This service can be called from:

- User-space thread operating in primary (i.e. harden) mode.

Rescheduling: always.

Note

"current" is valid here since the shadow runs with the properties of the Linux task.

References XNAFFSET, xnpod_suspend_thread(), XNPRIOSSET, XNRELAX, XNROOT, and XN-TRAPSW.

Referenced by xnpod_handle_exception(), and xnshadow_harden().

4.13 Thread synchronization services.

Files

- file [synch.c](#)
Thread synchronization services.

Functions

- void [xnsynch_init](#) (struct xnsynch *synch, xnflags_t flags, xnarch_atomic_t *fastlock)
Initialize a synchronization object.
- xnflags_t [xnsynch_sleep_on](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)
Sleep on an ownerless synchronization object.
- struct xnthread * [xnsynch_wakeup_one_sleeper](#) (struct xnsynch *synch)
Give the resource ownership to the next waiting thread.
- struct xnpholder * [xnsynch_wakeup_this_sleeper](#) (struct xnsynch *synch, struct xnpholder *holder)
Give the resource ownership to a given waiting thread.
- xnflags_t [xnsynch_acquire](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)
Acquire the ownership of a synchronization object.
- static void [xnsynch_clear_boost](#) (struct xnsynch *synch, struct xnthread *owner)
Clear the priority boost.
- void [xnsynch_requeue_sleeper](#) (struct xnthread *thread)
Change a sleeper's priority.

- struct xntthread * [xnsynch_peek_pendq](#) (struct xnsynch *synch)
Access the thread leading a synch object wait queue.
- int [xnsynch_flush](#) (struct xnsynch *synch, xnflags_t reason)
Unblock all waiters pending on a resource.
- void [xnsynch_forget_sleeper](#) (struct xntthread *thread)
Abort a wait for a resource.
- void [xnsynch_release_all_ownerships](#) (struct xntthread *thread)
Release all ownerships.
- static struct xntthread * [xnsynch_release](#) (struct xnsynch *synch, struct xntthread *thread)
Give the resource ownership to the next waiting thread.

4.13.1 Detailed Description

Thread synchronization services.

4.13.2 Function Documentation

4.13.2.1 xnflags_t xnsynch_acquire (struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode)

Acquire the ownership of a synchronization object.

This service should be called by upper interfaces wanting the current thread to acquire the ownership of the given resource. If the resource is already assigned to a thread, the caller is suspended.

This service must be used only with synchronization objects that track ownership (XNSYNCH_OWNER set).

Parameters

synch The descriptor address of the synchronization object to acquire.

timeout The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout_mode*. Passing XN_INFINITE **and** setting *mode* to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

timeout_mode The mode of the *timeout* parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also [xntimer_start\(\)](#)).

Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: possible.

References XNBOOST, XNBREAK, XNOTHER, XNPEND, `xnpod_suspend_thread()`, XNRMID, XNROBBED, XNTIMEO, and XNWAKEN.

4.13.2.2 `void xnsynch_clear_boost (struct xnsynch * synch, struct xnthread * owner)` `[static]`

Clear the priority boost.

For internal use only.

This service is called internally whenever a synchronization object is not claimed anymore by sleepers to reset the object owner's priority to its initial level.

Parameters

synch The descriptor address of the synchronization object.

owner The descriptor address of the thread which currently owns the synchronization object.

Note

This routine must be entered `nklock` locked, interrupts off.

References XNBOOST.

Referenced by `xnsynch_flush()`, and `xnsynch_forget_sleeper()`.

4.13.2.3 `int xnsynch_flush (struct xnsynch * synch, xnflags_t reason)`

Unblock all waiters pending on a resource.

This service atomically releases all threads which currently sleep on a given resource.

This service should be called by upper interfaces under circumstances requiring that the pending queue of a given resource is cleared, such as before the resource is deleted.

Parameters

synch The descriptor address of the synchronization object to be flushed.

reason Some flags to set in the information mask of every unblocked thread. Zero is an acceptable value. The following bits are pre-defined by the nucleus:

- XNRMID should be set to indicate that the synchronization object is about to be destroyed (see `xnpod_resume_thread()`).

- XNBREAK should be set to indicate that the wait has been forcibly interrupted (see [xnpod_unblock_thread\(\)](#)).

Returns

XNSYNCH_RESCHEDED is returned if at least one thread is unblocked, which means the caller should invoke [xnpod_schedule\(\)](#) for applying the new scheduling state. Otherwise, XNSYNCH_DONE is returned.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object is no more owned by any thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, [xnpod_resume_thread\(\)](#), and [xnsynch_clear_boost\(\)](#).

Referenced by [xnregistry_put\(\)](#).

4.13.2.4 void xnsynch_forget_sleeper (struct xnthread * *thread*)

Abort a wait for a resource.

For internal use only.

Performs all the necessary housekeeping chores to stop a thread from waiting on a given synchronization object.

Parameters

thread The descriptor address of the affected thread.

When the trace support is enabled (i.e. MVM), the idle state is posted to the synchronization object's state diagram (if any) whenever no thread remains blocked on it. The real-time interfaces must ensure that such condition (i.e. EMPTY/IDLE) is mapped to state #0.

Note

This routine must be entered `nklock` locked, interrupts off.

References XNPEND, and [xnsynch_clear_boost\(\)](#).

Referenced by [xnpod_delete_thread\(\)](#), [xnpod_resume_thread\(\)](#), and [xnpod_suspend_thread\(\)](#).

4.13.2.5 `void xnsynch_init (struct xnsynch * synch, xnflags_t flags, xnarch_atomic_t * fastlock)`

Initialize a synchronization object.

Initializes a new specialized object which can subsequently be used to synchronize real-time activities. The Xenomai nucleus provides a basic synchronization object which can be used to build higher resource objects. Nucleus threads can wait for and signal such objects in order to synchronize their activities.

This object has built-in support for priority inheritance.

Parameters

synch The address of a synchronization object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

flags A set of creation flags affecting the operation. The valid flags are:

- XNSYNCH_PRIO causes the threads waiting for the resource to pend in priority order. Otherwise, FIFO ordering is used (XNSYNCH_FIFO).
- XNSYNCH_OWNER indicates that the synchronization object shall track its owning thread (required if XNSYNCH_PIP is selected). Note that setting this flag implies the use `xnsynch_acquire` and `xnsynch_release` instead of `xnsynch_sleep_on` and `xnsynch_wakeup_one_sleeper/xnsynch_wakeup_this_sleeper`.
- XNSYNCH_PIP causes the priority inheritance mechanism to be automatically activated when a priority inversion is detected among threads using this object. Otherwise, no priority inheritance takes place upon priority inversion (XNSYNCH_NOPIP).
- XNSYNCH_DREORD (Disable REORdering) tells the nucleus that the wait queue should not be reordered whenever the priority of a blocked thread it holds is changed. If this flag is not specified, changing the priority of a blocked thread using `xnpod_set_thread_schedparam()` will cause this object's wait queue to be reordered according to the new priority level, provided the synchronization object makes the waiters wait by priority order on the awaited resource (XNSYNCH_PRIO).

Parameters

fastlock Address of the fast lock word to be associated with the synchronization object. If NULL is passed or XNSYNCH_OWNER is not set, fast-lock support is disabled.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

Referenced by `xnregistry_enter()`, and `xnselector_init()`.

4.13.2.6 `struct xnthread* xnsynch_peek_pendq (struct xnsynch * synch) [read]`

Access the thread leading a synch object wait queue.

This services returns the descriptor address of to the thread leading a synchronization object wait queue.

Parameters

synch The descriptor address of the target synchronization object.

Returns

The descriptor address of the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

4.13.2.7 `struct xnthread * xnsynch_release (struct xnsynch * synch, struct xnthread * owner) [static, read]`

Give the resource ownership to the next waiting thread.

This service releases the ownership of the given synchronization object. The thread which is currently leading the object's pending list, if any, is unblocked from its pending state. However, no reschedule is performed.

This service must be used only with synchronization objects that track ownership (XNSYNCH_OWNER set).

Parameters

synch The descriptor address of the synchronization object whose ownership is changed.

owner The descriptor address of the current owner.

Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNOTHER.

Referenced by `xnsynch_release_all_ownerships()`.

4.13.2.8 `void xnsynch_release_all_ownerships (struct xnthread * thread)`

Release all ownerships.

For internal use only.

This call is used internally to release all the ownerships obtained by a thread on synchronization objects. This routine must be entered interrupts off.

Parameters

thread The descriptor address of the affected thread.

Note

This routine must be entered `nklock` locked, interrupts off.

References `xnsynch_release()`.

Referenced by `__xnpod_reset_thread()`, and `xnpod_delete_thread()`.

4.13.2.9 `void xnsynch_requeue_sleeper (struct xnthread * thread)`

Change a sleeper's priority.

For internal use only.

This service is used by the PIP code to update the pending priority of a sleeping thread.

Parameters

thread The descriptor address of the affected thread.

Note

This routine must be entered `nklock` locked, interrupts off.

References XNBOOST.

4.13.2.10 `xnflags_t xnsynch_sleep_on (struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode)`

Sleep on an ownerless synchronization object.

Makes the calling thread sleep on the specified synchronization object, waiting for it to be signaled.

This service should be called by upper interfaces wanting the current thread to pend on the given resource. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER).

Parameters

synch The descriptor address of the synchronization object to sleep on.

timeout The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on *timeout_mode*. Passing XN_-INFINITE **and** setting *mode* to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.

timeout_mode The mode of the *timeout* parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also [xntimer_start\(\)](#)).

Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: always.

References XNBREAK, XNPEND, `xnpod_suspend_thread()`, XNRMID, and XNTIMEO.

Referenced by `xnregistry_bind()`, `xnregistry_remove_safe()`, and `xnselect()`.

4.13.2.11 `struct xnthread* xnsynch_wakeup_one_sleeper (struct xnsynch * synch) [read]`

Give the resource ownership to the next waiting thread.

This service wakes up the thread which is currently leading the synchronization object's pending list. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a single waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER not set).

Parameters

synch The descriptor address of the synchronization object whose ownership is changed.

Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod_resume_thread().

4.13.2.12 `struct xnpholder* xnsynch_wakeup_this_sleeper (struct xnsynch * synch, struct xnpholder * holder) [read]`

Give the resource ownership to a given waiting thread.

This service wakes up a specific thread which is currently pending on the given synchronization object. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a specific waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER not set).

Parameters

synch The descriptor address of the synchronization object whose ownership is changed.

holder The link holder address of the thread to unblock (&thread->plink) which MUST be currently linked to the synchronization object's pending queue (i.e. synch->pendq).

Returns

The link address of the unblocked thread in the synchronization object's pending queue.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transferred to the unblocked thread.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References XNPEND, and xnpod_resume_thread().

4.14 Timer services.

Files

- file [timer.h](#)
- file [timer.c](#)

Functions

- int [xntimer_start](#) (xntimer_t *timer, xnticks_t value, xnticks_t interval, xntmode_t mode)
Arm a timer.
- xnticks_t [xntimer_get_date](#) (xntimer_t *timer)
Return the absolute expiration date.
- xnticks_t [xntimer_get_timeout](#) (xntimer_t *timer)
Return the relative expiration date.
- xnticks_t [xntimer_get_interval](#) (xntimer_t *timer)
Return the timer interval value.
- static void [xntimer_stop](#) (xntimer_t *timer)
Disarm a timer.
- void [xntimer_tick](#) (void)
Process a timer tick.
- void [xntimer_init](#) (xntimer_t *timer, void(*handler)(xntimer_t *timer))
Initialize a timer object.
- void [xntimer_destroy](#) (xntimer_t *timer)
Release a timer object.
- unsigned long [xntimer_get_overruns](#) (xntimer_t *timer, xnticks_t now)
Get the count of overruns for the last tick.

- void `xntimer_freeze` (void)
Freeze all timers (from every time bases).

4.14.1 Detailed Description

The Xenomai timer facility always operate the timer hardware in oneshot mode, regardless of the time base in effect. Periodic timing is obtained through a software emulation, using cascading timers.

The timer object stores time as a count of CPU ticks (e.g. TSC values).

4.14.2 Function Documentation

4.14.2.1 void `xntimer_destroy` (`xntimer_t` * *timer*)

Release a timer object.

Destroys a timer. After it has been destroyed, all resources associated with the timer have been released. The timer is automatically deactivated before deletion if active on entry.

Parameters

timer The address of a valid timer descriptor.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task

Rescheduling: never.

References `xntimer_stop`().

Referenced by `xnpod_delete_thread`().

4.14.2.2 void `xntimer_freeze` (void)

Freeze all timers (from every time bases).

For internal use only.

This routine deactivates all active timers atomically.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task

Rescheduling: never.

References `xnsched::status`.

Referenced by `xnpod_disable_timesource()`.

4.14.2.3 `xnticks_t xntimer_get_date (xntimer_t * timer)`

Return the absolute expiration date.

Return the next expiration date of a timer as an absolute count of nanoseconds.

Parameters

timer The address of a valid timer descriptor.

Returns

The expiration date in nanoseconds. The special value `XN_INFINITE` is returned if *timer* is currently disabled.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

4.14.2.4 `xnticks_t xntimer_get_interval (xntimer_t * timer)`

Return the timer interval value.

Return the timer interval value in nanoseconds.

Parameters

timer The address of a valid timer descriptor.

Returns

The duration of a period in nanoseconds. The special value `XN_INFINITE` is returned if *timer* is currently disabled or one shot.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

4.14.2.5 unsigned long xntimer_get_overruns (xntimer_t * timer, xnticks_t now)

Get the count of overruns for the last tick.

This service returns the count of pending overruns for the last tick of a given timer, as measured by the difference between the expected expiry date of the timer and the date *now* passed as argument.

Parameters

timer The address of a valid timer descriptor.

now current date (in the monotonic time base)

Returns

the number of overruns of *timer* at date *now*

Referenced by `xnpod_wait_thread_period()`.

4.14.2.6 xnticks_t xntimer_get_timeout (xntimer_t * timer)

Return the relative expiration date.

This call returns the count of nanoseconds remaining until the timer expires.

Parameters

timer The address of a valid timer descriptor.

Returns

The count of nanoseconds until expiry. The special value `XN_INFINITE` is returned if *timer* is currently disabled. It might happen that the timer expires when this service runs (even if the associated handler has not been fired yet); in such a case, 1 is returned.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

4.14.2.7 void xntimer_init (xntimer_t * timer, void(*)(xntimer_t *timer) handler)

Initialize a timer object.

Creates a timer. When created, a timer is left disarmed; it must be started using `xntimer_start()` in order to be activated.

Parameters

timer The address of a timer descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.

handler The routine to call upon expiration of the timer.

There is no limitation on the number of timers which can be created/active concurrently.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

4.14.2.8 `int xntimer_start (xntimer_t * timer, xnticks_t value, xnticks_t interval, xntmode_t mode)`

Arm a timer.

Activates a timer so that the associated timeout handler will be fired after each expiration time. A timer can be either periodic or one-shot, depending on the reload value passed to this routine. The given timer must have been previously initialized.

Parameters

timer The address of a valid timer descriptor.

value The date of the initial timer shot, expressed in nanoseconds.

interval The reload value of the timer. It is a periodic interval value to be used for reprogramming the next timer shot, expressed in nanoseconds. If *interval* is equal to XN_INFINITE, the timer will not be reloaded after it has expired.

mode The timer mode. It can be XN_RELATIVE if *value* shall be interpreted as a relative date, XN_ABSOLUTE for an absolute date based on the monotonic clock of the related time base (as returned by `xnclock_read_monotonic()`), or XN_REALTIME if the absolute date is based on the adjustable real-time clock (obtained from `xnclock_read()`).

Returns

0 is returned upon success, or -ETIMEDOUT if an absolute date in the past has been given.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

Note

Must be called with `nklock` held, IRQs off.

Referenced by `xnpod_enable_timesource()`, `xnpod_set_thread_periodic()`, `xnpod_set_thread_timeslice()`, and `xnpod_suspend_thread()`.

4.14.2.9 int xntimer_stop (xntimer_t * timer) [inline, static]

Disarm a timer.

This service deactivates a timer previously armed using [xntimer_start\(\)](#). Once disarmed, the timer can be subsequently re-armed using the latter service.

Parameters

timer The address of a valid timer descriptor.

Environments:

This service can be called from:

- Any kernel context.

Rescheduling: never.

Note

Must be called with nklock held, IRQs off.

Referenced by xnpod_resume_thread(), xnpod_set_thread_periodic(), xnpod_set_thread_tslice(), and xntimer_destroy().

4.14.2.10 void xntimer_tick (void)

Process a timer tick.

For internal use only.

This routine informs all active timers that the clock has been updated by processing the outstanding timer list. Elapsed timer actions will be fired.

Environments:

This service can be called from:

- Interrupt service routine, nklock locked, interrupts off

Rescheduling: never.

References xnsched::htimer, xnsched::lflags, and xnsched::status.

4.15 Virtual file services**Data Structures**

- struct [xnvfile_lock_ops](#)
Vfile locking operations.

- struct [xnvmfile_regular_ops](#)
Regular vfile operation descriptor.
- struct [xnvmfile_regular_iterator](#)
Regular vfile iterator.
- struct [xnvmfile_snapshot_ops](#)
Snapshot vfile operation descriptor.
- struct [xnvmfile_rev_tag](#)
Snapshot revision tag.
- struct [xnvmfile_snapshot](#)
Snapshot vfile descriptor.
- struct [xnvmfile_snapshot_iterator](#)
Snapshot-driven vfile iterator.

Files

- file [vfile.h](#)
This file is part of the Xenomai project.

Functions

- int [xnvmfile_init_snapshot](#) (const char *name, struct [xnvmfile_snapshot](#) *vfile, struct xnvmfile_directory *parent)
Initialize a snapshot-driven vfile.
- int [xnvmfile_init_regular](#) (const char *name, struct xnvmfile_regular *vfile, struct xnvmfile_directory *parent)
Initialize a regular vfile.
- int [xnvmfile_init_dir](#) (const char *name, struct xnvmfile_directory *vdir, struct xnvmfile_directory *parent)
Initialize a virtual directory entry.
- int [xnvmfile_init_link](#) (const char *from, const char *to, struct xnvmfile_link *vlink, struct xnvmfile_directory *parent)
Initialize a virtual link entry.
- void [xnvmfile_destroy](#) (struct xnvmfile *vfile)
Removes a virtual file entry.
- ssize_t [xnvmfile_get_blob](#) (struct xnvmfile_input *input, void *data, size_t size)
Read in a data bulk written to the vfile.

- `ssize_t xnvfile_get_string` (struct `xnvfile_input` *input, char *s, `size_t` maxlen)
Read in a C-string written to the vfile.
- `ssize_t xnvfile_get_integer` (struct `xnvfile_input` *input, long *valp)
Evaluate the string written to the vfile as a long integer.

Variables

- struct `xnvfile_directory` `nkvfroot`
Xenomai vfile root directory.
- struct `xnvfile_directory` `nkvfroot`
Xenomai vfile root directory.

4.15.1 Detailed Description

Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces. This encapsulation is aimed at:

- supporting consistent collection of very large record-based output, without incurring latency peaks for undergoing real-time activities.
- in the future, hiding discrepancies between linux kernel releases, regarding the proper way to export kernel object states to userland, either via the `/proc` interface or by any other mean.

This virtual file implementation offers record-based read support based on `seq_files`, single-buffer write support, directory and link handling, all visible from the `/proc` namespace.

The vfile support exposes four filesystem object types:

- snapshot-driven file (struct `xnvfile_snapshot`). This is commonly used to export real-time object states via the `/proc` filesystem. To minimize the latency involved in protecting the vfile routines from changes applied by real-time code on such objects, a snapshot of the data to output is first taken under proper locking, before the collected data is formatted and sent out in a lockless manner.

Because a large number of records may have to be output, the data collection phase is not strictly atomic as a whole, but only protected at record level. The vfile implementation can be notified of updates to the underlying data set, and restart the collection from scratch until the snapshot is fully consistent.

- regular sequential file (struct `xnvfile_regular`). This is basically an encapsulated sequential file object as available from the host kernel (i.e. `seq_file`), with a few additional features to make it more handy in a Xenomai environment, like implicit locking support and shortened declaration for simplest, single-record output.
- virtual link (struct `xnvfile_link`). This is a symbolic link feature integrated with the vfile semantics. The link target is computed dynamically at creation time from a user-given helper routine.

- virtual directory (struct `xnvfile_directory`). A directory object, which can be used to create a hierarchy for ordering a set of vfile objects.

4.15.2 Function Documentation

4.15.2.1 void `xnvfile_destroy` (struct `xnvfile` * *vfile*)

Removes a virtual file entry.

Parameters

vfile A pointer to the virtual file descriptor to remove.

References `nkvfroot`.

4.15.2.2 ssize_t `xnvfile_get_blob` (struct `xnvfile_input` * *input*, void * *data*, size_t *size*)

Read in a data bulk written to the vfile.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_blob()` retrieves this data as an untyped binary blob, and copies it back to the caller's buffer.

Parameters

input A pointer to the input descriptor passed to the `store()` handler.

data The address of the destination buffer to copy the input data to.

size The maximum number of bytes to copy to the destination buffer. If *size* is larger than the actual data size, the input is truncated to *size*.

Returns

The number of bytes read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- `-EFAULT` indicates an invalid source buffer address.

Referenced by `xnvfile_get_integer()`, and `xnvfile_get_string()`.

4.15.2.3 ssize_t `xnvfile_get_integer` (struct `xnvfile_input` * *input*, long * *valp*)

Evaluate the string written to the vfile as a long integer.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_integer()` retrieves and interprets this data as a long integer, and copies the resulting value back to *valp*.

The long integer can be expressed in decimal, octal or hexadecimal bases depending on the prefix found.

Parameters

input A pointer to the input descriptor passed to the `store()` handler.

valp The address of a long integer variable to receive the value.

Returns

The number of characters read while evaluating the input as a long integer upon success. Otherwise, a negative error code is returned:

- -EINVAL indicates a parse error on the input stream; the written text cannot be evaluated as a long integer.
- -EFAULT indicates an invalid source buffer address.

References `xnvfile_get_blob()`.

4.15.2.4 `ssize_t xnvfile_get_string (struct xnvfile_input * input, char * s, size_t maxlen)`

Read in a C-string written to the vfile.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_string()` retrieves this data as a null-terminated character string, and copies it back to the caller's buffer.

Parameters

input A pointer to the input descriptor passed to the `store()` handler.

s The address of the destination string buffer to copy the input data to.

maxlen The maximum number of bytes to copy to the destination buffer, including the ending null character. If *maxlen* is larger than the actual string length, the input is truncated to *maxlen*.

Returns

The number of characters read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- -EFAULT indicates an invalid source buffer address.

References `xnvfile_get_blob()`.

4.15.2.5 `int xnvfile_init_dir (const char * name, struct xnvfile_directory * vdir, struct xnvfile_directory * parent)`

Initialize a virtual directory entry.

Parameters

name The name which should appear in the pseudo-filesystem, identifying the `vdir` entry.

vdir A pointer to the virtual directory descriptor to initialize.

parent A pointer to a virtual directory descriptor standing for the parent directory of the new `vdir`. If NULL, the `/proc` root directory will be used. `/proc/xenomai` is mapped on the globally available `nkvrroot` `vdir`.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual directory entry cannot be created in the /proc hierarchy.

4.15.2.6 `int xnvfile_init_link (const char * from, const char * to, struct xnvfile_link * vlink, struct xnvfile_directory * parent)`

Initialize a virtual link entry.

Parameters

from The name which should appear in the pseudo-filesystem, identifying the vlink entry.

to The target file name which should be referred to symbolically by *name*.

vlink A pointer to the virtual link descriptor to initialize.

parent A pointer to a virtual directory descriptor standing for the parent directory of the new vlink. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfroot* vdir.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual link entry cannot be created in the /proc hierarchy.

4.15.2.7 `int xnvfile_init_regular (const char * name, struct xnvfile_regular * vfile, struct xnvfile_directory * parent)`

Initialize a regular vfile.

Parameters

name The name which should appear in the pseudo-filesystem, identifying the vfile entry.

vfile A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- *.privsz* is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- *entry.lockops* is a pointer to a [locking](#) descriptor", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case no locking will be applied.
- *.ops* is a pointer to an [operation descriptor](#).

Parameters

parent A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvfroot* vdir.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

4.15.2.8 `int xnvfile_init_snapshot (const char * name, struct xnvfile_snapshot * vfile, struct xnvfile_directory * parent)`

Initialize a snapshot-driven vfile.

Parameters

name The name which should appear in the pseudo-filesystem, identifying the vfile entry.

vfile A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- .privsz is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- .datasz is the size (in bytes) of a single record to be collected by the [next\(\) handler](#) from the [operation descriptor](#).
- .tag is a pointer to a mandatory vfile revision tag structure (struct [xnvfile_rev_tag](#)). This tag will be monitored for changes by the vfile core while collecting data to output, so that any update detected will cause the current snapshot data to be dropped, and the collection to restart from the beginning. To this end, any change to the data which may be part of the collected records, should also invoke `xnvfile_touch()` on the associated tag.
- entry.lockops is a pointer to a [locking descriptor](#)", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case the operations on the nucleus lock (i.e. `nklock`) will be used internally around calls to data collection handlers (see [operation descriptor](#)).
- .ops is a pointer to an [operation descriptor](#).

Parameters

parent A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available *nkvrroot* vdir.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

References `xnvfile_snapshot_ops::store`.

4.15.3 Variable Documentation

4.15.3.1 struct `xnvfile_directory` `nkvfroot`

Xenomai vfile root directory.

This `vdir` maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by `xnvfile_destroy()`.

4.15.3.2 struct `xnvfile_directory` `nkvfroot`

Xenomai vfile root directory.

This `vdir` maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

Referenced by `xnvfile_destroy()`.

4.16 Sched

Data Structures

- struct `xnsched`
Scheduling information structure.

Files

- file `sched.h`
Scheduler interface header.
- file `sched-idle.c`
Idle scheduling class implementation (i.e. Linux placeholder).
- file `sched-rt.c`
Common real-time scheduling class implementation (FIFO + RR).
- file `sched-sporadic.c`
POSIX `SCHED_SPORADIC` scheduling class.
- file `sched-tp.c`
Temporal partitioning (typical of IMA systems).
- file `sched.c`

Typedefs

- typedef struct `xnsched` `xnsched_t`
Scheduling information structure.

Functions

- static void `xnsched_rotate` (struct `xnsched` *`sched`, struct `xnsched_class` *`sched_class`, const union `xnsched_policy_param` *`sched_param`)
Rotate a scheduler runqueue.

4.16.1 Function Documentation

4.16.1.1 void `xnsched_rotate` (struct `xnsched` * `sched`, struct `xnsched_class` * `sched_class`, const union `xnsched_policy_param` * `sched_param`) [inline, static]

Rotate a scheduler runqueue.

The specified scheduling class is requested to rotate its runqueue for the given scheduler. Rotation is performed according to the scheduling parameter specified by `sched_param`.

Note

The nucleus supports round-robin scheduling for the members of the RT class.

Parameters

`sched` The per-CPU scheduler hosting the target scheduling class.

`sched_class` The scheduling class which should rotate its runqueue.

`sched_param` The scheduling parameter providing rotation information to the specified scheduling class.

Environments:

This service should be called from:

- Kernel-based task
- Interrupt service routine
- User-space task (primary mode only)

Rescheduling: never.


Chapter 5

Data Structure Documentation

5.1 xnpod Struct Reference

Real-time pod descriptor.

Collaboration diagram for xnpod:



Data Fields

- xnflags_t [status](#)
- [xnsched_t](#) [sched](#) [XNARCH_NR_CPUS]
- xnqueue_t [threadq](#)
- xnqueue_t [tstartq](#)
- xnqueue_t [tswitchq](#)
- xnqueue_t [tdeleteq](#)
- atomic_counter_t [timerlck](#)
- int [refcnt](#)

5.1.1 Detailed Description

Real-time pod descriptor. The source of all Xenomai magic.

5.1.2 Field Documentation

5.1.2.1 `int xnpod::refcnt`

Reference count.

Referenced by `xnpod_init()`.

5.1.2.2 `xnsched_t xnpod::sched[XNARCH_NR_CPUS]`

Per-cpu scheduler slots.

Referenced by `xnpod_init()`.

5.1.2.3 `xnflags_t xnpod::status`

Status bitmask.

Referenced by `xnpod_init()`.

5.1.2.4 `xnqueue_t xnpod::tdeleteq`

Thread delete hook queue.

Referenced by `xnpod_init()`.

5.1.2.5 `xnqueue_t xnpod::threadq`

All existing threads.

Referenced by `xnpod_init()`.

5.1.2.6 `atomic_counter_t xnpod::timerlck`

Timer lock depth.

Referenced by `xnpod_init()`.

5.1.2.7 `xnqueue_t xnpod::tstartq`

Thread start hook queue.

Referenced by `xnpod_init()`.

5.1.2.8 xnqueue_t xnpod::tswitchq

Thread switch hook queue.

Referenced by xnpod_init().

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/[pod.h](#)

5.2 xnsched Struct Reference

Scheduling information structure.

Data Fields

- xnflags_t [status](#)
- xnflags_t [lflags](#)
- struct xnthread * [curr](#)
- struct xnsched_rt [rt](#)
- volatile unsigned [inesting](#)
- struct xntimer [htimer](#)
- struct xnthread [rootcb](#)

5.2.1 Detailed Description

Scheduling information structure.

5.2.2 Field Documentation

5.2.2.1 struct xnthread* xnsched::curr

Current thread.

Referenced by xnpod_delete_thread(), and xnpod_suspend_thread().

5.2.2.2 struct xntimer xnsched::htimer

Host timer.

Referenced by xnpod_enable_timesource(), and xntimer_tick().

5.2.2.3 volatile unsigned xnsched::inesting

Interrupt nesting level.

5.2.2.4 `xnflags_t xnsched::lflags`

Scheduler specific local flags bitmask.

Referenced by `xnpod_schedule()`, and `xntimer_tick()`.

5.2.2.5 `struct xnthread xnsched::rootcb`

Root thread control block.

Referenced by `xnpod_init()`.

5.2.2.6 `struct xnsched_rt xnsched::rt`

Context of built-in real-time class.

5.2.2.7 `xnflags_t xnsched::status`

Scheduler specific status bitmask.

Referenced by `xnpod_delete_thread()`, `xnpod_schedule()`, `xntimer_freeze()`, and `xntimer_tick()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/sched.h`

5.3 `xnthread_info` Struct Reference

Structure containing thread information.

Data Fields

- unsigned long `state`
Thread state.
- int `bprio`
Base priority.
- int `cprio`
Current priority.
- int `cpu`
CPU the thread currently runs on.
- unsigned long `affinity`
Thread's CPU affinity.
- unsigned long long `relpoint`
Time of next release.

- unsigned long long `exectime`
Execution time in primary mode in nanoseconds.
- unsigned long `modeswitches`
Number of primary->secondary mode switches.
- unsigned long `ctxswitches`
Number of context switches.
- unsigned long `pagefaults`
Number of triggered page faults.
- unsigned long `syscalls`
Number of Xenomai syscalls.
- char `name` [XNOBJECT_NAME_LEN]
Symbolic name assigned at creation.

5.3.1 Detailed Description

Structure containing thread information.

5.3.2 Field Documentation

5.3.2.1 unsigned long `xnthread_info::affinity`

Thread's CPU affinity.

5.3.2.2 int `xnthread_info::bprio`

Base priority.

5.3.2.3 int `xnthread_info::cprio`

Current priority.

May change through Priority Inheritance.

5.3.2.4 int `xnthread_info::cpu`

CPU the thread currently runs on.

5.3.2.5 unsigned long `xnthread_info::ctxswitches`

Number of context switches.

5.3.2.6 unsigned long long xntthread_info::exectime

Execution time in primary mode in nanoseconds.

5.3.2.7 unsigned long xntthread_info::modeswitches

Number of primary->secondary mode switches.

5.3.2.8 char xntthread_info::name[XNOBJECT_NAME_LEN]

Symbolic name assigned at creation.

5.3.2.9 unsigned long xntthread_info::pagefaults

Number of triggered page faults.

5.3.2.10 unsigned long long xntthread_info::relpoint

Time of next release.

5.3.2.11 unsigned long xntthread_info::state

Thread state,.

See also

[Thread state flags.](#)

5.3.2.12 unsigned long xntthread_info::syscalls

Number of Xenomai syscalls.

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/thread.h

5.4 xnvfile_lock_ops Struct Reference

Vfile locking operations.

Data Fields

- int(* [get](#))(struct xnvfile *vfile)
- void(* [put](#))(struct xnvfile *vfile)

5.4.1 Detailed Description

Vfile locking operations.

This structure describes the operations to be provided for implementing locking support on vfiles. They apply to both snapshot-driven and regular vfiles.

5.4.2 Field Documentation

5.4.2.1 `int(* xnvfile_lock_ops::get)(struct xnvfile *vfile)`

This handler should grab the desired lock.

Parameters

vfile A pointer to the virtual file which needs locking.

Returns

zero should be returned if the call succeeds. Otherwise, a negative error code can be returned; upon error, the current vfile operation is aborted, and the user-space caller is passed back the error value.

5.4.2.2 `void(* xnvfile_lock_ops::put)(struct xnvfile *vfile)`

This handler should release the lock previously grabbed by the [get\(\) handler](#).

Parameters

vfile A pointer to the virtual file which currently holds the lock to release.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

5.5 xnvfile_regular_iterator Struct Reference

Regular vfile iterator.

Data Fields

- `loff_t pos`
Current record position while iterating.
- `struct seq_file * seq`
Backlink to the host sequential file supporting the vfile.
- `struct xnvfile_regular * vfile`
Backlink to the vfile being read.

- char [private](#) [0]
Start of private area.

5.5.1 Detailed Description

Regular vfile iterator.

This structure defines an iterator over a regular vfile.

5.5.2 Field Documentation

5.5.2.1 `loff_t xnvfile_regular_iterator::pos`

Current record position while iterating.

5.5.2.2 `char xnvfile_regular_iterator::private[0]`

Start of private area.

Use `xnvfile_iterator_priv()` to address it.

5.5.2.3 `struct seq_file* xnvfile_regular_iterator::seq`

Backlink to the host sequential file supporting the vfile.

5.5.2.4 `struct xnvfile_regular* xnvfile_regular_iterator::vfile`

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

5.6 `xnvfile_regular_ops` Struct Reference

Regular vfile operation descriptor.

Data Fields

- `int(* rewind)(struct xnvfile_regular_iterator *it)`
- `void *(* begin)(struct xnvfile_regular_iterator *it)`
- `void *(* next)(struct xnvfile_regular_iterator *it)`
- `void(* end)(struct xnvfile_regular_iterator *it)`
- `int(* show)(struct xnvfile_regular_iterator *it, void *data)`
- `ssize_t(* store)(struct xnvfile_input *input)`

5.6.1 Detailed Description

Regular vfile operation descriptor.

This structure describes the operations available with a regular vfile. It defines handlers for sending back formatted kernel data upon a user-space read request, and for obtaining user data upon a user-space write request.

5.6.2 Field Documentation

5.6.2.1 `void>(* xnvfile_regular_ops::begin)(struct xnvfile_regular_iterator *it)`

This handler should prepare for iterating over the records upon a read request, starting from the specified position.

Parameters

- it* A pointer to the current vfile iterator. On entry, `it->pos` is set to the (0-based) position of the first record to output. This handler may be called multiple times with different position requests.

Returns

A pointer to the first record to format and output, to be passed to the [show\(\) handler](#) as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- `VFILE_SEQ_START`, a special value indicating that [the show\(\) handler](#) should receive a NULL data pointer first, in order to output a header.
- `ERR_PTR(errno)`, where `errno` is a negative error code; upon error, the current operation will be aborted immediately.

Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the [show\(\) handler](#) will be called only once for a read operation, with a NULL *data* parameter. This particular setting is convenient for simple regular vfiles having a single, fixed record to output.

5.6.2.2 `void(* xnvfile_regular_ops::end)(struct xnvfile_regular_iterator *it)`

This handler is called after all records have been output.

Parameters

- it* A pointer to the current vfile iterator.

Note

This handler is optional and the pointer may be NULL.

5.6.2.3 void>(* xnvfile_regular_ops::next)(struct xnvfile_regular_iterator *it)

This handler should return the address of the next record to format and output by the [show\(\)](#) handler".

Parameters

it A pointer to the current vfile iterator. On entry, *it->pos* is set to the (0-based) position of the next record to output.

Returns

A pointer to the next record to format and output, to be passed to the [show\(\)](#) handler as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- ERR_PTR(errno), where errno is a negative error code; upon error, the current operation will be aborted immediately.

Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the read operation will stop after the first invocation of the [show\(\)](#) handler.

5.6.2.4 int(* xnvfile_regular_ops::rewind)(struct xnvfile_regular_iterator *it)

This handler is called only once, when the virtual file is opened, before the [begin\(\)](#) handler is invoked.

Parameters

it A pointer to the vfile iterator which will be used to read the file contents.

Returns

Zero should be returned upon success. Otherwise, a negative error code aborts the operation, and is passed back to the reader.

Note

This handler is optional. It should not be used to allocate resources but rather to perform consistency checks, since no closure call is issued in case the open sequence eventually fails.

5.6.2.5 int(* xnvfile_regular_ops::show)(struct xnvfile_regular_iterator *it, void *data)

This handler should format and output a record.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

Parameters

it A pointer to the current vfile iterator.

data A pointer to the record to format then output. The first call to the handler may receive a NULL *data* pointer, depending on the presence and/or return of a [handler](#); the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- VFILE_SEQ_SKIP, a special value indicating that the current record should be skipped and will not be output.

5.6.2.6 ssize_t(* xnvfile_regular_ops::store)(struct xnvfile_input *input)

This handler receives data written to the vfile, likely for updating some kernel setting, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile_get_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile_get_integer\(\)](#) will attempt to return a long integer from the input data.

Parameters

input A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

Note

This handler is optional, and may be omitted for read-only vfiles.

The documentation for this struct was generated from the following file:

- include/cobalt/nucleus/[vfile.h](#)

5.7 xnvfile_rev_tag Struct Reference

Snapshot revision tag.

Data Fields

- int [rev](#)
Current revision number.

5.7.1 Detailed Description

Snapshot revision tag.

This structure defines a revision tag to be used with [snapshot-driven vfiles](#).

5.7.2 Field Documentation

5.7.2.1 `int xnvfile_rev_tag::rev`

Current revision number.


The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

5.8 `xnvfile_snapshot` Struct Reference

Snapshot vfile descriptor.

Collaboration diagram for `xnvfile_snapshot`:



5.8.1 Detailed Description

Snapshot vfile descriptor.

This structure describes a snapshot-driven vfile. Reading from such a vfile involves a preliminary data collection phase under lock protection, and a subsequent formatting and output phase of the collected data records. Locking is done in a way that does not increase worst-case latency, regardless of the number of records to be collected for output.

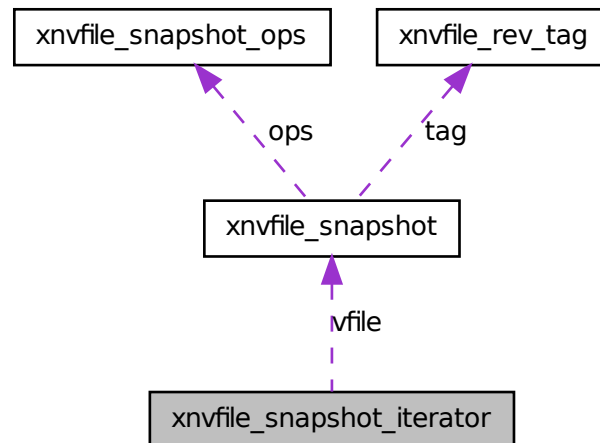
The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

5.9 xnvfile_snapshot_iterator Struct Reference

Snapshot-driven vfile iterator.

Collaboration diagram for xnvfile_snapshot_iterator:



Data Fields

- `int` `nrdata`
Number of collected records.
- `caddr_t` `databuf`
Address of record buffer.
- `struct seq_file *` `seq`
Backlink to the host sequential file supporting the vfile.
- `struct xnvfile_snapshot *` `vfile`
Backlink to the vfile being read.
- `void(* endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`
Buffer release handler.
- `char` `private` [0]
Start of private area.

5.9.1 Detailed Description

Snapshot-driven vfile iterator.

This structure defines an iterator over a snapshot-driven vfile.

5.9.2 Field Documentation

5.9.2.1 `caddr_t xnvfile_snapshot_iterator::databuf`

Address of record buffer.

5.9.2.2 `void(* xnvfile_snapshot_iterator::endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`

Buffer release handler.

5.9.2.3 `int xnvfile_snapshot_iterator::nrdata`

Number of collected records.

5.9.2.4 `char xnvfile_snapshot_iterator::private[0]`

Start of private area.

Use `xnvfile_iterator_priv()` to address it.

5.9.2.5 `struct seq_file* xnvfile_snapshot_iterator::seq`

Backlink to the host sequential file supporting the vfile.

5.9.2.6 `struct xnvfile_snapshot* xnvfile_snapshot_iterator::vfile`

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- `include/cobalt/nucleus/vfile.h`

5.10 `xnvfile_snapshot_ops` Struct Reference

Snapshot vfile operation descriptor.

Data Fields

- `int(* rewind)(struct xnvfile_snapshot_iterator *it)`
- `void(* begin)(struct xnvfile_snapshot_iterator *it)`

- void(* [end](#))(struct [xnvfile_snapshot_iterator](#) *it, void *buf)
- int(* [next](#))(struct [xnvfile_snapshot_iterator](#) *it, void *data)
- int(* [show](#))(struct [xnvfile_snapshot_iterator](#) *it, void *data)
- ssize_t(* [store](#))(struct xnvfile_input *input)

5.10.1 Detailed Description

Snapshot vfile operation descriptor.

This structure describes the operations available with a snapshot-driven vfile. It defines handlers for returning a printable snapshot of some Xenomai object contents upon a user-space read request, and for updating this object upon a user-space write request.

5.10.2 Field Documentation

5.10.2.1 void*(* xnvfile_snapshot_ops::begin)(struct xnvfile_snapshot_iterator *it)

This handler should allocate the snapshot buffer to hold records during the data collection phase. When specified, all records collected via the [next\(\)](#) handler" will be written to a cell from the memory area returned by [begin\(\)](#).

Parameters

it A pointer to the current snapshot iterator.

Returns

A pointer to the record buffer, if the call succeeds. Otherwise:

- NULL in case of allocation error. This will abort the data collection, and return -ENOMEM to the reader.
- VFILE_SEQ_EMPTY, a special value indicating that no record will be output. In such a case, the [next\(\) handler](#) will not be called, and the data collection will stop immediately. However, the [show\(\) handler](#) will still be called once, with a NULL data pointer (i.e. header display request).

Note

This handler is optional; if none is given, an internal allocation depending on the value returned by the [rewind\(\) handler](#) can be obtained.

5.10.2.2 void(* xnvfile_snapshot_ops::end)(struct xnvfile_snapshot_iterator *it, void *buf)

This handler releases the memory buffer previously obtained from [begin\(\)](#). It is usually called after the snapshot data has been output by [show\(\)](#), but it may also be called before rewinding the vfile after a revision change, to release the dropped buffer.

Parameters

it A pointer to the current snapshot iterator.

buf A pointer to the buffer to release.

Note

This routine is optional and the pointer may be NULL. It is not needed upon internal buffer allocation; see the description of the [rewind\(\)](#) handler".

5.10.2.3 `int(* xnvfile_snapshot_ops::next)(struct xnvfile_snapshot_iterator *it, void *data)`

This handler fetches the next record, as part of the snapshot data to be sent back to the reader via the [show\(\)](#).

Parameters

it A pointer to the current snapshot iterator.

data A pointer to the record to fill in.

Returns

a strictly positive value, if the call succeeds and leaves a valid record into *data*, which should be passed to the [show\(\) handler\(\)](#) during the formatting and output phase. Otherwise:

- A negative error code. This will abort the data collection, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped. In such a case, the *data* pointer is not advanced to the next position before the [next\(\) handler](#) is called anew.

Note

This handler is called with the vfile lock held. Before each invocation of this handler, the vfile core checks whether the revision tag has been touched, in which case the data collection is restarted from scratch. A data collection phase succeeds whenever all records can be fetched via the [next\(\) handler](#), while the revision tag remains unchanged, which indicates that a consistent snapshot of the object state was taken.

5.10.2.4 `int(* xnvfile_snapshot_ops::rewind)(struct xnvfile_snapshot_iterator *it)`

This handler (re-)initializes the data collection, moving the seek pointer at the first record. When the file revision tag is touched while collecting data, the current reading is aborted, all collected data dropped, and the vfile is eventually rewound.

Parameters

it A pointer to the current snapshot iterator. Two useful information can be retrieved from this iterator in this context:

- *it->vfile* is a pointer to the descriptor of the virtual file being rewound.
- `xnvfile_iterator_priv(it)` returns a pointer to the private data area, available from the descriptor, which size is *vfile->privsz*. If the latter size is zero, the returned pointer is meaningless and should not be used.

Returns

A negative error code aborts the data collection, and is passed back to the reader. Otherwise:

- a strictly positive value is interpreted as the total number of records which will be returned by the [next\(\) handler](#) during the data collection phase. If no [begin\(\) handler](#) is provided in the [operation descriptor](#), this value is used to allocate the snapshot buffer internally. The size of this buffer would then be `vfile->datasz * value`.
- zero leaves the allocation to the [begin\(\) handler](#) if present, or indicates that no record is to be output in case such handler is not given.

Note

This handler is optional; a NULL value indicates that nothing needs to be done for rewinding the vfile. It is called with the vfile lock held.

5.10.2.5 `int(* xnvfile_snapshot_ops::show)(struct xnvfile_snapshot_iterator *it, void *data)`

This handler should format and output a record from the collected data.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

Parameters

it A pointer to the current snapshot iterator.

data A pointer to the record to format then output. The first call to the handler is always passed a NULL *data* pointer; the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

5.10.2.6 `ssize_t(* xnvfile_snapshot_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating the associated Xenomai object's state, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the `/proc` entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile_get_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile_get_integer\(\)](#) will attempt to return a long integer from the input data.

Parameters

input A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.

Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

Note

This handler is optional, and may be omitted for read-only vfiles.

Referenced by `xnvfile_init_snapshot()`.

The documentation for this struct was generated from the following file:


- `include/cobalt/nucleus/vfile.h`

Chapter 6


File Documentation

6.1 include/cobalt/nucleus/bufd.h File Reference

Include dependency graph for bufd.h:



This graph shows which files directly or indirectly include this file:



Functions

- static void [xnbufd_map_uread](#) (struct xnbufd *bufd, const void __user *ptr, size_t len)
Initialize a buffer descriptor for reading from user memory.
- static void [xnbufd_map_uwrite](#) (struct xnbufd *bufd, void __user *ptr, size_t len)
Initialize a buffer descriptor for writing to user memory.
- ssize_t [xnbufd_unmap_uread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uread\(\)](#).
- ssize_t [xnbufd_unmap_uwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uwrite\(\)](#).
- static void [xnbufd_map_kread](#) (struct xnbufd *bufd, const void *ptr, size_t len)
Initialize a buffer descriptor for reading from kernel memory.
- static void [xnbufd_map_kwrite](#) (struct xnbufd *bufd, void *ptr, size_t len)
Initialize a buffer descriptor for writing to kernel memory.
- ssize_t [xnbufd_unmap_kread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kread\(\)](#).
- ssize_t [xnbufd_unmap_kwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kwrite\(\)](#).
- ssize_t [xnbufd_copy_to_kmem](#) (void *ptr, struct xnbufd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- ssize_t [xnbufd_copy_from_kmem](#) (struct xnbufd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- void [xnbufd_invalidate](#) (struct xnbufd *bufd)
Invalidate a buffer descriptor.
- static void [xnbufd_reset](#) (struct xnbufd *bufd)
Reset a buffer descriptor.

6.1.1 Detailed Description

Note

Copyright (C) 2009 Philippe Gerum <rpm@xenomai.org>.

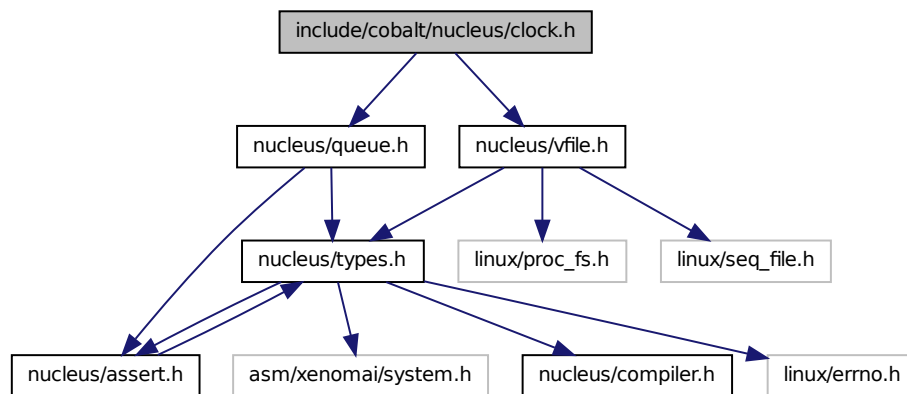
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.2 include/cobalt/nucleus/clock.h File Reference

Include dependency graph for clock.h:



This graph shows which files directly or indirectly include this file:



Functions

- void [xnclock_adjust](#) (xnsticks_t delta)
Adjust the clock time for the system.

6.2.1 Detailed Description

Note

Copyright (C) 2006,2007 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


6.3 include/cobalt/nucleus/hostrt.h File Reference

Definitions for global semaphore heap shared objects.

Include dependency graph for hostrt.h:



This graph shows which files directly or indirectly include this file:



6.3.1 Detailed Description

Definitions for global semaphore heap shared objects.

Author

Wolfgang Mauerer

Copyright (C) 2010 Wolfgang Mauerer <wolfgang.mauerer@siemens.com>.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.4 include/cobalt/nucleus/map.h File Reference

Include dependency graph for map.h:



This graph shows which files directly or indirectly include this file:



Functions

- `xnmap_t * xnmap_create` (`int nkeys`, `int reserve`, `int offset`)
Create a map.
- `void xnmap_delete` (`xnmap_t *map`)
Delete a map.
- `int xnmap_enter` (`xnmap_t *map`, `int key`, `void *objaddr`)
Index an object into a map.
- `int xnmap_remove` (`xnmap_t *map`, `int key`)
Remove an object reference from a map.

- static void * [xnmap_fetch_nocheck](#) (xnmap_t *map, int key)

Search an object into a map - unchecked form.

- static void * [xnmap_fetch](#) (xnmap_t *map, int key)

Search an object into a map.

6.4.1 Detailed Description

Note

Copyright (C) 2007 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.5 include/cobalt/nucleus/pod.h File Reference

Real-time pod interface header.

Include dependency graph for pod.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [xnpod](#)
Real-time pod descriptor.

Functions

- void [__xnpod_reset_thread](#) (struct xnthread *thread)
Reset the thread.
- int [xnpod_init](#) (void)
Initialize the core pod.
- int [xnpod_enable_timesource](#) (void)
Activate the core time source.
- void [xnpod_disable_timesource](#) (void)
Stop the core time source.
- void [xnpod_shutdown](#) (int xtype)
Shutdown the current pod.
- int [xnpod_init_thread](#) (struct xnthread *thread, const struct xnthread_init_attr *attr, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)
Initialize a new thread.
- int [xnpod_start_thread](#) (xnthread_t *thread, const struct xnthread_start_attr *attr)
Initial start of a newly created thread.
- void [xnpod_stop_thread](#) (xnthread_t *thread)
Stop a thread.
- void [xnpod_delete_thread](#) (xnthread_t *thread)
Delete a thread.
- void [xnpod_abort_thread](#) (xnthread_t *thread)
Abort a thread.
- xnflags_t [xnpod_set_thread_mode](#) (xnthread_t *thread, xnflags_t clrmask, xnflags_t set-mask)
Change a thread's control mode.

- void [xn timer suspend_thread](#) (xn timer_t *thread, xn timer_flags_t mask, xn timer_ticks_t timeout, xn timer_mode_t timeout_mode, struct xn timer_synch *wchan)
Suspend a thread.
- void [xn timer resume_thread](#) (xn timer_t *thread, xn timer_flags_t mask)
Resume a thread.
- int [xn timer unblock_thread](#) (xn timer_t *thread)
Unblock a thread.
- int [xn timer set_thread_schedparam](#) (struct xn timer_thread *thread, struct xn timer_sched_class *sched_class, const union xn timer_sched_policy_param *sched_param)
Change the base scheduling parameters of a thread.
- int [xn timer migrate_thread](#) (int cpu)
Migrate the current thread.
- void [xn timer dispatch_signals](#) (void)
Deliver pending asynchronous signals to the running thread.
- static void [xn timer schedule](#) (void)
Rescheduling procedure entry point.
- int [xn timer handle_exception](#) (struct ipipe_trap_data *d)
Exception handler.
- int [xn timer set_thread_periodic](#) (xn timer_t *thread, xn timer_ticks_t idate, xn timer_mode_t timeout_mode, xn timer_ticks_t period)
Make a thread periodic.
- int [xn timer wait_thread_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.
- int [xn timer set_thread_tslice](#) (struct xn timer_thread *thread, xn timer_ticks_t quantum)
Set thread time-slicing information.
- int [xn timer add_hook](#) (int type, void(*routine)(xn timer_t *))
Install a nucleus hook.
- int [xn timer remove_hook](#) (int type, void(*routine)(xn timer_t *))
Remove a nucleus hook.

6.5.1 Detailed Description

Real-time pod interface header.

Author

Philippe Gerum

Copyright (C) 2001-2007 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2004 The Xenomai project <<http://www.xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

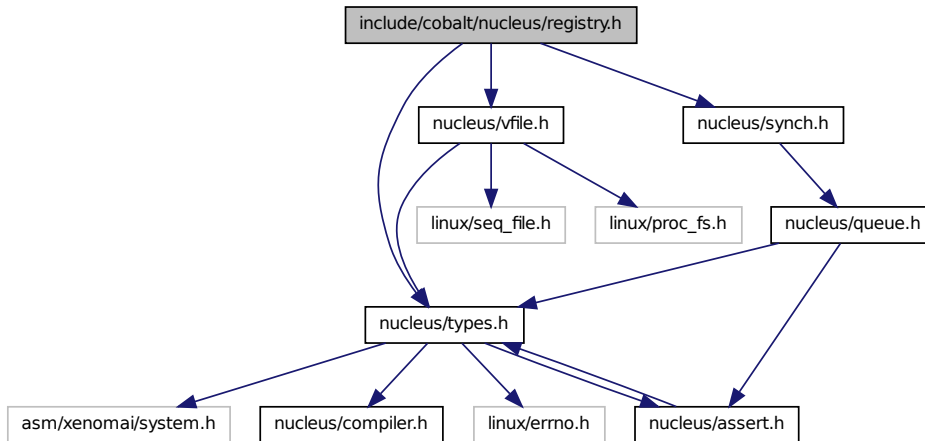
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


6.6 include/cobalt/nucleus/registry.h File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.h:



This graph shows which files directly or indirectly include this file:



Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, struct xninode *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_remove_safe](#) (xnhandle_t handle, xnticks_t timeout)
Unregister an idle real-time object.
- void * [xnregistry_get](#) (xnhandle_t handle)
Find and lock a real-time object into the registry.
- void * [xnregistry_fetch](#) (xnhandle_t handle)
Find a real-time object into the registry.
- u_long [xnregistry_put](#) (xnhandle_t handle)
Unlock a real-time object from the registry.

6.6.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2004 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.7 include/cobalt/nucleus/sched-idle.h File Reference

Definitions for the IDLE scheduling class.

This graph shows which files directly or indirectly include this file:



6.7.1 Detailed Description

Definitions for the IDLE scheduling class.

Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.8 include/cobalt/nucleus/sched-rt.h File Reference

Definitions for the RT scheduling class.

This graph shows which files directly or indirectly include this file:



6.8.1 Detailed Description

Definitions for the RT scheduling class.

Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.9 include/cobalt/nucleus/sched-sporadic.h File Reference

Definitions for the SSP scheduling class.

This graph shows which files directly or indirectly include this file:



6.9.1 Detailed Description

Definitions for the SSP scheduling class.

Author

Philippe Gerum

Copyright (C) 2009 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.10 include/cobalt/nucleus/sched-tp.h File Reference

Definitions for the TP scheduling class.

This graph shows which files directly or indirectly include this file:



6.10.1 Detailed Description

Definitions for the TP scheduling class.

Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

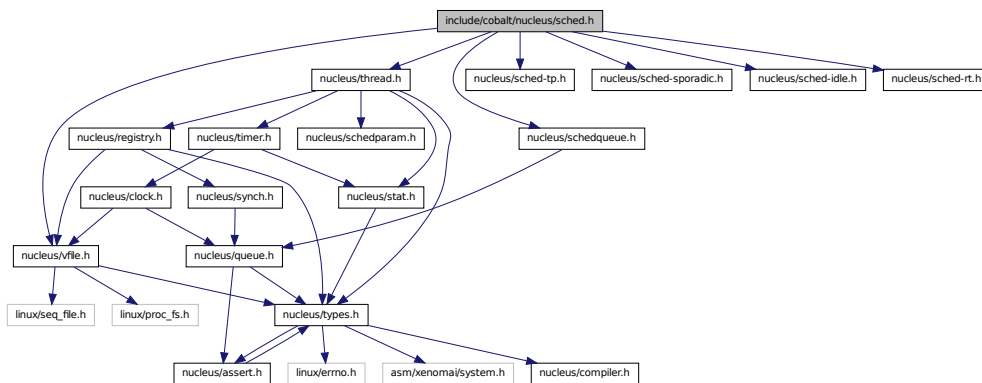
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.11 include/cobalt/nucleus/sched.h File Reference

Scheduler interface header.

Include dependency graph for sched.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [xnsched](#)

Scheduling information structure.

Typedefs

- typedef struct [xnsched](#) [xnsched_t](#)

Scheduling information structure.

Functions

- static void [xnsched_rotate](#) (struct [xnsched](#) *sched, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)

Rotate a scheduler runqueue.

6.11.1 Detailed Description

Scheduler interface header.

Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

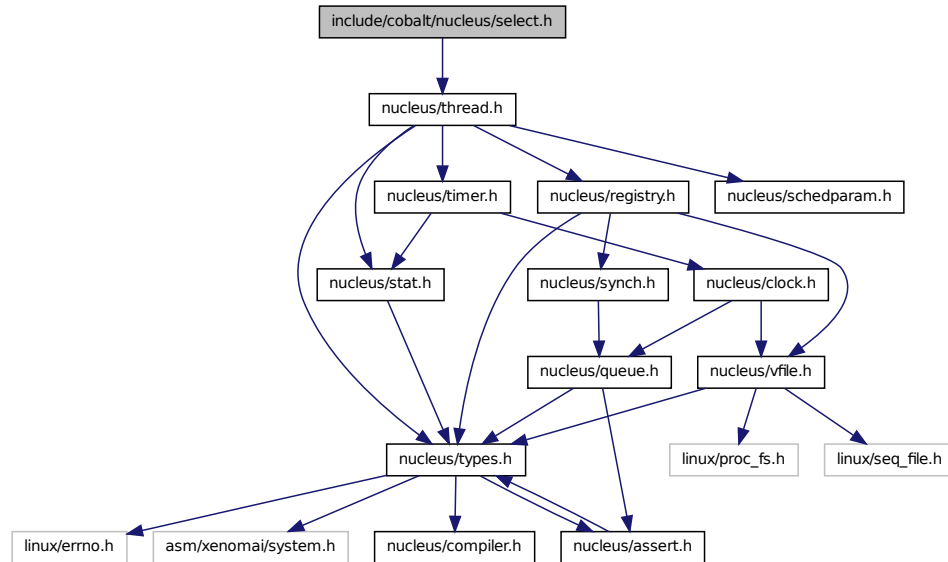
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


6.12 include/cobalt/nucleus/select.h File Reference

file descriptors events multiplexing header.

Include dependency graph for select.h:



This graph shows which files directly or indirectly include this file:



Functions

- void [xnselect_init](#) (struct xnselect *select_block)
Initialize a struct xnselect structure.
- int [xnselect_bind](#) (struct xnselect *select_block, struct xnselect_binding *binding, struct xnselector *selector, unsigned type, unsigned index, unsigned state)
Bind a file descriptor (represented by its xnselect structure) to a selector block.

- static int [xnselect_signal](#) (struct xnselect *select_block, unsigned state)

Signal a file descriptor state change.

- void [xnselect_destroy](#) (struct xnselect *select_block)

Destroy the xnselect structure associated with a file descriptor.

- int [xnselector_init](#) (struct xnselector *selector)

Initialize a selector structure.

- int [xnselect](#) (struct xnselector *selector, fd_set *out_fds[XNSELECT_MAX_TYPES], fd_set *in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)

Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.

- void [xnselector_destroy](#) (struct xnselector *selector)

Destroy a selector block.

6.12.1 Detailed Description

file descriptors events multiplexing header.

Author

Gilles Chanteperdrix

Copyright (C) 2008 Efixo <gilles.chanteperdrix@xenomai.org>

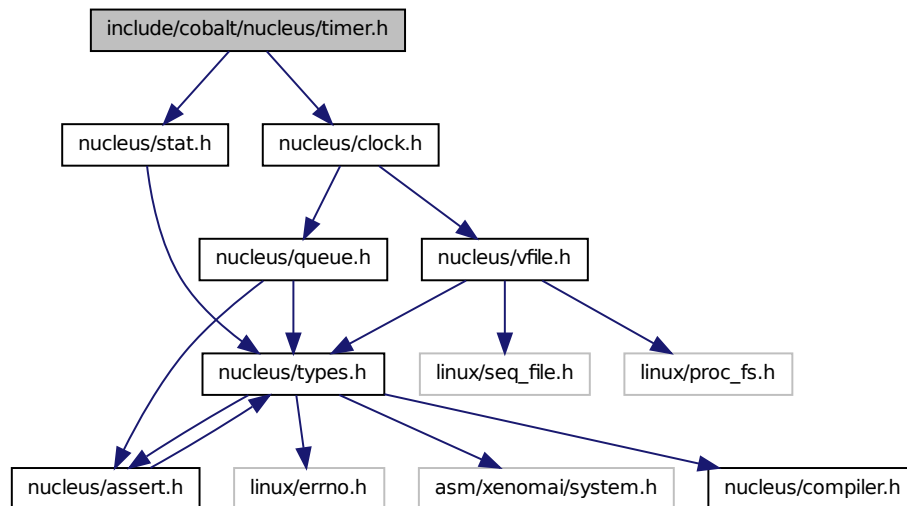
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.13 include/cobalt/nucleus/timer.h File Reference

Include dependency graph for timer.h:



This graph shows which files directly or indirectly include this file:



Functions

- void [`xntimer_destroy`](#) (`xntimer_t *timer`)
Release a timer object.
- int [`xntimer_start`](#) (`xntimer_t *timer`, `xnticks_t value`, `xnticks_t interval`, `xntmode_t mode`)
Arm a timer.
- `xnticks_t` [`xntimer_get_date`](#) (`xntimer_t *timer`)
Return the absolute expiration date.
- `xnticks_t` [`xntimer_get_timeout`](#) (`xntimer_t *timer`)
Return the relative expiration date.
- `xnticks_t` [`xntimer_get_interval`](#) (`xntimer_t *timer`)
Return the timer interval value.

- static void `xntimer_stop` (xntimer_t *timer)
Disarm a timer.
- unsigned long `xntimer_get_overruns` (xntimer_t *timer, xnticks_t now)
Get the count of overruns for the last tick.
- void `xntimer_freeze` (void)
Freeze all timers (from every time bases).
- void `xntimer_tick` (void)
Process a timer tick.

6.13.1 Detailed Description

Note

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


6.14 include/cobalt/nucleus/vdso.h File Reference

Definitions for global semaphore heap shared objects.

Include dependency graph for vdso.h:



This graph shows which files directly or indirectly include this file:



6.14.1 Detailed Description

Definitions for global semaphore heap shared objects.

Author

Wolfgang Mauerer

Copyright (C) 2009 Wolfgang Mauerer <wolfgang.mauerer@siemens.com>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

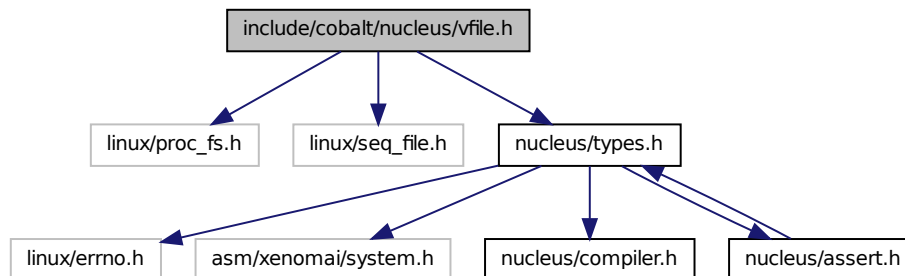
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


6.15 include/cobalt/nucleus/vfile.h File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [xnfile_lock_ops](#)
Vfile locking operations.
- struct [xnfile_regular_ops](#)
Regular vfile operation descriptor.
- struct [xnfile_regular_iterator](#)
Regular vfile iterator.
- struct [xnfile_snapshot_ops](#)
Snapshot vfile operation descriptor.
- struct [xnfile_rev_tag](#)
Snapshot revision tag.
- struct [xnfile_snapshot](#)
Snapshot vfile descriptor.
- struct [xnfile_snapshot_iterator](#)
Snapshot-driven vfile iterator.

Functions

- int [xnvmfile_init_snapshot](#) (const char *name, struct [xnvmfile_snapshot](#) *vfile, struct [xnvmfile_directory](#) *parent)
Initialize a snapshot-driven vfile.
- int [xnvmfile_init_regular](#) (const char *name, struct [xnvmfile_regular](#) *vfile, struct [xnvmfile_directory](#) *parent)
Initialize a regular vfile.
- int [xnvmfile_init_dir](#) (const char *name, struct [xnvmfile_directory](#) *vdir, struct [xnvmfile_directory](#) *parent)
Initialize a virtual directory entry.
- int [xnvmfile_init_link](#) (const char *from, const char *to, struct [xnvmfile_link](#) *vlink, struct [xnvmfile_directory](#) *parent)
Initialize a virtual link entry.
- void [xnvmfile_destroy](#) (struct [xnvmfile](#) *vfile)
Removes a virtual file entry.
- ssize_t [xnvmfile_get_blob](#) (struct [xnvmfile_input](#) *input, void *data, size_t size)
Read in a data bulk written to the vfile.
- ssize_t [xnvmfile_get_string](#) (struct [xnvmfile_input](#) *input, char *s, size_t maxlen)
Read in a C-string written to the vfile.
- ssize_t [xnvmfile_get_integer](#) (struct [xnvmfile_input](#) *input, long *valp)
Evaluate the string written to the vfile as a long integer.

Variables

- struct [xnvmfile_directory](#) [nkvfroot](#)
Xenomai vfile root directory.

6.15.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2010 Philippe Gerum <rpm@xenomai.org>

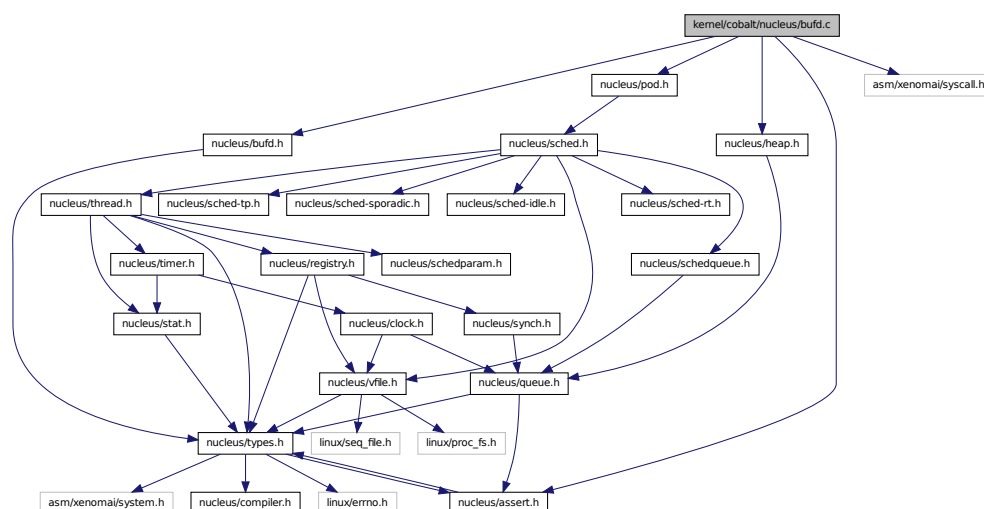
This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.16 kernel/cobalt/nucleus/bufd.c File Reference

Include dependency graph for bufd.c:



Functions

- `ssize_t xnbufd_copy_to_kmem` (void *ptr, struct xnbufd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- `ssize_t xnbufd_copy_from_kmem` (struct xnbufd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- `ssize_t xnbufd_unmap_uread` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from `xnbufd_map_uread()`.
- `ssize_t xnbufd_unmap_uwrite` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from `xnbufd_map_uwrite()`.
- `void xnbufd_invalidate` (struct xnbufd *bufd)
Invalidate a buffer descriptor.
- `ssize_t xnbufd_unmap_kread` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from `xnbufd_map_kread()`.
- `ssize_t xnbufd_unmap_kwrite` (struct xnbufd *bufd)

Finalize a buffer descriptor obtained from `xnbufd_map_kwrite()`.

6.16.1 Detailed Description

Note

Copyright (C) 2009 Philippe Gerum <rpm@xenomai.org>.

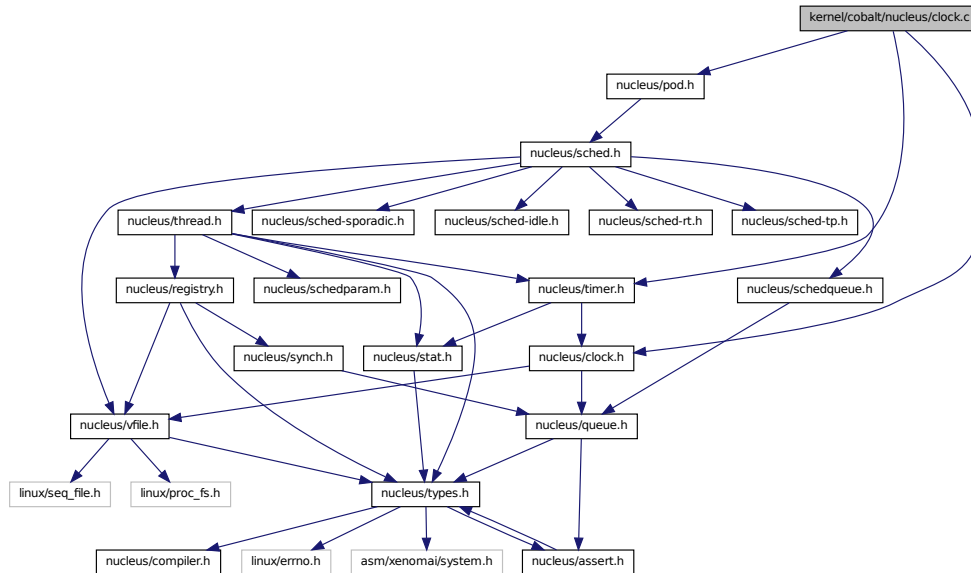
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.17 kernel/cobalt/nucleus/clock.c File Reference

Include dependency graph for clock.c:



Functions

- void `xnlock_adjust` (xnsticks_t delta)
Adjust the clock time for the system.

6.17.1 Detailed Description

Note

Copyright (C) 2006-2011 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

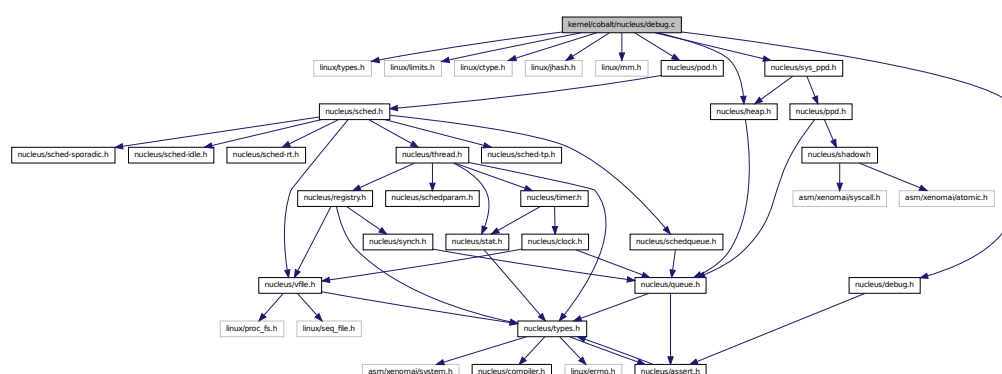
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.18 kernel/cobalt/nucleus/debug.c File Reference

Debug services.

Include dependency graph for debug.c:



6.18.1 Detailed Description

Debug services.

Author

Philippe Gerum

Copyright (C) 2010 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

- void [xnheap_schedule_free](#) (xnheap_t *heap, void *block, xnholder_t *link)

Schedule a memory block for release.

6.19.1 Detailed Description

Dynamic memory allocation services.

Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

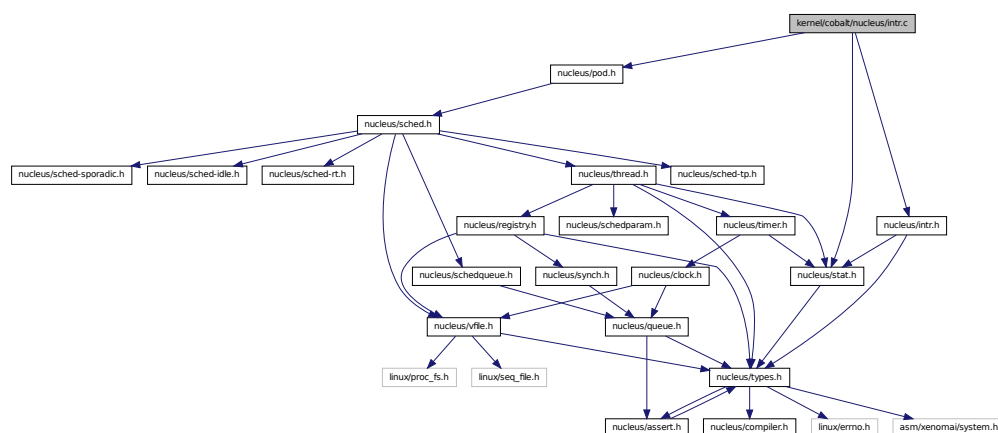
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.20 kernel/cobalt/nucleus/intr.c File Reference

Interrupt management.

Include dependency graph for intr.c:



Functions

- int [xnintr_init](#) (xnintr_t *intr, const char *name, unsigned irq, xnintr_t isr, xniack_t iack, xnflags_t flags)
Initialize an interrupt object.
- int [xnintr_destroy](#) (xnintr_t *intr)
Destroy an interrupt object.
- int [xnintr_attach](#) (xnintr_t *intr, void *cookie)
Attach an interrupt object.
- int [xnintr_detach](#) (xnintr_t *intr)
Detach an interrupt object.
- void [xnintr_enable](#) (xnintr_t *intr)
Enable an interrupt object.
- void [xnintr_disable](#) (xnintr_t *intr)
Disable an interrupt object.
- void [xnintr_affinity](#) (xnintr_t *intr, xnarch_cpumask_t cpumask)
Set interrupt's processor affinity.

6.20.1 Detailed Description

Interrupt management.

Author

Philippe Gerum

Copyright (C) 2001,2002,2003 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2005,2006 Dmitry Adamushko <dmitry.adamushko@gmail.com>. Copyright (C) 2007 Jan Kiszka <jan.kiszka@web.de>.

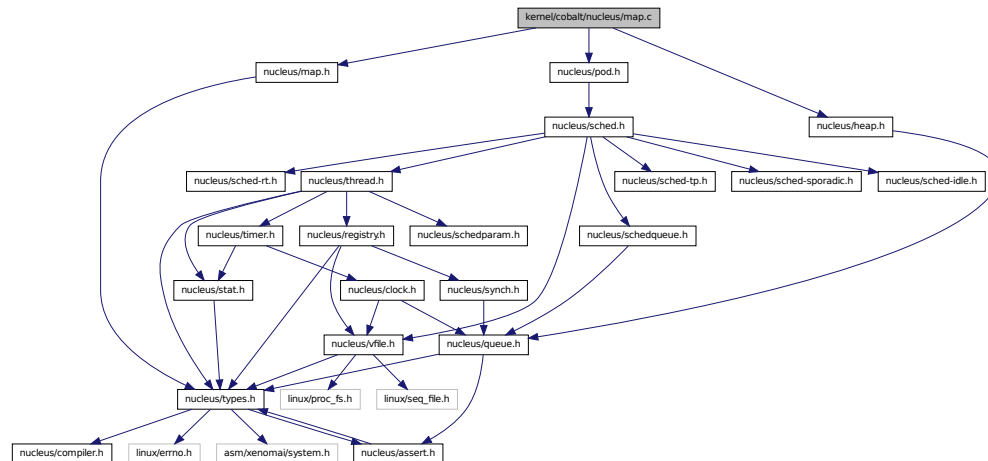
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.21 kernel/cobalt/nucleus/map.c File Reference

Include dependency graph for map.c:



Functions

- `xnmap_t * xnmap_create (int nkeys, int reserve, int offset)`
Create a map.
- `void xnmap_delete (xnmap_t *map)`
Delete a map.
- `int xnmap_enter (xnmap_t *map, int key, void *objaddr)`
Index an object into a map.
- `int xnmap_remove (xnmap_t *map, int key)`
Remove an object reference from a map.

6.21.1 Detailed Description

Note

Copyright (C) 2007 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

Delete a thread.

- void [xn timer abort thread](#) (xn timer_t *thread)
Abort a thread.
- void [xn timer suspend thread](#) (xn timer_t *thread, xn timer_flags_t mask, xn timer_ticks_t timeout, xn timer_mode_t timeout_mode, struct xn timer_synch *wchan)
Suspend a thread.
- void [xn timer resume thread](#) (xn timer_t *thread, xn timer_flags_t mask)
Resume a thread.
- int [xn timer unblock thread](#) (xn timer_t *thread)
Unblock a thread.
- int [xn timer set thread sched param](#) (struct xn timer_t *thread, struct xn timer_sched_class *sched_class, const union xn timer_sched_policy_param *sched_param)
Change the base scheduling parameters of a thread.
- int [xn timer migrate thread](#) (int cpu)
Migrate the current thread.
- void [xn timer dispatch signals](#) (void)
Deliver pending asynchronous signals to the running thread.
- void [xn timer welcome thread](#) (xn timer_t *thread, int imask)
Thread prologue.
- int [xn timer add hook](#) (int type, void(*routine)(xn timer_t *))
Install a nucleus hook.
- int [xn timer remove hook](#) (int type, void(*routine)(xn timer_t *))
Remove a nucleus hook.
- int [xn timer handle exception](#) (struct ipipe_trap_data *d)
Exception handler.
- int [xn timer enable timesource](#) (void)
Activate the core time source.
- void [xn timer disable timesource](#) (void)
Stop the core time source.
- int [xn timer set thread periodic](#) (xn timer_t *thread, xn timer_ticks_t idate, xn timer_mode_t timeout_mode, xn timer_ticks_t period)
Make a thread periodic.
- int [xn timer wait thread period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.

- int `xnpod_set_thread_tslice` (struct `xnthread` *thread, `xnticks_t` quantum)

Set thread time-slicing information.

6.22.1 Detailed Description

Real-time pod services.

Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.Xenomai.org>>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

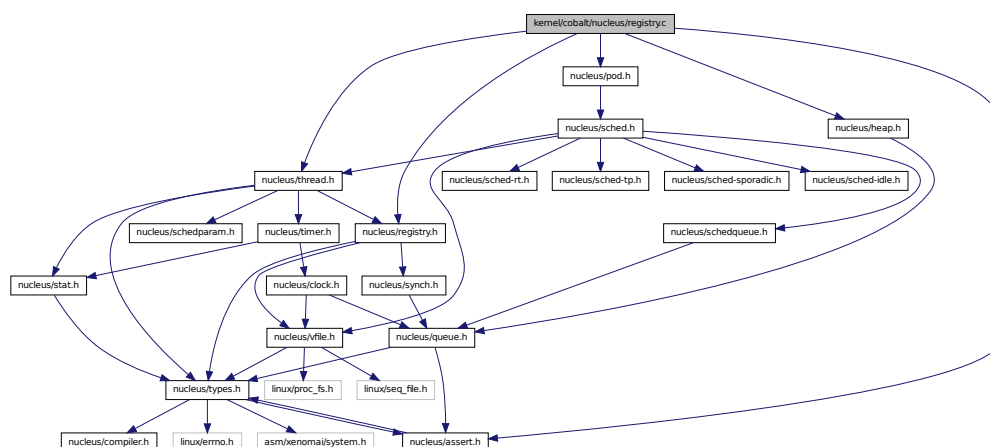
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.23 kernel/cobalt/nucleus/registry.c File Reference

This file is part of the Xenomai project.

Include dependency graph for `registry.c`:



Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, struct xninode *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_remove_safe](#) (xnhandle_t handle, xnticks_t timeout)
Unregister an idle real-time object.
- void * [xnregistry_get](#) (xnhandle_t handle)
Find and lock a real-time object into the registry.
- u_long [xnregistry_put](#) (xnhandle_t handle)
Unlock a real-time object from the registry.
- void * [xnregistry_fetch](#) (xnhandle_t handle)
Find a real-time object into the registry.

6.23.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2004 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.24 kernel/cobalt/nucleus/sched-idle.c File Reference

Idle scheduling class implementation (i.e. Linux placeholder).

Include dependency graph for sched-idle.c:



6.24.1 Detailed Description

Idle scheduling class implementation (i.e. Linux placeholder).

Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.25 kernel/cobalt/nucleus/sched-rt.c File Reference

Common real-time scheduling class implementation (FIFO + RR).

Include dependency graph for sched-rt.c:



6.25.1 Detailed Description

Common real-time scheduling class implementation (FIFO + RR).

Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

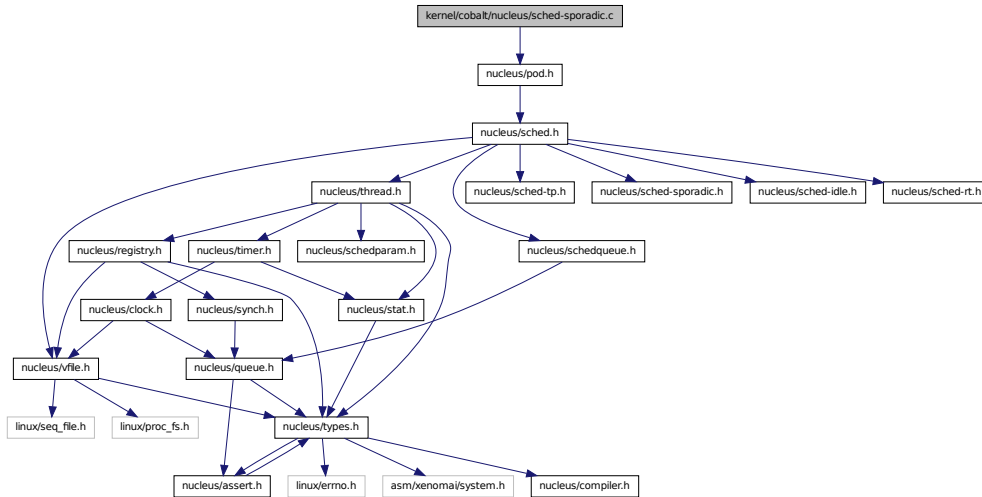
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.26 kernel/cobalt/nucleus/sched-sporadic.c File Reference

POSIX SCHED_SPORADIC scheduling class.

Include dependency graph for sched-sporadic.c:



6.26.1 Detailed Description

POSIX SCHED_SPORADIC scheduling class.

Author

Philippe Gerum Copyright (C) 2009 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

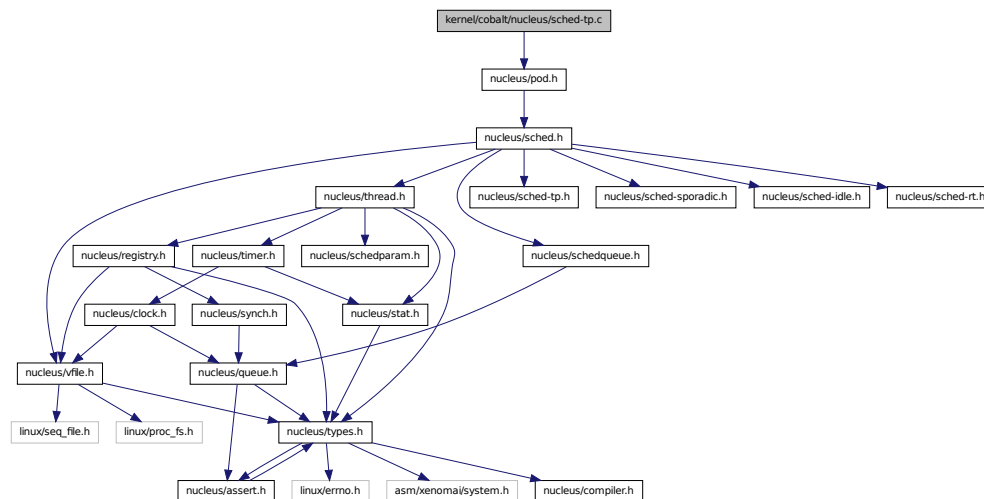
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.27 kernel/cobalt/nucleus/sched-tp.c File Reference

Temporal partitioning (typical of IMA systems).

Include dependency graph for sched-tp.c:



6.27.1 Detailed Description

Temporal partitioning (typical of IMA systems).

Author

Philippe Gerum Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

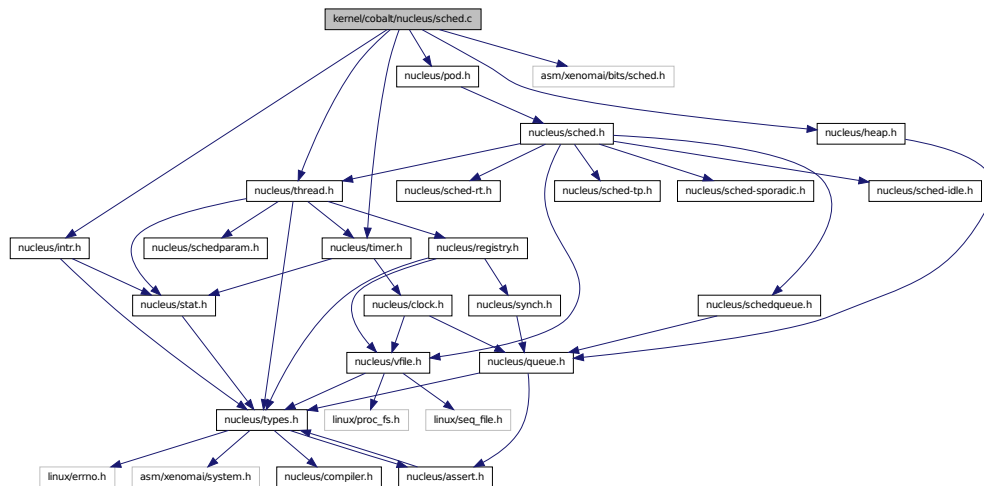
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.28 kernel/cobalt/nucleus/sched.c File Reference

Include dependency graph for sched.c:



6.28.1 Detailed Description

Author

Philippe Gerum

Copyright (C) 2008 Philippe Gerum <rpm@xenomai.org>.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

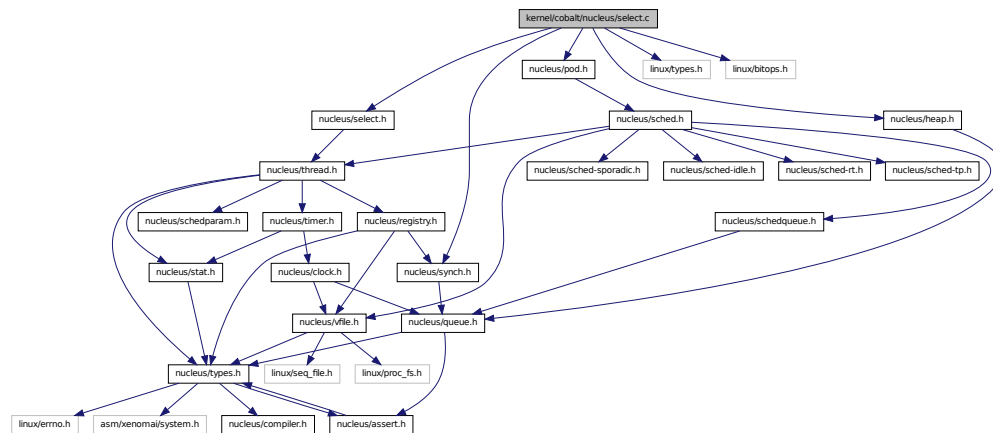
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.29 kernel/cobalt/nucleus/select.c File Reference

file descriptors events multiplexing.

Include dependency graph for select.c:



Functions

- void [xnselect_init](#) (struct xnselect *select_block)
Initialize a struct xnselect structure.
- int [xnselect_bind](#) (struct xnselect *select_block, struct xnselect_binding *binding, struct xnselector *selector, unsigned type, unsigned index, unsigned state)
Bind a file descriptor (represented by its xnselect structure) to a selector block.
- void [xnselect_destroy](#) (struct xnselect *select_block)
Destroy the xnselect structure associated with a file descriptor.
- int [xnselector_init](#) (struct xnselector *selector)
Initialize a selector structure.
- int [xnselect](#) (struct xnselector *selector, fd_set *out_fds[XNSELECT_MAX_TYPES], fd_set *in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)
Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.
- void [xnselector_destroy](#) (struct xnselector *selector)
Destroy a selector block.

6.29.1 Detailed Description

file descriptors events multiplexing.

Author

Gilles Chanteperdrix

Copyright (C) 2008 Efixo <gilles.chanteperdrix@xenomai.org>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.30 kernel/cobalt/nucleus/shadow.c File Reference

Real-time shadow services.

Include dependency graph for shadow.c:



Functions

- `int xnshadow_harden` (void)
Migrate a Linux task to the Xenomai domain.
- `void xnshadow_relax` (int notify, int reason)
Switch a shadow thread back to the Linux domain.
- `int xnshadow_map` (xnthread_t *thread, xncompletion_t __user *u_completion, unsigned long __user *u_window_offset)
Create a shadow thread context.
- `xnshadow_ppd_t * xnshadow_ppd_get` (unsigned int muxid)
Return the per-process data attached to the calling process.

6.30.1 Detailed Description

Real-time shadow services.

Author

Philippe Gerum

Copyright (C) 2001-2012 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 The RTAI project <<http://www.rtai.org>> Copyright (C) 2004 The HYADES project <<http://www.hyades-itea.org>> Copyright (C) 2005 The Xenomai project <<http://www.xenomai.org>> Copyright (C) 2006 Gilles Chantepredrix <gilles.chantepredrix@xenomai.org>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

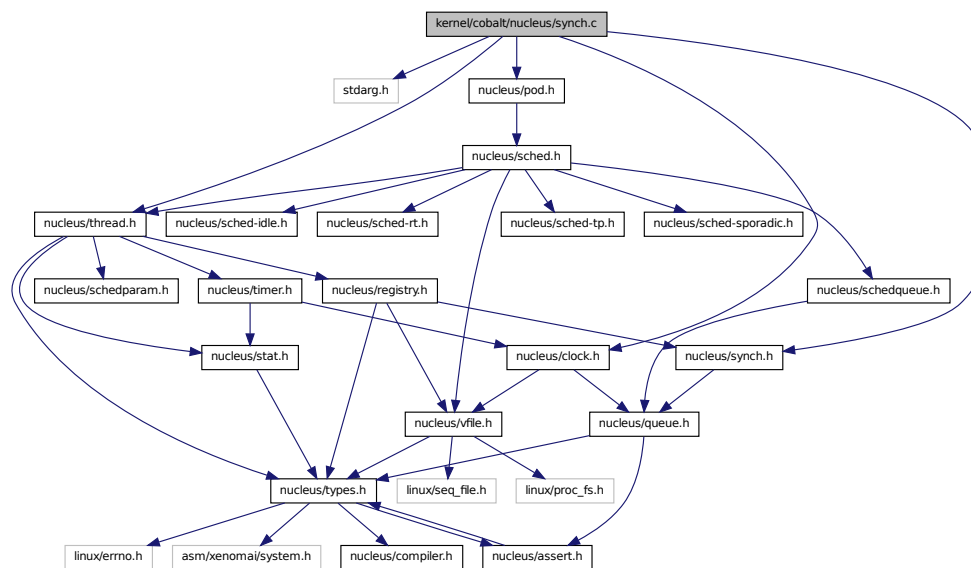
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.31 kernel/cobalt/nucleus/synch.c File Reference

Thread synchronization services.

Include dependency graph for synch.c:



Functions

- void [xnsynch_init](#) (struct xnsynch *synch, xnflags_t flags, xnarch_atomic_t *fastlock)
Initialize a synchronization object.
- xnflags_t [xnsynch_sleep_on](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)

Sleep on an ownerless synchronization object.

- struct xnthread * [xnsynch_wakeup_one_sleeper](#) (struct xnsynch *synch)
Give the resource ownership to the next waiting thread.
- struct xnpholder * [xnsynch_wakeup_this_sleeper](#) (struct xnsynch *synch, struct xnpholder *holder)
Give the resource ownership to a given waiting thread.
- xnflags_t [xnsynch_acquire](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)
Acquire the ownership of a synchronization object.
- static void [xnsynch_clear_boost](#) (struct xnsynch *synch, struct xnthread *owner)
Clear the priority boost.
- void [xnsynch_requeue_sleeper](#) (struct xnthread *thread)
Change a sleeper's priority.
- struct xnthread * [xnsynch_peek_pendq](#) (struct xnsynch *synch)
Access the thread leading a synch object wait queue.
- int [xnsynch_flush](#) (struct xnsynch *synch, xnflags_t reason)
Unblock all waiters pending on a resource.
- void [xnsynch_forget_sleeper](#) (struct xnthread *thread)
Abort a wait for a resource.
- void [xnsynch_release_all_ownerships](#) (struct xnthread *thread)
Release all ownerships.

6.31.1 Detailed Description

Thread synchronization services.

Author

Philippe Gerum

Copyright (C) 2001-2008 Philippe Gerum <rpm@xenomai.org>.

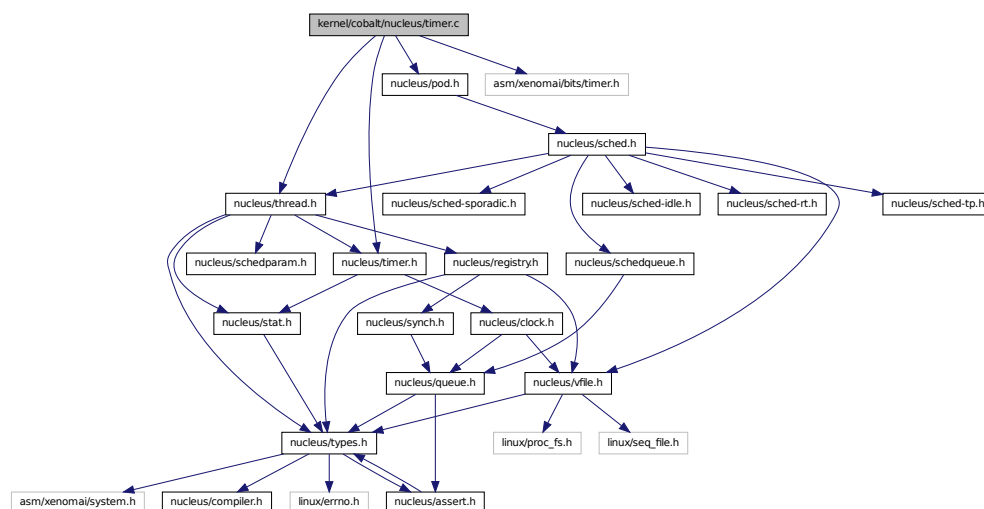
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.32 kernel/cobalt/nucleus/timer.c File Reference

Include dependency graph for timer.c:



Functions

- `int xntimer_start(xntimer_t *timer, xnticks_t value, xnticks_t interval, xntmode_t mode)`
Arm a timer.
- `xnticks_t xntimer_get_date(xntimer_t *timer)`
Return the absolute expiration date.
- `xnticks_t xntimer_get_timeout(xntimer_t *timer)`
Return the relative expiration date.
- `xnticks_t xntimer_get_interval(xntimer_t *timer)`
Return the timer interval value.
- `void xntimer_tick(void)`
Process a timer tick.
- `void xntimer_init(xntimer_t *timer, void(*handler)(xntimer_t *timer))`
Initialize a timer object.
- `void xntimer_destroy(xntimer_t *timer)`
Release a timer object.
- `unsigned long xntimer_get_overruns(xntimer_t *timer, xnticks_t now)`
Get the count of overruns for the last tick.

- void `xntimer_freeze` (void)

Freeze all timers (from every time bases).

6.32.1 Detailed Description

Note

Copyright (C) 2001,2002,2003,2007 Philippe Gerum <rpm@xenomai.org>. Copyright (C) 2004 Gilles Chantepredrix <gilles.chantepredrix@xenomai.org>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

6.33 kernel/cobalt/nucleus/vfile.c File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.c:



Functions

- int [xnvmfile_init_snapshot](#) (const char *name, struct [xnvmfile_snapshot](#) *vfile, struct [xnvmfile_directory](#) *parent)
Initialize a snapshot-driven vfile.
- int [xnvmfile_init_regular](#) (const char *name, struct [xnvmfile_regular](#) *vfile, struct [xnvmfile_directory](#) *parent)
Initialize a regular vfile.
- int [xnvmfile_init_dir](#) (const char *name, struct [xnvmfile_directory](#) *vdir, struct [xnvmfile_directory](#) *parent)
Initialize a virtual directory entry.
- int [xnvmfile_init_link](#) (const char *from, const char *to, struct [xnvmfile_link](#) *vlink, struct [xnvmfile_directory](#) *parent)
Initialize a virtual link entry.
- void [xnvmfile_destroy](#) (struct [xnvmfile](#) *vfile)
Removes a virtual file entry.
- ssize_t [xnvmfile_get_blob](#) (struct [xnvmfile_input](#) *input, void *data, size_t size)
Read in a data bulk written to the vfile.
- ssize_t [xnvmfile_get_string](#) (struct [xnvmfile_input](#) *input, char *s, size_t maxlen)
Read in a C-string written to the vfile.
- ssize_t [xnvmfile_get_integer](#) (struct [xnvmfile_input](#) *input, long *valp)
Evaluate the string written to the vfile as a long integer.

Variables

- struct [xnvmfile_directory](#) [nkvfroot](#)
Xenomai vfile root directory.

6.33.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2010 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Index

- [__xnpod_reset_thread](#)
 - [pod](#), [39](#)
- [affinity](#)
 - [xnthread_info](#), [97](#)
- [begin](#)
 - [xnfile_regular_ops](#), [101](#)
 - [xnfile_snapshot_ops](#), [107](#)
- [bprio](#)
 - [xnthread_info](#), [97](#)
- [bufd](#)
 - [xnbufd_copy_from_kmem](#), [13](#)
 - [xnbufd_copy_to_kmem](#), [14](#)
 - [xnbufd_invalidate](#), [15](#)
 - [xnbufd_map_kread](#), [15](#)
 - [xnbufd_map_kwrite](#), [16](#)
 - [xnbufd_map_uread](#), [16](#)
 - [xnbufd_map_uwrite](#), [17](#)
 - [xnbufd_reset](#), [17](#)
 - [xnbufd_unmap_kread](#), [17](#)
 - [xnbufd_unmap_kwrite](#), [18](#)
 - [xnbufd_unmap_uread](#), [18](#)
 - [xnbufd_unmap_uwrite](#), [19](#)
- [Buffer descriptors.](#), [10](#)
- [clock](#)
 - [xnclock_adjust](#), [20](#)
- [cprio](#)
 - [xnthread_info](#), [97](#)
- [cpu](#)
 - [xnthread_info](#), [97](#)
- [ctxswitches](#)
 - [xnthread_info](#), [97](#)
- [curr](#)
 - [xnsched](#), [95](#)
- [databuf](#)
 - [xnfile_snapshot_iterator](#), [106](#)
- [Debugging services.](#), [21](#)
- [Dynamic memory allocation services.](#), [21](#)
- [end](#)
 - [xnfile_regular_ops](#), [101](#)
 - [xnfile_snapshot_ops](#), [107](#)
- [endfn](#)
 - [xnfile_snapshot_iterator](#), [106](#)
- [exectime](#)
 - [xnthread_info](#), [97](#)
- [File descriptors events multiplexing services.](#), [62](#)
- [get](#)
 - [xnfile_lock_ops](#), [99](#)
- [heap](#)
 - [xnheap_alloc](#), [22](#)
 - [xnheap_destroy](#), [23](#)
 - [xnheap_extend](#), [23](#)
 - [xnheap_free](#), [24](#)
 - [xnheap_init](#), [24](#)
 - [xnheap_schedule_free](#), [25](#)
 - [xnheap_set_label](#), [26](#)
 - [xnheap_test_and_free](#), [26](#)
- [htimer](#)
 - [xnsched](#), [95](#)
- [include/cobalt/nucleus/bufd.h](#), [111](#)
- [include/cobalt/nucleus/clock.h](#), [113](#)
- [include/cobalt/nucleus/hostrt.h](#), [114](#)
- [include/cobalt/nucleus/map.h](#), [116](#)
- [include/cobalt/nucleus/pod.h](#), [117](#)
- [include/cobalt/nucleus/registry.h](#), [120](#)
- [include/cobalt/nucleus/sched-idle.h](#), [121](#)
- [include/cobalt/nucleus/sched-rt.h](#), [122](#)
- [include/cobalt/nucleus/sched-sporadic.h](#), [123](#)
- [include/cobalt/nucleus/sched-tp.h](#), [123](#)
- [include/cobalt/nucleus/sched.h](#), [124](#)
- [include/cobalt/nucleus/select.h](#), [125](#)
- [include/cobalt/nucleus/timer.h](#), [128](#)
- [include/cobalt/nucleus/vdso.h](#), [129](#)
- [include/cobalt/nucleus/vfile.h](#), [130](#)
- [inesting](#)
 - [xnsched](#), [95](#)
- [Interrupt management.](#), [27](#)
- [intr](#)
 - [xnintr_affinity](#), [28](#)
 - [xnintr_attach](#), [28](#)
 - [xnintr_destroy](#), [29](#)
 - [xnintr_detach](#), [29](#)

- xnintr_disable, 30
 - xnintr_enable, 30
 - xnintr_init, 31
- kernel/cobalt/nucleus/bufd.c, 133
- kernel/cobalt/nucleus/clock.c, 134
- kernel/cobalt/nucleus/debug.c, 135
- kernel/cobalt/nucleus/heap.c, 136
- kernel/cobalt/nucleus/intr.c, 137
- kernel/cobalt/nucleus/map.c, 139
- kernel/cobalt/nucleus/pod.c, 140
- kernel/cobalt/nucleus/registry.c, 142
- kernel/cobalt/nucleus/sched-idle.c, 143
- kernel/cobalt/nucleus/sched-rt.c, 144
- kernel/cobalt/nucleus/sched-sporadic.c, 145
- kernel/cobalt/nucleus/sched-tp.c, 146
- kernel/cobalt/nucleus/sched.c, 148
- kernel/cobalt/nucleus/select.c, 148
- kernel/cobalt/nucleus/shadow.c, 150
- kernel/cobalt/nucleus/synch.c, 151
- kernel/cobalt/nucleus/timer.c, 153
- kernel/cobalt/nucleus/vfile.c, 154
- lflags
 - xnsched, 95
- Lightweight key-to-object mapping service, 33
- map
 - xnmap_create, 33
 - xnmap_delete, 34
 - xnmap_enter, 34
 - xnmap_fetch, 35
 - xnmap_fetch_nocheck, 36
 - xnmap_remove, 36
- modeswitches
 - xnthread_info, 98
- name
 - xnthread_info, 98
- next
 - xnvfile_regular_ops, 101
 - xnvfile_snapshot_ops, 108
- nkvfroot
 - vfile, 90
- nrdata
 - xnvfile_snapshot_iterator, 106
- nucleus_state_flags
 - XNHeld, 8
 - XNLOCK, 8
 - XNMIGRATE, 9
 - XNPEND, 9
 - XNREADY, 9
 - XNSUSP, 9
- pagefaults
 - xnthread_info, 98
- pod
 - __xnpod_reset_thread, 39
 - xnpod_abort_thread, 39
 - xnpod_add_hook, 40
 - xnpod_delete_thread, 41
 - xnpod_disable_timesource, 41
 - xnpod_dispatch_signals, 42
 - xnpod_enable_timesource, 42
 - xnpod_handle_exception, 42
 - xnpod_init, 43
 - xnpod_init_thread, 43
 - xnpod_migrate_thread, 45
 - xnpod_remove_hook, 45
 - xnpod_resume_thread, 45
 - xnpod_schedule, 46
 - xnpod_set_thread_mode, 48
 - xnpod_set_thread_periodic, 48
 - xnpod_set_thread_schedparam, 49
 - xnpod_set_thread_tslice, 50
 - xnpod_shutdown, 51
 - xnpod_start_thread, 51
 - xnpod_stop_thread, 53
 - xnpod_suspend_thread, 53
 - xnpod_unblock_thread, 54
 - xnpod_wait_thread_period, 55
 - xnpod_welcome_thread, 56
- pos
 - xnvfile_regular_iterator, 100
- private
 - xnvfile_regular_iterator, 100
 - xnvfile_snapshot_iterator, 106
- put
 - xnvfile_lock_ops, 99
- Real-time pod services., 37
- Real-time shadow services., 66
- refcnt
 - xnpod, 94
- registry
 - xnregistry_bind, 57
 - xnregistry_enter, 58
 - xnregistry_fetch, 59
 - xnregistry_get, 59
 - xnregistry_put, 60
 - xnregistry_remove, 61
 - xnregistry_remove_safe, 61
- Registry services., 56
- relpoint
 - xnthread_info, 98
- rev
 - xnvfile_rev_tag, 104
- rewind
 - xnvfile_regular_ops, 102

- xnvfile_snapshot_ops, 108
- rootcb
 - xnsched, 96
- rt
 - xnsched, 96
- Sched, 90
- sched
 - xnpod, 94
 - xnsched_rotate, 91
- select
 - xnselect, 63
 - xnselect_bind, 64
 - xnselect_destroy, 64
 - xnselect_init, 65
 - xnselect_signal, 65
 - xnselector_destroy, 65
 - xnselector_init, 65
- seq
 - xnvfile_regular_iterator, 100
 - xnvfile_snapshot_iterator, 106
- shadow
 - xnshadow_harden, 66
 - xnshadow_map, 67
 - xnshadow_ppd_get, 68
 - xnshadow_relax, 68
- show
 - xnvfile_regular_ops, 102
 - xnvfile_snapshot_ops, 109
- state
 - xnthread_info, 98
- status
 - xnpod, 94
 - xnsched, 96
- store
 - xnvfile_regular_ops, 103
 - xnvfile_snapshot_ops, 109
- synch
 - xnsynch_acquire, 70
 - xnsynch_clear_boost, 71
 - xnsynch_flush, 71
 - xnsynch_forget_sleeper, 72
 - xnsynch_init, 72
 - xnsynch_peek_pendq, 73
 - xnsynch_release, 74
 - xnsynch_release_all_ownerships, 75
 - xnsynch_requeue_sleeper, 75
 - xnsynch_sleep_on, 75
 - xnsynch_wakeup_one_sleeper, 76
 - xnsynch_wakeup_this_sleeper, 77
- syscalls
 - xnthread_info, 98
- System clock services., 20
- tdeleteq
 - xnpod, 94
- Thread information flags., 9
- Thread state flags., 7
- Thread synchronization services., 69
- threadq
 - xnpod, 94
- timer
 - xntimer_destroy, 79
 - xntimer_freeze, 79
 - xntimer_get_date, 80
 - xntimer_get_interval, 80
 - xntimer_get_overruns, 80
 - xntimer_get_timeout, 81
 - xntimer_init, 81
 - xntimer_start, 82
 - xntimer_stop, 82
 - xntimer_tick, 83
- Timer services., 78
- timerlck
 - xnpod, 94
- tstartq
 - xnpod, 94
- tswitchq
 - xnpod, 94
- vfile
 - nkvfroot, 90
 - xnvfile_destroy, 86
 - xnvfile_get_blob, 86
 - xnvfile_get_integer, 86
 - xnvfile_get_string, 87
 - xnvfile_init_dir, 87
 - xnvfile_init_link, 88
 - xnvfile_init_regular, 88
 - xnvfile_init_snapshot, 89
 - xnvfile_regular_iterator, 100
 - xnvfile_snapshot_iterator, 106
- Virtual file services, 83
- xnbufd_copy_from_kmem
 - bufd, 13
- xnbufd_copy_to_kmem
 - bufd, 14
- xnbufd_invalidate
 - bufd, 15
- xnbufd_map_kread
 - bufd, 15
- xnbufd_map_kwrite
 - bufd, 16
- xnbufd_map_uread
 - bufd, 16
- xnbufd_map_uwrite
 - bufd, 17

- xnbufd_reset
 - bufd, [17](#)
- xnbufd_unmap_kread
 - bufd, [17](#)
- xnbufd_unmap_kwrite
 - bufd, [18](#)
- xnbufd_unmap_uread
 - bufd, [18](#)
- xnbufd_unmap_uwrite
 - bufd, [19](#)
- xnclock_adjust
 - clock, [20](#)
- xnheap_alloc
 - heap, [22](#)
- xnheap_destroy
 - heap, [23](#)
- xnheap_extend
 - heap, [23](#)
- xnheap_free
 - heap, [24](#)
- xnheap_init
 - heap, [24](#)
- xnheap_schedule_free
 - heap, [25](#)
- xnheap_set_label
 - heap, [26](#)
- xnheap_test_and_free
 - heap, [26](#)
- XNHELD
 - nucleus_state_flags, [8](#)
- xnintr_affinity
 - intr, [28](#)
- xnintr_attach
 - intr, [28](#)
- xnintr_destroy
 - intr, [29](#)
- xnintr_detach
 - intr, [29](#)
- xnintr_disable
 - intr, [30](#)
- xnintr_enable
 - intr, [30](#)
- xnintr_init
 - intr, [31](#)
- XNLOCK
 - nucleus_state_flags, [8](#)
- xnmap_create
 - map, [33](#)
- xnmap_delete
 - map, [34](#)
- xnmap_enter
 - map, [34](#)
- xnmap_fetch
 - map, [35](#)
- xnmap_fetch_nocheck
 - map, [36](#)
- xnmap_remove
 - map, [36](#)
- XNMIGRATE
 - nucleus_state_flags, [9](#)
- XNPEND
 - nucleus_state_flags, [9](#)
- xnpod, [93](#)
 - refcnt, [94](#)
 - sched, [94](#)
 - status, [94](#)
 - tdeleteq, [94](#)
 - threadq, [94](#)
 - timerlck, [94](#)
 - tstartq, [94](#)
 - tswitchq, [94](#)
- xnpod_abort_thread
 - pod, [39](#)
- xnpod_add_hook
 - pod, [40](#)
- xnpod_delete_thread
 - pod, [41](#)
- xnpod_disable_timesource
 - pod, [41](#)
- xnpod_dispatch_signals
 - pod, [42](#)
- xnpod_enable_timesource
 - pod, [42](#)
- xnpod_handle_exception
 - pod, [42](#)
- xnpod_init
 - pod, [43](#)
- xnpod_init_thread
 - pod, [43](#)
- xnpod_migrate_thread
 - pod, [45](#)
- xnpod_remove_hook
 - pod, [45](#)
- xnpod_resume_thread
 - pod, [45](#)
- xnpod_schedule
 - pod, [46](#)
- xnpod_set_thread_mode
 - pod, [48](#)
- xnpod_set_thread_periodic
 - pod, [48](#)
- xnpod_set_thread_schedparam
 - pod, [49](#)
- xnpod_set_thread_tslice
 - pod, [50](#)
- xnpod_shutdown
 - pod, [51](#)
- xnpod_start_thread

- pod, [51](#)
- xnpod_stop_thread
 - pod, [53](#)
- xnpod_suspend_thread
 - pod, [53](#)
- xnpod_unblock_thread
 - pod, [54](#)
- xnpod_wait_thread_period
 - pod, [55](#)
- xnpod_welcome_thread
 - pod, [56](#)
- XNREADY
 - nucleus_state_flags, [9](#)
- xnregistry_bind
 - registry, [57](#)
- xnregistry_enter
 - registry, [58](#)
- xnregistry_fetch
 - registry, [59](#)
- xnregistry_get
 - registry, [59](#)
- xnregistry_put
 - registry, [60](#)
- xnregistry_remove
 - registry, [61](#)
- xnregistry_remove_safe
 - registry, [61](#)
- xnsched, [95](#)
 - curr, [95](#)
 - htimer, [95](#)
 - inesting, [95](#)
 - lflags, [95](#)
 - rootcb, [96](#)
 - rt, [96](#)
 - status, [96](#)
- xnsched_rotate
 - sched, [91](#)
- xnselect
 - select, [63](#)
- xnselect_bind
 - select, [64](#)
- xnselect_destroy
 - select, [64](#)
- xnselect_init
 - select, [65](#)
- xnselect_signal
 - select, [65](#)
- xnselector_destroy
 - select, [65](#)
- xnselector_init
 - select, [65](#)
- xnshadow_harden
 - shadow, [66](#)
- xnshadow_map
 - shadow, [67](#)
- xnshadow_ppd_get
 - shadow, [68](#)
- xnshadow_relax
 - shadow, [68](#)
- XNSUSP
 - nucleus_state_flags, [9](#)
- xnsynch_acquire
 - synch, [70](#)
- xnsynch_clear_boost
 - synch, [71](#)
- xnsynch_flush
 - synch, [71](#)
- xnsynch_forget_sleeper
 - synch, [72](#)
- xnsynch_init
 - synch, [72](#)
- xnsynch_peek_pendq
 - synch, [73](#)
- xnsynch_release
 - synch, [74](#)
- xnsynch_release_all_ownerships
 - synch, [75](#)
- xnsynch_requeue_sleeper
 - synch, [75](#)
- xnsynch_sleep_on
 - synch, [75](#)
- xnsynch_wakeup_one_sleeper
 - synch, [76](#)
- xnsynch_wakeup_this_sleeper
 - synch, [77](#)
- xnthread_info, [96](#)
 - affinity, [97](#)
 - bprio, [97](#)
 - cprio, [97](#)
 - cpu, [97](#)
 - ctxswitches, [97](#)
 - exectime, [97](#)
 - modeswitches, [98](#)
 - name, [98](#)
 - pagefaults, [98](#)
 - relpoint, [98](#)
 - state, [98](#)
 - syscalls, [98](#)
- xntimer_destroy
 - timer, [79](#)
- xntimer_freeze
 - timer, [79](#)
- xntimer_get_date
 - timer, [80](#)
- xntimer_get_interval
 - timer, [80](#)
- xntimer_get_overruns
 - timer, [80](#)

xntimer_get_timeout
 timer, 81
xntimer_init
 timer, 81
xntimer_start
 timer, 82
xntimer_stop
 timer, 82
xntimer_tick
 timer, 83
xnvfile_destroy
 vfile, 86
xnvfile_get_blob
 vfile, 86
xnvfile_get_integer
 vfile, 86
xnvfile_get_string
 vfile, 87
xnvfile_init_dir
 vfile, 87
xnvfile_init_link
 vfile, 88
xnvfile_init_regular
 vfile, 88
xnvfile_init_snapshot
 vfile, 89
xnvfile_lock_ops, 98
 get, 99
 put, 99
xnvfile_regular_iterator, 99
 pos, 100
 private, 100
 seq, 100
 vfile, 100
xnvfile_regular_ops, 100
 begin, 101
 end, 101
 next, 101
 rewind, 102
 show, 102
 store, 103
xnvfile_rev_tag, 103
 rev, 104
xnvfile_snapshot, 104
xnvfile_snapshot_iterator, 105
 databuf, 106
 endfn, 106
 nrdata, 106
 private, 106
 seq, 106
 vfile, 106
xnvfile_snapshot_ops, 106
 begin, 107
 end, 107
 next, 108
 rewind, 108
 show, 109
 store, 109