Xenomai RTDM skin API

2.5.90

Generated by Doxygen 1.7.1

Contents

Chapter 1

Deprecated List

Global rtdm_device::open_rt Only use non-real-time open handler in new drivers.

Global rtdm_device::socket_rt Only use non-real-time socket creation handler in new drivers.

Global rtdm_operations::close_rt Only use non-real-time close handler in new drivers.

Global rtdm_task_sleep_until(nanosecs_abs_t wakeup_time) Use rtdm_task_sleep_abs instead!

Chapter 2

Module Index

2.1 Modules

Here is a list of all modules:

Chapter 3

Data Structure Index

3.1 Data Structures

Here are the data structures with brief descriptions:

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

Chapter 5

Module Documentation

5.1   CAN Devices

Collaboration diagram for CAN Devices:

Data Structures

- struct can_bittime_std

    Standard bit-time parameters according to Bosch.

- struct can_bittime_btr

    Hardware-specific BTR bit-times.

- struct can_bittime

    Custom CAN bit-time definition.

- struct can_filter

    Filter for reception of CAN messages.

- struct sockaddr_can

    Socket address structure for the CAN address family.

- struct can_frame

    Raw CAN frame.

Files

- file rtcan.h

  Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

Defines

- #define AF_CAN 29

  CAN address family.

- #define PF_CAN AF_CAN

  CAN protocol family.

- #define SOL_CAN_RAW 103

  CAN socket levels.

Typedefs

- typedef uint32_t can_id_t

  Type of CAN id (see CAN_xxx_MASK and CAN_xxx_FLAG).

- typedef can_id_t can_err_mask_t

  Type of CAN error mask.

- typedef uint32_t can_baudrate_t

  Baudrate definition in bits per second.

- typedef enum CAN_BITTIME_TYPE can_bittime_type_t

  See CAN_BITTIME_TYPE.

- typedef enum CAN_MODE can_mode_t

  See CAN_MODE.

- typedef int can_ctrlmode_t

  See CAN_CTRLMODE.

- typedef enum CAN_STATE can_state_t

  See CAN_STATE.

- typedef struct can_filter can_filter_t

  Filter for reception of CAN messages.

- typedef struct can_frame can_frame_t

  Raw CAN frame.

Enumerations

- enum CAN_BITTIME_TYPE { CAN_BITTIME_STD, CAN_BITTIME_BTR }

  Supported CAN bit-time types.

CAN operation modes

Modes into which CAN controllers can be set

- enum CAN_MODE { CAN_MODE_STOP = 0, CAN_MODE_START, CAN_MODE_SLEEP }

CAN controller states

States a CAN controller can be in.

- enum CAN_STATE {

  CAN_STATE_ACTIVE = 0, CAN_STATE_BUS_WARNING, CAN_STATE_BUS_PASSIVE, CAN_STATE_BUS_OFF,

  CAN_STATE_SCANNING_BAUDRATE,   CAN_STATE_STOPPED,   CAN_STATE_-SLEEPING }

CAN ID masks

Bit masks for masking CAN IDs

- #define CAN_EFF_MASK 0x1FFFFFFF

  Bit mask for extended CAN IDs.

- #define CAN_SFF_MASK 0x000007FF

  Bit mask for standard CAN IDs.

CAN ID flags

Flags within a CAN ID indicating special CAN frame attributes

- #define CAN_EFF_FLAG 0x80000000

  Extended frame.

- #define CAN_RTR_FLAG 0x40000000

  Remote transmission frame.

- #define CAN_ERR_FLAG 0x20000000

  Error frame (see Errors), not valid in struct can_filter.

- #define CAN_INV_FILTER CAN_ERR_FLAG

  Invert CAN filter definition, only valid in struct can_filter.

Particular CAN protocols

Possible protocols for the PF_CAN protocol family

Currently only the RAW protocol is supported.

- #define CAN_RAW 1

  Raw protocol of PF_CAN, applicable to socket type SOCK_RAW.

CAN controller modes

Special CAN controllers modes, which can be or'ed together.

Note

These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.

- #define CAN_CTRLMODE_LISTENONLY 0x1
- #define CAN_CTRLMODE_LOOPBACK 0x2

Timestamp switches

Arguments to pass to RTCAN_RTIOC_TAKE_TIMESTAMP

- #define RTCAN_TAKE_NO_TIMESTAMPS 0

  Switch off taking timestamps.

- #define RTCAN_TAKE_TIMESTAMPS 1

  Do take timestamps.

RAW socket options

Setting and getting CAN RAW socket options.

- #define CAN_RAW_FILTER 0x1

  CAN filter definition.

- #define CAN_RAW_ERR_FILTER 0x2

  CAN error mask.

- #define CAN_RAW_LOOPBACK 0x3

  CAN TX loopback.

- #define CAN_RAW_RECV_OWN_MSGS 0x4

  CAN receive own messages.

IOCTLs

CAN device IOCTLs

- #define SIOCGIFINDEX defined_by_kernel_header_file

  Get CAN interface index by name.

- #define SIOCSCANBAUDRATE _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)

  Set baud rate.

- #define SIOCGCANBAUDRATE _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)

  Get baud rate.

- #define SIOCSCANCUSTOMBITTIME _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)

  Set custom bit time parameter.

- #define SIOCGCANCUSTOMBITTIME _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)

  Get custum bit-time parameters.

- #define SIOCSCANMODE _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)

  Set operation mode of CAN controller.

- #define SIOCGCANSTATE _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)

  Get current state of CAN controller.

- #define SIOCSCANCTRLMODE _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)

  Set special controller modes.

- #define SIOCGCANCTRLMODE _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)

  Get special controller modes.

- #define RTCAN_RTIOC_TAKE_TIMESTAMP _IOW(RTIOC_TYPE_CAN, 0x09, int)

  Enable or disable storing a high precision timestamp upon reception of a CAN frame.

- #define RTCAN_RTIOC_RCV_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_rel_-t)

  Specify a reception timeout for a socket.

- #define RTCAN_RTIOC_SND_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_-t)

  Specify a transmission timeout for a socket.

Error mask

Error class (mask) in can_id field of struct can_frame to be used with CAN_RAW_ERR_FILTER.

Note: Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.

Note: In case of a bus-off error condition (CAN_ERR_BUSOFF), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling CAN_-MODE_START.

Note: Bus error interrupts (CAN_ERR_BUSERROR) are enabled when an application is calling a Recv function on a socket listening on bus errors (using CAN_RAW_ERR_FILTER). After one bus error has occured, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.

- #define CAN_ERR_TX_TIMEOUT 0x00000001U

  TX timeout (netdevice driver).

- #define CAN_ERR_LOSTARB 0x00000002U

  Lost arbitration (see data[0]).

- #define CAN_ERR_CRTL 0x00000004U

  Controller problems (see data[1]).

- #define CAN_ERR_PROT 0x00000008U

  Protocol violations (see data[2], data[3]).

- #define CAN_ERR_TRX 0x00000010U

  Transceiver status (see data[4]).

- #define CAN_ERR_ACK 0x00000020U

  Received no ACK on transmission.

- #define CAN_ERR_BUSOFF 0x00000040U

  Bus off.

- #define CAN_ERR_BUSERROR 0x00000080U

  Bus error (may flood!).

- #define CAN_ERR_RESTARTED 0x00000100U

  Controller restarted.

- #define CAN_ERR_MASK 0x1FFFFFFFU

  Omit EFF, RTR, ERR flags.

Arbitration lost error

Error in the data[0] field of struct can_frame.

- #define CAN_ERR_LOSTARB_UNSPEC 0x00

  unspecified

Controller problems

Error in the data[1] field of struct can_frame.

- #define CAN_ERR_CRTL_UNSPEC 0x00

    unspecified

- #define CAN_ERR_CRTL_RX_OVERFLOW 0x01

    RX buffer overflow.

- #define CAN_ERR_CRTL_TX_OVERFLOW 0x02

    TX buffer overflow.

- #define CAN_ERR_CRTL_RX_WARNING 0x04

    reached warning level for RX errors

- #define CAN_ERR_CRTL_TX_WARNING 0x08

    reached warning level for TX errors

- #define CAN_ERR_CRTL_RX_PASSIVE 0x10

    reached passive level for RX errors

- #define CAN_ERR_CRTL_TX_PASSIVE 0x20

    reached passive level for TX errors

Protocol error type

Error in the data[2] field of struct can_frame.

- #define CAN_ERR_PROT_UNSPEC 0x00

    unspecified

- #define CAN_ERR_PROT_BIT 0x01

    single bit error

- #define CAN_ERR_PROT_FORM 0x02

    frame format error

- #define CAN_ERR_PROT_STUFF 0x04

    bit stuffing error

- #define CAN_ERR_PROT_BIT0 0x08

    unable to send dominant bit

- #define CAN_ERR_PROT_BIT1 0x10

    unable to send recessive bit

- #define CAN_ERR_PROT_OVERLOAD 0x20

bus overload

- #define CAN_ERR_PROT_ACTIVE 0x40

  active error announcement

- #define CAN_ERR_PROT_TX 0x80

  error occured on transmission

Protocol error location

Error in the data[4] field of struct can_frame.

- #define CAN_ERR_PROT_LOC_UNSPEC 0x00

  unspecified

- #define CAN_ERR_PROT_LOC_SOF 0x03

  start of frame

- #define CAN_ERR_PROT_LOC_ID28_21 0x02

  ID bits 28 - 21 (SFF: 10 - 3).

- #define CAN_ERR_PROT_LOC_ID20_18 0x06

  ID bits 20 - 18 (SFF: 2 - 0).

- #define CAN_ERR_PROT_LOC_SRTR 0x04

  substitute RTR (SFF: RTR)

- #define CAN_ERR_PROT_LOC_IDE 0x05

  identifier extension

- #define CAN_ERR_PROT_LOC_ID17_13 0x07

  ID bits 17-13.

- #define CAN_ERR_PROT_LOC_ID12_05 0x0F

  ID bits 12-5.

- #define CAN_ERR_PROT_LOC_ID04_00 0x0E

  ID bits 4-0.

- #define CAN_ERR_PROT_LOC_RTR 0x0C

  RTR.

- #define CAN_ERR_PROT_LOC_RES1 0x0D

  reserved bit 1

- #define CAN_ERR_PROT_LOC_RES0 0x09

  reserved bit 0

- #define CAN_ERR_PROT_LOC_DLC 0x0B

    data length code

- #define CAN_ERR_PROT_LOC_DATA 0x0A

    data section

- #define CAN_ERR_PROT_LOC_CRC_SEQ 0x08

    CRC sequence.

- #define CAN_ERR_PROT_LOC_CRC_DEL 0x18

    CRC delimiter.

- #define CAN_ERR_PROT_LOC_ACK 0x19

    ACK slot.

- #define CAN_ERR_PROT_LOC_ACK_DEL 0x1B

    ACK delimiter.

- #define CAN_ERR_PROT_LOC_EOF 0x1A

    end of frame

- #define CAN_ERR_PROT_LOC_INTERM 0x12

    intermission

- #define CAN_ERR_TRX_UNSPEC 0x00

    0000 0000

- #define CAN_ERR_TRX_CANH_NO_WIRE 0x04

    0000 0100

- #define CAN_ERR_TRX_CANH_SHORT_TO_BAT 0x05

    0000 0101

- #define CAN_ERR_TRX_CANH_SHORT_TO_VCC 0x06

    0000 0110

- #define CAN_ERR_TRX_CANH_SHORT_TO_GND 0x07

    0000 0111

- #define CAN_ERR_TRX_CANL_NO_WIRE 0x40

    0100 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_BAT 0x50

    0101 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_VCC 0x60

    0110 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_GND 0x70

    0111 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_CANH 0x80

    1000 0000

5.1.1 Detailed Description

This is the common interface a RTDM-compliant CAN device has to provide. Feel free to report bugs and comments on this profile to the "Socketcan" mailing list (Socketcan-core@lists.berlios.de) or directly to the authors (wg@grandegger.com or Sebastian.Smolorz@stud.uni-hannover.de).

Profile Revision: 2

Device Characteristics

> Device Flags: RTDM_PROTOCOL_DEVICE
> Protocol Family: PF_CAN
> Socket Type: SOCK_RAW
> Device Class: RTDM_CLASS_CAN

Supported Operations

> **Socket**
> Environments: non-RT (RT optional, deprecated)
> Specific return values:
>
> - -EPROTONOSUPPORT (Protocol is not supported by the driver. See CAN protocols for possible protocols.)
>
> **Close**
> Blocking calls to any of the Send or Receive functions will be unblocked when the socket is closed and return with an error.
> Environments: non-RT (RT optional, deprecated)
> Specific return values: none
> **IOCTL**
> Mandatory Environments: see below
> Specific return values: see below
> **Bind**
> Binds a socket to one or all CAN devices (see struct sockaddr_can). If a filter list has been defined with setsockopt (see Sockopts), it will be used upon reception of CAN frames to decide whether the bound socket will receive a frame. If no filter has been defined, the socket will receive **all** CAN frames on the specified interface(s).
> Binding to special interface index **0** will make the socket receive CAN frames from all CAN interfaces.
> Binding to an interface index is also relevant for the Send functions because they will transmit a message over the interface the socket is bound to when no socket address is given to them.
> Environments: non-RT (RT optional)
> Specific return values:
>
> - -EFAULT (It was not possible to access user space memory area at the specified address.)
> - -ENOMEM (Not enough memory to fulfill the operation)
> - -EINVAL (Invalid address family, or invalid length of address structure)
> - -ENODEV (Invalid CAN interface index)
> - -ENOSPC (No enough space for filter list)

- -EBADF (Socket is about to be closed)
- -EAGAIN (Too many receivers. Old binding (if any) is still active. Close some sockets and try again.)

Setsockopt, Getsockopt
These functions allow to set and get various socket options. Currently, only CAN raw sockets are supported.
Supported Levels and Options:

- Level **SOL_CAN_RAW** : CAN RAW protocol (see CAN_RAW)
    - Option CAN_RAW_FILTER : CAN filter list
    - Option CAN_RAW_ERR_FILTER : CAN error mask
    - Option CAN_RAW_LOOPBACK : CAN TX loopback to local sockets

Environments: non-RT (RT optional)
Specific return values: see links to options above.

Recv, Recvfrom, Recvmsg
These functions receive CAN messages from a socket. Only one message per call can be received, so only one buffer with the correct length must be passed. For `SOCK_RAW`, this is the size of struct can_frame.
Unlike a call to one of the Send functions, a Recv function will not return with an error if an interface is down (due to bus-off or setting of stop mode) or in sleep mode. Moreover, in such a case there may still be some CAN messages in the socket buffer which could be read out successfully.
It is possible to receive a high precision timestamp with every CAN message. The condition is a former instruction to the socket via RTCAN_RTIOC_TAKE_TIMESTAMP. The timestamp will be copied to the `msg_control` buffer of `struct msghdr` if it points to a valid memory location with size of nanosecs_abs_t. If this is a NULL pointer the timestamp will be discarded silently.
Note: A `msg_controllen` of `0` upon completion of the function call indicates that no timestamp is available for that message.
Supported Flags [in]:

- MSG_DONTWAIT (By setting this flag the operation will only succeed if it would not block, i.e. if there is a message in the socket buffer. This flag takes precedence over a timeout specified by RTCAN_RTIOC_RCV_TIMEOUT.)
- MSG_PEEK (Receive a message but leave it in the socket buffer. The next receive operation will get that message again.)

Supported Flags [out]: none
Environments: RT (non-RT optional)
Specific return values:

- Non-negative value (Indicating the successful reception of a CAN message. For `SOCK_-RAW`, this is the size of struct can_frame regardless of the actual size of the payload.)
- -EFAULT (It was not possible to access user space memory area at one of the specified addresses.)
- -EINVAL (Unsupported flag detected, or invalid length of socket address buffer, or invalid length of message control buffer)
- -EMSGSIZE (Zero or more than one iovec buffer passed, or buffer too small)
- -EAGAIN (No data available in non-blocking mode)
- -EBADF (Socket was closed.)
- -EINTR (Operation was interrupted explicitly or by signal.)
- -ETIMEDOUT (Timeout)

Send, Sendto, Sendmsg

These functions send out CAN messages. Only one message per call can be transmitted, so only one buffer with the correct length must be passed. For `SOCK_RAW`, this is the size of struct can_frame.

The following only applies to `SOCK_RAW:` If a socket address of struct sockaddr_can is given, only `can_ifindex` is used. It is also possible to omit the socket address. Then the interface the socket is bound to will be used for sending messages.

If an interface goes down (due to bus-off or setting of stop mode) all senders that were blocked on this interface will be woken up.

Supported Flags:

- MSG_DONTWAIT (By setting this flag the transmit operation will only succeed if it would not block. This flag takes precedence over a timeout specified by RTCAN_-RTIOC_SND_TIMEOUT.)

Environments: RT (non-RT optional)

Specific return values:

- Non-negative value equal to given buffer size (Indicating the successful completion of the function call. See also note.)
- -EOPNOTSUPP (MSG_OOB flag is not supported.)
- -EINVAL (Unsupported flag detected *or:* Invalid length of socket address *or:* Invalid address family *or:* Data length code of CAN frame not between 0 and 15 *or:* CAN standard frame has got an ID not between 0 and 2031)
- -EMSGSIZE (Zero or more than one buffer passed or invalid size of buffer)
- -EFAULT (It was not possible to access user space memory area at one of the specified addresses.)
- -ENXIO (Invalid CAN interface index - `0` is not allowed here - or socket not bound or rather bound to all interfaces.)
- -ENETDOWN (Controller is bus-off or in stopped state.)
- -ECOMM (Controller is sleeping)
- -EAGAIN (Cannot transmit without blocking but a non-blocking call was requested.)
- -EINTR (Operation was interrupted explicitly or by signal)
- -EBADF (Socket was closed.)
- -ETIMEDOUT (Timeout)

Note: A successful completion of the function call does not implicate a successful transmission of the message.

5.1.2 Define Documentation

5.1.2.1 #define CAN_CTRLMODE_LISTENONLY 0x1

Listen-Only mode

In this mode the CAN controller would give no acknowledge to the CAN-bus, even if a message is received successfully and messages would not be transmitted. This mode might be useful for bus-monitoring, hot-plugging or throughput analysis.

5.1.2.2 #define CAN_CTRLMODE_LOOPBACK 0x2

Loopback mode

In this mode the CAN controller does an internal loop-back, a message is transmitted and simultaneously received. That mode can be used for self test operation.

5.1.2.3 #define CAN_ERR_LOSTARB_UNSPEC 0x00

unspecified

else bit number in bitstream

5.1.2.4 #define CAN_RAW_ERR_FILTER 0x2

CAN error mask.

A CAN error mask (see Errors) can be set with `setsockopt`. This mask is then used to decide if error frames are delivered to this socket in case of error condidtions. The error frames are marked with the CAN_ERR_FLAG of CAN_xxx_FLAG and must be handled by the application properly. A detailed description of the errors can be found in the `can_id` and the `data` fields of struct can_frame (see Errors for futher details).

Parameters

    [in] *level* **SOL_CAN_RAW**

    [in] *optname* **CAN_RAW_ERR_FILTER**

    [in] *optval* Pointer to error mask of type can_err_mask_t.

    [in] *optlen* Size of error mask: sizeof(can_err_mask_t).

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)

- -EINVAL (Invalid length "optlen")

Examples:

    rtcanrecv.c.

5.1.2.5 #define CAN_RAW_FILTER 0x1

CAN filter definition.

A CAN raw filter list with elements of struct can_filter can be installed with `setsockopt`. This list is used upon reception of CAN frames to decide whether the bound socket will receive a frame. An empty filter list can also be defined using optlen = 0, which is recommanded for write-only sockets.

If the socket was already bound with Bind, the old filter list gets replaced with the new one. Be aware that already received, but not read out CAN frames may stay in the socket buffer.

Parameters

    [in] *level* **SOL_CAN_RAW**

[in] *optname* **CAN_RAW_FILTER**

[in] *optval* Pointer to array of struct can_filter.

[in] *optlen* Size of filter list: count * sizeof(struct can_filter).

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -ENOMEM (Not enough memory to fulfill the operation)
- -EINVAL (Invalid length "optlen")
- -ENOSPC (No space to store filter list, check RT-Socket-CAN kernel parameters)

Examples:

rtcan_rtt.c, rtcanrecv.c, and rtcansend.c.

5.1.2.6 #define CAN_RAW_LOOPBACK 0x3

CAN TX loopback.

The TX loopback to other local sockets can be selected with this `setsockopt`.

Note

The TX loopback feature must be enabled in the kernel and then the loopback to other local TX sockets is enabled by default.

Parameters

[in] *level* **SOL_CAN_RAW**

[in] *optname* **CAN_RAW_LOOPBACK**

[in] *optval* Pointer to integer value.

[in] *optlen* Size of int: sizeof(int).

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -EINVAL (Invalid length "optlen")
- -EOPNOTSUPP (not supported, check RT-Socket-CAN kernel parameters).

Examples:

rtcansend.c.

5.1.2.7 #define CAN_RAW_RECV_OWN_MSGS 0x4

CAN receive own messages.

Not supported by RT-Socket-CAN, but defined for compatibility with Socket-CAN.

5.1.2.8 #define RTCAN_RTIOC_RCV_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_rel_t)

Specify a reception timeout for a socket.

Defines a timeout for all receive operations via a socket which will take effect when one of the receive functions is called without the `MSG_DONTWAIT` flag set.

The default value for a newly created socket is an infinite timeout.

Note

> The setting of the timeout value is not done atomically to avoid locks. Please set the value before receiving messages from the socket.

Parameters

> [in] *arg* Pointer to nanosecs_rel_t variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See Timeouts for special timeouts.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

> rtcanrecv.c.

5.1.2.9 #define RTCAN_RTIOC_SND_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_t)

Specify a transmission timeout for a socket.

Defines a timeout for all send operations via a socket which will take effect when one of the send functions is called without the `MSG_DONTWAIT` flag set.

The default value for a newly created socket is an infinite timeout.

Note

> The setting of the timeout value is not done atomically to avoid locks. Please set the value before sending messages to the socket.

Parameters

[in] *arg* Pointer to nanosecs_rel_t variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See Timeouts for special timeouts.

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

rtcansend.c.

5.1.2.10 #define RTCAN_RTIOC_TAKE_TIMESTAMP _IOW(RTIOC_TYPE_CAN, 0x09, int)

Enable or disable storing a high precision timestamp upon reception of a CAN frame.

A newly created socket takes no timestamps by default.

Parameters

[in] *arg* int variable, see Timestamp switches

Returns

0 on success.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Activating taking timestamps only has an effect on newly received CAN messages from the bus. Frames that already are in the socket buffer do not have timestamps if it was deactivated before. See Receive for more details.

Rescheduling: never.

Examples:

rtcanrecv.c.

5.1.2.11   #define SIOCGCANBAUDRATE _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)

Get baud rate.

Parameters

> [in,out] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_name` must hold a valid CAN interface name, `ifr_ifru` will be filled with an instance of can_baudrate_t.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.
> - -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.1.2.12   #define SIOCGCANCTRLMODE _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)

Get special controller modes.

Parameters

> [in] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_-name` must hold a valid CAN interface name, `ifr_ifru` must be filled with an instance of can_ctrlmode_t.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.
> - -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.1.2.13 #define SIOCGCANCUSTOMBITTIME _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)

Get custum bit-time parameters.

Parameters

> [in,out] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_name` must hold a valid CAN interface name, `ifr_ifru` will be filled with an instance of struct can_bittime.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.
> - -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.1.2.14 #define SIOCGCANSTATE _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)

Get current state of CAN controller.

States are divided into main states and additional error indicators. A CAN controller is always in exactly one main state. CAN bus errors are registered by the CAN hardware and collected by the driver. There is one error indicator (bit) per error type. If this IOCTL is triggered the error types which occured since the last call of this IOCTL are reported and thereafter the error indicators are cleared. See also CAN controller states.

Parameters

> [in,out] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_name` must hold a valid CAN interface name, `ifr_ifru` will be filled with an instance of can_mode_t.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.1.2.15 #define SIOCGIFINDEX defined_by_kernel_header_file

Get CAN interface index by name.

Parameters

> [in,out] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). If `ifr_name` holds a valid CAN interface name `ifr_ifindex` will be filled with the corresponding interface index.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

> rtcan_rtt.c, rtcanconfig.c, rtcanrecv.c, and rtcansend.c.

5.1.2.16 #define SIOCSCANBAUDRATE _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)

Set baud rate.

The baudrate must be specified in bits per second. The driver will try to calculate resonable CAN bit-timing parameters. You can use SIOCSCANCUSTOMBITTIME to set custom bit-timing.

Parameters

> [in] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_-name` must hold a valid CAN interface name, `ifr_ifru` must be filled with an instance of can_baudrate_t.

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see can_baudrate_t.
- -EDOM : Baud rate not possible.
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

Setting the baud rate is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

rtcanconfig.c.

5.1.2.17 #define SIOCSCANCTRLMODE _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)

Set special controller modes.

Various special controller modes could be or'ed together (see CAN_CTRLMODE for further information).

Parameters

[in] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_-name` must hold a valid CAN interface name, `ifr_ifru` must be filled with an instance of can_ctrlmode_t.

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see can_baudrate_t.
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

Setting special controller modes is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

rtcanconfig.c.

5.1.2.18 #define SIOCSCANCUSTOMBITTIME _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)

Set custom bit time parameter.

Custem-bit time could be defined in various formats (see struct can_bittime).

Parameters

[in] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_-name` must hold a valid CAN interface name, `ifr_ifru` must be filled with an instance of struct can_bittime.

Returns

0 on success, otherwise:
- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see can_baudrate_t.
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

Setting the bit-time is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

rtcanconfig.c.

5.1.2.19 #define SIOCSCANMODE _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)

Set operation mode of CAN controller.

See CAN controller modes for available modes.

Parameters

> [in] *arg* Pointer to interface request structure buffer (`struct ifreq` from linux/if.h). `ifr_-`
> `name` must hold a valid CAN interface name, `ifr_ifru` must be filled with an instance
> of can_mode_t.

Returns

> 0 on success, otherwise:
>
> - -EFAULT: It was not possible to access user space memory area at the specified address.
> - -ENODEV: No device with specified name exists.
> - -EAGAIN: (CAN_MODE_START, CAN_MODE_STOP) Could not successfully set
>   mode, hardware is busy. Try again.
> - -EINVAL: (CAN_MODE_START) Cannot start controller, set baud rate first.
> - -ENETDOWN: (CAN_MODE_SLEEP) Cannot go into sleep mode because controller is
>   stopped or bus off.
> - -EOPNOTSUPP: unknown mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

> Setting a CAN controller into normal operation after a bus-off can take some time (128
> occurrences of 11 consecutive recessive bits). In such a case, although this IOCTL will return
> immediately with success and SIOCGCANSTATE will report CAN_STATE_ACTIVE, bus-off
> recovery may still be in progress.
> If a controller is bus-off, setting it into stop mode will return no error but the controller
> remains bus-off.

Rescheduling: possible.

Examples:

> rtcanconfig.c.

5.1.2.20 #define SOL_CAN_RAW 103

CAN socket levels.

Used for Sockopts for the particular protocols.

Examples:

> rtcan_rtt.c, rtcanrecv.c, and rtcansend.c.

5.1.3 Typedef Documentation

5.1.3.1 typedef struct can_filter can_filter_t

Filter for reception of CAN messages.

This filter works as follows: A received CAN ID is AND'ed bitwise with `can_mask` and then compared to `can_id`. This also includes the CAN_EFF_FLAG and CAN_RTR_FLAG of CAN_-xxx_FLAG. If this comparison is true, the message will be received by the socket. The logic can be inverted with the `can_id` flag CAN_INV_FILTER :

```
if (can_id & CAN_INV_FILTER) {
   if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))
      accept-message;
} else {
   if ((received_can_id & can_mask) == can_id)
      accept-message;
}
```

Multiple filters can be arranged in a filter list and set with Sockopts. If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

5.1.3.2 typedef struct can_frame can_frame_t

Raw CAN frame.

Central structure for receiving and sending CAN frames.

5.1.4 Enumeration Type Documentation

5.1.4.1 enum CAN_BITTIME_TYPE

Supported CAN bit-time types.

Enumerator:

> *CAN_BITTIME_STD* Standard bit-time definition according to Bosch.
> *CAN_BITTIME_BTR* Hardware-specific BTR bit-time definition.

5.1.4.2 enum CAN_MODE

Enumerator:

> *CAN_MODE_STOP* Set controller in Stop mode (no reception / transmission possible)
> *CAN_MODE_START* Set controller into normal operation.
> > Coming from stopped mode or bus off, the controller begins with no errors in CAN_-STATE_ACTIVE.
> *CAN_MODE_SLEEP* Set controller into Sleep mode.
> > This is only possible if the controller is not stopped or bus-off.
> > Notice that sleep mode will only be entered when there is no bus activity. If the controller detects bus activity while "sleeping" it will go into operating mode again.
> > To actively leave sleep mode again trigger `CAN_MODE_START`.

5.1.4.3 enum CAN_STATE

Enumerator:

    CAN_STATE_ACTIVE   CAN controller is error active.

    CAN_STATE_BUS_WARNING   CAN controller is error active, warning level is reached.

    CAN_STATE_BUS_PASSIVE   CAN controller is error passive.

    CAN_STATE_BUS_OFF   CAN controller went into Bus Off.

    CAN_STATE_SCANNING_BAUDRATE   CAN controller is scanning to get the baudrate.

    CAN_STATE_STOPPED   CAN controller is in stopped mode.

    CAN_STATE_SLEEPING   CAN controller is in Sleep mode.

5.2 Real-Time Driver Model

Collaboration diagram for Real-Time Driver Model:

Modules

- User API
- Driver Development API
- Device Profiles

Typedefs

- typedef uint64_t nanosecs_abs_t

    RTDM type for representing absolute dates.

- typedef int64_t nanosecs_rel_t

    RTDM type for representing relative intervals.

API Versioning

- #define RTDM_API_VER 8

  Common user and driver API version.

- #define RTDM_API_MIN_COMPAT_VER 6

  Minimum API revision compatible with the current release.

RTDM_TIMEOUT_xxx

Special timeout values

- #define RTDM_TIMEOUT_INFINITE 0

  Block forever.

- #define RTDM_TIMEOUT_NONE (-1)

  Any negative timeout means non-blocking.

5.2.1 Detailed Description

The Real-Time Driver Model (RTDM) provides a unified interface to both users and developers of real-time device drivers. Specifically, it addresses the constraints of mixed RT/non-RT systems like Xenomai. RTDM conforms to POSIX semantics (IEEE Std 1003.1) where available and applicable.

API Revision: 8

5.2.2 Define Documentation

5.2.2.1 #define RTDM_TIMEOUT_INFINITE 0

Block forever.

5.2.2.2 #define RTDM_TIMEOUT_NONE (-1)

Any negative timeout means non-blocking.

5.2.3 Typedef Documentation

5.2.3.1 typedef uint64_t nanosecs_abs_t

RTDM type for representing absolute dates.

Its base type is a 64 bit unsigned integer. The unit is 1 nanosecond.

5.2.3.2 typedef int64_t nanosecs_rel_t

RTDM type for representing relative intervals.

Its base type is a 64 bit signed integer. The unit is 1 nanosecond. Relative intervals can also encode the special timeouts "infinite" and "non-blocking", see RTDM_TIMEOUT_xxx.

5.3 User API

Collaboration diagram for User API:

Files

- file rtdm.h

  Real-Time Driver Model for Xenomai, user API header.

Functions

- int rt_dev_open (const char ∗path, int oflag,...)

  Open a device.

- int rt_dev_socket (int protocol_family, int socket_type, int protocol)

  Create a socket.

- int rt_dev_close (int fd)

  Close a device or socket.

- int rt_dev_ioctl (int fd, int request,...)

  Issue an IOCTL.

- ssize_t rt_dev_read (int fd, void ∗buf, size_t nbyte)

  Read from device.

- ssize_t rt_dev_write (int fd, const void ∗buf, size_t nbyte)

  Write to device.

- ssize_t rt_dev_recvmsg (int fd, struct msghdr ∗msg, int flags)

Receive message from socket.

- ssize_t rt_dev_recvfrom (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)

  Receive message from socket.

- ssize_t rt_dev_recv (int fd, void *buf, size_t len, int flags)

  Receive message from socket.

- ssize_t rt_dev_sendmsg (int fd, const struct msghdr *msg, int flags)

  Transmit message to socket.

- ssize_t rt_dev_sendto (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)

  Transmit message to socket.

- ssize_t rt_dev_send (int fd, const void *buf, size_t len, int flags)

  Transmit message to socket.

- int rt_dev_bind (int fd, const struct sockaddr *my_addr, socklen_t addrlen)

  Bind to local address.

- int rt_dev_connect (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)

  Connect to remote address.

- int rt_dev_listen (int fd, int backlog)

  Listen for incomming connection requests.

- int rt_dev_accept (int fd, struct sockaddr *addr, socklen_t *addrlen)

  Accept a connection requests.

- int rt_dev_shutdown (int fd, int how)

  Shut down parts of a connection.

- int rt_dev_getsockopt (int fd, int level, int optname, void *optval, socklen_t *optlen)

  Get socket option.

- int rt_dev_setsockopt (int fd, int level, int optname, const void *optval, socklen_t optlen)

  Set socket option.

- int rt_dev_getsockname (int fd, struct sockaddr *name, socklen_t *namelen)

  Get local socket address.

- int rt_dev_getpeername (int fd, struct sockaddr *name, socklen_t *namelen)

  Get socket destination address.

5.3.1 Detailed Description

This is the upper interface of RTDM provided to application programs both in kernel and user space. Note that certain functions may not be implemented by every device. Refer to the Device Profiles for precise information.

5.3.2 Function Documentation

5.3.2.1 int rt_dev_accept (int *fd,* struct sockaddr ∗ *addr,* socklen_t ∗ *addrlen*)

Accept a connection requests.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [out] *addr* Buffer for remote address
>
> [in,out] *addrlen* Address buffer size

Returns

> 0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> accept() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.2 int rt_dev_bind (int *fd,* const struct sockaddr ∗ *my_addr,* socklen_t *addrlen*)

Bind to local address.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in] *my_addr* Address buffer
>
> [in] *addrlen* Address buffer size

Returns

> 0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> bind() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

Examples:

rtcanrecv.c, and rtcansend.c.

5.3.2.3 int rt_dev_close (int *fd*)

Close a device or socket.

Parameters

[in] *fd* File descriptor as returned by rt_dev_open() or rt_dev_socket()

Returns

0 on success, otherwise a negative error code.

Note

If the matching rt_dev_open() or rt_dev_socket() call took place in non-real-time context, rt_dev_close() must be issued within non-real-time as well. Otherwise, the call will fail.

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

close() in IEEE Std 1003.1, `http://www.opengroup.org/onlinepubs/009695399`

5.3.2.4 int rt_dev_connect (int *fd,* const struct sockaddr ∗ *serv_addr,* socklen_t *addrlen*)

Connect to remote address.

Parameters

[in] *fd* File descriptor as returned by rt_dev_socket()

[in] *serv_addr* Address buffer

[in] *addrlen* Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

connect() in IEEE Std 1003.1, `http://www.opengroup.org/onlinepubs/009695399`

5.3.2.5 int rt_dev_getpeername (int *fd,* struct sockaddr ∗ *name,* socklen_t ∗ *namelen*)

Get socket destination address.

Parameters

[in] *fd* File descriptor as returned by rt_dev_socket()

[out] *name* Address buffer

[in,out] *namelen* Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

getpeername() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.6 int rt_dev_getsockname (int *fd,* struct sockaddr ∗ *name,* socklen_t ∗ *namelen*)

Get local socket address.

Parameters

[in] *fd* File descriptor as returned by rt_dev_socket()

[out] *name* Address buffer

[in,out] *namelen* Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

getsockname() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.7 int rt_dev_getsockopt (int *fd,* int *level,* int *optname,* void ∗ *optval,* socklen_t ∗ *optlen*)

Get socket option.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in] *level* Addressed stack level
>
> [in] *optname* Option name ID
>
> [out] *optval* Value buffer
>
> [in,out] *optlen* Value buffer size

Returns

> 0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> getsockopt() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.8 int rt_dev_ioctl (int *fd,* int *request,* ...)

Issue an IOCTL.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_open() or rt_dev_socket()
>
> [in] *request* IOCTL code
>
> **...** Optional third argument, depending on IOCTL function (void * or unsigned long)

Returns

> Positiv value on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> ioctl() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.9 int rt_dev_listen (int *fd,* int *backlog*)

Listen for incomming connection requests.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()

[in] *backlog* Maximum queue length

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

lsiten() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.10 int rt_dev_open (const char ∗ *path,* int *oflag,* ...)

Open a device.

Parameters

[in] *path* Device name

[in] *oflag* Open flags

... Further parameters will be ignored.

Returns

Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

open() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.11 ssize_t rt_dev_read (int *fd,* void ∗ *buf,* size_t *nbyte*)

Read from device.

Parameters

[in] *fd* File descriptor as returned by rt_dev_open()

[out] *buf* Input buffer

[in] *nbyte* Number of bytes to read

Returns

Number of bytes read, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> `read()` in IEEE Std 1003.1, `http://www.opengroup.org/onlinepubs/009695399`

5.3.2.12 ssize_t rt_dev_recv (int *fd*, void * *buf*, size_t *len*, int *flags*)

Receive message from socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [out] *buf* Message buffer
>
> [in] *len* Message buffer size
>
> [in] *flags* Message flags

Returns

> Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> `recv()` in IEEE Std 1003.1, `http://www.opengroup.org/onlinepubs/009695399`

5.3.2.13 ssize_t rt_dev_recvfrom (int *fd*, void * *buf*, size_t *len*, int *flags*, struct sockaddr * *from*, socklen_t * *fromlen*)

Receive message from socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [out] *buf* Message buffer
>
> [in] *len* Message buffer size
>
> [in] *flags* Message flags
>
> [out] *from* Buffer for message sender address
>
> [in,out] *fromlen* Address buffer size

Returns

> Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> recvfrom() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

References rt_dev_recvfrom().

Referenced by rt_dev_recvfrom().

5.3.2.14 ssize_t rt_dev_recvmsg (int *fd,* struct msghdr ∗ *msg,* int *flags*)

Receive message from socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in,out] *msg* Message descriptor
>
> [in] *flags* Message flags

Returns

> Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> recvmsg() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.15 ssize_t rt_dev_send (int *fd,* const void ∗ *buf,* size_t *len,* int *flags*)

Transmit message to socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in] *buf* Message buffer
>
> [in] *len* Message buffer size
>
> [in] *flags* Message flags

Returns

> Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> send() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

Examples:

> rtcansend.c.

5.3.2.16 ssize_t rt_dev_sendmsg (int *fd,* const struct msghdr * *msg,* int *flags*)

Transmit message to socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in] *msg* Message descriptor
>
> [in] *flags* Message flags

Returns

> Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

> sendmsg() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.17 ssize_t rt_dev_sendto (int *fd,* const void * *buf,* size_t *len,* int *flags,* const struct sockaddr * *to,* socklen_t *tolen*)

Transmit message to socket.

Parameters

> [in] *fd* File descriptor as returned by rt_dev_socket()
>
> [in] *buf* Message buffer
>
> [in] *len* Message buffer size
>
> [in] *flags* Message flags
>
> [in] *to* Buffer for message destination address
>
> [in] *tolen* Address buffer size

Returns

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

sendto() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

Examples:

rtcansend.c.

5.3.2.18   int rt_dev_setsockopt (int *fd,* int *level,* int *optname,* const void * *optval,* socklen_t *optlen*)

Set socket option.

Parameters

[in] *fd* File descriptor as returned by rt_dev_socket()

[in] *level* Addressed stack level

[in] *optname* Option name ID

[in] *optval* Value buffer

[in] *optlen* Value buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

setsockopt() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

Examples:

rtcanrecv.c, and rtcansend.c.

5.3.2.19   int rt_dev_shutdown (int *fd,* int *how*)

Shut down parts of a connection.

Parameters

>   [in] *fd* File descriptor as returned by rt_dev_socket()

>   [in] *how* Specifies the part to be shut down (SHUT_xxx)

Returns

>   0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

>   shutdown() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.20   int rt_dev_socket (int *protocol_family,* int *socket_type,* int *protocol*)

Create a socket.

Parameters

>   [in] *protocol_family* Protocol family (PF_xxx)

>   [in] *socket_type* Socket type (SOCK_xxx)

>   [in] *protocol* Protocol ID, 0 for default

Returns

>   Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

>   socket() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.3.2.21   ssize_t rt_dev_write (int *fd,* const void ∗ *buf,* size_t *nbyte*)

Write to device.

Parameters

>   [in] *fd* File descriptor as returned by rt_dev_open()

>   [in] *buf* Output buffer

>   [in] *nbyte* Number of bytes to write

Returns

Number of bytes written, otherwise negative error code

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

See also

`write()` in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.4 Real-time IPC protocols

Profile Revision: 1

Collaboration diagram for Real-time IPC protocols:

Data Structures

- struct rtipc_port_label

  Port label information structure.

- struct sockaddr_ipc

  Socket address structure for the RTIPC address family.

Files

- file rtipc.h

  This file is part of the Xenomai project.

Typedefs

- typedef int16_t rtipc_port_t

  Port number type for the RTIPC address family.

RTIPC protocol list

protocols for the PF_RTIPC protocol family

- enum { IPCPROTO_IPC = 0, IPCPROTO_XDDP = 1, IPCPROTO_IDDP = 2, IPCPROTO_-BUFP = 3 }

Supported operations

Standard socket operations supported by the RTIPC protocols.

- int socket__AF_RTIPC (int domain=AF_RTIPC, int type=SOCK_DGRAM, int protocol)

  Create an endpoint for communication in the AF_RTIPC domain.

- int close__AF_RTIPC (int sockfd)

  Close a RTIPC socket descriptor.

- int bind__AF_RTIPC (int sockfd, const struct sockaddr_ipc *addr, socklen_t addrlen)

  Bind a RTIPC socket to a port.

- int connect__AF_RTIPC (int sockfd, const struct sockaddr_ipc *addr, socklen_t addrlen)

  Initiate a connection on a RTIPC socket.

- int setsockopt__AF_RTIPC (int sockfd, int level, int optname, const void *optval, socklen_t optlen)

  Set options on RTIPC sockets.

- int getsockopt__AF_RTIPC (int sockfd, int level, int optname, void *optval, socklen_t *optlen)

  Get options on RTIPC sockets.

- ssize_t sendmsg__AF_RTIPC (int sockfd, const struct msghdr *msg, int flags)

  Send a message on a RTIPC socket.

- ssize_t recvmsg__AF_RTIPC (int sockfd, struct msghdr *msg, int flags)

  Receive a message from a RTIPC socket.

- int getsockname__AF_RTIPC (int sockfd, struct sockaddr_ipc *addr, socklen_t *addrlen)

  Get socket name.

- int getpeername__AF_RTIPC (int sockfd, struct sockaddr_ipc *addr, socklen_t *addrlen)

  Get socket peer.

XDDP socket options

Setting and getting XDDP socket options.

- #define XDDP_LABEL 1

XDDP label assignment.

- #define XDDP_POOLSZ 2

  XDDP local pool size configuration.

- #define XDDP_BUFSZ 3

  XDDP streaming buffer size configuration.

- #define XDDP_MONITOR 4

  XDDP monitoring callback.

XDDP events

Specific events occurring on XDDP channels, which can be monitored via the XDDP_MONITOR socket option.

- #define XDDP_EVTIN 1

  Monitor writes to the non real-time endpoint.

- #define XDDP_EVTOUT 2

  Monitor reads from the non real-time endpoint.

- #define XDDP_EVTDOWN 3

  Monitor close from the non real-time endpoint.

- #define XDDP_EVTNOBUF 4

  Monitor memory shortage for non real-time datagrams.

IDDP socket options

Setting and getting IDDP socket options.

- #define IDDP_LABEL 1

  IDDP label assignment.

- #define IDDP_POOLSZ 2

  IDDP local pool size configuration.

BUFP socket options

Setting and getting BUFP socket options.

- #define BUFP_LABEL 1

  BUFP label assignment.

- #define BUFP_BUFSZ 2

  BUFP buffer size configuration.

Socket level options

Setting and getting supported standard socket level options.

- #define SO_SNDTIMEO defined_by_kernel_header_file

    IPCPROTO_IDDP and IPCPROTO_BUFP protocols support the standard SO_SNDTIMEO socket option, from the SOL_SOCKET *level.*

- #define SO_RCVTIMEO defined_by_kernel_header_file

    All RTIPC protocols support the standard SO_RCVTIMEO socket option, from the SOL_SOCKET *level.*

5.4.1   Detailed Description

Profile Revision: 1

Device Characteristics

    Device Flags: RTDM_PROTOCOL_DEVICE
    Protocol Family: PF_RTIPC
    Socket Type: SOCK_DGRAM
    Device Class: RTDM_CLASS_RTIPC

5.4.2   Define Documentation

5.4.2.1   #define BUFP_BUFSZ 2

BUFP buffer size configuration.

All messages written to a BUFP socket are buffered in a single per-socket memory area. Configuring the size of such buffer prior to binding the socket to a destination port is mandatory.

It is not allowed to configure a buffer size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

    : the buffer memory is obtained from the host allocator by the bind call.

Parameters

    [in] *level* SOL_BUFP

    [in] *optname* **BUFP_BUFSZ**

    [in] *optval* Pointer to a variable of type size_t, containing the required size of the buffer to reserve at binding time

    [in] *optlen* sizeof(size_t)

Returns

    0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* is invalid or *∗optval* is zero)

Calling context:

RT/non-RT

Examples:

bufp-label.c, and bufp-readwrite.c.

5.4.2.2 #define BUFP_LABEL 1

BUFP label assignment.

ASCII label strings can be attached to BUFP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see BUFP port binding).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

[in] *level* SOL_BUFP

[in] *optname* **BUFP_LABEL**

[in] *optval* Pointer to struct rtipc_port_label

[in] *optlen* sizeof(struct rtipc_port_label)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT

Examples:

bufp-label.c.

5.4.2.3 #define IDDP_LABEL 1

IDDP label assignment.

ASCII label strings can be attached to IDDP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see IDDP port binding).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

    [in] *level* SOL_IDDP

    [in] *optname* **IDDP_LABEL**

    [in] *optval* Pointer to struct rtipc_port_label

    [in] *optlen* sizeof(struct rtipc_port_label)

Returns

    0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* is invalid)

Calling context:

    RT/non-RT

Examples:

    iddp-label.c.

5.4.2.4 #define IDDP_POOLSZ 2

IDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

    : the pool memory is obtained from the host allocator by the bind call.

Parameters

    [in] *level* SOL_IDDP

[in] *optname* **IDDP_POOLSZ**

[in] *optval* Pointer to a variable of type size_t, containing the required size of the local pool to reserve at binding time

[in] *optlen* sizeof(size_t)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* is invalid or ∗*optval* is zero)

Calling context:

RT/non-RT

Examples:

iddp-sendrecv.c.

5.4.2.5 #define SO_RCVTIMEO defined_by_kernel_header_file

All RTIPC protocols support the standard SO_RCVTIMEO socket option, from the `SOL_SOCKET` level.

See also

setsockopt(), getsockopt() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399/

Examples:

xddp-label.c.

5.4.2.6 #define SO_SNDTIMEO defined_by_kernel_header_file

IPCPROTO_IDDP and IPCPROTO_BUFP protocols support the standard SO_SNDTIMEO socket option, from the `SOL_SOCKET` level.

See also

setsockopt(), getsockopt() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399/

5.4.2.7 #define XDDP_BUFSZ 3

XDDP streaming buffer size configuration.

In addition to sending datagrams, real-time threads may stream data in a byte-oriented mode through the port as well. This increases the bandwidth and reduces the overhead, when the

overall data to send to the Linux domain is collected by bits, and keeping the message boundaries is not required.

This feature is enabled when a non-zero buffer size is set for the socket. In that case, the real-time data accumulates into the streaming buffer when MSG_MORE is passed to any of the send functions, until:

- the receiver from the Linux domain wakes up and consumes it,

- a different source port attempts to send data to the same destination port,

- MSG_MORE is absent from the send flags,

- the buffer is full,

whichever comes first.

Setting *optval to zero disables the streaming buffer, in which case all sendings are conveyed in separate datagrams, regardless of MSG_MORE.

Note

> only a single streaming buffer exists per socket. When this buffer is full, the real-time data stops accumulating and sending operations resume in mere datagram mode. Accumulation may happen again after some or all data in the streaming buffer is consumed from the Linux domain endpoint.

The streaming buffer size may be adjusted multiple times during the socket lifetime; the latest configuration change will take effect when the accumulation resumes after the previous buffer was flushed.

Parameters

> [in] *level* SOL_XDDP
>
> [in] *optname* **XDDP_BUFSZ**
>
> [in] *optval* Pointer to a variable of type size_t, containing the required size of the streaming buffer
>
> [in] *optlen* sizeof(size_t)

Returns

> 0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -ENOMEM (Not enough memory)

- -EINVAL (*optlen* is invalid)

Calling context:

> RT/non-RT

Examples:

> xddp-stream.c.

5.4.2.8 #define XDDP_EVTDOWN 3

Monitor close from the non real-time endpoint.

XDDP_EVTDOWN is sent when the non real-time endpoint is closed. The argument is always 0.

5.4.2.9 #define XDDP_EVTIN 1

Monitor writes to the non real-time endpoint.

XDDP_EVTIN is sent when data is written to the non real-time endpoint the socket is bound to (i.e. via /dev/rtp*N*), which means that some input is pending for the real-time endpoint. The argument is the size of the incoming message.

5.4.2.10 #define XDDP_EVTNOBUF 4

Monitor memory shortage for non real-time datagrams.

XDDP_EVTNOBUF is sent when no memory is available from the pool to hold the message currently sent from the non real-time endpoint. The argument is the size of the failed allocation. Upon return from the callback, the caller will block and retry until enough space is available from the pool; during that process, the callback might be invoked multiple times, each time a new attempt to get the required memory fails.

5.4.2.11 #define XDDP_EVTOUT 2

Monitor reads from the non real-time endpoint.

XDDP_EVTOUT is sent when the non real-time endpoint successfully reads a complete message (i.e. via /dev/rtp*N*). The argument is the size of the outgoing message.

5.4.2.12 #define XDDP_LABEL 1

XDDP label assignment.

ASCII label strings can be attached to XDDP ports, so that opening the non-RT endpoint can be done by specifying this symbolic device name rather than referring to a raw pseudo-device entry (i.e. /dev/rtp*N*).

When available, this label will be registered when binding, in addition to the port number (see XDDP port binding).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

> [in] *level* SOL_XDDP
>
> [in] *optname* **XDDP_LABEL**
>
> [in] *optval* Pointer to struct rtipc_port_label
>
> [in] *optlen* sizeof(struct rtipc_port_label)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* invalid)

Calling context:

RT/non-RT

Examples:

xddp-label.c.

5.4.2.13 #define XDDP_MONITOR 4

XDDP monitoring callback.

Other RTDM drivers may install a user-defined callback via the rtdm_setsockopt call from the inter-driver API, in order to collect particular events occurring on the channel.

This notification mechanism is particularly useful to monitor a channel asynchronously while performing other tasks.

The user-provided routine will be passed the RTDM file descriptor of the socket receiving the event, the event code, and an optional argument. Four events are currently defined, see XDDP_-EVENTS.

The XDDP_EVTIN and XDDP_EVTOUT events are fired on behalf of a fully atomic context; therefore, care must be taken to keep their overhead low. In those cases, the Xenomai services that may be called from the callback are restricted to the set allowed to a real-time interrupt handler.

Parameters

[in] *level* SOL_XDDP

[in] *optname* **XDDP_MONITOR**

[in] *optval* Pointer to a pointer to function of type int (∗)(int fd, int event, long arg), containing the address of the user-defined callback.Passing a NULL callback pointer in *optval* disables monitoring.

[in] *optlen* sizeof(int (∗)(int fd, int event, long arg))

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EPERM (Operation not allowed from user-space)

- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT, kernel space only

5.4.2.14 #define XDDP_POOLSZ 2

XDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

: the pool memory is obtained from the host allocator by the bind call.

Parameters

[in] *level* SOL_XDDP

[in] *optname* **XDDP_POOLSZ**

[in] *optval* Pointer to a variable of type size_t, containing the required size of the local pool to reserve at binding time

[in] *optlen* sizeof(size_t)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)

- -EALREADY (socket already bound)

- -EINVAL (*optlen* invalid or ∗*optval* is zero)

Calling context:

RT/non-RT

Examples:

xddp-echo.c.

5.4.3 Enumeration Type Documentation

5.4.3.1 anonymous enum

Enumerator:

IPCPROTO_IPC Default protocol (IDDP).

IPCPROTO_XDDP Cross-domain datagram protocol (RT <-> non-RT). Real-time Xenomai threads and regular Linux threads may want to exchange data in a way that does not require the former to leave the real-time domain (i.e. secondary mode). The RTDM-based XDDP protocol is available for this purpose.

On the Linux domain side, pseudo-device files named /dev/rtp<*minor*> give regular POSIX threads access to non real-time communication endpoints, via the standard

character-based I/O interface. On the Xenomai domain side, sockets may be bound to XDDP ports, which act as proxies to send and receive data to/from the associated pseudo-device files. Ports and pseudo-device minor numbers are paired, meaning that e.g. port 7 will proxy the traffic to/from /dev/rtp7.

All data sent through a bound/connected XDDP socket via `sendto(2)` or `write(2)` will be passed to the peer endpoint in the Linux domain, and made available for reading via the standard `read(2)` system call. Conversely, all data sent using `write(2)` through the non real-time endpoint will be conveyed to the real-time socket endpoint, and made available to the `recvfrom(2)` or `read(2)` system calls.

IPCPROTO_IDDP Intra-domain datagram protocol (RT <-> RT). The RTDM-based IDDP protocol enables real-time threads to exchange datagrams within the Xenomai domain, via socket endpoints.

IPCPROTO_BUFP Buffer protocol (RT <-> RT, byte-oriented). The RTDM-based BUFP protocol implements a lightweight, byte-oriented, one-way Producer-Consumer data path. All messages written are buffered into a single memory area in strict FIFO order, until read by the consumer.

This protocol always prevents short writes, and only allows short reads when a potential deadlock situation arises (i.e. readers and writers waiting for each other indefinitely).

5.4.4 Function Documentation

5.4.4.1 int bind__AF_RTIPC (int *sockfd,* const struct sockaddr_ipc * *addr,* socklen_t *addrlen*)

Bind a RTIPC socket to a port.

Bind the socket to a destination port.

Parameters

> [in] *addr* The address to bind the socket to (see struct sockaddr_ipc). The meaning of such address depends on the RTIPC protocol in use for the socket:

- IPCPROTO_XDDP

This action creates an endpoint for channelling traffic between the Xenomai and Linux domains.

sipc_family must be AF_RTIPC, *sipc_port* is either -1, or a valid free port number between 0 and CONFIG_XENO_OPT_PIPE_NRDEV-1.

If *sipc_port* is -1, a free port will be assigned automatically.

Upon success, the pseudo-device /dev/rtp*N* will be reserved for this communication channel, where *N* is the assigned port number. The non real-time side shall open this device to exchange data over the bound socket.

If a label was assigned (see XDDP_LABEL) prior to binding the socket to a port, a registry link referring to the created pseudo-device will be automatically set up as `/proc/xenomai/registry/rtipc/xddp/label,` where *label* is the label string passed to setsockopt() for the XDDP_LABEL option.

- IPCPROTO_IDDP

This action creates an endpoint for exchanging datagrams within the Xenomai domain.

sipc_family must be AF_RTIPC, *sipc_port* is either -1, or a valid free port number between 0 and CONFIG_XENO_OPT_IDDP_NRPORT-1.

If *sipc_port* is -1, a free port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see IDDP_LABEL) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as /proc/xenomai/registry/rtipc/iddp/label, where *label* is the label string passed to setsockopt() for the IDDP_LABEL option.

- IPCPROTO_BUFP

This action creates an endpoint for a one-way byte stream within the Xenomai domain.

sipc_family must be AF_RTIPC, *sipc_port* is either -1, or a valid free port number between 0 and CONFIG_XENO_OPT_BUFP_NRPORT-1.

If *sipc_port* is -1, an available port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see BUFP_LABEL) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as /proc/xenomai/registry/rtipc/bufp/label, where *label* is the label string passed to setsockopt() for the *BUFP_LABEL* option.

Returns

> In addition to the standard error codes for `bind(2)`, the following specific error code may be returned:
>
> - -EFAULT (Invalid data address given)
> - -ENOMEM (Not enough memory)
> - -EINVAL (Invalid parameter)
> - -EADDRINUSE (Socket already bound to a port, or no port available)

Calling context:

> non-RT

5.4.4.2 int close__AF_RTIPC (int *sockfd*)

Close a RTIPC socket descriptor.

Blocking calls to any of the sendmsg or recvmsg functions will be unblocked when the socket is closed and return with an error.

Returns

> In addition to the standard error codes for `close(2)`, the following specific error code may be returned: none

Calling context:

> non-RT

5.4.4.3 int connect__AF_RTIPC (int *sockfd,* const struct sockaddr_ipc * *addr,* socklen_t *addrlen*)

Initiate a connection on a RTIPC socket.

Parameters

[in] *addr* The address to connect the socket to (see struct sockaddr_ipc).

- If sipc_port is a valid port for the protocol, it is used verbatim and the connection succeeds immediately, regardless of whether the destination is bound at the time of the call.

- If sipc_port is -1 and a label was assigned to the socket, connect() blocks for the requested amount of time (see SO_RCVTIMEO) until a socket is bound to the same label via bind(2) (see XDDP_LABEL, IDDP_LABEL, BUFP_LABEL), in which case a connection is established between both endpoints.

- If sipc_port is -1 and no label was assigned to the socket, the default destination address is cleared, meaning that any subsequent write to the socket will return -EDESTADDRREQ, until a valid destination address is set via connect(2) or bind(2).

Returns

In addition to the standard error codes for connect(2), the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.4.4.4 int getpeername__AF_RTIPC (int *sockfd,* struct sockaddr_ipc * *addr,* socklen_t * *addrlen*)

Get socket peer.

The name of the remote endpoint for the socket is copied back (see struct sockaddr_ipc). This is the default destination address for messages sent on the socket. It can be set either explicitly via connect(2), or implicitly via bind(2) if no connect(2) was called prior to binding the socket to a port, in which case both the local and remote names are equal.

Returns

In addition to the standard error codes for getpeername(2), the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.4.4.5 int getsockname__AF_RTIPC (int *sockfd,* struct sockaddr_ipc * *addr,* socklen_t * *addrlen*)

Get socket name.

The name of the local endpoint for the socket is copied back (see struct sockaddr_ipc).

Returns

In addition to the standard error codes for `getsockname(2)`, the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.4.4.6 int getsockopt__AF_RTIPC (int *sockfd,* int *level,* int *optname,* void ∗ *optval,* socklen_t ∗ *optlen*)

Get options on RTIPC sockets.

These functions allow to get various socket options. Supported Levels and Options:

- Level SOL_SOCKET

- Level SOL_XDDP

- Level SOL_IDDP

- Level SOL_BUFP

Returns

In addition to the standard error codes for `getsockopt(2)`, the following specific error code may be returned: follow the option links above.

Calling context:

RT/non-RT

5.4.4.7 ssize_t recvmsg__AF_RTIPC (int *sockfd,* struct msghdr ∗ *msg,* int *flags*)

Receive a message from a RTIPC socket.

Parameters

[`in`] *flags* Operation flags:

- MSG_DONTWAIT Non-blocking I/O operation. The caller will not be blocked whenever no message is immediately available for receipt at the time of the call, but will rather return with -EWOULDBLOCK.

Note

IPCPROTO_BUFP does not allow for short reads and always returns the requested amount of bytes, except in one situation: whenever some writer is waiting for sending data upon a buffer full condition, while the caller would have to wait for receiving a complete message. This is usually the sign of a pathological use of the BUFP socket, like defining an incorrect buffer size via BUFP_BUFSZ. In that case, a short read is allowed to prevent a deadlock.

Returns

In addition to the standard error codes for `recvmsg(2)`, the following specific error code may be returned: none.

Calling context:

RT

5.4.4.8 ssize_t sendmsg__AF_RTIPC (int *sockfd,* const struct msghdr ∗ *msg,* int *flags*)

Send a message on a RTIPC socket.

Parameters

[in] *flags* Operation flags:

- MSG_OOB Send out-of-band message. For all RTIPC protocols except IPCPROTO_BUFP, sending out-of-band data actually means pushing them to the head of the receiving queue, so that the reader will always receive them before normal messages. IPCPROTO_BUFP does not support out-of-band sending.

- MSG_DONTWAIT Non-blocking I/O operation. The caller will not be blocked whenever the message cannot be sent immediately at the time of the call (e.g. memory shortage), but will rather return with -EWOULDBLOCK. Unlike other RTIPC protocols, IPCPROTO_-XDDP accepts but never considers MSG_DONTWAIT since writing to a real-time XDDP endpoint is inherently a non-blocking operation.

- MSG_MORE Accumulate data before sending. This flag is accepted by the IPCPROTO_-XDDP protocol only, and tells the send service to accumulate the outgoing data into an internal streaming buffer, instead of issuing a datagram immediately for it. See XDDP_-BUFSZ for more.

Note

No RTIPC protocol allows for short writes, and only complete messages are sent to the peer.

Returns

In addition to the standard error codes for `sendmsg(2)`, the following specific error code may be returned: none.

Calling context:

RT

5.4.4.9 int setsockopt__AF_RTIPC (int *sockfd,* int *level,* int *optname,* const void ∗ *optval,* socklen_t *optlen*)

Set options on RTIPC sockets.

These functions allow to set various socket options. Supported Levels and Options:

- Level SOL_SOCKET

- Level SOL_XDDP

- Level SOL_IDDP

- Level SOL_BUFP

Returns

In addition to the standard error codes for `setsockopt(2)`, the following specific error code may be returned: follow the option links above.

Calling context:

non-RT

5.4.4.10   int socket__AF_RTIPC (int *domain* = *AF_RTIPC,* int *type* = *SOCK_DGRAM,* int *protocol*)

Create an endpoint for communication in the AF_RTIPC domain.

Parameters

[in] *protocol* Any of IPCPROTO_XDDP, IPCPROTO_IDDP, or IPCPROTO_BUFP. IPCPROTO_IPC is also valid, and refers to the default RTIPC protocol, namely IPCPROTO_IDDP.

Returns

In addition to the standard error codes for `socket(2)`, the following specific error code may be returned:

- -ENOPROTOOPT (Protocol is known, but not compiled in the RTIPC driver). See RTIPC protocols for available protocols.

Calling context:

non-RT

5.5   Serial Devices

Collaboration diagram for Serial Devices:

Data Structures

- struct rtser_config

    Serial device configuration.

- struct rtser_status

    Serial device status.

- struct rtser_event

    Additional information about serial device events.

Files

- file rtserial.h

    Real-Time Driver Model for Xenomai, serial device profile header.

Defines

- #define RTSER_RTIOC_BREAK_CTL _IOR(RTIOC_TYPE_SERIAL, 0x06, int)

    Set or clear break on UART output line.

RTSER_BREAK_xxx

Break control

- typedef struct rtser_config rtser_config_t

    Serial device configuration.

- typedef struct rtser_status rtser_status_t

    Serial device status.

- typedef struct rtser_event rtser_event_t

    Additional information about serial device events.

- #define RTSER_BREAK_CLR 0x00

    Serial device configuration.

- #define RTSER_BREAK_SET 0x01

    Serial device configuration.

- #define RTIOC_TYPE_SERIAL RTDM_CLASS_SERIAL

    Serial device configuration.

RTSER_DEF_BAUD

Default baud rate

- #define **RTSER_DEF_BAUD** 9600

RTSER_xxx_PARITY

Number of parity bits

- #define **RTSER_NO_PARITY** 0x00
- #define **RTSER_ODD_PARITY** 0x01
- #define **RTSER_EVEN_PARITY** 0x03
- #define **RTSER_DEF_PARITY** RTSER_NO_PARITY

RTSER_xxx_BITS

Number of data bits

- #define **RTSER_5_BITS** 0x00
- #define **RTSER_6_BITS** 0x01
- #define **RTSER_7_BITS** 0x02
- #define **RTSER_8_BITS** 0x03
- #define **RTSER_DEF_BITS** RTSER_8_BITS

RTSER_xxx_STOPB

Number of stop bits

- #define RTSER_1_STOPB 0x00

    valid only in combination with 5 data bits

- #define RTSER_1_5_STOPB 0x01

    valid only in combination with 5 data bits

- #define RTSER_2_STOPB 0x01

    valid only in combination with 5 data bits

- #define RTSER_DEF_STOPB RTSER_1_STOPB

    valid only in combination with 5 data bits

RTSER_xxx_HAND

Handshake mechanisms

- #define **RTSER_NO_HAND** 0x00
- #define **RTSER_RTSCTS_HAND** 0x01
- #define **RTSER_DEF_HAND** RTSER_NO_HAND

RTSER_FIFO_xxx

Reception FIFO interrupt threshold

- #define **RTSER_FIFO_DEPTH_1** 0x00
- #define **RTSER_FIFO_DEPTH_4** 0x40
- #define **RTSER_FIFO_DEPTH_8** 0x80
- #define **RTSER_FIFO_DEPTH_14** 0xC0
- #define **RTSER_DEF_FIFO_DEPTH** RTSER_FIFO_DEPTH_1

RTSER_TIMEOUT_xxx

Special timeout values, see also RTDM_TIMEOUT_xxx

- #define **RTSER_TIMEOUT_INFINITE** RTDM_TIMEOUT_INFINITE
- #define **RTSER_TIMEOUT_NONE** RTDM_TIMEOUT_NONE
- #define **RTSER_DEF_TIMEOUT** RTDM_TIMEOUT_INFINITE

RTSER_xxx_TIMESTAMP_HISTORY

Timestamp history control

- #define **RTSER_RX_TIMESTAMP_HISTORY** 0x01
- #define **RTSER_DEF_TIMESTAMP_HISTORY** 0x00

RTSER_EVENT_xxx

Events bits

- #define **RTSER_EVENT_RXPEND** 0x01
- #define **RTSER_EVENT_ERRPEND** 0x02
- #define **RTSER_EVENT_MODEMHI** 0x04
- #define **RTSER_EVENT_MODEMLO** 0x08
- #define **RTSER_DEF_EVENT_MASK** 0x00

RTSER_SET_xxx

Configuration mask bits

- #define **RTSER_SET_BAUD** 0x0001
- #define **RTSER_SET_PARITY** 0x0002
- #define **RTSER_SET_DATA_BITS** 0x0004
- #define **RTSER_SET_STOP_BITS** 0x0008
- #define **RTSER_SET_HANDSHAKE** 0x0010
- #define **RTSER_SET_FIFO_DEPTH** 0x0020
- #define **RTSER_SET_TIMEOUT_RX** 0x0100
- #define **RTSER_SET_TIMEOUT_TX** 0x0200
- #define **RTSER_SET_TIMEOUT_EVENT** 0x0400
- #define **RTSER_SET_TIMESTAMP_HISTORY** 0x0800
- #define **RTSER_SET_EVENT_MASK** 0x1000

RTSER_LSR_xxx

Line status bits

- #define **RTSER_LSR_DATA** 0x01
- #define **RTSER_LSR_OVERRUN_ERR** 0x02
- #define **RTSER_LSR_PARITY_ERR** 0x04
- #define **RTSER_LSR_FRAMING_ERR** 0x08
- #define **RTSER_LSR_BREAK_IND** 0x10
- #define **RTSER_LSR_THR_EMTPY** 0x20
- #define **RTSER_LSR_TRANSM_EMPTY** 0x40
- #define **RTSER_LSR_FIFO_ERR** 0x80
- #define **RTSER_SOFT_OVERRUN_ERR** 0x0100

RTSER_MSR_xxx

Modem status bits

- #define **RTSER_MSR_DCTS** 0x01
- #define **RTSER_MSR_DDSR** 0x02
- #define **RTSER_MSR_TERI** 0x04
- #define **RTSER_MSR_DDCD** 0x08
- #define **RTSER_MSR_CTS** 0x10
- #define **RTSER_MSR_DSR** 0x20
- #define **RTSER_MSR_RI** 0x40
- #define **RTSER_MSR_DCD** 0x80

RTSER_MCR_xxx

Modem control bits

- #define **RTSER_MCR_DTR** 0x01
- #define **RTSER_MCR_RTS** 0x02
- #define **RTSER_MCR_OUT1** 0x04
- #define **RTSER_MCR_OUT2** 0x08
- #define **RTSER_MCR_LOOP** 0x10

Sub-Classes of RTDM_CLASS_SERIAL

- #define **RTDM_SUBCLASS_16550A** 0

IOCTLs

Serial device IOCTLs

- #define RTSER_RTIOC_GET_CONFIG _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_-config)

    Get serial device configuration.

- #define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_-config)

  Set serial device configuration.

- #define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_-status)

  Get serial device status.

- #define RTSER_RTIOC_GET_CONTROL _IOR(RTIOC_TYPE_SERIAL, 0x03, int)

  Get serial device's modem contol register.

- #define RTSER_RTIOC_SET_CONTROL _IOW(RTIOC_TYPE_SERIAL, 0x04, int)

  Set serial device's modem contol register.

- #define RTSER_RTIOC_WAIT_EVENT _IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_-event)

  Wait on serial device events according to previously set mask.

5.5.1 Detailed Description

This is the common interface a RTDM-compliant serial device has to provide. Feel free to comment on this profile via the Xenomai mailing list (`Xenomai-core@gna.org`) or directly to the author (`jan.kiszka@web.de`).

Profile Revision: 3

Device Characteristics

Device Flags: `RTDM_NAMED_DEVICE`, `RTDM_EXCLUSIVE`
Device Name: `"rtser<N>"`, N >= 0
Device Class: `RTDM_CLASS_SERIAL`

Supported Operations

Open
Environments: non-RT (RT optional, deprecated)
Specific return values: none
Close
Environments: non-RT (RT optional, deprecated)
Specific return values: none
IOCTL
Mandatory Environments: see below
Specific return values: see below
Read
Environments: RT (non-RT optional)
Specific return values:

- -ETIMEDOUT

- -EINTR (interrupted explicitly or by signal)

- -EAGAIN (no data available in non-blocking mode)

- -EBADF (device has been closed while reading)

- -EIO (hardware error or broken bit stream)

Write

Environments: RT (non-RT optional)

Specific return values:

- -ETIMEDOUT

- -EINTR (interrupted explicitly or by signal)

- -EAGAIN (no data written in non-blocking mode)

- -EBADF (device has been closed while writing)

5.5.2 Define Documentation

5.5.2.1 #define RTSER_RTIOC_BREAK_CTL _IOR(RTIOC_TYPE_SERIAL, 0x06, int)

Set or clear break on UART output line.

Parameters

    [in] *arg* RTSER_BREAK_SET or RTSER_BREAK_CLR (int)

Returns

    0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

    A set break condition may also be cleared on UART line reconfiguration.

Rescheduling: never.

5.5.2.2 #define RTSER_RTIOC_GET_CONFIG _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_config)

Get serial device configuration.

Parameters

    [out] *arg* Pointer to configuration buffer (struct rtser_config)

Returns

    0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.5.2.3 #define RTSER_RTIOC_GET_CONTROL _IOR(RTIOC_TYPE_SERIAL, 0x03, int)

Get serial device's modem contol register.

Parameters

> [out] *arg* Pointer to variable receiving the content (int, see RTSER_MCR_xxx)

Returns

> 0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.5.2.4 #define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_status)

Get serial device status.

Parameters

> [out] *arg* Pointer to status buffer (struct rtser_status)

Returns

> 0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

> The error states `RTSER_LSR_OVERRUN_ERR`, `RTSER_LSR_PARITY_ERR`, `RTSER_LSR_FRAMING_ERR`, and `RTSER_SOFT_OVERRUN_ERR` that may have occured during previous read accesses to the device will be saved for being reported via this IOCTL. Upon return from `RTSER_RTIOC_-GET_STATUS`, the saved state will be cleared.

Rescheduling: never.

5.5.2.5 #define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_config)

Set serial device configuration.

Parameters

> [in] *arg* Pointer to configuration buffer (struct rtser_config)

Returns

> 0 on success, otherwise:

- -EPERM is returned if the caller's context is invalid, see note below.

- -ENOMEM is returned if a new history buffer for timestamps cannot be allocated.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Note

> If rtser_config contains a valid timestamp_history and the addressed device has been opened in non-real-time context, this IOCTL must be issued in non-real-time context as well. Otherwise, this command will fail.

Rescheduling: never.

Examples:

> cross-link.c.

5.5.2.6 #define RTSER_RTIOC_SET_CONTROL _IOW(RTIOC_TYPE_SERIAL, 0x04, int)

Set serial device's modem contol register.

Parameters

> [in] *arg* New control register content (int, see RTSER_MCR_xxx)

Returns

> 0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.5.2.7 #define RTSER_RTIOC_WAIT_EVENT _IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_event)

Wait on serial device events according to previously set mask.

Parameters

> [out] *arg* Pointer to event information buffer (struct rtser_event)

Returns

> 0 on success, otherwise:

- -EBUSY is returned if another task is already waiting on events of this device.

- -EBADF is returned if the file descriptor is invalid or the device has just been closed.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

Examples:

> cross-link.c.

5.6 Testing Devices

Collaboration diagram for Testing Devices:

Files

- file rttesting.h

  Real-Time Driver Model for Xenomai, testing device profile header.

Sub-Classes of RTDM_CLASS_TESTING

- #define RTDM_SUBCLASS_TIMERBENCH 0

  subclass name: "timerbench"

- #define RTDM_SUBCLASS_IRQBENCH 1

  subclass name: "irqbench"

- #define RTDM_SUBCLASS_SWITCHTEST 2

  subclass name: "switchtest"

- #define RTDM_SUBCLASS_RTDMTEST 3

  subclase name: "rtdm"

IOCTLs

Testing device IOCTLs

- #define **RTTST_RTIOC_INTERM_BENCH_RES** _IOWR(RTIOC_TYPE_TESTING, 0x00, struct rttst_interm_bench_res)
- #define **RTTST_RTIOC_TMBENCH_START** _IOW(RTIOC_TYPE_TESTING, 0x10, struct rttst_tmbench_config)
- #define **RTTST_RTIOC_TMBENCH_STOP** _IOWR(RTIOC_TYPE_TESTING, 0x11, struct rttst_overall_bench_res)
- #define **RTTST_RTIOC_IRQBENCH_START** _IOW(RTIOC_TYPE_TESTING, 0x20, struct rttst_irqbench_config)
- #define **RTTST_RTIOC_IRQBENCH_STOP** _IO(RTIOC_TYPE_TESTING, 0x21)

- #define **RTTST_RTIOC_IRQBENCH_GET_STATS** _IOR(RTIOC_TYPE_TESTING, 0x22, struct rttst_irqbench_stats)
- #define **RTTST_RTIOC_IRQBENCH_WAIT_IRQ** _IO(RTIOC_TYPE_TESTING, 0x23)
- #define **RTTST_RTIOC_IRQBENCH_REPLY_IRQ** _IO(RTIOC_TYPE_TESTING, 0x24)
- #define **RTTST_RTIOC_SWTEST_SET_TASKS_COUNT** _IOW(RTIOC_TYPE_TESTING, 0x30, unsigned long)
- #define **RTTST_RTIOC_SWTEST_SET_CPU** _IOW(RTIOC_TYPE_TESTING, 0x31, unsigned long)
- #define **RTTST_RTIOC_SWTEST_REGISTER_UTASK** _IOW(RTIOC_TYPE_TESTING, 0x32, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_CREATE_KTASK** _IOWR(RTIOC_TYPE_TESTING, 0x33, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_PEND** _IOR(RTIOC_TYPE_TESTING, 0x34, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_SWITCH_TO** _IOR(RTIOC_TYPE_TESTING, 0x35, struct rttst_swtest_dir)
- #define **RTTST_RTIOC_SWTEST_GET_SWITCHES_COUNT** _IOR(RTIOC_TYPE_-TESTING, 0x36, unsigned long)
- #define **RTTST_RTIOC_SWTEST_GET_LAST_ERROR** _IOR(RTIOC_TYPE_TESTING, 0x37, struct rttst_swtest_error)
- #define **RTTST_RTIOC_SWTEST_SET_PAUSE** _IOW(RTIOC_TYPE_TESTING, 0x38, unsigned long)
- #define **RTTST_RTIOC_RTDM_DEFER_CLOSE** _IOW(RTIOC_TYPE_TESTING, 0x40, unsigned long)

5.6.1 Detailed Description

This group of devices is intended to provide in-kernel testing results. Feel free to comment on this profile via the Xenomai mailing list (xenomai-core@gna.org) or directly to the author (jan.kiszka@web.de).

Profile Revision: 2

Device Characteristics

Device Flags: RTDM_NAMED_DEVICE
Device Name: "rttest[-<subclass>]<N>", N >= 0, optional subclass name to simplify device discovery
Device Class: RTDM_CLASS_TESTING

Supported Operations

Open
Environments: non-RT (RT optional, deprecated)
Specific return values: none
Close
Environments: non-RT (RT optional, deprecated)
Specific return values: none
IOCTL
Mandatory Environments: see TSTIOCTLs below
Specific return values: see TSTIOCTLs below

5.7   Inter-Driver API

Collaboration diagram for Inter-Driver API:

Functions

- struct rtdm_dev_context ∗ rtdm_context_get (int fd)

    Retrieve and lock a device context.

- int rtdm_select_bind (int fd, rtdm_selector_t ∗selector, enum rtdm_selecttype type, unsigned fd_index)

    Bind a selector to specified event types of a given file descriptor.

- void rtdm_context_lock (struct rtdm_dev_context ∗context)

    Increment context reference counter.

- void rtdm_context_unlock (struct rtdm_dev_context ∗context)

    Decrement context reference counter.

- void rtdm_context_put (struct rtdm_dev_context ∗context)

    Release a device context obtained via rtdm_context_get().

- int rtdm_open (const char ∗path, int oflag,...)

    Open a device.

- int rtdm_socket (int protocol_family, int socket_type, int protocol)

    Create a socket.

- int rtdm_close (int fd)

    Close a device or socket.

- int rtdm_ioctl (int fd, int request,...)

    Issue an IOCTL.

- ssize_t rtdm_read (int fd, void ∗buf, size_t nbyte)

    Read from device.

- ssize_t rtdm_write (int fd, const void ∗buf, size_t nbyte)

    Write to device.

- ssize_t rtdm_recvmsg (int fd, struct msghdr *msg, int flags)

  Receive message from socket.

- ssize_t rtdm_recvfrom (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)

  Receive message from socket.

- ssize_t rtdm_recv (int fd, void *buf, size_t len, int flags)

  Receive message from socket.

- ssize_t rtdm_sendmsg (int fd, const struct msghdr *msg, int flags)

  Transmit message to socket.

- ssize_t rtdm_sendto (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)

  Transmit message to socket.

- ssize_t rtdm_send (int fd, const void *buf, size_t len, int flags)

  Transmit message to socket.

- int rtdm_bind (int fd, const struct sockaddr *my_addr, socklen_t addrlen)

  Bind to local address.

- int rtdm_connect (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)

  Connect to remote address.

- int rtdm_listen (int fd, int backlog)

  Listen for incomming connection requests.

- int rtdm_accept (int fd, struct sockaddr *addr, socklen_t *addrlen)

  Accept a connection requests.

- int rtdm_shutdown (int fd, int how)

  Shut down parts of a connection.

- int rtdm_getsockopt (int fd, int level, int optname, void *optval, socklen_t *optlen)

  Get socket option.

- int rtdm_setsockopt (int fd, int level, int optname, const void *optval, socklen_t optlen)

  Set socket option.

- int rtdm_getsockname (int fd, struct sockaddr *name, socklen_t *namelen)

  Get local socket address.

- int rtdm_getpeername (int fd, struct sockaddr *name, socklen_t *namelen)

  Get socket destination address.

5.7.1 Function Documentation

5.7.1.1 int rtdm_accept (int *fd,* struct sockaddr * *addr,* socklen_t * *addrlen*)

Accept a connection requests.

Refer to rt_dev_accept() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.2 int rtdm_bind (int *fd,* const struct sockaddr * *my_addr,* socklen_t *addrlen*)

Bind to local address.

Refer to rt_dev_bind() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.3 int rtdm_close (int *fd*)

Close a device or socket.

Refer to rt_dev_close() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.4 int rtdm_connect (int *fd,* const struct sockaddr * *serv_addr,* socklen_t *addrlen*)

Connect to remote address.

Refer to rt_dev_connect() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.5 struct rtdm_dev_context* rtdm_context_get (int *fd*) **[read]**

Retrieve and lock a device context.

Parameters

    [in] *fd* File descriptor

Returns

    Pointer to associated device context, or NULL on error

Note

The device context has to be unlocked using rtdm_context_put() when it is no longer referenced.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

References rtdm_dev_context::close_lock_count.

Referenced by rtdm_select_bind().

5.7.1.6 void rtdm_context_lock (struct rtdm_dev_context * *context*)

Increment context reference counter.

Parameters

[in] *context* Device context

Note

rtdm_context_get() automatically increments the lock counter. You only need to call this function in special scenarios, e.g. when keeping additional references to the context structure that have different lifetimes. Only use rtdm_context_lock() on contexts that are currently locked via an earlier rtdm_context_get()/rtdm_contex_lock() or while running a device operation handler.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.7.1.7 void rtdm_context_put (struct rtdm_dev_context ∗ *context*)

Release a device context obtained via rtdm_context_get().

Parameters

    [in] *context* Device context

Note

    Every successful call to rtdm_context_get() must be matched by a rtdm_context_put() invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.7.1.8 void rtdm_context_unlock (struct rtdm_dev_context ∗ *context*)

Decrement context reference counter.

Parameters

    [in] *context* Device context

Note

    Every call to rtdm_context_locked() must be matched by a rtdm_context_unlock() invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

Referenced by rtdm_select_bind().

5.7.1.9   int rtdm_getpeername (int *fd,* struct sockaddr ∗ *name,* socklen_t ∗ *namelen*)

Get socket destination address.

Refer to rt_dev_getpeername() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.10   int rtdm_getsockname (int *fd,* struct sockaddr ∗ *name,* socklen_t ∗ *namelen*)

Get local socket address.

Refer to rt_dev_getsockname() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.11   int rtdm_getsockopt (int *fd,* int *level,* int *optname,* void ∗ *optval,* socklen_t ∗ *optlen*)

Get socket option.

Refer to rt_dev_getsockopt() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.12   int rtdm_ioctl (int *fd,* int *request,* ...)

Issue an IOCTL.

Refer to rt_dev_ioctl() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.13   int rtdm_listen (int *fd,* int *backlog*)

Listen for incomming connection requests.

Refer to rt_dev_listen() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.14 int rtdm_open (const char ∗ *path,* int *oflag,* ...)

Open a device.

Refer to rt_dev_open() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.15 ssize_t rtdm_read (int *fd,* void ∗ *buf,* size_t *nbyte*)

Read from device.

Refer to rt_dev_read() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.16 ssize_t rtdm_recv (int *fd,* void ∗ *buf,* size_t *len,* int *flags*)

Receive message from socket.

Refer to rt_dev_recv() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.17 ssize_t rtdm_recvfrom (int *fd,* void ∗ *buf,* size_t *len,* int *flags,* struct sockaddr ∗ *from,* socklen_t ∗ *fromlen*)

Receive message from socket.

Refer to rt_dev_recvfrom() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.18 ssize_t rtdm_recvmsg (int *fd,* struct msghdr ∗ *msg,* int *flags*)

Receive message from socket.

Refer to rt_dev_recvmsg() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.19 int rtdm_select_bind (int *fd,* rtdm_selector_t ∗ *selector,* enum rtdm_selecttype *type,* unsigned *fd_index*)

Bind a selector to specified event types of a given file descriptor.

For internal use only.

This function is invoked by higher RTOS layers implementing select-like services. It shall not be called directly by RTDM drivers.

Parameters

> [in] *fd* File descriptor to bind to
>
> [in,out] *selector* Selector object that shall be bound to the given event
>
> [in] *type* Event type the caller is interested in
>
> [in] *fd_index* Index in the file descriptor set of the caller

Returns

> 0 on success, otherwise:

- -EBADF is returned if the file descriptor *fd* cannot be resolved.

- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

References rtdm_dev_context::ops, rtdm_context_get(), rtdm_context_unlock(), and rtdm_-operations::select_bind.

5.7.1.20 ssize_t rtdm_send (int *fd,* const void ∗ *buf,* size_t *len,* int *flags*)

Transmit message to socket.

Refer to rt_dev_send() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.21 ssize_t rtdm_sendmsg (int *fd,* const struct msghdr ∗ *msg,* int *flags*)

Transmit message to socket.

Refer to rt_dev_sendmsg() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.22 ssize_t rtdm_sendto (int *fd,* const void ∗ *buf,* size_t *len,* int *flags,* const struct sockaddr ∗ *to,* socklen_t *tolen*)

Transmit message to socket.

Refer to rt_dev_sendto() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.23 int rtdm_setsockopt (int *fd,* int *level,* int *optname,* const void ∗ *optval,* socklen_t *optlen*)

Set socket option.

Refer to rt_dev_setsockopt() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.24 int rtdm_shutdown (int *fd,* int *how*)

Shut down parts of a connection.

Refer to rt_dev_shutdown() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.25 int rtdm_socket (int *protocol_family,* int *socket_type,* int *protocol*)

Create a socket.

Refer to rt_dev_socket() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.7.1.26   ssize_t rtdm_write (int *fd,* const void ∗ *buf,* size_t *nbyte*)

Write to device.

Refer to rt_dev_write() for parameters and return values

Environments:

Depends on driver implementation, see Device Profiles.

Rescheduling: possible.

5.8   Device Registration Services

Collaboration diagram for Device Registration Services:

Data Structures

- struct rtdm_operations

    Device operations.

- struct rtdm_dev_context

    Device context.

- struct rtdm_device

    RTDM device.

Modules

- Synchronisation Services

Functions

- static void ∗ rtdm_context_to_private (struct rtdm_dev_context ∗context)

    Locate the driver private area associated to a device context structure.

- static struct rtdm_dev_context ∗ rtdm_private_to_context (void ∗dev_private)

    Locate a device context structure from its driver private area.

- int rtdm_dev_register (struct rtdm_device ∗device)

    Register a RTDM device.

- int rtdm_dev_unregister (struct rtdm_device ∗device, unsigned int poll_delay)

    Unregisters a RTDM device.

Operation Handler Prototypes

- typedef int(∗ rtdm_open_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, int oflag)

    Named device open handler.

- typedef int(∗ rtdm_socket_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, int protocol)

    Socket creation handler for protocol devices.

- typedef int(∗ rtdm_close_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info)

    Close handler.

- typedef int(∗ rtdm_ioctl_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, unsigned int request, void __user ∗arg)

    IOCTL handler.

- typedef int(∗ rtdm_select_bind_handler_t)(struct rtdm_dev_context ∗context, rtdm_-selector_t ∗selector, enum rtdm_selecttype type, unsigned fd_index)

    Select binding handler.

- typedef ssize_t(∗ rtdm_read_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_-info_t ∗user_info, void ∗buf, size_t nbyte)

    Read handler.

- typedef ssize_t(∗ rtdm_write_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_-info_t ∗user_info, const void ∗buf, size_t nbyte)

    Write handler.

- typedef ssize_t(∗ rtdm_recvmsg_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_-info_t ∗user_info, struct msghdr ∗msg, int flags)

    Receive message handler.

- typedef ssize_t(∗ rtdm_sendmsg_handler_t)(struct rtdm_dev_context ∗context, rtdm_-user_info_t ∗user_info, const struct msghdr ∗msg, int flags)

    Transmit message handler.

Device Flags

Static flags describing a RTDM device

- #define RTDM_EXCLUSIVE 0x0001

If set, only a single instance of the device can be requested by an application.

- #define RTDM_NAMED_DEVICE 0x0010

  If set, the device is addressed via a clear-text name.

- #define RTDM_PROTOCOL_DEVICE 0x0020

  If set, the device is addressed via a combination of protocol ID and socket type.

- #define RTDM_DEVICE_TYPE_MASK 0x00F0

  Mask selecting the device type.

Context Flags

Dynamic flags describing the state of an open RTDM device (bit numbers)

- #define RTDM_CREATED_IN_NRT 0

  Set by RTDM if the device instance was created in non-real-time context.

- #define RTDM_CLOSING 1

  Set by RTDM when the device is being closed.

- #define RTDM_USER_CONTEXT_FLAG 8

  Lowest bit number the driver developer can use freely.

Driver Versioning

Current revisions of RTDM structures, encoding of driver versions. See API Versioning for the interface revision.

- #define RTDM_DEVICE_STRUCT_VER 5

  Version of struct rtdm_device.

- #define RTDM_CONTEXT_STRUCT_VER 3

  Version of struct rtdm_dev_context.

- #define RTDM_SECURE_DEVICE 0x80000000

  Flag indicating a secure variant of RTDM (not supported here).

- #define RTDM_DRIVER_VER(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))

  Version code constructor for driver revisions.

- #define RTDM_DRIVER_MAJOR_VER(ver) (((ver) >> 16) & 0xFF)

  Get major version number from driver revision code.

- #define RTDM_DRIVER_MINOR_VER(ver) (((ver) >> 8) & 0xFF)

  Get minor version number from driver revision code.

- #define RTDM_DRIVER_PATCH_VER(ver) ((ver) & 0xFF)

  Get patch version number from driver revision code.

5.8.1 Define Documentation

5.8.1.1 #define RTDM_CLOSING 1

Set by RTDM when the device is being closed.

5.8.1.2 #define RTDM_CREATED_IN_NRT 0

Set by RTDM if the device instance was created in non-real-time context.

5.8.1.3 #define RTDM_DEVICE_TYPE_MASK 0x00F0

Mask selecting the device type.

Referenced by rtdm_dev_register(), and rtdm_dev_unregister().

5.8.1.4 #define RTDM_EXCLUSIVE 0x0001

If set, only a single instance of the device can be requested by an application.

Referenced by rtdm_dev_register().

5.8.1.5 #define RTDM_NAMED_DEVICE 0x0010

If set, the device is addressed via a clear-text name.

Referenced by rtdm_dev_register(), and rtdm_dev_unregister().

5.8.1.6 #define RTDM_PROTOCOL_DEVICE 0x0020

If set, the device is addressed via a combination of protocol ID and socket type.

Referenced by rtdm_dev_register().

5.8.2 Typedef Documentation

5.8.2.1 typedef int(∗ rtdm_close_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info)

Close handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode or deferred user mode call

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, -EAGAIN to request a recall after a grace period, or a valid negative error code according to IEEE Std 1003.1.

Note

Drivers must be prepared for that case that the close handler is invoked more than once per open context (even if the handler already completed an earlier run successfully). The driver has to avoid releasing resources twice as well as returning false errors on successive close invocations.

See also

`close()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.2 typedef int(* rtdm_ioctl_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, unsigned int request, void __user *arg)

IOCTL handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *request* Request number as passed by the user

[in,out] *arg* Request argument as passed by the user

Returns

A positive value or 0 on success. On failure return either -ENOSYS, to request that the function be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

`ioctl()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.3 typedef int(* rtdm_open_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, int oflag)

Named device open handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *oflag* Open flags as passed by the user

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

open() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.2.4 typedef ssize_t(∗ rtdm_read_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, void ∗buf, size_t nbyte)

Read handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[out] *buf* Input buffer as passed by the user

[in] *nbyte* Number of bytes the user requests to read

Returns

On success, the number of bytes read. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

read() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.2.5 typedef ssize_t(∗ rtdm_recvmsg_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, struct msghdr ∗msg, int flags)

Receive message handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in,out] *msg* Message descriptor as passed by the user, automatically mirrored to safe kernel memory in case of user mode call

[in] *flags* Message flags as passed by the user

Returns

On success, the number of bytes received. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

recvmsg() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.2.6   typedef int(∗ rtdm_select_bind_handler_t)(struct rtdm_dev_context ∗context, rtdm_selector_t ∗selector, enum rtdm_selecttype type, unsigned fd_index)

Select binding handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in,out] *selector* Object that shall be bound to the given event

[in] *type* Event type the selector is interested in

[in] *fd_index* Opaque value, to be passed to rtdm_event_select_bind or rtdm_sem_select_-bind unmodfied

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

5.8.2.7   typedef ssize_t(∗ rtdm_sendmsg_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, const struct msghdr ∗msg, int flags)

Transmit message handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *msg* Message descriptor as passed by the user, automatically mirrored to safe kernel memory in case of user mode call

[in] *flags* Message flags as passed by the user

Returns

On success, the number of bytes transmitted. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

sendmsg() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.2.8 typedef int(∗ rtdm_socket_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, int protocol)

Socket creation handler for protocol devices.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *protocol* Protocol number as passed by the user

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

socket() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.2.9 typedef ssize_t(∗ rtdm_write_handler_t)(struct rtdm_dev_context ∗context, rtdm_user_info_t ∗user_info, const void ∗buf, size_t nbyte)

Write handler.

Parameters

[in] *context* Context structure associated with opened device instance

[in] *user_info* Opaque pointer to information about user mode caller, NULL if kernel mode call

[in] *buf* Output buffer as passed by the user

[in] *nbyte* Number of bytes the user requests to write

Returns

On success, the number of bytes written. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See also

write() in IEEE Std 1003.1, http://www.opengroup.org/onlinepubs/009695399

5.8.3 Function Documentation

5.8.3.1 static void∗ rtdm_context_to_private (struct rtdm_dev_context ∗ *context*) [inline, static]

Locate the driver private area associated to a device context structure.

Parameters

    [in] *context* Context structure associated with opened device instance

Returns

    The address of the private driver area associated to *context*.

References rtdm_dev_context::dev_private.

5.8.3.2 int rtdm_dev_register (struct rtdm_device ∗ *device*)

Register a RTDM device.

Parameters

    [in] *device* Pointer to structure describing the new device.

Returns

    0 is returned upon success. Otherwise:

- -EINVAL is returned if the device structure contains invalid entries. Check kernel log in this case.

- -ENOMEM is returned if the context for an exclusive device cannot be allocated.

- -EEXIST is returned if the specified device name of protocol ID is already in use.

- -EAGAIN is returned if some /proc entry cannot be created.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References rtdm_operations::close_nrt, rtdm_operations::close_rt, rtdm_device::context_-size, rtdm_device::device_class, rtdm_device::device_flags, rtdm_device::device_name, rtdm_device::device_sub_class, rtdm_device::driver_version, rtdm_device::open_rt, rtdm_-device::ops, rtdm_device::proc_name, rtdm_device::profile_version, rtdm_device::protocol_-family, rtdm_device::reserved, RTDM_DEVICE_STRUCT_VER, RTDM_DEVICE_TYPE_-MASK, RTDM_EXCLUSIVE, RTDM_NAMED_DEVICE, RTDM_PROTOCOL_DEVICE, rtdm_operations::select_bind, rtdm_device::socket_rt, rtdm_device::socket_type, and rtdm_-device::struct_version.

5.8.3.3 int rtdm_dev_unregister (struct rtdm_device ∗ *device,* unsigned int *poll_delay*)

Unregisters a RTDM device.

Parameters

    [in] *device* Pointer to structure describing the device to be unregistered.

[in] *poll_delay* Polling delay in milliseconds to check repeatedly for open instances of *device*, or 0 for non-blocking mode.

Returns

0 is returned upon success. Otherwise:

- -ENODEV is returned if the device was not registered.

- -EAGAIN is returned if the device is busy with open instances and 0 has been passed for *poll_delay*.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References rtdm_device::device_flags, rtdm_device::device_name, rtdm_device::protocol_family, rtdm_device::reserved, RTDM_DEVICE_TYPE_MASK, RTDM_NAMED_DEVICE, and rtdm_-device::socket_type.

5.8.3.4 static struct rtdm_dev_context* rtdm_private_to_context (void * *dev_private*) [static, read]

Locate a device context structure from its driver private area.

Parameters

[in] *dev_private* Address of a private context area

Returns

The address of the device context structure defining *dev_private*.

5.9 Driver Development API

Collaboration diagram for Driver Development API:

Modules

- Inter-Driver API
- Device Registration Services
- Clock Services
- Task Services
- Timer Services
- Synchronisation Services
- Interrupt Management Services
- Non-Real-Time Signalling Services
- Utility Services

Files

- file rtdm_driver.h

     Real-Time Driver Model for Xenomai, driver API header.

5.9.1 Detailed Description

This is the lower interface of RTDM provided to device drivers, currently limited to kernel-space. Real-time drivers should only use functions of this interface in order to remain portable.

5.10 Clock Services

Collaboration diagram for Clock Services:

Functions

- **nanosecs_abs_t rtdm_clock_read** (void)

  Get system time.

- **nanosecs_abs_t rtdm_clock_read_monotonic** (void)

  Get monotonic time.

5.10.1 Function Documentation

5.10.1.1 nanosecs_abs_t rtdm_clock_read (void)

Get system time.

Returns

The system time in nanoseconds is returned

Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period.

The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.10.1.2   nanosecs_abs_t rtdm_clock_read_monotonic (void)

Get monotonic time.

Returns

The monotonic time in nanoseconds is returned

Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period.
The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.11   Task Services

Collaboration diagram for Task Services:

Typedefs

- typedef void(∗ rtdm_task_proc_t)(void ∗arg)

    Real-time task procedure.

Functions

- int rtdm_task_init (rtdm_task_t ∗task, const char ∗name, rtdm_task_proc_t task_proc, void ∗arg, int priority, nanosecs_rel_t period)

    Intialise and start a real-time task.

- void rtdm_task_destroy (rtdm_task_t ∗task)

    Destroy a real-time task.

- void rtdm_task_set_priority (rtdm_task_t ∗task, int priority)

    Adjust real-time task priority.

- int rtdm_task_set_period (rtdm_task_t ∗task, nanosecs_rel_t period)

    Adjust real-time task period.

- int rtdm_task_wait_period (void)

    Wait on next real-time task period.

- int rtdm_task_unblock (rtdm_task_t ∗task)

    Activate a blocked real-time task.

- rtdm_task_t ∗ rtdm_task_current (void)

    Get current real-time task.

- int rtdm_task_sleep (nanosecs_rel_t delay)

    Sleep a specified amount of time.

- int rtdm_task_sleep_until (nanosecs_abs_t wakeup_time)

    Sleep until a specified absolute time.

- int rtdm_task_sleep_abs (nanosecs_abs_t wakeup_time, enum rtdm_timer_mode mode)

    Sleep until a specified absolute time.

- void rtdm_task_join_nrt (rtdm_task_t ∗task, unsigned int poll_delay)

    Wait on a real-time task to terminate.

- void rtdm_task_busy_sleep (nanosecs_rel_t delay)

    Busy-wait a specified amount of time.

Task Priority Range

Maximum and minimum task priorities

- #define **RTDM_TASK_LOWEST_PRIORITY** XNSCHED_LOW_PRIO
- #define **RTDM_TASK_HIGHEST_PRIORITY** XNSCHED_HIGH_PRIO

Task Priority Modification

Raise or lower task priorities by one level

- #define **RTDM_TASK_RAISE_PRIORITY** (+1)
- #define **RTDM_TASK_LOWER_PRIORITY** (-1)

5.11.1 Typedef Documentation

5.11.1.1 typedef void(* rtdm_task_proc_t)(void *arg)

Real-time task procedure.

Parameters

[in,out] *arg* argument as passed to rtdm_task_init()

5.11.2 Function Documentation

5.11.2.1 void rtdm_task_busy_sleep (nanosecs_rel_t *delay*)

Busy-wait a specified amount of time.

Parameters

[in] *delay* Delay in nanoseconds. Note that a zero delay does **not** have the meaning of
RTDM_TIMEOUT_INFINITE here.

Note

The caller must not be migratable to different CPUs while executing this service. Otherwise,
the actual delay will be undefined.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine (should be avoided or kept short)

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never (except due to external interruptions).

5.11.2.2 rtdm_task_t* rtdm_task_current (void)

Get current real-time task.

Returns

Pointer to task handle

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.11.2.3 void rtdm_task_destroy (rtdm_task_t * *task*)

Destroy a real-time task.

Parameters

[in,out] *task* Task handle as returned by rtdm_task_init()

Note

Passing the same task handle to RTDM services after the completion of this function is not allowed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.11.2.4 int rtdm_task_init (rtdm_task_t * *task*, const char * *name*, rtdm_task_proc_t *task_proc*, void * *arg*, int *priority*, nanosecs_rel_t *period*)

Intialise and start a real-time task.

After initialising a task, the task handle remains valid and can be passed to RTDM services until either rtdm_task_destroy() or rtdm_task_join_nrt() was invoked.

Parameters

[in,out] *task* Task handle

[in] *name* Optional task name

[in] *task_proc* Procedure to be executed by the task

[in] *arg* Custom argument passed to `task_proc()` on entry

[in] *priority* Priority of the task, see also Task Priority Range

[in] *period* Period in nanoseconds of a cyclic task, 0 for non-cyclic mode

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.11.2.5 void rtdm_task_join_nrt (rtdm_task_t ∗ *task,* unsigned int *poll_delay*)

Wait on a real-time task to terminate.

Parameters

[in,out] *task* Task handle as returned by rtdm_task_init()

[in] *poll_delay* Delay in milliseconds between periodic tests for the state of the real-time task. This parameter is ignored if the termination is internally realised without polling.

Note

Passing the same task handle to RTDM services after the completion of this function is not allowed.
This service does not trigger the termination of the targeted task. The user has to take of this, otherwise rtdm_task_join_nrt() will never return.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- User-space task (non-RT)

Rescheduling: possible.

5.11.2.6 int rtdm_task_set_period (rtdm_task_t ∗ *task,* nanosecs_rel_t *period*)

Adjust real-time task period.

Parameters

    [in,out] *task* Task handle as returned by rtdm_task_init()

    [in] *period* New period in nanoseconds of a cyclic task, 0 for non-cyclic mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.11.2.7 void rtdm_task_set_priority (rtdm_task_t ∗ *task,* int *priority*)

Adjust real-time task priority.

Parameters

    [in,out] *task* Task handle as returned by rtdm_task_init()

    [in] *priority* New priority of the task, see also Task Priority Range

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.11.2.8 int rtdm_task_sleep (nanosecs_rel_t *delay*)

Sleep a specified amount of time.

Parameters

    [in] *delay* Delay in nanoseconds, see RTDM_TIMEOUT_xxx for special values.

Returns

    0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: always.

5.11.2.9 int rtdm_task_sleep_abs (nanosecs_abs_t *wakeup_time,* enum rtdm_timer_mode *mode*)

Sleep until a specified absolute time.

Parameters

    [in] *wakeup_time* Absolute timeout in nanoseconds

    [in] *mode* Selects the timer mode, see RTDM_TIMERMODE_xxx for details

Returns

    0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().

- -EPERM *may* be returned if an illegal invocation environment is detected.

- -EINVAL is returned if an invalid parameter was passed.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

5.11.2.10 int rtdm_task_sleep_until (nanosecs_abs_t *wakeup_time*)

Sleep until a specified absolute time.

Deprecated

    Use rtdm_task_sleep_abs instead!

Parameters

[in] *wakeup_time* Absolute timeout in nanoseconds

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_- unblock().

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

5.11.2.11   int rtdm_task_unblock (rtdm_task_t ∗ *task*)

Activate a blocked real-time task.

Returns

Non-zero is returned if the task was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.11.2.12   int rtdm_task_wait_period (void   )

Wait on next real-time task period.

Returns

0 on success, otherwise:

- -EINVAL is returned if calling task is not in periodic mode.

- -ETIMEDOUT is returned if a timer overrun occurred, which indicates that a previous release point has been missed by the calling task.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: always, unless a timer overrun occured.

5.12 Timer Services

Collaboration diagram for Timer Services:

Typedefs

- typedef void(∗ rtdm_timer_handler_t)(rtdm_timer_t ∗timer)

    Timer handler.

Functions

- int rtdm_timer_init (rtdm_timer_t ∗timer, rtdm_timer_handler_t handler, const char ∗name)

    Initialise a timer.

- void rtdm_timer_destroy (rtdm_timer_t ∗timer)

    Destroy a timer.

- int rtdm_timer_start (rtdm_timer_t ∗timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)

    Start a timer.

- void rtdm_timer_stop (rtdm_timer_t ∗timer)

    Stop a timer.

- int rtdm_timer_start_in_handler (rtdm_timer_t ∗timer, nanosecs_abs_t expiry, nanosecs_-rel_t interval, enum rtdm_timer_mode mode)

    Start a timer from inside a timer handler.

- void rtdm_timer_stop_in_handler (rtdm_timer_t ∗timer)

    Stop a timer from inside a timer handler.

RTDM_TIMERMODE_xxx

Timer operation modes

- enum rtdm_timer_mode { RTDM_TIMERMODE_RELATIVE = XN_RELATIVE, RTDM_-TIMERMODE_ABSOLUTE = XN_ABSOLUTE, RTDM_TIMERMODE_REALTIME = XN_-REALTIME }

5.12.1 Typedef Documentation

5.12.1.1 typedef void(∗ rtdm_timer_handler_t)(rtdm_timer_t ∗timer)

Timer handler.

Parameters

    [in] *timer* Timer handle as returned by rtdm_timer_init()

5.12.2 Enumeration Type Documentation

5.12.2.1 enum rtdm_timer_mode

Enumerator:

    RTDM_TIMERMODE_RELATIVE Monotonic timer with relative timeout.
    RTDM_TIMERMODE_ABSOLUTE Monotonic timer with absolute timeout.
    RTDM_TIMERMODE_REALTIME Adjustable timer with absolute timeout.

5.12.3 Function Documentation

5.12.3.1 void rtdm_timer_destroy (rtdm_timer_t ∗ *timer*)

Destroy a timer.

Parameters

    [in,out] *timer* Timer handle as returned by rtdm_timer_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.12.3.2   int rtdm_timer_init (rtdm_timer_t ∗ *timer,* rtdm_timer_handler_t *handler,* const char ∗ *name*)

Initialise a timer.

Parameters

[in,out] *timer* Timer handle

[in] *handler* Handler to be called on timer expiry

[in] *name* Optional timer name

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.12.3.3   int rtdm_timer_start (rtdm_timer_t ∗ *timer,* nanosecs_abs_t *expiry,* nanosecs_rel_t *interval,* enum rtdm_timer_mode *mode*)

Start a timer.

Parameters

[in,out] *timer* Timer handle as returned by rtdm_timer_init()

[in] *expiry* Firing time of the timer, `mode` defines if relative or absolute

[in] *interval* Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers

[in] *mode* Defines the operation mode, see RTDM_TIMERMODE_xxx for possible values

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if `expiry` describes an absolute date in the past.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.12.3.4 int rtdm_timer_start_in_handler (rtdm_timer_t ∗ *timer,* nanosecs_abs_t *expiry,* nanosecs_rel_t *interval,* enum rtdm_timer_mode *mode*)

Start a timer from inside a timer handler.

Parameters

> [in,out] *timer* Timer handle as returned by rtdm_timer_init()
>
> [in] *expiry* Firing time of the timer, mode defines if relative or absolute
>
> [in] *interval* Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers
>
> [in] *mode* Defines the operation mode, see RTDM_TIMERMODE_xxx for possible values

Returns

> 0 on success, otherwise:

- -ETIMEDOUT is returned if expiry describes an absolute date in the past.

Environments:

This service can be called from:

- Timer handler

Rescheduling: never.

5.12.3.5 void rtdm_timer_stop (rtdm_timer_t ∗ *timer*)

Stop a timer.

Parameters

> [in,out] *timer* Timer handle as returned by rtdm_timer_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.12.3.6 void rtdm_timer_stop_in_handler (rtdm_timer_t ∗ *timer*)

Stop a timer from inside a timer handler.

Parameters

    [in,out] *timer* Timer handle as returned by rtdm_timer_init()

Environments:

This service can be called from:

- Timer handler

Rescheduling: never.

5.13 Synchronisation Services

Collaboration diagram for Synchronisation Services:

Functions

- int rtdm_select_bind (int fd, rtdm_selector_t ∗selector, enum rtdm_selecttype type, unsigned fd_index)

    Bind a selector to specified event types of a given file descriptor.

RTDM_SELECTTYPE_xxx

Event types select can bind to

- enum rtdm_selecttype { RTDM_SELECTTYPE_READ = XNSELECT_READ, RTDM_-SELECTTYPE_WRITE = XNSELECT_WRITE, RTDM_SELECTTYPE_EXCEPT = XNSELECT_EXCEPT }

Spinlock with Preemption Deactivation

- typedef rthal_spinlock_t rtdm_lock_t

  Lock variable.

- typedef unsigned long rtdm_lockctx_t

  Variable to save the context while holding a lock.

- #define RTDM_LOCK_UNLOCKED RTHAL_SPIN_LOCK_UNLOCKED

  Static lock initialisation.

- #define rtdm_lock_init(lock) rthal_spin_lock_init(lock)

  Dynamic lock initialisation.

- #define rtdm_lock_get(lock) rthal_spin_lock(lock)

  Acquire lock from non-preemptible contexts.

- #define rtdm_lock_put(lock) rthal_spin_unlock(lock)

  Release lock without preemption restoration.

- #define rtdm_lock_get_irqsave(lock, context) rthal_spin_lock_irqsave(lock, context)

  Acquire lock and disable preemption.

- #define rtdm_lock_put_irqrestore(lock, context) rthal_spin_unlock_irqrestore(lock, context)

  Release lock and restore preemption state.

- #define rtdm_lock_irqsave(context) rthal_local_irq_save(context)

  Disable preemption locally.

- #define rtdm_lock_irqrestore(context) rthal_local_irq_restore(context)

  Restore preemption state.

Timeout Sequence Management

- void rtdm_toseq_init (rtdm_toseq_t *timeout_seq, nanosecs_rel_t timeout)

  Initialise a timeout sequence.

- EXPORT_SYMBOL_GPL (rtdm_toseq_init)

  Initialise a timeout sequence.

Event Services

- void rtdm_event_init (rtdm_event_t ∗event, unsigned long pending)

  Initialise an event.

- EXPORT_SYMBOL_GPL (rtdm_event_init)

  Initialise an event.

- void rtdm_event_destroy (rtdm_event_t ∗event)

  Destroy an event.

- void rtdm_event_pulse (rtdm_event_t ∗event)

  Signal an event occurrence to currently listening waiters.

- void rtdm_event_signal (rtdm_event_t ∗event)

  Signal an event occurrence.

- EXPORT_SYMBOL_GPL (rtdm_event_signal)

  Initialise an event.

- int rtdm_event_wait (rtdm_event_t ∗event)

  Wait on event occurrence.

- EXPORT_SYMBOL_GPL (rtdm_event_wait)

  Initialise an event.

- int rtdm_event_timedwait (rtdm_event_t ∗event, nanosecs_rel_t timeout, rtdm_toseq_t ∗timeout_seq)

  Wait on event occurrence with timeout.

- EXPORT_SYMBOL_GPL (rtdm_event_timedwait)

  Initialise an event.

- void rtdm_event_clear (rtdm_event_t ∗event)

  Clear event state.

- EXPORT_SYMBOL_GPL (rtdm_event_clear)

  Initialise an event.

- int rtdm_event_select_bind (rtdm_event_t ∗event, rtdm_selector_t ∗selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to an event.

- EXPORT_SYMBOL_GPL (rtdm_event_select_bind)

  Initialise an event.

Semaphore Services

- void rtdm_sem_init (rtdm_sem_t ∗sem, unsigned long value)

  Initialise a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_init)

  Initialise a semaphore.

- void rtdm_sem_destroy (rtdm_sem_t ∗sem)

  Destroy a semaphore.

- int rtdm_sem_down (rtdm_sem_t ∗sem)

  Decrement a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_down)

  Initialise a semaphore.

- int rtdm_sem_timeddown (rtdm_sem_t ∗sem, nanosecs_rel_t timeout, rtdm_toseq_-t ∗timeout_seq)

  Decrement a semaphore with timeout.

- EXPORT_SYMBOL_GPL (rtdm_sem_timeddown)

  Initialise a semaphore.

- void rtdm_sem_up (rtdm_sem_t ∗sem)

  Increment a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_up)

  Initialise a semaphore.

- int rtdm_sem_select_bind (rtdm_sem_t ∗sem, rtdm_selector_t ∗selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_select_bind)

  Initialise a semaphore.

Mutex Services

- void rtdm_mutex_init (rtdm_mutex_t ∗mutex)

  Initialise a mutex.

- EXPORT_SYMBOL_GPL (rtdm_mutex_init)

  Initialise a mutex.

- void rtdm_mutex_destroy (rtdm_mutex_t ∗mutex)

  Destroy a mutex.

- void rtdm_mutex_unlock (rtdm_mutex_t ∗mutex)

  Release a mutex.

- int rtdm_mutex_lock (rtdm_mutex_t ∗mutex)

  Request a mutex.

- EXPORT_SYMBOL_GPL (rtdm_mutex_lock)

  Initialise a mutex.

- int rtdm_mutex_timedlock (rtdm_mutex_t ∗mutex, nanosecs_rel_t timeout, rtdm_toseq_t ∗timeout_seq)

  Request a mutex with timeout.

- EXPORT_SYMBOL_GPL (rtdm_mutex_timedlock)

  Initialise a mutex.

Global Lock across Scheduler Invocation

- #define RTDM_EXECUTE_ATOMICALLY(code_block)

  Execute code block atomically.

5.13.1 Define Documentation

5.13.1.1 #define RTDM_EXECUTE_ATOMICALLY(*code_block*)

Value:

```
{                                           \
        <ENTER_ATOMIC_SECTION>              \
        code_block;                         \
        <LEAVE_ATOMIC_SECTION>             \
}
```

Execute code block atomically.

Generally, it is illegal to suspend the current task by calling rtdm_task_sleep(), rtdm_event_-wait(), etc. while holding a spinlock. In contrast, this macro allows to combine several operations including a potentially rescheduling call to an atomic code block with respect to other RTDM_-EXECUTE_ATOMICALLY() blocks. The macro is a light-weight alternative for protecting code blocks via mutexes, and it can even be used to synchronise real-time and non-real-time contexts.

Parameters

    code_block Commands to be executed atomically

Note

It is not allowed to leave the code block explicitly by using break, return, goto, etc. This would leave the global lock held during the code block execution in an inconsistent state. Moreover, do not embed complex operations into the code bock. Consider that they will be executed under preemption lock with interrupts switched-off. Also note that invocation of rescheduling calls may break the atomicity until the task gains the CPU again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible, depends on functions called within *code_block*.

5.13.1.2 #define rtdm_lock_get(*lock*) rthal_spin_lock(lock)

Acquire lock from non-preemptible contexts.

Parameters

   lock Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.1.3 #define rtdm_lock_get_irqsave(*lock,*   *context*) rthal_spin_lock_irqsave(lock, context)

Acquire lock and disable preemption.

Parameters

   lock Address of lock variable
   context name of local variable to store the context in

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.1.4 #define rtdm_lock_init(*lock*) rthal_spin_lock_init(lock)

Dynamic lock initialisation.

Parameters

> *lock* Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.1.5 #define rtdm_lock_irqrestore(*context*) rthal_local_irq_restore(context)

Restore preemption state.

Parameters

> *context* name of local variable which stored the context

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.1.6 #define rtdm_lock_irqsave(*context*) rthal_local_irq_save(context)

Disable preemption locally.

Parameters

> *context* name of local variable to store the context in

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.1.7   #define rtdm_lock_put(*lock*) rthal_spin_unlock(lock)

Release lock without preemption restoration.

Parameters

   lock Address of lock variable

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.1.8   #define rtdm_lock_put_irqrestore(*lock*,   *context*) rthal_spin_unlock_- irqrestore(lock, context)

Release lock and restore preemption state.

Parameters

   lock Address of lock variable
   context name of local variable which stored the context

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.2 Enumeration Type Documentation

5.13.2.1 enum rtdm_selecttype

Enumerator:

    RTDM_SELECTTYPE_READ   Select input data availability events.

    RTDM_SELECTTYPE_WRITE   Select ouput buffer availability events.

    RTDM_SELECTTYPE_EXCEPT   Select exceptional events.

5.13.3 Function Documentation

5.13.3.1 EXPORT_SYMBOL_GPL (rtdm_toseq_init)

Initialise a timeout sequence.

This service initialises a timeout sequence handle according to the given timeout value. Timeout sequences allow to maintain a continuous *timeout* across multiple calls of blocking synchronisation services. A typical application scenario is given below.

Parameters

    [in,out] *timeout_seq*   Timeout sequence handle

    [in] *timeout*   Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values

Application Scenario:

```
int device_service_routine(...)
{
      rtdm_toseq_t timeout_seq;
      ...

      rtdm_toseq_init(&timeout_seq, timeout);
      ...
      while (received < requested) {
            ret = rtdm_event_timedwait(&data_available, timeout, &timeout_seq
);
            if (ret < 0) // including -ETIMEDOUT
                  break;

            // receive some data
            ...
      }
      ...
}
```

Using a timeout sequence in such a scenario avoids that the user-provided relative `timeout` is restarted on every call to rtdm_event_timedwait(), potentially causing an overall delay that is larger than specified by `timeout`. Moreover, all functions supporting timeout sequences also interpret special timeout values (infinite and non-blocking), disburdening the driver developer from handling them separately.

Environments:

This service can be called from:

- Kernel-based task

• User-space task (RT)

Rescheduling: never.

5.13.3.2 EXPORT_SYMBOL_GPL (rtdm_event_init)

Initialise an event.

Parameters

[in,out] *event* Event handle

[in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

• Kernel module initialization/cleanup code

• Kernel-based task

• User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.3 EXPORT_SYMBOL_GPL (rtdm_event_signal)

Initialise an event.

Parameters

[in,out] *event* Event handle

[in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

• Kernel module initialization/cleanup code

• Kernel-based task

• User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.4 EXPORT_SYMBOL_GPL (rtdm_event_select_bind)

Initialise an event.

Parameters

[in,out] *event* Event handle

[in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.5 EXPORT_SYMBOL_GPL (rtdm_mutex_lock)

Initialise a mutex.

This function initalises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

Parameters

[in,out] *mutex* Mutex handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.6 EXPORT_SYMBOL_GPL (rtdm_sem_init)

Initialise a semaphore.

Parameters

[in,out] *sem* Semaphore handle
[in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.7 EXPORT_SYMBOL_GPL (rtdm_mutex_timedlock)

Initialise a mutex.

This function initalises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

Parameters

[in,out] *mutex* Mutex handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.8 EXPORT_SYMBOL_GPL (rtdm_event_wait)

Initialise an event.

Parameters

[in,out] *event* Event handle
[in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.9 EXPORT_SYMBOL_GPL (rtdm_sem_down)

Initialise a semaphore.

Parameters

[in,out] *sem* Semaphore handle
[in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.10 EXPORT_SYMBOL_GPL (rtdm_event_timedwait)

Initialise an event.

Parameters

    [in,out] *event* Event handle

    [in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.11 EXPORT_SYMBOL_GPL (rtdm_sem_timeddown)

Initialise a semaphore.

Parameters

    [in,out] *sem* Semaphore handle

    [in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.12 EXPORT_SYMBOL_GPL (rtdm_sem_up)

Initialise a semaphore.

Parameters

> [in,out] *sem* Semaphore handle
>
> [in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.13 EXPORT_SYMBOL_GPL (rtdm_event_clear)

Initialise an event.

Parameters

> [in,out] *event* Event handle
>
> [in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.14 EXPORT_SYMBOL_GPL (rtdm_sem_select_bind)

Initialise a semaphore.

Parameters

> [in,out] *sem* Semaphore handle
>
> [in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.15 EXPORT_SYMBOL_GPL (rtdm_mutex_init)

Initialise a mutex.

This function initalises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

Parameters

    [in,out] *mutex* Mutex handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.16 void rtdm_event_clear (rtdm_event_t ∗ *event*)

Clear event state.

Parameters

    [in,out] *event* Event handle as returned by rtdm_event_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.17 void rtdm_event_destroy (rtdm_event_t * *event*)

Destroy an event.

Parameters

    [in,out] *event* Event handle as returned by rtdm_event_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.18 void rtdm_event_init (rtdm_event_t * *event,* unsigned long *pending*)

Initialise an event.

Parameters

    [in,out] *event* Event handle

    [in] *pending* Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.19 void rtdm_event_pulse (rtdm_event_t * *event*)

Signal an event occurrence to currently listening waiters.

This function wakes up all current waiters of the given event, but it does not change the event state. Subsequently callers of rtdm_event_wait() or rtdm_event_timedwait() will therefore be blocked first.

Parameters

    [in,out] *event* Event handle as returned by rtdm_event_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.20 int rtdm_event_select_bind (rtdm_event_t ∗ *event,* rtdm_selector_t ∗ *selector,* enum rtdm_selecttype *type,* unsigned *fd_index*)

Bind a selector to an event.

This functions binds the given selector to an event so that the former is notified when the event state changes. Typically the select binding handler will invoke this service.

Parameters

> [in,out] *event* Event handle as returned by rtdm_event_init()

> [in,out] *selector* Selector as passed to the select binding handler

> [in] *type* Type of the bound event as passed to the select binding handler

> [in] *fd_index* File descriptor index as passed to the select binding handler

Returns

> 0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.

- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.21 void rtdm_event_signal (rtdm_event_t ∗ *event*)

Signal an event occurrence.

This function sets the given event and wakes up all current waiters. If no waiter is presently registered, the next call to rtdm_event_wait() or rtdm_event_timedwait() will return immediately.

Parameters

> [in,out] *event* Event handle as returned by rtdm_event_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.22 int rtdm_event_timedwait (rtdm_event_t * *event,* nanosecs_rel_t *timeout,* rtdm_toseq_t * *timeout_seq*)

Wait on event occurrence with timeout.

This function waits or tests for the occurence of the given event, taking the provided timeout into account. On successful return, the event is reset.

Parameters

[in,out] *event* Event handle as returned by rtdm_event_init()

[in] *timeout* Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values

[in,out] *timeout_seq* Handle of a timeout sequence as returned by rtdm_toseq_init() or NULL

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().

- -EIDRM is returned if *event* has been destroyed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

- -EWOULDBLOCK is returned if a negative *timeout* (i.e., non-blocking operation) has been specified.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

Referenced by rtdm_event_wait().

5.13.3.23   int rtdm_event_wait (rtdm_event_t ∗ *event*)

Wait on event occurrence.

This is the light-weight version of rtdm_event_timedwait(), implying an infinite timeout.

Parameters

> [in,out] *event* Event handle as returned by rtdm_event_init()

Returns

> 0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().

- -EIDRM is returned if *event* has been destroyed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

References rtdm_event_timedwait().

5.13.3.24   void rtdm_mutex_destroy (rtdm_mutex_t ∗ *mutex*)

Destroy a mutex.

Parameters

> [in,out] *mutex* Mutex handle as returned by rtdm_mutex_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.25   void rtdm_mutex_init (rtdm_mutex_t ∗ *mutex*)

Initialise a mutex.

This function initalises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

Parameters

    [in,out] *mutex* Mutex handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.26   int rtdm_mutex_lock (rtdm_mutex_t ∗ *mutex*)

Request a mutex.

This is the light-weight version of rtdm_mutex_timedlock(), implying an infinite timeout.

Parameters

    [in,out] *mutex* Mutex handle as returned by rtdm_mutex_init()

Returns

    0 on success, otherwise:

- -EIDRM is returned if *mutex* has been destroyed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

References rtdm_mutex_timedlock().

5.13.3.27   int rtdm_mutex_timedlock (rtdm_mutex_t ∗ *mutex,* nanosecs_rel_t *timeout,* rtdm_toseq_t ∗ *timeout_seq*)

Request a mutex with timeout.

This function tries to acquire the given mutex. If it is not available, the caller is blocked unless non-blocking operation was selected.

Parameters

[in,out] *mutex* Mutex handle as returned by rtdm_mutex_init()

[in] *timeout* Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values

[in,out] *timeout_seq* Handle of a timeout sequence as returned by rtdm_toseq_init() or NULL

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.

- -EIDRM is returned if *mutex* has been destroyed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

Referenced by rtdm_mutex_lock().

5.13.3.28   void rtdm_mutex_unlock (rtdm_mutex_t ∗ *mutex*)

Release a mutex.

This function releases the given mutex, waking up a potential waiter which was blocked upon rtdm_mutex_lock() or rtdm_mutex_timedlock().

Parameters

[in,out] *mutex* Mutex handle as returned by rtdm_mutex_init()

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: possible.

5.13.3.29   int rtdm_select_bind (int *fd,* rtdm_selector_t ∗ *selector,* enum rtdm_selecttype *type,* unsigned *fd_index*)

Bind a selector to specified event types of a given file descriptor.

For internal use only.

This function is invoked by higher RTOS layers implementing select-like services. It shall not be called directly by RTDM drivers.

Parameters

    [in] *fd* File descriptor to bind to

    [in,out] *selector* Selector object that shall be bound to the given event

    [in] *type* Event type the caller is interested in

    [in] *fd_index* Index in the file descriptor set of the caller

Returns

    0 on success, otherwise:

- -EBADF is returned if the file descriptor *fd* cannot be resolved.

- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

References rtdm_dev_context::ops, rtdm_context_get(), rtdm_context_unlock(), and rtdm_-operations::select_bind.

5.13.3.30   void rtdm_sem_destroy (rtdm_sem_t ∗ *sem*)

Destroy a semaphore.

Parameters

    [in,out] *sem* Semaphore handle as returned by rtdm_sem_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.31 int rtdm_sem_down (rtdm_sem_t ∗ *sem*)

Decrement a semaphore.

This is the light-weight version of rtdm_sem_timeddown(), implying an infinite timeout.

Parameters

> [in,out] *sem* Semaphore handle as returned by rtdm_sem_init()

Returns

> 0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().
- -EIDRM is returned if *sem* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References rtdm_sem_timeddown().

5.13.3.32 void rtdm_sem_init (rtdm_sem_t ∗ *sem,* unsigned long *value*)

Initialise a semaphore.

Parameters

> [in,out] *sem* Semaphore handle
>
> [in] *value* Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.33 int rtdm_sem_select_bind (rtdm_sem_t * *sem,* rtdm_selector_t * *selector,* enum rtdm_selecttype *type,* unsigned *fd_index*)

Bind a selector to a semaphore.

This functions binds the given selector to the semaphore so that the former is notified when the semaphore state changes. Typically the select binding handler will invoke this service.

Parameters

    [in,out] *sem* Semaphore handle as returned by rtdm_sem_init()

    [in,out] *selector* Selector as passed to the select binding handler

    [in] *type* Type of the bound event as passed to the select binding handler

    [in] *fd_index* File descriptor index as passed to the select binding handler

Returns

    0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.

- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.13.3.34 int rtdm_sem_timeddown (rtdm_sem_t * *sem,* nanosecs_rel_t *timeout,* rtdm_toseq_t * *timeout_seq*)

Decrement a semaphore with timeout.

This function tries to decrement the given semphore's value if it is positive on entry. If not, the caller is blocked unless non-blocking operation was selected.

Parameters

> [in,out] *sem* Semaphore handle as returned by rtdm_sem_init()
>
> [in] *timeout* Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values
>
> [in,out] *timeout_seq* Handle of a timeout sequence as returned by rtdm_toseq_init() or NULL

Returns

> 0 on success, otherwise:

- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.

- -EINTR is returned if calling task has been unblock by a signal or explicitly via rtdm_task_-unblock().

- -EIDRM is returned if *sem* has been destroyed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

Referenced by rtdm_sem_down().

5.13.3.35 void rtdm_sem_up (rtdm_sem_t ∗ *sem*)

Increment a semaphore.

This function increments the given semphore's value, waking up a potential waiter which was blocked upon rtdm_sem_down().

Parameters

> [in,out] *sem* Semaphore handle as returned by rtdm_sem_init()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.13.3.36 void rtdm_toseq_init (rtdm_toseq_t ∗ *timeout_seq,* nanosecs_rel_t *timeout*)

Initialise a timeout sequence.

This service initialises a timeout sequence handle according to the given timeout value. Timeout sequences allow to maintain a continuous *timeout* across multiple calls of blocking synchronisation services. A typical application scenario is given below.

Parameters

> [in,out] *timeout_seq* Timeout sequence handle

> [in] *timeout* Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values

Application Scenario:

```
int device_service_routine(...)
{
        rtdm_toseq_t timeout_seq;
        ...

        rtdm_toseq_init(&timeout_seq, timeout);
        ...
        while (received < requested) {
                ret = rtdm_event_timedwait(&data_available, timeout, &timeout_seq
    );
                if (ret < 0) // including -ETIMEDOUT
                        break;

                // receive some data
                ...
        }
        ...
}
```

Using a timeout sequence in such a scenario avoids that the user-provided relative `timeout` is restarted on every call to rtdm_event_timedwait(), potentially causing an overall delay that is larger than specified by `timeout`. Moreover, all functions supporting timeout sequences also interpret special timeout values (infinite and non-blocking), disburdening the driver developer from handling them separately.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)

Rescheduling: never.

5.14   Interrupt Management Services

Collaboration diagram for Interrupt Management Services:

Defines

- #define rtdm_irq_get_arg(irq_handle, type) ((type *)irq_handle->cookie)

  Retrieve IRQ handler argument.

Typedefs

- typedef int(* rtdm_irq_handler_t)(rtdm_irq_t *irq_handle)

  Interrupt handler.

Functions

- int rtdm_irq_request (rtdm_irq_t *irq_handle, unsigned int irq_no, rtdm_irq_handler_t handler, unsigned long flags, const char *device_name, void *arg)

  Register an interrupt handler.

- int rtdm_irq_free (rtdm_irq_t *irq_handle)

  Release an interrupt handler.

- int rtdm_irq_enable (rtdm_irq_t *irq_handle)

  Enable interrupt line.

- int rtdm_irq_disable (rtdm_irq_t *irq_handle)

  Disable interrupt line.

RTDM_IRQTYPE_xxx

Interrupt registrations flags

- #define RTDM_IRQTYPE_SHARED XN_ISR_SHARED

  Enable IRQ-sharing with other real-time drivers.

- #define RTDM_IRQTYPE_EDGE XN_ISR_EDGE
   Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.

RTDM_IRQ_xxx

Return flags of interrupt handlers

- #define RTDM_IRQ_NONE XN_ISR_NONE
   Unhandled interrupt.

- #define RTDM_IRQ_HANDLED XN_ISR_HANDLED
   Denote handled interrupt.

5.14.1 Define Documentation

5.14.1.1 #define rtdm_irq_get_arg(*irq_handle*, *type*) ((type *)irq_handle->cookie)

Retrieve IRQ handler argument.

Parameters

    irq_handle IRQ handle
    type Type of the pointer to return

Returns

    The argument pointer registered on rtdm_irq_request() is returned, type-casted to the specified *type*.

Environments:

This service can be called from:

- Interrupt service routine

Rescheduling: never.

5.14.2 Typedef Documentation

5.14.2.1 typedef int(* rtdm_irq_handler_t)(rtdm_irq_t *irq_handle)

Interrupt handler.

Parameters

    [in] *irq_handle* IRQ handle as returned by rtdm_irq_request()

Returns

    0 or a combination of RTDM_IRQ_xxx flags

5.14.3 Function Documentation

5.14.3.1 int rtdm_irq_disable (rtdm_irq_t * *irq_handle*)

Disable interrupt line.

Parameters

    [in,out] *irq_handle* IRQ handle as returned by rtdm_irq_request()

Returns

    0 on success, otherwise negative error code

Note

    This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.14.3.2 int rtdm_irq_enable (rtdm_irq_t * *irq_handle*)

Enable interrupt line.

Parameters

    [in,out] *irq_handle* IRQ handle as returned by rtdm_irq_request()

Returns

    0 on success, otherwise negative error code

Note

    This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.14.3.3 int rtdm_irq_free (rtdm_irq_t * *irq_handle*)

Release an interrupt handler.

Parameters

>   [in,out] *irq_handle* IRQ handle as returned by rtdm_irq_request()

Returns

>   0 on success, otherwise negative error code

Note

>   The caller is responsible for shutting down the IRQ source at device level before invoking this service. In turn, rtdm_irq_free ensures that any pending event on the given IRQ line is fully processed on return from this service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.14.3.4 int rtdm_irq_request (rtdm_irq_t * *irq_handle,* unsigned int *irq_no,* rtdm_irq_handler_t *handler,* unsigned long *flags,* const char * *device_name,* void * *arg*)

Register an interrupt handler.

This function registers the provided handler with an IRQ line and enables the line.

Parameters

>   [in,out] *irq_handle* IRQ handle
>
>   [in] *irq_no* Line number of the addressed IRQ
>
>   [in] *handler* Interrupt handler
>
>   [in] *flags* Registration flags, see RTDM_IRQTYPE_xxx for details

[in] *device_name* Device name to show up in real-time IRQ lists

[in] *arg* Pointer to be passed to the interrupt handler on invocation

Returns

0 on success, otherwise:

- -EINVAL is returned if an invalid parameter was passed.

- -EBUSY is returned if the specified IRQ line is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.15 Non-Real-Time Signalling Services

Collaboration diagram for Non-Real-Time Signalling Services:

Typedefs

- typedef void(* rtdm_nrtsig_handler_t)(rtdm_nrtsig_t nrt_sig, void *arg)

  Non-real-time signal handler.

Functions

- int rtdm_nrtsig_init (rtdm_nrtsig_t *nrt_sig, rtdm_nrtsig_handler_t handler, void *arg)

  Register a non-real-time signal handler.

- void rtdm_nrtsig_destroy (rtdm_nrtsig_t *nrt_sig)

  Release a non-realtime signal handler.

- void rtdm_nrtsig_pend (rtdm_nrtsig_t ∗nrt_sig)

    Trigger non-real-time signal.

5.15.1 Detailed Description

These services provide a mechanism to request the execution of a specified handler in non-real-time context. The triggering can safely be performed in real-time context without suffering from unknown delays. The handler execution will be deferred until the next time the real-time subsystem releases the CPU to the non-real-time part.

5.15.2 Typedef Documentation

5.15.2.1 typedef void(∗ rtdm_nrtsig_handler_t)(rtdm_nrtsig_t nrt_sig, void ∗arg)

Non-real-time signal handler.

Parameters

[in] *nrt_sig* Signal handle as returned by rtdm_nrtsig_init()

[in] *arg* Argument as passed to rtdm_nrtsig_init()

Note

The signal handler will run in soft-IRQ context of the non-real-time subsystem. Note the implications of this context, e.g. no invocation of blocking operations.

5.15.3 Function Documentation

5.15.3.1 void rtdm_nrtsig_destroy (rtdm_nrtsig_t ∗ *nrt_sig*)

Release a non-realtime signal handler.

Parameters

[in,out] *nrt_sig* Signal handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.15.3.2 int rtdm_nrtsig_init (rtdm_nrtsig_t ∗ *nrt_sig,* rtdm_nrtsig_handler_t *handler,* void ∗ *arg*)

Register a non-real-time signal handler.

Parameters

> [in,out] *nrt_sig* Signal handle

> [in] *handler* Non-real-time signal handler

> [in] *arg* Custom argument passed to handler() on each invocation

Returns

> 0 on success, otherwise:

- -EAGAIN is returned if no free signal slot is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.15.3.3 void rtdm_nrtsig_pend (rtdm_nrtsig_t ∗ *nrt_sig*)

Trigger non-real-time signal.

Parameters

> [in,out] *nrt_sig* Signal handle

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

5.16 Utility Services

Collaboration diagram for Utility Services:

Functions

- int rtdm_mmap_to_user (rtdm_user_info_t *user_info, void *src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)

  Map a kernel memory range into the address space of the user.

- int rtdm_iomap_to_user (rtdm_user_info_t *user_info, phys_addr_t src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)

  Map an I/O memory range into the address space of the user.

- int rtdm_munmap (rtdm_user_info_t *user_info, void *ptr, size_t len)

  Unmap a user memory range.

- void rtdm_printk (const char *format,...)

  Real-time safe message printing on kernel console.

- void * rtdm_malloc (size_t size)

  Allocate memory block in real-time context.

- void rtdm_free (void *ptr)

  Release real-time memory block.

- int rtdm_read_user_ok (rtdm_user_info_t *user_info, const void __user *ptr, size_t size)

  Check if read access to user-space memory block is safe.

- int rtdm_rw_user_ok (rtdm_user_info_t *user_info, const void __user *ptr, size_t size)

  Check if read/write access to user-space memory block is safe.

- int rtdm_copy_from_user (rtdm_user_info_t *user_info, void *dst, const void __user *src, size_t size)

  Copy user-space memory block to specified buffer.

- int rtdm_safe_copy_from_user (rtdm_user_info_t *user_info, void *dst, const void __user *src, size_t size)

  Check if read access to user-space memory block and copy it to specified buffer.

- int rtdm_copy_to_user (rtdm_user_info_t ∗user_info, void __user ∗dst, const void ∗src, size_t size)

    Copy specified buffer to user-space memory block.

- int rtdm_safe_copy_to_user (rtdm_user_info_t ∗user_info, void __user ∗dst, const void ∗src, size_t size)

    Check if read/write access to user-space memory block is safe and copy specified buffer to it.

- int rtdm_strncpy_from_user (rtdm_user_info_t ∗user_info, char ∗dst, const char __user ∗src, size_t count)

    Copy user-space string to specified buffer.

- int rtdm_in_rt_context (void)

    Test if running in a real-time task.

- int rtdm_rt_capable (rtdm_user_info_t ∗user_info)

    Test if the caller is capable of running in real-time context.

5.16.1 Function Documentation

5.16.1.1 int rtdm_copy_from_user (rtdm_user_info_t ∗ *user_info,* void ∗ *dst,* const void __user ∗ *src,* size_t *size*)

Copy user-space memory block to specified buffer.

Parameters

> [in] *user_info* User information pointer as passed to the invoked device operation handler
>
> [in] *dst* Destination buffer address
>
> [in] *src* Address of the user-space memory block
>
> [in] *size* Size of the memory block

Returns

> 0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

> Before invoking this service, verify via rtdm_read_user_ok() that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.2 int rtdm_copy_to_user (rtdm_user_info_t ∗ *user_info,* **void __user** ∗ *dst,* **const void** ∗ *src,* **size_t** *size* **)**

Copy specified buffer to user-space memory block.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *dst* Address of the user-space memory block

[in] *src* Source buffer address

[in] *size* Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

Before invoking this service, verify via rtdm_rw_user_ok() that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.3 void rtdm_free (void ∗ *ptr* **)**

Release real-time memory block.

Parameters

[in] *ptr* Pointer to memory block as returned by rtdm_malloc()

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine (consider the overhead!)

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.4 int rtdm_in_rt_context (void)

Test if running in a real-time task.

Returns

Non-zero is returned if the caller resides in real-time context, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.5 int rtdm_iomap_to_user (rtdm_user_info_t ∗ *user_info*, phys_addr_t *src_addr*, size_t *len*, int *prot*, void ∗∗ *pptr*, struct vm_operations_struct ∗ *vm_ops*, void ∗ *vm_private_data*)

Map an I/O memory range into the address space of the user.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *src_addr* physical I/O address to be mapped

[in] *len* Length of the memory range

[in] *prot* Protection flags for the user's memory range, typically either PROT_READ or PROT_READ|PROT_WRITE

[in,out] *pptr* Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return

[in] *vm_ops* vm_operations to be executed on the vma_area of the user memory range or NULL

[in] *vm_private_data* Private data to be stored in the vma_area, primarily useful for vm_-operation handlers

Returns

0 on success, otherwise (most common values):

- -EINVAL is returned if an invalid start address, size, or destination address was passed.

- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.

- -EAGAIN is returned if too much memory has been already locked by the user process.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Note

> RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via rtdm_munmap(), either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard munmap() or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of rtdm_iomap_to_user() can pass a vm_operations_struct on invocation, defining a close handler for the vm_area. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- User-space task (non-RT)

Rescheduling: possible.

5.16.1.6   void∗ rtdm_malloc (size_t *size*)

Allocate memory block in real-time context.

Parameters

> [in] *size* Requested size of the memory block

Returns

> The pointer to the allocated block is returned on success, NULL otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine (consider the overhead!)

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.7   int rtdm_mmap_to_user (rtdm_user_info_t ∗ *user_info*, void ∗ *src_addr*, size_t *len*, int *prot*, void ∗∗ *pptr*, struct vm_operations_struct ∗ *vm_ops*, void ∗ *vm_private_data*)

Map a kernel memory range into the address space of the user.

Parameters

> [in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *src_addr* Kernel virtual address to be mapped

[in] *len* Length of the memory range

[in] *prot* Protection flags for the user's memory range, typically either PROT_READ or PROT_READ|PROT_WRITE

[in,out] *pptr* Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return

[in] *vm_ops* vm_operations to be executed on the vma_area of the user memory range or NULL

[in] *vm_private_data* Private data to be stored in the vma_area, primarily useful for vm_-operation handlers

Returns

0 on success, otherwise (most common values):

- -EINVAL is returned if an invalid start address, size, or destination address was passed.

- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.

- -EAGAIN is returned if too much memory has been already locked by the user process.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Note

This service only works on memory regions allocated via kmalloc() or vmalloc(). To map physical I/O memory to user-space use rtdm_iomap_to_user() instead.
RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via rtdm_munmap(), either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard munmap() or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of rtdm_mmap_to_user() can pass a vm_operations_struct on invocation, defining a close handler for the vm_area. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- User-space task (non-RT)

Rescheduling: possible.

5.16.1.8 int rtdm_munmap (rtdm_user_info_t ∗ *user_info*, void ∗ *ptr*, size_t *len*)

Unmap a user memory range.

Parameters

    [in] *user_info* User information pointer as passed to rtdm_mmap_to_user() when requesting to map the memory range

    [in] *ptr* User address or the memory range

    [in] *len* Length of the memory range

Returns

    0 on success, otherwise:

- -EINVAL is returned if an invalid address or size was passed.

- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- User-space task (non-RT)

Rescheduling: possible.

5.16.1.9 void rtdm_printk (const char * *format,* ...)

Real-time safe message printing on kernel console.

Parameters

    [in] *format* Format string (conforming standard `printf()`)

    ... Arguments referred by *format*

Returns

    On success, this service returns the number of characters printed. Otherwise, a negative error code is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Interrupt service routine (consider the overhead!)

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

5.16.1.10   int rtdm_read_user_ok (rtdm_user_info_t ∗ *user_info,* const void __user ∗ *ptr,* size_t *size*)

Check if read access to user-space memory block is safe.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *ptr* Address of the user-provided memory block

[in] *size* Size of the memory block

Returns

Non-zero is return when it is safe to read from the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.11   int rtdm_rt_capable (rtdm_user_info_t ∗ *user_info*)

Test if the caller is capable of running in real-time context.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

Returns

Non-zero is returned if the caller is able to execute in real-time context (independent of its current execution mode), 0 otherwise.

Note

This function can be used by drivers that provide different implementations for the same service depending on the execution mode of the caller. If a caller requests such a service in non-real-time context but is capable of running in real-time as well, it might be appropriate for the driver to reject the request via -ENOSYS so that RTDM can switch the caller and restart the request in real-time context.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.12 int rtdm_rw_user_ok (rtdm_user_info_t * *user_info,* const void __user * *ptr,* size_t *size*)

Check if read/write access to user-space memory block is safe.

Parameters

> [in] *user_info* User information pointer as passed to the invoked device operation handler
>
> [in] *ptr* Address of the user-provided memory block
>
> [in] *size* Size of the memory block

Returns

> Non-zero is return when it is safe to read from or write to the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.13 int rtdm_safe_copy_from_user (rtdm_user_info_t * *user_info,* void * *dst,* const void __user * *src,* size_t *size*)

Check if read access to user-space memory block and copy it to specified buffer.

Parameters

> [in] *user_info* User information pointer as passed to the invoked device operation handler
>
> [in] *dst* Destination buffer address
>
> [in] *src* Address of the user-space memory block
>
> [in] *size* Size of the memory block

Returns

> 0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

> This service is a combination of rtdm_read_user_ok and rtdm_copy_from_user.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.14 int rtdm_safe_copy_to_user (rtdm_user_info_t * *user_info*, void __user * *dst*, const void * *src*, size_t *size*)

Check if read/write access to user-space memory block is safe and copy specified buffer to it.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *dst* Address of the user-space memory block

[in] *src* Source buffer address

[in] *size* Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

This service is a combination of rtdm_rw_user_ok and rtdm_copy_to_user.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.16.1.15 int rtdm_strncpy_from_user (rtdm_user_info_t * *user_info*, char * *dst*, const char __user * *src*, size_t *count*)

Copy user-space string to specified buffer.

Parameters

[in] *user_info* User information pointer as passed to the invoked device operation handler

[in] *dst* Destination buffer address

[in] *src* Address of the user-space string

[in] *count* Maximum number of bytes to copy, including the trailing '0'

Returns

Length of the string on success (not including the trailing '0'), otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

This services already includes a check of the source address, calling rtdm_read_user_ok() for *src* explicitly is not required.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: never.

5.17 Device Profiles

Collaboration diagram for Device Profiles:

Data Structures

- struct rtdm_device_info

    Device information.

Modules

- CAN Devices
- Real-time IPC protocols
    - *Profile Revision: 1*

- Serial Devices
- Testing Devices

Typedefs

- typedef struct rtdm_device_info rtdm_device_info_t

    Device information.

RTDM_CLASS_xxx

Device classes

- #define **RTDM_CLASS_PARPORT** 1
- #define **RTDM_CLASS_SERIAL** 2
- #define **RTDM_CLASS_CAN** 3
- #define **RTDM_CLASS_NETWORK** 4
- #define **RTDM_CLASS_RTMAC** 5
- #define **RTDM_CLASS_TESTING** 6
- #define **RTDM_CLASS_RTIPC** 7
- #define **RTDM_CLASS_EXPERIMENTAL** 224
- #define **RTDM_CLASS_MAX** 255

Device Naming

Maximum length of device names (excluding the final null character)

- #define **RTDM_MAX_DEVNAME_LEN** 31

RTDM_PURGE_xxx_BUFFER

Flags selecting buffers to be purged

- #define **RTDM_PURGE_RX_BUFFER** 0x0001
- #define **RTDM_PURGE_TX_BUFFER** 0x0002

Common IOCTLs

The following IOCTLs are common to all device profiles.

- #define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_-device_info)

Retrieve information about a device or socket.

- #define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)

  Purge internal device or socket buffers.

5.17.1 Detailed Description

Device profiles define which operation handlers a driver of a certain class has to implement, which name or protocol it has to register, which IOCTLs it has to provide, and further details. Sub-classes can be defined in order to extend a device profile with more hardware-specific functions.

5.17.2 Define Documentation

5.17.2.1 #define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)

Retrieve information about a device or socket.

Parameters

[out] *arg* Pointer to information buffer (struct rtdm_device_info)

5.17.2.2 #define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)

Purge internal device or socket buffers.

Parameters

[in] *arg* Purge mask, see RTDM_PURGE_xxx_BUFFER

Chapter 6

Data Structure Documentation

6.1   can_bittime Struct Reference

Custom CAN bit-time definition.

Collaboration diagram for can_bittime:

Data Fields

- can_bittime_type_t **type**

    Type of bit-time definition.

- struct can_bittime_std **std**

    Standard bit-time.

- struct can_bittime_btr **btr**

    Hardware-spcific BTR bit-time.

6.1.1 Detailed Description

Custom CAN bit-time definition.

Examples:

rtcanconfig.c.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.2 can_bittime_btr Struct Reference

Hardware-specific BTR bit-times.

Data Fields

- uint8_t btr0

  Bus timing register 0.

- uint8_t btr1

  Bus timing register 1.

6.2.1 Detailed Description

Hardware-specific BTR bit-times.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.3 can_bittime_std Struct Reference

Standard bit-time parameters according to Bosch.

Data Fields

- uint32_t brp

  Baud rate prescaler.

- uint8_t prop_seg

  from 1 to 8

- uint8_t phase_seg1

  from 1 to 8

- uint8_t phase_seg2

  from 1 to 8

- uint8_t sjw:7

  from 1 to 4

- uint8_t sam:1

  1 - enable triple sampling

6.3.1 Detailed Description

Standard bit-time parameters according to Bosch.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.4 can_filter Struct Reference

Filter for reception of CAN messages.

Data Fields

- uint32_t can_id

  CAN ID which must match with incoming IDs after passing the mask.

- uint32_t can_mask

  Mask which is applied to incoming IDs.

6.4.1 Detailed Description

Filter for reception of CAN messages. This filter works as follows: A received CAN ID is AND'ed bitwise with `can_mask` and then compared to `can_id`. This also includes the CAN_EFF_FLAG and CAN_RTR_FLAG of CAN_xxx_FLAG. If this comparison is true, the message will be received by the socket. The logic can be inverted with the `can_id` flag CAN_INV_FILTER :

```
if (can_id & CAN_INV_FILTER) {
   if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))
      accept-message;
} else {
   if ((received_can_id & can_mask) == can_id)
      accept-message;
}
```

Multiple filters can be arranged in a filter list and set with Sockopts. If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

Examples:

rtcan_rtt.c, and rtcanrecv.c.

6.4.2 Field Documentation

6.4.2.1 uint32_t can_filter::can_id

CAN ID which must match with incoming IDs after passing the mask.

The filter logic can be inverted with the flag CAN_INV_FILTER.

6.4.2.2 uint32_t can_filter::can_mask

Mask which is applied to incoming IDs.

See CAN ID masks if exactly one CAN ID should come through.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.5 can_frame Struct Reference

Raw CAN frame.

Public Member Functions

- uint8_t data[8] __attribute__ ((aligned(8)))
  Payload data bytes.

Data Fields

- can_id_t can_id
  CAN ID of the frame.

- uint8_t can_dlc
  Size of the payload in bytes.

6.5.1 Detailed Description

Raw CAN frame. Central structure for receiving and sending CAN frames.

Examples:

rtcan_rtt.c, rtcanrecv.c, and rtcansend.c.

6.5.2 Field Documentation

6.5.2.1 can_id_t can_frame::can_id

CAN ID of the frame.

See CAN ID flags for special bits.

Examples:

rtcan_rtt.c.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.6 rtdm_dev_context Struct Reference

Device context.

Collaboration diagram for rtdm_dev_context:

Data Fields

- unsigned long context_flags

  Context flags, see Context Flags for details.

- int fd

  Associated file descriptor.

- atomic_t close_lock_count

    Lock counter of context, held while structure is referenced by an operation handler.

- struct rtdm_operations * ops

    Set of active device operation handlers.

- struct rtdm_device * device

    Reference to owning device.

- struct rtdm_devctx_reserved reserved

    Data stored by RTDM inside a device context (internal use only).

- char dev_private [0]

    Begin of driver defined context data structure.

6.6.1 Detailed Description

Device context. A device context structure is associated with every open device instance. RTDM takes care of its creation and destruction and passes it to the operation handlers when being invoked.

Drivers can attach arbitrary data immediately after the official structure. The size of this data is provided via rtdm_device.context_size during device registration.

The documentation for this struct was generated from the following file:

- include/rtdm/rtdm_driver.h

6.7 rtdm_device Struct Reference

RTDM device.

Collaboration diagram for rtdm_device:

Data Fields

- int struct_version

  Revision number of this structure, see Driver Versioning defines.

- int device_flags

  Device flags, see Device Flags for details.

- size_t context_size

  Size of driver defined appendix to struct rtdm_dev_context.

- char device_name [RTDM_MAX_DEVNAME_LEN+1]

  Named device identification (orthogonal to Linux device name space).

- int protocol_family

  Protocol device identification: protocol family (PF_xxx).

- int socket_type

  Protocol device identification: socket type (SOCK_xxx).

- rtdm_open_handler_t open_rt

  Named device instance creation for real-time contexts, optional (but deprecated) if open_nrt is non-NULL, ignored for protocol devices.

- rtdm_open_handler_t open_nrt

  Named device instance creation for non-real-time contexts, optional if open_rt is non-NULL, ignored for protocol devices.

- rtdm_socket_handler_t socket_rt

  Protocol socket creation for real-time contexts, optional (but deprecated) if socket_nrt is non-NULL, ignored for named devices.

- rtdm_socket_handler_t socket_nrt

  Protocol socket creation for non-real-time contexts, optional if socket_rt is non-NULL, ignored for named devices.

- struct rtdm_operations ops

  Default operations on newly opened device instance.

- int device_class

  Device class ID, see RTDM_CLASS_xxx.

- int device_sub_class

  Device sub-class, see RTDM_SUBCLASS_xxx definition in the Device Profiles.

- int profile_version

  Supported device profile version.

- const char * driver_name

  Informational driver name (reported via /proc).

- int driver_version

    Driver version, see Driver Versioning defines.

- const char ∗ peripheral_name

    Informational peripheral name the device is attached to (reported via /proc).

- const char ∗ provider_name

    Informational driver provider name (reported via /proc).

- const char ∗ proc_name

    Name of /proc entry for the device, must not be NULL.

- int device_id

    Driver definable device ID.

- void ∗ device_data

    Driver definable device data.

- struct rtdm_dev_reserved reserved

    Data stored by RTDM inside a registered device (internal use only).

6.7.1 Detailed Description

RTDM device. This structure specifies a RTDM device. As some fields, especially the reserved area, will be modified by RTDM during runtime, the structure must not reside in write-protected memory.

6.7.2 Field Documentation

6.7.2.1 rtdm_open_handler_t rtdm_device::open_rt

Named device instance creation for real-time contexts, optional (but deprecated) if open_nrt is non-NULL, ignored for protocol devices.

Deprecated

    Only use non-real-time open handler in new drivers.

Referenced by rtdm_dev_register().

6.7.2.2 rtdm_socket_handler_t rtdm_device::socket_rt

Protocol socket creation for real-time contexts, optional (but deprecated) if socket_nrt is non-NULL, ignored for named devices.

Deprecated

    Only use non-real-time socket creation handler in new drivers.

Referenced by rtdm_dev_register().

The documentation for this struct was generated from the following file:

- include/rtdm/rtdm_driver.h

6.8 rtdm_device_info Struct Reference

Device information.

Data Fields

- int device_flags

    Device flags, see Device Flags for details.

- int device_class

    Device class ID, see RTDM_CLASS_xxx.

- int device_sub_class

    Device sub-class, either RTDM_SUBCLASS_GENERIC or a RTDM_SUBCLASS_xxx definition of the related Device Profile.

- int profile_version

    Supported device profile version.

6.8.1 Detailed Description

Device information.

The documentation for this struct was generated from the following file:

- include/rtdm/rtdm.h

6.9 rtdm_operations Struct Reference

Device operations.

Data Fields

Common Operations

- rtdm_close_handler_t close_rt

    Close handler for real-time contexts (optional, deprecated).

- rtdm_close_handler_t close_nrt

    Close handler for non-real-time contexts (required).

- rtdm_ioctl_handler_t ioctl_rt

    IOCTL from real-time context (optional).

- rtdm_ioctl_handler_t ioctl_nrt

    IOCTL from non-real-time context (optional).

- rtdm_select_bind_handler_t select_bind

    Select binding handler for any context (optional).

Stream-Oriented Device Operations

- rtdm_read_handler_t read_rt

    Read handler for real-time context (optional).

- rtdm_read_handler_t read_nrt

    Read handler for non-real-time context (optional).

- rtdm_write_handler_t write_rt

    Write handler for real-time context (optional).

- rtdm_write_handler_t write_nrt

    Write handler for non-real-time context (optional).

Message-Oriented Device Operations

- rtdm_recvmsg_handler_t recvmsg_rt

    Receive message handler for real-time context (optional).

- rtdm_recvmsg_handler_t recvmsg_nrt

    Receive message handler for non-real-time context (optional).

- rtdm_sendmsg_handler_t sendmsg_rt

    Transmit message handler for real-time context (optional).

- rtdm_sendmsg_handler_t sendmsg_nrt

    Transmit message handler for non-real-time context (optional).

6.9.1 Detailed Description

Device operations.

6.9.2 Field Documentation

6.9.2.1 rtdm_close_handler_t rtdm_operations::close_rt

Close handler for real-time contexts (optional, deprecated).

Deprecated

> Only use non-real-time close handler in new drivers.

Referenced by rtdm_dev_register().

The documentation for this struct was generated from the following file:

- include/rtdm/rtdm_driver.h

6.10 rtipc_port_label Struct Reference

Port label information structure.

Data Fields

- char label [XNOBJECT_NAME_LEN]

  Port label string, null-terminated.

6.10.1 Detailed Description

Port label information structure.

Examples:

  bufp-label.c, iddp-label.c, and xddp-label.c.

6.10.2 Field Documentation

6.10.2.1 char rtipc_port_label::label[XNOBJECT_NAME_LEN]

Port label string, null-terminated.

The documentation for this struct was generated from the following file:

- include/rtdm/rtipc.h

6.11 rtser_config Struct Reference

Serial device configuration.

Data Fields

- int config_mask

  mask specifying valid fields, see RTSER_SET_xxx

- int baud_rate

  baud rate, default RTSER_DEF_BAUD

- int parity

number of parity bits, see RTSER_xxx_PARITY

- int data_bits

  number of data bits, see RTSER_xxx_BITS

- int stop_bits

  number of stop bits, see RTSER_xxx_STOPB

- int handshake

  handshake mechanisms, see RTSER_xxx_HAND

- int fifo_depth

  reception FIFO interrupt threshold, see RTSER_FIFO_xxx

- nanosecs_rel_t rx_timeout

  reception timeout, see RTSER_TIMEOUT_xxx for special values

- nanosecs_rel_t tx_timeout

  transmission timeout, see RTSER_TIMEOUT_xxx for special values

- nanosecs_rel_t event_timeout

  event timeout, see RTSER_TIMEOUT_xxx for special values

- int timestamp_history

  enable timestamp history, see RTSER_xxx_TIMESTAMP_HISTORY

- int event_mask

  event mask to be used with RTSER_RTIOC_WAIT_EVENT, see RTSER_EVENT_xxx

6.11.1 Detailed Description

Serial device configuration.

Examples:

cross-link.c.

The documentation for this struct was generated from the following file:

- include/rtdm/rtserial.h

6.12 rtser_event Struct Reference

Additional information about serial device events.

Data Fields

- int events

  signalled events, see RTSER_EVENT_xxx

- int rx_pending

  number of pending input characters

- nanosecs_abs_t last_timestamp

  last interrupt timestamp

- nanosecs_abs_t rxpend_timestamp

  reception timestamp of oldest character in input queue

6.12.1 Detailed Description

Additional information about serial device events.

Examples:

  cross-link.c.

The documentation for this struct was generated from the following file:

- include/rtdm/rtserial.h

6.13 rtser_status Struct Reference

Serial device status.

Data Fields

- int line_status

  line status register, see RTSER_LSR_xxx

- int modem_status

  modem status register, see RTSER_MSR_xxx

6.13.1 Detailed Description

Serial device status.

The documentation for this struct was generated from the following file:

- include/rtdm/rtserial.h

6.14 sockaddr_can Struct Reference

Socket address structure for the CAN address family.

Data Fields

- sa_family_t can_family

  CAN address family, must be `AF_CAN`.

- int can_ifindex

  Interface index of CAN controller.

6.14.1 Detailed Description

Socket address structure for the CAN address family.

Examples:

rtcan_rtt.c, rtcanrecv.c, and rtcansend.c.

6.14.2 Field Documentation

6.14.2.1 int sockaddr_can::can_ifindex

Interface index of CAN controller.

See SIOCGIFINDEX.

The documentation for this struct was generated from the following file:

- include/rtdm/rtcan.h

6.15 sockaddr_ipc Struct Reference

Socket address structure for the RTIPC address family.

Data Fields

- sa_family_t sipc_family

  RTIPC address family, must be `AF_RTIPC`.

- rtipc_port_t sipc_port

  Port number.

6.15.1 Detailed Description

Socket address structure for the RTIPC address family.

Examples:

bufp-label.c, bufp-readwrite.c, iddp-label.c, iddp-sendrecv.c, xddp-echo.c, xddp-label.c, and xddp-stream.c.

6.15.2 Field Documentation

6.15.2.1 rtipc_port_t sockaddr_ipc::sipc_port

Port number.

The documentation for this struct was generated from the following file:

- include/rtdm/rtipc.h

Chapter 7

File Documentation

7.1   include/rtdm/rtcan.h File Reference

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

Include dependency graph for rtcan.h:

Data Structures

- struct can_bittime_std

  Standard bit-time parameters according to Bosch.

- struct can_bittime_btr

  Hardware-specific BTR bit-times.

- struct can_bittime

  Custom CAN bit-time definition.

- struct can_filter

  Filter for reception of CAN messages.

- struct sockaddr_can

Socket address structure for the CAN address family.

- struct can_frame

  Raw CAN frame.

Defines

- #define AF_CAN 29

  CAN address family.

- #define PF_CAN AF_CAN

  CAN protocol family.

- #define SOL_CAN_RAW 103

  CAN socket levels.

CAN ID masks

Bit masks for masking CAN IDs

- #define CAN_EFF_MASK 0x1FFFFFFF

  Bit mask for extended CAN IDs.

- #define CAN_SFF_MASK 0x000007FF

  Bit mask for standard CAN IDs.

CAN ID flags

Flags within a CAN ID indicating special CAN frame attributes

- #define CAN_EFF_FLAG 0x80000000

  Extended frame.

- #define CAN_RTR_FLAG 0x40000000

  Remote transmission frame.

- #define CAN_ERR_FLAG 0x20000000

  Error frame (see Errors), not valid in struct can_filter.

- #define CAN_INV_FILTER CAN_ERR_FLAG

  Invert CAN filter definition, only valid in struct can_filter.

Particular CAN protocols

Possible protocols for the PF_CAN protocol family

Currently only the RAW protocol is supported.

- #define CAN_RAW 1

  Raw protocol of `PF_CAN`, applicable to socket type `SOCK_RAW`.

CAN controller modes

Special CAN controllers modes, which can be or'ed together.

Note

> *These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.*

- #define CAN_CTRLMODE_LISTENONLY 0x1
- #define CAN_CTRLMODE_LOOPBACK 0x2

Timestamp switches

Arguments to pass to RTCAN_RTIOC_TAKE_TIMESTAMP

- #define RTCAN_TAKE_NO_TIMESTAMPS 0

    Switch off taking timestamps.

- #define RTCAN_TAKE_TIMESTAMPS 1

    Do take timestamps.

RAW socket options

Setting and getting CAN RAW socket options.

- #define CAN_RAW_FILTER 0x1

    CAN filter definition.

- #define CAN_RAW_ERR_FILTER 0x2

    CAN error mask.

- #define CAN_RAW_LOOPBACK 0x3

    CAN TX loopback.

- #define CAN_RAW_RECV_OWN_MSGS 0x4

    CAN receive own messages.

IOCTLs

CAN device IOCTLs

- #define SIOCGIFINDEX defined_by_kernel_header_file

    Get CAN interface index by name.

- #define SIOCSCANBAUDRATE _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)

    Set baud rate.

- #define SIOCGCANBAUDRATE _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)

    Get baud rate.

- #define SIOCSCANCUSTOMBITTIME _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)

    Set custom bit time parameter.

- #define SIOCGCANCUSTOMBITTIME _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)

    Get custum bit-time parameters.

- #define SIOCSCANMODE _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)

    Set operation mode of CAN controller.

- #define SIOCGCANSTATE _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)

    Get current state of CAN controller.

- #define SIOCSCANCTRLMODE _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)

    Set special controller modes.

- #define SIOCGCANCTRLMODE _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)

    Get special controller modes.

- #define RTCAN_RTIOC_TAKE_TIMESTAMP _IOW(RTIOC_TYPE_CAN, 0x09, int)

    Enable or disable storing a high precision timestamp upon reception of a CAN frame.

- #define RTCAN_RTIOC_RCV_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_-
    rel_t)

    Specify a reception timeout for a socket.

- #define RTCAN_RTIOC_SND_TIMEOUT _IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_-
    rel_t)

    Specify a transmission timeout for a socket.

Error mask

Error class (mask) in `can_id` *field of struct can_frame to be used with CAN_RAW_ERR_FILTER.*

Note: *Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.*

Note: *In case of a bus-off error condition (CAN_ERR_BUSOFF), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling CAN_MODE_-START.*

Note: *Bus error interrupts (CAN_ERR_BUSERROR) are enabled when an application is calling a Recv function on a socket listening on bus errors (using CAN_RAW_ERR_FILTER). After one bus error has occured, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.*

- #define CAN_ERR_TX_TIMEOUT 0x00000001U

    TX timeout (netdevice driver).

- #define CAN_ERR_LOSTARB 0x00000002U

    Lost arbitration (see data[0]).

- #define CAN_ERR_CRTL 0x00000004U

    Controller problems (see data[1]).

- #define CAN_ERR_PROT 0x00000008U

    Protocol violations (see data[2], data[3]).

- #define CAN_ERR_TRX 0x00000010U

    Transceiver status (see data[4]).

- #define CAN_ERR_ACK 0x00000020U

    Received no ACK on transmission.

- #define CAN_ERR_BUSOFF 0x00000040U

  Bus off.

- #define CAN_ERR_BUSERROR 0x00000080U

  Bus error (may flood!).

- #define CAN_ERR_RESTARTED 0x00000100U

  Controller restarted.

- #define CAN_ERR_MASK 0x1FFFFFFFU

  Omit EFF, RTR, ERR flags.

Arbitration lost error

Error in the data[0] field of struct can_frame.

- #define CAN_ERR_LOSTARB_UNSPEC 0x00

  unspecified

Controller problems

Error in the data[1] field of struct can_frame.

- #define CAN_ERR_CRTL_UNSPEC 0x00

  unspecified

- #define CAN_ERR_CRTL_RX_OVERFLOW 0x01

  RX buffer overflow.

- #define CAN_ERR_CRTL_TX_OVERFLOW 0x02

  TX buffer overflow.

- #define CAN_ERR_CRTL_RX_WARNING 0x04

  reached warning level for RX errors

- #define CAN_ERR_CRTL_TX_WARNING 0x08

  reached warning level for TX errors

- #define CAN_ERR_CRTL_RX_PASSIVE 0x10

  reached passive level for RX errors

- #define CAN_ERR_CRTL_TX_PASSIVE 0x20

  reached passive level for TX errors

Protocol error type

Error in the data[2] field of struct can_frame.

- #define CAN_ERR_PROT_UNSPEC 0x00

  unspecified

- #define CAN_ERR_PROT_BIT 0x01

single bit error

- #define CAN_ERR_PROT_FORM 0x02

  frame format error

- #define CAN_ERR_PROT_STUFF 0x04

  bit stuffing error

- #define CAN_ERR_PROT_BIT0 0x08

  unable to send dominant bit

- #define CAN_ERR_PROT_BIT1 0x10

  unable to send recessive bit

- #define CAN_ERR_PROT_OVERLOAD 0x20

  bus overload

- #define CAN_ERR_PROT_ACTIVE 0x40

  active error announcement

- #define CAN_ERR_PROT_TX 0x80

  error occured on transmission

Protocol error location

Error in the data[4] field of struct can_frame.

- #define CAN_ERR_PROT_LOC_UNSPEC 0x00

  unspecified

- #define CAN_ERR_PROT_LOC_SOF 0x03

  start of frame

- #define CAN_ERR_PROT_LOC_ID28_21 0x02

  ID bits 28 - 21 (SFF: 10 - 3).

- #define CAN_ERR_PROT_LOC_ID20_18 0x06

  ID bits 20 - 18 (SFF: 2 - 0).

- #define CAN_ERR_PROT_LOC_SRTR 0x04

  substitute RTR (SFF: RTR)

- #define CAN_ERR_PROT_LOC_IDE 0x05

  identifier extension

- #define CAN_ERR_PROT_LOC_ID17_13 0x07

  ID bits 17-13.

- #define CAN_ERR_PROT_LOC_ID12_05 0x0F

  ID bits 12-5.

- #define CAN_ERR_PROT_LOC_ID04_00 0x0E

  ID bits 4-0.

- #define CAN_ERR_PROT_LOC_RTR 0x0C

    RTR.

- #define CAN_ERR_PROT_LOC_RES1 0x0D

    reserved bit 1

- #define CAN_ERR_PROT_LOC_RES0 0x09

    reserved bit 0

- #define CAN_ERR_PROT_LOC_DLC 0x0B

    data length code

- #define CAN_ERR_PROT_LOC_DATA 0x0A

    data section

- #define CAN_ERR_PROT_LOC_CRC_SEQ 0x08

    CRC sequence.

- #define CAN_ERR_PROT_LOC_CRC_DEL 0x18

    CRC delimiter.

- #define CAN_ERR_PROT_LOC_ACK 0x19

    ACK slot.

- #define CAN_ERR_PROT_LOC_ACK_DEL 0x1B

    ACK delimiter.

- #define CAN_ERR_PROT_LOC_EOF 0x1A

    end of frame

- #define CAN_ERR_PROT_LOC_INTERM 0x12

    intermission

- #define CAN_ERR_TRX_UNSPEC 0x00

    0000 0000

- #define CAN_ERR_TRX_CANH_NO_WIRE 0x04

    0000 0100

- #define CAN_ERR_TRX_CANH_SHORT_TO_BAT 0x05

    0000 0101

- #define CAN_ERR_TRX_CANH_SHORT_TO_VCC 0x06

    0000 0110

- #define CAN_ERR_TRX_CANH_SHORT_TO_GND 0x07

    0000 0111

- #define CAN_ERR_TRX_CANL_NO_WIRE 0x40

    0100 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_BAT 0x50

0101 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_VCC 0x60

  0110 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_GND 0x70

  0111 0000

- #define CAN_ERR_TRX_CANL_SHORT_TO_CANH 0x80

  1000 0000

Typedefs

- typedef uint32_t can_id_t

  Type of CAN id (see CAN_xxx_MASK and CAN_xxx_FLAG).

- typedef can_id_t can_err_mask_t

  Type of CAN error mask.

- typedef uint32_t can_baudrate_t

  Baudrate definition in bits per second.

- typedef enum CAN_BITTIME_TYPE can_bittime_type_t

  See CAN_BITTIME_TYPE.

- typedef enum CAN_MODE can_mode_t

  See CAN_MODE.

- typedef int can_ctrlmode_t

  See CAN_CTRLMODE.

- typedef enum CAN_STATE can_state_t

  See CAN_STATE.

- typedef struct can_filter can_filter_t

  Filter for reception of CAN messages.

- typedef struct can_frame can_frame_t

  Raw CAN frame.

Enumerations

- enum CAN_BITTIME_TYPE { CAN_BITTIME_STD, CAN_BITTIME_BTR }

  Supported CAN bit-time types.

CAN operation modes

Modes into which CAN controllers can be set

- enum CAN_MODE { CAN_MODE_STOP = 0, CAN_MODE_START, CAN_MODE_-SLEEP }

CAN controller states

States a CAN controller can be in.

- enum CAN_STATE {

  CAN_STATE_ACTIVE = 0, CAN_STATE_BUS_WARNING, CAN_STATE_BUS_-PASSIVE, CAN_STATE_BUS_OFF,

  CAN_STATE_SCANNING_BAUDRATE, CAN_STATE_STOPPED, CAN_STATE_-SLEEPING }

7.1.1 Detailed Description

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

Note

Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>
Copyright (C) 2005, 2006 Sebastian Smolorz <Sebastian.Smolorz@stud.uni-hannover.de>

This RTDM CAN device profile header is based on:

include/linux/can.h, include/linux/socket.h, net/can/pf_can.h in linux-can.patch, a CAN socket framework for Linux

Copyright (C) 2004, 2005, Robert Schwebel, Benedikt Spranger, Marc Kleine-Budde, Pengutronix

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU

General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.2 include/rtdm/rtdm.h File Reference

Real-Time Driver Model for Xenomai, user API header.

Include dependency graph for rtdm.h:

This graph shows which files directly or indirectly include this file:

Data Structures

- struct rtdm_device_info

    Device information.

Defines

API Versioning

- #define RTDM_API_VER 8

    Common user and driver API version.

- #define RTDM_API_MIN_COMPAT_VER 6

    Minimum API revision compatible with the current release.

RTDM_TIMEOUT_xxx

Special timeout values

- #define RTDM_TIMEOUT_INFINITE 0

    Block forever.

- #define RTDM_TIMEOUT_NONE (-1)

    Any negative timeout means non-blocking.

RTDM_CLASS_xxx

Device classes

- #define **RTDM_CLASS_PARPORT** 1
- #define **RTDM_CLASS_SERIAL** 2
- #define **RTDM_CLASS_CAN** 3
- #define **RTDM_CLASS_NETWORK** 4
- #define **RTDM_CLASS_RTMAC** 5
- #define **RTDM_CLASS_TESTING** 6
- #define **RTDM_CLASS_RTIPC** 7
- #define **RTDM_CLASS_EXPERIMENTAL** 224
- #define **RTDM_CLASS_MAX** 255

Device Naming

Maximum length of device names (excluding the final null character)

- #define **RTDM_MAX_DEVNAME_LEN** 31

RTDM_PURGE_xxx_BUFFER

Flags selecting buffers to be purged

- #define **RTDM_PURGE_RX_BUFFER** 0x0001
- #define **RTDM_PURGE_TX_BUFFER** 0x0002

Common IOCTLs

The following IOCTLs are common to all device profiles.

- #define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_-
  device_info)

    Retrieve information about a device or socket.

- #define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)

    Purge internal device or socket buffers.

Typedefs

- typedef uint64_t nanosecs_abs_t

    RTDM type for representing absolute dates.

- typedef int64_t nanosecs_rel_t

    RTDM type for representing relative intervals.

- typedef struct rtdm_device_info rtdm_device_info_t

    Device information.

7.2.1 Detailed Description

Real-Time Driver Model for Xenomai, user API header.

Note

> Copyright (C) 2005, 2006 Jan Kiszka <jan.kiszka@web.de>
> Copyright (C) 2005 Joerg Langenberg <joerg.langenberg@gmx.net>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.3 include/rtdm/rtdm_driver.h File Reference

Real-Time Driver Model for Xenomai, driver API header.

Include dependency graph for rtdm_driver.h:

This graph shows which files directly or indirectly include this file:

Data Structures

- struct rtdm_operations

    Device operations.

- struct rtdm_dev_context

    Device context.

- struct rtdm_device

    RTDM device.

Defines

- #define rtdm_irq_get_arg(irq_handle, type) ((type ∗)irq_handle->cookie)

    Retrieve IRQ handler argument.

Device Flags

Static flags describing a RTDM device

- #define RTDM_EXCLUSIVE 0x0001

    If set, only a single instance of the device can be requested by an application.

- #define RTDM_NAMED_DEVICE 0x0010

    If set, the device is addressed via a clear-text name.

- #define RTDM_PROTOCOL_DEVICE 0x0020

    If set, the device is addressed via a combination of protocol ID and socket type.

- #define RTDM_DEVICE_TYPE_MASK 0x00F0

    Mask selecting the device type.

Context Flags

Dynamic flags describing the state of an open RTDM device (bit numbers)

- #define RTDM_CREATED_IN_NRT 0

    Set by RTDM if the device instance was created in non-real-time context.

- #define RTDM_CLOSING 1

    Set by RTDM when the device is being closed.

- #define RTDM_USER_CONTEXT_FLAG 8

    Lowest bit number the driver developer can use freely.

Driver Versioning

Current revisions of RTDM structures, encoding of driver versions. See API Versioning for the interface revision.

- #define RTDM_DEVICE_STRUCT_VER 5

Version of struct rtdm_device.

- #define RTDM_CONTEXT_STRUCT_VER 3

  Version of struct rtdm_dev_context.

- #define RTDM_SECURE_DEVICE 0x80000000

  Flag indicating a secure variant of RTDM (not supported here).

- #define RTDM_DRIVER_VER(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))

  Version code constructor for driver revisions.

- #define RTDM_DRIVER_MAJOR_VER(ver) (((ver) >> 16) & 0xFF)

  Get major version number from driver revision code.

- #define RTDM_DRIVER_MINOR_VER(ver) (((ver) >> 8) & 0xFF)

  Get minor version number from driver revision code.

- #define RTDM_DRIVER_PATCH_VER(ver) ((ver) & 0xFF)

  Get patch version number from driver revision code.

Global Lock across Scheduler Invocation

- #define RTDM_EXECUTE_ATOMICALLY(code_block)

  Execute code block atomically.

RTDM_IRQTYPE_xxx

Interrupt registrations flags

- #define RTDM_IRQTYPE_SHARED XN_ISR_SHARED

  Enable IRQ-sharing with other real-time drivers.

- #define RTDM_IRQTYPE_EDGE XN_ISR_EDGE

  Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.

RTDM_IRQ_xxx

Return flags of interrupt handlers

- #define RTDM_IRQ_NONE XN_ISR_NONE

  Unhandled interrupt.

- #define RTDM_IRQ_HANDLED XN_ISR_HANDLED

  Denote handled interrupt.

Task Priority Range

Maximum and minimum task priorities

- #define **RTDM_TASK_LOWEST_PRIORITY** XNSCHED_LOW_PRIO
- #define **RTDM_TASK_HIGHEST_PRIORITY** XNSCHED_HIGH_PRIO

Task Priority Modification

Raise or lower task priorities by one level

- #define **RTDM_TASK_RAISE_PRIORITY** (+1)
- #define **RTDM_TASK_LOWER_PRIORITY** (-1)

Typedefs

- typedef int(* rtdm_irq_handler_t)(rtdm_irq_t *irq_handle)

    Interrupt handler.

- typedef void(* rtdm_nrtsig_handler_t)(rtdm_nrtsig_t nrt_sig, void *arg)

    Non-real-time signal handler.

- typedef void(* rtdm_timer_handler_t)(rtdm_timer_t *timer)

    Timer handler.

- typedef void(* rtdm_task_proc_t)(void *arg)

    Real-time task procedure.

Operation Handler Prototypes

- typedef int(* rtdm_open_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, int oflag)

    Named device open handler.

- typedef int(* rtdm_socket_handler_t)(struct rtdm_dev_context *context, rtdm_user_-info_t *user_info, int protocol)

    Socket creation handler for protocol devices.

- typedef int(* rtdm_close_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info)

    Close handler.

- typedef int(* rtdm_ioctl_handler_t)(struct rtdm_dev_context *context, rtdm_user_info_t *user_info, unsigned int request, void __user *arg)

    IOCTL handler.

- typedef int(* rtdm_select_bind_handler_t)(struct rtdm_dev_context *context, rtdm_-selector_t *selector, enum rtdm_selecttype type, unsigned fd_index)

    Select binding handler.

- typedef ssize_t(* rtdm_read_handler_t)(struct rtdm_dev_context *context, rtdm_user_-info_t *user_info, void *buf, size_t nbyte)

    Read handler.

- typedef ssize_t(* rtdm_write_handler_t)(struct rtdm_dev_context *context, rtdm_user_-info_t *user_info, const void *buf, size_t nbyte)

    Write handler.

- typedef ssize_t(* rtdm_recvmsg_handler_t)(struct rtdm_dev_context *context, rtdm_-user_info_t *user_info, struct msghdr *msg, int flags)

Receive message handler.

- typedef ssize_t(∗ rtdm_sendmsg_handler_t)(struct rtdm_dev_context ∗context, rtdm_-
  user_info_t ∗user_info, const struct msghdr ∗msg, int flags)

  Transmit message handler.

Enumerations

RTDM_SELECTTYPE_xxx

Event types select can bind to

- enum rtdm_selecttype { RTDM_SELECTTYPE_READ = XNSELECT_READ, RTDM_-
  SELECTTYPE_WRITE = XNSELECT_WRITE, RTDM_SELECTTYPE_EXCEPT =
  XNSELECT_EXCEPT }

RTDM_TIMERMODE_xxx

Timer operation modes

- enum rtdm_timer_mode { RTDM_TIMERMODE_RELATIVE = XN_RELATIVE, RTDM_-
  TIMERMODE_ABSOLUTE = XN_ABSOLUTE, RTDM_TIMERMODE_REALTIME =
  XN_REALTIME }

Functions

- static void ∗ rtdm_context_to_private (struct rtdm_dev_context ∗context)

  Locate the driver private area associated to a device context structure.

- static struct rtdm_dev_context ∗ rtdm_private_to_context (void ∗dev_private)

  Locate a device context structure from its driver private area.

- int rtdm_dev_register (struct rtdm_device ∗device)

  Register a RTDM device.

- int rtdm_dev_unregister (struct rtdm_device ∗device, unsigned int poll_delay)

  Unregisters a RTDM device.

- struct rtdm_dev_context ∗ rtdm_context_get (int fd)

  Retrieve and lock a device context.

- int rtdm_select_bind (int fd, rtdm_selector_t ∗selector, enum rtdm_selecttype type, unsigned
  fd_index)

  Bind a selector to specified event types of a given file descriptor.

- int rtdm_irq_request (rtdm_irq_t ∗irq_handle, unsigned int irq_no, rtdm_irq_handler_t han-
  dler, unsigned long flags, const char ∗device_name, void ∗arg)

  Register an interrupt handler.

- void rtdm_timer_destroy (rtdm_timer_t ∗timer)

  Destroy a timer.

- int rtdm_timer_start (rtdm_timer_t *timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)

  Start a timer.

- void rtdm_timer_stop (rtdm_timer_t *timer)

  Stop a timer.

- int rtdm_task_init (rtdm_task_t *task, const char *name, rtdm_task_proc_t task_proc, void *arg, int priority, nanosecs_rel_t period)

  Intialise and start a real-time task.

- void rtdm_task_busy_sleep (nanosecs_rel_t delay)

  Busy-wait a specified amount of time.

- void rtdm_toseq_init (rtdm_toseq_t *timeout_seq, nanosecs_rel_t timeout)

  Initialise a timeout sequence.

- void rtdm_event_init (rtdm_event_t *event, unsigned long pending)

  Initialise an event.

- int rtdm_event_select_bind (rtdm_event_t *event, rtdm_selector_t *selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to an event.

- int rtdm_event_wait (rtdm_event_t *event)

  Wait on event occurrence.

- int rtdm_event_timedwait (rtdm_event_t *event, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)

  Wait on event occurrence with timeout.

- void rtdm_event_signal (rtdm_event_t *event)

  Signal an event occurrence.

- void rtdm_event_clear (rtdm_event_t *event)

  Clear event state.

- void rtdm_sem_init (rtdm_sem_t *sem, unsigned long value)

  Initialise a semaphore.

- int rtdm_sem_select_bind (rtdm_sem_t *sem, rtdm_selector_t *selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to a semaphore.

- int rtdm_sem_down (rtdm_sem_t *sem)

  Decrement a semaphore.

- int rtdm_sem_timeddown (rtdm_sem_t *sem, nanosecs_rel_t timeout, rtdm_toseq_-t *timeout_seq)

Decrement a semaphore with timeout.

- void rtdm_sem_up (rtdm_sem_t *sem)

  Increment a semaphore.

- void rtdm_mutex_init (rtdm_mutex_t *mutex)

  Initialise a mutex.

- int rtdm_mutex_lock (rtdm_mutex_t *mutex)

  Request a mutex.

- int rtdm_mutex_timedlock (rtdm_mutex_t *mutex, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)

  Request a mutex with timeout.

Spinlock with Preemption Deactivation

- #define RTDM_LOCK_UNLOCKED RTHAL_SPIN_LOCK_UNLOCKED

  Static lock initialisation.

- #define rtdm_lock_init(lock) rthal_spin_lock_init(lock)

  Dynamic lock initialisation.

- #define rtdm_lock_get(lock) rthal_spin_lock(lock)

  Acquire lock from non-preemptible contexts.

- #define rtdm_lock_put(lock) rthal_spin_unlock(lock)

  Release lock without preemption restoration.

- #define rtdm_lock_get_irqsave(lock, context) rthal_spin_lock_irqsave(lock, context)

  Acquire lock and disable preemption.

- #define rtdm_lock_put_irqrestore(lock, context) rthal_spin_unlock_irqrestore(lock, context)

  Release lock and restore preemption state.

- #define rtdm_lock_irqsave(context) rthal_local_irq_save(context)

  Disable preemption locally.

- #define rtdm_lock_irqrestore(context) rthal_local_irq_restore(context)

  Restore preemption state.

- typedef rthal_spinlock_t rtdm_lock_t

  Lock variable.

- typedef unsigned long rtdm_lockctx_t

  Variable to save the context while holding a lock.

7.3.1 Detailed Description

Real-Time Driver Model for Xenomai, driver API header.

Note

Copyright (C) 2005-2007 Jan Kiszka <jan.kiszka@web.de>
Copyright (C) 2005 Joerg Langenberg <joerg.langenberg@gmx.net>
Copyright (C) 2008 Gilles Chanteperdrix <gilles.chanteperdrix@xenomai.org>

7.4 include/rtdm/rtipc.h File Reference

This file is part of the Xenomai project.

Include dependency graph for rtipc.h:

Data Structures

- struct rtipc_port_label

    Port label information structure.

- struct sockaddr_ipc

    Socket address structure for the RTIPC address family.

Defines

XDDP socket options

Setting and getting XDDP socket options.

- #define XDDP_LABEL 1
    XDDP label assignment.

- #define XDDP_POOLSZ 2
    XDDP local pool size configuration.

- #define XDDP_BUFSZ 3
    XDDP streaming buffer size configuration.

- #define XDDP_MONITOR 4
    XDDP monitoring callback.

XDDP events

Specific events occurring on XDDP channels, which can be monitored via the XDDP_MONITOR socket option.

- #define XDDP_EVTIN 1
    Monitor writes to the non real-time endpoint.

- #define XDDP_EVTOUT 2
    Monitor reads from the non real-time endpoint.

- #define XDDP_EVTDOWN 3
    Monitor close from the non real-time endpoint.

- #define XDDP_EVTNOBUF 4
    Monitor memory shortage for non real-time datagrams.

IDDP socket options

Setting and getting IDDP socket options.

- #define IDDP_LABEL 1
    IDDP label assignment.

- #define IDDP_POOLSZ 2
    IDDP local pool size configuration.

BUFP socket options

Setting and getting BUFP socket options.

- #define BUFP_LABEL 1
    BUFP label assignment.

- #define BUFP_BUFSZ 2
    BUFP buffer size configuration.

Socket level options

Setting and getting supported standard socket level options.

- #define SO_SNDTIMEO defined_by_kernel_header_file

  IPCPROTO_IDDP and IPCPROTO_BUFP protocols support the standard SO_SNDTIMEO socket option, from the `SOL_SOCKET` *level.*

- #define SO_RCVTIMEO defined_by_kernel_header_file

  All RTIPC protocols support the standard SO_RCVTIMEO socket option, from the `SOL_SOCKET` *level.*

Typedefs

- typedef int16_t rtipc_port_t

  Port number type for the RTIPC address family.

Enumerations

RTIPC protocol list

protocols for the PF_RTIPC protocol family

- enum { IPCPROTO_IPC = 0, IPCPROTO_XDDP = 1, IPCPROTO_IDDP = 2, IPCPROTO_-BUFP = 3 }

Functions

Supported operations

Standard socket operations supported by the RTIPC protocols.

- int socket__AF_RTIPC (int domain=AF_RTIPC, int type=SOCK_DGRAM, int protocol)

  Create an endpoint for communication in the AF_RTIPC domain.

- int close__AF_RTIPC (int sockfd)

  Close a RTIPC socket descriptor.

- int bind__AF_RTIPC (int sockfd, const struct sockaddr_ipc *addr, socklen_t addrlen)

  Bind a RTIPC socket to a port.

- int connect__AF_RTIPC (int sockfd, const struct sockaddr_ipc *addr, socklen_t addrlen)

  Initiate a connection on a RTIPC socket.

- int setsockopt__AF_RTIPC (int sockfd, int level, int optname, const void *optval, socklen_t optlen)

  Set options on RTIPC sockets.

- int getsockopt__AF_RTIPC (int sockfd, int level, int optname, void *optval, socklen_t *optlen)

  Get options on RTIPC sockets.

- ssize_t sendmsg__AF_RTIPC (int sockfd, const struct msghdr *msg, int flags)

  Send a message on a RTIPC socket.

- ssize_t recvmsg__AF_RTIPC (int sockfd, struct msghdr *msg, int flags)

    Receive a message from a RTIPC socket.

- int getsockname__AF_RTIPC (int sockfd, struct sockaddr_ipc *addr, socklen_t *addrlen)

    Get socket name.

- int getpeername__AF_RTIPC (int sockfd, struct sockaddr_ipc *addr, socklen_t *addrlen)

    Get socket peer.

7.4.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2009 Philippe Gerum <rpm@xenomai.org>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.5 include/rtdm/rtserial.h File Reference

Real-Time Driver Model for Xenomai, serial device profile header.

Include dependency graph for rtserial.h:

Data Structures

- struct [rtser_config](#)

  Serial device configuration.

- struct [rtser_status](#)

  Serial device status.

- struct [rtser_event](#)

  Additional information about serial device events.

Defines

- #define [RTSER_RTIOC_BREAK_CTL](#) _IOR(RTIOC_TYPE_SERIAL, 0x06, int)

  Set or clear break on UART output line.

RTSER_DEF_BAUD

Default baud rate

- #define **RTSER_DEF_BAUD** 9600

RTSER_xxx_PARITY

Number of parity bits

- #define **RTSER_NO_PARITY** 0x00
- #define **RTSER_ODD_PARITY** 0x01
- #define **RTSER_EVEN_PARITY** 0x03
- #define **RTSER_DEF_PARITY** RTSER_NO_PARITY

RTSER_xxx_BITS

Number of data bits

- #define **RTSER_5_BITS** 0x00
- #define **RTSER_6_BITS** 0x01
- #define **RTSER_7_BITS** 0x02
- #define **RTSER_8_BITS** 0x03
- #define **RTSER_DEF_BITS** RTSER_8_BITS

RTSER_xxx_STOPB

Number of stop bits

- #define [RTSER_1_STOPB](#) 0x00

  valid only in combination with 5 data bits

- #define [RTSER_1_5_STOPB](#) 0x01

  valid only in combination with 5 data bits

- #define [RTSER_2_STOPB](#) 0x01

  valid only in combination with 5 data bits

- #define RTSER_DEF_STOPB RTSER_1_STOPB

  valid only in combination with 5 data bits

RTSER_xxx_HAND

Handshake mechanisms

- #define **RTSER_NO_HAND** 0x00
- #define **RTSER_RTSCTS_HAND** 0x01
- #define **RTSER_DEF_HAND** RTSER_NO_HAND

RTSER_FIFO_xxx

Reception FIFO interrupt threshold

- #define **RTSER_FIFO_DEPTH_1** 0x00
- #define **RTSER_FIFO_DEPTH_4** 0x40
- #define **RTSER_FIFO_DEPTH_8** 0x80
- #define **RTSER_FIFO_DEPTH_14** 0xC0
- #define **RTSER_DEF_FIFO_DEPTH** RTSER_FIFO_DEPTH_1

RTSER_TIMEOUT_xxx

Special timeout values, see also RTDM_TIMEOUT_xxx

- #define **RTSER_TIMEOUT_INFINITE** RTDM_TIMEOUT_INFINITE
- #define **RTSER_TIMEOUT_NONE** RTDM_TIMEOUT_NONE
- #define **RTSER_DEF_TIMEOUT** RTDM_TIMEOUT_INFINITE

RTSER_xxx_TIMESTAMP_HISTORY

Timestamp history control

- #define **RTSER_RX_TIMESTAMP_HISTORY** 0x01
- #define **RTSER_DEF_TIMESTAMP_HISTORY** 0x00

RTSER_EVENT_xxx

Events bits

- #define **RTSER_EVENT_RXPEND** 0x01
- #define **RTSER_EVENT_ERRPEND** 0x02
- #define **RTSER_EVENT_MODEMHI** 0x04
- #define **RTSER_EVENT_MODEMLO** 0x08
- #define **RTSER_DEF_EVENT_MASK** 0x00

RTSER_SET_xxx

Configuration mask bits

- #define **RTSER_SET_BAUD** 0x0001
- #define **RTSER_SET_PARITY** 0x0002
- #define **RTSER_SET_DATA_BITS** 0x0004
- #define **RTSER_SET_STOP_BITS** 0x0008
- #define **RTSER_SET_HANDSHAKE** 0x0010
- #define **RTSER_SET_FIFO_DEPTH** 0x0020
- #define **RTSER_SET_TIMEOUT_RX** 0x0100
- #define **RTSER_SET_TIMEOUT_TX** 0x0200
- #define **RTSER_SET_TIMEOUT_EVENT** 0x0400
- #define **RTSER_SET_TIMESTAMP_HISTORY** 0x0800

- #define **RTSER_SET_EVENT_MASK** 0x1000

RTSER_LSR_xxx

Line status bits

- #define **RTSER_LSR_DATA** 0x01
- #define **RTSER_LSR_OVERRUN_ERR** 0x02
- #define **RTSER_LSR_PARITY_ERR** 0x04
- #define **RTSER_LSR_FRAMING_ERR** 0x08
- #define **RTSER_LSR_BREAK_IND** 0x10
- #define **RTSER_LSR_THR_EMTPY** 0x20
- #define **RTSER_LSR_TRANSM_EMPTY** 0x40
- #define **RTSER_LSR_FIFO_ERR** 0x80
- #define **RTSER_SOFT_OVERRUN_ERR** 0x0100

RTSER_MSR_xxx

Modem status bits

- #define **RTSER_MSR_DCTS** 0x01
- #define **RTSER_MSR_DDSR** 0x02
- #define **RTSER_MSR_TERI** 0x04
- #define **RTSER_MSR_DDCD** 0x08
- #define **RTSER_MSR_CTS** 0x10
- #define **RTSER_MSR_DSR** 0x20
- #define **RTSER_MSR_RI** 0x40
- #define **RTSER_MSR_DCD** 0x80

RTSER_MCR_xxx

Modem control bits

- #define **RTSER_MCR_DTR** 0x01
- #define **RTSER_MCR_RTS** 0x02
- #define **RTSER_MCR_OUT1** 0x04
- #define **RTSER_MCR_OUT2** 0x08
- #define **RTSER_MCR_LOOP** 0x10

Sub-Classes of RTDM_CLASS_SERIAL

- #define **RTDM_SUBCLASS_16550A** 0

IOCTLs

Serial device IOCTLs

- #define RTSER_RTIOC_GET_CONFIG _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_-config)

    Get serial device configuration.

- #define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_-config)

    Set serial device configuration.

- #define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_-status)

    Get serial device status.

- #define RTSER_RTIOC_GET_CONTROL _IOR(RTIOC_TYPE_SERIAL, 0x03, int)

  Get serial device's modem contol register.

- #define RTSER_RTIOC_SET_CONTROL _IOW(RTIOC_TYPE_SERIAL, 0x04, int)

  Set serial device's modem contol register.

- #define RTSER_RTIOC_WAIT_EVENT _IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_-event)

  Wait on serial device events according to previously set mask.

RTSER_BREAK_xxx

Break control

- #define RTSER_BREAK_CLR 0x00

  Serial device configuration.

- #define RTSER_BREAK_SET 0x01

  Serial device configuration.

- #define RTIOC_TYPE_SERIAL RTDM_CLASS_SERIAL

  Serial device configuration.

- typedef struct rtser_config rtser_config_t

  Serial device configuration.

- typedef struct rtser_status rtser_status_t

  Serial device status.

- typedef struct rtser_event rtser_event_t

  Additional information about serial device events.

7.5.1 Detailed Description

Real-Time Driver Model for Xenomai, serial device profile header.

Note

Copyright (C) 2005-2007 Jan Kiszka <jan.kiszka@web.de>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.6 include/rtdm/rttesting.h File Reference

Real-Time Driver Model for Xenomai, testing device profile header.

Include dependency graph for rttesting.h:

Defines

Sub-Classes of RTDM_CLASS_TESTING

- #define RTDM_SUBCLASS_TIMERBENCH 0

  subclass name: "timerbench"

- #define RTDM_SUBCLASS_IRQBENCH 1

  subclass name: "irqbench"

- #define RTDM_SUBCLASS_SWITCHTEST 2

  subclass name: "switchtest"

- #define RTDM_SUBCLASS_RTDMTEST 3

  subclase name: "rtdm"

IOCTLs

Testing device IOCTLs

- #define **RTTST_RTIOC_INTERM_BENCH_RES** _IOWR(RTIOC_TYPE_TESTING, 0x00, struct rttst_interm_bench_res)
- #define **RTTST_RTIOC_TMBENCH_START** _IOW(RTIOC_TYPE_TESTING, 0x10, struct rttst_tmbench_config)
- #define **RTTST_RTIOC_TMBENCH_STOP** _IOWR(RTIOC_TYPE_TESTING, 0x11, struct rttst_overall_bench_res)
- #define **RTTST_RTIOC_IRQBENCH_START** _IOW(RTIOC_TYPE_TESTING, 0x20, struct rttst_irqbench_config)
- #define **RTTST_RTIOC_IRQBENCH_STOP** _IO(RTIOC_TYPE_TESTING, 0x21)
- #define **RTTST_RTIOC_IRQBENCH_GET_STATS** _IOR(RTIOC_TYPE_TESTING, 0x22, struct rttst_irqbench_stats)

- #define **RTTST_RTIOC_IRQBENCH_WAIT_IRQ** _IO(RTIOC_TYPE_TESTING, 0x23)
- #define **RTTST_RTIOC_IRQBENCH_REPLY_IRQ** _IO(RTIOC_TYPE_TESTING, 0x24)
- #define **RTTST_RTIOC_SWTEST_SET_TASKS_COUNT** _IOW(RTIOC_TYPE_-TESTING, 0x30, unsigned long)
- #define **RTTST_RTIOC_SWTEST_SET_CPU** _IOW(RTIOC_TYPE_TESTING, 0x31, unsigned long)
- #define **RTTST_RTIOC_SWTEST_REGISTER_UTASK** _IOW(RTIOC_TYPE_-TESTING, 0x32, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_CREATE_KTASK** _IOWR(RTIOC_TYPE_TESTING, 0x33, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_PEND** _IOR(RTIOC_TYPE_TESTING, 0x34, struct rttst_swtest_task)
- #define **RTTST_RTIOC_SWTEST_SWITCH_TO** _IOR(RTIOC_TYPE_TESTING, 0x35, struct rttst_swtest_dir)
- #define **RTTST_RTIOC_SWTEST_GET_SWITCHES_COUNT** _IOR(RTIOC_TYPE_-TESTING, 0x36, unsigned long)
- #define **RTTST_RTIOC_SWTEST_GET_LAST_ERROR** _IOR(RTIOC_TYPE_-TESTING, 0x37, struct rttst_swtest_error)
- #define **RTTST_RTIOC_SWTEST_SET_PAUSE** _IOW(RTIOC_TYPE_TESTING, 0x38, unsigned long)
- #define **RTTST_RTIOC_RTDM_DEFER_CLOSE** _IOW(RTIOC_TYPE_TESTING, 0x40, unsigned long)

7.6.1 Detailed Description

Real-Time Driver Model for Xenomai, testing device profile header.

Note

>   Copyright (C) 2005 Jan Kiszka <jan.kiszka@web.de>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.7 ksrc/skins/rtdm/device.c File Reference

Real-Time Driver Model for Xenomai, device management.

Include dependency graph for device.c:

Functions

- int rtdm_dev_register (struct rtdm_device ∗device)

  Register a RTDM device.

- int rtdm_dev_unregister (struct rtdm_device ∗device, unsigned int poll_delay)

  Unregisters a RTDM device.

7.7.1 Detailed Description

Real-Time Driver Model for Xenomai, device management.

Note

> Copyright (C) 2005 Jan Kiszka <`jan.kiszka@web.de`>
> Copyright (C) 2005 Joerg Langenberg <`joerg.langenberg@gmx.net`>

7.8 ksrc/skins/rtdm/drvlib.c File Reference

Real-Time Driver Model for Xenomai, driver library.

Include dependency graph for drvlib.c:

Functions

- [nanosecs_abs_t rtdm_clock_read](#) (void)

  Get system time.

- [nanosecs_abs_t rtdm_clock_read_monotonic](#) (void)

  Get monotonic time.

- int [rtdm_task_init](#) (rtdm_task_t ∗task, const char ∗name, [rtdm_task_proc_t](#) task_proc, void ∗arg, int priority, [nanosecs_rel_t](#) period)

  Intialise and start a real-time task.

- void [rtdm_task_destroy](#) (rtdm_task_t ∗task)

  Destroy a real-time task.

- void [rtdm_task_set_priority](#) (rtdm_task_t ∗task, int priority)

  Adjust real-time task priority.

- int [rtdm_task_set_period](#) (rtdm_task_t ∗task, [nanosecs_rel_t](#) period)

  Adjust real-time task period.

- int [rtdm_task_wait_period](#) (void)

  Wait on next real-time task period.

- int [rtdm_task_unblock](#) (rtdm_task_t ∗task)

  Activate a blocked real-time task.

- rtdm_task_t ∗ [rtdm_task_current](#) (void)

  Get current real-time task.

- int [rtdm_task_sleep](#) ([nanosecs_rel_t](#) delay)

  Sleep a specified amount of time.

- int [rtdm_task_sleep_until](#) ([nanosecs_abs_t](#) wakeup_time)

  Sleep until a specified absolute time.

- int [rtdm_task_sleep_abs](#) ([nanosecs_abs_t](#) wakeup_time, enum [rtdm_timer_mode](#) mode)

  Sleep until a specified absolute time.

- void [rtdm_task_join_nrt](#) (rtdm_task_t ∗task, unsigned int poll_delay)

Wait on a real-time task to terminate.

- void rtdm_task_busy_sleep (nanosecs_rel_t delay)

  Busy-wait a specified amount of time.

- int rtdm_timer_init (rtdm_timer_t *timer, rtdm_timer_handler_t handler, const char *name)

  Initialise a timer.

- void rtdm_timer_destroy (rtdm_timer_t *timer)

  Destroy a timer.

- int rtdm_timer_start (rtdm_timer_t *timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)

  Start a timer.

- void rtdm_timer_stop (rtdm_timer_t *timer)

  Stop a timer.

- int rtdm_timer_start_in_handler (rtdm_timer_t *timer, nanosecs_abs_t expiry, nanosecs_-rel_t interval, enum rtdm_timer_mode mode)

  Start a timer from inside a timer handler.

- void rtdm_timer_stop_in_handler (rtdm_timer_t *timer)

  Stop a timer from inside a timer handler.

- int rtdm_irq_request (rtdm_irq_t *irq_handle, unsigned int irq_no, rtdm_irq_handler_t handler, unsigned long flags, const char *device_name, void *arg)

  Register an interrupt handler.

- int rtdm_irq_free (rtdm_irq_t *irq_handle)

  Release an interrupt handler.

- int rtdm_irq_enable (rtdm_irq_t *irq_handle)

  Enable interrupt line.

- int rtdm_irq_disable (rtdm_irq_t *irq_handle)

  Disable interrupt line.

- int rtdm_nrtsig_init (rtdm_nrtsig_t *nrt_sig, rtdm_nrtsig_handler_t handler, void *arg)

  Register a non-real-time signal handler.

- void rtdm_nrtsig_destroy (rtdm_nrtsig_t *nrt_sig)

  Release a non-realtime signal handler.

- void rtdm_nrtsig_pend (rtdm_nrtsig_t *nrt_sig)

  Trigger non-real-time signal.

- int rtdm_mmap_to_user (rtdm_user_info_t *user_info, void *src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)

Map a kernel memory range into the address space of the user.

- int rtdm_iomap_to_user (rtdm_user_info_t *user_info, phys_addr_t src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)

  Map an I/O memory range into the address space of the user.

- int rtdm_munmap (rtdm_user_info_t *user_info, void *ptr, size_t len)

  Unmap a user memory range.

- void rtdm_printk (const char *format,...)

  Real-time safe message printing on kernel console.

- void * rtdm_malloc (size_t size)

  Allocate memory block in real-time context.

- void rtdm_free (void *ptr)

  Release real-time memory block.

- int rtdm_read_user_ok (rtdm_user_info_t *user_info, const void __user *ptr, size_t size)

  Check if read access to user-space memory block is safe.

- int rtdm_rw_user_ok (rtdm_user_info_t *user_info, const void __user *ptr, size_t size)

  Check if read/write access to user-space memory block is safe.

- int rtdm_copy_from_user (rtdm_user_info_t *user_info, void *dst, const void __user *src, size_t size)

  Copy user-space memory block to specified buffer.

- int rtdm_safe_copy_from_user (rtdm_user_info_t *user_info, void *dst, const void __user *src, size_t size)

  Check if read access to user-space memory block and copy it to specified buffer.

- int rtdm_copy_to_user (rtdm_user_info_t *user_info, void __user *dst, const void *src, size_t size)

  Copy specified buffer to user-space memory block.

- int rtdm_safe_copy_to_user (rtdm_user_info_t *user_info, void __user *dst, const void *src, size_t size)

  Check if read/write access to user-space memory block is safe and copy specified buffer to it.

- int rtdm_strncpy_from_user (rtdm_user_info_t *user_info, char *dst, const char __user *src, size_t count)

  Copy user-space string to specified buffer.

- int rtdm_in_rt_context (void)

  Test if running in a real-time task.

- int rtdm_rt_capable (rtdm_user_info_t *user_info)

  Test if the caller is capable of running in real-time context.

Timeout Sequence Management

- void rtdm_toseq_init (rtdm_toseq_t *timeout_seq, nanosecs_rel_t timeout)

  Initialise a timeout sequence.

- EXPORT_SYMBOL_GPL (rtdm_toseq_init)

  Initialise a timeout sequence.

Event Services

- void rtdm_event_init (rtdm_event_t *event, unsigned long pending)

  Initialise an event.

- EXPORT_SYMBOL_GPL (rtdm_event_init)

  Initialise an event.

- void rtdm_event_destroy (rtdm_event_t *event)

  Destroy an event.

- void rtdm_event_pulse (rtdm_event_t *event)

  Signal an event occurrence to currently listening waiters.

- void rtdm_event_signal (rtdm_event_t *event)

  Signal an event occurrence.

- EXPORT_SYMBOL_GPL (rtdm_event_signal)

  Initialise an event.

- int rtdm_event_wait (rtdm_event_t *event)

  Wait on event occurrence.

- EXPORT_SYMBOL_GPL (rtdm_event_wait)

  Initialise an event.

- int rtdm_event_timedwait (rtdm_event_t *event, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)

  Wait on event occurrence with timeout.

- EXPORT_SYMBOL_GPL (rtdm_event_timedwait)

  Initialise an event.

- void rtdm_event_clear (rtdm_event_t *event)

  Clear event state.

- EXPORT_SYMBOL_GPL (rtdm_event_clear)

  Initialise an event.

- int rtdm_event_select_bind (rtdm_event_t *event, rtdm_selector_t *selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to an event.

- EXPORT_SYMBOL_GPL (rtdm_event_select_bind)

  Initialise an event.

Semaphore Services

- void rtdm_sem_init (rtdm_sem_t ∗sem, unsigned long value)

  Initialise a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_init)

  Initialise a semaphore.

- void rtdm_sem_destroy (rtdm_sem_t ∗sem)

  Destroy a semaphore.

- int rtdm_sem_down (rtdm_sem_t ∗sem)

  Decrement a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_down)

  Initialise a semaphore.

- int rtdm_sem_timeddown (rtdm_sem_t ∗sem, nanosecs_rel_t timeout, rtdm_toseq_t ∗timeout_seq)

  Decrement a semaphore with timeout.

- EXPORT_SYMBOL_GPL (rtdm_sem_timeddown)

  Initialise a semaphore.

- void rtdm_sem_up (rtdm_sem_t ∗sem)

  Increment a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_up)

  Initialise a semaphore.

- int rtdm_sem_select_bind (rtdm_sem_t ∗sem, rtdm_selector_t ∗selector, enum rtdm_-selecttype type, unsigned fd_index)

  Bind a selector to a semaphore.

- EXPORT_SYMBOL_GPL (rtdm_sem_select_bind)

  Initialise a semaphore.

Mutex Services

- void rtdm_mutex_init (rtdm_mutex_t ∗mutex)

  Initialise a mutex.

- EXPORT_SYMBOL_GPL (rtdm_mutex_init)

  Initialise a mutex.

- void rtdm_mutex_destroy (rtdm_mutex_t ∗mutex)

  Destroy a mutex.

- void rtdm_mutex_unlock (rtdm_mutex_t ∗mutex)

  Release a mutex.

- int rtdm_mutex_lock (rtdm_mutex_t ∗mutex)

  Request a mutex.

- EXPORT_SYMBOL_GPL (rtdm_mutex_lock)

    Initialise a mutex.

- int rtdm_mutex_timedlock (rtdm_mutex_t ∗mutex, nanosecs_rel_t timeout, rtdm_toseq_t ∗timeout_seq)

    Request a mutex with timeout.

- EXPORT_SYMBOL_GPL (rtdm_mutex_timedlock)

    Initialise a mutex.

7.8.1 Detailed Description

Real-Time Driver Model for Xenomai, driver library.

Note

Copyright (C) 2005-2007 Jan Kiszka <jan.kiszka@web.de>
Copyright (C) 2005 Joerg Langenberg <joerg.langenberg@gmx.net>
Copyright (C) 2008 Gilles Chanteperdrix <gilles.chanteperdrix@xenomai.org>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.9 ksrc/skins/rtdm/module.c File Reference

Real-Time Driver Model for Xenomai.

Include dependency graph for module.c:

7.9.1 Detailed Description

Real-Time Driver Model for Xenomai.

Note

> Copyright (C) 2005, 2006 Jan Kiszka <jan.kiszka@web.de>
> Copyright (C) 2005 Joerg Langenberg <joerg.langenberg@gmx.net>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.10 ksrc/skins/rtdm/core.c File Reference

Real-Time Driver Model for Xenomai, device operation multiplexing.

Include dependency graph for core.c:

Functions

- struct rtdm_dev_context ∗ rtdm_context_get (int fd)

  Retrieve and lock a device context.

- int rtdm_select_bind (int fd, rtdm_selector_t ∗selector, enum rtdm_selecttype type, unsigned fd_index)

  Bind a selector to specified event types of a given file descriptor.

- void rtdm_context_lock (struct rtdm_dev_context ∗context)

  Increment context reference counter.

- void rtdm_context_unlock (struct rtdm_dev_context ∗context)

Decrement context reference counter.

- void rtdm_context_put (struct rtdm_dev_context *context)

  Release a device context obtained via rtdm_context_get().

- int rtdm_open (const char *path, int oflag,...)

  Open a device.

- int rtdm_socket (int protocol_family, int socket_type, int protocol)

  Create a socket.

- int rtdm_close (int fd)

  Close a device or socket.

- int rtdm_ioctl (int fd, int request,...)

  Issue an IOCTL.

- ssize_t rtdm_read (int fd, void *buf, size_t nbyte)

  Read from device.

- ssize_t rtdm_write (int fd, const void *buf, size_t nbyte)

  Write to device.

- ssize_t rtdm_recvmsg (int fd, struct msghdr *msg, int flags)

  Receive message from socket.

- ssize_t rtdm_recvfrom (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)

  Receive message from socket.

- ssize_t rtdm_recv (int fd, void *buf, size_t len, int flags)

  Receive message from socket.

- ssize_t rtdm_sendmsg (int fd, const struct msghdr *msg, int flags)

  Transmit message to socket.

- ssize_t rtdm_sendto (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)

  Transmit message to socket.

- ssize_t rtdm_send (int fd, const void *buf, size_t len, int flags)

  Transmit message to socket.

- int rtdm_bind (int fd, const struct sockaddr *my_addr, socklen_t addrlen)

  Bind to local address.

- int rtdm_connect (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)

  Connect to remote address.

- int rtdm_listen (int fd, int backlog)

Listen for incomming connection requests.

- int rtdm_accept (int fd, struct sockaddr ∗addr, socklen_t ∗addrlen)

  Accept a connection requests.

- int rtdm_shutdown (int fd, int how)

  Shut down parts of a connection.

- int rtdm_getsockopt (int fd, int level, int optname, void ∗optval, socklen_t ∗optlen)

  Get socket option.

- int rtdm_setsockopt (int fd, int level, int optname, const void ∗optval, socklen_t optlen)

  Set socket option.

- int rtdm_getsockname (int fd, struct sockaddr ∗name, socklen_t ∗namelen)

  Get local socket address.

- int rtdm_getpeername (int fd, struct sockaddr ∗name, socklen_t ∗namelen)

  Get socket destination address.

- int rt_dev_open (const char ∗path, int oflag,...)

  Open a device.

- int rt_dev_socket (int protocol_family, int socket_type, int protocol)

  Create a socket.

- int rt_dev_close (int fd)

  Close a device or socket.

- int rt_dev_ioctl (int fd, int request,...)

  Issue an IOCTL.

- ssize_t rt_dev_read (int fd, void ∗buf, size_t nbyte)

  Read from device.

- ssize_t rt_dev_write (int fd, const void ∗buf, size_t nbyte)

  Write to device.

- ssize_t rt_dev_recvmsg (int fd, struct msghdr ∗msg, int flags)

  Receive message from socket.

- ssize_t rt_dev_recvfrom (int fd, void ∗buf, size_t len, int flags, struct sockaddr ∗from, socklen_t ∗fromlen)

  Receive message from socket.

- ssize_t rt_dev_recv (int fd, void ∗buf, size_t len, int flags)

  Receive message from socket.

- ssize_t rt_dev_sendmsg (int fd, const struct msghdr ∗msg, int flags)

  Transmit message to socket.

- ssize_t rt_dev_sendto (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)

    Transmit message to socket.

- ssize_t rt_dev_send (int fd, const void *buf, size_t len, int flags)

    Transmit message to socket.

- int rt_dev_bind (int fd, const struct sockaddr *my_addr, socklen_t addrlen)

    Bind to local address.

- int rt_dev_connect (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)

    Connect to remote address.

- int rt_dev_listen (int fd, int backlog)

    Listen for incomming connection requests.

- int rt_dev_accept (int fd, struct sockaddr *addr, socklen_t *addrlen)

    Accept a connection requests.

- int rt_dev_shutdown (int fd, int how)

    Shut down parts of a connection.

- int rt_dev_getsockopt (int fd, int level, int optname, void *optval, socklen_t *optlen)

    Get socket option.

- int rt_dev_setsockopt (int fd, int level, int optname, const void *optval, socklen_t optlen)

    Set socket option.

- int rt_dev_getsockname (int fd, struct sockaddr *name, socklen_t *namelen)

    Get local socket address.

- int rt_dev_getpeername (int fd, struct sockaddr *name, socklen_t *namelen)

    Get socket destination address.

7.10.1 Detailed Description

Real-Time Driver Model for Xenomai, device operation multiplexing.

Note

Copyright (C) 2005 Jan Kiszka <jan.kiszka@web.de>
Copyright (C) 2005 Joerg Langenberg <joerg.langenberg@gmx.net>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Chapter 8

Example Documentation

8.1 bufp-label.c

```
/*
 * BUFP-based client/server demo, using the read(2)/write(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created.  A server thread (reader)
 * is bound to a real-time port and receives a stream of bytes sent to
 * this port from a client thread (writer).
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t svtid, cltid;

#define BUFP_PORT_LABEL  "bufp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};
```

```
static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *server(void *arg)
{
        struct rtipc_port_label plabel;
        struct sockaddr_ipc saddr;
        char buf[128];
        size_t bufsz;
        int ret, s;

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUFP);
        if (s < 0)
                fail("socket");

        /*
         * Set a 16k buffer for the server endpoint. This
         * configuration must be done prior to binding the socket to a
         * port.
         */
        bufsz = 16384; /* bytes */
        ret = setsockopt(s, SOL_BUFP, BUFP_BUFSZ,
                        &bufsz, sizeof(bufsz));
        if (ret)
                fail("setsockopt");

        /*
         * Set a port label. This name will be registered when
         * binding, in addition to the port number (if given).
         */
        strcpy(plabel.label, BUFP_PORT_LABEL);
        ret = setsockopt(s, SOL_BUFP, BUFP_LABEL,
                        &plabel, sizeof(plabel));
        if (ret)
                fail("setsockopt");

        /*
         * Bind the socket to the port. Assign that port a label, so
         * that peers may use a descriptive information to locate
         * it. Labeled ports will appear in the
         * /proc/xenomai/registry/rtipc/bufp directory once the socket
         * is bound.
         *
         * saddr.sipc_port specifies the port number to use. If -1 is
         * passed, the BUFP driver will auto-select an idle port.
         */
        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = -1;
        ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("bind");

        for (;;) {
                ret = read(s, buf, sizeof(buf));
                if (ret < 0) {
                        close(s);
                        fail("read");
                }
                rt_printf("%s: received %d bytes, \"%.*s\"\n",
                        __FUNCTION__, ret, ret, buf);
        }

        return NULL;
}
```

```
void *client(void *arg)
{
        struct rtipc_port_label plabel;
        struct sockaddr_ipc svsaddr;
        int ret, s, n = 0, len;
        struct timespec ts;
        char buf[128];

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUFP);
        if (s < 0)
                fail("socket");

        /*
         * Set the port label. This name will be used to find the peer
         * when connecting, instead of the port number. The label must
         * be set _after_ the socket is bound to the port, so that
         * BUFP does not try to register this label for the client
         * port as well (like the server thread did).
         */
        strcpy(plabel.label, BUFP_PORT_LABEL);
        ret = setsockopt(s, SOL_BUFP, BUFP_LABEL,
                         &plabel, sizeof(plabel));
        if (ret)
                fail("setsockopt");

        memset(&svsaddr, 0, sizeof(svsaddr));
        svsaddr.sipc_family = AF_RTIPC;
        svsaddr.sipc_port = -1; /* Tell BUFP to search by label. */
        ret = connect(s, (struct sockaddr *)&svsaddr, sizeof(svsaddr));
        if (ret)
                fail("connect");

        for (;;) {
                len = strlen(msg[n]);
                ret = write(s, msg[n], len);
                if (ret < 0) {
                        close(s);
                        fail("write");
                }
                rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                          __FUNCTION__, ret, ret, msg[n]);
                n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                /*
                 * We run in full real-time mode (i.e. primary mode),
                 * so we have to let the system breathe between two
                 * iterations.
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(svtid);
        pthread_cancel(cltid);
        signal(sig, SIG_DFL);
        pthread_join(svtid, NULL);
        pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
```

```
        struct sched_param svparam = {.sched_priority = 71 };
        struct sched_param clparam = {.sched_priority = 70 };
        pthread_attr_t svattr, clattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary mode.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&svattr);
        pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
        pthread_attr_setschedparam(&svattr, &svparam);

        errno = pthread_create(&svtid, &svattr, &server, NULL);
        if (errno)
                fail("pthread_create");

        pthread_attr_init(&clattr);
        pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
        pthread_attr_setschedparam(&clattr, &clparam);

        errno = pthread_create(&cltid, &clattr, &client, NULL);
        if (errno)
                fail("pthread_create");

        sigsuspend(&oldmask);

        return 0;
}
```

8.2 bufp-readwrite.c

```
/*
 * BUFP-based client/server demo, using the read(2)/write(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created.  A server thread (reader)
 * is bound to a real-time port and receives a stream of bytes sent to
 * this port from a client thread (writer).
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
```

```
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t svtid, cltid;

#define BUFP_SVPORT 12

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *server(void *arg)
{
        struct sockaddr_ipc saddr;
        char buf[128];
        size_t bufsz;
        int ret, s;

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUFP);
        if (s < 0)
                fail("socket");

        /*
         * Set a 16k buffer for the server endpoint. This
         * configuration must be done prior to binding the socket to a
         * port.
         */
        bufsz = 16384; /* bytes */
        ret = setsockopt(s, SOL_BUFP, BUFP_BUFSZ,
                        &bufsz, sizeof(bufsz));
        if (ret)
                fail("setsockopt");

        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = BUFP_SVPORT;
        ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("bind");

        for (;;) {
                ret = read(s, buf, sizeof(buf));
                if (ret < 0) {
                        close(s);
```

```
                                        fail("read");
                        }
                        rt_printf("%s: received %d bytes, \"%.*s\"\n",
                                  __FUNCTION__, ret, ret, buf);
                }

        return NULL;
}

void *client(void *arg)
{
        struct sockaddr_ipc svsaddr;
        int ret, s, n = 0, len;
        struct timespec ts;
        char buf[128];

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUFP);
        if (s < 0)
                fail("socket");

        memset(&svsaddr, 0, sizeof(svsaddr));
        svsaddr.sipc_family = AF_RTIPC;
        svsaddr.sipc_port = BUFP_SVPORT;
        ret = connect(s, (struct sockaddr *)&svsaddr, sizeof(svsaddr));
        if (ret)
                fail("connect");

        for (;;) {
                len = strlen(msg[n]);
                ret = write(s, msg[n], len);
                if (ret < 0) {
                        close(s);
                        fail("write");
                }
                rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                          __FUNCTION__, ret, ret, msg[n]);
                n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                /*
                 * We run in full real-time mode (i.e. primary mode),
                 * so we have to let the system breathe between two
                 * iterations.
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(svtid);
        pthread_cancel(cltid);
        signal(sig, SIG_DFL);
        pthread_join(svtid, NULL);
        pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param svparam = {.sched_priority = 71 };
        struct sched_param clparam = {.sched_priority = 70 };
        pthread_attr_t svattr, clattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);
```

```
        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary mode.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&svattr);
        pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
        pthread_attr_setschedparam(&svattr, &svparam);

        errno = pthread_create(&svtid, &svattr, &server, NULL);
        if (errno)
                fail("pthread_create");

        pthread_attr_init(&clattr);
        pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
        pthread_attr_setschedparam(&clattr, &clparam);

        errno = pthread_create(&cltid, &clattr, &client, NULL);
        if (errno)
                fail("pthread_create");

        sigsuspend(&oldmask);

        return 0;
}
```

8.3 cross-link.c

```
/*
 * cross-link.c
 *
 * Userspace test program (Xenomai native skin) for RTDM-based UART drivers
 * Copyright 2005 by Joerg Langenberg <joergel75@gmx.net>
 *
 * Updates by Jan Kiszka <jan.kiszka@web.de>
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
```

```
 */
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <sys/mman.h>

#include <native/task.h>
#include <native/timer.h>

#include <rtdm/rtserial.h>

#define MAIN_PREFIX   "main : "
#define WTASK_PREFIX  "write_task: "
#define RTASK_PREFIX  "read_task: "

#define WRITE_FILE    "rtser0"
#define READ_FILE     "rtser1"

int read_fd  = -1;
int write_fd = -1;

#define STATE_FILE_OPENED       1
#define STATE_TASK_CREATED      2

unsigned int read_state = 0;
unsigned int write_state = 0;

/*                       --s-ms-us-ns */
RTIME write_task_period_ns =   100000000llu;
RT_TASK write_task;
RT_TASK read_task;

static const struct rtser_config read_config = {
        .config_mask       = 0xFFFF,
        .baud_rate         = 115200,
        .parity            = RTSER_DEF_PARITY,
        .data_bits         = RTSER_DEF_BITS,
        .stop_bits         = RTSER_DEF_STOPB,
        .handshake         = RTSER_DEF_HAND,
        .fifo_depth        = RTSER_DEF_FIFO_DEPTH,
        .rx_timeout        = RTSER_DEF_TIMEOUT,
        .tx_timeout        = RTSER_DEF_TIMEOUT,
        .event_timeout     = 1000000000, /* 1 s */
        .timestamp_history = RTSER_RX_TIMESTAMP_HISTORY,
        .event_mask        = RTSER_EVENT_RXPEND,
};

static const struct rtser_config write_config = {
        .config_mask       = RTSER_SET_BAUD | RTSER_SET_TIMESTAMP_HISTORY,
        .baud_rate         = 115200,
        .timestamp_history = RTSER_DEF_TIMESTAMP_HISTORY,
        /* the rest implicitely remains default */
};

static int close_file( int fd, char *name)
{
        int err, i=0;

        do {
                i++;
                err = rt_dev_close(fd);
                switch (err) {
                case -EAGAIN:
                        printf(MAIN_PREFIX "%s -> EAGAIN (%d times)\n",
                               name, i);
                        rt_task_sleep(50000); /* wait 50us */
                        break;
```

```
                        case 0:
                                printf(MAIN_PREFIX "%s -> closed\n", name);
                                break;
                        default:
                                printf(MAIN_PREFIX "%s -> %s\n", name,
                                        strerror(-err));
                                break;
                        }
                } while (err == -EAGAIN && i < 10);

        return err;
}

void cleanup_all(void)
{
        if (read_state & STATE_FILE_OPENED) {
                close_file(read_fd, READ_FILE" (read)");
                read_state &= ~STATE_FILE_OPENED;
        }

        if (write_state & STATE_FILE_OPENED) {
                close_file(write_fd, WRITE_FILE " (write)");
                write_state &= ~STATE_FILE_OPENED;
        }

        if (write_state & STATE_TASK_CREATED) {
                printf(MAIN_PREFIX "delete write_task\n");
                rt_task_delete(&write_task);
                write_state &= ~STATE_TASK_CREATED;
        }

        if (read_state & STATE_TASK_CREATED) {
                printf(MAIN_PREFIX "delete read_task\n");
                rt_task_delete(&read_task);
                read_state &= ~STATE_TASK_CREATED;
        }
}

void catch_signal(int sig)
{
        cleanup_all();
        printf(MAIN_PREFIX "exit\n");
        return;
}

void write_task_proc(void *arg)
{
        int err;
        RTIME write_time;
        ssize_t sz = sizeof(RTIME);
        ssize_t written = 0;

        err = rt_task_set_periodic(NULL, TM_NOW,
                                rt_timer_ns2ticks(write_task_period_ns));
        if (err) {
                printf(WTASK_PREFIX "error on set periodic, %s\n",
                        strerror(-err));
                goto exit_write_task;
        }

        while (1) {
                err = rt_task_wait_period(NULL);
                if (err) {
                        printf(WTASK_PREFIX
                                "error on rt_task_wait_period, %s\n",
                                strerror(-err));
                        break;
```

```
                          }

                          write_time = rt_timer_read();

                          written = rt_dev_write(write_fd, &write_time, sz);
                          if (written < 0 ) {
                                  printf(WTASK_PREFIX "error on rt_dev_write, %s\n",
                                          strerror(-err));
                                  break;
                          } else if (written != sz) {
                                  printf(WTASK_PREFIX "only %d / %d byte transmitted\n",
                                          written, sz);
                                  break;
                          }
                  }

 exit_write_task:
          if ((write_state & STATE_FILE_OPENED) &&
              close_file(write_fd, WRITE_FILE " (write)") == 0)
                  write_state &= ~STATE_FILE_OPENED;

          printf(WTASK_PREFIX "exit\n");
}

void read_task_proc(void *arg)
{
          int err;
          int nr = 0;
          RTIME read_time  = 0;
          RTIME write_time = 0;
          RTIME irq_time   = 0;
          ssize_t sz = sizeof(RTIME);
          ssize_t read = 0;
          struct rtser_event rx_event;

          printf(" Nr |   write->irq   |   irq->read    |   write->read   |\n");
          printf("-------------------------------------------------------\n");

          /*
           * We are in secondary mode now due to printf, the next
           * blocking Xenomai or driver call will switch us back
           * (here: RTSER_RTIOC_WAIT_EVENT).
           */

          while (1) {
                  /* waiting for event */
                  err = rt_dev_ioctl(read_fd, RTSER_RTIOC_WAIT_EVENT, &rx_event);
                  if (err) {
                          printf(RTASK_PREFIX
                                  "error on RTSER_RTIOC_WAIT_EVENT, %s\n",
                                  strerror(-err));
                          if (err == -ETIMEDOUT)
                                   continue;
                          break;
                  }

                  irq_time = rx_event.rxpend_timestamp;
                  read = rt_dev_read(read_fd, &write_time, sz);
                  if (read == sz) {
                          read_time = rt_timer_read();
                          printf("%3d |%16llu |%16llu |%16llu\n", nr,
                                  irq_time  - write_time,
                                  read_time - irq_time,
                                  read_time - write_time);
                          nr++;
                  } else if (read < 0 ) {
                          printf(RTASK_PREFIX "error on rt_dev_read, code %s\n",
```

```
                                    strerror(-err));
                            break;
                    } else {
                            printf(RTASK_PREFIX "only %d / %d byte received \n",
                                    read, sz);
                            break;
                    }
            }

            if ((read_state & STATE_FILE_OPENED) &&
                close_file(read_fd, READ_FILE " (read)") == 0)
                    read_state &= ~STATE_FILE_OPENED;

            printf(RTASK_PREFIX "exit\n");
    }

    int main(int argc, char* argv[])
    {
            int err = 0;

            signal(SIGTERM, catch_signal);
            signal(SIGINT, catch_signal);

            /* no memory-swapping for this programm */
            mlockall(MCL_CURRENT | MCL_FUTURE);

            /* open rtser0 */
            write_fd = rt_dev_open( WRITE_FILE, 0);
            if (write_fd < 0) {
                    printf(MAIN_PREFIX "can't open %s (write), %s\n", WRITE_FILE,
                            strerror(-write_fd));
                    goto error;
            }
            write_state |= STATE_FILE_OPENED;
            printf(MAIN_PREFIX "write-file opened\n");

            /* writing write-config */
            err = rt_dev_ioctl(write_fd, RTSER_RTIOC_SET_CONFIG, &write_config);
            if (err) {
                    printf(MAIN_PREFIX "error while RTSER_RTIOC_SET_CONFIG, %s\n",
                            strerror(-err));
                    goto error;
            }
            printf(MAIN_PREFIX "write-config written\n");

            /* open rtser1 */
            read_fd = rt_dev_open( READ_FILE, 0 );
            if (read_fd < 0) {
                    printf(MAIN_PREFIX "can't open %s (read), %s\n", READ_FILE,
                            strerror(-read_fd));
                    goto error;
            }
            read_state |= STATE_FILE_OPENED;
            printf(MAIN_PREFIX "read-file opened\n");

            /* writing read-config */
            err = rt_dev_ioctl(read_fd, RTSER_RTIOC_SET_CONFIG, &read_config);
            if (err) {
                    printf(MAIN_PREFIX "error while rt_dev_ioctl, %s\n",
                            strerror(-err));
                    goto error;
            }
            printf(MAIN_PREFIX "read-config written\n");

            /* create write_task */
            err = rt_task_create(&write_task, "write_task", 0, 50, 0);
            if (err) {
```

```
                    printf(MAIN_PREFIX "failed to create write_task, %s\n",
                           strerror(-err));
                    goto error;
            }
            write_state |= STATE_TASK_CREATED;
            printf(MAIN_PREFIX "write-task created\n");

            /* create read_task */
            err = rt_task_create(&read_task, "read_task", 0, 51, 0);
            if (err) {
                    printf(MAIN_PREFIX "failed to create read_task, %s\n",
                           strerror(-err));
                    goto error;
            }
            read_state |= STATE_TASK_CREATED;
            printf(MAIN_PREFIX "read-task created\n");

            /* start write_task */
            printf(MAIN_PREFIX "starting write-task\n");
            err = rt_task_start(&write_task, &write_task_proc, NULL);
            if (err) {
                    printf(MAIN_PREFIX "failed to start write_task, %s\n",
                           strerror(-err));
                    goto error;
            }

            /* start read_task */
            printf(MAIN_PREFIX "starting read-task\n");
            err = rt_task_start(&read_task,&read_task_proc,NULL);
            if (err) {
                    printf(MAIN_PREFIX "failed to start read_task, %s\n",
                           strerror(-err));
                    goto error;
            }

            pause();
            return 0;

 error:
        cleanup_all();
        return err;
}
```

8.4 iddp-label.c

```
/*
 * IDDP-based client/server demo, using the write(2)/recvfrom(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created.  A server thread (reader)
 * is bound to a labeled real-time port and receives datagrams sent to
 * this port from a client thread (writer). The client thread attaches
 * to the port opened by the server using a labeled connection
 * request. The client socket is bound to a different port, only to
 * provide a valid peer name; this is optional.
 *
 * ASCII labels can be attached to bound ports, in order to connect
 * sockets to them in a more descriptive way than using plain numeric
 * port values.
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
```

```
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t svtid, cltid;

#define IDDP_CLPORT  27

#define IDDP_PORT_LABEL  "iddp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *server(void *arg)
{
        struct sockaddr_ipc saddr, claddr;
        struct rtipc_port_label plabel;
        socklen_t addrlen;
        char buf[128];
        int ret, s;

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
        if (s < 0)
                fail("socket");

        /*
         * We will use Xenomai's system heap for datagram, so no
         * IDDP_POOLSZ required here.
         */

        /*
         * Set a port label. This name will be registered when
         * binding, in addition to the port number (if given).
         */
        strcpy(plabel.label, IDDP_PORT_LABEL);
        ret = setsockopt(s, SOL_IDDP, IDDP_LABEL,
                        &plabel, sizeof(plabel));
        if (ret)
                fail("setsockopt");

        /*
```

```
          * Bind the socket to the port. Assign that port a label, so
          * that peers may use a descriptive information to locate
          * it. Labeled ports will appear in the
          * /proc/xenomai/registry/rtipc/iddp directory once the socket
          * is bound.
          *
          * saddr.sipc_port specifies the port number to use. If -1 is
          * passed, the IDDP driver will auto-select an idle port.
          */
         saddr.sipc_family = AF_RTIPC;
         saddr.sipc_port = -1;    /* Pick next free */
         ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
         if (ret)
                 fail("bind");

         for (;;) {
                 addrlen = sizeof(saddr);
                 ret = recvfrom(s, buf, sizeof(buf), 0,
                                 (struct sockaddr *)&claddr, &addrlen);
                 if (ret < 0) {
                         close(s);
                         fail("recvfrom");
                 }
                 rt_printf("%s: received %d bytes, \"%.*s\" from port %d\n",
                           __FUNCTION__, ret, ret, buf, claddr.sipc_port);
         }

         return NULL;
}

void *client(void *arg)
{
         struct sockaddr_ipc svsaddr, clsaddr;
         struct rtipc_port_label plabel;
         int ret, s, n = 0, len;
         struct timespec ts;
         char buf[128];

         s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
         if (s < 0)
                 fail("socket");

         /*
          * Set a name on the client socket. This is strictly optional,
          * and only done here for the purpose of getting back a
          * different port number in recvfrom().
          */
         clsaddr.sipc_family = AF_RTIPC;
         clsaddr.sipc_port = IDDP_CLPORT;
         ret = bind(s, (struct sockaddr *)&clsaddr, sizeof(clsaddr));
         if (ret)
                 fail("bind");

         /*
          * Set the port label. This name will be used to find the peer
          * when connecting, instead of the port number. The label must
          * be set _after_ the socket is bound to the port, so that
          * IDDP does not try to register this label for the client
          * port as well (like the server thread did).
          */
         strcpy(plabel.label, IDDP_PORT_LABEL);
         ret = setsockopt(s, SOL_IDDP, IDDP_LABEL,
                         &plabel, sizeof(plabel));
         if (ret)
                 fail("setsockopt");

         memset(&svsaddr, 0, sizeof(svsaddr));
```

```
        svsaddr.sipc_family = AF_RTIPC;
        svsaddr.sipc_port = -1; /* Tell IDDP to search by label. */
        ret = connect(s, (struct sockaddr *)&svsaddr, sizeof(svsaddr));
        if (ret)
                fail("connect");

        for (;;) {
                len = strlen(msg[n]);
                /* Send to default destination we connected to. */
                ret = write(s, msg[n], len);
                if (ret < 0) {
                        close(s);
                        fail("sendto");
                }
                rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                          __FUNCTION__, ret, ret, msg[n]);
                n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                /*
                 * We run in full real-time mode (i.e. primary mode),
                 * so we have to let the system breathe between two
                 * iterations.
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(svtid);
        pthread_cancel(cltid);
        signal(sig, SIG_DFL);
        pthread_join(svtid, NULL);
        pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param svparam = {.sched_priority = 71 };
        struct sched_param clparam = {.sched_priority = 70 };
        pthread_attr_t svattr, clattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary mode.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&svattr);
        pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
```

```
        pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
        pthread_attr_setschedparam(&svattr, &svparam);

        errno = pthread_create(&svtid, &svattr, &server, NULL);
        if (errno)
                fail("pthread_create");

        pthread_attr_init(&clattr);
        pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
        pthread_attr_setschedparam(&clattr, &clparam);

        errno = pthread_create(&cltid, &clattr, &client, NULL);
        if (errno)
                fail("pthread_create");

        sigsuspend(&oldmask);

        return 0;
}
```

8.5 iddp-sendrecv.c

```
/*
 * IDDP-based client/server demo, using the sendto(2)/recvfrom(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created.  A server thread (reader)
 * is bound to a real-time port and receives datagrams sent to this
 * port from a client thread (writer). The client socket is bound to a
 * different port, only to provide a valid peer name; this is
 * optional.
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t svtid, cltid;

#define IDDP_SVPORT 12
#define IDDP_CLPORT 13

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
```

```
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *server(void *arg)
{
        struct sockaddr_ipc saddr, claddr;
        socklen_t addrlen;
        char buf[128];
        size_t poolsz;
        int ret, s;

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
        if (s < 0)
                fail("socket");

        /*
         * Set a local 32k pool for the server endpoint. Memory needed
         * to convey datagrams will be pulled from this pool, instead
         * of Xenomai's system pool.
         */
        poolsz = 32768; /* bytes */
        ret = setsockopt(s, SOL_IDDP, IDDP_POOLSZ,
                          &poolsz, sizeof(poolsz));
        if (ret)
                fail("setsockopt");

        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = IDDP_SVPORT;
        ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("bind");

        for (;;) {
                addrlen = sizeof(saddr);
                ret = recvfrom(s, buf, sizeof(buf), 0,
                                (struct sockaddr *)&claddr, &addrlen);
                if (ret < 0) {
                        close(s);
                        fail("recvfrom");
                }
                rt_printf("%s: received %d bytes, \"%.*s\" from port %d\n",
                           __FUNCTION__, ret, ret, buf, claddr.sipc_port);
        }

        return NULL;
}

void *client(void *arg)
{
        struct sockaddr_ipc svsaddr, clsaddr;
        int ret, s, n = 0, len;
        struct timespec ts;
        char buf[128];

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
        if (s < 0)
                fail("socket");
```

```
        clsaddr.sipc_family = AF_RTIPC;
        clsaddr.sipc_port = IDDP_CLPORT;
        ret = bind(s, (struct sockaddr *)&clsaddr, sizeof(clsaddr));
        if (ret)
                fail("bind");

        svsaddr.sipc_family = AF_RTIPC;
        svsaddr.sipc_port = IDDP_SVPORT;
        for (;;) {
                len = strlen(msg[n]);
                ret = sendto(s, msg[n], len, 0,
                             (struct sockaddr *)&svsaddr, sizeof(svsaddr));
                if (ret < 0) {
                        close(s);
                        fail("sendto");
                }
                rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                            __FUNCTION__, ret, ret, msg[n]);
                n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                /*
                 * We run in full real-time mode (i.e. primary mode),
                 * so we have to let the system breathe between two
                 * iterations.
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(svtid);
        pthread_cancel(cltid);
        signal(sig, SIG_DFL);
        pthread_join(svtid, NULL);
        pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param svparam = {.sched_priority = 71 };
        struct sched_param clparam = {.sched_priority = 70 };
        pthread_attr_t svattr, clattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary mode.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&svattr);
```

```
        pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
        pthread_attr_setschedparam(&svattr, &svparam);

        errno = pthread_create(&svtid, &svattr, &server, NULL);
        if (errno)
                fail("pthread_create");

        pthread_attr_init(&clattr);
        pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
        pthread_attr_setschedparam(&clattr, &clparam);

        errno = pthread_create(&cltid, &clattr, &client, NULL);
        if (errno)
                fail("pthread_create");

        sigsuspend(&oldmask);

        return 0;
}
```

8.6   rtcan_rtt.c

```
/*
 * Round-Trip-Time Test - sends and receives messages and measures the
 *                        time in between.
 *
 * Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>
 *
 * Based on RTnet's examples/xenomai/posix/rtt-sender.c.
 *
 * Copyright (C) 2002 Ulrich Marx <marx@kammer.uni-hannover.de>
 *               2002 Marc Kleine-Budde <kleine-budde@gmx.de>
 *               2006 Jan Kiszka <jan.kiszka@web.de>
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 *
 *
 * The program sends out CAN messages periodically and copies the current
 * time-stamp to the payload. At reception, that time-stamp is compared
 * with the current time to determine the round-trip time. The jitter
 * values are printer out regularly. Concurrent tests can be carried out
 * by starting the program with different message identifiers. It is also
 * possible to use this program on a remote system as simple repeater to
 * loopback messages.
 */

#include <errno.h>
#include <mqueue.h>
#include <signal.h>
```

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <limits.h>
#include <getopt.h>
#include <netinet/in.h>
#include <net/if.h>
#include <sys/ioctl.h>
#include <sys/mman.h>

#ifdef __XENO__
#include <rtdm/rtcan.h>
#else
#include <linux/can.h>
#include <linux/can/raw.h>
#endif

#define NSEC_PER_SEC 1000000000

static unsigned int cycle = 10000; /* 10 ms */
static canid_t can_id = 0x1;

static pthread_t txthread, rxthread;
static int txsock, rxsock;
static mqd_t mq;
static int txcount, rxcount;
static int overruns;
static int repeater;

struct rtt_stat {
    long long rtt;
    long long rtt_min;
    long long rtt_max;
    long long rtt_sum;
    long long rtt_sum_last;
    int counts_per_sec;
};

static void print_usage(char *prg)
{
    fprintf(stderr,
            "Usage: %s  [Options] <tx-can-interface> <rx-can-interface>\n"
            "Options:\n"
            " -h, --help     This help\n"
            " -r, --repeater Repeater, send back received messages\n"
            " -i, --id=ID    CAN Identifier (default = 0x1)\n"
            " -c, --cycle    Cycle time in us (default = 10000us)\n",
            prg);
}

void *transmitter(void *arg)
{
    struct sched_param  param = { .sched_priority = 80 };
    struct timespec next_period;
    struct timespec time;
    struct can_frame frame;
    long long *rtt_time = (long long *)&frame.data;

    /* Pre-fill CAN frame */
    frame.can_id = can_id;
    frame.can_dlc = sizeof(*rtt_time);

#ifdef __XENO__
    pthread_set_name_np(pthread_self(), "rtcan_rtt_transmitter");
#endif
```

```
        pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

        clock_gettime(CLOCK_MONOTONIC, &next_period);

        while(1) {
            next_period.tv_nsec += cycle * 1000;
            while (next_period.tv_nsec >= NSEC_PER_SEC) {
                    next_period.tv_nsec -= NSEC_PER_SEC;
                    next_period.tv_sec++;
            }

            clock_nanosleep(CLOCK_MONOTONIC, TIMER_ABSTIME, &next_period, NULL);

            if (rxcount != txcount) {
                overruns++;
                continue;
            }

            clock_gettime(CLOCK_MONOTONIC, &time);
            *rtt_time = (long long)time.tv_sec * NSEC_PER_SEC + time.tv_nsec;

            /* Transmit the message containing the local time */
            if (send(txsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
                if (errno == EBADF)
                    printf("terminating transmitter thread\n");
                else
                    perror("send failed");
                return NULL;
            }
            txcount++;
        }
}


void *receiver(void *arg)
{
    struct sched_param param = { .sched_priority = 82 };
    struct timespec time;
    struct can_frame frame;
    long long *rtt_time = (long long *)frame.data;
    struct rtt_stat rtt_stat = {0, 1000000000000000000LL, -1000000000000000000LL,

                                0, 0, 0};

#ifdef __XENO__
    pthread_set_name_np(pthread_self(), "rtcan_rtt_receiver");
#endif
    pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

    rtt_stat.counts_per_sec = 1000000 / cycle;

    while (1) {
        if (recv(rxsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
            if (errno == EBADF)
                printf("terminating receiver thread\n");
            else
                perror("recv failed");
            return NULL;
        }
        if (repeater) {
            /* Transmit the message back as is */
            if (send(txsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
                if (errno == EBADF)
                    printf("terminating transmitter thread\n");
                else
                    perror("send failed");
                return NULL;
```

```
            }
            txcount++;
        } else {
            clock_gettime(CLOCK_MONOTONIC, &time);
            if (rxcount > 0) {
                rtt_stat.rtt = ((long long)time.tv_sec * 1000000000LL +
                                time.tv_nsec - *rtt_time);
                rtt_stat.rtt_sum += rtt_stat.rtt;
                if (rtt_stat.rtt <  rtt_stat.rtt_min)
                    rtt_stat.rtt_min = rtt_stat.rtt;
                if (rtt_stat.rtt > rtt_stat.rtt_max)
                    rtt_stat.rtt_max = rtt_stat.rtt;
            }
        }
        rxcount++;

        if ((rxcount % rtt_stat.counts_per_sec) == 0) {
            mq_send(mq, (char *)&rtt_stat, sizeof(rtt_stat), 0);
            rtt_stat.rtt_sum_last = rtt_stat.rtt_sum;
        }
    }
}

void catch_signal(int sig)
{
    mq_close(mq);
}


int main(int argc, char *argv[])
{
    struct sched_param param = { .sched_priority = 1 };
    pthread_attr_t thattr;
    struct mq_attr mqattr;
    struct sockaddr_can rxaddr, txaddr;
    struct can_filter rxfilter[1];
    struct rtt_stat rtt_stat;
    char mqname[32];
    char *txdev, *rxdev;
    struct ifreq ifr;
    int ret, opt;

    struct option long_options[] = {
        { "id", required_argument, 0, 'i'},
        { "cycle", required_argument, 0, 'c'},
        { "repeater", no_argument, 0, 'r'},
        { "help", no_argument, 0, 'h'},
        { 0, 0, 0, 0},
    };

    while ((opt = getopt_long(argc, argv, "hri:c:",
                              long_options, NULL)) != -1) {
        switch (opt) {
        case 'c':
            cycle = atoi(optarg);
            break;

        case 'i':
            can_id = strtoul(optarg, NULL, 0);
            break;

        case 'r':
            repeater = 1;
            break;

        default:
            fprintf(stderr, "Unknown option %c\n", opt);
```

```
    case 'h':
        print_usage(argv[0]);
        exit(-1);
    }
}

printf("%d %d\n", optind, argc);
if (optind + 2 != argc) {
    print_usage(argv[0]);
    exit(0);
}

txdev = argv[optind];
rxdev = argv[optind + 1];

/* Create and configure RX socket */
if ((rxsock = socket(PF_CAN, SOCK_RAW, CAN_RAW)) < 0) {
    perror("RX socket failed");
    return -1;
}

strncpy(ifr.ifr_name, rxdev, IFNAMSIZ);
printf("RX rxsock=%d, ifr_name=%s\n", rxsock, ifr.ifr_name);

if (ioctl(rxsock, SIOCGIFINDEX, &ifr) < 0) {
    perror("RX ioctl SIOCGIFINDEX failed");
    goto failure1;
}

/* We only want to receive our own messages */
rxfilter[0].can_id = can_id;
rxfilter[0].can_mask = 0x3ff;
if (setsockopt(rxsock, SOL_CAN_RAW, CAN_RAW_FILTER,
               &rxfilter, sizeof(struct can_filter)) < 0) {
    perror("RX setsockopt CAN_RAW_FILTER failed");
    goto failure1;
}
memset(&rxaddr, 0, sizeof(rxaddr));
rxaddr.can_ifindex = ifr.ifr_ifindex;
rxaddr.can_family = AF_CAN;
if (bind(rxsock, (struct sockaddr *)&rxaddr, sizeof(rxaddr)) < 0) {
    perror("RX bind failed\n");
    goto failure1;
}

/* Create and configure TX socket */

if (strcmp(rxdev, txdev) == 0) {
    txsock = rxsock;
} else {
    if ((txsock = socket(PF_CAN, SOCK_RAW, 0)) < 0) {
        perror("TX socket failed");
        goto failure1;
    }

    strncpy(ifr.ifr_name, txdev, IFNAMSIZ);
    printf("TX txsock=%d, ifr_name=%s\n", txsock, ifr.ifr_name);

    if (ioctl(txsock, SIOCGIFINDEX, &ifr) < 0) {
        perror("TX ioctl SIOCGIFINDEX failed");
        goto failure2;
    }

    /* Suppress definiton of a default receive filter list */
    if (setsockopt(txsock, SOL_CAN_RAW, CAN_RAW_FILTER, NULL, 0) < 0) {
        perror("TX setsockopt CAN_RAW_FILTER failed");
        goto failure2;
```

```
        }

        memset(&txaddr, 0, sizeof(txaddr));
        txaddr.can_ifindex = ifr.ifr_ifindex;
        txaddr.can_family = AF_CAN;

        if (bind(txsock, (struct sockaddr *)&txaddr, sizeof(txaddr)) < 0) {
                perror("TX bind failed\n");
                goto failure2;
        }
    }

    signal(SIGTERM, catch_signal);
    signal(SIGINT, catch_signal);
    signal(SIGHUP, catch_signal);
    mlockall(MCL_CURRENT|MCL_FUTURE);

    printf("Round-Trip-Time test %s -> %s with CAN ID 0x%x\n",
           argv[optind], argv[optind + 1], can_id);
    printf("Cycle time: %d us\n", cycle);
    printf("All RTT timing figures are in us.\n");

    /* Create statistics message queue */
    snprintf(mqname, sizeof(mqname), "/rtcan_rtt-%d", getpid());
    mqattr.mq_flags   = 0;
    mqattr.mq_maxmsg  = 100;
    mqattr.mq_msgsize = sizeof(struct rtt_stat);
    mq = mq_open(mqname, O_RDWR | O_CREAT | O_EXCL, 0600, &mqattr);
    if (mq == (mqd_t)-1) {
        perror("opening mqueue failed");
        goto failure2;
    }

    /* Create receiver RT-thread */
    pthread_attr_init(&thattr);
    pthread_attr_setdetachstate(&thattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setstacksize(&thattr, PTHREAD_STACK_MIN);
    ret = pthread_create(&rxthread, &thattr, &receiver, NULL);
    if (ret) {
        fprintf(stderr, "%s: pthread_create(receiver) failed\n",
                strerror(-ret));
        goto failure3;
    }

    if (!repeater) {
        /* Create transitter RT-thread */
        ret = pthread_create(&txthread, &thattr, &transmitter, NULL);
        if (ret) {
            fprintf(stderr, "%s: pthread_create(transmitter) failed\n",
                    strerror(-ret));
            goto failure4;
        }
    }

    pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

    if (repeater)
        printf("Messages\n");
    else
        printf("Messages RTTlast RTT_avg RTT_min RTT_max Overruns\n");

    while (1) {
        long long rtt_avg;

        ret = mq_receive(mq, (char *)&rtt_stat, sizeof(rtt_stat), NULL);
        if (ret != sizeof(rtt_stat)) {
            if (ret < 0) {
```

```
                    if (errno == EBADF)
                        printf("terminating mq_receive\n");
                    else
                        perror("mq_receive failed");
                } else
                    fprintf(stderr,
                            "mq_receive returned invalid length %d\n", ret);
                break;
            }

            if (repeater) {
                printf("%8d\n", rxcount);
            } else {
                rtt_avg = ((rtt_stat.rtt_sum - rtt_stat.rtt_sum_last) /
                            rtt_stat.counts_per_sec);
                printf("%8d %7ld %7ld %7ld %7ld %8d\n", rxcount,
                        (long)(rtt_stat.rtt / 1000), (long)(rtt_avg / 1000),
                        (long)(rtt_stat.rtt_min / 1000),
                        (long)(rtt_stat.rtt_max / 1000),
                        overruns);
            }
        }

        /* This call also leaves primary mode, required for socket cleanup. */
        printf("shutting down\n");

        /* Important: First close the sockets! */
        while ((close(rxsock) < 0) && (errno == EAGAIN)) {
            printf("RX socket busy - waiting...\n");
            sleep(1);
        }
        while ((close(txsock) < 0) && (errno == EAGAIN)) {
            printf("TX socket busy - waiting...\n");
            sleep(1);
        }

        pthread_join(txthread, NULL);
        pthread_kill(rxthread, SIGHUP);
        pthread_join(rxthread, NULL);

        return 0;

 failure4:
        pthread_kill(rxthread, SIGHUP);
        pthread_join(rxthread, NULL);
 failure3:
        mq_close(mq);
 failure2:
        close(txsock);
 failure1:
        close(rxsock);

        return 1;
}
```

8.7 rtcanconfig.c

```
/*
 * Program to configuring the CAN controller
 *
 * Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>
 *
 * Copyright (C) 2005, 2006 Sebastian Smolorz
 *                        <Sebastian.Smolorz@stud.uni-hannover.de>
 *
```

```
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 */

#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <string.h>
#include <time.h>
#include <errno.h>
#include <getopt.h>
#include <sys/mman.h>

#include <rtdm/rtcan.h>

static void print_usage(char *prg)
{
    fprintf(stderr,
            "Usage: %s <can-interface> [Options] [up|down|start|stop|sleep]\n"
            "Options:\n"
            " -v, --verbose           be verbose\n"
            " -h, --help              this help\n"
            " -c, --ctrlmode=CTRLMODE  listenonly, loopback or none\n"
            " -b, --baudrate=BPS       baudrate in bits/sec\n"
            " -B, --bittime=BTR0:BTR1  BTR or standard bit-time\n"
            " -B, --bittime=BRP:PROP_SEG:PHASE_SEG1:PHASE_SEG2:SJW:SAM\n",
            prg);
}

can_baudrate_t string_to_baudrate(char *str)
{
    can_baudrate_t baudrate;
    if (sscanf(str, "%i", &baudrate) != 1)
        return -1;
    return baudrate;
}

int string_to_mode(char *str)
{
    if ( !strcmp(str, "up") || !strcmp(str, "start") )
        return CAN_MODE_START;
    else if ( !strcmp(str, "down") || !strcmp(str, "stop") )
        return CAN_MODE_STOP;
    else if ( !strcmp(str, "sleep") )
        return CAN_MODE_SLEEP;
    return -EINVAL;
}

int string_to_ctrlmode(char *str)
{
    if ( !strcmp(str, "listenonly") )
        return CAN_CTRLMODE_LISTENONLY;
    else if ( !strcmp(str, "loopback") )
        return CAN_CTRLMODE_LOOPBACK;
```

```
    else if ( !strcmp(str, "none") )
        return 0;

    return -1;
}

int main(int argc, char *argv[])
{
    char    ifname[16];
    int     can_fd = -1;
    int     new_baudrate = -1;
    int     new_mode = -1;
    int     new_ctrlmode = 0, set_ctrlmode = 0;
    int     verbose = 0;
    int     bittime_count = 0, bittime_data[6];
    struct  ifreq ifr;
    can_baudrate_t *baudrate;
    can_ctrlmode_t *ctrlmode;
    can_mode_t *mode;
    struct can_bittime *bittime;
    int opt, ret;
    char* ptr;

    struct option long_options[] = {
        { "help", no_argument, 0, 'h' },
        { "verbose", no_argument, 0, 'v'},
        { "baudrate", required_argument, 0, 'b'},
        { "bittime", required_argument, 0, 'B'},
        { "ctrlmode", required_argument, 0, 'c'},
        { 0, 0, 0, 0},
    };

    while ((opt = getopt_long(argc, argv, "hvb:B:c:",
                              long_options, NULL)) != -1) {
        switch (opt) {
        case 'h':
            print_usage(argv[0]);
            exit(0);

        case 'v':
            verbose = 1;
            break;

        case 'b':
            new_baudrate = string_to_baudrate(optarg);
            if (new_baudrate == -1) {
                print_usage(argv[0]);
                exit(0);
            }
            break;

        case 'B':
            ptr = optarg;
            while (1) {
                bittime_data[bittime_count++] = strtoul(ptr, NULL, 0);
                if (!(ptr = strchr(ptr, ':')))
                    break;
                ptr++;
            }
            if (bittime_count != 2 && bittime_count != 6) {
                print_usage(argv[0]);
                exit(0);
            }
            break;

        case 'c':
            ret = string_to_ctrlmode(optarg);
```

```
            if (ret == -1) {
                print_usage(argv[0]);
                exit(0);
            }
            new_ctrlmode |= ret;
            set_ctrlmode = 1;
            break;

            break;

        default:
            fprintf(stderr, "Unknown option %c\n", opt);
            break;
        }
    }

    /* Get CAN interface name */
    if (optind != argc - 1 && optind != argc - 2) {
        print_usage(argv[0]);
        return 0;
    }

    strncpy(ifname, argv[optind], IFNAMSIZ);
    strncpy(ifr.ifr_name, ifname, IFNAMSIZ);

    if (optind == argc - 2) {   /* Get mode setting */
        new_mode = string_to_mode(argv[optind + 1]);
        if (verbose)
            printf("mode: %s (%#x)\n", argv[optind + 1], new_mode);
        if (new_mode < 0) {
            print_usage(argv[0]);
            return 0;
        }
    }

    can_fd = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
    if (can_fd < 0) {
        fprintf(stderr, "Cannot open RTDM CAN socket. Maybe driver not loaded? \n
  ");
        return can_fd;
    }

    ret = rt_dev_ioctl(can_fd, SIOCGIFINDEX, &ifr);
    if (ret) {
        fprintf(stderr,"Can't get interface index for %s, code = %d\n", ifname, r
  et);
        return ret;
    }


    if (new_baudrate != -1) {
        if (verbose)
            printf("baudrate: %d\n", new_baudrate);
        baudrate = (can_baudrate_t *)&ifr.ifr_ifru;
        *baudrate = new_baudrate;
        ret = rt_dev_ioctl(can_fd, SIOCSCANBAUDRATE, &ifr);
        if (ret) {
            goto abort;
        }
    }

    if (bittime_count) {
        bittime = (struct can_bittime *)&ifr.ifr_ifru;
        if (bittime_count == 2) {
            bittime->type = CAN_BITTIME_BTR;
            bittime->btr.btr0 = bittime_data[0];
            bittime->btr.btr1 = bittime_data[1];
```

```
        if (verbose)
            printf("bit-time: btr0=0x%02x btr1=0x%02x\n",
                   bittime->btr.btr0, bittime->btr.btr1);
    } else {
        bittime->type = CAN_BITTIME_STD;
        bittime->std.brp = bittime_data[0];
        bittime->std.prop_seg = bittime_data[1];
        bittime->std.phase_seg1 = bittime_data[2];
        bittime->std.phase_seg2 = bittime_data[3];
        bittime->std.sjw = bittime_data[4];
        bittime->std.sam = bittime_data[5];
        if (verbose)
            printf("bit-time: brp=%d prop_seg=%d phase_seg1=%d "
                   "phase_seg2=%d sjw=%d sam=%d\n",
                   bittime->std.brp,
                   bittime->std.prop_seg,
                   bittime->std.phase_seg1,
                   bittime->std.phase_seg2,
                   bittime->std.sjw,
                   bittime->std.sam);
    }

    ret = rt_dev_ioctl(can_fd, SIOCSCANCUSTOMBITTIME, &ifr);
    if (ret) {
        goto abort;
    }

}

if (set_ctrlmode != 0) {
    ctrlmode = (can_ctrlmode_t *)&ifr.ifr_ifru;
    *ctrlmode = new_ctrlmode;
    if (verbose)
        printf("ctrlmode: %#x\n", new_ctrlmode);
    ret = rt_dev_ioctl(can_fd, SIOCSCANCTRLMODE, &ifr);
    if (ret) {
        goto abort;
    }
}

if (new_mode != -1) {
    mode = (can_mode_t *)&ifr.ifr_ifru;
    *mode = new_mode;
    ret = rt_dev_ioctl(can_fd, SIOCSCANMODE, &ifr);
    if (ret) {
        goto abort;
    }
}

rt_dev_close(can_fd);
return 0;

abort:
rt_dev_close(can_fd);
return ret;
}
```

8.8 rtcanrecv.c

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <time.h>
#include <errno.h>
```

```
#include <getopt.h>
#include <sys/mman.h>

#include <native/task.h>
#include <native/pipe.h>

#include <rtdm/rtcan.h>

static void print_usage(char *prg)
{
    fprintf(stderr,
            "Usage: %s [<can-interface>] [Options]\n"
            "Options:\n"
            " -f  --filter=id:mask[:id:mask]... apply filter\n"
            " -e  --error=mask       receive error messages\n"
            " -t, --timeout=MS       timeout in ms\n"
            " -T, --timestamp        with absolute timestamp\n"
            " -R, --timestamp-rel    with relative timestamp\n"
            " -v, --verbose          be verbose\n"
            " -p, --print=MODULO     print every MODULO message\n"
            " -h, --help             this help\n",
            prg);
}


extern int optind, opterr, optopt;

static int s = -1, verbose = 0, print = 1;
static nanosecs_rel_t timeout = 0, with_timestamp = 0, timestamp_rel = 0;

RT_TASK rt_task_desc;

#define BUF_SIZ 255
#define MAX_FILTER 16

struct sockaddr_can recv_addr;
struct can_filter recv_filter[MAX_FILTER];
static int filter_count = 0;

int add_filter(u_int32_t id, u_int32_t mask)
{
    if (filter_count >= MAX_FILTER)
        return -1;
    recv_filter[filter_count].can_id = id;
    recv_filter[filter_count].can_mask = mask;
    printf("Filter #%d: id=0x%08x mask=0x%08x\n", filter_count, id, mask);
    filter_count++;
    return 0;
}

void cleanup(void)
{
    int ret;

    if (verbose)
        printf("Cleaning up...\n");

    if (s >= 0) {
        ret = rt_dev_close(s);
        s = -1;
        if (ret) {
            fprintf(stderr, "rt_dev_close: %s\n", strerror(-ret));
        }
        rt_task_delete(&rt_task_desc);
    }
    }
```

```
void cleanup_and_exit(int sig)
{
    if (verbose)
        printf("Signal %d received\n", sig);
    cleanup();
    exit(0);
}

void rt_task(void)
{
    int i, ret, count = 0;
    struct can_frame frame;
    struct sockaddr_can addr;
    socklen_t addrlen = sizeof(addr);
    struct msghdr msg;
    struct iovec iov;
    nanosecs_abs_t timestamp, timestamp_prev = 0;

    if (with_timestamp) {
        msg.msg_iov = &iov;
        msg.msg_iovlen = 1;
        msg.msg_name = (void *)&addr;
        msg.msg_namelen = sizeof(struct sockaddr_can);
        msg.msg_control = (void *)&timestamp;
        msg.msg_controllen = sizeof(nanosecs_abs_t);
    }

    while (1) {
        if (with_timestamp) {
            iov.iov_base = (void *)&frame;
            iov.iov_len = sizeof(can_frame_t);
            ret = rt_dev_recvmsg(s, &msg, 0);
        } else
            ret = rt_dev_recvfrom(s, (void *)&frame, sizeof(can_frame_t), 0,
                                  (struct sockaddr *)&addr, &addrlen);
        if (ret < 0) {
            switch (ret) {
            case -ETIMEDOUT:
                if (verbose)
                    printf("rt_dev_recv: timed out");
                continue;
            case -EBADF:
                if (verbose)
                    printf("rt_dev_recv: aborted because socket was closed");
                break;
            default:
                fprintf(stderr, "rt_dev_recv: %s\n", strerror(-ret));
            }
            break;
        }

        if (print && (count % print) == 0) {
            printf("#%d: (%d) ", count, addr.can_ifindex);
            if (with_timestamp && msg.msg_controllen) {
                if (timestamp_rel) {
                printf("%lldns ", (long long)(timestamp - timestamp_prev));
                    timestamp_prev = timestamp;
                } else
                    printf("%lldns ", (long long)timestamp);
            }
            if (frame.can_id & CAN_ERR_FLAG)
                printf("!0x%08x!", frame.can_id & CAN_ERR_MASK);
            else if (frame.can_id & CAN_EFF_FLAG)
                printf("<0x%08x>", frame.can_id & CAN_EFF_MASK);
            else
                printf("<0x%03x>", frame.can_id & CAN_SFF_MASK);
```

```
                printf(" [%d]", frame.can_dlc);
                if (!(frame.can_id & CAN_RTR_FLAG))
                    for (i = 0; i < frame.can_dlc; i++) {
                        printf(" %02x", frame.data[i]);
                    }
                if (frame.can_id & CAN_ERR_FLAG) {
                    printf(" ERROR ");
                    if (frame.can_id & CAN_ERR_BUSOFF)
                        printf("bus-off");
                    if (frame.can_id & CAN_ERR_CRTL)
                        printf("controller problem");
                } else if (frame.can_id & CAN_RTR_FLAG)
                    printf(" remote request");
                printf("\n");
            }
            count++;
        }
}

int main(int argc, char **argv)
{
        int opt, ret;
        u_int32_t id, mask;
        u_int32_t err_mask = 0;
        struct ifreq ifr;
        char *ptr;
        char name[32];

        struct option long_options[] = {
            { "help", no_argument, 0, 'h' },
            { "verbose", no_argument, 0, 'v'},
            { "filter", required_argument, 0, 'f'},
            { "error", required_argument, 0, 'e'},
            { "timeout", required_argument, 0, 't'},
            { "timestamp", no_argument, 0, 'T'},
            { "timestamp-rel", no_argument, 0, 'R'},
            { 0, 0, 0, 0},
        };

        mlockall(MCL_CURRENT | MCL_FUTURE);

        signal(SIGTERM, cleanup_and_exit);
        signal(SIGINT, cleanup_and_exit);

        while ((opt = getopt_long(argc, argv, "hve:f:t:p:RT",
                                long_options, NULL)) != -1) {
            switch (opt) {
            case 'h':
                print_usage(argv[0]);
                exit(0);

            case 'p':
                print = strtoul(optarg, NULL, 0);
                break;

            case 'v':
                verbose = 1;
                break;

            case 'e':
                err_mask = strtoul(optarg, NULL, 0);
                break;

            case 'f':
                ptr = optarg;
                while (1) {
                    id = strtoul(ptr, NULL, 0);
```

```
            ptr = strchr(ptr, ':');
            if (!ptr) {
                fprintf(stderr, "filter must be applied in the form id:mask[:
  id:mask]...\n");
                exit(1);
            }
            ptr++;
            mask = strtoul(ptr, NULL, 0);
            ptr = strchr(ptr, ':');
            add_filter(id, mask);
            if (!ptr)
                break;
            ptr++;
        }
        break;

    case 't':
        timeout = (nanosecs_rel_t)strtoul(optarg, NULL, 0) * 1000000;
        break;

    case 'R':
        timestamp_rel = 1;
    case 'T':
        with_timestamp = 1;
        break;

    default:
        fprintf(stderr, "Unknown option %c\n", opt);
        break;
    }
}

ret = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
if (ret < 0) {
    fprintf(stderr, "rt_dev_socket: %s\n", strerror(-ret));
    return -1;
}
s = ret;

if (argv[optind] == NULL) {
    if (verbose)
        printf("interface all\n");

    ifr.ifr_ifindex = 0;
} else {
    if (verbose)
        printf("interface %s\n", argv[optind]);

    strncpy(ifr.ifr_name, argv[optind], IFNAMSIZ);
    if (verbose)
        printf("s=%d, ifr_name=%s\n", s, ifr.ifr_name);

    ret = rt_dev_ioctl(s, SIOCGIFINDEX, &ifr);
    if (ret < 0) {
        fprintf(stderr, "rt_dev_ioctl GET_IFINDEX: %s\n", strerror(-ret));
        goto failure;
    }
}

if (err_mask) {
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW, CAN_RAW_ERR_FILTER,
                            &err_mask, sizeof(err_mask));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }
    if (verbose)
```

```
            printf("Using err_mask=%#x\n", err_mask);
    }

    if (filter_count) {
        ret = rt_dev_setsockopt(s, SOL_CAN_RAW, CAN_RAW_FILTER,
                                &recv_filter, filter_count *
                                sizeof(struct can_filter));
        if (ret < 0) {
            fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
            goto failure;
        }
    }

    recv_addr.can_family = AF_CAN;
    recv_addr.can_ifindex = ifr.ifr_ifindex;
    ret = rt_dev_bind(s, (struct sockaddr *)&recv_addr,
                      sizeof(struct sockaddr_can));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_bind: %s\n", strerror(-ret));
        goto failure;
    }

    if (timeout) {
        if (verbose)
            printf("Timeout: %lld ns\n", (long long)timeout);
        ret = rt_dev_ioctl(s, RTCAN_RTIOC_RCV_TIMEOUT, &timeout);
        if (ret) {
            fprintf(stderr, "rt_dev_ioctl RCV_TIMEOUT: %s\n", strerror(-ret));
            goto failure;
        }
    }

    if (with_timestamp) {
        ret = rt_dev_ioctl(s, RTCAN_RTIOC_TAKE_TIMESTAMP, RTCAN_TAKE_TIMESTAMPS);

        if (ret) {
            fprintf(stderr, "rt_dev_ioctl TAKE_TIMESTAMP: %s\n", strerror(-ret));

            goto failure;
        }
    }

    snprintf(name, sizeof(name), "rtcanrecv-%d", getpid());
    ret = rt_task_shadow(&rt_task_desc, name, 0, 0);
    if (ret) {
        fprintf(stderr, "rt_task_shadow: %s\n", strerror(-ret));
        goto failure;
    }

    rt_task();
    /* never returns */

failure:
    cleanup();
    return -1;
}
```

8.9 rtcansend.c

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <time.h>
#include <errno.h>
```

```
#include <getopt.h>
#include <sys/mman.h>

#include <native/task.h>
#include <native/timer.h>
#include <native/pipe.h>

#include <rtdm/rtcan.h>

extern int optind, opterr, optopt;

static void print_usage(char *prg)
{
    fprintf(stderr,
            "Usage: %s <can-interface> [Options] <can-msg>\n"
            "<can-msg> can consist of up to 8 bytes given as a space separated li
    st\n"
            "Options:\n"
            " -i, --identifier=ID  CAN Identifier (default = 1)\n"
            " -r  --rtr            send remote request\n"
            " -e  --extended       send extended frame\n"
            " -l  --loop=COUNT     send message COUNT times\n"
            " -c, --count          message count in data[0-3]\n"
            " -d, --delay=MS       delay in ms (default = 1ms)\n"
            " -s, --send           use send instead of sendto\n"
            " -t, --timeout=MS     timeout in ms\n"
            " -L, --loopback=0|1   switch local loopback off or on\n"
            " -v, --verbose        be verbose\n"
            " -p, --print=MODULO   print every MODULO message\n"
            " -h, --help           this help\n",
            prg);
}


RT_TASK rt_task_desc;

static int s=-1, dlc=0, rtr=0, extended=0, verbose=0, loops=1;
static SRTIME delay=1000000;
static int count=0, print=1, use_send=0, loopback=-1;
static nanosecs_rel_t timeout = 0;
static struct can_frame frame;
static struct sockaddr_can to_addr;


void cleanup(void)
{
    int ret;

    if (verbose)
        printf("Cleaning up...\n");

    usleep(100000);

    if (s >= 0) {
        ret = rt_dev_close(s);
        s = -1;
        if (ret) {
            fprintf(stderr, "rt_dev_close: %s\n", strerror(-ret));
        }
        rt_task_delete(&rt_task_desc);
    }
}

void cleanup_and_exit(int sig)
{
    if (verbose)
        printf("Signal %d received\n", sig);
```

```
    cleanup();
    exit(0);
}

void rt_task(void)
{
    int i, j, ret;

    for (i = 0; i < loops; i++) {
        rt_task_sleep(rt_timer_ns2ticks(delay));
        if (count)
            memcpy(&frame.data[0], &i, sizeof(i));
        /* Note: sendto avoids the definiton of a receive filter list */
        if (use_send)
            ret = rt_dev_send(s, (void *)&frame, sizeof(can_frame_t), 0);
        else
            ret = rt_dev_sendto(s, (void *)&frame, sizeof(can_frame_t), 0,
                                (struct sockaddr *)&to_addr, sizeof(to_addr));
        if (ret < 0) {
            switch (ret) {
            case -ETIMEDOUT:
                if (verbose)
                    printf("rt_dev_send(to): timed out");
                break;
            case -EBADF:
                if (verbose)
                    printf("rt_dev_send(to): aborted because socket was closed");

                break;
            default:
                fprintf(stderr, "rt_dev_send: %s\n", strerror(-ret));
                break;
            }
            i = loops;          /* abort */
            break;
        }
        if (verbose && (i % print) == 0) {
            if (frame.can_id & CAN_EFF_FLAG)
                printf("<0x%08x>", frame.can_id & CAN_EFF_MASK);
            else
                printf("<0x%03x>", frame.can_id & CAN_SFF_MASK);
            printf(" [%d]", frame.can_dlc);
            for (j = 0; j < frame.can_dlc; j++) {
                printf(" %02x", frame.data[j]);
            }
            printf("\n");
        }
    }
}

int main(int argc, char **argv)
{
    int i, opt, ret;
    struct ifreq ifr;
    char name[32];

    struct option long_options[] = {
        { "help", no_argument, 0, 'h' },
        { "identifier", required_argument, 0, 'i'},
        { "rtr", no_argument, 0, 'r'},
        { "extended", no_argument, 0, 'e'},
        { "verbose", no_argument, 0, 'v'},
        { "count", no_argument, 0, 'c'},
        { "print", required_argument, 0, 'p'},
        { "loop", required_argument, 0, 'l'},
        { "delay", required_argument, 0, 'd'},
        { "send", no_argument, 0, 's'},
```

```
        { "timeout", required_argument, 0, 't'},
        { "loopback", required_argument, 0, 'L'},
        { 0, 0, 0, 0},
    };

    mlockall(MCL_CURRENT | MCL_FUTURE);

    signal(SIGTERM, cleanup_and_exit);
    signal(SIGINT, cleanup_and_exit);

    frame.can_id = 1;

    while ((opt = getopt_long(argc, argv, "hvi:l:red:t:cp:sL:",
                              long_options, NULL)) != -1) {
        switch (opt) {
        case 'h':
            print_usage(argv[0]);
            exit(0);

        case 'p':
            print = strtoul(optarg, NULL, 0);

        case 'v':
            verbose = 1;
            break;

        case 'c':
            count = 1;
            break;

        case 'l':
            loops = strtoul(optarg, NULL, 0);
            break;

        case 'i':
            frame.can_id = strtoul(optarg, NULL, 0);
            break;

        case 'r':
            rtr = 1;
            break;

        case 'e':
            extended = 1;
            break;

        case 'd':
            delay = strtoul(optarg, NULL, 0) * 1000000LL;
            break;

        case 's':
            use_send = 1;
            break;

        case 't':
            timeout = strtoul(optarg, NULL, 0) * 1000000LL;
            break;

        case 'L':
            loopback = strtoul(optarg, NULL, 0);
            break;

        default:
            fprintf(stderr, "Unknown option %c\n", opt);
            break;
        }
    }
```

```
if (optind == argc) {
    print_usage(argv[0]);
    exit(0);
}

if (argv[optind] == NULL) {
    fprintf(stderr, "No Interface supplied\n");
    exit(-1);
}

if (verbose)
    printf("interface %s\n", argv[optind]);

ret = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
if (ret < 0) {
    fprintf(stderr, "rt_dev_socket: %s\n", strerror(-ret));
    return -1;
}
s = ret;

if (loopback >= 0) {
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW, CAN_RAW_LOOPBACK,
                            &loopback, sizeof(loopback));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }
    if (verbose)
        printf("Using loopback=%d\n", loopback);
}

strncpy(ifr.ifr_name, argv[optind], IFNAMSIZ);
if (verbose)
    printf("s=%d, ifr_name=%s\n", s, ifr.ifr_name);

ret = rt_dev_ioctl(s, SIOCGIFINDEX, &ifr);
if (ret < 0) {
    fprintf(stderr, "rt_dev_ioctl: %s\n", strerror(-ret));
    goto failure;
}

memset(&to_addr, 0, sizeof(to_addr));
to_addr.can_ifindex = ifr.ifr_ifindex;
to_addr.can_family = AF_CAN;
if (use_send) {
    /* Suppress definiton of a default receive filter list */
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW, CAN_RAW_FILTER, NULL, 0);
    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }

    ret = rt_dev_bind(s, (struct sockaddr *)&to_addr, sizeof(to_addr));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_bind: %s\n", strerror(-ret));
        goto failure;
    }
}

if (count)
    frame.can_dlc = sizeof(int);
else {
    for (i = optind + 1; i < argc; i++) {
        frame.data[dlc] = strtoul(argv[i], NULL, 0);
        dlc++;
        if( dlc == 8 )
```

```
                    break;
            }
            frame.can_dlc = dlc;
    }

    if (rtr)
        frame.can_id |= CAN_RTR_FLAG;

    if (extended)
        frame.can_id |= CAN_EFF_FLAG;

    if (timeout) {
        if (verbose)
            printf("Timeout: %lld ns\n", (long long)timeout);
        ret = rt_dev_ioctl(s, RTCAN_RTIOC_SND_TIMEOUT, &timeout);
        if (ret) {
            fprintf(stderr, "rt_dev_ioctl SND_TIMEOUT: %s\n", strerror(-ret));
            goto failure;
        }
    }

    snprintf(name, sizeof(name), "rtcansend-%d", getpid());
    ret = rt_task_shadow(&rt_task_desc, name, 1, 0);
    if (ret) {
        fprintf(stderr, "rt_task_shadow: %s\n", strerror(-ret));
        goto failure;
    }

    rt_task();

    cleanup();
    return 0;

 failure:
    cleanup();
    return -1;
}
```

8.10 xddp-echo.c

```
/*
 * XDDP-based RT/NRT threads communication demo.
 *
 * Real-time Xenomai threads and regular Linux threads may want to
 * exchange data in a way that does not require the former to leave
 * the real-time domain (i.e. secondary mode). Message pipes - as
 * implemented by the RTDM-based XDDP protocol - are provided for this
 * purpose.
 *
 * On the Linux domain side, pseudo-device files named /dev/rtp<minor>
 * give regular POSIX threads access to non real-time communication
 * endpoints, via the standard character-based I/O interface. On the
 * Xenomai domain side, sockets may be bound to XDDP ports, which act
 * as proxies to send and receive data to/from the associated
 * pseudo-device files. Ports and pseudo-device minor numbers are
 * paired, meaning that e.g. port 7 will proxy the traffic for
 * /dev/rtp7. Therefore, port numbers may range from 0 to
 * CONFIG_XENO_OPT_PIPE_NRDEV - 1.
 *
 * All data sent through a bound/connected XDDP socket via sendto(2) or
 * write(2) will be passed to the peer endpoint in the Linux domain,
 * and made available for reading via the standard read(2) system
 * call. Conversely, all data sent using write(2) through the non
 * real-time endpoint will be conveyed to the real-time socket
 * endpoint, and made available to the recvfrom(2) or read(2) system
```

```
 * calls.
 *
 * Both threads can use the bi-directional data path to send and
 * receive datagrams in a FIFO manner, as illustrated by the simple
 * echoing process implemented by this program.
 *
 * realtime_thread----------------------------->-------+
 *   => get socket                                      |
 *   => bind socket to port 0                           v
 *   => write traffic to NRT domain via sendto()        |
 *   => read traffic from NRT domain via recvfrom() <--|--+
 *                                                      | |
 * regular_thread--------------------------------------+ |
 *   => open /dev/rtp0                                  | ^
 *   => read traffic from RT domain via read()          | |
 *   => echo traffic back to RT domain via write()     +--+
 *
 * See Makefile in this directory for build directives.
 *
 * NOTE: XDDP is a replacement for the legacy RT_PIPE interface
 * available from the native skin until Xenomai 3.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t rt, nrt;

#define XDDP_PORT 0     /* [0..CONFIG-XENO_OPT_PIPE_NRDEV - 1] */

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *realtime_thread(void *arg)
{
        struct sockaddr_ipc saddr;
```

```
int ret, s, n = 0, len;
struct timespec ts;
size_t poolsz;
char buf[128];

/*
 * Get a datagram socket to bind to the RT endpoint. Each
 * endpoint is represented by a port number within the XDDP
 * protocol namespace.
 */
s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
if (s < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
}

/*
 * Set a local 16k pool for the RT endpoint. Memory needed to
 * convey datagrams will be pulled from this pool, instead of
 * Xenomai's system pool.
 */
poolsz = 16384; /* bytes */
ret = setsockopt(s, SOL_XDDP, XDDP_POOLSZ,
                  &poolsz, sizeof(poolsz));
if (ret)
        fail("setsockopt");

/*
 * Bind the socket to the port, to setup a proxy to channel
 * traffic to/from the Linux domain.
 *
 * saddr.sipc_port specifies the port number to use.
 */
memset(&saddr, 0, sizeof(saddr));
saddr.sipc_family = AF_RTIPC;
saddr.sipc_port = XDDP_PORT;
ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
if (ret)
        fail("bind");

for (;;) {
        len = strlen(msg[n]);
        /*
         * Send a datagram to the NRT endpoint via the proxy.
         * We may pass a NULL destination address, since a
         * bound socket is assigned a default destination
         * address matching the binding address (unless
         * connect(2) was issued before bind(2), in which case
         * the former would prevail).
         */
        ret = sendto(s, msg[n], len, 0, NULL, 0);
        if (ret != len)
                fail("sendto");

        rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                  __FUNCTION__, ret, ret, msg[n]);

        /* Read back packets echoed by the regular thread */
        ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
        if (ret <= 0)
                fail("recvfrom");

        rt_printf("   => \"%.*s\" echoed by peer\n", ret, buf);

        n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
        /*
         * We run in full real-time mode (i.e. primary mode),
```

```
                 * so we have to let the system breathe between two
                 * iterations.
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void *regular_thread(void *arg)
{
        char buf[128], *devname;
        int fd, ret;

        if (asprintf(&devname, "/dev/rtp%d", XDDP_PORT) < 0)
                fail("asprintf");

        fd = open(devname, O_RDWR);
        free(devname);
        if (fd < 0)
                fail("open");

        for (;;) {
                /* Get the next message from realtime_thread. */
                ret = read(fd, buf, sizeof(buf));
                if (ret <= 0)
                        fail("read");

                /* Echo the message back to realtime_thread. */
                ret = write(fd, buf, ret);
                if (ret <= 0)
                        fail("write");
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(rt);
        pthread_cancel(nrt);
        signal(sig, SIG_DFL);
        pthread_join(rt, NULL);
        pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param rtparam = { .sched_priority = 42 };
        pthread_attr_t rtattr, regattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
```

```
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary (i.e. non real-time) mode when
     * writing output.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&rtattr);
    pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
    pthread_attr_setschedparam(&rtattr, &rtparam);

    errno = pthread_create(&rt, &rtattr, &realtime_thread, NULL);
    if (errno)
            fail("pthread_create");

    pthread_attr_init(&regattr);
    pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

    errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
    if (errno)
            fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}
```

8.11   xddp-label.c

```
/*
 * XDDP-based RT/NRT threads communication demo.
 *
 * Real-time Xenomai threads and regular Linux threads may want to
 * exchange data in a way that does not require the former to leave
 * the real-time domain (i.e. secondary mode). Message pipes - as
 * implemented by the RTDM-based XDDP protocol - are provided for this
 * purpose.
 *
 * On the Linux domain side, pseudo-device files named /dev/rtp<minor>
 * give regular POSIX threads access to non real-time communication
 * endpoints, via the standard character-based I/O interface. On the
 * Xenomai domain side, sockets may be bound to XDDP ports, which act
 * as proxies to send and receive data to/from the associated
 * pseudo-device files. Ports and pseudo-device minor numbers are
 * paired, meaning that e.g. port 7 will proxy the traffic for
 * /dev/rtp7. Therefore, port numbers may range from 0 to
 * CONFIG_XENO_OPT_PIPE_NRDEV - 1.
 *
 * All data sent through a bound/connected XDDP socket via sendto(2) or
 * write(2) will be passed to the peer endpoint in the Linux domain,
 * and made available for reading via the standard read(2) system
 * call. Conversely, all data sent using write(2) through the non
 * real-time endpoint will be conveyed to the real-time socket
 * endpoint, and made available to the recvfrom(2) or read(2) system
 * calls.
 *
 * ASCII labels can be attached to bound ports, in order to connect
 * sockets to them in a more descriptive way than using plain numeric
 * port values.
 *
 * The example code below illustrates the following process:
 *
```

```
 * realtime_thread1--------------------------->---------+
 *   => get socket                                      |
 *   => bind socket to port "xddp-demo"                 |
 *   => read traffic from NRT domain via recvfrom()   <--+--+
 *                                                       |  |
 * realtime_thread2--------------------------------------+  |
 *   => get socket                                       |  |
 *   => connect socket to port "xddp-demo"               |  |
 *   => write traffic to NRT domain via sendto()         v  |
 *                                                       |  ^
 * regular_thread----------------------------------------+  |
 *   => open /proc/xenomai/registry/rtipc/xddp/xddp-demo |  |
 *   => read traffic from RT domain via read()           |  |
 *   => mirror traffic to RT domain via write()          +--+
 *
 * See Makefile in this directory for build directives.
 *
 * NOTE: XDDP is a replacement for the legacy RT_PIPE interface
 * available from the native skin until Xenomai 3.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t rt1, rt2, nrt;

#define XDDP_PORT_LABEL  "xddp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *realtime_thread1(void *arg)
{
        struct rtipc_port_label plabel;
        struct sockaddr_ipc saddr;
        int ret, s, len;
```

```
        char buf[128];

        /*
         * Get a datagram socket to bind to the RT endpoint. Each
         * endpoint is represented by a port number within the XDDP
         * protocol namespace.
         */
        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
        if (s < 0) {
                perror("socket");
                exit(EXIT_FAILURE);
        }

        /*
         * Set a port label. This name will be registered when
         * binding, in addition to the port number (if given).
         */
        strcpy(plabel.label, XDDP_PORT_LABEL);
        ret = setsockopt(s, SOL_XDDP, XDDP_LABEL,
                         &plabel, sizeof(plabel));
        if (ret)
                fail("setsockopt");
        /*
         * Bind the socket to the port, to setup a proxy to channel
         * traffic to/from the Linux domain. Assign that port a label,
         * so that peers may use a descriptive information to locate
         * it. For instance, the pseudo-device matching our RT
         * endpoint will appear as
         * /proc/xenomai/registry/rtipc/xddp/<XDDP_PORT_LABEL> in the
         * Linux domain, once the socket is bound.
         *
         * saddr.sipc_port specifies the port number to use. If -1 is
         * passed, the XDDP driver will auto-select an idle port.
         */
        memset(&saddr, 0, sizeof(saddr));
        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = -1;
        ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("bind");

        for (;;) {
                /* Get packets relayed by the regular thread */
                ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
                if (ret <= 0)
                        fail("recvfrom");

                rt_printf("%s: \"%.*s\" relayed by peer\n", __FUNCTION__, ret, bu
f);
        }

        return NULL;
}

void *realtime_thread2(void *arg)
{
        struct rtipc_port_label plabel;
        struct sockaddr_ipc saddr;
        int ret, s, n = 0, len;
        struct timespec ts;
        struct timeval tv;
        socklen_t addrlen;
        char buf[128];

        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
        if (s < 0) {
                perror("socket");
```

```
                exit(EXIT_FAILURE);
        }

        /*
         * Set the socket timeout; it will apply when attempting to
         * connect to a labeled port, and to recvfrom() calls.  The
         * following setup tells the XDDP driver to wait for at most
         * one second until a socket is bound to a port using the same
         * label, or return with a timeout error.
         */
        tv.tv_sec = 1;
        tv.tv_usec = 0;
        ret = setsockopt(s, SOL_SOCKET, SO_RCVTIMEO,
                         &tv, sizeof(tv));
        if (ret)
                fail("setsockopt");

        /*
         * Set a port label. This name will be used to find the peer
         * when connecting, instead of the port number.
         */
        strcpy(plabel.label, XDDP_PORT_LABEL);
        ret = setsockopt(s, SOL_XDDP, XDDP_LABEL,
                         &plabel, sizeof(plabel));
        if (ret)
                fail("setsockopt");

        memset(&saddr, 0, sizeof(saddr));
        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = -1;    /* Tell XDDP to search by label. */
        ret = connect(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("connect");

        /*
         * We succeeded in making the port our default destination
         * address by using its label, but we don't know its actual
         * port number yet. Use getpeername() to retrieve it.
         */
        addrlen = sizeof(saddr);
        ret = getpeername(s, (struct sockaddr *)&saddr, &addrlen);
        if (ret || addrlen != sizeof(saddr))
                fail("getpeername");

        rt_printf("%s: NRT peer is reading from /dev/rtp%d\n",
                   __FUNCTION__, saddr.sipc_port);

        for (;;) {
                len = strlen(msg[n]);
                /*
                 * Send a datagram to the NRT endpoint via the proxy.
                 * We may pass a NULL destination address, since the
                 * socket was successfully assigned the proper default
                 * address via connect(2).
                 */
                ret = sendto(s, msg[n], len, 0, NULL, 0);
                if (ret != len)
                        fail("sendto");

                rt_printf("%s: sent %d bytes, \"%.*s\"\n",
                          __FUNCTION__, ret, ret, msg[n]);

                n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                /*
                 * We run in full real-time mode (i.e. primary mode),
                 * so we have to let the system breathe between two
                 * iterations.
```

```
                 */
                ts.tv_sec = 0;
                ts.tv_nsec = 500000000; /* 500 ms */
                clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
}

void *regular_thread(void *arg)
{
        char buf[128], *devname;
        int fd, ret;

        if (asprintf(&devname,
                        "/proc/xenomai/registry/rtipc/xddp/%s",
                        XDDP_PORT_LABEL) < 0)
                fail("asprintf");

        fd = open(devname, O_RDWR);
        free(devname);
        if (fd < 0)
                fail("open");

        for (;;) {
                /* Get the next message from realtime_thread2. */
                ret = read(fd, buf, sizeof(buf));
                if (ret <= 0)
                        fail("read");

                /* Relay the message to realtime_thread1. */
                ret = write(fd, buf, ret);
                if (ret <= 0)
                        fail("write");
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(rt1);
        pthread_cancel(rt2);
        pthread_cancel(nrt);
        signal(sig, SIG_DFL);
        pthread_join(rt1, NULL);
        pthread_join(rt2, NULL);
        pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param rtparam = { .sched_priority = 42 };
        pthread_attr_t rtattr, regattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);
```

```
       /*
        * This is a real-time compatible printf() package from
        * Xenomai's RT Development Kit (RTDK), that does NOT cause
        * any transition to secondary (i.e. non real-time) mode when
        * writing output.
        */
       rt_print_auto_init(1);

       pthread_attr_init(&rtattr);
       pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
       pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
       pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
       pthread_attr_setschedparam(&rtattr, &rtparam);

       /* Both real-time threads have the same attribute set. */

       errno = pthread_create(&rt1, &rtattr, &realtime_thread1, NULL);
       if (errno)
               fail("pthread_create");

       errno = pthread_create(&rt2, &rtattr, &realtime_thread2, NULL);
       if (errno)
               fail("pthread_create");

       pthread_attr_init(&regattr);
       pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
       pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
       pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

       errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
       if (errno)
               fail("pthread_create");

       sigsuspend(&oldmask);

       return 0;
}
```

8.12   xddp-stream.c

```
/*
 * XDDP-based RT/NRT threads communication demo.
 *
 * Real-time Xenomai threads and regular Linux threads may want to
 * exchange data in a way that does not require the former to leave
 * the real-time domain (i.e. secondary mode). Message pipes - as
 * implemented by the RTDM-based XDDP protocol - are provided for this
 * purpose.
 *
 * On the Linux domain side, pseudo-device files named /dev/rtp<minor>
 * give regular POSIX threads access to non real-time communication
 * endpoints, via the standard character-based I/O interface. On the
 * Xenomai domain side, sockets may be bound to XDDP ports, which act
 * as proxies to send and receive data to/from the associated
 * pseudo-device files. Ports and pseudo-device minor numbers are
 * paired, meaning that e.g. port 7 will proxy the traffic for
 * /dev/rtp7. Therefore, port numbers may range from 0 to
 * CONFIG_XENO_OPT_PIPE_NRDEV - 1.
 *
 * All data sent through a bound/connected XDDP socket via sendto(2) or
 * write(2) will be passed to the peer endpoint in the Linux domain,
 * and made available for reading via the standard read(2) system
 * call. Conversely, all data sent using write(2) through the non
 * real-time endpoint will be conveyed to the real-time socket
 * endpoint, and made available to the recvfrom(2) or read(2) system
```

```
 * calls.
 *
 * In addition to sending datagrams, real-time threads may stream data
 * in a byte-oriented mode through the proxy as well. This increases
 * the bandwidth and reduces the overhead, when a lot of data has to
 * flow down to the Linux domain, if keeping the message boundaries is
 * not required. The example code below illustrates such use.
 *
 * realtime_thread----------------------------------->----------+
 *   => get socket                                              |
 *   => bind socket to port 0                                   v
 *   => write scattered traffic to NRT domain via sendto()      |
 *   => read traffic from NRT domain via recvfrom()           <--|--+
 *                                                              | |
 * regular_thread-----------------------------------------------+ |
 *   => open /dev/rtp0                                          | ^
 *   => read traffic from RT domain via read()                 | |
 *   => echo traffic back to RT domain via write()             +--+
 *
 * See Makefile in this directory for build directives.
 *
 * NOTE: XDDP is a replacement for the legacy RT_PIPE interface
 * available from the native skin until Xenomai 3.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtdk.h>
#include <rtdm/rtipc.h>

pthread_t rt, nrt;

#define XDDP_PORT 0     /* [0..CONFIG-XENO_OPT_PIPE_NRDEV - 1] */

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
        perror(reason);
        exit(EXIT_FAILURE);
}

void *realtime_thread(void *arg)
```

```
{
        struct sockaddr_ipc saddr;
        int ret, s, n = 0, len, b;
        struct timespec ts;
        size_t streamsz;
        char buf[128];

        /*
         * Get a datagram socket to bind to the RT endpoint. Each
         * endpoint is represented by a port number within the XDDP
         * protocol namespace.
         */
        s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
        if (s < 0) {
                perror("socket");
                exit(EXIT_FAILURE);
        }

        /*
         * Tell the XDDP driver that we will use the streaming
         * capabilities on this socket. To this end, we have to
         * specify the size of the streaming buffer, as a count of
         * bytes. The real-time output will be buffered up to that
         * amount, and sent as a single datagram to the NRT endpoint
         * when fully gathered, or when another source port attempts
         * to send data to the same endpoint. Passing a null size
         * would disable streaming.
         */
        streamsz = 1024; /* bytes */
        ret = setsockopt(s, SOL_XDDP, XDDP_BUFSZ,
                         &streamsz, sizeof(streamsz));
        if (ret)
                fail("setsockopt");
        /*
         * Bind the socket to the port, to setup a proxy to channel
         * traffic to/from the Linux domain.
         *
         * saddr.sipc_port specifies the port number to use.
         */
        memset(&saddr, 0, sizeof(saddr));
        saddr.sipc_family = AF_RTIPC;
        saddr.sipc_port = XDDP_PORT;
        ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
        if (ret)
                fail("bind");

        for (;;) {
                len = strlen(msg[n]);
                /*
                 * Send a datagram to the NRT endpoint via the proxy.
                 * The output is artificially scattered in separate
                 * one-byte sendings, to illustrate the use of
                 * MSG_MORE.
                 */
                for (b = 0; b < len; b++) {
                        ret = sendto(s, msg[n] + b, 1, MSG_MORE, NULL, 0);
                        if (ret != 1)
                                fail("sendto");
                }

                rt_printf("%s: sent (scattered) %d-bytes message, \"%.*s\"\n",
                          __FUNCTION__, len, len, msg[n]);

                /* Read back packets echoed by the regular thread */
                ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
                if (ret <= 0)
                        fail("recvfrom");
```

```
                      rt_printf("  => \"%.*s\" echoed by peer\n", ret, buf);

                      n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
                      /*
                       * We run in full real-time mode (i.e. primary mode),
                       * so we have to let the system breathe between two
                       * iterations.
                       */
                      ts.tv_sec = 0;
                      ts.tv_nsec = 500000000; /* 500 ms */
                      clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
              }

              return NULL;
      }

void *regular_thread(void *arg)
{
        char buf[128], *devname;
        int fd, ret;

        if (asprintf(&devname, "/dev/rtp%d", XDDP_PORT) < 0)
                fail("asprintf");

        fd = open(devname, O_RDWR);
        free(devname);
        if (fd < 0)
                fail("open");

        for (;;) {
                /* Get the next message from realtime_thread. */
                ret = read(fd, buf, sizeof(buf));
                if (ret <= 0)
                        fail("read");

                /* Echo the message back to realtime_thread. */
                ret = write(fd, buf, ret);
                if (ret <= 0)
                        fail("write");
        }

        return NULL;
}

void cleanup_upon_sig(int sig)
{
        pthread_cancel(rt);
        pthread_cancel(nrt);
        signal(sig, SIG_DFL);
        pthread_join(rt, NULL);
        pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
        struct sched_param rtparam = { .sched_priority = 42 };
        pthread_attr_t rtattr, regattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
```

```
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary (i.e. non real-time) mode when
         * writing output.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&rtattr);
        pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
        pthread_attr_setschedparam(&rtattr, &rtparam);

        errno = pthread_create(&rt, &rtattr, &realtime_thread, NULL);
        if (errno)
                fail("pthread_create");

        pthread_attr_init(&regattr);
        pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

        errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
        if (errno)
                fail("pthread_create");

        sigsuspend(&oldmask);

        return 0;
}
```

Index