

Xenomai

2.99.0

Generated by Doxygen 1.8.6

Fri Jun 13 2014 12:09:58

Contents

1	Deprecated List	1
2	Module Index	3
2.1	Modules	3
3	Data Structure Index	5
3.1	Data Structures	5
4	File Index	7
4.1	File List	7
5	Module Documentation	15
5.1	Channels and ranges	15
5.1.1	Detailed Description	16
5.2	Big dual kernel lock	18
5.2.1	Detailed Description	18
5.2.2	Macro Definition Documentation	18
5.2.2.1	cobalt_atomic_enter	18
5.2.2.2	cobalt_atomic_leave	19
5.2.2.3	RTDM_EXECUTE_ATOMICALY	19
5.3	Spinlock with preemption deactivation	21
5.3.1	Detailed Description	21
5.3.2	Macro Definition Documentation	21
5.3.2.1	rtdm_lock_irqrestore	21
5.3.2.2	rtdm_lock_irqsave	22
5.3.3	Function Documentation	22
5.3.3.1	rtdm_lock_get	22
5.3.3.2	rtdm_lock_init	23
5.3.3.3	rtdm_lock_put	24
5.3.3.4	rtdm_lock_put_irqrestore	24
5.4	Thread state flags	25
5.4.1	Detailed Description	26
5.4.2	Macro Definition Documentation	26

5.4.2.1	XNHELD	26
5.4.2.2	XNLOCK	26
5.4.2.3	XNMIGRATE	26
5.4.2.4	XNPEND	26
5.4.2.5	XNREADY	26
5.4.2.6	XNSUSP	26
5.4.2.7	XNTRAPLB	27
5.5	Thread information flags	28
5.5.1	Detailed Description	28
5.6	CAN Devices	29
5.6.1	Detailed Description	35
5.6.2	Macro Definition Documentation	37
5.6.2.1	CAN_CTRLMODE_3_SAMPLES	37
5.6.2.2	CAN_CTRLMODE_LISTENONLY	38
5.6.2.3	CAN_CTRLMODE_LOOPBACK	38
5.6.2.4	CAN_ERR_LOSTARB_UNSPEC	38
5.6.2.5	CAN_RAW_ERR_FILTER	38
5.6.2.6	CAN_RAW_FILTER	39
5.6.2.7	CAN_RAW_LOOPBACK	39
5.6.2.8	CAN_RAW_RECV_OWN_MSGS	40
5.6.2.9	RTCAN_RTIOC_RCV_TIMEOUT	40
5.6.2.10	RTCAN_RTIOC_SND_TIMEOUT	40
5.6.2.11	RTCAN_RTIOC_TAKE_TIMESTAMP	41
5.6.2.12	SIOGCANBAUDRATE	42
5.6.2.13	SIOGCANCTRLMODE	42
5.6.2.14	SIOGCANCUSTOMBITTIME	43
5.6.2.15	SIOGCANSTATE	43
5.6.2.16	SIOCGIFINDEX	44
5.6.2.17	SIOCSCANBAUDRATE	44
5.6.2.18	SIOCSCANCTRLMODE	45
5.6.2.19	SIOCSCANCUSTOMBITTIME	46
5.6.2.20	SIOCSCANMODE	47
5.6.2.21	SOL_CAN_RAW	47
5.6.3	Typedef Documentation	48
5.6.3.1	can_filter_t	48
5.6.3.2	can_frame_t	48
5.6.4	Enumeration Type Documentation	48
5.6.4.1	CAN_BITTIME_TYPE	48
5.6.4.2	CAN_MODE	48
5.6.4.3	CAN_STATE	49

5.7	Real-Time Driver Model	50
5.7.1	Detailed Description	51
5.7.2	Macro Definition Documentation	51
5.7.2.1	RTDM_TIMEOUT_INFINITE	51
5.7.2.2	RTDM_TIMEOUT_NONE	51
5.7.3	Typedef Documentation	51
5.7.3.1	nanosecs_abs_t	51
5.7.3.2	nanosecs_rel_t	51
5.8	RTDM User API	52
5.8.1	Detailed Description	53
5.8.2	Function Documentation	53
5.8.2.1	rt_dev_accept	53
5.8.2.2	rt_dev_bind	53
5.8.2.3	rt_dev_close	54
5.8.2.4	rt_dev_connect	54
5.8.2.5	rt_dev_getpeername	55
5.8.2.6	rt_dev_getsockname	55
5.8.2.7	rt_dev_getsockopt	56
5.8.2.8	rt_dev_ioctl	56
5.8.2.9	rt_dev_listen	57
5.8.2.10	rt_dev_open	58
5.8.2.11	rt_dev_read	58
5.8.2.12	rt_dev_recv	59
5.8.2.13	rt_dev_recvfrom	59
5.8.2.14	rt_dev_recvmsg	60
5.8.2.15	rt_dev_send	60
5.8.2.16	rt_dev_sendmsg	61
5.8.2.17	rt_dev_sendto	61
5.8.2.18	rt_dev_setsockopt	62
5.8.2.19	rt_dev_shutdown	62
5.8.2.20	rt_dev_socket	63
5.8.2.21	rt_dev_write	63
5.9	Serial Devices	65
5.10	Testing Devices	66
5.11	Real-time IPC	67
5.11.1	Detailed Description	69
5.11.2	Macro Definition Documentation	69
5.11.2.1	BUFP_BUFSZ	69
5.11.2.2	BUFP_LABEL	70
5.11.2.3	IDDP_LABEL	70

5.11.2.4 IDDP_POOLSZ	71
5.11.2.5 SO_RCVTIMEO	72
5.11.2.6 SO_SNDTIMEO	72
5.11.2.7 XDDP_BUFSZ	72
5.11.2.8 XDDP_EVTDOWN	73
5.11.2.9 XDDP_EVTIN	73
5.11.2.10XDDP_EVTNOBUF	73
5.11.2.11XDDP_EVTOUT	73
5.11.2.12XDDP_LABEL	74
5.11.2.13XDDP_MONITOR	74
5.11.2.14XDDP_POOLSZ	75
5.11.3 Enumeration Type Documentation	76
5.11.3.1 anonymous enum	76
5.11.4 Function Documentation	76
5.11.4.1 bind__AF_RTIPC	76
5.11.4.2 close__AF_RTIPC	77
5.11.4.3 connect__AF_RTIPC	78
5.11.4.4 getpeername__AF_RTIPC	78
5.11.4.5 getsockname__AF_RTIPC	79
5.11.4.6 getsockopt__AF_RTIPC	79
5.11.4.7 recvmsg__AF_RTIPC	79
5.11.4.8 sendmsg__AF_RTIPC	80
5.11.4.9 setsockopt__AF_RTIPC	81
5.11.4.10socket__AF_RTIPC	81
5.12 Asynchronous Procedure Calls	82
5.12.1 Detailed Description	82
5.12.2 Function Documentation	82
5.12.2.1 xnapc_alloc	82
5.12.2.2 xnapc_free	83
5.12.2.3 xnapc_schedule	83
5.13 Helpers for in-kernel arithmetics	84
5.13.1 Detailed Description	84
5.14 Buffer descriptor	85
5.14.1 Detailed Description	85
5.14.2 Function Documentation	87
5.14.2.1 xnbufd_copy_from_kmem	87
5.14.2.2 xnbufd_copy_to_kmem	88
5.14.2.3 xnbufd_invalidate	89
5.14.2.4 xnbufd_map_kread	89
5.14.2.5 xnbufd_map_kwrite	89

5.14.2.6	xnbufd_map_uread	90
5.14.2.7	xnbufd_map_uwrite	90
5.14.2.8	xnbufd_reset	90
5.14.2.9	xnbufd_unmap_kread	91
5.14.2.10	xnbufd_unmap_kwrite	91
5.14.2.11	xnbufd_unmap_uread	91
5.14.2.12	xnbufd_unmap_uwrite	92
5.15	Clock services	93
5.15.1	Detailed Description	93
5.15.2	Function Documentation	93
5.15.2.1	xnclock_adjust	93
5.15.2.2	xnclock_deregister	94
5.15.2.3	xnclock_register	94
5.15.2.4	xnclock_tick	94
5.16	Debugging services	96
5.16.1	Detailed Description	96
5.17	Dynamic memory allocation services	97
5.17.1	Detailed Description	97
5.17.2	Function Documentation	98
5.17.2.1	xnheap_alloc	98
5.17.2.2	xnheap_destroy	98
5.17.2.3	xnheap_extend	99
5.17.2.4	xnheap_free	99
5.17.2.5	xnheap_init	100
5.17.2.6	xnheap_set_label	100
5.17.2.7	xnheap_test_and_free	101
5.18	Xenomai nucleus	102
5.18.1	Detailed Description	103
5.19	Interrupt management	104
5.19.1	Detailed Description	104
5.19.2	Function Documentation	104
5.19.2.1	xnintr_affinity	104
5.19.2.2	xnintr_attach	105
5.19.2.3	xnintr_destroy	105
5.19.2.4	xnintr_detach	105
5.19.2.5	xnintr_disable	106
5.19.2.6	xnintr_enable	106
5.19.2.7	xnintr_init	106
5.20	SMP locking services	108
5.20.1	Detailed Description	108

5.21 Lightweight key-to-object mapping service	109
5.21.1 Detailed Description	109
5.21.2 Function Documentation	109
5.21.2.1 xnmap_create	109
5.21.2.2 xnmap_delete	110
5.21.2.3 xnmap_enter	110
5.21.2.4 xnmap_fetch	111
5.21.2.5 xnmap_fetch_noccheck	111
5.21.2.6 xnmap_remove	112
5.22 Registry services	114
5.22.1 Detailed Description	114
5.22.2 Function Documentation	114
5.22.2.1 xnregistry_bind	114
5.22.2.2 xnregistry_enter	115
5.22.2.3 xnregistry_lookup	116
5.22.2.4 xnregistry_remove	116
5.22.2.5 xnregistry_unlink	117
5.23 Driver programming interface	118
5.23.1 Detailed Description	118
5.24 Wait queue	119
5.24.1 Detailed Description	120
5.24.2 Function Documentation	120
5.24.2.1 rtdm_for_each_waiter	120
5.24.2.2 rtdm_for_each_waiter_safe	120
5.24.2.3 rtdm_timedwait	120
5.24.2.4 rtdm_timedwait_condition	122
5.24.2.5 rtdm_timedwait_condition_locked	123
5.24.2.6 rtdm_timedwait_locked	123
5.24.2.7 rtdm_wait	124
5.24.2.8 rtdm_wait_condition	124
5.24.2.9 rtdm_wait_condition_locked	125
5.24.2.10 rtdm_wait_locked	125
5.24.2.11 rtdm_waitqueue_broadcast	125
5.24.2.12 rtdm_waitqueue_destroy	126
5.24.2.13 rtdm_waitqueue_flush	126
5.24.2.14 rtdm_waitqueue_init	126
5.24.2.15 rtdm_waitqueue_lock	126
5.24.2.16 rtdm_waitqueue_signal	127
5.24.2.17 rtdm_waitqueue_unlock	127
5.24.2.18 rtdm_waitqueue_wakeup	127

5.25 Driver to driver services	128
5.25.1 Detailed Description	129
5.25.2 Function Documentation	129
5.25.2.1 rtdm_accept	129
5.25.2.2 rtdm_bind	129
5.25.2.3 rtdm_close	129
5.25.2.4 rtdm_connect	129
5.25.2.5 rtdm_getpeername	130
5.25.2.6 rtdm_getsockname	130
5.25.2.7 rtdm_getsockopt	130
5.25.2.8 rtdm_ioctl	130
5.25.2.9 rtdm_listen	130
5.25.2.10 rtdm_open	131
5.25.2.11 rtdm_read	131
5.25.2.12 rtdm_recv	131
5.25.2.13 rtdm_recvfrom	131
5.25.2.14 rtdm_recvmsg	131
5.25.2.15 rtdm_send	132
5.25.2.16 rtdm_sendmsg	132
5.25.2.17 rtdm_sendto	132
5.25.2.18 rtdm_setsockopt	132
5.25.2.19 rtdm_shutdown	132
5.25.2.20 rtdm_socket	132
5.25.2.21 rtdm_write	133
5.26 Device Registration Services	134
5.26.1 Detailed Description	135
5.26.2 Macro Definition Documentation	135
5.26.2.1 RTDM_DEVICE_TYPE_MASK	135
5.26.2.2 RTDM_EXCLUSIVE	135
5.26.2.3 RTDM_NAMED_DEVICE	135
5.26.2.4 RTDM_PROTOCOL_DEVICE	135
5.26.3 Function Documentation	135
5.26.3.1 rtdm_dev_register	135
5.26.3.2 rtdm_dev_unregister	136
5.27 Clock Services	137
5.27.1 Detailed Description	137
5.27.2 Function Documentation	137
5.27.2.1 rtdm_clock_read	137
5.27.2.2 rtdm_clock_read_monotonic	138
5.28 Task Services	139

5.28.1 Detailed Description	140
5.28.2 Typedef Documentation	140
5.28.2.1 rtdm_task_proc_t	140
5.28.3 Function Documentation	140
5.28.3.1 rtdm_task_busy_sleep	140
5.28.3.2 rtdm_task_current	141
5.28.3.3 rtdm_task_destroy	141
5.28.3.4 rtdm_task_init	141
5.28.3.5 rtdm_task_join_nrt	142
5.28.3.6 rtdm_task_set_period	142
5.28.3.7 rtdm_task_set_priority	143
5.28.3.8 rtdm_task_should_stop	143
5.28.3.9 rtdm_task_sleep	143
5.28.3.10 rtdm_task_sleep_abs	144
5.28.3.11 rtdm_task_sleep_until	144
5.28.3.12 rtdm_task_unblock	145
5.28.3.13 rtdm_task_wait_period	145
5.29 Timer Services	146
5.29.1 Detailed Description	146
5.29.2 Typedef Documentation	146
5.29.2.1 rtdm_timer_handler_t	146
5.29.3 Enumeration Type Documentation	147
5.29.3.1 rtdm_timer_mode	147
5.29.4 Function Documentation	147
5.29.4.1 rtdm_timer_destroy	147
5.29.4.2 rtdm_timer_init	147
5.29.4.3 rtdm_timer_start	148
5.29.4.4 rtdm_timer_start_in_handler	148
5.29.4.5 rtdm_timer_stop	149
5.29.4.6 rtdm_timer_stop_in_handler	149
5.30 Synchronisation Services	150
5.30.1 Detailed Description	151
5.30.2 Enumeration Type Documentation	151
5.30.2.1 rtdm_selecttype	151
5.30.3 Function Documentation	151
5.30.3.1 rtdm_context_lock	151
5.30.3.2 rtdm_context_put	151
5.30.3.3 rtdm_context_unlock	152
5.30.3.4 rtdm_toseq_init	152
5.31 Event Services	154

5.31.1 Detailed Description	154
5.31.2 Function Documentation	154
5.31.2.1 rtdm_event_clear	154
5.31.2.2 rtdm_event_destroy	155
5.31.2.3 rtdm_event_init	155
5.31.2.4 rtdm_event_pulse	155
5.31.2.5 rtdm_event_select_bind	156
5.31.2.6 rtdm_event_signal	156
5.31.2.7 rtdm_event_timedwait	157
5.31.2.8 rtdm_event_wait	158
5.32 Semaphore Services	159
5.32.1 Detailed Description	159
5.32.2 Function Documentation	159
5.32.2.1 rtdm_sem_destroy	159
5.32.2.2 rtdm_sem_down	160
5.32.2.3 rtdm_sem_init	160
5.32.2.4 rtdm_sem_select_bind	160
5.32.2.5 rtdm_sem_timeddown	161
5.32.2.6 rtdm_sem_up	162
5.33 Mutex services	163
5.33.1 Detailed Description	163
5.33.2 Function Documentation	163
5.33.2.1 rtdm_mutex_destroy	163
5.33.2.2 rtdm_mutex_init	164
5.33.2.3 rtdm_mutex_lock	164
5.33.2.4 rtdm_mutex_timedlock	164
5.33.2.5 rtdm_mutex_unlock	165
5.34 Interrupt Management Services	166
5.34.1 Detailed Description	167
5.34.2 Macro Definition Documentation	167
5.34.2.1 rtdm_irq_get_arg	167
5.34.3 Typedef Documentation	167
5.34.3.1 rtdm_irq_handler_t	167
5.34.4 Function Documentation	167
5.34.4.1 rtdm_irq_disable	167
5.34.4.2 rtdm_irq_enable	168
5.34.4.3 rtdm_irq_free	169
5.34.4.4 rtdm_irq_request	170
5.35 Non-Real-Time Signalling Services	171
5.35.1 Detailed Description	171

5.35.2 Typedef Documentation	171
5.35.2.1 rtdm_nrtsig_handler_t	171
5.35.3 Function Documentation	172
5.35.3.1 rtdm_nrtsig_destroy	172
5.35.3.2 rtdm_nrtsig_init	172
5.35.3.3 rtdm_nrtsig_pend	172
5.36 Utility Services	174
5.36.1 Detailed Description	175
5.36.2 Function Documentation	175
5.36.2.1 rtdm_copy_from_user	175
5.36.2.2 rtdm_copy_to_user	175
5.36.2.3 rtdm_free	176
5.36.2.4 rtdm_in_rt_context	176
5.36.2.5 rtdm_iomap_to_user	176
5.36.2.6 rtdm_malloc	177
5.36.2.7 rtdm_mmap_to_user	178
5.36.2.8 rtdm_munmap	179
5.36.2.9 rtdm_printk	179
5.36.2.10 rtdm_printk_ratelimited	180
5.36.2.11 rtdm_ratelimit	180
5.36.2.12 rtdm_read_user_ok	181
5.36.2.13 rtdm_rt_capable	181
5.36.2.14 rtdm_rw_user_ok	182
5.36.2.15 rtdm_safe_copy_from_user	182
5.36.2.16 rtdm_safe_copy_to_user	183
5.36.2.17 rtdm_strncpy_from_user	184
5.37 Device Profiles	185
5.37.1 Detailed Description	186
5.37.2 Macro Definition Documentation	186
5.37.2.1 RTIOC_DEVICE_INFO	186
5.37.2.2 RTIOC_PURGE	186
5.38 SCHED_QUOTA scheduling policy	188
5.38.1 Detailed Description	188
5.39 File descriptors events multiplexing services	189
5.39.1 Detailed Description	189
5.39.2 Function Documentation	190
5.39.2.1 xselect	190
5.39.2.2 xselect_bind	190
5.39.2.3 xselect_destroy	191
5.39.2.4 xselect_init	191

5.39.2.5	xnselect_signal	192
5.39.2.6	xnselector_destroy	192
5.39.2.7	xnselector_init	192
5.40	Real-time shadow services	193
5.40.1	Detailed Description	193
5.40.2	Function Documentation	194
5.40.2.1	xnshadow_get_context	194
5.40.2.2	xnshadow_harden	194
5.40.2.3	xnshadow_map_kernel	194
5.40.2.4	xnshadow_map_user	195
5.40.2.5	xnshadow_pop_personality	196
5.40.2.6	xnshadow_push_personality	196
5.40.2.7	xnshadow_register_personality	196
5.40.2.8	xnshadow_relax	197
5.41	Thread synchronization services.	198
5.41.1	Detailed Description	198
5.41.2	Function Documentation	199
5.41.2.1	clear_boost	199
5.41.2.2	xnsynch_acquire	199
5.41.2.3	xnsynch_flush	200
5.41.2.4	xnsynch_forget_sleeper	200
5.41.2.5	xnsynch_init	201
5.41.2.6	xnsynch_peek_pendq	201
5.41.2.7	xnsynch_release	202
5.41.2.8	xnsynch_release_all_ownerships	202
5.41.2.9	xnsynch_requeue_sleeper	202
5.41.2.10	xnsynch_sleep_on	203
5.41.2.11	xnsynch_wakeup_one_sleeper	203
5.41.2.12	xnsynch_wakeup_this_sleeper	204
5.42	Thread services	205
5.42.1	Detailed Description	205
5.42.2	Function Documentation	205
5.42.2.1	xnthread_cancel	205
5.42.2.2	xnthread_init	206
5.42.2.3	xnthread_migrate	207
5.42.2.4	xnthread_resume	207
5.42.2.5	xnthread_set_mode	208
5.42.2.6	xnthread_set_periodic	209
5.42.2.7	xnthread_set_schedparam	210
5.42.2.8	xnthread_set_slice	210

5.42.2.9	xnthread_start	211
5.42.2.10	xnthread_suspend	212
5.42.2.11	xnthread_unblock	213
5.42.2.12	xnthread_wait_period	213
5.43	Timer services	214
5.43.1	Detailed Description	214
5.43.2	Function Documentation	215
5.43.2.1	program_htick_shot	215
5.43.2.2	switch_htick_mode	215
5.43.2.3	xntimer_destroy	216
5.43.2.4	xntimer_get_date	216
5.43.2.5	xntimer_get_overruns	216
5.43.2.6	xntimer_get_timeout	217
5.43.2.7	xntimer_grab_hardware	217
5.43.2.8	xntimer_init	218
5.43.2.9	xntimer_interval	218
5.43.2.10	xntimer_release_hardware	218
5.43.2.11	xntimer_start	219
5.43.2.12	xntimer_stop	219
5.44	Virtual file services	220
5.44.1	Detailed Description	221
5.44.2	Function Documentation	221
5.44.2.1	xnvfile_destroy	221
5.44.2.2	xnvfile_get_blob	222
5.44.2.3	xnvfile_get_integer	222
5.44.2.4	xnvfile_get_string	223
5.44.2.5	xnvfile_init_dir	223
5.44.2.6	xnvfile_init_link	224
5.44.2.7	xnvfile_init_regular	225
5.44.2.8	xnvfile_init_snapshot	225
5.44.3	Variable Documentation	226
5.44.3.1	nkvfroot	226
5.44.3.2	nkvfroot	226
5.45	Analogy interface	227
5.45.1	Detailed Description	227
5.46	Driver API	228
5.46.1	Detailed Description	228
5.47	Driver management services	229
5.47.1	Detailed Description	229
5.47.2	Function Documentation	229

5.47.2.1	a4l_register_drv	229
5.47.2.2	a4l_unregister_drv	229
5.48	Subdevice management services	231
5.48.1	Detailed Description	232
5.48.2	Function Documentation	233
5.48.2.1	a4l_add_subd	233
5.48.2.2	a4l_alloc_subd	233
5.48.2.3	a4l_get_subd	233
5.49	Buffer management services	235
5.49.1	Detailed Description	236
5.49.2	Function Documentation	236
5.49.2.1	a4l_buf_commit_absget	236
5.49.2.2	a4l_buf_commit_absput	237
5.49.2.3	a4l_buf_commit_get	237
5.49.2.4	a4l_buf_commit_put	238
5.49.2.5	a4l_buf_count	238
5.49.2.6	a4l_buf_evt	238
5.49.2.7	a4l_buf_get	239
5.49.2.8	a4l_buf_prepare_absget	239
5.49.2.9	a4l_buf_prepare_absput	239
5.49.2.10	a4l_buf_prepare_get	240
5.49.2.11	a4l_buf_prepare_put	240
5.49.2.12	a4l_buf_put	240
5.49.2.13	a4l_get_chan	242
5.49.2.14	a4l_get_cmd	242
5.50	Interrupt management services	243
5.50.1	Detailed Description	243
5.50.2	Function Documentation	243
5.50.2.1	a4l_free_irq	243
5.50.2.2	a4l_get_irq	243
5.50.2.3	a4l_request_irq	244
5.51	Misc services	246
5.51.1	Detailed Description	246
5.51.2	Function Documentation	246
5.51.2.1	a4l_get_time	246
5.52	Clocks and timers	247
5.52.1	Detailed Description	247
5.52.2	Function Documentation	247
5.52.2.1	timer_create	248
5.52.2.2	timer_gettime	248

5.52.2.3 timer_settime	249
5.53 Condition variables	251
5.53.1 Detailed Description	251
5.53.2 Function Documentation	251
5.53.2.1 pthread_cond_destroy	251
5.53.2.2 pthread_cond_init	252
5.54 Cobalt/POSIX real-time API	253
5.54.1 Detailed Description	253
5.55 Message queues	254
5.55.1 Detailed Description	254
5.55.2 Function Documentation	254
5.55.2.1 mq_close	254
5.55.2.2 mq_getattr	255
5.55.2.3 mq_notify	256
5.55.2.4 mq_open	256
5.55.2.5 mq_setattr	257
5.55.2.6 mq_unlink	258
5.56 Mutual exclusion	259
5.56.1 Detailed Description	259
5.57 Semaphores	260
5.57.1 Detailed Description	260
5.57.2 Function Documentation	260
5.57.2.1 sem_close	260
5.57.2.2 sem_destroy	261
5.57.2.3 sem_post	261
5.57.2.4 sem_timedwait	263
5.57.2.5 sem_trywait	264
5.57.2.6 sem_unlink	264
5.57.2.7 sem_wait	265
5.58 Thread management	266
5.58.1 Detailed Description	266
5.58.2 Function Documentation	266
5.58.2.1 pthread_create	266
5.58.2.2 pthread_getschedparam_ex	267
5.58.2.3 pthread_make_periodic_np	268
5.58.2.4 pthread_set_mode_np	268
5.58.2.5 pthread_setschedparam_ex	269
5.59 Scheduling management	271
5.59.1 Detailed Description	271
5.59.2 Function Documentation	271

5.59.2.1 sched_getconfig_np	271
5.59.2.2 sched_setconfig_np	272
5.60 Asynchronous acquisition API	274
5.60.1 Detailed Description	276
5.60.2 Function Documentation	276
5.60.2.1 a4l_get_bufsize	276
5.60.2.2 a4l_mark_bufw	276
5.60.2.3 a4l_mmap	277
5.60.2.4 a4l_poll	277
5.60.2.5 a4l_set_bufsize	278
5.60.2.6 a4l_snd_cancel	278
5.60.2.7 a4l_snd_command	279
5.61 Asynchronous acquisition API	280
5.61.1 Detailed Description	280
5.61.2 Function Documentation	280
5.61.2.1 a4l_async_read	280
5.61.2.2 a4l_async_write	281
5.62 Descriptor Syscall API	282
5.62.1 Detailed Description	282
5.62.2 Function Documentation	282
5.62.2.1 a4l_sys_desc	282
5.63 Descriptor API	284
5.63.1 Detailed Description	284
5.63.2 Function Documentation	284
5.63.2.1 a4l_close	284
5.63.2.2 a4l_fill_desc	285
5.63.2.3 a4l_get_chinfo	285
5.63.2.4 a4l_get_rnginfo	285
5.63.2.5 a4l_get_subdinfo	286
5.63.2.6 a4l_open	286
5.64 Range / conversion API	287
5.64.1 Detailed Description	287
5.64.2 Function Documentation	287
5.64.2.1 a4l_dtoraw	287
5.64.2.2 a4l_find_range	288
5.64.2.3 a4l_ftoraw	288
5.64.2.4 a4l_rawtod	289
5.64.2.5 a4l_rawtof	289
5.64.2.6 a4l_rawtoul	290
5.64.2.7 a4l_sizeof_chan	290

5.64.2.8	a4l_sizeof_subd	290
5.64.2.9	a4l_ultoraw	291
5.65	Level 1 API	292
5.65.1	Detailed Description	292
5.66	Synchronous acquisition API	293
5.66.1	Detailed Description	295
5.66.2	Function Documentation	295
5.66.2.1	a4l_snd_insn	295
5.66.2.2	a4l_snd_insnlist	295
5.67	Level 2 API	297
5.67.1	Detailed Description	297
5.68	Synchronous acquisition API	298
5.68.1	Detailed Description	298
5.68.2	Function Documentation	298
5.68.2.1	a4l_config_subd	298
5.68.2.2	a4l_sync_dio	299
5.68.2.3	a4l_sync_read	299
5.68.2.4	a4l_sync_write	300
5.69	Analogy user API	301
5.69.1	Detailed Description	301
5.70	Level 0 API	302
5.70.1	Detailed Description	302
5.71	Basic Syscall API	303
5.71.1	Detailed Description	303
5.71.2	Function Documentation	303
5.71.2.1	a4l_sys_close	303
5.71.2.2	a4l_sys_open	303
5.71.2.3	a4l_sys_read	304
5.71.2.4	a4l_sys_write	304
5.72	Attach / detach Syscall API	305
5.72.1	Detailed Description	305
5.72.2	Function Documentation	305
5.72.2.1	a4l_sys_attach	305
5.72.2.2	a4l_sys_bufcfg	305
5.72.2.3	a4l_sys_detach	306
5.73	Alarm services	307
5.73.1	Detailed Description	307
5.73.2	Function Documentation	307
5.73.2.1	rt_alarm_create	307
5.73.2.2	rt_alarm_delete	308

5.73.2.3	rt_alarm_inquire	308
5.73.2.4	rt_alarm_start	309
5.73.2.5	rt_alarm_stop	309
5.74	Buffer services	311
5.74.1	Detailed Description	312
5.74.2	Macro Definition Documentation	312
5.74.2.1	B_PRIO	312
5.74.3	Function Documentation	312
5.74.3.1	rt_buffer_bind	312
5.74.3.2	rt_buffer_clear	313
5.74.3.3	rt_buffer_create	313
5.74.3.4	rt_buffer_delete	314
5.74.3.5	rt_buffer_inquire	314
5.74.3.6	rt_buffer_read	315
5.74.3.7	rt_buffer_read_timed	315
5.74.3.8	rt_buffer_read_until	316
5.74.3.9	rt_buffer_unbind	317
5.74.3.10	rt_buffer_write	317
5.74.3.11	rt_buffer_write_timed	317
5.74.3.12	rt_buffer_write_until	318
5.75	Condition variable services	319
5.75.1	Detailed Description	319
5.75.2	Function Documentation	320
5.75.2.1	rt_cond_bind	320
5.75.2.2	rt_cond_broadcast	320
5.75.2.3	rt_cond_create	321
5.75.2.4	rt_cond_delete	321
5.75.2.5	rt_cond_inquire	322
5.75.2.6	rt_cond_signal	322
5.75.2.7	rt_cond_unbind	323
5.75.2.8	rt_cond_wait	323
5.75.2.9	rt_cond_wait_timed	323
5.75.2.10	rt_cond_wait_until	324
5.76	Event flag group services	325
5.76.1	Detailed Description	326
5.76.2	Macro Definition Documentation	326
5.76.2.1	EV_ANY	326
5.76.2.2	EV_PRIO	326
5.76.3	Function Documentation	326
5.76.3.1	rt_event_bind	326

5.76.3.2	rt_event_clear	327
5.76.3.3	rt_event_create	327
5.76.3.4	rt_event_delete	328
5.76.3.5	rt_event_inquire	328
5.76.3.6	rt_event_signal	329
5.76.3.7	rt_event_unbind	329
5.76.3.8	rt_event_wait	329
5.76.3.9	rt_event_wait_timed	330
5.76.3.10	rt_event_wait_until	331
5.77	Heap management services	332
5.77.1	Detailed Description	333
5.77.2	Macro Definition Documentation	333
5.77.2.1	H_PRIO	333
5.77.3	Function Documentation	333
5.77.3.1	rt_heap_alloc	333
5.77.3.2	rt_heap_alloc_timed	333
5.77.3.3	rt_heap_alloc_until	334
5.77.3.4	rt_heap_bind	334
5.77.3.5	rt_heap_create	335
5.77.3.6	rt_heap_delete	336
5.77.3.7	rt_heap_free	337
5.77.3.8	rt_heap_inquire	338
5.77.3.9	rt_heap_unbind	338
5.78	Alchemy API	339
5.78.1	Detailed Description	340
5.79	Mutex services	341
5.79.1	Detailed Description	341
5.79.2	Function Documentation	342
5.79.2.1	rt_mutex_acquire	342
5.79.2.2	rt_mutex_acquire_timed	342
5.79.2.3	rt_mutex_acquire_until	343
5.79.2.4	rt_mutex_bind	343
5.79.2.5	rt_mutex_create	344
5.79.2.6	rt_mutex_delete	344
5.79.2.7	rt_mutex_inquire	345
5.79.2.8	rt_mutex_release	345
5.79.2.9	rt_mutex_unbind	345
5.80	Message pipe services	347
5.80.1	Detailed Description	347
5.80.2	Macro Definition Documentation	348

5.80.2.1	P_MINOR_AUTO	348
5.80.2.2	P_URGENT	348
5.80.3	Function Documentation	348
5.80.3.1	rt_pipe_bind	348
5.80.3.2	rt_pipe_create	349
5.80.3.3	rt_pipe_delete	350
5.80.3.4	rt_pipe_read	350
5.80.3.5	rt_pipe_read_timed	351
5.80.3.6	rt_pipe_read_until	351
5.80.3.7	rt_pipe_unbind	352
5.80.3.8	rt_pipe_write	352
5.81	Message queue services	354
5.81.1	Detailed Description	355
5.81.2	Macro Definition Documentation	355
5.81.2.1	Q_PRIO	355
5.81.2.2	Q_SHARED	355
5.81.3	Function Documentation	355
5.81.3.1	rt_queue_alloc	355
5.81.3.2	rt_queue_bind	356
5.81.3.3	rt_queue_create	356
5.81.3.4	rt_queue_delete	357
5.81.3.5	rt_queue_flush	358
5.81.3.6	rt_queue_free	358
5.81.3.7	rt_queue_inquire	359
5.81.3.8	rt_queue_read	360
5.81.3.9	rt_queue_read_timed	360
5.81.3.10	rt_queue_read_until	361
5.81.3.11	rt_queue_receive	361
5.81.3.12	rt_queue_receive_timed	362
5.81.3.13	rt_queue_receive_until	363
5.81.3.14	rt_queue_send	363
5.81.3.15	rt_queue_unbind	364
5.82	Semaphore services	366
5.82.1	Detailed Description	367
5.82.2	Macro Definition Documentation	367
5.82.2.1	S_PRIO	367
5.82.3	Function Documentation	367
5.82.3.1	rt_sem_bind	367
5.82.3.2	rt_sem_broadcast	368
5.82.3.3	rt_sem_create	368

5.82.3.4	<code>rt_sem_delete</code>	369
5.82.3.5	<code>rt_sem_inquire</code>	369
5.82.3.6	<code>rt_sem_p</code>	370
5.82.3.7	<code>rt_sem_p_timed</code>	370
5.82.3.8	<code>rt_sem_p_until</code>	371
5.82.3.9	<code>rt_sem_unbind</code>	371
5.82.3.10	<code>rt_sem_v</code>	371
5.83	Task management services	372
5.83.1	Detailed Description	373
5.83.2	Macro Definition Documentation	374
5.83.2.1	<code>T_FPU</code>	374
5.83.2.2	<code>T_LOCK</code>	374
5.83.2.3	<code>T_LOPRIO</code>	374
5.83.2.4	<code>T_WARNSW</code>	374
5.83.3	Function Documentation	374
5.83.3.1	<code>rt_task_bind</code>	374
5.83.3.2	<code>rt_task_create</code>	375
5.83.3.3	<code>rt_task_delete</code>	376
5.83.3.4	<code>rt_task_inquire</code>	376
5.83.3.5	<code>rt_task_join</code>	377
5.83.3.6	<code>rt_task_receive</code>	378
5.83.3.7	<code>rt_task_receive_timed</code>	378
5.83.3.8	<code>rt_task_receive_until</code>	379
5.83.3.9	<code>rt_task_reply</code>	379
5.83.3.10	<code>rt_task_resume</code>	380
5.83.3.11	<code>rt_task_same</code>	380
5.83.3.12	<code>rt_task_self</code>	381
5.83.3.13	<code>rt_task_send</code>	381
5.83.3.14	<code>rt_task_send_timed</code>	381
5.83.3.15	<code>rt_task_send_until</code>	383
5.83.3.16	<code>rt_task_set_affinity</code>	383
5.83.3.17	<code>rt_task_set_mode</code>	384
5.83.3.18	<code>rt_task_set_periodic</code>	385
5.83.3.19	<code>rt_task_set_priority</code>	385
5.83.3.20	<code>rt_task_shadow</code>	386
5.83.3.21	<code>rt_task_sleep</code>	387
5.83.3.22	<code>rt_task_sleep_until</code>	388
5.83.3.23	<code>rt_task_slice</code>	389
5.83.3.24	<code>rt_task_spawn</code>	390
5.83.3.25	<code>rt_task_start</code>	390

5.83.3.26	rt_task_suspend	391
5.83.3.27	rt_task_unbind	391
5.83.3.28	rt_task_unblock	392
5.83.3.29	rt_task_wait_period	392
5.83.3.30	rt_task_yield	393
5.84	Timer management services	394
5.84.1	Detailed Description	394
5.84.2	Typedef Documentation	394
5.84.2.1	RT_TIMER_INFO	394
5.84.3	Function Documentation	395
5.84.3.1	rt_timer_inquire	395
5.84.3.2	rt_timer_ns2ticks	395
5.84.3.3	rt_timer_read	395
5.84.3.4	rt_timer_spin	396
5.84.3.5	rt_timer_ticks2ns	396
6	Data Structure Documentation	397
6.1	a4l_channel Struct Reference	397
6.1.1	Detailed Description	397
6.1.2	Field Documentation	397
6.1.2.1	flags	397
6.1.2.2	nb_bits	397
6.2	a4l_channels_desc Struct Reference	397
6.2.1	Detailed Description	398
6.2.2	Field Documentation	398
6.2.2.1	chans	398
6.2.2.2	length	398
6.2.2.3	mode	398
6.3	a4l_cmd_desc Struct Reference	398
6.3.1	Detailed Description	399
6.3.2	Field Documentation	399
6.3.2.1	idx_subd	399
6.4	a4l_descriptor Struct Reference	400
6.4.1	Detailed Description	400
6.4.2	Field Documentation	400
6.4.2.1	board_name	400
6.4.2.2	fd	400
6.4.2.3	idx_read_subd	400
6.4.2.4	idx_write_subd	401
6.4.2.5	magic	401

6.4.2.6	nb_subd	401
6.4.2.7	sbdata	401
6.4.2.8	sbsize	401
6.5	a4l_driver Struct Reference	401
6.5.1	Detailed Description	402
6.6	a4l_instruction Struct Reference	402
6.6.1	Detailed Description	402
6.6.2	Field Documentation	402
6.6.2.1	idx_subd	402
6.7	a4l_instruction_list Struct Reference	403
6.7.1	Detailed Description	403
6.8	a4l_range Struct Reference	403
6.8.1	Detailed Description	404
6.8.2	Field Documentation	404
6.8.2.1	flags	404
6.8.2.2	max	404
6.8.2.3	min	404
6.9	a4l_subdevice Struct Reference	404
6.9.1	Detailed Description	405
6.10	atomic_long_t Struct Reference	406
6.10.1	Detailed Description	406
6.11	can_bittime Struct Reference	406
6.11.1	Detailed Description	407
6.12	can_bittime_btr Struct Reference	407
6.12.1	Detailed Description	407
6.13	can_bittime_std Struct Reference	407
6.13.1	Detailed Description	408
6.14	can_filter Struct Reference	408
6.14.1	Detailed Description	408
6.14.2	Field Documentation	408
6.14.2.1	can_id	408
6.14.2.2	can_mask	409
6.15	can_frame Struct Reference	409
6.15.1	Detailed Description	409
6.15.2	Field Documentation	409
6.15.2.1	can_id	409
6.16	RT_ALARM_INFO Struct Reference	410
6.16.1	Detailed Description	410
6.17	RT_BUFFER_INFO Struct Reference	410
6.17.1	Detailed Description	410

6.18 RT_COND_INFO Struct Reference	411
6.18.1 Detailed Description	411
6.19 RT_EVENT_INFO Struct Reference	411
6.19.1 Detailed Description	411
6.20 RT_HEAP_INFO Struct Reference	411
6.20.1 Detailed Description	412
6.20.2 Field Documentation	412
6.20.2.1 usablemem	412
6.21 RT_MUTEX_INFO Struct Reference	412
6.21.1 Detailed Description	412
6.21.2 Field Documentation	413
6.21.2.1 owner	413
6.22 RT_QUEUE_INFO Struct Reference	413
6.22.1 Detailed Description	413
6.23 RT_SEM_INFO Struct Reference	413
6.23.1 Detailed Description	414
6.24 RT_TASK_INFO Struct Reference	414
6.24.1 Detailed Description	414
6.25 rt_timer_info Struct Reference	414
6.25.1 Detailed Description	415
6.25.2 Field Documentation	415
6.25.2.1 date	415
6.25.2.2 tsc	415
6.26 rtdm_dev_context Struct Reference	415
6.26.1 Detailed Description	416
6.26.2 Field Documentation	416
6.26.2.1 device	416
6.27 rtdm_device Struct Reference	416
6.27.1 Detailed Description	417
6.27.2 Field Documentation	418
6.27.2.1 open	418
6.27.2.2 socket	418
6.28 rtdm_device_info Struct Reference	418
6.28.1 Detailed Description	418
6.29 rtipc_port_label Struct Reference	418
6.29.1 Detailed Description	419
6.29.2 Field Documentation	419
6.29.2.1 label	419
6.30 rtser_config Struct Reference	419
6.30.1 Detailed Description	420

6.31 rtser_event Struct Reference	420
6.31.1 Detailed Description	420
6.32 rtser_status Struct Reference	420
6.32.1 Detailed Description	421
6.33 sockaddr_can Struct Reference	421
6.33.1 Detailed Description	421
6.33.2 Field Documentation	421
6.33.2.1 can_ifindex	421
6.34 sockaddr_ipc Struct Reference	421
6.34.1 Detailed Description	422
6.34.2 Field Documentation	422
6.34.2.1 sipc_port	422
6.35 xnsched Struct Reference	422
6.35.1 Detailed Description	422
6.35.2 Field Documentation	422
6.35.2.1 curr	422
6.35.2.2 htimer	422
6.35.2.3 inesting	423
6.35.2.4 lflags	423
6.35.2.5 rrbtimer	423
6.35.2.6 rt	423
6.35.2.7 status	423
6.36 xnfile_lock_ops Struct Reference	423
6.36.1 Detailed Description	423
6.36.2 Field Documentation	423
6.36.2.1 get	423
6.36.2.2 put	424
6.37 xnfile_regular_iterator Struct Reference	424
6.37.1 Detailed Description	424
6.37.2 Field Documentation	424
6.37.2.1 pos	424
6.37.2.2 private	425
6.37.2.3 seq	425
6.37.2.4 vfile	425
6.38 xnfile_regular_ops Struct Reference	425
6.38.1 Detailed Description	425
6.38.2 Field Documentation	425
6.38.2.1 begin	425
6.38.2.2 end	426
6.38.2.3 next	426

6.38.2.4	rewind	427
6.38.2.5	show	427
6.38.2.6	store	427
6.39	xnvfile_rev_tag Struct Reference	428
6.39.1	Detailed Description	428
6.39.2	Field Documentation	428
6.39.2.1	rev	428
6.40	xnvfile_snapshot Struct Reference	428
6.40.1	Detailed Description	429
6.41	xnvfile_snapshot_iterator Struct Reference	429
6.41.1	Detailed Description	430
6.41.2	Field Documentation	430
6.41.2.1	databuf	430
6.41.2.2	endfn	430
6.41.2.3	nrdata	431
6.41.2.4	private	431
6.41.2.5	seq	431
6.41.2.6	vfile	431
6.42	xnvfile_snapshot_ops Struct Reference	431
6.42.1	Detailed Description	431
6.42.2	Field Documentation	431
6.42.2.1	begin	431
6.42.2.2	end	432
6.42.2.3	next	432
6.42.2.4	rewind	433
6.42.2.5	show	433
6.42.2.6	store	434
7	File Documentation	435
7.1	include/cobalt/kernel/rtdm/analogy/buffer.h File Reference	435
7.1.1	Detailed Description	435
7.2	include/cobalt/kernel/timer.h File Reference	436
7.2.1	Detailed Description	437
7.3	include/cobalt/kernel/sched.h File Reference	437
7.3.1	Detailed Description	438
7.3.2	Function Documentation	438
7.3.2.1	xnsched_rotate	439
7.3.2.2	xnsched_run	439
7.4	include/cobalt/kernel/apc.h File Reference	440
7.4.1	Detailed Description	441

7.5	include/cobalt/kernel/bufd.h File Reference	441
7.5.1	Detailed Description	442
7.6	include/cobalt/kernel/clock.h File Reference	443
7.6.1	Detailed Description	444
7.7	include/cobalt/kernel/map.h File Reference	444
7.7.1	Detailed Description	445
7.8	include/cobalt/kernel/ppd.h File Reference	446
7.8.1	Detailed Description	446
7.9	include/cobalt/kernel/registry.h File Reference	447
7.9.1	Detailed Description	448
7.10	include/cobalt/kernel/rtdm/analogy/channel_range.h File Reference	448
7.10.1	Detailed Description	450
7.11	include/cobalt/kernel/rtdm/analogy/context.h File Reference	450
7.11.1	Detailed Description	451
7.12	include/cobalt/kernel/rtdm/analogy/device.h File Reference	451
7.12.1	Detailed Description	452
7.13	include/cobalt/kernel/rtdm/analogy/driver.h File Reference	452
7.13.1	Detailed Description	452
7.14	include/cobalt/kernel/rtdm/driver.h File Reference	453
7.14.1	Detailed Description	457
7.14.2	Typedef Documentation	457
7.14.2.1	rtdm_open_handler_t	457
7.14.2.2	rtdm_socket_handler_t	458
7.14.3	Function Documentation	458
7.14.3.1	rtdm_fd_device	458
7.14.3.2	rtdm_fd_is_user	458
7.14.3.3	rtdm_fd_to_private	459
7.14.3.4	rtdm_private_to_fd	459
7.15	include/cobalt/kernel/rtdm/analogy/instruction.h File Reference	459
7.15.1	Detailed Description	459
7.16	include/cobalt/kernel/rtdm/analogy/rtdm_helpers.h File Reference	460
7.16.1	Detailed Description	460
7.17	include/cobalt/kernel/rtdm/analogy/subdevice.h File Reference	461
7.17.1	Detailed Description	461
7.18	include/cobalt/kernel/rtdm/analogy/transfer.h File Reference	462
7.18.1	Detailed Description	462
7.19	include/cobalt/kernel/rtdm/can.h File Reference	463
7.19.1	Detailed Description	463
7.20	include/rtdm/can.h File Reference	463
7.20.1	Detailed Description	464

7.21 include/rtdm/uapi/can.h File Reference	464
7.21.1 Detailed Description	470
7.22 include/cobalt/kernel/rtdm/cobalt.h File Reference	471
7.22.1 Detailed Description	471
7.23 include/cobalt/kernel/rtdm/ipc.h File Reference	471
7.23.1 Detailed Description	472
7.24 include/rtdm/ipc.h File Reference	472
7.24.1 Detailed Description	472
7.25 include/rtdm/uapi/ipc.h File Reference	473
7.25.1 Detailed Description	475
7.26 include/cobalt/kernel/rtdm/rtdm.h File Reference	476
7.26.1 Detailed Description	476
7.27 include/rtdm/rtdm.h File Reference	476
7.27.1 Detailed Description	477
7.28 include/rtdm/uapi/rtdm.h File Reference	477
7.28.1 Detailed Description	479
7.29 include/cobalt/kernel/rtdm/serial.h File Reference	479
7.29.1 Detailed Description	479
7.30 include/rtdm/serial.h File Reference	479
7.30.1 Detailed Description	480
7.31 include/rtdm/uapi/serial.h File Reference	480
7.31.1 Detailed Description	484
7.31.2 Macro Definition Documentation	484
7.31.2.1 RTSER_RTIOC_BREAK_CTL	484
7.31.2.2 RTSER_RTIOC_GET_CONFIG	485
7.31.2.3 RTSER_RTIOC_GET_CONTROL	485
7.31.2.4 RTSER_RTIOC_GET_STATUS	486
7.31.2.5 RTSER_RTIOC_SET_CONFIG	487
7.31.2.6 RTSER_RTIOC_SET_CONTROL	488
7.31.2.7 RTSER_RTIOC_WAIT_EVENT	488
7.32 include/cobalt/kernel/rtdm/testing.h File Reference	489
7.32.1 Detailed Description	489
7.33 include/rtdm/testing.h File Reference	489
7.33.1 Detailed Description	490
7.34 include/rtdm/uapi/testing.h File Reference	490
7.34.1 Detailed Description	491
7.35 include/cobalt/kernel/sched-idle.h File Reference	492
7.35.1 Detailed Description	492
7.36 include/cobalt/kernel/sched-rt.h File Reference	493
7.36.1 Detailed Description	493

7.37 include/cobalt/kernel/sched-sporadic.h File Reference	493
7.37.1 Detailed Description	494
7.38 include/cobalt/kernel/select.h File Reference	494
7.38.1 Detailed Description	496
7.39 include/cobalt/kernel/vfile.h File Reference	496
7.39.1 Detailed Description	498
7.40 include/rtdm/analogy.h File Reference	498
7.40.1 Detailed Description	499
7.41 include/rtdm/uapi/analogy.h File Reference	499
7.41.1 Detailed Description	504
7.41.2 Macro Definition Documentation	504
7.41.2.1 A4L_RNG_FACTOR	504
7.42 kernel/cobalt/apc.c File Reference	505
7.42.1 Detailed Description	505
7.43 kernel/cobalt/posix/syscall.c File Reference	505
7.43.1 Detailed Description	506
7.44 kernel/cobalt/arith.c File Reference	506
7.44.1 Detailed Description	506
7.45 kernel/cobalt/bufd.c File Reference	507
7.45.1 Detailed Description	507
7.46 kernel/cobalt/rtdm/init.c File Reference	508
7.46.1 Detailed Description	508
7.47 kernel/cobalt/map.c File Reference	508
7.47.1 Detailed Description	509
7.48 lib/analogy/internal.h File Reference	509
7.48.1 Detailed Description	510
7.49 kernel/cobalt/posix/select.c File Reference	510
7.49.1 Detailed Description	510
7.50 kernel/cobalt/timer.c File Reference	511
7.50.1 Detailed Description	511
7.51 kernel/cobalt/registry.c File Reference	512
7.51.1 Detailed Description	512
7.52 kernel/cobalt/rtdm/core.c File Reference	513
7.52.1 Detailed Description	516
7.53 kernel/cobalt/rtdm/device.c File Reference	516
7.53.1 Detailed Description	516
7.54 kernel/drivers/analogy/device.c File Reference	517
7.54.1 Detailed Description	517
7.55 kernel/cobalt/rtdm/drvlib.c File Reference	517
7.55.1 Detailed Description	521

7.56 kernel/cobalt/vfile.c File Reference	521
7.56.1 Detailed Description	522
7.57 kernel/drivers/analogy/buffer.c File Reference	522
7.57.1 Detailed Description	523
7.58 kernel/drivers/analogy/command.c File Reference	524
7.58.1 Detailed Description	524
7.59 kernel/drivers/analogy/driver.c File Reference	524
7.59.1 Detailed Description	525
7.60 kernel/drivers/analogy/driver_facilities.c File Reference	525
7.60.1 Detailed Description	527
7.61 kernel/drivers/analogy/instruction.c File Reference	527
7.61.1 Detailed Description	527
7.62 kernel/drivers/analogy/intel/8255.c File Reference	528
7.62.1 Detailed Description	528
7.63 kernel/drivers/analogy/intel/8255.h File Reference	528
7.63.1 Detailed Description	529
7.64 kernel/drivers/analogy/intel/parport.c File Reference	529
7.64.1 Detailed Description	530
7.65 kernel/drivers/analogy/national_instruments/mio_common.c File Reference	530
7.65.1 Detailed Description	531
7.66 kernel/drivers/analogy/national_instruments/mite.c File Reference	531
7.66.1 Detailed Description	531
7.67 kernel/drivers/analogy/national_instruments/mite.h File Reference	532
7.67.1 Detailed Description	533
7.68 kernel/drivers/analogy/national_instruments/ni_mio.h File Reference	533
7.68.1 Detailed Description	533
7.69 kernel/drivers/analogy/national_instruments/ni_stc.h File Reference	534
7.69.1 Detailed Description	534
7.70 kernel/drivers/analogy/national_instruments/ni_tio.h File Reference	535
7.70.1 Detailed Description	535
7.71 kernel/drivers/analogy/national_instruments/pcmio.c File Reference	536
7.71.1 Detailed Description	536
7.72 kernel/drivers/analogy/national_instruments/tio_common.c File Reference	537
7.72.1 Detailed Description	538
7.73 kernel/drivers/analogy/proc.h File Reference	538
7.73.1 Detailed Description	539
7.74 kernel/drivers/analogy/rtdm_helpers.c File Reference	539
7.74.1 Detailed Description	540
7.75 kernel/drivers/analogy/rtdm_interface.c File Reference	540
7.75.1 Detailed Description	540

7.76 kernel/drivers/analog/sensoray/s526.c File Reference	540
7.76.1 Detailed Description	541
7.77 kernel/drivers/analog/subdevice.c File Reference	541
7.77.1 Detailed Description	542
7.78 kernel/drivers/analog/transfer.c File Reference	542
7.78.1 Detailed Description	543
7.79 lib/analog/async.c File Reference	543
7.79.1 Detailed Description	544
7.80 lib/analog/descriptor.c File Reference	544
7.80.1 Detailed Description	545
7.81 lib/analog/info.c File Reference	545
7.81.1 Detailed Description	545
7.82 lib/analog/range.c File Reference	546
7.82.1 Detailed Description	547
7.83 lib/analog/root_leaf.h File Reference	547
7.83.1 Detailed Description	547
7.84 lib/analog/sync.c File Reference	548
7.84.1 Detailed Description	548
7.85 lib/analog/sys.c File Reference	549
7.85.1 Detailed Description	549
8 Example Documentation	551
8.1 bufp-label.c	551
8.2 bufp-readwrite.c	553
8.3 can-rtt.c	556
8.4 cross-link.c	561
8.5 iddp-label.c	565
8.6 iddp-sendrecv.c	567
8.7 rtcanconfig.c	570
8.8 rtcanrecv.c	573
8.9 rtcanseend.c	577
8.10 xddp-echo.c	580
8.11 xddp-label.c	583
8.12 xddp-stream.c	587
Index	591

Chapter 1

Deprecated List

Global [rtdm_task_sleep_until](#) (nanosecs_abs_t wakeup_time)

Use rtdm_task_sleep_abs instead!

Chapter 2

Module Index

2.1 Modules

Here is a list of all modules:

Real-Time Driver Model	50
RTDM User API	52
Driver programming interface	118
Driver to driver services	128
Device Registration Services	134
Clock Services	137
Task Services	139
Timer Services	146
Synchronisation Services	150
Big dual kernel lock	18
Spinlock with preemption deactivation	21
Wait queue	119
Event Services	154
Semaphore Services	159
Mutex services	163
Interrupt Management Services	166
Non-Real-Time Signalling Services	171
Utility Services	174
Device Profiles	185
CAN Devices	29
Serial Devices	65
Testing Devices	66
Real-time IPC	67
Xenomai nucleus	102
Thread state flags	25
Thread information flags	28
Asynchronous Procedure Calls	82
Helpers for in-kernel arithmetics	84
Buffer descriptor	85
Clock services.	93
Debugging services	96
Dynamic memory allocation services.	97
Interrupt management	104
SMP locking services	108
Lightweight key-to-object mapping service	109
Registry services	114
File descriptors events multiplexing services	189

Real-time shadow services	193
Thread synchronization services.	198
Thread services	205
Timer services	214
Virtual file services	220
SCHED_QUOTA scheduling policy	188
Analogy interface	227
Driver API	228
Channels and ranges	15
Driver management services	229
Subdevice management services	231
Buffer management services	235
Interrupt management services	243
Misc services	246
Analogy user API	301
Level 1 API	292
Asynchronous acquisition API	274
Descriptor API	284
Synchronous acquisition API	293
Level 2 API	297
Asynchronous acquisition API	280
Range / conversion API	287
Synchronous acquisition API	298
Level 0 API	302
Descriptor Syscall API	282
Basic Syscall API	303
Attach / detach Syscall API	305
Cobalt/POSIX real-time API	253
Clocks and timers	247
Condition variables	251
Message queues	254
Mutual exclusion	259
Semaphores	260
Thread management	266
Scheduling management	271
Alchemy API	339
Alarm services	307
Buffer services	311
Condition variable services	319
Event flag group services	325
Heap management services	332
Mutex services	341
Message pipe services	347
Message queue services	354
Semaphore services	366
Task management services	372
Timer management services	394

Chapter 3

Data Structure Index

3.1 Data Structures

Here are the data structures with brief descriptions:

a4l_channel	Structure describing some channel's characteristics	397
a4l_channels_desc	Structure describing a channels set	397
a4l_cmd_desc	Structure describing the asynchronous instruction	398
a4l_descriptor	Structure containing device-information useful to users	400
a4l_driver	Structure containing driver declaration data	401
a4l_instruction	Structure describing the synchronous instruction	402
a4l_instruction_list	Structure describing the list of synchronous instructions	403
a4l_range	Structure describing a (unique) range	403
a4l_subdevice	Structure describing the subdevice	404
atomic_long_t	Copyright © 2011 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org	406
can_bittime	Custom CAN bit-time definition	406
can_bittime_btr	Hardware-specific BTR bit-times	407
can_bittime_std	Standard bit-time parameters according to Bosch	407
can_filter	Filter for reception of CAN messages	408
can_frame	Raw CAN frame	409
RT_ALARM_INFO	Alarm status descriptor	410
RT_BUFFER_INFO	Buffer status descriptor	410
RT_COND_INFO	Condition variable status descriptor	411
RT_EVENT_INFO	Event status descriptor	411

RT_HEAP_INFO	
Heap status descriptor	411
RT_MUTEX_INFO	
Mutex status descriptor	412
RT_QUEUE_INFO	
Queue status descriptor	413
RT_SEM_INFO	
Semaphore status descriptor	413
RT_TASK_INFO	
Task status descriptor	414
rt_timer_info	
Timer status descriptor	414
rtdm_dev_context	
Device context	415
rtdm_device	
RTDM device	416
rtdm_device_info	
Device information	418
rtipc_port_label	
Port label information structure	418
rtser_config	
Serial device configuration	419
rtser_event	
Additional information about serial device events	420
rtser_status	
Serial device status	420
sockaddr_can	
Socket address structure for the CAN address family	421
sockaddr_ipc	
Socket address structure for the RTIPC address family	421
xnsched	
Scheduling information structure	422
xnvfile_lock_ops	
Vfile locking operations	423
xnvfile_regular_iterator	
Regular vfile iterator	424
xnvfile_regular_ops	
Regular vfile operation descriptor	425
xnvfile_rev_tag	
Snapshot revision tag	428
xnvfile_snapshot	
Snapshot vfile descriptor	428
xnvfile_snapshot_iterator	
Snapshot-driven vfile iterator	429
xnvfile_snapshot_ops	
Snapshot vfile operation descriptor	431

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

include/ version.h	??
include/alchemy/ alarm.h	??
include/alchemy/ buffer.h	??
include/alchemy/ cond.h	??
include/alchemy/ event.h	??
include/alchemy/ heap.h	??
include/alchemy/ mutex.h	??
include/alchemy/ pipe.h	??
include/alchemy/ queue.h	??
include/alchemy/ sem.h	??
include/alchemy/ task.h	??
include/alchemy/ timer.h	??
include/boilerplate/ ancillaries.h	??
include/boilerplate/ compiler.h	??
include/boilerplate/ debug.h	??
include/boilerplate/ hash.h	??
include/boilerplate/ list.h	??
include/boilerplate/ lock.h	??
include/boilerplate/ obstack.h	??
include/boilerplate/ private-list.h	??
include/boilerplate/ scope.h	??
include/boilerplate/ shared-list.h	??
include/boilerplate/ time.h	??
include/cobalt/ arith.h	??
include/cobalt/ fcntl.h	??
include/cobalt/ mqueue.h	??
include/cobalt/ pthread.h	??
include/cobalt/ sched.h	??
include/cobalt/ semaphore.h	??
include/cobalt/ signal.h	??
include/cobalt/ stdio.h	??
include/cobalt/ stdlib.h	??
include/cobalt/ syslog.h	??
include/cobalt/ ticks.h	??
include/cobalt/ time.h	??
include/cobalt/ trace.h	??
include/cobalt/ unistd.h	??
include/cobalt/ wrappers.h	??

include/cobalt/boilerplate/ limits.h	??
include/cobalt/boilerplate/ sched.h	??
include/cobalt/boilerplate/ signal.h	??
include/cobalt/boilerplate/ trace.h	??
include/cobalt/boilerplate/ wrappers.h	??
include/cobalt/kernel/ ancillaries.h	??
include/cobalt/kernel/ apc.h	440
include/cobalt/kernel/ arith.h	??
include/cobalt/kernel/ assert.h	??
include/cobalt/kernel/ bheap.h	??
include/cobalt/kernel/ bufd.h	441
include/cobalt/kernel/ clock.h	443
include/cobalt/kernel/ heap.h	??
include/cobalt/kernel/ init.h	??
include/cobalt/kernel/ intr.h	??
include/cobalt/kernel/ list.h	??
include/cobalt/kernel/ lock.h	??
include/cobalt/kernel/ map.h	444
include/cobalt/kernel/ pipe.h	??
include/cobalt/kernel/ ppd.h	446
include/cobalt/kernel/ registry.h	
This file is part of the Xenomai project	447
include/cobalt/kernel/ sched-idle.h	
Definitions for the IDLE scheduling class	492
include/cobalt/kernel/ sched-quota.h	??
include/cobalt/kernel/ sched-rt.h	
Definitions for the RT scheduling class	493
include/cobalt/kernel/ sched-sporadic.h	
Definitions for the SSP scheduling class	493
include/cobalt/kernel/ sched-tp.h	??
include/cobalt/kernel/ sched-weak.h	??
include/cobalt/kernel/ sched.h	
Scheduler interface header	437
include/cobalt/kernel/ schedparam.h	??
include/cobalt/kernel/ schedqueue.h	??
include/cobalt/kernel/ select.h	
File descriptors events multiplexing header	494
include/cobalt/kernel/ shadow.h	??
include/cobalt/kernel/ stat.h	??
include/cobalt/kernel/ synch.h	??
include/cobalt/kernel/ thread.h	??
include/cobalt/kernel/ timer.h	436
include/cobalt/kernel/ trace.h	??
include/cobalt/kernel/ tree.h	??
include/cobalt/kernel/ vdso.h	??
include/cobalt/kernel/ vfile.h	
This file is part of the Xenomai project	496
include/cobalt/kernel/rtdm/ can.h	463
include/cobalt/kernel/rtdm/ cobalt.h	
This file is part of the Xenomai project	471
include/cobalt/kernel/rtdm/ driver.h	
Real-Time Driver Model for Xenomai, driver API header	453
include/cobalt/kernel/rtdm/ fd.h	??
include/cobalt/kernel/rtdm/ ipc.h	
This file is part of the Xenomai project	471
include/cobalt/kernel/rtdm/ rtdm.h	476
include/cobalt/kernel/rtdm/ serial.h	479
include/cobalt/kernel/rtdm/ testing.h	489

include/cobalt/kernel/rtdm/analogy/ buffer.h	
Analogy for Linux, buffer related features	435
include/cobalt/kernel/rtdm/analogy/ channel_range.h	
Analogy for Linux, channel, range related features	448
include/cobalt/kernel/rtdm/analogy/ command.h	??
include/cobalt/kernel/rtdm/analogy/ context.h	
Analogy for Linux, context structure / macros declarations	450
include/cobalt/kernel/rtdm/analogy/ device.h	
Analogy for Linux, device related features	451
include/cobalt/kernel/rtdm/analogy/ driver.h	
Analogy for Linux, driver facilities	452
include/cobalt/kernel/rtdm/analogy/ instruction.h	
Analogy for Linux, instruction related features	459
include/cobalt/kernel/rtdm/analogy/ rtdm_helpers.h	
Analogy for Linux, Operation system facilities	460
include/cobalt/kernel/rtdm/analogy/ subdevice.h	
Analogy for Linux, subdevice related features	461
include/cobalt/kernel/rtdm/analogy/ transfer.h	
Analogy for Linux, transfer related features	462
include/cobalt/sys/ ioctl.h	??
include/cobalt/sys/ select.h	??
include/cobalt/sys/ socket.h	??
include/cobalt/sys/ time.h	??
include/cobalt/sys/ timerfd.h	??
include/cobalt/uapi/ cond.h	??
include/cobalt/uapi/ event.h	??
include/cobalt/uapi/ monitor.h	??
include/cobalt/uapi/ mutex.h	??
include/cobalt/uapi/ sched.h	??
include/cobalt/uapi/ sem.h	??
include/cobalt/uapi/ signal.h	??
include/cobalt/uapi/ syscall.h	??
include/cobalt/uapi/ thread.h	??
include/cobalt/uapi/ time.h	??
include/cobalt/uapi/asm-generic/ arith.h	??
include/cobalt/uapi/asm-generic/ features.h	??
include/cobalt/uapi/asm-generic/ syscall.h	??
include/cobalt/uapi/kernel/ heap.h	??
include/cobalt/uapi/kernel/ limits.h	??
include/cobalt/uapi/kernel/ synch.h	??
include/cobalt/uapi/kernel/ thread.h	??
include/cobalt/uapi/kernel/ trace.h	??
include/cobalt/uapi/kernel/ types.h	??
include/cobalt/uapi/kernel/ urw.h	??
include/cobalt/uapi/kernel/ vdso.h	??
include/cobalt/uapi/rtdm/ syscall.h	??
include/copperplate/ clockobj.h	??
include/copperplate/ cluster.h	??
include/copperplate/ debug.h	??
include/copperplate/ eventobj.h	??
include/copperplate/ heapobj.h	??
include/copperplate/ init.h	??
include/copperplate/ reference.h	??
include/copperplate/ registry-obstack.h	??
include/copperplate/ registry.h	??
include/copperplate/ semobj.h	??
include/copperplate/ syncobj.h	??
include/copperplate/ threadobj.h	??

include/copperplate/ timerobj.h	??
include/copperplate/ traceobj.h	??
include/copperplate/ wrappers.h	??
include/mercury/boilerplate/ limits.h	??
include/mercury/boilerplate/ sched.h	??
include/mercury/boilerplate/ signal.h	??
include/mercury/boilerplate/ trace.h	??
include/mercury/boilerplate/ wrappers.h	??
include/nocore/ atomic.h	??
include/psos/ psos.h	??
include/rtdm/ analogy.h	
Analogy for Linux, library facilities	498
include/rtdm/ can.h	463
include/rtdm/ ipc.h	
This file is part of the Xenomai project	472
include/rtdm/ rtdm.h	476
include/rtdm/ serial.h	
Real-Time Driver Model for Xenomai, serial device profile header	479
include/rtdm/ testing.h	
Real-Time Driver Model for Xenomai, testing device profile header	489
include/rtdm/uapi/ analogy.h	
Analogy for Linux, UAPI bits	499
include/rtdm/uapi/ can.h	
Real-Time Driver Model for RT-Socket-CAN, CAN device profile header	464
include/rtdm/uapi/ ipc.h	
This file is part of the Xenomai project	473
include/rtdm/uapi/ rtdm.h	
Real-Time Driver Model for Xenomai, user API header	477
include/rtdm/uapi/ serial.h	
Real-Time Driver Model for Xenomai, serial device profile header	480
include/rtdm/uapi/ testing.h	
Real-Time Driver Model for Xenomai, testing device profile header	490
include/vxworks/ errnoLib.h	??
include/vxworks/ intLib.h	??
include/vxworks/ kernLib.h	??
include/vxworks/ lstLib.h	??
include/vxworks/ memPartLib.h	??
include/vxworks/ msgQLib.h	??
include/vxworks/ rngLib.h	??
include/vxworks/ semLib.h	??
include/vxworks/ sysLib.h	??
include/vxworks/ taskInfo.h	??
include/vxworks/ taskLib.h	??
include/vxworks/ tickLib.h	??
include/vxworks/ types.h	??
include/vxworks/ wdLib.h	??
kernel/cobalt/ apc.c	505
kernel/cobalt/ arith.c	506
kernel/cobalt/ bufd.c	507
kernel/cobalt/ debug.h	??
kernel/cobalt/ map.c	508
kernel/cobalt/ procfs.h	??
kernel/cobalt/ registry.c	
This file is part of the Xenomai project	512
kernel/cobalt/ timer.c	511
kernel/cobalt/ vfile.c	
This file is part of the Xenomai project	521
kernel/cobalt/arch/arm/include/asm/xenomai/ calibration.h	??

kernel/cobalt/arch/arm/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/ fpctest.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/arch/arm/include/asm/xenomai/uapi/ tsc.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ calibration.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ fpctest.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/blackfin/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ calibration.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ fpctest.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/nios2/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ calibration.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ fpctest.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/powerpc/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ calibration.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ fpctest.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/sh/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ calibration.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ features.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ fpctest.h	??

kernel/cobalt/arch/x86/include/asm/xenomai/ machine.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ smi.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ syscall.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ thread.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/ wrappers.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/uapi/ arith.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/uapi/ features.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/uapi/ fpctest.h	??
kernel/cobalt/arch/x86/include/asm/xenomai/uapi/ syscall.h	??
kernel/cobalt/include/asm-generic/xenomai/ machine.h	??
kernel/cobalt/include/asm-generic/xenomai/ mayday.h	??
kernel/cobalt/include/asm-generic/xenomai/ pci_ids.h	??
kernel/cobalt/include/asm-generic/xenomai/ syscall.h	??
kernel/cobalt/include/asm-generic/xenomai/ thread.h	??
kernel/cobalt/include/asm-generic/xenomai/ wrappers.h	??
kernel/cobalt/include/ipipe/ thread_info.h	??
kernel/cobalt/posix/ clock.h	??
kernel/cobalt/posix/ cond.h	??
kernel/cobalt/posix/ event.h	??
kernel/cobalt/posix/ extension.h	??
kernel/cobalt/posix/ internal.h	??
kernel/cobalt/posix/ monitor.h	??
kernel/cobalt/posix/ mqueue.h	??
kernel/cobalt/posix/ mutex.h	??
kernel/cobalt/posix/ process.h	??
kernel/cobalt/posix/ sched.h	??
kernel/cobalt/posix/ select.c	
This file is part of the Xenomai project	510
kernel/cobalt/posix/ select.h	??
kernel/cobalt/posix/ sem.h	??
kernel/cobalt/posix/ signal.h	??
kernel/cobalt/posix/ syscall.c	
This file is part of the Xenomai project	505
kernel/cobalt/posix/ thread.h	??
kernel/cobalt/posix/ timer.h	??
kernel/cobalt/posix/ timerfd.h	??
kernel/cobalt/rtdm/ core.c	
Real-Time Driver Model for Xenomai, device operation multiplexing	513
kernel/cobalt/rtdm/ device.c	
Real-Time Driver Model for Xenomai, device management	516
kernel/cobalt/rtdm/ drvlib.c	
Real-Time Driver Model for Xenomai, driver library	517
kernel/cobalt/rtdm/ init.c	
Real-Time Driver Model for Xenomai	508
kernel/cobalt/rtdm/ internal.h	??
kernel/cobalt/rtdm/ syscall.h	??
kernel/cobalt/trace/ cobalt-core.h	??
kernel/cobalt/trace/ cobalt-posix.h	??
kernel/cobalt/trace/ cobalt-rtdm.h	??
kernel/drivers/analogy/ buffer.c	
Analogy for Linux, buffer related features	522
kernel/drivers/analogy/ command.c	
Analogy for Linux, command related features	524
kernel/drivers/analogy/ device.c	
Analogy for Linux, device related features	517
kernel/drivers/analogy/ driver.c	
Analogy for Linux, driver related features	524

kernel/drivers/analog/driver_facilities.c	525
Analogy for Linux, driver facilities	
kernel/drivers/analog/instruction.c	527
Analogy for Linux, instruction related features	
kernel/drivers/analog/proc.h	538
Analogy for Linux, procfs related features	
kernel/drivers/analog/rtdm_helpers.c	539
Analogy for Linux, RTDM helpers	
kernel/drivers/analog/rtdm_interface.c	540
Analogy for Linux, user interface (open, read, write, ioctl, proc)	
kernel/drivers/analog/subdevice.c	541
Analogy for Linux, subdevice, channel and range related features	
kernel/drivers/analog/transfer.c	542
Analogy for Linux, transfer related features	
kernel/drivers/analog/intel/8255.c	528
Analogy subdevice driver for 8255 chip	
kernel/drivers/analog/intel/8255.h	528
Hardware driver for 8255 chip	
kernel/drivers/analog/intel/parport.c	529
Analogy driver for standard parallel port	
kernel/drivers/analog/national_instruments/mio_common.c	530
Hardware driver for DAQ-STC based boards	
kernel/drivers/analog/national_instruments/mite.c	531
Hardware driver for NI Mite PCI interface chip	
kernel/drivers/analog/national_instruments/mite.h	532
Hardware driver for NI Mite PCI interface chip	
kernel/drivers/analog/national_instruments/ni_mio.h	533
Hardware driver for NI Mite PCI interface chip	
kernel/drivers/analog/national_instruments/ni_stc.h	534
Register descriptions for NI DAQ-STC chip	
kernel/drivers/analog/national_instruments/ni_tio.h	535
Hardware driver for NI general purpose counter	
kernel/drivers/analog/national_instruments/pcimio.c	536
Hardware driver for NI PCI-MIO E series cards	
kernel/drivers/analog/national_instruments/tio_common.c	537
Hardware driver for NI general purpose counter	
kernel/drivers/analog/sensoray/s526.c	540
Analogy driver for Sensoray Model 526 board	
kernel/drivers/can/rtdcan_dev.h	??
kernel/drivers/can/rtdcan_internal.h	??
kernel/drivers/can/rtdcan_list.h	??
kernel/drivers/can/rtdcan_raw.h	??
kernel/drivers/can/rtdcan_socket.h	??
kernel/drivers/can/rtdcan_version.h	??
kernel/drivers/can/mscan/rtdcan_mscan.h	??
kernel/drivers/can/mscan/rtdcan_mscan_regs.h	??
kernel/drivers/can/sja1000/rtdcan_sja1000.h	??
kernel/drivers/can/sja1000/rtdcan_sja1000_regs.h	??
kernel/drivers/ipc/internal.h	??
kernel/drivers/serial/16550A_io.h	??
kernel/drivers/serial/16550A_pci.h	??
kernel/drivers/serial/16550A_pnp.h	??
lib/alchemy/alarm.h	??
lib/alchemy/buffer.h	??
lib/alchemy/cond.h	??
lib/alchemy/event.h	??
lib/alchemy/heap.h	??
lib/alchemy/init.h	??

lib/alchemy/ internal.h	??
lib/alchemy/ mutex.h	??
lib/alchemy/ pipe.h	??
lib/alchemy/ queue.h	??
lib/alchemy/ reference.h	??
lib/alchemy/ sem.h	??
lib/alchemy/ task.h	??
lib/alchemy/ timer.h	??
lib/analogy/ async.c	
Analogy for Linux, command, transfer, etc	543
lib/analogy/ descriptor.c	
Analogy for Linux, descriptor related features	544
lib/analogy/ info.c	
Analogy for Linux, device, subdevice, etc	545
lib/analogy/ internal.h	
Analogy for Linux, internal declarations	509
lib/analogy/ range.c	
Analogy for Linux, range related features	546
lib/analogy/ root_leaf.h	
Analogy for Linux, root / leaf system	547
lib/analogy/ sync.c	
Analogy for Linux, instruction related features	548
lib/analogy/ sys.c	
Analogy for Linux, descriptor related features	549
lib/cobalt/ current.h	??
lib/cobalt/ init.h	??
lib/cobalt/ internal.h	??
lib/cobalt/ sem_heap.h	??
lib/cobalt/arch/arm/include/asm/xenomai/ features.h	??
lib/cobalt/arch/arm/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/arm/include/asm/xenomai/ tsc.h	??
lib/cobalt/arch/blackfin/include/asm/xenomai/ features.h	??
lib/cobalt/arch/blackfin/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/blackfin/include/asm/xenomai/ tsc.h	??
lib/cobalt/arch/nios2/include/asm/xenomai/ features.h	??
lib/cobalt/arch/nios2/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/nios2/include/asm/xenomai/ tsc.h	??
lib/cobalt/arch/powerpc/include/asm/xenomai/ features.h	??
lib/cobalt/arch/powerpc/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/powerpc/include/asm/xenomai/ tsc.h	??
lib/cobalt/arch/sh/include/asm/xenomai/ features.h	??
lib/cobalt/arch/sh/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/sh/include/asm/xenomai/ tsc.h	??
lib/cobalt/arch/x86/include/asm/xenomai/ features.h	??
lib/cobalt/arch/x86/include/asm/xenomai/ syscall.h	??
lib/cobalt/arch/x86/include/asm/xenomai/ tsc.h	??


Chapter 5

Module Documentation

5.1 Channels and ranges

Channels.

Collaboration diagram for Channels and ranges:



Data Structures

- struct [a4l_channel](#)
Structure describing some channel's characteristics.
- struct [a4l_channels_desc](#)
Structure describing a channels set.
- struct [a4l_range](#)
Structure describing a (unique) range.

Macros

- #define [A4L_CHAN_GLOBAL](#) 0x10
Internal use flag (must not be used by driver developer)
- #define [A4L_RNG_GLOBAL](#) 0x8
Internal use flag (must not be used by driver developer)
- #define [RANGE](#)(x, y)
Macro to declare a (unique) range with no unit defined.
- #define [RANGE_V](#)(x, y)
Macro to declare a (unique) range in Volt.
- #define [RANGE_mA](#)(x, y)
Macro to declare a (unique) range in milliAmpere.

- `#define RANGE_ext(x, y)`
Macro to declare a (unique) range in some external reference.
- `#define A4L_RNG_GLOBAL_RNGDESC 0`
Constant to define a ranges descriptor as global (inter-channel)
- `#define A4L_RNG_PERCHAN_RNGDESC 1`
Constant to define a ranges descriptor as specific for a channel.
- `#define RNG_GLOBAL(x)`
Macro to declare a ranges global descriptor in one line.

Channel reference

Flags to define the channel's reference

- `#define A4L_CHAN_AREF_GROUND 0x1`
Ground reference.
- `#define A4L_CHAN_AREF_COMMON 0x2`
Common reference.
- `#define A4L_CHAN_AREF_DIFF 0x4`
Differential reference.
- `#define A4L_CHAN_AREF_OTHER 0x8`
Misc reference.

Channels declaration mode

Constant to define whether the channels in a descriptor are identical

- `#define A4L_CHAN_GLOBAL_CHANDESC 0`
Global declaration, the set contains channels with similar characteristics.
- `#define A4L_CHAN_PERCHAN_CHANDESC 1`
Per channel declaration, the descriptor gathers different channels.

5.1.1 Detailed Description

Channels. According to the Analogy nomenclature, the channel is the elementary acquisition entity. One channel is supposed to acquire one data at a time. A channel can be:

- an analog input or an analog output;
- a digital input or a digital output;

Channels are defined by their type and by some other characteristics like:

- their resolutions for analog channels (which usually ranges from 8 to 32 bits);
- their references;

Such parameters must be declared for each channel composing a subdevice. The structure `a4l_channel` (struct `a4l_channel`) is used to define one channel.

Another structure named `a4l_channels_desc` (struct `a4l_channels_desc`) gathers all channels for a specific subdevice. This latter structure also stores :

- the channels count;

- the channels declaration mode (A4L_CHAN_GLOBAL_CHANDESC or A4L_CHAN_PERCHAN_CHANDESC): if all the channels composing a subdevice are identical, there is no need to declare the parameters for each channel; the global declaration mode eases the structure composition.

Usually the channels descriptor looks like this:

```
struct a4l_channels_desc example_chan = {
    mode: A4L_CHAN_GLOBAL_CHANDESC, -> Global declaration
                                     mode is set
    length: 8, -> 8 channels
    chans: {
        {A4L_CHAN_AREF_GROUND, 16}, -> Each channel is 16 bits
                                     wide with the ground as
                                     reference
    },
};
```

Ranges

So as to perform conversion from logical values acquired by the device to physical units, some range structure(s) must be declared on the driver side.

Such structures contain:

- the physical unit type (Volt, Ampere, none);
- the minimal and maximal values;

These range structures must be associated with the channels at subdevice registration time as a channel can work with many ranges. At configuration time (thanks to an Analogy command), one range will be selected for each enabled channel.

Consequently, for each channel, the developer must declare all the possible ranges in a structure called struct a4l_rngtab. Here is an example:

```
struct a4l_rngtab example_tab = {
    length: 2,
    rngs: {
        RANGE_V(-5,5),
        RANGE_V(-10,10),
    },
};
```

For each subdevice, a specific structure is designed to gather all the ranges tabs of all the channels. In this structure, called struct a4l_rngdesc, three fields must be filled:

- the declaration mode (A4L_RNG_GLOBAL_RNGDESC or A4L_RNG_PERCHAN_RNGDESC);
- the number of ranges tab;
- the tab of ranges tabs pointers;

Most of the time, the channels which belong to the same subdevice use the same set of ranges. So, there is no need to declare the same ranges for each channel. A macro is defined to prevent redundant declarations: [RNG_GLOBAL\(\)](#).

Here is an example:

```
struct a4l_rngdesc example_rng = RNG_GLOBAL(example_tab);
```

5.2 Big dual kernel lock

Collaboration diagram for Big dual kernel lock:



Macros

- `#define cobalt_atomic_enter(context)`
Enter atomic section (dual kernel only)
- `#define cobalt_atomic_leave(context)`
Leave atomic section (dual kernel only)
- `#define RTDM_EXECUTE_ATOMICALY(code_block)`
Execute code block atomically (DEPRECATED)

5.2.1 Detailed Description

5.2.2 Macro Definition Documentation

5.2.2.1 `#define cobalt_atomic_enter(context)`

Value:

```
do {
    xnlock_get_irqsave(&nklock, (context)); \
    __xnsched_lock(); \
} while (0)
```

Enter atomic section (dual kernel only)

This call opens a fully atomic section, serializing execution with respect to all interrupt handlers (including for real-time IRQs) and Xenomai threads running on all CPUs.

Parameters

<i>context</i>	name of local variable to store the context in. This variable updated by the real-time core will hold the information required to leave the atomic section properly.
----------------	--

Note

Atomic sections may be nested.

Since the strongest lock is acquired by this service, it can be used to synchronize real-time and non-real-time contexts.

Warning

This service is not portable to the Mercury core, and should be restricted to Cobalt-specific use cases.

5.2.2.2 #define cobalt_atomic_leave(context)

Value:

```
do {
    __xnsched_unlock();
    xnlock_put_irqrestore(&nklock, (context));
} while (0)
```

Leave atomic section (dual kernel only)

This call closes an atomic section previously opened by a call to [cobalt_atomic_enter\(\)](#), restoring the preemption and interrupt state which prevailed prior to entering the exited section.

Parameters

<i>context</i>	name of local variable which stored the context.
----------------	--

Warning

This service is not portable to the Mercury core, and should be restricted to Cobalt-specific use cases.

5.2.2.3 #define RTDM_EXECUTE_ATOMICALY(code_block)

Value:

```
{
    <ENTER_ATOMIC_SECTION>
    code_block;
    <LEAVE_ATOMIC_SECTION>
}
```

Execute code block atomically (DEPRECATED)

Generally, it is illegal to suspend the current task by calling [rt dm_task_sleep\(\)](#), [rt dm_event_wait\(\)](#), etc. while holding a spinlock. In contrast, this macro allows to combine several operations including a potentially rescheduling call to an atomic code block with respect to other [RTDM_EXECUTE_ATOMICALY\(\)](#) blocks. The macro is a light-weight alternative for protecting code blocks via mutexes, and it can even be used to synchronise real-time and non-real-time contexts.

Parameters

<i>code_block</i>	Commands to be executed atomically
-------------------	------------------------------------

Note

It is not allowed to leave the code block explicitly by using `break`, `return`, `goto`, etc. This would leave the global lock held during the code block execution in an inconsistent state. Moreover, do not embed complex operations into the code block. Consider that they will be executed under preemption lock with interrupts switched-off. Also note that invocation of rescheduling calls may break the atomicity until the task gains the CPU again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task

- User-space task (RT, non-RT)


Rescheduling: possible, depends on functions called within *code_block*.

Warning

This construct is deprecated and will be phased out in Xenomai 3.0. Please use `rtdm_waitqueue` services instead.

5.3 Spinlock with preemption deactivation

Collaboration diagram for Spinlock with preemption deactivation:



Macros

- `#define RTDM_LOCK_UNLOCKED(__name) IPIPE_SPIN_LOCK_UNLOCKED`
Static lock initialisation.
- `#define rtdm_lock_irqsave(context) splhigh(context)`
Disable preemption locally.
- `#define rtdm_lock_irqrestore(context) splexit(context)`
Restore preemption state.

Typedefs

- `typedef ipipe_spinlock_t rtdm_lock_t`
Lock variable.
- `typedef unsigned long rtdm_lockctx_t`
Variable to save the context while holding a lock.

Functions

- `static void rtdm_lock_init (rtdm_lock_t *lock)`
Dynamic lock initialisation.
- `static void rtdm_lock_get (rtdm_lock_t *lock)`
Acquire lock from non-preemptible contexts.
- `static void rtdm_lock_put (rtdm_lock_t *lock)`
Release lock without preemption restoration.
- `static void rtdm_lock_put_irqrestore (rtdm_lock_t *lock, rtdm_lockctx_t context)`
Release lock and restore preemption state.

5.3.1 Detailed Description

5.3.2 Macro Definition Documentation

5.3.2.1 `#define rtdm_lock_irqrestore(context) splexit(context)`

Restore preemption state.

Parameters

<i>context</i>	name of local variable which stored the context
----------------	---

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.3.2.2 `#define rtdm_lock_irqsave(context) splhigh(context)`

Disable preemption locally.

Parameters

<i>context</i>	name of local variable to store the context in
----------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.3.3 Function Documentation

5.3.3.1 `static void rtdm_lock_get (rtdm_lock_t * lock) [inline], [static]`

Acquire lock from non-preemptible contexts.

Parameters

<i>lock</i>	Address of lock variable
-------------	--------------------------

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.3.3.2 static void rtdm_lock_init (**rtdm_lock_t** * lock) [inline], [static]

Dynamic lock initialisation.

Parameters

<i>lock</i>	Address of lock variable
-------------	--------------------------

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.3.3.3 static void rtdm_lock_put (**rtdm_lock_t** * lock) [inline], [static]

Release lock without preemption restoration.

Parameters

<i>lock</i>	Address of lock variable
-------------	--------------------------

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.3.3.4 static void rtdm_lock_put_irqrestore (**rtdm_lock_t** * lock, **rtdm_lockctx_t** context) [inline], [static]

Release lock and restore preemption state.

Parameters

<i>lock</i>	Address of lock variable
<i>context</i>	name of local variable which stored the context

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

Referenced by a4l_request_irq(), and rtdm_ratelimit().

5.4 Thread state flags

Bits reporting permanent or transient states of threads.

Collaboration diagram for Thread state flags:



Macros

- #define **XNSUSP** 0x00000001
Suspended.
- #define **XNPEND** 0x00000002
Sleep-wait for a resource.
- #define **XNDELAY** 0x00000004
Delayed.
- #define **XNREADY** 0x00000008
Linked to the ready queue.
- #define **XNDORMANT** 0x00000010
Not started yet.
- #define **XNZOMBIE** 0x00000020
Zombie thread in deletion process.
- #define **XNMAPPED** 0x00000040
Thread is mapped to a linux task.
- #define **XNRELAX** 0x00000080
Relaxed shadow thread (blocking bit)
- #define **XNMIGRATE** 0x00000100
Thread is currently migrating to another CPU.
- #define **XNHELD** 0x00000200
Thread is held to process emergency.
- #define **XNBOOST** 0x00000400
Undergoes a PIP boost.
- #define **XNDEBUG** 0x00000800
Hit a debugger breakpoint.
- #define **XNLOCK** 0x00001000
Holds the scheduler lock (i.e.
- #define **XNRRB** 0x00002000
Undergoes a round-robin scheduling.
- #define **XNTRAPSW** 0x00004000
Trap execution mode switches.
- #define **XNFPU** 0x00008000
Thread uses FPU.
- #define **XNROOT** 0x00010000
Root thread (that is, Linux/IDLE)

- `#define XNWEAK 0x00020000`
Non real-time shadow (from the WEAK class)
- `#define XNUSER 0x00040000`
Shadow thread running in userland.
- `#define XNJOINED 0x00080000`
Another thread waits for joining this thread.
- `#define XNTRAPLB 0x00100000`
Trap lock break (i.e.

5.4.1 Detailed Description

Bits reporting permanent or transient states of threads.

5.4.2 Macro Definition Documentation

5.4.2.1 `#define XNHELD 0x00000200`

Thread is held to process emergency.

Referenced by `xnthread_resume()`, and `xnthread_suspend()`.

5.4.2.2 `#define XNLOCK 0x00001000`

Holds the scheduler lock (i.e.

not preemptible)

Referenced by `xnthread_set_mode()`, and `xnthread_suspend()`.

5.4.2.3 `#define XNMIGRATE 0x00000100`

Thread is currently migrating to another CPU.

5.4.2.4 `#define XNPEND 0x00000002`

Sleep-wait for a resource.

Referenced by `xnsynch_acquire()`, `xnsynch_flush()`, `xnsynch_forget_sleeper()`, `xnsynch_sleep_on()`, `xnsynch_wakeup_one_sleeper()`, `xnsynch_wakeup_this_sleeper()`, `xnthread_resume()`, and `xnthread_unblock()`.

5.4.2.5 `#define XNREADY 0x00000008`

Linked to the ready queue.

Referenced by `xnthread_resume()`, and `xnthread_suspend()`.

5.4.2.6 `#define XNSUSP 0x00000001`

Suspended.

Referenced by `xnthread_init()`, `xnthread_start()`, and `xnthread_suspend()`.

5.4.2.7 `#define XNTRAPLB 0x00100000`

Trap lock break (i.e.

may not sleep with XNLOCK)

Referenced by `xnthread_suspend()`.

5.5 Thread information flags

Bits reporting events notified to threads.

Collaboration diagram for Thread information flags:



Macros

- `#define XNTIMEO 0x00000001`
Woken up due to a timeout condition.
- `#define XNRMID 0x00000002`
Pending on a removed resource.
- `#define XNBREAK 0x00000004`
Forcibly awoken from a wait state.
- `#define XNKICKED 0x00000008`
Forced out of primary mode.
- `#define XNWAKEN 0x00000010`
Thread waken up upon resource availability.
- `#define XNROBBED 0x00000020`
Robbed from resource ownership.
- `#define XNCANCELLED 0x00000040`
Cancellation request is pending.
- `#define XNSWREP 0x00000080`
Mode switch already reported.
- `#define XNMOVED 0x00000100`
CPU migration in primary mode occurred.

5.5.1 Detailed Description

Bits reporting events notified to threads.

5.6 CAN Devices

This is the common interface a RTDM-compliant CAN device has to provide.

Collaboration diagram for CAN Devices:



Data Structures

- struct [can_bittime_std](#)
Standard bit-time parameters according to Bosch.
- struct [can_bittime_btr](#)
Hardware-specific BTR bit-times.
- struct [can_bittime](#)
Custom CAN bit-time definition.
- struct [can_filter](#)
Filter for reception of CAN messages.
- struct [sockaddr_can](#)
Socket address structure for the CAN address family.
- struct [can_frame](#)
Raw CAN frame.

Macros

- `#define` [AF_CAN](#) 29
CAN address family.
- `#define` [PF_CAN](#) [AF_CAN](#)
CAN protocol family.
- `#define` [SOL_CAN_RAW](#) 103
CAN socket levels.

Typedefs

- typedef uint32_t [can_id_t](#)
Type of CAN id (see [CAN_xxx_MASK](#) and [CAN_xxx_FLAG](#))
- typedef [can_id_t](#) [can_err_mask_t](#)
Type of CAN error mask.
- typedef uint32_t [can_baudrate_t](#)
Baudrate definition in bits per second.
- typedef enum [CAN_BITTIME_TYPE](#) [can_bittime_type_t](#)
See [CAN_BITTIME_TYPE](#).
- typedef enum [CAN_MODE](#) [can_mode_t](#)

- See [CAN_MODE](#).
- typedef int [can_ctrlmode_t](#)
See [CAN_CTRLMODE](#).
- typedef enum [CAN_STATE](#) [can_state_t](#)
See [CAN_STATE](#).
- typedef struct [can_filter](#) [can_filter_t](#)
Filter for reception of CAN messages.
- typedef struct [can_frame](#) [can_frame_t](#)
Raw CAN frame.

Enumerations

- enum [CAN_BITTIME_TYPE](#) { [CAN_BITTIME_STD](#), [CAN_BITTIME_BTR](#) }
Supported CAN bit-time types.

CAN ID masks

Bit masks for masking CAN IDs

- #define [CAN_EFF_MASK](#) 0x1FFFFFFF
Bit mask for extended CAN IDs.
- #define [CAN_SFF_MASK](#) 0x000007FF
Bit mask for standard CAN IDs.

CAN ID flags

Flags within a CAN ID indicating special CAN frame attributes

- #define [CAN_EFF_FLAG](#) 0x80000000
Extended frame.
- #define [CAN_RTR_FLAG](#) 0x40000000
Remote transmission frame.
- #define [CAN_ERR_FLAG](#) 0x20000000
Error frame (see [Errors](#)), not valid in struct [can_filter](#).
- #define [CAN_INV_FILTER](#) [CAN_ERR_FLAG](#)
Invert CAN filter definition, only valid in struct [can_filter](#).

Particular CAN protocols

Possible protocols for the PF_CAN protocol family

Currently only the RAW protocol is supported.

- #define [CAN_RAW](#) 1
Raw protocol of PF_CAN, applicable to socket type [SOCK_RAW](#).

CAN operation modes

Modes into which CAN controllers can be set

- enum [CAN_MODE](#) { [CAN_MODE_STOP](#) = 0, [CAN_MODE_START](#), [CAN_MODE_SLEEP](#) }

CAN controller modes

Special CAN controllers modes, which can be or'ed together.

Note

These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.

- #define `CAN_CTRLMODE_LISTENONLY` 0x1
- #define `CAN_CTRLMODE_LOOPBACK` 0x2
- #define `CAN_CTRLMODE_3_SAMPLES` 0x4

CAN controller states

States a CAN controller can be in.

- enum `CAN_STATE` {
`CAN_STATE_ERROR_ACTIVE = 0, CAN_STATE_ACTIVE = 0, CAN_STATE_ERROR_WARNI-`
`NG = 1, CAN_STATE_BUS_WARNING = 1,`
`CAN_STATE_ERROR_PASSIVE = 2, CAN_STATE_BUS_PASSIVE = 2, CAN_STATE_BUS_O-`
`FF, CAN_STATE_SCANNING_BAUDRATE,`
`CAN_STATE_STOPPED, CAN_STATE_SLEEPING }`

Timestamp switches

Arguments to pass to `RTCAN_RTIOC_TAKE_TIMESTAMP`

- #define `RTCAN_TAKE_NO_TIMESTAMPS` 0
Switch off taking timestamps.
- #define `RTCAN_TAKE_TIMESTAMPS` 1
Do take timestamps.

RAW socket options

Setting and getting CAN RAW socket options.

- #define `CAN_RAW_FILTER` 0x1
CAN filter definition.
- #define `CAN_RAW_ERR_FILTER` 0x2
CAN error mask.
- #define `CAN_RAW_LOOPBACK` 0x3
CAN TX loopback.
- #define `CAN_RAW_RECV_OWN_MSGS` 0x4
CAN receive own messages.

IOCTLs

CAN device IOCTLs

- #define `SIOCGIFINDEX` `defined_by_kernel_header_file`

- Get CAN interface index by name.*

 - #define [SIOCSCANBAUDRATE](#) _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)

Set baud rate.

 - #define [SIOCGCANBAUDRATE](#) _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)

Get baud rate.

 - #define [SIOCSCANCUSTOMBITTIME](#) _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)

Set custom bit time parameter.

 - #define [SIOCGCANCUSTOMBITTIME](#) _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)

Get custom bit-time parameters.

 - #define [SIOCSCANMODE](#) _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)

Set operation mode of CAN controller.

 - #define [SIOCGCANSTATE](#) _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)

Get current state of CAN controller.

 - #define [SIOCSCANCTRLMODE](#) _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)

Set special controller modes.

 - #define [SIOCGCANCTRLMODE](#) _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)

Get special controller modes.

 - #define [RTCAN_RTIOC_TAKE_TIMESTAMP](#) _IOW(RTIOC_TYPE_CAN, 0x09, int)

Enable or disable storing a high precision timestamp upon reception of a CAN frame.

 - #define [RTCAN_RTIOC_RCV_TIMEOUT](#) _IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_rel_t)

Specify a reception timeout for a socket.

 - #define [RTCAN_RTIOC_SND_TIMEOUT](#) _IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_t)

Specify a transmission timeout for a socket.

Error mask

Error class (mask) in `can_id` field of struct [can_frame](#) to be used with [CAN_RAW_ERR_FILTER](#).

Note: Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.

Note: In case of a bus-off error condition ([CAN_ERR_BUSOFF](#)), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling [CAN_MODE_START](#).

Note: Bus error interrupts ([CAN_ERR_BUSERROR](#)) are enabled when an application is calling a [Recv](#) function on a socket listening on bus errors (using [CAN_RAW_ERR_FILTER](#)). After one bus error has occurred, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.

- #define [CAN_ERR_TX_TIMEOUT](#) 0x00000001U
- TX timeout (netdevice driver)*
- #define [CAN_ERR_LOSTARB](#) 0x00000002U
- Lost arbitration (see [data\[0\]](#))*
- #define [CAN_ERR_CRTL](#) 0x00000004U
- Controller problems (see [data\[1\]](#))*
- #define [CAN_ERR_PROT](#) 0x00000008U
- Protocol violations (see [data\[2\]](#), [data\[3\]](#))*
- #define [CAN_ERR_TRX](#) 0x00000010U
- Transceiver status (see [data\[4\]](#))*
- #define [CAN_ERR_ACK](#) 0x00000020U
- Received no ACK on transmission.*
- #define [CAN_ERR_BUSOFF](#) 0x00000040U

Bus off.

- #define `CAN_ERR_BUSERROR` 0x00000080U

Bus error (may flood!)

- #define `CAN_ERR_RESTARTED` 0x00000100U

Controller restarted.

- #define `CAN_ERR_MASK` 0x1FFFFFFFU

Omit EFF, RTR, ERR flags.

Arbitration lost error

Error in the data[0] field of struct `can_frame`.

- #define `CAN_ERR_LOSTARB_UNSPEC` 0x00
unspecified

Controller problems

Error in the data[1] field of struct `can_frame`.

- #define `CAN_ERR_CRTL_UNSPEC` 0x00
unspecified
- #define `CAN_ERR_CRTL_RX_OVERFLOW` 0x01
RX buffer overflow.
- #define `CAN_ERR_CRTL_TX_OVERFLOW` 0x02
TX buffer overflow.
- #define `CAN_ERR_CRTL_RX_WARNING` 0x04
reached warning level for RX errors
- #define `CAN_ERR_CRTL_TX_WARNING` 0x08
reached warning level for TX errors
- #define `CAN_ERR_CRTL_RX_PASSIVE` 0x10
reached passive level for RX errors
- #define `CAN_ERR_CRTL_TX_PASSIVE` 0x20
reached passive level for TX errors

Protocol error type

Error in the data[2] field of struct `can_frame`.

- #define `CAN_ERR_PROT_UNSPEC` 0x00
unspecified
- #define `CAN_ERR_PROT_BIT` 0x01
single bit error
- #define `CAN_ERR_PROT_FORM` 0x02
frame format error
- #define `CAN_ERR_PROT_STUFF` 0x04
bit stuffing error
- #define `CAN_ERR_PROT_BIT0` 0x08
unable to send dominant bit
- #define `CAN_ERR_PROT_BIT1` 0x10

- *unable to send recessive bit*
- #define [CAN_ERR_PROT_OVERLOAD](#) 0x20
- *bus overload*
- #define [CAN_ERR_PROT_ACTIVE](#) 0x40
- *active error announcement*
- #define [CAN_ERR_PROT_TX](#) 0x80
- *error occurred on transmission*

Protocol error location

Error in the data[4] field of struct [can_frame](#).

- #define [CAN_ERR_PROT_LOC_UNSPEC](#) 0x00
- *unspecified*
- #define [CAN_ERR_PROT_LOC_SOF](#) 0x03
- *start of frame*
- #define [CAN_ERR_PROT_LOC_ID28_21](#) 0x02
- *ID bits 28 - 21 (SFF: 10 - 3)*
- #define [CAN_ERR_PROT_LOC_ID20_18](#) 0x06
- *ID bits 20 - 18 (SFF: 2 - 0)*
- #define [CAN_ERR_PROT_LOC_SRTR](#) 0x04
- *substitute RTR (SFF: RTR)*
- #define [CAN_ERR_PROT_LOC_IDE](#) 0x05
- *identifier extension*
- #define [CAN_ERR_PROT_LOC_ID17_13](#) 0x07
- *ID bits 17-13.*
- #define [CAN_ERR_PROT_LOC_ID12_05](#) 0x0F
- *ID bits 12-5.*
- #define [CAN_ERR_PROT_LOC_ID04_00](#) 0x0E
- *ID bits 4-0.*
- #define [CAN_ERR_PROT_LOC_RTR](#) 0x0C
- *RTR.*
- #define [CAN_ERR_PROT_LOC_RES1](#) 0x0D
- *reserved bit 1*
- #define [CAN_ERR_PROT_LOC_RES0](#) 0x09
- *reserved bit 0*
- #define [CAN_ERR_PROT_LOC_DLC](#) 0x0B
- *data length code*
- #define [CAN_ERR_PROT_LOC_DATA](#) 0x0A
- *data section*
- #define [CAN_ERR_PROT_LOC_CRC_SEQ](#) 0x08
- *CRC sequence.*
- #define [CAN_ERR_PROT_LOC_CRC_DEL](#) 0x18
- *CRC delimiter.*
- #define [CAN_ERR_PROT_LOC_ACK](#) 0x19
- *ACK slot.*
- #define [CAN_ERR_PROT_LOC_ACK_DEL](#) 0x1B
- *ACK delimiter.*
- #define [CAN_ERR_PROT_LOC_EOF](#) 0x1A
- *end of frame*

- #define [CAN_ERR_PROT_LOC_INTERM](#) 0x12
intermission
- #define [CAN_ERR_TRX_UNSPEC](#) 0x00
0000 0000
- #define [CAN_ERR_TRX_CANH_NO_WIRE](#) 0x04
0000 0100
- #define [CAN_ERR_TRX_CANH_SHORT_TO_BAT](#) 0x05
0000 0101
- #define [CAN_ERR_TRX_CANH_SHORT_TO_VCC](#) 0x06
0000 0110
- #define [CAN_ERR_TRX_CANH_SHORT_TO_GND](#) 0x07
0000 0111
- #define [CAN_ERR_TRX_CANL_NO_WIRE](#) 0x40
0100 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_BAT](#) 0x50
0101 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_VCC](#) 0x60
0110 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_GND](#) 0x70
0111 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_CANH](#) 0x80
1000 0000

5.6.1 Detailed Description

This is the common interface a RTDM-compliant CAN device has to provide. Feel free to report bugs and comments on this profile to the "Socketcan" mailing list (Socketcan-core@lists.berlios.de) or directly to the authors (wg@grandegger.com or Sebastian.Smolorz@stud.uni-hannover.de).

Profile Revision: 2

Device Characteristics

Device Flags: RTDM_PROTOCOL_DEVICE
Protocol Family: PF_CAN
Socket Type: SOCK_RAW
Device Class: RTDM_CLASS_CAN

Supported Operations

Socket

Environments: non-RT (RT optional, deprecated)

Specific return values:

- -EPROTONOSUPPORT (Protocol is not supported by the driver. See [CAN protocols](#) for possible protocols.)

Close

Blocking calls to any of the [Send](#) or [Receive](#) functions will be unblocked when the socket is closed and return with an error.

Environments: non-RT (RT optional, deprecated)

Specific return values: none

IOCTL

Mandatory Environments: see [below](#)

Specific return values: see [below](#)

Bind

Binds a socket to one or all CAN devices (see struct [sockaddr_can](#)). If a filter list has been defined with `setsockopt` (see [Sockopts](#)), it will be used upon reception of CAN frames to decide whether the bound socket will receive a frame. If no filter has been defined, the socket will receive **all** CAN frames on the specified interface(s).

Binding to special interface index 0 will make the socket receive CAN frames from all CAN interfaces. Binding to an interface index is also relevant for the [Send](#) functions because they will transmit a message over the interface the socket is bound to when no socket address is given to them.

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -ENOMEM (Not enough memory to fulfill the operation)
- -EINVAL (Invalid address family, or invalid length of address structure)
- -ENODEV (Invalid CAN interface index)
- -ENOSPC (No enough space for filter list)
- -EBADF (Socket is about to be closed)
- -EAGAIN (Too many receivers. Old binding (if any) is still active. Close some sockets and try again.)

Setsockopt, Getsockopt

These functions allow to set and get various socket options. Currently, only CAN raw sockets are supported.

Supported Levels and Options:

- Level **SOL_CAN_RAW** : CAN RAW protocol (see [CAN_RAW](#))
 - Option [CAN_RAW_FILTER](#) : CAN filter list
 - Option [CAN_RAW_ERR_FILTER](#) : CAN error mask
 - Option [CAN_RAW_LOOPBACK](#) : CAN TX loopback to local sockets

Environments: non-RT (RT optional)

Specific return values: see links to options above.

Recv, Recvfrom, Recvmsg

These functions receive CAN messages from a socket. Only one message per call can be received, so only one buffer with the correct length must be passed. For `SOCK_RAW`, this is the size of struct [can_frame](#).

Unlike a call to one of the [Send](#) functions, a `Recv` function will not return with an error if an interface is down (due to bus-off or setting of stop mode) or in sleep mode. Moreover, in such a case there may still be some CAN messages in the socket buffer which could be read out successfully.

It is possible to receive a high precision timestamp with every CAN message. The condition is a former instruction to the socket via [RTCAN_RTIOC_TAKE_TIMESTAMP](#). The timestamp will be copied to the `msg_control` buffer of struct `msghdr` if it points to a valid memory location with size of [nanosecs_abs_t](#). If this is a NULL pointer the timestamp will be discarded silently.

Note: A `msg_controllen` of 0 upon completion of the function call indicates that no timestamp is available for that message.

Supported Flags [in]:

- `MSG_DONTWAIT` (By setting this flag the operation will only succeed if it would not block, i.e. if there is a message in the socket buffer. This flag takes precedence over a timeout specified by [RTCAN_RTIOC_RCV_TIMEOUT](#).)
- `MSG_PEEK` (Receive a message but leave it in the socket buffer. The next receive operation will get that message again.)

Supported Flags [out]: none

Environments: RT (non-RT optional)

Specific return values:

- Non-negative value (Indicating the successful reception of a CAN message. For `SOCK_RAW`, this is the size of struct [can_frame](#) regardless of the actual size of the payload.)
- -EFAULT (It was not possible to access user space memory area at one of the specified addresses.)

- -EINVAL (Unsupported flag detected, or invalid length of socket address buffer, or invalid length of message control buffer)
- -EMSGSIZE (Zero or more than one iovec buffer passed, or buffer too small)
- -EAGAIN (No data available in non-blocking mode)
- -EBADF (Socket was closed.)
- -EINTR (Operation was interrupted explicitly or by signal.)
- -ETIMEDOUT (Timeout)

Send, Sendto, Sendmsg

These functions send out CAN messages. Only one message per call can be transmitted, so only one buffer with the correct length must be passed. For SOCK_RAW, this is the size of struct [can_frame](#). The following only applies to SOCK_RAW: If a socket address of struct [sockaddr_can](#) is given, only `can_ifindex` is used. It is also possible to omit the socket address. Then the interface the socket is bound to will be used for sending messages.

If an interface goes down (due to bus-off or setting of stop mode) all senders that were blocked on this interface will be woken up.

Supported Flags:

- MSG_DONTWAIT (By setting this flag the transmit operation will only succeed if it would not block. This flag takes precedence over a timeout specified by [RTCAN_RTIOC_SND_TIMEOUT](#).)

Environments: RT (non-RT optional)

Specific return values:

- Non-negative value equal to given buffer size (Indicating the successful completion of the function call. See also note.)
- -EOPNOTSUPP (MSG_OOB flag is not supported.)
- -EINVAL (Unsupported flag detected *or*: Invalid length of socket address *or*: Invalid address family *or*: Data length code of CAN frame not between 0 and 15 *or*: CAN standard frame has got an ID not between 0 and 2031)
- -EMSGSIZE (Zero or more than one buffer passed or invalid size of buffer)
- -EFAULT (It was not possible to access user space memory area at one of the specified addresses.)
- -ENXIO (Invalid CAN interface index - 0 is not allowed here - or socket not bound or rather bound to all interfaces.)
- -ENETDOWN (Controller is bus-off or in stopped state.)
- -ECOMM (Controller is sleeping)
- -EAGAIN (Cannot transmit without blocking but a non-blocking call was requested.)
- -EINTR (Operation was interrupted explicitly or by signal)
- -EBADF (Socket was closed.)
- -ETIMEDOUT (Timeout)

Note: A successful completion of the function call does not implicate a successful transmission of the message.

5.6.2 Macro Definition Documentation

5.6.2.1 #define CAN_CTRLMODE_3_SAMPLES 0x4

Triple sampling mode

In this mode the CAN controller uses Triple sampling.

5.6.2.2 #define CAN_CTRLMODE_LISTENONLY 0x1

Listen-Only mode

In this mode the CAN controller would give no acknowledge to the CAN-bus, even if a message is received successfully and messages would not be transmitted. This mode might be useful for bus-monitoring, hot-plugging or throughput analysis.

Examples:

[rtcanconfig.c](#).

5.6.2.3 #define CAN_CTRLMODE_LOOPBACK 0x2

Loopback mode

In this mode the CAN controller does an internal loop-back, a message is transmitted and simultaneously received. That mode can be used for self test operation.

Examples:

[rtcanconfig.c](#).

5.6.2.4 #define CAN_ERR_LOSTARB_UNSPEC 0x00

unspecified

else bit number in bitstream

5.6.2.5 #define CAN_RAW_ERR_FILTER 0x2

CAN error mask.

A CAN error mask (see [Errors](#)) can be set with `setsockopt`. This mask is then used to decide if error frames are delivered to this socket in case of error conditions. The error frames are marked with the `CAN_ERR_FLAG` of `CAN_XXX_FLAG` and must be handled by the application properly. A detailed description of the errors can be found in the `can_id` and the data fields of struct `can_frame` (see [Errors](#) for further details).

Parameters

in	level	SOL_CAN_RAW
in	optname	CAN_RAW_ERR_FILTER
in	optval	Pointer to error mask of type <code>can_err_mask_t</code> .
in	optlen	Size of error mask: <code>sizeof(can_err_mask_t)</code> .

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -EINVAL (Invalid length "optlen")

Examples:

[rtcanrecv.c](#).

5.6.2.6 #define CAN_RAW_FILTER 0x1

CAN filter definition.

A CAN raw filter list with elements of struct [can_filter](#) can be installed with `setsockopt`. This list is used upon reception of CAN frames to decide whether the bound socket will receive a frame. An empty filter list can also be defined using `optlen = 0`, which is recommended for write-only sockets.

If the socket was already bound with [Bind](#), the old filter list gets replaced with the new one. Be aware that already received, but not read out CAN frames may stay in the socket buffer.

Parameters

in	level	SOL_CAN_RAW
in	optname	CAN_RAW_FILTER
in	optval	Pointer to array of struct can_filter .
in	optlen	Size of filter list: <code>count * sizeof(struct can_filter)</code> . Environments: non-RT (RT optional) Specific return values: <ul style="list-style-type: none"> • -EFAULT (It was not possible to access user space memory area at the specified address.) • -ENOMEM (Not enough memory to fulfill the operation) • -EINVAL (Invalid length "optlen") • -ENOSPC (No space to store filter list, check RT-Socket-CAN kernel parameters)

Examples:

[can-rtt.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.6.2.7 #define CAN_RAW_LOOPBACK 0x3

CAN TX loopback.

The TX loopback to other local sockets can be selected with this `setsockopt`.

Note

The TX loopback feature must be enabled in the kernel and then the loopback to other local TX sockets is enabled by default.

Parameters

in	level	SOL_CAN_RAW
in	optname	CAN_RAW_LOOPBACK
in	optval	Pointer to integer value.
in	optlen	Size of int: <code>sizeof(int)</code> .

Environments: non-RT (RT optional)

Specific return values:

- -EFAULT (It was not possible to access user space memory area at the specified address.)
- -EINVAL (Invalid length "optlen")
- -EOPNOTSUPP (not supported, check RT-Socket-CAN kernel parameters).

Examples:

[rtcansend.c](#).

5.6.2.8 `#define CAN_RAW_RECV_OWN_MSGS 0x4`

CAN receive own messages.

Not supported by RT-Socket-CAN, but defined for compatibility with Socket-CAN.

5.6.2.9 `#define RTCAN_RTIOC_RCV_TIMEOUT_IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_rel_t)`

Specify a reception timeout for a socket.

Defines a timeout for all receive operations via a socket which will take effect when one of the [receive functions](#) is called without the MSG_DONTWAIT flag set.

The default value for a newly created socket is an infinite timeout.

Note

The setting of the timeout value is not done atomically to avoid locks. Please set the value before receiving messages from the socket.

Parameters

in	arg	Pointer to nanosecs_rel_t variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See Timeouts for special timeouts.
-----------	------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

[rtcanrecv.c](#).

5.6.2.10 `#define RTCAN_RTIOC_SND_TIMEOUT_IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_t)`

Specify a transmission timeout for a socket.

Defines a timeout for all send operations via a socket which will take effect when one of the [send functions](#) is called without the MSG_DONTWAIT flag set.

The default value for a newly created socket is an infinite timeout.

Note

The setting of the timeout value is not done atomically to avoid locks. Please set the value before sending messages to the socket.

Parameters

<i>in</i>	<i>arg</i>	Pointer to nanosecs_rel_t variable. The value is interpreted as relative timeout in nanoseconds in case of a positive value. See Timeouts for special timeouts.
-----------	------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

[rtcanseend.c](#).

5.6.2.11 `#define RTCAN_RTIOC_TAKE_TIMESTAMP_IOW(RTIOC_TYPE_CAN, 0x09, int)`

Enable or disable storing a high precision timestamp upon reception of a CAN frame.

A newly created socket takes no timestamps by default.

Parameters

<i>in</i>	<i>arg</i>	int variable, see Timestamp switches
-----------	------------	--

Returns

0 on success.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Activating taking timestamps only has an effect on newly received CAN messages from the bus. Frames that already are in the socket buffer do not have timestamps if it was deactivated before. See [Receive](#) for more details.

Rescheduling: never.

Examples:

[rtcanrecv.c](#).

5.6.2.12 `#define SIOCGCANBAUDRATE _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)`

Get baud rate.

Parameters

<code>in,out</code>	<code>arg</code>	Pointer to interface request structure buffer (struct ifreq from linux/if.h). <code>ifr_name</code> must hold a valid CAN interface name, <code>ifr_ifru</code> will be filled with an instance of can_baudrate_t .
---------------------	------------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.6.2.13 `#define SIOCGCANCTRLMODE _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)`

Get special controller modes.

Parameters

<code>in</code>	<code>arg</code>	Pointer to interface request structure buffer (struct ifreq from linux/if.h). <code>ifr_name</code> must hold a valid CAN interface name, <code>ifr_ifru</code> must be filled with an instance of can_ctrlmode_t .
-----------------	------------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.6.2.14 `#define SIOCGCANCUSTOMBITTIME _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)`

Get custom bit-time parameters.

Parameters

in,out	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr_name must hold a valid CAN interface name, ifr_ifru will be filled with an instance of struct can_bittime .
--------	-----	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No baud rate was set yet.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.6.2.15 `#define SIOCGCANSTATE _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)`

Get current state of CAN controller.

States are divided into main states and additional error indicators. A CAN controller is always in exactly one main state. CAN bus errors are registered by the CAN hardware and collected by the driver. There is one error indicator (bit) per error type. If this IOCTL is triggered the error types which occurred since the last call of this IOCTL are reported and thereafter the error indicators are cleared. See also [CAN controller states](#).

Parameters

in,out	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr_name must hold a valid CAN interface name, ifr_ifru will be filled with an instance of can_mode_t .
--------	-----	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.6.2.16 `#define SIOCGIFINDEX` defined_by_kernel_header_file

Get CAN interface index by name.

Parameters

in,out	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). If ifr_name holds a valid CAN interface name ifr_ifindex will be filled with the corresponding interface index.
--------	-----	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Examples:

[can-rtt.c](#), [rtcanconfig.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.6.2.17 `#define SIOCSCANBAUDRATE` _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)

Set baud rate.

The baudrate must be specified in bits per second. The driver will try to calculate resonable CAN bit-timing parameters. You can use [SIOCSCANCUSTOMBITTIME](#) to set custom bit-timing.

Parameters

in	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr_name must hold a valid CAN interface name, ifr_ifru must be filled with an instance of can_baudrate_t .
-----------	------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can_baudrate_t](#).
- -EDOM : Baud rate not possible.
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Setting the baud rate is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

[rtcanconfig.c](#).

5.6.2.18 #define SIOCSCANCTRLMODE _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)

Set special controller modes.

Various special controller modes could be or'ed together (see [CAN_CTRLMODE](#) for further information).

Parameters

in	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr_name must hold a valid CAN interface name, ifr_ifru must be filled with an instance of can_ctrlmode_t .
-----------	------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can_baudrate_t](#).
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Setting special controller modes is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

[rtcanconfig.c](#).

5.6.2.19 #define SIOCSCANCUSTOMBITTIME _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)

Set custom bit time parameter.

Custum-bit time could be defined in various formats (see struct [can_bittime](#)).

Parameters

in	arg	Pointer to interface request structure buffer (struct ifreq from linux/if.h). ifr_name must hold a valid CAN interface name, ifr_ifru must be filled with an instance of struct can_bittime .
----	-----	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EINVAL: No valid baud rate, see [can_baudrate_t](#).
- -EAGAIN: Request could not be successfully fulfilled. Try again.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Setting the bit-time is a configuration task. It should be done deliberately or otherwise CAN messages will likely be lost.

Rescheduling: possible.

Examples:

[rtcanconfig.c](#).

5.6.2.20 `#define SIOCSCANMODE _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)`

Set operation mode of CAN controller.

See [CAN controller modes](#) for available modes.

Parameters

<code>in</code>	<code>arg</code>	Pointer to interface request structure buffer (<code>struct ifreq</code> from <code>linux/if.h</code>). <code>ifr_name</code> must hold a valid CAN interface name, <code>ifr_ifru</code> must be filled with an instance of can_mode_t .
-----------------	------------------	---

Returns

0 on success, otherwise:

- -EFAULT: It was not possible to access user space memory area at the specified address.
- -ENODEV: No device with specified name exists.
- -EAGAIN: ([CAN_MODE_START](#), [CAN_MODE_STOP](#)) Could not successfully set mode, hardware is busy. Try again.
- -EINVAL: ([CAN_MODE_START](#)) Cannot start controller, set baud rate first.
- -ENETDOWN: ([CAN_MODE_SLEEP](#)) Cannot go into sleep mode because controller is stopped or bus off.
- -EOPNOTSUPP: unknown mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

Setting a CAN controller into normal operation after a bus-off can take some time (128 occurrences of 11 consecutive recessive bits). In such a case, although this IOCTL will return immediately with success and [SIOCGCANSTATE](#) will report [CAN_STATE_ACTIVE](#), bus-off recovery may still be in progress.

If a controller is bus-off, setting it into stop mode will return no error but the controller remains bus-off.

Rescheduling: possible.

Examples:

[rtcanconfig.c](#).

5.6.2.21 `#define SOL_CAN_RAW 103`

CAN socket levels.

Used for [Sockopts](#) for the particular protocols.

Examples:

[can-rtt.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.6.3 Typedef Documentation

5.6.3.1 typedef struct **can_filter** **can_filter_t**

Filter for reception of CAN messages.

This filter works as follows: A received CAN ID is AND'ed bitwise with `can_mask` and then compared to `can_id`. This also includes the `CAN EFF_FLAG` and `CAN_RTR_FLAG` of `CAN_xxx_FLAG`. If this comparison is true, the message will be received by the socket. The logic can be inverted with the `can_id` flag `CAN_INV_FILTER` :

```
if (can_id & CAN_INV_FILTER) {
    if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))
        accept-message;
} else {
    if ((received_can_id & can_mask) == can_id)
        accept-message;
}
```

Multiple filters can be arranged in a filter list and set with `Sockopts`. If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

5.6.3.2 typedef struct **can_frame** **can_frame_t**

Raw CAN frame.

Central structure for receiving and sending CAN frames.

Examples:

[rtcanrecv.c](#).

5.6.4 Enumeration Type Documentation

5.6.4.1 enum **CAN_BITTIME_TYPE**

Supported CAN bit-time types.

Enumerator

CAN_BITTIME_STD Standard bit-time definition according to Bosch.

CAN_BITTIME_BTR Hardware-specific BTR bit-time definition.

5.6.4.2 enum **CAN_MODE**

Enumerator

CAN_MODE_STOP Set controller in Stop mode (no reception / transmission possible)

CAN_MODE_START Set controller into normal operation.

Coming from stopped mode or bus off, the controller begins with no errors in `CAN_STATE_ACTIVE`.

CAN_MODE_SLEEP Set controller into Sleep mode.

This is only possible if the controller is not stopped or bus-off.

Notice that sleep mode will only be entered when there is no bus activity. If the controller detects bus activity while "sleeping" it will go into operating mode again.

To actively leave sleep mode again trigger `CAN_MODE_START`.

5.6.4.3 enum **CAN_STATE**

Enumerator

CAN_STATE_ERROR_ACTIVE CAN controller is error active.

CAN_STATE_ACTIVE CAN controller is active.

CAN_STATE_ERROR_WARNING CAN controller is error active, warning level is reached.

CAN_STATE_BUS_WARNING CAN controller is error active, warning level is reached.

CAN_STATE_ERROR_PASSIVE CAN controller is error passive.

CAN_STATE_BUS_PASSIVE CAN controller is error passive.

CAN_STATE_BUS_OFF CAN controller went into Bus Off.

CAN_STATE_SCANNING_BAUDRATE CAN controller is scanning to get the baudrate.


CAN_STATE_STOPPED CAN controller is in stopped mode.

CAN_STATE_SLEEPING CAN controller is in Sleep mode.

5.7 Real-Time Driver Model

The Real-Time Driver Model (RTDM) provides a unified interface to both users and developers of real-time device drivers.

Collaboration diagram for Real-Time Driver Model:



Modules

- [RTDM User API](#)
Application interface to RTDM services.
- [Driver programming interface](#)
RTDM driver programming interface.
- [Device Profiles](#)
Pre-defined classes of real-time devices.

Typedefs

- typedef uint64_t [nanosecs_abs_t](#)
RTDM type for representing absolute dates.
- typedef int64_t [nanosecs_rel_t](#)
RTDM type for representing relative intervals.

API Versioning

- #define [RTDM_API_VER](#) 9
Common user and driver API version.
- #define [RTDM_API_MIN_COMPAT_VER](#) 9
Minimum API revision compatible with the current release.

RTDM_TIMEOUT_XXX

Special timeout values

- `#define RTDM_TIMEOUT_INFINITE 0`
Block forever.
- `#define RTDM_TIMEOUT_NONE (-1)`
Any negative timeout means non-blocking.

5.7.1 Detailed Description

The Real-Time Driver Model (RTDM) provides a unified interface to both users and developers of real-time device drivers. Specifically, it addresses the constraints of mixed RT/non-RT systems like Xenomai. RTDM conforms to POSIX semantics (IEEE Std 1003.1) where available and applicable.

API Revision: 8

5.7.2 Macro Definition Documentation

5.7.2.1 `#define RTDM_TIMEOUT_INFINITE 0`

Block forever.

5.7.2.2 `#define RTDM_TIMEOUT_NONE (-1)`

Any negative timeout means non-blocking.

5.7.3 Typedef Documentation

5.7.3.1 `typedef uint64_t nanosecs_abs_t`

RTDM type for representing absolute dates.

Its base type is a 64 bit unsigned integer. The unit is 1 nanosecond.

Examples:

[rtcanrecv.c](#).

5.7.3.2 `typedef int64_t nanosecs_rel_t`

RTDM type for representing relative intervals.

Its base type is a 64 bit signed integer. The unit is 1 nanosecond. Relative intervals can also encode the special timeouts "infinite" and "non-blocking", see [RTDM_TIMEOUT_xxx](#).

Examples:

[rtcanrecv.c](#).

5.8 RTDM User API

Application interface to RTDM services.

Collaboration diagram for RTDM User API:



Files

- file [rtdm.h](#)
Real-Time Driver Model for Xenomai, user API header.

Functions

- int [rt_dev_open](#) (const char *path, int oflag,...)
Open a device.
- int [rt_dev_socket](#) (int protocol_family, int socket_type, int protocol)
Create a socket.
- int [rt_dev_close](#) (int fd)
Close a device or socket.
- int [rt_dev_ioctl](#) (int fd, int request,...)
Issue an IOCTL.
- ssize_t [rt_dev_read](#) (int fd, void *buf, size_t nbyte)
Read from device.
- ssize_t [rt_dev_write](#) (int fd, const void *buf, size_t nbyte)
Write to device.
- ssize_t [rt_dev_recvmsg](#) (int fd, struct msghdr *msg, int flags)
Receive message from socket.
- ssize_t [rt_dev_recvfrom](#) (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)
Receive message from socket.
- ssize_t [rt_dev_recv](#) (int fd, void *buf, size_t len, int flags)
Receive message from socket.
- ssize_t [rt_dev_sendmsg](#) (int fd, const struct msghdr *msg, int flags)
Transmit message to socket.
- ssize_t [rt_dev_sendto](#) (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)
Transmit message to socket.
- ssize_t [rt_dev_send](#) (int fd, const void *buf, size_t len, int flags)
Transmit message to socket.
- int [rt_dev_bind](#) (int fd, const struct sockaddr *my_addr, socklen_t addrlen)
Bind to local address.
- int [rt_dev_connect](#) (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)

- *Connect to remote address.*
- int `rt_dev_listen` (int fd, int backlog)
Listen for incoming connection requests.
- int `rt_dev_accept` (int fd, struct sockaddr *addr, socklen_t *addrlen)
Accept a connection requests.
- int `rt_dev_shutdown` (int fd, int how)
Shut down parts of a connection.
- int `rt_dev_getsockopt` (int fd, int level, int optname, void *optval, socklen_t *optlen)
Get socket option.
- int `rt_dev_setsockopt` (int fd, int level, int optname, const void *optval, socklen_t optlen)
Set socket option.
- int `rt_dev_getsockname` (int fd, struct sockaddr *name, socklen_t *namelen)
Get local socket address.
- int `rt_dev_getpeername` (int fd, struct sockaddr *name, socklen_t *namelen)
Get socket destination address.

5.8.1 Detailed Description

Application interface to RTDM services. This is the upper interface of RTDM provided to application programs both in kernel and user space. Note that certain functions may not be implemented by every device. Refer to the [Device Profiles](#) for precise information.

5.8.2 Function Documentation

5.8.2.1 int `rt_dev_accept` (int fd, struct sockaddr * addr, socklen_t * addrlen)

Accept a connection requests.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
out	<i>addr</i>	Buffer for remote address
in,out	<i>addrlen</i>	Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`accept()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.2 int `rt_dev_bind` (int fd, const struct sockaddr * my_addr, socklen_t addrlen)

Bind to local address.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>my_addr</i>	Address buffer
in	<i>addrlen</i>	Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`bind()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcanrecv.c](#), and [rtcansend.c](#).

5.8.2.3 int rt_dev_close (int fd)

Close a device or socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_open() or rt_dev_socket()
----	-----------	---

Returns

0 on success, otherwise a negative error code.

Note

If the matching [rt_dev_open\(\)](#) or [rt_dev_socket\(\)](#) call took place in non-real-time context, [rt_dev_close\(\)](#) must be issued within non-real-time as well. Otherwise, the call will fail.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`close()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[cross-link.c](#), [rtcanconfig.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.8.2.4 int rt_dev_connect (int fd, const struct sockaddr * serv_addr, socklen_t addrlen)

Connect to remote address.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>serv_addr</i>	Address buffer
in	<i>addrlen</i>	Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`connect()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.5 `int rt_dev_getpeername (int fd, struct sockaddr * name, socklen_t * namelen)`

Get socket destination address.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
out	<i>name</i>	Address buffer
in,out	<i>namelen</i>	Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`getpeername()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.6 `int rt_dev_getsockname (int fd, struct sockaddr * name, socklen_t * namelen)`

Get local socket address.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
out	<i>name</i>	Address buffer
in,out	<i>namelen</i>	Address buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

getsockname() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.7 `int rt_dev_getsockopt (int fd, int level, int optname, void * optval, socklen_t * optlen)`

Get socket option.

Parameters

<code>in</code>	<code>fd</code>	File descriptor as returned by rt_dev_socket()
<code>in</code>	<code>level</code>	Addressed stack level
<code>in</code>	<code>optname</code>	Option name ID
<code>out</code>	<code>optval</code>	Value buffer
<code>in,out</code>	<code>optlen</code>	Value buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

getsockopt() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.8 `int rt_dev_ioctl (int fd, int request, ...)`

Issue an IOCTL.

Parameters

<code>in</code>	<code>fd</code>	File descriptor as returned by rt_dev_open() or rt_dev_socket()
<code>in</code>	<code>request</code>	IOCTL code
	<code>...</code>	Optional third argument, depending on IOCTL function (void * or unsigned long)

Returns

Positiv value on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

ioctl() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[cross-link.c](#), [rtcanconfig.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.8.2.9 `int rt_dev_listen (int fd, int backlog)`

Listen for incoming connection requests.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>backlog</i>	Maximum queue length

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`lsiten()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.10 `int rt_dev_open (const char * path, int oflag, ...)`

Open a device.

Parameters

in	<i>path</i>	Device name
in	<i>oflag</i>	Open flags
	<i>...</i>	Further parameters will be ignored.

Returns

Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`open()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[cross-link.c](#).

5.8.2.11 `ssize_t rt_dev_read (int fd, void * buf, size_t nbyte)`

Read from device.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_open()
----	-----------	--

out	<i>buf</i>	Input buffer
in	<i>nbyte</i>	Number of bytes to read

Returns

Number of bytes read, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

read() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[cross-link.c](#).

5.8.2.12 `ssize_t rt_dev_recv (int fd, void * buf, size_t len, int flags)`

Receive message from socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
out	<i>buf</i>	Message buffer
in	<i>len</i>	Message buffer size
in	<i>flags</i>	Message flags

Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

recv() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.13 `ssize_t rt_dev_recvfrom (int fd, void * buf, size_t len, int flags, struct sockaddr * from, socklen_t * fromlen)`

Receive message from socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
----	-----------	--

out	<i>buf</i>	Message buffer
in	<i>len</i>	Message buffer size
in	<i>flags</i>	Message flags
out	<i>from</i>	Buffer for message sender address
in,out	<i>fromlen</i>	Address buffer size

Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`recvfrom()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcanrecv.c](#).

5.8.2.14 `ssize_t rt_dev_recvmsg (int fd, struct msghdr * msg, int flags)`

Receive message from socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in,out	<i>msg</i>	Message descriptor
in	<i>flags</i>	Message flags

Returns

Number of bytes received, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`recvmsg()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcanrecv.c](#).

5.8.2.15 `ssize_t rt_dev_send (int fd, const void * buf, size_t len, int flags)`

Transmit message to socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>buf</i>	Message buffer
in	<i>len</i>	Message buffer size
in	<i>flags</i>	Message flags

Returns

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`send()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcansend.c](#).

5.8.2.16 `ssize_t rt_dev_sendmsg (int fd, const struct msghdr * msg, int flags)`

Transmit message to socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>msg</i>	Message descriptor
in	<i>flags</i>	Message flags

Returns

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`sendmsg()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.17 `ssize_t rt_dev_sendto (int fd, const void * buf, size_t len, int flags, const struct sockaddr * to, socklen_t tolen)`

Transmit message to socket.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>buf</i>	Message buffer
in	<i>len</i>	Message buffer size
in	<i>flags</i>	Message flags
in	<i>to</i>	Buffer for message destination address
in	<i>tolen</i>	Address buffer size

Returns

Number of bytes sent, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`sendto()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcansend.c](#).

5.8.2.18 `int rt_dev_setsockopt (int fd, int level, int optname, const void * optval, socklen_t optlen)`

Set socket option.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>level</i>	Addressed stack level
in	<i>optname</i>	Option name ID
in	<i>optval</i>	Value buffer
in	<i>optlen</i>	Value buffer size

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`setsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcanrecv.c](#), and [rtcansend.c](#).

5.8.2.19 `int rt_dev_shutdown (int fd, int how)`

Shut down parts of a connection.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_socket()
in	<i>how</i>	Specifies the part to be shut down (SHUT_XXX)

Returns

0 on success, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`shutdown()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

5.8.2.20 `int rt_dev_socket (int protocol_family, int socket_type, int protocol)`

Create a socket.

Parameters

in	<i>protocol_family</i>	Protocol family (PF_XXX)
in	<i>socket_type</i>	Socket type (SOCK_XXX)
in	<i>protocol</i>	Protocol ID, 0 for default

Returns

Positive file descriptor value on success, otherwise a negative error code.

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`socket()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[rtcanconfig.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

5.8.2.21 `ssize_t rt_dev_write (int fd, const void * buf, size_t nbyte)`

Write to device.

Parameters

in	<i>fd</i>	File descriptor as returned by rt_dev_open()
in	<i>buf</i>	Output buffer

<code>in</code>	<code>nbyte</code>	Number of bytes to write
-----------------	--------------------	--------------------------

Returns

Number of bytes written, otherwise negative error code

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

See Also

`write()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

Examples:

[cross-link.c](#).

5.9 Serial Devices

This is the common interface a RTDM-compliant serial device has to provide.

Collaboration diagram for Serial Devices:



This is the common interface a RTDM-compliant serial device has to provide. Feel free to comment on this profile via the Xenomai mailing list (Xenomai-core@gnat.org) or directly to the author (jan.kiszka@web.de).

Profile Revision: 3

Device Characteristics

Device Flags: RTDM_NAMED_DEVICE, RTDM_EXCLUSIVE

Device Name: "rtser<N>", N >= 0

Device Class: RTDM_CLASS_SERIAL

Supported Operations

Open

Environments: non-RT (RT optional, deprecated)

Specific return values: none

Close

Environments: non-RT (RT optional, deprecated)

Specific return values: none

IOCTL

Mandatory Environments: see [below](#)

Specific return values: see [below](#)

Read

Environments: RT (non-RT optional)

Specific return values:

- -ETIMEDOUT
- -EINTR (interrupted explicitly or by signal)
- -EAGAIN (no data available in non-blocking mode)
- -EBADF (device has been closed while reading)
- -EIO (hardware error or broken bit stream)

Write

Environments: RT (non-RT optional)

Specific return values:

- -ETIMEDOUT
- -EINTR (interrupted explicitly or by signal)
- -EAGAIN (no data written in non-blocking mode)
- -EBADF (device has been closed while writing)

5.10 Testing Devices

This group of devices is intended to provide in-kernel testing results.

Collaboration diagram for Testing Devices:



This group of devices is intended to provide in-kernel testing results. Feel free to comment on this profile via the Xenomai mailing list (xenomai-core@gna.org) or directly to the author (jan.kiszka@web.de).

Profile Revision: 2

Device Characteristics

Device Flags: RTDM_NAMED_DEVICE

Device Name: "rttest[-<subclass>]<N>", N >= 0, optional subclass name to simplify device discovery

Device Class: RTDM_CLASS_TESTING

Supported Operations

Open

Environments: non-RT (RT optional, deprecated)

Specific return values: none

Close

Environments: non-RT (RT optional, deprecated)

Specific return values: none

IOCTL

Mandatory Environments: see [TSTIOCTLs](#) below

Specific return values: see [TSTIOCTLs](#) below

5.11 Real-time IPC

Profile Revision: 1

Collaboration diagram for Real-time IPC:



Data Structures

- struct [rtipc_port_label](#)
Port label information structure.
- struct [sockaddr_ipc](#)
Socket address structure for the RTIPC address family.

Typedefs

- typedef int16_t [rtipc_port_t](#)
Port number type for the RTIPC address family.

Supported operations

Standard socket operations supported by the RTIPC protocols.

- int [socket__AF_RTIPC](#) (int domain=AF_RTIPC, int type=SOCK_DGRAM, int protocol)
Create an endpoint for communication in the AF_RTIPC domain.
- int [close__AF_RTIPC](#) (int sockfd)
Close a RTIPC socket descriptor.
- int [bind__AF_RTIPC](#) (int sockfd, const struct [sockaddr_ipc](#) *addr, socklen_t addrlen)
Bind a RTIPC socket to a port.
- int [connect__AF_RTIPC](#) (int sockfd, const struct [sockaddr_ipc](#) *addr, socklen_t addrlen)
Initiate a connection on a RTIPC socket.
- int [setsockopt__AF_RTIPC](#) (int sockfd, int level, int optname, const void *optval, socklen_t optlen)
Set options on RTIPC sockets.
- int [getsockopt__AF_RTIPC](#) (int sockfd, int level, int optname, void *optval, socklen_t *optlen)
Get options on RTIPC sockets.
- ssize_t [sendmsg__AF_RTIPC](#) (int sockfd, const struct msghdr *msg, int flags)
Send a message on a RTIPC socket.
- ssize_t [recvmsg__AF_RTIPC](#) (int sockfd, struct msghdr *msg, int flags)
Receive a message from a RTIPC socket.
- int [getsockname__AF_RTIPC](#) (int sockfd, struct [sockaddr_ipc](#) *addr, socklen_t *addrlen)
Get socket name.
- int [getpeername__AF_RTIPC](#) (int sockfd, struct [sockaddr_ipc](#) *addr, socklen_t *addrlen)
Get socket peer.

RTIPC protocol list

protocols for the PF_RTIPC protocol family

- enum { [IPCPROTO_IPC](#) = 0, [IPCPROTO_XDDP](#) = 1, [IPCPROTO_IDDP](#) = 2, [IPCPROTO_BUFP](#) = 3 }

XDDP socket options

Setting and getting XDDP socket options.

- #define [XDDP_LABEL](#) 1
XDDP label assignment.
- #define [XDDP_POOLSZ](#) 2
XDDP local pool size configuration.
- #define [XDDP_BUFSZ](#) 3
XDDP streaming buffer size configuration.
- #define [XDDP_MONITOR](#) 4
XDDP monitoring callback.

XDDP events

Specific events occurring on XDDP channels, which can be monitored via the [XDDP_MONITOR](#) socket option.

- #define [XDDP_EVTIN](#) 1
Monitor writes to the non real-time endpoint.
- #define [XDDP_EVTOUT](#) 2
Monitor reads from the non real-time endpoint.
- #define [XDDP_EVTDOWN](#) 3
Monitor close from the non real-time endpoint.
- #define [XDDP_EVTNOBUF](#) 4
Monitor memory shortage for non real-time datagrams.

IDDP socket options

Setting and getting IDDP socket options.

- #define [IDDP_LABEL](#) 1
IDDP label assignment.
- #define [IDDP_POOLSZ](#) 2
IDDP local pool size configuration.

BUFP socket options

Setting and getting BUFP socket options.

- #define [BUFP_LABEL](#) 1
BUFP label assignment.
- #define [BUFP_BUFSZ](#) 2
BUFP buffer size configuration.

Socket level options

Setting and getting supported standard socket level options.

- `#define SO_SNDTIMEO` defined_by_kernel_header_file
IPPROTO_IDDP and IPPROTO_BUF protocols support the standard SO_SNDTIMEO socket option, from the SOL_SOCKET level.
- `#define SO_RCVTIMEO` defined_by_kernel_header_file
All RTIPC protocols support the standard SO_RCVTIMEO socket option, from the SOL_SOCKET level.

5.11.1 Detailed Description

Profile Revision: 1

Device Characteristics

Device Flags: RTDM_PROTOCOL_DEVICE
Protocol Family: PF_RTIPC
Socket Type: SOCK_DGRAM
Device Class: RTDM_CLASS_RTIPC

5.11.2 Macro Definition Documentation

5.11.2.1 `#define BUFP_BUFSZ 2`

BUFP buffer size configuration.

All messages written to a BUFP socket are buffered in a single per-socket memory area. Configuring the size of such buffer prior to binding the socket to a destination port is mandatory.

It is not allowed to configure a buffer size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

: the buffer memory is obtained from the host allocator by the [bind call](#).

Parameters

in	level	SOL_BUF
in	optname	BUFP_BUFSZ
in	optval	Pointer to a variable of type <code>size_t</code> , containing the required size of the buffer to reserve at binding time
in	optlen	<code>sizeof(size_t)</code>

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid or **optval* is zero)

Calling context:

RT/non-RT

Examples:

[bufp-label.c](#), and [bufp-readwrite.c](#).

5.11.2.2 #define BUFP_LABEL 1

BUFP label assignment.

ASCII label strings can be attached to BUFP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see [BUFP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

in	level	SOL_BUFP
in	optname	BUFP_LABEL
in	optval	Pointer to struct rtipc_port_label
in	optlen	sizeof(struct rtipc_port_label)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT

Examples:

[bufp-label.c](#).

5.11.2.3 #define IDDP_LABEL 1

IDDP label assignment.

ASCII label strings can be attached to IDDP ports, in order to connect sockets to them in a more descriptive way than using plain numeric port values.

When available, this label will be registered when binding, in addition to the port number (see [IDDP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

in	level	SOL_IDDP
in	optname	IDDP_LABEL
in	optval	Pointer to struct rtipc_port_label
in	optlen	sizeof(struct rtipc_port_label)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT

Examples:

[iddp-label.c](#).

5.11.2.4 #define IDDP_POOLSZ 2

IDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

: the pool memory is obtained from the host allocator by the [bind call](#).

Parameters

in	level	SOL_IDDP
in	optname	IDDP_POOLSZ
in	optval	Pointer to a variable of type <code>size_t</code> , containing the required size of the local pool to reserve at binding time
in	optlen	sizeof(size_t)

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* is invalid or **optval* is zero)

Calling context:

RT/non-RT

Examples:

[iddp-sendrecv.c](#).

5.11.2.5 `#define SO_RCVTIMEO` defined_by_kernel_header_file

All RTIPC protocols support the standard `SO_RCVTIMEO` socket option, from the `SOL_SOCKET` level.

See Also

`setsockopt()`, `getsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399/>

Examples:

[xddp-label.c](#).

Referenced by `rt_pipe_read_timed()`.

5.11.2.6 `#define SO_SNDTIMEO` defined_by_kernel_header_file

[IPPROTO_IDDP](#) and [IPPROTO_BUF](#) protocols support the standard `SO_SNDTIMEO` socket option, from the `SOL_SOCKET` level.

See Also

`setsockopt()`, `getsockopt()` in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399/>

5.11.2.7 `#define XDDP_BUFSZ` 3

XDDP streaming buffer size configuration.

In addition to sending datagrams, real-time threads may stream data in a byte-oriented mode through the port as well. This increases the bandwidth and reduces the overhead, when the overall data to send to the Linux domain is collected by bits, and keeping the message boundaries is not required.

This feature is enabled when a non-zero buffer size is set for the socket. In that case, the real-time data accumulates into the streaming buffer when `MSG_MORE` is passed to any of the [send functions](#), until:

- the receiver from the Linux domain wakes up and consumes it,
- a different source port attempts to send data to the same destination port,
- `MSG_MORE` is absent from the send flags,
- the buffer is full,

whichever comes first.

Setting **optval* to zero disables the streaming buffer, in which case all sendings are conveyed in separate datagrams, regardless of `MSG_MORE`.

Note

only a single streaming buffer exists per socket. When this buffer is full, the real-time data stops accumulating and sending operations resume in mere datagram mode. Accumulation may happen again after some or all data in the streaming buffer is consumed from the Linux domain endpoint.

The streaming buffer size may be adjusted multiple times during the socket lifetime; the latest configuration change will take effect when the accumulation resumes after the previous buffer was flushed.

Parameters

in	<i>level</i>	SOL_XDDP
in	<i>optname</i>	XDDP_BUFSZ
in	<i>optval</i>	Pointer to a variable of type <code>size_t</code> , containing the required size of the streaming buffer
in	<i>optlen</i>	<code>sizeof(size_t)</code>

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -ENOMEM (Not enough memory)
- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT

Examples:

[xddp-stream.c](#).

Referenced by `rt_pipe_create()`.

5.11.2.8 `#define XDDP_EVTDOWN 3`

[Monitor](#) close from the non real-time endpoint.

XDDP_EVTDOWN is sent when the non real-time endpoint is closed. The argument is always 0.

5.11.2.9 `#define XDDP_EVTIN 1`

[Monitor](#) writes to the non real-time endpoint.

XDDP_EVTIN is sent when data is written to the non real-time endpoint the socket is bound to (i.e. via `/dev/rtpN`), which means that some input is pending for the real-time endpoint. The argument is the size of the incoming message.

5.11.2.10 `#define XDDP_EVTNOBUF 4`

[Monitor](#) memory shortage for non real-time datagrams.

XDDP_EVTNOBUF is sent when no memory is available from the pool to hold the message currently sent from the non real-time endpoint. The argument is the size of the failed allocation. Upon return from the callback, the caller will block and retry until enough space is available from the pool; during that process, the callback might be invoked multiple times, each time a new attempt to get the required memory fails.

5.11.2.11 `#define XDDP_EVTOUT 2`

[Monitor](#) reads from the non real-time endpoint.

XDDP_EVTOUT is sent when the non real-time endpoint successfully reads a complete message (i.e. via `/dev/rtpN`). The argument is the size of the outgoing message.

5.11.2.12 `#define XDDP_LABEL 1`

XDDP label assignment.

ASCII label strings can be attached to XDDP ports, so that opening the non-RT endpoint can be done by specifying this symbolic device name rather than referring to a raw pseudo-device entry (i.e. `/dev/rtpN`).

When available, this label will be registered when binding, in addition to the port number (see [XDDP port binding](#)).

It is not allowed to assign a label after the socket was bound. However, multiple assignment calls are allowed prior to the binding; the last label set will be used.

Parameters

in	<i>level</i>	SOL_XDDP
in	<i>optname</i>	XDDP_LABEL
in	<i>optval</i>	Pointer to struct rtipc_port_label
in	<i>optlen</i>	<code>sizeof(struct rtipc_port_label)</code>

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* invalid)

Calling context:

RT/non-RT

Examples:

[xddp-label.c](#).

Referenced by `rt_pipe_create()`.

5.11.2.13 `#define XDDP_MONITOR 4`

XDDP monitoring callback.

Other RTDM drivers may install a user-defined callback via the [rtdm_setsockopt](#) call from the inter-driver API, in order to collect particular events occurring on the channel.

This notification mechanism is particularly useful to monitor a channel asynchronously while performing other tasks.

The user-provided routine will be passed the RTDM file descriptor of the socket receiving the event, the event code, and an optional argument. Four events are currently defined, see [XDDP_EVENTS](#).

The XDDP_EVTIN and XDDP_EVTOUT events are fired on behalf of a fully atomic context; therefore, care must be taken to keep their overhead low. In those cases, the Xenomai services that may be called from the callback are restricted to the set allowed to a real-time interrupt handler.

Parameters

in	level	SOL_XDDP
in	optname	XDDP_MONITOR
in	optval	Pointer to a pointer to function of type <code>int (*)(int fd, int event, long arg)</code> , containing the address of the user-defined callback. Passing a NULL callback pointer in <i>optval</i> disables monitoring.
in	optlen	<code>sizeof(int (*)(int fd, int event, long arg))</code>

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EPERM (Operation not allowed from user-space)
- -EINVAL (*optlen* is invalid)

Calling context:

RT/non-RT, kernel space only

5.11.2.14 #define XDDP_POOLSZ 2

XDDP local pool size configuration.

By default, the memory needed to convey the data is pulled from Xenomai's system pool. Setting a local pool size overrides this default for the socket.

If a non-zero size was configured, a local pool is allocated at binding time. This pool will provide storage for pending datagrams.

It is not allowed to configure a local pool size after the socket was bound. However, multiple configuration calls are allowed prior to the binding; the last value set will be used.

Note

: the pool memory is obtained from the host allocator by the [bind call](#).

Parameters

in	level	SOL_XDDP
in	optname	XDDP_POOLSZ
in	optval	Pointer to a variable of type <code>size_t</code> , containing the required size of the local pool to reserve at binding time
in	optlen	<code>sizeof(size_t)</code>

Returns

0 is returned upon success. Otherwise:

- -EFAULT (Invalid data address given)
- -EALREADY (socket already bound)
- -EINVAL (*optlen* invalid or **optval* is zero)

Calling context:

RT/non-RT

Examples:

[xddp-echo.c](#).

Referenced by `rt_pipe_create()`.

5.11.3 Enumeration Type Documentation

5.11.3.1 anonymous enum

Enumerator

IPCPROTO_IPC Default protocol (IDDP)

IPCPROTO_XDDP Cross-domain datagram protocol (RT <-> non-RT). Real-time Xenomai threads and regular Linux threads may want to exchange data in a way that does not require the former to leave the real-time domain (i.e. secondary mode). The RTDM-based XDDP protocol is available for this purpose.

On the Linux domain side, pseudo-device files named `/dev/rtp<minor>` give regular POSIX threads access to non real-time communication endpoints, via the standard character-based I/O interface. On the Xenomai domain side, sockets may be bound to XDDP ports, which act as proxies to send and receive data to/from the associated pseudo-device files. Ports and pseudo-device minor numbers are paired, meaning that e.g. port 7 will proxy the traffic to/from `/dev/rtp7`.

All data sent through a bound/connected XDDP socket via `sendto(2)` or `write(2)` will be passed to the peer endpoint in the Linux domain, and made available for reading via the standard `read(2)` system call. Conversely, all data sent using `write(2)` through the non real-time endpoint will be conveyed to the real-time socket endpoint, and made available to the `recvfrom(2)` or `read(2)` system calls.

IPCPROTO_IDDP Intra-domain datagram protocol (RT <-> RT). The RTDM-based IDDP protocol enables real-time threads to exchange datagrams within the Xenomai domain, via socket endpoints.

IPCPROTO_BUFP Buffer protocol (RT <-> RT, byte-oriented). The RTDM-based BUFP protocol implements a lightweight, byte-oriented, one-way Producer-Consumer data path. All messages written are buffered into a single memory area in strict FIFO order, until read by the consumer.

This protocol always prevents short writes, and only allows short reads when a potential deadlock situation arises (i.e. readers and writers waiting for each other indefinitely).

5.11.4 Function Documentation

5.11.4.1 `int bind__AF_RTIPC (int sockfd, const struct sockaddr_ipc * addr, socklen_t addrlen)`

Bind a RTIPC socket to a port.

Bind the socket to a destination port.

Parameters

<code>in</code>	<code>sockfd</code>	The RTDM file descriptor obtained from the socket creation call.
<code>in</code>	<code>addr</code>	The address to bind the socket to (see struct sockaddr_ipc). The meaning of such address depends on the RTIPC protocol in use for the socket:

- **IPCPROTO_XDDP**

This action creates an endpoint for channelling traffic between the Xenomai and Linux domains.

sipc_family must be `AF_RTIPC`, *sipc_port* is either -1, or a valid free port number between 0 and `CONFIG_XENO_OPT_PIPE_NRDEV-1`.

If *sipc_port* is -1, a free port will be assigned automatically.

Upon success, the pseudo-device `/dev/rtpN` will be reserved for this communication channel, where *N* is the assigned port number. The non real-time side shall open this device to exchange data over the bound socket.

If a label was assigned (see [XDDP_LABEL](#)) prior to binding the socket to a port, a registry link referring to the created pseudo-device will be automatically set up as `/proc/xenomai/registry/rtipc/xddp/label`, where *label* is the label string passed to `setsockopt()` for the [XDDP_LABEL](#) option.

- `IPCPROTO_IDDP`

This action creates an endpoint for exchanging datagrams within the Xenomai domain.

sipc_family must be `AF_RTIPC`, *sipc_port* is either -1, or a valid free port number between 0 and `CONFIG_XENO_OPT_IDDP_NRPORT-1`.

If *sipc_port* is -1, a free port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see [IDDP_LABEL](#)) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as `/proc/xenomai/registry/rtipc/iddp/label`, where *label* is the label string passed to `setsockopt()` for the [IDDP_LABEL](#) option.

- `IPCPROTO_BUFP`

This action creates an endpoint for a one-way byte stream within the Xenomai domain.

sipc_family must be `AF_RTIPC`, *sipc_port* is either -1, or a valid free port number between 0 and `CONFIG_XENO_OPT_BUFP_NRPORT-1`.

If *sipc_port* is -1, an available port will be assigned automatically. The real-time peer shall connect to the same port for exchanging data over the bound socket.

If a label was assigned (see [BUF_P_LABEL](#)) prior to binding the socket to a port, a registry link referring to the assigned port number will be automatically set up as `/proc/xenomai/registry/rtipc/bufp/label`, where *label* is the label string passed to `setsockopt()` for the [BUF_P_LABEL](#) option.

Parameters

<code>in</code>	<code>addrlen</code>	The size in bytes of the structure pointed to by <i>addr</i> .
-----------------	----------------------	--

Returns

In addition to the standard error codes for `bind(2)`, the following specific error code may be returned:

- `-EFAULT` (Invalid data address given)
- `-ENOMEM` (Not enough memory)
- `-EINVAL` (Invalid parameter)
- `-EADDRINUSE` (Socket already bound to a port, or no port available)

Calling context:

non-RT

5.11.4.2 `int close__AF_RTIPC (int sockfd)`

Close a RTIPC socket descriptor.

Blocking calls to any of the [sendmsg](#) or [recvmsg](#) functions will be unblocked when the socket is closed and return with an error.

Parameters

in	<i>sockfd</i>	The socket descriptor to close.
-----------	---------------	---------------------------------

Returns

In addition to the standard error codes for `close(2)`, the following specific error code may be returned: none

Calling context:

non-RT

5.11.4.3 `int connect__AF_RTIPC (int sockfd, const struct sockaddr_ipc * addr, socklen_t addrlen)`

Initiate a connection on a RTIPC socket.

Parameters

in	<i>sockfd</i>	The RTDM file descriptor obtained from the socket creation call.
in	<i>addr</i>	The address to connect the socket to (see struct sockaddr_ipc).

- If `sipc_port` is a valid port for the protocol, it is used verbatim and the connection succeeds immediately, regardless of whether the destination is bound at the time of the call.
- If `sipc_port` is -1 and a label was assigned to the socket, `connect()` blocks for the requested amount of time (see [SO_RCVTIMEO](#)) until a socket is bound to the same label via `bind(2)` (see [XDDP_LABEL](#), [IDDP_LABEL](#), [BUFP_LABEL](#)), in which case a connection is established between both endpoints.
- If `sipc_port` is -1 and no label was assigned to the socket, the default destination address is cleared, meaning that any subsequent write to the socket will return `-EDESTADDRREQ`, until a valid destination address is set via `connect(2)` or `bind(2)`.

Parameters

in	<i>addrlen</i>	The size in bytes of the structure pointed to by <i>addr</i> .
-----------	----------------	--

Returns

In addition to the standard error codes for `connect(2)`, the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.11.4.4 `int getpeername__AF_RTIPC (int sockfd, struct sockaddr_ipc * addr, socklen_t * addrlen)`

Get socket peer.

The name of the remote endpoint for the socket is copied back (see struct [sockaddr_ipc](#)). This is the default destination address for messages sent on the socket. It can be set either explicitly via `connect(2)`, or implicitly via `bind(2)` if no `connect(2)` was called prior to binding the socket to a port, in which case both the local and remote names are equal.

Returns

In addition to the standard error codes for `getpeername(2)`, the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.11.4.5 `int getsockname__AF_RTIPC (int sockfd, struct sockaddr_ipc * addr, socklen_t * addrlen)`

Get socket name.

The name of the local endpoint for the socket is copied back (see struct [sockaddr_ipc](#)).

Returns

In addition to the standard error codes for `getsockname(2)`, the following specific error code may be returned: none.

Calling context:

RT/non-RT

5.11.4.6 `int getsockopt__AF_RTIPC (int sockfd, int level, int optname, void * optval, socklen_t * optlen)`

Get options on RTIPC sockets.

These functions allow to get various socket options. Supported Levels and Options:

- Level [SOL_SOCKET](#)
- Level [SOL_XDDP](#)
- Level [SOL_IDDP](#)
- Level [SOL_BUFP](#)

Returns

In addition to the standard error codes for `getsockopt(2)`, the following specific error code may be returned: follow the option links above.

Calling context:

RT/non-RT

5.11.4.7 `ssize_t recvmsg__AF_RTIPC (int sockfd, struct msghdr * msg, int flags)`

Receive a message from a RTIPC socket.

Parameters

in	<i>sockfd</i>	The RTDM file descriptor obtained from the socket creation call.
out	<i>msg</i>	The address the message header will be copied at.
in	<i>flags</i>	Operation flags:

- **MSG_DONTWAIT** Non-blocking I/O operation. The caller will not be blocked whenever no message is immediately available for receipt at the time of the call, but will rather return with **-EWOULDBLOCK**.

Note

IPPROTO_BUF does not allow for short reads and always returns the requested amount of bytes, except in one situation: whenever some writer is waiting for sending data upon a buffer full condition, while the caller would have to wait for receiving a complete message. This is usually the sign of a pathological use of the **BUF** socket, like defining an incorrect buffer size via **BUF_BUFSZ**. In that case, a short read is allowed to prevent a deadlock.

Returns

In addition to the standard error codes for `recvmsg(2)`, the following specific error code may be returned: none.

Calling context:

RT

5.11.4.8 `ssize_t sendmsg__AF_RTIPC (int sockfd, const struct msghdr * msg, int flags)`

Send a message on a RTIPC socket.

Parameters

in	<i>sockfd</i>	The RTDM file descriptor obtained from the socket creation call.
in	<i>msg</i>	The address of the message header conveying the datagram.
in	<i>flags</i>	Operation flags:

- **MSG_OOB** Send out-of-band message. For all RTIPC protocols except **IPPROTO_BUF**, sending out-of-band data actually means pushing them to the head of the receiving queue, so that the reader will always receive them before normal messages. **IPPROTO_BUF** does not support out-of-band sending.
- **MSG_DONTWAIT** Non-blocking I/O operation. The caller will not be blocked whenever the message cannot be sent immediately at the time of the call (e.g. memory shortage), but will rather return with **-EWOULDBLOCK**. Unlike other RTIPC protocols, **IPPROTO_XDDP** accepts but never considers **MSG_DONTWAIT** since writing to a real-time XDDP endpoint is inherently a non-blocking operation.
- **MSG_MORE** Accumulate data before sending. This flag is accepted by the **IPPROTO_XDDP** protocol only, and tells the send service to accumulate the outgoing data into an internal streaming buffer, instead of issuing a datagram immediately for it. See **XDDP_BUFSZ** for more.

Note

No RTIPC protocol allows for short writes, and only complete messages are sent to the peer.

Returns

In addition to the standard error codes for `sendmsg(2)`, the following specific error code may be returned: none.

Calling context:

RT

5.11.4.9 `int setsockopt__AF_RTIPC (int sockfd, int level, int optname, const void * optval, socklen_t optlen)`

Set options on RTIPC sockets.

These functions allow to set various socket options. Supported Levels and Options:

- Level [SOL_SOCKET](#)
- Level [SOL_XDDP](#)
- Level [SOL_IDDP](#)
- Level [SOL_BUFP](#)

Returns

In addition to the standard error codes for `setsockopt(2)`, the following specific error code may be returned: follow the option links above.

Calling context:

non-RT

5.11.4.10 `int socket__AF_RTIPC (int domain = AF_RTIPC, int type = SOCK_DGRAM, int protocol)`

Create an endpoint for communication in the `AF_RTIPC` domain.

Parameters

in	<i>domain</i>	The communication domain. Must be <code>AF_RTIPC</code> .
in	<i>type</i>	The socket type. Must be <code>SOCK_DGRAM</code> .
in	<i>protocol</i>	Any of IPCPROTO_XDDP , IPCPROTO_IDDP , or IPCPROTO_BUFP . IPCPROTO_IPC is also valid, and refers to the default RTIPC protocol, namely IPCPROTO_IDDP .

Returns

In addition to the standard error codes for `socket(2)`, the following specific error code may be returned:

- `-ENOPROTOOPT` (Protocol is known, but not compiled in the RTIPC driver). See [RTIPC protocols](#) for available protocols.


Calling context:

non-RT

5.12 Asynchronous Procedure Calls

Services for scheduling function calls in the Linux domain.

Collaboration diagram for Asynchronous Procedure Calls:



Functions

- `int xnapc_alloc (const char *name, void(*handler)(void *cookie), void *cookie)`
Allocate an APC slot.
- `void xnapc_free (int apc)`
Releases an APC slot.
- `static void xnapc_schedule (int apc)`
Schedule an APC invocation.

5.12.1 Detailed Description

Services for scheduling function calls in the Linux domain. APC is the acronym for Asynchronous Procedure Call, a mean by which activities from the Xenomai domain can schedule deferred invocations of handlers to be run into the Linux domain, as soon as possible when the Linux kernel gets back in control.

Up to `BITS_PER_LONG` APC slots can be active at any point in time.

APC support is built upon the interrupt pipeline's virtual interrupt support.

5.12.2 Function Documentation

5.12.2.1 `int xnapc_alloc (const char * name, void(*) (void *cookie) handler, void * cookie)`

Allocate an APC slot.

APC is the acronym for Asynchronous Procedure Call, a mean by which activities from the Xenomai domain can schedule deferred invocations of handlers to be run into the Linux domain, as soon as possible when the Linux kernel gets back in control. Up to `BITS_PER_LONG` APC slots can be active at any point in time. APC support is built upon the interrupt pipeline's virtual interrupt support.

Any Linux kernel service which is callable from a regular Linux interrupt handler is in essence available to APC handlers.

Parameters

<i>name</i>	is a symbolic name identifying the APC which will get reported through the /proc/xenomai/apc interface. Passing NULL to create an anonymous APC is allowed.
<i>handler</i>	The address of the fault handler to call upon exception condition. The handle will be passed the <i>cookie</i> value unmodified.
<i>cookie</i>	A user-defined opaque pointer the APC handler receives as its sole argument.

Returns

a valid APC identifier is returned upon success, or a negative error code otherwise:

- -EINVAL is returned if *handler* is invalid.
- -EBUSY is returned if no more APC slots are available.

Remarks

Tags: none.

5.12.2.2 void xnapc_free (int apc)

Releases an APC slot.

This service deallocates an APC slot obtained by [xnapc_alloc\(\)](#).

Parameters

<i>apc</i>	The APC id. to release, as returned by a successful call to the xnapc_alloc() service.
------------	--

Remarks

Tags: none.

5.12.2.3 static inline int xnapc_schedule (int apc) [inline], [static]

Schedule an APC invocation.

This service marks the APC as pending for the Linux domain, so that its handler will be called as soon as possible, when the Linux domain gets back in control.

When posted from the Linux domain, the APC handler is fired as soon as the interrupt mask is explicitly cleared by some kernel code. When posted from the Xenomai domain, the APC handler is fired as soon as the Linux domain is resumed, i.e. after Xenomai has completed all its pending duties.

Parameters

<i>apc</i>	The APC id. to schedule.
------------	--------------------------

This service can be called from:

- Any domain context, albeit the usual calling place is from the Xenomai domain.

5.13 Helpers for in-kernel arithmetics

Collaboration diagram for Helpers for in-kernel arithmetics:



5.13.1 Detailed Description

5.14 Buffer descriptor

Abstraction for copying data to/from different address spaces.

Collaboration diagram for Buffer descriptor:



Functions

- static void [xnbufd_map_uread](#) (struct xnbufd *bufd, const void __user *ptr, size_t len)
Initialize a buffer descriptor for reading from user memory.
- static void [xnbufd_map_uwrite](#) (struct xnbufd *bufd, void __user *ptr, size_t len)
Initialize a buffer descriptor for writing to user memory.
- ssize_t [xnbufd_unmap_uread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uread\(\)](#).
- ssize_t [xnbufd_unmap_uwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uwrite\(\)](#).
- static void [xnbufd_map_kread](#) (struct xnbufd *bufd, const void *ptr, size_t len)
Initialize a buffer descriptor for reading from kernel memory.
- static void [xnbufd_map_kwrite](#) (struct xnbufd *bufd, void *ptr, size_t len)
Initialize a buffer descriptor for writing to kernel memory.
- ssize_t [xnbufd_unmap_kread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kread\(\)](#).
- ssize_t [xnbufd_unmap_kwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kwrite\(\)](#).
- ssize_t [xnbufd_copy_to_kmem](#) (void *ptr, struct xnbufd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- ssize_t [xnbufd_copy_from_kmem](#) (struct xnbufd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- void [xnbufd_invalidate](#) (struct xnbufd *bufd)
Invalidate a buffer descriptor.
- static void [xnbufd_reset](#) (struct xnbufd *bufd)
Reset a buffer descriptor.

5.14.1 Detailed Description

Abstraction for copying data to/from different address spaces. A buffer descriptor is a simple abstraction dealing with copy operations to/from memory buffers which may belong to different address spaces.

To this end, the buffer descriptor library provides a small set of copy routines which are aware of address space restrictions when moving data, and a generic container type which can hold a reference to - or cover - a particular memory area, either present in kernel space, or in any of the existing user memory contexts.

The goal of the buffer descriptor abstraction is to hide address space specifics from Xenomai services dealing with memory areas, allowing them to operate on multiple address spaces seamlessly.

The common usage patterns are as follows:

- Implementing a Xenomai syscall returning a bulk of data to the caller, which may have to be copied back to either kernel or user space:

```
[Syscall implementation]
ssize_t rt_bulk_read_inner(struct xnbuid *buid)
{
    ssize_t ret;
    size_t len;
    void *bulk;

    bulk = get_next_readable_bulk(&len);
    ret = xnbuid_copy_from_kmem(buid, bulk, min(buid->b_len, len));
    free_bulk(bulk);

    ret = this_may_fail();
    if (ret)
        xnbuid_invalidate(buid);

    return ret;
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_read(void *ptr, size_t len)
{
    struct xnbuid buid;
    ssize_t ret;

    xnbuid_map_kwrite(&buid, ptr, len);
    ret = rt_bulk_read_inner(&buid);
    xnbuid_unmap_kwrite(&buid);

    return ret;
}

[Userland trampoline for user syscalls]
int __rt_bulk_read(struct pt_regs *regs)
{
    struct xnbuid buid;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbuid_map_uwrite(&buid, ptr, len);
    ret = rt_bulk_read_inner(&buid);
    xnbuid_unmap_uwrite(&buid);

    return ret;
}
```

- Implementing a Xenomai syscall receiving a bulk of data from the caller, which may have to be read from either kernel or user space:

```
[Syscall implementation]
ssize_t rt_bulk_write_inner(struct xnbuid *buid)
{
    void *bulk = get_free_bulk(buid->b_len);
    return xnbuid_copy_to_kmem(bulk, buid, buid->b_len);
}

[Kernel wrapper for in-kernel calls]
int rt_bulk_write(const void *ptr, size_t len)
{
    struct xnbuid buid;
    ssize_t ret;

    xnbuid_map_kread(&buid, ptr, len);
    ret = rt_bulk_write_inner(&buid);
    xnbuid_unmap_kread(&buid);

    return ret;
}

[Userland trampoline for user syscalls]
```

```

int __rt_bulk_write(struct pt_regs *regs)
{
    struct xnbuid bufd;
    void __user *ptr;
    ssize_t ret;
    size_t len;

    ptr = (void __user *)__xn_reg_arg1(regs);
    len = __xn_reg_arg2(regs);

    xnbuid_map_uread(&bufd, ptr, len);
    ret = rt_bulk_write_inner(&bufd);
    xnbuid_unmap_uread(&bufd);

    return ret;
}

```

5.14.2 Function Documentation

5.14.2.1 `ssize_t xnbuid_copy_from_kmem (struct xnbuid * bufd, void * from, size_t len)`

Copy kernel memory to the area covered by a buffer descriptor.

This routine copies *len* bytes from the kernel memory starting at *from* to the area referred to by the buffer descriptor *bufd*. `xnbuid_copy_from_kmem()` tracks the write offset within the destination memory internally, so that it may be called several times in a loop, until the entire memory area is stored.

The destination address space is dealt with, according to the following rules:

- if *bufd* refers to a writable kernel area (i.e. see `xnbuid_map_kwrite()`), the copy is immediately and fully performed with no restriction.
- if *bufd* refers to a writable user area (i.e. see `xnbuid_map_uwrite()`), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while writing to that memory.
- if *bufd* refers to a user area which may not be immediately written to from the current context, the copy is postponed until `xnbuid_unmap_uwrite()` is invoked for *ubufd*, at which point the copy will take place. In such a case, the source memory is transferred to a carry over buffer allocated internally; this operation may lead to request dynamic memory from the nucleus heap if *len* is greater than 64 bytes.

Parameters

<i>bufd</i>	The address of the buffer descriptor covering the user memory to copy data to.
<i>from</i>	The start address of the kernel memory to copy from.
<i>len</i>	The length of the kernel memory to copy to <i>bufd</i> .

Returns

The number of bytes written so far to the memory area covered by *ubufd*. Otherwise,

- -ENOMEM is returned when no memory is available from the nucleus heap to allocate the carry over buffer.

Remarks

Tags: none.

Note

Calling this routine while holding the `nklock` and/or running with interrupts disabled is invalid, and doing so will trigger a debug assertion.

This routine may switch the caller to secondary mode if a page fault occurs while reading from the user area. For that reason, `xnbufd_copy_to_kmem()` may only be called from a preemptible section (Linux-wise).

5.14.2.2 `ssize_t xnbufd_copy_to_kmem (void * to, struct xnbufd * bufd, size_t len)`

Copy memory covered by a buffer descriptor to kernel memory.

This routine copies `len` bytes from the area referred to by the buffer descriptor `bufd` to the kernel memory area `to`. `xnbufd_copy_to_kmem()` tracks the read offset within the source memory internally, so that it may be called several times in a loop, until the entire memory area is loaded.

The source address space is dealt with, according to the following rules:

- if `bufd` refers to readable kernel area (i.e. see `xnbufd_map_kread()`), the copy is immediately and fully performed with no restriction.
- if `bufd` refers to a readable user area (i.e. see `xnbufd_map_uread()`), the copy is performed only if that area lives in the currently active address space, and only if the caller may sleep Linux-wise to process any potential page fault which may arise while reading from that memory.
- any attempt to read from `bufd` from a non-suitable context is considered as a bug, and will raise a panic assertion when the nucleus is compiled in debug mode.

Parameters

<code>to</code>	The start address of the kernel memory to copy to.
<code>bufd</code>	The address of the buffer descriptor covering the user memory to copy data from.
<code>len</code>	The length of the user memory to copy from <code>bufd</code> .

Returns

The number of bytes read so far from the memory area covered by `ubufd`. Otherwise:

- `-EINVAL` is returned upon attempt to read from the user area from an invalid context. This error is only returned when the debug mode is disabled; otherwise a panic assertion is raised.

Remarks

Tags: none.

Note

Calling this routine while holding the `nklock` and/or running with interrupts disabled is invalid, and doing so will trigger a debug assertion.

This routine may switch the caller to secondary mode if a page fault occurs while reading from the user area. For that reason, `xnbufd_copy_to_kmem()` may only be called from a preemptible section (Linux-wise).

5.14.2.3 void `xnbufd_invalidate` (struct `xnbufd` * `bufd`)

Invalidate a buffer descriptor.

The buffer descriptor is invalidated, making it unusable for further copy operations. If an outstanding carry over buffer was allocated by a previous call to `xnbufd_copy_from_kmem()`, it is immediately freed so that no data transfer will happen when the descriptor is finalized.

The only action that may subsequently be performed on an invalidated descriptor is calling the relevant unmapping routine for it. For that reason, `xnbufd_invalidate()` should be invoked on the error path when data may have been transferred to the carry over buffer.

Parameters

<i>bufd</i>	The address of the buffer descriptor to invalidate.
-------------	---

Remarks

Tags: isr-allowed.

5.14.2.4 void `xnbufd_map_kread` (struct `xnbufd` * `bufd`, const void * `ptr`, size_t `len`) [inline], [static]

Initialize a buffer descriptor for reading from kernel memory.

The new buffer descriptor may be used to copy data from kernel memory. This routine should be used in pair with `xnbufd_unmap_kread()`.

Parameters

<i>bufd</i>	The address of the buffer descriptor which will map a <i>len</i> bytes kernel memory area, starting from <i>ptr</i> .
<i>ptr</i>	The start of the kernel buffer to map.
<i>len</i>	The length of the kernel buffer starting at <i>ptr</i> .

Remarks

Tags: isr-allowed.

5.14.2.5 void `xnbufd_map_kwrite` (struct `xnbufd` * `bufd`, void * `ptr`, size_t `len`) [inline], [static]

Initialize a buffer descriptor for writing to kernel memory.

The new buffer descriptor may be used to copy data to kernel memory. This routine should be used in pair with `xnbufd_unmap_kwrite()`.

Parameters

<i>bufd</i>	The address of the buffer descriptor which will map a <i>len</i> bytes kernel memory area, starting from <i>ptr</i> .
<i>ptr</i>	The start of the kernel buffer to map.
<i>len</i>	The length of the kernel buffer starting at <i>ptr</i> .

Remarks

Tags: isr-allowed.

5.14.2.6 `void xnbuid_map_uread (struct xnbuid * bufd, const void __user * ptr, size_t len)`
`[inline], [static]`

Initialize a buffer descriptor for reading from user memory.

The new buffer descriptor may be used to copy data from user memory. This routine should be used in pair with [xnbuid_unmap_uread\(\)](#).

Parameters

<i>bufd</i>	The address of the buffer descriptor which will map a <i>len</i> bytes user memory area, starting from <i>ptr</i> . <i>ptr</i> is never dereferenced directly, since it may refer to a buffer that lives in another address space.
<i>ptr</i>	The start of the user buffer to map.
<i>len</i>	The length of the user buffer starting at <i>ptr</i> .

Remarks

Tags: none.

5.14.2.7 `void xnbuid_map_uwrite (struct xnbuid * bufd, void __user * ptr, size_t len)` `[inline],`
`[static]`

Initialize a buffer descriptor for writing to user memory.

The new buffer descriptor may be used to copy data to user memory. This routine should be used in pair with [xnbuid_unmap_uwrite\(\)](#).

Parameters

<i>bufd</i>	The address of the buffer descriptor which will map a <i>len</i> bytes user memory area, starting from <i>ptr</i> . <i>ptr</i> is never dereferenced directly, since it may refer to a buffer that lives in another address space.
<i>ptr</i>	The start of the user buffer to map.
<i>len</i>	The length of the user buffer starting at <i>ptr</i> .

Remarks

Tags: none.

5.14.2.8 `void xnbuid_reset (struct xnbuid * bufd)` `[inline], [static]`

Reset a buffer descriptor.

The buffer descriptor is reset, so that all data already copied is forgotten. Any carry over buffer allocated is kept, though.

Parameters

<i>bufd</i>	The address of the buffer descriptor to reset.
-------------	--

Remarks

Tags: isr-allowed.

5.14.2.9 `ssize_t xnbuid_unmap_kread (struct xnbuid * bufd)`

Finalize a buffer descriptor obtained from [xnbuid_map_kread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_kread\(\)](#), to read data from a kernel area.

Parameters

<i>bufd</i>	The address of the buffer descriptor to finalize.
-------------	---

Returns

The number of bytes read so far from the memory area covered by *ubufd*.

Remarks

Tags: isr-allowed.

5.14.2.10 `ssize_t xnbuid_unmap_kwrite (struct xnbuid * bufd)`

Finalize a buffer descriptor obtained from [xnbuid_map_kwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_kwrite\(\)](#), to write data to a kernel area.

Parameters

<i>bufd</i>	The address of the buffer descriptor to finalize.
-------------	---

Returns

The number of bytes written so far to the memory area covered by *ubufd*.

Remarks

Tags: isr-allowed.

5.14.2.11 `ssize_t xnbuid_unmap_uread (struct xnbuid * bufd)`

Finalize a buffer descriptor obtained from [xnbuid_map_uread\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_uread\(\)](#), to read data from a user area.

Parameters

<i>bufd</i>	The address of the buffer descriptor to finalize.
-------------	---

Returns

The number of bytes read so far from the memory area covered by *ubufd*.

Remarks

Tags: none.

Note

Calling this routine while holding the `nklock` and/or running with interrupts disabled is invalid, and doing so will trigger a debug assertion.

5.14.2.12 `ssize_t xnbuid_unmap_uwrite (struct xnbuid * bufd)`

Finalize a buffer descriptor obtained from [xnbuid_map_uwrite\(\)](#).

This routine finalizes a buffer descriptor previously initialized by a call to [xnbuid_map_uwrite\(\)](#), to write data to a user area.

The main action taken is to write the contents of the kernel memory area passed to [xnbuid_copy_from_kmem\(\)](#) whenever the copy operation was postponed at that time; the carry over buffer is eventually released as needed. If [xnbuid_copy_from_kmem\(\)](#) was allowed to copy to the destination user memory at once, then [xnbuid_unmap_uwrite\(\)](#) leads to a no-op.

Parameters

<i>bufd</i>	The address of the buffer descriptor to finalize.
-------------	---

Returns

The number of bytes written so far to the memory area covered by *ubufd*.

Remarks

Tags: none.

Note

Calling this routine while holding the `nklock` and/or running with interrupts disabled is invalid, and doing so will trigger a debug assertion.

5.15 Clock services.

Collaboration diagram for Clock services.:



Functions

- int [xnclock_register](#) (struct xnclock *clock)
Register a Xenomai clock.
- void [xnclock_deregister](#) (struct xnclock *clock)
Deregister a Xenomai clock.
- void [xnclock_tick](#) (struct xnclock *clock)
Process a clock tick.
- void [xnclock_adjust](#) (struct xnclock *clock, xnsticks_t delta)
Adjust a clock time.

5.15.1 Detailed Description

Note

Copyright (C) 2006-2011 Philippe Gerum rpm@xenomai.org.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

5.15.2 Function Documentation

5.15.2.1 void xnclock_adjust (struct xnclock * clock, xnsticks_t delta)

Adjust a clock time.

This service changes the epoch for the given clock by applying the specified tick delta on its wallclock offset.

Parameters

<i>clock</i>	The clock to adjust.
<i>delta</i>	The adjustment value expressed in nanoseconds.

Note

This routine must be entered *nklock* locked, interrupts off.

Remarks

Tags: none.

Note

Xenomai tracks the system time in *nkclock*, as a monotonously increasing count of ticks since the epoch. The epoch is initially the same as the underlying machine time.

5.15.2.2 void xnclock_deregister (struct xnclock * clock)

Deregister a Xenomai clock.

This service uninstalls a Xenomai clock previously registered with [xnclock_register\(\)](#).

This service may be called once all timers driven by *clock* have been stopped.

Parameters

<i>clock</i>	The clock to deregister.
--------------	--------------------------

Remarks

Tags: secondary-only.

5.15.2.3 int xnclock_register (struct xnclock * clock)

Register a Xenomai clock.

This service installs a new clock which may be used to drive Xenomai timers.

Parameters

<i>clock</i>	The new clock to register.
--------------	----------------------------

Remarks

Tags: secondary-only.

5.15.2.4 void xnclock_tick (struct xnclock * clock)

Process a clock tick.

This routine processes an incoming *clock* event, firing elapsed timers as appropriate.

Parameters

<i>clock</i>	The clock for which a new event was received.
--------------	---

Remarks

Tags: primary-only, isr-only, atomic-entry.

Note

The current CPU must be part of the real-time affinity set, otherwise weird things may happen.

References `xnsched::htimer`, `xnsched::lflags`, and `xnsched::status`.

5.16 Debugging services

Copyright (C) 2010 Philippe Gerum rpm@xenomai.org.

Collaboration diagram for Debugging services:



5.16.1 Detailed Description

Copyright (C) 2010 Philippe Gerum rpm@xenomai.org. Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

5.17 Dynamic memory allocation services.

Copyright (C) 2001,2002,2003 Philippe Gerum rpm@xenomai.org.

Collaboration diagram for Dynamic memory allocation services.:



Functions

- int [xnheap_init](#) (struct xnheap *heap, void *heapaddr, unsigned long heapsize, unsigned long pagesize)
Initialize a memory heap.
- void [xnheap_set_label](#) (struct xnheap *heap, const char *label,...)
Set the heap's label string.
- void [xnheap_destroy](#) (struct xnheap *heap, void(*flushfn)(struct xnheap *heap, void *extaddr, unsigned long extsize, void *cookie), void *cookie)
Destroys a memory heap.
- void * [xnheap_alloc](#) (struct xnheap *heap, unsigned long size)
Allocate a memory block from a memory heap.
- int [xnheap_test_and_free](#) (struct xnheap *heap, void *block, int(*ckfn)(void *block))
Test and release a memory block to a memory heap.
- int [xnheap_free](#) (struct xnheap *heap, void *block)
Release a memory block to a memory heap.
- int [xnheap_extend](#) (struct xnheap *heap, void *extaddr, unsigned long extsize)
Extend a memory heap.

5.17.1 Detailed Description

Copyright (C) 2001,2002,2003 Philippe Gerum rpm@xenomai.org. Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

The implementation of the memory allocator follows the algorithm described in a USENIX 1988 paper called "Design of a General Purpose Memory Allocator for the 4.3BSD Unix Kernel" by Marshall K. McKusick and Michael J. Karels. You can find it at various locations on the net, including <http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf>. A minor variation allows this implementation to have 'extendable' heaps when needed, with multiple memory extents providing autonomous page address spaces.

The data structures hierarchy is as follows:

```

HEAP {
    block_buckets[]
    extent_queue -----+
}
                        |
                        V
                EXTENT #1 {
                        {static header}
                        page_map[npages]
                        page_array[npages][pagesize]
                } --+
                    |
                    |
                    V
                EXTENT #n {
                        {static header}
                        page_map[npages]
                        page_array[npages][pagesize]
                }

```

5.17.2 Function Documentation

5.17.2.1 void* xnheap_alloc (struct xnheap * heap, unsigned long size)

Allocate a memory block from a memory heap.

Allocates a contiguous region of memory from an active memory heap. Such allocation is guaranteed to be time-bounded.

Parameters

<i>heap</i>	The descriptor address of the heap to get memory from.
<i>size</i>	The size in bytes of the requested block. Sizes lower or equal to the page size are rounded either to the minimum allocation size if lower than this value, or to the minimum alignment size if greater or equal to this value. In the current implementation, with MINALLOC = 8 and MINALIGN = 16, a 7 bytes request will be rounded to 8 bytes, and a 17 bytes request will be rounded to 32.

Returns

The address of the allocated region upon success, or NULL if no memory is available from the specified heap.

Remarks

Tags: isr-allowed.

Referenced by xnshadow_map_user().

5.17.2.2 void xnheap_destroy (struct xnheap * heap, void(*) (struct xnheap *heap, void *extaddr, unsigned long extsize, void *cookie) flushfn, void * cookie)

Destroys a memory heap.

Destroys a memory heap.

Parameters

<i>heap</i>	The descriptor address of the destroyed heap.
<i>flushfn</i>	If non-NULL, the address of a flush routine which will be called for each extent attached to the heap. This routine can be used by the calling code to further release the heap memory.
<i>cookie</i>	If <i>flushfn</i> is non-NULL, <i>cookie</i> is an opaque pointer which will be passed unmodified to <i>flushfn</i> .

Remarks

Tags: none.

5.17.2.3 `int xnheap_extend (struct xnheap * heap, void * extaddr, unsigned long extsize)`

Extend a memory heap.

Add a new extent to an existing memory heap.

Parameters

<i>heap</i>	The descriptor address of the heap to add an extent to.
<i>extaddr</i>	The address of the extent memory.
<i>extsize</i>	The size of the extent memory (in bytes). In the current implementation, this size must match the one of the initial extent passed to xnheap_init() .

Returns

0 is returned upon success, or -EINVAL is returned if *extsize* differs from the initial extent's size.

Remarks

Tags: isr-allowed.

5.17.2.4 `int xnheap_free (struct xnheap * heap, void * block)`

Release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from.

Parameters

<i>heap</i>	The descriptor address of the heap to release memory to.
<i>block</i>	The address of the region to be returned to the heap.

Returns

0 is returned upon success, or one of the following error codes:

- -EFAULT is returned whenever the memory address is outside the heap address space.
- -EINVAL is returned whenever the memory address does not represent a valid block.

Remarks

Tags: isr-allowed.

References [xnheap_test_and_free\(\)](#).

5.17.2.5 `int xnheap_init (struct xnheap * heap, void * heapaddr, unsigned long heapsize, unsigned long pagesize)`

Initialize a memory heap.

Initializes a memory heap suitable for time-bounded allocation requests of dynamic memory.

Parameters

<i>heap</i>	The address of a heap descriptor which will be used to store the allocation data. This descriptor must always be valid while the heap is active therefore it must be allocated in permanent memory.
<i>heapaddr</i>	The address of the heap storage area. All allocations will be made from the given area in time-bounded mode. Since additional extents can be added to a heap, this parameter is also known as the "initial extent".
<i>heapsize</i>	The size in bytes of the initial extent pointed at by <i>heapaddr</i> . <i>heapsize</i> must be a multiple of <i>pagesize</i> and lower than 16 Mbytes. <i>heapsize</i> must be large enough to contain a dynamically-sized internal header. The following formula gives the size of this header: $H = \text{heapsize}, P = \text{pagesize}, M = \text{sizeof}(\text{struct pagemap}), E = \text{sizeof}(\text{struct xnextent})$ $\text{hdrsize} = ((H - E) * M) / (M + 1)$ This value is then aligned on the next 16-byte boundary. The routine <code>xnheap_overhead()</code> computes the corrected heap size according to the previous formula.
<i>pagesize</i>	The size in bytes of the fundamental memory page which will be used to subdivide the heap internally. Choosing the right page size is important regarding performance and memory fragmentation issues, so it might be a good idea to take a look at http://docs.FreeBSD.org/44doc/papers/kernmalloc.pdf to pick the best one for your needs. In the current implementation, <i>pagesize</i> must be a power of two in the range [8 .. 32768] inclusive.

Returns

0 is returned upon success, or one of the following error codes:

- -EINVAL is returned whenever a parameter is invalid.

Remarks

Tags: none.

5.17.2.6 `void xnheap_set_label (struct xnheap * heap, const char * label, ...)`

Set the heap's label string.

Set the heap label that will be used in statistic outputs.

Parameters

<i>heap</i>	The address of a heap descriptor.
<i>label</i>	Label string displayed in statistic outputs. This parameter can be a format string, in which case succeeding parameters will be used to resolve the final label.

Remarks

Tags: none.

5.17.2.7 `int xnheap_test_and_free (struct xnheap * heap, void * block, int(*)(void *block) ckfn)`

Test and release a memory block to a memory heap.

Releases a memory region to the memory heap it was previously allocated from. Before the actual release is performed, an optional user-defined can be invoked to check for additional criteria with respect to the request consistency.

Parameters

<i>heap</i>	The descriptor address of the heap to release memory to.
<i>block</i>	The address of the region to be returned to the heap.
<i>ckfn</i>	The address of a user-supplied verification routine which is to be called after the memory address specified by <i>block</i> has been checked for validity. The routine is expected to proceed to further consistency checks, and either return zero upon success, or non-zero upon error. In the latter case, the release process is aborted, and <i>ckfn</i> 's return value is passed back to the caller of this service as its error return code.

Warning

ckfn must not reschedule either directly or indirectly.

Returns

0 is returned upon success, or -EINVAL is returned whenever the block is not a valid region of the specified heap. Additional return codes can also be defined locally by the *ckfn* routine.

Remarks

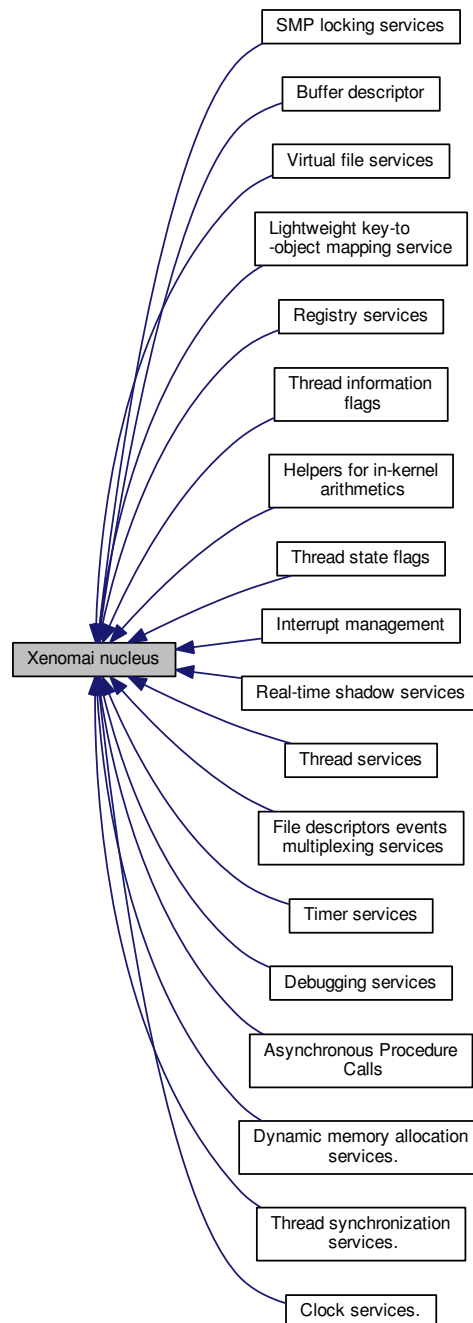
Tags: isr-allowed.

Referenced by `xnheap_free()`.

5.18 Xenomai nucleus

The Xenomai nucleus supplements the native Linux kernel in dual kernel configurations.

Collaboration diagram for Xenomai nucleus:



Modules

- [Thread state flags](#)

- Bits reporting permanent or transient states of threads.*
- [Thread information flags](#)
 - Bits reporting events notified to threads.*
- [Asynchronous Procedure Calls](#)
 - Services for scheduling function calls in the Linux domain.*
- [Helpers for in-kernel arithmetics](#)
- [Buffer descriptor](#)
 - Abstraction for copying data to/from different address spaces.*
- [Clock services.](#)
- [Debugging services](#)
 - Copyright (C) 2010 Philippe Gerum rpm@xenomai.org.*
- [Dynamic memory allocation services.](#)
 - Copyright (C) 2001,2002,2003 Philippe Gerum rpm@xenomai.org.*
- [Interrupt management](#)
- [SMP locking services](#)
 - Copyright (C) 2001-2012 Philippe Gerum rpm@xenomai.org.*
- [Lightweight key-to-object mapping service](#)
 - A map is a simple indexing structure which associates unique integer keys with pointers to objects.*
- [Registry services](#)
 - The registry provides a mean to index object descriptors on unique alphanumeric keys.*
- [File descriptors events multiplexing services](#)
 - File descriptors events multiplexing services.*
- [Real-time shadow services](#)
 - Copyright (C) 2001-2013 Philippe Gerum rpm@xenomai.org.*
- [Thread synchronization services.](#)
 - Copyright (C) 2001-2008 Philippe Gerum rpm@xenomai.org.*
- [Thread services](#)
- [Timer services](#)
 - The Xenomai timer facility depends on a clock source (xnclock) for scheduling the next activation times.*
- [Virtual file services](#)
 - Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces.*

5.18.1 Detailed Description

The Xenomai nucleus supplements the native Linux kernel in dual kernel configurations. It deals with all time-critical activities, such as handling interrupts, and scheduling real-time threads. The nucleus has higher priority over all the native kernel activities.

5.19 Interrupt management

Collaboration diagram for Interrupt management:



Functions

- int [xnintr_init](#) (struct xnintr *intr, const char *name, unsigned int irq, xnintr_t isr, xniack_t iack, int flags)
Initialize an interrupt object.
- void [xnintr_destroy](#) (struct xnintr *intr)
Destroy an interrupt object.
- int [xnintr_attach](#) (struct xnintr *intr, void *cookie)
Attach an interrupt object.
- void [xnintr_detach](#) (struct xnintr *intr)
Detach an interrupt object.
- void [xnintr_enable](#) (struct xnintr *intr)
Enable an interrupt object.
- void [xnintr_disable](#) (struct xnintr *intr)
Disable an interrupt object.
- void [xnintr_affinity](#) (struct xnintr *intr, cpumask_t cpumask)
Set interrupt's processor affinity.

5.19.1 Detailed Description

5.19.2 Function Documentation

5.19.2.1 void xnintr_affinity (struct xnintr * intr, cpumask_t cpumask)

Set interrupt's processor affinity.

Restricts the IRQ associated with the interrupt object *intr* to be received only on processors which bits are set in *cpumask*.

Parameters

<i>intr</i>	The descriptor address of the interrupt object which affinity is to be changed.
<i>cpumask</i>	The new processor affinity of the interrupt object.

Note

Depending on architectures, setting more than one bit in *cpumask* could be meaningless.

Remarks

Tags: secondary-only.

5.19.2.2 `int xnintr_attach (struct xnintr * intr, void * cookie)`

Attach an interrupt object.

Attach an interrupt object previously initialized by [xnintr_init\(\)](#). After this operation is completed, all IRQs received from the corresponding interrupt channel are directed to the object's ISR.

Parameters

<i>intr</i>	The descriptor address of the interrupt object to attach.
<i>cookie</i>	A user-defined opaque value which is stored into the interrupt object descriptor for further retrieval by the ISR/ISR handlers.

Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if a low-level error occurred while attaching the interrupt.
- -EBUSY is returned if the interrupt object was already attached.

Note

The caller **must not** hold `nklock` when invoking this service, this would cause deadlocks.

Remarks

Tags: secondary-only.

Note

Attaching an interrupt resets the tracked number of receipts to zero.

Referenced by `rtdm_irq_request()`.

5.19.2.3 `void xnintr_destroy (struct xnintr * intr)`

Destroy an interrupt object.

Destroys an interrupt object previously initialized by [xnintr_init\(\)](#). The interrupt object is automatically detached by a call to [xnintr_detach\(\)](#). No more IRQs will be dispatched by this object after this service has returned.

Parameters

<i>intr</i>	The descriptor address of the interrupt object to destroy.
-------------	--

Remarks

Tags: secondary-only.

References `xnintr_detach()`.

5.19.2.4 `void xnintr_detach (struct xnintr * intr)`

Detach an interrupt object.

Detach an interrupt object previously attached by [xnintr_attach\(\)](#). After this operation is completed, no more IRQs are directed to the object's ISR, but the interrupt object itself remains valid. A detached interrupt object can be attached again by a subsequent call to [xnintr_attach\(\)](#).

Parameters

<i>intr</i>	The descriptor address of the interrupt object to detach.
-------------	---

Note

The caller **must not** hold `nklock` when invoking this service, this would cause deadlocks.

Remarks

Tags: secondary-only.

Referenced by `xnintr_destroy()`.

5.19.2.5 `void xnintr_disable (struct xnintr * intr)`

Disable an interrupt object.

Disables the hardware interrupt line associated with an interrupt object. This operation invalidates further interrupt requests from the given source until the IRQ line is re-enabled anew.

Parameters

<i>intr</i>	The descriptor address of the interrupt object to disable.
-------------	--

Remarks

Tags: secondary-only.

5.19.2.6 `void xnintr_enable (struct xnintr * intr)`

Enable an interrupt object.

Enables the hardware interrupt line associated with an interrupt object.

Parameters

<i>intr</i>	The descriptor address of the interrupt object to enable.
-------------	---

Remarks

Tags: secondary-only.

Referenced by `rtm_irq_request()`.

5.19.2.7 `int xnintr_init (struct xnintr * intr, const char * name, unsigned int irq, xnintr_t isr, xniack_t iack, int flags)`

Initialize an interrupt object.

Associates an interrupt object with an IRQ line.

When an interrupt occurs on the given *irq* line, the ISR is fired in order to deal with the hardware event. The interrupt service code may call any non-blocking service from the nucleus.

Upon receipt of an IRQ, the ISR is immediately called on behalf of the interrupted stack context, the rescheduling procedure is locked, and the interrupt source is masked at hardware level. The status value returned by the ISR is then checked for the following values:

- `XN_ISR_HANDLED` indicates that the interrupt request has been fulfilled by the ISR.
- `XN_ISR_NONE` indicates the opposite to `XN_ISR_HANDLED`. The ISR must always return this value when it determines that the interrupt request has not been issued by the dedicated hardware device.

In addition, one of the following bits may be set by the ISR :

Warning

Use these bits with care and only when you do understand their effect on the system. The ISR is not encouraged to use these bits in case it shares the IRQ line with other ISRs in the real-time domain.

- `XN_ISR_NOENABLE` prevents the IRQ line from being re-enabled after the ISR has returned.
- `XN_ISR_PROPAGATE` causes the IRQ event to be propagated down the pipeline to Linux. This is the regular way to share interrupts between the nucleus and the regular Linux kernel. In effect, `XN_ISR_PROPAGATE` implies `XN_ISR_NOENABLE` since it would make no sense to re-enable the IRQ line before the Linux kernel had a chance to process the propagated interrupt.

A count of interrupt receipts is tracked into the interrupt descriptor, and reset to zero each time the interrupt object is attached. Since this count could wrap around, it should be used as an indication of interrupt activity only.

Parameters

<i>intr</i>	The address of a interrupt object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.
<i>name</i>	An ASCII string standing for the symbolic name of the interrupt object or NULL.
<i>irq</i>	The hardware interrupt channel associated with the interrupt object. This value is architecture-dependent. An interrupt object must then be attached to the hardware interrupt vector using the <code>xnintr_attach()</code> service for the associated IRQs to be directed to this object.
<i>isr</i>	The address of a valid low-level interrupt service routine if this parameter is non-zero. This handler will be called each time the corresponding IRQ is delivered on behalf of an interrupt context. When called, the ISR is passed the descriptor address of the interrupt object.
<i>iack</i>	The address of an optional interrupt acknowledge routine, aimed at replacing the default one. Only very specific situations actually require to override the default setting for this parameter, like having to acknowledge non-standard PIC hardware. <i>iack</i> should return a non-zero value to indicate that the interrupt has been properly acknowledged. If <i>iack</i> is NULL, the default routine will be used instead.
<i>flags</i>	A set of creation flags affecting the operation. The valid flags are:

- `XN_ISR_SHARED` enables IRQ-sharing with other interrupt objects.
- `XN_ISR_EDGE` is an additional flag need to be set together with `XN_ISR_SHARED` to enable IRQ-sharing of edge-triggered interrupts.

Returns

0 is returned on success. Otherwise, `-EINVAL` is returned if *irq* is not a valid interrupt number.

Remarks

Tags: secondary-only.

Referenced by `rtdm_irq_request()`.

5.20 SMP locking services

Copyright (C) 2001-2012 Philippe Gerum rpm@xenomai.org.

Collaboration diagram for SMP locking services:



5.20.1 Detailed Description

Copyright (C) 2001-2012 Philippe Gerum rpm@xenomai.org. Copyright (C) 2001-2008,2012 Philippe Gerum rpm@xenomai.org.

Copyright (C) 2004,2005 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

5.21 Lightweight key-to-object mapping service

A map is a simple indexing structure which associates unique integer keys with pointers to objects.

Collaboration diagram for Lightweight key-to-object mapping service:



Functions

- `struct xnmap * xnmap_create (int nkeys, int reserve, int offset)`
Create a map.
- `void xnmap_delete (struct xnmap *map)`
Delete a map.
- `int xnmap_enter (struct xnmap *map, int key, void *objaddr)`
Index an object into a map.
- `int xnmap_remove (struct xnmap *map, int key)`
Remove an object reference from a map.
- `static void * xnmap_fetch_nocheck (struct xnmap *map, int key)`
Search an object into a map - unchecked form.
- `static void * xnmap_fetch (struct xnmap *map, int key)`
Search an object into a map.

5.21.1 Detailed Description

A map is a simple indexing structure which associates unique integer keys with pointers to objects. The current implementation supports reservation, for naming/indexing objects, either on a fixed, user-provided integer (i.e. a reserved key value), or by drawing the next available key internally if the caller did not specify any fixed key. For instance, in some given map, the key space ranging from 0 to 255 could be reserved for fixed keys, whilst the range from 256 to 511 could be available for drawing free keys dynamically.

A maximum of 1024 unique keys per map is supported on 32bit machines.

(This implementation should not be confused with C++ STL maps, which are dynamically expandable and allow arbitrary key types; Xenomai maps don't).

5.21.2 Function Documentation

5.21.2.1 `struct xnmap * xnmap_create (int nkeys, int reserve, int offset)`

Create a map.

Allocates a new map with the specified addressing capabilities. The memory is obtained from the Xenomai system heap.

Parameters

<i>nkeys</i>	The maximum number of unique keys the map will be able to hold. This value cannot exceed the static limit represented by <code>XNMAP_MAX_KEYS</code> , and must be a power of two.
<i>reserve</i>	The number of keys which should be kept for reservation within the index space. Reserving a key means to specify a valid key to the <code>xnmap_enter()</code> service, which will then attempt to register this exact key, instead of drawing the next available key from the unreserved index space. When reservation is in effect, the unreserved index space will hold key values greater than <i>reserve</i> , keeping the low key values for the reserved space. For instance, passing <i>reserve</i> = 32 would cause the index range [0 .. 31] to be kept for reserved keys. When non-zero, <i>reserve</i> is rounded to the next multiple of <code>BITS_PER_LONG</code> . If <i>reserve</i> is zero no reservation will be available from the map.
<i>offset</i>	The lowest key value <code>xnmap_enter()</code> will return to the caller. Key values will be in the range [0 + offset .. <i>nkeys</i> + offset - 1]. Negative offsets are valid.

Returns

the address of the new map is returned on success; otherwise, NULL is returned if *nkeys* is invalid.

Remarks

Tags: none.

5.21.2.2 void xnmap_delete (struct xnmap * map)

Delete a map.

Deletes a map, freeing any associated memory back to the Xenomai system heap.

Parameters

<i>map</i>	The address of the map to delete.
------------	-----------------------------------

Remarks

Tags: none.

5.21.2.3 int xnmap_enter (struct xnmap * map, int key, void * objaddr)

Index an object into a map.

Insert a new object into the given map.

Parameters

<i>map</i>	The address of the map to insert into.
<i>key</i>	The key to index the object on. If this key is within the valid index range [0 - offset .. <i>nkeys</i> - offset - 1], then an attempt to reserve this exact key is made. If <i>key</i> has an out-of-range value lower or equal to 0 - offset - 1, then an attempt is made to draw a free key from the unreserved index space.

<i>objaddr</i>	The address of the object to index on the key. This value will be returned by a successful call to xnmap_fetch() with the same key.
----------------	---

Returns

a valid key is returned on success, either *key* if reserved, or the next free key. Otherwise:

- -EEXIST is returned upon attempt to reserve a busy key.
- -ENOSPC when no more free key is available.

Remarks

Tags: isr-allowed.

5.21.2.4 void xnmap_fetch (struct xnmap * map, int key) [inline], [static]

Search an object into a map.

Retrieve an object reference from the given map by its index key.

Parameters

<i>map</i>	The address of the map to retrieve from.
<i>key</i>	The key to be searched for in the map index.

Returns

The indexed object address is returned on success, otherwise NULL is returned when *key* is invalid or no object is currently indexed on it.

Remarks

Tags: isr-allowed.

5.21.2.5 void xnmap_fetch_nocheck (struct xnmap * map, int key) [inline], [static]

Search an object into a map - unchecked form.

Retrieve an object reference from the given map by its index key, but does not perform any sanity check on the provided key.

Parameters

<i>map</i>	The address of the map to retrieve from.
<i>key</i>	The key to be searched for in the map index.

Returns

The indexed object address is returned on success, otherwise NULL is returned when no object is currently indexed on *key*.

Remarks

Tags: isr-allowed.

5.21.2.6 `int xnmap_remove (struct xnmap * map, int key)`

Remove an object reference from a map.

Removes an object reference from the given map, releasing the associated key.

Parameters

<i>map</i>	The address of the map to remove from.
<i>key</i>	The key the object reference to be removed is indexed on.

Returns

0 is returned on success. Otherwise:

- -ESRCH is returned if *key* is invalid.

Remarks

Tags: isr-allowed.

5.22 Registry services

The registry provides a mean to index object descriptors on unique alphanumeric keys.

Collaboration diagram for Registry services:



Functions

- int [xnregistry_enter](#) (const char *key, void *objaddr, xnhandle_t *phandle, struct xnnode *pnode)
Register a real-time object.
- int [xnregistry_bind](#) (const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)
Bind to a real-time object.
- int [xnregistry_remove](#) (xnhandle_t handle)
Forcibly unregister a real-time object.
- int [xnregistry_unlink](#) (const char *key)
Turn a named object into an anonymous object.
- static void * [xnregistry_lookup](#) (xnhandle_t handle, unsigned long *cstamp_r)
Find a real-time object into the registry.

5.22.1 Detailed Description

The registry provides a mean to index object descriptors on unique alphanumeric keys. When labeled this way, an object is globally exported; it can be searched for, and its descriptor returned to the caller for further use; the latter operation is called a "binding". When no object has been registered under the given name yet, the registry can be asked to set up a rendez-vous, blocking the caller until the object is eventually registered.

5.22.2 Function Documentation

5.22.2.1 int xnregistry_bind (const char * key, xnticks_t timeout, int timeout_mode, xnhandle_t * phandle)

Bind to a real-time object.

This service retrieves the registry handle of a given object identified by its key. Unless otherwise specified, this service will block the caller if the object is not registered yet, waiting for such registration to occur.

Parameters

<i>key</i>	A valid NULL-terminated string which identifies the object to bind to.
<i>timeout</i>	The timeout which may be used to limit the time the thread wait for the object to be registered. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic (XN_ABSOLUTE), or absolute adjustable (XN_REALTIME) depending on <i>timeout_mode</i> . Passing XN_INFINITE and setting <i>timeout_mode</i> to XN_RELATIVE specifies an unbounded wait. Passing XN_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry. All other values are used as a wait limit.
<i>timeout_mode</i>	The mode of the <i>timeout</i> parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also xntimer_start()).
<i>phandle</i>	A pointer to a memory location which will be written upon success with the generic handle defined by the registry for the retrieved object. Contents of this memory is undefined upon failure.

Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *key* is NULL.
- -EINTR is returned if [xnthread_unblock\(\)](#) has been called for the waiting thread before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to XN_NONBLOCK and the searched object is not registered on entry. As a special exception, this error is also returned if this service should block, but was called from a context which cannot sleep (e.g. interrupt, non-realtime or scheduler locked).
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.

Remarks

Tags: primary-only, might-switch.

References XNBREAK, xnsynch_sleep_on(), and XNTIMEO.

```
5.22.2.2 int xnregistry_enter ( const char * key, void * objaddr, xnhandle_t * phandle, struct xninode
    * pnode )
```

Register a real-time object.

This service allocates a new registry slot for an associated object, and indexes it by an alphanumeric key for later retrieval.

Parameters

<i>key</i>	A valid NULL-terminated string by which the object will be indexed and later retrieved in the registry. Since it is assumed that such key is stored into the registered object, it will <i>not</i> be copied but only kept by reference in the registry. Pass an empty or NULL string if the object shall only occupy a registry slot for handle-based lookups.
<i>objaddr</i>	An opaque pointer to the object to index by <i>key</i> .
<i>phandle</i>	A pointer to a generic handle defined by the registry which will uniquely identify the indexed object, until the latter is unregistered using the xnregistry_remove() service.

<i>pnode</i>	A pointer to an optional /proc node class descriptor. This structure provides the information needed to export all objects from the given class through the /proc filesystem, under the /proc/xenomai/registry entry. Passing NULL indicates that no /proc support is available for the newly registered object.
--------------	--

Returns

0 is returned upon success. Otherwise:

- -EINVAL is returned if *objaddr* is NULL, or if *key* is non-NULL and contains an invalid '/' character.
- -ENOMEM is returned if the system fails to get enough dynamic memory from the global real-time heap in order to register the object.
- -EEXIST is returned if the *key* is already in use.

Remarks

Tags: might-switch.

References `xnsched_run()`.

Referenced by `rtdm_dev_register()`.

```
5.22.2.3 void * xnregistry_lookup ( xnhandle_t handle, unsigned long * cstamp_r ) [inline],
[static]
```

Find a real-time object into the registry.

This service retrieves an object from its handle into the registry and returns the memory address of its descriptor. Optionally, it also copies back the object's creation stamp which is unique across object registration calls.

Parameters

<i>handle</i>	The generic handle of the object to fetch.
<i>cstamp_r</i>	If not-NULL, the object's creation stamp will be copied to this memory area.

Returns

The memory address of the object's descriptor is returned on success. Otherwise, NULL is returned if *handle* does not reference a registered object.

Remarks

Tags: isr-allowed.

```
5.22.2.4 int xnregistry_remove ( xnhandle_t handle )
```

Forcibly unregister a real-time object.

This service forcibly removes an object from the registry. The removal is performed regardless of the current object's locking status.

Parameters

<i>handle</i>	The generic handle of the object to remove.
---------------	---

Returns

0 is returned upon success. Otherwise:

- -ESRCH is returned if *handle* does not reference a registered object.

Remarks

Tags: none.

Referenced by `rtdm_dev_unregister()`.

5.22.2.5 `int xnregistry_unlink (const char * key)`

Turn a named object into an anonymous object.

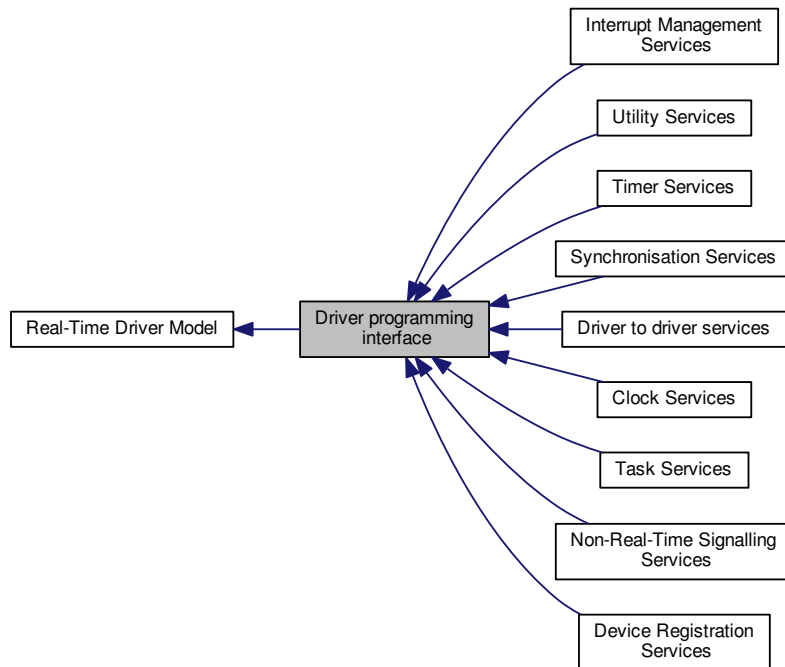
Remarks

Tags: none.

5.23 Driver programming interface

RTDM driver programming interface.

Collaboration diagram for Driver programming interface:



Modules

- [Driver to driver services](#)
Inter-driver interface.
- [Device Registration Services](#)
- [Clock Services](#)
- [Task Services](#)
- [Timer Services](#)
- [Synchronisation Services](#)
- [Interrupt Management Services](#)
- [Non-Real-Time Signalling Services](#)

These services provide a mechanism to request the execution of a specified handler in non-real-time context.

- [Utility Services](#)


5.23.1 Detailed Description

RTDM driver programming interface.

5.24 Wait queue

RTDM wait queues.

Collaboration diagram for Wait queue:



Functions

- void [rtdm_waitqueue_init](#) (struct rtdm_waitqueue *wq)
Initialize a RTDM wait queue.
- void [rtdm_waitqueue_destroy](#) (struct rtdm_waitqueue *wq)
Deletes a RTDM wait queue.
- [rtdm_timedwait_condition_locked](#) (struct rtdm_wait_queue *wq, C_expr condition, [nanosecs_rel_t](#) timeout, rtdm_toseq_t *toseq)
Timed sleep on a locked waitqueue until a condition gets true.
- [rtdm_wait_condition_locked](#) (struct rtdm_wait_queue *wq, C_expr condition)
Sleep on a locked waitqueue until a condition gets true.
- [rtdm_timedwait_condition](#) (struct rtdm_wait_queue *wq, C_expr condition, [nanosecs_rel_t](#) timeout, rtdm_toseq_t *toseq)
Timed sleep on a waitqueue until a condition gets true.
- void [rtdm_timedwait](#) (struct rtdm_wait_queue *wq, [nanosecs_rel_t](#) timeout, rtdm_toseq_t *toseq)
Timed sleep on a waitqueue unconditionally.
- void [rtdm_timedwait_locked](#) (struct rtdm_wait_queue *wq, [nanosecs_rel_t](#) timeout, rtdm_toseq_t *toseq)
Timed sleep on a locked waitqueue unconditionally.
- [rtdm_wait_condition](#) (struct rtdm_wait_queue *wq, C_expr condition)
Sleep on a waitqueue until a condition gets true.
- void [rtdm_wait](#) (struct rtdm_wait_queue *wq)
Sleep on a waitqueue unconditionally.
- void [rtdm_wait_locked](#) (struct rtdm_wait_queue *wq)
Sleep on a locked waitqueue unconditionally.
- void [rtdm_waitqueue_lock](#) (struct rtdm_wait_queue *wq, [rtdm_lockctx_t](#) context)
Lock a waitqueue.
- void [rtdm_waitqueue_unlock](#) (struct rtdm_wait_queue *wq, [rtdm_lockctx_t](#) context)
Unlock a waitqueue.
- void [rtdm_waitqueue_signal](#) (struct rtdm_wait_queue *wq)
Signal a waitqueue.
- void [rtdm_waitqueue_broadcast](#) (struct rtdm_wait_queue *wq)
Broadcast a waitqueue.
- void [rtdm_waitqueue_flush](#) (struct rtdm_wait_queue *wq)
Flush a waitqueue.
- void [rtdm_waitqueue_wakeup](#) (struct rtdm_wait_queue *wq, rtdm_task_t waiter)

Signal a particular waiter on a waitqueue.

- [rtdm_for_each_waiter](#) (rtdm_task_t pos, struct rtdm_wait_queue *wq)

Simple iterator for waitqueues.

- [rtdm_for_each_waiter_safe](#) (rtdm_task_t pos, rtdm_task_t tmp, struct rtdm_wait_queue *wq)

Safe iterator for waitqueues.

5.24.1 Detailed Description

RTDM wait queues.

5.24.2 Function Documentation

5.24.2.1 [rtdm_for_each_waiter](#) (rtdm_task_t pos, struct rtdm_wait_queue * wq)

Simple iterator for waitqueues.

This construct traverses the wait list of a given waitqueue *wq*, assigning each RTDM task pointer to the cursor variable *pos*, which must be of type `rtdm_task_t`.

wq must have been locked by a call to [rtdm_waitqueue_lock\(\)](#) prior to traversing its wait list.

Parameters

<i>pos</i>	cursor variable holding a pointer to the RTDM task being fetched.
<i>wq</i>	waitqueue to scan.

Note

The waitqueue should not be signaled, broadcast or flushed during the traversal, unless the loop is aborted immediately after. Should multiple waiters be readied while iterating, the safe form [rtdm_for_each_waiter_safe\(\)](#) must be used for traversal instead.

5.24.2.2 [rtdm_for_each_waiter_safe](#) (rtdm_task_t pos, rtdm_task_t tmp, struct rtdm_wait_queue * wq)

Safe iterator for waitqueues.

This construct traverses the wait list of a given waitqueue *wq*, assigning each RTDM task pointer to the cursor variable *pos*, which must be of type `rtdm_task_t`.

Unlike with [rtdm_for_each_waiter\(\)](#), the waitqueue may be signaled, broadcast or flushed during the traversal.

wq must have been locked by a call to [rtdm_waitqueue_lock\(\)](#) prior to traversing its wait list.

Parameters

<i>pos</i>	cursor variable holding a pointer to the RTDM task being fetched.
<i>tmp</i>	temporary cursor variable.
<i>wq</i>	waitqueue to scan.

5.24.2.3 `void rtdm_timedwait (struct rtdm_wait_queue * wq, nanosecs_rel_t timeout, rtdm_toseq_t * toseq)`

Timed sleep on a waitqueue unconditionally.

The calling task is put to sleep until the waitqueue is signaled by either `rtdm_waitqueue_signal()` or `rtdm_waitqueue_broadcast()`, or flushed by a call to `rtdm_waitqueue_flush()`, or a timeout occurs.

Parameters

	<i>wq</i>	waitqueue to wait on.
	<i>timeout</i>	relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values.
in,out	<i>toseq</i>	handle of a timeout sequence as returned by rtdm_toseq_init() or NULL.

Returns

0 on success, otherwise:

- -EINTR is returned if the waitqueue has been flushed, or the calling task has received a Linux signal or has been forcibly unblocked by a call to [rtdm_task_unblock\(\)](#).
- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

Note

Passing `RTDM_TIMEOUT_NONE` to *timeout* makes no sense for such service, and might cause unexpected behavior.

5.24.2.4 `rtdm_timedwait_condition (struct rtdm_wait_queue * wq, C_expr condition, nanosecs_rel_t timeout, rtdm_toseq_t * toseq)`

Timed sleep on a waitqueue until a condition gets true.

The calling task is put to sleep until *condition* evaluates to true or a timeout occurs. The condition is checked each time the waitqueue *wq* is signaled.

Parameters

	<i>wq</i>	waitqueue to wait on.
	<i>condition</i>	C expression for the event to wait for.
	<i>timeout</i>	relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values.
in,out	<i>toseq</i>	handle of a timeout sequence as returned by rtdm_toseq_init() or NULL.

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has received a Linux signal or has been forcibly unblocked by a call to [rtdm_task_unblock\(\)](#).
- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

Note

[rtdm_waitqueue_signal\(\)](#) has to be called after changing any variable that could change the result of the wait condition.

Passing `RTDM_TIMEOUT_NONE` to *timeout* makes no sense for such service, and might cause unexpected behavior.

5.24.2.5 `rtm_timedwait_condition_locked (struct rtdm_wait_queue * wq, C_expr condition, nanosecs_rel_t timeout, rtdm_toseq_t * toseq)`

Timed sleep on a locked waitqueue until a condition gets true.

The calling task is put to sleep until *condition* evaluates to true or a timeout occurs. The condition is checked each time the waitqueue *wq* is signaled.

The waitqueue must have been locked by a call to [rtm_waitqueue_lock\(\)](#) prior to calling this service.

Parameters

	<i>wq</i>	locked waitqueue to wait on. The waitqueue lock is dropped when sleeping, then reacquired before this service returns to the caller.
	<i>condition</i>	C expression for the event to wait for.
	<i>timeout</i>	relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values.
in,out	<i>toseq</i>	handle of a timeout sequence as returned by rtm_toseq_init() or NULL.

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has received a Linux signal or has been forcibly unblocked by a call to [rtm_task_unblock\(\)](#).
- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

Note

[rtm_waitqueue_signal\(\)](#) has to be called after changing any variable that could change the result of the wait condition.

Passing [RTDM_TIMEOUT_NONE](#) to *timeout* makes no sense for such service, and might cause unexpected behavior.

5.24.2.6 `void rtdm_timedwait_locked (struct rtdm_wait_queue * wq, nanosecs_rel_t timeout, rtdm_toseq_t * toseq)`

Timed sleep on a locked waitqueue unconditionally.

The calling task is put to sleep until the waitqueue is signaled by either [rtm_waitqueue_signal\(\)](#) or [rtm_waitqueue_broadcast\(\)](#), or flushed by a call to [rtm_waitqueue_flush\(\)](#), or a timeout occurs.

The waitqueue must have been locked by a call to [rtm_waitqueue_lock\(\)](#) prior to calling this service.

Parameters

	<i>wq</i>	locked waitqueue to wait on. The waitqueue lock is dropped when sleeping, then reacquired before this service returns to the caller.
	<i>timeout</i>	relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values.
in,out	<i>toseq</i>	handle of a timeout sequence as returned by rtm_toseq_init() or NULL.

Returns

0 on success, otherwise:

- -EINTR is returned if the waitqueue has been flushed, or the calling task has received a Linux signal or has been forcibly unblocked by a call to [rtm_task_unblock\(\)](#).
- -ETIMEDOUT is returned if the if the request has not been satisfied within the specified amount of time.

Note

Passing RTDM_TIMEOUT_NONE to *timeout* makes no sense for such service, and might cause unexpected behavior.

5.24.2.7 void rtdm_wait (struct rtdm_wait_queue * wq)

Sleep on a waitqueue unconditionally.

The calling task is put to sleep until the waitqueue is signaled by either [rtm_waitqueue_signal\(\)](#) or [rtm_waitqueue_broadcast\(\)](#), or flushed by a call to [rtm_waitqueue_flush\(\)](#).

Parameters

<i>wq</i>	waitqueue to wait on.
-----------	-----------------------

Returns

0 on success, otherwise:

- -EINTR is returned if the waitqueue has been flushed, or the calling task has received a Linux signal or has been forcibly unblocked by a call to [rtm_task_unblock\(\)](#).

5.24.2.8 rtdm_wait_condition (struct rtdm_wait_queue * wq, C_expr condition)

Sleep on a waitqueue until a condition gets true.

The calling task is put to sleep until *condition* evaluates to true. The condition is checked each time the waitqueue *wq* is signaled.

Parameters

<i>wq</i>	waitqueue to wait on
<i>condition</i>	C expression for the event to wait for.

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has received a Linux signal or has been forcibly unblocked by a call to [rtm_task_unblock\(\)](#).

Note

[rtm_waitqueue_signal\(\)](#) has to be called after changing any variable that could change the result of the wait condition.

5.24.2.9 `rtdm_wait_condition_locked (struct rtdm_wait_queue * wq, C_expr condition)`

Sleep on a locked waitqueue until a condition gets true.

The calling task is put to sleep until *condition* evaluates to true. The condition is checked each time the waitqueue *wq* is signaled.

The waitqueue must have been locked by a call to [rtdm_waitqueue_lock\(\)](#) prior to calling this service.

Parameters

<i>wq</i>	locked waitqueue to wait on. The waitqueue lock is dropped when sleeping, then reacquired before this service returns to the caller.
<i>condition</i>	C expression for the event to wait for.

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has received a Linux signal or has been forcibly unblocked by a call to [rtdm_task_unblock\(\)](#).

Note

[rtdm_waitqueue_signal\(\)](#) has to be called after changing any variable that could change the result of the wait condition.

5.24.2.10 `void rtdm_wait_locked (struct rtdm_wait_queue * wq)`

Sleep on a locked waitqueue unconditionally.

The calling task is put to sleep until the waitqueue is signaled by either [rtdm_waitqueue_signal\(\)](#) or [rtdm_waitqueue_broadcast\(\)](#), or flushed by a call to [rtdm_waitqueue_flush\(\)](#).

The waitqueue must have been locked by a call to [rtdm_waitqueue_lock\(\)](#) prior to calling this service.

Parameters

<i>wq</i>	locked waitqueue to wait on. The waitqueue lock is dropped when sleeping, then reacquired before this service returns to the caller.
-----------	--

Returns

0 on success, otherwise:

- -EINTR is returned if the waitqueue has been flushed, or the calling task has received a Linux signal or has been forcibly unblocked by a call to [rtdm_task_unblock\(\)](#).

5.24.2.11 `void rtdm_waitqueue_broadcast (struct rtdm_wait_queue * wq)`

Broadcast a waitqueue.

Broadcast the waitqueue *wq*, waking up all waiters. Each readied task may assume to have received the wake up event.

Parameters

<i>wq</i>	waitqueue to broadcast.
-----------	-------------------------

Returns

non-zero if at least one task has been readied as a result of this call, zero otherwise.

5.24.2.12 `void rtdm_waitqueue_destroy (struct rtdm_waitqueue * wq)`

Deletes a RTDM wait queue.

Dismantles a wait queue structure, releasing all resources attached to it.

Parameters

<i>wq</i>	waitqueue to delete.
-----------	----------------------

5.24.2.13 `void rtdm_waitqueue_flush (struct rtdm_wait_queue * wq)`

Flush a waitqueue.

Flushes the waitqueue *wq*, unblocking all waiters with an error status (-EINTR).

Parameters

<i>wq</i>	waitqueue to flush.
-----------	---------------------

Returns

non-zero if at least one task has been readied as a result of this call, zero otherwise.

5.24.2.14 `void rtdm_waitqueue_init (struct rtdm_waitqueue * wq)`

Initialize a RTDM wait queue.

Sets up a wait queue structure for further use.

Parameters

<i>wq</i>	waitqueue to initialize.
-----------	--------------------------

5.24.2.15 `void rtdm_waitqueue_lock (struct rtdm_wait_queue * wq, rtdm_lockctx_t context)`

Lock a waitqueue.

Acquires the lock on the waitqueue *wq*.

Parameters

<i>wq</i>	waitqueue to lock.
<i>context</i>	name of local variable to store the context in.

Note

Recursive locking might lead to unexpected behavior, including lock up.

5.24.2.16 void rtdm_waitqueue_signal (struct rtdm_wait_queue * wq)

Signal a waitqueue.

Signals the waitqueue *wq*, waking up a single waiter (if any).

Parameters

<i>wq</i>	waitqueue to signal.
-----------	----------------------

Returns

non-zero if a task has been readied as a result of this call, zero otherwise.

5.24.2.17 void rtdm_waitqueue_unlock (struct rtdm_wait_queue * wq, **rtdm_lockctx_t** context)

Unlock a waitqueue.

Releases the lock on the waitqueue *wq*.

Parameters

<i>wq</i>	waitqueue to unlock.
<i>context</i>	name of local variable to retrieve the context from.

5.24.2.18 void rtdm_waitqueue_wakeup (struct rtdm_wait_queue * wq, rtdm_task_t waiter)

Signal a particular waiter on a waitqueue.

Signals the waitqueue *wq*, waking up waiter *waiter* only, which must be currently sleeping on the waitqueue.


Parameters

<i>wq</i>	waitqueue to signal.
<i>waiter</i>	RTDM task to wake up.

5.25 Driver to driver services

Inter-driver interface.

Collaboration diagram for Driver to driver services:



Functions

- `int rtdm_open` (const char *path, int oflag,...)
Open a device.
- `int rtdm_socket` (int protocol_family, int socket_type, int protocol)
Create a socket.
- `int rtdm_close` (int fd)
Close a device or socket.
- `int rtdm_ioctl` (int fd, int request,...)
Issue an IOCTL.
- `ssize_t rtdm_read` (int fd, void *buf, size_t nbyte)
Read from device.
- `ssize_t rtdm_write` (int fd, const void *buf, size_t nbyte)
Write to device.
- `ssize_t rtdm_recvmmsg` (int fd, struct msghdr *msg, int flags)
Receive message from socket.
- `ssize_t rtdm_recvfrom` (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)
Receive message from socket.
- `ssize_t rtdm_recv` (int fd, void *buf, size_t len, int flags)
Receive message from socket.
- `ssize_t rtdm_sendmsg` (int fd, const struct msghdr *msg, int flags)
Transmit message to socket.
- `ssize_t rtdm_sendto` (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)
Transmit message to socket.
- `ssize_t rtdm_send` (int fd, const void *buf, size_t len, int flags)
Transmit message to socket.
- `int rtdm_bind` (int fd, const struct sockaddr *my_addr, socklen_t addrlen)
Bind to local address.
- `int rtdm_connect` (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)
Connect to remote address.
- `int rtdm_listen` (int fd, int backlog)
Listen for incoming connection requests.
- `int rtdm_accept` (int fd, struct sockaddr *addr, socklen_t *addrlen)

Accept a connection request.

- int [rt dm_shutdown](#) (int fd, int how)

Shut down parts of a connection.

- int [rt dm_getsockopt](#) (int fd, int level, int optname, void *optval, socklen_t *optlen)

Get socket option.

- int [rt dm_setsockopt](#) (int fd, int level, int optname, const void *optval, socklen_t optlen)

Set socket option.

- int [rt dm_getsockname](#) (int fd, struct sockaddr *name, socklen_t *namelen)

Get local socket address.

- int [rt dm_getpeername](#) (int fd, struct sockaddr *name, socklen_t *namelen)

Get socket destination address.

5.25.1 Detailed Description

Inter-driver interface.

5.25.2 Function Documentation

5.25.2.1 int [rt dm_accept](#) (int fd, struct sockaddr * addr, socklen_t * addrlen)

Accept a connection request.

Refer to [rt_dev_accept\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.2 int [rt dm_bind](#) (int fd, const struct sockaddr * my_addr, socklen_t addrlen)

Bind to local address.

Refer to [rt_dev_bind\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.3 int [rt dm_close](#) (int fd)

Close a device or socket.

Refer to [rt_dev_close\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.4 int [rt dm_connect](#) (int fd, const struct sockaddr * serv_addr, socklen_t addrlen)

Connect to remote address.

Refer to [rt_dev_connect\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.5 `int rtdm_getpeername (int fd, struct sockaddr * name, socklen_t * namelen)`

Get socket destination address.

Refer to [rt_dev_getpeername\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.6 `int rtdm_getsockname (int fd, struct sockaddr * name, socklen_t * namelen)`

Get local socket address.

Refer to [rt_dev_getsockname\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.7 `int rtdm_getsockopt (int fd, int level, int optname, void * optval, socklen_t * optlen)`

Get socket option.

Refer to [rt_dev_getsockopt\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.8 `int rtdm_ioctl (int fd, int request, ...)`

Issue an IOCTL.

Refer to [rt_dev_ioctl\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.9 `int rtdm_listen (int fd, int backlog)`

Listen for incoming connection requests.

Refer to [rt_dev_listen\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.10 `int rtdm_open (const char * path, int oflag, ...)`

Open a device.

Refer to [rt_dev_open\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.11 `ssize_t rtdm_read (int fd, void * buf, size_t nbyte)`

Read from device.

Refer to [rt_dev_read\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.12 `ssize_t rtdm_recv (int fd, void * buf, size_t len, int flags)`

Receive message from socket.

Refer to [rt_dev_recv\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.13 `ssize_t rtdm_recvfrom (int fd, void * buf, size_t len, int flags, struct sockaddr * from, socklen_t * fromlen)`

Receive message from socket.

Refer to [rt_dev_recvfrom\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.14 `ssize_t rtdm_recvmsg (int fd, struct msghdr * msg, int flags)`

Receive message from socket.

Refer to [rt_dev_recvmsg\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.15 `ssize_t rtdm_send (int fd, const void * buf, size_t len, int flags)`

Transmit message to socket.

Refer to [rt_dev_send\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.16 `ssize_t rtdm_sendmsg (int fd, const struct msghdr * msg, int flags)`

Transmit message to socket.

Refer to [rt_dev_sendmsg\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.17 `ssize_t rtdm_sendto (int fd, const void * buf, size_t len, int flags, const struct sockaddr * to, socklen_t tolen)`

Transmit message to socket.

Refer to [rt_dev_sendto\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.18 `int rtdm_setsockopt (int fd, int level, int optname, const void * optval, socklen_t optlen)`

Set socket option.

Refer to [rt_dev_setsockopt\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.19 `int rtdm_shutdown (int fd, int how)`

Shut down parts of a connection.

Refer to [rt_dev_shutdown\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.20 `int rtdm_socket (int protocol_family, int socket_type, int protocol)`

Create a socket.

Refer to [rt_dev_socket\(\)](#) for parameters and return values

Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.25.2.21 `ssize_t rtdm_write (int fd, const void * buf, size_t nbyte)`

Write to device.

Refer to [rt_dev_write\(\)](#) for parameters and return values


Environments:

Depends on driver implementation, see [Device Profiles](#).

Rescheduling: possible.

5.26 Device Registration Services

Collaboration diagram for Device Registration Services:



Functions

- int `rtdm_dev_register` (struct `rtdm_device` *device)
Register a RTDM device.
- int `rtdm_dev_unregister` (struct `rtdm_device` *device, unsigned int poll_delay)
Unregisters a RTDM device.

Device Flags

Static flags describing a RTDM device

- #define `RTDM_EXCLUSIVE` 0x0001
If set, only a single instance of the device can be requested by an application.
- #define `RTDM_NAMED_DEVICE` 0x0010
If set, the device is addressed via a clear-text name.
- #define `RTDM_PROTOCOL_DEVICE` 0x0020
If set, the device is addressed via a combination of protocol ID and socket type.
- #define `RTDM_DEVICE_TYPE_MASK` 0x00F0
Mask selecting the device type.

Driver Versioning

Current revisions of RTDM structures, encoding of driver versions. See [API Versioning](#) for the interface revision.

- #define `RTDM_DEVICE_STRUCT_VER` 6
Version of struct `rtdm_device`.
- #define `RTDM_CONTEXT_STRUCT_VER` 4
Version of struct `rtdm_dev_context`.
- #define `RTDM_SECURE_DEVICE` 0x80000000
Flag indicating a secure variant of RTDM (not supported here)
- #define `RTDM_DRIVER_VER`(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))
Version code constructor for driver revisions.
- #define `RTDM_DRIVER_MAJOR_VER`(ver) (((ver) >> 16) & 0xFF)
Get major version number from driver revision code.

- `#define RTDM_DRIVER_MINOR_VER(ver) (((ver) >> 8) & 0xFF)`
Get minor version number from driver revision code.
- `#define RTDM_DRIVER_PATCH_VER(ver) ((ver) & 0xFF)`
Get patch version number from driver revision code.

5.26.1 Detailed Description

5.26.2 Macro Definition Documentation

5.26.2.1 `#define RTDM_DEVICE_TYPE_MASK 0x00F0`

Mask selecting the device type.

Referenced by `rtdm_dev_register()`, and `rtdm_dev_unregister()`.

5.26.2.2 `#define RTDM_EXCLUSIVE 0x0001`

If set, only a single instance of the device can be requested by an application.

Referenced by `rtdm_dev_register()`.

5.26.2.3 `#define RTDM_NAMED_DEVICE 0x0010`

If set, the device is addressed via a clear-text name.

Referenced by `rtdm_dev_register()`, and `rtdm_dev_unregister()`.

5.26.2.4 `#define RTDM_PROTOCOL_DEVICE 0x0020`

If set, the device is addressed via a combination of protocol ID and socket type.

Referenced by `rtdm_dev_register()`.

5.26.3 Function Documentation

5.26.3.1 `int rtdm_dev_register (struct rtdm_device * device)`

Register a RTDM device.

Parameters

<code>in</code>	<code>device</code>	Pointer to structure describing the new device.
-----------------	---------------------	---

Returns

0 is returned upon success. Otherwise:

- `-EINVAL` is returned if the device structure contains invalid entries. Check kernel log in this case.
- `-ENOMEM` is returned if the context for an exclusive device cannot be allocated.
- `-EEXIST` is returned if the specified device name or protocol ID is already in use.
- `-EAGAIN` is returned if some `/proc` entry cannot be created.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

Rescheduling: never.

References `rtm_device::context_size`, `rtm_device::device_flags`, `rtm_device::device_name`, `rtm_device::open`, `rtm_device::ops`, `rtm_device::proc_name`, `rtm_device::protocol_family`, `rtm_device::reserved`, `RTDM_DEVICE_STRUCT_VER`, `RTDM_DEVICE_TYPE_MASK`, `RTDM_EXCLUSIVE`, `RTDM_NAMED_DEVICE`, `RTDM_PROTOCOL_DEVICE`, `rtm_device::socket`, `rtm_device::socket_type`, `rtm_device::struct_version`, and `xnregistry_enter()`.

5.26.3.2 `int rtdm_dev_unregister (struct rtdm_device * device, unsigned int poll_delay)`

Unregisters a RTDM device.

Parameters

<code>in</code>	<code>device</code>	Pointer to structure describing the device to be unregistered.
<code>in</code>	<code>poll_delay</code>	Polling delay in milliseconds to check repeatedly for open instances of <code>device</code> , or 0 for non-blocking mode.

Returns

0 is returned upon success. Otherwise:

- `-ENODEV` is returned if the device was not registered.
- `-EAGAIN` is returned if the device is busy with open instances and 0 has been passed for `poll_delay`.

Environments:

This service can be called from:


- Kernel module initialization/cleanup code

Rescheduling: never.

References `rtm_device::device_flags`, `rtm_device::device_name`, `rtm_device::protocol_family`, `rtm_device::reserved`, `RTDM_DEVICE_TYPE_MASK`, `RTDM_NAMED_DEVICE`, `rtm_device::socket_type`, and `xnregistry_remove()`.

5.27 Clock Services

Collaboration diagram for Clock Services:



Functions

- [nanosecs_abs_t rtdm_clock_read](#) (void)
Get system time.
- [nanosecs_abs_t rtdm_clock_read_monotonic](#) (void)
Get monotonic time.

5.27.1 Detailed Description

5.27.2 Function Documentation

5.27.2.1 **nanosecs_abs_t** rtdm_clock_read (void)

Get system time.

Returns

The system time in nanoseconds is returned

Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period. The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Referenced by `a4l_get_time()`, and `rtdm_ratelimit()`.

5.27.2.2 `nanosecs_abs_t rtdm_clock_read_monotonic (void)`

Get monotonic time.

Returns

The monotonic time in nanoseconds is returned

Note

The resolution of this service depends on the system timer. In particular, if the system timer is running in periodic mode, the return value will be limited to multiples of the timer tick period.

The system timer may have to be started to obtain valid results. Whether this happens automatically (as on Xenomai) or is controlled by the application depends on the RTDM host environment.

Environments:


This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.28 Task Services

Collaboration diagram for Task Services:



Typedefs

- typedef void(* [rtdm_task_proc_t](#))(void *arg)
Real-time task procedure.

Functions

- int [rtdm_task_init](#) (rtdm_task_t *task, const char *name, [rtdm_task_proc_t](#) task_proc, void *arg, int priority, [nanosecs_rel_t](#) period)
Initialise and start a real-time task.
- void [rtdm_task_destroy](#) (rtdm_task_t *task)
Destroy a real-time task.
- int [rtdm_task_should_stop](#) (void)
Check for pending termination request.
- void [rtdm_task_set_priority](#) (rtdm_task_t *task, int priority)
Adjust real-time task priority.
- int [rtdm_task_set_period](#) (rtdm_task_t *task, [nanosecs_rel_t](#) period)
Adjust real-time task period.
- int [rtdm_task_wait_period](#) (void)
Wait on next real-time task period.
- int [rtdm_task_unblock](#) (rtdm_task_t *task)
Activate a blocked real-time task.
- rtdm_task_t * [rtdm_task_current](#) (void)
Get current real-time task.
- int [rtdm_task_sleep](#) ([nanosecs_rel_t](#) delay)
Sleep a specified amount of time.
- int [rtdm_task_sleep_until](#) ([nanosecs_abs_t](#) wakeup_time)
Sleep until a specified absolute time.
- int [rtdm_task_sleep_abs](#) ([nanosecs_abs_t](#) wakeup_time, enum [rtdm_timer_mode](#) mode)
Sleep until a specified absolute time.
- void [rtdm_task_join_nrt](#) (rtdm_task_t *task, unsigned int poll_delay)
Wait on a real-time task to terminate.
- void [rtdm_task_busy_sleep](#) ([nanosecs_rel_t](#) delay)
Busy-wait a specified amount of time.

Task Priority Range

Maximum and minimum task priorities

- `#define RTDM_TASK_LOWEST_PRIORITY 0`
- `#define RTDM_TASK_HIGHEST_PRIORITY 99`

Task Priority Modification

Raise or lower task priorities by one level

- `#define RTDM_TASK_RAISE_PRIORITY (+1)`
- `#define RTDM_TASK_LOWER_PRIORITY (-1)`

5.28.1 Detailed Description

5.28.2 Typedef Documentation

5.28.2.1 `typedef void(* rtdm_task_proc_t)(void *arg)`

Real-time task procedure.

Parameters

in, out	<i>arg</i>	argument as passed to rtdm_task_init()
----------------	------------	--

5.28.3 Function Documentation

5.28.3.1 `void rtdm_task_busy_sleep (nanosecs_rel_t delay)`

Busy-wait a specified amount of time.

Parameters

in	<i>delay</i>	Delay in nanoseconds. Note that a zero delay does not have the meaning of <code>RTDM_TIMEOUT_INFINITE</code> here.
-----------	--------------	---

Note

The caller must not be migratable to different CPUs while executing this service. Otherwise, the actual delay will be undefined.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (should be avoided or kept short)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never (except due to external interruptions).

5.28.3.2 `rtm_task_t* rtdm_task_current (void)`

Get current real-time task.

Returns

Pointer to task handle

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.28.3.3 `void rtdm_task_destroy (rtdm_task_t * task)`

Destroy a real-time task.

This call sends a termination request to *task*, then waits for it to exit. All RTDM task should check for pending termination requests by calling [rtdm_task_should_stop\(\)](#) from their work loop.

If *task* is current, [rtdm_task_destroy\(\)](#) terminates the current context, and does not return to the caller.

Parameters

in, out	<i>task</i>	Task handle as returned by rtdm_task_init()
---------	-------------	---

Note

Passing the same task handle to RTDM services after the completion of this function is not allowed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task

Rescheduling: yes.

5.28.3.4 `int rtdm_task_init (rtdm_task_t * task, const char * name, rtdm_task_proc_t task_proc, void * arg, int priority, nanosecs_rel_t period)`

Initialise and start a real-time task.

After initialising a task, the task handle remains valid and can be passed to RTDM services until either [rtdm_task_destroy\(\)](#) or [rtdm_task_join_nrt\(\)](#) was invoked.

Parameters

in, out	<i>task</i>	Task handle
---------	-------------	-------------

in	<i>name</i>	Optional task name
in	<i>task_proc</i>	Procedure to be executed by the task
in	<i>arg</i>	Custom argument passed to <code>task_proc()</code> on entry
in	<i>priority</i>	Priority of the task, see also Task Priority Range
in	<i>period</i>	Period in nanoseconds of a cyclic task, 0 for non-cyclic mode

Returns

0 on success, otherwise negative error code

Calling context: This service can be called from secondary mode only.

Rescheduling: possible.

References `xnthread_cancel()`, `xnthread_init()`, `xnthread_set_periodic()`, and `xnthread_start()`.

5.28.3.5 `void rtdm_task_join_nrt (rtdm_task_t * task, unsigned int poll_delay)`

Wait on a real-time task to terminate.

Parameters

in,out	<i>task</i>	Task handle as returned by rtdm_task_init()
in	<i>poll_delay</i>	Delay in milliseconds between periodic tests for the state of the real-time task. This parameter is ignored if the termination is internally realised without polling.

Note

Passing the same task handle to RTDM services after the completion of this function is not allowed. This service does not trigger the termination of the targeted task. The user has to take of this, otherwise [rtdm_task_join_nrt\(\)](#) will never return.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

5.28.3.6 `int rtdm_task_set_period (rtdm_task_t * task, nanosecs_rel_t period)`

Adjust real-time task period.

Parameters

in,out	<i>task</i>	Task handle as returned by rtdm_task_init()
in	<i>period</i>	New period in nanoseconds of a cyclic task, 0 for non-cyclic mode

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task

- User-space task (RT, non-RT)

Rescheduling: possible.

5.28.3.7 void `rtdm_task_set_priority` (`rtdm_task_t * task`, `int priority`)

Adjust real-time task priority.

Parameters

<code>in, out</code>	<code>task</code>	Task handle as returned by rtdm_task_init()
<code>in</code>	<code>priority</code>	New priority of the task, see also Task Priority Range

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.28.3.8 int `rtdm_task_should_stop` (void)

Check for pending termination request.

Check whether a termination request was received by the current RTDM task. Termination requests are sent by calling [rtdm_task_destroy\(\)](#).

Returns

Non-zero indicates that a termination request is pending, in which case the caller should wrap up and exit.

Calling context: kernel task

Rescheduling: none.

5.28.3.9 int `rtdm_task_sleep` (**nanosecs_rel_t** `delay`)

Sleep a specified amount of time.

Parameters

<code>in</code>	<code>delay</code>	Delay in nanoseconds, see RTDM_TIMEOUT_xxx for special values.
-----------------	--------------------	--

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtdm_task_unblock\(\)](#).
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always.

5.28.3.10 `int rtdm_task_sleep_abs (nanosecs_abs_t wakeup_time, enum rtdm_timer_mode mode)`

Sleep until a specified absolute time.

Parameters

<code>in</code>	<code>wakeup_time</code>	Absolute timeout in nanoseconds
<code>in</code>	<code>mode</code>	Selects the timer mode, see <code>RTDM_TIMERMODE_XXX</code> for details

Returns

0 on success, otherwise:

- `-EINTR` is returned if calling task has been unblock by a signal or explicitly via [rtdm_task_unblock\(\)](#).
- `-EPERM` *may* be returned if an illegal invocation environment is detected.
- `-EINVAL` is returned if an invalid parameter was passed.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

5.28.3.11 `int rtdm_task_sleep_until (nanosecs_abs_t wakeup_time)`

Sleep until a specified absolute time.

Deprecated Use `rtdm_task_sleep_abs` instead!

Parameters

<code>in</code>	<code>wakeup_time</code>	Absolute timeout in nanoseconds
-----------------	--------------------------	---------------------------------

Returns

0 on success, otherwise:

- `-EINTR` is returned if calling task has been unblock by a signal or explicitly via [rtdm_task_unblock\(\)](#).
- `-EPERM` *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless the specified time already passed.

5.28.3.12 `int rtdm_task_unblock (rtdm_task_t * task)`

Activate a blocked real-time task.

Returns

Non-zero is returned if the task was actually unblocked from a pending wait state, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.28.3.13 `int rtdm_task_wait_period (void)`

Wait on next real-time task period.

Returns

0 on success, otherwise:

- -EINVAL is returned if calling task is not in periodic mode.
- -ETIMEDOUT is returned if a timer overrun occurred, which indicates that a previous release point has been missed by the calling task.

Environments:


This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: always, unless a timer overrun occurred.

5.29 Timer Services

Collaboration diagram for Timer Services:



Typedefs

- typedef void(* [rtdm_timer_handler_t](#))(rtdm_timer_t *timer)
Timer handler.

Functions

- int [rtdm_timer_init](#) (rtdm_timer_t *timer, [rtdm_timer_handler_t](#) handler, const char *name)
Initialise a timer.
- void [rtdm_timer_destroy](#) (rtdm_timer_t *timer)
Destroy a timer.
- int [rtdm_timer_start](#) (rtdm_timer_t *timer, [nanosecs_abs_t](#) expiry, [nanosecs_rel_t](#) interval, enum [rtdm_timer_mode](#) mode)
Start a timer.
- void [rtdm_timer_stop](#) (rtdm_timer_t *timer)
Stop a timer.
- int [rtdm_timer_start_in_handler](#) (rtdm_timer_t *timer, [nanosecs_abs_t](#) expiry, [nanosecs_rel_t](#) interval, enum [rtdm_timer_mode](#) mode)
Start a timer from inside a timer handler.
- void [rtdm_timer_stop_in_handler](#) (rtdm_timer_t *timer)
Stop a timer from inside a timer handler.

RTDM_TIMERMODE_XXX

Timer operation modes

- enum [rtdm_timer_mode](#) { [RTDM_TIMERMODE_RELATIVE](#) = XN_RELATIVE, [RTDM_TIMERMODE_ABSOLUTE](#) = XN_ABSOLUTE, [RTDM_TIMERMODE_REALTIME](#) = XN_REALTIME }

5.29.1 Detailed Description

5.29.2 Typedef Documentation

5.29.2.1 typedef void(* rtdm_timer_handler_t)(rtdm_timer_t *timer)

Timer handler.

Parameters

in	<i>timer</i>	Timer handle as returned by rt dm_timer_init()
-----------	--------------	--

5.29.3 Enumeration Type Documentation

5.29.3.1 enum **rt dm_timer_mode**

Enumerator

RTDM_TIMERMODE_RELATIVE Monotonic timer with relative timeout.

RTDM_TIMERMODE_ABSOLUTE Monotonic timer with absolute timeout.

RTDM_TIMERMODE_REALTIME Adjustable timer with absolute timeout.

5.29.4 Function Documentation

5.29.4.1 void **rt dm_timer_destroy** (**rt dm_timer_t** * timer)

Destroy a timer.

Parameters

in, out	<i>timer</i>	Timer handle as returned by rt dm_timer_init()
----------------	--------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xntimer_destroy\(\)](#).

5.29.4.2 int **rt dm_timer_init** (**rt dm_timer_t** * timer, **rt dm_timer_handler_t** handler, const char * name)

Initialise a timer.

Parameters

in, out	<i>timer</i>	Timer handle
in	<i>handler</i>	Handler to be called on timer expiry
in	<i>name</i>	Optional timer name

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.29.4.3 `int rtdm_timer_start (rtdm_timer_t * timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)`

Start a timer.

Parameters

in,out	<i>timer</i>	Timer handle as returned by rtdm_timer_init()
in	<i>expiry</i>	Firing time of the timer, <i>mode</i> defines if relative or absolute
in	<i>interval</i>	Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers
in	<i>mode</i>	Defines the operation mode, see RTDM_TIMERMODE_XXX for possible values

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if *expiry* describes an absolute date in the past.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xntimer_start\(\)](#).

5.29.4.4 `int rtdm_timer_start_in_handler (rtdm_timer_t * timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)`

Start a timer from inside a timer handler.

Parameters

in,out	<i>timer</i>	Timer handle as returned by rtdm_timer_init()
in	<i>expiry</i>	Firing time of the timer, <i>mode</i> defines if relative or absolute
in	<i>interval</i>	Relative reload value, > 0 if the timer shall work in periodic mode with the specific interval, 0 for one-shot timers

<i>in</i>	<i>mode</i>	Defines the operation mode, see RTDM_TIMERMODE_XXX for possible values
-----------	-------------	--

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if expiry describes an absolute date in the past.

Environments:

This service can be called from:

- Timer handler

Rescheduling: never.

5.29.4.5 void `rtdm_timer_stop` (`rtdm_timer_t * timer`)

Stop a timer.

Parameters

<i>in,out</i>	<i>timer</i>	Timer handle as returned by rtdm_timer_init()
---------------	--------------	---

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xntimer_stop()`.

5.29.4.6 void `rtdm_timer_stop_in_handler` (`rtdm_timer_t * timer`)

Stop a timer from inside a timer handler.

Parameters

<i>in,out</i>	<i>timer</i>	Timer handle as returned by rtdm_timer_init()
---------------	--------------	---

Environments:


This service can be called from:

- Timer handler

Rescheduling: never.

5.30 Synchronisation Services

Collaboration diagram for Synchronisation Services:



Modules

- [Big dual kernel lock](#)
- [Spinlock with preemption deactivation](#)
- [Wait queue](#)
RTDM wait queues.
- [Event Services](#)
- [Semaphore Services](#)
- [Mutex services](#)

Functions

- void [rtdm_context_lock](#) (struct [rtdm_dev_context](#) *context)
Increment context reference counter.
- void [rtdm_context_unlock](#) (struct [rtdm_dev_context](#) *context)
Decrement context reference counter.
- void [rtdm_context_put](#) (struct [rtdm_dev_context](#) *context)
Release a device context obtained via [rtdm_context_get\(\)](#)

RTDM_SELECTTYPE_xxx

Event types select can bind to

- enum [rtdm_selecttype](#) { [RTDM_SELECTTYPE_READ](#) = XNSELECT_READ, [RTDM_SELECTTYPE_WRITE](#) = XNSELECT_WRITE, [RTDM_SELECTTYPE_EXCEPT](#) = XNSELECT_EXCEPT }

Timeout Sequence Management

- void [rtdm_toseq_init](#) ([rtdm_toseq_t](#) *timeout_seq, [nanosecs_rel_t](#) timeout)
Initialise a timeout sequence.

5.30.1 Detailed Description

5.30.2 Enumeration Type Documentation

5.30.2.1 enum **rtdm_selecttype**

Enumerator

RTDM_SELECTTYPE_READ Select input data availability events.

RTDM_SELECTTYPE_WRITE Select output buffer availability events.

RTDM_SELECTTYPE_EXCEPT Select exceptional events.

5.30.3 Function Documentation

5.30.3.1 void **rtdm_context_lock** (struct **rtdm_dev_context** * context)

Increment context reference counter.

Parameters

in	<i>context</i>	Device context
-----------	----------------	----------------

Note

rtdm_context_get() automatically increments the lock counter. You only need to call this function in special scenarios, e.g. when keeping additional references to the context structure that have different lifetimes. Only use [rtdm_context_lock\(\)](#) on contexts that are currently locked via an earlier **rtdm_context_get()/rtdm_context_lock()** or while running a device operation handler.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.30.3.2 void **rtdm_context_put** (struct **rtdm_dev_context** * context)

Release a device context obtained via **rtdm_context_get()**

Parameters

in	<i>context</i>	Device context
-----------	----------------	----------------

Note

Every successful call to **rtdm_context_get()** must be matched by a [rtdm_context_put\(\)](#) invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.30.3.3 void `rtm_context_unlock` (struct **rtm_dev_context** * context)

Decrement context reference counter.

Parameters

in	<i>context</i>	Device context
----	----------------	----------------

Note

Every call to `rtm_context_locked()` must be matched by a [rtm_context_unlock\(\)](#) invocation.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.30.3.4 void `rtm_toseq_init` (rtm_toseq_t * timeout_seq, **nanosecs_rel_t** timeout)

Initialise a timeout sequence.

This service initialises a timeout sequence handle according to the given timeout value. Timeout sequences allow to maintain a continuous *timeout* across multiple calls of blocking synchronisation services. A typical application scenario is given below.

Parameters

in,out	<i>timeout_seq</i>	Timeout sequence handle
in	<i>timeout</i>	Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values

Application Scenario:

```
int device_service_routine(...)
{
    rtm_toseq_t timeout_seq;
    ...

    rtm_toseq_init(&timeout_seq, timeout);
    ...
    while (received < requested) {
        ret = rtm_event_timedwait(&data_available, timeout, &timeout_seq);
        if (ret < 0) // including -ETIMEDOUT
            break;

        // receive some data
        ...
    }
    ...
}
```


Using a timeout sequence in such a scenario avoids that the user-provided relative `timeout` is restarted on every call to `rtdm_event_timedwait()`, potentially causing an overall delay that is larger than specified by `timeout`. Moreover, all functions supporting timeout sequences also interpret special timeout values (infinite and non-blocking), disburdening the driver developer from handling them separately.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: never.

5.31 Event Services

Collaboration diagram for Event Services:



Functions

- void [rt dm_event_init](#) (rt dm_event_t *event, unsigned long pending)
Initialise an event.
- void [rt dm_event_destroy](#) (rt dm_event_t *event)
Destroy an event.
- void [rt dm_event_pulse](#) (rt dm_event_t *event)
Signal an event occurrence to currently listening waiters.
- void [rt dm_event_signal](#) (rt dm_event_t *event)
Signal an event occurrence.
- int [rt dm_event_wait](#) (rt dm_event_t *event)
Wait on event occurrence.
- int [rt dm_event_timedwait](#) (rt dm_event_t *event, [nanosecs_rel_t](#) timeout, rt dm_toseq_t *timeout_seq)
Wait on event occurrence with timeout.
- void [rt dm_event_clear](#) (rt dm_event_t *event)
Clear event state.
- int [rt dm_event_select_bind](#) (rt dm_event_t *event, rt dm_selector_t *selector, enum [rt dm_selecttype](#) type, unsigned int fd_index)
Bind a selector to an event.

5.31.1 Detailed Description

5.31.2 Function Documentation

5.31.2.1 void [rt dm_event_clear](#) (rt dm_event_t * event)

Clear event state.

Parameters

in,out	<i>event</i>	Event handle as returned by rt dm_event_init()
--------	--------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_signal()`.

5.31.2.2 `void rtdm_event_destroy (rtdm_event_t * event)`

Destroy an event.

Parameters

<i>in, out</i>	<i>event</i>	Event handle as returned by rtdm_event_init()
----------------	--------------	---

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References `XNRMID`, and `xnselect_destroy()`.

5.31.2.3 `void rtdm_event_init (rtdm_event_t * event, unsigned long pending)`

Initialise an event.

Parameters

<i>in, out</i>	<i>event</i>	Event handle
<i>in</i>	<i>pending</i>	Non-zero if event shall be initialised as set, 0 otherwise

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_init()`, and `xnsynch_init()`.

5.31.2.4 `void rtdm_event_pulse (rtdm_event_t * event)`

Signal an event occurrence to currently listening waiters.

This function wakes up all current waiters of the given event, but it does not change the event state. Subsequently callers of [rtdm_event_wait\(\)](#) or [rtdm_event_timedwait\(\)](#) will therefore be blocked first.

Parameters

in,out	<i>event</i>	Event handle as returned by rt dm_event_init()
--------	--------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.31.2.5 `int rtdm_event_select_bind (rtdm_event_t * event, rtdm_selector_t * selector, enum rtdm_selecttype type, unsigned int fd_index)`

Bind a selector to an event.

This functions binds the given selector to an event so that the former is notified when the event state changes. Typically the select binding handler will invoke this service.

Parameters

in,out	<i>event</i>	Event handle as returned by rt dm_event_init()
in,out	<i>selector</i>	Selector as passed to the select binding handler
in	<i>type</i>	Type of the bound event as passed to the select binding handler
in	<i>fd_index</i>	File descriptor index as passed to the select binding handler

Returns

0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.
- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References `xnselect_bind()`.

5.31.2.6 `void rtdm_event_signal (rtdm_event_t * event)`

Signal an event occurrence.

This function sets the given event and wakes up all current waiters. If no waiter is presently registered, the next call to [rt dm_event_wait\(\)](#) or [rt dm_event_timedwait\(\)](#) will return immediately.

Parameters

in,out	<i>event</i>	Event handle as returned by rt dm_event_init()
--------	--------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References [xnsched_run\(\)](#), [xnselect_signal\(\)](#), and [xnsynch_flush\(\)](#).

5.31.2.7 `int rtdm_event_timedwait (rtdm_event_t * event, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq)`

Wait on event occurrence with timeout.

This function waits or tests for the occurrence of the given event, taking the provided timeout into account. On successful return, the event is reset.

Parameters

in,out	<i>event</i>	Event handle as returned by rt dm_event_init()
in	<i>timeout</i>	Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values
in,out	<i>timeout_seq</i>	Handle of a timeout sequence as returned by rt dm_toseq_init() or NULL

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the request has not been satisfied within the specified amount of time.
- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rt dm_task_unblock\(\)](#).
- -EIDRM is returned if *event* has been destroyed.
- -EPERM may be returned if an illegal invocation environment is detected.
- -EWOULDBLOCK is returned if a negative *timeout* (i.e., non-blocking operation) has been specified.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [XNBREAK](#), [XNRMID](#), [xnselect_signal\(\)](#), [xnsynch_sleep_on\(\)](#), and [XNTIMEO](#).

Referenced by [rt dm_event_wait\(\)](#).

5.31.2.8 `int rtdm_event_wait (rtdm_event_t * event)`

Wait on event occurrence.

This is the light-weight version of [rtdm_event_timedwait\(\)](#), implying an infinite timeout.

Parameters

in,out	<i>event</i>	Event handle as returned by rtdm_event_init()
--------	--------------	---

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtdm_task_unblock\(\)](#).
- -EIDRM is returned if *event* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [rtdm_event_timedwait\(\)](#).

5.32 Semaphore Services

Collaboration diagram for Semaphore Services:



Functions

- void [rtdm_sem_init](#) (rtdm_sem_t *sem, unsigned long value)
Initialise a semaphore.
- void [rtdm_sem_destroy](#) (rtdm_sem_t *sem)
Destroy a semaphore.
- int [rtdm_sem_down](#) (rtdm_sem_t *sem)
Decrement a semaphore.
- int [rtdm_sem_timeddown](#) (rtdm_sem_t *sem, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)
Decrement a semaphore with timeout.
- void [rtdm_sem_up](#) (rtdm_sem_t *sem)
Increment a semaphore.
- int [rtdm_sem_select_bind](#) (rtdm_sem_t *sem, rtdm_selector_t *selector, enum [rtdm_selecttype](#) type, unsigned int fd_index)
Bind a selector to a semaphore.

5.32.1 Detailed Description

5.32.2 Function Documentation

5.32.2.1 void rtdm_sem_destroy (rtdm_sem_t * sem)

Destroy a semaphore.

Parameters

in,out	sem	Semaphore handle as returned by rtdm_sem_init()
--------	-----	---

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References XNRMID, and xselect_destroy().

5.32.2.2 int rtdm_sem_down (rtdm_sem_t * sem)

Decrement a semaphore.

This is the light-weight version of [rtdm_sem_timeddown\(\)](#), implying an infinite timeout.

Parameters

in,out	<i>sem</i>	Semaphore handle as returned by rtdm_sem_init()
--------	------------	---

Returns

0 on success, otherwise:

- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtdm_task_unblock\(\)](#).
- -EIDRM is returned if *sem* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [rtdm_sem_timeddown\(\)](#).

5.32.2.3 void rtdm_sem_init (rtdm_sem_t * sem, unsigned long value)

Initialise a semaphore.

Parameters

in,out	<i>sem</i>	Semaphore handle
in	<i>value</i>	Initial value of the semaphore

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xnselect_init\(\)](#), and [xnsynch_init\(\)](#).

5.32.2.4 int rtdm_sem_select_bind (rtdm_sem_t * sem, rtdm_selector_t * selector, enum **rtdm_selecttype** type, unsigned int fd_index)

Bind a selector to a semaphore.

This functions binds the given selector to the semaphore so that the former is notified when the semaphore state changes. Typically the select binding handler will invoke this service.

Parameters

in,out	<i>sem</i>	Semaphore handle as returned by rtm_sem_init()
in,out	<i>selector</i>	Selector as passed to the select binding handler
in	<i>type</i>	Type of the bound event as passed to the select binding handler
in	<i>fd_index</i>	File descriptor index as passed to the select binding handler

Returns

0 on success, otherwise:

- -ENOMEM is returned if there is insufficient memory to establish the dynamic binding.
- -EINVAL is returned if *type* or *fd_index* are invalid.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References [xnselect_bind\(\)](#).

5.32.2.5 `int rtdm_sem_timeddown (rtdm_sem_t * sem, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq)`

Decrement a semaphore with timeout.

This function tries to decrement the given semaphore's value if it is positive on entry. If not, the caller is blocked unless non-blocking operation was selected.

Parameters

in,out	<i>sem</i>	Semaphore handle as returned by rtm_sem_init()
in	<i>timeout</i>	Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values
in,out	<i>timeout_seq</i>	Handle of a timeout sequence as returned by rtm_toseq_init() or NULL

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the request has not been satisfied within the specified amount of time.
- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.
- -EINTR is returned if calling task has been unblock by a signal or explicitly via [rtm_task_unblock\(\)](#).
- -EIDRM is returned if *sem* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References XNBREAK, XNRMID, xnsselect_signal(), xnsynch_sleep_on(), and XNTIMEO.

Referenced by rtdm_sem_down().

5.32.2.6 void rtdm_sem_up (rtdm_sem_t * sem)

Increment a semaphore.

This function increments the given semaphore's value, waking up a potential waiter which was blocked upon [rtdm_sem_down\(\)](#).

Parameters

in, out	sem	Semaphore handle as returned by rtdm_sem_init()
---------	-----	---

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References xnsched_run(), xnsselect_signal(), and xnsynch_wakeup_one_sleeper().

5.33 Mutex services

Collaboration diagram for Mutex services:



Functions

- void [rt dm_mutex_init](#) (rt dm_mutex_t *mutex)
Initialise a mutex.
- void [rt dm_mutex_destroy](#) (rt dm_mutex_t *mutex)
Destroy a mutex.
- void [rt dm_mutex_unlock](#) (rt dm_mutex_t *mutex)
Release a mutex.
- int [rt dm_mutex_lock](#) (rt dm_mutex_t *mutex)
Request a mutex.
- int [rt dm_mutex_timedlock](#) (rt dm_mutex_t *mutex, [nanosecs_rel_t](#) timeout, rt dm_toseq_t *timeout-
_seq)
Request a mutex with timeout.

5.33.1 Detailed Description

5.33.2 Function Documentation

5.33.2.1 void rt dm_mutex_destroy (rt dm_mutex_t * mutex)

Destroy a mutex.

Parameters

in,out	<i>mutex</i>	Mutex handle as returned by rt dm_mutex_init()
--------	--------------	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References XNRMID.

5.33.2.2 void rtdm_mutex_init (rtdm_mutex_t * mutex)

Initialise a mutex.

This function initialises a basic mutex with priority inversion protection. "Basic", as it does not allow a mutex owner to recursively lock the same mutex again.

Parameters

in,out	<i>mutex</i>	Mutex handle
--------	--------------	--------------

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

References xnsynch_init().

5.33.2.3 int rtdm_mutex_lock (rtdm_mutex_t * mutex)

Request a mutex.

This is the light-weight version of [rtdm_mutex_timedlock\(\)](#), implying an infinite timeout.

Parameters

in,out	<i>mutex</i>	Mutex handle as returned by rtdm_mutex_init()
--------	--------------	---

Returns

0 on success, otherwise:

- -EIDRM is returned if *mutex* has been destroyed.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [rtdm_mutex_timedlock\(\)](#).

5.33.2.4 int rtdm_mutex_timedlock (rtdm_mutex_t * mutex, nanosecs_rel_t timeout, rtdm_toseq_t * timeout_seq)

Request a mutex with timeout.

This function tries to acquire the given mutex. If it is not available, the caller is blocked unless non-blocking operation was selected.

Parameters

in,out	<i>mutex</i>	Mutex handle as returned by rtm_mutex_init()
in	<i>timeout</i>	Relative timeout in nanoseconds, see RTDM_TIMEOUT_xxx for special values
in,out	<i>timeout_seq</i>	Handle of a timeout sequence as returned by rtm_toseq_init() or NULL

Returns

0 on success, otherwise:

- -ETIMEDOUT is returned if the request has not been satisfied within the specified amount of time.
- -EWOULDBLOCK is returned if *timeout* is negative and the semaphore value is currently not positive.
- -EIDRM is returned if *mutex* has been destroyed.
- -EPERM may be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References XNBREAK, XNRMID, xnsynch_acquire(), and XNTIMEO.

Referenced by [rtm_mutex_lock\(\)](#).

5.33.2.5 void [rtm_mutex_unlock](#) ([rtm_mutex_t](#) * mutex)

Release a mutex.

This function releases the given mutex, waking up a potential waiter which was blocked upon [rtm_mutex_lock\(\)](#) or [rtm_mutex_timedlock\(\)](#).

Parameters

in,out	<i>mutex</i>	Mutex handle as returned by rtm_mutex_init()
--------	--------------	--

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT)

Rescheduling: possible.

References [xnsched_run\(\)](#), and [xnsynch_release\(\)](#).

5.34 Interrupt Management Services

Collaboration diagram for Interrupt Management Services:



Macros

- `#define rtdm_irq_get_arg(irq_handle, type) ((type *)irq_handle->cookie)`
Retrieve IRQ handler argument.

Typedefs

- `typedef int(* rtdm_irq_handler_t)(rtdm_irq_t *irq_handle)`
Interrupt handler.

Functions

- `int rtdm_irq_request (rtdm_irq_t *irq_handle, unsigned int irq_no, rtdm_irq_handler_t handler, unsigned long flags, const char *device_name, void *arg)`
Register an interrupt handler.
- `int rtdm_irq_free (rtdm_irq_t *irq_handle)`
Release an interrupt handler.
- `int rtdm_irq_enable (rtdm_irq_t *irq_handle)`
Enable interrupt line.
- `int rtdm_irq_disable (rtdm_irq_t *irq_handle)`
Disable interrupt line.

RTDM_IRQTYPE_XXX

Interrupt registrations flags

- `#define RTDM_IRQTYPE_SHARED XN_ISR_SHARED`
Enable IRQ-sharing with other real-time drivers.
- `#define RTDM_IRQTYPE_EDGE XN_ISR_EDGE`
Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.

RTDM_IRQ_XXX

Return flags of interrupt handlers

- #define [RTDM_IRQ_NONE](#) XN_ISR_NONE
Unhandled interrupt.
- #define [RTDM_IRQ_HANDLED](#) XN_ISR_HANDLED
Denote handled interrupt.

5.34.1 Detailed Description

5.34.2 Macro Definition Documentation

5.34.2.1 #define rtdm_irq_get_arg(irq_handle, type) ((type *)irq_handle->cookie)

Retrieve IRQ handler argument.

Parameters

<i>irq_handle</i>	IRQ handle
<i>type</i>	Type of the pointer to return

Returns

The argument pointer registered on [rtdm_irq_request\(\)](#) is returned, type-casted to the specified *type*.

Environments:

This service can be called from:

- Interrupt service routine

Rescheduling: never.

5.34.3 Typedef Documentation

5.34.3.1 typedef int(* rtdm_irq_handler_t)(rtdm_irq_t *irq_handle)

Interrupt handler.

Parameters

<i>in</i>	<i>irq_handle</i>	IRQ handle as returned by rtdm_irq_request()
-----------	-------------------	--

Returns

0 or a combination of [RTDM_IRQ_XXX](#) flags

5.34.4 Function Documentation

5.34.4.1 int rtdm_irq_disable (rtdm_irq_t * irq_handle)

Disable interrupt line.

Parameters

in,out	<i>irq_handle</i>	IRQ handle as returned by rtm_irq_request()
--------	-------------------	---

Returns

0 on success, otherwise negative error code

Note

This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization. Also, certain IRQ types may not allow the invocation over RT and interrupt contexts. The caller is responsible for excluding such conflicts.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.34.4.2 int rtm_irq_enable (rtm_irq_t * irq_handle)

Enable interrupt line.

Parameters

in,out	<i>irq_handle</i>	IRQ handle as returned by rtm_irq_request()
--------	-------------------	---

Returns

0 on success, otherwise negative error code

Note

This service is for exceptional use only. Drivers should always prefer interrupt masking at device level (via corresponding control registers etc.) over masking at line level. Keep in mind that the latter is incompatible with IRQ line sharing and can also be more costly as interrupt controller access requires broader synchronization. Also, certain IRQ types may not allow the invocation over RT and interrupt contexts. The caller is responsible for excluding such conflicts.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

5.34.4.3 `int rtdm_irq_free (rtdm_irq_t * irq_handle)`

Release an interrupt handler.

Parameters

in,out	<i>irq_handle</i>	IRQ handle as returned by rt dm_irq_request()
--------	-------------------	---

Returns

0 on success, otherwise negative error code

Note

The caller is responsible for shutting down the IRQ source at device level before invoking this service. In turn, `rt dm_irq_free` ensures that any pending event on the given IRQ line is fully processed on return from this service.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: never.

5.34.4.4 `int rt dm_irq_request (rt dm_irq_t * irq_handle, unsigned int irq_no, rt dm_irq_handler_t handler, unsigned long flags, const char * device_name, void * arg)`

Register an interrupt handler.

This function registers the provided handler with an IRQ line and enables the line.

Parameters

in,out	<i>irq_handle</i>	IRQ handle
in	<i>irq_no</i>	Line number of the addressed IRQ
in	<i>handler</i>	Interrupt handler
in	<i>flags</i>	Registration flags, see RTDM_IRQTYPE_xxx for details
in	<i>device_name</i>	Device name to show up in real-time IRQ lists
in	<i>arg</i>	Pointer to be passed to the interrupt handler on invocation

Returns

0 on success, otherwise:

- -EINVAL is returned if an invalid parameter was passed.
- -EBUSY is returned if the specified IRQ line is already in use.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)


Rescheduling: never.

References `xnintr_attach()`, `xnintr_enable()`, and `xnintr_init()`.

5.35 Non-Real-Time Signalling Services

These services provide a mechanism to request the execution of a specified handler in non-real-time context.

Collaboration diagram for Non-Real-Time Signalling Services:



Typedefs

- typedef void(* [rtm_nrtsig_handler_t](#))(rtm_nrtsig_t nrt_sig, void *arg)
Non-real-time signal handler.

Functions

- int [rtm_nrtsig_init](#) (rtm_nrtsig_t *nrt_sig, [rtm_nrtsig_handler_t](#) handler, void *arg)
Register a non-real-time signal handler.
- void [rtm_nrtsig_destroy](#) (rtm_nrtsig_t *nrt_sig)
Release a non-realtime signal handler.
- void [rtm_nrtsig_pend](#) (rtm_nrtsig_t *nrt_sig)
Trigger non-real-time signal.

5.35.1 Detailed Description

These services provide a mechanism to request the execution of a specified handler in non-real-time context. The triggering can safely be performed in real-time context without suffering from unknown delays. The handler execution will be deferred until the next time the real-time subsystem releases the CPU to the non-real-time part.

5.35.2 Typedef Documentation

5.35.2.1 typedef void(* [rtm_nrtsig_handler_t](#))(rtm_nrtsig_t nrt_sig, void *arg)

Non-real-time signal handler.

Parameters

in	<i>nrt_sig</i>	Signal handle as returned by rtm_nrtsig_init()
in	<i>arg</i>	Argument as passed to rtm_nrtsig_init()

Note

The signal handler will run in soft-IRQ context of the non-real-time subsystem. Note the implications of this context, e.g. no invocation of blocking operations.

5.35.3 Function Documentation

5.35.3.1 void rtdm_nrtsig_destroy (rtdm_nrtsig_t * nrt_sig)

Release a non-realtime signal handler.

Parameters

in,out	<i>nrt_sig</i>	Signal handle
--------	----------------	---------------

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.35.3.2 int rtdm_nrtsig_init (rtdm_nrtsig_t * nrt_sig, **rtdm_nrtsig_handler_t** handler, void * arg)

Register a non-real-time signal handler.

Parameters

in,out	<i>nrt_sig</i>	Signal handle
in	<i>handler</i>	Non-real-time signal handler
in	<i>arg</i>	Custom argument passed to handler() on each invocation

Returns

0 on success, otherwise:

- -EAGAIN is returned if no free signal slot is available.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.35.3.3 void rtdm_nrtsig_pend (rtdm_nrtsig_t * nrt_sig)

Trigger non-real-time signal.

Parameters

in,out	<i>nrt_sig</i>	Signal handle
--------	----------------	---------------

Environments:


This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

5.36 Utility Services

Collaboration diagram for Utility Services:



Functions

- `int rtdm_mmap_to_user` (struct rtdm_fd *fd, void *src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)
Map a kernel memory range into the address space of the user.
- `int rtdm_iomap_to_user` (struct rtdm_fd *fd, phys_addr_t src_addr, size_t len, int prot, void **pptr, struct vm_operations_struct *vm_ops, void *vm_private_data)
Map an I/O memory range into the address space of the user.
- `int rtdm_munmap` (struct rtdm_fd *fd, void *ptr, size_t len)
Unmap a user memory range.
- `int rtdm_ratelimit` (struct rtdm_ratelimit_state *rs, const char *func)
Enforces a rate limit.
- `void rtdm_printk_ratelimited` (const char *format,...)
Real-time safe rate-limited message printing on kernel console.
- `void rtdm_printk` (const char *format,...)
Real-time safe message printing on kernel console.
- `void * rtdm_malloc` (size_t size)
Allocate memory block in real-time context.
- `void rtdm_free` (void *ptr)
Release real-time memory block.
- `int rtdm_read_user_ok` (struct rtdm_fd *fd, const void __user *ptr, size_t size)
Check if read access to user-space memory block is safe.
- `int rtdm_rw_user_ok` (struct rtdm_fd *fd, const void __user *ptr, size_t size)
Check if read/write access to user-space memory block is safe.
- `int rtdm_copy_from_user` (struct rtdm_fd *fd, void *dst, const void __user *src, size_t size)
Copy user-space memory block to specified buffer.
- `int rtdm_safe_copy_from_user` (struct rtdm_fd *fd, void *dst, const void __user *src, size_t size)
Check if read access to user-space memory block and copy it to specified buffer.
- `int rtdm_copy_to_user` (struct rtdm_fd *fd, void __user *dst, const void *src, size_t size)
Copy specified buffer to user-space memory block.
- `int rtdm_safe_copy_to_user` (struct rtdm_fd *fd, void __user *dst, const void *src, size_t size)
Check if read/write access to user-space memory block is safe and copy specified buffer to it.
- `int rtdm_strncpy_from_user` (struct rtdm_fd *fd, char *dst, const char __user *src, size_t count)
Copy user-space string to specified buffer.
- `int rtdm_in_rt_context` (void)
Test if running in a real-time task.
- `int rtdm_rt_capable` (struct rtdm_fd *fd)
Test if the caller is capable of running in real-time context.

5.36.1 Detailed Description

5.36.2 Function Documentation

5.36.2.1 `int rtdm_copy_from_user (struct rtdm_fd * fd, void * dst, const void __user * src, size_t size)`

Copy user-space memory block to specified buffer.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>dst</i>	Destination buffer address
in	<i>src</i>	Address of the user-space memory block
in	<i>size</i>	Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

Before invoking this service, verify via [rtdm_read_user_ok\(\)](#) that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.2 `int rtdm_copy_to_user (struct rtdm_fd * fd, void __user * dst, const void * src, size_t size)`

Copy specified buffer to user-space memory block.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>dst</i>	Address of the user-space memory block
in	<i>src</i>	Source buffer address
in	<i>size</i>	Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

Before invoking this service, verify via [rtm_rw_user_ok\(\)](#) that the provided user-space address can securely be accessed.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.3 void rtdm_free (void * ptr)

Release real-time memory block.

Parameters

in	ptr	Pointer to memory block as returned by rtdm_malloc()
----	-----	--

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.4 int rtdm_in_rt_context (void)

Test if running in a real-time task.

Returns

Non-zero is returned if the caller resides in real-time context, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.5 int rtdm_iomap_to_user (struct rtdm_fd * fd, phys_addr_t src_addr, size_t len, int prot, void ** pptr, struct vm_operations_struct * vm_ops, void * vm_private_data)

Map an I/O memory range into the address space of the user.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>src_addr</i>	physical I/O address to be mapped
in	<i>len</i>	Length of the memory range
in	<i>prot</i>	Protection flags for the user's memory range, typically either PROT_READ or PROT_READ PROT_WRITE
in,out	<i>pptr</i>	Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return
in	<i>vm_ops</i>	vm_operations to be executed on the vma_area of the user memory range or NULL
in	<i>vm_private_data</i>	Private data to be stored in the vma_area, primarily useful for vm_operation handlers

Returns

0 on success, otherwise (most common values):

- -EINVAL is returned if an invalid start address, size, or destination address was passed.
- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.
- -EAGAIN is returned if too much memory has been already locked by the user process.
- -EPERM *may* be returned if an illegal invocation environment is detected.

Note

RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via [rtdm_munmap\(\)](#), either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard `munmap()` or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of [rtdm_iomap_to_user\(\)](#) can pass a `vm_operations_struct` on invocation, defining a close handler for the `vm_area`. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

5.36.2.6 void* rtdm_malloc (size_t size)

Allocate memory block in real-time context.

Parameters

<i>in</i>	<i>size</i>	Requested size of the memory block
-----------	-------------	------------------------------------

Returns

The pointer to the allocated block is returned on success, NULL otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

Referenced by `a4l_alloc_subd()`.

```
5.36.2.7 int rtdm_mmap_to_user ( struct rtdm_fd * fd, void * src_addr, size_t len, int prot, void **
    pptr, struct vm_operations_struct * vm_ops, void * vm_private_data )
```

Map a kernel memory range into the address space of the user.

Parameters

<i>in</i>	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
<i>in</i>	<i>src_addr</i>	Kernel virtual address to be mapped
<i>in</i>	<i>len</i>	Length of the memory range
<i>in</i>	<i>prot</i>	Protection flags for the user's memory range, typically either PROT_READ or PROT_READ PROT_WRITE
<i>in,out</i>	<i>pptr</i>	Address of a pointer containing the desired user address or NULL on entry and the finally assigned address on return
<i>in</i>	<i>vm_ops</i>	vm_operations to be executed on the vma_area of the user memory range or NULL
<i>in</i>	<i>vm_private_data</i>	Private data to be stored in the vma_area, primarily useful for vm_operation handlers

Returns

0 on success, otherwise (most common values):

- -EINVAL is returned if an invalid start address, size, or destination address was passed.
- -ENOMEM is returned if there is insufficient free memory or the limit of memory mapping for the user process was reached.
- -EAGAIN is returned if too much memory has been already locked by the user process.
- -EPERM may be returned if an illegal invocation environment is detected.

Note

This service only works on memory regions allocated via `kmalloc()` or `vmalloc()`. To map physical I/O memory to user-space use `rtm_iomap_to_user()` instead.

RTDM supports two models for unmapping the user memory range again. One is explicit unmapping via `rtm_munmap()`, either performed when the user requests it via an IOCTL etc. or when the related device is closed. The other is automatic unmapping, triggered by the user invoking standard `munmap()` or by the termination of the related process. To track release of the mapping and therefore relinquishment of the referenced physical memory, the caller of `rtm_mmap_to_user()` can pass a `vm_operations_struct` on invocation, defining a close handler for the `vm_area`. See Linux documentaion (e.g. Linux Device Drivers book) on virtual memory management for details.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

5.36.2.8 `int rtm_munmap (struct rtm_fd * fd, void * ptr, size_t len)`

Unmap a user memory range.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to <code>rtm_mmap_to_user()</code> when requesting to map the memory range
in	<i>ptr</i>	User address or the memory range
in	<i>len</i>	Length of the memory range

Returns

0 on success, otherwise:

- `-EINVAL` is returned if an invalid address or size was passed.
- `-EPERM` may be returned if an illegal invocation environment is detected.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- User-space task (non-RT)

Rescheduling: possible.

5.36.2.9 `void rtm_printk (const char * format, ...)`

Real-time safe message printing on kernel console.

Parameters

in	<i>format</i>	Format string (conforming standard <code>printf()</code>)
	...	Arguments referred by <i>format</i>

Returns

On success, this service returns the number of characters printed. Otherwise, a negative error code is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

5.36.2.10 `void rtdm_printk_ratelimited (const char * format, ...)`

Real-time safe rate-limited message printing on kernel console.

Parameters

in	<i>format</i>	Format string (conforming standard <code>printf()</code>)
	...	Arguments referred by <i>format</i>

Returns

On success, this service returns the number of characters printed. Otherwise, a negative error code is returned.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Interrupt service routine (consider the overhead!)
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never in real-time context, possible in non-real-time environments.

5.36.2.11 `int rtdm_ratelimit (struct rtdm_ratelimit_state * rs, const char * func)`

Enforces a rate limit.

This function enforces a rate limit: not more than *rs->burst* callbacks in every *rs->interval*.

Parameters

<i>in, out</i>	<i>rs</i>	rtm_ratelimit_state data
<i>in</i>	<i>func</i>	name of calling function

Returns

0 means callback will be suppressed and 1 means go ahead and do it

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: possible.

References `rtm_clock_read()`, and `rtm_lock_put_irqrestore()`.

5.36.2.12 `int rtm_read_user_ok (struct rtm_fd * fd, const void __user * ptr, size_t size)`

Check if read access to user-space memory block is safe.

Parameters

<i>in</i>	<i>fd</i>	RTM file descriptor as passed to the invoked device operation handler
<i>in</i>	<i>ptr</i>	Address of the user-provided memory block
<i>in</i>	<i>size</i>	Size of the memory block

Returns

Non-zero is return when it is safe to read from the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.13 `int rtm_rt_capable (struct rtm_fd * fd)`

Test if the caller is capable of running in real-time context.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
-----------	-----------	--

Returns

Non-zero is returned if the caller is able to execute in real-time context (independent of its current execution mode), 0 otherwise.

Note

This function can be used by drivers that provide different implementations for the same service depending on the execution mode of the caller. If a caller requests such a service in non-real-time context but is capable of running in real-time as well, it might be appropriate for the driver to reject the request via -ENOSYS so that RTDM can switch the caller and restart the request in real-time context.

Environments:

This service can be called from:

- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.14 `int rtdm_rw_user_ok (struct rtdm_fd * fd, const void __user * ptr, size_t size)`

Check if read/write access to user-space memory block is safe.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>ptr</i>	Address of the user-provided memory block
in	<i>size</i>	Size of the memory block

Returns

Non-zero is return when it is safe to read from or write to the specified memory block, 0 otherwise.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.15 `int rtdm_safe_copy_from_user (struct rtdm_fd * fd, void * dst, const void __user * src, size_t size)`

Check if read access to user-space memory block and copy it to specified buffer.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>dst</i>	Destination buffer address
in	<i>src</i>	Address of the user-space memory block
in	<i>size</i>	Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

This service is a combination of `rtdm_read_user_ok` and `rtdm_copy_from_user`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.16 `int rtdm_safe_copy_to_user (struct rtdm_fd * fd, void __user * dst, const void * src, size_t size)`

Check if read/write access to user-space memory block is safe and copy specified buffer to it.

Parameters

in	<i>fd</i>	RTDM file descriptor as passed to the invoked device operation handler
in	<i>dst</i>	Address of the user-space memory block
in	<i>src</i>	Source buffer address
in	<i>size</i>	Size of the memory block

Returns

0 on success, otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

This service is a combination of `rtdm_rw_user_ok` and `rtdm_copy_to_user`.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.36.2.17 `int rtdm_strncpy_from_user (struct rtdm_fd * fd, char * dst, const char __user * src, size_t count)`

Copy user-space string to specified buffer.

Parameters

<code>in</code>	<code>fd</code>	RTDM file descriptor as passed to the invoked device operation handler
<code>in</code>	<code>dst</code>	Destination buffer address
<code>in</code>	<code>src</code>	Address of the user-space string
<code>in</code>	<code>count</code>	Maximum number of bytes to copy, including the trailing '0'

Returns

Length of the string on success (not including the trailing '0'), otherwise:

- -EFAULT is returned if an invalid memory area was accessed.

Note

This services already includes a check of the source address, calling [rtdm_read_user_ok\(\)](#) for `src` explicitly is not required.

Environments:

This service can be called from:


- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

5.37 Device Profiles

Pre-defined classes of real-time devices.

Collaboration diagram for Device Profiles:



Modules

- [CAN Devices](#)

This is the common interface a RTDM-compliant CAN device has to provide.

- [Serial Devices](#)

This is the common interface a RTDM-compliant serial device has to provide.

- [Testing Devices](#)

This group of devices is intended to provide in-kernel testing results.

- [Real-time IPC](#)

Profile Revision: 1

Data Structures

- struct [rtdm_device_info](#)

Device information.

Typedefs

- typedef struct [rtdm_device_info](#) [rtdm_device_info_t](#)

Device information.

RTDM_CLASS_xxx

Device classes

- **#define RTDM_CLASS_PARPORT 1**
- **#define RTDM_CLASS_SERIAL 2**
- **#define RTDM_CLASS_CAN 3**
- **#define RTDM_CLASS_NETWORK 4**
- **#define RTDM_CLASS_RTMAC 5**

- `#define RTDM_CLASS_TESTING 6`
- `#define RTDM_CLASS_RTIPC 7`
- `#define RTDM_CLASS_COBALT 8`
- `#define RTDM_CLASS_EXPERIMENTAL 224`
- `#define RTDM_CLASS_MAX 255`

Device Naming

Maximum length of device names (excluding the final null character)

- `#define RTDM_MAX_DEVNAME_LEN 31`

RTDM_PURGE_xxx_BUFFER

Flags selecting buffers to be purged

- `#define RTDM_PURGE_RX_BUFFER 0x0001`
- `#define RTDM_PURGE_TX_BUFFER 0x0002`

Common IOCTLs

The following IOCTLs are common to all device `rtdm_profiles`.

- `#define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)`
Retrieve information about a device or socket.
- `#define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)`
Purge internal device or socket buffers.

5.37.1 Detailed Description

Pre-defined classes of real-time devices. Device profiles define which operation handlers a driver of a certain class has to implement, which name or protocol it has to register, which IOCTLs it has to provide, and further details. Sub-classes can be defined in order to extend a device profile with more hardware-specific functions.

5.37.2 Macro Definition Documentation

5.37.2.1 `#define RTIOC_DEVICE_INFO _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)`

Retrieve information about a device or socket.

Parameters

out	arg	Pointer to information buffer (struct rtdm_device_info)
-----	-----	--

5.37.2.2 `#define RTIOC_PURGE _IOW(RTIOC_TYPE_COMMON, 0x10, int)`

Purge internal device or socket buffers.

Parameters

in	<i>arg</i>	Purge mask, see RTDM_PURGE_xxx_BUFFER
----	------------	---

5.38 SCHED_QUOTA scheduling policy

This file implements the SCHED_QUOTA scheduling policy for Xenomai.

5.38.1 Detailed Description


This file implements the SCHED_QUOTA scheduling policy for Xenomai. This policy enforces a limitation on the CPU consumption of threads over a globally defined period, known as the quota interval. This is done by pooling threads with common requirements in groups, and giving each group a share of the global period (CONFIG_XENO_OPT_SCHED_QUOTA_PERIOD).

When threads have entirely consumed the quota allotted to the group they belong to, the latter is suspended as a whole, until the next quota interval starts. At this point, a new runtime budget is given to each group, in accordance with its share.

5.39 File descriptors events multiplexing services

File descriptors events multiplexing services.

Collaboration diagram for File descriptors events multiplexing services:



Functions

- void [xnselect_init](#) (struct [xnselect](#) *select_block)
Initialize a struct xnselect structure.
- static int [xnselect_signal](#) (struct [xnselect](#) *select_block, unsigned int state)
Signal a file descriptor state change.
- void [xnselect_destroy](#) (struct [xnselect](#) *select_block)
Destroy the xnselect structure associated with a file descriptor.
- int [xnselector_init](#) (struct xnselector *selector)
Initialize a selector structure.
- int [xnselect](#) (struct xnselector *selector, fd_set *out_fds[XNSELECT_MAX_TYPES], fd_set *in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)
Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.
- void [xnselector_destroy](#) (struct xnselector *selector)
Destroy a selector block.
- int [xnselect_bind](#) (struct [xnselect](#) *select_block, struct xnselect_binding *binding, struct xnselector *selector, unsigned type, unsigned index, unsigned state)
Bind a file descriptor (represented by its xnselect structure) to a selector block.

5.39.1 Detailed Description

File descriptors events multiplexing services.

Author

Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org

Copyright (C) 2008 Efixo

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

This module implements the services needed for implementing the posix "select" service, or any other events multiplexing services.

Following the implementation of the posix select service, this module defines three types of events:

- *XNSELECT_READ* meaning that a file descriptor is ready for reading;
- *XNSELECT_WRITE* meaning that a file descriptor is ready for writing;
- *XNSELECT_EXCEPT* meaning that a file descriptor received an exceptional event.

It works by defining two structures:

- a *struct xnselect* structure, which should be added to every file descriptor for every event type (read, write, or except);
- a *struct xnselector* structure, the selection structure, passed by the thread calling the xnselect service, where this service does all its housekeeping.

5.39.2 Function Documentation

5.39.2.1 `int xnselect (struct xnselector * selector, fd_set * out_fds[XNSELECT_MAX_TYPES], fd_set * in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)`

Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.

Parameters

<i>selector</i>	structure to check for pending events
<i>out_fds</i>	The set of descriptors with pending events if a strictly positive number is returned, or the set of descriptors not yet bound if -ECHRNG is returned;
<i>in_fds</i>	the set of descriptors which events should be checked
<i>nfds</i>	the highest-numbered descriptor in any of the <i>in_fds</i> sets, plus 1;
<i>timeout</i>	the timeout, whose meaning depends on <i>timeout_mode</i> , note that xnselect() pass <i>timeout</i> and <i>timeout_mode</i> unchanged to <i>xnsynch_sleep_on</i> , so passing a relative value different from XN_INFINITE as a timeout with <i>timeout_mode</i> set to XN_RELATIVE, will cause a longer sleep than expected if the sleep is interrupted.
<i>timeout_mode</i>	the mode of <i>timeout</i> .

Return values

<i>-EINVAL</i>	if <i>nfds</i> is negative;
<i>-ECHRNG</i>	if some of the descriptors passed in <i>in_fds</i> have not yet been registered with xnselect_bind() , <i>out_fds</i> contains the set of such descriptors;
<i>-EINTR</i>	if <i>xnselect</i> was interrupted while waiting;
<i>0</i>	in case of timeout.
<i>the</i>	number of file descriptors having received an event.

Remarks

Tags: might-switch.

References XNBREAK, *xnsynch_sleep_on()*, and XNTIMEO.

5.39.2.2 `int xnselect_bind (struct xnselect * select_block, struct xnselect_binding * binding, struct xnselector * selector, unsigned type, unsigned index, unsigned state)`

Bind a file descriptor (represented by its *xnselect* structure) to a selector block.

Parameters

<i>select_block</i>	pointer to the <i>struct xselect</i> to be bound;
<i>binding</i>	pointer to a newly allocated (using <code>xnmalloc</code>) <i>struct xselect_binding</i> ;
<i>selector</i>	pointer to the selector structure;
<i>type</i>	type of events (<code>XNSELECT_READ</code> , <code>XNSELECT_WRITE</code> , or <code>XNSELECT_EXCEPT</code>);
<i>index</i>	index of the file descriptor (represented by <i>select_block</i>) in the bit fields used by the <i>selector</i> structure;
<i>state</i>	current state of the file descriptor.

select_block must have been initialized with `xnselect_init()`, the *xnselector* structure must have been initialized with `xnselector_init()`, *binding* may be uninitialized.

This service must be called with `nklock` locked, `irqs` off. For this reason, the *binding* parameter must have been allocated by the caller outside the locking section.

Return values

<code>-EINVAL</code>	if <i>type</i> or <i>index</i> is invalid;
<code>0</code>	otherwise.

Remarks

Tags: might-switch, atomic-entry.

References `xnsched_run()`.

Referenced by `rtdm_event_select_bind()`, and `rtdm_sem_select_bind()`.

5.39.2.3 void xnselect_destroy (struct **xnselect** * select_block)

Destroy the *xnselect* structure associated with a file descriptor.

Any binding with a *xnselector* block is destroyed.

Parameters

<i>select_block</i>	pointer to the <i>xnselect</i> structure associated with a file descriptor
---------------------	--

Remarks

Tags: might-switch.

References `xnsched_run()`.

Referenced by `rtdm_event_destroy()`, and `rtdm_sem_destroy()`.

5.39.2.4 void xnselect_init (struct **xnselect** * select_block)

Initialize a *struct xnselect* structure.

This service must be called to initialize a *struct xnselect* structure before it is bound to a selector by the means of `xnselect_bind()`.

Parameters

<i>select_block</i>	pointer to the <i>xnselect</i> structure to be initialized
---------------------	--

Remarks

Tags: none.

Referenced by `rtdm_event_init()`, and `rtdm_sem_init()`.

5.39.2.5 `static int xnselect_signal (struct xnselect * select_block, unsigned int state) [inline], [static]`

Signal a file descriptor state change.

Parameters

<i>select_block</i>	pointer to an <i>xnselect</i> structure representing the file descriptor whose state changed;
<i>state</i>	new value of the state.

Return values

<i>1</i>	if rescheduling is needed;
<i>0</i>	otherwise.

Referenced by `rtdm_event_clear()`, `rtdm_event_signal()`, `rtdm_event_timedwait()`, `rtdm_sem_timeddown()`, and `rtdm_sem_up()`.

5.39.2.6 `void xnselector_destroy (struct xnselector * selector)`

Destroy a selector block.

All bindings with file descriptor are destroyed.

Parameters

<i>selector</i>	the selector block to be destroyed
-----------------	------------------------------------

Remarks

Tags: none.

5.39.2.7 `int xnselector_init (struct xnselector * selector)`

Initialize a selector structure.

Parameters

<i>selector</i>	The selector structure to be initialized.
-----------------	---

Return values

<i>0</i>	
----------	--

Remarks

Tags: none.

References `xnsynch_init()`.

5.40 Real-time shadow services

Copyright (C) 2001-2013 Philippe Gerum rpm@xenomai.org.

Collaboration diagram for Real-time shadow services:



Functions

- int [xnshadow_harden](#) (void)
Migrate a Linux task to the Xenomai domain.
- void [xnshadow_relax](#) (int notify, int reason)
Switch a shadow thread back to the Linux domain.
- int [xnshadow_map_user](#) (struct xnthread *thread, unsigned long __user *u_window_offset)
Create a shadow thread context over a user task.
- int [xnshadow_map_kernel](#) (struct xnthread *thread, struct completion *done)
Create a shadow thread context over a kernel task.
- int [xnshadow_register_personality](#) (struct xnpersonality *personality)
Register a new interface personality.
- void * [xnshadow_get_context](#) (unsigned int muxid)
Return the per-process data attached to the calling process.
- struct xnpersonality * [xnshadow_push_personality](#) (int muxid)
Stack a new personality over the current thread.
- void [xnshadow_pop_personality](#) (struct xnpersonality *prev)
Pop the topmost personality from the current thread.

5.40.1 Detailed Description

Copyright (C) 2001-2013 Philippe Gerum rpm@xenomai.org. Copyright (C) 2001-2013 The Xenomai project <http://www.xenomai.org> Copyright (C) 2006 Gilles Chanteperrin gilles.-chanteperrin@xenomai.org

SMP support Copyright (C) 2004 The HYADES project <http://www.hyades-itea.org> RTAI/fusion Copyright (C) 2004 The RTAI project <http://www.rtai.org>

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Real-time shadow services.

5.40.2 Function Documentation

5.40.2.1 void* xnshadow_get_context (unsigned int muxid)

Return the per-process data attached to the calling process.

This service returns the per-process data attached to the calling process for the personality whose muxid is *muxid*. It must be called with nklock locked, irqs off.

See [xnshadow_register_personality\(\)](#) documentation for information on the way to attach a per-process data to a process.

Parameters

<i>muxid</i>	the personality muxid.
--------------	------------------------

Returns

the per-process data if the current context is a user-space process;
NULL otherwise.

Remarks

Tags: atomic-entry.

5.40.2.2 int xnshadow_harden (void)

Migrate a Linux task to the Xenomai domain.

This service causes the transition of "current" from the Linux domain to Xenomai. The shadow will resume in the Xenomai domain as returning from `schedule()`.

Remarks

Tags: secondary-only, might-switch.

References XNDEBUG, XNRELAX, and `xnshadow_relax()`.

Referenced by `xnshadow_map_kernel()`.

5.40.2.3 int xnshadow_map_kernel (struct xnthread * thread, struct completion * done)

Create a shadow thread context over a kernel task.

This call maps a nucleus thread to the "current" Linux task running in kernel space. The priority and scheduling class of the underlying Linux task are not affected; it is assumed that the caller did set them appropriately before issuing the shadow mapping request.

This call immediately moves the calling kernel thread to the Xenomai domain.

Parameters

<i>thread</i>	The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to xnthread_init() .
---------------	---

<i>done</i>	A completion object to be signaled when <i>thread</i> is fully mapped over the current Linux context, waiting for xnthread_start() .
-------------	--

Returns

0 is returned on success. Otherwise:

- -ERESTARTSYS is returned if the current Linux task has received a signal, thus preventing the final migration to the Xenomai domain (i.e. in order to process the signal in the Linux domain). This error should not be considered as fatal.
- -EPERM is returned if the shadow thread has been killed before the current task had a chance to return to the caller. In such a case, the real-time mapping operation has failed globally, and no Xenomai resource remains attached to it.
- -EINVAL is returned if the thread control block bears the XNUSER bit.
- -EBUSY is returned if either the current Linux task or the associated shadow thread is already involved in a shadow mapping.

Remarks

Tags: secondary-only, might-switch.

References XNDORMANT, XNMAPPED, XNRELAX, [xnshadow_harden\(\)](#), [xnthread_resume\(\)](#), [xnthread_suspend\(\)](#), and XNUSER.

5.40.2.4 `int xnshadow_map_user (struct xnthread * thread, unsigned long __user * u_window_offset)`

Create a shadow thread context over a user task.

This call maps a nucleus thread to the "current" Linux task running in userland. The priority and scheduling class of the underlying Linux task are not affected; it is assumed that the interface library did set them appropriately before issuing the shadow mapping request.

Parameters

<i>thread</i>	The descriptor address of the new shadow thread to be mapped to "current". This descriptor must have been previously initialized by a call to xnthread_init() .
<i>u_window_offset</i>	will receive the offset of the per-thread "u_window" structure in the process shared heap associated to <i>thread</i> . This structure reflects thread state information visible from userland through a shared memory window.

Returns

0 is returned on success. Otherwise:

- -EINVAL is returned if the thread control block does not bear the XNUSER bit.
- -EBUSY is returned if either the current Linux task or the associated shadow thread is already involved in a shadow mapping.

Remarks

Tags: secondary-only.

References [xnheap_alloc\(\)](#), XNMAPPED, XNRELAX, [xnthread_start\(\)](#), [xnthread_suspend\(\)](#), and XNUSER.

5.40.2.5 void xnshadow_pop_personality (struct xnpersonality * prev)

Pop the topmost personality from the current thread.

This service pops the topmost personality off the current thread.

Parameters

<i>prev</i>	the previous personality which was returned by the latest call to xnshadow_push_personality() for the current thread.
-------------	---

Remarks

Tags: secondary-only.

5.40.2.6 struct xnpersonality* xnshadow_push_personality (int muxid)

Stack a new personality over the current thread.

This service registers the current thread as a member of the additional personality identified by *muxid*. If the current thread is already assigned this personality, the call returns successfully with no effect.

Parameters

<i>muxid</i>	the identifier of the additional personality.
--------------	---

Returns

A handle to the previous personality. The caller should save this handle for unstacking *muxid* when applicable via a call to [xnshadow_pop_personality\(\)](#).

Remarks

Tags: secondary-only.

5.40.2.7 int xnshadow_register_personality (struct xnpersonality * personality)

Register a new interface personality.

- `personality->ops.attach_process()` is called when a user-space process binds to the personality, on behalf of one of its threads. The `attach_process()` handler may return:
 - . an opaque pointer, representing the context of the calling process for this personality;
 - . a NULL pointer, meaning that no per-process structure should be attached to this process for this personality;
 - . `ERR_PTR(negative value)` indicating an error, the binding process will then abort.
- `personality->ops.detach()` is called on behalf of an exiting user-space process which has previously attached to the personality. This handler is passed a pointer to the per-process data received earlier from the `ops->attach_process()` handler.

Remarks

Tags: none.

5.40.2.8 void xnshadow_relax (int notify, int reason)

Switch a shadow thread back to the Linux domain.

This service yields the control of the running shadow back to Linux. This is obtained by suspending the shadow and scheduling a wake up call for the mated user task inside the Linux domain. The Linux task will resume on return from [xnthread_suspend\(\)](#) on behalf of the root thread.

Parameters

<i>notify</i>	A boolean flag indicating whether threads monitored from secondary mode switches should be sent a SIGDEBUG signal. For instance, some internal operations like task exit should not trigger such signal.
<i>reason</i>	The reason to report along with the SIGDEBUG signal.

Remarks

Tags: primary-only, might-switch.

Note

"current" is valid here since the shadow runs with the properties of the Linux task.


References XNMOVED, XNRELAX, xnthread_suspend(), XNTRAPSW, and XNUSER.

Referenced by xnshadow_harden().

5.41 Thread synchronization services.

Copyright (C) 2001-2008 Philippe Gerum rpm@xenomai.org.

Collaboration diagram for Thread synchronization services.:



Functions

- void [xnsynch_init](#) (struct xnsynch *synch, int flags, [atomic_long_t](#) *fastlock)
Initialize a synchronization object.
- int [xnsynch_sleep_on](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)
Sleep on an ownerless synchronization object.
- struct xnthread * [xnsynch_wakeup_one_sleeper](#) (struct xnsynch *synch)
Unblock the heading thread from wait.
- void [xnsynch_wakeup_this_sleeper](#) (struct xnsynch *synch, struct xnthread *sleeper)
Unblock a particular thread from wait.
- int [xnsynch_acquire](#) (struct xnsynch *synch, xnticks_t timeout, xntmode_t timeout_mode)
Acquire the ownership of a synchronization object.
- static void [clear_boost](#) (struct xnsynch *synch, struct xnthread *owner)
Clear the priority boost.
- struct xnthread * [xnsynch_release](#) (struct xnsynch *synch, struct xnthread *thread)
Give the resource ownership to the next waiting thread.
- void [xnsynch_requeue_sleeper](#) (struct xnthread *thread)
Change a sleeper's priority.
- struct xnthread * [xnsynch_peek_pendq](#) (struct xnsynch *synch)
Access the thread leading a synch object wait queue.
- int [xnsynch_flush](#) (struct xnsynch *synch, int reason)
Unblock all waiters pending on a resource.
- void [xnsynch_forget_sleeper](#) (struct xnthread *thread)
Abort a wait for a resource.
- void [xnsynch_release_all_ownerships](#) (struct xnthread *thread)
Release all ownerships.

5.41.1 Detailed Description

Copyright (C) 2001-2008 Philippe Gerum rpm@xenomai.org. Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

5.41.2 Function Documentation

5.41.2.1 void clear_boost (struct xnsynch * synch, struct xntthread * owner) [static]

Clear the priority boost.

This service is called internally whenever a synchronization object is not claimed anymore by sleepers to reset the object owner's priority to its initial level.

Parameters

<i>synch</i>	The descriptor address of the synchronization object.
<i>owner</i>	The descriptor address of the thread which currently owns the synchronization object.

Remarks

Tags: atomic-entry.

References XNBOOST, and XNZOMBIE.

Referenced by xnsynch_flush(), and xnsynch_forget_sleeper().

5.41.2.2 int xnsynch_acquire (struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode)

Acquire the ownership of a synchronization object.

This service should be called by upper interfaces wanting the current thread to acquire the ownership of the given resource. If the resource is already assigned to another thread, the caller is suspended.

This service must be used only with synchronization objects that track ownership (XNSYNCH_OWNER set).

Parameters

<i>synch</i>	The descriptor address of the synchronization object to acquire.
<i>timeout</i>	The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on <i>timeout_mode</i> . Passing XN_INFINITE and setting <i>mode</i> to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.
<i>timeout_mode</i>	The mode of the <i>timeout</i> parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also xntimer_start()).

Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Remarks

Tags: might-switch.

References XNBOOST, XNBREAK, XNPEND, XNRMID, XNROBBED, xntthread_suspend(), XNTIMEO, XNWAKEN, and XNWEAK.

Referenced by rtdm_mutex_timedlock().

5.41.2.3 `int xnsynch_flush (struct xnsynch * synch, int reason)`

Unblock all waiters pending on a resource.

This service atomically releases all threads which currently sleep on a given resource.

This service should be called by upper interfaces under circumstances requiring that the pending queue of a given resource is cleared, such as before the resource is deleted.

Parameters

<i>synch</i>	The descriptor address of the synchronization object to be flushed.
<i>reason</i>	Some flags to set in the information mask of every unblocked thread. Zero is an acceptable value. The following bits are pre-defined by the nucleus:

- XNRMID should be set to indicate that the synchronization object is about to be destroyed (see [xnthread_resume\(\)](#)).
- XNBREAK should be set to indicate that the wait has been forcibly interrupted (see [xnthread_unblock\(\)](#)).

Returns

XNSYNCH_RESCHEDED is returned if at least one thread is unblocked, which means the caller should invoke [xnsched_run\(\)](#) for applying the new scheduling state. Otherwise, XNSYNCH_DONE is returned.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object is no more owned by any thread.

Remarks

Tags: isr-allowed.

References [clear_boost\(\)](#), [XNPEND](#), and [xnthread_resume\(\)](#).

Referenced by [rtdm_event_signal\(\)](#).

5.41.2.4 `void xnsynch_forget_sleeper (struct xnthread * thread)`

Abort a wait for a resource.

Performs all the necessary housekeeping chores to stop a thread from waiting on a given synchronization object.

Parameters

<i>thread</i>	The descriptor address of the affected thread.
---------------	--

Remarks

Tags: atomic-entry.

References [clear_boost\(\)](#), and [XNPEND](#).

Referenced by [xnthread_resume\(\)](#), and [xnthread_suspend\(\)](#).

5.41.2.5 void xnsynch_init (struct xnsynch * synch, int flags, **atomic_long_t** * fastlock)

Initialize a synchronization object.

Initializes a synchronization object. Xenomai threads can wait on and signal such objects for serializing access to resources. This object has built-in support for priority inheritance.

Parameters

<i>synch</i>	The address of a synchronization object descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.
<i>flags</i>	A set of creation flags affecting the operation. The valid flags are:

- XNSYNCH_PRIO causes the threads waiting for the resource to pend in priority order. Otherwise, FIFO ordering is used (XNSYNCH_FIFO).
- XNSYNCH_OWNER indicates that the synchronization object shall track the resource ownership, allowing a single owner at most at any point in time. Note that setting this flag implies the use of [xnsynch_acquire\(\)](#) and [xnsynch_release\(\)](#) instead of [xnsynch_sleep_on\(\)](#) and [xnsynch_wakeup_*](#)().
- XNSYNCH_PIP enables priority inheritance when a priority inversion is detected among threads using this object. XNSYNCH_PIP enables XNSYNCH_OWNER and XNSYNCH_PRIO implicitly.
- XNSYNCH_DREORD (Disable REORdering) tells the nucleus that the wait queue should not be reordered whenever the priority of a blocked thread it holds is changed. If this flag is not specified, changing the priority of a blocked thread using [xnthread_set_schedparam\(\)](#) will cause this object's wait queue to be reordered according to the new priority level, provided the synchronization object makes the waiters wait by priority order on the awaited resource (XNSYNCH_PRIO).

Parameters

<i>fastlock</i>	Address of the fast lock word to be associated with a synchronization object with ownership tracking. Therefore, a valid fast-lock address is required if XNSYNCH_OWNER is set in <i>flags</i> .
-----------------	--

Remarks

Tags: none.

Referenced by [rtdm_event_init\(\)](#), [rtdm_mutex_init\(\)](#), [rtdm_sem_init\(\)](#), and [xnselector_init\(\)](#).

5.41.2.6 struct xnthread* xnsynch_peek_pendq (struct xnsynch * synch)

Access the thread leading a synch object wait queue.

This services returns the descriptor address of to the thread leading a synchronization object wait queue.

Parameters

<i>synch</i>	The descriptor address of the target synchronization object.
--------------	--

Returns

The descriptor address of the unblocked thread.

Remarks

Tags: isr-allowed.

5.41.2.7 `struct xnthread* xnsynch_release (struct xnsynch * synch, struct xnthread * thread)`

Give the resource ownership to the next waiting thread.

This service releases the ownership of the given synchronization object. The thread which is currently leading the object's pending list, if any, is unblocked from its pending state. However, no reschedule is performed.

This service must be used only with synchronization objects that track ownership (XNSYNCH_OWNER set).

Parameters

<i>synch</i>	The descriptor address of the synchronization object whose ownership is changed.
<i>thread</i>	The descriptor address of the current owner.

Returns

The descriptor address of the unblocked thread.

Side-effects:

- The effective priority of the previous resource owner might be lowered to its base priority value as a consequence of the priority inheritance boost being cleared.
- The synchronization object ownership is transfered to the unblocked thread.

Remarks

Tags: none.

References XNWEAK.

Referenced by `rtdm_mutex_unlock()`, and `xnsynch_release_all_ownerships()`.

5.41.2.8 `void xnsynch_release_all_ownerships (struct xnthread * thread)`

Release all ownerships.

This call is used internally to release all the ownerships obtained by a thread on synchronization objects. This routine must be entered interrupts off.

Parameters

<i>thread</i>	The descriptor address of the affected thread.
---------------	--

Remarks

Tags: atomic-entry.

References `xnsynch_release()`.

5.41.2.9 `void xnsynch_requeue_sleeper (struct xnthread * thread)`

Change a sleeper's priority.

This service is used by the PIP code to update the pending priority of a sleeping thread.

Parameters

<i>thread</i>	The descriptor address of the affected thread.
---------------	--

Remarks

Tags: atomic-entry.

References XNBOOST.

5.41.2.10 `int xnsynch_sleep_on (struct xnsynch * synch, xnticks_t timeout, xntmode_t timeout_mode)`

Sleep on an ownerless synchronization object.

Makes the calling thread sleep on the specified synchronization object, waiting for it to be signaled.

This service should be called by upper interfaces wanting the current thread to pend on the given resource. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER).

Parameters

<i>synch</i>	The descriptor address of the synchronization object to sleep on.
<i>timeout</i>	The timeout which may be used to limit the time the thread pends on the resource. This value is a wait time given as a count of nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on <i>timeout_mode</i> . Passing XN_INFINITE and setting <i>mode</i> to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer.
<i>timeout_mode</i>	The mode of the <i>timeout</i> parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also xntimer_start()).

Returns

A bitmask which may include zero or one information bit among XNRMID, XNTIMEO and XNBREAK, which should be tested by the caller, for detecting respectively: object deletion, timeout or signal/unblock conditions which might have happened while waiting.

Remarks

Tags: might-switch.

References XNBREAK, XNPEND, XNRMID, `xnthread_suspend()`, and XNTIMEO.

Referenced by `rtdm_event_timedwait()`, `rtdm_sem_timeddown()`, `xnregistry_bind()`, and `xnselect()`.

5.41.2.11 `struct xnthread* xnsynch_wakeup_one_sleeper (struct xnsynch * synch)`

Unblock the heading thread from wait.

This service wakes up the thread which is currently leading the synchronization object's pending list. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a single waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER not set).

Parameters

<i>synch</i>	The descriptor address of the synchronization object whose ownership is changed.
--------------	--

Returns

The descriptor address of the unblocked thread.

Remarks

Tags: isr-allowed.

References XNPEND, and `xnthread_resume()`.

Referenced by `rtdm_sem_up()`.

5.41.2.12 `void xnsynch_wakeup_this_sleeper (struct xnsynch * synch, struct xnthread * sleeper)`

Unblock a particular thread from wait.

This service wakes up a specific thread which is currently pending on the given synchronization object. The sleeping thread is unblocked from its pending state, but no reschedule is performed.

This service should be called by upper interfaces wanting to signal the given resource so that a specific waiter is resumed. It must not be used with synchronization objects that are supposed to track ownership (XNSYNCH_OWNER not set).

Parameters

<i>synch</i>	The descriptor address of the synchronization object whose ownership is changed.
<i>sleeper</i>	The thread to unblock which MUST be currently linked to the synchronization object's pending queue (i.e. <code>synch->pendq</code>).

Remarks

Tags: isr-allowed.

References XNPEND, and `xnthread_resume()`.

5.42 Thread services

Collaboration diagram for Thread services:



Functions

- int [xnthread_init](#) (struct xnthread *thread, const struct xnthread_init_attr *attr, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)
Initialize a new thread.
- int [xnthread_start](#) (struct xnthread *thread, const struct xnthread_start_attr *attr)
Start a newly created thread.
- int [xnthread_set_mode](#) (struct xnthread *thread, int clrmask, int setmask)
Change a thread's control mode.
- void [xnthread_suspend](#) (struct xnthread *thread, int mask, xnticks_t timeout, xntmode_t timeout_mode, struct xnsynch *wchan)
Suspend a thread.
- void [xnthread_resume](#) (struct xnthread *thread, int mask)
Resume a thread.
- int [xnthread_unblock](#) (struct xnthread *thread)
Unblock a thread.
- int [xnthread_set_periodic](#) (struct xnthread *thread, xnticks_t idate, xntmode_t timeout_mode, xnticks_t period)
Make a thread periodic.
- int [xnthread_wait_period](#) (unsigned long *overruns_r)
Wait for the next periodic release point.
- int [xnthread_set_slice](#) (struct xnthread *thread, xnticks_t quantum)
Set thread time-slicing information.
- void [xnthread_cancel](#) (struct xnthread *thread)
Cancel a thread.
- int [xnthread_set_schedparam](#) (struct xnthread *thread, struct xnsched_class *sched_class, const union xnsched_policy_param *sched_param)
Change the base scheduling parameters of a thread.
- static int [xnthread_migrate](#) (int cpu)
Migrate the current thread.

5.42.1 Detailed Description

5.42.2 Function Documentation

5.42.2.1 void xnthread_cancel (struct xnthread * thread)

Cancel a thread.

Request cancellation of a thread. This service forces *thread* to exit from any blocking call. *thread* will terminate as soon as it reaches a cancellation point. Cancellation points are defined for the following situations:

- *thread* self-cancels by a call to [xnthread_cancel\(\)](#).
- *thread* invokes a Linux syscall (user-space shadow only).
- *thread* receives a Linux signal (user-space shadow only).
- *thread* explicitly calls [xnthread_test_cancel\(\)](#).

Parameters

<i>thread</i>	The descriptor address of the thread to terminate.
---------------	--

Remarks

Tags: might-switch.

References XNCANCELLED, XNDORMANT, XNKICKED, [xnsched_run\(\)](#), and [xnthread_resume\(\)](#).

Referenced by [rtdm_task_init\(\)](#).

5.42.2.2 `int xnthread_init (struct xnthread * thread, const struct xnthread_init_attr * attr, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param)`

Initialize a new thread.

Initializes a new thread. The thread is left dormant until it is actually started by [xnthread_start\(\)](#).

Parameters

<i>thread</i>	The address of a thread descriptor the nucleus will use to store the thread-specific data. This descriptor must always be valid while the thread is active therefore it must be allocated in permanent memory.
---------------	--

Warning

Some architectures may require the descriptor to be properly aligned in memory; this is an additional reason for descriptors not to be laid in the program stack where alignment constraints might not always be satisfied.

Parameters

<i>attr</i>	A pointer to an attribute block describing the initial properties of the new thread. Members of this structure are defined as follows:
-------------	--

- **name:** An ASCII string standing for the symbolic name of the thread. This name is copied to a safe place into the thread descriptor. This name might be used in various situations by the nucleus for issuing human-readable diagnostic messages, so it is usually a good idea to provide a sensible value here. NULL is fine though and means "anonymous".
- **flags:** A set of creation flags affecting the operation. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the created thread:
 - XNSUSP creates the thread in a suspended state. In such a case, the thread shall be explicitly resumed using the [xnthread_resume\(\)](#) service for its execution to actually begin, additionally to issuing [xnthread_start\(\)](#) for it. This flag can also be specified when invoking [xnthread_start\(\)](#) as a starting mode.

- XNUSER shall be set if *thread* will be mapped over an existing user-space task. Otherwise, a new kernel host task is created, then paired with the new Xenomai thread.
- XNFPU (enable FPU) tells the nucleus that the new thread may use the floating-point unit. XNFPU is implicitly assumed for user-space threads even if not set in *flags*.
- affinity: The processor affinity of this thread. Passing CPU_MASK_ALL means "any cpu" from the allowed core affinity mask (*nkaffinity*). Passing an empty set is invalid.

Parameters

<i>sched_class</i>	The initial scheduling class the new thread should be assigned to.
<i>sched_param</i>	The initial scheduling parameters to set for the new thread; <i>sched_param</i> must be valid within the context of <i>sched_class</i> .

Returns

0 is returned on success. Otherwise, the following error code indicates the cause of the failure:

- -EINVAL is returned if *attr->flags* has invalid bits set, or *attr->affinity* is invalid (e.g. empty).

Remarks

Tags: secondary-only.

References XNFPU, XNSUSP, and XNUSER.

Referenced by *rdm_task_init()*.

5.42.2.3 `int xnthread_migrate (int cpu) [inline], [static]`

Migrate the current thread.

This call makes the current thread migrate to another (real-time) CPU if its affinity allows it. This call is available from primary mode only.

Parameters

<i>cpu</i>	The destination CPU.
------------	----------------------

Return values

0	if the thread could migrate ;
-EPERM	if the calling context is invalid, or the scheduler is locked.
-EINVAL	if the current thread affinity forbids this migration.

Remarks

Tags: might-switch.

5.42.2.4 `void xnthread_resume (struct xnthread * thread, int mask)`

Resume a thread.

Resumes the execution of a thread previously suspended by one or more calls to [xnthread_suspend\(\)](#). This call removes a suspensive condition affecting the target thread. When all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters

<i>thread</i>	The descriptor address of the resumed thread.
<i>mask</i>	The suspension mask specifying the suspensive condition to remove from the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag removes the explicit suspension condition. This condition might be additive to the XNPEND condition.
- XNDELAY. This flag removes the counted delay wait condition.
- XNPEND. This flag removes the resource wait condition. If a watchdog is armed, it is automatically disarmed by this call. Unlike the two previous conditions, only the current thread can set this condition for itself, i.e. no thread can force another one to pend on a resource.

When the thread is eventually resumed by one or more calls to [xnthread_resume\(\)](#), the caller of [xnthread_suspend\(\)](#) in the awakened thread that suspended itself should check for the following bits in its own information mask to determine what caused its wake up:

- XNRMID means that the caller must assume that the pended synchronization object has been destroyed (see [xnsynch_flush\(\)](#)).
- XNTIMEO means that the delay elapsed, or the watchdog went off before the corresponding synchronization object was signaled.
- XNBREAK means that the wait has been forcibly broken by a call to [xnthread_unblock\(\)](#).

Remarks

Tags: isr-allowed.

References XNDELAY, XNHELD, XNPEND, XNREADY, [xnsynch_forget_sleeper\(\)](#), and [xntimer_stop\(\)](#).

Referenced by [xnshadow_map_kernel\(\)](#), [xnsynch_flush\(\)](#), [xnsynch_wakeup_one_sleeper\(\)](#), [xnsynch_wakeup_this_sleeper\(\)](#), [xnthread_cancel\(\)](#), [xnthread_start\(\)](#), and [xnthread_unblock\(\)](#).

5.42.2.5 `int xnthread_set_mode (struct xnthread * thread, int clrmask, int setmask)`

Change a thread's control mode.

Change the control mode of a given thread. The control mode affects the behaviour of the nucleus regarding the specified thread.

Parameters

<i>thread</i>	The descriptor address of the affected thread.
<i>clrmask</i>	Clears the corresponding bits from the control field before <i>setmask</i> is applied. The scheduler lock held by the current thread can be forcibly released by passing the XNLOCK bit in this mask. In this case, the lock nesting count is also reset to zero.
<i>setmask</i>	The new thread mode. The following flags may be set in this bitmask:

- XNLOCK makes *thread* non-preemptible by other threads when running on a CPU. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in. If *thread* is current, the scheduler is immediately locked, otherwise such lock will take effect next time *thread* resumes on a CPU.
- XNTRAPSW causes the thread to receive a SIGDEBUG signal when it switches to secondary mode. This is a debugging aid for detecting spurious relaxes.

- XNTRAPLB disallows breaking the scheduler lock. In the default case, a thread which holds the scheduler lock is allowed to drop it temporarily for sleeping. If this mode bit is set, such thread would return immediately with XNBREAK set from [xnthread_suspend\(\)](#).

Remarks

Tags: might-switch.

Note

Setting *clrmask* and *setmask* to zero leads to a nop, only returning the previous mode if *mode_r* is a valid address.

References XNLOCK.

5.42.2.6 `int xnthread_set_periodic (struct xnthread * thread, xnticks_t idate, xntmode_t timeout_mode, xnticks_t period)`

Make a thread periodic.

Make a thread periodic by programming its first release point and its period in the processor time line. Subsequent calls to [xnthread_wait_period\(\)](#) will delay the thread until the next periodic release point in the processor timeline is reached.

Parameters

<i>thread</i>	The core thread to make periodic.
<i>idate</i>	The initial (absolute) date of the first release point, expressed in nanoseconds. The affected thread will be delayed by the first call to xnthread_wait_period() until this point is reached. If <i>idate</i> is equal to XN_INFINITE, the current system date is used, and no initial delay takes place. In the latter case, <i>timeout_mode</i> is not considered and can have any valid value.
<i>timeout_mode</i>	The mode of the <i>idate</i> parameter. It can either be set to XN_ABSOLUTE or XN_REALTIME with <i>idate</i> different from XN_INFINITE (see also xntimer_start()).
<i>period</i>	The period of the thread, expressed in nanoseconds. As a side-effect, passing XN_INFINITE attempts to stop the thread's periodic timer; in the latter case, the routine always exits successfully, regardless of the previous state of this timer.

Returns

0 is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *idate* is different from XN_INFINITE and represents a date in the past.
- -EINVAL is returned if *period* is different from XN_INFINITE but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`. -EINVAL is also returned if *timeout_mode* is not compatible with *idate*, such as XN_RELATIVE with *idate* different from XN_INFINITE.

Remarks

Tags: none.

References [xntimer_start\(\)](#), and [xntimer_stop\(\)](#).

Referenced by [rtdm_task_init\(\)](#).

```
5.42.2.7 int xnthread_set_schedparam ( struct xnthread * thread, struct xnsched_class *
      sched_class, const union xnsched_policy_param * sched_param )
```

Change the base scheduling parameters of a thread.

Changes the base scheduling policy and parameters of a thread. If the thread is currently blocked, waiting in priority-pending mode (XNSYNCH_PRIO) for a synchronization object to be signaled, the nucleus will attempt to reorder the object's wait queue so that it reflects the new sleeper's priority, unless the XNSYNCH_DREORD flag has been set for the pended object.

Parameters

<i>thread</i>	The descriptor address of the affected thread. See note.
<i>sched_class</i>	The new scheduling class the thread should be assigned to.
<i>sched_param</i>	The scheduling parameters to set for the thread; <i>sched_param</i> must be valid within the context of <i>sched_class</i> .

It is absolutely required to use this service to change a thread priority, in order to have all the needed housekeeping chores correctly performed. i.e. Do *not* call `xnsched_set_policy()` directly or worse, change the `thread.cprio` field by hand in any case.

Returns

0 is returned on success. Otherwise, a negative error code indicates the cause of a failure that happened in the scheduling class implementation for *sched_class*. Invalid parameters passed into *sched_param* are common causes of error.

Side-effects:

- This service does not call the rescheduling procedure but may affect the state of the runnable queue for the previous and new scheduling classes.
- Assigning the same scheduling class and parameters to a running or ready thread moves it to the end of the runnable queue, thus causing a manual round-robin.

Remarks

Tags: none.

Note

The changes only apply to the Xenomai scheduling parameters for *thread*. There is no propagation/translation of such changes to the Linux scheduler for the task mated to the Xenomai target thread.

```
5.42.2.8 int xnthread_set_slice ( struct xnthread * thread, xnticks_t quantum )
```

Set thread time-slicing information.

Update the time-slicing information for a given thread. This service enables or disables round-robin scheduling for the thread, depending on the value of *quantum*. By default, times-slicing is disabled for a new thread initialized by a call to [xnthread_init\(\)](#).

Parameters

<i>thread</i>	The descriptor address of the affected thread.
<i>quantum</i>	The time quantum assigned to the thread expressed in nanoseconds. If <i>quantum</i> is different from XN_INFINITE, the time-slice for the thread is set to that value and its current time credit is refilled (i.e. the thread is given a full time-slice to run next). Otherwise, if <i>quantum</i> equals XN_INFINITE, time-slicing is stopped for that thread.

Returns

0 is returned upon success. Otherwise, -EINVAL is returned if *quantum* is not XN_INFINITE and:

- the base scheduling class of the target thread does not support time-slicing,
- *quantum* is smaller than the master clock gravity, which denotes a spurious value.

Remarks

Tags: none.

References xnsched::curr, xnsched::rrbtimer, XNRRB, xntimer_start(), and xntimer_stop().

5.42.2.9 int xnthread_start (struct xnthread * thread, const struct xnthread_start_attr * attr)

Start a newly created thread.

Starts a (newly) created thread, scheduling it for the first time. This call releases the target thread from the XNDORMANT state. This service also sets the initial mode for the new thread.

Parameters

<i>thread</i>	The descriptor address of the started thread which must have been previously initialized by a call to xnthread_init() .
<i>attr</i>	A pointer to an attribute block describing the execution properties of the new thread. Members of this structure are defined as follows:

- mode: The initial thread mode. The following flags can be part of this bitmask, each of them affecting the nucleus behaviour regarding the started thread:
 - XNLOCK causes the thread to lock the scheduler when it starts. The target thread will have to call the xnsched_unlock() service to unlock the scheduler. A non-preemptible thread may still block, in which case, the lock is reasserted when the thread is scheduled back in.
 - XNSUSP makes the thread start in a suspended state. In such a case, the thread will have to be explicitly resumed using the [xnthread_resume\(\)](#) service for its execution to actually begin.
- entry: The address of the thread's body routine. In other words, it is the thread entry point.
- cookie: A user-defined opaque cookie the nucleus will pass to the emerging thread as the sole argument of its entry point.

Return values

0	if <i>thread</i> could be started ;
-EBUSY	if <i>thread</i> was not dormant or stopped ;

Remarks

Tags: might-switch.

References XNDORMANT, xnsched_run(), XNSUSP, and xnthread_resume().

Referenced by rtdm_task_init(), and xnshadow_map_user().

5.42.2.10 void `xnthread_suspend` (struct `xnthread` * `thread`, int `mask`, `xnticks_t` `timeout`, `xntmode_t` `timeout_mode`, struct `xnsynch` * `wchan`)

Suspend a thread.

Suspends the execution of a thread according to a given suspensive condition. This thread will not be eligible for scheduling until it all the pending suspensive conditions set by this service are removed by one or more calls to [xnthread_resume\(\)](#).

Parameters

<i>thread</i>	The descriptor address of the suspended thread.
<i>mask</i>	The suspension mask specifying the suspensive condition to add to the thread's wait mask. Possible values usable by the caller are:

- XNSUSP. This flag forcibly suspends a thread, regardless of any resource to wait for. A reverse call to [xnthread_resume\(\)](#) specifying the XNSUSP bit must be issued to remove this condition, which is cumulative with other suspension bits. *wchan* should be NULL when using this suspending mode.
- XNDELAY. This flag denotes a counted delay wait (in ticks) which duration is defined by the value of the *timeout* parameter.
- XNPEND. This flag denotes a wait for a synchronization object to be signaled. The *wchan* argument must point to this object. A *timeout* value can be passed to bound the wait. This suspending mode should not be used directly by the client interface, but rather through the [xnsynch_sleep_on\(\)](#) call.

Parameters

<i>timeout</i>	The timeout which may be used to limit the time the thread pends on a resource. This value is a wait time given in nanoseconds. It can either be relative, absolute monotonic, or absolute adjustable depending on <i>timeout_mode</i> .
----------------	--

Passing XN_INFINITE **and** setting *timeout_mode* to XN_RELATIVE specifies an unbounded wait. All other values are used to initialize a watchdog timer. If the current operation mode of the system timer is oneshot and *timeout* elapses before [xnthread_suspend\(\)](#) has completed, then the target thread will not be suspended, and this routine leads to a null effect.

Parameters

<i>timeout_mode</i>	The mode of the <i>timeout</i> parameter. It can either be set to XN_RELATIVE, XN_ABSOLUTE, or XN_REALTIME (see also xntimer_start()).
<i>wchan</i>	The address of a pended resource. This parameter is used internally by the synchronization object implementation code to specify on which object the suspended thread pends. NULL is a legitimate value when this parameter does not apply to the current suspending mode (e.g. XNSUSP).

Note

If the target thread has received a Linux-originated signal, then this service immediately exits without suspending the thread, but raises the XNBREAK condition in its information mask.

Remarks

Tags: isr-allowed, might-switch.

References `xnsched::curr`, `xnsched::lflags`, XNBREAK, XNDELAY, XNHELD, XNKICKED, XNLOCK, XNREADY, XNRELAX, XNRMID, XNROBBED, `xnsched_run()`, XNSUSP, `xnsynch_forget_sleeper()`, XN-TIMEO, `xntimer_start()`, XNTRAPLB, XNUSER, and XNWAKEN.

Referenced by `xnshadow_map_kernel()`, `xnshadow_map_user()`, `xnshadow_relax()`, `xnsynch_acquire()`, `xnsynch_sleep_on()`, and `xnthread_wait_period()`.

5.42.2.11 `int xntthread_unblock (struct xntthread * thread)`

Unblock a thread.

Breaks the thread out of any wait it is currently in. This call removes the XNDELAY and XNPENDING suspensive conditions previously put by [xntthread_suspend\(\)](#) on the target thread. If all suspensive conditions are gone, the thread is left in a READY state at which point it becomes eligible anew for scheduling.

Parameters

<i>thread</i>	The descriptor address of the unblocked thread.
---------------	---

This call neither releases the thread from the XNSUSP, XNRELAX, XNDORMANT or XNHOLD suspensive conditions.

When the thread resumes execution, the XNBREAK bit is set in the unblocked thread's information mask. Unblocking a non-blocked thread is perfectly harmless.

Returns

non-zero is returned if the thread was actually unblocked from a pending wait state, 0 otherwise.

Remarks

Tags: isr-allowed.

References XNBREAK, XNDELAY, XNPEND, and [xntthread_resume\(\)](#).

5.42.2.12 `int xntthread_wait_period (unsigned long * overruns_r)`

Wait for the next periodic release point.

Make the current thread wait for the next periodic release point in the processor time line.

Parameters

<i>overruns_r</i>	If non-NULL, <i>overruns_r</i> must be a pointer to a memory location which will be written with the count of pending overruns. This value is copied only when xntthread_wait_period() returns -ETIMEDOUT or success; the memory location remains unmodified otherwise. If NULL, this count will never be copied back.
-------------------	--

Returns

0 is returned upon success; if *overruns_r* is valid, zero is copied to the pointed memory location. Otherwise:

- -EWOULDBLOCK is returned if [xntthread_set_periodic\(\)](#) has not previously been called for the calling thread.
- -EINTR is returned if [xntthread_unblock\(\)](#) has been called for the waiting thread before the next periodic release point has been reached. In this case, the overrun counter is reset too.
- -ETIMEDOUT is returned if the timer has overrun, which indicates that one or more previous release points have been missed by the calling thread. If *overruns_r* is valid, the count of pending overruns is copied to the pointed memory location.

Remarks

Tags: primary-only, might-sleep.

References XNBREAK, XNDELAY, [xntthread_suspend\(\)](#), and [xntimer_get_overruns\(\)](#).

5.43 Timer services

The Xenomai timer facility depends on a clock source (xnclock) for scheduling the next activation times. Collaboration diagram for Timer services:



Functions

- static `xnticks_t xntimer_interval` (struct `xntimer` *timer)
Return the timer interval value.
- int `xntimer_start` (struct `xntimer` *timer, `xnticks_t` value, `xnticks_t` interval, `xntmode_t` mode)
Arm a timer.
- `xnticks_t xntimer_get_date` (struct `xntimer` *timer)
Return the absolute expiration date.
- `xnticks_t xntimer_get_timeout` (struct `xntimer` *timer)
Return the relative expiration date.
- static void `xntimer_stop` (struct `xntimer` *timer)
Disarm a timer.
- void `xntimer_init` (struct `xntimer` *timer, struct `xnclock` *clock, void(*handler)(struct `xntimer` *timer), struct `xnthread` *thread)
Initialize a timer object.
- void `xntimer_destroy` (struct `xntimer` *timer)
Release a timer object.
- unsigned long long `xntimer_get_overruns` (struct `xntimer` *timer, `xnticks_t` now)
Get the count of overruns for the last tick.
- static int `program_htick_shot` (unsigned long delay, struct `clock_event_device` *cdev)
Program next host tick as a Xenomai timer event.
- static void `switch_htick_mode` (enum `clock_event_mode` mode, struct `clock_event_device` *cdev)
Tick mode switch emulation callback.
- int `xntimer_grab_hardware` (int cpu)
Grab the hardware timer.
- void `xntimer_release_hardware` (int cpu)
Release the hardware timer.

5.43.1 Detailed Description

The Xenomai timer facility depends on a clock source (xnclock) for scheduling the next activation times. The core provides and depends on a monotonic clock source (nkclock) with nanosecond resolution, driving the platform timer hardware exposed by the interrupt pipeline.

5.43.2 Function Documentation

5.43.2.1 `static int program_htick_shot (unsigned long delay, struct clock_event_device * cdev)`
`[static]`

Program next host tick as a Xenomai timer event.

Program the next shot for the host tick on the current CPU. Emulation is done using a nucleus timer attached to the master timebase.

Parameters

<i>delay</i>	The time delta from the current date to the next tick, expressed as a count of nanoseconds.
<i>cdev</i>	An pointer to the clock device which notifies us.

Environment:

This routine is a callback invoked from the kernel's clock event handlers.

Note

GENERIC_CLOCKEVENTS is required from the host kernel.

Remarks

Tags: none.

References `xnsched::htimer`, and `xntimer_start()`.

Referenced by `xntimer_grab_hardware()`.

5.43.2.2 `void switch_htick_mode (enum clock_event_mode mode, struct clock_event_device * cdev)`
`[static]`

Tick mode switch emulation callback.

Changes the host tick mode for the tick device of the current CPU.

Parameters

<i>mode</i>	The new mode to switch to. The possible values are:
-------------	---

- `CLOCK_EVT_MODE_ONESHOT`, for a switch to oneshot mode.
- `CLOCK_EVT_MODE_PERIODIC`, for a switch to periodic mode. The current implementation for the generic clockevent layer Linux exhibits should never downgrade from a oneshot to a periodic tick mode, so this mode should not be encountered. This said, the associated code is provided, basically for illustration purposes.
- `CLOCK_EVT_MODE_SHUTDOWN`, indicates the removal of the current tick device. Normally, the nucleus only interposes on tick devices which should never be shut down, so this mode should not be encountered.

Parameters

<i>cdev</i>	An opaque pointer to the clock device which notifies us.
-------------	--

Remarks

Tags: none.

Note

GENERIC_CLOCKEVENTS is required from the host kernel.

References xnsched::htimer, xntimer_start(), and xntimer_stop().

Referenced by xntimer_grab_hardware().

5.43.2.3 void xntimer_destroy (struct xntimer * timer)

Release a timer object.

Destroys a timer. After it has been destroyed, all resources associated with the timer have been released. The timer is automatically deactivated before deletion if active on entry.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
--------------	--

Remarks

Tags: none.

References xntimer_stop().

Referenced by rtdm_timer_destroy().

5.43.2.4 xnticks_t xntimer_get_date (struct xntimer * timer)

Return the absolute expiration date.

Return the next expiration date of a timer as an absolute count of nanoseconds.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
--------------	--

Returns

The expiration date in nanoseconds. The special value XN_INFINITE is returned if *timer* is currently disabled.

Remarks

Tags: isr-allowed.

5.43.2.5 unsigned long long xntimer_get_overruns (struct xntimer * timer, xnticks_t now)

Get the count of overruns for the last tick.

This service returns the count of pending overruns for the last tick of a given timer, as measured by the difference between the expected expiry date of the timer and the date *now* passed as argument.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
<i>now</i>	current date (as <code>xnclock_read_raw(xntimer_clock(timer))</code>)

Returns

the number of overruns of *timer* at date *now*

Remarks

Tags: irqs-off, isr-allowed.

Referenced by `xnthread_wait_period()`.

5.43.2.6 `xnticks_t xntimer_get_timeout (struct xntimer * timer)`

Return the relative expiration date.

This call returns the count of nanoseconds remaining until the timer expires.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
--------------	--

Returns

The count of nanoseconds until expiry. The special value `XN_INFINITE` is returned if *timer* is currently disabled. It might happen that the timer expires when this service runs (even if the associated handler has not been fired yet); in such a case, 1 is returned.

Remarks

Tags: irqs-off, isr-allowed.

5.43.2.7 `int xntimer_grab_hardware (int cpu)`

Grab the hardware timer.

[xntimer_grab_hardware\(\)](#) grabs and tunes the hardware timer in oneshot mode in order to clock the master time base. `GENERIC_CLOCKEVENTS` is required from the host kernel.

Host tick emulation is performed for sharing the clockchip hardware between Linux and Xenomai, when the former provides support for oneshot timing (i.e. high resolution timers and no-HZ scheduler ticking).

Parameters

<i>cpu</i>	The CPU number to grab the timer from.
------------	--

Returns

a positive value is returned on success, representing the duration of a Linux periodic tick expressed as a count of nanoseconds; zero should be returned when the Linux kernel does not undergo periodic timing on the given CPU (e.g. oneshot mode). Otherwise:

- `-EBUSY` is returned if the hardware timer has already been grabbed. [xntimer_release_hardware\(\)](#) must be issued before [xntimer_grab_hardware\(\)](#) is called again.
- `-ENODEV` is returned if the hardware timer cannot be used. This situation may occur after the kernel disabled the timer due to invalid calibration results; in such a case, such hardware is unusable for any timing duties.

Remarks

Tags: secondary-only.

References `program_htick_shot()`, and `switch_htick_mode()`.

```
5.43.2.8 void xntimer_init ( struct xntimer * timer, struct xnclock * clock, void(*) (struct xntimer * timer)
        handler, struct xnthread * thread )
```

Initialize a timer object.

Creates a timer. When created, a timer is left disarmed; it must be started using `xntimer_start()` in order to be activated.

Parameters

<i>timer</i>	The address of a timer descriptor the nucleus will use to store the object-specific data. This descriptor must always be valid while the object is active therefore it must be allocated in permanent memory.
<i>clock</i>	The clock the timer relates to. Xenomai defines a monotonic system clock, with nanosecond resolution, named <code>nkclock</code> . In addition, external clocks driven by other tick sources may be created dynamically if <code>CONFIG_XENO_OPT_EXTCLOCK</code> is defined.
<i>handler</i>	The routine to call upon expiration of the timer.
<i>thread</i>	The optional thread object the new timer is affine to. If non-NULL, the timer will fire on the same CPU <i>thread</i> currently runs on by default. A call to <code>xntimer_set_sched()</code> may change this setting.

There is no limitation on the number of timers which can be created/active concurrently.

Remarks

Tags: none.

```
5.43.2.9 xnticks_t xntimer_interval ( struct xntimer * timer ) [inline], [static]
```

Return the timer interval value.

Return the timer interval value in nanoseconds.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
--------------	--

Returns

The duration of a period in nanoseconds. The special value `XN_INFINITE` is returned if *timer* is currently disabled or one shot.

Remarks

Tags: isr-allowed.

```
5.43.2.10 void xntimer_release_hardware ( int cpu )
```

Release the hardware timer.

Releases the hardware timer, thus reverting the effect of a previous call to `xntimer_grab_hardware()`. In case the timer hardware is shared with Linux, a periodic setup suitable for the Linux kernel is reset.

Parameters

<i>cpu</i>	The CPU number the timer was grabbed from.
------------	--

Remarks

Tags: secondary-only.

5.43.2.11 `int xntimer_start (struct xntimer * timer, xnticks_t value, xnticks_t interval, xntmode_t mode)`

Arm a timer.

Activates a timer so that the associated timeout handler will be fired after each expiration time. A timer can be either periodic or one-shot, depending on the reload value passed to this routine. The given timer must have been previously initialized.

A timer is attached to the clock specified in [xntimer_init\(\)](#).

Parameters

<i>timer</i>	The address of a valid timer descriptor.
<i>value</i>	The date of the initial timer shot, expressed in nanoseconds.
<i>interval</i>	The reload value of the timer. It is a periodic interval value to be used for reprogramming the next timer shot, expressed in nanoseconds. If <i>interval</i> is equal to <code>XN_INFINITE</code> , the timer will not be reloaded after it has expired.
<i>mode</i>	The timer mode. It can be <code>XN_RELATIVE</code> if <i>value</i> shall be interpreted as a relative date, <code>XN_ABSOLUTE</code> for an absolute date based on the monotonic clock of the related time base (as returned by <code>xnclock_read_monotonic()</code>), or <code>XN_REALTIME</code> if the absolute date is based on the adjustable real-time date for the relevant clock (obtained from <code>xnclock_read_realtime()</code>).

Returns

0 is returned upon success, or `-ETIMEDOUT` if an absolute date in the past has been given.

Remarks

Tags: atomic-entry.

Referenced by `program_htick_shot()`, `rtdm_timer_start()`, `switch_htick_mode()`, `xnthread_set_periodic()`, `xnthread_set_slice()`, and `xnthread_suspend()`.

5.43.2.12 `int xntimer_stop (struct xntimer * timer) [inline], [static]`

Disarm a timer.

This service deactivates a timer previously armed using [xntimer_start\(\)](#). Once disarmed, the timer can be subsequently re-armed using the latter service.

Parameters

<i>timer</i>	The address of a valid timer descriptor.
--------------	--

Remarks

Tags: atomic-entry, isr-allowed.

Referenced by `rtdm_timer_stop()`, `switch_htick_mode()`, `xnthread_resume()`, `xnthread_set_periodic()`, `xnthread_set_slice()`, and `xntimer_destroy()`.

5.44 Virtual file services

Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces.

Collaboration diagram for Virtual file services:



Data Structures

- struct [xnvfile_lock_ops](#)
Vfile locking operations.
- struct [xnvfile_regular_ops](#)
Regular vfile operation descriptor.
- struct [xnvfile_regular_iterator](#)
Regular vfile iterator.
- struct [xnvfile_snapshot_ops](#)
Snapshot vfile operation descriptor.
- struct [xnvfile_rev_tag](#)
Snapshot revision tag.
- struct [xnvfile_snapshot](#)
Snapshot vfile descriptor.
- struct [xnvfile_snapshot_iterator](#)
Snapshot-driven vfile iterator.

Functions

- int [xnvfile_init_snapshot](#) (const char *name, struct [xnvfile_snapshot](#) *vfile, struct [xnvfile_directory](#) *parent)
Initialize a snapshot-driven vfile.
- int [xnvfile_init_regular](#) (const char *name, struct [xnvfile_regular](#) *vfile, struct [xnvfile_directory](#) *parent)
Initialize a regular vfile.
- int [xnvfile_init_dir](#) (const char *name, struct [xnvfile_directory](#) *vdir, struct [xnvfile_directory](#) *parent)
Initialize a virtual directory entry.
- int [xnvfile_init_link](#) (const char *from, const char *to, struct [xnvfile_link](#) *vlink, struct [xnvfile_directory](#) *parent)
Initialize a virtual link entry.
- void [xnvfile_destroy](#) (struct [xnvfile](#) *vfile)
Removes a virtual file entry.
- ssize_t [xnvfile_get_blob](#) (struct [xnvfile_input](#) *input, void *data, size_t size)
Read in a data bulk written to the vfile.
- ssize_t [xnvfile_get_string](#) (struct [xnvfile_input](#) *input, char *s, size_t maxlen)

Read in a C-string written to the vfile.

- `ssize_t xnvfile_get_integer` (struct `xnvfile_input` *input, long *valp)

Evaluate the string written to the vfile as a long integer.

Variables

- struct `xnvfile_directory` `nkvfroot`

Xenomai vfile root directory.

- struct `xnvfile_directory` `nkvfroot`

Xenomai vfile root directory.

5.44.1 Detailed Description

Virtual files provide a mean to export Xenomai object states to user-space, based on common kernel interfaces. This encapsulation is aimed at:

- supporting consistent collection of very large record-based output, without incurring latency peaks for undergoing real-time activities.
- in the future, hiding discrepancies between linux kernel releases, regarding the proper way to export kernel object states to userland, either via the `/proc` interface or by any other mean.

This virtual file implementation offers record-based read support based on `seq_files`, single-buffer write support, directory and link handling, all visible from the `/proc` namespace.

The vfile support exposes four filesystem object types:

- snapshot-driven file (struct `xnvfile_snapshot`). This is commonly used to export real-time object states via the `/proc` filesystem. To minimize the latency involved in protecting the vfile routines from changes applied by real-time code on such objects, a snapshot of the data to output is first taken under proper locking, before the collected data is formatted and sent out in a lockless manner.

Because a large number of records may have to be output, the data collection phase is not strictly atomic as a whole, but only protected at record level. The vfile implementation can be notified of updates to the underlying data set, and restart the collection from scratch until the snapshot is fully consistent.

- regular sequential file (struct `xnvfile_regular`). This is basically an encapsulated sequential file object as available from the host kernel (i.e. `seq_file`), with a few additional features to make it more handy in a Xenomai environment, like implicit locking support and shortened declaration for simplest, single-record output.
- virtual link (struct `xnvfile_link`). This is a symbolic link feature integrated with the vfile semantics. The link target is computed dynamically at creation time from a user-given helper routine.
- virtual directory (struct `xnvfile_directory`). A directory object, which can be used to create a hierarchy for ordering a set of vfile objects.

5.44.2 Function Documentation

5.44.2.1 `void xnvfile_destroy (struct xnvfile * vfile)`

Removes a virtual file entry.

Parameters

<i>vfile</i>	A pointer to the virtual file descriptor to remove.
--------------	---

Remarks

Tags: secondary-only.

5.44.2.2 `ssize_t xnvfile_get_blob (struct xnvfile_input * input, void * data, size_t size)`

Read in a data bulk written to the vfile.

When writing to a vfile, the associated store() handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_blob()` retrieves this data as an untyped binary blob, and copies it back to the caller's buffer.

Parameters

<i>input</i>	A pointer to the input descriptor passed to the store() handler.
<i>data</i>	The address of the destination buffer to copy the input data to.
<i>size</i>	The maximum number of bytes to copy to the destination buffer. If <i>size</i> is larger than the actual data size, the input is truncated to <i>size</i> .

Returns

The number of bytes read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- -EFAULT indicates an invalid source buffer address.

Remarks

Tags: secondary-only.

Referenced by `xnvfile_get_integer()`, and `xnvfile_get_string()`.

5.44.2.3 `ssize_t xnvfile_get_integer (struct xnvfile_input * input, long * valp)`

Evaluate the string written to the vfile as a long integer.

When writing to a vfile, the associated store() handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_integer()` retrieves and interprets this data as a long integer, and copies the resulting value back to *valp*.

The long integer can be expressed in decimal, octal or hexadecimal bases depending on the prefix found.

Parameters

<i>input</i>	A pointer to the input descriptor passed to the store() handler.
<i>valp</i>	The address of a long integer variable to receive the value.

Returns

The number of characters read while evaluating the input as a long integer upon success. Otherwise, a negative error code is returned:

- -EINVAL indicates a parse error on the input stream; the written text cannot be evaluated as a long integer.
- -EFAULT indicates an invalid source buffer address.

Remarks

Tags: secondary-only.

References `xnvfile_get_blob()`.

5.44.2.4 `ssize_t xnvfile_get_string (struct xnvfile_input * input, char * s, size_t maxlen)`

Read in a C-string written to the vfile.

When writing to a vfile, the associated `store()` handler from the [snapshot-driven vfile](#) or [regular vfile](#) is called, with a single argument describing the input data. `xnvfile_get_string()` retrieves this data as a null-terminated character string, and copies it back to the caller's buffer.

Parameters

<i>input</i>	A pointer to the input descriptor passed to the <code>store()</code> handler.
<i>s</i>	The address of the destination string buffer to copy the input data to.
<i>maxlen</i>	The maximum number of bytes to copy to the destination buffer, including the ending null character. If <i>maxlen</i> is larger than the actual string length, the input is truncated to <i>maxlen</i> .

Returns

The number of characters read and copied to the destination buffer upon success. Otherwise, a negative error code is returned:

- `-EFAULT` indicates an invalid source buffer address.

Remarks

Tags: secondary-only.

References `xnvfile_get_blob()`.

5.44.2.5 `int xnvfile_init_dir (const char * name, struct xnvfile_directory * vdir, struct xnvfile_directory * parent)`

Initialize a virtual directory entry.

Parameters

<i>name</i>	The name which should appear in the pseudo-filesystem, identifying the vdir entry.
<i>vdir</i>	A pointer to the virtual directory descriptor to initialize.
<i>parent</i>	A pointer to a virtual directory descriptor standing for the parent directory of the new vdir. If NULL, the <code>/proc</code> root directory will be used. <code>/proc/xenomai</code> is mapped on the globally available <code>nkvfroot</code> vdir.

Returns

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if the virtual directory entry cannot be created in the `/proc` hierarchy.

Remarks

Tags: secondary-only.

5.44.2.6 `int xnvfile_init_link (const char * from, const char * to, struct xnvfile_link * vlink, struct xnvfile_directory * parent)`

Initialize a virtual link entry.

Parameters

<i>from</i>	The name which should appear in the pseudo-filesystem, identifying the vlink entry.
<i>to</i>	The target file name which should be referred to symbolically by <i>name</i> .
<i>vlink</i>	A pointer to the virtual link descriptor to initialize.
<i>parent</i>	A pointer to a virtual directory descriptor standing for the parent directory of the new vlink. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available <i>nkvfroot</i> vdir.

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual link entry cannot be created in the /proc hierarchy.

Remarks

Tags: secondary-only.

5.44.2.7 `int xnvfile_init_regular (const char * name, struct xnvfile_regular * vfile, struct xnvfile_directory * parent)`

Initialize a regular vfile.

Parameters

<i>name</i>	The name which should appear in the pseudo-filesystem, identifying the vfile entry.
<i>vfile</i>	A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- .privsz is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- entry.lockops is a pointer to a [locking descriptor](#)", defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case no locking will be applied.
- .ops is a pointer to an [operation descriptor](#).

Parameters

<i>parent</i>	A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the /proc root directory will be used. /proc/xenomai is mapped on the globally available <i>nkvfroot</i> vdir.
---------------	---

Returns

0 is returned on success. Otherwise:

- -ENOMEM is returned if the virtual file entry cannot be created in the /proc hierarchy.

Remarks

Tags: secondary-only.

5.44.2.8 `int xnvfile_init_snapshot (const char * name, struct xnvfile_snapshot * vfile, struct xnvfile_directory * parent)`

Initialize a snapshot-driven vfile.

Parameters

<i>name</i>	The name which should appear in the pseudo-filesystem, identifying the vfile entry.
<i>vfile</i>	A pointer to a vfile descriptor to initialize from. The following fields in this structure should be filled in prior to call this routine:

- `.privsz` is the size (in bytes) of the private data area to be reserved in the [vfile iterator](#). A NULL value indicates that no private area should be reserved.
- `.datasz` is the size (in bytes) of a single record to be collected by the [next\(\) handler](#) from the [operation descriptor](#).
- `.tag` is a pointer to a mandatory vfile revision tag structure (struct [xnvmfile_rev_tag](#)). This tag will be monitored for changes by the vfile core while collecting data to output, so that any update detected will cause the current snapshot data to be dropped, and the collection to restart from the beginning. To this end, any change to the data which may be part of the collected records, should also invoke `xnvmfile_touch()` on the associated tag.
- `entry.lockops` is a pointer to a [locking descriptor](#), defining the lock and unlock operations for the vfile. This pointer may be left to NULL, in which case the operations on the nucleus lock (i.e. `nklock`) will be used internally around calls to data collection handlers (see [operation descriptor](#)).
- `.ops` is a pointer to an [operation descriptor](#).

Parameters

<i>parent</i>	A pointer to a virtual directory descriptor; the vfile entry will be created into this directory. If NULL, the <code>/proc</code> root directory will be used. <code>/proc/xenomai</code> is mapped on the globally available <code>nkvmfroot</code> vdir.
---------------	--

Returns

0 is returned on success. Otherwise:

- `-ENOMEM` is returned if the virtual file entry cannot be created in the `/proc` hierarchy.

Remarks

Tags: secondary-only.

References `xnvmfile_snapshot_ops::store`.

5.44.3 Variable Documentation

5.44.3.1 struct `xnvmfile_directory` `nkvmfroot`

Xenomai vfile root directory.

This vdir maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

5.44.3.2 struct `xnvmfile_directory` `nkvmfroot`


Xenomai vfile root directory.

This vdir maps the `/proc/xenomai` directory. It can be used to create a hierarchy of Xenomai-related vfiles under this root.

5.45 Analogy interface

A RTDM-based interface for implementing DAQ card drivers.

Collaboration diagram for Analogy interface:



Modules

- [Driver API](#)
Programming interface provided to DAQ card drivers.
- [Analogy user API](#)

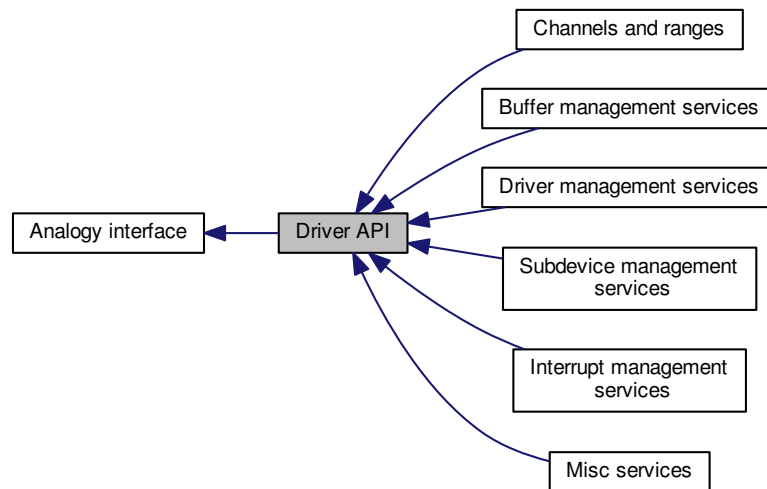
5.45.1 Detailed Description

A RTDM-based interface for implementing DAQ card drivers.

5.46 Driver API

Programming interface provided to DAQ card drivers.

Collaboration diagram for Driver API:



Modules

- [Channels and ranges](#)
Channels.
- [Driver management services](#)
Analogy driver registration / unregistration.
- [Subdevice management services](#)
Subdevice declaration in a driver.
- [Buffer management services](#)
Buffer management services.
- [Interrupt management services](#)
- [Misc services](#)


5.46.1 Detailed Description

Programming interface provided to DAQ card drivers.

5.47 Driver management services

Analogy driver registration / unregistration.

Collaboration diagram for Driver management services:



Functions

- int [a4l_register_drv](#) (struct [a4l_driver](#) *drv)
Register an Analogy driver.
- int [a4l_unregister_drv](#) (struct [a4l_driver](#) *drv)
Unregister an Analogy driver.

5.47.1 Detailed Description

Analogy driver registration / unregistration. In a common Linux char driver, the developer has to register a fops structure filled with callbacks for read / write / mmap / ioctl operations.

Analogy drivers do not have to implement read / write / mmap / ioctl functions, these procedures are implemented in the Analogy generic layer. Then, the transfers between user-space and kernel-space are already managed. Analogy drivers work with commands and instructions which are some kind of more dedicated read / write operations. And, instead of registering a fops structure, a Analogy driver must register some [a4l_driver](#) structure.

5.47.2 Function Documentation

5.47.2.1 int [a4l_register_drv](#) (struct **a4l_driver** * drv)

Register an Analogy driver.

After initialising a driver structure, the driver must be made available so as to be attached.

Parameters

in	<i>drv</i>	Driver descriptor structure
----	------------	-----------------------------

Returns

0 on success, otherwise negative error code.

References [a4l_driver::board_name](#), and [a4l_driver::list](#).

5.47.2.2 int [a4l_unregister_drv](#) (struct **a4l_driver** * drv)

Unregister an Analogy driver.

This function removes the driver descriptor from the Analogy driver list. The driver cannot be attached anymore.

Parameters

<i>in</i>	<i>drv</i>	Driver descriptor structure
-----------	------------	-----------------------------

Returns


0 on success, otherwise negative error code.

References `a4l_driver::board_name`, and `a4l_driver::list`.

5.48 Subdevice management services

Subdevice declaration in a driver.

Collaboration diagram for Subdevice management services:



Functions

- struct `a4l_subdevice` * `a4l_alloc_subd` (int sizeof_priv, void(*setup)(struct `a4l_subdevice` *))
Allocate a subdevice descriptor.
- int `a4l_add_subd` (struct `a4l_device` *dev, struct `a4l_subdevice` *subd)
Add a subdevice to the driver descriptor.
- struct `a4l_subdevice` * `a4l_get_subd` (struct `a4l_device` *dev, int idx)
Get a pointer to the subdevice descriptor referenced by its registration index.

Subdevices types

Flags to define the subdevice type

- #define `A4L_SUBD_UNUSED` (`A4L_SUBD_MASK_SPECIAL|0x1`)
Unused subdevice.
- #define `A4L_SUBD_AI` (`A4L_SUBD_MASK_READ|0x2`)
Analog input subdevice.
- #define `A4L_SUBD_AO` (`A4L_SUBD_MASK_WRITE|0x4`)
Analog output subdevice.
- #define `A4L_SUBD_DI` (`A4L_SUBD_MASK_READ|0x8`)
Digital input subdevice.
- #define `A4L_SUBD_DO` (`A4L_SUBD_MASK_WRITE|0x10`)
Digital output subdevice.
- #define `A4L_SUBD_DIO` (`A4L_SUBD_MASK_SPECIAL|0x20`)
Digital input/output subdevice.
- #define `A4L_SUBD_COUNTER` (`A4L_SUBD_MASK_SPECIAL|0x40`)
Counter subdevice.
- #define `A4L_SUBD_TIMER` (`A4L_SUBD_MASK_SPECIAL|0x80`)
Timer subdevice.
- #define `A4L_SUBD_MEMORY` (`A4L_SUBD_MASK_SPECIAL|0x100`)
Memory, EEPROM, DPRAM.
- #define `A4L_SUBD_CALIB` (`A4L_SUBD_MASK_SPECIAL|0x200`)
Calibration subdevice DACs.
- #define `A4L_SUBD_PROC` (`A4L_SUBD_MASK_SPECIAL|0x400`)

Processor, DSP.

- `#define A4L_SUBD_SERIAL (A4L_SUBD_MASK_SPECIAL|0x800)`

Serial IO subdevice.

- `#define A4L_SUBD_TYPES`

Mask which gathers all the types.

Subdevice features

Flags to define the subdevice's capabilities

- `#define A4L_SUBD_CMD 0x1000`

The subdevice can handle command (i.e it can perform asynchronous acquisition)

- `#define A4L_SUBD_MMAP 0x8000`

The subdevice support mmap operations (technically, any driver can do it; however, the developer might want that his driver must be accessed through read / write.

Subdevice status

Flags to define the subdevice's status

- `#define A4L_SUBD_BUSY_NR 0`

The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.

- `#define A4L_SUBD_BUSY (1 << A4L_SUBD_BUSY_NR)`

The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.

- `#define A4L_SUBD_CLEAN_NR 1`

The subdevice is about to be cleaned in the middle of the detach procedure.

- `#define A4L_SUBD_CLEAN (1 << A4L_SUBD_CLEAN_NR)`

The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.

5.48.1 Detailed Description

Subdevice declaration in a driver. The subdevice structure is the most complex one in the Analogy driver layer. It contains some description fields to fill and some callbacks to declare.

The description fields are:

- `flags`: to define the subdevice type and its capabilities;
- `chan_desc`: to describe the channels which compose the subdevice;
- `rng_desc`: to declare the usable ranges;

The functions callbakcs are:

- `do_cmd()` and `do_cmdtest()`: to performe asynchronous acquisitions thanks to commands;
- `cancel()`: to abort a working asynchronous acquisition;
- `munge()`: to apply modifications on the data freshly acquired during an asynchronous transfer. Warning: using this feature with can significantly reduce the performances (if the munge operation is complex, it will trigger high CPU charge and if the acquisition device is DMA capable, many cache-misses and cache-replaces will occur (the benefits of the DMA controller will vanish);
- `trigger()`: optionnaly to launch an asynchronous acquisition;

- `insn_read()`, `insn_write()`, `insn_bits()`, `insn_config()`: to perform synchronous acquisition operations.

Once the subdevice is filled, it must be inserted into the driver structure thanks to `a4l_add_subd()`.

5.48.2 Function Documentation

5.48.2.1 `int a4l_add_subd (struct a4l_device * dev, struct a4l_subdevice * subd)`

Add a subdevice to the driver descriptor.

Once the driver descriptor structure is initialized, the function `a4l_add_subd()` must be used so to add some subdevices to the driver.

Parameters

<i>in</i>	<i>dev</i>	Device descriptor structure
<i>in</i>	<i>subd</i>	Subdevice descriptor structure

Returns

the index with which the subdevice has been registered, in case of error a negative error code is returned.

References `a4l_subdevice::dev`, `a4l_subdevice::idx`, and `a4l_subdevice::list`.

5.48.2.2 `struct a4l_subdevice* a4l_alloc_subd (int sizeof_priv, void(*)(struct a4l_subdevice *) setup)`

Allocate a subdevice descriptor.

This is a helper function so as to get a suitable subdevice descriptor

Parameters

<i>in</i>	<i>sizeof_priv</i>	Size of the subdevice's private data
<i>in</i>	<i>setup</i>	Setup function to be called after the allocation

Returns

the index with which the subdevice has been registered, in case of error a negative error code is returned.

References `rtdm_malloc()`.

5.48.2.3 `struct a4l_subdevice* a4l_get_subd (struct a4l_device * dev, int idx)`

Get a pointer to the subdevice descriptor referenced by its registration index.

This function is scarcely useful as all the drivers callbacks get the related subdevice descriptor as first argument. This function is not optimized, it goes through a linked list to get the proper pointer. So it must not be used in real-time context but at initialization / cleanup time (attach / detach).

Parameters

<code>in</code>	<code>dev</code>	Device descriptor structure
<code>in</code>	<code>idx</code>	Subdevice index

Returns

0 on success, otherwise negative error code.

5.49 Buffer management services

Buffer management services.

Collaboration diagram for Buffer management services:



Functions

- `int a4l_buf_prepare_absput (struct a4l_subdevice *subd, unsigned long count)`
Update the absolute count of data sent from the device to the buffer since the start of the acquisition and after the next DMA shot.
- `int a4l_buf_commit_absput (struct a4l_subdevice *subd, unsigned long count)`
Set the absolute count of data which was sent from the device to the buffer since the start of the acquisition and until the last DMA shot.
- `int a4l_buf_prepare_put (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data which is to be sent to the buffer at the next DMA shot.
- `int a4l_buf_commit_put (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data sent to the buffer during the last completed DMA shots.
- `int a4l_buf_put (struct a4l_subdevice *subd, void *bufdata, unsigned long count)`
Copy some data from the device driver to the buffer.
- `int a4l_buf_prepare_absget (struct a4l_subdevice *subd, unsigned long count)`
Update the absolute count of data sent from the buffer to the device since the start of the acquisition and after the next DMA shot.
- `int a4l_buf_commit_absget (struct a4l_subdevice *subd, unsigned long count)`
Set the absolute count of data which was sent from the buffer to the device since the start of the acquisition and until the last DMA shot.
- `int a4l_buf_prepare_get (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data which is to be sent from the buffer to the device at the next DMA shot.
- `int a4l_buf_commit_get (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data sent from the buffer to the device during the last completed DMA shots.
- `int a4l_buf_get (struct a4l_subdevice *subd, void *bufdata, unsigned long count)`
Copy some data from the buffer to the device driver.
- `int a4l_buf_evt (struct a4l_subdevice *subd, unsigned long evts)`
Signal some event(s) to a user-space program involved in some read / write operation.
- `unsigned long a4l_buf_count (struct a4l_subdevice *subd)`
Get the data amount available in the Analogy buffer.
- `struct a4l_cmd_desc * a4l_get_cmd (struct a4l_subdevice *subd)`
Get the current Analogy command descriptor.
- `int a4l_get_chan (struct a4l_subdevice *subd)`
Get the channel index according to its type.

5.49.1 Detailed Description

Buffer management services. The buffer is the key component of the Analogy infrastructure. It manages transfers between the user-space and the Analogy drivers thanks to generic functions which are described hereafter. Thanks to the buffer subsystem, the driver developer does not have to care about the way the user program retrieves or sends data.

To write a classical char driver, the developer has to fill a fops structure so as to provide transfer operations to the user program (read, write, ioctl and mmap if need be).

The Analogy infrastructure manages the whole interface with the userspace; the common read, write, mmap, etc. callbacks are generic Analogy functions. These functions manage (and perform, if need be) transfers between the user-space and an asynchronous buffer thanks to lockless mechanisms.

Consequently, the developer has to use the proper buffer functions in order to write / read acquired data into / from the asynchronous buffer.

Here are listed the functions:

- `a4l_buf_prepare_(abs)put()` and `a4l_buf_commit_(abs)put()`
- `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()`
- [`a4l_buf_put\(\)`](#)
- [`a4l_buf_get\(\)`](#)
- [`a4l_buf_evt\(\)`](#).

The functions count might seem high; however, the developer needs a few of them to write a driver. Having so many functions enables to manage any transfer cases:

- If some DMA controller is available, there is no need to make the driver copy the acquired data into the asynchronous buffer, the DMA controller must directly trigger DMA shots into / from the buffer. In that case, a function `a4l_buf_prepare_*()` must be used so as to set up the DMA transfer and a function `a4l_buf_commit_*()` has to be called to complete the transfer().
- For DMA controllers which need to work with global counter (the transferred data count since the beginning of the acquisition), the functions `a4l_buf_*_abs_*()` have been made available.
- If no DMA controller is available, the driver has to perform the copy between the hardware component and the asynchronous buffer. In such cases, the functions [`a4l_buf_get\(\)`](#) and [`a4l_buf_put\(\)`](#) are useful.

5.49.2 Function Documentation

5.49.2.1 `int a4l_buf_commit_absget (struct a4l_subdevice * subd, unsigned long count)`

Set the absolute count of data which was sent from the buffer to the device since the start of the acquisition and until the last DMA shot.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the transfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count transferred to the device during the last DMA shot plus the data count which have been sent since the beginning of the acquisition

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.2 `int a4l_buf_commit_absput (struct a4l_subdevice * subd, unsigned long count)`

Set the absolute count of data which was sent from the device to the buffer since the start of the acquisition and until the last DMA shot.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count transferred to the buffer during the last DMA shot plus the data count which have been sent / retrieved since the beginning of the acquisition

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.3 `int a4l_buf_commit_get (struct a4l_subdevice * subd, unsigned long count)`

Set the count of data sent from the buffer to the device during the last completed DMA shots.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The amount of data transferred

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.4 `int a4l_buf_commit_put (struct a4l_subdevice * subd, unsigned long count)`

Set the count of data sent to the buffer during the last completed DMA shots.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The amount of data transferred

Returns

0 on success, otherwise negative error code.

References `A4L_SUBD_BUSY_NR`, `a4l_subdevice::buf`, and `a4l_subdevice::status`.

5.49.2.5 `unsigned long a4l_buf_count (struct a4l_subdevice * subd)`

Get the data amount available in the Analogy buffer.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
-----------	-------------	--------------------------------

Returns

the amount of data available in the Analogy buffer.

References `A4L_SUBD_BUSY_NR`, `a4l_subdevice::buf`, and `a4l_subdevice::status`.

5.49.2.6 `int a4l_buf_evt (struct a4l_subdevice * subd, unsigned long evts)`

Signal some event(s) to a user-space program involved in some read / write operation.

The function `a4l_buf_evt()` is useful in many cases:

- To wake-up a process waiting for some data to read.
- To wake-up a process waiting for some data to write.
- To notify the user-process an error has occured during the acquisition.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>evts</i>	Some specific event to notify: <ul style="list-style-type: none"> • <code>A4L_BUF_ERROR</code> to indicate some error has occured during the transfer • <code>A4L_BUF_EOA</code> to indicate the acquisition is complete (this event is automatically set, it should not be used).

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.7 int a4l_buf_get (struct **a4l_subdevice** * subd, void * bufdata, unsigned long count)

Copy some data from the buffer to the device driver.

The function [a4l_buf_get\(\)](#) must copy data coming from the Analogy buffer to some acquisition device. This ring-buffer is an intermediate area between the device driver and the user-space program, which is supposed to provide the data to send to the device.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>bufdata</i>	The data buffer to copy into the Analogy buffer
in	<i>count</i>	The amount of data to copy

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.8 int a4l_buf_prepare_absget (struct **a4l_subdevice** * subd, unsigned long count)

Update the absolute count of data sent from the buffer to the device since the start of the acquisition and after the next DMA shot.

The functions [a4l_buf_prepare_\(abs\)put\(\)](#), [a4l_buf_commit_\(abs\)put\(\)](#), [a4l_buf_prepare_\(abs\)get\(\)](#) and [a4l_buf_commit_\(abs\)get\(\)](#) have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count to be transferred during the next DMA shot plus the data count which have been copied since the start of the acquisition

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.9 int a4l_buf_prepare_absput (struct **a4l_subdevice** * subd, unsigned long count)

Update the absolute count of data sent from the device to the buffer since the start of the acquisition and after the next DMA shot.

The functions [a4l_buf_prepare_\(abs\)put\(\)](#), [a4l_buf_commit_\(abs\)put\(\)](#), [a4l_buf_prepare_\(abs\)get\(\)](#) and [a4l_buf_commit_\(abs\)get\(\)](#) have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count to be transferred during the next DMA shot plus the data count which have been copied since the start of the acquisition

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.10 `int a4l_buf_prepare_get (struct a4l_subdevice * subd, unsigned long count)`

Set the count of data which is to be sent from the buffer to the device at the next DMA shot.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count to be transferred

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.11 `int a4l_buf_prepare_put (struct a4l_subdevice * subd, unsigned long count)`

Set the count of data which is to be sent to the buffer at the next DMA shot.

The functions `a4l_buf_prepare_(abs)put()`, `a4l_buf_commit_(abs)put()`, `a4l_buf_prepare_(abs)get()` and `a4l_buf_commit_(abs)get()` have been made available for DMA transfers. In such situations, no data copy is needed between the Analogy buffer and the device as some DMA controller is in charge of performing data shots from / to the Analogy buffer. However, some pointers still have to be updated so as to monitor the tranfers.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>count</i>	The data count to be transferred

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.12 `int a4l_buf_put (struct a4l_subdevice * subd, void * bufdata, unsigned long count)`

Copy some data from the device driver to the buffer.

The function `a4l_buf_put()` must copy data coming from some acquisition device to the Analogy buffer. This ring-buffer is an intermediate area between the device driver and the user-space program, which is supposed to recover the acquired data.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
in	<i>bufdata</i>	The data buffer to copy into the Analogy buffer
in	<i>count</i>	The amount of data to copy

Returns

0 on success, otherwise negative error code.

References A4L_SUBD_BUSY_NR, a4l_subdevice::buf, and a4l_subdevice::status.

5.49.2.13 int a4l_get_chan (struct **a4l_subdevice** * subd)

Get the channel index according to its type.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
----	-------------	--------------------------------

Returns

the channel index.

References A4L_CHAN_GLOBAL_CHANDESC, a4l_get_cmd(), a4l_subdevice::buf, a4l_subdevice::chan_desc, a4l_cmd_desc::chan_descs, a4l_channels_desc::chans, a4l_channels_desc::mode, a4l_channel::nb_bits, and a4l_cmd_desc::nb_chan.

5.49.2.14 struct **a4l_cmd_desc*** a4l_get_cmd (struct **a4l_subdevice** * subd)

Get the current Analogy command descriptor.

Parameters

in	<i>subd</i>	Subdevice descriptor structure
----	-------------	--------------------------------


Returns

the command descriptor.

Referenced by a4l_get_chan().

5.50 Interrupt management services

Collaboration diagram for Interrupt management services:



Functions

- unsigned int [a4l_get_irq](#) (struct a4l_device *dev)
Get the interrupt number in use for a specific device.
- int [a4l_request_irq](#) (struct a4l_device *dev, unsigned int irq, a4l_irq_hdlr_t handler, unsigned long flags, void *cookie)
Register an interrupt handler for a specific device.
- int [a4l_free_irq](#) (struct a4l_device *dev, unsigned int irq)
Release an interrupt handler for a specific device.

5.50.1 Detailed Description

5.50.2 Function Documentation

5.50.2.1 int [a4l_free_irq](#) (struct a4l_device * dev, unsigned int irq)

Release an interrupt handler for a specific device.

Parameters

in	<i>dev</i>	Device descriptor structure
in	<i>irq</i>	Line number of the addressed IRQ

Returns

0 on success, otherwise negative error code.

5.50.2.2 unsigned int [a4l_get_irq](#) (struct a4l_device * dev)

Get the interrupt number in use for a specific device.

Parameters

in	<i>dev</i>	Device descriptor structure
----	------------	-----------------------------

Returns

the line number used or A4L_IRQ_UNUSED if no interrupt is registered.

5.50.2.3 `int a4l_request_irq (struct a4l_device * dev, unsigned int irq, a4l_irq_hdlr_t handler,
unsigned long flags, void * cookie)`

Register an interrupt handler for a specific device.

Parameters

in	<i>dev</i>	Device descriptor structure
in	<i>irq</i>	Line number of the addressed IRQ
in	<i>handler</i>	Interrupt handler
in	<i>flags</i>	Registration flags: <ul style="list-style-type: none">• RTDM_IRQTYPE_SHARED: enable IRQ-sharing with other drivers (Warning: real-time drivers and non-real-time drivers cannot share an interrupt line).• RTDM_IRQTYPE_EDGE: mark IRQ as edge-triggered (Warning: this flag is meaningless in RTDM-less context).• A4L_IRQ_DISABLED: keep IRQ disabled when calling the action handler (Warning: this flag is ignored in RTDM-enabled configuration).
in	<i>cookie</i>	Pointer to be passed to the interrupt handler on invocation

Returns

0 on success, otherwise negative error code.

References `rtdm_lock_put_irqrestore()`.

5.51 Misc services

Collaboration diagram for Misc services:



Functions

- unsigned long long [a4l_get_time](#) (void)
Get the absolute time in nanoseconds.

5.51.1 Detailed Description

5.51.2 Function Documentation

5.51.2.1 unsigned long long a4l_get_time (void)

Get the absolute time in nanoseconds.

Returns

the absolute time expressed in nanoseconds

References [rtdm_clock_read\(\)](#).

5.52 Clocks and timers

Cobalt/POSIX clock and timer services.

Collaboration diagram for Clocks and timers:



Functions

- int [timer_create](#) (clockid_t clockid, const struct sigevent *__restrict__ evp, timer_t *__restrict__ timerid)
Create a timer.
- int [timer_settime](#) (timer_t timerid, int flags, const struct itimerspec *__restrict__ value, struct itimerspec *__restrict__ ovalue)
Start or stop a timer.
- int [timer_gettime](#) (timer_t timerid, struct itimerspec *value)
Get timer next expiration date and reload value.

5.52.1 Detailed Description

Cobalt/POSIX clock and timer services. Cobalt supports three built-in clocks:

CLOCK_REALTIME maps to the nucleus system clock, keeping time as the amount of time since the Epoch, with a resolution of one nanosecond.

CLOCK_MONOTONIC maps to an architecture-dependent high resolution counter, so is suitable for measuring short time intervals. However, when used for sleeping (with [clock_nanosleep\(\)](#)), the CLOCK_MONOTONIC clock has a resolution of one nanosecond, like the CLOCK_REALTIME clock.

CLOCK_MONOTONIC_RAW is Linux-specific, and provides monotonic time values from a hardware timer which is not adjusted by NTP. This is strictly equivalent to CLOCK_MONOTONIC with Xenomai, which is not NTP adjusted either.

In addition, external clocks can be dynamically registered using the [cobalt_clock_register\(\)](#) service. These clocks are fully managed by Cobalt extension code, which should advertise each incoming tick by calling [xnclock_tick\(\)](#) for the relevant clock, from an interrupt context.

Timer objects may be created with the [timer_create\(\)](#) service using any of the built-in or external clocks. The resolution of these timers is clock-specific. However, built-in clocks all have nanosecond resolution, as specified for [clock_nanosleep\(\)](#).

See Also

[Specification.](#)

5.52.2 Function Documentation

```
5.52.2.1 int timer_create ( clockid_t clockid, const struct sigevent *__restrict__ evp, timer_t
        *__restrict__ timerid )
```

Create a timer.

This service creates a timer based on the clock *clockid*.

If *evp* is not *NULL*, it describes the notification mechanism used on timer expiration. Only thread-directed notification is supported (*evp->sigev_notify* set to *SIGEV_THREAD_ID*).

If *evp* is *NULL*, the current Cobalt thread will receive the notifications with signal *SIGALRM*.

The recipient thread is delivered notifications when it calls any of the *sigwait()*, *sigtimedwait()* or *sigwait-info()* services.

If this service succeeds, an identifier for the created timer is returned at the address *timerid*. The timer is unarmed until started with the [timer_settime\(\)](#) service.

Parameters

<i>clockid</i>	clock used as a timing base;
<i>evp</i>	description of the asynchronous notification to occur when the timer expires;
<i>timerid</i>	address where the identifier of the created timer will be stored on success.

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • <i>EINVAL</i>, the clock <i>clockid</i> is invalid; • <i>EINVAL</i>, the member <i>sigev_notify</i> of the sigevent structure at the address <i>evp</i> is not <i>SIGEV_THREAD_ID</i>; • <i>EINVAL</i>, the member <i>sigev_signo</i> of the sigevent structure is an invalid signal number; • <i>EAGAIN</i>, the maximum number of timers was exceeded, recompile with a larger value.

See Also

[Specification.](#)

```
5.52.2.2 int timer_gettime ( timer_t timerid, struct itimerspec * value )
```

Get timer next expiration date and reload value.

This service stores, at the address *value*, the expiration date (member *it_value*) and reload value (member *it_interval*) of the timer *timerid*. The values are returned as time intervals, and as multiples of the system clock tick duration (see note in section [Clocks and timers services](#) for details on the duration of the system clock tick). If the timer was not started, the returned members *it_value* and *it_interval* of *value* are zero.

Parameters

<i>timerid</i>	timer identifier;
----------------	-------------------

<i>value</i>	address where the timer expiration date and reload value are stored on success.
--------------	---

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> EINVAL, <i>timerid</i> is invalid. For <i>timerid</i> to be valid, it must belong to the current process.

See Also

[Specification.](#)

5.52.2.3 `timer_settime (timer_t timerid, int flags, const struct itimerspec *__restrict__ value, struct itimerspec *__restrict__ ovalue)`

Start or stop a timer.

This service sets a timer expiration date and reload value of the timer *timerid*. If *ovalue* is not *NULL*, the current expiration date and reload value are stored at the address *ovalue* as with [timer_gettime\(\)](#).

If the member *it_value* of the **itimerspec** structure at *value* is zero, the timer is stopped, otherwise the timer is started. If the member *it_interval* is not zero, the timer is periodic. The current thread must be a Cobalt thread (created with [pthread_create\(\)](#)) and will be notified via signal of timer expirations. Note that these notifications will cause user-space threads to switch to secondary mode.

When starting the timer, if *flags* is *TIMER_ABSTIME*, the expiration value is interpreted as an absolute date of the clock passed to the [timer_create\(\)](#) service. Otherwise, the expiration value is interpreted as a time interval.

Expiration date and reload value are rounded to an integer count of nanoseconds.

Parameters

<i>timerid</i>	identifier of the timer to be started or stopped;
<i>flags</i>	one of 0 or <i>TIMER_ABSTIME</i> ;
<i>value</i>	address where the specified timer expiration date and reload value are read;
<i>ovalue</i>	address where the specified timer previous expiration date and reload value are stored if not <i>NULL</i> .

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> EINVAL, the specified timer identifier, expiration date or reload value is invalid. For <i>timerid</i> to be valid, it must belong to the current process.

Valid contexts:

- Cobalt kernel-space thread,
- kernel-space thread cancellation cleanup routine,
- Cobalt user-space thread (switches to primary mode),
- user-space thread cancellation cleanup routine.

See Also

[Specification.](#)

5.53 Condition variables

Cobalt/POSIX condition variable services.

Collaboration diagram for Condition variables:



Functions

- int [pthread_cond_init](#) (pthread_cond_t *cond, const pthread_condattr_t *attr)
Initialize a condition variable.
- int [pthread_cond_destroy](#) (pthread_cond_t *cond)
Destroy a condition variable.

5.53.1 Detailed Description

Cobalt/POSIX condition variable services. A condition variable is a synchronization object that allows threads to suspend execution until some predicate on shared data is satisfied. The basic operations on conditions are: signal the condition (when the predicate becomes true), and wait for the condition, suspending the thread execution until another thread signals the condition.

A condition variable must always be associated with a mutex, to avoid the race condition where a thread prepares to wait on a condition variable and another thread signals the condition just before the first thread actually waits on it.

Before it can be used, a condition variable has to be initialized with [pthread_cond_init\(\)](#). An attribute object, which reference may be passed to this service, allows to select the features of the created condition variable, namely the *clock* used by the [pthread_cond_timedwait\(\)](#) service (*CLOCK_REALTIME* is used by default), and whether it may be shared between several processes (it may not be shared by default, see [pthread_condattr_setpshared\(\)](#)).

Note that only [pthread_cond_init\(\)](#) may be used to initialize a condition variable, using the static initializer *PTHREAD_COND_INITIALIZER* is not supported.

5.53.2 Function Documentation

5.53.2.1 int pthread_cond_destroy (pthread_cond_t * cond)

Destroy a condition variable.

This service destroys the condition variable *cond*, if no thread is currently blocked on it. The condition variable becomes invalid for all condition variable services (they all return the EINVAL error) except [pthread_cond_init\(\)](#).

Parameters

<i>cond</i>	the condition variable to be destroyed.
-------------	---

Returns

0 on succes,
an error number if:

- EINVAL, the condition variable *cond* is invalid;
- EPERM, the condition variable is not process-shared and does not belong to the current process;
- EBUSY, some thread is currently using the condition variable.

See Also

[Specification.](#)

Referenced by `rt_cond_create()`, and `rt_cond_delete()`.

5.53.2.2 `int pthread_cond_init (pthread_cond_t * cond, const pthread_condattr_t * attr)`

Initialize a condition variable.

This service initializes the condition variable *cond*, using the condition variable attributes object *attr*. If *attr* is *NULL* or this service is called from user-space, default attributes are used (see `pthread_condattr_init()`).

Parameters

<i>cond</i>	the condition variable to be initialized;
<i>attr</i>	the condition variable attributes object.

Returns

0 on succes,
an error number if:

- EINVAL, the condition variable attributes object *attr* is invalid or uninitialized;
- EBUSY, the condition variable *cond* was already initialized;
- ENOMEM, insufficient memory exists in the system heap to initialize the condition variable, increase `CONFIG_XENO_OPT_SYS_HEAPSZ`.

See Also

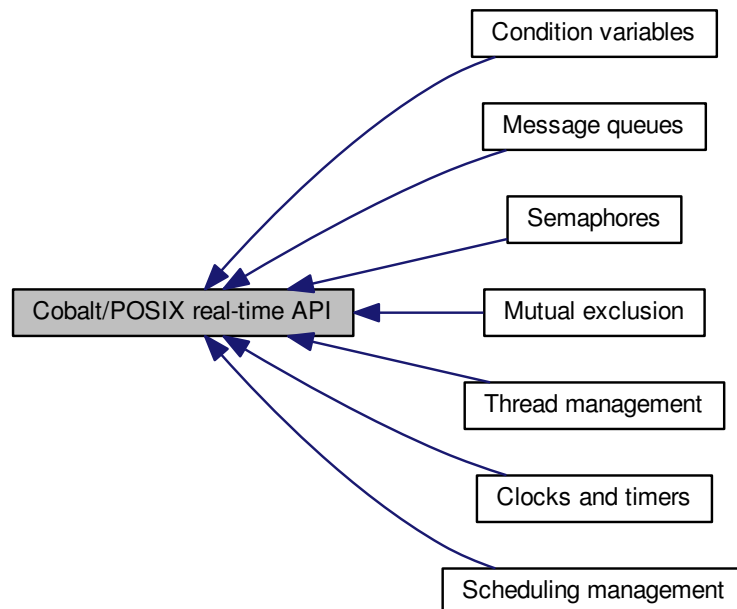
[Specification.](#)

Referenced by `rt_cond_create()`.

5.54 Cobalt/POSIX real-time API

The Cobalt interface is an implementation of a small subset of the [Single Unix specification](#) over the Xenomai generic RTOS core.

Collaboration diagram for Cobalt/POSIX real-time API:



Modules

- [Clocks and timers](#)
Cobalt/POSIX clock and timer services.
- [Condition variables](#)
Cobalt/POSIX condition variable services.
- [Message queues](#)
Cobalt/POSIX message queue services.
- [Mutual exclusion](#)
Cobalt/POSIX mutual exclusion services.
- [Semaphores](#)
Cobalt/POSIX semaphore services.
- [Thread management](#)
Cobalt/POSIX thread management services.
- [Scheduling management](#)
Cobalt/POSIX scheduling management services.

5.54.1 Detailed Description

The Cobalt interface is an implementation of a small subset of the [Single Unix specification](#) over the Xenomai generic RTOS core.

5.55 Message queues

Cobalt/POSIX message queue services.

Collaboration diagram for Message queues:



Functions

- `mqd_t mq_open (const char *name, int oflags,...)`
Open a message queue.
- `int mq_close (mqd_t mqd)`
Close a message queue.
- `int mq_unlink (const char *name)`
Unlink a message queue.
- `int mq_getattr (mqd_t mqd, struct mq_attr *attr)`
Get message queue attributes.
- `int mq_setattr (mqd_t mqd, const struct mq_attr *__restrict__ attr, struct mq_attr *__restrict__ oattr)`
Set message queue attributes.
- `int mq_notify (mqd_t mqd, const struct sigevent *evp)`
Enable notification on message arrival.

5.55.1 Detailed Description

Cobalt/POSIX message queue services. A message queue allow exchanging data between real-time threads. For a POSIX message queue, maximum message length and maximum number of messages are fixed when it is created with `mq_open()`.

5.55.2 Function Documentation

5.55.2.1 `int mq_close (mqd_t mqd)`

Close a message queue.

This service closes the message queue descriptor `mqd`. The message queue is destroyed only when all open descriptors are closed, and when unlinked with a call to the `mq_unlink()` service.

Parameters

<code>mqd</code>	message queue descriptor.
------------------	---------------------------

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EBADF, <i>mqd</i> is an invalid message queue descriptor; • EPERM, the caller context is invalid.

Valid contexts:

- kernel module initialization or cleanup routine;
- kernel-space cancellation cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode);
- user-space cancellation cleanup routine.

See Also

[Specification.](#)

5.55.2.2 int mq_getattr (mqd_t mqd, struct mq_attr * attr)

Get message queue attributes.

This service stores, at the address *attr*, the attributes of the messages queue descriptor *mqd*.

The following attributes are set:

- *mq_flags*, flags of the message queue descriptor *mqd*;
- *mq_maxmsg*, maximum number of messages in the message queue;
- *mq_msgsize*, maximum message size;
- *mq_curmsgs*, number of messages currently in the queue.

Parameters

<i>mqd</i>	message queue descriptor;
<i>attr</i>	address where the message queue attributes will be stored on success.

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EBADF, <i>mqd</i> is not a valid descriptor.

See Also

[Specification.](#)

5.55.2.3 `int mq_notify (mqd_t mqd, const struct sigevent * evp)`

Enable notification on message arrival.

If *evp* is not *NULL* and is the address of a **sigevent** structure with the *sigev_notify* member set to *SIGEV_SIGNAL*, the current thread will be notified by a signal when a message is sent to the message queue *mqd*, the queue is empty, and no thread is blocked in call to *mq_receive()* or *mq_timedreceive()*. After the notification, the thread is unregistered.

If *evp* is *NULL* or the *sigev_notify* member is *SIGEV_NONE*, the current thread is unregistered.

Only one thread may be registered at a time.

If the current thread is not a Cobalt thread (created with [pthread_create\(\)](#)), this service fails.

Note that signals sent to user-space Cobalt threads will cause them to switch to secondary mode.

Parameters

<i>mqd</i>	message queue descriptor;
<i>evp</i>	pointer to an event notification structure.

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • <i>EINVAL</i>, <i>evp</i> is invalid; • <i>EPERM</i>, the caller context is invalid; • <i>EBADF</i>, <i>mqd</i> is not a valid message queue descriptor; • <i>EBUSY</i>, another thread is already registered.

Valid contexts:

- Xenomai kernel-space Cobalt thread,
- Xenomai user-space Cobalt thread (switches to primary mode).

See Also

[Specification.](#)

5.55.2.4 `mqd_t mq_open (const char * name, int oflags, ...)`

Open a message queue.

This service establishes a connection between the message queue named *name* and the calling context (kernel-space as a whole, or user-space process).

One of the following values should be set in *oflags*:

- *O_RDONLY*, meaning that the returned queue descriptor may only be used for receiving messages;
- *O_WRONLY*, meaning that the returned queue descriptor may only be used for sending messages;
- *O_RDWR*, meaning that the returned queue descriptor may be used for both sending and receiving messages.

If no message queue named *name* exists, and *oflags* has the *O_CREAT* bit set, the message queue is created by this function, taking two more arguments:

- a *mode* argument, of type **mode_t**, currently ignored;
- an *attr* argument, pointer to an **mq_attr** structure, specifying the attributes of the new message queue.

If *oflags* has the two bits *O_CREAT* and *O_EXCL* set and the message queue already exists, this service fails.

If the *O_NONBLOCK* bit is set in *oflags*, the *mq_send()*, *mq_receive()*, *mq_timedsend()* and *mq_timedreceive()* services return -1 with *errno* set to *EAGAIN* instead of blocking their caller.

The following arguments of the **mq_attr** structure at the address *attr* are used when creating a message queue:

- *mq_maxmsg* is the maximum number of messages in the queue (128 by default);
- *mq_msgsize* is the maximum size of each message (128 by default).

name may be any arbitrary string, in which slashes have no particular meaning. However, for portability, using a name which starts with a slash and contains no other slash is recommended.

Parameters

<i>name</i>	name of the message queue to open;
<i>oflags</i>	flags.

Returns

a message queue descriptor on success;

-1 with *errno* set if:

- *ENAMETOOLONG*, the length of the *name* argument exceeds 64 characters;
- *EEXIST*, the bits *O_CREAT* and *O_EXCL* were set in *oflags* and the message queue already exists;
- *ENOENT*, the bit *O_CREAT* is not set in *oflags* and the message queue does not exist;
- *ENOSPC*, allocation of system memory failed, or insufficient memory exists in the system heap to create the queue, try increasing *CONFIG_XENO_OPT_SYS_HEAPSZ*;
- *EPERM*, attempting to create a message queue from an invalid context;
- *EINVAL*, the *attr* argument is invalid;
- *EMFILE*, too many descriptors are currently open.

Valid contexts:

When creating a message queue, only the following contexts are valid:

- kernel module initialization or cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode).

See Also

[Specification.](#)

```
5.55.2.5 int mq_setattr ( mqd_t mqd, const struct mq_attr *__restrict__ attr, struct mq_attr
    *__restrict__ oattr )
```

Set message queue attributes.

This service sets the flags of the *mqd* descriptor to the value of the member *mq_flags* of the **mq_attr** structure pointed to by *attr*.

The previous value of the message queue attributes are stored at the address *oattr* if it is not *NULL*.

Only setting or clearing the *O_NONBLOCK* flag has an effect.

Parameters

<i>mqd</i>	message queue descriptor;
<i>attr</i>	pointer to new attributes (only <i>mq_flags</i> is used);
<i>oattr</i>	if not <i>NULL</i> , address where previous message queue attributes will be stored on success.

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> EBADF, <i>mqd</i> is not a valid message queue descriptor.

See Also

[Specification.](#)

5.55.2.6 int mq_unlink (const char * name)

Unlink a message queue.

This service unlinks the message queue named *name*. The message queue is not destroyed until all queue descriptors obtained with the [mq_open\(\)](#) service are closed with the [mq_close\(\)](#) service. However, after a call to this service, the unlinked queue may no longer be reached with the [mq_open\(\)](#) service.

Parameters

<i>name</i>	name of the message queue to be unlinked.
-------------	---

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> EPERM, the caller context is invalid; ENAMETOOLONG, the length of the <i>name</i> argument exceeds 64 characters; ENOENT, the message queue does not exist.

Valid contexts:

- kernel module initialization or cleanup routine;
- kernel-space cancellation cleanup routine;
- user-space thread (Xenomai threads switch to secondary mode);
- user-space cancellation cleanup routine.

See Also

[Specification.](#)

5.56 Mutual exclusion

Cobalt/POSIX mutual exclusion services.

Collaboration diagram for Mutual exclusion:



5.56.1 Detailed Description

Cobalt/POSIX mutual exclusion services. A mutex is a MUTual EXclusion device, and is useful for protecting shared data structures from concurrent modifications, and implementing critical sections and monitors.

A mutex has two possible states: unlocked (not owned by any thread), and locked (owned by one thread). A mutex can never be owned by two different threads simultaneously. A thread attempting to lock a mutex that is already locked by another thread is suspended until the owning thread unlocks the mutex first.

Before it can be used, a mutex has to be initialized with `pthread_mutex_init()`. An attribute object, which reference may be passed to this service, allows to select the features of the created mutex, namely its *type* (see `pthread_mutexattr_settype()`), the priority *protocol* it uses (see `pthread_mutexattr_setprotocol()`) and whether it may be shared between several processes (see `pthread_mutexattr_setpshared()`).


By default, Cobalt mutexes are of the normal type, use no priority protocol and may not be shared between several processes.

Note that only `pthread_mutex_init()` may be used to initialize a mutex, using the static initializer `PTHREAD_MUTEX_INITIALIZER` is not supported.

5.57 Semaphores

Cobalt/POSIX semaphore services.

Collaboration diagram for Semaphores:



Functions

- int [sem_destroy](#) (sem_t *sem)
Destroy an unnamed semaphore.
- int [sem_post](#) (sem_t *sem)
Post a semaphore.
- int [sem_trywait](#) (sem_t *sem)
Attempt to decrement a semaphore.
- int [sem_wait](#) (sem_t *sem)
Decrement a semaphore.
- int [sem_timedwait](#) (sem_t *sem, const struct timespec *abs_timeout)
Attempt to decrement a semaphore with a time limit.
- int [sem_close](#) (sem_t *sem)
Close a named semaphore.
- int [sem_unlink](#) (const char *name)
Unlink a named semaphore.

5.57.1 Detailed Description

Cobalt/POSIX semaphore services. Semaphores are counters for resources shared between threads. The basic operations on semaphores are: increment the counter atomically, and wait until the counter is non-null and decrement it atomically.

Semaphores have a maximum value past which they cannot be incremented. The macro `SEM_VALUE_MAX` is defined to be this maximum value.

5.57.2 Function Documentation

5.57.2.1 int sem_close (sem_t * sem)

Close a named semaphore.

This service closes the semaphore *sem*. The semaphore is destroyed only when unlinked with a call to the [sem_unlink\(\)](#) service and when each call to `sem_open()` matches a call to this service.

When a semaphore is destroyed, the memory it used is returned to the system heap, so that further references to this semaphore are not guaranteed to fail, as is the case for unnamed semaphores.

This service fails if *sem* is an unnamed semaphore.

Parameters

<i>sem</i>	the semaphore to be closed.
------------	-----------------------------

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EINVAL, the semaphore <i>sem</i> is invalid or is an unnamed semaphore.

See Also

[Specification.](#)

5.57.2.2 int sem_destroy (sem_t * sem)

Destroy an unnamed semaphore.

This service destroys the semaphore *sem*. Threads currently blocked on *sem* are unblocked and the service they called return -1 with *errno* set to EINVAL. The semaphore is then considered invalid by all semaphore services (they all fail with *errno* set to EINVAL) except *sem_init()*.

This service fails if *sem* is a named semaphore.

Parameters

<i>sem</i>	the semaphore to be destroyed.
------------	--------------------------------

Return values

<i>always</i>	0 on success. If SEM_WARNDEL was mentioned in <i>sem_init_np()</i> , the semaphore is deleted as requested and a strictly positive value is returned to warn the caller if threads were pending on it, otherwise zero is returned. If SEM_NOBUSYDEL was mentioned in <i>sem_init_np()</i> , sem_destroy() may succeed only if no thread is waiting on the semaphore to delete, otherwise -EBUSY is returned.
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EINVAL, the semaphore <i>sem</i> is invalid or a named semaphore; • EPERM, the semaphore <i>sem</i> is not process-shared and does not belong to the current process. • EBUSY, a thread is currently waiting on the semaphore <i>sem</i> with SEM_NOBUSYDEL set.

See Also

[Specification.](#)

5.57.2.3 int sem_post (sem_t * sem)

Post a semaphore.

This service posts the semaphore *sem*.

If no thread is currently blocked on this semaphore, its count is incremented unless "pulse" mode is enabled for it (see `sem_init_np()`, `SEM_PULSE`). If a thread is blocked on the semaphore, the thread heading the wait queue is unblocked.

Parameters

<i>sem</i>	the semaphore to be signaled.
------------	-------------------------------

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EINVAL, the specified semaphore is invalid or uninitialized; • EPERM, the semaphore <i>sm</i> is not process-shared and does not belong to the current process; • EAGAIN, the semaphore count is <i>SEM_VALUE_MAX</i>.

See Also

[Specification.](#)

5.57.2.4 int sem_timedwait (sem_t * sem, const struct timespec * abs_timeout)

Attempt to decrement a semaphore with a time limit.

This service is equivalent to [sem_wait\(\)](#), except that the caller is only blocked until the timeout *abs_timeout* expires.

Parameters

<i>sem</i>	the semaphore to be decremented;
<i>abs_timeout</i>	the timeout, expressed as an absolute value of the relevant clock for the semaphore, either CLOCK_MONOTONIC if SEM_RAWCLOCK was mentioned via sem_init_np(), or CLOCK_REALTIME otherwise.

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EPERM, the caller context is invalid; • EINVAL, the semaphore is invalid or uninitialized; • EINVAL, the specified timeout is invalid; • EPERM, the semaphore <i>sm</i> is not process-shared and does not belong to the current process; • EINTR, the caller was interrupted by a signal while blocked in this service; • ETIMEDOUT, the semaphore could not be decremented and the specified timeout expired.

Valid contexts:

- Xenomai kernel-space thread,
- Xenomai user-space thread (switches to primary mode).

See Also

[Specification.](#)

References `sem_trywait()`.

5.57.2.5 `int sem_trywait (sem_t * sem)`

Attempt to decrement a semaphore.

This service is equivalent to `sem_wait()`, except that it returns immediately if the semaphore *sem* is currently depleted, and that it is not a cancellation point.

Parameters

<i>sem</i>	the semaphore to be decremented.
------------	----------------------------------

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EINVAL, the specified semaphore is invalid or uninitialized; • EPERM, the semaphore <i>sem</i> is not process-shared and does not belong to the current process; • EAGAIN, the specified semaphore is currently fully depleted. •

See Also

[Specification.](#)

Referenced by `sem_timedwait()`, and `sem_wait()`.

5.57.2.6 `int sem_unlink (const char * name)`

Unlink a named semaphore.

This service unlinks the semaphore named *name*. This semaphore is not destroyed until all references obtained with `sem_open()` are closed by calling `sem_close()`. However, the unlinked semaphore may no longer be reached with the `sem_open()` service.

When a semaphore is destroyed, the memory it used is returned to the system heap, so that further references to this semaphore are not guaranteed to fail, as is the case for unnamed semaphores.

Parameters

<i>name</i>	the name of the semaphore to be unlinked.
-------------	---

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • ENAMETOOLONG, the length of the <i>name</i> argument exceeds 64 characters; • ENOENT, the named semaphore does not exist.

See Also

[Specification.](#)

5.57.2.7 `int sem_wait (sem_t * sem)`

Decrement a semaphore.

This service decrements the semaphore *sem* if it is currently if its value is greater than 0. If the semaphore's value is currently zero, the calling thread is suspended until the semaphore is posted, or a signal is delivered to the calling thread.

This service is a cancellation point for Cobalt threads (created with the [pthread_create\(\)](#) service). When such a thread is cancelled while blocked in a call to this service, the semaphore state is left unchanged before the cancellation cleanup handlers are called.

Parameters

<i>sem</i>	the semaphore to be decremented.
------------	----------------------------------

Return values

0	on success;
-1	with <i>errno</i> set if: <ul style="list-style-type: none"> • EPERM, the caller context is invalid; • EINVAL, the semaphore is invalid or uninitialized; • EPERM, the semaphore <i>sem</i> is not process-shared and does not belong to the current process; • EINTR, the caller was interrupted by a signal while blocked in this service.

Valid contexts:

- Xenomai kernel-space thread,
- Xenomai user-space thread (switches to primary mode).

See Also

[Specification.](#)

References `sem_trywait()`.

5.58 Thread management

Cobalt/POSIX thread management services.

Collaboration diagram for Thread management:



Functions

- int [pthread_setschedparam_ex](#) (pthread_t thread, int policy, const struct sched_param_ex *param_ex)
Set extended scheduling policy of thread.
- int [pthread_getschedparam_ex](#) (pthread_t thread, int *__restrict__ policy_r, struct sched_param_ex *__restrict__ param_ex)
Get extended scheduling policy of thread.
- int [pthread_create](#) (pthread_t *ptid_r, const pthread_attr_t *attr, void *(*start)(void *), void *arg)
Create a new thread.
- int [pthread_make_periodic_np](#) (pthread_t thread, clockid_t clk_id, const struct timespec *__restrict__ starttp, const struct timespec *__restrict__ periodtp)
Make a thread periodic.
- int [pthread_set_mode_np](#) (int clrmask, int setmask, int *mode_r)
Set the mode of the current thread.

5.58.1 Detailed Description

Cobalt/POSIX thread management services.

See Also

[Specification.](#)

5.58.2 Function Documentation

5.58.2.1 int pthread_create (pthread_t * ptid_r, const pthread_attr_t * attr, void (*)(void *) start, void * arg)

Create a new thread.

This service creates a thread managed by the Xenomai nucleus in dual kernel configuration.

The new thread signal mask is inherited from the current thread, if it was also created with [pthread_create\(\)](#), otherwise the new thread signal mask is empty.

Other attributes of the new thread depend on the *attr* argument. If *attr* is NULL, default values for these attributes are used.

Returning from the *start* routine has the same effect as calling pthread_exit() with the return value.

Parameters

<i>ptid_r</i>	address where the identifier of the new thread will be stored on success;
<i>attr</i>	thread attributes;
<i>start</i>	thread start routine;
<i>arg</i>	opaque user-supplied argument passed to <i>start</i> ;

Returns

0 on success;

an error number if:

- EINVAL, *attr* is invalid;
- EAGAIN, insufficient memory exists in the system heap to create a new thread, increase CONFIG_XENO_OPT_SYS_HEAPSZ;
- EINVAL, thread attribute *inheritsched* is set to PTHREAD_INHERIT_SCHED and the calling thread does not belong to the Cobalt interface;

See Also

[Specification.](#)

Note

When creating or shadowing a Xenomai thread for the first time in user-space, Xenomai installs a handler for the SIGSHADOW signal. If you had installed a handler before that, it will be automatically called by Xenomai for SIGSHADOW signals that it has not sent.

If, however, you install a signal handler for SIGSHADOW after creating or shadowing the first Xenomai thread, you have to explicitly call the function `cobalt_sigshadow_handler` at the beginning of your signal handler, using its return to know if the signal was in fact an internal signal of Xenomai (in which case it returns 1), or if you should handle the signal (in which case it returns 0). `cobalt_sigshadow_handler` prototype is:

int cobalt_sigshadow_handler(int sig, struct siginfo *si, void *ctxt);

Which means that you should register your handler with `sigaction`, using the SA_SIGINFO flag, and pass all the arguments you received to `cobalt_sigshadow_handler`.

5.58.2.2 `int pthread_getschedparam_ex (pthread_t thread, int *__restrict__ policy_r, struct sched_param_ex *__restrict__ param_ex)`

Get extended scheduling policy of thread.

This service is an extended version of the regular `pthread_getschedparam()` service, which also supports Xenomai-specific or additional POSIX scheduling policies, not available with the host Linux environment.

Parameters

<i>thread</i>	target thread;
<i>policy_r</i>	address where the scheduling policy of <i>thread</i> is stored on success;
<i>param_ex</i>	address where the scheduling parameters of <i>thread</i> are stored on success.

Returns

- 0 on success;
 an error number if:
- ESRCH, *thread* is invalid.

See Also

[Specification.](#)

```
5.58.2.3 int pthread_make_periodic_np ( pthread_t thread, clockid_t clk_id, const struct timespec
      *__restrict__ starttp, const struct timespec *__restrict__ periodtp )
```

Make a thread periodic.

This service make the Cobalt interface *thread* periodic.

This service is a non-portable extension of the POSIX interface.

Parameters

<i>thread</i>	thread identifier.
<i>clk_id</i>	clock identifier, either CLOCK_REALTIME, CLOCK_MONOTONIC or CLOCK_MONOTONIC_RAW.
<i>starttp</i>	start time, expressed as an absolute value of the clock <i>clock_id</i> . The affected thread will be delayed until this point is reached.
<i>periodtp</i>	period, expressed as a time interval.

Returns

- 0 on success;
 an error number if:
- ESRCH, *thread* is invalid;
 - ETIMEDOUT, the start time has already passed.
 - EINVAL, the specified clock is unsupported;

```
5.58.2.4 int pthread_set_mode_np ( int clrmask, int setmask, int * mode_r )
```

Set the mode of the current thread.

This service sets the mode of the calling thread. *clrmask* and *setmask* are two bit masks which are respectively cleared and set in the calling thread status. They are a bitwise OR of the following values:

- PTHREAD_LOCK_SCHED, when set, locks the scheduler, which prevents the current thread from being switched out until the scheduler is unlocked;
- PTHREAD_WARNSW, when set, causes the signal SIGXCPU to be sent to the current thread, whenever it involuntary switches to secondary mode;
- PTHREAD_CONFORMING can be passed in *setmask* to switch the current user-space task to its preferred runtime mode. The only meaningful use of this switch is to force a real-time shadow back to primary mode. Any other use leads to a nop.
- PTHREAD_DISABLE_LOCKBREAK disallows breaking the scheduler lock. In the default case, a thread which holds the scheduler lock is allowed to drop it temporarily for sleeping. If this mode bit is set, such thread would return with EINTR immediately from any blocking call.

PTHREAD_LOCK_SCHED and PTHREAD_DISABLE_LOCKBREAK are valid for any Xenomai thread, other bits are valid for Xenomai user-space threads only.

This service is a non-portable extension of the POSIX interface.

Parameters

<i>clrmask</i>	set of bits to be cleared;
<i>setmask</i>	set of bits to be set.
<i>mode_r</i>	If non-NULL, <i>mode_r</i> must be a pointer to a memory location which will be written upon success with the previous set of active mode bits. If NULL, the previous set of active mode bits will not be returned.

Returns

- 0 on success;
 an error number if:
- EINVAL, some bit in *clrmask* or *setmask* is invalid.

Note

Setting *clrmask* and *setmask* to zero leads to a nop, only returning the previous mode if *mode_r* is a valid address.

5.58.2.5 `int pthread_setschedparam_ex (pthread_t thread, int policy, const struct sched_param_ex * param_ex)`

Set extended scheduling policy of thread.

This service is an extended version of the regular `pthread_setschedparam()` service, which supports Xenomai-specific or additional scheduling policies, not available with the host Linux environment.

This service set the scheduling policy of the Xenomai thread *thread* to the value *policy*, and its scheduling parameters (e.g. its priority) to the value pointed to by *param_ex*.

If *thread* does not match the identifier of a Xenomai thread, this action falls back to the regular `pthread_setschedparam()` service.

Parameters

<i>thread</i>	target Cobalt thread;
<i>policy</i>	scheduling policy, one of SCHED_WEAK, SCHED_FIFO, SCHED_COBALT, SCHED_RR, SCHED_SPORADIC, SCHED_TP, SCHED_QUOTA or SCHED_NORMAL;
<i>param_ex</i>	scheduling parameters address. As a special exception, a negative <code>sched_priority</code> value is interpreted as if SCHED_WEAK was given in <i>policy</i> , using the absolute value of this parameter as the weak priority level.

When `CONFIG_XENO_OPT_SCHED_WEAK` is enabled, SCHED_WEAK exhibits priority levels in the [0..99] range (inclusive). Otherwise, `sched_priority` must be zero for the SCHED_WEAK policy.

Returns

- 0 on success;
 an error number if:
- ESRCH, *thread* is invalid;
 - EINVAL, *policy* or *param_ex->sched_priority* is invalid;
 - EAGAIN, in user-space, insufficient memory exists in the system heap, increase `CONFIG_XENO_OPT_SYS_HEAPSZ`;
 - EFAULT, in user-space, *param_ex* is an invalid address;
 - EPERM, in user-space, the calling process does not have superuser permissions.

See Also

[Specification.](#)

Note

When creating or shadowing a Xenomai thread for the first time in user-space, Xenomai installs a handler for the SIGSHADOW signal. If you had installed a handler before that, it will be automatically called by Xenomai for SIGSHADOW signals that it has not sent.

If, however, you install a signal handler for SIGSHADOW after creating or shadowing the first Xenomai thread, you have to explicitly call the function `cobalt_sigshadow_handler` at the beginning of your signal handler, using its return to know if the signal was in fact an internal signal of Xenomai (in which case it returns 1), or if you should handle the signal (in which case it returns 0). `cobalt_sigshadow_handler` prototype is:

`int cobalt_sigshadow_handler(int sig, struct siginfo *si, void *ctxt);`

Which means that you should register your handler with `sigaction`, using the `SA_SIGINFO` flag, and pass all the arguments you received to `cobalt_sigshadow_handler`.

[pthread_setschedparam_ex\(\)](#) may switch the caller to secondary mode.

5.59 Scheduling management

Cobalt/POSIX scheduling management services.

Collaboration diagram for Scheduling management:



Functions

- int [sched_setconfig_np](#) (int cpu, int policy, const union sched_config *config, size_t len)
Set CPU-specific scheduler settings for a policy.
- ssize_t [sched_getconfig_np](#) (int cpu, int policy, union sched_config *config, size_t *len_r)
Retrieve CPU-specific scheduler settings for a policy.

5.59.1 Detailed Description

Cobalt/POSIX scheduling management services.

5.59.2 Function Documentation

5.59.2.1 `ssize_t sched_getconfig_np (int cpu, int policy, union sched_config * config, size_t * len_r)`

Retrieve CPU-specific scheduler settings for a policy.

A configuration is strictly local to the target *cpu*, and may differ from other processors.

Parameters

<i>cpu</i>	processor to retrieve the configuration of.
<i>policy</i>	scheduling policy to which the configuration data applies. Currently, SCHED_TP and SCHED_QUOTA are valid.
<i>config</i>	a pointer to a memory area where the configuration data will be copied back. This area must be at least <i>*len_r</i> bytes long.
<i>len_r</i>	overall length of the configuration data returned (in bytes).

Returns

the number of bytes copied to *config* on success;
a negative error number if:

- EINVAL, *cpu* is invalid, or *policy* is unsupported by the current kernel configuration, or *len* cannot hold the retrieved configuration data.
- ESRCH, with *policy* equal to SCHED_QUOTA, if the group identifier required to perform the operation is not valid.
- ENOMEM, lack of memory to perform the operation.

- ENOSPC, *len* is too short.

5.59.2.2 `int sched_setconfig_np (int cpu, int policy, const union sched_config * config, size_t len)`

Set CPU-specific scheduler settings for a policy.

A configuration is strictly local to the target *cpu*, and may differ from other processors.

Parameters

<i>cpu</i>	processor to load the configuration of.
<i>policy</i>	scheduling policy to which the configuration data applies. Currently, SCHED_TP and SCHED_QUOTA are valid.
<i>config</i>	a pointer to the configuration data to load on <i>cpu</i> , applicable to <i>policy</i> .

Settings applicable to SCHED_TP

This call installs the temporal partitions for *cpu*.

- `config.tp.windows` should be a non-null set of time windows, defining the scheduling time slots for *cpu*. Each window defines its offset from the start of the global time frame (`windows[].offset`), a duration (`windows[].duration`), and the partition id it applies to (`windows[].ptid`).

Time windows must be strictly contiguous, i.e. `windows[n].offset + windows[n].duration` shall equal `windows[n + 1].offset`. If `windows[].ptid` is in the range `[0..CONFIG_XENO_OPT_SCHED_TP_NRPART-1]`, SCHED_TP threads which belong to the partition being referred to may run for the duration of the time window.

Time holes may be defined using windows assigned to the pseudo partition #-1, during which no SCHED_TP threads may be scheduled.

- `config.tp.nr_windows` should define the number of elements present in the `config.tp.windows[]` array.

Settings applicable to SCHED_QUOTA

This call manages thread groups running on *cpu*.

- `config.quota.op` should define the operation to be carried out. Valid operations are:
 - `sched_quota_add` for creating a new thread group on *cpu*. The new group identifier will be written back to `config.quota.add.tgid_r` upon success. A new group is given no initial runtime budget when created. `sched_quota_set` should be issued to enable it.
 - `sched_quota_remove` for deleting a thread group on *cpu*. The group identifier should be passed in `config.quota.remove.tgid`.
 - `sched_quota_set` for updating the scheduling parameters of a thread group defined on *cpu*. The group identifier should be passed in `config.quota.set.tgid`, along with the allotted percentage of the quota interval (`config.quota.set.quota`), and the peak percentage allowed (`config.quota.set.quota_peak`).
 - `sched_quota_get` for retrieving the scheduling parameters of a thread group defined on *cpu*. The group identifier should be passed in `config.quota.get.tgid`. The allotted percentage of the quota interval (`config.quota.get.quota_r`), and the peak percentage (`config.quota.get.quota_peak_r`) will be written to the given output variables. The result of this operation is identical to calling `sched_getconfig_np()`.

Parameters

<i>len</i>	overall length of the configuration data (in bytes).
------------	--


Returns

0 on success;
an error number if:

- EINVAL, *cpu* is invalid, or *policy* is unsupported by the current kernel configuration, *len* is invalid, or *config* contains invalid parameters.
- ENOMEM, lack of memory to perform the operation.
- EBUSY, with *policy* equal to SCHED_QUOTA, if an attempt is made to remove a thread group which still manages threads.
- ESRCH, with *policy* equal to SCHED_QUOTA, if the group identifier required to perform the operation is not valid.

5.60 Asynchronous acquisition API

Collaboration diagram for Asynchronous acquisition API:



Data Structures

- struct [a4l_cmd_desc](#)
Structure describing the asynchronous instruction.

Functions

- int [a4l_snd_command](#) ([a4l_desc_t](#) *dsc, [a4l_cmd_t](#) *cmd)
Send a command to an Analogy device.
- int [a4l_snd_cancel](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd)
Cancel an asynchronous acquisition.
- int [a4l_set_bufsize](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned long size)
Change the size of the asynchronous buffer.
- int [a4l_get_bufsize](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned long *size)
Get the size of the asynchronous buffer.
- int [a4l_mark_bufw](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned long cur, unsigned long *new)
Update the asynchronous buffer state.
- int [a4l_poll](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned long ms_timeout)
Get the available data count.
- int [a4l_mmap](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned long size, void **ptr)
Map the asynchronous ring-buffer into a user-space.

ANALOGY_CMD_XXX

Common command flags definitions

- #define [A4L_CMD_SIMUL](#) 0x1
Do not execute the command, just check it.
- #define [A4L_CMD_BULK](#) 0x2
Perform data recovery / transmission in bulk mode.
- #define [A4L_CMD_WRITE](#) 0x4
Perform a command which will write data to the device.

TRIG_XXX

Command triggers flags definitions

- #define **TRIG_NONE** 0x00000001
Never trigger.
- #define **TRIG_NOW** 0x00000002
Trigger now + N ns.
- #define **TRIG_FOLLOW** 0x00000004
Trigger on next lower level trig.
- #define **TRIG_TIME** 0x00000008
Trigger at time N ns.
- #define **TRIG_TIMER** 0x00000010
Trigger at rate N ns.
- #define **TRIG_COUNT** 0x00000020
Trigger when count reaches N.
- #define **TRIG_EXT** 0x00000040
Trigger on external signal N.
- #define **TRIG_INT** 0x00000080
Trigger on analogy-internal signal N.
- #define **TRIG_OTHER** 0x00000100
Driver defined trigger.
- #define **TRIG_WAKE_EOS** 0x0020
Wake up on end-of-scan.
- #define **TRIG_ROUND_MASK** 0x00030000
Trigger not implemented yet.
- #define **TRIG_ROUND_NEAREST** 0x00000000
Trigger not implemented yet.
- #define **TRIG_ROUND_DOWN** 0x00010000
Trigger not implemented yet.
- #define **TRIG_ROUND_UP** 0x00020000
Trigger not implemented yet.
- #define **TRIG_ROUND_UP_NEXT** 0x00030000
Trigger not implemented yet.

Channel macros

Specific precompilation macros and constants useful for the channels descriptors tab located in the command structure

- #define **CHAN**(a) ((a) & 0xffff)
Channel indication macro.
- #define **RNG**(a) (((a) & 0xff) << 16)
Range definition macro.
- #define **AREF**(a) (((a) & 0xf) << 24)
Reference definition macro.
- #define **FLAGS**(a) ((a) & CR_FLAGS_MASK)
Flags definition macro.
- #define **PACK**(a, b, c) (**CHAN**(a) | **RNG**(b) | **AREF**(c))
Channel + range + reference definition macro.

- #define `PACK_FLAGS(a, b, c, d)` (`CHAN(a) | RNG(b) | AREF(c) | FLAGS(d)`)
Channel + range + reference + flags definition macro.
- #define `AREF_GROUND` 0x00
Analog reference is analog ground.
- #define `AREF_COMMON` 0x01
Analog reference is analog common.
- #define `AREF_DIFF` 0x02
Analog reference is differential.
- #define `AREF_OTHER` 0x03
Analog reference is undefined.

5.60.1 Detailed Description

5.60.2 Function Documentation

5.60.2.1 `int a4l_get_bufsize (a4l_desc_t * dsc, unsigned int idx_subd, unsigned long * size)`

Get the size of the asynchronous buffer.

During asynchronous acquisition, a ring-buffer enables the transfers from / to user-space. Functions like `a4l_read()` or `a4l_write()` recovers / sends data through this intermediate buffer. Please note, there is one ring-buffer per subdevice capable of asynchronous acquisition. By default, each buffer size is set to 64 KB.

Parameters

in	<i>dsc</i>	Device descriptor filled by <code>a4l_open()</code> (and optionally <code>a4l_fill_desc()</code>)
in	<i>idx_subd</i>	Index of the concerned subdevice
out	<i>size</i>	Buffer size

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong

References `a4l_descriptor::fd`.

5.60.2.2 `int a4l_mark_bufwr (a4l_desc_t * dsc, unsigned int idx_subd, unsigned long cur, unsigned long * new)`

Update the asynchronous buffer state.

When the mapping of the asynchronous ring-buffer (thanks to `a4l_mmap()` is disabled, common read / write syscalls have to be used. In input case, `a4l_read()` must be used for:

- the retrieval of the acquired data.
- the notification to the Analogy layer that the acquired data have been consumed, then the area in the ring-buffer which was containing becomes available. In output case, `a4l_write()` must be called to:
- send some data to the Analogy layer.
- signal the Analogy layer that a chunk of data in the ring-buffer must be used by the driver.

In mmap configuration, these features are provided by unique function named [a4l_mark_bufwr\(\)](#). In input case, [a4l_mark_bufwr\(\)](#) can :

- recover the count of data newly available in the ring-buffer.
- notify the Analogy layer how many bytes have been consumed. In output case, [a4l_mark_bufwr\(\)](#) can:
- recover the count of data available for writing.
- notify Analogy that some bytes have been written.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>cur</i>	Amount of consumed data
out	<i>new</i>	Amount of available data

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; the descriptor and the new pointer should be checked; check also the kernel log ("dmesg")
- -EFAULT is returned if a user <-> kernel transfer went wrong

References [a4l_descriptor::fd](#).

5.60.2.3 int [a4l_mmap](#) ([a4l_desc_t](#) * dsc, unsigned int idx_subd, unsigned long size, void ** ptr)

Map the asynchronous ring-buffer into a user-space.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>size</i>	Size of the buffer to map
out	<i>ptr</i>	Address of the pointer containing the assigned address on return

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong, the descriptor and the pointer should be checked; check also the kernel log
- -EPERM is returned if the function is called in an RT context or if the buffer to resize is mapped in user-space (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EBUSY is returned if the buffer is already mapped in user-space

References [a4l_descriptor::fd](#).

5.60.2.4 int [a4l_poll](#) ([a4l_desc_t](#) * dsc, unsigned int idx_subd, unsigned long ms_timeout)

Get the available data count.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>ms_timeout</i>	The number of miliseconds to wait for some data to be available. Passing A4L_INFINITE causes the caller to block indefinitely until some data is available. Passing A4L_NONBLOCK causes the function to return immediately without waiting for any available data

Returns

the available data count. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EINTR is returned if calling task has been unblocked by a signal

References [a4l_descriptor::fd](#).

Referenced by [a4l_async_read\(\)](#), and [a4l_async_write\(\)](#).

5.60.2.5 `int a4l_set_bufsize (a4l_desc_t * dsc, unsigned int idx_subd, unsigned long size)`

Change the size of the asynchronous buffer.

During asynchronous acquisition, a ring-buffer enables the transfers from / to user-space. Functions like [a4l_read\(\)](#) or [a4l_write\(\)](#) recovers / sends data through this intermediate buffer. The function [a4l_set_bufsize\(\)](#) can change the size of the ring-buffer. Please note, there is one ring-buffer per subdevice capable of asynchronous acquisition. By default, each buffer size is set to 64 KB.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>size</i>	New buffer size, the maximal tolerated value is 16MB (A4L_BUF_MAXSIZE)

Returns

0 on success. Otherwise:

- -EINVAL is returned if the analogy descriptor is not correct or if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EPERM is returned if the function is called in an RT context or if the buffer to resize is mapped in user-space (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EBUSY is returned if the selected subdevice is already processing an asynchronous operation
- -ENOMEM is returned if the system is out of memory

References [a4l_sys_bufcfg\(\)](#), and [a4l_descriptor::fd](#).

5.60.2.6 `int a4l_snd_cancel (a4l_desc_t * dsc, unsigned int idx_subd)`

Cancel an asynchronous acquisition.

The function [a4l_snd_cancel\(\)](#) is devoted to stop an asynchronous acquisition configured thanks to an Analogy command.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Subdevice index

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EIO is returned if the selected subdevice does not support asynchronous operation

References `a4l_descriptor::fd`.

5.60.2.7 `int a4l_snd_command (a4l_desc_t * dsc, a4l_cmd_t * cmd)`

Send a command to an Analogy device.

The function [a4l_snd_command\(\)](#) triggers asynchronous acquisition.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>cmd</i>	Command structure

Returns


0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -ENOMEM is returned if the system is out of memory
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EIO is returned if the selected subdevice cannot handle command
- -EBUSY is returned if the selected subdevice is already processing an asynchronous operation

References `a4l_descriptor::fd`.

5.61 Asynchronous acquisition API

Collaboration diagram for Asynchronous acquisition API:



Functions

- int [a4l_async_read](#) ([a4l_desc_t](#) *dsc, void *buf, size_t nbyte, unsigned long ms_timeout)
Perform asynchronous read operation on the analog input subdevice.
- int [a4l_async_write](#) ([a4l_desc_t](#) *dsc, void *buf, size_t nbyte, unsigned long ms_timeout)
Perform asynchronous write operation on the analog input subdevice.

5.61.1 Detailed Description

5.61.2 Function Documentation

5.61.2.1 int [a4l_async_read](#) ([a4l_desc_t](#) * dsc, void * buf, size_t nbyte, unsigned long ms_timeout)

Perform asynchronous read operation on the analog input subdevice.

The function [a4l_async_read\(\)](#) is only useful for acquisition configured through an Analogy command.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
out	<i>buf</i>	Input buffer
in	<i>nbyte</i>	Number of bytes to read
in	<i>ms_timeout</i>	The number of milliseconds to wait for some data to be available. Passing A4L_INFINITE causes the caller to block indefinitely until some data is available. Passing A4L_NONBLOCK causes the function to return immediately without waiting for any available data

Returns

Number of bytes read, otherwise negative error code:

- -EINVAL is returned if some argument is missing or wrong, the descriptor should be checked; check also the kernel log
- -ENOENT is returned if the device's reading subdevice is idle (no command was sent)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EINTR is returned if calling task has been unblocked by a signal

References [a4l_poll\(\)](#), [a4l_sys_read\(\)](#), [a4l_descriptor::fd](#), and [a4l_descriptor::idx_read_subd](#).


```
5.61.2.2 int a4l_async_write ( a4l_desc_t * dsc, void * buf, size_t nbyte, unsigned long ms_timeout
    )
```

Perform asynchronous write operation on the analog input subdevice.

The function [a4l_async_write\(\)](#) is only useful for acquisition configured through an Analogy command.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>buf</i>	Output buffer
in	<i>nbyte</i>	Number of bytes to write
in	<i>ms_timeout</i>	The number of milliseconds to wait for some free area to be available. Passing A4L_INFINITE causes the caller to block indefinitely until some data is available. Passing A4L_NONBLOCK causes the function to return immediately without waiting any available space to write data.

Returns

Number of bytes written, otherwise negative error code:

- -EINVAL is returned if some argument is missing or wrong, the descriptor should be checked; check also the kernel log
- -ENOENT is returned if the device's reading subdevice is idle (no command was sent)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EINTR is returned if calling task has been unblocked by a signal

References [a4l_poll\(\)](#), [a4l_sys_write\(\)](#), [a4l_descriptor::fd](#), and [a4l_descriptor::idx_write_subd](#).

5.62 Descriptor Syscall API

Collaboration diagram for Descriptor Syscall API:



Data Structures

- struct [a4l_descriptor](#)
Structure containing device-information useful to users.

Functions

- int [a4l_sys_desc](#) (int fd, [a4l_desc_t](#) *dsc, int pass)
Get a descriptor on an attached device.

ANALOGY_XXX_DESC

Constants used as argument so as to define the description depth to recover

- #define [A4L_BSC_DESC](#) 0x0
BSC stands for basic descriptor (device data)
- #define [A4L_CPLX_DESC](#) 0x1
CPLX stands for complex descriptor (subdevice + channel + range data)

5.62.1 Detailed Description

5.62.2 Function Documentation

5.62.2.1 int a4l_sys_desc (int fd, [a4l_desc_t](#) * dsc, int pass)

Get a descriptor on an attached device.

Once the device has been attached, the function `a4l_get_desc()` retrieves various information on the device (subdevices, channels, ranges, etc.). The function `a4l_get_desc()` can be called twice:

- The first time, almost all the fields, except `sbddata`, are set (`board_name`, `nb_subd`, `idx_read_subd`, `idx_write_subd`, `magic`, `sbsize`); the last field, `sbddata`, is supposed to be a pointer on a buffer, which size is defined by the field `sbsize`.
- The second time, the buffer pointed by `sbddata` is filled with data about the subdevices, the channels and the ranges.

Between the two calls, an allocation must be performed in order to recover a buffer large enough to contain all the data. These data are set up according a root-leaf organization (device -> subdevice -> channel -> range). They cannot be accessed directly; specific functions are available so as to retrieve them:

- [a4l_get_subdinfo\(\)](#) to get some subdevice's characteristics.
- [a4l_get_chaninfo\(\)](#) to get some channel's characteristics.
- [a4l_get_rnginfo\(\)](#) to get some range's characteristics.

Parameters

in	<i>fd</i>	Driver file descriptor
out	<i>dsc</i>	Device descriptor
in	<i>pass</i>	Description level to retrieve: <ul style="list-style-type: none"> • A4L_BSC_DESC to get the basic descriptor (notably the size of the data buffer to allocate). • A4L_CPLX_DESC to get the complex descriptor, the data buffer is filled with characteristics about the subdevices, the channels and the ranges.

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; the pass argument should be checked; check also the kernel log ("dmesg")
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -ENODEV is returned if the descriptor is incoherent (the device may be unattached)

References `A4L_BSC_DESC`, `a4l_descriptor::magic`, `a4l_descriptor::sbdata`, and `a4l_descriptor::sbsize`.

Referenced by `a4l_fill_desc()`, and `a4l_open()`.

5.63 Descriptor API

Collaboration diagram for Descriptor API:



Functions

- `int a4l_open (a4l_desc_t *dsc, const char *fname)`
Open an Analogy device and basically fill the descriptor.
- `int a4l_close (a4l_desc_t *dsc)`
Close the Analogy device related with the descriptor.
- `int a4l_fill_desc (a4l_desc_t *dsc)`
Fill the descriptor with subdevices, channels and ranges data.
- `int a4l_get_subdinfo (a4l_desc_t *dsc, unsigned int subd, a4l_sbinfo_t **info)`
Get an information structure on a specified subdevice.
- `int a4l_get_chinfo (a4l_desc_t *dsc, unsigned int subd, unsigned int chan, a4l_chinfo_t **info)`
Get an information structure on a specified channel.
- `int a4l_get_rnginfo (a4l_desc_t *dsc, unsigned int subd, unsigned int chan, unsigned int rng, a4l_rnginfo_t **info)`
Get an information structure on a specified range.

5.63.1 Detailed Description

This is the API interface used to fill and use Analogy device descriptor structure

5.63.2 Function Documentation

5.63.2.1 `int a4l_close (a4l_desc_t * dsc)`

Close the Analogy device related with the descriptor.

The file descriptor is associated with a context. The context is one of the enabler of asynchronous transfers. So, by closing the file descriptor, the programmer must keep in mind that the currently occurring asynchronous transfer will be cancelled.

Parameters

in	dsc	Device descriptor
----	-----	-------------------

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; the the dsc pointer should be checked; check also the kernel log ("dmesg")

References `a4l_sys_close()`, and `a4l_descriptor::fd`.

5.63.2.2 `int a4l_fill_desc (a4l_desc_t * dsc)`

Fill the descriptor with subdevices, channels and ranges data.

Parameters

in	<i>dsc</i>	Device descriptor partly filled by a4l_open() .
-----------	------------	---

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; the the dsc pointer should be checked; check also the kernel log ("dmesg")
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -ENODEV is returned if the descriptor is incoherent (the device may be unattached)

References `A4L_CPLX_DESC`, `a4l_sys_desc()`, `a4l_descriptor::fd`, and `a4l_descriptor::magic`.

5.63.2.3 `int a4l_get_chinfo (a4l_desc_t * dsc, unsigned int subd, unsigned int chan, a4l_chinfo_t ** info)`

Get an information structure on a specified channel.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() and a4l_fill_desc()
in	<i>subd</i>	Subdevice index
in	<i>chan</i>	Channel index
out	<i>info</i>	Channel information structure

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; subd, chan and the dsc pointer should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_get_chinfo\(\)](#)

References `a4l_descriptor::magic`, `a4l_descriptor::nb_subd`, and `a4l_descriptor::sbdata`.

Referenced by `a4l_find_range()`.

5.63.2.4 `int a4l_get_rnginfo (a4l_desc_t * dsc, unsigned int subd, unsigned int chan, unsigned int rng, a4l_rnginfo_t ** info)`

Get an information structure on a specified range.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() and a4l_fill_desc()
in	<i>subd</i>	Subdevice index
in	<i>chan</i>	Channel index
in	<i>rng</i>	Range index
out	<i>info</i>	Range information structure

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; subd, chan, rng and the dsc pointer should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_get_rnginfo\(\)](#)

References [a4l_descriptor::magic](#), [a4l_descriptor::nb_subd](#), and [a4l_descriptor::sbddata](#).

Referenced by [a4l_find_range\(\)](#).

5.63.2.5 int [a4l_get_subdinfo](#) (**a4l_desc_t** * dsc, unsigned int subd, **a4l_sbinfo_t** ** info)

Get an information structure on a specified subdevice.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() and a4l_fill_desc()
in	<i>subd</i>	Subdevice index
out	<i>info</i>	Subdevice information structure

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; subd and the dsc pointer should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_get_subdinfo\(\)](#).

References [a4l_descriptor::magic](#), [a4l_descriptor::nb_subd](#), and [a4l_descriptor::sbddata](#).

Referenced by [a4l_sync_dio\(\)](#).

5.63.2.6 int [a4l_open](#) (**a4l_desc_t** * dsc, const char * fname)

Open an Analogy device and basically fill the descriptor.

Parameters

out	<i>dsc</i>	Device descriptor
in	<i>fname</i>	Device name

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong; the fname and the dsc pointer should be checked; check also the kernel log ("dmesg")
- -EFAULT is returned if a user <-> kernel transfer went wrong

References [A4L_BSC_DESC](#), [a4l_sys_close\(\)](#), [a4l_sys_desc\(\)](#), [a4l_sys_open\(\)](#), and [a4l_descriptor::fd](#).

5.64 Range / conversion API

Collaboration diagram for Range / conversion API:



Functions

- int [a4l_sizeof_chan](#) (a4l_chinfo_t *chan)
Get the size in memory of an acquired element.
- int [a4l_sizeof_subd](#) (a4l_sbinfo_t *subd)
Get the size in memory of a digital acquired element.
- int [a4l_find_range](#) (a4l_desc_t *dsc, unsigned int idx_subd, unsigned int idx_chan, unsigned long unit, double min, double max, a4l_rnginfo_t **rng)
Find the most suitable range.
- int [a4l_rawtoul](#) (a4l_chinfo_t *chan, unsigned long *dst, void *src, int cnt)
Unpack raw data (from the driver) into unsigned long values.
- int [a4l_rawtof](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, float *dst, void *src, int cnt)
Convert raw data (from the driver) to float-typed samples.
- int [a4l_rawtod](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, double *dst, void *src, int cnt)
Convert raw data (from the driver) to double-typed samples.
- int [a4l_ultoraw](#) (a4l_chinfo_t *chan, void *dst, unsigned long *src, int cnt)
Pack unsigned long values into raw data (for the driver)
- int [a4l_ftoraw](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, void *dst, float *src, int cnt)
Convert float-typed samples to raw data (for the driver)
- int [a4l_dtoraw](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, void *dst, double *src, int cnt)
Convert double-typed samples to raw data (for the driver)

5.64.1 Detailed Description

5.64.2 Function Documentation

5.64.2.1 int a4l_dtoraw (a4l_chinfo_t * chan, a4l_rnginfo_t * rng, void * dst, double * src, int cnt)

Convert double-typed samples to raw data (for the driver)

Parameters

in	chan	Channel descriptor
in	rng	Range descriptor

out	<i>dst</i>	Output buffer
in	<i>src</i>	Input buffer
in	<i>cnt</i>	Count of conversion to perform

Returns

the count of conversion performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; *chan*, *rng* and the pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_dtoraw\(\)](#)

References `A4L_RNG_FACTOR`, and `a4l_sizeof_chan()`.

5.64.2.2 `int a4l_find_range (a4l_desc_t * dsc, unsigned int idx_subd, unsigned int idx_chan, unsigned long unit, double min, double max, a4l_rnginfo_t ** rng)`

Find the most suitable range.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() and a4l_fill_desc()
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>idx_chan</i>	Index of the concerned channel
in	<i>unit</i>	Unit type used in the range
in	<i>min</i>	Minimal limit value
in	<i>max</i>	Maximal limit value
out	<i>rng</i>	Found range

Returns

The index of the most suitable range on success. Otherwise:

- -ENOENT is returned if a suitable range is not found.
- -EINVAL is returned if some argument is missing or wrong; *idx_subd*, *idx_chan* and the *dsc* pointer should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_find_range\(\)](#)

References `a4l_get_chinfo()`, `a4l_get_rnginfo()`, `A4L_RNG_FACTOR`, `A4L_RNG_UNIT`, and `a4l_descriptor::magic`.

5.64.2.3 `int a4l_ftoraw (a4l_chinfo_t * chan, a4l_rnginfo_t * rng, void * dst, float * src, int cnt)`

Convert float-typed samples to raw data (for the driver)

Parameters

in	<i>chan</i>	Channel descriptor
in	<i>rng</i>	Range descriptor
out	<i>dst</i>	Output buffer
in	<i>src</i>	Input buffer

<i>in</i>	<i>cnt</i>	Count of conversion to perform
-----------	------------	--------------------------------

Returns

the count of conversion performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; *chan*, *rng* and the pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_ftoraw\(\)](#)

References A4L_RNG_FACTOR, and [a4l_sizeof_chan\(\)](#).

5.64.2.4 `int a4l_rawtod (a4l_chinfo_t * chan, a4l_rnginfo_t * rng, double * dst, void * src, int cnt)`

Convert raw data (from the driver) to double-typed samples.

Parameters

<i>in</i>	<i>chan</i>	Channel descriptor
<i>in</i>	<i>rng</i>	Range descriptor
<i>out</i>	<i>dst</i>	Ouput buffer
<i>in</i>	<i>src</i>	Input buffer
<i>in</i>	<i>cnt</i>	Count of conversion to perform

Returns

the count of conversion performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; *chan*, *rng* and the pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_rawtod\(\)](#)

References A4L_RNG_FACTOR, and [a4l_sizeof_chan\(\)](#).

5.64.2.5 `int a4l_rawtof (a4l_chinfo_t * chan, a4l_rnginfo_t * rng, float * dst, void * src, int cnt)`

Convert raw data (from the driver) to float-typed samples.

Parameters

<i>in</i>	<i>chan</i>	Channel descriptor
<i>in</i>	<i>rng</i>	Range descriptor
<i>out</i>	<i>dst</i>	Ouput buffer
<i>in</i>	<i>src</i>	Input buffer
<i>in</i>	<i>cnt</i>	Count of conversion to perform

Returns

the count of conversion performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; *chan*, *rng* and the pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_rawtof\(\)](#)

References A4L_RNG_FACTOR, and [a4l_sizeof_chan\(\)](#).

5.64.2.6 `int a4l_rawtoul (a4l_chinfo_t * chan, unsigned long * dst, void * src, int cnt)`

Unpack raw data (from the driver) into unsigned long values.

This function takes as input driver-specific data and scatters each element into an entry of an unsigned long table. It is a convenience routine which performs no conversion, just copy.

Parameters

in	<i>chan</i>	Channel descriptor
out	<i>dst</i>	Output buffer
in	<i>src</i>	Input buffer
in	<i>cnt</i>	Count of transfer to copy

Returns

the count of copy performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; *chan*, *dst* and *src* pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_ultoraw\(\)](#)

References [a4l_sizeof_chan\(\)](#).

5.64.2.7 `int a4l_sizeof_chan (a4l_chinfo_t * chan)`

Get the size in memory of an acquired element.

According to the board, the channels have various acquisition widths. With values like 8, 16 or 32, there is no problem finding out the size in memory (1, 2, 4); however with widths like 12 or 24, this function might be helpful to guess the size needed in RAM for a single acquired element.

Parameters

in	<i>chan</i>	Channel descriptor
----	-------------	--------------------

Returns

the size in memory of an acquired element, otherwise a negative error code:

- -EINVAL is returned if the argument *chan* is NULL

Referenced by [a4l_dtoraw\(\)](#), [a4l_ftoraw\(\)](#), [a4l_rawtod\(\)](#), [a4l_rawtof\(\)](#), [a4l_rawtoul\(\)](#), and [a4l_ultoraw\(\)](#).

5.64.2.8 `int a4l_sizeof_subd (a4l_sbinfo_t * subd)`

Get the size in memory of a digital acquired element.

This function is only useful for DIO subdevices. Digital subdevices are a specific kind of subdevice on which channels are regarded as bits composing the subdevice's bitfield. During a DIO acquisition, all bits are sampled. Therefore, [a4l_sizeof_chan\(\)](#) is useless in this case and we have to use [a4l_sizeof_subd\(\)](#). With bitfields which sizes are 8, 16 or 32, there is no problem finding out the size in memory (1, 2, 4); however with widths like 12 or 24, this function might be helpful to guess the size needed in RAM for a single acquired element.

Parameters

in	<i>subd</i>	Subdevice descriptor
-----------	-------------	----------------------

Returns

the size in memory of an acquired element, otherwise a negative error code:

- -EINVAL is returned if the argument chan is NULL or if the subdevice is not a digital subdevice

References A4L_SUBD_DI, A4L_SUBD_DIO, A4L_SUBD_DO, and A4L_SUBD_TYPES.

Referenced by a4l_sync_dio().

5.64.2.9 int a4l_ultoraw (a4l_chinfo_t * chan, void * dst, unsigned long * src, int cnt)

Pack unsigned long values into raw data (for the driver)

This function takes as input a table of unsigned long values and gather them according to the channel width. It is a convenience routine which performs no conversion, just formatting.

Parameters

in	<i>chan</i>	Channel descriptor
out	<i>dst</i>	Ouput buffer
in	<i>src</i>	Input buffer
in	<i>cnt</i>	Count of transfer to copy

Returns


the count of copy performed, otherwise a negative error code:

- -EINVAL is returned if some argument is missing or wrong; chan, dst and src pointers should be checked; check also the kernel log ("dmesg"); WARNING: [a4l_fill_desc\(\)](#) should be called before using [a4l_ultoraw\(\)](#)

References a4l_sizeof_chan().

5.65 Level 1 API

Collaboration diagram for Level 1 API:




Modules

- [Asynchronous acquisition API](#)
- [Descriptor API](#)
- [Synchronous acquisition API](#)

5.65.1 Detailed Description

5.66 Synchronous acquisition API

Collaboration diagram for Synchronous acquisition API:



Data Structures

- struct [a4l_instruction](#)
Structure describing the synchronous instruction.
- struct [a4l_instruction_list](#)
Structure describing the list of synchronous instructions.

Macros

- #define [A4L_INSN_WAIT_MAX](#) 100000
Maximal wait duration.

Functions

- int [a4l_snd_insnlist](#) ([a4l_desc_t](#) *dsc, [a4l_insnlst_t](#) *arg)
Perform a list of synchronous acquisition misc operations.
- int [a4l_snd_insn](#) ([a4l_desc_t](#) *dsc, [a4l_insn_t](#) *arg)
Perform a synchronous acquisition misc operation.

Instruction type

Flags to define the type of instruction

- #define [A4L_INSN_READ](#) (0 | A4L_INSN_MASK_READ)
Read instruction.
- #define [A4L_INSN_WRITE](#) (1 | A4L_INSN_MASK_WRITE)
Write instruction.
- #define [A4L_INSN_BITS](#)
"Bits" instruction
- #define [A4L_INSN_CONFIG](#)
Configuration instruction.
- #define [A4L_INSN_GTOD](#)
Get time instruction.
- #define [A4L_INSN_WAIT](#)
Wait instruction.
- #define [A4L_INSN_INTTRIG](#)
Trigger instruction (to start asynchronous acquisition)

Configuration instruction type

Values to define the type of configuration instruction

- **#define A4L_INSN_CONFIG_DIO_INPUT** 0
- **#define A4L_INSN_CONFIG_DIO_OUTPUT** 1
- **#define A4L_INSN_CONFIG_DIO_OPENDRAIN** 2
- **#define A4L_INSN_CONFIG_ANALOG_TRIG** 16
- **#define A4L_INSN_CONFIG_ALT_SOURCE** 20
- **#define A4L_INSN_CONFIG_DIGITAL_TRIG** 21
- **#define A4L_INSN_CONFIG_BLOCK_SIZE** 22
- **#define A4L_INSN_CONFIG_TIMER_1** 23
- **#define A4L_INSN_CONFIG_FILTER** 24
- **#define A4L_INSN_CONFIG_CHANGE_NOTIFY** 25
- **#define A4L_INSN_CONFIG_SERIAL_CLOCK** 26
- **#define A4L_INSN_CONFIG_BIDIRECTIONAL_DATA** 27
- **#define A4L_INSN_CONFIG_DIO_QUERY** 28
- **#define A4L_INSN_CONFIG_PWM_OUTPUT** 29
- **#define A4L_INSN_CONFIG_GET_PWM_OUTPUT** 30
- **#define A4L_INSN_CONFIG_ARM** 31
- **#define A4L_INSN_CONFIG_DISARM** 32
- **#define A4L_INSN_CONFIG_GET_COUNTER_STATUS** 33
- **#define A4L_INSN_CONFIG_RESET** 34
- **#define A4L_INSN_CONFIG_GPCT_SINGLE_PULSE_GENERATOR** 1001 /* Use CTR as single pulsegenerator */
- **#define A4L_INSN_CONFIG_GPCT_PULSE_TRAIN_GENERATOR** 1002 /* Use CTR as pulse-traingenerator */
- **#define A4L_INSN_CONFIG_GPCT_QUADRATURE_ENCODER** 1003 /* Use the counter as encoder */
- **#define A4L_INSN_CONFIG_SET_GATE_SRC** 2001 /* Set gate source */
- **#define A4L_INSN_CONFIG_GET_GATE_SRC** 2002 /* Get gate source */
- **#define A4L_INSN_CONFIG_SET_CLOCK_SRC** 2003 /* Set master clock source */
- **#define A4L_INSN_CONFIG_GET_CLOCK_SRC** 2004 /* Get master clock source */
- **#define A4L_INSN_CONFIG_SET_OTHER_SRC** 2005 /* Set other source */
- **#define A4L_INSN_CONFIG_SET_COUNTER_MODE** 4097
- **#define A4L_INSN_CONFIG_SET_ROUTING** 4099
- **#define A4L_INSN_CONFIG_GET_ROUTING** 4109

Counter status bits

Status bits for INSN_CONFIG_GET_COUNTER_STATUS

- **#define A4L_COUNTER_ARMED** 0x1
- **#define A4L_COUNTER_COUNTING** 0x2
- **#define A4L_COUNTER_TERMINAL_COUNT** 0x4

IO direction

Values to define the IO polarity

- **#define A4L_INPUT** 0
- **#define A4L_OUTPUT** 1
- **#define A4L_OPENDRAIN** 2

Events types

Values to define the Analogy events. They might used to send some specific events through the instruction interface.

- `#define A4L_EV_START 0x00040000`
- `#define A4L_EV_SCAN_BEGIN 0x00080000`
- `#define A4L_EV_CONVERT 0x00100000`
- `#define A4L_EV_SCAN_END 0x00200000`
- `#define A4L_EV_STOP 0x00400000`

5.66.1 Detailed Description

5.66.2 Function Documentation

5.66.2.1 `int a4l_snd_insn (a4l_desc_t * dsc, a4l_insn_t * arg)`

Perform a synchronous acquisition misc operation.

The function `a4l_snd_insn()` triggers a synchronous acquisition.

Parameters

<code>in</code>	<code>dsc</code>	Device descriptor filled by <code>a4l_open()</code> (and optionally <code>a4l_fill_desc()</code>)
<code>in</code>	<code>arg</code>	Instruction structure

Returns

0 on success. Otherwise:

- `-EINVAL` is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- `-EFAULT` is returned if a user <-> kernel transfer went wrong
- `-ENOMEM` is returned if the system is out of memory

References `a4l_descriptor::fd`.

Referenced by `a4l_config_subd()`, `a4l_sync_dio()`, `a4l_sync_read()`, and `a4l_sync_write()`.

5.66.2.2 `int a4l_snd_insnlist (a4l_desc_t * dsc, a4l_insnlst_t * arg)`

Perform a list of synchronous acquisition misc operations.

The function `a4l_snd_insnlist()` is able to send many synchronous instructions on a various set of sub-devices, channels, etc.

Parameters

<code>in</code>	<code>dsc</code>	Device descriptor filled by <code>a4l_open()</code> (and optionally <code>a4l_fill_desc()</code>)
<code>in</code>	<code>arg</code>	Instructions list structure

Returns

0 on success. Otherwise:

- `-EINVAL` is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- `-EFAULT` is returned if a user <-> kernel transfer went wrong


- -ENOMEM is returned if the system is out of memory

References `a4l_descriptor::fd`.

Referenced by `a4l_sync_read()`, and `a4l_sync_write()`.

5.67 Level 2 API

Collaboration diagram for Level 2 API:




Modules

- [Asynchronous acquisition API](#)
- [Range / conversion API](#)
- [Synchronous acquisition API](#)

5.67.1 Detailed Description

5.68 Synchronous acquisition API

Collaboration diagram for Synchronous acquisition API:



Functions

- int [a4l_sync_write](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void *buf, size_t nbyte)
Perform a synchronous acquisition write operation.
- int [a4l_sync_read](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void *buf, size_t nbyte)
Perform a synchronous acquisition read operation.
- int [a4l_sync_dio](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, void *mask, void *buf)
Perform a synchronous acquisition digital acquisition.
- int [a4l_config_subd](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int type,...)
Configure a subdevice.

5.68.1 Detailed Description

5.68.2 Function Documentation

5.68.2.1 int [a4l_config_subd](#) (**a4l_desc_t** * dsc, unsigned int idx_subd, unsigned int type, ...)

Configure a subdevice.

[a4l_config_subd\(\)](#) takes a variable count of arguments. According to the configuration type, some additional argument is necessary:

- A4L_INSN_CONFIG_DIO_INPUT: the channel index (unsigned int)
- A4L_INSN_CONFIG_DIO_OUTPUT: the channel index (unsigned int)
- A4L_INSN_CONFIG_DIO_QUERY: the returned DIO polarity (unsigned int *)

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>type</i>	Configuration parameter

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -ENOSYS is returned if the configuration parameter is not supported

References A4L_INSN_CONFIG, a4l_snd_insn(), CHAN, a4l_instruction::chan_desc, a4l_instruction::data_size, and a4l_instruction::type.

5.68.2.2 int a4l_sync_dio (**a4l_desc_t** * dsc, unsigned int idx_subd, void * mask, void * buf)

Perform a synchronous acquisition digital acquisition.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>mask</i>	Write mask which indicates which bit(s) must be modified
in,out	<i>buf</i>	Input / output buffer

Returns

Number of bytes read, otherwise negative error code:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -ENOMEM is returned if the system is out of memory
- -ENOSYS is returned if the driver does not provide any handler "instruction bits"

References a4l_get_subinfo(), A4L_INSN_BITS, a4l_sizeof_subd(), a4l_snd_insn(), a4l_instruction::data_size, and a4l_instruction::type.

5.68.2.3 int a4l_sync_read (**a4l_desc_t** * dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void * buf, size_t nbyte)

Perform a synchronous acquisition read operation.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>chan_desc</i>	Channel descriptor (channel, range and reference)
in	<i>ns_delay</i>	Optional delay (in nanoseconds) to wait between the setting of the input channel and sample(s) acquisition(s).
in	<i>buf</i>	Input buffer
in	<i>nbyte</i>	Number of bytes to read

Returns

Number of bytes read, otherwise negative error code:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong

- -ENOMEM is returned if the system is out of memory

References A4L_INSN_READ, A4L_INSN_WAIT, a4l_snd_insn(), a4l_snd_insnlist(), a4l_instruction_list::count, a4l_instruction::data, a4l_instruction::data_size, and a4l_instruction::type.

5.68.2.4 int a4l_sync_write (**a4l_desc_t** * dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void * buf, size_t nbyte)

Perform a synchronous acquisition write operation.

Parameters

in	<i>dsc</i>	Device descriptor filled by a4l_open() (and optionally a4l_fill_desc())
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>chan_desc</i>	Channel descriptor (channel, range and reference)
in	<i>ns_delay</i>	Optional delay (in nanoseconds) to wait between the setting of the input channel and sample(s) acquisition(s).
in	<i>buf</i>	Output buffer
in	<i>nbyte</i>	Number of bytes to write

Returns


Number of bytes written, otherwise negative error code:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -ENOMEM is returned if the system is out of memory

References A4L_INSN_WAIT, A4L_INSN_WRITE, a4l_snd_insn(), a4l_snd_insnlist(), a4l_instruction_list::count, a4l_instruction::data, a4l_instruction::data_size, and a4l_instruction::type.

5.69 Analogy user API

Collaboration diagram for Analogy user API:



Modules


- [Level 1 API](#)
- [Level 2 API](#)
- [Level 0 API](#)

5.69.1 Detailed Description

This is the API interface of Analogy library

5.70 Level 0 API

Collaboration diagram for Level 0 API:



Modules

- [Descriptor Syscall API](#)
- [Basic Syscall API](#)
- [Attach / detach Syscall API](#)

5.70.1 Detailed Description

System call interface to core Analogy services

This interface should not be used directly by applications.

5.71 Basic Syscall API

Collaboration diagram for Basic Syscall API:



Functions

- int [a4l_sys_open](#) (const char *fname)
Open an Analogy device.
- int [a4l_sys_close](#) (int fd)
Close an Analogy device.
- int [a4l_sys_read](#) (int fd, void *buf, size_t nbyte)
Read from an Analogy device.
- int [a4l_sys_write](#) (int fd, void *buf, size_t nbyte)
Write to an Analogy device.

5.71.1 Detailed Description

5.71.2 Function Documentation

5.71.2.1 int [a4l_sys_close](#) (int fd)

Close an Analogy device.

Parameters

in	<i>fd</i>	File descriptor as returned by a4l_sys_open()
----	-----------	---

Returns

0 on success, otherwise a negative error code.

Referenced by [a4l_close\(\)](#), and [a4l_open\(\)](#).

5.71.2.2 int [a4l_sys_open](#) (const char * fname)

Open an Analogy device.

Parameters

in	<i>fname</i>	Device name
----	--------------	-------------

Returns

Positive file descriptor value on success, otherwise a negative error code.

Referenced by [a4l_open\(\)](#).

5.71.2.3 `int a4l_sys_read (int fd, void * buf, size_t nbyte)`

Read from an Analogy device.

The function `a4l_read()` is only useful for acquisition configured through an Analogy command.

Parameters

<code>in</code>	<code>fd</code>	File descriptor as returned by a4l_sys_open()
<code>out</code>	<code>buf</code>	Input buffer
<code>in</code>	<code>nbyte</code>	Number of bytes to read

Returns

Number of bytes read. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -ENOENT is returned if the device's reading subdevice is idle (no command was sent)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EINTR is returned if calling task has been unblocked by a signal

Referenced by `a4l_async_read()`.

5.71.2.4 `int a4l_sys_write (int fd, void * buf, size_t nbyte)`

Write to an Analogy device.

The function `a4l_write()` is only useful for acquisition configured through an Analogy command.

Parameters

<code>in</code>	<code>fd</code>	File descriptor as returned by a4l_sys_open()
<code>in</code>	<code>buf</code>	Output buffer
<code>in</code>	<code>nbyte</code>	Number of bytes to write

Returns


Number of bytes written. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -ENOENT is returned if the device's writing subdevice is idle (no command was sent)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EINTR is returned if calling task has been unblocked by a signal

Referenced by `a4l_async_write()`.

5.72 Attach / detach Syscall API

Collaboration diagram for Attach / detach Syscall API:



Functions

- int [a4l_sys_attach](#) (int fd, a4l_lnkdesc_t *arg)
Attach an Analogy device to a driver.
- int [a4l_sys_detach](#) (int fd)
Detach an Analogy device from a driver.
- int [a4l_sys_bufcfg](#) (int fd, unsigned int idx_subd, unsigned long size)
Configure the buffer size.

5.72.1 Detailed Description

5.72.2 Function Documentation

5.72.2.1 int a4l_sys_attach (int fd, a4l_lnkdesc_t * arg)

Attach an Analogy device to a driver.

Parameters

in	<i>fd</i>	File descriptor as returned by a4l_sys_open()
in	<i>arg</i>	Link descriptor argument

Returns

0 on success. Otherwise:

- -ENOMEM is returned if the system is out of memory
- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -ENODEV is returned in case of internal error (Please, type "dmesg" for more info)
- -ENXIO is returned in case of internal error (Please, type "dmesg" for more info)

5.72.2.2 int a4l_sys_bufcfg (int fd, unsigned int idx_subd, unsigned long size)

Configure the buffer size.

This function can configure the buffer size of the file descriptor currently in use. If the subdevice index is set to A4L_BUF_DEFMAGIC, it can also define the default buffer size at open time.

Parameters

in	<i>fd</i>	File descriptor as returned by a4l_sys_open()
in	<i>idx_subd</i>	Index of the concerned subdevice
in	<i>size</i>	Buffer size to be set

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EPERM is returned if the function is called in an RT context or if the buffer to resize is mapped in user-space (Please, type "dmesg" for more info)
- -EFAULT is returned if a user <-> kernel transfer went wrong
- -EBUSY is returned if the selected subdevice is already processing an asynchronous operation
- -ENOMEM is returned if the system is out of memory

Referenced by `a4l_set_bufsize()`.

5.72.2.3 `int a4l_sys_detach (int fd)`

Detach an Analogy device from a driver.

Parameters

in	<i>fd</i>	File descriptor as returned by a4l_sys_open()
----	-----------	---

Returns

0 on success. Otherwise:

- -EINVAL is returned if some argument is missing or wrong (Please, type "dmesg" for more info)
- -EBUSY is returned if the device to be detached is in use
- -EPERM is returned if the device to be detached still has some buffer mapped in user-space
- -ENODEV is returned in case of internal error (Please, type "dmesg" for more info)
- -ENXIO is returned in case of internal error (Please, type "dmesg" for more info)

5.73 Alarm services

General-purpose watchdog timers.

Collaboration diagram for Alarm services:



Data Structures

- struct [RT_ALARM_INFO](#)
Alarm status descriptor.

Functions

- int [rt_alarm_create](#) (RT_ALARM *alarm, const char *name, void(*handler)(void *arg), void *arg)
Create an alarm object.
- int [rt_alarm_delete](#) (RT_ALARM *alarm)
Delete an alarm.
- int [rt_alarm_start](#) (RT_ALARM *alarm, RTIME value, RTIME interval)
Start an alarm.
- int [rt_alarm_stop](#) (RT_ALARM *alarm)
Stop an alarm.
- int [rt_alarm_inquire](#) (RT_ALARM *alarm, [RT_ALARM_INFO](#) *info)
Query alarm status.

5.73.1 Detailed Description

General-purpose watchdog timers. Alarms are general-purpose watchdog timers. Alchemy tasks may create any number of alarms and use them to run a user-defined handler, after a specified initial delay has elapsed. Alarms can be either one shot or periodic; in the latter case, the real-time system automatically reprograms the alarm for the next shot according to a user-defined interval value.

5.73.2 Function Documentation

5.73.2.1 int [rt_alarm_create](#) (RT_ALARM * alarm, const char * name, void(*) (void *arg) handler, void * arg)

Create an alarm object.

This routine creates an object triggering an alarm routine at a specified time in the future. Alarms can be periodic or oneshot, depending on the reload interval value passed to [rt_alarm_start\(\)](#).

Parameters

<i>alarm</i>	The address of an alarm descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the alarm. When non-NULL and non-empty, a copy of this string is used for indexing the created alarm into the object registry.
<i>handler</i>	The address of the routine to call when the alarm expires. This routine is passed the <i>arg</i> value.
<i>arg</i>	A user-defined opaque argument passed to the <i>handler</i> .

Returns

Zero is returned upon success. Otherwise:

- -ENOMEM is returned if the system fails to get memory from the local pool in order to create the alarm.
- -EEXIST is returned if the *name* is conflicting with an already registered alarm.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Alarms are process-private objects and thus cannot be shared by multiple processes, even if they belong to the same Xenomai session.

5.73.2.2 `int rt_alarm_delete (RT_ALARM * alarm)`

Delete an alarm.

This routine deletes an alarm object previously created by a call to [rt_alarm_create\(\)](#).

Parameters

<i>alarm</i>	The descriptor address of the deleted alarm.
--------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid alarm descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling contexts:

- Regular POSIX threads
- Xenomai threads

5.73.2.3 `int rt_alarm_inquire (RT_ALARM * alarm, RT_ALARM_INFO * info)`

Query alarm status.

This routine returns the status information about the specified *alarm*.

Parameters

<i>alarm</i>	The descriptor address of the alarm to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid alarm descriptor.

Valid calling context: any.

References RT_ALARM_INFO::active, RT_ALARM_INFO::expiries, and RT_ALARM_INFO::name.

5.73.2.4 int rt_alarm_start (RT_ALARM * alarm, RTIME value, RTIME interval)

Start an alarm.

This routine programs the trigger date of an alarm object. An alarm can be either periodic or oneshot, depending on the *interval* value.

Alarm handlers are always called on behalf of Xenomai's internal timer event routine. Therefore, Xenomai routines which can be called from such handlers are restricted to the set of services available on behalf of an asynchronous context.

This service overrides any previous setup of the expiry date and reload interval for the alarm.

Parameters

<i>alarm</i>	The descriptor address of the started alarm.
<i>value</i>	The relative date of the first expiry, expressed in clock ticks (see note).
<i>interval</i>	The reload value of the alarm. It is a periodic interval value to be used for reprogramming the next alarm shot, expressed in clock ticks (see note). If <i>interval</i> is equal to TM_INFINITE, the alarm will not be reloaded after it has expired.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid alarm descriptor.

Valid calling context: any.

Note

Each of the initial *value* and *interval* is interpreted as a multiple of the Alchemy clock resolution (see `-alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.73.2.5 int rt_alarm_stop (RT_ALARM * alarm)

Stop an alarm.

This routine disables an alarm object, preventing any further expiry until it is re-enabled via [rt_alarm_start\(\)](#).

Parameters

<i>alarm</i>	The descriptor address of the stopped alarm.
--------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid alarm descriptor.

Valid calling context: any.

5.74 Buffer services

Lightweight FIFO IPC mechanism.

Collaboration diagram for Buffer services:



Data Structures

- struct [RT_BUFFER_INFO](#)
Buffer status descriptor.

Macros

- #define [B_PRIO](#) 0x1 /* Pend by task priority order. */
Creation flags.

Functions

- int [rt_buffer_create](#) (RT_BUFFER *bf, const char *name, size_t bufsz, int mode)
Create an IPC buffer.
- int [rt_buffer_delete](#) (RT_BUFFER *bf)
Delete an IPC buffer.
- ssize_t [rt_buffer_write_timed](#) (RT_BUFFER *bf, const void *ptr, size_t size, const struct timespec *abs_timeout)
Write to an IPC buffer.
- static ssize_t [rt_buffer_write_until](#) (RT_BUFFER *bf, const void *ptr, size_t size, RTIME timeout)
Write to an IPC buffer (with absolute scalar timeout).
- static ssize_t [rt_buffer_write](#) (RT_BUFFER *bf, const void *ptr, size_t size, RTIME timeout)
Write to an IPC buffer (with relative scalar timeout).
- ssize_t [rt_buffer_read_timed](#) (RT_BUFFER *bf, void *ptr, size_t size, const struct timespec *abs_timeout)
Read from an IPC buffer.
- static ssize_t [rt_buffer_read_until](#) (RT_BUFFER *bf, void *ptr, size_t size, RTIME timeout)
Read from an IPC buffer (with absolute scalar timeout).
- static ssize_t [rt_buffer_read](#) (RT_BUFFER *bf, void *ptr, size_t size, RTIME timeout)
Read from an IPC buffer (with relative scalar timeout).
- int [rt_buffer_clear](#) (RT_BUFFER *bf)
Clear an IPC buffer.
- int [rt_buffer_inquire](#) (RT_BUFFER *bf, [RT_BUFFER_INFO](#) *info)
Query buffer status.
- int [rt_buffer_bind](#) (RT_BUFFER *bf, const char *name, RTIME timeout)

Bind to an IPC buffer.

- `int rt_buffer_unbind (RT_BUFFER *bf)`

Unbind from an IPC buffer.

5.74.1 Detailed Description

Lightweight FIFO IPC mechanism. A buffer is a lightweight IPC mechanism, implementing a fast, one-way producer-consumer data path. All messages written are buffered in a single memory area in strict FIFO order, until read either in blocking or non-blocking mode.

Message are always atomically handled on the write side (i.e. no interleave, no short writes), whilst only complete messages are normally returned to the read side. However, short reads may happen under a well-defined situation (see note in `rt_buffer_read()`), albeit they can be fully avoided by proper use of the buffer.

5.74.2 Macro Definition Documentation

5.74.2.1 `#define B_PRIO 0x1 /* Pend by task priority order. */`

Creation flags.

Referenced by `rt_buffer_create()`.

5.74.3 Function Documentation

5.74.3.1 `int rt_buffer_bind (RT_BUFFER * bf, const char * name, RTIME timeout)`

Bind to an IPC buffer.

This routine creates a new descriptor to refer to an existing IPC buffer identified by its symbolic name. If the object does not exist on entry, the caller may block until a buffer of the given name is created.

Parameters

<i>bf</i>	The address of a buffer descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the buffer to bind to. This string should match the object name argument passed to <code>rt_buffer_create()</code> .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing <code>TM_INFINITE</code> causes the caller to block indefinitely until the object is registered. Passing <code>TM_NONBLOCK</code> causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if `rt_task_unblock()` was called for the current task before the retrieval has completed.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `TM_NONBLOCK` and the searched object is not registered on entry.
- `-ETIMEDOUT` is returned if the object cannot be retrieved within the specified amount of time.
- `-EPERM` is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.74.3.2 `int rt_buffer_clear (RT_BUFFER * bf)`

Clear an IPC buffer.

This routine empties a buffer from any data.

Parameters

<i>bf</i>	The descriptor address of the buffer to clear.
-----------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *bf* is not a valid buffer descriptor.

Valid calling context: any.

5.74.3.3 `int rt_buffer_create (RT_BUFFER * bf, const char * name, size_t bufsz, int mode)`

Create an IPC buffer.

This routine creates an IPC object that allows tasks to send and receive data asynchronously via a memory buffer. Data may be of an arbitrary length, albeit this IPC is best suited for small to medium-sized messages, since data always have to be copied to the buffer during transit. Large messages may be more efficiently handled by message queues (RT_QUEUE).

Parameters

<i>bf</i>	The address of a buffer descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the buffer. When non-NULL and non-empty, a copy of this string is used for indexing the created buffer into the object registry.
<i>bufsz</i>	The size of the buffer space available to hold data. The required memory is obtained from the main heap.
<i>mode</i>	The buffer creation mode. The following flags can be OR'ed into this bitmask, each of them affecting the new buffer:

- B_FIFO makes tasks pend in FIFO order for reading data from the buffer.
- B_PRIO makes tasks pend in priority order for reading data from the buffer.

This parameter also applies to tasks blocked on the buffer's write side (see [rt_buffer_write\(\)](#)).

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *mode* is invalid or *bufsz* is zero.
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the buffer.
- -EEXIST is returned if the *name* is conflicting with an already registered buffer.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Buffers can be shared by multiple processes which belong to the same Xenomai session.

References B_PRIO.

5.74.3.4 int rt_buffer_delete (RT_BUFFER * bf)

Delete an IPC buffer.

This routine deletes a buffer object previously created by a call to [rt_buffer_create\(\)](#).

Parameters

<i>bf</i>	The descriptor address of the deleted buffer.
-----------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *bf* is not a valid buffer descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.74.3.5 int rt_buffer_inquire (RT_BUFFER * bf, RT_BUFFER_INFO * info)

Query buffer status.

This routine returns the status information about the specified buffer.

Parameters

<i>bf</i>	The descriptor address of the buffer to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *bf* is not a valid buffer descriptor.

Valid calling context: any.

References `RT_BUFFER_INFO::availmem`, `RT_BUFFER_INFO::iwaiters`, `RT_BUFFER_INFO::name`, `RT_BUFFER_INFO::owaiters`, and `RT_BUFFER_INFO::totalmem`.

5.74.3.6 `ssize_t rt_buffer_read (RT_BUFFER * bf, void * ptr, size_t len, RTIME timeout)`
`[inline], [static]`

Read from an IPC buffer (with relative scalar timeout).

This routine is a variant of [rt_buffer_read_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>bf</i>	The descriptor address of the buffer to read from.
<i>ptr</i>	A pointer to a memory area which will be written upon success with the received data.
<i>len</i>	The length in bytes of the memory area pointed to by <i>ptr</i> .
<i>timeout</i>	A delay expressed in clock ticks.

References `rt_buffer_read_timed()`.

5.74.3.7 `ssize_t rt_buffer_read_timed (RT_BUFFER * bf, void * ptr, size_t len, const struct timespec * abs_timeout)`

Read from an IPC buffer.

This routine reads the next message from the specified buffer. If no message is available on entry, the caller is allowed to block until enough data is written to the buffer, or a timeout elapses.

Parameters

<i>bf</i>	The descriptor address of the buffer to read from.
<i>ptr</i>	A pointer to a memory area which will be written upon success with the received data.
<i>len</i>	The length in bytes of the memory area pointed to by <i>ptr</i> . Under normal circumstances, rt_buffer_read_timed() only returns entire messages as specified by the <i>len</i> argument, or an error value. However, short reads are allowed when a potential deadlock situation is detected (see note below).

<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for a message to be available from the buffer (see note). Passing NULL causes the caller to block indefinitely until enough data is available. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case not enough data is available.
--------------------	---

Returns

The number of bytes read from the buffer is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before a complete message arrives.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and not enough data is immediately available on entry to form a complete message.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before enough data became available to form a complete message.
- -EINVAL is returned if *bf* is not a valid buffer descriptor, or *len* is greater than the actual buffer length.
- -EIDRM is returned if *bf* is deleted while the caller was waiting for data. In such event, *bf* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Note

A short read (i.e. fewer bytes returned than requested by *len*) may happen whenever a pathological use of the buffer is encountered. This condition only arises when the system detects that one or more writers are waiting for sending data, while a reader would have to wait for receiving a complete message at the same time. For instance, consider the following sequence, involving a 1024-byte buffer (*bf*) and two threads:

```
writer thread > rt_write_buffer(&bf, ptr, 1, TM_INFINITE); (one byte to read, 1023 bytes available for sending)
writer thread > rt_write_buffer(&bf, ptr, 1024, TM_INFINITE); (writer blocks - no space for another 1024-byte message)
reader thread > rt_read_buffer(&bf, ptr, 1024, TM_INFINITE); (short read - a truncated (1-byte) message is returned)
```

In order to prevent both threads to wait for each other indefinitely, a short read is allowed, which may be completed by a subsequent call to [rt_buffer_read\(\)](#) or [rt_buffer_read_until\(\)](#). If that case arises, thread priorities, buffer and/or message lengths should likely be fixed, in order to eliminate such condition.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by [rt_buffer_read\(\)](#), and [rt_buffer_read_until\(\)](#).

```
5.74.3.8 ssize_t rt_buffer_read_until ( RT_BUFFER * bf, void * ptr, size_t len, RTIME abs_timeout
) [inline], [static]
```

Read from an IPC buffer (with absolute scalar timeout).

This routine is a variant of [rt_buffer_read_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>bf</i>	The descriptor address of the buffer to read from.
<i>ptr</i>	A pointer to a memory area which will be written upon success with the received data.
<i>len</i>	The length in bytes of the memory area pointed to by <i>ptr</i> .
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_buffer_read_timed()`.

5.74.3.9 `int rt_buffer_unbind (RT_BUFFER * bf)`

Unbind from an IPC buffer.

Parameters

<i>bf</i>	The descriptor address of the buffer to unbind from.
-----------	--

This routine releases a previous binding to an IPC buffer. After this call has returned, the descriptor is no more valid for referencing this object.

5.74.3.10 `ssize_t rt_buffer_write (RT_BUFFER * bf, const void * ptr, size_t len, RTIME timeout)`
[inline], [static]

Write to an IPC buffer (with relative scalar timeout).

This routine is a variant of [rt_buffer_write_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>bf</i>	The descriptor address of the buffer to write to.
<i>ptr</i>	The address of the message data to be written to the buffer.
<i>len</i>	The length in bytes of the message data.
<i>timeout</i>	A delay expressed in clock ticks.

References `rt_buffer_write_timed()`.

5.74.3.11 `ssize_t rt_buffer_write_timed (RT_BUFFER * bf, const void * ptr, size_t len, const struct timespec * abs_timeout)`

Write to an IPC buffer.

This routine writes a message to the specified buffer. If not enough buffer space is available on entry to hold the message, the caller is allowed to block until enough room is freed, or a timeout elapses, whichever comes first.

Parameters

<i>bf</i>	The descriptor address of the buffer to write to.
<i>ptr</i>	The address of the message data to be written to the buffer.
<i>len</i>	The length in bytes of the message data. Zero is a valid value, in which case the buffer is left untouched, and zero is returned to the caller.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for enough buffer space to be available to hold the message (see note). Passing NULL causes the caller to block indefinitely until enough buffer space is available. Passing { .tv_nsec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case of buffer space shortage.

Returns

The number of bytes written to the buffer is returned upon success. Otherwise:

- -ETIMEDOUT is returned if the absolute *abs_timeout* date is reached before enough buffer space is available to hold the message.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and no buffer space is immediately available on entry to hold the message.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before enough buffer space became available to hold the message.
- -EINVAL is returned if *bf* is not a valid buffer descriptor, or *len* is greater than the actual buffer length.
- -EIDRM is returned if *bf* is deleted while the caller was waiting for buffer space. In such event, *bf* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } .

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `-alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by [rt_buffer_write\(\)](#), and [rt_buffer_write_until\(\)](#).

```
5.74.3.12  ssize_t rt_buffer_write_until ( RT_BUFFER * bf, const void * ptr, size_t len, RTIME
          abs_timeout ) [inline], [static]
```

Write to an IPC buffer (with absolute scalar timeout).

This routine is a variant of [rt_buffer_write_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters


<i>bf</i>	The descriptor address of the buffer to write to.
<i>ptr</i>	The address of the message data to be written to the buffer.
<i>len</i>	The length in bytes of the message data.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References [rt_buffer_write_timed\(\)](#).

5.75 Condition variable services

POSIXish condition variable mechanism.

Collaboration diagram for Condition variable services:



Data Structures

- struct [RT_COND_INFO](#)
Condition variable status descriptor.

Functions

- int [rt_cond_create](#) (RT_COND *cond, const char *name)
Create a condition variable.
- int [rt_cond_delete](#) (RT_COND *cond)
Delete a condition variable.
- int [rt_cond_signal](#) (RT_COND *cond)
Signal a condition variable.
- int [rt_cond_broadcast](#) (RT_COND *cond)
Broadcast a condition variable.
- int [rt_cond_wait_timed](#) (RT_COND *cond, RT_MUTEX *mutex, const struct timespec *abs_timeout)
Wait on a condition variable.
- static int [rt_cond_wait_until](#) (RT_COND *cond, RT_MUTEX *mutex, RTIME timeout)
Wait on a condition variable (with absolute scalar timeout).
- static int [rt_cond_wait](#) (RT_COND *cond, RT_MUTEX *mutex, RTIME timeout)
Wait on a condition variable (with relative scalar timeout).
- int [rt_cond_inquire](#) (RT_COND *cond, [RT_COND_INFO](#) *info)
Query condition variable status.
- int [rt_cond_bind](#) (RT_COND *cond, const char *name, RTIME timeout)
Bind to a condition variable.
- int [rt_cond_unbind](#) (RT_COND *cond)
Unbind from a condition variable.

5.75.1 Detailed Description

POSIXish condition variable mechanism. A condition variable is a synchronization mechanism which allows tasks to suspend execution until some predicate on some arbitrary shared data is satisfied.

The basic operations on conditions are: signal the condition (when the predicate becomes true), and wait for the condition, blocking the task execution until another task signals the condition. A condition variable must always be associated with a mutex, to avoid a well-known race condition where a task prepares to wait on a condition variable and another task signals the condition just before the first task actually waits on it.

5.75.2 Function Documentation

5.75.2.1 `int rt_cond_bind (RT_COND * cond, const char * name, RTIME timeout)`

Bind to a condition variable.

This routine creates a new descriptor to refer to an existing condition variable identified by its symbolic name. If the object not exist on entry, the caller may block until a condition variable of the given name is created.

Parameters

<i>cond</i>	The address of a condition variable descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the condition variable to bind to. This string should match the object name argument passed to rt_cond_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing TM_INFINITE causes the caller to block indefinitely until the object is registered. Passing TM_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to TM_NONBLOCK and the searched object is not registered on entry.
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.75.2.2 `int rt_cond_broadcast (RT_COND * cond)`

Broadcast a condition variable.

All tasks currently waiting on the condition variable are immediately unblocked.

Parameters

<i>cond</i>	The descriptor address of the condition variable to broadcast.
-------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *cond* is not a valid condition variable descriptor.

Valid calling context: any.

5.75.2.3 int rt_cond_create (RT_COND * cond, const char * name)

Create a condition variable.

Create a synchronization object which allows tasks to suspend execution until some predicate on shared data is satisfied.

Parameters

<i>cond</i>	The address of a condition variable descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the condition variable. When non-NULL and non-empty, a copy of this string is used for indexing the created condition variable into the object registry.

Returns

Zero is returned upon success. Otherwise:

- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the condition variable.
- -EEXIST is returned if the *name* is conflicting with an already registered condition variable.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Condition variables can be shared by multiple processes which belong to the same Xenomai session.

References pthread_cond_destroy(), and pthread_cond_init().

5.75.2.4 int rt_cond_delete (RT_COND * cond)

Delete a condition variable.

This routine deletes a condition variable object previously created by a call to [rt_cond_create\(\)](#).

Parameters

<i>cond</i>	The descriptor address of the deleted condition variable.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid condition variable descriptor.
- -EPERM is returned if this service was called from an asynchronous context.
- -EBUSY is returned upon an attempt to destroy the object referenced by *cond* while it is referenced (for example, while being used in a [rt_cond_wait\(\)](#), [rt_cond_wait_timed\(\)](#) or [rt_cond_wait_until\(\)](#) by another task).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

References [pthread_cond_destroy\(\)](#).

5.75.2.5 `int rt_cond_inquire (RT_COND * cond, RT_COND_INFO * info)`

Query condition variable status.

This routine returns the status information about the specified condition variable.

Parameters

<i>cond</i>	The descriptor address of the condition variable to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *cond* is not a valid condition variable descriptor.

Valid calling context: any.

References [RT_COND_INFO::name](#).

5.75.2.6 `int rt_cond_signal (RT_COND * cond)`

Signal a condition variable.

If the condition variable *cond* is pended, this routine immediately unblocks the first waiting task (by queuing priority order).

Parameters

<i>cond</i>	The descriptor address of the condition variable to signal.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *cond* is not a valid condition variable descriptor.

Valid calling context: any.

5.75.2.7 int rt_cond_unbind (RT_COND * cond)

Unbind from a condition variable.

Parameters

<i>cond</i>	The descriptor address of the condition variable to unbind from.
-------------	--

This routine releases a previous binding to a condition variable. After this call has returned, the descriptor is no more valid for referencing this object.

5.75.2.8 int rt_cond_wait (RT_COND * cond, RT_MUTEX * mutex, RTIME timeout) [inline], [static]

Wait on a condition variable (with relative scalar timeout).

This routine is a variant of [rt_cond_wait_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>cond</i>	The descriptor address of the condition variable to wait on.
<i>mutex</i>	The address of the mutex serializing the access to the shared data.
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_cond_wait_timed\(\)](#).

5.75.2.9 int rt_cond_wait_timed (RT_COND * cond, RT_MUTEX * mutex, const struct timespec * abs_timeout)

Wait on a condition variable.

This service atomically releases the mutex and blocks the calling task, until the condition variable *cond* is signaled or a timeout occurs, whichever comes first. The mutex is re-acquired before returning from this service.

Parameters

<i>cond</i>	The descriptor address of the condition variable to wait on.
<i>mutex</i>	The address of the mutex serializing the access to the shared data.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for the condition variable to be signaled (see note). Passing NULL causes the caller to block indefinitely.

Returns

Zero is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before the condition variable is signaled.

- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } .
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task.
- -EINVAL is returned if *cond* is not a valid condition variable descriptor.
- -EIDRM is returned if *cond* is deleted while the caller was waiting on the condition variable. In such event, *cond* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by [rt_cond_wait\(\)](#), and [rt_cond_wait_until\(\)](#).

```
5.75.2.10 int rt_cond_wait_until ( RT_COND * cond, RT_MUTEX * mutex, RTIME abs_timeout )
           [inline], [static]
```

Wait on a condition variable (with absolute scalar timeout).

This routine is a variant of [rt_cond_wait_timed\(\)](#) accepting an *abs_timeout* specification expressed as a scalar value.

Parameters

<i>cond</i>	The descriptor address of the condition variable to wait on.
<i>mutex</i>	The address of the mutex serializing the access to the shared data.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References [rt_cond_wait_timed\(\)](#).

5.76 Event flag group services

Inter-task notification mechanism based on discrete flags.

Collaboration diagram for Event flag group services:



Data Structures

- struct [RT_EVENT_INFO](#)
Event status descriptor.

Macros

- #define [EV_PRIO](#) 0x1 /* Pend by task priority order. */
Creation flags.
- #define [EV_ANY](#) 0x1 /* Disjunctive wait. */
Operation flags.

Functions

- int [rt_event_create](#) (RT_EVENT *event, const char *name, unsigned long ivalue, int mode)
Create an event flag group.
- int [rt_event_delete](#) (RT_EVENT *event)
Delete an event flag group.
- int [rt_event_signal](#) (RT_EVENT *event, unsigned long mask)
Signal an event.
- int [rt_event_wait_timed](#) (RT_EVENT *event, unsigned long mask, unsigned long *mask_r, int mode, const struct timespec *abs_timeout)
Wait for an arbitrary set of events.
- static int [rt_event_wait_until](#) (RT_EVENT *event, unsigned long mask, unsigned long *mask_r, int mode, RTIME timeout)
Wait for an arbitrary set of events (with absolute scalar timeout).
- static int [rt_event_wait](#) (RT_EVENT *event, unsigned long mask, unsigned long *mask_r, int mode, RTIME timeout)
Wait for an arbitrary set of events (with relative scalar timeout).
- int [rt_event_clear](#) (RT_EVENT *event, unsigned long mask, unsigned long *mask_r)
Clear event flags.
- int [rt_event_inquire](#) (RT_EVENT *event, [RT_EVENT_INFO](#) *info)
Query event flag group status.
- int [rt_event_bind](#) (RT_EVENT *event, const char *name, RTIME timeout)
Bind to an event flag group.
- int [rt_event_unbind](#) (RT_EVENT *event)
Unbind from an event flag group.

5.76.1 Detailed Description

Inter-task notification mechanism based on discrete flags. An event flag group is a synchronization object represented by a long-word structure; every available bit in this word represents a user-defined event flag.

When a bit is set, the associated event is said to have occurred. Xenomai tasks can use this mechanism to signal the occurrence of particular events to other tasks.

Tasks can either wait for events to occur in a conjunctive manner (all awaited events must have occurred to satisfy the wait request), or in a disjunctive way (at least one of the awaited events must have occurred to satisfy the wait request).

5.76.2 Macro Definition Documentation

5.76.2.1 `#define EV_ANY 0x1 /* Disjunctive wait. */`

Operation flags.

Referenced by `rt_event_wait_timed()`.

5.76.2.2 `#define EV_PRIO 0x1 /* Pend by task priority order. */`

Creation flags.

Referenced by `rt_event_create()`.

5.76.3 Function Documentation

5.76.3.1 `int rt_event_bind (RT_EVENT * event, const char * name, RTIME timeout)`

Bind to an event flag group.

This routine creates a new descriptor to refer to an existing event flag group identified by its symbolic name. If the object does not exist on entry, the caller may block until an event flag group of the given name is created.

Parameters

<i>event</i>	The address of an event flag group descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the event flag group to bind to. This string should match the object name argument passed to rt_event_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing <code>TM_INFINITE</code> causes the caller to block indefinitely until the object is registered. Passing <code>TM_NONBLOCK</code> causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `TM_NONBLOCK` and the searched object is not registered on entry.
- `-ETIMEDOUT` is returned if the object cannot be retrieved within the specified amount of time.

- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.76.3.2 `int rt_event_clear (RT_EVENT * event, unsigned long mask, unsigned long * mask_r)`

Clear event flags.

This routine clears a set of flags from *event*.

Parameters

<i>event</i>	The descriptor address of the affected event.
<i>mask</i>	The set of event flags to be cleared.
<i>mask_r</i>	If non-NULL, <i>mask_r</i> is the address of a memory location which will receive the previous value of the event flag group before the flags are cleared.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *event* is not a valid event flag group descriptor.

Valid calling context: any.

5.76.3.3 `int rt_event_create (RT_EVENT * event, const char * name, unsigned long ivalue, int mode)`

Create an event flag group.

Event groups provide for task synchronization by allowing a set of flags (or "events") to be waited for and posted atomically. An event group contains a mask of received events; an arbitrary set of event flags can be pended or posted in a single operation.

Parameters

<i>event</i>	The address of an event descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the event. When non-NULL and non-empty, a copy of this string is used for indexing the created event into the object registry.

<i>ivalue</i>	The initial value of the group's event mask.
<i>mode</i>	The event group creation mode. The following flags can be OR'ed into this bitmask:

- EV_FIFO makes tasks pend in FIFO order on the event flag group.
- EV_PRIO makes tasks pend in priority order on the event flag group.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *mode* is invalid.
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the event flag group.
- -EEXIST is returned if the *name* is conflicting with an already registered event flag group.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Event flag groups can be shared by multiple processes which belong to the same Xenomai session.

References EV_PRIO.

5.76.3.4 int rt_event_delete (RT_EVENT * event)

Delete an event flag group.

This routine deletes a event flag group previously created by a call to [rt_event_create\(\)](#).

Parameters

<i>event</i>	The descriptor address of the deleted object.
--------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *event* is not a valid event flag group descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.76.3.5 int rt_event_inquire (RT_EVENT * event, RT_EVENT_INFO * info)

Query event flag group status.

This routine returns the status information about *event*.

Parameters

<i>event</i>	The descriptor address of the event flag group to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *event* is not a valid event flag group descriptor.

Valid calling context: any.

References RT_EVENT_INFO::name, RT_EVENT_INFO::nwaiters, and RT_EVENT_INFO::value.

5.76.3.6 int rt_event_signal (RT_EVENT * event, unsigned long mask)

Signal an event.

Post a set of flags to *event*. All tasks having their wait request satisfied as a result of this operation are immediately readied.

Parameters

<i>event</i>	The descriptor address of the event flag group to signal.
<i>mask</i>	The set of events to be posted.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *event* is not an event flag group descriptor.

Valid calling context: any.

5.76.3.7 int rt_event_unbind (RT_EVENT * event)

Unbind from an event flag group.

Parameters

<i>event</i>	The descriptor address of the object to unbind from.
--------------	--

This routine releases a previous binding to an event flag group. After this call has returned, the descriptor is no more valid for referencing this object.

5.76.3.8 int rt_event_wait (RT_EVENT * event, unsigned long mask, unsigned long * mask_r, int mode, RTIME timeout) [inline], [static]

Wait for an arbitrary set of events (with relative scalar timeout).

This routine is a variant of [rt_event_wait_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>event</i>	The descriptor address of the event flag group to wait on.
<i>mask</i>	The set of bits to wait for.
<i>mask_r</i>	The value of the event mask at the time the task was readied.
<i>mode</i>	The pend mode.
<i>timeout</i>	A delay expressed in clock ticks,

References `rt_event_wait_timed()`.

5.76.3.9 `int rt_event_wait_timed (RT_EVENT * event, unsigned long mask, unsigned long * mask_r, int mode, const struct timespec * abs_timeout)`

Wait for an arbitrary set of events.

Waits for one or more events to be signaled in *event*, or until a timeout elapses.

Parameters

<i>event</i>	The descriptor address of the event flag group to wait on.
<i>mask</i>	The set of bits to wait for. Passing zero causes this service to return immediately with a success value; the current value of the event mask is also copied to <i>mask_r</i> .
<i>mask_r</i>	The value of the event mask at the time the task was readied.
<i>mode</i>	The pend mode. The following flags can be OR'ed into this bitmask, each of them affecting the operation:

- `EV_ANY` makes the task pend in disjunctive mode (i.e. OR); this means that the request is fulfilled when at least one bit set into *mask* is set in the current event mask.
- `EV_ALL` makes the task pend in conjunctive mode (i.e. AND); this means that the request is fulfilled when at all bits set into *mask* are set in the current event mask.

Parameters

<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for the request to be satisfied (see note). Passing NULL causes the caller to block indefinitely until the request is satisfied. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return without blocking in case the request cannot be satisfied immediately.
--------------------	--

Returns

Zero is returned upon success. Otherwise:

- `-ETIMEDOUT` is returned if *abs_timeout* is reached before the request is satisfied.
- `-EWOULDBLOCK` is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and the requested flags are not set on entry to the call.
- `-EINTR` is returned if `rt_task_unblock()` was called for the current task before the request is satisfied.
- `-EINVAL` is returned if *mode* is invalid, or *event* is not a valid event flag group descriptor.
- `-EIDRM` is returned if *event* is deleted while the caller was sleeping on it. In such a case, *event* is no more valid upon return of this service.
- `-EPERM` is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads.

- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

References EV_ANY.

Referenced by `rt_event_wait()`, and `rt_event_wait_until()`.

5.76.3.10 `int rt_event_wait_until (RT_EVENT * event, unsigned long mask, unsigned long * mask_r, int mode, RTIME abs_timeout) [inline], [static]`

Wait for an arbitrary set of events (with absolute scalar timeout).

This routine is a variant of [rt_event_wait_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>event</i>	The descriptor address of the event flag group to wait on.
<i>mask</i>	The set of bits to wait for.
<i>mask_r</i>	The value of the event mask at the time the task was readied.
<i>mode</i>	The pend mode.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_event_wait_timed()`.

5.77 Heap management services

Region of memory dedicated to real-time allocation.

Collaboration diagram for Heap management services:



Data Structures

- struct [RT_HEAP_INFO](#)
Heap status descriptor.

Macros

- `#define H_PRIO 0x1` /* Pend by task priority order. */
Creation flags.

Functions

- int [rt_heap_create](#) (RT_HEAP *heap, const char *name, size_t heapsize, int mode)
Create a heap.
- int [rt_heap_delete](#) (RT_HEAP *heap)
Delete a heap.
- int [rt_heap_alloc_timed](#) (RT_HEAP *heap, size_t size, const struct timespec *abs_timeout, void **blockp)
Allocate a block from a heap.
- static int [rt_heap_alloc_until](#) (RT_HEAP *heap, size_t size, RTIME timeout, void **blockp)
Allocate a block from a heap (with absolute scalar timeout).
- static int [rt_heap_alloc](#) (RT_HEAP *heap, size_t size, RTIME timeout, void **blockp)
Allocate a block from a heap (with relative scalar timeout).
- int [rt_heap_free](#) (RT_HEAP *heap, void *block)
Release a block to a heap.
- int [rt_heap_inquire](#) (RT_HEAP *heap, [RT_HEAP_INFO](#) *info)
Query heap status.
- int [rt_heap_bind](#) (RT_HEAP *heap, const char *name, RTIME timeout)
Bind to a heap.
- int [rt_heap_unbind](#) (RT_HEAP *heap)
Unbind from a heap.

5.77.1 Detailed Description

Region of memory dedicated to real-time allocation. Heaps are regions of memory used for dynamic memory allocation in a time-bounded fashion. Blocks of memory are allocated and freed in an arbitrary order and the pattern of allocation and size of blocks is not known until run time.

5.77.2 Macro Definition Documentation

5.77.2.1 `#define H_PRIO 0x1 /* Pend by task priority order. */`

Creation flags.

Referenced by `rt_heap_create()`.

5.77.3 Function Documentation

5.77.3.1 `int rt_heap_alloc (RT_HEAP * heap, size_t size, RTIME timeout, void ** blockp)`
`[inline], [static]`

Allocate a block from a heap (with relative scalar timeout).

This routine is a variant of `rt_heap_alloc_timed()` accepting a relative timeout specification expressed as a scalar value.

References `rt_heap_alloc_timed()`.

5.77.3.2 `int rt_heap_alloc_timed (RT_HEAP * heap, size_t size, const struct timespec * abs_timeout, void ** blockp)`

Allocate a block from a heap.

This service allocates a block from a given heap, or returns the address of the single memory segment if `H_SINGLE` was mentioned in the creation mode to `rt_heap_create()`. When not enough memory is available on entry to this service, tasks may be blocked until their allocation request can be fulfilled.

Parameters

<i>heap</i>	The descriptor address of the heap to allocate from.
<i>size</i>	The requested size (in bytes) of the block. If the heap is managed as a single-block area (<code>H_SINGLE</code>), this value can be either zero, or the same value given to <code>rt_heap_create()</code> . In that case, the same block covering the entire heap space is returned to all callers of this service.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for a block of the requested size to be available from the heap (see note). Passing <code>NULL</code> causes the caller to block indefinitely until a block is available. Passing <code>{ .tv_sec = 0, .tv_nsec = 0 }</code> causes the service to return immediately without blocking in case not block is available.
<i>blockp</i>	A pointer to a memory location which will be written upon success with the address of the allocated block, or the start address of the single memory segment. In the former case, the block can be freed using <code>rt_heap_free()</code> .

Returns

Zero is returned upon success. Otherwise:

- `-ETIMEDOUT` is returned if *abs_timeout* is reached before a block is available.

- -EWOULDBLOCK is returned if *abs_timeout* is equal to { .tv_sec = 0, .tv_nsec = 0 } and no block is immediately available on entry to fulfill the allocation request.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before a block became available.
- -EINVAL is returned if *heap* is not a valid heap descriptor, or *heap* is managed as a single-block area (i.e. H_SINGLE mode) and *size* is non-zero but does not match the original heap size passed to [rt_heap_create\(\)](#).
- -EIDRM is returned if *heap* is deleted while the caller was waiting for a block. In such event, *heap* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by [rt_heap_alloc\(\)](#), and [rt_heap_alloc_until\(\)](#).

```
5.77.3.3 int rt_heap_alloc_until ( RT_HEAP * heap, size_t size, RTIME abs_timeout, void ** blockp
) [inline], [static]
```

Allocate a block from a heap (with absolute scalar timeout).

This routine is a variant of [rt_heap_alloc_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

References [rt_heap_alloc_timed\(\)](#).

```
5.77.3.4 int rt_heap_bind ( RT_HEAP * heap, const char * name, RTIME timeout )
```

Bind to a heap.

This routine creates a new descriptor to refer to an existing heap identified by its symbolic name. If the object does not exist on entry, the caller may block until a heap of the given name is created.

Parameters

<i>heap</i>	The address of a heap descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the heap to bind to. This string should match the object name argument passed to rt_heap_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing TM_INFINITE causes the caller to block indefinitely until the object is registered. Passing TM_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- -EINTR is returned if `rt_task_unblock()` was called for the current task before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to TM_NONBLOCK and the searched object is not registered on entry.
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.77.3.5 `int rt_heap_create (RT_HEAP * heap, const char * name, size_t heapsz, int mode)`

Create a heap.

This routine creates a memory heap suitable for time-bounded allocation requests of RAM chunks. When not enough memory is available, tasks may be blocked until their allocation request can be fulfilled.

By default, heaps support allocation of multiple blocks of memory in an arbitrary order. However, it is possible to ask for single-block management by passing the H_SINGLE flag into the *mode* parameter, in which case the entire memory space managed by the heap is made available as a unique block. In this mode, all allocation requests made through `rt_heap_alloc()` will return the same block address, pointing at the beginning of the heap memory.

Parameters

<i>heap</i>	The address of a heap descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the heap. When non-NULL and non-empty, a copy of this string is used for indexing the created heap into the object registry.
<i>heapsz</i>	The size (in bytes) of the memory pool, blocks will be claimed and released to. This area is not extensible, so this value must be compatible with the highest memory pressure that could be expected. The valid range is between 2k and 2Gb.
<i>mode</i>	The heap creation mode. The following flags can be OR'ed into this bitmask, each of them affecting the new heap:

- H_FIFO makes tasks pend in FIFO order on the heap when waiting for available blocks.
- H_PRIO makes tasks pend in priority order on the heap when waiting for available blocks.
- H_SINGLE causes the entire heap space to be managed as a single memory block.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *mode* is invalid, or *heapsz* is not in the range [2k..2Gb].
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the heap.
- -EEXIST is returned if the *name* is conflicting with an already registered heap.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Heaps can be shared by multiple processes which belong to the same Xenomai session.

References H_PRIO.

5.77.3.6 int rt_heap_delete (RT_HEAP * heap)

Delete a heap.

This routine deletes a heap object previously created by a call to [rt_heap_create\(\)](#), releasing all tasks currently blocked on it.

Parameters

<i>heap</i>	The descriptor address of the deleted heap.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *heap* is not a valid heap descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.77.3.7 int rt_heap_free (RT_HEAP * heap, void * block)

Release a block to a heap.

This service should be used to release a block to the heap it belongs to. An attempt to fulfill the request of every task blocked on [rt_heap_alloc\(\)](#) is made once *block* is returned to the memory pool.

Parameters

<i>heap</i>	The descriptor address of the heap to release the block to.
<i>block</i>	The address of the block to free.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *heap* is not a valid heap descriptor, or *block* is not a valid block previously allocated by the [rt_heap_alloc\(\)](#) service from *heap*.

Valid calling contexts: any.

5.77.3.8 int rt_heap_inquire (RT_HEAP * heap, RT_HEAP_INFO * info)

Query heap status.

This routine returns the status information about *heap*.

Parameters

<i>heap</i>	The descriptor address of the heap to get the status of.
<i>info</i>	A pointer to the return buffer" to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *heap* is not a valid heap descriptor.

Valid calling context: any.

References RT_HEAP_INFO::heapsize, RT_HEAP_INFO::name, RT_HEAP_INFO::nwaiters, RT_HEAP_INFO::usablemem, and RT_HEAP_INFO::usedmem.

5.77.3.9 int rt_heap_unbind (RT_HEAP * heap)

Unbind from a heap.

Parameters


<i>heap</i>	The descriptor address of the heap to unbind from.
-------------	--

This routine releases a previous binding to a heap. After this call has returned, the descriptor is no more valid for referencing this object.

5.78 Alchemy API

A programming interface reminiscent from traditional RTOS APIs.

Collaboration diagram for Alchemy API:



Modules

- [Alarm services](#)
General-purpose watchdog timers.
- [Buffer services](#)
Lightweight FIFO IPC mechanism.
- [Condition variable services](#)
POSIXish condition variable mechanism.

- [Event flag group services](#)
Inter-task notification mechanism based on discrete flags.
- [Heap management services](#)
Region of memory dedicated to real-time allocation.
- [Mutex services](#)
POSIXish mutual exclusion services.
- [Message pipe services](#)
Two-way communication channel between Xenomai & Linux domains.
- [Message queue services](#)
real-time IPC mechanism for sending messages of arbitrary size
- [Semaphore services](#)
Counting semaphore IPC mechanism.
- [Task management services](#)
Services dealing with preemptive multi-tasking.
- [Timer management services](#)
Services for reading and spinning on the hardware timer.

5.78.1 Detailed Description

A programming interface reminiscent from traditional RTOS APIs. This interface is an evolution of the former *native* API available with the Xenomai 2.x series.

5.79 Mutex services

POSIXish mutual exclusion servicesl.

Collaboration diagram for Mutex services:



Data Structures

- struct [RT_MUTEX_INFO](#)
Mutex status descriptor.

Functions

- int [rt_mutex_create](#) (RT_MUTEX *mutex, const char *name)
Create a mutex.
- int [rt_mutex_delete](#) (RT_MUTEX *mutex)
Delete a mutex.
- int [rt_mutex_acquire_timed](#) (RT_MUTEX *mutex, const struct timespec *abs_timeout)
Acquire/lock a mutex (with absolute timeout date).
- static int [rt_mutex_acquire_until](#) (RT_MUTEX *mutex, RTIME timeout)
Acquire/lock a mutex (with absolute scalar timeout).
- static int [rt_mutex_acquire](#) (RT_MUTEX *mutex, RTIME timeout)
Acquire/lock a mutex (with relative scalar timeout).
- int [rt_mutex_release](#) (RT_MUTEX *mutex)
Release/unlock a mutex.
- int [rt_mutex_inquire](#) (RT_MUTEX *mutex, [RT_MUTEX_INFO](#) *info)
Query mutex status.
- int [rt_mutex_bind](#) (RT_MUTEX *mutex, const char *name, RTIME timeout)
Bind to a mutex.
- int [rt_mutex_unbind](#) (RT_MUTEX *mutex)
Unbind from a mutex.

5.79.1 Detailed Description

POSIXish mutual exclusion servicesl. A mutex is a MUTual EXclusion object, and is useful for protecting shared data structures from concurrent modifications, and implementing critical sections and monitors.

A mutex has two possible states: unlocked (not owned by any task), and locked (owned by one task). A mutex can never be owned by two different tasks simultaneously. A task attempting to lock a mutex that is already locked by another task is blocked until the latter unlocks the mutex first.

Xenomai mutex services enforce a priority inheritance protocol in order to solve priority inversions.

5.79.2 Function Documentation

5.79.2.1 `int rt_mutex_acquire (RT_MUTEX * mutex, RTIME timeout) [inline], [static]`

Acquire/lock a mutex (with relative scalar timeout).

This routine is a variant of [rt_mutex_acquire_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>mutex</i>	The descriptor address of the mutex to acquire.
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_mutex_acquire_timed\(\)](#).

5.79.2.2 `int rt_mutex_acquire_timed (RT_MUTEX * mutex, const struct timespec * abs_timeout)`

Acquire/lock a mutex (with absolute timeout date).

Attempt to lock a mutex. The calling task is blocked until the mutex is available, in which case it is locked again before this service returns. Xenomai mutexes are implicitly recursive and implement the priority inheritance protocol.

Parameters

<i>mutex</i>	The descriptor address of the mutex to acquire.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for the mutex to be available (see note). Passing NULL the caller to block indefinitely. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case <i>mutex</i> is already locked by another task.

Returns

Zero is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before the mutex is available.
- -EWOULDBLOCK is returned if *timeout* is { .tv_sec = 0, .tv_nsec = 0 } and the mutex is not immediately available.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task.
- -EINVAL is returned if *mutex* is not a valid mutex descriptor.
- -EIDRM is returned if *mutex* is deleted while the caller was waiting on it. In such event, *mutex* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads

Core specifics:

Over the Cobalt core, a real-time task with effective priority zero keeps running in primary mode until it releases the mutex.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by [rt_mutex_acquire\(\)](#), and [rt_mutex_acquire_until\(\)](#).

```
5.79.2.3 int rt_mutex_acquire_until ( RT_MUTEX * mutex, RTIME abs_timeout ) [inline],
        [static]
```

Acquire/lock a mutex (with absolute scalar timeout).

This routine is a variant of [rt_mutex_acquire_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>mutex</i>	The descriptor address of the mutex to acquire.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References [rt_mutex_acquire_timed\(\)](#).

```
5.79.2.4 int rt_mutex_bind ( RT_MUTEX * mutex, const char * name, RTIME timeout )
```

Bind to a mutex.

This routine creates a new descriptor to refer to an existing mutex identified by its symbolic name. If the object not exist on entry, the caller may block until a mutex of the given name is created.

Parameters

<i>mutex</i>	The address of a mutex descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the mutex to bind to. This string should match the object name argument passed to rt_mutex_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing TM_INFINITE causes the caller to block indefinitely until the object is registered. Passing TM_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to TM_NONBLOCK and the searched object is not registered on entry.
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.79.2.5 `int rt_mutex_create (RT_MUTEX * mutex, const char * name)`

Create a mutex.

Create a mutual exclusion object that allows multiple tasks to synchronize access to a shared resource. A mutex is left in an unlocked state after creation.

Parameters

<i>mutex</i>	The address of a mutex descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the mutex. When non-NULL and non-empty, a copy of this string is used for indexing the created mutex into the object registry.

Returns

Zero is returned upon success. Otherwise:

- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the mutex.
- -EEXIST is returned if the *name* is conflicting with an already registered mutex.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Mutexes can be shared by multiple processes which belong to the same Xenomai session.

5.79.2.6 `int rt_mutex_delete (RT_MUTEX * mutex)`

Delete a mutex.

This routine deletes a mutex object previously created by a call to [rt_mutex_create\(\)](#).

Parameters

<i>mutex</i>	The descriptor address of the deleted mutex.
--------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid mutex descriptor.
- -EPERM is returned if this service was called from an asynchronous context.
- -EBUSY is returned upon an attempt to destroy the object referenced by *mutex* while it is referenced (for example, while being used in a [rt_mutex_acquire\(\)](#), [rt_mutex_acquire_timed\(\)](#) or [rt_mutex_acquire_until\(\)](#) by another task).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.79.2.7 `int rt_mutex_inquire (RT_MUTEX * mutex, RT_MUTEX_INFO * info)`

Query mutex status.

This routine returns the status information about the specified mutex.

Parameters

<i>mutex</i>	The descriptor address of the mutex to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *mutex* is not a valid mutex descriptor.
- -EPERM is returned if this service is called from an interrupt context.

Valid calling context:

- Xenomai threads

References `RT_MUTEX_INFO::name`, and `RT_MUTEX_INFO::owner`.

5.79.2.8 `int rt_mutex_release (RT_MUTEX * mutex)`

Release/unlock a mutex.

This routine releases a mutex object previously locked by a call to [rt_mutex_acquire\(\)](#) or [rt_mutex_acquire_until\(\)](#). If the mutex is pended, the first waiting task (by priority order) is immediately unblocked and transferred the ownership of the mutex; otherwise, the mutex is left in an unlocked state.

Parameters

<i>mutex</i>	The descriptor address of the deleted mutex.
--------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *alarm* is not a valid mutex descriptor.
- -EPERM is returned if *mutex* is not owned by the current task, or more generally if this service was called from a context which cannot own any mutex (e.g. interrupt context).

Valid calling context:

- Xenomai threads

5.79.2.9 `int rt_mutex_unbind (RT_MUTEX * mutex)`

Unbind from a mutex.

Parameters

<i>mutex</i>	The descriptor address of the mutex to unbind from.
--------------	---

This routine releases a previous binding to a mutex. After this call has returned, the descriptor is no more valid for referencing this object.

5.80 Message pipe services

Two-way communication channel between Xenomai & Linux domains.

Collaboration diagram for Message pipe services:



Macros

- `#define P_MINOR_AUTO XNPIPE_MINOR_AUTO`
Creation flags.
- `#define P_URGENT 0x1`
Operation flags.

Functions

- `int rt_pipe_create (RT_PIPE *pipe, const char *name, int minor, size_t poolsize)`
Create a message pipe.
- `int rt_pipe_delete (RT_PIPE *pipe)`
Delete a message pipe.
- `ssize_t rt_pipe_read_timed (RT_PIPE *pipe, void *buf, size_t size, const struct timespec *abs_timeout)`
Read a message from a pipe.
- `static ssize_t rt_pipe_read_until (RT_PIPE *pipe, void *buf, size_t size, RTIME timeout)`
Read from a pipe (with absolute scalar timeout).
- `static ssize_t rt_pipe_read (RT_PIPE *pipe, void *buf, size_t size, RTIME timeout)`
Read from a pipe (with relative scalar timeout).
- `ssize_t rt_pipe_write (RT_PIPE *pipe, const void *buf, size_t size, int mode)`
Write a message to a pipe.
- `int rt_pipe_bind (RT_PIPE *pipe, const char *name, RTIME timeout)`
Bind to a message pipe.
- `int rt_pipe_unbind (RT_PIPE *pipe)`
Unbind from a message pipe.

5.80.1 Detailed Description

Two-way communication channel between Xenomai & Linux domains. A message pipe is a two-way communication channel between Xenomai threads and normal Linux threads using regular file I/O operations on a pseudo-device. Pipes can be operated in a message-oriented fashion so that message boundaries are preserved, and also in byte-oriented streaming mode from real-time to normal Linux threads for optimal throughput.

Xenomai threads open their side of the pipe using the [rt_pipe_create\(\)](#) service; regular Linux threads do the same by opening one of the `/dev/rtpN` special devices, where N is the minor number agreed upon between both ends of each pipe.

In addition, named pipes are available through the registry support, which automatically creates a symbolic link from entries under `/proc/xenomai/registry/rtpc/xddp/` to the corresponding special device file.

Note

Alchemy's message pipes are fully based on the [XDDP protocol](#) available from the RTDM/ipc driver.

5.80.2 Macro Definition Documentation

5.80.2.1 `#define P_MINOR_AUTO XNPIPE_MINOR_AUTO`

Creation flags.

5.80.2.2 `#define P_URGENT 0x1`

Operation flags.

Referenced by `rt_pipe_write()`.

5.80.3 Function Documentation

5.80.3.1 `int rt_pipe_bind (RT_PIPE * pipe, const char * name, RTIME timeout)`

Bind to a message pipe.

This routine creates a new descriptor to refer to an existing message pipe identified by its symbolic name. If the object does not exist on entry, the caller may block until a pipe of the given name is created.

Parameters

<i>pipe</i>	The address of a pipe descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the pipe to bind to. This string should match the object name argument passed to rt_pipe_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing <code>TM_INFINITE</code> causes the caller to block indefinitely until the object is registered. Passing <code>TM_NONBLOCK</code> causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `TM_NONBLOCK` and the searched object is not registered on entry.
- `-ETIMEDOUT` is returned if the object cannot be retrieved within the specified amount of time.
- `-EPERM` is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals `TM_NONBLOCK`.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.80.3.2 `int rt_pipe_create (RT_PIPE * pipe, const char * name, int minor, size_t poolsize)`

Create a message pipe.

This service opens a bi-directional communication channel for exchanging messages between Xenomai threads and regular Linux threads. Pipes natively preserve message boundaries, but can also be used in byte-oriented streaming mode from Xenomai to Linux.

`rt_pipe_create()` always returns immediately, even if no thread has opened the associated special device file yet. On the contrary, the non real-time side could block upon attempt to open the special device file until `rt_pipe_create()` is issued on the same pipe from a Xenomai thread, unless `O_NONBLOCK` was given to the `open(2)` system call.

Parameters

<i>pipe</i>	The address of a pipe descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the pipe. When non-NULL and non-empty, a copy of this string is used for indexing the created pipe into the object registry.

Named pipes are supported through the use of the registry. Passing a valid *name* parameter when creating a message pipe causes a symbolic link to be created from `/proc/xenomai/registry/rtpc/xddp/name` to the associated special device (i.e. `/dev/rtp*`), so that the specific *minor* information does not need to be known from those processes for opening the proper device file. In such a case, both sides of the pipe only need to agree upon a symbolic name to refer to the same data path, which is especially useful whenever the *minor* number is picked up dynamically using an adaptive algorithm, such as passing `P_MINOR_AUTO` as *minor* value.

Parameters

<i>minor</i>	The minor number of the device associated with the pipe. Passing <code>P_MINOR_AUTO</code> causes the minor number to be auto-allocated. In such a case, a symbolic link will be automatically created from <code>/proc/xenomai/registry/rtpc/xddp/name</code> to the allocated pipe device entry. Valid minor numbers range from 0 to <code>CONFIG_XENO_OPT_PIPE_NRDEV-1</code> .
<i>poolsize</i>	Specifies the size of a dedicated buffer pool for the pipe. Passing 0 means that all message allocations for this pipe are performed on the Cobalt core heap.

Returns

Zero is returned upon success. Otherwise:

- `-ENOMEM` is returned if the system fails to get memory from the main heap in order to create the pipe.
- `-ENODEV` is returned if *minor* is different from `P_MINOR_AUTO` and is not a valid minor number.
- `-EEXIST` is returned if the *name* is conflicting with an already registered pipe.
- `-EBUSY` is returned if *minor* is already open.

- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

References IPCPROTO_XDDP, rtipc_port_label::label, sockaddr_ipc::sipc_family, sockaddr_ipc::sipc_port, XDDP_BUFSZ, XDDP_LABEL, and XDDP_POOLSZ.

5.80.3.3 int rt_pipe_delete (RT_PIPE * pipe)

Delete a message pipe.

This routine deletes a pipe object previously created by a call to [rt_pipe_create\(\)](#). All resources attached to that pipe are automatically released, all pending data is flushed.

Parameters

<i>pipe</i>	The descriptor address of the deleted pipe.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *pipe* is not a valid pipe descriptor.
- -EIDRM is returned if *pipe* is a closed pipe descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.80.3.4 ssize_t rt_pipe_read (RT_PIPE * pipe, void * buf, size_t size, RTIME timeout) [inline], [static]

Read from a pipe (with relative scalar timeout).

This routine is a variant of [rt_queue_read_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>pipe</i>	The descriptor address of the message pipe to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the message received.
<i>size</i>	The count of bytes from the received message to read up into <i>buf</i> . If <i>size</i> is lower than the actual message size, -ENOBUS is returned since the incompletely received message would be lost. If <i>size</i> is zero, this call returns immediately with no other action.
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_pipe_read_timed\(\)](#).

5.80.3.5 `ssize_t rt_pipe_read_timed (RT_PIPE * pipe, void * buf, size_t size, const struct timespec * abs_timeout)`

Read a message from a pipe.

This service reads the next available message from a given pipe.

Parameters

<i>pipe</i>	The descriptor address of the message pipe to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the message received.
<i>size</i>	The count of bytes from the received message to read up into <i>buf</i> . If <i>size</i> is lower than the actual message size, -ENOBUFFS is returned since the incompletely received message would be lost. If <i>size</i> is zero, this call returns immediately with no other action.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for a message to be available from the pipe (see note). Passing NULL causes the caller to block indefinitely until a message is available. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case no message is available.

Returns

The number of bytes available from the received message is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before a message arrives.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and no message is immediately available on entry to the call.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before a message was available.
- -EINVAL is returned if *pipe* is not a valid pipe descriptor.
- -EIDRM is returned if *pipe* is deleted while the caller was waiting for a message. In such event, *pipe* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

References `SO_RCVTIMEO`.

Referenced by `rt_pipe_read()`, and `rt_pipe_read_until()`.

5.80.3.6 `ssize_t rt_pipe_read_until (RT_PIPE * pipe, void * buf, size_t size, RTIME abs_timeout)`
[inline], [static]

Read from a pipe (with absolute scalar timeout).

This routine is a variant of [rt_queue_read_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>pipe</i>	The descriptor address of the message pipe to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the message received.
<i>size</i>	The count of bytes from the received message to read up into <i>buf</i> . If <i>size</i> is lower than the actual message size, -ENOBUFS is returned since the incompletely received message would be lost. If <i>size</i> is zero, this call returns immediately with no other action.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_pipe_read_timed()`.

5.80.3.7 `int rt_pipe_unbind (RT_PIPE * pipe)`

Unbind from a message pipe.

Parameters

<i>pipe</i>	The descriptor address of the pipe to unbind from.
-------------	--

This routine releases a previous binding to a message pipe. After this call has returned, the descriptor is no more valid for referencing this object.

5.80.3.8 `ssize_t rt_pipe_write (RT_PIPE * pipe, const void * buf, size_t size, int mode)`

Write a message to a pipe.

This service writes a complete message to be received from the associated special device. `rt_pipe_write()` always preserves message boundaries, which means that all data sent through a single call of this service will be gathered in a single `read(2)` operation from the special device.

This service differs from `rt_pipe_send()` in that it accepts a pointer to the raw data to be sent, instead of a canned message buffer.

Parameters

<i>pipe</i>	The descriptor address of the pipe to write to.
<i>buf</i>	The address of the first data byte to send. The data will be copied to an internal buffer before transmission.
<i>size</i>	The size in bytes of the message (payload data only). Zero is a valid value, in which case the service returns immediately without sending any message.
<i>mode</i>	A set of flags affecting the operation:

- `P_URGENT` causes the message to be prepended to the output queue, ensuring a LIFO ordering.
- `P_NORMAL` causes the message to be appended to the output queue, ensuring a FIFO ordering.

Returns

Upon success, this service returns *size*. Upon error, one of the following error codes is returned:

- `-EINVAL` is returned if *mode* is invalid or *pipe* is not a pipe descriptor.
- `-ENOMEM` is returned if not enough buffer space is available to complete the operation.
- `-EIDRM` is returned if *pipe* is a closed pipe descriptor.

Note

Writing data to a pipe before any peer has opened the associated special device is allowed. The output will be buffered until then, only restricted by the available memory in the associated buffer pool (see [rt_pipe_create\(\)](#)).

References P_URGENT.

5.81 Message queue services

real-time IPC mechanism for sending messages of arbitrary size

Collaboration diagram for Message queue services:



Data Structures

- struct [RT_QUEUE_INFO](#)
Queue status descriptor.

Macros

- #define [Q_PRIO](#) 0x1 /* Pend by task priority order. */
Creation flags.
- #define [Q_SHARED](#) 0x0
Deprecated, compat only.

Functions

- int [rt_queue_create](#) (RT_QUEUE *queue, const char *name, size_t poolsize, size_t qlimit, int mode)
Create a message queue.
- int [rt_queue_delete](#) (RT_QUEUE *queue)
Delete a message queue.
- void * [rt_queue_alloc](#) (RT_QUEUE *queue, size_t size)
Allocate a message buffer.
- int [rt_queue_free](#) (RT_QUEUE *queue, void *buf)
Free a message buffer.
- int [rt_queue_send](#) (RT_QUEUE *queue, const void *buf, size_t size, int mode)
Send a message to a queue.
- ssize_t [rt_queue_receive_timed](#) (RT_QUEUE *queue, void **bufp, const struct timespec *abs_timeout)
Receive a message from a queue (with absolute timeout date).
- static ssize_t [rt_queue_receive_until](#) (RT_QUEUE *queue, void **bufp, RTIME timeout)
Receive from a queue (with absolute scalar timeout).
- static ssize_t [rt_queue_receive](#) (RT_QUEUE *queue, void **bufp, RTIME timeout)
Receive from a queue (with relative scalar timeout).
- ssize_t [rt_queue_read_timed](#) (RT_QUEUE *queue, void *buf, size_t size, const struct timespec *abs_timeout)
Read from a queue.

- static ssize_t [rt_queue_read_until](#) (RT_QUEUE *queue, void *buf, size_t size, RTIME timeout)
Read from a queue (with absolute scalar timeout).
- static ssize_t [rt_queue_read](#) (RT_QUEUE *queue, void *buf, size_t size, RTIME timeout)
Read from a queue (with relative scalar timeout).
- int [rt_queue_flush](#) (RT_QUEUE *queue)
Flush pending messages from a queue.
- int [rt_queue_inquire](#) (RT_QUEUE *queue, [RT_QUEUE_INFO](#) *info)
Query queue status.
- int [rt_queue_bind](#) (RT_QUEUE *queue, const char *name, RTIME timeout)
Bind to a message queue.
- int [rt_queue_unbind](#) (RT_QUEUE *queue)
Unbind from a message queue.

5.81.1 Detailed Description

real-time IPC mechanism for sending messages of arbitrary size Message queueing is a method by which real-time tasks can exchange or pass data through a Xenomai-managed queue of messages. Messages can vary in length and be assigned different types or usages. A message queue can be created by one task and used by multiple tasks that send and/or receive messages to the queue.

5.81.2 Macro Definition Documentation

5.81.2.1 `#define Q_PRIO 0x1 /* Pend by task priority order. */`

Creation flags.

Referenced by `rt_queue_create()`.

5.81.2.2 `#define Q_SHARED 0x0`

Deprecated, compat only.

5.81.3 Function Documentation

5.81.3.1 `void * rt_queue_alloc (RT_QUEUE * q, size_t size)`

Allocate a message buffer.

This service allocates a message buffer from the queue's internal pool. This buffer can be filled in with payload information, prior enqueueing it by a call to [rt_queue_send\(\)](#). When used in pair, these services provide a zero-copy interface for sending messages.

Parameters

<i>q</i>	The descriptor address of the queue to allocate a buffer from.
<i>size</i>	The requested size in bytes of the buffer. Zero is an acceptable value, which means that the message conveys no payload; in this case, the receiver will get a zero-sized message.

Returns

The address of the allocated buffer upon success, or NULL if the call fails.

Valid calling context: any.

5.81.3.2 `int rt_queue_bind (RT_QUEUE * q, const char * name, RTIME timeout)`

Bind to a message queue.

This routine creates a new descriptor to refer to an existing message queue identified by its symbolic name. If the object does not exist on entry, the caller may block until a queue of the given name is created.

Parameters

<i>q</i>	The address of a queue descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the queue to bind to. This string should match the object name argument passed to rt_queue_create() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing TM_INFINITE causes the caller to block indefinitely until the object is registered. Passing TM_NONBLOCK causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- -EWOULDBLOCK is returned if *timeout* is equal to TM_NONBLOCK and the searched object is not registered on entry.
- -ETIMEDOUT is returned if the object cannot be retrieved within the specified amount of time.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.81.3.3 `int rt_queue_create (RT_QUEUE * q, const char * name, size_t poolsize, size_t qlimit, int mode)`

Create a message queue.

Create a message queue object which allows multiple tasks to exchange data through the use of variable-sized messages. A message queue is created empty.

This service needs the special character device `/dev/rtheap (10,254)` when called from user-space tasks.

Parameters

<i>q</i>	The address of a queue descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the queue. When non-NULL and non-empty, a copy of this string is used for indexing the created queue into the object registry.
<i>poolsize</i>	The size (in bytes) of the message buffer pool to be pre-allocated for holding messages. Message buffers will be claimed and released to this pool. The buffer pool memory cannot be extended. See note.
<i>qlimit</i>	This parameter allows to limit the maximum number of messages which can be queued at any point in time, sending to a full queue begets an error. The special value Q_UNLIMITED can be passed to disable the limit check.
<i>mode</i>	The queue creation mode. The following flags can be OR'ed into this bitmask, each of them affecting the new queue:

- Q_FIFO makes tasks pend in FIFO order on the queue for consuming messages.
- Q_PRIO makes tasks pend in priority order on the queue.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *mode* is invalid or *poolsize* is zero.
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the queue.
- -EEXIST is returned if the *name* is conflicting with an already registered queue.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Queues can be shared by multiple processes which belong to the same Xenomai session. Each message pending into the queue consumes four long words plus the actual payload size, aligned to the next long word boundary. e.g. a 6 byte message on a 32 bit platform would require 24 bytes of storage into the pool.

When *qlimit* is given (i.e. different from Q_UNLIMITED), this overhead is accounted for automatically, so that *qlimit* messages of *poolsize* / *qlimit* bytes can be stored into the pool concurrently. Otherwise, *poolsize* is increased by 5% internally to cope with such overhead.

References Q_PRIO.

5.81.3.4 `int rt_queue_delete (RT_QUEUE * q)`

Delete a message queue.

This routine deletes a queue object previously created by a call to [rt_queue_create\(\)](#). All resources attached to that queue are automatically released, including all pending messages.

Parameters

<i>q</i>	The descriptor address of the deleted queue.
----------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *q* is not a valid queue descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.81.3.5 int rt_queue_flush (RT_QUEUE * q)

Flush pending messages from a queue.

This routine flushes all messages currently pending in a queue, releasing all message buffers appropriately.

Parameters

<i>q</i>	The descriptor address of the queue to flush.
----------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *q* is not a valid queue descriptor.

Valid calling context: any.

5.81.3.6 int rt_queue_free (RT_QUEUE * q, void * buf)

Free a message buffer.

This service releases a message buffer to the queue's internal pool.

Parameters

<i>q</i>	The descriptor address of the queue to release a buffer to.
<i>buf</i>	The address of the message buffer to free. Even zero-sized messages carrying no payload data must be freed, since they are assigned a valid memory space to store internal information.

Returns

Zero is returned upon success, or -EINVAL if *buf* is not a valid message buffer previously allocated by the [rt_queue_alloc\(\)](#) service, or the caller did not get ownership of the message through a successful return from [rt_queue_receive\(\)](#).

Valid calling context: any.

5.81.3.7 `int rt_queue_inquire (RT_QUEUE * q, RT_QUEUE_INFO * info)`

Query queue status.

This routine returns the status information about the specified queue.

Parameters

<i>q</i>	The descriptor address of the queue to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *q* is not a valid queue descriptor.

Valid calling context: any.

References RT_QUEUE_INFO::mode, RT_QUEUE_INFO::name, RT_QUEUE_INFO::nmessages, RT_QUEUE_INFO::nwaiters, RT_QUEUE_INFO::poolsize, RT_QUEUE_INFO::qlimit, and RT_QUEUE_INFO::usedmem.

```
5.81.3.8  ssize_t rt_queue_read ( RT_QUEUE * q, void * buf, size_t size, RTIME timeout )
          [inline], [static]
```

Read from a queue (with relative scalar timeout).

This routine is a variant of [rt_queue_read_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>q</i>	The descriptor address of the message queue to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the received message payload.
<i>size</i>	The length in bytes of the memory area pointed to by <i>buf</i> .
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_queue_read_timed\(\)](#).

```
5.81.3.9  ssize_t rt_queue_read_timed ( RT_QUEUE * q, void * buf, size_t size, const struct
          timespec * abs_timeout )
```

Read from a queue.

This service reads the next available message from a given queue.

Parameters

<i>q</i>	The descriptor address of the message queue to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the received message payload. The internal message buffer conveying the data is automatically freed by this call.
<i>size</i>	The length in bytes of the memory area pointed to by <i>buf</i> . Messages larger than <i>size</i> are truncated appropriately.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for a message to be available from the queue (see note). Passing NULL causes the caller to block indefinitely until a message is available. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case no message is available.

Returns

The number of bytes copied to *buf* is returned upon success. Zero is a possible value corresponding to a zero-sized message passed to `rt_queue_send()` or `rt_queue_write()`. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before a message arrives.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and no message is immediately available on entry to the call.
- -EINTR is returned if `rt_task_unblock()` was called for the current task before a message was available.
- -EINVAL is returned if *q* is not a valid queue descriptor.
- -EIDRM is returned if *q* is deleted while the caller was waiting for a message. In such event, *q* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by `rt_queue_read()`, and `rt_queue_read_until()`.

```
5.81.3.10 ssize_t rt_queue_read_until ( RT_QUEUE * q, void * buf, size_t size, RTIME abs_timeout
) [inline], [static]
```

Read from a queue (with absolute scalar timeout).

This routine is a variant of `rt_queue_read_timed()` accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>q</i>	The descriptor address of the message queue to read from.
<i>buf</i>	A pointer to a memory area which will be written upon success with the received message payload.
<i>size</i>	The length in bytes of the memory area pointed to by <i>buf</i> .
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_queue_read_timed()`.

```
5.81.3.11 ssize_t rt_queue_receive ( RT_QUEUE * q, void ** bufp, RTIME timeout ) [inline],
[static]
```

Receive from a queue (with relative scalar timeout).

This routine is a variant of `rt_queue_receive_timed()` accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>q</i>	The descriptor address of the message queue to receive from.
<i>bufp</i>	A pointer to a memory location which will be written with the address of the received message.
<i>timeout</i>	A delay expressed in clock ticks.

References `rt_queue_receive_timed()`.

5.81.3.12 `ssize_t rt_queue_receive_timed (RT_QUEUE * q, void ** bufp, const struct timespec * abs_timeout)`

Receive a message from a queue (with absolute timeout date).

This service receives the next available message from a given queue.

Parameters

<i>q</i>	The descriptor address of the message queue to receive from.
<i>bufp</i>	A pointer to a memory location which will be written with the address of the received message, upon success. Once consumed, the message space should be freed using <code>rt_queue_free()</code> .
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for a message to be available from the queue (see note). Passing NULL causes the caller to block indefinitely until a message is available. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return immediately without blocking in case no message is available.

Returns

The number of bytes available from the received message is returned upon success. Zero is a possible value corresponding to a zero-sized message passed to `rt_queue_send()` or `rt_queue_write()`. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before a message arrives.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and no message is immediately available on entry to the call.
- -EINTR is returned if `rt_task_unblock()` was called for the current task before a message was available.
- -EINVAL is returned if *q* is not a valid queue descriptor.
- -EIDRM is returned if *q* is deleted while the caller was waiting for a message. In such event, *q* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } .

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `-alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by `rt_queue_receive()`, and `rt_queue_receive_until()`.

5.81.3.13 `ssize_t rt_queue_receive_until (RT_QUEUE * q, void ** bufp, RTIME abs_timeout)`
`[inline], [static]`

Receive from a queue (with absolute scalar timeout).

This routine is a variant of [rt_queue_receive_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>q</i>	The descriptor address of the message queue to receive from.
<i>bufp</i>	A pointer to a memory location which will be written with the address of the received message.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References [rt_queue_receive_timed\(\)](#).

5.81.3.14 `int rt_queue_send (RT_QUEUE * q, const void * buf, size_t size, int mode)`

Send a message to a queue.

This service sends a complete message to a given queue. The message must have been allocated by a previous call to [rt_queue_alloc\(\)](#).

Parameters

<i>q</i>	The descriptor address of the message queue to send to.
<i>buf</i>	The address of the message buffer to be sent. The message buffer must have been allocated using the rt_queue_alloc() service. Once passed to rt_queue_send() , the memory pointed to by <i>buf</i> is no more under the control of the sender and thus should not be referenced by it anymore; deallocation of this memory must be handled on the receiving side.
<i>size</i>	The actual size in bytes of the message, which may be lower than the allocated size for the buffer obtained from rt_queue_alloc() . Zero is a valid value, in which case an empty message will be sent.
<i>mode</i>	A set of flags affecting the operation:

- `Q_URGENT` causes the message to be prepended to the message queue, ensuring a LIFO ordering.
- `Q_NORMAL` causes the message to be appended to the message queue, ensuring a FIFO ordering.
- `Q_BROADCAST` causes the message to be sent to all tasks currently waiting for messages. The message is not copied; a reference count is maintained instead so that the message will remain valid until the last receiver releases its own reference using [rt_queue_free\(\)](#), after which the message space will be returned to the queue's internal pool.

Returns

Upon success, this service returns the number of receivers which got awoken as a result of the operation. If zero is returned, no task was waiting on the receiving side of the queue, and the message has been enqueued. Upon error, one of the following error codes is returned:

- `-EINVAL` is returned if *q* is not a message queue descriptor, *mode* is invalid, or *buf* is `NULL`.
- `-ENOMEM` is returned if queuing the message would exceed the limit defined for the queue at creation.

Valid calling context: any.

5.81.3.15 `int rt_queue_unbind (RT_QUEUE * q)`

Unbind from a message queue.

Parameters

<i>q</i>	The descriptor address of the queue to unbind from.
----------	---

This routine releases a previous binding to a message queue. After this call has returned, the descriptor is no more valid for referencing this object.

5.82 Semaphore services

Counting semaphore IPC mechanism.

Collaboration diagram for Semaphore services:



Data Structures

- struct `RT_SEM_INFO`
Semaphore status descriptor.

Macros

- `#define S_PRIO 0x1` /* Pend by task priority order. */
Creation flags.

Functions

- int `rt_sem_create` (RT_SEM *sem, const char *name, unsigned long icount, int mode)
Create a counting semaphore.
- int `rt_sem_delete` (RT_SEM *sem)
Delete a semaphore.
- int `rt_sem_p_timed` (RT_SEM *sem, const struct timespec *abs_timeout)
Pend on a semaphore.
- static int `rt_sem_p_until` (RT_SEM *sem, RTIME timeout)
Pend on a semaphore (with absolute scalar timeout).
- static int `rt_sem_p` (RT_SEM *sem, RTIME timeout)
Pend on a semaphore (with relative scalar timeout).
- int `rt_sem_v` (RT_SEM *sem)
Signal a semaphore.
- int `rt_sem_broadcast` (RT_SEM *sem)
Broadcast a semaphore.
- int `rt_sem_inquire` (RT_SEM *sem, RT_SEM_INFO *info)
Query semaphore status.
- int `rt_sem_bind` (RT_SEM *sem, const char *name, RTIME timeout)
Bind to a semaphore.
- int `rt_sem_unbind` (RT_SEM *sem)
Unbind from a semaphore.

5.82.1 Detailed Description

Counting semaphore IPC mechanism. A counting semaphore is a synchronization object for controlling the concurrency level allowed in accessing a resource from multiple real-time tasks, based on the value of a count variable accessed atomically. The semaphore is used through the P ("Proberen", from the Dutch "test and decrement") and V ("Verhogen", increment) operations. The P operation decrements the semaphore count by one if non-zero, or waits until a V operation is issued by another task. Conversely, the V operation releases a resource by incrementing the count by one, unblocking the heading task waiting on the P operation if any. Waiting on a semaphore may cause a priority inversion.

If no more than a single resource is made available at any point in time, the semaphore enforces mutual exclusion and thus can be used to serialize access to a critical section. However, mutexes should be used instead in order to prevent priority inversions, based on the priority inheritance protocol.

5.82.2 Macro Definition Documentation

5.82.2.1 `#define S_PRIO 0x1 /* Pend by task priority order. */`

Creation flags.

Referenced by `rt_sem_create()`.

5.82.3 Function Documentation

5.82.3.1 `int rt_sem_bind (RT_SEM * sem, const char * name, RTIME timeout)`

Bind to a semaphore.

This routine creates a new descriptor to refer to an existing semaphore identified by its symbolic name. If the object does not exist on entry, the caller may block until a semaphore of the given name is created.

Parameters

<i>sem</i>	The address of a semaphore descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the semaphore to bind to. This string should match the object name argument passed to <code>rt_sem_create()</code> .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing <code>TM_INFINITE</code> causes the caller to block indefinitely until the object is registered. Passing <code>TM_NONBLOCK</code> causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if `rt_task_unblock()` was called for the current task before the retrieval has completed.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `TM_NONBLOCK` and the searched object is not registered on entry.
- `-ETIMEDOUT` is returned if the object cannot be retrieved within the specified amount of time.
- `-EPERM` is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads

- Any other context if *timeout* equals TM_NONBLOCK.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `--alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.82.3.2 `int rt_sem_broadcast (RT_SEM * sem)`

Broadcast a semaphore.

All tasks currently waiting on the semaphore are immediately unblocked. The semaphore count is set to zero.

Parameters

<i>sem</i>	The descriptor address of the semaphore to broadcast.
------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *sem* is not a valid semaphore descriptor.

Valid calling context: any.

5.82.3.3 `int rt_sem_create (RT_SEM * sem, const char * name, unsigned long icount, int mode)`

Create a counting semaphore.

Parameters

<i>sem</i>	The address of a semaphore descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the semaphore. When non-NULL and non-empty, a copy of this string is used for indexing the created semaphore into the object registry.
<i>icount</i>	The initial value of the counting semaphore.
<i>mode</i>	The semaphore creation mode. The following flags can be OR'ed into this bitmask:

- S_FIFO makes tasks pend in FIFO order on the semaphore.
- S_PRIO makes tasks pend in priority order on the semaphore.
- S_PULSE causes the semaphore to behave in "pulse" mode. In this mode, the V (signal) operation attempts to release a single waiter each time it is called, without incrementing the semaphore count, even if no waiter is pending. For this reason, the semaphore count in pulse mode remains zero.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *icount* is non-zero and S_PULSE is set in *mode*, or *mode* is otherwise invalid.
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the semaphore.

- -EEXIST is returned if the *name* is conflicting with an already registered semaphore.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

Semaphores can be shared by multiple processes which belong to the same Xenomai session.

References S_PRIO.

5.82.3.4 `int rt_sem_delete (RT_SEM * sem)`

Delete a semaphore.

This routine deletes a semaphore previously created by a call to [rt_sem_create\(\)](#).

Parameters

<i>sem</i>	The descriptor address of the deleted object.
------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *sem* is not a valid semaphore descriptor.
- -EPERM is returned if this service was called from an asynchronous context.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.82.3.5 `int rt_sem_inquire (RT_SEM * sem, RT_SEM_INFO * info)`

Query semaphore status.

This routine returns the status information about the specified semaphore.

Parameters

<i>sem</i>	The descriptor address of the semaphore to get the status of.
<i>info</i>	A pointer to the returnbuffer to copy the information to.

Returns

Zero is returned and status information is written to the structure pointed at by *info* upon success. Otherwise:

- -EINVAL is returned if *sem* is not a valid semaphore descriptor.

Valid calling context: any.

References RT_SEM_INFO::count, RT_SEM_INFO::name, and RT_SEM_INFO::nwaiters.

5.82.3.6 `int rt_sem_p (RT_SEM * sem, RTIME timeout) [inline], [static]`

Pend on a semaphore (with relative scalar timeout).

This routine is a variant of `rt_sem_p_timed()` accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>sem</i>	The descriptor address of the semaphore to wait on.
<i>timeout</i>	A delay expressed in clock ticks.

References `rt_sem_p_timed()`.

5.82.3.7 `int rt_sem_p_timed (RT_SEM * sem, const struct timespec * abs_timeout)`

Pend on a semaphore.

Test and decrement the semaphore count. If the semaphore value is greater than zero, it is decremented by one and the service immediately returns to the caller. Otherwise, the caller is blocked until the semaphore is either signaled or destroyed, unless a non-blocking operation was required.

Parameters

<i>sem</i>	The descriptor address of the semaphore to wait on.
<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for the request to be satisfied (see note). Passing NULL causes the caller to block indefinitely until the request is satisfied. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return without blocking in case the request cannot be satisfied immediately.

Returns

Zero is returned upon success. Otherwise:

- -ETIMEDOUT is returned if *abs_timeout* is reached before the request is satisfied.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and the semaphore count is zero on entry.
- -EINTR is returned if `rt_task_unblock()` was called for the current task before the request is satisfied.
- -EINVAL is returned if *sem* is not a valid semaphore descriptor.
- -EIDRM is returned if *sem* is deleted while the caller was sleeping on it. In such a case, *sem* is no more valid upon return of this service.
- -EPERM is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 }.

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by `rt_sem_p()`, and `rt_sem_p_until()`.

5.82.3.8 `int rt_sem_p_until (RT_SEM * sem, RTIME abs_timeout) [inline], [static]`

Pend on a semaphore (with absolute scalar timeout).

This routine is a variant of [rt_sem_p_timed\(\)](#) accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>sem</i>	The descriptor address of the semaphore to wait on.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References [rt_sem_p_timed\(\)](#).

5.82.3.9 `int rt_sem_unbind (RT_SEM * sem)`

Unbind from a semaphore.

Parameters

<i>sem</i>	The descriptor address of the semaphore to unbind from.
------------	---

This routine releases a previous binding to a semaphore. After this call has returned, the descriptor is no more valid for referencing this object.

5.82.3.10 `int rt_sem_v (RT_SEM * sem)`

Signal a semaphore.

If the semaphore is pended, the task heading the wait queue is immediately unblocked. Otherwise, the semaphore count is incremented by one, unless the semaphore is used in "pulse" mode (see [rt_sem_create\(\)](#)).

Parameters

<i>sem</i>	The descriptor address of the semaphore to signal.
------------	--

Returns

Zero is returned upon success. Otherwise:


- -EINVAL is returned if *sem* is not a valid semaphore descriptor.

Valid calling context: any.

5.83 Task management services

Services dealing with preemptive multi-tasking.

Collaboration diagram for Task management services:



Data Structures

- struct [RT_TASK_INFO](#)
Task status descriptor.

Macros

- #define [T_LOPRIO](#) 0
Task priorities.
- #define [T_LOCK](#) __THREAD_M_LOCK
Task mode bits.
- #define [T_WARNSW](#) __THREAD_M_WARNSW
Cobalt only, nop over Mercury.
- #define [T_FPU](#) 0x0
Deprecated, compat only.

Functions

- int [rt_task_create](#) (RT_TASK *task, const char *name, int stksize, int prio, int mode)
Create a real-time task.
- int [rt_task_delete](#) (RT_TASK *task)
Delete a real-time task.
- int [rt_task_set_affinity](#) (RT_TASK *task, const cpu_set_t *cpus)
Set CPU affinity of real-time task.
- int [rt_task_start](#) (RT_TASK *task, void(*entry)(void *arg), void *arg)
Start a real-time task.
- int [rt_task_spawn](#) (RT_TASK *task, const char *name, int stksize, int prio, int mode, void(*entry)(void *arg), void *arg)
Create and start a real-time task.
- int [rt_task_shadow](#) (RT_TASK *task, const char *name, int prio, int mode)
Turn caller into a real-time task.
- int [rt_task_join](#) (RT_TASK *task)
Wait on the termination of a real-time task.
- int [rt_task_set_periodic](#) (RT_TASK *task, RTIME idate, RTIME period)
Make a real-time task periodic.

- int `rt_task_wait_period` (unsigned long *overruns_r)
Wait for the next periodic release point.
- int `rt_task_sleep` (RTIME delay)
Delay the current real-time task (with relative delay).
- int `rt_task_sleep_until` (RTIME date)
Delay the current real-time task (with absolute wakeup date).
- int `rt_task_same` (RT_TASK *task1, RT_TASK *task2)
Compare real-time task descriptors.
- int `rt_task_suspend` (RT_TASK *task)
Suspend a real-time task.
- int `rt_task_resume` (RT_TASK *task)
Resume a real-time task.
- RT_TASK * `rt_task_self` (void)
Retrieve the current task descriptor.
- int `rt_task_set_priority` (RT_TASK *task, int prio)
Change the base priority of a real-time task.
- int `rt_task_set_mode` (int clrmask, int setmask, int *mode_r)
Change the current task mode.
- int `rt_task_yield` (void)
Manual round-robin.
- int `rt_task_unblock` (RT_TASK *task)
Unblock a real-time task.
- int `rt_task_slice` (RT_TASK *task, RTIME quantum)
Set a task's round-robin quantum.
- int `rt_task_inquire` (RT_TASK *task, RT_TASK_INFO *info)
Retrieve information about a real-time task.
- ssize_t `rt_task_send_timed` (RT_TASK *task, RT_TASK_MCB *mcb_s, RT_TASK_MCB *mcb_r, const struct timespec *abs_timeout)
Send a message to a real-time task.
- static ssize_t `rt_task_send_until` (RT_TASK *task, RT_TASK_MCB *mcb_s, RT_TASK_MCB *mcb_r, RTIME timeout)
Send a message to a real-time task (with absolute scalar timeout).
- static ssize_t `rt_task_send` (RT_TASK *task, RT_TASK_MCB *mcb_s, RT_TASK_MCB *mcb_r, RTIME timeout)
Send a message to a real-time task (with relative scalar timeout).
- int `rt_task_receive_timed` (RT_TASK_MCB *mcb_r, const struct timespec *abs_timeout)
Receive a message from a real-time task.
- static int `rt_task_receive_until` (RT_TASK_MCB *mcb_r, RTIME timeout)
Receive a message from a real-time task (with absolute scalar timeout).
- static int `rt_task_receive` (RT_TASK_MCB *mcb_r, RTIME timeout)
Receive a message from a real-time task (with relative scalar timeout).
- int `rt_task_reply` (int flowid, RT_TASK_MCB *mcb_s)
Reply to a remote task message.
- int `rt_task_bind` (RT_TASK *task, const char *name, RTIME timeout)
Bind to a task.
- int `rt_task_unbind` (RT_TASK *task)
Unbind from a task.

5.83.1 Detailed Description

Services dealing with preemptive multi-tasking. Each Alchemy task is an independent portion of the overall application code embodied in a C procedure, which executes on its own stack context.

5.83.2 Macro Definition Documentation

5.83.2.1 `#define T_FPU 0x0`

Deprecated, compat only.

5.83.2.2 `#define T_LOCK __THREAD_M_LOCK`

Task mode bits.

Referenced by `rt_task_create()`, `rt_task_set_mode()`, and `rt_task_shadow()`.

5.83.2.3 `#define T_LOPRIO 0`

Task priorities.

5.83.2.4 `#define T_WARNSW __THREAD_M_WARNSW`

Cobalt only, nop over Mercury.

Referenced by `rt_task_create()`, `rt_task_set_mode()`, and `rt_task_shadow()`.

5.83.3 Function Documentation

5.83.3.1 `int rt_task_bind (RT_TASK * task, const char * name, RTIME timeout)`

Bind to a task.

This routine creates a new descriptor to refer to an existing Alchemy task identified by its symbolic name. If the object does not exist on entry, the caller may block until a task of the given name is created.

Parameters

<i>task</i>	The address of a task descriptor filled in by the operation. Contents of this memory is undefined upon failure.
<i>name</i>	A valid NULL-terminated name which identifies the task to bind to. This string should match the object name argument passed to rt_task_create() , or rt_task_shadow() .
<i>timeout</i>	The number of clock ticks to wait for the registration to occur (see note). Passing <code>TM_INFINITE</code> causes the caller to block indefinitely until the object is registered. Passing <code>TM_NONBLOCK</code> causes the service to return immediately without waiting if the object is not registered on entry.

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if [rt_task_unblock\(\)](#) was called for the current task before the retrieval has completed.
- `-EWOULDBLOCK` is returned if *timeout* is equal to `TM_NONBLOCK` and the searched object is not registered on entry.
- `-ETIMEDOUT` is returned if the object cannot be retrieved within the specified amount of time.
- `-EPERM` is returned if this service should block, but was not called from a Xenomai thread.

Valid calling contexts:

- Xenomai threads
- Any other context if *timeout* equals `TM_NONBLOCK`.

Note

The *timeout* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.83.3.2 `int rt_task_create (RT_TASK * task, const char * name, int stksize, int prio, int mode)`

Create a real-time task.

This service creates a task with access to the full set of Xenomai real-time services. If *prio* is non-zero, the new task belongs to Xenomai's real-time FIFO scheduling class, aka `SCHED_FIFO`. If *prio* is zero, the task belongs to the regular `SCHED_OTHER` class.

Creating tasks with zero priority is useful for running non real-time processes which may invoke blocking real-time services, such as pending on a semaphore, reading from a message queue or a buffer, and so on.

Once created, the task is left dormant until it is actually started by [rt_task_start\(\)](#).

Parameters

<i>task</i>	The address of a task descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the task. When non-NULL and non-empty, a copy of this string is used for indexing the created task into the object registry.
<i>stksize</i>	The size of the stack (in bytes) for the new task. If zero is passed, a system-dependent default size will be substituted.
<i>prio</i>	The base priority of the new task. This value must be in the <code>[0 .. 99]</code> range, where 0 is the lowest effective priority.
<i>mode</i>	The task creation mode. The following flags can be OR'ed into this bitmask:

- `T_JOINABLE` allows another task to wait on the termination of the new task. [rt_task_join\(\)](#) shall be called for this task to clean up any resources after its termination.
- `T_LOCK` causes the new task to lock the scheduler prior to entering the user routine specified by [rt_task_start\(\)](#). A call to [rt_task_set_mode\(\)](#) from the new task is required to drop this lock.
- When running over the Cobalt core, `T_WARNSW` causes the `SIGXCPU` signal to be sent to the current task whenever it switches to the secondary mode. This feature is useful to detect unwanted migrations to the Linux domain. This flag has no effect over the Mercury core.

Returns

Zero is returned upon success. Otherwise:

- `-EINVAL` is returned if either *prio*, *mode* or *stksize* are invalid.
- `-ENOMEM` is returned if the system fails to get memory from the main heap in order to create the task.
- `-EEXIST` is returned if the *name* is conflicting with an already registered task.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Core specifics:

When running over the Cobalt core:

- calling [rt_task_create\(\)](#) causes SCHED_FIFO tasks to switch to secondary mode.
- members of Xenomai's SCHED_FIFO class running in the primary domain have utmost priority over all Linux activities in the system, including Linux interrupt handlers.

When running over the Mercury core, the new task belongs to the regular POSIX SCHED_FIFO class.

Note

Tasks can be referred to from multiple processes which all belong to the same Xenomai session.

Examples:

[cross-link.c](#).

References T_LOCK, and T_WARN SW.

Referenced by [rt_task_spawn\(\)](#).

5.83.3.3 int [rt_task_delete](#) (RT_TASK * task)

Delete a real-time task.

This call terminates a task previously created by [rt_task_create\(\)](#).

Tasks created with the T_JOINABLE flag shall be joined by a subsequent call to [rt_task_join\(\)](#) once successfully deleted, to reclaim all resources.

Parameters

<i>task</i>	The descriptor address of the deleted task, or NULL for self-deletion.
-------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.
- -EPERM is returned if *task* is NULL and this service was called from an invalid context. In addition, this error is always raised when this service is called from asynchronous context, such as a timer/alarm handler.

Valid calling context:

- Alchemy tasks only if *task* is NULL, any thread context otherwise.

Examples:

[cross-link.c](#).

5.83.3.4 int [rt_task_inquire](#) (RT_TASK * task, RT_TASK_INFO * info)

Retrieve information about a real-time task.

Return various information about an Alchemy task. This service may also be used to probe for task existence.

Parameters

<i>task</i>	The descriptor address of the task. If <i>task</i> is NULL, information about the current task is returned.
<i>info</i>	The address of a structure the task information will be written to. Passing NULL is valid, in which case the system is only probed for existence of the specified task.

Returns

Zero is returned if the task exists. In addition, if *info* is non-NULL, it is filled in with task information.

- -EINVAL is returned if *task* is not a valid task descriptor, or if *prio* is invalid.
- -EPERM is returned if *task* is NULL and this service was called from an invalid context.

Valid calling context:

- Alchemy tasks if *task* is NULL, any otherwise.

References RT_TASK_INFO::name, RT_TASK_INFO::pid, RT_TASK_INFO::prio, and RT_TASK_INFO::stat.

5.83.3.5 int rt_task_join (RT_TASK * task)

Wait on the termination of a real-time task.

This service blocks the caller in non-real-time context until *task* has terminated. All resources are released after successful completion of this service.

The specified task must have been created by the same process that wants to join it, and the T_JOINABLE mode flag must have been set on creation to [rt_task_create\(\)](#).

Tasks created with the T_JOINABLE flag shall be joined by a subsequent call to [rt_task_join\(\)](#) once successfully deleted, to reclaim all resources.

Parameters

<i>task</i>	The descriptor address of the task to join.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.
- -EINVAL is returned if the task was not created with T_JOINABLE set or some other task is already waiting on the termination.
- -EDEADLK is returned if *task* refers to the caller.
- -ESRCH is returned if *task* no longer exists or refers to task created by a different process.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

After successful completion of this service, it is neither required nor valid to additionally invoke [rt_task_delete\(\)](#) on the same task.

5.83.3.6 `ssize_t rt_task_receive (RT_TASK_MCB * mcb_r, RTIME timeout) [inline], [static]`

Receive a message from a real-time task (with relative scalar timeout).

This routine is a variant of [rt_task_receive_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>mcb_r</i>	The address of a message control block referring to the receive message area.
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_task_receive_timed\(\)](#).

5.83.3.7 `int rt_task_receive_timed (RT_TASK_MCB * mcb_r, const struct timespec * abs_timeout)`

Receive a message from a real-time task.

This service is part of the synchronous message passing support available to Alchemy tasks. The caller receives a variable-sized message from another task. The sender is blocked until the caller invokes [rt_task_reply\(\)](#) to finish the transaction.

A basic message control block is used to store the location and size of the data area to receive from the client, in addition to a user-defined operation code.

Parameters

<i>mcb_r</i>	The address of a message control block referring to the receive message area. The fields from this control block should be set as follows:
--------------	--

- *mcb_r->data* should contain the address of a buffer large enough to collect the data sent by the remote task;
- *mcb_r->size* should contain the size in bytes of the buffer space pointed at by *mcb_r->data*. If *mcb_r->size* is lower than the actual size of the received message, no data copy takes place and `-ENOBUFS` is returned to the caller. See note.

Upon return, *mcb_r->opcode* will contain the operation code sent from the remote task using [rt_task_send\(\)](#).

Parameters

<i>abs_timeout</i>	The number of clock ticks to wait for receiving a message (see note). Passing NULL causes the caller to block indefinitely until a remote task eventually sends a message. Passing { <i>.tv_sec</i> = 0, <i>.tv_nsec</i> = 0 } causes the service to return immediately without waiting if no remote task is currently waiting for sending a message.
--------------------	---

Returns

A strictly positive value is returned upon success, representing a flow identifier for the opening transaction; this token should be passed to [rt_task_reply\(\)](#), in order to send back a reply to and unblock the remote task appropriately. Otherwise:

- `-EPERM` is returned if this service was called from an invalid context.
- `-EINTR` is returned if [rt_task_unblock\(\)](#) was called for the current task before a message was received.
- `-ENOBUFS` is returned if *mcb_r* does not point at a message area large enough to collect the remote task's message.
- `-EWOULDBLOCK` is returned if *abs_timeout* is { *.tv_sec* = 0, *.tv_nsec* = 0 } and no remote task is currently waiting for sending a message to the caller.

- -ETIMEDOUT is returned if no message was received within the *timeout*.

Valid calling context:

- Alchemy tasks

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by `rt_task_receive()`, and `rt_task_receive_until()`.

```
5.83.3.8 ssize_t rt_task_receive_until ( RT_TASK_MCB * mcb_r, RTIME abs_timeout ) [inline],
[static]
```

Receive a message from a real-time task (with absolute scalar timeout).

This routine is a variant of `rt_task_receive_timed()` accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>mcb_r</i>	The address of a message control block referring to the receive message area.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_task_receive_timed()`.

```
5.83.3.9 int rt_task_reply ( int flowid, RT_TASK_MCB * mcb_s )
```

Reply to a remote task message.

This service is part of the synchronous message passing support available to Alchemy tasks. The caller sends a variable-sized message back to a remote task, in response to this task's initial message received by a call to `rt_task_receive()`. As a consequence of calling `rt_task_reply()`, the remote task will be unblocked from the `rt_task_send()` service.

A basic message control block is used to store the location and size of the data area to send back, in addition to a user-defined status code.

Parameters

<i>flowid</i>	The flow identifier returned by a previous call to <code>rt_task_receive()</code> which uniquely identifies the current transaction.
<i>mcb_s</i>	The address of an optional message control block referring to the message to be sent back. If <i>mcb_s</i> is NULL, the remote will be unblocked without getting any reply data. When <i>mcb_s</i> is valid, the fields from this control block should be set as follows:

- *mcb_s->data* should contain the address of the payload data to send to the remote task.
- *mcb_s->size* should contain the size in bytes of the payload data pointed at by *mcb_s->data*. Zero is a legitimate value, and indicates that no payload data will be transferred. In the latter case, *mcb_s->data* will be ignored.
- *mcb_s->opcode* is an opaque status code carried during the message transfer the caller can fill with any appropriate value. It will be made available "as is" to the remote task into the status code field by the `rt_task_send()` service. If *mcb_s* is NULL, Zero will be returned to the remote task into the status code field.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *flowid* is invalid.
- -ENXIO is returned if *flowid* does not match the expected identifier returned from the latest call of the current task to [rt_task_receive\(\)](#), or if the remote task stopped waiting for the reply in the meantime (e.g. the remote could have been deleted or forcibly unblocked).
- -ENOBUFS is returned if the reply data referred to by *mcb_s* is larger than the reply area mentioned by the remote task when calling [rt_task_send\(\)](#). In such a case, the remote task also receives -ENOBUFS on return from [rt_task_send\(\)](#).
- -EPERM is returned if this service was called from an invalid context.

Valid calling context:

- Alchemy tasks

5.83.3.10 int rt_task_resume (RT_TASK * task)

Resume a real-time task.

Forcibly resume the execution of a task which was previously suspended by a call to [rt_task_suspend\(\)](#), if the suspend nesting count decrements to zero.

Parameters

<i>task</i>	The descriptor address of the task to resume.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.

Valid calling context: any.

Note

Blocked and suspended task states are cumulative. Therefore, resuming a task currently waiting on a synchronization object (e.g. semaphore, queue) does not make it eligible for scheduling until the awaited resource is eventually acquired, or a timeout elapses.

5.83.3.11 int rt_task_same (RT_TASK * task1, RT_TASK * task2)

Compare real-time task descriptors.

This predicate returns true if *task1* and *task2* refer to the same task.

Parameters

<i>task1</i>	First task descriptor to compare.
--------------	-----------------------------------

<i>task2</i>	Second task descriptor to compare.
--------------	------------------------------------

Returns

A non-zero value is returned if both descriptors refer to the same task, zero otherwise.

5.83.3.12 `RT_TASK * rt_task_self (void)`

Retrieve the current task descriptor.

Return the address of the current Alchemy task descriptor.

Returns

The address of the task descriptor referring to the current Alchemy task is returned upon success, or NULL if not called from an valid task context.

Valid calling context:

- Alchemy tasks.

5.83.3.13 `ssize_t rt_task_send (RT_TASK * task, RT_TASK_MCB * mcb_s, RT_TASK_MCB * mcb_r, RTIME timeout) [inline], [static]`

Send a message to a real-time task (with relative scalar timeout).

This routine is a variant of [rt_task_send_timed\(\)](#) accepting a relative timeout specification expressed as a scalar value.

Parameters

<i>task</i>	The descriptor address of the recipient task.
<i>mcb_s</i>	The address of the message control block referring to the message to be sent.
<i>mcb_r</i>	The address of an optional message control block referring to the reply message area.
<i>timeout</i>	A delay expressed in clock ticks.

References [rt_task_send_timed\(\)](#).

5.83.3.14 `ssize_t rt_task_send_timed (RT_TASK * task, RT_TASK_MCB * mcb_s, RT_TASK_MCB * mcb_r, const struct timespec * abs_timeout)`

Send a message to a real-time task.

This service is part of the synchronous message passing support available to Alchemy tasks. The caller sends a variable-sized message to another task, waiting for the remote to receive the initial message by a call to [rt_task_receive\(\)](#), then reply to it using [rt_task_reply\(\)](#).

A basic message control block is used to store the location and size of the data area to send or retrieve upon reply, in addition to a user-defined operation code.

Parameters

<i>task</i>	The descriptor address of the recipient task.
<i>mcb_s</i>	The address of the message control block referring to the message to be sent. The fields from this control block should be set as follows:

- *mcb_s->data* should contain the address of the payload data to send to the remote task.
- *mcb_s->size* should contain the size in bytes of the payload data pointed at by *mcb_s->data*. Zero is a legitimate value, and indicates that no payload data will be transferred. In the latter case, *mcb_s->data* will be ignored.
- *mcb_s->opcode* is an opaque operation code carried during the message transfer, the caller can fill with any appropriate value. It will be made available "as is" to the remote task into the operation code field by the [rt_task_receive\(\)](#) service.

Parameters

<i>mcb_r</i>	The address of an optional message control block referring to the reply message area. If <i>mcb_r</i> is NULL and a reply is sent back by the remote task, the reply message will be discarded, and -ENOBUFS will be returned to the caller. When <i>mcb_r</i> is valid, the fields from this control block should be set as follows:
--------------	---

- *mcb_r->data* should contain the address of a buffer large enough to collect the reply data from the remote task.
- *mcb_r->size* should contain the size in bytes of the buffer space pointed at by *mcb_r->data*. If *mcb_r->size* is lower than the actual size of the reply message, no data copy takes place and -ENOBUFS is returned to the caller.

Upon return, *mcb_r->opcode* will contain the status code sent back from the remote task using [rt_task_reply\(\)](#), or zero if unspecified.

Parameters

<i>abs_timeout</i>	An absolute date expressed in clock ticks, specifying a time limit to wait for the recipient task to reply to the initial message (see note). Passing NULL causes the caller to block indefinitely until a reply is received. Passing { .tv_sec = 0, .tv_nsec = 0 } causes the service to return without blocking in case the recipient task is not waiting for messages at the time of the call.
--------------------	---

Returns

A positive value is returned upon success, representing the length (in bytes) of the reply message returned by the remote task. Zero is a success status, meaning either that *mcb_r* was NULL on entry, or that no actual message was passed to the remote call to [rt_task_reply\(\)](#). Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.
- -EPERM is returned if this service was called from an invalid context.
- -ENOBUFS is returned if *mcb_r* does not point at a message area large enough to collect the remote task's reply. This includes the case where *mcb_r* is NULL on entry, despite the remote task attempts to send a reply message.
- -EWOULDBLOCK is returned if *abs_timeout* is { .tv_sec = 0, .tv_nsec = 0 } and the recipient *task* is not currently waiting for a message on the [rt_task_receive\(\)](#) service.
- -EIDRM is returned if *task* has been deleted while waiting for a reply.
- -EINTR is returned if [rt_task_unblock\(\)](#) was called for the current task before any reply was received from the recipient *task*.

Valid calling context:

- Xenomai threads

Note

abs_timeout is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Referenced by `rt_task_send()`, and `rt_task_send_until()`.

```
5.83.3.15  ssize_t rt_task_send_until ( RT_TASK * task, RT_TASK_MCB * mcb_s, RT_TASK_MCB
          * mcb_r, RTIME abs_timeout )  [inline], [static]
```

Send a message to a real-time task (with absolute scalar timeout).

This routine is a variant of `rt_task_send_timed()` accepting an absolute timeout specification expressed as a scalar value.

Parameters

<i>task</i>	The descriptor address of the recipient task.
<i>mcb_s</i>	The address of the message control block referring to the message to be sent.
<i>mcb_r</i>	The address of an optional message control block referring to the reply message area.
<i>abs_timeout</i>	An absolute date expressed in clock ticks.

References `rt_task_send_timed()`.

```
5.83.3.16  int rt_task_set_affinity ( RT_TASK * task, const cpu_set_t * cpus )
```

Set CPU affinity of real-time task.

This calls makes *task* affine to the set of CPUs defined by *cpus*.

Parameters

<i>task</i>	The descriptor address of the task. If <i>task</i> is NULL, the CPU affinity of the current task is changed.
<i>cpus</i>	The set of CPUs <i>task</i> should be affine to.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is NULL but the caller is not a Xenomai task, or if *task* is non-NULL but not a valid task descriptor.
- -EINVAL is returned if *cpus* contains no processors that are currently physically on the system and permitted to the process according to any restrictions that may be imposed by the "cpuset" mechanism described in `cpuset(7)`.

Valid calling context:

- Alchemy tasks if *task* is NULL, any otherwise.

5.83.3.17 `int rt_task_set_mode (int clrmask, int setmask, int * mode_r)`

Change the current task mode.

Each Alchemy task has a set of internal flags determining several operating conditions. `rt_task_set_mode()` takes a bitmask of mode bits to clear for disabling the corresponding modes for the current task, and another one to set for enabling them. The mode bits which were previously in effect before the change can be returned upon request.

The following bits can be part of the bitmask:

- T_LOCK causes the current task to lock the scheduler on the current CPU, preventing all further involuntary task switches on this CPU. Clearing this bit unlocks the scheduler.
- Only when running over the Cobalt core:
 - T_WARNSW causes the SIGXCPU signal to be sent to the current task whenever it switches to the secondary mode. This feature is useful to detect unwanted migrations to the Linux domain.
 - T_CONFORMING can be passed in *setmask* to switch the current Alchemy task to its preferred runtime mode. The only meaningful use of this switch is to force a real-time task back to primary mode (see note). Any other use leads to a nop.

These two last flags have no effect over the Mercury core, and are simply ignored.

Parameters

<i>clrmask</i>	A bitmask of mode bits to clear for the current task, before <i>setmask</i> is applied. Zero is an acceptable value which leads to a no-op.
<i>setmask</i>	A bitmask of mode bits to set for the current task. Zero is an acceptable value which leads to a no-op.
<i>mode_r</i>	If non-NULL, <i>mode_r</i> must be a pointer to a memory location which will be written upon success with the previous set of active mode bits. If NULL, the previous set of active mode bits will not be returned.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor, or if any bit from *clrmask* or *setmask* is invalid.
- -EPERM is returned if this service was called from an invalid context.

Valid calling context:

- Alchemy tasks.

Note

Forcing the task mode using the T_CONFORMING bit from user code is almost always wrong, since the Xenomai/cobalt core handles mode switches internally when/if required. Most often, manual mode switching from applications introduces useless overhead. This mode bit is part of the API only to cover rare use cases in middleware code based on the Alchemy interface.

References T_LOCK, and T_WARNSW.

5.83.3.18 `int rt_task_set_periodic (RT_TASK * task, RTIME idate, RTIME period)`

Make a real-time task periodic.

Make a task periodic by programing its first release point and its period in the processor time line. *task* should then call [rt_task_wait_period\(\)](#) to sleep until the next periodic release point in the processor timeline is reached.

Parameters

<i>task</i>	The descriptor address of the periodic task. If <i>task</i> is NULL, the current task is made periodic. <i>task</i> must belong the current process.
<i>idate</i>	The initial (absolute) date of the first release point, expressed in clock ticks (see note). If <i>idate</i> is equal to TM_NOW, the current system date is used.
<i>period</i>	The period of the task, expressed in clock ticks (see note). Passing TM_INFINITE stops the task's periodic timer if enabled, then returns successfully.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is NULL but the caller is not a Xenomai task, or if *task* is non-NULL but not a valid task descriptor.
- -ETIMEDOUT is returned if *idate* is different from TM_INFINITE and represents a date in the past.

Valid calling contexts:

- Alchemy tasks if *task* is NULL, any otherwise.

Core specifics:

Over Cobalt, -EINVAL is returned if *period* is different from TM_INFINITE but shorter than the scheduling latency value for the target system, as available from `/proc/xenomai/latency`.

Note

The *idate* and *period* values are interpreted as a multiple of the Alchemy clock resolution (see `-alchemy-clock-resolution` option, defaults to 1 nanosecond).

Examples:

[cross-link.c](#).

5.83.3.19 `int rt_task_set_priority (RT_TASK * task, int prio)`

Change the base priority of a real-time task.

The base priority of a task defines the relative importance of the work being done by each task, which gains control of the CPU accordingly.

Changing the base priority of a task does not affect the priority boost the target task might have obtained as a consequence of a priority inheritance undergoing.

Parameters

<i>task</i>	The descriptor address of the task to update. If <i>task</i> is NULL, the priority of the current task is changed.
<i>prio</i>	The new priority. This value must range from [T_LOPRIO .. T_HIPRIO] (inclusive) where T_LOPRIO is the lowest effective priority.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor, or if *prio* is invalid.
- -EPERM is returned if *task* is NULL and this service was called from an invalid context.

Valid calling context:

- Alchemy tasks if *task* is NULL, any otherwise.

Note

Assigning the same priority to a running or ready task moves it to the end of its priority group, thus causing a manual round-robin.

5.83.3.20 int rt_task_shadow (RT_TASK * task, const char * name, int prio, int mode)

Turn caller into a real-time task.

Extends the calling Linux task with Xenomai capabilities, with access to the full set of Xenomai real-time services. This service is typically used for turning the main() thread of an application process into a Xenomai-enabled task.

If *prio* is non-zero, the new task moves to Xenomai's real-time FIFO scheduling class, aka SCHED_FIFO. If *prio* is zero, the task moves to the regular SCHED_OTHER class.

Running Xenomai tasks with zero priority is useful for running non real-time processes which may invoke blocking real-time services, such as pending on a semaphore, reading from a message queue or a buffer, and so on.

Once shadowed with the Xenomai extension, the calling task returns and resumes execution normally from the call site.

Parameters

<i>task</i>	If non-NULL, the address of a task descriptor which can be later used to identify uniquely the task, upon success of this call. If NULL, no descriptor is returned.
<i>name</i>	An ASCII string standing for the symbolic name of the task. When non-NULL and non-empty, a copy of this string is used for indexing the task into the object registry.
<i>prio</i>	The base priority of the task. This value must be in the [0 .. 99] range, where 0 is the lowest effective priority.
<i>mode</i>	The task shadowing mode. The following flags can be OR'ed into this bitmask:

- T_LOCK causes the current task to lock the scheduler before returning to the caller, preventing all further involuntary task switches on the current CPU. A call to [rt_task_set_mode\(\)](#) from the current task is required to drop this lock.
- When running over the Cobalt core, T_WARNSW causes the SIGXCPU signal to be sent to the current task whenever it switches to the secondary mode. This feature is useful to detect unwanted migrations to the Linux domain. This flag has no effect over the Mercury core.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *prio* is invalid.
- -ENOMEM is returned if the system fails to get memory from the main heap in order to create the task extension.
- -EEXIST is returned if the *name* is conflicting with an already registered task.
- -EBUSY is returned if the caller is already mapped to a Xenomai task context.
- -EPERM is returned if this service was called from an invalid context.

Valid calling context:

- Regular POSIX threads

Core specifics:

When running over the Cobalt core:

- the caller always returns from this service in primary mode.

Note

Tasks can be referred to from multiple processes which all belong to the same Xenomai session.

Examples:

[rtcanrecv.c](#), and [rtcansend.c](#).

References T_LOCK, and T_WARNSW.

5.83.3.21 `int rt_task_sleep (RTIME delay)`

Delay the current real-time task (with relative delay).

This routine is a variant of [rt_task_sleep_until\(\)](#) accepting a relative timeout specification.

Parameters

<i>delay</i>	A relative delay expressed in clock ticks (see note). A zero delay causes this service to return immediately to the caller with a success status.
--------------	---

Returns

See [rt_task_sleep_until\(\)](#).

Note

The *delay* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Examples:

[cross-link.c](#), and [rtcansend.c](#).

5.83.3.22 `int rt_task_sleep_until (RTIME date)`

Delay the current real-time task (with absolute wakeup date).

Delay the execution of the calling task until a given date is reached. The caller is put to sleep, and does not consume any CPU time in such a state.

Parameters

<i>date</i>	An absolute date expressed in clock ticks, specifying a wakeup date (see note). As a special case, <code>TM_INFINITE</code> is an acceptable value that causes the caller to block indefinitely, until <code>rt_task_unblock()</code> is called against it. Otherwise, any wake up date in the past causes the task to return immediately.
-------------	--

Returns

Zero is returned upon success. Otherwise:

- `-EINTR` is returned if `rt_task_unblock()` was called for the current task.
- `-ETIMEDOUT` is returned if *date* has already elapsed.
- `-EPERM` is returned if this service was called from an invalid context.

Valid calling context:

- Xenomai threads

Note

The *date* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.83.3.23 `int rt_task_slice (RT_TASK * task, RTIME quantum)`

Set a task's round-robin quantum.

Set the time credit allotted to a task undergoing the round-robin scheduling. If *quantum* is non-zero, `rt_task_slice()` also refills the current quantum for the target task, otherwise, time-slicing is stopped for that task.

In other words, `rt_task_slice()` should be used to toggle round-robin scheduling for an Alchemy task.

Parameters

<i>task</i>	The descriptor address of the task to update. If <i>task</i> is <code>NULL</code> , the time credit of the current task is changed. <i>task</i> must belong to the current process.
<i>quantum</i>	The round-robin quantum for the task expressed in clock ticks (see note).

Returns

Zero is returned upon success. Otherwise:

- `-EINVAL` is returned if *task* is not a valid task descriptor, or if *prio* is invalid.
- `-EPERM` is returned if *task* is `NULL` and this service was called from an invalid context.

Valid calling context:

- Alchemy tasks if *task* is `NULL`, any otherwise.

Note

The *quantum* value is interpreted as a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

5.83.3.24 `int rt_task_spawn (RT_TASK * task, const char * name, int stksize, int prio, int mode, void(*) (void *arg) entry, void * arg)`

Create and start a real-time task.

This service spawns a task by combining calls to [rt_task_create\(\)](#) and [rt_task_start\(\)](#) for the new task.

Parameters

<i>task</i>	The address of a task descriptor which can be later used to identify uniquely the created object, upon success of this call.
<i>name</i>	An ASCII string standing for the symbolic name of the task. When non-NULL and non-empty, a copy of this string is used for indexing the created task into the object registry.
<i>stksize</i>	The size of the stack (in bytes) for the new task. If zero is passed, a system-dependent default size will be substituted.
<i>prio</i>	The base priority of the new task. This value must be in the [0 .. 99] range, where 0 is the lowest effective priority.
<i>mode</i>	The task creation mode. See rt_task_create() .
<i>entry</i>	The address of the task entry point.
<i>arg</i>	A user-defined opaque argument <i>entry</i> will receive.

Returns

See [rt_task_create\(\)](#).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Core specifics: see [rt_task_create\(\)](#).

References [rt_task_create\(\)](#), and [rt_task_start\(\)](#).

5.83.3.25 `int rt_task_start (RT_TASK * task, void(*) (void *arg) entry, void * arg)`

Start a real-time task.

This call starts execution of a task previously created by [rt_task_create\(\)](#). This service causes the started task to leave the initial dormant state.

Parameters

<i>task</i>	The descriptor address of the task to be started.
<i>entry</i>	The address of the task entry point.
<i>arg</i>	A user-defined opaque argument <i>entry</i> will receive.

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.

Valid calling context: any.

Note

Starting an already started task leads to a nop, returning a success status.

Examples:

[cross-link.c](#).

Referenced by `rt_task_spawn()`.

5.83.3.26 `int rt_task_suspend (RT_TASK * task)`

Suspend a real-time task.

Forcibly suspend the execution of a task. This task will not be eligible for scheduling until it is explicitly resumed by a call to `rt_task_resume()`. In other words, the suspended state caused by a call to `rt_task_suspend()` is cumulative with respect to the delayed and blocked states caused by other services, and is managed separately from them.

A nesting count is maintained so that `rt_task_suspend()` and `rt_task_resume()` must be used in pairs.

Receiving a Linux signal causes the suspended task to resume immediately.

Parameters

<i>task</i>	The descriptor address of the task to suspend. If <i>task</i> is NULL, the current task is suspended.
-------------	---

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is NULL but the caller is not a Xenomai task, or if *task* is non-NULL but not a valid task descriptor.
- -EINTR is returned if a Linux signal has been received by the caller if suspended.
- -EPERM is returned if *task* is NULL and this service was called from an invalid context.

Valid calling context:

- Alchemy tasks if *task* is NULL, any otherwise.

Note

Blocked and suspended task states are cumulative. Therefore, suspending a task currently waiting on a synchronization object (e.g. semaphore, queue) holds its execution until it is resumed, despite the awaited resource may have been acquired, or a timeout has elapsed in the meantime.

5.83.3.27 `int rt_task_unbind (RT_TASK * task)`

Unbind from a task.

Parameters

<i>task</i>	The descriptor address of the task to unbind from.
-------------	--

This routine releases a previous binding to an Alchemy task. After this call has returned, the descriptor is no more valid for referencing this object.

5.83.3.28 `int rt_task_unblock (RT_TASK * task)`

Unblock a real-time task.

Break the task out of any wait it is currently in. This call clears all delay and/or resource wait condition for the target task.

However, `rt_task_unblock()` does not resume a task which has been forcibly suspended by a previous call to `rt_task_suspend()`. If all suspensive conditions are gone, the task becomes eligible anew for scheduling.

Parameters

<i>task</i>	The descriptor address of the task to unblock.
-------------	--

Returns

Zero is returned upon success. Otherwise:

- -EINVAL is returned if *task* is not a valid task descriptor.

Valid calling context: any.

5.83.3.29 `int rt_task_wait_period (unsigned long * overruns_r)`

Wait for the next periodic release point.

Delay the current task until the next periodic release point is reached. The periodic timer should have been previously started for *task* by a call to `rt_task_set_periodic()`.

Parameters

<i>overruns_r</i>	If non-NULL, <i>overruns_r</i> shall be a pointer to a memory location which will be written with the count of pending overruns. This value is written to only when <code>rt_task_wait_period()</code> returns -ETIMEDOUT or success. The memory location remains unmodified otherwise. If NULL, this count will not be returned.
-------------------	---

Returns

Zero is returned upon success. If *overruns_r* is non-NULL, zero is written to the pointed memory location. Otherwise:

- -EWOULDBLOCK is returned if `rt_task_set_periodic()` was not called for the current task.
- -EINTR is returned if `rt_task_unblock()` was called for the waiting task before the next periodic release point was reached. In this case, the overrun counter is also cleared.
- -ETIMEDOUT is returned if a timer overrun occurred, which indicates that a previous release point was missed by the calling task. If *overruns_r* is non-NULL, the count of pending overruns is written to the pointed memory location.
- -EPERM is returned if this service was called from an invalid context.

Valid calling context:

- Alchemy tasks.

Note

If the current release point has already been reached at the time of the call, the current task immediately returns from this service with no delay.

Examples:

[cross-link.c](#).

5.83.3.30 `int rt_task_yield (void)`

Manual round-robin.

Move the current task to the end of its priority group, so that the next equal-priority task in ready state is switched in.

Returns

Zero is returned upon success. Otherwise:

- -EPERM is returned if this service was called from an invalid context.

Valid calling context:

- Xenomai threads.

5.84 Timer management services

Services for reading and spinning on the hardware timer.

Collaboration diagram for Timer management services:



Data Structures

- struct [rt_timer_info](#)
Timer status descriptor.

Typedefs

- typedef struct [rt_timer_info](#) [RT_TIMER_INFO](#)
Timer status descriptor.

Functions

- SRTIME [rt_timer_ns2ticks](#) (SRTIME ns)
Convert nanoseconds to Alchemy clock ticks.
- SRTIME [rt_timer_ticks2ns](#) (SRTIME ticks)
Convert Alchemy clock ticks to nanoseconds.
- RTIME [rt_timer_read](#) (void)
Return the current system time.
- int [rt_timer_inquire](#) ([RT_TIMER_INFO](#) *info)
Inquire about the Xenomai core timer.
- void [rt_timer_spin](#) (RTIME ns)
Busy wait burning CPU cycles.

5.84.1 Detailed Description

Services for reading and spinning on the hardware timer.

5.84.2 Typedef Documentation

5.84.2.1 typedef struct [rt_timer_info](#) [RT_TIMER_INFO](#)

Timer status descriptor.

This structure reports various static and runtime information about the timer, returned by a call to [rt_timer_inquire\(\)](#).

5.84.3 Function Documentation

5.84.3.1 `int rt_timer_inquire (RT_TIMER_INFO * info)`

Inquire about the Xenomai core timer.

Return status information of the Xenomai core timer.

Parameters

<i>info</i>	The address of a structure the status data will be written to.
-------------	--

Returns

This service always returns 0.

Valid calling context:

- Regular POSIX threads
- Xenomai threads

References `rt_timer_info::date`, `rt_timer_info::period`, and `rt_timer_info::tsc`.

5.84.3.2 `SRTIME rt_timer_ns2ticks (SRTIME ns)`

Convert nanoseconds to Alchemy clock ticks.

Convert a count of nanoseconds to Alchemy clock ticks. This routine operates on signed nanosecond values. This is the converse call to [rt_timer_ticks2ns\(\)](#).

Parameters

<i>ns</i>	The count of nanoseconds to convert.
-----------	--------------------------------------

Returns

The corresponding value expressed in clock ticks of the Alchemy clock. The resolution of the Alchemy clock can be set using the `–alchemy-clock-resolution` option when starting the application process (defaults to 1 nanosecond).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Examples:

[cross-link.c](#), and [rtcansend.c](#).

5.84.3.3 `RTIME rt_timer_read (void)`

Return the current system time.

Return the current time maintained by the Xenomai core clock.

Returns

The current time expressed in clock ticks (see note).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Note

The *time* value is a multiple of the Alchemy clock resolution (see `–alchemy-clock-resolution` option, defaults to 1 nanosecond).

Examples:

[cross-link.c](#).

5.84.3.4 void `rt_timer_spin` (RTIME ns)

Busy wait burning CPU cycles.

Enter a busy waiting loop for a count of nanoseconds.

Since this service is always called with interrupts enabled, the caller might be preempted by other real-time activities, therefore the actual delay might be longer than specified.

Parameters

<i>ns</i>	The time to wait expressed in nanoseconds.
-----------	--

Valid calling context:

- Regular POSIX threads
- Xenomai threads

5.84.3.5 SRTIME `rt_timer_ticks2ns` (SRTIME ns)

Convert Alchemy clock ticks to nanoseconds.

Convert a count of Alchemy clock ticks to nanoseconds. This routine operates on signed nanosecond values. This is the converse call to [rt_timer_ns2ticks\(\)](#).

Parameters

<i>ns</i>	The count of nanoseconds to convert.
-----------	--------------------------------------

Returns

The corresponding value expressed in nanoseconds. The resolution of the Alchemy clock can be set using the `–alchemy-clock-resolution` option when starting the application process (defaults to 1 nanosecond).

Valid calling context:

- Regular POSIX threads
- Xenomai threads

Chapter 6

Data Structure Documentation

6.1 a4l_channel Struct Reference

Structure describing some channel's characteristics.

Data Fields

- unsigned long [flags](#)
- unsigned long [nb_bits](#)

6.1.1 Detailed Description

Structure describing some channel's characteristics.

6.1.2 Field Documentation

6.1.2.1 unsigned long a4l_channel::flags

Channel flags to define the reference.

6.1.2.2 unsigned long a4l_channel::nb_bits

Channel resolution.

Referenced by `a4l_get_chan()`.


The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/rtdm/analogy/channel_range.h`

6.2 a4l_channels_desc Struct Reference

Structure describing a channels set.

Collaboration diagram for `a4l_channels_desc`:



Data Fields

- unsigned long `mode`
- unsigned long `length`
- struct `a4l_channel` `chans` []

6.2.1 Detailed Description

Structure describing a channels set.

6.2.2 Field Documentation

6.2.2.1 struct **`a4l_channel`** `a4l_channels_desc::chans` []

Channels tab

Referenced by `a4l_get_chan()`.

6.2.2.2 unsigned long `a4l_channels_desc::length`

Channels count

6.2.2.3 unsigned long `a4l_channels_desc::mode`

Declaration mode (global or per channel)

Referenced by `a4l_get_chan()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/rtdm/analogy/channel_range.h`

6.3 `a4l_cmd_desc` Struct Reference

Structure describing the asynchronous instruction.

Data Fields

- unsigned char [idx_subd](#)
Subdevice to which the command will be applied.
- unsigned long [flags](#)
Command flags.
- unsigned int [start_src](#)
Start trigger type.
- unsigned int [start_arg](#)
Start trigger argument.
- unsigned int [scan_begin_src](#)
Scan begin trigger type.
- unsigned int [scan_begin_arg](#)
Scan begin trigger argument.
- unsigned int [convert_src](#)
Convert trigger type.
- unsigned int [convert_arg](#)
Convert trigger argument.
- unsigned int [scan_end_src](#)
Scan end trigger type.
- unsigned int [scan_end_arg](#)
Scan end trigger argument.
- unsigned int [stop_src](#)
Stop trigger type.
- unsigned int [stop_arg](#)
Stop trigger argument.
- unsigned char [nb_chan](#)
Count of channels related with the command.
- unsigned int * [chan_descs](#)
Tab containing channels descriptors.
- unsigned int [data_len](#)
Driver specific buffer size.
- `sample_t` * [data](#)
Driver specific buffer pointer.

6.3.1 Detailed Description

Structure describing the asynchronous instruction.

See Also

[a4l_snd_command\(\)](#)

6.3.2 Field Documentation

6.3.2.1 unsigned char a4l_cmd_desc::idx_subd

Subdevice to which the command will be applied.

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/analog.h`

6.4 a4l_descriptor Struct Reference

Structure containing device-information useful to users.

Data Fields

- char `board_name` [A4L_NAMELEN]
Board name.
- int `nb_subd`
Subdevices count.
- int `idx_read_subd`
Input subdevice index.
- int `idx_write_subd`
Output subdevice index.
- int `fd`
File descriptor.
- unsigned int `magic`
Opaque field.
- int `sbsize`
Data buffer size.
- void * `sbddata`
Data buffer pointer.

6.4.1 Detailed Description

Structure containing device-information useful to users.

See Also

`a4l_get_desc()`

6.4.2 Field Documentation

6.4.2.1 char a4l_descriptor::board_name[A4L_NAMELEN]

Board name.

6.4.2.2 int a4l_descriptor::fd

File descriptor.

Referenced by `a4l_async_read()`, `a4l_async_write()`, `a4l_close()`, `a4l_fill_desc()`, `a4l_get_bufsize()`, `a4l_mark_bufrw()`, `a4l_mmap()`, `a4l_open()`, `a4l_poll()`, `a4l_set_bufsize()`, `a4l_snd_cancel()`, `a4l_snd_command()`, `a4l_snd_insn()`, and `a4l_snd_insnlist()`.

6.4.2.3 int a4l_descriptor::idx_read_subd

Input subdevice index.

Referenced by `a4l_async_read()`.

6.4.2.4 int a4l_descriptor::idx_write_subd

Output subdevice index.

Referenced by a4l_async_write().

6.4.2.5 unsigned int a4l_descriptor::magic

Opaque field.

Referenced by a4l_fill_desc(), a4l_find_range(), a4l_get_chinfo(), a4l_get_rnginfo(), a4l_get_subdinfo(), and a4l_sys_desc().

6.4.2.6 int a4l_descriptor::nb_subd

Subdevices count.

Referenced by a4l_get_chinfo(), a4l_get_rnginfo(), and a4l_get_subdinfo().

6.4.2.7 void* a4l_descriptor::sbddata

Data buffer pointer.

Referenced by a4l_get_chinfo(), a4l_get_rnginfo(), a4l_get_subdinfo(), and a4l_sys_desc().

6.4.2.8 int a4l_descriptor::sbsize

Data buffer size.

Referenced by a4l_sys_desc().

The documentation for this struct was generated from the following file:

- include/rtdm/[analogy.h](#)

6.5 a4l_driver Struct Reference

Structure containing driver declaration data.

Data Fields

- struct list_head [list](#)
List stuff.
- struct module * [owner](#)
Pointer to module containing the code.
- unsigned int [flags](#)
Type / status driver's flags.
- char * [board_name](#)
Board name.
- int [privdata_size](#)
Size of the driver's private data.
- int(* [attach](#))(struct a4l_device *, struct a4l_link_desc *)

Attach procedure.

- `int(* detach)(struct a4l_device *)`

Detach procedure.

6.5.1 Detailed Description

Structure containing driver declaration data.

See Also

[rt_task_inquire\(\)](#)

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/rtdm/analogy/driver.h`

6.6 a4l_instruction Struct Reference

Structure describing the synchronous instruction.

Data Fields

- unsigned int [type](#)
Instruction type.
- unsigned int [idx_subd](#)
Subdevice to which the instruction will be applied.
- unsigned int [chan_desc](#)
Channel descriptor.
- unsigned int [data_size](#)
Size of the instruction data.
- void * [data](#)
Instruction data.

6.6.1 Detailed Description

Structure describing the synchronous instruction.

See Also

[a4l_snd_insn\(\)](#)

6.6.2 Field Documentation

6.6.2.1 unsigned int a4l_instruction::idx_subd

Subdevice to which the instruction will be applied.


The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/analogy.h`

6.7 a4l_instruction_list Struct Reference

Structure describing the list of synchronous instructions.

Collaboration diagram for a4l_instruction_list:



Data Fields

- unsigned int [count](#)
Instructions count.
- [a4l_insn_t](#) * [insns](#)
Tab containing the instructions pointers.

6.7.1 Detailed Description

Structure describing the list of synchronous instructions.

See Also

[a4l_snd_insnlist\(\)](#)

The documentation for this struct was generated from the following file:

- include/rtdm/uapi/[analogy.h](#)

6.8 a4l_range Struct Reference

Structure describing a (unique) range.

Data Fields

- long [min](#)
- long [max](#)
- unsigned long [flags](#)

6.8.1 Detailed Description

Structure describing a (unique) range.

6.8.2 Field Documentation

6.8.2.1 unsigned long a4l_range::flags

Range flags (unit, etc.)

6.8.2.2 long a4l_range::max

Maximal value

6.8.2.3 long a4l_range::min

Minimal value


The documentation for this struct was generated from the following file:

- [include/cobalt/kernel/rtdm/analogy/channel_range.h](#)

6.9 a4l_subdevice Struct Reference

Structure describing the subdevice.

Collaboration diagram for a4l_subdevice:



Data Fields

- struct list_head [list](#)

- *List stuff.*
- struct a4l_device * [dev](#)
 - *Containing device.*
- unsigned int [idx](#)
 - *Subdevice index.*
- struct a4l_buffer * [buf](#)
 - *Linked buffer.*
- unsigned long [status](#)
 - *Subdevice's status.*
- unsigned long [flags](#)
 - *Type flags.*
- struct [a4l_channels_desc](#) * [chan_desc](#)
 - *Tab of channels descriptors pointers.*
- struct a4l_rngdesc * [rng_desc](#)
 - *Tab of ranges descriptors pointers.*
- struct [a4l_cmd_desc](#) * [cmd_mask](#)
 - *Command capabilities mask.*
- int(* [insn_read](#))(struct [a4l_subdevice](#) *, struct a4l_kernel_instruction *)
 - *Callback for the instruction "read".*
- int(* [insn_write](#))(struct [a4l_subdevice](#) *, struct a4l_kernel_instruction *)
 - *Callback for the instruction "write".*
- int(* [insn_bits](#))(struct [a4l_subdevice](#) *, struct a4l_kernel_instruction *)
 - *Callback for the instruction "bits".*
- int(* [insn_config](#))(struct [a4l_subdevice](#) *, struct a4l_kernel_instruction *)
 - *Callback for the configuration instruction.*
- int(* [do_cmd](#))(struct [a4l_subdevice](#) *, struct [a4l_cmd_desc](#) *)
 - *Callback for command handling.*
- int(* [do_cmdtest](#))(struct [a4l_subdevice](#) *, struct [a4l_cmd_desc](#) *)
 - *Callback for command checking.*
- void(* [cancel](#))(struct [a4l_subdevice](#) *)
 - *Callback for asynchronous transfer cancellation.*
- void(* [munge](#))(struct [a4l_subdevice](#) *, void *, unsigned long)
 - *Callback for munge operation.*
- int(* [trigger](#))(struct [a4l_subdevice](#) *, lsampl_t)
 - *Callback for trigger operation.*
- char [priv](#) [0]
 - *Private data.*

6.9.1 Detailed Description

Structure describing the subdevice.

See Also

[a4l_add_subd\(\)](#)

The documentation for this struct was generated from the following file:

- [include/cobalt/kernel/rtdm/analogy/subdevice.h](#)

6.10 atomic_long_t Struct Reference

Copyright © 2011 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org.

6.10.1 Detailed Description

Copyright © 2011 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org.

Copyright © 2013 Philippe Gerum rpm@xenomai.org.

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.


The documentation for this struct was generated from the following file:

- `include/nocore/atomic.h`

6.11 can_bittime Struct Reference

Custom CAN bit-time definition.

Collaboration diagram for `can_bittime`:



Data Fields

- `can_bittime_type_t` type
Type of bit-time definition.
- struct `can_bittime_std` `std`
Standard bit-time.
- struct `can_bittime_btr` `btr`
Hardware-specific BTR bit-time.

6.11.1 Detailed Description

Custom CAN bit-time definition.

Examples:

[rtcanconfig.c](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/can.h](#)

6.12 can_bittime_btr Struct Reference

Hardware-specific BTR bit-times.

Data Fields

- [uint8_t btr0](#)
Bus timing register 0.
- [uint8_t btr1](#)
Bus timing register 1.

6.12.1 Detailed Description

Hardware-specific BTR bit-times.

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/can.h](#)

6.13 can_bittime_std Struct Reference

Standard bit-time parameters according to Bosch.

Data Fields

- [uint32_t brp](#)
Baud rate prescaler.
- [uint8_t prop_seg](#)
from 1 to 8
- [uint8_t phase_seg1](#)
from 1 to 8
- [uint8_t phase_seg2](#)
from 1 to 8
- [uint8_t sjw:7](#)
from 1 to 4
- [uint8_t sam:1](#)
1 - enable triple sampling

6.13.1 Detailed Description

Standard bit-time parameters according to Bosch.

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/can.h`

6.14 `can_filter` Struct Reference

Filter for reception of CAN messages.

Data Fields

- `uint32_t can_id`
CAN ID which must match with incoming IDs after passing the mask.
- `uint32_t can_mask`
Mask which is applied to incoming IDs.

6.14.1 Detailed Description

Filter for reception of CAN messages.

This filter works as follows: A received CAN ID is AND'ed bitwise with `can_mask` and then compared to `can_id`. This also includes the `CAN_EFF_FLAG` and `CAN_RTR_FLAG` of `CAN_XXX_FLAG`. If this comparison is true, the message will be received by the socket. The logic can be inverted with the `can_id` flag `CAN_INV_FILTER`:

```
if (can_id & CAN_INV_FILTER) {
    if ((received_can_id & can_mask) != (can_id & ~CAN_INV_FILTER))
        accept-message;
} else {
    if ((received_can_id & can_mask) == can_id)
        accept-message;
}
```

Multiple filters can be arranged in a filter list and set with `Sockopts`. If one of these filters matches a CAN ID upon reception of a CAN frame, this frame is accepted.

Examples:

`can-rtt.c`, and `rtcanrecv.c`.

6.14.2 Field Documentation

6.14.2.1 `uint32_t can_filter::can_id`

CAN ID which must match with incoming IDs after passing the mask.

The filter logic can be inverted with the flag `CAN_INV_FILTER`.

Examples:

`rtcanrecv.c`.

6.14.2.2 `uint32_t can_filter::can_mask`

Mask which is applied to incoming IDs.

See [CAN ID masks](#) if exactly one CAN ID should come through.

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/can.h`

6.15 can_frame Struct Reference

Raw CAN frame.

Public Member Functions

- `uint8_t data[8] __attribute__((aligned(8)))`
Payload data bytes.

Data Fields

- `can_id_t can_id`
CAN ID of the frame.
- `uint8_t can_dlc`
Size of the payload in bytes.

6.15.1 Detailed Description

Raw CAN frame.

Central structure for receiving and sending CAN frames.

Examples:

[can-rtt.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

6.15.2 Field Documentation

6.15.2.1 `can_id_t can_frame::can_id`

CAN ID of the frame.

See [CAN ID flags](#) for special bits.

Examples:

[can-rtt.c](#).

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/can.h`

6.16 RT_ALARM_INFO Struct Reference

Alarm status descriptor.

Data Fields

- unsigned long [expiries](#)
Number of past expiries.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of alarm object.
- int [active](#)
Active flag.

6.16.1 Detailed Description

Alarm status descriptor.

This structure reports various static and runtime information about a real-time alarm, returned by a call to [rt_alarm_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/alarm.h

6.17 RT_BUFFER_INFO Struct Reference

Buffer status descriptor.

Data Fields

- int [iwaiters](#)
Number of tasks waiting on the read side of the buffer for input data.
- int [owaiters](#)
Number of tasks waiting on the write side of the buffer for sending out data.
- size_t [totalmem](#)
Overall size of buffer (in bytes).
- size_t [availmem](#)
Amount of memory currently available for holding more data.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of the buffer.

6.17.1 Detailed Description

Buffer status descriptor.

This structure reports various static and runtime information about a real-time buffer, returned by a call to [rt_buffer_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/buffer.h

6.18 RT_COND_INFO Struct Reference

Condition variable status descriptor.

Data Fields

- char [name](#) [XNOBJECT_NAME_LEN]
Name of condition variable.

6.18.1 Detailed Description

Condition variable status descriptor.

This structure reports various static and runtime information about a condition variable, returned by a call to [rt_cond_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/cond.h

6.19 RT_EVENT_INFO Struct Reference

Event status descriptor.

Data Fields

- unsigned long [value](#)
Current value of the event flag group.
- int [nwaiters](#)
Number of tasks currently waiting for events.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of event flag group.

6.19.1 Detailed Description

Event status descriptor.

This structure reports various static and runtime information about an event flag group, returned by a call to [rt_event_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/event.h

6.20 RT_HEAP_INFO Struct Reference

Heap status descriptor.

Data Fields

- int [nwaiters](#)
Number of tasks waiting for available memory in [rt_heap_alloc\(\)](#).
- int [mode](#)
Creation mode flags as given to [rt_heap_create\(\)](#).
- size_t [heapsize](#)
Overall size of heap (in bytes).
- size_t [usablemem](#)
Maximum amount of memory available from the heap.
- size_t [usedmem](#)
Amount of heap memory currently consumed.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of heap.

6.20.1 Detailed Description

Heap status descriptor.

This structure reports various static and runtime information about a real-time heap, returned by a call to [rt_heap_inquire\(\)](#).

6.20.2 Field Documentation

6.20.2.1 size_t RT_HEAP_INFO::usablemem

Maximum amount of memory available from the heap.

This value accounts for the overhead of internal data structures required to maintain the heap.

Referenced by [rt_heap_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- [include/alchemy/heap.h](#)

6.21 RT_MUTEX_INFO Struct Reference

Mutex status descriptor.

Data Fields

- RT_TASK [owner](#)
Current mutex owner, or null if unlocked.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of mutex.

6.21.1 Detailed Description

Mutex status descriptor.

This structure reports various static and runtime information about a mutex, returned by a call to [rt_mutex_inquire\(\)](#).

6.21.2 Field Documentation

6.21.2.1 RT_TASK RT_MUTEX_INFO::owner

Current mutex owner, or null if unlocked.

This information is in essence transient, and may not be valid anymore once used by the caller.

Referenced by `rt_mutex_inquire()`.

The documentation for this struct was generated from the following file:

- `include/alchemy/mutex.h`

6.22 RT_QUEUE_INFO Struct Reference

Queue status descriptor.

Data Fields

- `int nwaiters`
Number of tasks currently waiting on the queue for messages.
- `int nmessages`
Number of messages pending in queue.
- `int mode`
Queue mode bits, as given to `rt_queue_create()`.
- `size_t qlimit`
Maximum number of messages in queue, zero if unlimited.
- `size_t poolsize`
Size of memory pool for holding message buffers (in bytes).
- `size_t usedmem`
Amount of memory consumed from the buffer pool.
- `char name [XNOBJECT_NAME_LEN]`
Name of message queue.

6.22.1 Detailed Description

Queue status descriptor.

This structure reports various static and runtime information about a real-time queue, returned by a call to `rt_queue_inquire()`.

The documentation for this struct was generated from the following file:

- `include/alchemy/queue.h`

6.23 RT_SEM_INFO Struct Reference

Semaphore status descriptor.

Data Fields

- unsigned long [count](#)
Current semaphore value.
- int [nwaiters](#)
Number of tasks waiting on the semaphore.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of semaphore.

6.23.1 Detailed Description

Semaphore status descriptor.

This structure reports various static and runtime information about a semaphore, returned by a call to [rt_sem_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/sem.h

6.24 RT_TASK_INFO Struct Reference

Task status descriptor.

Data Fields

- int [prio](#)
Task priority.
- struct threadobj_stat [stat](#)
Task status.
- char [name](#) [XNOBJECT_NAME_LEN]
Name of task.
- pid_t [pid](#)
Host pid.

6.24.1 Detailed Description

Task status descriptor.

This structure reports various static and runtime information about a real-time task, returned by a call to [rt_task_inquire\(\)](#).

The documentation for this struct was generated from the following file:

- include/alchemy/task.h

6.25 rt_timer_info Struct Reference

Timer status descriptor.

Data Fields

- RTIME [period](#)

Clock resolution in nanoseconds.

- RTIME [tsc](#)

Current time stamp counter value.

- RTIME [date](#)

Current monotonic date, based on the time stamp counter value.

6.25.1 Detailed Description

Timer status descriptor.

This structure reports various static and runtime information about the timer, returned by a call to [rt_timer_inquire\(\)](#).

6.25.2 Field Documentation

6.25.2.1 RTIME rt_timer_info::date

Current monotonic date, based on the time stamp counter value.

The date is expressed in clock ticks, therefore depends on the Alchemy clock resolution applicable to the current process.

Referenced by [rt_timer_inquire\(\)](#).

6.25.2.2 RTIME rt_timer_info::tsc

Current time stamp counter value.

The source of this information is hardware-dependent, and does not depend on the per-process clock settings. Consecutive readings from a single CPU are guaranteed to be monotonically incrementing, however readings may not be synchronized on multi-core hardware if the time stamp counter is local to each CPU. Therefore, whether consecutive readings from different CPUs are consistent and monotonically incrementing depends on the underlying TSC source.

Referenced by [rt_timer_inquire\(\)](#).


The documentation for this struct was generated from the following file:

- `include/alchemy/timer.h`

6.26 rtdm_dev_context Struct Reference

Device context.

Collaboration diagram for `rtm_dev_context`:



Data Fields

- struct [rtm_device](#) * [device](#)
Set of active device operation handlers.
- struct `rtm_devctx_reserved` [reserved](#)
Data stored by RTDM inside a device context (internal use only)
- char [dev_private](#) [0]
Begin of driver defined context data structure.

6.26.1 Detailed Description

Device context.

A device context structure is associated with every open device instance. RTDM takes care of its creation and destruction and passes it to the operation handlers when being invoked.

Drivers can attach arbitrary data immediately after the official structure. The size of this data is provided via [rtm_device.context_size](#) during device registration.

6.26.2 Field Documentation

6.26.2.1 struct **rtm_device*** `rtm_dev_context::device`

Set of active device operation handlers.

Reference to owning device

Referenced by `rtm_fd_device()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/rtm/driver.h`

6.27 rtm_device Struct Reference

RTDM device.

Data Fields

- struct rtdm_dev_reserved [reserved](#)
Data stored by RTDM inside a registered device (internal use only)
- int [struct_version](#)
Revision number of this structure, see [Driver Versioning](#) defines.
- int [device_flags](#)
Device flags, see [Device Flags](#) for details.
- size_t [context_size](#)
Size of driver defined appendix to struct [rtdm_dev_context](#).
- char [device_name](#) [RTDM_MAX_DEVNAME_LEN+1]
Named device identification (orthogonal to Linux device name space)
- int [protocol_family](#)
Protocol device identification: protocol family (PF_XXX)
- int [socket_type](#)
Protocol device identification: socket type (SOCK_XXX)
- [rtdm_open_handler_t](#) [open](#)
Named device instance creation for real-time contexts.
- [rtdm_socket_handler_t](#) [socket](#)
Protocol socket creation for real-time contexts.
- struct rtdm_fd_ops [ops](#)
Default operations on newly opened device instance.
- int [device_class](#)
Device class ID, see [RTDM_CLASS_XXX](#).
- int [device_sub_class](#)
Device sub-class, see [RTDM_SUBCLASS_XXX](#) definition in the [Device Profiles](#).
- int [profile_version](#)
Supported device profile version.
- const char * [driver_name](#)
Informational driver name (reported via /proc)
- int [driver_version](#)
Driver version, see [Driver Versioning](#) defines.
- const char * [peripheral_name](#)
Informational peripheral name the device is attached to (reported via /proc)
- const char * [provider_name](#)
Informational driver provider name (reported via /proc)
- const char * [proc_name](#)
Name of /proc entry for the device, must not be NULL.
- int [device_id](#)
Driver definable device ID.
- void * [device_data](#)
Driver definable device data.

6.27.1 Detailed Description

RTDM device.

This structure specifies a RTDM device. As some fields, especially the reserved area, will be modified by RTDM during runtime, the structure must not reside in write-protected memory.

6.27.2 Field Documentation

6.27.2.1 `rtdm_open_handler_t` `rtdm_device::open`

Named device instance creation for real-time contexts.

Referenced by `rtdm_dev_register()`.

6.27.2.2 `rtdm_socket_handler_t` `rtdm_device::socket`

Protocol socket creation for real-time contexts.

Referenced by `rtdm_dev_register()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/rtdm/driver.h`

6.28 `rtdm_device_info` Struct Reference

Device information.

Data Fields

- int `device_flags`
Device flags, see [Device Flags](#) for details.
- int `device_class`
Device class ID, see [RTDM_CLASS_xxx](#).
- int `device_sub_class`
Device sub-class, either [RTDM_SUBCLASS_GENERIC](#) or a [RTDM_SUBCLASS_xxx](#) definition of the related [Device Profile](#).
- int `profile_version`
Supported device profile version.

6.28.1 Detailed Description

Device information.

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/rtdm.h`

6.29 `rtipc_port_label` Struct Reference

Port label information structure.

Data Fields

- char `label` [`XNOBJECT_NAME_LEN`]
Port label string, null-terminated.

6.29.1 Detailed Description

Port label information structure.

Examples:

[bufp-label.c](#), [iddp-label.c](#), and [xddp-label.c](#).

6.29.2 Field Documentation

6.29.2.1 char rtipc_port_label::label[XNOBJECT_NAME_LEN]

Port label string, null-terminated.

Referenced by `rt_pipe_create()`.

The documentation for this struct was generated from the following file:

- `include/rtdm/uapi/ipc.h`

6.30 rtser_config Struct Reference

Serial device configuration.

Data Fields

- int [config_mask](#)
mask specifying valid fields, see [RTSER_SET_xxx](#)
- int [baud_rate](#)
baud rate, default [RTSER_DEF_BAUD](#)
- int [parity](#)
number of parity bits, see [RTSER_xxx_PARITY](#)
- int [data_bits](#)
number of data bits, see [RTSER_xxx_BITS](#)
- int [stop_bits](#)
number of stop bits, see [RTSER_xxx_STOPB](#)
- int [handshake](#)
handshake mechanisms, see [RTSER_xxx_HAND](#)
- int [fifo_depth](#)
reception FIFO interrupt threshold, see [RTSER_FIFO_xxx](#)
- [nanosecs_rel_t rx_timeout](#)
reception timeout, see [RTSER_TIMEOUT_xxx](#) for special values
- [nanosecs_rel_t tx_timeout](#)
transmission timeout, see [RTSER_TIMEOUT_xxx](#) for special values
- [nanosecs_rel_t event_timeout](#)
event timeout, see [RTSER_TIMEOUT_xxx](#) for special values
- int [timestamp_history](#)
enable timestamp history, see [RTSER_xxx_TIMESTAMP_HISTORY](#)
- int [event_mask](#)
event mask to be used with [RTSER_RTIOC_WAIT_EVENT](#), see [RTSER_EVENT_xxx](#)
- int [rs485](#)
enable RS485 mode, see [RTSER_RS485_xxx](#)

6.30.1 Detailed Description

Serial device configuration.

Examples:

[cross-link.c](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/serial.h](#)

6.31 rtser_event Struct Reference

Additional information about serial device events.

Data Fields

- int [events](#)
signalled events, see [RTSER_EVENT_XXX](#)
- int [rx_pending](#)
number of pending input characters
- [nanosecs_abs_t last_timestamp](#)
last interrupt timestamp
- [nanosecs_abs_t rxpend_timestamp](#)
reception timestamp of oldest character in input queue

6.31.1 Detailed Description

Additional information about serial device events.

Examples:

[cross-link.c](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/serial.h](#)

6.32 rtser_status Struct Reference

Serial device status.

Data Fields

- int [line_status](#)
line status register, see [RTSER_LSR_XXX](#)
- int [modem_status](#)
modem status register, see [RTSER_MSR_XXX](#)

6.32.1 Detailed Description

Serial device status.

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/serial.h](#)

6.33 sockaddr_can Struct Reference

Socket address structure for the CAN address family.

Data Fields

- [sa_family_t can_family](#)
CAN address family, must be AF_CAN.
- [int can_ifindex](#)
Interface index of CAN controller.

6.33.1 Detailed Description

Socket address structure for the CAN address family.

Examples:

[can-rtt.c](#), [rtcanrecv.c](#), and [rtcansend.c](#).

6.33.2 Field Documentation

6.33.2.1 int sockaddr_can::can_ifindex

Interface index of CAN controller.

See [SIOCGIFINDEX](#).

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/can.h](#)

6.34 sockaddr_ipc Struct Reference

Socket address structure for the RTIPC address family.

Data Fields

- [sa_family_t sipc_family](#)
RTIPC address family, must be AF_RTIPC.
- [rtipc_port_t sipc_port](#)
Port number.

6.34.1 Detailed Description

Socket address structure for the RTIPC address family.

Examples:

[bufp-label.c](#), [bufp-readwrite.c](#), [iddp-label.c](#), [iddp-sendrecv.c](#), [xddp-echo.c](#), [xddp-label.c](#), and [xddp-stream.c](#).

6.34.2 Field Documentation

6.34.2.1 **rtipc_port_t** sockaddr_ipc::sipc_port

Port number.

Referenced by `rt_pipe_create()`.

The documentation for this struct was generated from the following file:

- [include/rtdm/uapi/ipc.h](#)

6.35 xnsched Struct Reference

Scheduling information structure.

Data Fields

- unsigned long [status](#)
- unsigned long [lflags](#)
- struct xntthread * [curr](#)
- struct xnsched_rt [rt](#)
- volatile unsigned [inesting](#)
- struct xntimer [htimer](#)
- struct xntimer [rrbtimer](#)

6.35.1 Detailed Description

Scheduling information structure.

6.35.2 Field Documentation

6.35.2.1 struct xntthread* xnsched::curr

Context of built-in real-time class.

Referenced by `xntthread_set_slice()`, and `xntthread_suspend()`.

6.35.2.2 struct xntimer xnsched::htimer

Round-robin timer.

Referenced by `program_htick_shot()`, `switch_htick_mode()`, and `xnclock_tick()`.

6.35.2.3 volatile unsigned xnsched::inesting

Host timer.

6.35.2.4 unsigned long xnsched::lflags

Current thread.

Referenced by `xnclock_tick()`, `xnsched_run()`, and `xnthread_suspend()`.

6.35.2.5 struct xntimer xnsched::rrbtimer

Root thread control block.

Referenced by `xnthread_set_slice()`.

6.35.2.6 struct xnsched_rt xnsched::rt

Interrupt nesting level.

6.35.2.7 unsigned long xnsched::status

< Scheduler specific status bitmask. Scheduler specific local flags bitmask.

Referenced by `xnclock_tick()`, and `xnsched_run()`.

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/sched.h`

6.36 xnvfile_lock_ops Struct Reference

Vfile locking operations.

Data Fields

- `int(* get)(struct xnvfile *vfile)`
- `void(* put)(struct xnvfile *vfile)`

6.36.1 Detailed Description

Vfile locking operations.

This structure describes the operations to be provided for implementing locking support on vfiles. They apply to both snapshot-driven and regular vfiles.

6.36.2 Field Documentation

6.36.2.1 `int(* xnvfile_lock_ops::get)(struct xnvfile *vfile)`

This handler should grab the desired lock.

Parameters

<i>vfile</i>	A pointer to the virtual file which needs locking.
--------------	--

Returns

zero should be returned if the call succeeds. Otherwise, a negative error code can be returned; upon error, the current vfile operation is aborted, and the user-space caller is passed back the error value.

6.36.2.2 void(* xnvfile_lock_ops::put)(struct xnvfile *vfile)

This handler should release the lock previously grabbed by the [get\(\) handler](#).

Parameters

<i>vfile</i>	A pointer to the virtual file which currently holds the lock to release.
--------------	--

The documentation for this struct was generated from the following file:

- include/cobalt/kernel/[vfile.h](#)

6.37 xnvfile_regular_iterator Struct Reference

Regular vfile iterator.

Data Fields

- loff_t [pos](#)
Current record position while iterating.
- struct seq_file * [seq](#)
Backlink to the host sequential file supporting the vfile.
- struct xnvfile_regular * [vfile](#)
Backlink to the vfile being read.
- char [private](#) [0]
Start of private area.

6.37.1 Detailed Description

Regular vfile iterator.

This structure defines an iterator over a regular vfile.

6.37.2 Field Documentation

6.37.2.1 loff_t xnvfile_regular_iterator::pos

Current record position while iterating.

6.37.2.2 `char xnvfile_regular_iterator::private[0]`

Start of private area.

Use `xnvfile_iterator_priv()` to address it.

6.37.2.3 `struct seq_file* xnvfile_regular_iterator::seq`

Backlink to the host sequential file supporting the vfile.

6.37.2.4 `struct xnvfile_regular* xnvfile_regular_iterator::vfile`

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/vfile.h`

6.38 xnvfile_regular_ops Struct Reference

Regular vfile operation descriptor.

Data Fields

- `int(* rewind)(struct xnvfile_regular_iterator *it)`
- `void *(* begin)(struct xnvfile_regular_iterator *it)`
- `void *(* next)(struct xnvfile_regular_iterator *it)`
- `void(* end)(struct xnvfile_regular_iterator *it)`
- `int(* show)(struct xnvfile_regular_iterator *it, void *data)`
- `ssize_t(* store)(struct xnvfile_input *input)`

6.38.1 Detailed Description

Regular vfile operation descriptor.

This structure describes the operations available with a regular vfile. It defines handlers for sending back formatted kernel data upon a user-space read request, and for obtaining user data upon a user-space write request.

6.38.2 Field Documentation

6.38.2.1 `void *(* xnvfile_regular_ops::begin)(struct xnvfile_regular_iterator *it)`

This handler should prepare for iterating over the records upon a read request, starting from the specified position.

Parameters

<i>it</i>	A pointer to the current vfile iterator. On entry, <code>it->pos</code> is set to the (0-based) position of the first record to output. This handler may be called multiple times with different position requests.
-----------	--

Returns

A pointer to the first record to format and output, to be passed to the [show\(\) handler](#) as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- `VFILE_SEQ_START`, a special value indicating that the [show\(\) handler](#) should receive a NULL data pointer first, in order to output a header.
- `ERR_PTR(errno)`, where `errno` is a negative error code; upon error, the current operation will be aborted immediately.

Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the [show\(\) handler\(\)](#) will be called only once for a read operation, with a NULL *data* parameter. This particular setting is convenient for simple regular vfiles having a single, fixed record to output.

6.38.2.2 `void(* xnvfile_regular_ops::end)(struct xnvfile_regular_iterator *it)`

This handler is called after all records have been output.

Parameters

<i>it</i>	A pointer to the current vfile iterator.
-----------	--

Note

This handler is optional and the pointer may be NULL.

6.38.2.3 `void*(xnvfile_regular_ops::next)(struct xnvfile_regular_iterator *it)`

This handler should return the address of the next record to format and output by the [show\(\) handler](#)".

Parameters

<i>it</i>	A pointer to the current vfile iterator. On entry, <code>it->pos</code> is set to the (0-based) position of the next record to output.
-----------	---

Returns

A pointer to the next record to format and output, to be passed to the [show\(\) handler](#) as its *data* parameter, if the call succeeds. Otherwise:

- NULL in case no record is available, in which case the read operation will terminate immediately with no output.
- `ERR_PTR(errno)`, where `errno` is a negative error code; upon error, the current operation will be aborted immediately.

Note

This handler is optional; if none is given in the operation descriptor (i.e. NULL value), the read operation will stop after the first invocation of the [show\(\) handler](#).

6.38.2.4 `int(* xnvfile_regular_ops::rewind)(struct xnvfile_regular_iterator *it)`

This handler is called only once, when the virtual file is opened, before the [begin\(\) handler](#) is invoked.

Parameters

<i>it</i>	A pointer to the vfile iterator which will be used to read the file contents.
-----------	---

Returns

Zero should be returned upon success. Otherwise, a negative error code aborts the operation, and is passed back to the reader.

Note

This handler is optional. It should not be used to allocate resources but rather to perform consistency checks, since no closure call is issued in case the open sequence eventually fails.

6.38.2.5 `int(* xnvfile_regular_ops::show)(struct xnvfile_regular_iterator *it, void *data)`

This handler should format and output a record.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

Parameters

<i>it</i>	A pointer to the current vfile iterator.
<i>data</i>	A pointer to the record to format then output. The first call to the handler may receive a NULL <i>data</i> pointer, depending on the presence and/or return of a handler ; the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.

Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

6.38.2.6 `ssize_t(* xnvfile_regular_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating some kernel setting, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the /proc entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile_get_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile_get_integer\(\)](#) will attempt to return a long integer from the input data.

Parameters

<i>input</i>	A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.
--------------	--

Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

Note

This handler is optional, and may be omitted for read-only vfiles.

The documentation for this struct was generated from the following file:

- [include/cobalt/kernel/vfile.h](#)

6.39 xnvfile_rev_tag Struct Reference

Snapshot revision tag.

Data Fields

- [int rev](#)
Current revision number.

6.39.1 Detailed Description

Snapshot revision tag.

This structure defines a revision tag to be used with [snapshot-driven vfiles](#).

6.39.2 Field Documentation

6.39.2.1 `int xnvfile_rev_tag::rev`

Current revision number.


The documentation for this struct was generated from the following file:

- [include/cobalt/kernel/vfile.h](#)

6.40 xnvfile_snapshot Struct Reference

Snapshot vfile descriptor.

Collaboration diagram for xnvfile_snapshot:



6.40.1 Detailed Description

Snapshot vfile descriptor.

This structure describes a snapshot-driven vfile. Reading from such a vfile involves a preliminary data collection phase under lock protection, and a subsequent formatting and output phase of the collected data records. Locking is done in a way that does not increase worst-case latency, regardless of the number of records to be collected for output.


The documentation for this struct was generated from the following file:

- `include/cobalt/kernel/vfile.h`

6.41 xnvfile_snapshot_iterator Struct Reference

Snapshot-driven vfile iterator.

Collaboration diagram for `xnvfile_snapshot_iterator`:



Data Fields

- `int` `nrdata`
Number of collected records.
- `caddr_t` `databuf`
Address of record buffer.
- `struct seq_file *` `seq`
Backlink to the host sequential file supporting the vfile.
- `struct xnvfile_snapshot *` `vfile`
Backlink to the vfile being read.
- `void(* endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`
Buffer release handler.
- `char` `private` `[0]`
Start of private area.

6.41.1 Detailed Description

Snapshot-driven vfile iterator.

This structure defines an iterator over a snapshot-driven vfile.

6.41.2 Field Documentation

6.41.2.1 `caddr_t xnvfile_snapshot_iterator::databuf`

Address of record buffer.

6.41.2.2 `void(* xnvfile_snapshot_iterator::endfn)(struct xnvfile_snapshot_iterator *it, void *buf)`

Buffer release handler.

6.41.2.3 int xnvfile_snapshot_iterator::nrdata

Number of collected records.

6.41.2.4 char xnvfile_snapshot_iterator::private[0]

Start of private area.

Use xnvfile_iterator_priv() to address it.

6.41.2.5 struct seq_file* xnvfile_snapshot_iterator::seq

Backlink to the host sequential file supporting the vfile.

6.41.2.6 struct **xnvfile_snapshot*** xnvfile_snapshot_iterator::vfile

Backlink to the vfile being read.

The documentation for this struct was generated from the following file:

- include/cobalt/kernel/vfile.h

6.42 xnvfile_snapshot_ops Struct Reference

Snapshot vfile operation descriptor.

Data Fields

- int(* [rewind](#))(struct [xnvfile_snapshot_iterator](#) *it)
- void (* [begin](#))(struct [xnvfile_snapshot_iterator](#) *it)
- void(* [end](#))(struct [xnvfile_snapshot_iterator](#) *it, void *buf)
- int(* [next](#))(struct [xnvfile_snapshot_iterator](#) *it, void *data)
- int(* [show](#))(struct [xnvfile_snapshot_iterator](#) *it, void *data)
- ssize_t(* [store](#))(struct xnvfile_input *input)

6.42.1 Detailed Description

Snapshot vfile operation descriptor.

This structure describes the operations available with a snapshot-driven vfile. It defines handlers for returning a printable snapshot of some Xenomai object contents upon a user-space read request, and for updating this object upon a user-space write request.

6.42.2 Field Documentation

6.42.2.1 void(* [xnvfile_snapshot_ops::begin](#))(struct **xnvfile_snapshot_iterator** *it)

This handler should allocate the snapshot buffer to hold records during the data collection phase. When specified, all records collected via the [next\(\)](#) handler" will be written to a cell from the memory area returned by [begin\(\)](#).

Parameters

<i>it</i>	A pointer to the current snapshot iterator.
-----------	---

Returns

A pointer to the record buffer, if the call succeeds. Otherwise:

- NULL in case of allocation error. This will abort the data collection, and return -ENOMEM to the reader.
- VFILE_SEQ_EMPTY, a special value indicating that no record will be output. In such a case, the [next\(\) handler](#) will not be called, and the data collection will stop immediately. However, the [show\(\) handler](#) will still be called once, with a NULL data pointer (i.e. header display request).

Note

This handler is optional; if none is given, an internal allocation depending on the value returned by the [rewind\(\) handler](#) can be obtained.

6.42.2.2 void(* xnvfile_snapshot_ops::end)(struct **xnvfile_snapshot_iterator** *it, void *buf)

This handler releases the memory buffer previously obtained from [begin\(\)](#). It is usually called after the snapshot data has been output by [show\(\)](#), but it may also be called before rewinding the vfile after a revision change, to release the dropped buffer.

Parameters

<i>it</i>	A pointer to the current snapshot iterator.
<i>buf</i>	A pointer to the buffer to release.

Note

This routine is optional and the pointer may be NULL. It is not needed upon internal buffer allocation; see the description of the [rewind\(\) handler](#)".

6.42.2.3 int(* xnvfile_snapshot_ops::next)(struct **xnvfile_snapshot_iterator** *it, void *data)

This handler fetches the next record, as part of the snapshot data to be sent back to the reader via the [show\(\)](#).

Parameters

<i>it</i>	A pointer to the current snapshot iterator.
<i>data</i>	A pointer to the record to fill in.

Returns

a strictly positive value, if the call succeeds and leaves a valid record into *data*, which should be passed to the [show\(\) handler](#) during the formatting and output phase. Otherwise:

- A negative error code. This will abort the data collection, and return this status to the reader.
- VFILE_SEQ_SKIP, a special value indicating that the current record should be skipped. In such a case, the *data* pointer is not advanced to the next position before the [next\(\) handler](#) is called anew.

Note

This handler is called with the vfile lock held. Before each invocation of this handler, the vfile core checks whether the revision tag has been touched, in which case the data collection is restarted from scratch. A data collection phase succeeds whenever all records can be fetched via the [next\(\) handler](#), while the revision tag remains unchanged, which indicates that a consistent snapshot of the object state was taken.

6.42.2.4 `int(* xnvfile_snapshot_ops::rewind)(struct xnvfile_snapshot_iterator *it)`

This handler (re-)initializes the data collection, moving the seek pointer at the first record. When the file revision tag is touched while collecting data, the current reading is aborted, all collected data dropped, and the vfile is eventually rewound.

Parameters

<i>it</i>	A pointer to the current snapshot iterator. Two useful information can be retrieved from this iterator in this context:
-----------	---

- `it->vfile` is a pointer to the descriptor of the virtual file being rewound.
- `xnvfile_iterator_priv(it)` returns a pointer to the private data area, available from the descriptor, which size is `vfile->privsz`. If the latter size is zero, the returned pointer is meaningless and should not be used.

Returns

A negative error code aborts the data collection, and is passed back to the reader. Otherwise:

- a strictly positive value is interpreted as the total number of records which will be returned by the [next\(\) handler](#) during the data collection phase. If no [begin\(\) handler](#) is provided in the [operation descriptor](#), this value is used to allocate the snapshot buffer internally. The size of this buffer would then be `vfile->datasz * value`.
- zero leaves the allocation to the [begin\(\) handler](#) if present, or indicates that no record is to be output in case such handler is not given.

Note

This handler is optional; a NULL value indicates that nothing needs to be done for rewinding the vfile. It is called with the vfile lock held.

6.42.2.5 `int(* xnvfile_snapshot_ops::show)(struct xnvfile_snapshot_iterator *it, void *data)`

This handler should format and output a record from the collected data.

`xnvfile_printf()`, `xnvfile_write()`, `xnvfile_puts()` and `xnvfile_putc()` are available to format and/or emit the output. All routines take the iterator argument *it* as their first parameter.

Parameters

<i>it</i>	A pointer to the current snapshot iterator.
-----------	---

<i>data</i>	A pointer to the record to format then output. The first call to the handler is always passed a NULL <i>data</i> pointer; the show handler should test this special value to output any header that fits, prior to receiving more calls with actual records.
-------------	--

Returns

zero if the call succeeds, also indicating that the handler should be called for the next record if any. Otherwise:

- A negative error code. This will abort the output phase, and return this status to the reader.
- `VFILE_SEQ_SKIP`, a special value indicating that the current record should be skipped and will not be output.

6.42.2.6 `ssize_t(* xnvfile_snapshot_ops::store)(struct xnvfile_input *input)`

This handler receives data written to the vfile, likely for updating the associated Xenomai object's state, or triggering any other action which fits. This is the only handler which deals with the write-side of a vfile. It is called when writing to the `/proc` entry of the vfile from a user-space process.

The input data is described by a descriptor passed to the handler, which may be subsequently passed to parsing helper routines. For instance, [xnvfile_get_string\(\)](#) will accept the input descriptor for returning the written data as a null-terminated character string. On the other hand, [xnvfile_get_integer\(\)](#) will attempt to return a long integer from the input data.

Parameters

<i>input</i>	A pointer to an input descriptor. It refers to an opaque data from the handler's standpoint.
--------------	--

Returns

the number of bytes read from the input descriptor if the call succeeds. Otherwise, a negative error code. Return values from parsing helper routines are commonly passed back to the caller by the [store\(\) handler](#).

Note

This handler is optional, and may be omitted for read-only vfiles.

Referenced by `xnvfile_init_snapshot()`.

The documentation for this struct was generated from the following file:

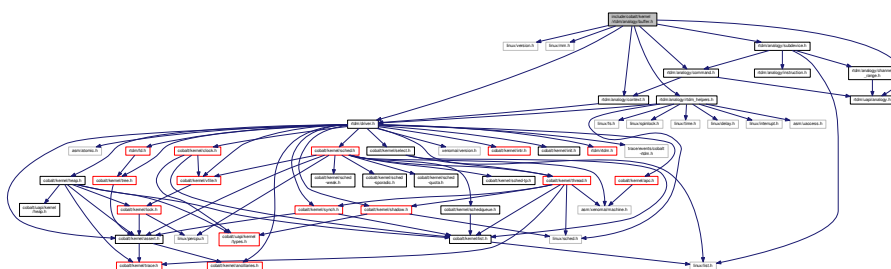
- `include/cobalt/kernel/vfile.h`

File Documentation


7.1 include/cobalt/kernel/rtdm/analogry/buffer.h File Reference

Analogy for Linux, buffer related features.

Include dependency graph for buffer.h:



This graph shows which files directly or indirectly include this file:



7.1.1 Detailed Description

Analogy for Linux, buffer related features. Copyright (C) 1997-2000 David A. Schlee ds@schlee.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

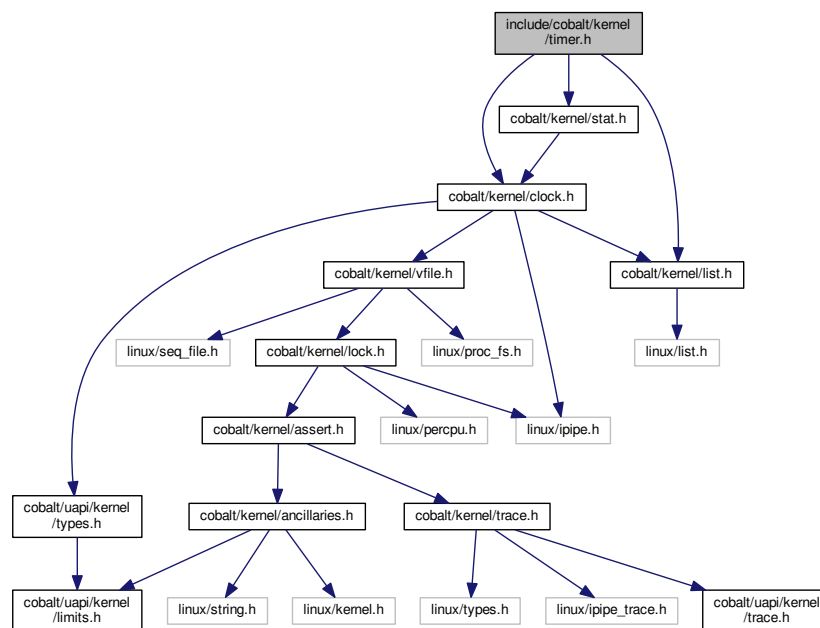
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

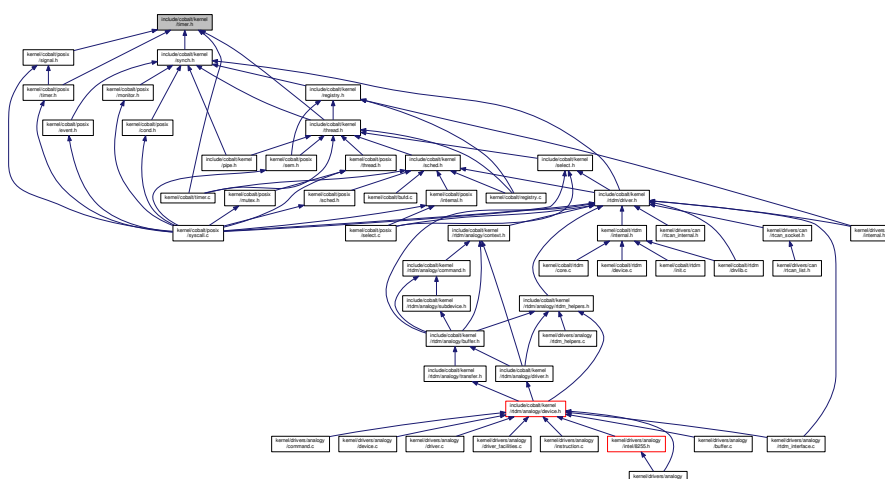
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.2 include/cobalt/kernel/timer.h File Reference

Include dependency graph for timer.h:



This graph shows which files directly or indirectly include this file:



Functions

- void [xntimer_destroy](#) (struct xntimer *timer)

- Generated on Fri Jun 13 2014 12:09:58 for Xenomai by Doxygen

7.3.2.1 void xnsched_rotate (struct **xnsched** * sched, struct xnsched_class * sched_class, const union xnsched_policy_param * sched_param) [inline], [static]

Rotate a scheduler runqueue.

The specified scheduling class is requested to rotate its runqueue for the given scheduler. Rotation is performed according to the scheduling parameter specified by *sched_param*.

Note

The nucleus supports round-robin scheduling for the members of the RT class.

Parameters

<i>sched</i>	The per-CPU scheduler hosting the target scheduling class.
<i>sched_class</i>	The scheduling class which should rotate its runqueue.
<i>sched_param</i>	The scheduling parameter providing rotation information to the specified scheduling class.

Environments:

This service should be called from:

- Kernel-based task
- Interrupt service routine
- User-space task (primary mode only)

Rescheduling: never.

7.3.2.2 int xnsched_run (void) [inline], [static]

The rescheduling procedure.

This is the central rescheduling routine which should be called to validate and apply changes which have previously been made to the nucleus scheduling state, such as suspending, resuming or changing the priority of threads. This call performs context switches as needed. [xnsched_run\(\)](#) schedules out the current thread if:

- the current thread is about to block.
- a runnable thread from a higher priority scheduling class is waiting for the CPU.
- the current thread does not lead the runnable threads from its own scheduling class (i.e. round-robin).

The Cobalt core implements a lazy rescheduling scheme so that most of the services affecting the threads state MUST be followed by a call to the rescheduling procedure for the new scheduling state to be applied.

In other words, multiple changes on the scheduler state can be done in a row, waking threads up, blocking others, without being immediately translated into the corresponding context switches. When all changes have been applied, [xnsched_run\(\)](#) should be called for considering those changes, and possibly switching context.

As a notable exception to the previous principle however, every action which ends up suspending the current thread begets an implicit call to the rescheduling procedure on behalf of the blocking service.

Typically, self-suspension or sleeping on a synchronization object automatically leads to a call to the rescheduling procedure, therefore the caller does not need to explicitly issue [xnsched_run\(\)](#) after such operations.

Functions

- int [xnapc_alloc](#) (const char *name, void(*handler)(void *cookie), void *cookie)

Allocate an APC slot.

- void [xnapc_free](#) (int apc)

Releases an APC slot.

- static void [xnapc_schedule](#) (int apc)

Schedule an APC invocation.

7.4.1 Detailed Description

Note

Copyright (C) 2012 Philippe Gerum rpm@xenomai.org.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.5 include/cobalt/kernel/bufd.h File Reference

Include dependency graph for bufd.h:



This graph shows which files directly or indirectly include this file:



Functions

- static void [xnbufd_map_uread](#) (struct xnbufd *bufd, const void __user *ptr, size_t len)
Initialize a buffer descriptor for reading from user memory.
- static void [xnbufd_map_uwrite](#) (struct xnbufd *bufd, void __user *ptr, size_t len)
Initialize a buffer descriptor for writing to user memory.
- ssize_t [xnbufd_unmap_uread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uread\(\)](#).
- ssize_t [xnbufd_unmap_uwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uwrite\(\)](#).
- static void [xnbufd_map_kread](#) (struct xnbufd *bufd, const void *ptr, size_t len)
Initialize a buffer descriptor for reading from kernel memory.
- static void [xnbufd_map_kwrite](#) (struct xnbufd *bufd, void *ptr, size_t len)
Initialize a buffer descriptor for writing to kernel memory.
- ssize_t [xnbufd_unmap_kread](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kread\(\)](#).
- ssize_t [xnbufd_unmap_kwrite](#) (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kwrite\(\)](#).
- ssize_t [xnbufd_copy_to_kmem](#) (void *ptr, struct xnbufd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- ssize_t [xnbufd_copy_from_kmem](#) (struct xnbufd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- void [xnbufd_invalidate](#) (struct xnbufd *bufd)
Invalidate a buffer descriptor.
- static void [xnbufd_reset](#) (struct xnbufd *bufd)
Reset a buffer descriptor.

7.5.1 Detailed Description

Note

Copyright (C) 2009 Philippe Gerum rpm@xenomai.org.

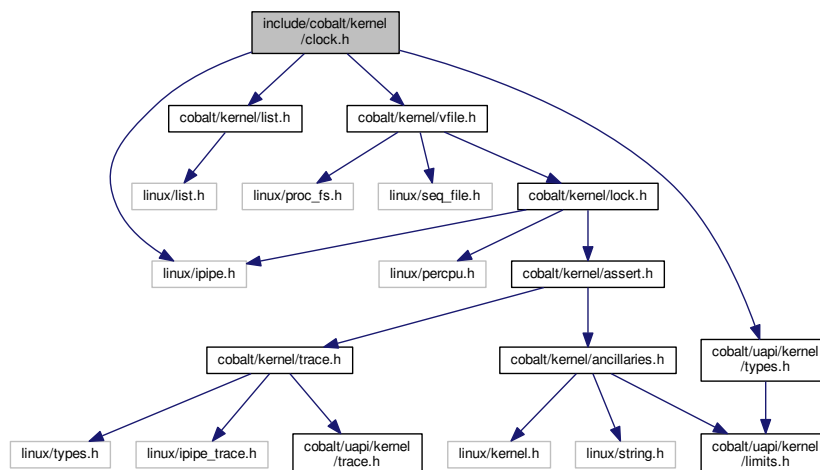
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.6 include/cobalt/kernel/clock.h File Reference

Include dependency graph for clock.h:



This graph shows which files directly or indirectly include this file:



Functions

- int `xnclock_register` (struct xnclock *clock)
Register a Xenomai clock.
- void `xnclock_deregister` (struct xnclock *clock)
Deregister a Xenomai clock.
- void `xnclock_tick` (struct xnclock *clock)
Process a clock tick.
- void `xnclock_adjust` (struct xnclock *clock, xnsticks_t delta)
Adjust a clock time.

7.6.1 Detailed Description

Note

Copyright (C) 2006,2007 Philippe Gerum rpm@xenomai.org.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.7 include/cobalt/kernel/map.h File Reference

Include dependency graph for map.h:



This graph shows which files directly or indirectly include this file:



Functions

- `struct xnmap * xnmap_create` (int nkeys, int reserve, int offset)
Create a map.
- `void xnmap_delete` (struct xnmap *map)
Delete a map.
- `int xnmap_enter` (struct xnmap *map, int key, void *objaddr)
Index an object into a map.
- `int xnmap_remove` (struct xnmap *map, int key)
Remove an object reference from a map.
- `static void * xnmap_fetch_nocheck` (struct xnmap *map, int key)
Search an object into a map - unchecked form.
- `static void * xnmap_fetch` (struct xnmap *map, int key)
Search an object into a map.

7.7.1 Detailed Description

Note

Copyright (C) 2007 Philippe Gerum rpm@xenomai.org.


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

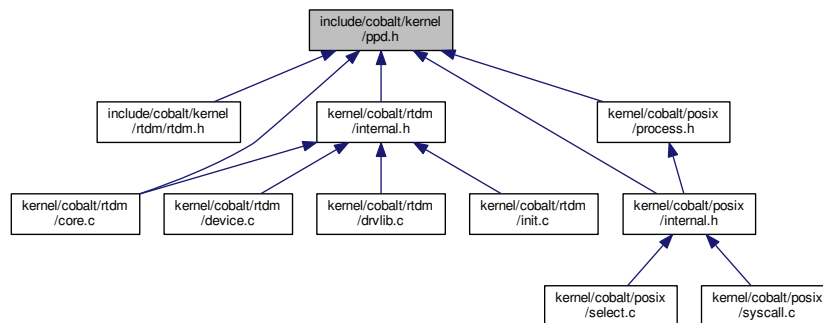
You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.8 include/cobalt/kernel/ppd.h File Reference

Include dependency graph for ppd.h:



This graph shows which files directly or indirectly include this file:



Functions

- `void * xnshadow_get_context` (unsigned int muxid)
Return the per-process data attached to the calling process.

7.8.1 Detailed Description

Note

Copyright © 2006 Gilles Chantepierdrix gch@xenomai.org Per-process data.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, Inc., 675 Mass Ave, Cambridge MA 02139, USA; either version 2 of the License, or (at your option) any later version.

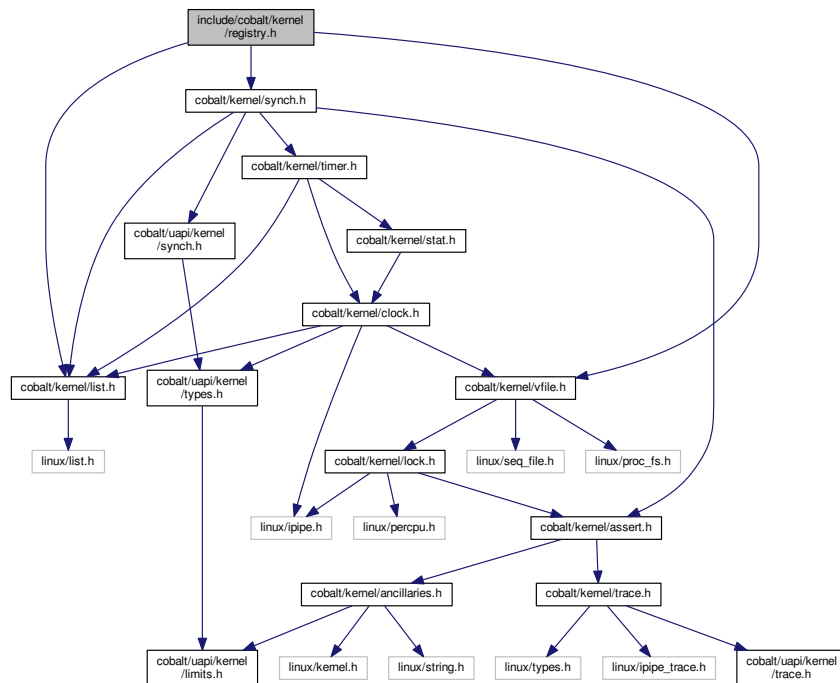
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

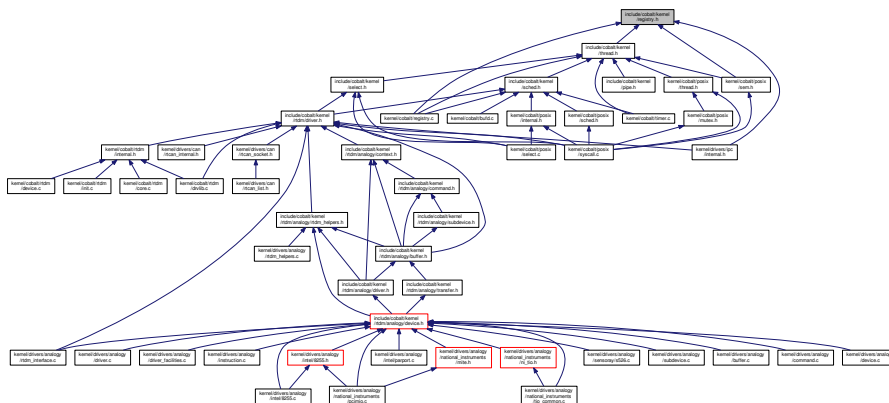
7.9 include/cobalt/kernel/registry.h File Reference

This file is part of the Xenomai project.

Include dependency graph for registry.h:



This graph shows which files directly or indirectly include this file:



Functions

- `int xnregistry_enter(const char *key, void *objaddr, xnhandle_t *phandle, struct xnnode *pnode)`
Register a real-time object.
- `int xnregistry_bind(const char *key, xnticks_t timeout, int timeout_mode, xnhandle_t *phandle)`
Bind to a real-time object.

- int `xnregistry_remove` (xnhandle_t handle)

Forcibly unregister a real-time object.

- static void * `xnregistry_lookup` (xnhandle_t handle, unsigned long *cstamp_r)

Find a real-time object into the registry.

- int `xnregistry_unlink` (const char *key)

Turn a named object into an anonymous object.

7.9.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2004 Philippe Gerum rpm@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

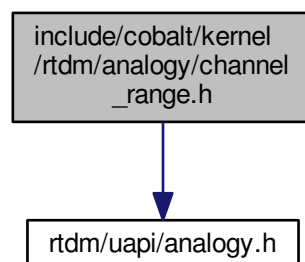
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

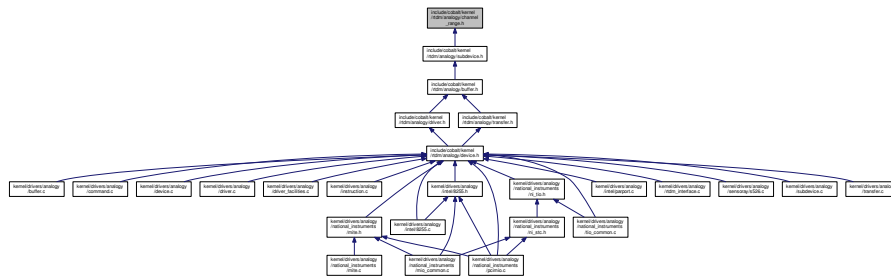
7.10 include/cobalt/kernel/rtdm/analogy/channel_range.h File Reference

Analogy for Linux, channel, range related features.

Include dependency graph for channel_range.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [a4l_channel](#)
Structure describing some channel's characteristics.
- struct [a4l_channels_desc](#)
Structure describing a channels set.
- struct [a4l_range](#)
Structure describing a (unique) range.

Macros

- #define [A4L_CHAN_GLOBAL](#) 0x10
Internal use flag (must not be used by driver developer)
- #define [A4L_RNG_GLOBAL](#) 0x8
Internal use flag (must not be used by driver developer)
- #define [RANGE](#)(x, y)
Macro to declare a (unique) range with no unit defined.
- #define [RANGE_V](#)(x, y)
Macro to declare a (unique) range in Volt.
- #define [RANGE_mA](#)(x, y)
Macro to declare a (unique) range in milliAmpere.
- #define [RANGE_ext](#)(x, y)
Macro to declare a (unique) range in some external reference.
- #define [A4L_RNG_GLOBAL_RNGDESC](#) 0
Constant to define a ranges descriptor as global (inter-channel)
- #define [A4L_RNG_PERCHAN_RNGDESC](#) 1
Constant to define a ranges descriptor as specific for a channel.
- #define [RNG_GLOBAL](#)(x)
Macro to declare a ranges global descriptor in one line.

Channel reference

Flags to define the channel's reference

- #define [A4L_CHAN_AREF_GROUND](#) 0x1
Ground reference.
- #define [A4L_CHAN_AREF_COMMON](#) 0x2
Common reference.
- #define [A4L_CHAN_AREF_DIFF](#) 0x4

The diagram illustrates the knowledge structure for the 'Knowledge of the basic concepts of the theory of the development of the human body' (KBC). It is a hierarchical tree structure. At the top is the root node 'KBC'. Below it are several intermediate nodes, including 'KBC-1', 'KBC-2', 'KBC-3', 'KBC-4', 'KBC-5', 'KBC-6', 'KBC-7', 'KBC-8', 'KBC-9', 'KBC-10', 'KBC-11', 'KBC-12', 'KBC-13', 'KBC-14', 'KBC-15', 'KBC-16', 'KBC-17', 'KBC-18', 'KBC-19', 'KBC-20', 'KBC-21', 'KBC-22', 'KBC-23', 'KBC-24', 'KBC-25', 'KBC-26', 'KBC-27', 'KBC-28', 'KBC-29', 'KBC-30', 'KBC-31', 'KBC-32', 'KBC-33', 'KBC-34', 'KBC-35', 'KBC-36', 'KBC-37', 'KBC-38', 'KBC-39', 'KBC-40', 'KBC-41', 'KBC-42', 'KBC-43', 'KBC-44', 'KBC-45', 'KBC-46', 'KBC-47', 'KBC-48', 'KBC-49', 'KBC-50', 'KBC-51', 'KBC-52', 'KBC-53', 'KBC-54', 'KBC-55', 'KBC-56', 'KBC-57', 'KBC-58', 'KBC-59', 'KBC-60', 'KBC-61', 'KBC-62', 'KBC-63', 'KBC-64', 'KBC-65', 'KBC-66', 'KBC-67', 'KBC-68', 'KBC-69', 'KBC-70', 'KBC-71', 'KBC-72', 'KBC-73', 'KBC-74', 'KBC-75', 'KBC-76', 'KBC-77', 'KBC-78', 'KBC-79', 'KBC-80', 'KBC-81', 'KBC-82', 'KBC-83', 'KBC-84', 'KBC-85', 'KBC-86', 'KBC-87', 'KBC-88', 'KBC-89', 'KBC-90', 'KBC-91', 'KBC-92', 'KBC-93', 'KBC-94', 'KBC-95', 'KBC-96', 'KBC-97', 'KBC-98', 'KBC-99', 'KBC-100'. The terminal nodes are organized into four main groups: 'Basic concepts of the theory of the development of the human body', 'Basic concepts of the theory of the development of the human body', 'Basic concepts of the theory of the development of the human body', and 'Basic concepts of the theory of the development of the human body'. The diagram illustrates the relationships between different levels of knowledge, from general concepts to specific details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

```

graph TD
    IS[Information Systems] --> ISE[Information Systems Engineering]
    IS --> ISM[Information Systems Management]
    ISE --> ISE_S[Information Systems Engineering Subdomains]
    ISE --> ISE_P[Information Systems Engineering Practices]
    ISE_S --> ISE_S_S[Information Systems Engineering Subdomains]
    ISE_S --> ISE_S_P[Information Systems Engineering Practices]
    ISE_P --> ISE_P_S[Information Systems Engineering Practices]
    ISE_P --> ISE_P_P[Information Systems Engineering Practices]
    
```


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

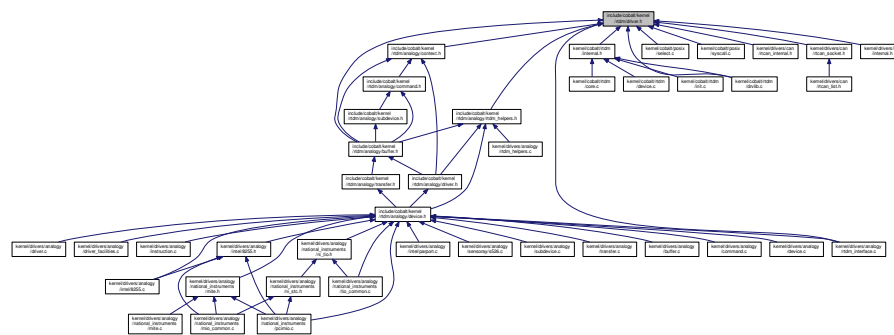
7.14 include/cobalt/kernel/rtdm/driver.h File Reference

Real-Time Driver Model for Xenomai, driver API header.

Include dependency graph for driver.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct `rtdm_dev_context`
Device context.
- struct `rtdm_device`
RTDM device.

Macros

- `#define cobalt_atomic_enter(context)`
Enter atomic section (dual kernel only)
- `#define cobalt_atomic_leave(context)`

- *Leave atomic section (dual kernel only)*
- #define `RTDM_EXECUTE_ATOMICALLY`(code_block)
Execute code block atomically (DEPRECATED)
- #define `RTDM_LOCK_UNLOCKED`(__name) `IPIPE_SPIN_LOCK_UNLOCKED`
Static lock initialisation.
- #define `rtdm_lock_irqsave`(context) `splhigh`(context)
Disable preemption locally.
- #define `rtdm_lock_irqrestore`(context) `splexit`(context)
Restore preemption state.
- #define `rtdm_irq_get_arg`(irq_handle, type) ((type *)irq_handle->cookie)
Retrieve IRQ handler argument.

Device Flags

Static flags describing a RTDM device

- #define `RTDM_EXCLUSIVE` 0x0001
If set, only a single instance of the device can be requested by an application.
- #define `RTDM_NAMED_DEVICE` 0x0010
If set, the device is addressed via a clear-text name.
- #define `RTDM_PROTOCOL_DEVICE` 0x0020
If set, the device is addressed via a combination of protocol ID and socket type.
- #define `RTDM_DEVICE_TYPE_MASK` 0x00F0
Mask selecting the device type.

Driver Versioning

Current revisions of RTDM structures, encoding of driver versions. See [API Versioning](#) for the interface revision.

- #define `RTDM_DEVICE_STRUCT_VER` 6
Version of struct `rtdm_device`.
- #define `RTDM_CONTEXT_STRUCT_VER` 4
Version of struct `rtdm_dev_context`.
- #define `RTDM_SECURE_DEVICE` 0x80000000
Flag indicating a secure variant of RTDM (not supported here)
- #define `RTDM_DRIVER_VER`(major, minor, patch) (((major & 0xFF) << 16) | ((minor & 0xFF) << 8) | (patch & 0xFF))
Version code constructor for driver revisions.
- #define `RTDM_DRIVER_MAJOR_VER`(ver) (((ver) >> 16) & 0xFF)
Get major version number from driver revision code.
- #define `RTDM_DRIVER_MINOR_VER`(ver) (((ver) >> 8) & 0xFF)
Get minor version number from driver revision code.
- #define `RTDM_DRIVER_PATCH_VER`(ver) ((ver) & 0xFF)
Get patch version number from driver revision code.

RTDM_IRQTYPE_xxx

Interrupt registrations flags

- #define `RTDM_IRQTYPE_SHARED` `XN_ISR_SHARED`
Enable IRQ-sharing with other real-time drivers.
- #define `RTDM_IRQTYPE_EDGE` `XN_ISR_EDGE`
Mark IRQ as edge-triggered, relevant for correct handling of shared edge-triggered IRQs.

RTDM_IRQ_xxx

Return flags of interrupt handlers

- #define [RTDM_IRQ_NONE](#) XN_ISR_NONE
Unhandled interrupt.
- #define [RTDM_IRQ_HANDLED](#) XN_ISR_HANDLED
Denote handled interrupt.

Task Priority Range

Maximum and minimum task priorities

- #define [RTDM_TASK_LOWEST_PRIORITY](#) 0
- #define [RTDM_TASK_HIGHEST_PRIORITY](#) 99

Task Priority Modification

Raise or lower task priorities by one level

- #define [RTDM_TASK_RAISE_PRIORITY](#) (+1)
- #define [RTDM_TASK_LOWER_PRIORITY](#) (-1)

Typedefs

- typedef ipipe_spinlock_t [rtdm_lock_t](#)
Lock variable.
- typedef unsigned long [rtdm_lockctx_t](#)
Variable to save the context while holding a lock.
- typedef int(* [rtdm_irq_handler_t](#))(rtdm_irq_t *irq_handle)
Interrupt handler.
- typedef void(* [rtdm_nrtsig_handler_t](#))(rtdm_nrtsig_t nrt_sig, void *arg)
Non-real-time signal handler.
- typedef void(* [rtdm_timer_handler_t](#))(rtdm_timer_t *timer)
Timer handler.
- typedef void(* [rtdm_task_proc_t](#))(void *arg)
Real-time task procedure.

Operation Handler Prototypes

- typedef int(* [rtdm_open_handler_t](#))(struct rtdm_fd *fd, int oflag)
Named device open handler.
- typedef int(* [rtdm_socket_handler_t](#))(struct rtdm_fd *fd, int protocol)
Socket creation handler for protocol devices.

Enumerations

[RTDM_SELECTTYPE_xxx](#)

Event types select can bind to

- enum [rtdm_selecttype](#) { [RTDM_SELECTTYPE_READ](#) = XNSELECT_READ, [RTDM_SELECTTYPE_WRITE](#) = XNSELECT_WRITE, [RTDM_SELECTTYPE_EXCEPT](#) = XNSELECT_EXCEPT }

[RTDM_TIMERMODE_xxx](#)

Timer operation modes

- enum [rtdm_timer_mode](#) { [RTDM_TIMERMODE_RELATIVE](#) = XN_RELATIVE, [RTDM_TIMERMODE_ABSOLUTE](#) = XN_ABSOLUTE, [RTDM_TIMERMODE_REALTIME](#) = XN_REALTIME }

Functions

- static void * [rtdm_fd_to_private](#) (struct rtdm_fd *fd)
Locate the driver private area associated to a device context structure.
- static struct rtdm_fd * [rtdm_private_to_fd](#) (void *dev_private)
Locate a device file descriptor structure from its driver private area.
- static bool [rtdm_fd_is_user](#) (struct rtdm_fd *fd)
Tell whether the passed file descriptor belongs to an application.
- static struct [rtdm_device](#) * [rtdm_fd_device](#) (struct rtdm_fd *fd)
Locate a device structure from a file descriptor.
- int [rtdm_dev_register](#) (struct [rtdm_device](#) *device)
Register a RTDM device.
- int [rtdm_dev_unregister](#) (struct [rtdm_device](#) *device, unsigned int poll_delay)
Unregisters a RTDM device.
- void [rtdm_toseq_init](#) (rtdm_toseq_t *timeout_seq, [nanosecs_rel_t](#) timeout)
Initialise a timeout sequence.
- static void [rtdm_lock_init](#) (rtdm_lock_t *lock)
Dynamic lock initialisation.
- static void [rtdm_lock_get](#) (rtdm_lock_t *lock)
Acquire lock from non-preemptible contexts.
- static void [rtdm_lock_put](#) (rtdm_lock_t *lock)
Release lock without preemption restoration.
- static void [rtdm_lock_put_irqrestore](#) (rtdm_lock_t *lock, [rtdm_lockctx_t](#) context)
Release lock and restore preemption state.
- int [rtdm_irq_request](#) (rtdm_irq_t *irq_handle, unsigned int irq_no, [rtdm_irq_handler_t](#) handler, unsigned long flags, const char *device_name, void *arg)
Register an interrupt handler.
- void [rtdm_timer_destroy](#) (rtdm_timer_t *timer)
Destroy a timer.
- int [rtdm_timer_start](#) (rtdm_timer_t *timer, [nanosecs_abs_t](#) expiry, [nanosecs_rel_t](#) interval, enum [rtdm_timer_mode](#) mode)
Start a timer.
- void [rtdm_timer_stop](#) (rtdm_timer_t *timer)
Stop a timer.
- int [rtdm_task_init](#) (rtdm_task_t *task, const char *name, [rtdm_task_proc_t](#) task_proc, void *arg, int priority, [nanosecs_rel_t](#) period)
Initialise and start a real-time task.
- void [rtdm_task_busy_sleep](#) ([nanosecs_rel_t](#) delay)
Busy-wait a specified amount of time.
- void [rtdm_event_init](#) (rtdm_event_t *event, unsigned long pending)
Initialise an event.
- int [rtdm_event_wait](#) (rtdm_event_t *event)
Wait on event occurrence.
- int [rtdm_event_timedwait](#) (rtdm_event_t *event, [nanosecs_rel_t](#) timeout, rtdm_toseq_t *timeout_seq)
Wait on event occurrence with timeout.
- void [rtdm_event_signal](#) (rtdm_event_t *event)
Signal an event occurrence.
- void [rtdm_event_clear](#) (rtdm_event_t *event)
Clear event state.
- void [rtdm_event_pulse](#) (rtdm_event_t *event)

- Signal an event occurrence to currently listening waiters.*

 - void `rtdm_event_destroy` (`rtdm_event_t *event`)

Destroy an event.
- void `rtdm_sem_init` (`rtdm_sem_t *sem`, unsigned long value)

Initialise a semaphore.
- int `rtdm_sem_down` (`rtdm_sem_t *sem`)

Decrement a semaphore.
- int `rtdm_sem_timeddown` (`rtdm_sem_t *sem`, `nanosecs_rel_t` timeout, `rtdm_toseq_t *timeout_seq`)

Decrement a semaphore with timeout.
- void `rtdm_sem_up` (`rtdm_sem_t *sem`)

Increment a semaphore.
- void `rtdm_sem_destroy` (`rtdm_sem_t *sem`)

Destroy a semaphore.
- void `rtdm_mutex_init` (`rtdm_mutex_t *mutex`)

Initialise a mutex.
- int `rtdm_mutex_lock` (`rtdm_mutex_t *mutex`)

Request a mutex.
- int `rtdm_mutex_timedlock` (`rtdm_mutex_t *mutex`, `nanosecs_rel_t` timeout, `rtdm_toseq_t *timeout_seq`)

Request a mutex with timeout.
- void `rtdm_mutex_unlock` (`rtdm_mutex_t *mutex`)

Release a mutex.
- void `rtdm_mutex_destroy` (`rtdm_mutex_t *mutex`)

Destroy a mutex.
- int `rtdm_ratelimit` (`struct rtdm_ratelimit_state *rs`, `const char *func`)

Enforces a rate limit.

7.14.1 Detailed Description

Real-Time Driver Model for Xenomai, driver API header.

Note

Copyright (C) 2005-2007 Jan Kiszka jan.kiszka@web.de
 Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net
 Copyright (C) 2008 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.14.2 Typedef Documentation

7.14.2.1 `typedef int(* rtdm_open_handler_t)(struct rtdm_fd *fd, int oflag)`

Named device open handler.

Parameters

in	<i>fd</i>	File descriptor structure associated with opened device instance
in	<i>oflag</i>	Open flags as passed by the user

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See Also

open() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

7.14.2.2 typedef int(* rtdm_socket_handler_t)(struct rtdm_fd *fd, int protocol)

Socket creation handler for protocol devices.

Parameters

in	<i>fd</i>	File descriptor structure associated with opened device instance
in	<i>protocol</i>	Protocol number as passed by the user

Returns

0 on success. On failure return either -ENOSYS, to request that this handler be called again from the opposite realtime/non-realtime context, or another negative error code.

See Also

socket() in IEEE Std 1003.1, <http://www.opengroup.org/onlinepubs/009695399>

7.14.3 Function Documentation

7.14.3.1 static struct **rtdm_device*** rtdm_fd_device (struct rtdm_fd * fd) [static]

Locate a device structure from a file descriptor.

Parameters

in	<i>fd</i>	File descriptor
----	-----------	-----------------

Returns

The address of the device structure to which this file descriptor is attached.

References rtdm_dev_context::device.

7.14.3.2 static bool rtdm_fd_is_user (struct rtdm_fd * fd) [inline], [static]

Tell whether the passed file descriptor belongs to an application.

Parameters

<code>in</code>	<code>fd</code>	File descriptor
-----------------	-----------------	-----------------

Returns

true if passed file descriptor belongs to an application, false otherwise.

7.14.3.3 static void* rtdm_fd_to_private (struct rtdm_fd * fd) [inline], [static]

Locate the driver private area associated to a device context structure.

Parameters

<code>in</code>	<code>fd</code>	File descriptor structure associated with opened device instance
-----------------	-----------------	--

Returns

The address of the private driver area associated to *file* descriptor.

References rtdm_dev_context::dev_private.

```
7.14.3.4 static struct rtdm fd* rtdm_private_to_fd ( void *dev_private ) [static]
```

Locate a device file descriptor structure from its driver private area.

Parameters

in	<i>dev_private</i>	Address of a private context area
----	--------------------	-----------------------------------

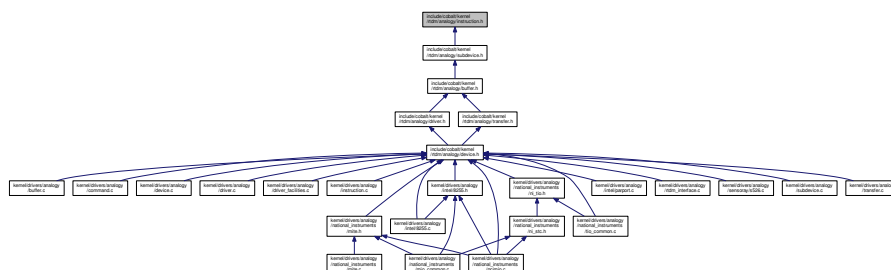
Returns

The address of the file descriptor structure defining *dev_private*.

7.15 include/cobalt/kernel/rtdm/analogry/instruction.h File Reference

Analogy for Linux, instruction related features.

This graph shows which files directly or indirectly include this file:

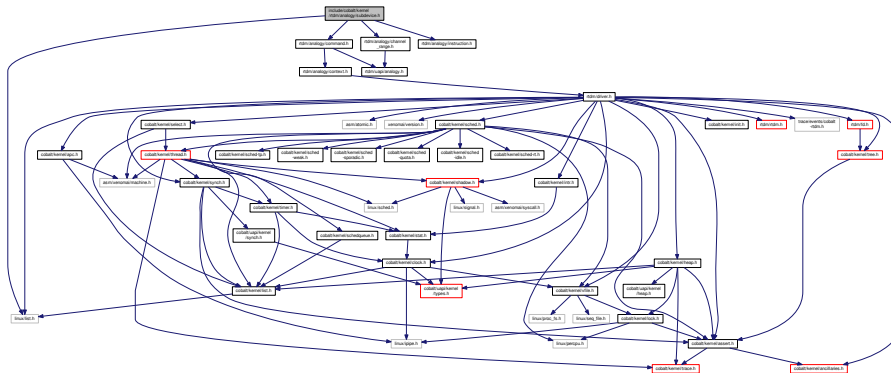


7.15.1 Detailed Description

Analogy for Linux, instruction related features.

7.17 include/cobalt/kernel/rtdm/analog/subdevice.h File Reference

Include dependency graph for subdevice.h:


[illegible]

- struct `a4l_subdevice`
Structure describing the subdevice.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.19 include/cobalt/kernel/rtdm/can.h File Reference

Include dependency graph for can.h:



7.19.1 Detailed Description

Note

Copyright (C) 2006 Wolfgang Grandegger wg@grandegger.com

Copyright (C) 2005, 2006 Sebastian Smolorz Sebastian.Smolorz@stud.uni-hannover.de

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU

General Public License for more details.


You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.20 include/rtdm/can.h File Reference

Include dependency graph for can.h:



This graph shows which files directly or indirectly include this file:



7.20.1 Detailed Description

Note

Copyright (C) 2006 Wolfgang Grandegger wg@grandegger.com

Copyright (C) 2005, 2006 Sebastian Smolorz Sebastian.Smolorz@stud.uni-hannover.de

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU


General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.21 include/rtdm/uapi/can.h File Reference

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

This graph shows which files directly or indirectly include this file:



Data Structures

- struct [can_bittime_std](#)
Standard bit-time parameters according to Bosch.
- struct [can_bittime_btr](#)
Hardware-specific BTR bit-times.
- struct [can_bittime](#)
Custom CAN bit-time definition.
- struct [can_filter](#)
Filter for reception of CAN messages.
- struct [sockaddr_can](#)
Socket address structure for the CAN address family.
- struct [can_frame](#)
Raw CAN frame.

Macros

- `#define` [AF_CAN](#) 29
CAN address family.
- `#define` [PF_CAN](#) [AF_CAN](#)
CAN protocol family.
- `#define` [SOL_CAN_RAW](#) 103
CAN socket levels.

CAN ID masks

Bit masks for masking CAN IDs

- #define `CAN_EFF_MASK` 0x1FFFFFFF
Bit mask for extended CAN IDs.
- #define `CAN_SFF_MASK` 0x000007FF
Bit mask for standard CAN IDs.

CAN ID flags

Flags within a CAN ID indicating special CAN frame attributes

- #define `CAN_EFF_FLAG` 0x80000000
Extended frame.
- #define `CAN_RTR_FLAG` 0x40000000
Remote transmission frame.
- #define `CAN_ERR_FLAG` 0x20000000
Error frame (see [Errors](#)), not valid in struct `can_filter`.
- #define `CAN_INV_FILTER CAN_ERR_FLAG`
Invert CAN filter definition, only valid in struct `can_filter`.

Particular CAN protocols

Possible protocols for the PF_CAN protocol family

Currently only the RAW protocol is supported.

- #define `CAN_RAW` 1
Raw protocol of PF_CAN, applicable to socket type `SOCK_RAW`.

CAN controller modes

Special CAN controllers modes, which can be or'ed together.

Note

These modes are hardware-dependent. Please consult the hardware manual of the CAN controller for more detailed information.

- #define `CAN_CTRLMODE_LISTENONLY` 0x1
- #define `CAN_CTRLMODE_LOOPBACK` 0x2
- #define `CAN_CTRLMODE_3_SAMPLES` 0x4

Timestamp switches

Arguments to pass to `RTCAN_RTIOC_TAKE_TIMESTAMP`

- #define `RTCAN_TAKE_NO_TIMESTAMPS` 0
Switch off taking timestamps.
- #define `RTCAN_TAKE_TIMESTAMPS` 1
Do take timestamps.

RAW socket options

Setting and getting CAN RAW socket options.

- #define `CAN_RAW_FILTER` 0x1
CAN filter definition.
- #define `CAN_RAW_ERR_FILTER` 0x2
CAN error mask.
- #define `CAN_RAW_LOOPBACK` 0x3
CAN TX loopback.
- #define `CAN_RAW_RECV_OWN_MSGS` 0x4

CAN receive own messages.

IOCTLs

CAN device IOCTLs

- #define [SIOCGIFINDEX](#) defined_by_kernel_header_file
Get CAN interface index by name.
- #define [SIOCSCANBAUDRATE](#) _IOW(RTIOC_TYPE_CAN, 0x01, struct ifreq)
Set baud rate.
- #define [SIOCGCANBAUDRATE](#) _IOWR(RTIOC_TYPE_CAN, 0x02, struct ifreq)
Get baud rate.
- #define [SIOCSCANCUSTOMBITTIME](#) _IOW(RTIOC_TYPE_CAN, 0x03, struct ifreq)
Set custom bit time parameter.
- #define [SIOCGCANCUSTOMBITTIME](#) _IOWR(RTIOC_TYPE_CAN, 0x04, struct ifreq)
Get custom bit-time parameters.
- #define [SIOCSCANMODE](#) _IOW(RTIOC_TYPE_CAN, 0x05, struct ifreq)
Set operation mode of CAN controller.
- #define [SIOCGCANSTATE](#) _IOWR(RTIOC_TYPE_CAN, 0x06, struct ifreq)
Get current state of CAN controller.
- #define [SIOCSCANCTRLMODE](#) _IOW(RTIOC_TYPE_CAN, 0x07, struct ifreq)
Set special controller modes.
- #define [SIOCGCANCTRLMODE](#) _IOWR(RTIOC_TYPE_CAN, 0x08, struct ifreq)
Get special controller modes.
- #define [RTCAN_RTIOC_TAKE_TIMESTAMP](#) _IOW(RTIOC_TYPE_CAN, 0x09, int)
Enable or disable storing a high precision timestamp upon reception of a CAN frame.
- #define [RTCAN_RTIOC_RCV_TIMEOUT](#) _IOW(RTIOC_TYPE_CAN, 0x0A, nanosecs_rel_t)
Specify a reception timeout for a socket.
- #define [RTCAN_RTIOC_SND_TIMEOUT](#) _IOW(RTIOC_TYPE_CAN, 0x0B, nanosecs_rel_t)
Specify a transmission timeout for a socket.

Error mask

Error class (mask) in `can_id` field of struct [can_frame](#) to be used with [CAN_RAW_ERR_FILTER](#).

Note: Error reporting is hardware dependent and most CAN controllers report less detailed error conditions than the SJA1000.

Note: In case of a bus-off error condition ([CAN_ERR_BUSOFF](#)), the CAN controller is **not** restarted automatically. It is the application's responsibility to react appropriately, e.g. calling [CAN_MODE_START](#).

Note: Bus error interrupts ([CAN_ERR_BUSERROR](#)) are enabled when an application is calling a [Recv](#) function on a socket listening on bus errors (using [CAN_RAW_ERR_FILTER](#)). After one bus error has occurred, the interrupt will be disabled to allow the application time for error processing and to efficiently avoid bus error interrupt flooding.

- #define [CAN_ERR_TX_TIMEOUT](#) 0x00000001U
TX timeout (netdevice driver)
- #define [CAN_ERR_LOSTARB](#) 0x00000002U
Lost arbitration (see [data\[0\]](#))
- #define [CAN_ERR_CRTL](#) 0x00000004U
Controller problems (see [data\[1\]](#))
- #define [CAN_ERR_PROT](#) 0x00000008U
Protocol violations (see [data\[2\]](#), [data\[3\]](#))
- #define [CAN_ERR_TRX](#) 0x00000010U
Transceiver status (see [data\[4\]](#))
- #define [CAN_ERR_ACK](#) 0x00000020U
Received no ACK on transmission.
- #define [CAN_ERR_BUSOFF](#) 0x00000040U
Bus off.

- #define `CAN_ERR_BUSERROR` 0x00000080U
Bus error (may flood!)
- #define `CAN_ERR_RESTARTED` 0x00000100U
Controller restarted.
- #define `CAN_ERR_MASK` 0x1FFFFFFFU
Omit EFF, RTR, ERR flags.

Arbitration lost error

Error in the `data[0]` field of struct `can_frame`.

- #define `CAN_ERR_LOSTARB_UNSPEC` 0x00
unspecified

Controller problems

Error in the `data[1]` field of struct `can_frame`.

- #define `CAN_ERR_CRTL_UNSPEC` 0x00
unspecified
- #define `CAN_ERR_CRTL_RX_OVERFLOW` 0x01
RX buffer overflow.
- #define `CAN_ERR_CRTL_TX_OVERFLOW` 0x02
TX buffer overflow.
- #define `CAN_ERR_CRTL_RX_WARNING` 0x04
reached warning level for RX errors
- #define `CAN_ERR_CRTL_TX_WARNING` 0x08
reached warning level for TX errors
- #define `CAN_ERR_CRTL_RX_PASSIVE` 0x10
reached passive level for RX errors
- #define `CAN_ERR_CRTL_TX_PASSIVE` 0x20
reached passive level for TX errors

Protocol error type

Error in the `data[2]` field of struct `can_frame`.

- #define `CAN_ERR_PROT_UNSPEC` 0x00
unspecified
- #define `CAN_ERR_PROT_BIT` 0x01
single bit error
- #define `CAN_ERR_PROT_FORM` 0x02
frame format error
- #define `CAN_ERR_PROT_STUFF` 0x04
bit stuffing error
- #define `CAN_ERR_PROT_BIT0` 0x08
unable to send dominant bit
- #define `CAN_ERR_PROT_BIT1` 0x10
unable to send recessive bit
- #define `CAN_ERR_PROT_OVERLOAD` 0x20
bus overload
- #define `CAN_ERR_PROT_ACTIVE` 0x40
active error announcement
- #define `CAN_ERR_PROT_TX` 0x80
error occurred on transmission

Protocol error location

Error in the `data[4]` field of struct `can_frame`.

- #define [CAN_ERR_PROT_LOC_UNSPEC](#) 0x00
unspecified
- #define [CAN_ERR_PROT_LOC_SOF](#) 0x03
start of frame
- #define [CAN_ERR_PROT_LOC_ID28_21](#) 0x02
ID bits 28 - 21 (SFF: 10 - 3)
- #define [CAN_ERR_PROT_LOC_ID20_18](#) 0x06
ID bits 20 - 18 (SFF: 2 - 0)
- #define [CAN_ERR_PROT_LOC_SRTR](#) 0x04
substitute RTR (SFF: RTR)
- #define [CAN_ERR_PROT_LOC_IDE](#) 0x05
identifier extension
- #define [CAN_ERR_PROT_LOC_ID17_13](#) 0x07
ID bits 17-13.
- #define [CAN_ERR_PROT_LOC_ID12_05](#) 0x0F
ID bits 12-5.
- #define [CAN_ERR_PROT_LOC_ID04_00](#) 0x0E
ID bits 4-0.
- #define [CAN_ERR_PROT_LOC_RTR](#) 0x0C
RTR.
- #define [CAN_ERR_PROT_LOC_RES1](#) 0x0D
reserved bit 1
- #define [CAN_ERR_PROT_LOC_RES0](#) 0x09
reserved bit 0
- #define [CAN_ERR_PROT_LOC_DLC](#) 0x0B
data length code
- #define [CAN_ERR_PROT_LOC_DATA](#) 0x0A
data section
- #define [CAN_ERR_PROT_LOC_CRC_SEQ](#) 0x08
CRC sequence.
- #define [CAN_ERR_PROT_LOC_CRC_DEL](#) 0x18
CRC delimiter.
- #define [CAN_ERR_PROT_LOC_ACK](#) 0x19
ACK slot.
- #define [CAN_ERR_PROT_LOC_ACK_DEL](#) 0x1B
ACK delimiter.
- #define [CAN_ERR_PROT_LOC_EOF](#) 0x1A
end of frame
- #define [CAN_ERR_PROT_LOC_INTERM](#) 0x12
intermission
- #define [CAN_ERR_TRX_UNSPEC](#) 0x00
0000 0000
- #define [CAN_ERR_TRX_CANH_NO_WIRE](#) 0x04
0000 0100
- #define [CAN_ERR_TRX_CANH_SHORT_TO_BAT](#) 0x05
0000 0101
- #define [CAN_ERR_TRX_CANH_SHORT_TO_VCC](#) 0x06
0000 0110
- #define [CAN_ERR_TRX_CANH_SHORT_TO_GND](#) 0x07
0000 0111
- #define [CAN_ERR_TRX_CANL_NO_WIRE](#) 0x40
0100 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_BAT](#) 0x50
0101 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_VCC](#) 0x60
0110 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_GND](#) 0x70
0111 0000
- #define [CAN_ERR_TRX_CANL_SHORT_TO_CANH](#) 0x80
1000 0000

Typedefs

- typedef uint32_t [can_id_t](#)
Type of CAN id (see [CAN_xxx_MASK](#) and [CAN_xxx_FLAG](#))
- typedef [can_id_t](#) [can_err_mask_t](#)
Type of CAN error mask.
- typedef uint32_t [can_baudrate_t](#)
Baudrate definition in bits per second.
- typedef enum [CAN_BITTIME_TYPE](#) [can_bittime_type_t](#)
See [CAN_BITTIME_TYPE](#).
- typedef enum [CAN_MODE](#) [can_mode_t](#)
See [CAN_MODE](#).
- typedef int [can_ctrlmode_t](#)
See [CAN_CTRLMODE](#).
- typedef enum [CAN_STATE](#) [can_state_t](#)
See [CAN_STATE](#).
- typedef struct [can_filter](#) [can_filter_t](#)
Filter for reception of CAN messages.
- typedef struct [can_frame](#) [can_frame_t](#)
Raw CAN frame.

Enumerations

- enum [CAN_BITTIME_TYPE](#) { [CAN_BITTIME_STD](#), [CAN_BITTIME_BTR](#) }
Supported CAN bit-time types.

CAN operation modes

Modes into which CAN controllers can be set

- enum [CAN_MODE](#) { [CAN_MODE_STOP](#) = 0, [CAN_MODE_START](#), [CAN_MODE_SLEEP](#) }

CAN controller states

States a CAN controller can be in.

- enum [CAN_STATE](#) {
[CAN_STATE_ERROR_ACTIVE](#) = 0, [CAN_STATE_ACTIVE](#) = 0, [CAN_STATE_ERROR_WARNING](#) = 1, [CAN_STATE_BUS_WARNING](#) = 1,
[CAN_STATE_ERROR_PASSIVE](#) = 2, [CAN_STATE_BUS_PASSIVE](#) = 2, [CAN_STATE_BUS_OFF](#), [CAN_STATE_SCANNING_BAUDRATE](#),
[CAN_STATE_STOPPED](#), [CAN_STATE_SLEEPING](#) }

7.21.1 Detailed Description

Real-Time Driver Model for RT-Socket-CAN, CAN device profile header.

Note

Copyright (C) 2006 Wolfgang Grandegger wg@grandegger.com

Copyright (C) 2005, 2006 Sebastian Smolorz Sebastian.Smolorz@stud.uni-hannover.de

This RTDM CAN device profile header is based on:

include/linux/can.h, include/linux/socket.h, net/can/pf_can.h in linux-can.patch, a CAN socket framework for Linux

Copyright (C) 2004, 2005, Robert Schwebel, Benedikt Spranger, Marc Kleine-Budde, Pengutronix

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU

General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.22 include/cobalt/kernel/rtdm/cobalt.h File Reference

This file is part of the Xenomai project.

Include dependency graph for cobalt.h:



7.22.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2013 Philippe Gerum rpm@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

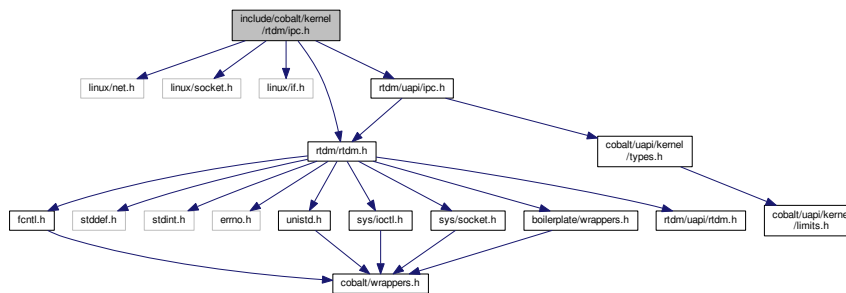
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.23 include/cobalt/kernel/rtdm/ipc.h File Reference

This file is part of the Xenomai project.

Include dependency graph for ipc.h:



7.23.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2009 Philippe Gerum rpm@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

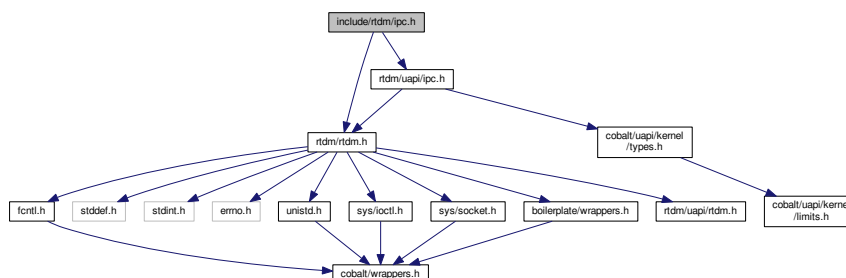
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.24 include/rtdm/ipc.h File Reference

This file is part of the Xenomai project.

Include dependency graph for ipc.h:



7.24.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2009 Philippe Gerum rpm@xenomai.org

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.


7.25 include/rtdm/uapi/ipc.h File Reference

This file is part of the Xenomai project.

Include dependency graph for ipc.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [rtipc_port_label](#)
Port label information structure.
- struct [sockaddr_ipc](#)
Socket address structure for the RTIPC address family.

Macros

XDDP socket options

Setting and getting XDDP socket options.

- #define [XDDP_LABEL](#) 1
XDDP label assignment.
- #define [XDDP_POOLSZ](#) 2
XDDP local pool size configuration.
- #define [XDDP_BUFSZ](#) 3
XDDP streaming buffer size configuration.
- #define [XDDP_MONITOR](#) 4
XDDP monitoring callback.

XDDP events

Specific events occurring on XDDP channels, which can be monitored via the [XDDP_MONITOR](#) socket option.

- #define [XDDP_EVTIN](#) 1
Monitor writes to the non real-time endpoint.
- #define [XDDP_EVTOUT](#) 2
Monitor reads from the non real-time endpoint.
- #define [XDDP_EVTDOWN](#) 3
Monitor close from the non real-time endpoint.
- #define [XDDP_EVTNOBUF](#) 4
Monitor memory shortage for non real-time datagrams.

IDDP socket options

Setting and getting IDDP socket options.

- #define [IDDP_LABEL](#) 1
IDDP label assignment.
- #define [IDDP_POOLSZ](#) 2
IDDP local pool size configuration.

BUFP socket options

Setting and getting BUFP socket options.

- #define [BUFP_LABEL](#) 1
BUFP label assignment.
- #define [BUFP_BUFSZ](#) 2
BUFP buffer size configuration.

Socket level options

Setting and getting supported standard socket level options.

- #define [SO_SNDBTIMEO](#) defined_by_kernel_header_file
IPPROTO_IDDP and IPPROTO_BUFP protocols support the standard SO_SNDBTIMEO socket option, from the SOL_SOCKET level.
- #define [SO_RCVTIMEO](#) defined_by_kernel_header_file
All RTIPC protocols support the standard SO_RCVTIMEO socket option, from the SOL_SOCKET level.

Typedefs

- typedef int16_t [rtipc_port_t](#)
Port number type for the RTIPC address family.

Enumerations

RTIPC protocol list

protocols for the PF_RTIPC protocol family

- enum { `IPCPROTO_IPC` = 0, `IPCPROTO_XDDP` = 1, `IPCPROTO_IDDP` = 2, `IPCPROTO_BUFP` = 3 }

Functions

Supported operations

Standard socket operations supported by the RTIPC protocols.

- int `socket__AF_RTIPC` (int domain=AF_RTIPC, int type=SOCK_DGRAM, int protocol)
Create an endpoint for communication in the AF_RTIPC domain.
- int `close__AF_RTIPC` (int sockfd)
Close a RTIPC socket descriptor.
- int `bind__AF_RTIPC` (int sockfd, const struct `sockaddr_ipc` *addr, socklen_t addrlen)
Bind a RTIPC socket to a port.
- int `connect__AF_RTIPC` (int sockfd, const struct `sockaddr_ipc` *addr, socklen_t addrlen)
Initiate a connection on a RTIPC socket.
- int `setsockopt__AF_RTIPC` (int sockfd, int level, int optname, const void *optval, socklen_t optlen)
Set options on RTIPC sockets.
- int `getsockopt__AF_RTIPC` (int sockfd, int level, int optname, void *optval, socklen_t *optlen)
Get options on RTIPC sockets.
- ssize_t `sendmsg__AF_RTIPC` (int sockfd, const struct `msghdr` *msg, int flags)
Send a message on a RTIPC socket.
- ssize_t `recvmsg__AF_RTIPC` (int sockfd, struct `msghdr` *msg, int flags)
Receive a message from a RTIPC socket.
- int `getsockname__AF_RTIPC` (int sockfd, struct `sockaddr_ipc` *addr, socklen_t *addrlen)
Get socket name.
- int `getpeername__AF_RTIPC` (int sockfd, struct `sockaddr_ipc` *addr, socklen_t *addrlen)
Get socket peer.

7.25.1 Detailed Description

This file is part of the Xenomai project.

Note


Copyright (C) 2009 Philippe Gerum rpm@xenomai.org

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

This graph shows which files directly or indirectly include this file:



7.27.1 Detailed Description

Note

Copyright (C) 2005, 2006 Jan Kiszka jan.kiszka@web.de
 Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

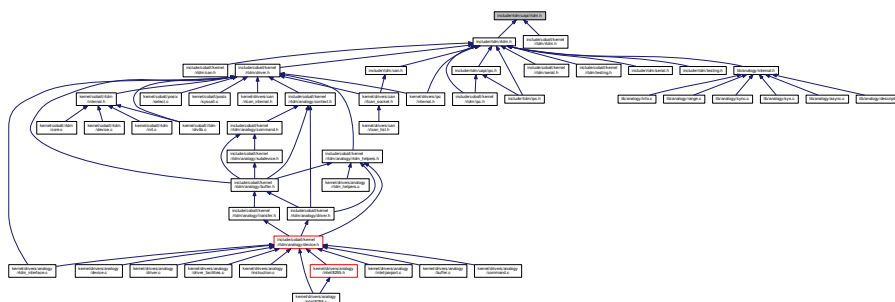
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.28 include/rtdm/uapi/rtdm.h File Reference

Real-Time Driver Model for Xenomai, user API header.

This graph shows which files directly or indirectly include this file:



Data Structures

- struct [rtdm_device_info](#)
Device information.

Macros

API Versioning

- #define `RTDM_API_VER` 9
Common user and driver API version.
- #define `RTDM_API_MIN_COMPAT_VER` 9
Minimum API revision compatible with the current release.

RTDM_TIMEOUT_xxx

Special timeout values

- #define `RTDM_TIMEOUT_INFINITE` 0
Block forever.
- #define `RTDM_TIMEOUT_NONE` (-1)
Any negative timeout means non-blocking.

RTDM_CLASS_xxx

Device classes

- #define `RTDM_CLASS_PARPORT` 1
- #define `RTDM_CLASS_SERIAL` 2
- #define `RTDM_CLASS_CAN` 3
- #define `RTDM_CLASS_NETWORK` 4
- #define `RTDM_CLASS_RTMAC` 5
- #define `RTDM_CLASS_TESTING` 6
- #define `RTDM_CLASS_RTIPC` 7
- #define `RTDM_CLASS_COBALT` 8
- #define `RTDM_CLASS_EXPERIMENTAL` 224
- #define `RTDM_CLASS_MAX` 255

Device Naming

Maximum length of device names (excluding the final null character)

- #define `RTDM_MAX_DEVNAME_LEN` 31

RTDM_PURGE_xxx_BUFFER

Flags selecting buffers to be purged

- #define `RTDM_PURGE_RX_BUFFER` 0x0001
- #define `RTDM_PURGE_TX_BUFFER` 0x0002

Common IOCTLs

The following IOCTLs are common to all device rtdm_profiles.

- #define `RTIOC_DEVICE_INFO` _IOR(RTIOC_TYPE_COMMON, 0x00, struct rtdm_device_info)
Retrieve information about a device or socket.
- #define `RTIOC_PURGE` _IOW(RTIOC_TYPE_COMMON, 0x10, int)
Purge internal device or socket buffers.

Typedefs

- typedef uint64_t `nanosecs_abs_t`
RTDM type for representing absolute dates.
- typedef int64_t `nanosecs_rel_t`
RTDM type for representing relative intervals.
- typedef struct `rtdm_device_info` `rtdm_device_info_t`
Device information.

7.28.1 Detailed Description

Real-Time Driver Model for Xenomai, user API header.

Note

Copyright (C) 2005, 2006 Jan Kiszka jan.kiszka@web.de
 Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net


This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.29 include/cobalt/kernel/rtdm/serial.h File Reference

Include dependency graph for serial.h:



7.29.1 Detailed Description

Note

Copyright (C) 2005-2007 Jan Kiszka jan.kiszka@web.de

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.30 include/rtdm/serial.h File Reference

Real-Time Driver Model for Xenomai, serial device profile header.

Include dependency graph for serial.h:



7.30.1 Detailed Description

Real-Time Driver Model for Xenomai, serial device profile header.

Note

Copyright (C) 2005-2007 Jan Kiszka jan.kiszka@web.de

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.31 include/rtdm/uapi/serial.h File Reference

Real-Time Driver Model for Xenomai, serial device profile header.

This graph shows which files directly or indirectly include this file:



Data Structures

- struct [rtser_config](#)
Serial device configuration.
- struct [rtser_status](#)
Serial device status.
- struct [rtser_event](#)
Additional information about serial device events.

Macros

- #define [RTSER_RTIOC_BREAK_CTL](#) _IOR(RTIOC_TYPE_SERIAL, 0x06, int)
Set or clear break on UART output line.

RTSER_DEF_BAUD

Default baud rate

- #define **RTSER_DEF_BAUD** 9600

RTSER_xxx_PARITY

Number of parity bits

- #define **RTSER_NO_PARITY** 0x00
- #define **RTSER_ODD_PARITY** 0x01
- #define **RTSER_EVEN_PARITY** 0x03
- #define **RTSER_DEF_PARITY** RTSER_NO_PARITY

RTSER_xxx_BITS

Number of data bits

- #define **RTSER_5_BITS** 0x00
- #define **RTSER_6_BITS** 0x01
- #define **RTSER_7_BITS** 0x02
- #define **RTSER_8_BITS** 0x03
- #define **RTSER_DEF_BITS** RTSER_8_BITS

RTSER_xxx_STOPB

Number of stop bits

- #define [RTSER_1_STOPB](#) 0x00
valid only in combination with 5 data bits
- #define [RTSER_1_5_STOPB](#) 0x01
valid only in combination with 5 data bits
- #define [RTSER_2_STOPB](#) 0x01
valid only in combination with 5 data bits
- #define [RTSER_DEF_STOPB](#) [RTSER_1_STOPB](#)
valid only in combination with 5 data bits

RTSER_xxx_HAND

Handshake mechanisms

- #define **RTSER_NO_HAND** 0x00
- #define **RTSER_RTCTS_HAND** 0x01
- #define **RTSER_DEF_HAND** RTSER_NO_HAND

RTSER_RS485_xxx*RS485 mode with automatic RTS handling*

- #define **RTSER_RS485_DISABLE** 0x00
- #define **RTSER_RS485_ENABLE** 0x01
- #define **RTSER_DEF_RS485** RTSER_RS485_DISABLE

RTSER_FIFO_xxx*Reception FIFO interrupt threshold*

- #define **RTSER_FIFO_DEPTH_1** 0x00
- #define **RTSER_FIFO_DEPTH_4** 0x40
- #define **RTSER_FIFO_DEPTH_8** 0x80
- #define **RTSER_FIFO_DEPTH_14** 0xC0
- #define **RTSER_DEF_FIFO_DEPTH** RTSER_FIFO_DEPTH_1

RTSER_TIMEOUT_xxx*Special timeout values, see also [RTDM_TIMEOUT_xxx](#)*

- #define **RTSER_TIMEOUT_INFINITE** [RTDM_TIMEOUT_INFINITE](#)
- #define **RTSER_TIMEOUT_NONE** [RTDM_TIMEOUT_NONE](#)
- #define **RTSER_DEF_TIMEOUT** [RTDM_TIMEOUT_INFINITE](#)

RTSER_xxx_TIMESTAMP_HISTORY*Timestamp history control*

- #define **RTSER_RX_TIMESTAMP_HISTORY** 0x01
- #define **RTSER_DEF_TIMESTAMP_HISTORY** 0x00

RTSER_EVENT_xxx*Events bits*

- #define **RTSER_EVENT_RXPEND** 0x01
- #define **RTSER_EVENT_ERRPEND** 0x02
- #define **RTSER_EVENT_MODEMHI** 0x04
- #define **RTSER_EVENT_MODEMLO** 0x08
- #define **RTSER_EVENT_TXEMPTY** 0x10
- #define **RTSER_DEF_EVENT_MASK** 0x00

RTSER_SET_xxx*Configuration mask bits*

- #define **RTSER_SET_BAUD** 0x0001
- #define **RTSER_SET_PARITY** 0x0002
- #define **RTSER_SET_DATA_BITS** 0x0004
- #define **RTSER_SET_STOP_BITS** 0x0008
- #define **RTSER_SET_HANDSHAKE** 0x0010
- #define **RTSER_SET_FIFO_DEPTH** 0x0020
- #define **RTSER_SET_TIMEOUT_RX** 0x0100
- #define **RTSER_SET_TIMEOUT_TX** 0x0200
- #define **RTSER_SET_TIMEOUT_EVENT** 0x0400
- #define **RTSER_SET_TIMESTAMP_HISTORY** 0x0800
- #define **RTSER_SET_EVENT_MASK** 0x1000
- #define **RTSER_SET_RS485** 0x2000

RTSER_LSR_xxx*Line status bits*

- #define **RTSER_LSR_DATA** 0x01
- #define **RTSER_LSR_OVERRUN_ERR** 0x02

- #define **RTSER_LSR_PARITY_ERR** 0x04
- #define **RTSER_LSR_FRAMING_ERR** 0x08
- #define **RTSER_LSR_BREAK_IND** 0x10
- #define **RTSER_LSR_THR_EMPTY** 0x20
- #define **RTSER_LSR_TRANSM_EMPTY** 0x40
- #define **RTSER_LSR_FIFO_ERR** 0x80
- #define **RTSER_SOFT_OVERRUN_ERR** 0x0100

RTSER_MSR_xxx

Modem status bits

- #define **RTSER_MSR_DCTS** 0x01
- #define **RTSER_MSR_DDSD** 0x02
- #define **RTSER_MSR_TERI** 0x04
- #define **RTSER_MSR_DDCD** 0x08
- #define **RTSER_MSR_CTS** 0x10
- #define **RTSER_MSR_DSR** 0x20
- #define **RTSER_MSR_RI** 0x40
- #define **RTSER_MSR_DCD** 0x80

RTSER_MCR_xxx

Modem control bits

- #define **RTSER_MCR_DTR** 0x01
- #define **RTSER_MCR_RTS** 0x02
- #define **RTSER_MCR_OUT1** 0x04
- #define **RTSER_MCR_OUT2** 0x08
- #define **RTSER_MCR_LOOP** 0x10

Sub-Classes of RTDM_CLASS_SERIAL

- #define **RTDM_SUBCLASS_16550A** 0

IOCTLs

Serial device IOCTLs

- #define **RTSER_RTIOC_GET_CONFIG** _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_config)
Get serial device configuration.
- #define **RTSER_RTIOC_SET_CONFIG** _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_config)
Set serial device configuration.
- #define **RTSER_RTIOC_GET_STATUS** _IOR(RTIOC_TYPE_SERIAL, 0x02, struct rtser_status)
Get serial device status.
- #define **RTSER_RTIOC_GET_CONTROL** _IOR(RTIOC_TYPE_SERIAL, 0x03, int)
Get serial device's modem control register.
- #define **RTSER_RTIOC_SET_CONTROL** _IOW(RTIOC_TYPE_SERIAL, 0x04, int)
Set serial device's modem control register.
- #define **RTSER_RTIOC_WAIT_EVENT** _IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_event)
Wait on serial device events according to previously set mask.

RTSER_BREAK_XXX

Break control

- #define [RTSER_BREAK_CLR](#) 0x00
Serial device configuration.
- #define [RTSER_BREAK_SET](#) 0x01
Serial device configuration.
- #define [RTIOC_TYPE_SERIAL](#) RTDM_CLASS_SERIAL
Serial device configuration.
- typedef struct [rtser_config](#) [rtser_config_t](#)
Serial device configuration.
- typedef struct [rtser_status](#) [rtser_status_t](#)
Serial device status.
- typedef struct [rtser_event](#) [rtser_event_t](#)
Additional information about serial device events.

7.31.1 Detailed Description

Real-Time Driver Model for Xenomai, serial device profile header.

Note

Copyright (C) 2005-2007 Jan Kiszka jan.kiszka@web.de

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.31.2 Macro Definition Documentation

7.31.2.1 #define RTSER_RTIOC_BREAK_CTL_IOR([RTIOC_TYPE_SERIAL](#), 0x06, int)

Set or clear break on UART output line.

Parameters

in	arg	RTSER_BREAK_SET or RTSER_BREAK_CLR (int)
-----------	------------	--

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code

- Kernel-based task
- User-space task (RT, non-RT)

Note

A set break condition may also be cleared on UART line reconfiguration.

Rescheduling: never.

7.31.2.2 `#define RTSER_RTIOC_GET_CONFIG _IOR(RTIOC_TYPE_SERIAL, 0x00, struct rtser_config)`

Get serial device configuration.

Parameters

out	arg	Pointer to configuration buffer (struct rtser_config)
-----	-----	--

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

7.31.2.3 `#define RTSER_RTIOC_GET_CONTROL _IOR(RTIOC_TYPE_SERIAL, 0x03, int)`

Get serial device's modem control register.

Parameters

out	arg	Pointer to variable receiving the content (int, see RTSER_MCR_xxx)
-----	-----	---

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

7.31.2.4 `#define RTSER_RTIOC_GET_STATUS _IOR(RTIOC_TYPE_SERIAL, 0x02, struct
rtser_status)`

Get serial device status.

Parameters

out	arg	Pointer to status buffer (struct rtser_status)
-----	-----	---

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

The error states `RTSER_LSR_OVERRUN_ERR`, `RTSER_LSR_PARITY_ERR`, `RTSER_LSR_FRAMING_ERR`, and `RTSER_SOFT_OVERRUN_ERR` that may have occurred during previous read accesses to the device will be saved for being reported via this IOCTL. Upon return from `RTSER_RTIOC_GET_STATUS`, the saved state will be cleared.

Rescheduling: never.

7.31.2.5 `#define RTSER_RTIOC_SET_CONFIG _IOW(RTIOC_TYPE_SERIAL, 0x01, struct rtser_config)`

Set serial device configuration.

Parameters

in	arg	Pointer to configuration buffer (struct rtser_config)
----	-----	--

Returns

0 on success, otherwise:

- `-EPERM` is returned if the caller's context is invalid, see note below.
- `-ENOMEM` is returned if a new history buffer for timestamps cannot be allocated.

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Note

If [rtser_config](#) contains a valid `timestamp_history` and the addressed device has been opened in non-real-time context, this IOCTL must be issued in non-real-time context as well. Otherwise, this command will fail.

Rescheduling: never.

Examples:

[cross-link.c](#).

7.31.2.6 `#define RTSER_RTIOC_SET_CONTROL_IOW(RTIOC_TYPE_SERIAL, 0x04, int)`

Set serial device's modem control register.

Parameters

<code>in</code>	<code>arg</code>	New control register content (int, see RTSER_MCR_xxx)
-----------------	------------------	--

Returns

0 on success, otherwise negative error code

Environments:

This service can be called from:

- Kernel module initialization/cleanup code
- Kernel-based task
- User-space task (RT, non-RT)

Rescheduling: never.

7.31.2.7 `#define RTSER_RTIOC_WAIT_EVENT_IOR(RTIOC_TYPE_SERIAL, 0x05, struct rtser_event)`

Wait on serial device events according to previously set mask.

Parameters

<code>out</code>	<code>arg</code>	Pointer to event information buffer (struct rtser_event)
------------------	------------------	---

Returns

0 on success, otherwise:

- -EBUSY is returned if another task is already waiting on events of this device.
- -EBADF is returned if the file descriptor is invalid or the device has just been closed.

Environments:

This service can be called from:

- Kernel-based task

- User-space task (RT)


Rescheduling: possible.

Examples:

[cross-link.c](#).

7.32 include/cobalt/kernel/rtdm/testing.h File Reference

Include dependency graph for testing.h:



7.32.1 Detailed Description

Note

Copyright (C) 2005 Jan Kiszka jan.kiszka@web.de

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.33 include/rtdm/testing.h File Reference

Real-Time Driver Model for Xenomai, testing device profile header.

Include dependency graph for testing.h:



7.33.1 Detailed Description

Real-Time Driver Model for Xenomai, testing device profile header.

Note

Copyright (C) 2005 Jan Kiszka jan.kiszka@web.de

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.34 include/rtdm/uapi/testing.h File Reference

Real-Time Driver Model for Xenomai, testing device profile header.

This graph shows which files directly or indirectly include this file:



Macros

Sub-Classes of RTDM_CLASS_TESTING

- #define [RTDM_SUBCLASS_TIMERBENCH](#) 0
subclass name: "timerbench"
- #define [RTDM_SUBCLASS_IRQBENCH](#) 1
subclass name: "irqbench"
- #define [RTDM_SUBCLASS_SWITCHTEST](#) 2
subclass name: "switchtest"
- #define [RTDM_SUBCLASS_RTDMTTEST](#) 3
subclass name: "rtdm"

IOCTLs

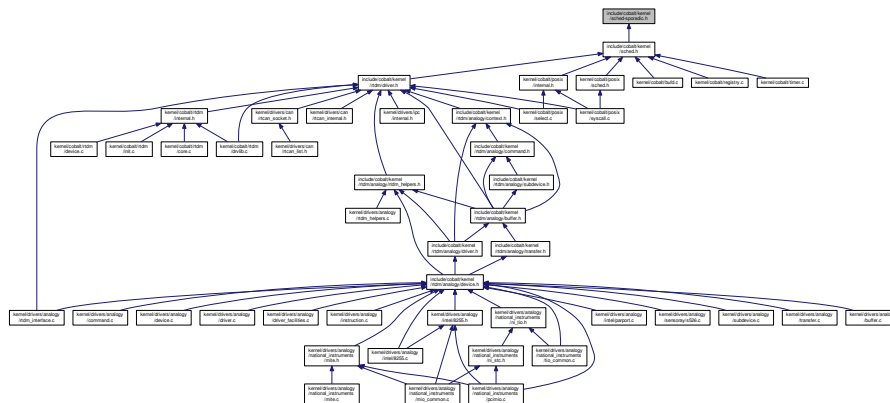
Testing device IOCTLs

- #define [RTTST_RTIOC_INTERM_BENCH_RES](#) _IOW(RTIOC_TYPE_TESTING, 0x00, struct rttst_interm_bench_res)
- #define [RTTST_RTIOC_TMBENCH_START](#) _IOW(RTIOC_TYPE_TESTING, 0x10, struct rttst_tmbench_config)
- #define [RTTST_RTIOC_TMBENCH_STOP](#) _IOW(RTIOC_TYPE_TESTING, 0x11, struct rttst_overall_bench_res)
- #define [RTTST_RTIOC_IRQBENCH_START](#) _IOW(RTIOC_TYPE_TESTING, 0x20, struct rttst_irqbench_config)
- #define [RTTST_RTIOC_IRQBENCH_STOP](#) _IO(RTIOC_TYPE_TESTING, 0x21)
- #define [RTTST_RTIOC_IRQBENCH_GET_STATS](#) _IOR(RTIOC_TYPE_TESTING, 0x22, struct rttst_irqbench_stats)
- #define [RTTST_RTIOC_IRQBENCH_WAIT_IRQ](#) _IO(RTIOC_TYPE_TESTING, 0x23)
- #define [RTTST_RTIOC_IRQBENCH_REPLY_IRQ](#) _IO(RTIOC_TYPE_TESTING, 0x24)
- #define [RTTST_RTIOC_SWTEST_SET_TASKS_COUNT](#) _IOW(RTIOC_TYPE_TESTING, 0x30, unsigned long)
- #define [RTTST_RTIOC_SWTEST_SET_CPU](#) _IOW(RTIOC_TYPE_TESTING, 0x31, unsigned long)
- #define [RTTST_RTIOC_SWTEST_REGISTER_UTASK](#) _IOW(RTIOC_TYPE_TESTING, 0x32, struct rttst_swtest_task)
- #define [RTTST_RTIOC_SWTEST_CREATE_KTASK](#) _IOW(RTIOC_TYPE_TESTING, 0x33, struct rttst_swtest_task)
- #define [RTTST_RTIOC_SWTEST_PEND](#) _IOR(RTIOC_TYPE_TESTING, 0x34, struct rttst_swtest_task)
- #define [RTTST_RTIOC_SWTEST_SWITCH_TO](#) _IOR(RTIOC_TYPE_TESTING, 0x35, struct rttst_swtest_dir)
- #define [RTTST_RTIOC_SWTEST_GET_SWITCHES_COUNT](#) _IOR(RTIOC_TYPE_TESTING, 0x36, unsigned long)
- #define [RTTST_RTIOC_SWTEST_GET_LAST_ERROR](#) _IOR(RTIOC_TYPE_TESTING, 0x37, struct rttst_swtest_error)
- #define [RTTST_RTIOC_SWTEST_SET_PAUSE](#) _IOW(RTIOC_TYPE_TESTING, 0x38, unsigned long)
- #define [RTTST_RTIOC_RTDMT_DEFER_CLOSE](#) _IOW(RTIOC_TYPE_TESTING, 0x40, unsigned long)

7.34.1 Detailed Description

Real-Time Driver Model for Xenomai, testing device profile header.

This graph shows which files directly or indirectly include this file:



7.37.1 Detailed Description

Definitions for the SSP scheduling class.

Author

Philippe Gerum

Copyright (C) 2009 Philippe Gerum rpm@xenomai.org.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

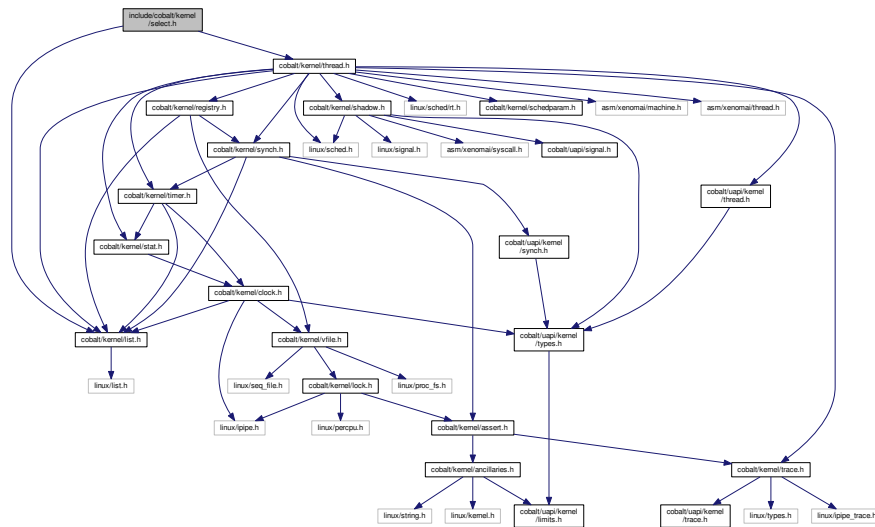
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

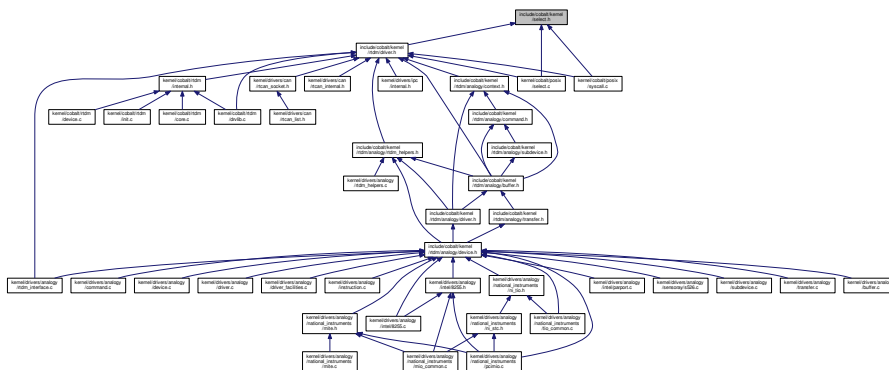
7.38 include/cobalt/kernel/select.h File Reference

file descriptors events multiplexing header.

Include dependency graph for select.h:



This graph shows which files directly or indirectly include this file:



Functions

- void `xnselect_init` (struct `xnselect` *select_block)
Initialize a struct xnselect structure.
- static int `xnselect_signal` (struct `xnselect` *select_block, unsigned int state)
Signal a file descriptor state change.
- void `xnselect_destroy` (struct `xnselect` *select_block)
Destroy the xnselect structure associated with a file descriptor.
- int `xnselector_init` (struct `xnselector` *selector)
Initialize a selector structure.
- int `xnselect` (struct `xnselector` *selector, fd_set *out_fds[XNSELECT_MAX_TYPES], fd_set *in_fds[XNSELECT_MAX_TYPES], int nfds, xnticks_t timeout, xntmode_t timeout_mode)
Check the state of a number of file descriptors, wait for a state change if no descriptor is ready.
- void `xnselector_destroy` (struct `xnselector` *selector)
Destroy a selector block.

7.38.1 Detailed Description

file descriptors events multiplexing header.

Author

Gilles Chanteperdrix

Copyright (C) 2008 Efixo gilles.chanteperdrix@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

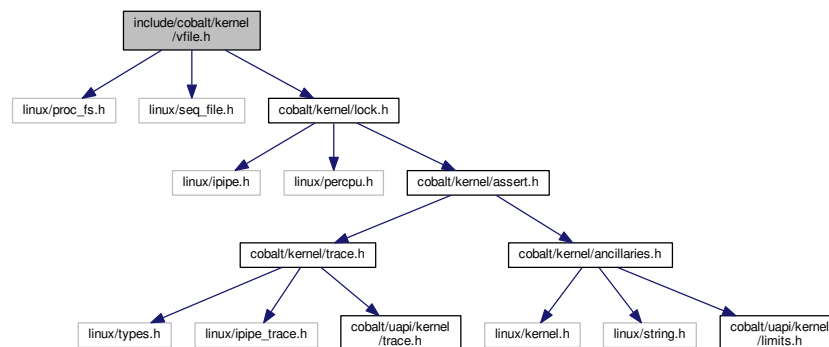
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

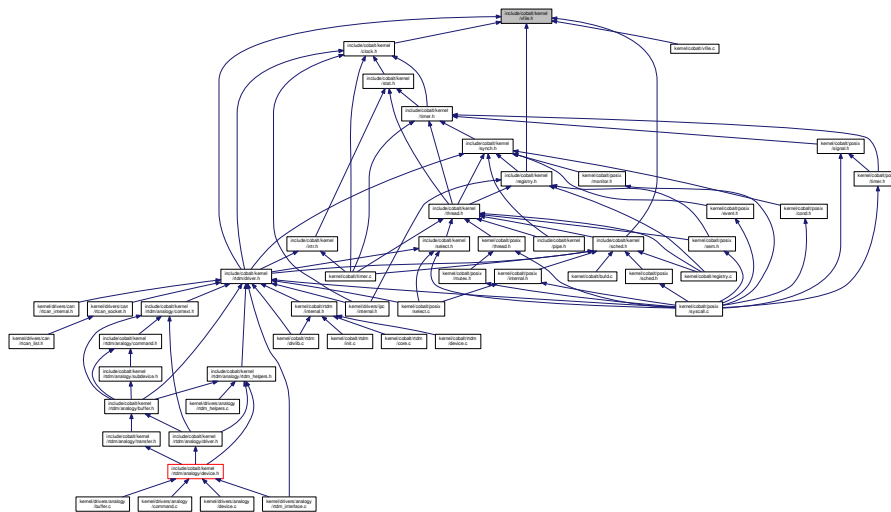
7.39 include/cobalt/kernel/vfile.h File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [xnfile_lock_ops](#)
Vfile locking operations.
- struct [xnfile_regular_ops](#)
Regular vfile operation descriptor.
- struct [xnfile_regular_iterator](#)
Regular vfile iterator.
- struct [xnfile_snapshot_ops](#)
Snapshot vfile operation descriptor.
- struct [xnfile_rev_tag](#)
Snapshot revision tag.
- struct [xnfile_snapshot](#)
Snapshot vfile descriptor.
- struct [xnfile_snapshot_iterator](#)
Snapshot-driven vfile iterator.

Functions

- int [xnfile_init_snapshot](#) (const char *name, struct [xnfile_snapshot](#) *vfile, struct [xnfile_directory](#) *parent)
Initialize a snapshot-driven vfile.
- int [xnfile_init_regular](#) (const char *name, struct [xnfile_regular](#) *vfile, struct [xnfile_directory](#) *parent)
Initialize a regular vfile.
- int [xnfile_init_dir](#) (const char *name, struct [xnfile_directory](#) *vdir, struct [xnfile_directory](#) *parent)
Initialize a virtual directory entry.
- int [xnfile_init_link](#) (const char *from, const char *to, struct [xnfile_link](#) *vlink, struct [xnfile_directory](#) *parent)
Initialize a virtual link entry.
- void [xnfile_destroy](#) (struct [xnfile](#) *vfile)
Removes a virtual file entry.

- `ssize_t xnvfile_get_blob` (struct xnvfile_input *input, void *data, size_t size)
Read in a data bulk written to the vfile.
- `ssize_t xnvfile_get_string` (struct xnvfile_input *input, char *s, size_t maxlen)
Read in a C-string written to the vfile.
- `ssize_t xnvfile_get_integer` (struct xnvfile_input *input, long *valp)
Evaluate the string written to the vfile as a long integer.

Variables

- struct xnvfile_directory `nkvrroot`
Xenomai vfile root directory.

7.39.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2010 Philippe Gerum rpm@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


7.40 include/rtdm/analogy.h File Reference

Analogy for Linux, library facilities.

Include dependency graph for analogy.h:



This graph shows which files directly or indirectly include this file:



Data Structures

- struct [a4l_descriptor](#)

Structure containing device-information useful to users.

Macros

ANALOGY_xxx_DESC

Constants used as argument so as to define the description depth to recover

- #define [A4L_BSC_DESC](#) 0x0
BSC stands for basic descriptor (device data)
- #define [A4L_CPLX_DESC](#) 0x1
CPLX stands for complex descriptor (subdevice + channel + range data)

7.40.1 Detailed Description

Analogy for Linux, library facilities.

Note

Copyright (C) 1997-2000 David A. Schleeef ds@schleeef.org
 Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.41 include/rtdm/uapi/analogy.h File Reference

Analogy for Linux, UAPI bits.

This graph shows which files directly or indirectly include this file:



Data Structures

- struct [a4l_cmd_desc](#)
Structure describing the asynchronous instruction.
- struct [a4l_instruction](#)
Structure describing the synchronous instruction.
- struct [a4l_instruction_list](#)
Structure describing the list of synchronous instructions.

Macros

- `#define A4L_RNG_FACTOR 1000000`
Constant for internal use only (must not be used by driver developer).
- `#define A4L_RNG_VOLT_UNIT 0x0`
Volt unit range flag.
- `#define A4L_RNG_MAMP_UNIT 0x1`
MilliAmpere unit range flag.
- `#define A4L_RNG_NO_UNIT 0x2`
No unit range flag.
- `#define A4L_RNG_EXT_UNIT 0x4`
External unit range flag.
- `#define A4L_RNG_UNIT(x)`
Macro to retrieve the range unit from the range flags.
- `#define A4L_INSN_WAIT_MAX 100000`
Maximal wait duration.

ANALOGY_CMD_xxx

Common command flags definitions

- `#define A4L_CMD_SIMUL 0x1`
Do not execute the command, just check it.
- `#define A4L_CMD_BULK 0x2`
Perform data recovery / transmission in bulk mode.
- `#define A4L_CMD_WRITE 0x4`
Perform a command which will write data to the device.

TRIG_xxx

Command triggers flags definitions

- #define **TRIG_NONE** 0x00000001
Never trigger.
- #define **TRIG_NOW** 0x00000002
Trigger now + N ns.
- #define **TRIG_FOLLOW** 0x00000004
Trigger on next lower level trig.
- #define **TRIG_TIME** 0x00000008
Trigger at time N ns.
- #define **TRIG_TIMER** 0x00000010
Trigger at rate N ns.
- #define **TRIG_COUNT** 0x00000020
Trigger when count reaches N.
- #define **TRIG_EXT** 0x00000040
Trigger on external signal N.
- #define **TRIG_INT** 0x00000080
Trigger on analogy-internal signal N.
- #define **TRIG_OTHER** 0x00000100
Driver defined trigger.
- #define **TRIG_WAKE_EOS** 0x0020
Wake up on end-of-scan.
- #define **TRIG_ROUND_MASK** 0x00030000
Trigger not implemented yet.
- #define **TRIG_ROUND_NEAREST** 0x00000000
Trigger not implemented yet.
- #define **TRIG_ROUND_DOWN** 0x00010000
Trigger not implemented yet.
- #define **TRIG_ROUND_UP** 0x00020000
Trigger not implemented yet.
- #define **TRIG_ROUND_UP_NEXT** 0x00030000
Trigger not implemented yet.

Channel macros

Specific precompilation macros and constants useful for the channels descriptors tab located in the command structure

- #define **CHAN**(a) ((a) & 0xffff)
Channel indication macro.
- #define **RNG**(a) (((a) & 0xff) << 16)
Range definition macro.
- #define **AREF**(a) (((a) & 0xf) << 24)
Reference definition macro.
- #define **FLAGS**(a) ((a) & CR_FLAGS_MASK)
Flags definition macro.
- #define **PACK**(a, b, c) (**CHAN**(a) | **RNG**(b) | **AREF**(c))
Channel + range + reference definition macro.
- #define **PACK_FLAGS**(a, b, c, d) (**CHAN**(a) | **RNG**(b) | **AREF**(c) | **FLAGS**(d))
Channel + range + reference + flags definition macro.
- #define **AREF_GROUND** 0x00
Analog reference is analog ground.
- #define **AREF_COMMON** 0x01
Analog reference is analog common.
- #define **AREF_DIFF** 0x02
Analog reference is differential.
- #define **AREF_OTHER** 0x03
Analog reference is undefined.

Subdevices types

Flags to define the subdevice type

- #define **A4L_SUBD_UNUSED** (A4L_SUBD_MASK_SPECIAL|0x1)
Unused subdevice.
- #define **A4L_SUBD_AI** (A4L_SUBD_MASK_READ|0x2)
Analog input subdevice.
- #define **A4L_SUBD_AO** (A4L_SUBD_MASK_WRITE|0x4)
Analog output subdevice.
- #define **A4L_SUBD_DI** (A4L_SUBD_MASK_READ|0x8)
Digital input subdevice.
- #define **A4L_SUBD_DO** (A4L_SUBD_MASK_WRITE|0x10)
Digital output subdevice.
- #define **A4L_SUBD_DIO** (A4L_SUBD_MASK_SPECIAL|0x20)
Digital input/output subdevice.
- #define **A4L_SUBD_COUNTER** (A4L_SUBD_MASK_SPECIAL|0x40)
Counter subdevice.
- #define **A4L_SUBD_TIMER** (A4L_SUBD_MASK_SPECIAL|0x80)
Timer subdevice.
- #define **A4L_SUBD_MEMORY** (A4L_SUBD_MASK_SPECIAL|0x100)
Memory, EEPROM, DPRAM.
- #define **A4L_SUBD_CALIB** (A4L_SUBD_MASK_SPECIAL|0x200)
Calibration subdevice DACs.
- #define **A4L_SUBD_PROC** (A4L_SUBD_MASK_SPECIAL|0x400)
Processor, DSP.
- #define **A4L_SUBD_SERIAL** (A4L_SUBD_MASK_SPECIAL|0x800)
Serial IO subdevice.
- #define **A4L_SUBD_TYPES**
Mask which gathers all the types.

Subdevice features

Flags to define the subdevice's capabilities

- #define **A4L_SUBD_CMD** 0x1000
The subdevice can handle command (i.e it can perform asynchronous acquisition)
- #define **A4L_SUBD_MMAP** 0x8000
The subdevice support mmap operations (technically, any driver can do it; however, the developer might want that his driver must be accessed through read / write.

Subdevice status

Flags to define the subdevice's status

- #define **A4L_SUBD_BUSY_NR** 0
The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.
- #define **A4L_SUBD_BUSY** (1 << **A4L_SUBD_BUSY_NR**)
The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.
- #define **A4L_SUBD_CLEAN_NR** 1
The subdevice is about to be cleaned in the middle of the detach procedure.
- #define **A4L_SUBD_CLEAN** (1 << **A4L_SUBD_CLEAN_NR**)
The subdevice is busy, a synchronous or an asynchronous acquisition is occurring.

Instruction type

Flags to define the type of instruction

- #define **A4L_INSN_READ** (0 | A4L_INSN_MASK_READ)

- Read instruction.*
- #define **A4L_INSN_WRITE** (1 | A4L_INSN_MASK_WRITE)
- Write instruction.*
- #define **A4L_INSN_BITS**
- "Bits" instruction*
- #define **A4L_INSN_CONFIG**
- Configuration instruction.*
- #define **A4L_INSN_GTOD**
- Get time instruction.*
- #define **A4L_INSN_WAIT**
- Wait instruction.*
- #define **A4L_INSN_INTTRIG**
- Trigger instruction (to start asynchronous acquisition)*

Configuration instruction type

Values to define the type of configuration instruction

- #define **A4L_INSN_CONFIG_DIO_INPUT** 0
- #define **A4L_INSN_CONFIG_DIO_OUTPUT** 1
- #define **A4L_INSN_CONFIG_DIO_OPENDRAIN** 2
- #define **A4L_INSN_CONFIG_ANALOG_TRIG** 16
- #define **A4L_INSN_CONFIG_ALT_SOURCE** 20
- #define **A4L_INSN_CONFIG_DIGITAL_TRIG** 21
- #define **A4L_INSN_CONFIG_BLOCK_SIZE** 22
- #define **A4L_INSN_CONFIG_TIMER_1** 23
- #define **A4L_INSN_CONFIG_FILTER** 24
- #define **A4L_INSN_CONFIG_CHANGE_NOTIFY** 25
- #define **A4L_INSN_CONFIG_SERIAL_CLOCK** 26
- #define **A4L_INSN_CONFIG_BIDIRECTIONAL_DATA** 27
- #define **A4L_INSN_CONFIG_DIO_QUERY** 28
- #define **A4L_INSN_CONFIG_PWM_OUTPUT** 29
- #define **A4L_INSN_CONFIG_GET_PWM_OUTPUT** 30
- #define **A4L_INSN_CONFIG_ARM** 31
- #define **A4L_INSN_CONFIG_DISARM** 32
- #define **A4L_INSN_CONFIG_GET_COUNTER_STATUS** 33
- #define **A4L_INSN_CONFIG_RESET** 34
- #define **A4L_INSN_CONFIG_GPCT_SINGLE_PULSE_GENERATOR** 1001 /* Use CTR as single pulse generator */
- #define **A4L_INSN_CONFIG_GPCT_PULSE_TRAIN_GENERATOR** 1002 /* Use CTR as pulsetraingenerator */
- #define **A4L_INSN_CONFIG_GPCT_QUADRATURE_ENCODER** 1003 /* Use the counter as encoder */
- #define **A4L_INSN_CONFIG_SET_GATE_SRC** 2001 /* Set gate source */
- #define **A4L_INSN_CONFIG_GET_GATE_SRC** 2002 /* Get gate source */
- #define **A4L_INSN_CONFIG_SET_CLOCK_SRC** 2003 /* Set master clock source */
- #define **A4L_INSN_CONFIG_GET_CLOCK_SRC** 2004 /* Get master clock source */
- #define **A4L_INSN_CONFIG_SET_OTHER_SRC** 2005 /* Set other source */
- #define **A4L_INSN_CONFIG_SET_COUNTER_MODE** 4097
- #define **A4L_INSN_CONFIG_SET_ROUTING** 4099
- #define **A4L_INSN_CONFIG_GET_ROUTING** 4109

Counter status bits

Status bits for **INSN_CONFIG_GET_COUNTER_STATUS**

- #define **A4L_COUNTER_ARMED** 0x1
- #define **A4L_COUNTER_COUNTING** 0x2
- #define **A4L_COUNTER_TERMINAL_COUNT** 0x4

IO direction

Values to define the IO polarity

- `#define A4L_INPUT 0`
- `#define A4L_OUTPUT 1`
- `#define A4L_OPENDRAIN 2`

Events types

Values to define the Analogy events. They might used to send some specific events through the instruction interface.

- `#define A4L_EV_START 0x00040000`
- `#define A4L_EV_SCAN_BEGIN 0x00080000`
- `#define A4L_EV_CONVERT 0x00100000`
- `#define A4L_EV_SCAN_END 0x00200000`
- `#define A4L_EV_STOP 0x00400000`

7.41.1 Detailed Description

Analogy for Linux, UAPI bits.

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.41.2 Macro Definition Documentation


7.41.2.1 `#define A4L_RNG_FACTOR 1000000`

Constant for internal use only (must not be used by driver developer).

Referenced by `a4l_dtoraw()`, `a4l_find_range()`, `a4l_ftoraw()`, `a4l_rawtod()`, and `a4l_rawtof()`.

7.42 kernel/cobalt/apc.c File Reference

Include dependency graph for apc.c:



Functions

- `int xnapc_alloc` (`const char *name`, `void(*handler)(void *cookie)`, `void *cookie`)
Allocate an APC slot.
- `void xnapc_free` (`int apc`)
Releases an APC slot.

7.42.1 Detailed Description

Note

Copyright (C) 2007,2012 Philippe Gerum rpm@xenomai.org.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.43 kernel/cobalt/posix/syscall.c File Reference

This file is part of the Xenomai project.

Include dependency graph for syscall.c:



7.43.1 Detailed Description

This file is part of the Xenomai project. Copyright (C) 2005 Philippe Gerum rpm@xenomai.org Copyright (C) 2005 Gilles Chantepedrix gilles.chantepedrix@xenomai.org


This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.44 kernel/cobalt/arith.c File Reference

Include dependency graph for arith.c:



7.44.1 Detailed Description

Note

Copyright © 2005 Gilles Chantepedrix.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

7.45 kernel/cobalt/bufd.c File Reference

[illegible]

- `ssize_t xnbufd_copy_to_kmem` (void *ptr, struct xnbufd *bufd, size_t len)
Copy memory covered by a buffer descriptor to kernel memory.
- `ssize_t xnbufd_copy_from_kmem` (struct xnbufd *bufd, void *from, size_t len)
Copy kernel memory to the area covered by a buffer descriptor.
- `ssize_t xnbufd_unmap_uread` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uread\(\)](#).
- `ssize_t xnbufd_unmap_uwrite` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_uwrite\(\)](#).
- `void xnbufd_invalidate` (struct xnbufd *bufd)
Invalidate a buffer descriptor.
- `ssize_t xnbufd_unmap_kread` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kread\(\)](#).
- `ssize_t xnbufd_unmap_kwrite` (struct xnbufd *bufd)
Finalize a buffer descriptor obtained from [xnbufd_map_kwrite\(\)](#).


Copyright (C) 2009 Philippe Gerum rpm@xenomai.org.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.46 kernel/cobalt/rtdm/init.c File Reference

Real-Time Driver Model for Xenomai.

Include dependency graph for init.c:



7.46.1 Detailed Description

Real-Time Driver Model for Xenomai.

Note

Copyright (C) 2005, 2006 Jan Kiszka jan.kiszka@web.de

Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net

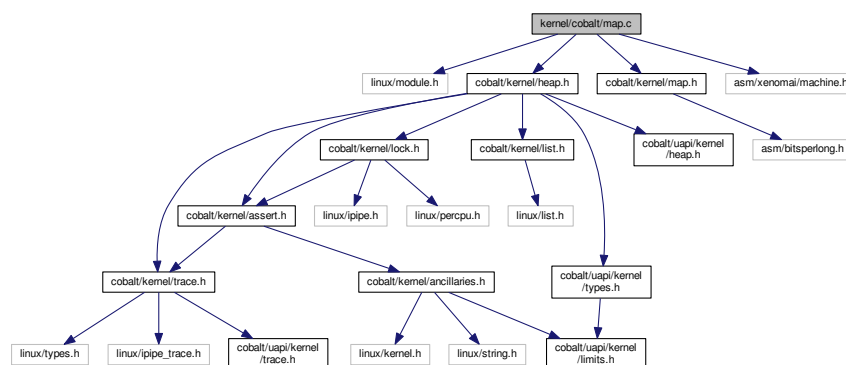
Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.47 kernel/cobalt/map.c File Reference

Include dependency graph for map.c:



Functions

- struct xnmap * [xnmap_create](#) (int nkeys, int reserve, int offset)
Create a map.
- void [xnmap_delete](#) (struct xnmap *map)
Delete a map.
- int [xnmap_enter](#) (struct xnmap *map, int key, void *objaddr)
Index an object into a map.
- int [xnmap_remove](#) (struct xnmap *map, int key)
Remove an object reference from a map.

7.47.1 Detailed Description

Note

Copyright (C) 2007 Philippe Gerum rpm@xenomai.org.

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


7.48 lib/analogy/internal.h File Reference

Analogy for Linux, internal declarations.

Include dependency graph for internal.h:



This graph shows which files directly or indirectly include this file:



7.48.1 Detailed Description

Analogy for Linux, internal declarations.

Note

Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

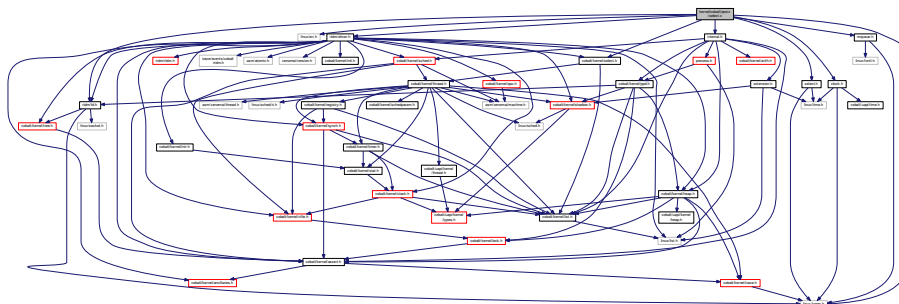
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.49 kernel/cobalt/posix/select.c File Reference

This file is part of the Xenomai project.

Include dependency graph for select.c:



7.49.1 Detailed Description

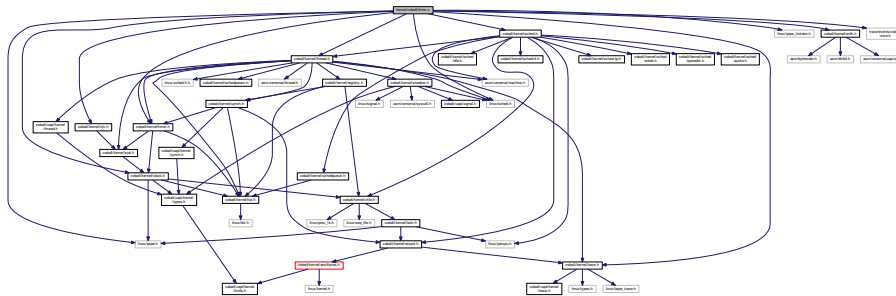
This file is part of the Xenomai project. Copyright (C) 2005 Philippe Gerum rpm@xenomai.org Copyright (C) 2005 Gilles Chanteperdrix gilles.chanteperdrix@xenomai.org

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Include dependency graph for timer.c:



- int [xntimer_start](#) (struct xntimer *timer, xnticks_t value, xnticks_t interval, xntmode_t mode)
Arm a timer.
- xnticks_t [xntimer_get_date](#) (struct xntimer *timer)
Return the absolute expiration date.
- xnticks_t [xntimer_get_timeout](#) (struct xntimer *timer)
Return the relative expiration date.
- void [xntimer_init](#) (struct xntimer *timer, struct xnclock *clock, void(*handler)(struct xntimer *timer), struct xnthread *thread)
Initialize a timer object.
- void [xntimer_destroy](#) (struct xntimer *timer)
Release a timer object.
- unsigned long long [xntimer_get_overruns](#) (struct xntimer *timer, xnticks_t now)
Get the count of overruns for the last tick.
- static int [program_htick_shot](#) (unsigned long delay, struct clock_event_device *cdev)
Program next host tick as a Xenomai timer event.
- static void [switch_htick_mode](#) (enum clock_event_mode mode, struct clock_event_device *cdev)
Tick mode switch emulation callback.
- int [xntimer_grab_hardware](#) (int cpu)
Grab the hardware timer.
- void [xntimer_release_hardware](#) (int cpu)
Release the hardware timer.

Note


Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Generated on Fri Jun 13 2014 12:09:58 for Xenomai by Doxygen

7.52 kernel/cobalt/rtdm/core.c File Reference

Real-Time Driver Model for Xenomai, device operation multiplexing.

Include dependency graph for core.c:



Functions

- void [rtdm_context_lock](#) (struct [rtdm_dev_context](#) *context)
Increment context reference counter.
- void [rtdm_context_unlock](#) (struct [rtdm_dev_context](#) *context)
Decrement context reference counter.
- void [rtdm_context_put](#) (struct [rtdm_dev_context](#) *context)
Release a device context obtained via [rtdm_context_get\(\)](#)
- void [rtdm_waitqueue_init](#) (struct [rtdm_waitqueue](#) *wq)
Initialize a RTDM wait queue.
- void [rtdm_waitqueue_destroy](#) (struct [rtdm_waitqueue](#) *wq)
Deletes a RTDM wait queue.
- [rtdm_timedwait_condition_locked](#) (struct [rtdm_wait_queue](#) *wq, C_expr condition, [nanosecs_rel_t](#) timeout, [rtdm_toseq_t](#) *toseq)
Timed sleep on a locked waitqueue until a condition gets true.
- [rtdm_wait_condition_locked](#) (struct [rtdm_wait_queue](#) *wq, C_expr condition)
Sleep on a locked waitqueue until a condition gets true.
- [rtdm_timedwait_condition](#) (struct [rtdm_wait_queue](#) *wq, C_expr condition, [nanosecs_rel_t](#) timeout, [rtdm_toseq_t](#) *toseq)
Timed sleep on a waitqueue until a condition gets true.
- void [rtdm_timedwait](#) (struct [rtdm_wait_queue](#) *wq, [nanosecs_rel_t](#) timeout, [rtdm_toseq_t](#) *toseq)
Timed sleep on a waitqueue unconditionally.
- void [rtdm_timedwait_locked](#) (struct [rtdm_wait_queue](#) *wq, [nanosecs_rel_t](#) timeout, [rtdm_toseq_t](#) *toseq)
Timed sleep on a locked waitqueue unconditionally.
- [rtdm_wait_condition](#) (struct [rtdm_wait_queue](#) *wq, C_expr condition)
Sleep on a waitqueue until a condition gets true.
- void [rtdm_wait](#) (struct [rtdm_wait_queue](#) *wq)
Sleep on a waitqueue unconditionally.
- void [rtdm_wait_locked](#) (struct [rtdm_wait_queue](#) *wq)
Sleep on a locked waitqueue unconditionally.
- void [rtdm_waitqueue_lock](#) (struct [rtdm_wait_queue](#) *wq, [rtdm_lockctx_t](#) context)
Lock a waitqueue.
- void [rtdm_waitqueue_unlock](#) (struct [rtdm_wait_queue](#) *wq, [rtdm_lockctx_t](#) context)
Unlock a waitqueue.
- void [rtdm_waitqueue_signal](#) (struct [rtdm_wait_queue](#) *wq)

- *Signal a waitqueue.*
- void [rt dm_waitqueue_broadcast](#) (struct rtdm_wait_queue *wq)
- *Broadcast a waitqueue.*
- void [rt dm_waitqueue_flush](#) (struct rtdm_wait_queue *wq)
- *Flush a waitqueue.*
- void [rt dm_waitqueue_wakeup](#) (struct rtdm_wait_queue *wq, rtdm_task_t waiter)
- *Signal a particular waiter on a waitqueue.*
- [rt dm_for_each_waiter](#) (rtdm_task_t pos, struct rtdm_wait_queue *wq)
- *Simple iterator for waitqueues.*
- [rt dm_for_each_waiter_safe](#) (rtdm_task_t pos, rtdm_task_t tmp, struct rtdm_wait_queue *wq)
- *Safe iterator for waitqueues.*
- int [rt dm_open](#) (const char *path, int oflag,...)
- *Open a device.*
- int [rt dm_socket](#) (int protocol_family, int socket_type, int protocol)
- *Create a socket.*
- int [rt dm_close](#) (int fd)
- *Close a device or socket.*
- int [rt dm_ioctl](#) (int fd, int request,...)
- *Issue an IOCTL.*
- ssize_t [rt dm_read](#) (int fd, void *buf, size_t nbyte)
- *Read from device.*
- ssize_t [rt dm_write](#) (int fd, const void *buf, size_t nbyte)
- *Write to device.*
- ssize_t [rt dm_recvm sg](#) (int fd, struct msghdr *msg, int flags)
- *Receive message from socket.*
- ssize_t [rt dm_recvfrom](#) (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)
- *Receive message from socket.*
- ssize_t [rt dm_recv](#) (int fd, void *buf, size_t len, int flags)
- *Receive message from socket.*
- ssize_t [rt dm_sendmsg](#) (int fd, const struct msghdr *msg, int flags)
- *Transmit message to socket.*
- ssize_t [rt dm_sendto](#) (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)
- *Transmit message to socket.*
- ssize_t [rt dm_send](#) (int fd, const void *buf, size_t len, int flags)
- *Transmit message to socket.*
- int [rt dm_bind](#) (int fd, const struct sockaddr *my_addr, socklen_t addrlen)
- *Bind to local address.*
- int [rt dm_connect](#) (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)
- *Connect to remote address.*
- int [rt dm_listen](#) (int fd, int backlog)
- *Listen for incoming connection requests.*
- int [rt dm_accept](#) (int fd, struct sockaddr *addr, socklen_t *addrlen)
- *Accept a connection request.*
- int [rt dm_shutdown](#) (int fd, int how)
- *Shut down parts of a connection.*
- int [rt dm_getsockopt](#) (int fd, int level, int optname, void *optval, socklen_t *optlen)
- *Get socket option.*
- int [rt dm_setsockopt](#) (int fd, int level, int optname, const void *optval, socklen_t optlen)
- *Set socket option.*

- `int rtdm_getsockname` (int fd, struct sockaddr *name, socklen_t *namelen)
Get local socket address.
- `int rtdm_getpeername` (int fd, struct sockaddr *name, socklen_t *namelen)
Get socket destination address.
- `int rt_dev_open` (const char *path, int oflag,...)
Open a device.
- `int rt_dev_socket` (int protocol_family, int socket_type, int protocol)
Create a socket.
- `int rt_dev_close` (int fd)
Close a device or socket.
- `int rt_dev_ioctl` (int fd, int request,...)
Issue an IOCTL.
- `ssize_t rt_dev_read` (int fd, void *buf, size_t nbyte)
Read from device.
- `ssize_t rt_dev_write` (int fd, const void *buf, size_t nbyte)
Write to device.
- `ssize_t rt_dev_recvmsg` (int fd, struct msghdr *msg, int flags)
Receive message from socket.
- `ssize_t rt_dev_recvfrom` (int fd, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)
Receive message from socket.
- `ssize_t rt_dev_recv` (int fd, void *buf, size_t len, int flags)
Receive message from socket.
- `ssize_t rt_dev_sendmsg` (int fd, const struct msghdr *msg, int flags)
Transmit message to socket.
- `ssize_t rt_dev_sendto` (int fd, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)
Transmit message to socket.
- `ssize_t rt_dev_send` (int fd, const void *buf, size_t len, int flags)
Transmit message to socket.
- `int rt_dev_bind` (int fd, const struct sockaddr *my_addr, socklen_t addrlen)
Bind to local address.
- `int rt_dev_connect` (int fd, const struct sockaddr *serv_addr, socklen_t addrlen)
Connect to remote address.
- `int rt_dev_listen` (int fd, int backlog)
Listen for incoming connection requests.
- `int rt_dev_accept` (int fd, struct sockaddr *addr, socklen_t *addrlen)
Accept a connection requests.
- `int rt_dev_shutdown` (int fd, int how)
Shut down parts of a connection.
- `int rt_dev_getsockopt` (int fd, int level, int optname, void *optval, socklen_t *optlen)
Get socket option.
- `int rt_dev_setsockopt` (int fd, int level, int optname, const void *optval, socklen_t optlen)
Set socket option.
- `int rt_dev_getsockname` (int fd, struct sockaddr *name, socklen_t *namelen)
Get local socket address.
- `int rt_dev_getpeername` (int fd, struct sockaddr *name, socklen_t *namelen)
Get socket destination address.

7.52.1 Detailed Description

Real-Time Driver Model for Xenomai, device operation multiplexing.

Note

Copyright (C) 2005 Jan Kiszka jan.kiszka@web.de

Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.53 kernel/cobalt/rtdm/device.c File Reference

Real-Time Driver Model for Xenomai, device management.

Include dependency graph for device.c:



Functions

- `int rtdm_dev_register (struct rtdm_device *device)`
Register a RTDM device.
- `int rtdm_dev_unregister (struct rtdm_device *device, unsigned int poll_delay)`
Unregisters a RTDM device.

7.53.1 Detailed Description

Real-Time Driver Model for Xenomai, device management.

Note

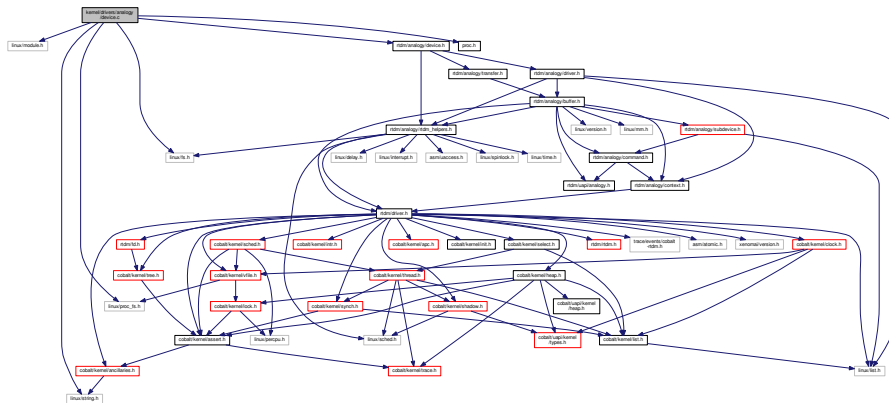
Copyright (C) 2005 Jan Kiszka jan.kiszka@web.de

Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

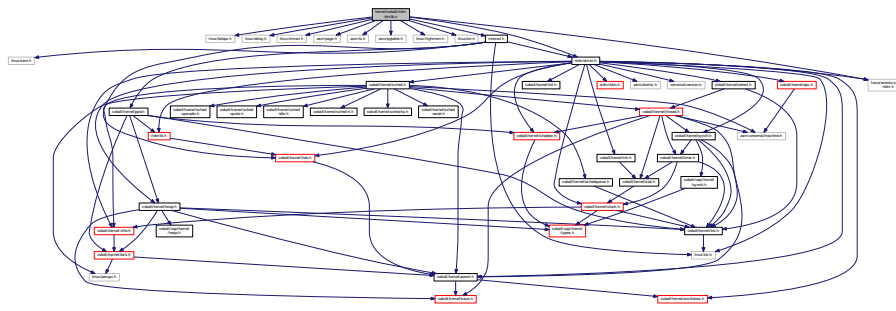
Include dependency graph for device.c:



You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Real-Time Driver Model for Xenomai, driver library.

Include dependency graph for `drvlib.c`:



Functions

- `nanosecs_abs_t rtdm_clock_read` (void)
Get system time.
- `nanosecs_abs_t rtdm_clock_read_monotonic` (void)
Get monotonic time.
- `int rtdm_task_init` (rtdm_task_t *task, const char *name, rtdm_task_proc_t task_proc, void *arg, int priority, nanosecs_rel_t period)
Initialise and start a real-time task.
- `void rtdm_task_destroy` (rtdm_task_t *task)
Destroy a real-time task.
- `int rtdm_task_should_stop` (void)
Check for pending termination request.
- `void rtdm_task_set_priority` (rtdm_task_t *task, int priority)
Adjust real-time task priority.
- `int rtdm_task_set_period` (rtdm_task_t *task, nanosecs_rel_t period)
Adjust real-time task period.
- `int rtdm_task_wait_period` (void)
Wait on next real-time task period.
- `int rtdm_task_unblock` (rtdm_task_t *task)
Activate a blocked real-time task.
- `rtdm_task_t * rtdm_task_current` (void)
Get current real-time task.
- `int rtdm_task_sleep` (nanosecs_rel_t delay)
Sleep a specified amount of time.
- `int rtdm_task_sleep_until` (nanosecs_abs_t wakeup_time)
Sleep until a specified absolute time.
- `int rtdm_task_sleep_abs` (nanosecs_abs_t wakeup_time, enum rtdm_timer_mode mode)
Sleep until a specified absolute time.
- `void rtdm_task_join_nrt` (rtdm_task_t *task, unsigned int poll_delay)
Wait on a real-time task to terminate.
- `void rtdm_task_busy_sleep` (nanosecs_rel_t delay)
Busy-wait a specified amount of time.
- `int rtdm_timer_init` (rtdm_timer_t *timer, rtdm_timer_handler_t handler, const char *name)
Initialise a timer.
- `void rtdm_timer_destroy` (rtdm_timer_t *timer)
Destroy a timer.

- int `rtdm_timer_start` (rtdm_timer_t *timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)
Start a timer.
- void `rtdm_timer_stop` (rtdm_timer_t *timer)
Stop a timer.
- int `rtdm_timer_start_in_handler` (rtdm_timer_t *timer, nanosecs_abs_t expiry, nanosecs_rel_t interval, enum rtdm_timer_mode mode)
Start a timer from inside a timer handler.
- void `rtdm_timer_stop_in_handler` (rtdm_timer_t *timer)
Stop a timer from inside a timer handler.
- void `rtdm_event_init` (rtdm_event_t *event, unsigned long pending)
Initialise an event.
- void `rtdm_event_destroy` (rtdm_event_t *event)
Destroy an event.
- void `rtdm_event_pulse` (rtdm_event_t *event)
Signal an event occurrence to currently listening waiters.
- void `rtdm_event_signal` (rtdm_event_t *event)
Signal an event occurrence.
- int `rtdm_event_wait` (rtdm_event_t *event)
Wait on event occurrence.
- int `rtdm_event_timedwait` (rtdm_event_t *event, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)
Wait on event occurrence with timeout.
- void `rtdm_event_clear` (rtdm_event_t *event)
Clear event state.
- int `rtdm_event_select_bind` (rtdm_event_t *event, rtdm_selector_t *selector, enum rtdm_selecttype type, unsigned int fd_index)
Bind a selector to an event.
- void `rtdm_sem_init` (rtdm_sem_t *sem, unsigned long value)
Initialise a semaphore.
- void `rtdm_sem_destroy` (rtdm_sem_t *sem)
Destroy a semaphore.
- int `rtdm_sem_down` (rtdm_sem_t *sem)
Decrement a semaphore.
- int `rtdm_sem_timeddown` (rtdm_sem_t *sem, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)
Decrement a semaphore with timeout.
- void `rtdm_sem_up` (rtdm_sem_t *sem)
Increment a semaphore.
- int `rtdm_sem_select_bind` (rtdm_sem_t *sem, rtdm_selector_t *selector, enum rtdm_selecttype type, unsigned int fd_index)
Bind a selector to a semaphore.
- void `rtdm_mutex_init` (rtdm_mutex_t *mutex)
Initialise a mutex.
- void `rtdm_mutex_destroy` (rtdm_mutex_t *mutex)
Destroy a mutex.
- void `rtdm_mutex_unlock` (rtdm_mutex_t *mutex)
Release a mutex.
- int `rtdm_mutex_lock` (rtdm_mutex_t *mutex)
Request a mutex.
- int `rtdm_mutex_timedlock` (rtdm_mutex_t *mutex, nanosecs_rel_t timeout, rtdm_toseq_t *timeout_seq)

- Request a mutex with timeout.*

 - int `rt dm_irq_request` (rt dm_irq_t *irq_handle, unsigned int irq_no, `rt dm_irq_handler_t` handler, unsigned long flags, const char *device_name, void *arg)
- Register an interrupt handler.*

 - int `rt dm_irq_free` (rt dm_irq_t *irq_handle)
- Release an interrupt handler.*

 - int `rt dm_irq_enable` (rt dm_irq_t *irq_handle)
- Enable interrupt line.*

 - int `rt dm_irq_disable` (rt dm_irq_t *irq_handle)
- Disable interrupt line.*

 - int `rt dm_nrt sig_init` (rt dm_nrt sig_t *nrt_sig, `rt dm_nrt sig_handler_t` handler, void *arg)
- Register a non-real-time signal handler.*

 - void `rt dm_nrt sig_destroy` (rt dm_nrt sig_t *nrt_sig)
- Release a non-realtime signal handler.*

 - void `rt dm_nrt sig_pend` (rt dm_nrt sig_t *nrt_sig)
- Trigger non-real-time signal.*

 - int `rt dm_mmap_to_user` (struct `rt dm_fd` *fd, void *src_addr, size_t len, int prot, void **pptr, struct `vm_operations_struct` *vm_ops, void *vm_private_data)
- Map a kernel memory range into the address space of the user.*

 - int `rt dm_iomap_to_user` (struct `rt dm_fd` *fd, phys_addr_t src_addr, size_t len, int prot, void **pptr, struct `vm_operations_struct` *vm_ops, void *vm_private_data)
- Map an I/O memory range into the address space of the user.*

 - int `rt dm_munmap` (struct `rt dm_fd` *fd, void *ptr, size_t len)
- Unmap a user memory range.*

 - int `rt dm_ratelimit` (struct `rt dm_ratelimit_state` *rs, const char *func)
- Enforces a rate limit.*

 - void `rt dm_printk_ratelimited` (const char *format,...)
- Real-time safe rate-limited message printing on kernel console.*

 - void `rt dm_printk` (const char *format,...)
- Real-time safe message printing on kernel console.*

 - void * `rt dm_malloc` (size_t size)
- Allocate memory block in real-time context.*

 - void `rt dm_free` (void *ptr)
- Release real-time memory block.*

 - int `rt dm_read_user_ok` (struct `rt dm_fd` *fd, const void __user *ptr, size_t size)
- Check if read access to user-space memory block is safe.*

 - int `rt dm_rw_user_ok` (struct `rt dm_fd` *fd, const void __user *ptr, size_t size)
- Check if read/write access to user-space memory block is safe.*

 - int `rt dm_copy_from_user` (struct `rt dm_fd` *fd, void *dst, const void __user *src, size_t size)
- Copy user-space memory block to specified buffer.*

 - int `rt dm_safe_copy_from_user` (struct `rt dm_fd` *fd, void *dst, const void __user *src, size_t size)
- Check if read access to user-space memory block and copy it to specified buffer.*

 - int `rt dm_copy_to_user` (struct `rt dm_fd` *fd, void __user *dst, const void *src, size_t size)
- Copy specified buffer to user-space memory block.*

 - int `rt dm_safe_copy_to_user` (struct `rt dm_fd` *fd, void __user *dst, const void *src, size_t size)
- Check if read/write access to user-space memory block is safe and copy specified buffer to it.*

 - int `rt dm_strncpy_from_user` (struct `rt dm_fd` *fd, char *dst, const char __user *src, size_t count)
- Copy user-space string to specified buffer.*

 - int `rt dm_in_rt_context` (void)
- Test if running in a real-time task.*

 - int `rt dm_rt_capable` (struct `rt dm_fd` *fd)

Test if the caller is capable of running in real-time context.

Timeout Sequence Management

- void `rtdm_toseq_init` (`rtdm_toseq_t *timeout_seq`, `nanosecs_rel_t timeout`)
Initialise a timeout sequence.

7.55.1 Detailed Description

Real-Time Driver Model for Xenomai, driver library.

Note

Copyright (C) 2005-2007 Jan Kiszka jan.kiszka@web.de
 Copyright (C) 2005 Joerg Langenberg joerg.langenberg@gmx.net
 Copyright (C) 2008 Gilles Chantepredrix gilles.chantepredrix@xenomai.org

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

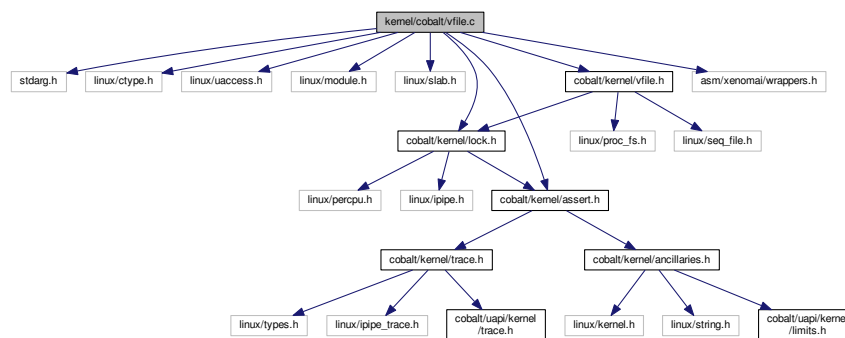
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.56 kernel/cobalt/vfile.c File Reference

This file is part of the Xenomai project.

Include dependency graph for vfile.c:



Functions

- int `xnvfile_init_snapshot` (`const char *name`, `struct xnvfile_snapshot *vfile`, `struct xnvfile_directory *parent`)
Initialize a snapshot-driven vfile.
- int `xnvfile_init_regular` (`const char *name`, `struct xnvfile_regular *vfile`, `struct xnvfile_directory *parent`)

Initialize a regular vfile.

- int [xnvmfile_init_dir](#) (const char *name, struct xnvmfile_directory *vdir, struct xnvmfile_directory *parent)

Initialize a virtual directory entry.

- int [xnvmfile_init_link](#) (const char *from, const char *to, struct xnvmfile_link *vlink, struct xnvmfile_directory *parent)

Initialize a virtual link entry.

- void [xnvmfile_destroy](#) (struct xnvmfile *vfile)

Removes a virtual file entry.

- ssize_t [xnvmfile_get_blob](#) (struct xnvmfile_input *input, void *data, size_t size)

Read in a data bulk written to the vfile.

- ssize_t [xnvmfile_get_string](#) (struct xnvmfile_input *input, char *s, size_t maxlen)

Read in a C-string written to the vfile.

- ssize_t [xnvmfile_get_integer](#) (struct xnvmfile_input *input, long *valp)

Evaluate the string written to the vfile as a long integer.

Variables

- struct xnvmfile_directory [nkvmfroot](#)

Xenomai vfile root directory.

7.56.1 Detailed Description

This file is part of the Xenomai project.

Note

Copyright (C) 2010 Philippe Gerum rpm@xenomai.org

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

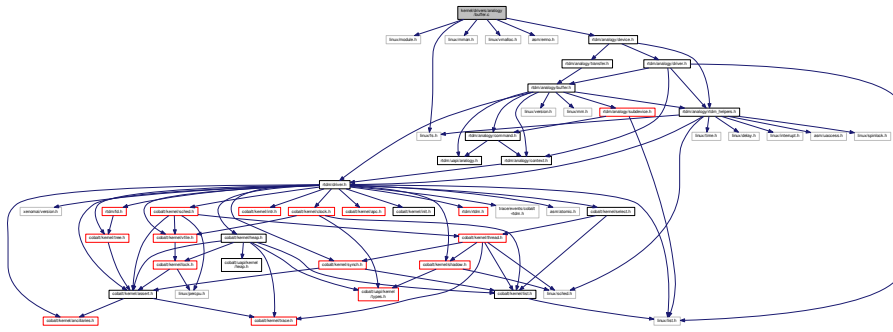
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.57 kernel/drivers/analog/buffer.c File Reference

Analogy for Linux, buffer related features.

Include dependency graph for buffer.c:



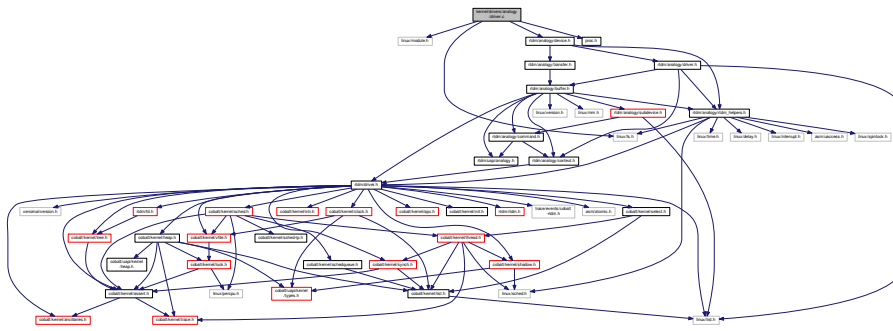
Functions

- `int a4l_get_chan (struct a4l_subdevice *subd)`
Get the channel index according to its type.
- `int a4l_buf_prepare_absput (struct a4l_subdevice *subd, unsigned long count)`
Update the absolute count of data sent from the device to the buffer since the start of the acquisition and after the next DMA shot.
- `int a4l_buf_commit_absput (struct a4l_subdevice *subd, unsigned long count)`
Set the absolute count of data which was sent from the device to the buffer since the start of the acquisition and until the last DMA shot.
- `int a4l_buf_prepare_put (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data which is to be sent to the buffer at the next DMA shot.
- `int a4l_buf_commit_put (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data sent to the buffer during the last completed DMA shots.
- `int a4l_buf_put (struct a4l_subdevice *subd, void *bufdata, unsigned long count)`
Copy some data from the device driver to the buffer.
- `int a4l_buf_prepare_absget (struct a4l_subdevice *subd, unsigned long count)`
Update the absolute count of data sent from the buffer to the device since the start of the acquisition and after the next DMA shot.
- `int a4l_buf_commit_absget (struct a4l_subdevice *subd, unsigned long count)`
Set the absolute count of data which was sent from the buffer to the device since the start of the acquisition and until the last DMA shot.
- `int a4l_buf_prepare_get (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data which is to be sent from the buffer to the device at the next DMA shot.
- `int a4l_buf_commit_get (struct a4l_subdevice *subd, unsigned long count)`
Set the count of data sent from the buffer to the device during the last completed DMA shots.
- `int a4l_buf_get (struct a4l_subdevice *subd, void *bufdata, unsigned long count)`
Copy some data from the buffer to the device driver.
- `int a4l_buf_evt (struct a4l_subdevice *subd, unsigned long evts)`
Signal some event(s) to a user-space program involved in some read / write operation.
- `unsigned long a4l_buf_count (struct a4l_subdevice *subd)`
Get the data amount available in the Analogy buffer.

7.57.1 Detailed Description

Analogy for Linux, buffer related features.

Include dependency graph for driver.c:



Functions

- `int a4l_register_drv (struct a4l_driver *drv)`
Register an Analogy driver.
- `int a4l_unregister_drv (struct a4l_driver *drv)`
Unregister an Analogy driver.

7.59.1 Detailed Description

Analogy for Linux, driver related features. Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.60 kernel/drivers/analogy/driver_facilities.c File Reference

Analogy for Linux, driver facilities.

Include dependency graph for driver_facilities.c:



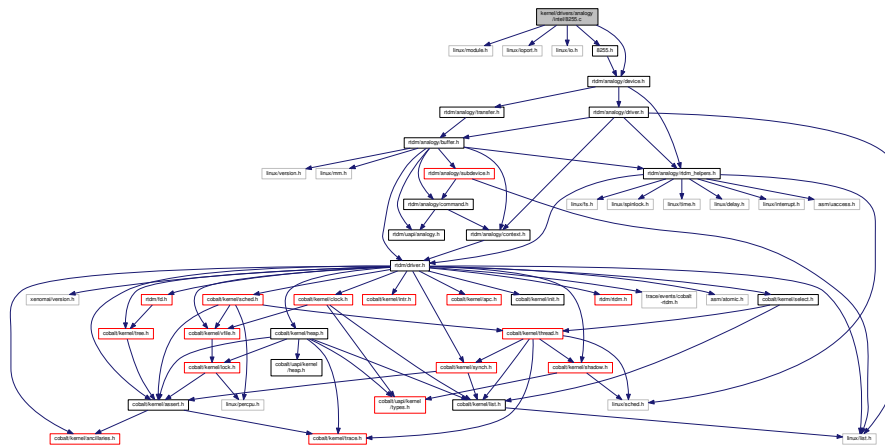
Functions

- int [a4l_register_drv](#) (struct [a4l_driver](#) *drv)
Register an Analogy driver.
- int [a4l_unregister_drv](#) (struct [a4l_driver](#) *drv)
Unregister an Analogy driver.
- struct [a4l_subdevice](#) * [a4l_alloc_subd](#) (int sizeof_priv, void(*setup)(struct [a4l_subdevice](#) *))
Allocate a subdevice descriptor.
- int [a4l_add_subd](#) (struct [a4l_device](#) *dev, struct [a4l_subdevice](#) *subd)
Add a subdevice to the driver descriptor.
- struct [a4l_subdevice](#) * [a4l_get_subd](#) (struct [a4l_device](#) *dev, int idx)
Get a pointer to the subdevice descriptor referenced by its registration index.
- int [a4l_buf_prepare_absput](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Update the absolute count of data sent from the device to the buffer since the start of the acquisition and after the next DMA shot.
- int [a4l_buf_commit_absput](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the absolute count of data which was sent from the device to the buffer since the start of the acquisition and until the last DMA shot.
- int [a4l_buf_prepare_put](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the count of data which is to be sent to the buffer at the next DMA shot.
- int [a4l_buf_commit_put](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the count of data sent to the buffer during the last completed DMA shots.
- int [a4l_buf_put](#) (struct [a4l_subdevice](#) *subd, void *bufdata, unsigned long count)
Copy some data from the device driver to the buffer.
- int [a4l_buf_prepare_absget](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Update the absolute count of data sent from the buffer to the device since the start of the acquisition and after the next DMA shot.
- int [a4l_buf_commit_absget](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the absolute count of data which was sent from the buffer to the device since the start of the acquisition and until the last DMA shot.
- int [a4l_buf_prepare_get](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the count of data which is to be sent from the buffer to the device at the next DMA shot.
- int [a4l_buf_commit_get](#) (struct [a4l_subdevice](#) *subd, unsigned long count)
Set the count of data sent from the buffer to the device during the last completed DMA shots.
- int [a4l_buf_get](#) (struct [a4l_subdevice](#) *subd, void *bufdata, unsigned long count)
Copy some data from the buffer to the device driver.
- int [a4l_buf_evt](#) (struct [a4l_subdevice](#) *subd, unsigned long evts)
Signal some event(s) to a user-space program involved in some read / write operation.
- unsigned long [a4l_buf_count](#) (struct [a4l_subdevice](#) *subd)
Get the data amount available in the Analogy buffer.
- struct [a4l_cmd_desc](#) * [a4l_get_cmd](#) (struct [a4l_subdevice](#) *subd)
Get the current Analogy command descriptor.
- int [a4l_get_chan](#) (struct [a4l_subdevice](#) *subd)
Get the channel index according to its type.
- unsigned int [a4l_get_irq](#) (struct [a4l_device](#) *dev)
Get the interrupt number in use for a specific device.
- int [a4l_request_irq](#) (struct [a4l_device](#) *dev, unsigned int irq, [a4l_irq_hdlr_t](#) handler, unsigned long flags, void *cookie)
Register an interrupt handler for a specific device.
- int [a4l_free_irq](#) (struct [a4l_device](#) *dev, unsigned int irq)
Release an interrupt handler for a specific device.
- unsigned long long [a4l_get_time](#) (void)
Get the absolute time in nanoseconds.

7.62 kernel/drivers/analogy/intel/8255.c File Reference

Analogy subdevice driver for 8255 chip.

Include dependency graph for 8255.c:



7.62.1 Detailed Description

Analogy subdevice driver for 8255 chip.

Note

Copyright (C) 1999 David A. Schleef ds@schleef.org

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

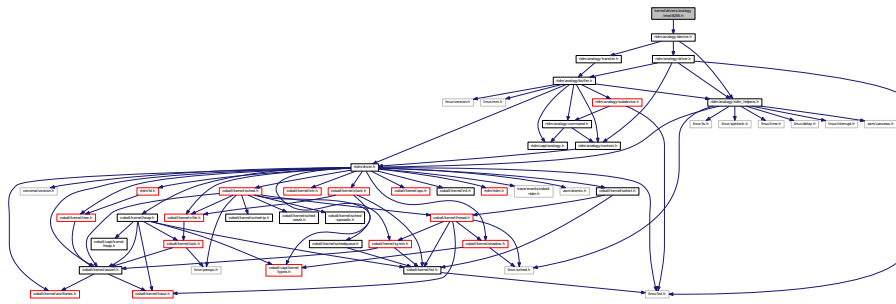
This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.


7.63 kernel/drivers/analogy/intel/8255.h File Reference

Hardware driver for 8255 chip.

Include dependency graph for 8255.h:



This graph shows which files directly or indirectly include this file:



7.63.1 Detailed Description

Hardware driver for 8255 chip.

Note

Copyright (C) 1999 David A. Schlee ds@schlee.org

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

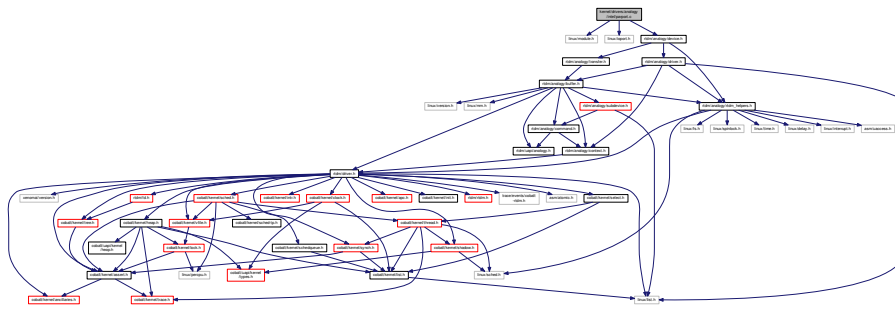
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.64 kernel/drivers/analog/intel/parport.c File Reference

Analogy driver for standard parallel port.

Include dependency graph for parport.c:



7.64.1 Detailed Description

Analogy driver for standard parallel port.

Note

Copyright (C) 1998,2001 David A. Schleef ds@schleef.org

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

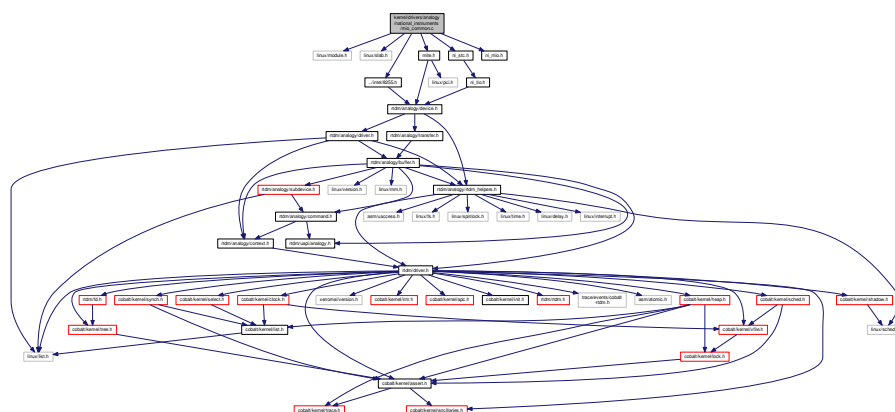
This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.65 kernel/drivers/analogy/national_instruments/mio_common.c File Reference

Hardware driver for DAQ-STC based boards.

Include dependency graph for mio_common.c:



7.65.1 Detailed Description

Hardware driver for DAQ-STC based boards. Copyright (C) 1997-2001 David A. Schleeef ds@schleeef.-org Copyright (C) 2002-2006 Frank Mori Hess fmhess@users.sourceforge.net

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Description: DAQ-STC systems

References: 340747b.pdf AT-MIO E series Register-Level Programmer Manual 341079b.pdf PCI E Series Register-Level Programmer Manual 340934b.pdf DAQ-STC reference manual 322080b.pdf 6711/6713/6715 User Manual 320945c.pdf PCI E Series User Manual 322138a.pdf PCI-6052E and D-AQPad-6052E User Manual 320517c.pdf AT E Series User manual (obsolete) 320517f.pdf AT E Series User manual 320906c.pdf Maximum signal ratings 321066a.pdf About 16x 321791a.pdf Discontinuation of at-mio-16e-10 rev. c 321808a.pdf About at-mio-16e-10 rev P 321837a.pdf Discontinuation of at-mio-16de-10 rev d 321838a.pdf About at-mio-16de-10 rev N

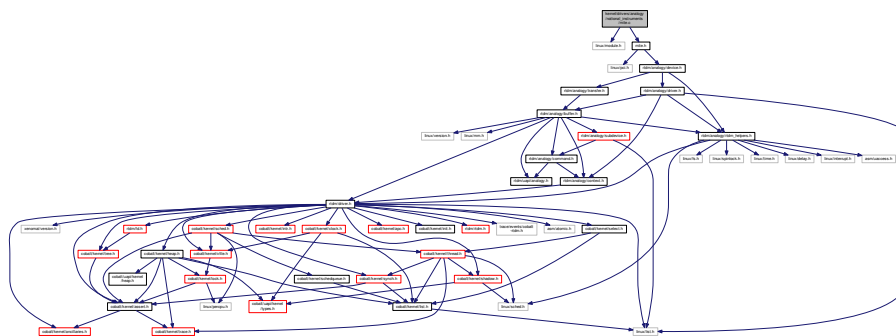
ISSUES:

- The interrupt routine needs to be cleaned up
- S-Series PCI-6143 support has been added but is not fully tested as yet. Terry Barnaby, BEAM Ltd.

7.66	kernel/drivers/analog/national_instruments/mite.c	File	Reference
------	---	------	-----------

Hardware driver for NI Mite PCI interface chip.

Include dependency graph for mite.c:



7.66.1 Detailed Description

Hardware driver for NI Mite PCI interface chip. Copyright (C) 1999 David A. Schlee ds@schlee.org

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your

7.67.1 Detailed Description

Hardware driver for NI Mite PCI interface chip.

Note

Copyright (C) 1999 David A. Schleef ds@schleef.org

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.68 kernel/drivers/analogy/national_instruments/ni_mio.h File Reference

Hardware driver for NI Mite PCI interface chip.

This graph shows which files directly or indirectly include this file:



7.68.1 Detailed Description

Hardware driver for NI Mite PCI interface chip.

Note

Copyright (C) 1999 David A. Schleef ds@schleef.org

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

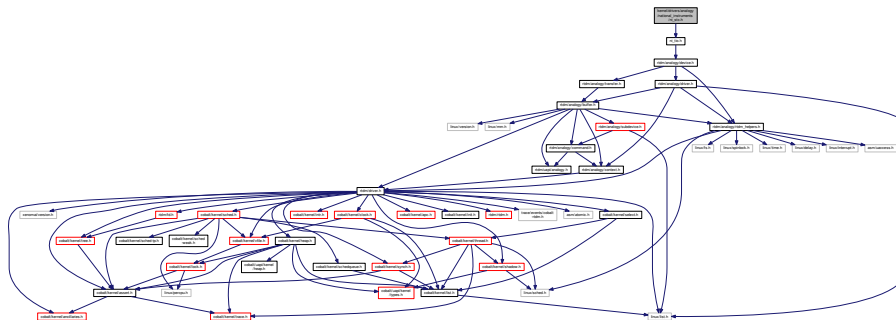
This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

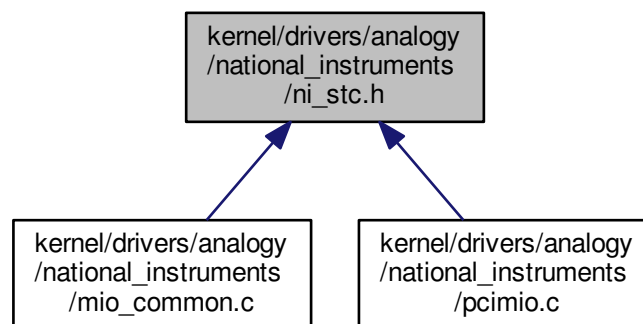
7.69 kernel/drivers/analogy/national_instruments/ni_stc.h File Reference

Register descriptions for NI DAQ-STC chip.

Include dependency graph for ni_stc.h:



This graph shows which files directly or indirectly include this file:



7.69.1 Detailed Description

Register descriptions for NI DAQ-STC chip. Copyright (C) 1998-9 David A. Schlee ds@schlee.org

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.


You should have received a copy of the GNU General Public License along with this code; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

References: 340934b.pdf DAQ-STC reference manual

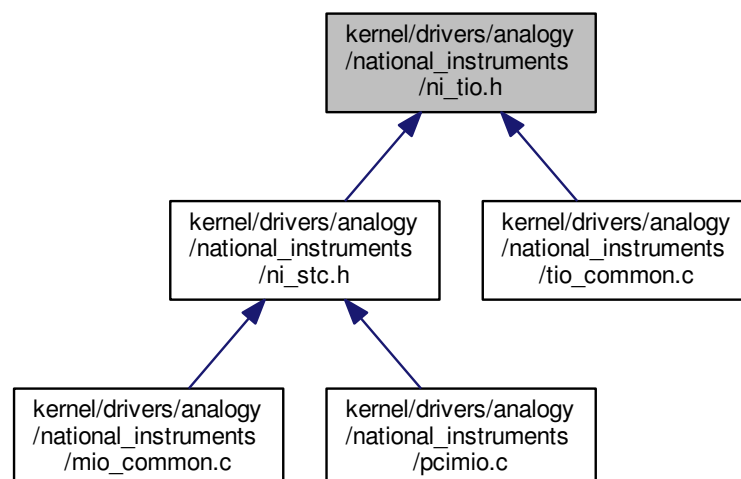
7.70 kernel/drivers/analogy/national_instruments/ni_tio.h File Reference

Hardware driver for NI general purpose counter.

Include dependency graph for ni_tio.h:



This graph shows which files directly or indirectly include this file:



7.70.1 Detailed Description

Hardware driver for NI general purpose counter.

I-6229, PCI-6250, PCI-6251, PCIe-6251, PCI-6254, PCI-6259, PCIe-6259, PCI-6280, PCI-6281, PXI-6281, PCI-6284, PCI-6289, PCI-6711, PXI-6711, PCI-6713, PXI-6713, PXI-6071E, PCI-6070E, PXI-6070E, PXI-6052E, PCI-6036E, PCI-6731, PCI-6733, PXI-6733, PCI-6143, PXI-6143

These boards are almost identical to the AT-MIO E series, except that they use the PCI bus instead of ISA (i.e., AT). See the notes for the `ni_atmio.o` driver for additional information about these boards.

By default, the driver uses DMA to transfer analog input data to memory. When DMA is enabled, not all triggering features are supported.

Note that the PCI-6143 is a simultaneous sampling device with 8 convertors. With this board all of the convertors perform one simultaneous sample during a scan interval. The period for a scan is used for the convert time in an `AnalogCmd`. The convert trigger source is normally set to `TRIG_NOW` by default.

The RTSI trigger bus is supported on these cards on subdevice

1. See the Analog library documentation for details.

References: 341079b.pdf PCI E Series Register-Level Programmer Manual 340934b.pdf DAQ-STC reference manual 322080b.pdf 6711/6713/6715 User Manual 320945c.pdf PCI E Series User Manual 322138a.pdf PCI-6052E and DAQPad-6052E User Manual

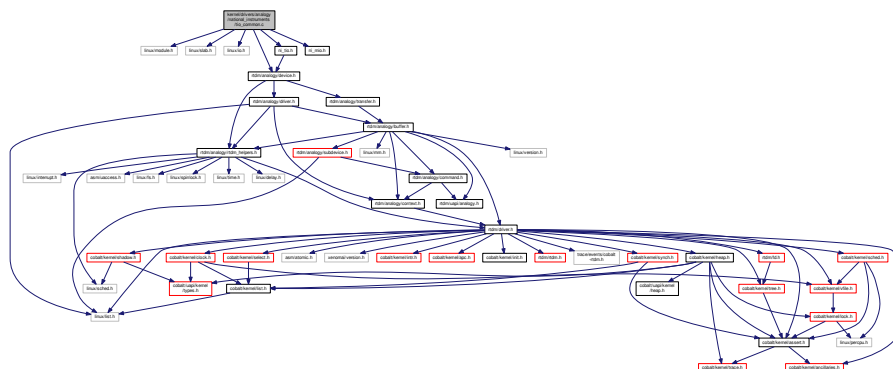
ISSUES:

- When DMA is enabled, `XXX_EV_CONVERT` does not work correctly.
- Calibration is not fully implemented
- SCXI is probably broken for m-series boards
- Digital I/O may not work on 673x.
- Information (number of channels, bits, etc.) for some devices may be incorrect. Please check this and submit a bug if there are problems for your device.
- Need to deal with external reference for DAC, and other DAC properties in board properties
- Deal with at-mio-16de-10 revision D to N changes, etc.
- Need to add other CALDAC type
- Need to slow down DAC loading. I don't trust NI's claim that two writes to the PCI bus slows IO enough. I would prefer to use `a4l_udelay()`. Timing specs: (clock) AD8522 30ns DAC8043 120ns DAC8800 60ns MB88341 ?

7.72 kernel/drivers/analog/national_instruments/tio_common.c File Reference

Hardware driver for NI general purpose counter.

Include dependency graph for `tiocommon.c`:



7.72.1 Detailed Description

Hardware driver for NI general purpose counter.

Note

Copyright (C) 2006 Frank Mori Hess fmhess@users.sourceforge.net

This code is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This code is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Description: National Instruments general purpose counters This module is not used directly by end-users. Rather, it is used by other drivers (for example `ni_660x` and `ni_pcmio`) to provide support for NI's general purpose counters. It was originally based on the counter code from `ni_660x.c` and `ni_mio-common.c`.

Author: J.P. Mellor jpmellor@rose-hulman.edu Herman Bruyninckx Herman.Bruyninckx@mech.kuleuven.ac.be Wim.-Meeussen Wim.-Meeussen@mech.kuleuven.ac.be, Klaas Gadeyne Klaas.Gadeyne@mech.kuleuven.ac.be, Frank Mori Hess fmhess@users.sourceforge.net

References: DAQ 660x Register-Level Programmer Manual (NI 370505A-01) DAQ 6601/6602 User Manual (NI 322137B-01) 340934b.pdf DAQ-STC reference manual


TODO:

- Support use of both banks X and Y

7.73 kernel/drivers/analog/proc.h File Reference

Analogy for Linux, procs related features.

This graph shows which files directly or indirectly include this file:



7.73.1 Detailed Description

Analogy for Linux, procs related features. Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
 Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

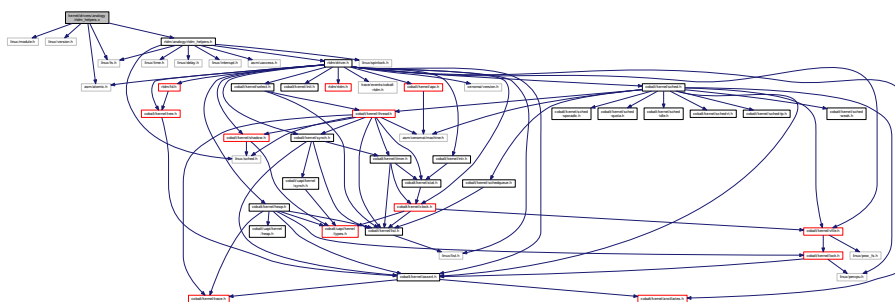
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.74 kernel/drivers/analogy/rtdm_helpers.c File Reference

Analogy for Linux, RTDM helpers.

Include dependency graph for rtdm_helpers.c:



Functions

- [nanosecs_abs_t a4l_get_time](#) (void)
Get the absolute time in nanoseconds.

7.74.1 Detailed Description

Analogy for Linux, RTDM helpers. Copyright (C) 1997-2000 David A. Schleef ds@schleef.org Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

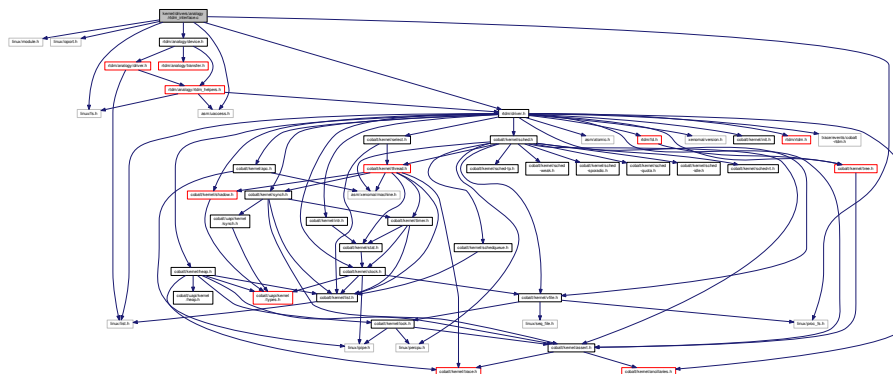
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.75 kernel/drivers/analogy/rtdm_interface.c File Reference

Analogy for Linux, user interface (open, read, write, ioctl, proc)

Include dependency graph for rtdm_interface.c:



7.75.1 Detailed Description

Analogy for Linux, user interface (open, read, write, ioctl, proc) Copyright (C) 1997-2000 David A. Schleef ds@schleef.org Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.76 kernel/drivers/analogy/sensoray/s526.c File Reference

Analogy driver for Sensoray Model 526 board.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Functions

- `struct a4l_subdevice * a4l_alloc_subd (int sizeof_priv, void(*setup)(struct a4l_subdevice *))`
Allocate a subdevice descriptor.
- `int a4l_add_subd (struct a4l_device *dev, struct a4l_subdevice *subd)`
Add a subdevice to the driver descriptor.
- `struct a4l_subdevice * a4l_get_subd (struct a4l_device *dev, int idx)`
Get a pointer to the subdevice descriptor referenced by its registration index.

7.77.1 Detailed Description

Analogy for Linux, subdevice, channel and range related features. Copyright (C) 1997-2000 David A. Schleef ds@schleef.org Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

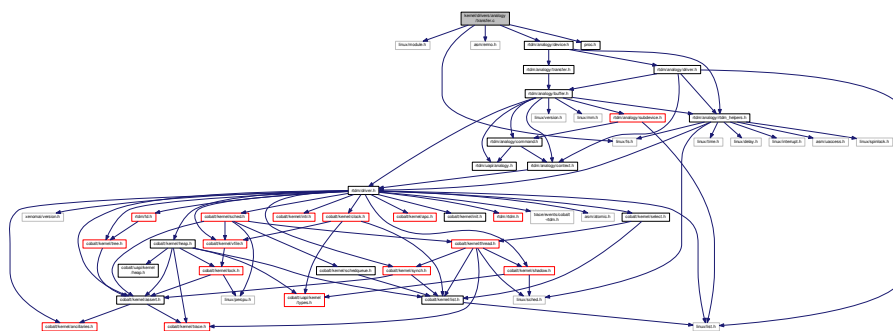
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.78 kernel/drivers/analogy/transfer.c File Reference

Analogy for Linux, transfer related features.

Include dependency graph for transfer.c:



Functions

- `int a4l_request_irq (struct a4l_device *dev, unsigned int irq, a4l_irq_hdlr_t handler, unsigned long flags, void *cookie)`
Register an interrupt handler for a specific device.
- `int a4l_free_irq (struct a4l_device *dev, unsigned int irq)`
Release an interrupt handler for a specific device.
- `unsigned int a4l_get_irq (struct a4l_device *dev)`
Get the interrupt number in use for a specific device.

7.78.1 Detailed Description

Analogy for Linux, transfer related features. Copyright (C) 1997-2000 David A. Schleeef ds@schleeef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

Xenomai is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

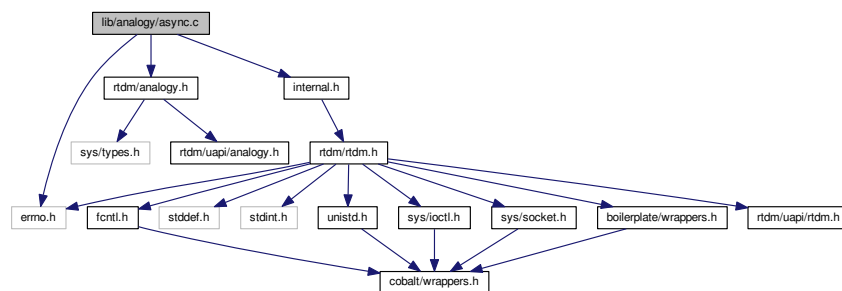
Xenomai is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with Xenomai; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

7.79 lib/analogy/async.c File Reference

Analogy for Linux, command, transfer, etc.

Include dependency graph for async.c:



Functions

- `int a4l_snd_command (a4l_desc_t *dsc, a4l_cmd_t *cmd)`
Send a command to an Analogy device.
- `int a4l_snd_cancel (a4l_desc_t *dsc, unsigned int idx_subd)`
Cancel an asynchronous acquisition.
- `int a4l_set_bufsize (a4l_desc_t *dsc, unsigned int idx_subd, unsigned long size)`
Change the size of the asynchronous buffer.
- `int a4l_get_bufsize (a4l_desc_t *dsc, unsigned int idx_subd, unsigned long *size)`
Get the size of the asynchronous buffer.
- `int a4l_mark_bufwr (a4l_desc_t *dsc, unsigned int idx_subd, unsigned long cur, unsigned long *new)`
Update the asynchronous buffer state.
- `int a4l_poll (a4l_desc_t *dsc, unsigned int idx_subd, unsigned long ms_timeout)`
Get the available data count.
- `int a4l_mmap (a4l_desc_t *dsc, unsigned int idx_subd, unsigned long size, void **ptr)`
Map the asynchronous ring-buffer into a user-space.
- `int a4l_async_read (a4l_desc_t *dsc, void *buf, size_t nbyte, unsigned long ms_timeout)`
Perform asynchronous read operation on the analog input subdevice.

- int [a4l_async_write](#) ([a4l_desc_t](#) *dsc, void *buf, size_t nbyte, unsigned long ms_timeout)
Perform asynchronous write operation on the analog input subdevice.

7.79.1 Detailed Description

Analogy for Linux, command, transfer, etc. related features

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.80 lib/analogy/descriptor.c File Reference

Analogy for Linux, descriptor related features.

Include dependency graph for descriptor.c:



Functions

- int [a4l_sys_desc](#) (int fd, [a4l_desc_t](#) *dsc, int pass)
Get a descriptor on an attached device.
- int [a4l_open](#) ([a4l_desc_t](#) *dsc, const char *fname)
Open an Analogy device and basically fill the descriptor.
- int [a4l_close](#) ([a4l_desc_t](#) *dsc)
Close the Analogy device related with the descriptor.
- int [a4l_fill_desc](#) ([a4l_desc_t](#) *dsc)
Fill the descriptor with subdevices, channels and ranges data.
- int [a4l_get_subdinfo](#) ([a4l_desc_t](#) *dsc, unsigned int subd, [a4l_sbinfo_t](#) **info)
Get an information structure on a specified subdevice.
- int [a4l_get_chinfo](#) ([a4l_desc_t](#) *dsc, unsigned int subd, unsigned int chan, [a4l_chinfo_t](#) **info)

Get an information structure on a specified channel.

- int `a4l_get_rnginfo` (`a4l_desc_t` *dsc, unsigned int subd, unsigned int chan, unsigned int rng, `a4l_rnginfo_t` **info)

Get an information structure on a specified range.

7.80.1 Detailed Description

Analogy for Linux, descriptor related features.

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org

Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

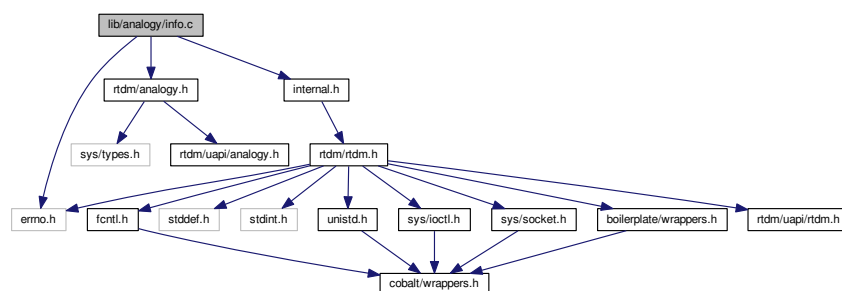
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.81 lib/analogy/info.c File Reference

Analogy for Linux, device, subdevice, etc.

Include dependency graph for info.c:



7.81.1 Detailed Description

Analogy for Linux, device, subdevice, etc. related features

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org

Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

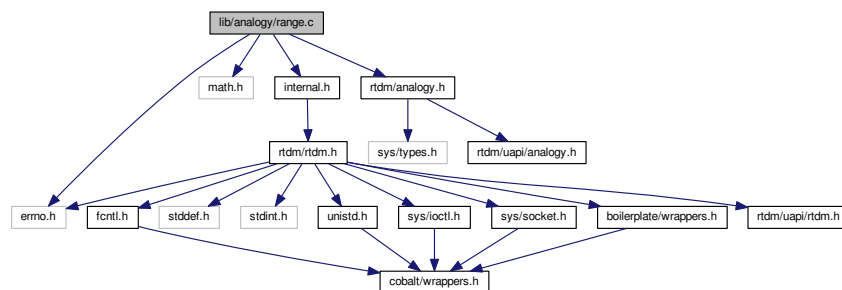
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.82 lib/analogy/range.c File Reference

Analogy for Linux, range related features.

Include dependency graph for range.c:



Functions

- int [a4l_sizeof_chan](#) (a4l_chinfo_t *chan)
Get the size in memory of an acquired element.
- int [a4l_sizeof_subd](#) (a4l_sbinfo_t *subd)
Get the size in memory of a digital acquired element.
- int [a4l_find_range](#) (a4l_desc_t *dsc, unsigned int idx_subd, unsigned int idx_chan, unsigned long unit, double min, double max, a4l_rnginfo_t **rng)
Find the most suitable range.
- int [a4l_rawtoul](#) (a4l_chinfo_t *chan, unsigned long *dst, void *src, int cnt)
Unpack raw data (from the driver) into unsigned long values.
- int [a4l_rawtof](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, float *dst, void *src, int cnt)
Convert raw data (from the driver) to float-typed samples.
- int [a4l_rawtod](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, double *dst, void *src, int cnt)
Convert raw data (from the driver) to double-typed samples.
- int [a4l_ultoraw](#) (a4l_chinfo_t *chan, void *dst, unsigned long *src, int cnt)
Pack unsigned long values into raw data (for the driver)
- int [a4l_ftoraw](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, void *dst, float *src, int cnt)
Convert float-typed samples to raw data (for the driver)
- int [a4l_dtoraw](#) (a4l_chinfo_t *chan, a4l_rnginfo_t *rng, void *dst, double *src, int cnt)
Convert double-typed samples to raw data (for the driver)

7.82.1 Detailed Description

Analogy for Linux, range related features.

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.


This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.83 lib/analogy/root_leaf.h File Reference

Analogy for Linux, root / leaf system.

This graph shows which files directly or indirectly include this file:



7.83.1 Detailed Description

Analogy for Linux, root / leaf system.

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org
Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

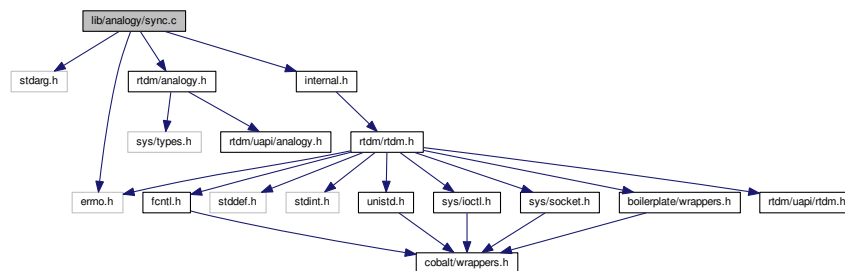
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.84 lib/analogy/sync.c File Reference

Analogy for Linux, instruction related features.

Include dependency graph for sync.c:



Functions

- int [a4l_snd_insnlist](#) ([a4l_desc_t](#) *dsc, [a4l_insnlst_t](#) *arg)
Perform a list of synchronous acquisition misc operations.
- int [a4l_snd_insn](#) ([a4l_desc_t](#) *dsc, [a4l_insn_t](#) *arg)
Perform a synchronous acquisition misc operation.
- int [a4l_sync_write](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void *buf, size_t nbyte)
Perform a synchronous acquisition write operation.
- int [a4l_sync_read](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int chan_desc, unsigned int ns_delay, void *buf, size_t nbyte)
Perform a synchronous acquisition read operation.
- int [a4l_sync_dio](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, void *mask, void *buf)
Perform a synchronous acquisition digital acquisition.
- int [a4l_config_subd](#) ([a4l_desc_t](#) *dsc, unsigned int idx_subd, unsigned int type,...)
Configure a subdevice.

7.84.1 Detailed Description

Analogy for Linux, instruction related features.

Note

Copyright (C) 1997-2000 David A. Schlee ds@schleef.org

Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

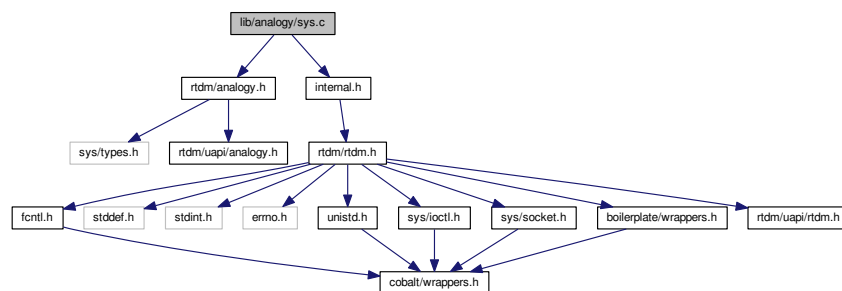
This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

7.85 lib/analogy/sys.c File Reference

Analogy for Linux, descriptor related features.

Include dependency graph for sys.c:



Functions

- int [a4l_sys_open](#) (const char *fname)
Open an Analogy device.
- int [a4l_sys_close](#) (int fd)
Close an Analogy device.
- int [a4l_sys_read](#) (int fd, void *buf, size_t nbyte)
Read from an Analogy device.
- int [a4l_sys_write](#) (int fd, void *buf, size_t nbyte)
Write to an Analogy device.
- int [a4l_sys_attach](#) (int fd, a4l_inkdsc_t *arg)
Attach an Analogy device to a driver.
- int [a4l_sys_detach](#) (int fd)
Detach an Analogy device from a driver.
- int [a4l_sys_bufcfg](#) (int fd, unsigned int idx_subd, unsigned long size)
Configure the buffer size.

7.85.1 Detailed Description

Analogy for Linux, descriptor related features.

Note

Copyright (C) 1997-2000 David A. Schleef ds@schleef.org

Copyright (C) 2008 Alexis Berlemont alexis.berlemont@free.fr

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA.

Chapter 8

Example Documentation

8.1 bufp-label.c

```
/*
 * BUFP-based client/server demo, using the read(2)/write(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created. A server thread (reader)
 * is bound to a real-time port and receives a stream of bytes sent to
 * this port from a client thread (writer).
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdm/ipc.h>

pthread_t svtid, cltid;

#define BUFP_PORT_LABEL "bufp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *server(void *arg)
{
    struct rtipc_port_label plabel;
    struct sockaddr_ipc saddr;
    char buf[128];
    size_t bufsz;
    int ret, s;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUFP);
    if (s < 0)
```

```

        fail("socket");

    /*
     * Set a 16k buffer for the server endpoint. This
     * configuration must be done prior to binding the socket to a
     * port.
     */
    bufsz = 16384; /* bytes */
    ret = setsockopt(s, SOL_BUF, BUFP_BUFSZ,
                    &bufsz, sizeof(bufsz));
    if (ret)
        fail("setsockopt");

    /*
     * Set a port label. This name will be registered when
     * binding, in addition to the port number (if given).
     */
    strcpy(plabel.label, BUFP_PORT_LABEL);
    ret = setsockopt(s, SOL_BUF, BUFP_LABEL,
                    &plabel, sizeof(plabel));
    if (ret)
        fail("setsockopt");

    /*
     * Bind the socket to the port. Assign that port a label, so
     * that peers may use a descriptive information to locate
     * it. Labeled ports will appear in the
     * /proc/xenomai/registry/rtpc/bufp directory once the socket
     * is bound.
     *
     * saddr.sipc_port specifies the port number to use. If -1 is
     * passed, the BUFP driver will auto-select an idle port.
     */
    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = -1;
    ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
    if (ret)
        fail("bind");

    for (;;) {
        ret = read(s, buf, sizeof(buf));
        if (ret < 0) {
            close(s);
            fail("read");
        }
        rt_printf("%s: received %d bytes, \"%s\"\n",
                  __FUNCTION__, ret, buf);
    }

    return NULL;
}

static void *client(void *arg)
{
    struct rtpc_port_label plabel;
    struct sockaddr_ipc svaddr;
    int ret, s, n = 0, len;
    struct timespec ts;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_BUF);
    if (s < 0)
        fail("socket");

    /*
     * Set the port label. This name will be used to find the peer
     * when connecting, instead of the port number. The label must
     * be set _after_ the socket is bound to the port, so that
     * BUFP does not try to register this label for the client
     * port as well (like the server thread did).
     */
    strcpy(plabel.label, BUFP_PORT_LABEL);
    ret = setsockopt(s, SOL_BUF, BUFP_LABEL,
                    &plabel, sizeof(plabel));
    if (ret)
        fail("setsockopt");

    memset(&svaddr, 0, sizeof(svaddr));
    svaddr.sipc_family = AF_RTIPC;
    svaddr.sipc_port = -1; /* Tell BUFP to search by label. */
    ret = connect(s, (struct sockaddr *)&svaddr, sizeof(svaddr));
    if (ret)
        fail("connect");

    for (;;) {
        len = strlen(msg[n]);
        ret = write(s, msg[n], len);
        if (ret < 0) {

```



```

        close(s);
        fail("write");
    }
    rt_printf("%s: sent %d bytes, \".*s\"\\n",
        __FUNCTION__, ret, ret, msg[n]);
    n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
    /*
     * We run in full real-time mode (i.e. primary mode),
     * so we have to let the system breathe between two
     * iterations.
     */
    ts.tv_sec = 0;
    ts.tv_nsec = 500000000; /* 500 ms */
    clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
}

return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(svtid);
    pthread_cancel(cltid);
    signal(sig, SIG_DFL);
    pthread_join(svtid, NULL);
    pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param svparam = {.sched_priority = 71 };
    struct sched_param clparam = {.sched_priority = 70 };
    pthread_attr_t svattr, clattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary mode.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&svattr);
    pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
    pthread_attr_setschedparam(&svattr, &svparam);

    errno = pthread_create(&svtid, &svattr, &server, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&clattr);
    pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
    pthread_attr_setschedparam(&clattr, &clparam);

    errno = pthread_create(&cltid, &clattr, &client, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

8.2 bufp-readwrite.c

```

/*
 * BUFp-based client/server demo, using the read(2)/write(2)
 * system calls to exchange data over a socket.

```

```

*
* In this example, two sockets are created. A server thread (reader)
* is bound to a real-time port and receives a stream of bytes sent to
* this port from a client thread (writer).
*
* See Makefile in this directory for build directives.
*/
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdev/ipc.h>

pthread_t svtid, cltid;

#define BUFP_SVPORT 12

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *server(void *arg)
{
    struct sockaddr_ipc saddr;
    char buf[128];
    size_t bufsz;
    int ret, s;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_BUF);
    if (s < 0)
        fail("socket");

    /*
     * Set a 16k buffer for the server endpoint. This
     * configuration must be done prior to binding the socket to a
     * port.
     */
    bufsz = 16384; /* bytes */
    ret = setsockopt(s, SOL_BUF, BUFP_BUFSZ,
                    &bufsz, sizeof(bufsz));
    if (ret)
        fail("setsockopt");

    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = BUFP_SVPORT;
    ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
    if (ret)
        fail("bind");

    for (;;) {
        ret = read(s, buf, sizeof(buf));
        if (ret < 0) {
            close(s);
            fail("read");
        }
        rt_printf("%s: received %d bytes, \"%s\"\n",
                __FUNCTION__, ret, buf);
    }

    return NULL;
}

```

```

static void *client(void *arg)
{
    struct sockaddr_ipc svaddr;
    int ret, s, n = 0, len;
    struct timespec ts;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_BUF);
    if (s < 0)
        fail("socket");

    memset(&svaddr, 0, sizeof(svaddr));
    svaddr.sipc_family = AF_RTIPC;
    svaddr.sipc_port = BUFP_SVPORT;
    ret = connect(s, (struct sockaddr *)&svaddr, sizeof(svaddr));
    if (ret)
        fail("connect");

    for (;;) {
        len = strlen(msg[n]);
        ret = write(s, msg[n], len);
        if (ret < 0) {
            close(s);
            fail("write");
        }
        rt_printf("%s: sent %d bytes, \"%s\"\n",
            __FUNCTION__, ret, msg[n]);
        n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
        /*
         * We run in full real-time mode (i.e. primary mode),
         * so we have to let the system breathe between two
         * iterations.
         */
        ts.tv_sec = 0;
        ts.tv_nsec = 500000000; /* 500 ms */
        clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
    }

    return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(svtid);
    pthread_cancel(cltid);
    signal(sig, SIG_DFL);
    pthread_join(svtid, NULL);
    pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param svparam = {.sched_priority = 71 };
    struct sched_param clparam = {.sched_priority = 70 };
    pthread_attr_t svattr, clattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary mode.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&svattr);
    pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
    pthread_attr_setschedparam(&svattr, &svparam);

    errno = pthread_create(&svtid, &svattr, &server, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&clattr);
    pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);

```

```

pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
pthread_attr_setschedparam(&clattr, &clparam);

errno = pthread_create(&cltid, &clattr, &client, NULL);
if (errno)
    fail("pthread_create");

sigsuspend(&oldmask);

return 0;
}

```

8.3 can-rtt.c

```

/*
 * Round-Trip-Time Test - sends and receives messages and measures the
 *                        time in between.
 *
 * Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>
 *
 * Based on RTnet's examples/xenomai/posix/rtt-sender.c.
 *
 * Copyright (C) 2002 Ulrich Marx <marx@kammer.uni-hannover.de>
 *      2002 Marc Kleine-Budde <kleine-budde@gmx.de>
 *      2006 Jan Kiszka <jan.kiszka@web.de>
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 *
 * The program sends out CAN messages periodically and copies the current
 * time-stamp to the payload. At reception, that time-stamp is compared
 * with the current time to determine the round-trip time. The jitter
 * values are printed out regularly. Concurrent tests can be carried out
 * by starting the program with different message identifiers. It is also
 * possible to use this program on a remote system as simple repeater to
 * loopback messages.
 */

#include <errno.h>
#include <mqueue.h>
#include <signal.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <limits.h>
#include <getopt.h>
#include <netinet/in.h>
#include <net/if.h>
#include <sys/ioctl.h>
#include <sys/mman.h>

#ifdef __XENO__
#include <rtdm/can.h>
#else
#include <linux/can.h>
#include <linux/can/raw.h>
#endif

#define NSEC_PER_SEC 1000000000

static unsigned int cycle = 10000; /* 10 ms */
static canid_t can_id = 0x1;

static pthread_t txthread, rxthread;
static int txsock, rxsock;
static mqd_t mq;
static int txcount, rxcount;
static int overruns;
static int repeater;

```

```

struct rtt_stat {
    long long rtt;
    long long rtt_min;
    long long rtt_max;
    long long rtt_sum;
    long long rtt_sum_last;
    int counts_per_sec;
};

static void print_usage(char *prg)
{
    fprintf(stderr,
        "Usage: %s [Options] <tx-can-interface> <rx-can-interface>\n"
        "Options:\n"
        "  -h, --help      This help\n"
        "  -r, --repeater  Repeater, send back received messages\n"
        "  -i, --id=ID     CAN Identifier (default = 0x1)\n"
        "  -c, --cycle     Cycle time in us (default = 10000us)\n",
        prg);
}

static void *transmitter(void *arg)
{
    struct sched_param param = { .sched_priority = 80 };
    struct timespec next_period;
    struct timespec time;
    struct can_frame frame;
    long long *rtt_time = (long long *)&frame.data;

    /* Pre-fill CAN frame */
    frame.can_id = can_id;
    frame.can_dlc = sizeof(*rtt_time);

#ifdef __XENO__
    pthread_set_name_np(pthread_self(), "rtcan_rtt_transmitter");
#endif
    pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

    clock_gettime(CLOCK_MONOTONIC, &next_period);

    while(1) {
        next_period.tv_nsec += cycle * 1000;
        while (next_period.tv_nsec >= NSEC_PER_SEC) {
            next_period.tv_nsec -= NSEC_PER_SEC;
            next_period.tv_sec++;
        }

        clock_nanosleep(CLOCK_MONOTONIC, TIMER_ABSTIME, &next_period, NULL);

        if (rxcount != txcount) {
            overruns++;
            continue;
        }

        clock_gettime(CLOCK_MONOTONIC, &time);
        *rtt_time = (long long)time.tv_sec * NSEC_PER_SEC + time.tv_nsec;

        /* Transmit the message containing the local time */
        if (send(txsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
            if (errno == EBADF)
                printf("terminating transmitter thread\n");
            else
                perror("send failed");
            return NULL;
        }
        txcount++;
    }
}

static void *receiver(void *arg)
{
    struct sched_param param = { .sched_priority = 82 };
    struct timespec time;
    struct can_frame frame;
    long long *rtt_time = (long long *)&frame.data;
    struct rtt_stat rtt_stat = {0, 10000000000000000LL, -10000000000000000LL,
                                0, 0, 0};

#ifdef __XENO__
    pthread_set_name_np(pthread_self(), "rtcan_rtt_receiver");
#endif
    pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

    rtt_stat.counts_per_sec = 1000000 / cycle;
}

```

```

while (1) {
    if (recv(rxsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
        if (errno == EBADF)
            printf("terminating receiver thread\n");
        else
            perror("recv failed");
        return NULL;
    }
    if (repeater) {
        /* Transmit the message back as is */
        if (send(txsock, (void *)&frame, sizeof(struct can_frame), 0) < 0) {
            if (errno == EBADF)
                printf("terminating transmitter thread\n");
            else
                perror("send failed");
            return NULL;
        }
        txcount++;
    } else {
        clock_gettime(CLOCK_MONOTONIC, &time);
        if (rxcount > 0) {
            rtt_stat.rtt = ((long long)time.tv_sec * 1000000000LL +
                           time.tv_nsec - *rtt_time);
            rtt_stat.rtt_sum += rtt_stat.rtt;
            if (rtt_stat.rtt < rtt_stat.rtt_min)
                rtt_stat.rtt_min = rtt_stat.rtt;
            if (rtt_stat.rtt > rtt_stat.rtt_max)
                rtt_stat.rtt_max = rtt_stat.rtt;
        }
        rxcount++;

        if ((rxcount % rtt_stat.counts_per_sec) == 0) {
            mq_send(mq, (char *)&rtt_stat, sizeof(rtt_stat), 0);
            rtt_stat.rtt_sum_last = rtt_stat.rtt_sum;
        }
    }
}

static void catch_signal(int sig)
{
    mq_close(mq);
}

int main(int argc, char *argv[])
{
    struct sched_param param = { .sched_priority = 1 };
    pthread_attr_t thattr;
    struct mq_attr mquattr;
    struct sockaddr_can rxaddr, txaddr;
    struct can_filter rxfilter[1];
    struct rtt_stat rtt_stat;
    char mqname[32];
    char *txdev, *rxdev;
    struct ifreq ifr;
    int ret, opt;

    struct option long_options[] = {
        { "id", required_argument, 0, 'i' },
        { "cycle", required_argument, 0, 'c' },
        { "repeater", no_argument, 0, 'r' },
        { "help", no_argument, 0, 'h' },
        { 0, 0, 0, 0 },
    };

    while ((opt = getopt_long(argc, argv, "hri:c:",
                             long_options, NULL)) != -1) {
        switch (opt) {
            case 'c':
                cycle = atoi(optarg);
                break;

            case 'i':
                can_id = strtoul(optarg, NULL, 0);
                break;

            case 'r':
                repeater = 1;
                break;

            default:
                fprintf(stderr, "Unknown option %c\n", opt);
            case 'h':
                print_usage(argv[0]);
                exit(-1);
        }
    }
}

```

```

}

printf("%d %d\n", optind, argc);
if (optind + 2 != argc) {
    print_usage(argv[0]);
    exit(0);
}

txdev = argv[optind];
rxdev = argv[optind + 1];

/* Create and configure RX socket */
if ((rxsock = socket(PF_CAN, SOCK_RAW, CAN_RAW)) < 0) {
    perror("RX socket failed");
    return -1;
}

strncpy(ifr.ifr_name, rxdev, IFNAMSIZ);
printf("RX rxsock=%d, ifr_name=%s\n", rxsock, ifr.ifr_name);

if (ioctl(rxsock, SIOCGIFINDEX, &ifr) < 0) {
    perror("RX ioctl SIOCGIFINDEX failed");
    goto failure1;
}

/* We only want to receive our own messages */
rxfilter[0].can_id = can_id;
rxfilter[0].can_mask = 0x3ff;
if (setsockopt(rxsock, SOL_CAN_RAW, CAN_RAW_FILTER,
               &rxfilter, sizeof(struct can_filter)) < 0) {
    perror("RX setsockopt CAN_RAW_FILTER failed");
    goto failure1;
}
memset(&rxaddr, 0, sizeof(rxaddr));
rxaddr.can_ifindex = ifr.ifr_ifindex;
rxaddr.can_family = AF_CAN;
if (bind(rxsock, (struct sockaddr *)&rxaddr, sizeof(rxaddr)) < 0) {
    perror("RX bind failed\n");
    goto failure1;
}

/* Create and configure TX socket */
if (strcmp(rxdev, txdev) == 0) {
    txsock = rxsock;
} else {
    if ((txsock = socket(PF_CAN, SOCK_RAW, 0)) < 0) {
        perror("TX socket failed");
        goto failure1;
    }

    strncpy(ifr.ifr_name, txdev, IFNAMSIZ);
    printf("TX txsock=%d, ifr_name=%s\n", txsock, ifr.ifr_name);

    if (ioctl(txsock, SIOCGIFINDEX, &ifr) < 0) {
        perror("TX ioctl SIOCGIFINDEX failed");
        goto failure2;
    }

    /* Suppress definition of a default receive filter list */
    if (setsockopt(txsock, SOL_CAN_RAW, CAN_RAW_FILTER, NULL, 0) < 0) {
        perror("TX setsockopt CAN_RAW_FILTER failed");
        goto failure2;
    }

    memset(&txaddr, 0, sizeof(txaddr));
    txaddr.can_ifindex = ifr.ifr_ifindex;
    txaddr.can_family = AF_CAN;

    if (bind(txsock, (struct sockaddr *)&txaddr, sizeof(txaddr)) < 0) {
        perror("TX bind failed\n");
        goto failure2;
    }
}

signal(SIGTERM, catch_signal);
signal(SIGINT, catch_signal);
signal(SIGHUP, catch_signal);
mlockall(MCL_CURRENT|MCL_FUTURE);

printf("Round-Trip-Time test %s -> %s with CAN ID 0x%x\n",
       argv[optind], argv[optind + 1], can_id);
printf("Cycle time: %d us\n", cycle);
printf("All RTT timing figures are in us.\n");

/* Create statistics message queue */
snprintf(mqname, sizeof(mqname), "/rtcan_rtt-%d", getpid());

```

```

mqattr.mq_flags = 0;
mqattr.mq_maxmsg = 100;
mqattr.mq_msgsize = sizeof(struct rtt_stat);
mq = mq_open(mqname, O_RDWR | O_CREAT | O_EXCL, 0600, &mqattr);
if (mq == (mqd_t)-1) {
    perror("opening mqueue failed");
    goto failure2;
}

/* Create receiver RT-thread */
pthread_attr_init(&thattr);
pthread_attr_setdetachstate(&thattr, PTHREAD_CREATE_JOINABLE);
pthread_attr_setstacksize(&thattr, PTHREAD_STACK_MIN);
ret = pthread_create(&rxthread, &thattr, &receiver, NULL);
if (ret) {
    fprintf(stderr, "%s: pthread_create(receiver) failed\n",
            strerror(-ret));
    goto failure3;
}

if (!repeater) {
    /* Create transmitter RT-thread */
    ret = pthread_create(&txthread, &thattr, &transmitter, NULL);
    if (ret) {
        fprintf(stderr, "%s: pthread_create(transmitter) failed\n",
                strerror(-ret));
        goto failure4;
    }
}

pthread_setschedparam(pthread_self(), SCHED_FIFO, &param);

if (repeater)
    printf("Messages\n");
else
    printf("Messages RTTlast RTT_avg RTT_min RTT_max Overruns\n");

while (1) {
    long long rtt_avg;

    ret = mq_receive(mq, (char *)&rtt_stat, sizeof(rtt_stat), NULL);
    if (ret != sizeof(rtt_stat)) {
        if (ret < 0) {
            if (errno == EBADF)
                printf("terminating mq_receive\n");
            else
                perror("mq_receive failed");
        } else
            fprintf(stderr,
                    "mq_receive returned invalid length %d\n", ret);
        break;
    }

    if (repeater) {
        printf("%8d\n", rxcount);
    } else {
        rtt_avg = ((rtt_stat.rtt_sum - rtt_stat.rtt_sum_last) /
                  rtt_stat.counts_per_sec);
        printf("%8d %7ld %7ld %7ld %7ld %8d\n", rxcount,
              (long)(rtt_stat.rtt / 1000), (long)(rtt_avg / 1000),
              (long)(rtt_stat.rtt_min / 1000),
              (long)(rtt_stat.rtt_max / 1000),
              overruns);
    }
}

/* This call also leaves primary mode, required for socket cleanup. */
printf("shutting down\n");

/* Important: First close the sockets! */
while ((close(rxsock) < 0) && (errno == EAGAIN)) {
    printf("RX socket busy - waiting...\n");
    sleep(1);
}
while ((close(txsock) < 0) && (errno == EAGAIN)) {
    printf("TX socket busy - waiting...\n");
    sleep(1);
}

pthread_join(txthread, NULL);
pthread_kill(rxthread, SIGHUP);
pthread_join(rxthread, NULL);

return 0;

failure4:
pthread_kill(rxthread, SIGHUP);

```



```

    pthread_join(rxtthread, NULL);
failure3:
    mq_close(mq);
failure2:
    close(txsock);
failure1:
    close(rxsock);

    return 1;
}

```

8.4 cross-link.c

```

/*
 * cross-link.c
 *
 * Userspace test program (Xenomai alchemy skin) for RTDM-based UART drivers
 * Copyright 2005 by Joerg Langenberg <joergel75@gmx.net>
 *
 * Updates by Jan Kiszka <jan.kiszka@web.de>
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 */
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <sys/mman.h>
#include <alchemy/task.h>
#include <alchemy/timer.h>
#include <rtdm/serial.h>

#define MAIN_PREFIX    "main : "
#define WTASK_PREFIX   "write_task: "
#define RTASK_PREFIX   "read_task: "

#define WRITE_FILE      "rtser0"
#define READ_FILE       "rtser1"

int read_fd = -1;
int write_fd = -1;

#define STATE_FILE_OPENED      1
#define STATE_TASK_CREATED    2

unsigned int read_state = 0;
unsigned int write_state = 0;

/*
 * --s-ms-us-ns */
RTIME write_task_period_ns = 1000000000ll;
RT_TASK write_task;
RT_TASK read_task;

static const struct rtser_config read_config = {
    .config_mask    = 0xFFFF,
    .baud_rate      = 115200,
    .parity         = RTSER_DEF_PARITY,
    .data_bits      = RTSER_DEF_BITS,
    .stop_bits      = RTSER_DEF_STOPB,
    .handshake       = RTSER_DEF_HAND,
    .fifo_depth     = RTSER_DEF_FIFO_DEPTH,
    .rx_timeout     = RTSER_DEF_TIMEOUT,
    .tx_timeout     = RTSER_DEF_TIMEOUT,
    .event_timeout  = 1000000000, /* 1 s */
    .timestamp_history = RTSER_RX_TIMESTAMP_HISTORY,
    .event_mask     = RTSER_EVENT_RXPEND,
};

static const struct rtser_config write_config = {
    .config_mask    = RTSER_SET_BAUD | RTSER_SET_TIMESTAMP_HISTORY,
    .baud_rate      = 115200,
    .timestamp_history = RTSER_DEF_TIMESTAMP_HISTORY,
};

```

```

        /* the rest implicitly remains default */
};

static int close_file( int fd, char *name)
{
    int err, i=0;

    do {
        i++;
        err = rt_dev_close(fd);
        switch (err) {
            case -EAGAIN:
                printf(MAIN_PREFIX "%s -> EAGAIN (%d times)\n",
                       name, i);
                rt_task_sleep(50000); /* wait 50us */
                break;
            case 0:
                printf(MAIN_PREFIX "%s -> closed\n", name);
                break;
            default:
                printf(MAIN_PREFIX "%s -> %s\n", name,
                       strerror(-err));
                break;
        }
    } while (err == -EAGAIN && i < 10);

    return err;
}

static void cleanup_all(void)
{
    if (read_state & STATE_FILE_OPENED) {
        close_file(read_fd, READ_FILE " (read)");
        read_state &= ~STATE_FILE_OPENED;
    }

    if (write_state & STATE_FILE_OPENED) {
        close_file(write_fd, WRITE_FILE " (write)");
        write_state &= ~STATE_FILE_OPENED;
    }

    if (write_state & STATE_TASK_CREATED) {
        printf(MAIN_PREFIX "delete write_task\n");
        rt_task_delete(&write_task);
        write_state &= ~STATE_TASK_CREATED;
    }

    if (read_state & STATE_TASK_CREATED) {
        printf(MAIN_PREFIX "delete read_task\n");
        rt_task_delete(&read_task);
        read_state &= ~STATE_TASK_CREATED;
    }
}

static void catch_signal(int sig)
{
    cleanup_all();
    printf(MAIN_PREFIX "exit\n");
    return;
}

static void write_task_proc(void *arg)
{
    int err;
    RTIME write_time;
    ssize_t sz = sizeof(RTIME);
    int written = 0;

    err = rt_task_set_periodic(NULL, TM_NOW,
                               rt_timer_ns2ticks(write_task_period_ns));
    if (err) {
        printf(WTASK_PREFIX "error on set periodic, %s\n",
               strerror(-err));
        goto exit_write_task;
    }

    while (1) {
        err = rt_task_wait_period(NULL);
        if (err) {
            printf(WTASK_PREFIX
                   "error on rt_task_wait_period, %s\n",
                   strerror(-err));
            break;
        }

        write_time = rt_timer_read();
    }
}

```

```

        written = rt_dev_write(write_fd, &write_time, sz);
        if (written < 0) {
            printf(WTASK_PREFIX "error on rt_dev_write, %s\n",
                   strerror(-err));
            break;
        } else if (written != sz) {
            printf(WTASK_PREFIX "only %d / %zd byte transmitted\n",
                   written, sz);
            break;
        }
    }

exit_write_task:
    if ((write_state & STATE_FILE_OPENED) &&
        close_file(write_fd, WRITE_FILE " (write)") == 0)
        write_state &= ~STATE_FILE_OPENED;

    printf(WTASK_PREFIX "exit\n");
}

static void read_task_proc(void *arg)
{
    int err;
    int nr = 0;
    RTIME read_time = 0;
    RTIME write_time = 0;
    RTIME irq_time = 0;
    ssize_t sz = sizeof(RTIME);
    int read = 0;
    struct rtser_event rx_event;

    printf("Nr | write->irq | irq->read | write->read |\n");
    printf("-----\n");

    /*
     * We are in secondary mode now due to printf, the next
     * blocking Xenomai or driver call will switch us back
     * (here: RTSER_RTIOC_WAIT_EVENT).
     */

    while (1) {
        /* waiting for event */
        err = rt_dev_ioctl(read_fd, RTSER_RTIOC_WAIT_EVENT, &
rx_event);
        if (err) {
            printf(RTASK_PREFIX
                   "error on RTSER_RTIOC_WAIT_EVENT, %s\n",
                   strerror(-err));
            if (err == -ETIMEDOUT)
                continue;
            break;
        }

        irq_time = rx_event.rxpend_timestamp;
        read = rt_dev_read(read_fd, &write_time, sz);
        if (read == sz) {
            read_time = rt_timer_read();
            printf("%3d |%16llu |%16llu |%16llu\n", nr,
                   irq_time - write_time,
                   read_time - irq_time,
                   read_time - write_time);
            nr++;
        } else if (read < 0) {
            printf(RTASK_PREFIX "error on rt_dev_read, code %s\n",
                   strerror(-err));
            break;
        } else {
            printf(RTASK_PREFIX "only %d / %zd byte received \n",
                   read, sz);
            break;
        }
    }

    if ((read_state & STATE_FILE_OPENED) &&
        close_file(read_fd, READ_FILE " (read)") == 0)
        read_state &= ~STATE_FILE_OPENED;

    printf(RTASK_PREFIX "exit\n");
}

int main(int argc, char* argv[])
{
    int err = 0;

    signal(SIGTERM, catch_signal);
    signal(SIGINT, catch_signal);

```

```

/* no memory-swapping for this programm */
mlockall(MCL_CURRENT | MCL_FUTURE);

/* open rtser0 */
write_fd = rt_dev_open( WRITE_FILE, 0);
if (write_fd < 0) {
    printf(MAIN_PREFIX "can't open %s (write), %s\n", WRITE_FILE,
           strerror(-write_fd));
    goto error;
}
write_state |= STATE_FILE_OPENED;
printf(MAIN_PREFIX "write-file opened\n");

/* writing write-config */
err = rt_dev_ioctl(write_fd, RTSER_RTIOC_SET_CONFIG, &
write_config);
if (err) {
    printf(MAIN_PREFIX "error while RTSER_RTIOC_SET_CONFIG, %s\n",
           strerror(-err));
    goto error;
}
printf(MAIN_PREFIX "write-config written\n");

/* open rtser1 */
read_fd = rt_dev_open( READ_FILE, 0 );
if (read_fd < 0) {
    printf(MAIN_PREFIX "can't open %s (read), %s\n", READ_FILE,
           strerror(-read_fd));
    goto error;
}
read_state |= STATE_FILE_OPENED;
printf(MAIN_PREFIX "read-file opened\n");

/* writing read-config */
err = rt_dev_ioctl(read_fd, RTSER_RTIOC_SET_CONFIG, &read_config)
;
if (err) {
    printf(MAIN_PREFIX "error while rt_dev_ioctl, %s\n",
           strerror(-err));
    goto error;
}
printf(MAIN_PREFIX "read-config written\n");

/* create write_task */
err = rt_task_create(&write_task, "write_task", 0, 50, 0);
if (err) {
    printf(MAIN_PREFIX "failed to create write_task, %s\n",
           strerror(-err));
    goto error;
}
write_state |= STATE_TASK_CREATED;
printf(MAIN_PREFIX "write-task created\n");

/* create read_task */
err = rt_task_create(&read_task, "read_task", 0, 51, 0);
if (err) {
    printf(MAIN_PREFIX "failed to create read_task, %s\n",
           strerror(-err));
    goto error;
}
read_state |= STATE_TASK_CREATED;
printf(MAIN_PREFIX "read-task created\n");

/* start write_task */
printf(MAIN_PREFIX "starting write-task\n");
err = rt_task_start(&write_task, &write_task_proc, NULL);
if (err) {
    printf(MAIN_PREFIX "failed to start write_task, %s\n",
           strerror(-err));
    goto error;
}

/* start read_task */
printf(MAIN_PREFIX "starting read-task\n");
err = rt_task_start(&read_task, &read_task_proc, NULL);
if (err) {
    printf(MAIN_PREFIX "failed to start read_task, %s\n",
           strerror(-err));
    goto error;
}

pause();
return 0;

error:
cleanup_all();
return err;

```

```
}

```

8.5 iddp-label.c

```
/*
 * IDDP-based client/server demo, using the write(2)/recvfrom(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created. A server thread (reader)
 * is bound to a labeled real-time port and receives datagrams sent to
 * this port from a client thread (writer). The client thread attaches
 * to the port opened by the server using a labeled connection
 * request. The client socket is bound to a different port, only to
 * provide a valid peer name; this is optional.
 *
 * ASCII labels can be attached to bound ports, in order to connect
 * sockets to them in a more descriptive way than using plain numeric
 * port values.
 *
 * See Makefile in this directory for build directives.
 */
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtm/ipc.h>

pthread_t svtid, cltid;

#define IDDP_CLPORT 27

#define IDDP_PORT_LABEL "iddp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *server(void *arg)
{
    struct sockaddr_ipc saddr, claddr;
    struct rtipc_port_label plabel;
    socklen_t addrlen;
    char buf[128];
    int ret, s;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
    if (s < 0)
        fail("socket");

    /*
     * We will use Xenomai's system heap for datagram, so no
     * IDDP_POOLSZ required here.
     */

    /*
     * Set a port label. This name will be registered when
     * binding, in addition to the port number (if given).
     */

```

```

strcpy(plabel.label, IDDP_PORT_LABEL);
ret = setsockopt(s, SOL_IDDP, IDDP_LABEL,
                &plabel, sizeof(plabel));
if (ret)
    fail("setsockopt");

/*
 * Bind the socket to the port. Assign that port a label, so
 * that peers may use a descriptive information to locate
 * it. Labeled ports will appear in the
 * /proc/xenomai/registry/rtipc/iddp directory once the socket
 * is bound.
 */
saddr.sipc_port specifies the port number to use. If -1 is
passed, the IDDP driver will auto-select an idle port.
*/
saddr.sipc_family = AF_RTIPC;
saddr.sipc_port = -1; /* Pick next free */
ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
if (ret)
    fail("bind");

for (;;) {
    addrlen = sizeof(saddr);
    ret = recvfrom(s, buf, sizeof(buf), 0,
                  (struct sockaddr *)&claddr, &addrlen);
    if (ret < 0) {
        close(s);
        fail("recvfrom");
    }
    rt_printf("%s: received %d bytes, \"%s\" from port %d\n",
              __FUNCTION__, ret, ret, buf, claddr.sipc_port);
}

return NULL;
}

static void *client(void *arg)
{
    struct sockaddr_ipc svaddr, claddr;
    struct rtipc_port_label plabel;
    int ret, s, n = 0, len;
    struct timespec ts;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_IDDP);
    if (s < 0)
        fail("socket");

    /*
     * Set a name on the client socket. This is strictly optional,
     * and only done here for the purpose of getting back a
     * different port number in recvfrom().
     */
    claddr.sipc_family = AF_RTIPC;
    claddr.sipc_port = IDDP_CLPORT;
    ret = bind(s, (struct sockaddr *)&claddr, sizeof(claddr));
    if (ret)
        fail("bind");

    /*
     * Set the port label. This name will be used to find the peer
     * when connecting, instead of the port number. The label must
     * be set _after_ the socket is bound to the port, so that
     * IDDP does not try to register this label for the client
     * port as well (like the server thread did).
     */
    strcpy(plabel.label, IDDP_PORT_LABEL);
    ret = setsockopt(s, SOL_IDDP, IDDP_LABEL,
                    &plabel, sizeof(plabel));
    if (ret)
        fail("setsockopt");

    memset(&svaddr, 0, sizeof(svaddr));
    svaddr.sipc_family = AF_RTIPC;
    svaddr.sipc_port = -1; /* Tell IDDP to search by label. */
    ret = connect(s, (struct sockaddr *)&svaddr, sizeof(svaddr));
    if (ret)
        fail("connect");

    for (;;) {
        len = strlen(msg[n]);
        /* Send to default destination we connected to. */
        ret = write(s, msg[n], len);
        if (ret < 0) {
            close(s);
            fail("sendto");
        }
    }
}

```

```

        rt_printf("%s: sent %d bytes, \".*s\\n\",
            __FUNCTION__, ret, ret, msg[n]);
        n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
    /*
     * We run in full real-time mode (i.e. primary mode),
     * so we have to let the system breathe between two
     * iterations.
     */
    ts.tv_sec = 0;
    ts.tv_nsec = 500000000; /* 500 ms */
    clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
}

return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(svtid);
    pthread_cancel(cltid);
    signal(sig, SIG_DFL);
    pthread_join(svtid, NULL);
    pthread_join(cltid, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param svparam = {.sched_priority = 71 };
    struct sched_param clparam = {.sched_priority = 70 };
    pthread_attr_t svattr, clattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary mode.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&svattr);
    pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
    pthread_attr_setschedparam(&svattr, &svparam);

    errno = pthread_create(&svtid, &svattr, &server, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&clattr);
    pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
    pthread_attr_setschedparam(&clattr, &clparam);

    errno = pthread_create(&cltid, &clattr, &client, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

8.6 iddp-sendrecv.c

```

/*
 * IDDP-based client/server demo, using the sendto(2)/recvfrom(2)
 * system calls to exchange data over a socket.
 *
 * In this example, two sockets are created. A server thread (reader)
 * is bound to a real-time port and receives datagrams sent to this

```

```

* port from a client thread (writer). The client socket is bound to a
* different port, only to provide a valid peer name; this is
* optional.
*
* See Makefile in this directory for build directives.
*/
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <pthread.h>
#include <errno.h>
#include <rtdm/ipc.h>

pthread_t svtid, cltid;

#define IDDP_SVPORT 12
#define IDDP_CLPORT 13

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *server(void *arg)
{
    struct sockaddr_ipc saddr, claddr;
    socklen_t addrlen;
    char buf[128];
    size_t poolsz;
    int ret, s;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_IDDP);
    if (s < 0)
        fail("socket");

    /*
     * Set a local 32k pool for the server endpoint. Memory needed
     * to convey datagrams will be pulled from this pool, instead
     * of Xenomai's system pool.
     */
    poolsz = 32768; /* bytes */
    ret = setsockopt(s, SOL_IDDP, IDDP_POOLSZ,
                    &poolsz, sizeof(poolsz));
    if (ret)
        fail("setsockopt");

    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = IDDP_SVPORT;
    ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
    if (ret)
        fail("bind");

    for (;;) {
        addrlen = sizeof(saddr);
        ret = recvfrom(s, buf, sizeof(buf), 0,
                      (struct sockaddr *)&claddr, &addrlen);
        if (ret < 0) {
            close(s);
            fail("recvfrom");
        }
        rt_printf("%s: received %d bytes, \"%s\" from port %d\n",
                  __FUNCTION__, ret, buf, claddr.sipc_port);
    }
}

```



```

        return NULL;
    }

    static void *client(void *arg)
    {
        struct sockaddr_ipc svaddr, claddr;
        int ret, s, n = 0, len;
        struct timespec ts;

        s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_IDDP);
        if (s < 0)
            fail("socket");

        claddr.sipc_family = AF_RTIPC;
        claddr.sipc_port = IDDP_CLPORT;
        ret = bind(s, (struct sockaddr *)&claddr, sizeof(claddr));
        if (ret)
            fail("bind");

        svaddr.sipc_family = AF_RTIPC;
        svaddr.sipc_port = IDDP_SVPORT;
        for (;;) {
            len = strlen(msg[n]);
            ret = sendto(s, msg[n], len, 0,
                (struct sockaddr *)&svaddr, sizeof(svaddr));
            if (ret < 0) {
                close(s);
                fail("sendto");
            }
            rt_printf("%s: sent %d bytes, \"%s\"\n",
                __FUNCTION__, ret, ret, msg[n]);
            n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
            /*
             * We run in full real-time mode (i.e. primary mode),
             * so we have to let the system breathe between two
             * iterations.
             */
            ts.tv_sec = 0;
            ts.tv_nsec = 500000000; /* 500 ms */
            clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
        }

        return NULL;
    }

    static void cleanup_upon_sig(int sig)
    {
        pthread_cancel(svtid);
        pthread_cancel(cltid);
        signal(sig, SIG_DFL);
        pthread_join(svtid, NULL);
        pthread_join(cltid, NULL);
    }

    int main(int argc, char **argv)
    {
        struct sched_param svparam = {.sched_priority = 71 };
        struct sched_param clparam = {.sched_priority = 70 };
        pthread_attr_t svattr, clattr;
        sigset_t mask, oldmask;

        mlockall(MCL_CURRENT | MCL_FUTURE);

        sigemptyset(&mask);
        sigaddset(&mask, SIGINT);
        signal(SIGINT, cleanup_upon_sig);
        sigaddset(&mask, SIGTERM);
        signal(SIGTERM, cleanup_upon_sig);
        sigaddset(&mask, SIGHUP);
        signal(SIGHUP, cleanup_upon_sig);
        pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

        /*
         * This is a real-time compatible printf() package from
         * Xenomai's RT Development Kit (RTDK), that does NOT cause
         * any transition to secondary mode.
         */
        rt_print_auto_init(1);

        pthread_attr_init(&svattr);
        pthread_attr_setdetachstate(&svattr, PTHREAD_CREATE_JOINABLE);
        pthread_attr_setinheritsched(&svattr, PTHREAD_EXPLICIT_SCHED);
        pthread_attr_setschedpolicy(&svattr, SCHED_FIFO);
        pthread_attr_setschedparam(&svattr, &svparam);

        errno = pthread_create(&svtid, &svattr, &server, NULL);
        if (errno)

```

```

        fail("pthread_create");

    pthread_attr_init(&clattr);
    pthread_attr_setdetachstate(&clattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&clattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&clattr, SCHED_FIFO);
    pthread_attr_setschedparam(&clattr, &clparam);

    errno = pthread_create(&cltid, &clattr, &client, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

8.7 rtcanconfig.c

```

/*
 * Program to configuring the CAN controller
 *
 * Copyright (C) 2006 Wolfgang Grandegger <wg@grandegger.com>
 *
 * Copyright (C) 2005, 2006 Sebastian Smolorz
 *      <Sebastian.Smolorz@stud.uni-hannover.de>
 *
 *
 * This program is free software; you can redistribute it and/or modify
 * it under the terms of the GNU General Public License as published by
 * the Free Software Foundation; either version 2 of the License, or
 * (at your option) any later version.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software
 * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 */

#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <string.h>
#include <time.h>
#include <errno.h>
#include <getopt.h>
#include <sys/mman.h>

#include <rtdev/can.h>

static void print_usage(char *prg)
{
    fprintf(stderr,
        "Usage: %s <can-interface> [Options] [up|down|start|stop|sleep]\n"
        "Options:\n"
        "  -v, --verbose           be verbose\n"
        "  -h, --help             this help\n"
        "  -c, --ctrlmode=CTRLMODE listenonly, loopback or none\n"
        "  -b, --baudrate=BPS      baudrate in bits/sec\n"
        "  -B, --bittime=BTR0:BTR1 BTR or standard bit-time\n"
        "  -B, --bittime=BRP:PROP_SEG:PHASE_SEG1:PHASE_SEG2:SJW:SAM\n",
        prg);
}

static can_baudrate_t string_to_baudrate(char *str)
{
    can_baudrate_t baudrate;
    if (sscanf(str, "%i", &baudrate) != 1)
        return -1;
    return baudrate;
}

static int string_to_mode(char *str)
{
    if ( ! strcmp(str, "up") || ! strcmp(str, "start") )
        return CAN_MODE_START;
    else if ( ! strcmp(str, "down") || ! strcmp(str, "stop") )
        return CAN_MODE_STOP;
}

```

```

    else if ( !strcmp(str, "sleep") )
        return CAN_MODE_SLEEP;
    return -EINVAL;
}

static int string_to_ctrlmode(char *str)
{
    if ( !strcmp(str, "listenonly") )
        return CAN_CTRLMODE_LISTENONLY;
    else if ( !strcmp(str, "loopback") )
        return CAN_CTRLMODE_LOOPBACK;
    else if ( !strcmp(str, "none") )
        return 0;

    return -1;
}

int main(int argc, char *argv[])
{
    char    ifname[16];
    int     can_fd = -1;
    int     new_baudrate = -1;
    int     new_mode = -1;
    int     new_ctrlmode = 0, set_ctrlmode = 0;
    int     verbose = 0;
    int     bittime_count = 0, bittime_data[6];
    can_baudrate_t *baudrate;
    can_ctrlmode_t *ctrlmode;
    can_mode_t *mode;
    union {
        struct ifreq ifr;
        struct can_bittime bittime;
        can_baudrate_t baudrate;
        can_ctrlmode_t ctrlmode;
        can_mode_t mode;
    } u;
    struct can_bittime *bittime;
    int opt, ret;
    char* ptr;

    struct option long_options[] = {
        { "help", no_argument, 0, 'h' },
        { "verbose", no_argument, 0, 'v' },
        { "baudrate", required_argument, 0, 'b' },
        { "bittime", required_argument, 0, 'B' },
        { "ctrlmode", required_argument, 0, 'c' },
        { 0, 0, 0, 0 },
    };

    while ((opt = getopt_long(argc, argv, "hvb:B:c:",
                             long_options, NULL)) != -1) {
        switch (opt) {
            case 'h':
                print_usage(argv[0]);
                exit(0);

            case 'v':
                verbose = 1;
                break;

            case 'b':
                new_baudrate = string_to_baudrate(optarg);
                if (new_baudrate == -1) {
                    print_usage(argv[0]);
                    exit(0);
                }
                break;

            case 'B':
                ptr = optarg;
                while (1) {
                    bittime_data[bittime_count++] = strtoul(ptr, NULL, 0);
                    if (!(ptr = strchr(ptr, ':')))
                        break;
                    ptr++;
                }
                if (bittime_count != 2 && bittime_count != 6) {
                    print_usage(argv[0]);
                    exit(0);
                }
                break;

            case 'c':
                ret = string_to_ctrlmode(optarg);
                if (ret == -1) {
                    print_usage(argv[0]);
                    exit(0);
                }

```

```

    }
    new_ctrlmode |= ret;
    set_ctrlmode = 1;
    break;

    break;

default:
    fprintf(stderr, "Unknown option %c\n", opt);
    break;
}
}

/* Get CAN interface name */
if (optind != argc - 1 && optind != argc - 2) {
    print_usage(argv[0]);
    return 0;
}

strncpy(iframe, argv[optind], IFNAMSIZ);
strncpy(u.ifr.ifr_name, iframe, IFNAMSIZ);

if (optind == argc - 2) { /* Get mode setting */
    new_mode = string_to_mode(argv[optind + 1]);
    if (verbose)
        printf("mode: %s (%#x)\n", argv[optind + 1], new_mode);
    if (new_mode < 0) {
        print_usage(argv[0]);
        return 0;
    }
}

can_fd = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
if (can_fd < 0) {
    fprintf(stderr, "Cannot open RTDM CAN socket. Maybe driver not loaded? \n");
    return can_fd;
}

ret = rt_dev_ioctl(can_fd, SIOCGIFINDEX, &u.ifr);
if (ret) {
    fprintf(stderr, "Can't get interface index for %s, code = %d\n", iframe, ret);
    return ret;
}

if (new_baudrate != -1) {
    if (verbose)
        printf("baudrate: %d\n", new_baudrate);
    baudrate = &u.baudrate;
    *baudrate = new_baudrate;
    ret = rt_dev_ioctl(can_fd, SIOCSKANBAUDRATE, &u.ifr);
    if (ret) {
        goto abort;
    }
}

if (bittime_count) {
    bittime = &u.bittime;
    if (bittime_count == 2) {
        bittime->type = CAN_BITTIME_BTR;
        bittime->btr.btr0 = bittime_data[0];
        bittime->btr.btr1 = bittime_data[1];
        if (verbose)
            printf("bit-time: btr0=0x%02x btr1=0x%02x\n",
                    bittime->btr.btr0, bittime->btr.btr1);
    } else {
        bittime->type = CAN_BITTIME_STD;
        bittime->std.brp = bittime_data[0];
        bittime->std.prop_seg = bittime_data[1];
        bittime->std.phase_seg1 = bittime_data[2];
        bittime->std.phase_seg2 = bittime_data[3];
        bittime->std.sjw = bittime_data[4];
        bittime->std.sam = bittime_data[5];
        if (verbose)
            printf("bit-time: brp=%d prop_seg=%d phase_seg1=%d "
                    "phase_seg2=%d sjw=%d sam=%d\n",
                    bittime->std.brp,
                    bittime->std.prop_seg,
                    bittime->std.phase_seg1,
                    bittime->std.phase_seg2,
                    bittime->std.sjw,
                    bittime->std.sam);
    }
}

ret = rt_dev_ioctl(can_fd, SIOCSKANCUSTOMBITTIME, &u.ifr);
if (ret) {
    goto abort;
}

```

```

    }

}

if (set_ctrlmode != 0) {
    ctrlmode = &u.ctrlmode;
    *ctrlmode = new_ctrlmode;
    if (verbose)
        printf("ctrlmode: %#x\n", new_ctrlmode);
    ret = rt_dev_ioctl(can_fd, SIOCSCANCTRLMODE, &u.ifr);
    if (ret) {
        goto abort;
    }
}

if (new_mode != -1) {
    mode = &u.mode;
    *mode = new_mode;
    ret = rt_dev_ioctl(can_fd, SIOCSCANMODE, &u.ifr);
    if (ret) {
        goto abort;
    }
}

rt_dev_close(can_fd);
return 0;

abort:
rt_dev_close(can_fd);
return ret;
}

```

8.8 rtcanrecv.c

```

#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <time.h>
#include <errno.h>
#include <getopt.h>
#include <sys/mman.h>

#include <alchemy/task.h>

#include <rtdev/can.h>

static void print_usage(char *prg)
{
    fprintf(stderr,
        "Usage: %s [<can-interface>] [Options]\n"
        "Options:\n"
        " -f --filter=id:mask[:id:mask]... apply filter\n"
        " -e --error=mask      receive error messages\n"
        " -t, --timeout=MS      timeout in ms\n"
        " -T, --timestamp      with absolute timestamp\n"
        " -R, --timestamp-rel   with relative timestamp\n"
        " -v, --verbose         be verbose\n"
        " -p, --print=MODULO    print every MODULO message\n"
        " -h, --help           this help\n",
        prg);
}

extern int optind, opterr, optopt;

static int s = -1, verbose = 0, print = 1;
static nanosecs_rel_t timeout = 0, with_timestamp = 0, timestamp_rel = 0;

RT_TASK rt_task_desc;

#define BUF_SIZ 255
#define MAX_FILTER 16

struct sockaddr_can recv_addr;
struct can_filter recv_filter[MAX_FILTER];
static int filter_count = 0;

static int add_filter(u_int32_t id, u_int32_t mask)
{
    if (filter_count >= MAX_FILTER)
        return -1;
    recv_filter[filter_count].can_id = id;
}

```

```

    recv_filter[filter_count].can_mask = mask;
    printf("Filter #d: id=0x%08x mask=0x%08x\n", filter_count, id, mask);
    filter_count++;
    return 0;
}

static void cleanup(void)
{
    int ret;

    if (verbose)
        printf("Cleaning up...\n");

    if (s >= 0) {
        ret = rt_dev_close(s);
        s = -1;
        if (ret) {
            fprintf(stderr, "rt_dev_close: %s\n", strerror(-ret));
        }
        exit(EXIT_SUCCESS);
    }
}

static void cleanup_and_exit(int sig)
{
    if (verbose)
        printf("Signal %d received\n", sig);
    cleanup();
    exit(0);
}

static void rt_task(void)
{
    int i, ret, count = 0;
    struct can_frame frame;
    struct sockaddr_can addr;
    socklen_t addrlen = sizeof(addr);
    struct msghdr msg;
    struct iovec iov;
    nanosecs_abs_t timestamp, timestamp_prev = 0;

    if (with_timestamp) {
        msg.msg_iov = &iov;
        msg.msg_iovlen = 1;
        msg.msg_name = (void *)&addr;
        msg.msg_namelen = sizeof(struct sockaddr_can);
        msg.msg_control = (void *)&timestamp;
        msg.msg_controllen = sizeof(nanosecs_abs_t);
    }

    while (1) {
        if (with_timestamp) {
            iov.iov_base = (void *)&frame;
            iov.iov_len = sizeof(can_frame_t);
            ret = rt_dev_recvmmsg(s, &msg, 0);
        } else
            ret = rt_dev_recvfrom(s, (void *)&frame, sizeof(
can_frame_t), 0,
                                (struct sockaddr *)&addr, &addrlen);

        if (ret < 0) {
            switch (ret) {
                case -ETIMEDOUT:
                    if (verbose)
                        printf("rt_dev_recv: timed out");
                    continue;
                case -EBADF:
                    if (verbose)
                        printf("rt_dev_recv: aborted because socket was closed");
                    break;
                default:
                    fprintf(stderr, "rt_dev_recv: %s\n", strerror(-ret));
            }
            break;
        }

        if (print && (count % print) == 0) {
            printf("#d: (%d) ", count, addr.can_ifindex);
            if (with_timestamp && msg.msg_controllen) {
                if (timestamp_rel) {
                    printf("%lldns ", (long long)(timestamp - timestamp_prev));
                    timestamp_prev = timestamp;
                } else
                    printf("%lldns ", (long long)timestamp);
            }
            if (frame.can_id & CAN_ERR_FLAG)
                printf("!0x%08x!", frame.can_id & CAN_ERR_MASK);
            else if (frame.can_id & CAN_EFF_FLAG)

```

```

        printf("<0x%08x>", frame.can_id & CAN_EFF_MASK);
    else
        printf("<0x%03x>", frame.can_id & CAN_SFF_MASK);

    printf(" [%d]", frame.can_dlc);
    if (!(frame.can_id & CAN_RTR_FLAG))
        for (i = 0; i < frame.can_dlc; i++) {
            printf(" %02x", frame.data[i]);
        }
    if (frame.can_id & CAN_ERR_FLAG) {
        printf(" ERROR ");
        if (frame.can_id & CAN_ERR_BUSOFF)
            printf("bus-off");
        if (frame.can_id & CAN_ERR_CRTL)
            printf("controller problem");
    } else if (frame.can_id & CAN_RTR_FLAG)
        printf(" remote request");
    printf("\n");
}
count++;
}
}

int main(int argc, char **argv)
{
    int opt, ret;
    u_int32_t id, mask;
    u_int32_t err_mask = 0;
    struct ifreq ifr;
    char *ptr;
    char name[32];

    struct option long_options[] = {
        { "help", no_argument, 0, 'h' },
        { "verbose", no_argument, 0, 'v' },
        { "filter", required_argument, 0, 'f' },
        { "error", required_argument, 0, 'e' },
        { "timeout", required_argument, 0, 't' },
        { "timestamp", no_argument, 0, 'T' },
        { "timestamp-rel", no_argument, 0, 'R' },
        { 0, 0, 0, 0 },
    };

    mlockall(MCL_CURRENT | MCL_FUTURE);

    signal(SIGTERM, cleanup_and_exit);
    signal(SIGINT, cleanup_and_exit);

    while ((opt = getopt_long(argc, argv, "hve:f:t:p:RT",
                               long_options, NULL)) != -1) {
        switch (opt) {
            case 'h':
                print_usage(argv[0]);
                exit(0);

            case 'p':
                print = strtoul(optarg, NULL, 0);
                break;

            case 'v':
                verbose = 1;
                break;

            case 'e':
                err_mask = strtoul(optarg, NULL, 0);
                break;

            case 'f':
                ptr = optarg;
                while (1) {
                    id = strtoul(ptr, NULL, 0);
                    ptr = strchr(ptr, ':');
                    if (!ptr) {
                        fprintf(stderr, "filter must be applied in the form id:mask[:id:mask]...\n");
                        exit(1);
                    }
                    ptr++;
                    mask = strtoul(ptr, NULL, 0);
                    ptr = strchr(ptr, ':');
                    add_filter(id, mask);
                    if (!ptr)
                        break;
                    ptr++;
                }
                break;

            case 't':

```

```

        timeout = (nanosecs_rel_t)strtoul(optarg, NULL, 0) * 1000000;
        break;

    case 'R':
        timestamp_rel = 1;
    case 'T':
        with_timestamp = 1;
        break;

    default:
        fprintf(stderr, "Unknown option %c\n", opt);
        break;
    }
}

ret = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
if (ret < 0) {
    fprintf(stderr, "rt_dev_socket: %s\n", strerror(-ret));
    return -1;
}
s = ret;

if (argv[optind] == NULL) {
    if (verbose)
        printf("interface all\n");

    ifr.ifr_ifindex = 0;
} else {
    if (verbose)
        printf("interface %s\n", argv[optind]);

    strncpy(ifr.ifr_name, argv[optind], IFNAMSIZ);
    if (verbose)
        printf("s=%d, ifr_name=%s\n", s, ifr.ifr_name);

    ret = rt_dev_ioctl(s, SIOCGIFINDEX, &ifr);
    if (ret < 0) {
        fprintf(stderr, "rt_dev_ioctl GET_IFINDEX: %s\n", strerror(-ret));
        goto failure;
    }
}

if (err_mask) {
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW,
        CAN_RAW_ERR_FILTER,
        &err_mask, sizeof(err_mask));

    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }
    if (verbose)
        printf("Using err_mask=%#x\n", err_mask);
}

if (filter_count) {
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW,
        CAN_RAW_FILTER,
        &recv_filter, filter_count *
        sizeof(struct can_filter));

    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }
}

recv_addr.can_family = AF_CAN;
recv_addr.can_ifindex = ifr.ifr_ifindex;
ret = rt_dev_bind(s, (struct sockaddr *)&recv_addr,
    sizeof(struct sockaddr_can));
if (ret < 0) {
    fprintf(stderr, "rt_dev_bind: %s\n", strerror(-ret));
    goto failure;
}

if (timeout) {
    if (verbose)
        printf("Timeout: %lld ns\n", (long long)timeout);
    ret = rt_dev_ioctl(s, RTCAN_RTIOC_RCV_TIMEOUT, &timeout);
    if (ret) {
        fprintf(stderr, "rt_dev_ioctl RCV_TIMEOUT: %s\n", strerror(-ret));
        goto failure;
    }
}

if (with_timestamp) {
    ret = rt_dev_ioctl(s, RTCAN_RTIOC_TAKE_TIMESTAMP,
        RTCAN_TAKE_TIMESTAMPS);
}

```



```

        if (ret) {
            fprintf(stderr, "rt_dev_ioctl TAKE_TIMESTAMP: %s\n", strerror(-ret));
            goto failure;
        }
    }

    snprintf(name, sizeof(name), "rtcanrecv-%d", getpid());
    ret = rt_task_shadow(&rt_task_desc, name, 0, 0);
    if (ret) {
        fprintf(stderr, "rt_task_shadow: %s\n", strerror(-ret));
        goto failure;
    }

    rt_task();
    /* never returns */

failure:
    cleanup();
    return -1;
}

```

8.9 rtcanse.c

```

#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>
#include <time.h>
#include <errno.h>
#include <getopt.h>
#include <sys/mman.h>

#include <alchemy/task.h>
#include <alchemy/timer.h>

#include <rtdev/can.h>

extern int optind, opterr, optopt;

static void print_usage(char *prg)
{
    fprintf(stderr,
        "Usage: %s <can-interface> [Options] <can-msg>\n"
        "<can-msg> can consist of up to 8 bytes given as a space separated list\n"
        "Options:\n"
        " -i, --identifier=ID    CAN Identifier (default = 1)\n"
        " -r --rtr               send remote request\n"
        " -e --extended          send extended frame\n"
        " -l --loop=COUNT       send message COUNT times\n"
        " -c, --count            message count in data[0-3]\n"
        " -d, --delay=MS         delay in ms (default = 1ms)\n"
        " -s, --send             use send instead of sendto\n"
        " -t, --timeout=MS       timeout in ms\n"
        " -L, --loopback=0|1     switch local loopback off or on\n"
        " -v, --verbose          be verbose\n"
        " -p, --print=MODULO     print every MODULO message\n"
        " -h, --help             this help\n",
        prg);
}

RT_TASK rt_task_desc;

static int s=-1, dlc=0, rtr=0, extended=0, verbose=0, loops=1;
static SRTIME delay=1000000;
static int count=0, print=1, use_send=0, loopback=-1;
static nanosecs_rel_t timeout = 0;
static struct can_frame frame;
static struct sockaddr_can to_addr;

static void cleanup(void)
{
    int ret;

    if (verbose)
        printf("Cleaning up...\n");

    usleep(100000);

    if (s >= 0) {
        ret = rt_dev_close(s);
        s = -1;
    }
}

```

```

        if (ret) {
            fprintf(stderr, "rt_dev_close: %s\n", strerror(-ret));
        }
        exit(EXIT_SUCCESS);
    }
}

static void cleanup_and_exit(int sig)
{
    if (verbose)
        printf("Signal %d received\n", sig);
    cleanup();
    exit(0);
}

static void rt_task(void)
{
    int i, j, ret;

    for (i = 0; i < loops; i++) {
        rt_task_sleep(rt_timer_ns2ticks(delay));
        if (count)
            memcpy(&frame.data[0], &i, sizeof(i));
        /* Note: sendto avoids the definition of a receive filter list */
        if (use_send)
            ret = rt_dev_send(s, (void *)&frame, sizeof(can_frame_t), 0);
        else
            ret = rt_dev_sendto(s, (void *)&frame, sizeof(
can_frame_t), 0,
                                (struct sockaddr *)&to_addr, sizeof(to_addr));
        if (ret < 0) {
            switch (ret) {
                case -ETIMEDOUT:
                    if (verbose)
                        printf("rt_dev_send(to): timed out");
                    break;
                case -EBADF:
                    if (verbose)
                        printf("rt_dev_send(to): aborted because socket was closed");
                    break;
                default:
                    fprintf(stderr, "rt_dev_send: %s\n", strerror(-ret));
                    break;
            }
            i = loops;          /* abort */
            break;
        }
        if (verbose && (i % print) == 0) {
            if (frame.can_id & CAN_EFF_FLAG)
                printf("<0x%08x>", frame.can_id & CAN_EFF_MASK);
            else
                printf("<0x%03x>", frame.can_id & CAN_SFF_MASK);
            printf(" [%d]", frame.can_dlc);
            for (j = 0; j < frame.can_dlc; j++) {
                printf(" %02x", frame.data[j]);
            }
            printf("\n");
        }
    }
}

int main(int argc, char **argv)
{
    int i, opt, ret;
    struct ifreq ifr;
    char name[32];

    struct option long_options[] = {
        { "help", no_argument, 0, 'h' },
        { "identifier", required_argument, 0, 'i' },
        { "rtr", no_argument, 0, 'r' },
        { "extended", no_argument, 0, 'e' },
        { "verbose", no_argument, 0, 'v' },
        { "count", no_argument, 0, 'c' },
        { "print", required_argument, 0, 'p' },
        { "loop", required_argument, 0, 'l' },
        { "delay", required_argument, 0, 'd' },
        { "send", no_argument, 0, 's' },
        { "timeout", required_argument, 0, 't' },
        { "loopback", required_argument, 0, 'L' },
        { 0, 0, 0, 0 },
    };

    mlockall(MCL_CURRENT | MCL_FUTURE);

    signal(SIGTERM, cleanup_and_exit);
    signal(SIGINT, cleanup_and_exit);

```

```

frame.can_id = 1;

while ((opt = getopt_long(argc, argv, "hvi:l:red:t:cp:sL:",
                           long_options, NULL)) != -1) {
    switch (opt) {
        case 'h':
            print_usage(argv[0]);
            exit(0);

        case 'p':
            print = strtoul(optarg, NULL, 0);

        case 'v':
            verbose = 1;
            break;

        case 'c':
            count = 1;
            break;

        case 'l':
            loops = strtoul(optarg, NULL, 0);
            break;

        case 'i':
            frame.can_id = strtoul(optarg, NULL, 0);
            break;

        case 'r':
            rtr = 1;
            break;

        case 'e':
            extended = 1;
            break;

        case 'd':
            delay = strtoul(optarg, NULL, 0) * 1000000LL;
            break;

        case 's':
            use_send = 1;
            break;

        case 't':
            timeout = strtoul(optarg, NULL, 0) * 1000000LL;
            break;

        case 'L':
            loopback = strtoul(optarg, NULL, 0);
            break;

        default:
            fprintf(stderr, "Unknown option %c\n", opt);
            break;
    }
}

if (optind == argc) {
    print_usage(argv[0]);
    exit(0);
}

if (argv[optind] == NULL) {
    fprintf(stderr, "No Interface supplied\n");
    exit(-1);
}

if (verbose)
    printf("interface %s\n", argv[optind]);

ret = rt_dev_socket(PF_CAN, SOCK_RAW, CAN_RAW);
if (ret < 0) {
    fprintf(stderr, "rt_dev_socket: %s\n", strerror(-ret));
    return -1;
}
s = ret;

if (loopback >= 0) {
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW,
                           CAN_RAW_LOOPBACK,
                           &loopback, sizeof(loopback));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }
}

```

```

    if (verbose)
        printf("Using loopback=%d\n", loopback);
}

strncpy(ifr.ifr_name, argv[optind], IFNAMSIZ);
if (verbose)
    printf("s=%d, ifr_name=%s\n", s, ifr.ifr_name);

ret = rt_dev_ioctl(s, SIOCGIFINDEX, &ifr);
if (ret < 0) {
    fprintf(stderr, "rt_dev_ioctl: %s\n", strerror(-ret));
    goto failure;
}

memset(&to_addr, 0, sizeof(to_addr));
to_addr.can_ifindex = ifr.ifr_ifindex;
to_addr.can_family = AF_CAN;
if (use_send) {
    /* Suppress definition of a default receive filter list */
    ret = rt_dev_setsockopt(s, SOL_CAN_RAW,
        CAN_RAW_FILTER, NULL, 0);
    if (ret < 0) {
        fprintf(stderr, "rt_dev_setsockopt: %s\n", strerror(-ret));
        goto failure;
    }

    ret = rt_dev_bind(s, (struct sockaddr *)&to_addr, sizeof(to_addr));
    if (ret < 0) {
        fprintf(stderr, "rt_dev_bind: %s\n", strerror(-ret));
        goto failure;
    }
}

if (count)
    frame.can_dlc = sizeof(int);
else {
    for (i = optind + 1; i < argc; i++) {
        frame.data[dlc] = strtoul(argv[i], NULL, 0);
        dlc++;
        if (dlc == 8)
            break;
    }
    frame.can_dlc = dlc;
}

if (rtr)
    frame.can_id |= CAN_RTR_FLAG;

if (extended)
    frame.can_id |= CAN_EFF_FLAG;

if (timeout) {
    if (verbose)
        printf("Timeout: %lld ns\n", (long long)timeout);
    ret = rt_dev_ioctl(s, RTCAN_RTIOC_SND_TIMEOUT, &timeout);
    if (ret) {
        fprintf(stderr, "rt_dev_ioctl SND_TIMEOUT: %s\n", strerror(-ret));
        goto failure;
    }
}

snprintf(name, sizeof(name), "rtcansend-%d", getpid());
ret = rt_task_shadow(&rt_task_desc, name, 1, 0);
if (ret) {
    fprintf(stderr, "rt_task_shadow: %s\n", strerror(-ret));
    goto failure;
}

rt_task();

cleanup();
return 0;

failure:
cleanup();
return -1;
}

```

8.10 xddp-echo.c

```

/*
 * XDDP-based RT/NRT threads communication demo.
 */

```

```

* Real-time Xenomai threads and regular Linux threads may want to
* exchange data in a way that does not require the former to leave
* the real-time domain (i.e. secondary mode). Message pipes - as
* implemented by the RTDM-based XDDP protocol - are provided for this
* purpose.
*
* On the Linux domain side, pseudo-device files named /dev/rtp<minor>
* give regular POSIX threads access to non real-time communication
* endpoints, via the standard character-based I/O interface. On the
* Xenomai domain side, sockets may be bound to XDDP ports, which act
* as proxies to send and receive data to/from the associated
* pseudo-device files. Ports and pseudo-device minor numbers are
* paired, meaning that e.g. port 7 will proxy the traffic for
* /dev/rtp7. Therefore, port numbers may range from 0 to
* CONFIG_XENO_OPT_PIPE_NRDEV - 1.
*
* All data sent through a bound/connected XDDP socket via sendto(2) or
* write(2) will be passed to the peer endpoint in the Linux domain,
* and made available for reading via the standard read(2) system
* call. Conversely, all data sent using write(2) through the non
* real-time endpoint will be conveyed to the real-time socket
* endpoint, and made available to the recvfrom(2) or read(2) system
* calls.
*
* Both threads can use the bi-directional data path to send and
* receive datagrams in a FIFO manner, as illustrated by the simple
* echoing process implemented by this program.
*
* realtime_thread----->-----+
* => get socket |
* => bind socket to port 0 v
* => write traffic to NRT domain via sendto() |
* => read traffic from NRT domain via recvfrom() <--|---+
* | |
* regular_thread-----+ |
* => open /dev/rtp0 | ^
* => read traffic from RT domain via read() | |
* => echo traffic back to RT domain via write() +---+
*
* See Makefile in this directory for build directives.
*
* NOTE: XDDP is a replacement for the legacy RT_PIPE interface
* available from the native skin until Xenomai 3.
*/
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtdm/ipc.h>

pthread_t rt, nrt;

#define XDDP_PORT 0 /* [0..CONFIG-XENO_OPT_PIPE_NRDEV - 1] */

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *realtime_thread(void *arg)
{

```

```

struct sockaddr_ipc saddr;
int ret, s, n = 0, len;
struct timespec ts;
size_t poolsz;
char buf[128];

/*
 * Get a datagram socket to bind to the RT endpoint. Each
 * endpoint is represented by a port number within the XDDP
 * protocol namespace.
 */
s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_XDDP);
if (s < 0) {
    perror("socket");
    exit(EXIT_FAILURE);
}

/*
 * Set a local 16k pool for the RT endpoint. Memory needed to
 * convey datagrams will be pulled from this pool, instead of
 * Xenomai's system pool.
 */
poolsz = 16384; /* bytes */
ret = setsockopt(s, SOL_XDDP, XDDP_POOLSZ,
                &poolsz, sizeof(poolsz));
if (ret)
    fail("setsockopt");

/*
 * Bind the socket to the port, to setup a proxy to channel
 * traffic to/from the Linux domain.
 *
 * saddr.sipc_port specifies the port number to use.
 */
memset(&saddr, 0, sizeof(saddr));
saddr.sipc_family = AF_RTIPC;
saddr.sipc_port = XDDP_PORT;
ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
if (ret)
    fail("bind");

for (;;) {
    len = strlen(msg[n]);
    /*
     * Send a datagram to the NRT endpoint via the proxy.
     * We may pass a NULL destination address, since a
     * bound socket is assigned a default destination
     * address matching the binding address (unless
     * connect(2) was issued before bind(2), in which case
     * the former would prevail).
     */
    ret = sendto(s, msg[n], len, 0, NULL, 0);
    if (ret != len)
        fail("sendto");

    rt_printf("%s: sent %d bytes, \"%s\"\n",
              __FUNCTION__, ret, msg[n]);

    /* Read back packets echoed by the regular thread */
    ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
    if (ret <= 0)
        fail("recvfrom");

    rt_printf("  => \"%s\" echoed by peer\n", ret, buf);

    n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
    /*
     * We run in full real-time mode (i.e. primary mode),
     * so we have to let the system breathe between two
     * iterations.
     */
    ts.tv_sec = 0;
    ts.tv_nsec = 500000000; /* 500 ms */
    clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
}

return NULL;
}

static void *regular_thread(void *arg)
{
    char buf[128], *devname;
    int fd, ret;

    if (asprintf(&devname, "/dev/rtp%d", XDDP_PORT) < 0)
        fail("asprintf");

```

```

    fd = open(devname, O_RDWR);
    free(devname);
    if (fd < 0)
        fail("open");

    for (;;) {
        /* Get the next message from realtime_thread. */
        ret = read(fd, buf, sizeof(buf));
        if (ret <= 0)
            fail("read");

        /* Echo the message back to realtime_thread. */
        ret = write(fd, buf, ret);
        if (ret <= 0)
            fail("write");
    }

    return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(rt);
    pthread_cancel(nrt);
    signal(sig, SIG_DFL);
    pthread_join(rt, NULL);
    pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param rtparam = { .sched_priority = 42 };
    pthread_attr_t rtattr, regattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary (i.e. non real-time) mode when
     * writing output.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&rtattr);
    pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
    pthread_attr_setschedparam(&rtattr, &rtparam);

    errno = pthread_create(&rt, &rtattr, &realtime_thread, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&regattr);
    pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

    errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

8.11 xddp-label.c

```

/*
 * XDDP-based RT/NRT threads communication demo.
 */

```

```

* Real-time Xenomai threads and regular Linux threads may want to
* exchange data in a way that does not require the former to leave
* the real-time domain (i.e. secondary mode). Message pipes - as
* implemented by the RTDM-based XDDP protocol - are provided for this
* purpose.
*
* On the Linux domain side, pseudo-device files named /dev/rtp<minor>
* give regular POSIX threads access to non real-time communication
* endpoints, via the standard character-based I/O interface. On the
* Xenomai domain side, sockets may be bound to XDDP ports, which act
* as proxies to send and receive data to/from the associated
* pseudo-device files. Ports and pseudo-device minor numbers are
* paired, meaning that e.g. port 7 will proxy the traffic for
* /dev/rtp7. Therefore, port numbers may range from 0 to
* CONFIG_XENO_OPT_PIPE_NRDEV - 1.
*
* All data sent through a bound/connected XDDP socket via sendto(2) or
* write(2) will be passed to the peer endpoint in the Linux domain,
* and made available for reading via the standard read(2) system
* call. Conversely, all data sent using write(2) through the non
* real-time endpoint will be conveyed to the real-time socket
* endpoint, and made available to the recvfrom(2) or read(2) system
* calls.
*
* ASCII labels can be attached to bound ports, in order to connect
* sockets to them in a more descriptive way than using plain numeric
* port values.
*
* The example code below illustrates the following process:
*
* realtime_thread1----->-----+
* => get socket                      |
* => bind socket to port "xddp-demo  |
* => read traffic from NRT domain via recvfrom()  <---+
*                                     | |
* realtime_thread2-----+-----+ |
* => get socket                      | |
* => connect socket to port "xddp-demo" | |
* => write traffic to NRT domain via sendto()    v |
*                                     | ^
* regular_thread-----+-----+ |
* => open /proc/xenomai/registry/rtpc/xddp/xddp-demo | |
* => read traffic from RT domain via read()          | |
* => mirror traffic to RT domain via write()        +---+
*
* See Makefile in this directory for build directives.
*
* NOTE: XDDP is a replacement for the legacy RT_PIPE interface
* available from the native skin until Xenomai 3.
*/
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtm/ipc.h>

pthread_t rt1, rt2, nrt;

#define XDDP_PORT_LABEL "xddp-demo"

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{

```



```

        perror(reason);
        exit(EXIT_FAILURE);
    }

static void *realtime_thread1(void *arg)
{
    struct rtipc_port_label plabel;
    struct sockaddr_ipc saddr;
    char buf[128];
    int ret, s;

    /*
     * Get a datagram socket to bind to the RT endpoint. Each
     * endpoint is represented by a port number within the XDDP
     * protocol namespace.
     */
    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
    if (s < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
    }

    /*
     * Set a port label. This name will be registered when
     * binding, in addition to the port number (if given).
     */
    strcpy(plabel.label, XDDP_PORT_LABEL);
    ret = setsockopt(s, SOL_XDDP, XDDP_LABEL,
                    &plabel, sizeof(plabel));
    if (ret)
        fail("setsockopt");

    /*
     * Bind the socket to the port, to setup a proxy to channel
     * traffic to/from the Linux domain. Assign that port a label,
     * so that peers may use a descriptive information to locate
     * it. For instance, the pseudo-device matching our RT
     * endpoint will appear as
     * /proc/xenomai/registry/rtipc/xddp/<XDDP_PORT_LABEL> in the
     * Linux domain, once the socket is bound.
     */
    /* saddr.sipc_port specifies the port number to use. If -1 is
     * passed, the XDDP driver will auto-select an idle port.
     */
    memset(&saddr, 0, sizeof(saddr));
    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = -1;
    ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
    if (ret)
        fail("bind");

    for (;;) {
        /* Get packets relayed by the regular thread */
        ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
        if (ret <= 0)
            fail("recvfrom");

        rt_printf("%s: \"%s\" relayed by peer\n", __FUNCTION__, ret, buf);
    }

    return NULL;
}

static void *realtime_thread2(void *arg)
{
    struct rtipc_port_label plabel;
    struct sockaddr_ipc saddr;
    int ret, s, n = 0, len;
    struct timespec ts;
    struct timeval tv;
    socklen_t addrlen;

    s = socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
    if (s < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
    }

    /*
     * Set the socket timeout; it will apply when attempting to
     * connect to a labeled port, and to recvfrom() calls. The
     * following setup tells the XDDP driver to wait for at most
     * one second until a socket is bound to a port using the same
     * label, or return with a timeout error.
     */
    tv.tv_sec = 1;
    tv.tv_usec = 0;
    ret = setsockopt(s, SOL_SOCKET, SO_RCVTIMEO,

```

```

        &tv, sizeof(tv));
    if (ret)
        fail("setsockopt");

    /*
     * Set a port label. This name will be used to find the peer
     * when connecting, instead of the port number.
     */
    strcpy(plabel.label, XDDP_PORT_LABEL);
    ret = setsockopt(s, SOL_XDDP, XDDP_LABEL,
                    &plabel, sizeof(plabel));
    if (ret)
        fail("setsockopt");

    memset(&saddr, 0, sizeof(saddr));
    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = -1; /* Tell XDDP to search by label. */
    ret = connect(s, (struct sockaddr *)&saddr, sizeof(saddr));
    if (ret)
        fail("connect");

    /*
     * We succeeded in making the port our default destination
     * address by using its label, but we don't know its actual
     * port number yet. Use getpeername() to retrieve it.
     */
    addrlen = sizeof(saddr);
    ret = getpeername(s, (struct sockaddr *)&saddr, &addrlen);
    if (ret || addrlen != sizeof(saddr))
        fail("getpeername");

    rt_printf("%s: NRT peer is reading from /dev/rtp%d\n",
              __FUNCTION__, saddr.sipc_port);

    for (;;) {
        len = strlen(msg[n]);
        /*
         * Send a datagram to the NRT endpoint via the proxy.
         * We may pass a NULL destination address, since the
         * socket was successfully assigned the proper default
         * address via connect(2).
         */
        ret = sendto(s, msg[n], len, 0, NULL, 0);
        if (ret != len)
            fail("sendto");

        rt_printf("%s: sent %d bytes, \"%s\"\n",
                  __FUNCTION__, ret, msg[n]);

        n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
        /*
         * We run in full real-time mode (i.e. primary mode),
         * so we have to let the system breathe between two
         * iterations.
         */
        ts.tv_sec = 0;
        ts.tv_nsec = 500000000; /* 500 ms */
        clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
    }

    return NULL;
}

static void *regular_thread(void *arg)
{
    char buf[128], *devname;
    int fd, ret;

    if (asprintf(&devname,
                "/proc/xenomai/registry/rtipc/xddp/%s",
                XDDP_PORT_LABEL) < 0)
        fail("asprintf");

    fd = open(devname, O_RDWR);
    free(devname);
    if (fd < 0)
        fail("open");

    for (;;) {
        /* Get the next message from realtime_thread2. */
        ret = read(fd, buf, sizeof(buf));
        if (ret <= 0)
            fail("read");

        /* Relay the message to realtime_thread1. */
        ret = write(fd, buf, ret);
        if (ret <= 0)

```

```

        fail("write");
    }

    return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(rt1);
    pthread_cancel(rt2);
    pthread_cancel(nrt);
    signal(sig, SIG_DFL);
    pthread_join(rt1, NULL);
    pthread_join(rt2, NULL);
    pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param rtparam = { .sched_priority = 42 };
    pthread_attr_t rtattr, regattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary (i.e. non real-time) mode when
     * writing output.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&rtattr);
    pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
    pthread_attr_setschedparam(&rtattr, &rtparam);

    /* Both real-time threads have the same attribute set. */

    errno = pthread_create(&rt1, &rtattr, &realtime_thread1, NULL);
    if (errno)
        fail("pthread_create");

    errno = pthread_create(&rt2, &rtattr, &realtime_thread2, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&regattr);
    pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

    errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

8.12 xddp-stream.c

```

/*
 * XDDP-based RT/NRT threads communication demo.
 *
 * Real-time Xenomai threads and regular Linux threads may want to
 * exchange data in a way that does not require the former to leave
 * the real-time domain (i.e. secondary mode). Message pipes - as
 * implemented by the RTDM-based XDDP protocol - are provided for this
 * purpose.
 */

```

```

* On the Linux domain side, pseudo-device files named /dev/rtp<minor>
* give regular POSIX threads access to non real-time communication
* endpoints, via the standard character-based I/O interface. On the
* Xenomai domain side, sockets may be bound to XDDP ports, which act
* as proxies to send and receive data to/from the associated
* pseudo-device files. Ports and pseudo-device minor numbers are
* paired, meaning that e.g. port 7 will proxy the traffic for
* /dev/rtp7. Therefore, port numbers may range from 0 to
* CONFIG_XENO_OPT_PIPE_NRDEV - 1.
*
* All data sent through a bound/connected XDDP socket via sendto(2) or
* write(2) will be passed to the peer endpoint in the Linux domain,
* and made available for reading via the standard read(2) system
* call. Conversely, all data sent using write(2) through the non
* real-time endpoint will be conveyed to the real-time socket
* endpoint, and made available to the recvfrom(2) or read(2) system
* calls.
*
* In addition to sending datagrams, real-time threads may stream data
* in a byte-oriented mode through the proxy as well. This increases
* the bandwidth and reduces the overhead, when a lot of data has to
* flow down to the Linux domain, if keeping the message boundaries is
* not required. The example code below illustrates such use.
*
* realtime_thread----->-----+
* => get socket                      |
* => bind socket to port 0           v
* => write scattered traffic to NRT domain via sendto() |
* => read traffic from NRT domain via recvfrom() <--|---+
*                                     | |
* regular_thread-----+ |
* => open /dev/rtp0        | ^
* => read traffic from RT domain via read() | |
* => echo traffic back to RT domain via write() +---+
*
* See Makefile in this directory for build directives.
*
* NOTE: XDDP is a replacement for the legacy RT_PIPE interface
* available from the native skin until Xenomai 3.
*/
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>
#include <malloc.h>
#include <pthread.h>
#include <fcntl.h>
#include <errno.h>
#include <rtdm/ipc.h>

pthread_t rt, nrt;

#define XDDP_PORT 0 /* [0..CONFIG-XENO_OPT_PIPE_NRDEV - 1] */

static const char *msg[] = {
    "Surfing With The Alien",
    "Lords of Karma",
    "Banana Mango",
    "Psycho Monkey",
    "Luminous Flesh Giants",
    "Moroccan Sunset",
    "Satch Boogie",
    "Flying In A Blue Dream",
    "Ride",
    "Summer Song",
    "Speed Of Light",
    "Crystal Planet",
    "Raspberry Jam Delta-V",
    "Champagne?",
    "Clouds Race Across The Sky",
    "Engines Of Creation"
};

static void fail(const char *reason)
{
    perror(reason);
    exit(EXIT_FAILURE);
}

static void *realtime_thread(void *arg)
{
    struct sockaddr_ipc saddr;
    int ret, s, n = 0, len, b;
    struct timespec ts;
    size_t streams;

```

```

char buf[128];

/*
 * Get a datagram socket to bind to the RT endpoint. Each
 * endpoint is represented by a port number within the XDDP
 * protocol namespace.
 */
s = socket(AF_RTIPC, SOCK_DGRAM, IPPROTO_XDDP);
if (s < 0) {
    perror("socket");
    exit(EXIT_FAILURE);
}

/*
 * Tell the XDDP driver that we will use the streaming
 * capabilities on this socket. To this end, we have to
 * specify the size of the streaming buffer, as a count of
 * bytes. The real-time output will be buffered up to that
 * amount, and sent as a single datagram to the NRT endpoint
 * when fully gathered, or when another source port attempts
 * to send data to the same endpoint. Passing a null size
 * would disable streaming.
 */
streamsz = 1024; /* bytes */
ret = setsockopt(s, SOL_XDDP, XDDP_BUFSZ,
                &streamsz, sizeof(streamsz));
if (ret)
    fail("setsockopt");

/*
 * Bind the socket to the port, to setup a proxy to channel
 * traffic to/from the Linux domain.
 *
 * saddr.sipc_port specifies the port number to use.
 */
memset(&saddr, 0, sizeof(saddr));
saddr.sipc_family = AF_RTIPC;
saddr.sipc_port = XDDP_PORT;
ret = bind(s, (struct sockaddr *)&saddr, sizeof(saddr));
if (ret)
    fail("bind");

for (;;) {
    len = strlen(msg[n]);
    /*
     * Send a datagram to the NRT endpoint via the proxy.
     * The output is artificially scattered in separate
     * one-byte sendings, to illustrate the use of
     * MSG_MORE.
     */
    for (b = 0; b < len; b++) {
        ret = sendto(s, msg[n] + b, 1, MSG_MORE, NULL, 0);
        if (ret != 1)
            fail("sendto");
    }

    rt_printf("%s: sent (scattered) %d-bytes message, \"%s\\n",
              __FUNCTION__, len, len, msg[n]);

    /* Read back packets echoed by the regular thread */
    ret = recvfrom(s, buf, sizeof(buf), 0, NULL, 0);
    if (ret <= 0)
        fail("recvfrom");

    rt_printf("  => \"%s\\n" echoed by peer\\n", ret, buf);

    n = (n + 1) % (sizeof(msg) / sizeof(msg[0]));
    /*
     * We run in full real-time mode (i.e. primary mode),
     * so we have to let the system breathe between two
     * iterations.
     */
    ts.tv_sec = 0;
    ts.tv_nsec = 500000000; /* 500 ms */
    clock_nanosleep(CLOCK_REALTIME, 0, &ts, NULL);
}

return NULL;
}

static void *regular_thread(void *arg)
{
    char buf[128], *devname;
    int fd, ret;

    if (asprintf(&devname, "/dev/rtp%d", XDDP_PORT) < 0)
        fail("asprintf");

```

```

    fd = open(devname, O_RDWR);
    free(devname);
    if (fd < 0)
        fail("open");

    for (;;) {
        /* Get the next message from realtime_thread. */
        ret = read(fd, buf, sizeof(buf));
        if (ret <= 0)
            fail("read");

        /* Echo the message back to realtime_thread. */
        ret = write(fd, buf, ret);
        if (ret <= 0)
            fail("write");
    }

    return NULL;
}

static void cleanup_upon_sig(int sig)
{
    pthread_cancel(rt);
    pthread_cancel(nrt);
    signal(sig, SIG_DFL);
    pthread_join(rt, NULL);
    pthread_join(nrt, NULL);
}

int main(int argc, char **argv)
{
    struct sched_param rtparam = { .sched_priority = 42 };
    pthread_attr_t rtattr, regattr;
    sigset_t mask, oldmask;

    mlockall(MCL_CURRENT | MCL_FUTURE);

    sigemptyset(&mask);
    sigaddset(&mask, SIGINT);
    signal(SIGINT, cleanup_upon_sig);
    sigaddset(&mask, SIGTERM);
    signal(SIGTERM, cleanup_upon_sig);
    sigaddset(&mask, SIGHUP);
    signal(SIGHUP, cleanup_upon_sig);
    pthread_sigmask(SIG_BLOCK, &mask, &oldmask);

    /*
     * This is a real-time compatible printf() package from
     * Xenomai's RT Development Kit (RTDK), that does NOT cause
     * any transition to secondary (i.e. non real-time) mode when
     * writing output.
     */
    rt_print_auto_init(1);

    pthread_attr_init(&rtattr);
    pthread_attr_setdetachstate(&rtattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&rtattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&rtattr, SCHED_FIFO);
    pthread_attr_setschedparam(&rtattr, &rtparam);

    errno = pthread_create(&rt, &rtattr, &realtime_thread, NULL);
    if (errno)
        fail("pthread_create");

    pthread_attr_init(&regattr);
    pthread_attr_setdetachstate(&regattr, PTHREAD_CREATE_JOINABLE);
    pthread_attr_setinheritsched(&regattr, PTHREAD_EXPLICIT_SCHED);
    pthread_attr_setschedpolicy(&regattr, SCHED_OTHER);

    errno = pthread_create(&nrt, &regattr, &regular_thread, NULL);
    if (errno)
        fail("pthread_create");

    sigsuspend(&oldmask);

    return 0;
}

```

Index

- A4L_RNG_FACTOR
 - uapi/analogy.h, [504](#)
- a4l_add_subd
 - Subdevice management services, [233](#)
- a4l_alloc_subd
 - Subdevice management services, [233](#)
- a4l_async_read
 - Asynchronous acquisition API, [280](#)
- a4l_async_write
 - Asynchronous acquisition API, [280](#)
- a4l_buf_commit_absget
 - Buffer management services, [236](#)
- a4l_buf_commit_absput
 - Buffer management services, [237](#)
- a4l_buf_commit_get
 - Buffer management services, [237](#)
- a4l_buf_commit_put
 - Buffer management services, [237](#)
- a4l_buf_count
 - Buffer management services, [238](#)
- a4l_buf_evt
 - Buffer management services, [238](#)
- a4l_buf_get
 - Buffer management services, [239](#)
- a4l_buf_prepare_absget
 - Buffer management services, [239](#)
- a4l_buf_prepare_absput
 - Buffer management services, [239](#)
- a4l_buf_prepare_get
 - Buffer management services, [240](#)
- a4l_buf_prepare_put
 - Buffer management services, [240](#)
- a4l_buf_put
 - Buffer management services, [240](#)
- a4l_channel, [397](#)
 - flags, [397](#)
 - nb_bits, [397](#)
- a4l_channels_desc, [397](#)
 - chans, [398](#)
 - length, [398](#)
 - mode, [398](#)
- a4l_close
 - Descriptor API, [284](#)
- a4l_cmd_desc, [398](#)
 - idx_subd, [399](#)
- a4l_config_subd
 - Synchronous acquisition API, [298](#)
- a4l_descriptor, [400](#)
 - board_name, [400](#)
 - fd, [400](#)
 - idx_read_subd, [400](#)
 - idx_write_subd, [400](#)
 - magic, [401](#)
 - nb_subd, [401](#)
 - sbddata, [401](#)
 - sbsize, [401](#)
- a4l_driver, [401](#)
- a4l_dtoraw
 - Range / conversion API, [287](#)
- a4l_fill_desc
 - Descriptor API, [285](#)
- a4l_find_range
 - Range / conversion API, [288](#)
- a4l_free_irq
 - Interrupt management services, [243](#)
- a4l_ftoraw
 - Range / conversion API, [288](#)
- a4l_get_bufsize
 - Asynchronous acquisition API, [276](#)
- a4l_get_chan
 - Buffer management services, [242](#)
- a4l_get_chinfo
 - Descriptor API, [285](#)
- a4l_get_cmd
 - Buffer management services, [242](#)
- a4l_get_irq
 - Interrupt management services, [243](#)
- a4l_get_rnginfo
 - Descriptor API, [285](#)
- a4l_get_subd
 - Subdevice management services, [233](#)
- a4l_get_subdinfo
 - Descriptor API, [286](#)
- a4l_get_time
 - Misc services, [246](#)
- a4l_instruction, [402](#)
 - idx_subd, [402](#)
- a4l_instruction_list, [403](#)
- a4l_mark_bufwr
 - Asynchronous acquisition API, [276](#)
- a4l_mmap
 - Asynchronous acquisition API, [277](#)
- a4l_open
 - Descriptor API, [286](#)
- a4l_poll
 - Asynchronous acquisition API, [277](#)
- a4l_range, [403](#)
 - flags, [404](#)

- max, [404](#)
 - min, [404](#)
- a4l_rawtod
 - Range / conversion API, [289](#)
- a4l_rawtof
 - Range / conversion API, [289](#)
- a4l_rawtoul
 - Range / conversion API, [289](#)
- a4l_register_drv
 - Driver management services, [229](#)
- a4l_request_irq
 - Interrupt management services, [243](#)
- a4l_set_bufsize
 - Asynchronous acquisition API, [278](#)
- a4l_sizeof_chan
 - Range / conversion API, [290](#)
- a4l_sizeof_subd
 - Range / conversion API, [290](#)
- a4l_snd_cancel
 - Asynchronous acquisition API, [278](#)
- a4l_snd_command
 - Asynchronous acquisition API, [279](#)
- a4l_snd_insn
 - Synchronous acquisition API, [295](#)
- a4l_snd_insnlist
 - Synchronous acquisition API, [295](#)
- a4l_subdevice, [404](#)
- a4l_sync_dio
 - Synchronous acquisition API, [299](#)
- a4l_sync_read
 - Synchronous acquisition API, [299](#)
- a4l_sync_write
 - Synchronous acquisition API, [300](#)
- a4l_sys_attach
 - Attach / detach Syscall API, [305](#)
- a4l_sys_bufcfg
 - Attach / detach Syscall API, [305](#)
- a4l_sys_close
 - Basic Syscall API, [303](#)
- a4l_sys_desc
 - Descriptor Syscall API, [282](#)
- a4l_sys_detach
 - Attach / detach Syscall API, [306](#)
- a4l_sys_open
 - Basic Syscall API, [303](#)
- a4l_sys_read
 - Basic Syscall API, [303](#)
- a4l_sys_write
 - Basic Syscall API, [304](#)
- a4l_ultoraw
 - Range / conversion API, [291](#)
- a4l_unregister_drv
 - Driver management services, [229](#)
- Alarm services, [307](#)
 - rt_alarm_create, [307](#)
 - rt_alarm_delete, [308](#)
 - rt_alarm_inquire, [308](#)
 - rt_alarm_start, [309](#)
 - rt_alarm_stop, [309](#)
- Alchemy API, [339](#)
- Analogy interface, [227](#)
- Analogy user API, [301](#)
- Asynchronous acquisition API, [274](#), [280](#)
 - a4l_async_read, [280](#)
 - a4l_async_write, [280](#)
 - a4l_get_bufsize, [276](#)
 - a4l_mark_bufw, [276](#)
 - a4l_mmap, [277](#)
 - a4l_poll, [277](#)
 - a4l_set_bufsize, [278](#)
 - a4l_snd_cancel, [278](#)
 - a4l_snd_command, [279](#)
- Asynchronous Procedure Calls, [82](#)
 - xnipc_alloc, [82](#)
 - xnipc_free, [83](#)
 - xnipc_schedule, [83](#)
- atomic_long_t, [406](#)
- Attach / detach Syscall API, [305](#)
 - a4l_sys_attach, [305](#)
 - a4l_sys_bufcfg, [305](#)
 - a4l_sys_detach, [306](#)
- B_PRIO
 - Buffer services, [312](#)
- BUFP_BUFSZ
 - Real-time IPC, [69](#)
- BUFP_LABEL
 - Real-time IPC, [70](#)
- Basic Syscall API, [303](#)
 - a4l_sys_close, [303](#)
 - a4l_sys_open, [303](#)
 - a4l_sys_read, [303](#)
 - a4l_sys_write, [304](#)
- begin
 - xnfile_regular_ops, [425](#)
 - xnfile_snapshot_ops, [431](#)
- Big dual kernel lock, [18](#)
 - cobalt_atomic_enter, [18](#)
 - cobalt_atomic_leave, [18](#)
- bind__AF_RTIPC
 - Real-time IPC, [76](#)
- board_name
 - a4l_descriptor, [400](#)
- Buffer descriptor, [85](#)
 - xnbufd_copy_from_kmem, [87](#)
 - xnbufd_copy_to_kmem, [88](#)
 - xnbufd_invalidate, [88](#)
 - xnbufd_map_kread, [89](#)
 - xnbufd_map_kwrite, [89](#)
 - xnbufd_map_uread, [89](#)
 - xnbufd_map_uwrite, [90](#)
 - xnbufd_reset, [90](#)
 - xnbufd_unmap_kread, [90](#)
 - xnbufd_unmap_kwrite, [91](#)
 - xnbufd_unmap_uread, [91](#)
 - xnbufd_unmap_uwrite, [91](#)
- Buffer management services, [235](#)

- a4l_buf_commit_absget, 236
- a4l_buf_commit_absput, 237
- a4l_buf_commit_get, 237
- a4l_buf_commit_put, 237
- a4l_buf_count, 238
- a4l_buf_evt, 238
- a4l_buf_get, 239
- a4l_buf_prepare_absget, 239
- a4l_buf_prepare_absput, 239
- a4l_buf_prepare_get, 240
- a4l_buf_prepare_put, 240
- a4l_buf_put, 240
- a4l_get_chan, 242
- a4l_get_cmd, 242
- Buffer services, 311
 - B_PRIO, 312
 - rt_buffer_bind, 312
 - rt_buffer_clear, 313
 - rt_buffer_create, 313
 - rt_buffer_delete, 314
 - rt_buffer_inquire, 314
 - rt_buffer_read, 315
 - rt_buffer_read_timed, 315
 - rt_buffer_read_until, 316
 - rt_buffer_unbind, 317
 - rt_buffer_write, 317
 - rt_buffer_write_timed, 317
 - rt_buffer_write_until, 318
- CAN Devices
 - CAN_BITTIME_BTR, 48
 - CAN_BITTIME_STD, 48
 - CAN_MODE_SLEEP, 48
 - CAN_MODE_START, 48
 - CAN_MODE_STOP, 48
 - CAN_STATE_ACTIVE, 49
 - CAN_STATE_BUS_OFF, 49
 - CAN_STATE_BUS_PASSIVE, 49
 - CAN_STATE_BUS_WARNING, 49
 - CAN_STATE_ERROR_ACTIVE, 49
 - CAN_STATE_ERROR_PASSIVE, 49
 - CAN_STATE_ERROR_WARNING, 49
 - CAN_STATE_SCANNING_BAUDRATE, 49
 - CAN_STATE_SLEEPING, 49
 - CAN_STATE_STOPPED, 49
- CAN_BITTIME_BTR
 - CAN Devices, 48
- CAN_BITTIME_STD
 - CAN Devices, 48
- CAN_MODE_SLEEP
 - CAN Devices, 48
- CAN_MODE_START
 - CAN Devices, 48
- CAN_MODE_STOP
 - CAN Devices, 48
- CAN_STATE_ACTIVE
 - CAN Devices, 49
- CAN_STATE_BUS_OFF
 - CAN Devices, 49
- CAN_STATE_BUS_PASSIVE
 - CAN Devices, 49
- CAN_STATE_BUS_WARNING
 - CAN Devices, 49
- CAN_STATE_ERROR_ACTIVE
 - CAN Devices, 49
- CAN_STATE_ERROR_PASSIVE
 - CAN Devices, 49
- CAN_STATE_ERROR_WARNING
 - CAN Devices, 49
- CAN_STATE_SCANNING_BAUDRATE
 - CAN Devices, 49
- CAN_STATE_SLEEPING
 - CAN Devices, 49
- CAN_STATE_STOPPED
 - CAN Devices, 49
- CAN Devices, 29
 - CAN_BITTIME_TYPE, 48
 - CAN_MODE, 48
 - CAN_RAW_FILTER, 38
 - CAN_RAW_LOOPBACK, 39
 - CAN_STATE, 48
 - can_filter_t, 48
 - can_frame_t, 48
 - SIOCGCANBAUDRATE, 42
 - SIOCGCANCTRLMODE, 42
 - SIOCGCANSTATE, 43
 - SIOCGIFINDEX, 44
 - SIOCSCANBAUDRATE, 44
 - SIOCSCANCTRLMODE, 45
 - SIOCSCANMODE, 46
 - SOL_CAN_RAW, 47
- CAN_BITTIME_TYPE
 - CAN Devices, 48
- CAN_MODE
 - CAN Devices, 48
- CAN_RAW_ERR_FILTER
 - CAN Devices, 38
- CAN_RAW_FILTER
 - CAN Devices, 38
- CAN_RAW_LOOPBACK
 - CAN Devices, 39
- CAN_STATE
 - CAN Devices, 48
- can_bittime, 406
- can_bittime_btr, 407
- can_bittime_std, 407
- can_filter, 408
 - can_id, 408
 - can_mask, 408
- can_filter_t
 - CAN Devices, 48
- can_frame, 409
 - can_id, 409
- can_frame_t
 - CAN Devices, 48
- can_id
 - can_filter, 408

- can_frame, 409
- can_ifindex
 - sockaddr_can, 421
- can_mask
 - can_filter, 408
- Channels and ranges, 15
- chans
 - a4l_channels_desc, 398
- clear_boost
 - Thread synchronization services., 199
- Clock Services, 137
 - rtdm_clock_read, 137
 - rtdm_clock_read_monotonic, 137
- Clock services., 93
 - xnclock_adjust, 93
 - xnclock_deregister, 94
 - xnclock_register, 94
 - xnclock_tick, 94
- Clocks and timers, 247
 - timer_create, 247
 - timer_gettime, 248
 - timer_settime, 249
- close__AF_RTIPC
 - Real-time IPC, 77
- Cobalt/POSIX real-time API, 253
- cobalt_atomic_enter
 - Big dual kernel lock, 18
- cobalt_atomic_leave
 - Big dual kernel lock, 18
- Condition variable services, 319
 - rt_cond_bind, 320
 - rt_cond_broadcast, 320
 - rt_cond_create, 321
 - rt_cond_delete, 321
 - rt_cond_inquire, 322
 - rt_cond_signal, 322
 - rt_cond_unbind, 323
 - rt_cond_wait, 323
 - rt_cond_wait_timed, 323
 - rt_cond_wait_until, 324
- Condition variables, 251
 - pthread_cond_destroy, 251
 - pthread_cond_init, 252
- connect__AF_RTIPC
 - Real-time IPC, 78
- curr
 - xnsched, 422
- databuf
 - xnvfile_snapshot_iterator, 430
- date
 - rt_timer_info, 415
- Debugging services, 96
- Descriptor API, 284
 - a4l_close, 284
 - a4l_fill_desc, 285
 - a4l_get_chinfo, 285
 - a4l_get_rnginfo, 285
 - a4l_get_subdinfo, 286
 - a4l_open, 286
- Descriptor Syscall API, 282
 - a4l_sys_desc, 282
- device
 - rtdm_dev_context, 416
- Device Profiles, 185
 - RTIOC_DEVICE_INFO, 186
 - RTIOC_PURGE, 186
- Device Registration Services, 134
 - RTDM_EXCLUSIVE, 135
 - RTDM_NAMED_DEVICE, 135
 - rtdm_dev_register, 135
 - rtdm_dev_unregister, 136
- Driver API, 228
- Driver management services, 229
 - a4l_register_drv, 229
 - a4l_unregister_drv, 229
- Driver programming interface, 118
- Driver to driver services, 128
 - rtdm_accept, 129
 - rtdm_bind, 129
 - rtdm_close, 129
 - rtdm_connect, 129
 - rtdm_getpeername, 130
 - rtdm_getsockname, 130
 - rtdm_getsockopt, 130
 - rtdm_ioctl, 130
 - rtdm_listen, 130
 - rtdm_open, 131
 - rtdm_read, 131
 - rtdm_recv, 131
 - rtdm_recvfrom, 131
 - rtdm_recvmsg, 131
 - rtdm_send, 131
 - rtdm_sendmsg, 132
 - rtdm_sendto, 132
 - rtdm_setsockopt, 132
 - rtdm_shutdown, 132
 - rtdm_socket, 132
 - rtdm_write, 133
- driver.h
 - rtdm_fd_device, 458
 - rtdm_fd_is_user, 458
 - rtdm_fd_to_private, 459
 - rtdm_open_handler_t, 457
 - rtdm_private_to_fd, 459
 - rtdm_socket_handler_t, 458
- Dynamic memory allocation services., 97
 - xnheap_alloc, 98
 - xnheap_destroy, 98
 - xnheap_extend, 99
 - xnheap_free, 99
 - xnheap_init, 99
 - xnheap_set_label, 100
 - xnheap_test_and_free, 100
- EV_ANY
 - Event flag group services, 326
- EV_PRIO

- Event flag group services, 326
- end
 - xnvfile_regular_ops, 426
 - xnvfile_snapshot_ops, 432
- endfn
 - xnvfile_snapshot_iterator, 430
- Event flag group services, 325
 - EV_ANY, 326
 - EV_PRIO, 326
 - rt_event_bind, 326
 - rt_event_clear, 327
 - rt_event_create, 327
 - rt_event_delete, 328
 - rt_event_inquire, 328
 - rt_event_signal, 329
 - rt_event_unbind, 329
 - rt_event_wait, 329
 - rt_event_wait_timed, 330
 - rt_event_wait_until, 331
- Event Services, 154
 - rtdm_event_clear, 154
 - rtdm_event_destroy, 155
 - rtdm_event_init, 155
 - rtdm_event_pulse, 155
 - rtdm_event_select_bind, 156
 - rtdm_event_signal, 156
 - rtdm_event_timedwait, 157
 - rtdm_event_wait, 157
- fd
 - a4l_descriptor, 400
- File descriptors events multiplexing services, 189
 - xnselect, 190
 - xnselect_bind, 190
 - xnselect_destroy, 191
 - xnselect_init, 191
 - xnselect_signal, 192
 - xnselector_destroy, 192
 - xnselector_init, 192
- flags
 - a4l_channel, 397
 - a4l_range, 404
- get
 - xnvfile_lock_ops, 423
- getpeername__AF_RTIPC
 - Real-time IPC, 78
- getsockname__AF_RTIPC
 - Real-time IPC, 79
- getsockopt__AF_RTIPC
 - Real-time IPC, 79
- H_PRIO
 - Heap management services, 333
- Heap management services, 332
 - H_PRIO, 333
 - rt_heap_alloc, 333
 - rt_heap_alloc_timed, 333
 - rt_heap_alloc_until, 334
 - rt_heap_bind, 334
 - rt_heap_create, 335
 - rt_heap_delete, 336
 - rt_heap_free, 336
 - rt_heap_inquire, 338
 - rt_heap_unbind, 338
- Helpers for in-kernel arithmetics, 84
- htimer
 - xnsched, 422
- IPCPROTO_BUF
 - Real-time IPC, 76
- IPCPROTO_IDDP
 - Real-time IPC, 76
- IPCPROTO_IPC
 - Real-time IPC, 76
- IPCPROTO_XDDP
 - Real-time IPC, 76
- IDDP_LABEL
 - Real-time IPC, 70
- IDDP_POOLSZ
 - Real-time IPC, 71
- idx_read_subd
 - a4l_descriptor, 400
- idx_subd
 - a4l_cmd_desc, 399
 - a4l_instruction, 402
- idx_write_subd
 - a4l_descriptor, 400
- include/cobalt/kernel/apc.h, 440
- include/cobalt/kernel/bufd.h, 441
- include/cobalt/kernel/clock.h, 443
- include/cobalt/kernel/map.h, 444
- include/cobalt/kernel/ppd.h, 446
- include/cobalt/kernel/registry.h, 447
- include/cobalt/kernel/rtdm/analogy/buffer.h, 435
- include/cobalt/kernel/rtdm/analogy/channel_ -
range.h, 448
- include/cobalt/kernel/rtdm/analogy/context.h, 450
- include/cobalt/kernel/rtdm/analogy/device.h, 451
- include/cobalt/kernel/rtdm/analogy/driver.h, 452
- include/cobalt/kernel/rtdm/analogy/instruction.h,
459
- include/cobalt/kernel/rtdm/analogy/rtdm_helpers.-
h, 460
- include/cobalt/kernel/rtdm/analogy/subdevice.h,
461
- include/cobalt/kernel/rtdm/analogy/transfer.h, 462
- include/cobalt/kernel/rtdm/can.h, 463
- include/cobalt/kernel/rtdm/cobalt.h, 471
- include/cobalt/kernel/rtdm/driver.h, 453
- include/cobalt/kernel/rtdm/ipc.h, 471
- include/cobalt/kernel/rtdm/rtdm.h, 476
- include/cobalt/kernel/rtdm/serial.h, 479
- include/cobalt/kernel/rtdm/testing.h, 489
- include/cobalt/kernel/sched-idle.h, 492
- include/cobalt/kernel/sched-rt.h, 493
- include/cobalt/kernel/sched-sporadic.h, 493
- include/cobalt/kernel/sched.h, 437

- xnsched_rotate, 438
 - xnsched_run, 439
- include/cobalt/kernel/select.h, 494
- include/cobalt/kernel/timer.h, 436
- include/cobalt/kernel/vfile.h, 496
- include/rtdm/analog.h, 498
- include/rtdm/can.h, 463
- include/rtdm/ipc.h, 472
- include/rtdm/rtdm.h, 476
- include/rtdm/serial.h, 479
- include/rtdm/testing.h, 489
- include/rtdm/uapi/analog.h, 499
- include/rtdm/uapi/can.h, 464
- include/rtdm/uapi/ipc.h, 473
- include/rtdm/uapi/rtdm.h, 477
- include/rtdm/uapi/serial.h, 480
- include/rtdm/uapi/testing.h, 490
- inesting
 - xnsched, 422
- Interrupt management, 104
 - xnintr_affinity, 104
 - xnintr_attach, 104
 - xnintr_destroy, 105
 - xnintr_detach, 105
 - xnintr_disable, 106
 - xnintr_enable, 106
 - xnintr_init, 106
- Interrupt Management Services, 166
 - rtdm_irq_disable, 167
 - rtdm_irq_enable, 168
 - rtdm_irq_free, 168
 - rtdm_irq_get_arg, 167
 - rtdm_irq_handler_t, 167
 - rtdm_irq_request, 170
- Interrupt management services, 243
 - a4l_free_irq, 243
 - a4l_get_irq, 243
 - a4l_request_irq, 243
- kernel/cobalt/apc.c, 505
- kernel/cobalt/arith.c, 506
- kernel/cobalt/bufd.c, 507
- kernel/cobalt/map.c, 508
- kernel/cobalt/posix/select.c, 510
- kernel/cobalt/posix/syscall.c, 505
- kernel/cobalt/registry.c, 512
- kernel/cobalt/rtdm/core.c, 513
- kernel/cobalt/rtdm/device.c, 516
- kernel/cobalt/rtdm/drvlib.c, 517
- kernel/cobalt/rtdm/init.c, 508
- kernel/cobalt/timer.c, 511
- kernel/cobalt/vfile.c, 521
- kernel/drivers/analog/buffer.c, 522
- kernel/drivers/analog/command.c, 524
- kernel/drivers/analog/device.c, 517
- kernel/drivers/analog/driver.c, 524
- kernel/drivers/analog/driver_facilities.c, 525
- kernel/drivers/analog/instruction.c, 527
- kernel/drivers/analog/intel/8255.c, 528
- kernel/drivers/analog/intel/8255.h, 528
- kernel/drivers/analog/intel/parport.c, 529
- kernel/drivers/analog/national_instruments/mio_common.c, 530
- kernel/drivers/analog/national_instruments/mite.c, 531
- kernel/drivers/analog/national_instruments/mite.h, 532
- kernel/drivers/analog/national_instruments/nio_mio.h, 533
- kernel/drivers/analog/national_instruments/nio_stc.h, 534
- kernel/drivers/analog/national_instruments/nio_tio.h, 535
- kernel/drivers/analog/national_instruments/pcmio.c, 536
- kernel/drivers/analog/national_instruments/tio_common.c, 537
- kernel/drivers/analog/proc.h, 538
- kernel/drivers/analog/rtdm_helpers.c, 539
- kernel/drivers/analog/rtdm_interface.c, 540
- kernel/drivers/analog/sensoray/s526.c, 540
- kernel/drivers/analog/subdevice.c, 541
- kernel/drivers/analog/transfer.c, 542
- label
 - rtipc_port_label, 419
- length
 - a4l_channels_desc, 398
- Level 0 API, 302
- Level 1 API, 292
- Level 2 API, 297
- lflags
 - xnsched, 423
- lib/analog/async.c, 543
- lib/analog/descriptor.c, 544
- lib/analog/info.c, 545
- lib/analog/internal.h, 509
- lib/analog/range.c, 546
- lib/analog/root_leaf.h, 547
- lib/analog/sync.c, 548
- lib/analog/sys.c, 549
- Lightweight key-to-object mapping service, 109
 - xnmap_create, 109
 - xnmap_delete, 110
 - xnmap_enter, 110
 - xnmap_fetch, 111
 - xnmap_fetch_nocheck, 111
 - xnmap_remove, 111
- magic
 - a4l_descriptor, 401
- max
 - a4l_range, 404
- Message pipe services, 347
 - P_MINOR_AUTO, 348
 - P_URGENT, 348
 - rt_pipe_bind, 348
 - rt_pipe_create, 349

- rt_pipe_delete, 350
- rt_pipe_read, 350
- rt_pipe_read_timed, 350
- rt_pipe_read_until, 351
- rt_pipe_unbind, 352
- rt_pipe_write, 352
- Message queue services, 354
 - Q_PRIO, 355
 - Q_SHARED, 355
 - rt_queue_alloc, 355
 - rt_queue_bind, 355
 - rt_queue_create, 356
 - rt_queue_delete, 357
 - rt_queue_flush, 358
 - rt_queue_free, 358
 - rt_queue_inquire, 358
 - rt_queue_read, 360
 - rt_queue_read_timed, 360
 - rt_queue_read_until, 361
 - rt_queue_receive, 361
 - rt_queue_receive_timed, 362
 - rt_queue_receive_until, 362
 - rt_queue_send, 363
 - rt_queue_unbind, 363
- Message queues, 254
 - mq_close, 254
 - mq_getattr, 255
 - mq_notify, 255
 - mq_open, 256
 - mq_setattr, 257
 - mq_unlink, 258
- min
 - a4l_range, 404
- Misc services, 246
 - a4l_get_time, 246
- mode
 - a4l_channels_desc, 398
- mq_close
 - Message queues, 254
- mq_getattr
 - Message queues, 255
- mq_notify
 - Message queues, 255
- mq_open
 - Message queues, 256
- mq_setattr
 - Message queues, 257
- mq_unlink
 - Message queues, 258
- Mutex services, 163, 341
 - rt_mutex_acquire, 342
 - rt_mutex_acquire_timed, 342
 - rt_mutex_acquire_until, 342
 - rt_mutex_bind, 343
 - rt_mutex_create, 343
 - rt_mutex_delete, 344
 - rt_mutex_inquire, 344
 - rt_mutex_release, 345
 - rt_mutex_unbind, 345
 - rtdm_mutex_destroy, 163
 - rtdm_mutex_init, 163
 - rtdm_mutex_lock, 164
 - rtdm_mutex_timedlock, 164
 - rtdm_mutex_unlock, 165
- Mutual exclusion, 259
- nanosecs_abs_t
 - Real-Time Driver Model, 51
- nanosecs_rel_t
 - Real-Time Driver Model, 51
- nb_bits
 - a4l_channel, 397
- nb_subd
 - a4l_descriptor, 401
- next
 - xnvfile_regular_ops, 426
 - xnvfile_snapshot_ops, 432
- nkvfroot
 - Virtual file services, 226
- Non-Real-Time Signalling Services, 171
 - rtdm_nrtsig_destroy, 172
 - rtdm_nrtsig_handler_t, 171
 - rtdm_nrtsig_init, 172
 - rtdm_nrtsig_pend, 172
- nrddata
 - xnvfile_snapshot_iterator, 430
- open
 - rtdm_device, 418
- owner
 - RT_MUTEX_INFO, 413
- P_MINOR_AUTO
 - Message pipe services, 348
- P_URGENT
 - Message pipe services, 348
- pos
 - xnvfile_regular_iterator, 424
- private
 - xnvfile_regular_iterator, 424
 - xnvfile_snapshot_iterator, 431
- program_htick_shot
 - Timer services, 215
- pthread_cond_destroy
 - Condition variables, 251
- pthread_cond_init
 - Condition variables, 252
- pthread_create
 - Thread management, 266
- pthread_getschedparam_ex
 - Thread management, 267
- pthread_make_periodic_np
 - Thread management, 268
- pthread_set_mode_np
 - Thread management, 268
- pthread_setschedparam_ex
 - Thread management, 269

- put
 - xnvfile_lock_ops, [424](#)
- Q_PRIO
 - Message queue services, [355](#)
- Q_SHARED
 - Message queue services, [355](#)
- RTDM_SELECTTYPE_EXCEPT
 - Synchronisation Services, [151](#)
- RTDM_SELECTTYPE_READ
 - Synchronisation Services, [151](#)
- RTDM_SELECTTYPE_WRITE
 - Synchronisation Services, [151](#)
- RTDM_TIMERMODE_ABSOLUTE
 - Timer Services, [147](#)
- RTDM_TIMERMODE_REALTIME
 - Timer Services, [147](#)
- RTDM_TIMERMODE_RELATIVE
 - Timer Services, [147](#)
- RT_ALARM_INFO, [410](#)
- RT_BUFFER_INFO, [410](#)
- RT_COND_INFO, [411](#)
- RT_EVENT_INFO, [411](#)
- RT_HEAP_INFO, [411](#)
 - usablemem, [412](#)
- RT_MUTEX_INFO, [412](#)
 - owner, [413](#)
- RT_QUEUE_INFO, [413](#)
- RT_SEM_INFO, [413](#)
- RT_TASK_INFO, [414](#)
- RT_TIMER_INFO
 - Timer management services, [394](#)
- RTDM User API, [52](#)
 - rt_dev_accept, [53](#)
 - rt_dev_bind, [53](#)
 - rt_dev_close, [54](#)
 - rt_dev_connect, [54](#)
 - rt_dev_getpeername, [55](#)
 - rt_dev_getsockname, [55](#)
 - rt_dev_getsockopt, [56](#)
 - rt_dev_ioctl, [56](#)
 - rt_dev_listen, [56](#)
 - rt_dev_open, [58](#)
 - rt_dev_read, [58](#)
 - rt_dev_recv, [59](#)
 - rt_dev_recvfrom, [59](#)
 - rt_dev_recvmsg, [60](#)
 - rt_dev_send, [60](#)
 - rt_dev_sendmsg, [61](#)
 - rt_dev_sendto, [61](#)
 - rt_dev_setsockopt, [62](#)
 - rt_dev_shutdown, [62](#)
 - rt_dev_socket, [63](#)
 - rt_dev_write, [63](#)
- RTDM_EXCLUSIVE
 - Device Registration Services, [135](#)
- RTDM_NAMED_DEVICE
 - Device Registration Services, [135](#)
- RTDM_TIMEOUT_NONE
 - Real-Time Driver Model, [51](#)
- RTIOC_DEVICE_INFO
 - Device Profiles, [186](#)
- RTIOC_PURGE
 - Device Profiles, [186](#)
- Range / conversion API, [287](#)
 - a4l_dtoraw, [287](#)
 - a4l_find_range, [288](#)
 - a4l_ftoraw, [288](#)
 - a4l_rawtod, [289](#)
 - a4l_rawtof, [289](#)
 - a4l_rawtoul, [289](#)
 - a4l_sizeof_chan, [290](#)
 - a4l_sizeof_subd, [290](#)
 - a4l_ultoraw, [291](#)
- Real-Time Driver Model, [50](#)
 - nanosecs_abs_t, [51](#)
 - nanosecs_rel_t, [51](#)
 - RTDM_TIMEOUT_NONE, [51](#)
- Real-time IPC
 - IPCPROTO_BUFP, [76](#)
 - IPCPROTO_IDDP, [76](#)
 - IPCPROTO_IPC, [76](#)
 - IPCPROTO_XDDP, [76](#)
- Real-time IPC, [67](#)
 - BUFP_BUFSZ, [69](#)
 - BUFP_LABEL, [70](#)
 - bind__AF_RTIPC, [76](#)
 - close__AF_RTIPC, [77](#)
 - connect__AF_RTIPC, [78](#)
 - getpeername__AF_RTIPC, [78](#)
 - getsockname__AF_RTIPC, [79](#)
 - getsockopt__AF_RTIPC, [79](#)
 - IDDP_LABEL, [70](#)
 - IDDP_POOLSZ, [71](#)
 - recvmsg__AF_RTIPC, [79](#)
 - SO_RCVTIMEO, [72](#)
 - SO_SNDTIMEO, [72](#)
 - sendmsg__AF_RTIPC, [80](#)
 - setsockopt__AF_RTIPC, [81](#)
 - socket__AF_RTIPC, [81](#)
 - XDDP_BUFSZ, [72](#)
 - XDDP_EVTDOWN, [73](#)
 - XDDP_EVTIN, [73](#)
 - XDDP_EVTNOBUF, [73](#)
 - XDDP_EVTOUT, [73](#)
 - XDDP_LABEL, [73](#)
 - XDDP_MONITOR, [74](#)
 - XDDP_POOLSZ, [75](#)
- Real-time shadow services, [193](#)
 - xnshadow_get_context, [194](#)
 - xnshadow_harden, [194](#)
 - xnshadow_map_kernel, [194](#)
 - xnshadow_map_user, [195](#)
 - xnshadow_pop_personality, [195](#)
 - xnshadow_push_personality, [196](#)
 - xnshadow_register_personality, [196](#)

- xnshadow_relax, [196](#)
- recvmsg__AF_RTIPC
 - Real-time IPC, [79](#)
- Registry services, [114](#)
 - xnregistry_bind, [114](#)
 - xnregistry_enter, [115](#)
 - xnregistry_lookup, [116](#)
 - xnregistry_remove, [116](#)
 - xnregistry_unlink, [117](#)
- rev
 - xnvfile_rev_tag, [428](#)
- rewind
 - xnvfile_regular_ops, [426](#)
 - xnvfile_snapshot_ops, [433](#)
- rrbtimer
 - xnsched, [423](#)
- rt
 - xnsched, [423](#)
- rt_alarm_create
 - Alarm services, [307](#)
- rt_alarm_delete
 - Alarm services, [308](#)
- rt_alarm_inquire
 - Alarm services, [308](#)
- rt_alarm_start
 - Alarm services, [309](#)
- rt_alarm_stop
 - Alarm services, [309](#)
- rt_buffer_bind
 - Buffer services, [312](#)
- rt_buffer_clear
 - Buffer services, [313](#)
- rt_buffer_create
 - Buffer services, [313](#)
- rt_buffer_delete
 - Buffer services, [314](#)
- rt_buffer_inquire
 - Buffer services, [314](#)
- rt_buffer_read
 - Buffer services, [315](#)
- rt_buffer_read_timed
 - Buffer services, [315](#)
- rt_buffer_read_until
 - Buffer services, [316](#)
- rt_buffer_unbind
 - Buffer services, [317](#)
- rt_buffer_write
 - Buffer services, [317](#)
- rt_buffer_write_timed
 - Buffer services, [317](#)
- rt_buffer_write_until
 - Buffer services, [318](#)
- rt_cond_bind
 - Condition variable services, [320](#)
- rt_cond_broadcast
 - Condition variable services, [320](#)
- rt_cond_create
 - Condition variable services, [321](#)
- rt_cond_delete
 - Condition variable services, [321](#)
- rt_cond_inquire
 - Condition variable services, [322](#)
- rt_cond_signal
 - Condition variable services, [322](#)
- rt_cond_unbind
 - Condition variable services, [323](#)
- rt_cond_wait
 - Condition variable services, [323](#)
- rt_cond_wait_timed
 - Condition variable services, [323](#)
- rt_cond_wait_until
 - Condition variable services, [324](#)
- rt_dev_accept
 - RTDM User API, [53](#)
- rt_dev_bind
 - RTDM User API, [53](#)
- rt_dev_close
 - RTDM User API, [54](#)
- rt_dev_connect
 - RTDM User API, [54](#)
- rt_dev_getpeername
 - RTDM User API, [55](#)
- rt_dev_getsockname
 - RTDM User API, [55](#)
- rt_dev_getsockopt
 - RTDM User API, [56](#)
- rt_dev_ioctl
 - RTDM User API, [56](#)
- rt_dev_listen
 - RTDM User API, [56](#)
- rt_dev_open
 - RTDM User API, [58](#)
- rt_dev_read
 - RTDM User API, [58](#)
- rt_dev_recv
 - RTDM User API, [59](#)
- rt_dev_recvfrom
 - RTDM User API, [59](#)
- rt_dev_recvmsg
 - RTDM User API, [60](#)
- rt_dev_send
 - RTDM User API, [60](#)
- rt_dev_sendmsg
 - RTDM User API, [61](#)
- rt_dev_sendto
 - RTDM User API, [61](#)
- rt_dev_setsockopt
 - RTDM User API, [62](#)
- rt_dev_shutdown
 - RTDM User API, [62](#)
- rt_dev_socket
 - RTDM User API, [63](#)
- rt_dev_write
 - RTDM User API, [63](#)
- rt_event_bind
 - Event flag group services, [326](#)

- rt_event_clear
 - Event flag group services, [327](#)
- rt_event_create
 - Event flag group services, [327](#)
- rt_event_delete
 - Event flag group services, [328](#)
- rt_event_inquire
 - Event flag group services, [328](#)
- rt_event_signal
 - Event flag group services, [329](#)
- rt_event_unbind
 - Event flag group services, [329](#)
- rt_event_wait
 - Event flag group services, [329](#)
- rt_event_wait_timed
 - Event flag group services, [330](#)
- rt_event_wait_until
 - Event flag group services, [331](#)
- rt_heap_alloc
 - Heap management services, [333](#)
- rt_heap_alloc_timed
 - Heap management services, [333](#)
- rt_heap_alloc_until
 - Heap management services, [334](#)
- rt_heap_bind
 - Heap management services, [334](#)
- rt_heap_create
 - Heap management services, [335](#)
- rt_heap_delete
 - Heap management services, [336](#)
- rt_heap_free
 - Heap management services, [336](#)
- rt_heap_inquire
 - Heap management services, [338](#)
- rt_heap_unbind
 - Heap management services, [338](#)
- rt_mutex_acquire
 - Mutex services, [342](#)
- rt_mutex_acquire_timed
 - Mutex services, [342](#)
- rt_mutex_acquire_until
 - Mutex services, [342](#)
- rt_mutex_bind
 - Mutex services, [343](#)
- rt_mutex_create
 - Mutex services, [343](#)
- rt_mutex_delete
 - Mutex services, [344](#)
- rt_mutex_inquire
 - Mutex services, [344](#)
- rt_mutex_release
 - Mutex services, [345](#)
- rt_mutex_unbind
 - Mutex services, [345](#)
- rt_pipe_bind
 - Message pipe services, [348](#)
- rt_pipe_create
 - Message pipe services, [349](#)
- rt_pipe_delete
 - Message pipe services, [350](#)
- rt_pipe_read
 - Message pipe services, [350](#)
- rt_pipe_read_timed
 - Message pipe services, [350](#)
- rt_pipe_read_until
 - Message pipe services, [351](#)
- rt_pipe_unbind
 - Message pipe services, [352](#)
- rt_pipe_write
 - Message pipe services, [352](#)
- rt_queue_alloc
 - Message queue services, [355](#)
- rt_queue_bind
 - Message queue services, [355](#)
- rt_queue_create
 - Message queue services, [356](#)
- rt_queue_delete
 - Message queue services, [357](#)
- rt_queue_flush
 - Message queue services, [358](#)
- rt_queue_free
 - Message queue services, [358](#)
- rt_queue_inquire
 - Message queue services, [358](#)
- rt_queue_read
 - Message queue services, [360](#)
- rt_queue_read_timed
 - Message queue services, [360](#)
- rt_queue_read_until
 - Message queue services, [361](#)
- rt_queue_receive
 - Message queue services, [361](#)
- rt_queue_receive_timed
 - Message queue services, [362](#)
- rt_queue_receive_until
 - Message queue services, [362](#)
- rt_queue_send
 - Message queue services, [363](#)
- rt_queue_unbind
 - Message queue services, [363](#)
- rt_sem_bind
 - Semaphore services, [367](#)
- rt_sem_broadcast
 - Semaphore services, [368](#)
- rt_sem_create
 - Semaphore services, [368](#)
- rt_sem_delete
 - Semaphore services, [369](#)
- rt_sem_inquire
 - Semaphore services, [369](#)
- rt_sem_p
 - Semaphore services, [369](#)
- rt_sem_p_timed
 - Semaphore services, [370](#)
- rt_sem_p_until
 - Semaphore services, [370](#)

- rt_sem_unbind
 - Semaphore services, [371](#)
- rt_sem_v
 - Semaphore services, [371](#)
- rt_task_bind
 - Task management services, [374](#)
- rt_task_create
 - Task management services, [375](#)
- rt_task_delete
 - Task management services, [376](#)
- rt_task_inquire
 - Task management services, [376](#)
- rt_task_join
 - Task management services, [377](#)
- rt_task_receive
 - Task management services, [377](#)
- rt_task_receive_timed
 - Task management services, [378](#)
- rt_task_receive_until
 - Task management services, [379](#)
- rt_task_reply
 - Task management services, [379](#)
- rt_task_resume
 - Task management services, [380](#)
- rt_task_same
 - Task management services, [380](#)
- rt_task_self
 - Task management services, [381](#)
- rt_task_send
 - Task management services, [381](#)
- rt_task_send_timed
 - Task management services, [381](#)
- rt_task_send_until
 - Task management services, [383](#)
- rt_task_set_affinity
 - Task management services, [383](#)
- rt_task_set_mode
 - Task management services, [383](#)
- rt_task_set_periodic
 - Task management services, [384](#)
- rt_task_set_priority
 - Task management services, [385](#)
- rt_task_shadow
 - Task management services, [386](#)
- rt_task_sleep
 - Task management services, [387](#)
- rt_task_sleep_until
 - Task management services, [387](#)
- rt_task_slice
 - Task management services, [389](#)
- rt_task_spawn
 - Task management services, [389](#)
- rt_task_start
 - Task management services, [390](#)
- rt_task_suspend
 - Task management services, [391](#)
- rt_task_unbind
 - Task management services, [391](#)
- rt_task_unblock
 - Task management services, [392](#)
- rt_task_wait_period
 - Task management services, [392](#)
- rt_task_yield
 - Task management services, [393](#)
- rt_timer_info, [414](#)
 - date, [415](#)
 - tsc, [415](#)
- rt_timer_inquire
 - Timer management services, [395](#)
- rt_timer_ns2ticks
 - Timer management services, [395](#)
- rt_timer_read
 - Timer management services, [395](#)
- rt_timer_spin
 - Timer management services, [396](#)
- rt_timer_ticks2ns
 - Timer management services, [396](#)
- rtdm_accept
 - Driver to driver services, [129](#)
- rtdm_bind
 - Driver to driver services, [129](#)
- rtdm_clock_read
 - Clock Services, [137](#)
- rtdm_clock_read_monotonic
 - Clock Services, [137](#)
- rtdm_close
 - Driver to driver services, [129](#)
- rtdm_connect
 - Driver to driver services, [129](#)
- rtdm_context_lock
 - Synchronisation Services, [151](#)
- rtdm_context_put
 - Synchronisation Services, [151](#)
- rtdm_context_unlock
 - Synchronisation Services, [152](#)
- rtdm_copy_from_user
 - Utility Services, [175](#)
- rtdm_copy_to_user
 - Utility Services, [175](#)
- rtdm_dev_context, [415](#)
 - device, [416](#)
- rtdm_dev_register
 - Device Registration Services, [135](#)
- rtdm_dev_unregister
 - Device Registration Services, [136](#)
- rtdm_device, [416](#)
 - open, [418](#)
 - socket, [418](#)
- rtdm_device_info, [418](#)
- rtdm_event_clear
 - Event Services, [154](#)
- rtdm_event_destroy
 - Event Services, [155](#)
- rtdm_event_init
 - Event Services, [155](#)
- rtdm_event_pulse

- Event Services, [155](#)
- `rtdm_event_select_bind`
 - Event Services, [156](#)
- `rtdm_event_signal`
 - Event Services, [156](#)
- `rtdm_event_timedwait`
 - Event Services, [157](#)
- `rtdm_event_wait`
 - Event Services, [157](#)
- `rtdm_fd_device`
 - `driver.h`, [458](#)
- `rtdm_fd_is_user`
 - `driver.h`, [458](#)
- `rtdm_fd_to_private`
 - `driver.h`, [459](#)
- `rtdm_for_each_waiter`
 - Wait queue, [120](#)
- `rtdm_for_each_waiter_safe`
 - Wait queue, [120](#)
- `rtdm_free`
 - Utility Services, [176](#)
- `rtdm_getpeername`
 - Driver to driver services, [130](#)
- `rtdm_getsockname`
 - Driver to driver services, [130](#)
- `rtdm_getsockopt`
 - Driver to driver services, [130](#)
- `rtdm_in_rt_context`
 - Utility Services, [176](#)
- `rtdm_ioctl`
 - Driver to driver services, [130](#)
- `rtdm_iomap_to_user`
 - Utility Services, [176](#)
- `rtdm_irq_disable`
 - Interrupt Management Services, [167](#)
- `rtdm_irq_enable`
 - Interrupt Management Services, [168](#)
- `rtdm_irq_free`
 - Interrupt Management Services, [168](#)
- `rtdm_irq_get_arg`
 - Interrupt Management Services, [167](#)
- `rtdm_irq_handler_t`
 - Interrupt Management Services, [167](#)
- `rtdm_irq_request`
 - Interrupt Management Services, [170](#)
- `rtdm_listen`
 - Driver to driver services, [130](#)
- `rtdm_lock_get`
 - Spinlock with preemption deactivation, [22](#)
- `rtdm_lock_init`
 - Spinlock with preemption deactivation, [22](#)
- `rtdm_lock_irqrestore`
 - Spinlock with preemption deactivation, [21](#)
- `rtdm_lock_irqsave`
 - Spinlock with preemption deactivation, [22](#)
- `rtdm_lock_put`
 - Spinlock with preemption deactivation, [24](#)
- `rtdm_lock_put_irqrestore`
 - Spinlock with preemption deactivation, [24](#)
- `rtdm_malloc`
 - Utility Services, [177](#)
- `rtdm_mmap_to_user`
 - Utility Services, [178](#)
- `rtdm_munmap`
 - Utility Services, [179](#)
- `rtdm_mutex_destroy`
 - Mutex services, [163](#)
- `rtdm_mutex_init`
 - Mutex services, [163](#)
- `rtdm_mutex_lock`
 - Mutex services, [164](#)
- `rtdm_mutex_timedlock`
 - Mutex services, [164](#)
- `rtdm_mutex_unlock`
 - Mutex services, [165](#)
- `rtdm_nrtsig_destroy`
 - Non-Real-Time Signalling Services, [172](#)
- `rtdm_nrtsig_handler_t`
 - Non-Real-Time Signalling Services, [171](#)
- `rtdm_nrtsig_init`
 - Non-Real-Time Signalling Services, [172](#)
- `rtdm_nrtsig_pend`
 - Non-Real-Time Signalling Services, [172](#)
- `rtdm_open`
 - Driver to driver services, [131](#)
- `rtdm_open_handler_t`
 - `driver.h`, [457](#)
- `rtdm_printk`
 - Utility Services, [179](#)
- `rtdm_printk_ratelimited`
 - Utility Services, [180](#)
- `rtdm_private_to_fd`
 - `driver.h`, [459](#)
- `rtdm_ratelimit`
 - Utility Services, [180](#)
- `rtdm_read`
 - Driver to driver services, [131](#)
- `rtdm_read_user_ok`
 - Utility Services, [181](#)
- `rtdm_recv`
 - Driver to driver services, [131](#)
- `rtdm_recvfrom`
 - Driver to driver services, [131](#)
- `rtdm_recvmsg`
 - Driver to driver services, [131](#)
- `rtdm_rt_capable`
 - Utility Services, [181](#)
- `rtdm_rw_user_ok`
 - Utility Services, [182](#)
- `rtdm_safe_copy_from_user`
 - Utility Services, [182](#)
- `rtdm_safe_copy_to_user`
 - Utility Services, [183](#)
- `rtdm_selecttype`
 - Synchronisation Services, [151](#)
- `rtdm_sem_destroy`

- Semaphore Services, [159](#)
- rtm_sem_down
 - Semaphore Services, [159](#)
- rtm_sem_init
 - Semaphore Services, [160](#)
- rtm_sem_select_bind
 - Semaphore Services, [160](#)
- rtm_sem_timeddown
 - Semaphore Services, [161](#)
- rtm_sem_up
 - Semaphore Services, [162](#)
- rtm_send
 - Driver to driver services, [131](#)
- rtm_sendmsg
 - Driver to driver services, [132](#)
- rtm_sendto
 - Driver to driver services, [132](#)
- rtm_setsockopt
 - Driver to driver services, [132](#)
- rtm_shutdown
 - Driver to driver services, [132](#)
- rtm_socket
 - Driver to driver services, [132](#)
- rtm_socket_handler_t
 - driver.h, [458](#)
- rtm_strncpy_from_user
 - Utility Services, [183](#)
- rtm_task_busy_sleep
 - Task Services, [140](#)
- rtm_task_current
 - Task Services, [140](#)
- rtm_task_destroy
 - Task Services, [141](#)
- rtm_task_init
 - Task Services, [141](#)
- rtm_task_join_nrt
 - Task Services, [142](#)
- rtm_task_proc_t
 - Task Services, [140](#)
- rtm_task_set_period
 - Task Services, [142](#)
- rtm_task_set_priority
 - Task Services, [143](#)
- rtm_task_should_stop
 - Task Services, [143](#)
- rtm_task_sleep
 - Task Services, [143](#)
- rtm_task_sleep_abs
 - Task Services, [144](#)
- rtm_task_sleep_until
 - Task Services, [144](#)
- rtm_task_unblock
 - Task Services, [145](#)
- rtm_task_wait_period
 - Task Services, [145](#)
- rtm_timedwait
 - Wait queue, [120](#)
- rtm_timedwait_condition
 - Wait queue, [122](#)
- rtm_timedwait_condition_locked
 - Wait queue, [122](#)
- rtm_timedwait_locked
 - Wait queue, [123](#)
- rtm_timer_destroy
 - Timer Services, [147](#)
- rtm_timer_handler_t
 - Timer Services, [146](#)
- rtm_timer_init
 - Timer Services, [147](#)
- rtm_timer_mode
 - Timer Services, [147](#)
- rtm_timer_start
 - Timer Services, [148](#)
- rtm_timer_start_in_handler
 - Timer Services, [148](#)
- rtm_timer_stop
 - Timer Services, [149](#)
- rtm_timer_stop_in_handler
 - Timer Services, [149](#)
- rtm_toseq_init
 - Synchronisation Services, [152](#)
- rtm_wait
 - Wait queue, [124](#)
- rtm_wait_condition
 - Wait queue, [124](#)
- rtm_wait_condition_locked
 - Wait queue, [124](#)
- rtm_wait_locked
 - Wait queue, [125](#)
- rtm_waitqueue_broadcast
 - Wait queue, [125](#)
- rtm_waitqueue_destroy
 - Wait queue, [126](#)
- rtm_waitqueue_flush
 - Wait queue, [126](#)
- rtm_waitqueue_init
 - Wait queue, [126](#)
- rtm_waitqueue_lock
 - Wait queue, [126](#)
- rtm_waitqueue_signal
 - Wait queue, [126](#)
- rtm_waitqueue_unlock
 - Wait queue, [127](#)
- rtm_waitqueue_wakeup
 - Wait queue, [127](#)
- rtm_write
 - Driver to driver services, [133](#)
- rtipc_port_label, [418](#)
 - label, [419](#)
- rtser_config, [419](#)
- rtser_event, [420](#)
- rtser_status, [420](#)
- S_PRIO
 - Semaphore services, [367](#)
- SCHED_QUOTA scheduling policy, [188](#)
- SIOCGCANBAUDRATE

- CAN Devices, [42](#)
- SIOCGCANCTRLMODE
 - CAN Devices, [42](#)
- SIOCGCANCUSTOMBITTIME
 - CAN Devices, [43](#)
- SIOCGCANSTATE
 - CAN Devices, [43](#)
- SIOCGIFINDEX
 - CAN Devices, [44](#)
- SIOCSCANBAUDRATE
 - CAN Devices, [44](#)
- SIOCSCANCTRLMODE
 - CAN Devices, [45](#)
- SIOCSCANCUSTOMBITTIME
 - CAN Devices, [46](#)
- SIOCSCANMODE
 - CAN Devices, [46](#)
- SMP locking services, [108](#)
- SO_RCVTIMEO
 - Real-time IPC, [72](#)
- SO_SNDTIMEO
 - Real-time IPC, [72](#)
- SOL_CAN_RAW
 - CAN Devices, [47](#)
- sbddata
 - a4l_descriptor, [401](#)
- sbsize
 - a4l_descriptor, [401](#)
- sched_getconfig_np
 - Scheduling management, [271](#)
- sched_setconfig_np
 - Scheduling management, [272](#)
- Scheduling management, [271](#)
 - sched_getconfig_np, [271](#)
 - sched_setconfig_np, [272](#)
- sem_close
 - Semaphores, [260](#)
- sem_destroy
 - Semaphores, [261](#)
- sem_post
 - Semaphores, [261](#)
- sem_timedwait
 - Semaphores, [263](#)
- sem_trywait
 - Semaphores, [264](#)
- sem_unlink
 - Semaphores, [264](#)
- sem_wait
 - Semaphores, [265](#)
- Semaphore Services, [159](#)
 - rtdm_sem_destroy, [159](#)
 - rtdm_sem_down, [159](#)
 - rtdm_sem_init, [160](#)
 - rtdm_sem_select_bind, [160](#)
 - rtdm_sem_timeddown, [161](#)
 - rtdm_sem_up, [162](#)
- Semaphore services, [366](#)
 - rt_sem_bind, [367](#)
 - rt_sem_broadcast, [368](#)
 - rt_sem_create, [368](#)
 - rt_sem_delete, [369](#)
 - rt_sem_inquire, [369](#)
 - rt_sem_p, [369](#)
 - rt_sem_p_timed, [370](#)
 - rt_sem_p_until, [370](#)
 - rt_sem_unbind, [371](#)
 - rt_sem_v, [371](#)
 - S_PRIO, [367](#)
- Semaphores, [260](#)
 - sem_close, [260](#)
 - sem_destroy, [261](#)
 - sem_post, [261](#)
 - sem_timedwait, [263](#)
 - sem_trywait, [264](#)
 - sem_unlink, [264](#)
 - sem_wait, [265](#)
- sendmsg_AF_RTIPC
 - Real-time IPC, [80](#)
- seq
 - xnvfile_regular_iterator, [425](#)
 - xnvfile_snapshot_iterator, [431](#)
- Serial Devices, [65](#)
- setsockopt_AF_RTIPC
 - Real-time IPC, [81](#)
- show
 - xnvfile_regular_ops, [427](#)
 - xnvfile_snapshot_ops, [433](#)
- sipc_port
 - sockaddr_ipc, [422](#)
- sockaddr_can, [421](#)
 - can_ifindex, [421](#)
- sockaddr_ipc, [421](#)
 - sipc_port, [422](#)
- socket
 - rtdm_device, [418](#)
- socket_AF_RTIPC
 - Real-time IPC, [81](#)
- Spinlock with preemption deactivation, [21](#)
 - rtdm_lock_get, [22](#)
 - rtdm_lock_init, [22](#)
 - rtdm_lock_irqrestore, [21](#)
 - rtdm_lock_irqsave, [22](#)
 - rtdm_lock_put, [24](#)
 - rtdm_lock_put_irqrestore, [24](#)
- status
 - xnsched, [423](#)
- store
 - xnvfile_regular_ops, [427](#)
 - xnvfile_snapshot_ops, [434](#)
- Subdevice management services, [231](#)
 - a4l_add_subd, [233](#)
 - a4l_alloc_subd, [233](#)
 - a4l_get_subd, [233](#)
- switch_htick_mode
 - Timer services, [215](#)
- Synchronisation Services, [150](#)

- RTDM_SELECTTYPE_EXCEPT, 151
- RTDM_SELECTTYPE_READ, 151
- RTDM_SELECTTYPE_WRITE, 151
- rtdm_context_lock, 151
- rtdm_context_put, 151
- rtdm_context_unlock, 152
- rtdm_selecttype, 151
- rtdm_toseq_init, 152
- Synchronous acquisition API, 293, 298
 - a4l_config_subd, 298
 - a4l_snd_insn, 295
 - a4l_snd_insnlist, 295
 - a4l_sync_dio, 299
 - a4l_sync_read, 299
 - a4l_sync_write, 300
- T_FPU
 - Task management services, 374
- T_LOCK
 - Task management services, 374
- T_LOPRIO
 - Task management services, 374
- T_WARNSW
 - Task management services, 374
- Task management services, 372
 - rt_task_bind, 374
 - rt_task_create, 375
 - rt_task_delete, 376
 - rt_task_inquire, 376
 - rt_task_join, 377
 - rt_task_receive, 377
 - rt_task_receive_timed, 378
 - rt_task_receive_until, 379
 - rt_task_reply, 379
 - rt_task_resume, 380
 - rt_task_same, 380
 - rt_task_self, 381
 - rt_task_send, 381
 - rt_task_send_timed, 381
 - rt_task_send_until, 383
 - rt_task_set_affinity, 383
 - rt_task_set_mode, 383
 - rt_task_set_periodic, 384
 - rt_task_set_priority, 385
 - rt_task_shadow, 386
 - rt_task_sleep, 387
 - rt_task_sleep_until, 387
 - rt_task_slice, 389
 - rt_task_spawn, 389
 - rt_task_start, 390
 - rt_task_suspend, 391
 - rt_task_unbind, 391
 - rt_task_unblock, 392
 - rt_task_wait_period, 392
 - rt_task_yield, 393
 - T_FPU, 374
 - T_LOCK, 374
 - T_LOPRIO, 374
 - T_WARNSW, 374
- Task Services, 139
 - rtdm_task_busy_sleep, 140
 - rtdm_task_current, 140
 - rtdm_task_destroy, 141
 - rtdm_task_init, 141
 - rtdm_task_join_nrt, 142
 - rtdm_task_proc_t, 140
 - rtdm_task_set_period, 142
 - rtdm_task_set_priority, 143
 - rtdm_task_should_stop, 143
 - rtdm_task_sleep, 143
 - rtdm_task_sleep_abs, 144
 - rtdm_task_sleep_until, 144
 - rtdm_task_unblock, 145
 - rtdm_task_wait_period, 145
- Testing Devices, 66
- Thread information flags, 28
- Thread management, 266
 - pthread_create, 266
 - pthread_getschedparam_ex, 267
 - pthread_make_periodic_np, 268
 - pthread_set_mode_np, 268
 - pthread_setschedparam_ex, 269
- Thread services, 205
 - xnthread_cancel, 205
 - xnthread_init, 206
 - xnthread_migrate, 207
 - xnthread_resume, 207
 - xnthread_set_mode, 208
 - xnthread_set_periodic, 209
 - xnthread_set_schedparam, 209
 - xnthread_set_slice, 210
 - xnthread_start, 211
 - xnthread_suspend, 211
 - xnthread_unblock, 212
 - xnthread_wait_period, 213
- Thread state flags, 25
 - XNHELD, 26
 - XNLOCK, 26
 - XNMIGRATE, 26
 - XNPEND, 26
 - XNREADY, 26
 - XNSUSP, 26
 - XNTRAPLB, 26
- Thread synchronization services., 198
 - clear_boost, 199
 - xnsynch_acquire, 199
 - xnsynch_flush, 199
 - xnsynch_forget_sleeper, 200
 - xnsynch_init, 200
 - xnsynch_peek_pendq, 201
 - xnsynch_release, 201
 - xnsynch_release_all_ownerships, 202
 - xnsynch_requeue_sleeper, 202
 - xnsynch_sleep_on, 203
 - xnsynch_wakeup_one_sleeper, 203
 - xnsynch_wakeup_this_sleeper, 204
- Timer management services, 394

- RT_TIMER_INFO, 394
- rt_timer_inquire, 395
- rt_timer_ns2ticks, 395
- rt_timer_read, 395
- rt_timer_spin, 396
- rt_timer_ticks2ns, 396
- Timer Services, 146
 - RTDM_TIMERMODE_ABSOLUTE, 147
 - RTDM_TIMERMODE_REALTIME, 147
 - RTDM_TIMERMODE_RELATIVE, 147
 - rtdm_timer_destroy, 147
 - rtdm_timer_handler_t, 146
 - rtdm_timer_init, 147
 - rtdm_timer_mode, 147
 - rtdm_timer_start, 148
 - rtdm_timer_start_in_handler, 148
 - rtdm_timer_stop, 149
 - rtdm_timer_stop_in_handler, 149
- Timer services, 214
 - program_htick_shot, 215
 - switch_htick_mode, 215
 - xntimer_destroy, 216
 - xntimer_get_date, 216
 - xntimer_get_overruns, 216
 - xntimer_get_timeout, 217
 - xntimer_grab_hardware, 217
 - xntimer_init, 217
 - xntimer_interval, 218
 - xntimer_release_hardware, 218
 - xntimer_start, 218
 - xntimer_stop, 219
- timer_create
 - Clocks and timers, 247
- timer_gettime
 - Clocks and timers, 248
- timer_settime
 - Clocks and timers, 249
- tsc
 - rt_timer_info, 415
- uapi/analog.h
 - A4L_RNG_FACTOR, 504
- usablemem
 - RT_HEAP_INFO, 412
- Utility Services, 174
 - rtdm_copy_from_user, 175
 - rtdm_copy_to_user, 175
 - rtdm_free, 176
 - rtdm_in_rt_context, 176
 - rtdm_iomap_to_user, 176
 - rtdm_malloc, 177
 - rtdm_mmap_to_user, 178
 - rtdm_munmap, 179
 - rtdm_printk, 179
 - rtdm_printk_ratelimited, 180
 - rtdm_ratelimit, 180
 - rtdm_read_user_ok, 181
 - rtdm_rt_capable, 181
 - rtdm_rw_user_ok, 182
 - rtdm_safe_copy_from_user, 182
 - rtdm_safe_copy_to_user, 183
 - rtdm_strncpy_from_user, 183
- vfile
 - xnvfile_regular_iterator, 425
 - xnvfile_snapshot_iterator, 431
- Virtual file services, 220
 - nkvfroot, 226
 - xnvfile_destroy, 221
 - xnvfile_get_blob, 222
 - xnvfile_get_integer, 222
 - xnvfile_get_string, 223
 - xnvfile_init_dir, 223
 - xnvfile_init_link, 223
 - xnvfile_init_regular, 225
 - xnvfile_init_snapshot, 225
- Wait queue, 119
 - rtdm_for_each_waiter, 120
 - rtdm_for_each_waiter_safe, 120
 - rtdm_timedwait, 120
 - rtdm_timedwait_condition, 122
 - rtdm_timedwait_condition_locked, 122
 - rtdm_timedwait_locked, 123
 - rtdm_wait, 124
 - rtdm_wait_condition, 124
 - rtdm_wait_condition_locked, 124
 - rtdm_wait_locked, 125
 - rtdm_waitqueue_broadcast, 125
 - rtdm_waitqueue_destroy, 126
 - rtdm_waitqueue_flush, 126
 - rtdm_waitqueue_init, 126
 - rtdm_waitqueue_lock, 126
 - rtdm_waitqueue_signal, 126
 - rtdm_waitqueue_unlock, 127
 - rtdm_waitqueue_wakeup, 127
- XDDP_BUFSZ
 - Real-time IPC, 72
- XDDP_EVTDOWN
 - Real-time IPC, 73
- XDDP_EVTIN
 - Real-time IPC, 73
- XDDP_EVTNOBUF
 - Real-time IPC, 73
- XDDP_EVTOUT
 - Real-time IPC, 73
- XDDP_LABEL
 - Real-time IPC, 73
- XDDP_MONITOR
 - Real-time IPC, 74
- XDDP_POOLSZ
 - Real-time IPC, 75
- XNHeld
 - Thread state flags, 26
- XNLOCK
 - Thread state flags, 26
- XNMIGRATE

- Thread state flags, [26](#)
- XNPEND
 - Thread state flags, [26](#)
- XNREADY
 - Thread state flags, [26](#)
- XNSUSP
 - Thread state flags, [26](#)
- XNTRAPLB
 - Thread state flags, [26](#)
- Xenomai nucleus, [102](#)
- xnipc_alloc
 - Asynchronous Procedure Calls, [82](#)
- xnipc_free
 - Asynchronous Procedure Calls, [83](#)
- xnipc_schedule
 - Asynchronous Procedure Calls, [83](#)
- xnbufd_copy_from_kmem
 - Buffer descriptor, [87](#)
- xnbufd_copy_to_kmem
 - Buffer descriptor, [88](#)
- xnbufd_invalidate
 - Buffer descriptor, [88](#)
- xnbufd_map_kread
 - Buffer descriptor, [89](#)
- xnbufd_map_kwrite
 - Buffer descriptor, [89](#)
- xnbufd_map_uread
 - Buffer descriptor, [89](#)
- xnbufd_map_uwrite
 - Buffer descriptor, [90](#)
- xnbufd_reset
 - Buffer descriptor, [90](#)
- xnbufd_unmap_kread
 - Buffer descriptor, [90](#)
- xnbufd_unmap_kwrite
 - Buffer descriptor, [91](#)
- xnbufd_unmap_uread
 - Buffer descriptor, [91](#)
- xnbufd_unmap_uwrite
 - Buffer descriptor, [91](#)
- xnclock_adjust
 - Clock services., [93](#)
- xnclock_deregister
 - Clock services., [94](#)
- xnclock_register
 - Clock services., [94](#)
- xnclock_tick
 - Clock services., [94](#)
- xnheap_alloc
 - Dynamic memory allocation services., [98](#)
- xnheap_destroy
 - Dynamic memory allocation services., [98](#)
- xnheap_extend
 - Dynamic memory allocation services., [99](#)
- xnheap_free
 - Dynamic memory allocation services., [99](#)
- xnheap_init
 - Dynamic memory allocation services., [99](#)
- xnheap_set_label
 - Dynamic memory allocation services., [100](#)
- xnheap_test_and_free
 - Dynamic memory allocation services., [100](#)
- xnintr_affinity
 - Interrupt management, [104](#)
- xnintr_attach
 - Interrupt management, [104](#)
- xnintr_destroy
 - Interrupt management, [105](#)
- xnintr_detach
 - Interrupt management, [105](#)
- xnintr_disable
 - Interrupt management, [106](#)
- xnintr_enable
 - Interrupt management, [106](#)
- xnintr_init
 - Interrupt management, [106](#)
- xnmap_create
 - Lightweight key-to-object mapping service, [109](#)
- xnmap_delete
 - Lightweight key-to-object mapping service, [110](#)
- xnmap_enter
 - Lightweight key-to-object mapping service, [110](#)
- xnmap_fetch
 - Lightweight key-to-object mapping service, [111](#)
- xnmap_fetch_nocheck
 - Lightweight key-to-object mapping service, [111](#)
- xnmap_remove
 - Lightweight key-to-object mapping service, [111](#)
- xnregistry_bind
 - Registry services, [114](#)
- xnregistry_enter
 - Registry services, [115](#)
- xnregistry_lookup
 - Registry services, [116](#)
- xnregistry_remove
 - Registry services, [116](#)
- xnregistry_unlink
 - Registry services, [117](#)
- xnsched, [422](#)
 - curr, [422](#)
 - htimer, [422](#)
 - inesting, [422](#)
 - lflags, [423](#)
 - rrbtimer, [423](#)
 - rt, [423](#)
 - status, [423](#)
- xnsched_rotate
 - include/cobalt/kernel/sched.h, [438](#)
- xnsched_run
 - include/cobalt/kernel/sched.h, [439](#)

- xselect
 - File descriptors events multiplexing services, [190](#)
- xselect_bind
 - File descriptors events multiplexing services, [190](#)
- xselect_destroy
 - File descriptors events multiplexing services, [191](#)
- xselect_init
 - File descriptors events multiplexing services, [191](#)
- xselect_signal
 - File descriptors events multiplexing services, [192](#)
- xselector_destroy
 - File descriptors events multiplexing services, [192](#)
- xselector_init
 - File descriptors events multiplexing services, [192](#)
- xnshadow_get_context
 - Real-time shadow services, [194](#)
- xnshadow_harden
 - Real-time shadow services, [194](#)
- xnshadow_map_kernel
 - Real-time shadow services, [194](#)
- xnshadow_map_user
 - Real-time shadow services, [195](#)
- xnshadow_pop_personality
 - Real-time shadow services, [195](#)
- xnshadow_push_personality
 - Real-time shadow services, [196](#)
- xnshadow_register_personality
 - Real-time shadow services, [196](#)
- xnshadow_relax
 - Real-time shadow services, [196](#)
- xsynch_acquire
 - Thread synchronization services., [199](#)
- xsynch_flush
 - Thread synchronization services., [199](#)
- xsynch_forget_sleeper
 - Thread synchronization services., [200](#)
- xsynch_init
 - Thread synchronization services., [200](#)
- xsynch_peek_pendq
 - Thread synchronization services., [201](#)
- xsynch_release
 - Thread synchronization services., [201](#)
- xsynch_release_all_ownerships
 - Thread synchronization services., [202](#)
- xsynch_requeue_sleeper
 - Thread synchronization services., [202](#)
- xsynch_sleep_on
 - Thread synchronization services., [203](#)
- xsynch_wakeup_one_sleeper
 - Thread synchronization services., [203](#)
- xsynch_wakeup_this_sleeper
 - Thread synchronization services., [204](#)
- xnthread_cancel
 - Thread services, [205](#)
- xnthread_init
 - Thread services, [206](#)
- xnthread_migrate
 - Thread services, [207](#)
- xnthread_resume
 - Thread services, [207](#)
- xnthread_set_mode
 - Thread services, [208](#)
- xnthread_set_periodic
 - Thread services, [209](#)
- xnthread_set_schedparam
 - Thread services, [209](#)
- xnthread_set_slice
 - Thread services, [210](#)
- xnthread_start
 - Thread services, [211](#)
- xnthread_suspend
 - Thread services, [211](#)
- xnthread_unblock
 - Thread services, [212](#)
- xnthread_wait_period
 - Thread services, [213](#)
- xntimer_destroy
 - Timer services, [216](#)
- xntimer_get_date
 - Timer services, [216](#)
- xntimer_get_overruns
 - Timer services, [216](#)
- xntimer_get_timeout
 - Timer services, [217](#)
- xntimer_grab_hardware
 - Timer services, [217](#)
- xntimer_init
 - Timer services, [217](#)
- xntimer_interval
 - Timer services, [218](#)
- xntimer_release_hardware
 - Timer services, [218](#)
- xntimer_start
 - Timer services, [218](#)
- xntimer_stop
 - Timer services, [219](#)
- xnfile_destroy
 - Virtual file services, [221](#)
- xnfile_get_blob
 - Virtual file services, [222](#)
- xnfile_get_integer
 - Virtual file services, [222](#)
- xnfile_get_string
 - Virtual file services, [223](#)
- xnfile_init_dir
 - Virtual file services, [223](#)
- xnfile_init_link
 - Virtual file services, [223](#)
- xnfile_init_regular

- Virtual file services, [225](#)
- xnvfile_init_snapshot
 - Virtual file services, [225](#)
- xnvfile_lock_ops, [423](#)
 - get, [423](#)
 - put, [424](#)
- xnvfile_regular_iterator, [424](#)
 - pos, [424](#)
 - private, [424](#)
 - seq, [425](#)
 - vfile, [425](#)
- xnvfile_regular_ops, [425](#)
 - begin, [425](#)
 - end, [426](#)
 - next, [426](#)
 - rewind, [426](#)
 - show, [427](#)
 - store, [427](#)
- xnvfile_rev_tag, [428](#)
 - rev, [428](#)
- xnvfile_snapshot, [428](#)
- xnvfile_snapshot_iterator, [429](#)
 - databuf, [430](#)
 - endfn, [430](#)
 - nrdata, [430](#)
 - private, [431](#)
 - seq, [431](#)
 - vfile, [431](#)
- xnvfile_snapshot_ops, [431](#)
 - begin, [431](#)
 - end, [432](#)
 - next, [432](#)
 - rewind, [433](#)
 - show, [433](#)
 - store, [434](#)